# GNU Compiler Collection Internals

For GCC version 10.3.1

(GNU Tools for STM32 10.3-2021.10.20211105-1100)

Richard M. Stallman and the GCC Developer Community

Copyright © 1988-2020 Free Software Foundation, Inc.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with the Invariant Sections being "Funding Free Software", the Front-Cover Texts being (a) (see below), and with the Back-Cover Texts being (b) (see below). A copy of the license is included in the section entitled "GNU Free Documentation License".

(a) The FSF's Front-Cover Text is:

A GNU Manual

(b) The FSF's Back-Cover Text is:

You have freedom to copy and modify this GNU Manual, like GNU software. Copies published by the Free Software Foundation raise funds for GNU development.

## **Short Contents**

Intr	oduction	. 1
1	Contributing to GCC Development	3
2	GCC and Portability	. 5
3	Interfacing to GCC Output	. 7
4	The GCC low-level runtime library	9
5	Language Front Ends in GCC	59
6	Source Tree Structure and Build System	. 61
7	Testsuites	79
8	Option specification files	119
9	Passes and Files of the Compiler	127
10	Sizes and offsets as runtime invariants	147
11	GENERIC	161
12	GIMPLE	209
13	Analysis and Optimization of GIMPLE tuples	247
14	RTL Representation	259
15	Control Flow Graph	317
16	Analysis and Representation of Loops	327
17	Machine Descriptions	337
18	Target Description Macros and Functions	479
19	Host Configuration	
20	Makefile Fragments	667
21	collect2	671
22	Standard Header File Directories	673
23	Memory Management and Type Information	675
24	Plugins	685
25	Link Time Optimization	693
26	Match and Simplify	701
27	Static Analyzer	707
28	User Experience Guidelines	715
Fun	ding Free Software	723
The	GNU Project and GNU/Linux	725
GNI	U General Public License	727
GNI	U Free Documentation License	739
Con	tributors to GCC	747
Opt	ion Index	765

ii	GNU Compiler Collection	(GCC) Internals
Concept Index	70	67

## Table of Contents

Ir	$\operatorname{ntroduction} \dots \dots \dots \dots \dots \dots \dots \dots$	1
1	Contributing to GCC Development	3
2	GCC and Portability	5
3	Interfacing to GCC Output	7
4		
	4.1 Routines for integer arithmetic	9
	4.1.1 Arithmetic functions	
	4.1.2 Comparison functions	10
	4.1.3 Trapping arithmetic functions	
	4.1.4 Bit operations	
	4.2 Routines for floating point emulation	
	4.2.1 Arithmetic functions	
	4.2.2 Conversion functions	
	4.2.3 Comparison functions	
	4.2.4 Other floating-point functions	
	4.3 Routines for decimal floating point emulation	
	4.3.1 Arithmetic functions	
	4.3.2 Conversion functions	
	4.3.3 Comparison functions	
	4.4 Routines for fixed-point fractional emulation	
	4.4.1 Arithmetic functions	
	4.4.2 Comparison functions	
	4.5 Language-independent routines for exception handling	
	4.6 Miscellaneous runtime library routines	
	4.6.1 Cache control functions	
	4.6.2 Split stack functions and variables	
5	Language Front Ends in GCC	59
6	Source Tree Structure and Build System	61
-	6.1 Configure Terms and History	
	6.2 Top Level Source Directory	
	6.3 The 'gcc' Subdirectory	
	6.3.1 Subdirectories of 'gcc'	
	6.3.2 Configuration in the 'gcc' Directory	
	6.3.2.1 Scripts Used by 'configure'	

		The 'config.build'; 'config.host'; and 'config	
	File	es	
	6.3.2.3	Files Created by configure	
		ld System in the 'gcc' Directory	
		xefile Targets	
		rary Source Files and Headers under the 'gcc' Direc	·
		ders Installed by GCC	
		Iding Documentation	
	6.3.7.1	Texinfo Manuals	
	6.3.7.2	Man Page Generation	
	6.3.7.3	Miscellaneous Documentation	
		atomy of a Language Front End	
		The Front End 'language' Directory	
	6.3.8.2	The Front End 'config-lang.in' File	
	6.3.8.3	The Front End 'Make-lang.in' File	
	6.3.9 Ana	atomy of a Target Back End	75
_	TD 4 *4		=0
7	Testsuit	$\operatorname{es}$	. 79
	7.1 Idioms U	Jsed in Testsuite Code	79
		es used within DejaGnu tests	
	7.2.1 Syn	tax and Descriptions of test directives	80
	7.2.1.1	Specify how to build the test	
	7.2.1.2	Specify additional compiler options	
	7.2.1.3	Modify the test timeout value	
	7.2.1.4	Skip a test for some targets	
	7.2.1.5	Expect a test to fail for some targets	
	7.2.1.6	Expect the test executable to fail	
	7.2.1.7	Verify compiler messages	
	7.2.1.8	Verify output of the test executable	
	7.2.1.9	Specify environment variables for a test	
	7.2.1.10	1 0	
	7.2.1.11		
		cting targets to which a test applies	
	v	words describing target attributes	
	7.2.3.1	Endianness	
	7.2.3.2	Data type sizes	
	7.2.3.3	Fortran-specific attributes	
	7.2.3.4 $7.2.3.5$	Vector-specific attributes	
		Thread Local Storage attributes	
	7.2.3.6 $7.2.3.7$	Decimal floating point attributes	
	7.2.3.7 $7.2.3.8$	ARM-specific attributes	
	7.2.3.8 $7.2.3.9$	AArch64-specific attributes	
	7.2.3.9 $7.2.3.10$	MIPS-specific attributes	
	7.2.3.10 $7.2.3.11$	1	
	7.2.3.11 $7.2.3.12$		
	7.2.3.12 $7.2.3.13$		
		3.70013/1 (MODELLIA III) 33	

	7.2.3.14 Local to tests in gcc.target/i386	. 103
	7.2.3.15 Local to tests in gcc.test-framework	. 104
	7.2.4 Features for dg-add-options	. 104
	7.2.5 Variants of dg-require-support	. 106
	7.2.6 Commands for use in dg-final	107
	7.2.6.1 Scan a particular file	107
	7.2.6.2 Scan the assembly output	107
	7.2.6.3 Scan optimization dump files	. 109
	7.2.6.4 Check for output files	109
	7.2.6.5 Checks for gcov tests	
	7.2.6.6 Clean up generated test files	
	7.3 Ada Language Testsuites	. 111
	7.4 C Language Testsuites	
	7.5 Support for testing link-time optimizations	
	7.6 Support for testing gcov	
	7.7 Support for testing profile-directed optimizations	. 114
	7.8 Support for testing binary compatibility	
	7.9 Support for torture testing using multiple options	116
	7.10 Support for testing GIMPLE passes	. 117
	7.11 Support for testing RTL passes	118
8	Option specification files	119
	8.1 Option file format	. 119
	8.2 Option properties	. 121
9	Passes and Files of the Compiler	127
	9.1 Parsing pass	
	9.2 Gimplification pass	
	9.3 Pass manager	
	9.4 Inter-procedural optimization passes	
	9.4.1 Small IPA passes	
	9.4.2 Regular IPA passes	
	9.4.3 Late IPA passes	
	9.5 Tree SSA passes	
	1	- 130
	9.6 RTL passes	
	9.6 RTL passes	143
	9.6 RTL passes	. 143 143
	9.6 RTL passes	<ul><li>143</li><li>143</li><li>143</li></ul>
	9.6 RTL passes 9.7 Optimization info 9.7.1 Dump setup 9.7.2 Optimization groups 9.7.3 Dump files and streams	143 143 143 143
	9.6 RTL passes 9.7 Optimization info 9.7.1 Dump setup 9.7.2 Optimization groups 9.7.3 Dump files and streams 9.7.4 Dump output verbosity	. 143 143 143 143 144
	9.6 RTL passes 9.7 Optimization info 9.7.1 Dump setup 9.7.2 Optimization groups 9.7.3 Dump files and streams	<ul><li>143</li><li>143</li><li>143</li><li>144</li><li>144</li></ul>

10	Sizes	${f s}$ and offsets as runtime invariants $\dots$	. 147
10.	1 Ove	rview of poly_int	147
10.5	2 Con	sequences of using poly_int	148
10.3	3 Con	nparisons involving poly_int	
	10.3.1	Comparison functions for poly_int	
	10.3.2	Properties of the poly_int comparisons	
	10.3.3	Comparing potentially-unordered poly_ints	
	10.3.4	Comparing ordered poly_ints	
	10.3.5	Checking for a poly_int marker value	
	10.3.6	Range checks on poly_ints	
	10.3.7	Sorting poly_ints	
	$\frac{4}{10.4.1}$	hmetic on poly_ints	
	10.4.1 $10.4.2$	Using poly_int with C++ arithmetic operators	
	10.4.2 $10.4.3$	wi arithmetic on poly_ints  Division of poly_ints	
	10.4.3 $10.4.4$	Other poly_int arithmetic	
10.		mment of poly_ints	
10.0	_	puting bounds on poly_ints	
10.		verting poly_ints	
10.		cellaneous poly_int routines	
10.9		delines for using poly_int	
101	o Gan	demies for dams polyline from the first terms and the first terms are the first terms and the first terms are the first terms and the first terms are the first terms	100
11	GEN	VERIC	161
11.		ciencies	
11.		rview	
	11.2.1	Trees	
	11.2.1 $11.2.2$	Identifiers	
	11.2.2	Containers	
11.3		es	
11.4		larations	
	11.4.1	Working with declarations	
	11.4.2	Internal structure	
	11.4		
	11.4	.2.2 Adding new DECL node types	
11.		ributes in trees	
11.	6 Exp	ressions	173
	11.6.1	Constant expressions	173
	11.6.2	References to storage	175
	11.6.3	Unary and Binary Expressions	177
	11.6.4	Vectors	184
11.	7 Stat	ements	186
	11.7.1	Basic Statements	186
	11.7.2	Blocks	
	11.7.3	Statement Sequences	
	11.7.4	Empty Statements	
	11.7.5	Jumps	
	11.7.6	Cleanups	
	11.7.7	OpenMP	190

11.7.8	OpenACC	192
11.8 Fund	ctions	193
11.8.1	Function Basics	193
11.8.2	Function Properties	194
11.9 Lang	guage-dependent trees	195
11.10 C a	and C++ Trees	196
11.10.1	Types for C++	196
11.10.2	Namespaces	198
11.10.3	Classes	199
11.10.4	Functions for C++	201
11.10.5	Statements for C++	203
11.1	0.5.1 Statements	204
11.10.6	C++ Expressions	206
12 GIM	PLE	209
12.1 Tup	le representation	210
12.1.1	gimple (gsbase)	
12.1.2	gimple_statement_with_ops	211
12.1.3	gimple_statement_with_memory_ops	
12.2 Clas	s hierarchy of GIMPLE statements	
	IPLE instruction set	
12.4 Exce	eption Handling	215
	poraries	
	rands	
12.6.1	Compound Expressions	
12.6.2	Compound Lvalues	
12.6.3	Conditional Expressions	
12.6.4	Logical Operators	
12.6.5	Manipulating operands	
12.6.6	Operand vector allocation	
12.6.7	Operand validation	
12.6.8	Statement validation	
12.7 Man	ipulating GIMPLE statements	
12.7.1	Common accessors	
12.8 Tup	le specific accessors	223
12.8.1	GIMPLE_ASM	223
12.8.2	GIMPLE_ASSIGN	224
12.8.3	GIMPLE_BIND	225
12.8.4	GIMPLE_CALL	226
12.8.5	GIMPLE_CATCH	227
12.8.6	GIMPLE_COND	228
12.8.7	GIMPLE_DEBUG	229
12.8.8	GIMPLE_EH_FILTER	230
12.8.9	GIMPLE_LABEL	
12.8.10	GIMPLE_GOTO	231
12.8.11	GIMPLE_NOP	
12.8.12	GIMPLE_OMP_ATOMIC_LOAD	
	GIMPLE_OMP_ATOMIC_STORE	

12.8.14	GIMPLE_OMP_CONTINUE	$\dots 232$
12.8.15	GIMPLE_OMP_CRITICAL	233
12.8.16	GIMPLE_OMP_FOR	233
12.8.17	GIMPLE_OMP_MASTER	234
12.8.18	GIMPLE_OMP_ORDERED	235
12.8.19	GIMPLE_OMP_PARALLEL	235
12.8.20	GIMPLE_OMP_RETURN	236
12.8.21	GIMPLE_OMP_SECTION	236
12.8.22	GIMPLE_OMP_SECTIONS	236
12.8.23	GIMPLE_OMP_SINGLE	$\dots 237$
12.8.24	GIMPLE_PHI	
12.8.25	GIMPLE_RESX	
12.8.26	GIMPLE_RETURN	
12.8.27	GIMPLE_SWITCH	
12.8.28	GIMPLE_TRY	
12.8.29	GIMPLE_WITH_CLEANUP_EXPR	
	PLE sequences	
	uence iterators	
	ding a new GIMPLE statement code	
12.12 Stat	tement and operand traversals	$\dots 245$
		_
13 Analy	ysis and Optimization of GIMPLI	E tuples
		247
13.1 Anno	otations	247
	Operands	
	Operand Iterators And Access Routines	
	Immediate Uses	
	c Single Assignment	
	Preserving the SSA form	
	Examining SSA_NAME nodes	
	Walking the dominator tree	
	s analysis	
	ory model	
14 RTL	Representation	259
	Object Types	
	Object Types	
	Classes and Formats	260
14.4 Acce	Classes and Formatsss to Operands	$\begin{array}{ccc} \dots & 260 \\ \dots & 262 \end{array}$
	Classes and Formats	260 262 263
14.5 Flags	Classes and Formatsss to Operandsss to Special Operandsss in an RTL Expressionss	260 262 263 266
14.5 Flags 14.6 Mack	Classes and Formatsss to Operandsss to Special Operandsss in an RTL Expressionsnine Modes	
<ul><li>14.5 Flags</li><li>14.6 Mach</li><li>14.7 Cons</li></ul>	Classes and Formats ss to Operands ss to Special Operands s in an RTL Expression nine Modes stant Expression Types	
<ul><li>14.5 Flags</li><li>14.6 Mach</li><li>14.7 Cons</li><li>14.8 Regis</li></ul>	Classes and Formats ss to Operands ss to Special Operands s in an RTL Expression nine Modes stant Expression Types sters and Memory	
<ul><li>14.5 Flags</li><li>14.6 Mach</li><li>14.7 Cons</li><li>14.8 Regis</li><li>14.9 RTL</li></ul>	Classes and Formats ss to Operands ss to Special Operands s in an RTL Expression nine Modes stant Expression Types sters and Memory Expressions for Arithmetic	
14.5 Flags 14.6 Mach 14.7 Cons 14.8 Regis 14.9 RTL 14.10 Con	Classes and Formats ss to Operands ss to Special Operands s in an RTL Expression nine Modes stant Expression Types sters and Memory Expressions for Arithmetic mparison Operations	
14.5 Flags 14.6 Mach 14.7 Cons 14.8 Regis 14.9 RTL 14.10 Con 14.11 Bit-	Classes and Formats ss to Operands ss to Special Operands s in an RTL Expression nine Modes stant Expression Types sters and Memory Expressions for Arithmetic	

	14.13	Conversions	. 295
	14.14	Declarations	. 297
	14.15	Side Effect Expressions	. 297
	14.16	Embedded Side-Effects on Addresses	. 302
	14.17	Assembler Instructions as Expressions	. 303
	14.18	Variable Location Debug Information in RTL	
	14.19	Insns	
	14.20	RTL Representation of Function-Call Insns	
	14.21	Structure Sharing Assumptions	
	14.22	Reading RTL	
15	6 C	Control Flow Graph	317
	15.1	Basic Blocks	. 317
	15.2	Edges	
	15.3	Profile information	
	15.4	Maintaining the CFG	
	15.5	Liveness information	
16	6 A	analysis and Representation of Loops	327
	16.1	Loop representation	
	16.2	Loop querying	
	16.3	Loop manipulation	
	16.4	Loop-closed SSA form	
	16.5	Scalar evolutions	
	16.6	IV analysis on RTL	
	16.7	Number of iterations analysis	
	16.8	Data Dependency Analysis	
17	7 N.	Iachine Descriptions	227
1 (		_	
	17.1	Overview of How the Machine Description is Used	
	17.2	Everything about Instruction Patterns	
	17.3	Example of define_insn	
	17.4	RTL Template	
	17.5	Output Templates and Operand Substitution	
	17.6	C Statements for Assembler Output	
	17.7	Predicates	
		.7.1 Machine-Independent Predicates	
	17	.7.2 Defining Machine-Specific Predicates	
	17.8	Operand Constraints	. 350
		.8.1 Simple Constraints	
		.8.2 Multiple Alternative Constraints	
		.8.3 Register Class Preferences	
		.8.4 Constraint Modifier Characters	
		.8.5 Constraints for Particular Machines	
		.8.6 Disable insn alternatives using the ${\tt enabled}$ attribute	
		.8.7 Defining Machine-Specific Constraints	
	17	.8.8 Testing constraints from C	201

17.9 Standard Pattern Names For Generation	392
17.10 When the Order of Patterns Matters	432
17.11 Interdependence of Patterns	. 433
17.12 Defining Jump Instruction Patterns	433
17.13 Defining Looping Instruction Patterns	434
17.14 Canonicalization of Instructions	436
17.15 Defining RTL Sequences for Code Generation	438
17.16 Defining How to Split Instructions	440
17.17 Including Patterns in Machine Descriptions	446
17.17.1 RTL Generation Tool Options for Directory Search	446
17.18 Machine-Specific Peephole Optimizers	446
17.18.1 RTL to Text Peephole Optimizers	447
17.18.2 RTL to RTL Peephole Optimizers	449
17.19 Instruction Attributes	. 450
17.19.1 Defining Attributes and their Values	451
17.19.2 Attribute Expressions	452
17.19.3 Assigning Attribute Values to Insns	455
17.19.4 Example of Attribute Specifications	456
17.19.5 Computing the Length of an Insn	457
17.19.6 Constant Attributes	. 458
17.19.7 Mnemonic Attribute	459
17.19.8 Delay Slot Scheduling	459
17.19.9 Specifying processor pipeline description	. 460
17.20 Conditional Execution	466
17.21 RTL Templates Transformations	467
17.21.1 define_subst Example	468
17.21.2 Pattern Matching in define_subst	. 469
17.21.3 Generation of output template in define_subst	469
17.22 Constant Definitions	470
17.23 Iterators	472
17.23.1 Mode Iterators	472
17.23.1.1 Defining Mode Iterators	472
17.23.1.2 Substitution in Mode Iterators	473
17.23.1.3 Mode Iterator Examples	. 473
17.23.2 Code Iterators	474
17.23.3 Int Iterators	475
17.23.4 Subst Iterators	476
17.23.5 Parameterized Names	477
18 Target Description Macros and Functions	
	479
18.1 The Global targetm Variable	
18.3 Run-time Target Specification	
18.4 Defining data structures for per-function information	
O V	
18.6 Layout of Source Language Data Types	
10.1 Register Usage	505

Basic Characteristics of Registers	. 505
Order of Allocation of Registers	. 507
How Values Fit in Registers	. 508
9	
9	
9	
· ·	
9	
9 9	
9	
9	
· · · · · · · · · · · · · · · · · · ·	
9	
~ ·	
· ·	
9	
Representation of condition codes using registers	. 574
justing the Instruction Scheduler	. 583
viding the Output into Sections (Texts, Data,)	. 590
sition Independent Code	. 595
fining the Output Assembler Language	. 596
The Overall Framework of an Assembler File	. 596
Output of Data	. 600
Output of Uninitialized Variables	. 602
How Initialization Functions Are Handled	. 612
Macros Controlling Initialization Routines	. 614
_	
<del>-</del>	
8	
	Basic Characteristics of Registers Order of Allocation of Registers How Values Fit in Registers Handling Leaf Functions Registers That Form a Stack ister Classes k Layout and Calling Conventions Basic Stack Layout Exception Handling Support Specifying How Stack Checking is Done Registers That Address the Stack Frame Eliminating Frame Pointer and Arg Pointer Passing Function Arguments on the Stack Passing Arguments in Registers How Scalar Function Values Are Returned How Large Values Are Returned Caller-Saves Register Allocation Function Entry and Exit Generating Code for Profiling Permitting tail calls Shrink-wrapping separate components Stack smashing protection Miscellaneous register hooks plementing the Varargs Macros oport for Nested Functions plicit Calls to Library Routines dressing Modes chored Addresses Indition Code Status Representation of condition codes using registers scribing Relative Costs of Operations justing the Instruction Scheduler iriding the Output into Sections (Texts, Data, ) sition Independent Code fining the Output Assembler Language The Overall Framework of an Assembler File Output of Data Output of Uninitialized Variables Output and Generation of Labels How Initialization Functions Are Handled Macros Controlling Initialization Routines Output of Assembler Instructions Output of Dispatch Tables Assembler Commands for Exception Regions O Assembler Commands for Exception Regions

	18.2	21.2 Specific Options for DBX Output	626
	18.2	21.3 Open-Ended Hooks for DBX Format	628
	18.2	21.4 File Names in DBX Format	628
	18.2	21.5 Macros for DWARF Output	629
	18.2	21.6 Macros for VMS Debug Format	631
18.	22	Cross Compilation and Floating Point	631
18.	23	Mode Switching Instructions	
18.	24	Defining target-specific uses ofattribute	
18.		Emulating TLS	
18.	26	Defining coprocessor specifics for MIPS targets	
18.		Parameters for Precompiled Header Validity Checking	
18.		C++ ABI parameters	
18.		D ABI parameters	
18.		Adding support for named address spaces	
18.	31	Miscellaneous Parameters	642
19	H	ost Configuration	663
19.	1 :	Host Common	663
19.	2	Host Filesystem	664
19.	3	Host Misc	. 665
20	$\mathbf{M}$	[akefile Fragments	667
20.	1 '	Target Makefile Fragments	667
20.	2	Host Makefile Fragments	670
21	СО	ollect2	671
22	St	andard Header File Directories	673
		difficulties of the Difference of the second	010
23	ъл	Iomory Management and Type Informati	ion
<b>4</b> 3	TVI	lemory Management and Type Informati	
	• •		
23.	1 '	The Inside of a GTY(())	. 676
23.		Support for inheritance	
23.	3	Support for user-provided GC marking routines	
	23.3	1 01	
23.4		Marking Roots for the Garbage Collector	
23.		Source Files Containing Type Information	
23.		How to invoke the garbage collector	
23.	7 '	Troubleshooting the garbage collector	684

24	$\mathbf{P}$	Plugins68	<b>85</b>
2	24.1	Loading Plugins	385
	24.2	Plugin API 6	
	24	4.2.1 Plugin license check	
	24	4.2.2 Plugin initialization	
	24	4.2.3 Plugin callbacks 6	687
2	24.3	Interacting with the pass manager	688
2	24.4	Interacting with the GCC Garbage Collector	689
2	24.5	Giving information about a plugin	689
2	24.6	Registering custom attributes or pragmas	
2	24.7	Recording information about pass execution	390
	24.8	Controlling which passes are being run	
2	24.9	Keeping track of available passes	
2	24.10	Building GCC plugins	591
<b>25</b>	$\mathbf{L}$	Link Time Optimization 69	93
2	5.1	Design Overview	393
	25	5.1.1 LTO modes of operation	394
2	25.2	LTO file sections	694
2	25.3	Using summary information in IPA passes	696
	25	5.3.1 Virtual clones	597
	25	5.3.2 IPA references	398
	25	5.3.3 Jump functions	598
2	25.4	Whole program assumptions, linker plugin and symbol visibilities	ies 598
2	25.5	Internal flags controlling 1to1	
<b>26</b>		Match and Simplify	JI
_	26.1		701
2	26.2	The Language	702
27	$\mathbf{S}$	Static Analyzer	<b>07</b>
2	27.1	Analyzer Internals	707
	27	·	707
	27		708
	27	•	708
		9	710
		9	711
		•	712
2	27.2		712
		98 9 1	712
			713

28 User Experience Guideline	es
28.1 Guidelines for Diagnostics	715
28.1.1 Talk in terms of the user's code	·
28.1.2 Diagnostics are actionable	715
28.1.3 The user's attention is important	nt 715
28.1.4 Precision of Wording	715
28.1.5 Try the diagnostic on real-world	
28.1.6 Make mismatches clear	
28.1.7 Location Information	
28.1.8 Coding Conventions	
28.1.9 Group logically-related diagnost	
28.1.10 Quoting	
28.1.11 Spelling and Terminology	
28.1.12 Fix-it hints	$ \begin{array}{cccccccccccccccccccccccccccccccccccc$
	of deletion, not replacement
20.1.12.2 Express defending terms	·
28.1.12.3 Multiple suggestions	
28.2 Guidelines for Options	
r	
Funding Free Software	723
The GNU Project and GNU/L	inuv 795
The Give Project and Give/L	mux
GNU General Public License	727
GNU Free Documentation Lice	ense
ADDENDUM: How to use this License for	your documents 746
TID D DI TION TO THE THEORY TO	jour decuments
Contributors to GCC	747
Ontion Indon	705
$ Option \ Index \dots \dots $	705
Concept Index	

Introduction 1

#### Introduction

This manual documents the internals of the GNU compilers, including how to port them to new targets and some information about how to write front ends for new languages. It corresponds to the compilers (GNU Tools for STM32 10.3-2021.10.20211105-1100) version 10.3.1. The use of the GNU compilers is documented in a separate manual. See Section "Introduction" in *Using the GNU Compiler Collection (GCC)*.

This manual is mainly a reference manual rather than a tutorial. It discusses how to contribute to GCC (see Chapter 1 [Contributing], page 3), the characteristics of the machines supported by GCC as hosts and targets (see Chapter 2 [Portability], page 5), how GCC relates to the ABIs on such systems (see Chapter 3 [Interface], page 7), and the characteristics of the languages for which GCC front ends are written (see Chapter 5 [Languages], page 59). It then describes the GCC source tree structure and build system, some of the interfaces to GCC front ends, and how support for a target system is implemented in GCC.

Additional tutorial information is linked to from http://gcc.gnu.org/readings.html.

## 1 Contributing to GCC Development

If you would like to help pretest GCC releases to assure they work well, current development sources are available via Git (see http://gcc.gnu.org/git.html). Source and binary snapshots are also available for FTP; see http://gcc.gnu.org/snapshots.html.

If you would like to work on improvements to GCC, please read the advice at these URLs:

```
http://gcc.gnu.org/contribute.html
http://gcc.gnu.org/contributewhy.html
```

for information on how to make useful contributions and avoid duplication of effort. Suggested projects are listed at http://gcc.gnu.org/projects/.

### 2 GCC and Portability

GCC itself aims to be portable to any machine where int is at least a 32-bit type. It aims to target machines with a flat (non-segmented) byte addressed data address space (the code address space can be separate). Target ABIs may have 8, 16, 32 or 64-bit int type. char can be wider than 8 bits.

GCC gets most of the information about the target machine from a machine description which gives an algebraic formula for each of the machine's instructions. This is a very clean way to describe the target. But when the compiler needs information that is difficult to express in this fashion, ad-hoc parameters have been defined for machine descriptions. The purpose of portability is to reduce the total work needed on the compiler; it was not of interest for its own sake.

GCC does not contain machine dependent code, but it does contain code that depends on machine parameters such as endianness (whether the most significant byte has the highest or lowest address of the bytes in a word) and the availability of autoincrement addressing. In the RTL-generation pass, it is often necessary to have multiple strategies for generating code for a particular kind of syntax tree, strategies that are usable for different combinations of parameters. Often, not all possible cases have been addressed, but only the common ones or only the ones that have been encountered. As a result, a new target may require additional strategies. You will know if this happens because the compiler will call abort. Fortunately, the new strategies can be added in a machine-independent fashion, and will affect only the target machines that need them.

### 3 Interfacing to GCC Output

GCC is normally configured to use the same function calling convention normally in use on the target system. This is done with the machine-description macros described (see Chapter 18 [Target Macros], page 479).

However, returning of structure and union values is done differently on some target machines. As a result, functions compiled with PCC returning such types cannot be called from code compiled with GCC, and vice versa. This does not cause trouble often because few Unix library routines return structures or unions.

GCC code returns structures and unions that are 1, 2, 4 or 8 bytes long in the same registers used for int or double return values. (GCC typically allocates variables of such types in registers also.) Structures and unions of other sizes are returned by storing them into an address passed by the caller (usually in a register). The target hook TARGET\_STRUCT\_VALUE\_RTX tells GCC where to pass this address.

By contrast, PCC on most target machines returns structures and unions of any size by copying the data into an area of static storage, and then returning the address of that storage as if it were a pointer value. The caller must copy the data from that memory area to the place where the value is wanted. This is slower than the method used by GCC, and fails to be reentrant.

On some target machines, such as RISC machines and the 80386, the standard system convention is to pass to the subroutine the address of where to return the value. On these machines, GCC has been configured to be compatible with the standard compiler, when this method is used. It may not be compatible for structures of 1, 2, 4 or 8 bytes.

GCC uses the system's standard convention for passing arguments. On some machines, the first few arguments are passed in registers; in others, all are passed on the stack. It would be possible to use registers for argument passing on any machine, and this would probably result in a significant speedup. But the result would be complete incompatibility with code that follows the standard convention. So this change is practical only if you are switching to GCC as the sole C compiler for the system. We may implement register argument passing on certain machines once we have a complete GNU system so that we can compile the libraries with GCC.

On some machines (particularly the SPARC), certain types of arguments are passed "by invisible reference". This means that the value is stored in memory, and the address of the memory location is passed to the subroutine.

If you use longjmp, beware of automatic variables. ISO C says that automatic variables that are not declared volatile have undefined values after a longjmp. And this is all GCC promises to do, because it is very difficult to restore register variables correctly, and one of GCC's features is that it can put variables in registers without your asking it to.

### 4 The GCC low-level runtime library

GCC provides a low-level runtime library, 'libgcc.a' or 'libgcc\_s.so.1' on some platforms. GCC generates calls to routines in this library automatically, whenever it needs to perform some operation that is too complicated to emit inline code for.

Most of the routines in libgcc handle arithmetic operations that the target processor cannot perform directly. This includes integer multiply and divide on some machines, and all floating-point and fixed-point operations on other machines. libgcc also includes routines for exception handling, and a handful of miscellaneous operations.

Some of these routines can be defined in mostly machine-independent C. Others must be hand-written in assembly language for each processor that needs them.

GCC will also generate calls to C library routines, such as memcpy and memset, in some cases. The set of routines that GCC may possibly use is documented in Section "Other Builtins" in *Using the GNU Compiler Collection (GCC)*.

These routines take arguments and return values of a specific machine mode, not a specific C type. See Section 14.6 [Machine Modes], page 271, for an explanation of this concept. For illustrative purposes, in this chapter the floating point type float is assumed to correspond to SFmode; double to DFmode; and long double to both TFmode and XFmode. Similarly, the integer types int and unsigned int correspond to SImode; long and unsigned long to DImode; and long long and unsigned long to TImode.

#### 4.1 Routines for integer arithmetic

The integer arithmetic routines are used on platforms that don't provide hardware support for arithmetic operations on some modes.

#### 4.1.1 Arithmetic functions

```
int __ashlsi3 (int a, int b)
                                                                     [Runtime Function]
long __ashldi3 (long a, int b)
                                                                     [Runtime Function]
long long __ashlti3 (long long a, int b)
                                                                     [Runtime Function]
     These functions return the result of shifting a left by b bits.
int __ashrsi3 (int a, int b)
                                                                     [Runtime Function]
long __ashrdi3 (long a, int b)
                                                                     [Runtime Function]
long long __ashrti3 (long long a, int b)
                                                                     [Runtime Function]
     These functions return the result of arithmetically shifting a right by b bits.
int __divsi3 (int a, int b)
                                                                     [Runtime Function]
long __divdi3 (long a, long b)
                                                                     [Runtime Function]
                                                                     [Runtime Function]
long long __divti3 (long long a, long long b)
     These functions return the quotient of the signed division of a and b.
int __lshrsi3 (int a, int b)
                                                                     [Runtime Function]
long __lshrdi3 (long a, int b)
                                                                     [Runtime Function]
long long __lshrti3 (long long a, int b)
                                                                     [Runtime Function]
     These functions return the result of logically shifting a right by b bits.
```

[Runtime Function]

[Runtime Function]

```
int __modsi3 (int a, int b)
                                                                   [Runtime Function]
long __moddi3 (long a, long b)
                                                                   [Runtime Function]
long long __modti3 (long long a, long long b)
                                                                   [Runtime Function]
     These functions return the remainder of the signed division of a and b.
int __mulsi3 (int a, int b)
                                                                   [Runtime Function]
long __muldi3 (long a, long b)
                                                                   [Runtime Function]
long long __multi3 (long long a, long long b)
                                                                   [Runtime Function]
     These functions return the product of a and b.
long __negdi2 (long a)
                                                                   [Runtime Function]
long long __negti2 (long long a)
                                                                   [Runtime Function]
     These functions return the negation of a.
unsigned int __udivsi3 (unsigned int a, unsigned int b)
                                                                   [Runtime Function]
unsigned long __udivdi3 (unsigned long a, unsigned long b)
                                                                   [Runtime Function]
unsigned long long __udivti3 (unsigned long long a,
                                                                   [Runtime Function]
         unsigned long long b)
     These functions return the quotient of the unsigned division of a and b.
unsigned long __udivmoddi4 (unsigned long a, unsigned long
                                                                   [Runtime Function]
          b, unsigned long *c)
unsigned long long __udivmodti4 (unsigned long long a,
                                                                   [Runtime Function]
         unsigned long long b, unsigned long long *c)
     These functions calculate both the quotient and remainder of the unsigned division
     of a and b. The return value is the quotient, and the remainder is placed in variable
     pointed to by c.
unsigned int __umodsi3 (unsigned int a, unsigned int b)
                                                                   [Runtime Function]
```

unsigned long long b)

These functions return the remainder of the unsigned division of a and b.

unsigned long \_\_umoddi3 (unsigned long a, unsigned long b)

unsigned long long \_\_umodti3 (unsigned long long a,

#### 4.1.2 Comparison functions

The following functions implement integral comparisons. These functions implement a low-level compare, upon which the higher level comparison operators (such as less than and greater than or equal to) can be constructed. The returned values lie in the range zero to two, to allow the high-level operators to be implemented by testing the returned result using either signed or unsigned comparison.

```
    int __cmpdi2 (long a, long b) [Runtime Function]
    int __cmpti2 (long long a, long long b) [Runtime Function]
    These functions perform a signed comparison of a and b. If a is less than b, they return 0; if a is greater than b, they return 2; and if a and b are equal they return 1.
    int __ucmpdi2 (unsigned long a, unsigned long b) [Runtime Function]
    int __ucmpti2 (unsigned long long a, unsigned long long b) [Runtime Function]
    These functions perform an unsigned comparison of a and b. If a is less than b, they return 0; if a is greater than b, they return 2; and if a and b are equal they return 1.
```

#### 4.1.3 Trapping arithmetic functions

The following functions implement trapping arithmetic. These functions call the libc function abort upon signed arithmetic overflow.

```
int __absvsi2 (int a)
                                                                 [Runtime Function]
long __absvdi2 (long a)
                                                                 [Runtime Function]
```

These functions return the absolute value of a.

```
int __addvsi3 (int a, int b)
                                                                  [Runtime Function]
long __addvdi3 (long a, long b)
                                                                   [Runtime Function]
```

These functions return the sum of a and b; that is a + b.

```
int __mulvsi3 (int a, int b)
                                                                   [Runtime Function]
long __mulvdi3 (long a, long b)
                                                                   [Runtime Function]
```

The functions return the product of a and b; that is a \* b.

```
int __negvsi2 (int a)
                                                                 [Runtime Function]
long __negvdi2 (long a)
                                                                 [Runtime Function]
```

These functions return the negation of a; that is -a.

```
int __subvsi3 (int a, int b)
                                                                  [Runtime Function]
long __subvdi3 (long a, long b)
                                                                  [Runtime Function]
```

These functions return the difference between b and a; that is a - b.

#### 4.1.4 Bit operations

```
int __clzsi2 (unsigned int a)
                                                                    [Runtime Function]
int __clzdi2 (unsigned long a)
                                                                    [Runtime Function]
int __clzti2 (unsigned long long a)
                                                                    [Runtime Function]
     These functions return the number of leading 0-bits in a, starting at the most signif-
```

icant bit position. If a is zero, the result is undefined.

```
int __ctzsi2 (unsigned int a)
                                                                  [Runtime Function]
int __ctzdi2 (unsigned long a)
                                                                  [Runtime Function]
int __ctzti2 (unsigned long long a)
                                                                  [Runtime Function]
```

These functions return the number of trailing 0-bits in a, starting at the least significant bit position. If a is zero, the result is undefined.

```
int __ffsdi2 (unsigned long a)
                                                                     [Runtime Function]
int __ffsti2 (unsigned long long a)
                                                                     [Runtime Function]
     These functions return the index of the least significant 1-bit in a, or the value zero
```

if a is zero. The least significant bit is index one.

```
int __paritysi2 (unsigned int a)
                                                                 [Runtime Function]
int __paritydi2 (unsigned long a)
                                                                 [Runtime Function]
int __parityti2 (unsigned long long a)
                                                                 [Runtime Function]
```

These functions return the value zero if the number of bits set in a is even, and the value one otherwise.

```
int __popcountsi2 (unsigned int a) [Runtime Function]
int __popcountdi2 (unsigned long a) [Runtime Function]
int __popcountti2 (unsigned long long a) [Runtime Function]
These functions return the number of bits set in a.

int32_t __bswapsi2 (int32_t a) [Runtime Function]
int64_t __bswapdi2 (int64_t a) [Runtime Function]
These functions return the a byteswapped.
```

#### 4.2 Routines for floating point emulation

The software floating point library is used on machines which do not have hardware support for floating point. It is also used whenever '-msoft-float' is used to disable generation of floating point instructions. (Not all targets support this switch.)

For compatibility with other compilers, the floating point emulation routines can be renamed with the DECLARE\_LIBRARY\_RENAMES macro (see Section 18.12 [Library Calls], page 560). In this section, the default names are used.

Presently the library does not support XFmode, which is used for long double on some architectures.

#### 4.2.1 Arithmetic functions

```
float __addsf3 (float a, float b)
                                                                   [Runtime Function]
double __adddf3 (double a, double b)
                                                                   [Runtime Function]
long double __addtf3 (long double a, long double b)
                                                                   [Runtime Function]
long double __addxf3 (long double a, long double b)
                                                                  [Runtime Function]
     These functions return the sum of a and b.
float __subsf3 (float a, float b)
                                                                  [Runtime Function]
double __subdf3 (double a, double b)
                                                                   [Runtime Function]
long double __subtf3 (long double a, long double b)
                                                                   [Runtime Function]
long double __subxf3 (long double a, long double b)
                                                                  [Runtime Function]
     These functions return the difference between b and a; that is, a - b.
float __mulsf3 (float a, float b)
                                                                   [Runtime Function]
double __muldf3 (double a, double b)
                                                                   [Runtime Function]
long double __multf3 (long double a, long double b)
                                                                   [Runtime Function]
long double __mulxf3 (long double a, long double b)
                                                                   [Runtime Function]
     These functions return the product of a and b.
float __divsf3 (float a, float b)
                                                                   [Runtime Function]
double __divdf3 (double a, double b)
                                                                   [Runtime Function]
long double __divtf3 (long double a, long double b)
                                                                   [Runtime Function]
long double __divxf3 (long double a, long double b)
                                                                   [Runtime Function]
     These functions return the quotient of a and b; that is, a/b.
float __negsf2 (float a)
                                                                  [Runtime Function]
double __negdf2 (double a)
                                                                   [Runtime Function]
long double __negtf2 (long double a)
                                                                   [Runtime Function]
```

unsigned long \_\_fixunstfdi (long double a)

[Runtime Function]

long double \_\_negxf2 (long double a) [Runtime Function]

These functions return the negation of a. They simply flip the sign bit, so they can produce negative zero and negative NaN.

#### 4.2.2 Conversion functions

```
double __extendsfdf2 (float a)
                                                                 [Runtime Function]
long double __extendsftf2 (float a)
                                                                 [Runtime Function]
long double __extendsfxf2 (float a)
                                                                 [Runtime Function]
long double __extenddftf2 (double a)
                                                                 [Runtime Function]
long double __extenddfxf2 (double a)
                                                                 [Runtime Function]
     These functions extend a to the wider mode of their return type.
double __truncxfdf2 (long double a)
                                                                 [Runtime Function]
double __trunctfdf2 (long double a)
                                                                 [Runtime Function]
float __truncxfsf2 (long double a)
                                                                 [Runtime Function]
float __trunctfsf2 (long double a)
                                                                 [Runtime Function]
float __truncdfsf2 (double a)
                                                                 [Runtime Function]
     These functions truncate a to the narrower mode of their return type, rounding toward
     zero.
int __fixsfsi (float a)
                                                                 [Runtime Function]
int __fixdfsi (double a)
                                                                 [Runtime Function]
int __fixtfsi (long double a)
                                                                 [Runtime Function]
int __fixxfsi (long double a)
                                                                 [Runtime Function]
     These functions convert a to a signed integer, rounding toward zero.
long __fixsfdi (float a)
                                                                 [Runtime Function]
long __fixdfdi (double a)
                                                                 [Runtime Function]
long __fixtfdi (long double a)
                                                                 [Runtime Function]
long __fixxfdi (long double a)
                                                                 [Runtime Function]
     These functions convert a to a signed long, rounding toward zero.
long long __fixsfti (float a)
                                                                 [Runtime Function]
long long __fixdfti (double a)
                                                                 [Runtime Function]
long long __fixtfti (long double a)
                                                                 [Runtime Function]
long long __fixxfti (long double a)
                                                                 [Runtime Function]
     These functions convert a to a signed long long, rounding toward zero.
unsigned int __fixunssfsi (float a)
                                                                 [Runtime Function]
unsigned int __fixunsdfsi (double a)
                                                                 [Runtime Function]
unsigned int __fixunstfsi (long double a)
                                                                 [Runtime Function]
unsigned int __fixunsxfsi (long double a)
                                                                 [Runtime Function]
     These functions convert a to an unsigned integer, rounding toward zero. Negative
     values all become zero.
                                                                 [Runtime Function]
unsigned long __fixunssfdi (float a)
unsigned long __fixunsdfdi (double a)
                                                                 [Runtime Function]
```

```
unsigned long __fixunsxfdi (long double a)
                                                                 [Runtime Function]
     These functions convert a to an unsigned long, rounding toward zero. Negative values
     all become zero.
unsigned long long __fixunssfti (float a)
                                                                 [Runtime Function]
unsigned long long __fixunsdfti (double a)
                                                                 [Runtime Function]
unsigned long long __fixunstfti (long double a)
                                                                 [Runtime Function]
unsigned long long __fixunsxfti (long double a)
                                                                 [Runtime Function]
     These functions convert a to an unsigned long long, rounding toward zero. Negative
     values all become zero.
float __floatsisf (int i)
                                                                 [Runtime Function]
double __floatsidf (int i)
                                                                 [Runtime Function]
long double __floatsitf (int i)
                                                                 [Runtime Function]
long double __floatsixf (int i)
                                                                 [Runtime Function]
     These functions convert i, a signed integer, to floating point.
float __floatdisf (long i)
                                                                 [Runtime Function]
double __floatdidf (long i)
                                                                 [Runtime Function]
long double __floatditf (long i)
                                                                 [Runtime Function]
long double __floatdixf (long i)
                                                                 [Runtime Function]
     These functions convert i, a signed long, to floating point.
float __floattisf (long long i)
                                                                 [Runtime Function]
double __floattidf (long long i)
                                                                 [Runtime Function]
long double __floattitf (long long i)
                                                                 [Runtime Function]
long double __floattixf (long long i)
                                                                 [Runtime Function]
     These functions convert i, a signed long long, to floating point.
float __floatunsisf (unsigned int i)
                                                                 [Runtime Function]
double __floatunsidf (unsigned int i)
                                                                 [Runtime Function]
long double __floatunsitf (unsigned int i)
                                                                 [Runtime Function]
long double __floatunsixf (unsigned int i)
                                                                 [Runtime Function]
     These functions convert i, an unsigned integer, to floating point.
float __floatundisf (unsigned long i)
                                                                 [Runtime Function]
double __floatundidf (unsigned long i)
                                                                 [Runtime Function]
long double __floatunditf (unsigned long i)
                                                                 [Runtime Function]
long double __floatundixf (unsigned long i)
                                                                 [Runtime Function]
     These functions convert i, an unsigned long, to floating point.
float __floatuntisf (unsigned long long i)
                                                                 [Runtime Function]
double __floatuntidf (unsigned long long i)
                                                                 [Runtime Function]
long double __floatuntitf (unsigned long long i)
                                                                 [Runtime Function]
long double __floatuntixf (unsigned long long i)
                                                                 [Runtime Function]
     These functions convert i, an unsigned long long, to floating point.
```

#### 4.2.3 Comparison functions

There are two sets of basic comparison functions.

```
int __cmpsf2 (float a, float b) [Runtime Function] int __cmpdf2 (double a, double b) [Runtime Function] int __cmptf2 (long double a, long double b) [Runtime Function] These functions calculate a <=> b. That is, if a is less than b, they return -1; if a is greater than b, they return 1; and if a and b are equal they return 0. If either argument is NaN they return 1, but you should not rely on this; if NaN is a possibility, use one of the higher-level comparison functions.
```

```
int __unordsf2 (float a, float b) [Runtime Function]
int __unorddf2 (double a, double b) [Runtime Function]
int __unordtf2 (long double a, long double b) [Runtime Function]
These functions return a nonzero value if either argument is NaN, otherwise 0.
```

There is also a complete group of higher level functions which correspond directly to comparison operators. They implement the ISO C semantics for floating-point comparisons, taking NaN into account. Pay careful attention to the return values defined for each set. Under the hood, all of these routines are implemented as

```
if (__unordXf2 (a, b))
  return E;
return __cmpXf2 (a, b);
```

where E is a constant chosen to give the proper behavior for NaN. Thus, the meaning of the return value is different for each set. Do not rely on this implementation; only the semantics documented below are guaranteed.

```
int __eqsf2 (float a, float b)
                                                                    [Runtime Function]
int __eqdf2 (double a, double b)
                                                                     [Runtime Function]
int __eqtf2 (long double a, long double b)
                                                                     [Runtime Function]
     These functions return zero if neither argument is NaN, and a and b are equal.
int __nesf2 (float a, float b)
                                                                     [Runtime Function]
int __nedf2 (double a, double b)
                                                                     [Runtime Function]
int __netf2 (long double a, long double b)
                                                                     [Runtime Function]
     These functions return a nonzero value if either argument is NaN, or if a and b are
     unequal.
int __gesf2 (float a, float b)
                                                                     [Runtime Function]
int __gedf2 (double a, double b)
                                                                     [Runtime Function]
int __getf2 (long double a, long double b)
                                                                    [Runtime Function]
     These functions return a value greater than or equal to zero if neither argument is
     NaN, and a is greater than or equal to b.
```

```
int __ltsf2 (float a, float b) [Runtime Function]
int __ltdf2 (double a, double b) [Runtime Function]
int __lttf2 (long double a, long double b) [Runtime Function]
```

These functions return a value less than zero if neither argument is NaN, and a is strictly less than b.

```
int __lesf2 (float a, float b)
                                                                    [Runtime Function]
int __ledf2 (double a, double b)
                                                                    [Runtime Function]
int __letf2 (long double a, long double b)
                                                                    [Runtime Function]
     These functions return a value less than or equal to zero if neither argument is NaN,
     and a is less than or equal to b.
int __gtsf2 (float a, float b)
                                                                    [Runtime Function]
int __gtdf2 (double a, double b)
                                                                    [Runtime Function]
                                                                    [Runtime Function]
int __gttf2 (long double a, long double b)
     These functions return a value greater than zero if neither argument is NaN, and a is
     strictly greater than b.
4.2.4 Other floating-point functions
```

```
float __powisf2 (float a, int b)
                                                                   [Runtime Function]
double __powidf2 (double a, int b)
                                                                   [Runtime Function]
long double __powitf2 (long double a, int b)
                                                                   [Runtime Function]
long double __powixf2 (long double a, int b)
                                                                   [Runtime Function]
     These functions convert raise a to the power b.
complex float __mulsc3 (float a, float b, float c, float d)
                                                                   [Runtime Function]
complex double __muldc3 (double a, double b, double c,
                                                                   [Runtime Function]
         double d)
complex long double __multc3 (long double a, long double
                                                                  [Runtime Function]
          b, long double c, long double d)
complex long double __mulxc3 (long double a, long double
                                                                  [Runtime Function]
          b, long double c, long double d)
     These functions return the product of a + ib and c + id, following the rules of C99
     Annex G.
complex float __divsc3 (float a, float b, float c, float d)
                                                                  [Runtime Function]
complex double __divdc3 (double a, double b, double c,
                                                                  [Runtime Function]
         double d)
complex long double __divtc3 (long double a, long double
                                                                  [Runtime Function]
          b, long double c, long double d)
complex long double __divxc3 (long double a, long double
                                                                  [Runtime Function]
          b, long double c, long double d)
     These functions return the quotient of a + ib and c + id (i.e., (a + ib)/(c + id)),
     following the rules of C99 Annex G.
```

#### 4.3 Routines for decimal floating point emulation

The software decimal floating point library implements IEEE 754-2008 decimal floating point arithmetic and is only activated on selected targets.

The software decimal floating point library supports either DPD (Densely Packed Decimal) or BID (Binary Integer Decimal) encoding as selected at configure time.

#### 4.3.1 Arithmetic functions

```
_Decimal32 __dpd_addsd3 (_Decimal32 a, _Decimal32 b)
                                                                [Runtime Function]
_Decimal32 __bid_addsd3 (_Decimal32 a, _Decimal32 b)
                                                                [Runtime Function]
_Decimal64 __dpd_adddd3 (_Decimal64 a, _Decimal64 b)
                                                                [Runtime Function]
_Decimal64 __bid_adddd3 (_Decimal64 a, _Decimal64 b)
                                                                [Runtime Function]
_Decimal128 __dpd_addtd3 (_Decimal128 a, _Decimal128 b)
                                                                [Runtime Function]
_Decimal128 __bid_addtd3 (_Decimal128 a, _Decimal128 b)
                                                                [Runtime Function]
     These functions return the sum of a and b.
_Decimal32 __dpd_subsd3 (_Decimal32 a, _Decimal32 b)
                                                                [Runtime Function]
_Decimal32 __bid_subsd3 (_Decimal32 a, _Decimal32 b)
                                                                [Runtime Function]
_Decimal64 __dpd_subdd3 (_Decimal64 a, _Decimal64 b)
                                                                [Runtime Function]
_Decimal64 __bid_subdd3 (_Decimal64 a, _Decimal64 b)
                                                                [Runtime Function]
_Decimal128 __dpd_subtd3 (_Decimal128 a, _Decimal128 b)
                                                                [Runtime Function]
_Decimal128 __bid_subtd3 (_Decimal128 a, _Decimal128 b)
                                                                [Runtime Function]
     These functions return the difference between b and a; that is, a - b.
_Decimal32 __dpd_mulsd3 (_Decimal32 a, _Decimal32 b)
                                                                [Runtime Function]
_Decimal32 __bid_mulsd3 (_Decimal32 a, _Decimal32 b)
                                                                [Runtime Function]
_Decimal64 __dpd_muldd3 (_Decimal64 a, _Decimal64 b)
                                                                [Runtime Function]
_Decimal64 __bid_muldd3 (_Decimal64 a, _Decimal64 b)
                                                                [Runtime Function]
_Decimal128 __dpd_multd3 (_Decimal128 a, _Decimal128 b)
                                                                [Runtime Function]
_Decimal128 __bid_multd3 (_Decimal128 a, _Decimal128 b)
                                                                [Runtime Function]
     These functions return the product of a and b.
_Decimal32 __dpd_divsd3 (_Decimal32 a, _Decimal32 b)
                                                                [Runtime Function]
_Decimal32 __bid_divsd3 (_Decimal32 a, _Decimal32 b)
                                                                [Runtime Function]
_Decimal64 __dpd_divdd3 (_Decimal64 a, _Decimal64 b)
                                                                [Runtime Function]
_Decimal64 __bid_divdd3 (_Decimal64 a, _Decimal64 b)
                                                                [Runtime Function]
_Decimal128 __dpd_divtd3 (_Decimal128 a, _Decimal128 b)
                                                                [Runtime Function]
_Decimal128 __bid_divtd3 (_Decimal128 a, _Decimal128 b)
                                                                [Runtime Function]
     These functions return the quotient of a and b; that is, a/b.
_Decimal32 __dpd_negsd2 (_Decimal32 a)
                                                                [Runtime Function]
_Decimal32 __bid_negsd2 (_Decimal32 a)
                                                                [Runtime Function]
_Decimal64 __dpd_negdd2 (_Decimal64 a)
                                                                [Runtime Function]
_Decimal64 __bid_negdd2 (_Decimal64 a)
                                                                [Runtime Function]
_Decimal128 __dpd_negtd2 (_Decimal128 a)
                                                                [Runtime Function]
_Decimal128 __bid_negtd2 (_Decimal128 a)
                                                                [Runtime Function]
     These functions return the negation of a. They simply flip the sign bit, so they can
     produce negative zero and negative NaN.
```

#### 4.3.2 Conversion functions

_Decimal64dpd_extendsddd2 (_Decimal32 a)	[Runtime Function]
_Decimal64bid_extendsddd2 (_Decimal32 a)	[Runtime Function]
_Decimal128dpd_extendsdtd2 (_Decimal32 a)	[Runtime Function]
_Decimal128bid_extendsdtd2 (_Decimal32 a)	[Runtime Function]

```
_Decimal128 __dpd_extendddtd2 (_Decimal64 a)
                                                               [Runtime Function]
_Decimal128 __bid_extendddtd2 (_Decimal64 a)
                                                               [Runtime Function]
_Decimal32 __dpd_truncddsd2 (_Decimal64 a)
                                                               [Runtime Function]
_Decimal32 __bid_truncddsd2 (_Decimal64 a)
                                                               [Runtime Function]
_Decimal32 __dpd_trunctdsd2 (_Decimal128 a)
                                                               [Runtime Function]
_Decimal32 __bid_trunctdsd2 (_Decimal128 a)
                                                               [Runtime Function]
_Decimal64 __dpd_trunctddd2 (_Decimal128 a)
                                                               [Runtime Function]
_Decimal64 __bid_trunctddd2 (_Decimal128 a)
                                                               [Runtime Function]
     These functions convert the value a from one decimal floating type to another.
_Decimal64 __dpd_extendsfdd (float a)
                                                               [Runtime Function]
_Decimal64 __bid_extendsfdd (float a)
                                                               [Runtime Function]
_Decimal128 __dpd_extendsftd (float a)
                                                               [Runtime Function]
_Decimal128 __bid_extendsftd (float a)
                                                               [Runtime Function]
_Decimal128 __dpd_extenddftd (double a)
                                                               [Runtime Function]
_Decimal128 __bid_extenddftd (double a)
                                                               [Runtime Function]
_Decimal128 __dpd_extendxftd (long double a)
                                                               [Runtime Function]
_Decimal128 __bid_extendxftd (long double a)
                                                               [Runtime Function]
_Decimal32 __dpd_truncdfsd (double a)
                                                               [Runtime Function]
_Decimal32 __bid_truncdfsd (double a)
                                                               [Runtime Function]
_Decimal32 __dpd_truncxfsd (long double a)
                                                               [Runtime Function]
_Decimal32 __bid_truncxfsd (long double a)
                                                               [Runtime Function]
_Decimal32 __dpd_trunctfsd (long double a)
                                                               [Runtime Function]
_Decimal32 __bid_trunctfsd (long double a)
                                                               [Runtime Function]
_Decimal64 __dpd_truncxfdd (long double a)
                                                               [Runtime Function]
_Decimal64 __bid_truncxfdd (long double a)
                                                               [Runtime Function]
_Decimal64 __dpd_trunctfdd (long double a)
                                                               [Runtime Function]
_Decimal64 __bid_trunctfdd (long double a)
                                                               [Runtime Function]
     These functions convert the value of a from a binary floating type to a decimal floating
     type of a different size.
float __dpd_truncddsf (_Decimal64 a)
                                                               [Runtime Function]
float __bid_truncddsf (_Decimal64 a)
                                                               [Runtime Function]
float __dpd_trunctdsf (_Decimal128 a)
                                                               [Runtime Function]
float __bid_trunctdsf (_Decimal128 a)
                                                               [Runtime Function]
double __dpd_extendsddf (_Decimal32 a)
                                                               [Runtime Function]
double __bid_extendsddf (_Decimal32 a)
                                                               [Runtime Function]
double __dpd_trunctddf (_Decimal128 a)
                                                               [Runtime Function]
double __bid_trunctddf (_Decimal128 a)
                                                               [Runtime Function]
long double __dpd_extendsdxf (_Decimal32 a)
                                                               [Runtime Function]
long double __bid_extendsdxf (_Decimal32 a)
                                                               [Runtime Function]
long double __dpd_extendddxf (_Decimal64 a)
                                                               [Runtime Function]
long double __bid_extendddxf (_Decimal64 a)
                                                               [Runtime Function]
long double __dpd_trunctdxf (_Decimal128 a)
                                                               [Runtime Function]
long double __bid_trunctdxf (_Decimal128 a)
                                                               [Runtime Function]
long double __dpd_extendsdtf (_Decimal32 a)
                                                               [Runtime Function]
long double __bid_extendsdtf (_Decimal32 a)
                                                               [Runtime Function]
long double __dpd_extendddtf (_Decimal64 a)
                                                               [Runtime Function]
```

```
long double __bid_extendddtf (_Decimal64 a)
                                                               [Runtime Function]
     These functions convert the value of a from a decimal floating type to a binary floating
     type of a different size.
_Decimal32 __dpd_extendsfsd (float a)
                                                                [Runtime Function]
_Decimal32 __bid_extendsfsd (float a)
                                                                [Runtime Function]
_Decimal64 __dpd_extenddfdd (double a)
                                                                [Runtime Function]
_Decimal64 __bid_extenddfdd (double a)
                                                                [Runtime Function]
_Decimal128 __dpd_extendtftd (long double a)
                                                                [Runtime Function]
_Decimal128 __bid_extendtftd (long double a)
                                                                [Runtime Function]
float __dpd_truncsdsf (_Decimal32 a)
                                                                [Runtime Function]
float __bid_truncsdsf (_Decimal32 a)
                                                                [Runtime Function]
double __dpd_truncdddf (_Decimal64 a)
                                                                [Runtime Function]
double __bid_truncdddf (_Decimal64 a)
                                                                [Runtime Function]
long double __dpd_trunctdtf (_Decimal128 a)
                                                                [Runtime Function]
long double __bid_trunctdtf (_Decimal128 a)
                                                                [Runtime Function]
     These functions convert the value of a between decimal and binary floating types of
     the same size.
int __dpd_fixsdsi (_Decimal32 a)
                                                                [Runtime Function]
int __bid_fixsdsi (_Decimal32 a)
                                                                [Runtime Function]
int __dpd_fixddsi (_Decimal64 a)
                                                                [Runtime Function]
int __bid_fixddsi (_Decimal64 a)
                                                                [Runtime Function]
int __dpd_fixtdsi (_Decimal128 a)
                                                                [Runtime Function]
int __bid_fixtdsi (_Decimal128 a)
                                                                [Runtime Function]
     These functions convert a to a signed integer.
                                                               [Runtime Function]
long __dpd_fixsddi (_Decimal32 a)
long __bid_fixsddi (_Decimal32 a)
                                                                [Runtime Function]
long __dpd_fixdddi (_Decimal64 a)
                                                                [Runtime Function]
long __bid_fixdddi (_Decimal64 a)
                                                                [Runtime Function]
long __dpd_fixtddi (_Decimal128 a)
                                                                [Runtime Function]
long __bid_fixtddi (_Decimal128 a)
                                                                [Runtime Function]
     These functions convert a to a signed long.
unsigned int __dpd_fixunssdsi (_Decimal32 a)
                                                               [Runtime Function]
unsigned int __bid_fixunssdsi (_Decimal32 a)
                                                                [Runtime Function]
unsigned int __dpd_fixunsddsi (_Decimal64 a)
                                                                [Runtime Function]
unsigned int __bid_fixunsddsi (_Decimal64 a)
                                                                [Runtime Function]
unsigned int __dpd_fixunstdsi (_Decimal128 a)
                                                                [Runtime Function]
unsigned int __bid_fixunstdsi (_Decimal128 a)
                                                               [Runtime Function]
     These functions convert a to an unsigned integer. Negative values all become zero.
unsigned long __dpd_fixunssddi (_Decimal32 a)
                                                               [Runtime Function]
unsigned long __bid_fixunssddi (_Decimal32 a)
                                                                [Runtime Function]
unsigned long __dpd_fixunsdddi (_Decimal64 a)
                                                                [Runtime Function]
unsigned long __bid_fixunsdddi (_Decimal64 a)
                                                                [Runtime Function]
unsigned long __dpd_fixunstddi (_Decimal128 a)
                                                               [Runtime Function]
```

```
unsigned long __bid_fixunstddi (_Decimal128 a)
                                                                [Runtime Function]
     These functions convert a to an unsigned long. Negative values all become zero.
_Decimal32 __dpd_floatsisd (int i)
                                                                [Runtime Function]
_Decimal32 __bid_floatsisd (int i)
                                                                [Runtime Function]
_Decimal64 __dpd_floatsidd (int i)
                                                                [Runtime Function]
_Decimal64 __bid_floatsidd (int i)
                                                                [Runtime Function]
_Decimal128 __dpd_floatsitd (int i)
                                                                [Runtime Function]
_Decimal128 __bid_floatsitd (int i)
                                                                [Runtime Function]
     These functions convert i, a signed integer, to decimal floating point.
_Decimal32 __dpd_floatdisd (long i)
                                                                [Runtime Function]
_Decimal32 __bid_floatdisd (long i)
                                                                [Runtime Function]
_Decimal64 __dpd_floatdidd (long i)
                                                                [Runtime Function]
_Decimal64 __bid_floatdidd (long i)
                                                                [Runtime Function]
_Decimal128 __dpd_floatditd (long i)
                                                                [Runtime Function]
_Decimal128 __bid_floatditd (long i)
                                                                [Runtime Function]
     These functions convert i, a signed long, to decimal floating point.
_Decimal32 __dpd_floatunssisd (unsigned int i)
                                                                [Runtime Function]
_Decimal32 __bid_floatunssisd (unsigned int i)
                                                                [Runtime Function]
_Decimal64 __dpd_floatunssidd (unsigned int i)
                                                                [Runtime Function]
_Decimal64 __bid_floatunssidd (unsigned int i)
                                                                [Runtime Function]
_Decimal128 __dpd_floatunssitd (unsigned int i)
                                                                [Runtime Function]
_Decimal128 __bid_floatunssitd (unsigned int i)
                                                                [Runtime Function]
     These functions convert i, an unsigned integer, to decimal floating point.
_Decimal32 __dpd_floatunsdisd (unsigned long i)
                                                                [Runtime Function]
_Decimal32 __bid_floatunsdisd (unsigned long i)
                                                                [Runtime Function]
_Decimal64 __dpd_floatunsdidd (unsigned long i)
                                                                [Runtime Function]
_Decimal64 __bid_floatunsdidd (unsigned long i)
                                                                [Runtime Function]
_Decimal128 __dpd_floatunsditd (unsigned long i)
                                                                [Runtime Function]
_Decimal128 __bid_floatunsditd (unsigned long i)
                                                                [Runtime Function]
     These functions convert i, an unsigned long, to decimal floating point.
```

#### 4.3.3 Comparison functions

These functions return a nonzero value if either argument is NaN, otherwise 0.

There is also a complete group of higher level functions which correspond directly to comparison operators. They implement the ISO C semantics for floating-point comparisons, taking NaN into account. Pay careful attention to the return values defined for each set. Under the hood, all of these routines are implemented as

[Runtime Function]

```
if (__bid_unordXd2 (a, b))
 return E;
return __bid_cmpXd2 (a, b);
```

where E is a constant chosen to give the proper behavior for NaN. Thus, the meaning of the return value is different for each set. Do not rely on this implementation; only the semantics documented below are guaranteed.

```
int __dpd_eqsd2 (_Decimal32 a, _Decimal32 b)
                                                                 [Runtime Function]
int __bid_egsd2 (_Decimal32 a, _Decimal32 b)
                                                                 [Runtime Function]
int __dpd_eqdd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __bid_eqdd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __dpd_eqtd2 (_Decimal128 a, _Decimal128 b)
                                                                  [Runtime Function]
int __bid_eqtd2 (_Decimal128 a, _Decimal128 b)
                                                                 [Runtime Function]
     These functions return zero if neither argument is NaN, and a and b are equal.
int __dpd_nesd2 (_Decimal32 a, _Decimal32 b)
                                                                  [Runtime Function]
int __bid_nesd2 (_Decimal32 a, _Decimal32 b)
                                                                  [Runtime Function]
int __dpd_nedd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __bid_nedd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __dpd_netd2 (_Decimal128 a, _Decimal128 b)
                                                                  [Runtime Function]
int __bid_netd2 (_Decimal128 a, _Decimal128 b)
                                                                 [Runtime Function]
     These functions return a nonzero value if either argument is NaN, or if a and b are
     unequal.
int __dpd_gesd2 (_Decimal32 a, _Decimal32 b)
                                                                 [Runtime Function]
int __bid_gesd2 (_Decimal32 a, _Decimal32 b)
                                                                  [Runtime Function]
int __dpd_gedd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __bid_gedd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __dpd_getd2 (_Decimal128 a, _Decimal128 b)
                                                                  [Runtime Function]
int __bid_getd2 (_Decimal128 a, _Decimal128 b)
                                                                 [Runtime Function]
     These functions return a value greater than or equal to zero if neither argument is
     NaN, and a is greater than or equal to b.
int __dpd_ltsd2 (_Decimal32 a, _Decimal32 b)
                                                                 [Runtime Function]
int __bid_ltsd2 (_Decimal32 a, _Decimal32 b)
                                                                  [Runtime Function]
int __dpd_ltdd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __bid_ltdd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __dpd_1ttd2 (_Decimal128 a, _Decimal128 b)
                                                                  [Runtime Function]
int __bid_1ttd2 (_Decimal128 a, _Decimal128 b)
                                                                 [Runtime Function]
     These functions return a value less than zero if neither argument is NaN, and a is
     strictly less than b.
int __dpd_lesd2 (_Decimal32 a, _Decimal32 b)
                                                                 [Runtime Function]
int __bid_lesd2 (_Decimal32 a, _Decimal32 b)
                                                                  [Runtime Function]
int __dpd_ledd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __bid_ledd2 (_Decimal64 a, _Decimal64 b)
                                                                  [Runtime Function]
int __dpd_letd2 (_Decimal128 a, _Decimal128 b)
```

```
int __bid_letd2 (_Decimal128 a, _Decimal128 b) [Runtime Function] These functions return a value less than or equal to zero if neither argument is NaN, and a is less than or equal to b.
```

## 4.4 Routines for fixed-point fractional emulation

The software fixed-point library implements fixed-point fractional arithmetic, and is only activated on selected targets.

For ease of comprehension fract is an alias for the \_Fract type, accum an alias for \_Accum, and sat an alias for \_Sat.

For illustrative purposes, in this section the fixed-point fractional type short fract is assumed to correspond to machine mode QQmode; unsigned short fract to UQQmode; fract to HQmode; unsigned fract to UHQmode; long fract to SQmode; unsigned long fract to USQmode; long long fract to UDQmode. Similarly the fixed-point accumulator type short accum corresponds to HAmode; unsigned short accum to UHAmode; accum to SAmode; unsigned accum to USAmode; long accum to DAmode; unsigned long accum to UDAmode; long long accum to TAmode; and unsigned long long accum to UTAmode.

#### 4.4.1 Arithmetic functions

strictly greater than b.

```
short fract __addqq3 (short fract a, short fract b)
                                                                 [Runtime Function]
fract __addhq3 (fract a, fract b)
                                                                 [Runtime Function]
long fract __addsq3 (long fract a, long fract b)
                                                                 [Runtime Function]
long long fract __adddq3 (long long fract a, long long fract
                                                                 [Runtime Function]
unsigned short fract __adduqq3 (unsigned short fract a,
                                                                 [Runtime Function]
         unsigned short fract b)
unsigned fract __adduhq3 (unsigned fract a, unsigned fract
                                                                 [Runtime Function]
unsigned long fract __addusq3 (unsigned long fract a,
                                                                 [Runtime Function]
         unsigned long fract b)
unsigned long long fract __addudq3 (unsigned long long
                                                                 [Runtime Function]
         fract a, unsigned long long fract b)
short accum __addha3 (short accum a, short accum b)
                                                                 [Runtime Function]
accum __addsa3 (accum a, accum b)
                                                                 [Runtime Function]
long accum __addda3 (long accum a, long accum b)
                                                                 [Runtime Function]
long long accum __addta3 (long long accum a, long long
                                                                 [Runtime Function]
         accum b)
```

unsigned short accumadduha3 (unsigned short accum a, unsigned short accum b)	[Runtime Function]
unsigned accumaddusa3 (unsigned accum a, unsigned accum b)	[Runtime Function]
unsigned long accumadduda3 (unsigned long accum a, unsigned long accum b)	[Runtime Function]
unsigned long long accumadduta3 (unsigned long long accum a, unsigned long long accum b)  These functions return the sum of a and b.	[Runtime Function]
<pre>short fractssaddqq3 (short fract a, short fract b) fractssaddhq3 (fract a, fract b) long fractssaddsq3 (long fract a, long fract b)</pre>	[Runtime Function] [Runtime Function] [Runtime Function]
<pre>long long fractssadddq3 (long long fract a, long long</pre>	[Runtime Function]
short accumssaddha3 (short accum a, short accum b) accumssaddsa3 (accum a, accum b) long accumssaddda3 (long accum a, long accum b) long long accumssaddta3 (long long accum a, long long accum b)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
These functions return the sum of $a$ and $b$ with signed saturation	n.
<pre>unsigned short fractusadduqq3 (unsigned short fract a,</pre>	[Runtime Function]
unsigned fractusadduhq3 (unsigned fract a, unsigned fract b)	[Runtime Function]
unsigned long fractusaddusq3 (unsigned long fract a, unsigned long fract b)	[Runtime Function]
unsigned long long fractusaddudq3 (unsigned long long fract a, unsigned long long fract b)	[Runtime Function]
<pre>unsigned short accumusadduha3 (unsigned short accum a, unsigned short accum b)</pre>	[Runtime Function]
<pre>unsigned accumusaddusa3 (unsigned accum a, unsigned</pre>	[Runtime Function]
<pre>unsigned long accumusadduda3 (unsigned long accum a,</pre>	[Runtime Function]
unsigned long long accumusadduta3 (unsigned long long accum a, unsigned long long accum b)  These functions return the sum of a and b with unsigned satura	[Runtime Function] tion.
<pre>short fractsubqq3 (short fract a, short fract b) fractsubhq3 (fract a, fract b) long fractsubsq3 (long fract a, long fract b) long long fractsubdq3 (long long fract a, long long fract b)</pre>	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
<pre>unsigned short fractsubuqq3 (unsigned short fract a,</pre>	[Runtime Function]

unsigned fractsubuhq3 (unsigned fract a, unsigned fract b)	[Runtime Function]
unsigned long fractsubusq3 (unsigned long fract a, unsigned long fract b)	[Runtime Function]
unsigned long long fractsubudq3 (unsigned long long fract a, unsigned long long fract b)	[Runtime Function]
short accumsubha3 (short accum a, short accum b) accumsubsa3 (accum a, accum b) long accumsubda3 (long accum a, long accum b) long long accumsubta3 (long long accum a, long long accum b)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
unsigned short accumsubuha3 (unsigned short accum a, unsigned short accum b)	[Runtime Function]
unsigned accumsubusa3 (unsigned accum a, unsigned accum b)	[Runtime Function]
unsigned long accumsubuda3 (unsigned long accum a, unsigned long accum b)	[Runtime Function]
unsigned long long accumsubuta3 (unsigned long long accum a, unsigned long long accum b)  These functions return the difference of a and b; that is, a - b.	[Runtime Function]
short fractsssubq3 (short fract a, short fract b)  fractsssubhq3 (fract a, fract b)  long fractsssubsq3 (long fract a, long fract b)  long long fractsssubdq3 (long long fract a, long long	[Runtime Function]
<pre>unsigned short fractussubuqq3 (unsigned short fract a,</pre>	[Runtime Function]
fract b) unsigned long fractussubusq3 (unsigned long fract a,	[Runtime Function]
unsigned long fract b) unsigned long long fractussubudq3 (unsigned long long fract a, unsigned long long fract b)	[Runtime Function]
unsigned short accumussubuha3 (unsigned short accum a, unsigned short accum b)	[Runtime Function]
unsigned accumussubusa3 (unsigned accum a, unsigned accum b)	[Runtime Function]
unsigned long accumussubuda3 (unsigned long accum a, unsigned long accum b)	[Runtime Function]

[Runtime Function]

unsigned long long accum \_\_ussubuta3 (unsigned long [Runtime Function] long accum a, unsigned long long accum b) These functions return the difference of a and b with unsigned saturation; that is, a - b. short fract \_\_mulqq3 (short fract a, short fract b) [Runtime Function] fract \_\_mulhq3 (fract a, fract b) [Runtime Function] long fract \_\_mulsq3 (long fract a, long fract b) [Runtime Function] long long fract \_\_muldq3 (long long fract a, long long fract [Runtime Function] b) unsigned short fract \_\_muluqq3 (unsigned short fract a, [Runtime Function] unsigned short fract b) unsigned fract \_\_muluhq3 (unsigned fract a, unsigned fract [Runtime Function] unsigned long fract \_\_mulusq3 (unsigned long fract a, [Runtime Function] unsigned long fract b) unsigned long long fract \_\_muludq3 (unsigned long long [Runtime Function] fract a, unsigned long long fract b) short accum \_\_mulha3 (short accum a, short accum b) [Runtime Function] accum \_\_mulsa3 (accum a, accum b) [Runtime Function] long accum \_\_mulda3 (long accum a, long accum b) [Runtime Function] long long accum \_\_multa3 (long long accum a, long long [Runtime Function] accum b) unsigned short accum \_\_muluha3 (unsigned short accum a, [Runtime Function] unsigned short accum b) unsigned accum \_\_mulusa3 (unsigned accum a, unsigned [Runtime Function] accum b) unsigned long accum \_\_muluda3 (unsigned long accum a, [Runtime Function] unsigned long accum b) unsigned long long accum \_\_muluta3 (unsigned long long [Runtime Function] accum a, unsigned long long accum b) These functions return the product of a and b. short fract \_\_ssmulqq3 (short fract a, short fract b) [Runtime Function] fract \_\_ssmulhq3 (fract a, fract b) [Runtime Function] long fract \_\_ssmulsq3 (long fract a, long fract b) [Runtime Function] long long fract \_\_ssmuldq3 (long long fract a, long long [Runtime Function] fract b) short accum \_\_ssmulha3 (short accum a, short accum b) [Runtime Function] accum \_\_ssmulsa3 (accum a, accum b) [Runtime Function] long accum \_\_ssmulda3 (long accum a, long accum b) [Runtime Function] long long accum \_\_ssmulta3 (long long accum a, long long

These functions return the product of a and b with signed saturation.

accum b)

```
unsigned short fract __usmuluqq3 (unsigned short fract a,
                                                                [Runtime Function]
         unsigned short fract b)
unsigned fract __usmuluhq3 (unsigned fract a, unsigned
                                                                [Runtime Function]
         fract b)
unsigned long fract __usmulusq3 (unsigned long fract a,
                                                                [Runtime Function]
         unsigned long fract b)
unsigned long long fract __usmuludq3 (unsigned long
                                                                [Runtime Function]
         long fract a, unsigned long long fract b)
unsigned short accum __usmuluha3 (unsigned short accum
                                                                [Runtime Function]
         a, unsigned short accum b)
unsigned accum __usmulusa3 (unsigned accum a, unsigned
                                                                [Runtime Function]
         accum b)
unsigned long accum __usmuluda3 (unsigned long accum a,
                                                                [Runtime Function]
         unsigned long accum b)
unsigned long long accum __usmuluta3 (unsigned long
                                                                [Runtime Function]
         long accum a, unsigned long long accum b)
     These functions return the product of a and b with unsigned saturation.
short fract __divqq3 (short fract a, short fract b)
                                                                 [Runtime Function]
fract __divhq3 (fract a, fract b)
                                                                 [Runtime Function]
long fract __divsq3 (long fract a, long fract b)
                                                                 [Runtime Function]
long long fract __divdq3 (long long fract a, long long fract
                                                                 [Runtime Function]
short accum __divha3 (short accum a, short accum b)
                                                                 [Runtime Function]
accum __divsa3 (accum a, accum b)
                                                                 [Runtime Function]
long accum __divda3 (long accum a, long accum b)
                                                                 [Runtime Function]
long long accum __divta3 (long long accum a, long long
                                                                 [Runtime Function]
         accum b)
     These functions return the quotient of the signed division of a and b.
unsigned short fract __udivuqq3 (unsigned short fract a,
                                                                [Runtime Function]
         unsigned short fract b)
unsigned fract __udivuhq3 (unsigned fract a, unsigned fract
                                                                [Runtime Function]
unsigned long fract __udivusq3 (unsigned long fract a,
                                                                [Runtime Function]
         unsigned long fract b)
unsigned long long fract __udivudq3 (unsigned long long
                                                                [Runtime Function]
         fract a, unsigned long long fract b)
unsigned short accum __udivuha3 (unsigned short accum a,
                                                                [Runtime Function]
         unsigned short accum b)
unsigned accum __udivusa3 (unsigned accum a, unsigned
                                                                [Runtime Function]
         accum b)
unsigned long accum __udivuda3 (unsigned long accum a,
                                                                [Runtime Function]
         unsigned long accum b)
unsigned long long accum __udivuta3 (unsigned long long
                                                                [Runtime Function]
         accum a, unsigned long long accum b)
     These functions return the quotient of the unsigned division of a and b.
```

<pre>short fractssdivqq3 (short fract a, short fract b) fractssdivhq3 (fract a, fract b) long fractssdivsq3 (long fract a, long fract b) long long fractssdivdq3 (long long fract a, long long fract b)</pre>	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
<pre>short accumssdivha3 (short accum a, short accum b) accumssdivsa3 (accum a, accum b) long accumssdivda3 (long accum a, long accum b)</pre>	[Runtime Function] [Runtime Function] [Runtime Function]
long long accumssdivta3 (long long accum a, long long accum b)  These functions return the quotient of the signed division of a	[Runtime Function] $a$ and $b$ with signed
saturation.	
unsigned short fractusdivuqq3 (unsigned short fract a, unsigned short fract b)	[Runtime Function]
<pre>unsigned fractusdivuhq3 (unsigned fract a, unsigned</pre>	[Runtime Function]
unsigned long fractusdivusq3 (unsigned long fract a, unsigned long fract b)	[Runtime Function]
unsigned long long fractusdivudq3 (unsigned long long fract a, unsigned long long fract b)	[Runtime Function]
unsigned short accumusdivuha3 (unsigned short accum a, unsigned short accum b)	[Runtime Function]
unsigned accumusdivusa3 (unsigned accum a, unsigned accum b)	[Runtime Function]
unsigned long accumusdivuda3 (unsigned long accum a, unsigned long accum b)	[Runtime Function]
unsigned long long accumusdivuta3 (unsigned long long accum a, unsigned long long accum b)	[Runtime Function]
These functions return the quotient of the unsigned division of a saturation.	and $b$ with unsigned
<pre>short fractnegqq2 (short fract a)</pre>	[Runtime Function]
fractneghq2 (fract a)	[Runtime Function]
long fractnegsq2 (long fract a)	[Runtime Function]
long long fractnegdq2 (long long fract a)	[Runtime Function]
unsigned short fractneguqq2 (unsigned short fract a)	[Runtime Function]
unsigned fractneguhq2 (unsigned fract a)	[Runtime Function]
unsigned long fractnegusq2 (unsigned long fract a)	[Runtime Function]
<pre>unsigned long long fractnegudq2 (unsigned long long fract a)</pre>	[Runtime Function]
short accumnegha2 (short accum a)	[Runtime Function]
accumnegsa2 (accum a)	[Runtime Function]
long accumnegda2 (long accum a)	[Runtime Function]
long long accumnegta2 (long long accum a)	[Runtime Function]
unsigned short accumneguha2 (unsigned short accum a)	[Runtime Function]
unsigned accumnegusa2 (unsigned accum a)	[Runtime Function]
unsigned long accumneguda2 (unsigned long accum a)	[Runtime Function]

int b)

```
unsigned long long accum __neguta2 (unsigned long long
                                                               [Runtime Function]
         accum a)
     These functions return the negation of a.
short fract __ssnegqq2 (short fract a)
                                                                [Runtime Function]
fract __ssneghq2 (fract a)
                                                                [Runtime Function]
long fract __ssnegsq2 (long fract a)
                                                                [Runtime Function]
long long fract __ssnegdq2 (long long fract a)
                                                                [Runtime Function]
short accum __ssnegha2 (short accum a)
                                                                [Runtime Function]
accum __ssnegsa2 (accum a)
                                                                [Runtime Function]
long accum __ssnegda2 (long accum a)
                                                                [Runtime Function]
long long accum __ssnegta2 (long long accum a)
                                                                [Runtime Function]
     These functions return the negation of a with signed saturation.
unsigned short fract __usneguqq2 (unsigned short fract a)
                                                                [Runtime Function]
unsigned fract __usneguhq2 (unsigned fract a)
                                                                [Runtime Function]
unsigned long fract __usnegusq2 (unsigned long fract a)
                                                                [Runtime Function]
unsigned long long fract __usnegudq2 (unsigned long
                                                                [Runtime Function]
         long fract a)
unsigned short accum __usneguha2 (unsigned short accum
                                                               [Runtime Function]
unsigned accum __usnegusa2 (unsigned accum a)
                                                                [Runtime Function]
unsigned long accum __usneguda2 (unsigned long accum a)
                                                                [Runtime Function]
unsigned long long accum __usneguta2 (unsigned long
                                                                [Runtime Function]
         long accum a)
     These functions return the negation of a with unsigned saturation.
short fract __ashlqq3 (short fract a, int b)
                                                                [Runtime Function]
fract __ashlhq3 (fract a, int b)
                                                                [Runtime Function]
long fract __ashlsq3 (long fract a, int b)
                                                                [Runtime Function]
long long fract __ashldq3 (long long fract a, int b)
                                                                [Runtime Function]
unsigned short fract __ashlugq3 (unsigned short fract a,
                                                                [Runtime Function]
         int b)
unsigned fract __ashluhq3 (unsigned fract a, int b)
                                                               [Runtime Function]
unsigned long fract __ashlusq3 (unsigned long fract a, int
                                                               [Runtime Function]
unsigned long long fract __ashludq3 (unsigned long long
                                                               [Runtime Function]
         fract a, int b)
short accum __ashlha3 (short accum a, int b)
                                                                [Runtime Function]
accum __ashlsa3 (accum a, int b)
                                                                [Runtime Function]
long accum __ashlda3 (long accum a, int b)
                                                                [Runtime Function]
long long accum __ashlta3 (long long accum a, int b)
                                                                [Runtime Function]
unsigned short accum __ashluha3 (unsigned short accum a,
                                                                [Runtime Function]
unsigned accum __ashlusa3 (unsigned accum a, int b)
                                                                [Runtime Function]
unsigned long accum __ashluda3 (unsigned long accum a,
                                                                [Runtime Function]
```

a, int b)

```
unsigned long long accum __ashluta3 (unsigned long long
                                                                 [Runtime Function]
         accum a, int b)
     These functions return the result of shifting a left by b bits.
short fract __ashrqq3 (short fract a, int b)
                                                                 [Runtime Function]
fract __ashrhq3 (fract a, int b)
                                                                 [Runtime Function]
long fract __ashrsq3 (long fract a, int b)
                                                                 [Runtime Function]
long long fract __ashrdq3 (long long fract a, int b)
                                                                 [Runtime Function]
short accum __ashrha3 (short accum a, int b)
                                                                 [Runtime Function]
accum __ashrsa3 (accum a, int b)
                                                                 [Runtime Function]
long accum __ashrda3 (long accum a, int b)
                                                                 [Runtime Function]
long long accum __ashrta3 (long long accum a, int b)
                                                                 [Runtime Function]
     These functions return the result of arithmetically shifting a right by b bits.
unsigned short fract __lshruqq3 (unsigned short fract a,
                                                                 [Runtime Function]
         int b)
unsigned fract __lshruhq3 (unsigned fract a, int b)
                                                                 [Runtime Function]
unsigned long fract __lshrusq3 (unsigned long fract a, int
                                                                 [Runtime Function]
unsigned long long fract __lshrudq3 (unsigned long long
                                                                 [Runtime Function]
         fract a, int b)
unsigned short accum __lshruha3 (unsigned short accum a,
                                                                 [Runtime Function]
unsigned accum __lshrusa3 (unsigned accum a, int b)
                                                                 [Runtime Function]
unsigned long accum __lshruda3 (unsigned long accum a,
                                                                 [Runtime Function]
         int b)
unsigned long long accum __lshruta3 (unsigned long long
                                                                 [Runtime Function]
         accum a, int b)
     These functions return the result of logically shifting a right by b bits.
fract __ssashlhq3 (fract a, int b)
                                                                 [Runtime Function]
long fract __ssashlsq3 (long fract a, int b)
                                                                 [Runtime Function]
long long fract __ssashldq3 (long long fract a, int b)
                                                                 [Runtime Function]
short accum __ssashlha3 (short accum a, int b)
                                                                 [Runtime Function]
accum __ssashlsa3 (accum a, int b)
                                                                 [Runtime Function]
long accum __ssashlda3 (long accum a, int b)
                                                                 [Runtime Function]
long long accum __ssashlta3 (long long accum a, int b)
                                                                 [Runtime Function]
     These functions return the result of shifting a left by b bits with signed saturation.
unsigned short fract __usashluqq3 (unsigned short fract
                                                                 [Runtime Function]
         a, int b)
unsigned fract __usashluhq3 (unsigned fract a, int b)
                                                                 [Runtime Function]
unsigned long fract __usashlusq3 (unsigned long fract a,
                                                                 [Runtime Function]
         int b)
unsigned long long fract __usashludq3 (unsigned long
                                                                 [Runtime Function]
         long fract a, int b)
unsigned short accum __usashluha3 (unsigned short accum
                                                                 [Runtime Function]
```

```
unsigned accum __usashlusa3 (unsigned accum a, int b)

unsigned long accum __usashluda3 (unsigned long accum
a, int b)

unsigned long long accum __usashluta3 (unsigned long long accum a, int b)

[Runtime Function]
[Runtime Function]
```

These functions return the result of shifting a left by b bits with unsigned saturation.

## 4.4.2 Comparison functions

The following functions implement fixed-point comparisons. These functions implement a low-level compare, upon which the higher level comparison operators (such as less than and greater than or equal to) can be constructed. The returned values lie in the range zero to two, to allow the high-level operators to be implemented by testing the returned result using either signed or unsigned comparison.

```
int __cmpqq2 (short fract a, short fract b)
                                                                   [Runtime Function]
int __cmphq2 (fract a, fract b)
                                                                   [Runtime Function]
int __cmpsq2 (long fract a, long fract b)
                                                                   [Runtime Function]
int __cmpdq2 (long long fract a, long long fract b)
                                                                   [Runtime Function]
int __cmpuqq2 (unsigned short fract a, unsigned short fract b)
                                                                   [Runtime Function]
                                                                   [Runtime Function]
int __cmpuhq2 (unsigned fract a, unsigned fract b)
int __cmpusq2 (unsigned long fract a, unsigned long fract b)
                                                                   [Runtime Function]
int __cmpudq2 (unsigned long long fract a, unsigned long long
                                                                   [Runtime Function]
         fract b)
int __cmpha2 (short accum a, short accum b)
                                                                   [Runtime Function]
int __cmpsa2 (accum a, accum b)
                                                                   [Runtime Function]
int __cmpda2 (long accum a, long accum b)
                                                                   [Runtime Function]
int __cmpta2 (long long accum a, long long accum b)
                                                                   [Runtime Function]
int __cmpuha2 (unsigned short accum a, unsigned short accum
                                                                   [Runtime Function]
          b)
int __cmpusa2 (unsigned accum a, unsigned accum b)
                                                                   [Runtime Function]
int __cmpuda2 (unsigned long accum a, unsigned long accum b)
                                                                   [Runtime Function]
int __cmputa2 (unsigned long long accum a, unsigned long long
                                                                   [Runtime Function]
```

These functions perform a signed or unsigned comparison of a and b (depending on the selected machine mode). If a is less than b, they return 0; if a is greater than b, they return 2; and if a and b are equal they return 1.

### 4.4.3 Conversion functions

```
fract __fractqqhq2 (short fract a)
                                                                [Runtime Function]
long fract __fractqqsq2 (short fract a)
                                                                [Runtime Function]
long long fract __fractqqdq2 (short fract a)
                                                                [Runtime Function]
short accum __fractqqha (short fract a)
                                                                [Runtime Function]
accum __fractqqsa (short fract a)
                                                                [Runtime Function]
long accum __fractqqda (short fract a)
                                                                [Runtime Function]
long long accum __fractqqta (short fract a)
                                                                [Runtime Function]
unsigned short fract __fractqquqq (short fract a)
                                                                [Runtime Function]
unsigned fract __fractqquhq (short fract a)
                                                                [Runtime Function]
```

```
unsigned long fract __fractqqusq (short fract a)
                                                               [Runtime Function]
unsigned long long fract __fractqqudq (short fract a)
                                                               [Runtime Function]
unsigned short accum __fractqquha (short fract a)
                                                               [Runtime Function]
unsigned accum __fractqqusa (short fract a)
                                                               [Runtime Function]
unsigned long accum __fractqquda (short fract a)
                                                               [Runtime Function]
unsigned long long accum __fractqquta (short fract a)
                                                               [Runtime Function]
signed char __fractqqqi (short fract a)
                                                               [Runtime Function]
short __fractqqhi (short fract a)
                                                               [Runtime Function]
int __fractqqsi (short fract a)
                                                               [Runtime Function]
long __fractqqdi (short fract a)
                                                               [Runtime Function]
                                                               [Runtime Function]
long long __fractqqti (short fract a)
float __fractqqsf (short fract a)
                                                               [Runtime Function]
double __fractqqdf (short fract a)
                                                               [Runtime Function]
short fract __fracthqqq2 (fract a)
                                                               [Runtime Function]
long fract __fracthqsq2 (fract a)
                                                               [Runtime Function]
long long fract __fracthqdq2 (fract a)
                                                               [Runtime Function]
short accum __fracthqha (fract a)
                                                               [Runtime Function]
accum __fracthqsa (fract a)
                                                               [Runtime Function]
long accum __fracthqda (fract a)
                                                               [Runtime Function]
long long accum __fracthqta (fract a)
                                                               [Runtime Function]
unsigned short fract __fracthquqq (fract a)
                                                               [Runtime Function]
unsigned fract __fracthquhq (fract a)
                                                               [Runtime Function]
unsigned long fract __fracthqusq (fract a)
                                                               [Runtime Function]
unsigned long long fract __fracthqudq (fract a)
                                                               [Runtime Function]
unsigned short accum __fracthquha (fract a)
                                                               [Runtime Function]
unsigned accum __fracthqusa (fract a)
                                                               [Runtime Function]
unsigned long accum __fracthquda (fract a)
                                                               [Runtime Function]
unsigned long long accum __fracthquta (fract a)
                                                               [Runtime Function]
signed char __fracthqqi (fract a)
                                                               [Runtime Function]
short __fracthqhi (fract a)
                                                               [Runtime Function]
int __fracthqsi (fract a)
                                                               [Runtime Function]
long __fracthqdi (fract a)
                                                               [Runtime Function]
long long __fracthqti (fract a)
                                                               [Runtime Function]
float __fracthqsf (fract a)
                                                               [Runtime Function]
double __fracthqdf (fract a)
                                                               [Runtime Function]
short fract __fractsqqq2 (long fract a)
                                                               [Runtime Function]
fract __fractsqhq2 (long fract a)
                                                               [Runtime Function]
long long fract __fractsqdq2 (long fract a)
                                                               [Runtime Function]
short accum __fractsqha (long fract a)
                                                               [Runtime Function]
accum __fractsqsa (long fract a)
                                                               [Runtime Function]
long accum __fractsqda (long fract a)
                                                               [Runtime Function]
long long accum __fractsqta (long fract a)
                                                               [Runtime Function]
unsigned short fract __fractsquqq (long fract a)
                                                               [Runtime Function]
unsigned fract __fractsquhq (long fract a)
                                                               [Runtime Function]
unsigned long fract __fractsqusq (long fract a)
                                                               [Runtime Function]
unsigned long long fract __fractsqudq (long fract a)
                                                               [Runtime Function]
unsigned short accum __fractsquha (long fract a)
                                                               [Runtime Function]
```

```
unsigned accum __fractsqusa (long fract a)
                                                               [Runtime Function]
unsigned long accum __fractsquda (long fract a)
                                                               [Runtime Function]
unsigned long long accum __fractsquta (long fract a)
                                                               [Runtime Function]
signed char __fractsqqi (long fract a)
                                                               [Runtime Function]
short __fractsqhi (long fract a)
                                                               [Runtime Function]
int __fractsqsi (long fract a)
                                                               [Runtime Function]
long __fractsqdi (long fract a)
                                                               [Runtime Function]
long long __fractsqti (long fract a)
                                                               [Runtime Function]
float __fractsqsf (long fract a)
                                                               [Runtime Function]
double __fractsqdf (long fract a)
                                                               [Runtime Function]
short fract __fractdqqq2 (long long fract a)
                                                               [Runtime Function]
fract __fractdqhq2 (long long fract a)
                                                               [Runtime Function]
long fract __fractdqsq2 (long long fract a)
                                                               [Runtime Function]
short accum __fractdqha (long long fract a)
                                                               [Runtime Function]
accum __fractdqsa (long long fract a)
                                                               [Runtime Function]
long accum __fractdqda (long long fract a)
                                                               [Runtime Function]
long long accum __fractdqta (long long fract a)
                                                               [Runtime Function]
unsigned short fract __fractdquqq (long long fract a)
                                                               [Runtime Function]
unsigned fract __fractdquhq (long long fract a)
                                                               [Runtime Function]
unsigned long fract __fractdqusq (long long fract a)
                                                               [Runtime Function]
unsigned long long fract __fractdqudq (long long fract
                                                               [Runtime Function]
         a)
unsigned short accum __fractdquha (long long fract a)
                                                               [Runtime Function]
unsigned accum __fractdqusa (long long fract a)
                                                               [Runtime Function]
unsigned long accum __fractdquda (long long fract a)
                                                               [Runtime Function]
unsigned long long accum __fractdquta (long long fract
                                                               [Runtime Function]
signed char __fractdqqi (long long fract a)
                                                               [Runtime Function]
short __fractdqhi (long long fract a)
                                                               [Runtime Function]
int __fractdqsi (long long fract a)
                                                               [Runtime Function]
long __fractdqdi (long long fract a)
                                                               [Runtime Function]
long long __fractdqti (long long fract a)
                                                               [Runtime Function]
float __fractdqsf (long long fract a)
                                                               [Runtime Function]
double __fractdqdf (long long fract a)
                                                               [Runtime Function]
short fract __fracthaqq (short accum a)
                                                               [Runtime Function]
fract __fracthang (short accum a)
                                                               [Runtime Function]
long fract __fracthasq (short accum a)
                                                               [Runtime Function]
long long fract __fracthadq (short accum a)
                                                               [Runtime Function]
accum __fracthasa2 (short accum a)
                                                               [Runtime Function]
long accum __fracthada2 (short accum a)
                                                               [Runtime Function]
long long accum __fracthata2 (short accum a)
                                                               [Runtime Function]
unsigned short fract __fracthauqq (short accum a)
                                                               [Runtime Function]
unsigned fract __fracthauhq (short accum a)
                                                               [Runtime Function]
unsigned long fract __fracthausq (short accum a)
                                                               [Runtime Function]
unsigned long long fract __fracthaudq (short accum a)
                                                               [Runtime Function]
unsigned short accum __fracthauha (short accum a)
                                                               [Runtime Function]
unsigned accum __fracthausa (short accum a)
                                                               [Runtime Function]
```

unsigned long accumfracthauda (short accum a)	[Runtime Function]
unsigned long long accumfracthauta (short accum a)	[Runtime Function]
signed charfracthaqi (short accum a)	[Runtime Function]
shortfracthahi (short accum a)	[Runtime Function]
intfracthasi (short accum a)	[Runtime Function]
longfracthadi (short accum a)	[Runtime Function]
long longfracthati (short accum a)	[Runtime Function]
floatfracthasf (short accum a)	[Runtime Function]
doublefracthadf (short accum a)	[Runtime Function]
short fractfractsaqq (accum a)	[Runtime Function]
fractfractsahq (accum a)	[Runtime Function]
long fractfractsasq (accum a)	[Runtime Function]
long long fractfractsadq (accum a)	[Runtime Function]
short accumfractsaha2 (accum a)	[Runtime Function]
long accumfractsada2 (accum a)	[Runtime Function]
long long accumfractsata2 (accum a)	[Runtime Function]
unsigned short fractfractsauqq (accum a)	[Runtime Function]
unsigned fractfractsauhq (accum a)	[Runtime Function]
unsigned long fractfractsausq (accum a)	[Runtime Function]
unsigned long long fractfractsaudq (accum a)	[Runtime Function]
unsigned short accumfractsauha (accum a)	[Runtime Function]
unsigned accumfractsausa (accum a)	[Runtime Function]
unsigned long accumfractsauda (accum a)	[Runtime Function]
unsigned long long accumfractsauta (accum a)	[Runtime Function]
signed charfractsaqi (accum a)	[Runtime Function]
shortfractsahi (accum a)	[Runtime Function]
intfractsasi (accum a)	[Runtime Function]
longfractsadi (accum a)	[Runtime Function]
<del></del>	[Runtime Function]
<pre>long longfractsati (accum a) floatfractsasf (accum a)</pre>	[Runtime Function]
,	[Runtime Function]
doublefractsadf (accum a)	[Runtime Function]
short fractfractdaqq (long accum a)	
fractfractdahq (long accum a)	[Runtime Function]
long fractfractdasq (long accum a)	[Runtime Function]
long long fractfractdadq (long accum a)	[Runtime Function]
short accumfractdaha2 (long accum a)	[Runtime Function]
accumfractdasa2 (long accum a)	[Runtime Function]
long long accumfractdata2 (long accum a)	[Runtime Function]
unsigned short fractfractdauqq (long accum a)	[Runtime Function]
unsigned fractfractdauhq (long accum a)	[Runtime Function]
unsigned long fractfractdausq (long accum a)	[Runtime Function]
unsigned long long fractfractdaudq (long accum a)	[Runtime Function]
unsigned short accumfractdauha (long accum a)	[Runtime Function]
unsigned accumfractdausa (long accum a)	[Runtime Function]
unsigned long accumfractdauda (long accum a)	[Runtime Function]
unsigned long long accumfractdauta (long accum a)	[Runtime Function]
signed charfractdaqi (long accum a)	[Runtime Function]

shortfractdahi (long accum a)	[Runtime Function]
intfractdasi (long accum a)	[Runtime Function]
<pre>longfractdadi (long accum a)</pre>	[Runtime Function]
<pre>long longfractdati (long accum a)</pre>	[Runtime Function]
floatfractdasf (long accum a)	[Runtime Function]
doublefractdadf (long accum a)	[Runtime Function]
short fractfracttaqq (long long accum a)	[Runtime Function]
<pre>fractfracttahq (long long accum a)</pre>	[Runtime Function]
<pre>long fractfracttasq (long long accum a)</pre>	[Runtime Function]
<pre>long long fractfracttadq (long long accum a)</pre>	[Runtime Function]
short accumfracttaha2 (long long accum a)	[Runtime Function]
accumfracttasa2 (long long accum a)	[Runtime Function]
<pre>long accumfracttada2 (long long accum a)</pre>	[Runtime Function]
unsigned short fractfracttauqq (long long accum a)	[Runtime Function]
unsigned fractfracttauhq (long long accum a)	[Runtime Function]
unsigned long fractfracttausq (long long accum a)	[Runtime Function]
unsigned long long fractfracttaudq (long long accum	[Runtime Function]
a)	i j
unsigned short accumfracttauha (long long accum a)	[Runtime Function]
unsigned accumfracttausa (long long accum a)	[Runtime Function]
unsigned long accumfracttauda (long long accum a)	[Runtime Function]
unsigned long long accumfracttauta (long long accum	[Runtime Function]
a)	[ ]
signed charfracttaqi (long long accum a)	[Runtime Function]
shortfracttahi (long long accum a)	[Runtime Function]
intfracttasi (long long accum a)	[Runtime Function]
longfracttadi (long long accum a)	[Runtime Function]
long longfracttati (long long accum a)	[Runtime Function]
floatfracttasf (long long accum a)	[Runtime Function]
doublefracttadf (long long accum a)	[Runtime Function]
short fractfractuqqq (unsigned short fract a)	[Runtime Function]
fractfractuqqqq (unsigned short fract a)	[Runtime Function]
long fractfractuqqsq (unsigned short fract a)	[Runtime Function]
	[Runtime Function]
long long fractfractuqqdq (unsigned short fract a)	[Runtime Function]
short accumfractuagqha (unsigned short fract a)	
accumfractuqqsa (unsigned short fract a)	[Runtime Function]
long accumfractuqqda (unsigned short fract a)	[Runtime Function]
long long accumfractuqqta (unsigned short fract a)	[Runtime Function]
unsigned fractfractuqquhq2 (unsigned short fract a)	[Runtime Function]
unsigned long fractfractuqqusq2 (unsigned short fract a)	[Runtime Function]
unsigned long fractfractuqqudq2 (unsigned	[Runtime Function]
short fract a)	
unsigned short accumfractuqquha (unsigned short fract	[Runtime Function]
a) unsigned accumfractuqqusa (unsigned short fract a)	[Runtime Function]
	. ,

```
unsigned long accum __fractuqquda (unsigned short fract
                                                               [Runtime Function]
unsigned long long accum __fractuqquta (unsigned short
                                                               [Runtime Function]
         fract a)
signed char __fractuqqqi (unsigned short fract a)
                                                               [Runtime Function]
short __fractuqqhi (unsigned short fract a)
                                                                [Runtime Function]
int __fractuqqsi (unsigned short fract a)
                                                                [Runtime Function]
long __fractuqqdi (unsigned short fract a)
                                                                [Runtime Function]
long long __fractuqqti (unsigned short fract a)
                                                                [Runtime Function]
float __fractuqqsf (unsigned short fract a)
                                                                [Runtime Function]
double __fractuqqdf (unsigned short fract a)
                                                                [Runtime Function]
short fract __fractuhqqq (unsigned fract a)
                                                                [Runtime Function]
fract __fractuhqhq (unsigned fract a)
                                                                [Runtime Function]
long fract __fractuhqsq (unsigned fract a)
                                                                [Runtime Function]
long long fract __fractuhqdq (unsigned fract a)
                                                                [Runtime Function]
short accum __fractuhqha (unsigned fract a)
                                                                [Runtime Function]
accum __fractuhqsa (unsigned fract a)
                                                                [Runtime Function]
long accum __fractuhqda (unsigned fract a)
                                                                [Runtime Function]
long long accum __fractuhqta (unsigned fract a)
                                                                [Runtime Function]
unsigned short fract __fractuhquqq2 (unsigned fract a)
                                                                [Runtime Function]
unsigned long fract __fractuhqusq2 (unsigned fract a)
                                                                [Runtime Function]
unsigned long long fract __fractuhqudq2 (unsigned
                                                                [Runtime Function]
         fract a)
unsigned short accum __fractuhquha (unsigned fract a)
                                                                [Runtime Function]
unsigned accum __fractuhqusa (unsigned fract a)
                                                                [Runtime Function]
unsigned long accum __fractuhquda (unsigned fract a)
                                                                [Runtime Function]
unsigned long long accum __fractuhquta (unsigned fract
                                                               [Runtime Function]
signed char __fractuhqqi (unsigned fract a)
                                                               [Runtime Function]
short __fractuhqhi (unsigned fract a)
                                                                [Runtime Function]
int __fractuhqsi (unsigned fract a)
                                                                [Runtime Function]
long __fractuhqdi (unsigned fract a)
                                                                [Runtime Function]
long long __fractuhqti (unsigned fract a)
                                                                [Runtime Function]
float __fractuhqsf (unsigned fract a)
                                                                [Runtime Function]
double __fractuhqdf (unsigned fract a)
                                                                [Runtime Function]
short fract __fractusqqq (unsigned long fract a)
                                                                [Runtime Function]
fract __fractusqhq (unsigned long fract a)
                                                                [Runtime Function]
long fract __fractusqsq (unsigned long fract a)
                                                                [Runtime Function]
long long fract __fractusqdq (unsigned long fract a)
                                                                [Runtime Function]
short accum __fractusqha (unsigned long fract a)
                                                                [Runtime Function]
accum __fractusqsa (unsigned long fract a)
                                                                [Runtime Function]
long accum __fractusqda (unsigned long fract a)
                                                                [Runtime Function]
long long accum __fractusqta (unsigned long fract a)
                                                                [Runtime Function]
unsigned short fract __fractusquqq2 (unsigned long fract
                                                               [Runtime Function]
unsigned fract __fractusquhq2 (unsigned long fract a)
                                                               [Runtime Function]
```

unsigned long long fractfractusqudq2 (unsigned long	[Runtime Function]
fract a) unsigned short accumfractusquha (unsigned long fract	[Runtime Function]
unsigned accumfractusqusa (unsigned long fract a) unsigned long accumfractusquda (unsigned long fract a) unsigned long long accumfractusquta (unsigned long	[Runtime Function] [Runtime Function]
fract a) signed charfractusqqi (unsigned long fract a) shortfractusqhi (unsigned long fract a) intfractusqsi (unsigned long fract a) longfractusqdi (unsigned long fract a)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
<pre>long longfractusqti (unsigned long fract a) floatfractusqsf (unsigned long fract a) doublefractusqdf (unsigned long fract a) short fractfractudqqq (unsigned long long fract a)</pre>	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
fractfractudqhq (unsigned long long fract a) long fractfractudqsq (unsigned long long fract a) long long fractfractudqdq (unsigned long long fract a) short accumfractudqha (unsigned long long fract a)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
accumfractudqsa (unsigned long long fract a) long accumfractudqda (unsigned long long fract a) long long accumfractudqta (unsigned long long fract a) unsigned short fractfractudquqq2 (unsigned long long fract a)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
unsigned fractfractudquhq2 (unsigned long long fract a) unsigned long fractfractudqusq2 (unsigned long long fract a)	[Runtime Function] [Runtime Function]
unsigned short accumfractudquha (unsigned long long fract a)	[Runtime Function]
unsigned accumfractudqusa (unsigned long long fract a) unsigned long accumfractudquda (unsigned long long fract a)	[Runtime Function] [Runtime Function]
unsigned long long accumfractudquta (unsigned long long fract a)	[Runtime Function]
signed charfractudqqi (unsigned long long fract a) shortfractudqhi (unsigned long long fract a) intfractudqsi (unsigned long long fract a)	[Runtime Function] [Runtime Function] [Runtime Function]
<pre>longfractudqdi (unsigned long long fract a) long longfractudqti (unsigned long long fract a) floatfractudqsf (unsigned long long fract a) doublefractudqdf (unsigned long long fract a)</pre>	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
short fractfractuhaqq (unsigned short accum a) fractfractuhahq (unsigned short accum a) long fractfractuhasq (unsigned short accum a) long long fractfractuhadq (unsigned short accum a) short accumfractuhaha (unsigned short accum a)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]

accumfractuhasa (unsigned short accum a) long accumfractuhada (unsigned short accum a) long long accumfractuhata (unsigned short accum a) unsigned short fractfractuhauqq (unsigned short accum a)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
unsigned fractfractuhauhq (unsigned short accum a) unsigned long fractfractuhausq (unsigned short accum a)	[Runtime Function] [Runtime Function]
unsigned long long fractfractuhaudq (unsigned short accum a)	[Runtime Function]
unsigned accumfractuhausa2 (unsigned short accum a) unsigned long accumfractuhauda2 (unsigned short accum a)	[Runtime Function] [Runtime Function]
unsigned long long accumfractuhauta2 (unsigned short accum a)	[Runtime Function]
signed charfractuhaqi (unsigned short accum a)	[Runtime Function]
shortfractuhahi (unsigned short accum a)	[Runtime Function]
intfractuhasi (unsigned short accum a)	[Runtime Function]
longfractuhadi (unsigned short accum a)	[Runtime Function]
long longfractuhati (unsigned short accum a)	[Runtime Function]
floatfractuhasf (unsigned short accum a)	[Runtime Function]
doublefractuhadf (unsigned short accum a)	[Runtime Function]
,	•
short fractfractusaqq (unsigned accum a)	[Runtime Function]
fractfractusahq (unsigned accum a)	[Runtime Function]
long fractfractusasq (unsigned accum a)	[Runtime Function]
long long fractfractusadq (unsigned accum a)	[Runtime Function]
short accumfractusaha (unsigned accum a)	[Runtime Function]
accumfractusasa (unsigned accum a)	[Runtime Function]
long accumfractusada (unsigned accum a)	[Runtime Function]
long long accumfractusata (unsigned accum a)	[Runtime Function]
unsigned short fractfractusauqq (unsigned accum a)	[Runtime Function]
unsigned fractfractusauhq (unsigned accum a)	[Runtime Function]
unsigned long fractfractusausq (unsigned accum a)	[Runtime Function]
unsigned long long fractfractusaudq (unsigned accum a)	[Runtime Function]
unsigned short accumfractusauha2 (unsigned accum a)	[Runtime Function]
unsigned long accumfractusauda2 (unsigned accum a)	[Runtime Function]
unsigned long long accumfractusauta2 (unsigned	[Runtime Function]
accum a)	
signed charfractusaqi (unsigned accum a)	[Runtime Function]
shortfractusahi (unsigned accum a)	[Runtime Function]
intfractusasi (unsigned accum a)	[Runtime Function]
longfractusadi (unsigned accum a)	[Runtime Function]
long longfractusati (unsigned accum a)	[Runtime Function]
floatfractusasf (unsigned accum a)	[Runtime Function]
doublefractusadf (unsigned accum a)	[Runtime Function]
<pre>short fractfractudaqq (unsigned long accum a)</pre>	[Runtime Function]

<pre>fractfractudahq (unsigned long accum a)</pre>	[Runtime Function]
long fractfractudasq (unsigned long accum a)	[Runtime Function]
long long fractfractudadq (unsigned long accum a)	[Runtime Function]
short accumfractudaha (unsigned long accum a)	[Runtime Function]
accumfractudasa (unsigned long accum a)	[Runtime Function]
long accumfractudada (unsigned long accum a)	[Runtime Function]
long long accumfractudata (unsigned long accum a)	[Runtime Function]
unsigned short fractfractudauqq (unsigned long	[Runtime Function]
accum a)	
unsigned fractfractudauhq (unsigned long accum a)	[Runtime Function]
unsigned long fractfractudausq (unsigned long accum	[Runtime Function]
a)	
unsigned long long fractfractudaudq (unsigned long accum a)	[Runtime Function]
unsigned short accumfractudauha2 (unsigned long	[Runtime Function]
accum a)	[realistific r directori]
unsigned accumfractudausa2 (unsigned long accum a)	[Runtime Function]
unsigned long long accumfractudauta2 (unsigned long	[Runtime Function]
accum a)	
signed charfractudaqi (unsigned long accum a)	[Runtime Function]
shortfractudahi (unsigned long accum a)	[Runtime Function]
intfractudasi (unsigned long accum a)	[Runtime Function]
longfractudadi (unsigned long accum a)	[Runtime Function]
long longfractudati (unsigned long accum a)	[Runtime Function]
floatfractudasf (unsigned long accum a)	[Runtime Function]
doublefractudadf (unsigned long accum a)	[Runtime Function]
short fractfractutaqq (unsigned long long accum a)	[Runtime Function]
fractfractutahq (unsigned long long accum a)	[Runtime Function]
<pre>long fractfractutasq (unsigned long long accum a)</pre>	[Runtime Function]
long long fractfractutadq (unsigned long long accum a)	[Runtime Function]
short accumfractutaha (unsigned long long accum a)	[Runtime Function]
accumfractutasa (unsigned long long accum a)	[Runtime Function]
long accumfractutada (unsigned long long accum a)	[Runtime Function]
long long accumfractutata (unsigned long long accum a)	[Runtime Function]
unsigned short fractfractutauqq (unsigned long long	[Runtime Function]
accum a)	
unsigned fractfractutauhq (unsigned long long accum a)	[Runtime Function]
unsigned long fractfractutausq (unsigned long long	[Runtime Function]
accum a)	
<pre>unsigned long long fractfractutaudq (unsigned long</pre>	[Runtime Function]
unsigned short accumfractutauha2 (unsigned long long accum a)	[Runtime Function]
unsigned accumfractutausa2 (unsigned long long accum a)	[Runtime Function]
unsigned long accumfractutauda2 (unsigned long long accum a)	[Runtime Function]

signed charfractutaqi (unsigned long long accum a)	[Runtime Function]
<b>-</b> \	[Runtime Function]
, – – ,	[Runtime Function]
longfractutadi (unsigned long long accum a)	[Runtime Function]
• • • • • • • • • • • • • • • • • • • •	[Runtime Function]
	[Runtime Function]
, – – ,	[Runtime Function]
` ,	[Runtime Function]
fractfractqihq (signed char a)	[Runtime Function]
long fractfractqisq (signed char a)	[Runtime Function]
long long fractfractqidq (signed char a)	[Runtime Function]
short accumfractqiha (signed char a)	[Runtime Function]
accumfractqisa (signed char a)	[Runtime Function]
long accumfractqida (signed char a)	[Runtime Function]
long long accumfractqita (signed char a)	[Runtime Function]
unsigned short fractfractqiuqq (signed char a)	[Runtime Function]
unsigned fractfractqiuhq (signed char a)	[Runtime Function]
unsigned long fractfractqiusq (signed char a)	[Runtime Function]
unsigned long long fractfractqiudq (signed char a)	[Runtime Function]
unsigned short accumfractqiuha (signed char a)	[Runtime Function]
unsigned accumfractqiusa (signed char a)	[Runtime Function]
unsigned long accumfractqiuda (signed char a)	[Runtime Function]
unsigned long long accumfractqiuta (signed char a)	[Runtime Function]
short fractfracthiqq (short a)	[Runtime Function]
<pre>fractfracthihq (short a)</pre>	[Runtime Function]
<pre>long fractfracthisq (short a)</pre>	[Runtime Function]
long long fractfracthidq (short a)	[Runtime Function]
short accumfracthiha (short a)	[Runtime Function]
accumfracthisa (short a)	[Runtime Function]
long accumfracthida (short a)	[Runtime Function]
long long accumfracthita (short a)	[Runtime Function]
	[Runtime Function]
<u> </u>	[Runtime Function]
unsigned long fractfracthiusq (short a)	[Runtime Function]
	[Runtime Function]
unsigned short accumfracthiuha (short a)	[Runtime Function]
` ,	[Runtime Function]
` , ,	[Runtime Function]
` ,	[Runtime Function]
\ /	[Runtime Function]
fractfractsihq (int a)	[Runtime Function]
<u> </u>	[Runtime Function]
<u> </u>	[Runtime Function]
· · ·	[Runtime Function]
accumfractsisa (int a)	[Runtime Function]
long accumfractsida (int a)	[Runtime Function]
long long accumfractsita (int a)	[Runtime Function]

	r=	
unsigned short fractfractsiuqq (int a)	[Runtime	Function]
unsigned fractfractsiuhq (int a)	[Runtime	Function]
unsigned long fractfractsiusq (int a)	Runtime	Function]
unsigned long long fractfractsiudq (int a)	Runtime	Function
unsigned short accumfractsiuha (int a)	L	Function]
` ,	L	Function]
unsigned accumfractsiusa (int a)	L	
unsigned long accumfractsiuda (int a)	L	Function]
unsigned long long accumfractsiuta (int a)	_	Function]
short fractfractdiqq (long a)	[Runtime	Function]
<pre>fractfractdihq (long a)</pre>	[Runtime	Function]
long fractfractdisq (long a)	Runtime	Function
long long fractfractdidq (long a)	L	Function
short accumfractdiha (long a)	L	Function]
· = /	L .	
accumfractdisa (long a)	L	Function]
long accumfractdida (long a)	L	Function]
long long accumfractdita (long a)	[Runtime	Function]
unsigned short fractfractdiuqq (long a)	[Runtime	Function]
unsigned fractfractdiuhq (long a)	Runtime	Function
unsigned long fractfractdiusq (long a)	L	Function
unsigned long long fractfractdiudq (long a)	L	Function]
<b>-</b> \ - /	L	
unsigned short accumfractdiuha (long a)	L	Function]
unsigned accumfractdiusa (long a)	L	Function]
unsigned long accumfractdiuda (long a)	-	Function]
unsigned long long accumfractdiuta (long a)	[Runtime	Function]
short fractfracttiqq (long long a)	[Runtime	Function]
fractfracttihq (long long a)	Runtime	Function
long fractfracttisq (long long a)	-	Function
long long fractfracttidq (long long a)	_	Function]
short accumfracttiha (long long a)	-	Function]
( 9 9 )	L	,
accumfracttisa (long long a)	L	Function]
long accumfracttida (long long a)	L	Function]
long long accumfracttita (long long a)	L	Function]
unsigned short fractfracttiuqq (long long a)	[Runtime	Function]
unsigned fractfracttiuhq (long long a)	[Runtime	Function]
unsigned long fractfracttiusq (long long a)	Runtime	Function
unsigned long long fractfracttiudq (long long a)	_	Function
unsigned short accumfracttiuha (long long a)	L	Function]
unsigned accumfracttiusa (long long a)	L	Function]
• • • • • • • • • • • • • • • • • • • •	_	•
unsigned long accumfracttiuda (long long a)	_	Function]
unsigned long long accumfracttiuta (long long a)	_	Function]
short fractfractsfqq (float a)	[Runtime	Function]
fractfractsfhq (float a)	[Runtime	Function]
long fractfractsfsq (float a)	-	Function
long long fractfractsfdq (float a)	L	Function]
short accumfractsfha (float a)	L	Function]
·	_	•
accumfractsfsa (float a)	_	Function]
long accumfractsfda (float a)	[Kuntime	Function]

signed non-fractionals, without saturation.

long long accumfractsfta (float a)	[Runtime Function]
unsigned short fractfractsfuqq (float a)	[Runtime Function]
unsigned fractfractsfuhq (float a)	[Runtime Function]
unsigned long fractfractsfusq (float a)	[Runtime Function]
unsigned long long fractfractsfudq (float a)	[Runtime Function]
unsigned short accumfractsfuha (float a)	[Runtime Function]
unsigned accumfractsfusa (float a)	[Runtime Function]
unsigned long accumfractsfuda (float a)	[Runtime Function]
unsigned long long accumfractsfuta (float a)	[Runtime Function]
short fractfractdfqq (double a)	[Runtime Function]
fractfractdfhq (double a)	[Runtime Function]
long fractfractdfsq (double a)	[Runtime Function]
long long fractfractdfdq (double a)	[Runtime Function]
short accumfractdfha (double a)	[Runtime Function]
accumfractdfsa (double a)	[Runtime Function]
long accumfractdfda (double a)	[Runtime Function]
long long accumfractdfta (double a)	[Runtime Function]
unsigned short fractfractdfuqq (double a)	[Runtime Function]
unsigned fractfractdfuhq (double a)	[Runtime Function]
unsigned long fractfractdfusq (double a)	[Runtime Function]
unsigned long long fractfractdfudq (double a)	[Runtime Function]
unsigned short accumfractdfuha (double a)	[Runtime Function]
unsigned accumfractdfusa (double a)	[Runtime Function]
unsigned long accumfractdfuda (double a)	[Runtime Function]
unsigned long long accumfractdfuta (double a)	[Runtime Function]
These functions convert from fractional and signed non-fractionals	s to fractionals and

<pre>fractsatfractqqhq2 (short fract a)</pre>	[Runtime Function]
<pre>long fractsatfractqqsq2 (short fract a)</pre>	[Runtime Function]
<pre>long long fractsatfractqqdq2 (short fract a)</pre>	[Runtime Function]
short accumsatfractqqha (short fract a)	[Runtime Function]
accumsatfractqqsa (short fract a)	[Runtime Function]
long accumsatfractqqda (short fract a)	[Runtime Function]
long long accumsatfractqqta (short fract a)	[Runtime Function]
unsigned short fractsatfractqquqq (short fract a)	[Runtime Function]
unsigned fractsatfractqquhq (short fract a)	[Runtime Function]
unsigned long fractsatfractqqusq (short fract a)	[Runtime Function]
unsigned long long fractsatfractqqudq (short fract	[Runtime Function]
a)	
unsigned short accumsatfractqquha (short fract a)	[Runtime Function]
unsigned accumsatfractqqusa (short fract a)	[Runtime Function]
unsigned long accumsatfractqquda (short fract a)	[Runtime Function]
unsigned long long accumsatfractqquta (short fract	[Runtime Function]
a)	
short fractsatfracthqqq2 (fract a)	[Runtime Function]
<pre>long fractsatfracthqsq2 (fract a)</pre>	[Runtime Function]

```
long long fract __satfracthqdq2 (fract a)
                                                              [Runtime Function]
short accum __satfracthqha (fract a)
                                                              [Runtime Function]
accum __satfracthqsa (fract a)
                                                              [Runtime Function]
long accum __satfracthqda (fract a)
                                                              [Runtime Function]
long long accum __satfracthqta (fract a)
                                                              [Runtime Function]
unsigned short fract __satfracthquqq (fract a)
                                                              [Runtime Function]
unsigned fract __satfracthquhq (fract a)
                                                              [Runtime Function]
unsigned long fract __satfracthqusq (fract a)
                                                              [Runtime Function]
unsigned long long fract __satfracthqudq (fract a)
                                                              [Runtime Function]
unsigned short accum __satfracthquha (fract a)
                                                              [Runtime Function]
                                                              [Runtime Function]
unsigned accum __satfracthqusa (fract a)
unsigned long accum __satfracthquda (fract a)
                                                              [Runtime Function]
unsigned long long accum __satfracthquta (fract a)
                                                              [Runtime Function]
short fract __satfractsqqq2 (long fract a)
                                                              [Runtime Function]
fract __satfractsqhq2 (long fract a)
                                                              [Runtime Function]
long long fract __satfractsqdq2 (long fract a)
                                                              [Runtime Function]
short accum __satfractsqha (long fract a)
                                                              [Runtime Function]
accum __satfractsqsa (long fract a)
                                                              [Runtime Function]
long accum __satfractsqda (long fract a)
                                                              [Runtime Function]
long long accum __satfractsqta (long fract a)
                                                              [Runtime Function]
unsigned short fract __satfractsquqq (long fract a)
                                                              [Runtime Function]
unsigned fract __satfractsquhq (long fract a)
                                                              [Runtime Function]
unsigned long fract __satfractsqusq (long fract a)
                                                              [Runtime Function]
unsigned long long fract __satfractsqudq (long fract a)
                                                              [Runtime Function]
unsigned short accum __satfractsquha (long fract a)
                                                              [Runtime Function]
unsigned accum __satfractsqusa (long fract a)
                                                              [Runtime Function]
unsigned long accum __satfractsquda (long fract a)
                                                              [Runtime Function]
unsigned long long accum __satfractsquta (long fract a)
                                                              [Runtime Function]
short fract __satfractdqqq2 (long long fract a)
                                                              [Runtime Function]
fract __satfractdqhq2 (long long fract a)
                                                              [Runtime Function]
long fract __satfractdqsq2 (long long fract a)
                                                              [Runtime Function]
short accum __satfractdqha (long long fract a)
                                                              [Runtime Function]
accum __satfractdqsa (long long fract a)
                                                              [Runtime Function]
long accum __satfractdqda (long long fract a)
                                                              [Runtime Function]
long long accum __satfractdqta (long long fract a)
                                                              [Runtime Function]
unsigned short fract __satfractdquqq (long long fract a)
                                                              [Runtime Function]
unsigned fract __satfractdquhq (long long fract a)
                                                              [Runtime Function]
unsigned long fract __satfractdqusq (long long fract a)
                                                              [Runtime Function]
unsigned long long fract __satfractdqudq (long long
                                                              [Runtime Function]
         fract a)
unsigned short accum __satfractdquha (long long fract a)
                                                              [Runtime Function]
unsigned accum __satfractdqusa (long long fract a)
                                                              [Runtime Function]
unsigned long accum __satfractdquda (long long fract a)
                                                              [Runtime Function]
unsigned long long accum __satfractdquta (long long
                                                              [Runtime Function]
         fract a)
short fract __satfracthaqq (short accum a)
                                                              [Runtime Function]
fract __satfracthahq (short accum a)
                                                              [Runtime Function]
```

```
long fract __satfracthasq (short accum a)
                                                             [Runtime Function]
long long fract __satfracthadq (short accum a)
                                                             [Runtime Function]
accum __satfracthasa2 (short accum a)
                                                             [Runtime Function]
long accum __satfracthada2 (short accum a)
                                                             [Runtime Function]
long long accum __satfracthata2 (short accum a)
                                                             [Runtime Function]
unsigned short fract __satfracthauqq (short accum a)
                                                             [Runtime Function]
unsigned fract __satfracthauhg (short accum a)
                                                             [Runtime Function]
unsigned long fract __satfracthausq (short accum a)
                                                             [Runtime Function]
unsigned long long fract __satfracthaudg (short accum
                                                             [Runtime Function]
unsigned short accum __satfracthauha (short accum a)
                                                             [Runtime Function]
unsigned accum __satfracthausa (short accum a)
                                                             [Runtime Function]
unsigned long accum __satfracthauda (short accum a)
                                                             [Runtime Function]
unsigned long long accum __satfracthauta (short accum
                                                             [Runtime Function]
short fract __satfractsaqq (accum a)
                                                             [Runtime Function]
fract __satfractsahq (accum a)
                                                             [Runtime Function]
long fract __satfractsasq (accum a)
                                                             [Runtime Function]
long long fract __satfractsadq (accum a)
                                                             [Runtime Function]
short accum __satfractsaha2 (accum a)
                                                             [Runtime Function]
long accum __satfractsada2 (accum a)
                                                             [Runtime Function]
long long accum __satfractsata2 (accum a)
                                                             [Runtime Function]
unsigned short fract __satfractsauqq (accum a)
                                                             [Runtime Function]
unsigned fract __satfractsauhq (accum a)
                                                             [Runtime Function]
unsigned long fract __satfractsausq (accum a)
                                                             [Runtime Function]
unsigned long long fract __satfractsaudq (accum a)
                                                             [Runtime Function]
unsigned short accum __satfractsauha (accum a)
                                                             [Runtime Function]
unsigned accum __satfractsausa (accum a)
                                                             [Runtime Function]
unsigned long accum __satfractsauda (accum a)
                                                             [Runtime Function]
unsigned long long accum __satfractsauta (accum a)
                                                             [Runtime Function]
short fract __satfractdaqq (long accum a)
                                                             [Runtime Function]
fract __satfractdahq (long accum a)
                                                             [Runtime Function]
long fract __satfractdasq (long accum a)
                                                             [Runtime Function]
long long fract __satfractdadq (long accum a)
                                                             [Runtime Function]
short accum __satfractdaha2 (long accum a)
                                                             [Runtime Function]
accum __satfractdasa2 (long accum a)
                                                             [Runtime Function]
long long accum __satfractdata2 (long accum a)
                                                             [Runtime Function]
unsigned short fract __satfractdauqq (long accum a)
                                                             [Runtime Function]
unsigned fract __satfractdauhq (long accum a)
                                                             [Runtime Function]
unsigned long fract __satfractdausq (long accum a)
                                                             [Runtime Function]
unsigned long long fract __satfractdaudq (long accum
                                                             [Runtime Function]
unsigned short accum __satfractdauha (long accum a)
                                                             [Runtime Function]
unsigned accum __satfractdausa (long accum a)
                                                             [Runtime Function]
unsigned long accum __satfractdauda (long accum a)
                                                             [Runtime Function]
unsigned long long accum __satfractdauta (long accum
                                                             [Runtime Function]
         a)
```

short fractsatfracttaqq (long long accum a) fractsatfracttahq (long long accum a) long fractsatfracttasq (long long accum a) long long fractsatfracttadq (long long accum a) short accumsatfracttaha2 (long long accum a) accumsatfracttasa2 (long long accum a) long accumsatfracttada2 (long long accum a) unsigned short fractsatfracttauqq (long long accum a)	[Runtime Function]
unsigned fractsatfracttauhq (long long accum a) unsigned long fractsatfracttausq (long long accum a) unsigned long long fractsatfracttaudq (long long accum a)	[Runtime Function] [Runtime Function] [Runtime Function]
unsigned short accumsatfracttauha (long long accum a)	[Runtime Function]
unsigned accumsatfracttausa (long long accum a) unsigned long accumsatfracttauda (long long accum a) unsigned long long accumsatfracttauta (long long accum a)	[Runtime Function] [Runtime Function] [Runtime Function]
short fractsatfractuqqqq (unsigned short fract a) fractsatfractuqqqq (unsigned short fract a) long fractsatfractuqqqq (unsigned short fract a) long long fractsatfractuqqdq (unsigned short fract a) short accumsatfractuqqha (unsigned short fract a) accumsatfractuqqsa (unsigned short fract a) long accumsatfractuqqda (unsigned short fract a) long long accumsatfractuqqta (unsigned short fract a) unsigned fractsatfractuqquhq2 (unsigned short fract a) unsigned long fractsatfractuqqusq2 (unsigned short fract a)	[Runtime Function]
unsigned long long fractsatfractuqqudq2 (unsigned short fract a)	[Runtime Function]
unsigned short accumsatfractuqquha (unsigned short fract a)	[Runtime Function]
unsigned accumsatfractuqqusa (unsigned short fract a) unsigned long accumsatfractuqquda (unsigned short fract a)	[Runtime Function] [Runtime Function]
unsigned long long accumsatfractuqquta (unsigned short fract a)	[Runtime Function]
short fractsatfractuhqqq (unsigned fract a) fractsatfractuhqhq (unsigned fract a) long fractsatfractuhqsq (unsigned fract a) long long fractsatfractuhqdq (unsigned fract a) short accumsatfractuhqha (unsigned fract a)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
accumsatfractuhqsa (unsigned fract a) long accumsatfractuhqda (unsigned fract a) long long accumsatfractuhqta (unsigned fract a)	[Runtime Function] [Runtime Function] [Runtime Function]

<pre>unsigned short fractsatfractuhquqq2 (unsigned fract a)</pre>	[Runtime Function]
unsigned long fractsatfractuhqusq2 (unsigned fract a)	[Runtime Function]
unsigned long long fractsatfractuhqudq2 (unsigned fract a)	[Runtime Function]
unsigned short accumsatfractuhquha (unsigned fract a)	[Runtime Function]
unsigned accumsatfractuhqusa (unsigned fract a)	[Runtime Function]
unsigned long accumsatfractuhquda (unsigned fract a) unsigned long long accumsatfractuhquta (unsigned fract a)	[Runtime Function] [Runtime Function]
short fractsatfractusqqq (unsigned long fract a)	[Runtime Function]
fractsatfractusqhq (unsigned long fract a)	[Runtime Function]
long fractsatfractusqsq (unsigned long fract a)	[Runtime Function]
long long fractsatfractusqdq (unsigned long fract a)	[Runtime Function]
short accumsatfractusqha (unsigned long fract a)	[Runtime Function]
accumsatfractusqsa (unsigned long fract a)	[Runtime Function]
long accumsatfractusqda (unsigned long fract a)	[Runtime Function]
long long accumsatfractusqta (unsigned long fract a)	[Runtime Function]
unsigned short fractsatfractusquqq2 (unsigned long	[Runtime Function]
fract a)	[realitime ranction]
unsigned fractsatfractusquhq2 (unsigned long fract a)	[Runtime Function]
unsigned long long fractsatfractusqudq2 (unsigned	[Runtime Function]
long fract a)	
unsigned short accumsatfractusquha (unsigned long fract a)	[Runtime Function]
unsigned accumsatfractusqusa (unsigned long fract a)	[Runtime Function]
unsigned long accumsatfractusquda (unsigned long fract a)	[Runtime Function]
unsigned long long accumsatfractusquta (unsigned long fract a)	[Runtime Function]
short fractsatfractudqqq (unsigned long long fract a)	[Runtime Function]
fractsatfractudqhq (unsigned long long fract a)	[Runtime Function]
long fractsatfractudqsq (unsigned long long fract a)	[Runtime Function]
<pre>long long fractsatfractudqdq (unsigned long long fract a)</pre>	[Runtime Function]
short accumsatfractudqha (unsigned long long fract a)	[Runtime Function]
accumsatfractudqsa (unsigned long long fract a)	[Runtime Function]
long accumsatfractudqda (unsigned long long fract a)	[Runtime Function]
long long accumsatfractudqta (unsigned long long fract a)	[Runtime Function]
unsigned short fractsatfractudquqq2 (unsigned long long fract a)	[Runtime Function]
unsigned fractsatfractudquhq2 (unsigned long long fract a)	[Runtime Function]
,	

<pre>unsigned long fractsatfractudqusq2 (unsigned long</pre>	[Runtime Function]
unsigned short accumsatfractudquha (unsigned long long fract a)	[Runtime Function]
<pre>unsigned accumsatfractudqusa (unsigned long long fract a)</pre>	[Runtime Function]
unsigned long accumsatfractudquda (unsigned long long fract a)	[Runtime Function]
unsigned long long accumsatfractudquta (unsigned long long fract a)	[Runtime Function]
short fractsatfractuhaqq (unsigned short accum a)	[Runtime Function]
fractsatfractuhahq (unsigned short accum a)	[Runtime Function]
long fractsatfractuhasq (unsigned short accum a)	[Runtime Function]
long long fractsatfractuhadq (unsigned short accum a)	[Runtime Function]
short accumsatfractuhaha (unsigned short accum a)	[Runtime Function]
accumsatfractuhasa (unsigned short accum a)	[Runtime Function]
long accumsatfractuhada (unsigned short accum a)	[Runtime Function]
long long accumsatfractuhata (unsigned short accum a)	[Runtime Function]
unsigned short fractsatfractuhauqq (unsigned short accum a)	[Runtime Function]
<pre>unsigned fractsatfractuhauhq (unsigned short accum a)</pre>	[Runtime Function]
unsigned long fractsatfractuhausq (unsigned short accum a)	[Runtime Function]
unsigned long long fractsatfractuhaudq (unsigned short accum a)	[Runtime Function]
unsigned accumsatfractuhausa2 (unsigned short accum a)	[Runtime Function]
unsigned long accumsatfractuhauda2 (unsigned short accum a)	[Runtime Function]
unsigned long long accumsatfractuhauta2 (unsigned short accum a)	[Runtime Function]
<pre>short fractsatfractusaqq (unsigned accum a)</pre>	[Runtime Function]
<pre>fractsatfractusahq (unsigned accum a)</pre>	[Runtime Function]
<pre>long fractsatfractusasq (unsigned accum a)</pre>	[Runtime Function]
<pre>long long fractsatfractusadq (unsigned accum a)</pre>	[Runtime Function]
short accumsatfractusaha (unsigned accum a)	[Runtime Function]
accumsatfractusasa (unsigned accum a)	[Runtime Function]
long accumsatfractusada (unsigned accum a)	[Runtime Function]
long long accumsatfractusata (unsigned accum a)	[Runtime Function]
<pre>unsigned short fractsatfractusauqq (unsigned accum a)</pre>	[Runtime Function]
<pre>unsigned fractsatfractusauhq (unsigned accum a)</pre>	[Runtime Function]
<pre>unsigned long fractsatfractusausq (unsigned accum a)</pre>	[Runtime Function]
unsigned long long fractsatfractusaudq (unsigned accum a)	[Runtime Function]

unsigned short accumsatfractusauha2 (unsigned accum a)	[Runtime Function]
unsigned long accumsatfractusauda2 (unsigned accum a)	[Runtime Function]
unsigned long long accumsatfractusauta2 (unsigned accum a)	[Runtime Function]
short fractsatfractudaqq (unsigned long accum a) fractsatfractudahq (unsigned long accum a) long fractsatfractudasq (unsigned long accum a) long long fractsatfractudadq (unsigned long accum a) short accumsatfractudaha (unsigned long accum a) accumsatfractudasa (unsigned long accum a)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
long accumsatfractudada (unsigned long accum a)	[Runtime Function]
long long accumsatfractudata (unsigned long accum a) unsigned short fractsatfractudauqq (unsigned long accum a)	[Runtime Function] [Runtime Function]
unsigned fractsatfractudauhq (unsigned long accum a)	[Runtime Function]
unsigned long fractsatfractudausq (unsigned long accum a)	[Runtime Function]
unsigned long long fractsatfractudaudq (unsigned long accum a)	[Runtime Function]
unsigned short accumsatfractudauha2 (unsigned long accum a)	[Runtime Function]
<pre>unsigned accumsatfractudausa2 (unsigned long accum a)</pre>	[Runtime Function]
unsigned long long accumsatfractudauta2 (unsigned long accum a)	[Runtime Function]
short fractsatfractutaqq (unsigned long long accum a)	[Runtime Function]
<pre>fractsatfractutahq (unsigned long long accum a) long fractsatfractutasq (unsigned long long accum a)</pre>	[Runtime Function] [Runtime Function]
long long fractsatfractutadq (unsigned long long accum a)  accum a)	[Runtime Function]
short accumsatfractutaha (unsigned long long accum a) accumsatfractutasa (unsigned long long accum a) long accumsatfractutada (unsigned long long accum a) long long accumsatfractutata (unsigned long long accum a)	[Runtime Function] [Runtime Function] [Runtime Function] [Runtime Function]
unsigned short fractsatfractutauqq (unsigned long long accum a)	[Runtime Function]
unsigned fractsatfractutauhq (unsigned long long accum a)	[Runtime Function]
unsigned long fractsatfractutausq (unsigned long long accum a)	[Runtime Function]
unsigned long long fractsatfractutaudq (unsigned long long accum a)	[Runtime Function]
unsigned short accumsatfractutauha2 (unsigned long long accum a)	[Runtime Function]

```
unsigned accum __satfractutausa2 (unsigned long long
                                                             [Runtime Function]
         accum a)
unsigned long accum __satfractutauda2 (unsigned long
                                                             [Runtime Function]
         long accum a)
short fract __satfractqiqq (signed char a)
                                                              [Runtime Function]
fract __satfractqihq (signed char a)
                                                              [Runtime Function]
long fract __satfractgisg (signed char a)
                                                              [Runtime Function]
long long fract __satfractqidq (signed char a)
                                                              [Runtime Function]
short accum __satfractqiha (signed char a)
                                                              [Runtime Function]
accum __satfractqisa (signed char a)
                                                              [Runtime Function]
long accum __satfractqida (signed char a)
                                                              [Runtime Function]
long long accum __satfractqita (signed char a)
                                                              [Runtime Function]
unsigned short fract __satfractqiuqq (signed char a)
                                                              [Runtime Function]
unsigned fract __satfractqiuhq (signed char a)
                                                              [Runtime Function]
unsigned long fract __satfractgiusg (signed char a)
                                                              [Runtime Function]
unsigned long long fract __satfractqiudq (signed char
                                                              [Runtime Function]
unsigned short accum __satfractqiuha (signed char a)
                                                              [Runtime Function]
unsigned accum __satfractqiusa (signed char a)
                                                              [Runtime Function]
unsigned long accum __satfractqiuda (signed char a)
                                                              [Runtime Function]
unsigned long long accum __satfractqiuta (signed char
                                                              [Runtime Function]
         a)
short fract __satfracthiqq (short a)
                                                              [Runtime Function]
fract __satfracthing (short a)
                                                              [Runtime Function]
long fract __satfracthisq (short a)
                                                              [Runtime Function]
long long fract __satfracthidq (short a)
                                                              [Runtime Function]
short accum __satfracthiha (short a)
                                                              [Runtime Function]
accum __satfracthisa (short a)
                                                              [Runtime Function]
long accum __satfracthida (short a)
                                                              [Runtime Function]
long long accum __satfracthita (short a)
                                                              [Runtime Function]
unsigned short fract __satfracthiuqq (short a)
                                                              [Runtime Function]
unsigned fract __satfracthiuhq (short a)
                                                              [Runtime Function]
unsigned long fract __satfracthiusq (short a)
                                                              [Runtime Function]
unsigned long long fract __satfracthiudg (short a)
                                                              [Runtime Function]
unsigned short accum __satfracthiuha (short a)
                                                              [Runtime Function]
unsigned accum __satfracthiusa (short a)
                                                              [Runtime Function]
unsigned long accum __satfracthiuda (short a)
                                                              [Runtime Function]
unsigned long long accum __satfracthiuta (short a)
                                                              [Runtime Function]
short fract __satfractsiqq (int a)
                                                              [Runtime Function]
fract __satfractsihq (int a)
                                                              [Runtime Function]
long fract __satfractsisq (int a)
                                                              [Runtime Function]
long long fract __satfractsidq (int a)
                                                              [Runtime Function]
short accum __satfractsiha (int a)
                                                              [Runtime Function]
accum __satfractsisa (int a)
                                                              [Runtime Function]
long accum __satfractsida (int a)
                                                              [Runtime Function]
long long accum __satfractsita (int a)
                                                              [Runtime Function]
unsigned short fract __satfractsiugg (int a)
                                                              [Runtime Function]
```

unsigned fractsatfractsiuhq (int a)	[Runtime	Function]
unsigned long fractsatfractsiusq (int a)	[Runtime	Function]
unsigned long long fractsatfractsiudq (int a)	[Runtime	Function]
unsigned short accumsatfractsiuha (int a)	[Runtime	Function]
unsigned accumsatfractsiusa (int a)	Runtime	Function]
unsigned long accumsatfractsiuda (int a)	Runtime	Function
unsigned long long accumsatfractsiuta (int a)	Runtime	Function
short fractsatfractdiqq (long a)	Runtime	Function
fractsatfractdihq (long a)	-	Function
long fractsatfractdisq (long a)	Runtime	Function
long long fractsatfractdidq (long a)	Runtime	Function
short accumsatfractdiha (long a)	_	Function
accumsatfractdisa (long a)	L	Function
long accumsatfractdida (long a)	_	Function
long long accumsatfractdita (long a)	L	Function
unsigned short fractsatfractdiuqq (long a)	L	Function
unsigned fractsatfractdiuhq (long a)	L	Function
unsigned long fractsatfractdiusq (long a)	L	Function
unsigned long long fractsatfractdiudq (long a)	_	Function]
unsigned short accumsatfractdiuha (long a)	_	Function
unsigned accumsatfractdiusa (long a)	L	Function]
unsigned long accumsatfractdiuda (long a)	_	Function]
unsigned long long accumsatfractdiuta (long a)	L	Function]
short fractsatfracttiqq (long long a)	L	Function]
fractsatfracttihq (long long a)	_	Function]
long fractsatfracttisq (long long a)	L	Function]
long long fractsatfracttidq (long long a)	_	Function]
short accumsatfracttiha (long long a)	_	Function]
accumsatfracttisa (long long a)	_	Function]
long accumsatfracttida (long long a)	L	Function]
long long accumsatfracttita (long long a)	L	Function]
unsigned short fractsatfracttiuqq (long long a)	_	Function]
unsigned fractsatfracttiuhq (long long a)	L	Function]
unsigned long fractsatfracttiusq (long long a)	_	Function]
unsigned long long fractsatfracttiudq (long long a)	-	Function]
unsigned short accumsatfracttiuha (long long a)	-	Function]
unsigned accumsatfracttiusa (long long a)	_	Function]
unsigned long accumsatfracttiuda (long long a)	-	Function]
unsigned long long accumsatfracttiuta (long long a)	-	Function]
short fractsatfractsfqq (float a)	-	Function]
fractsatfractsfhq (float a)	L	Function]
long fractsatfractsfsq (float a)	L	Function]
long long fractsatfractsfdq (float a)	-	Function]
short accumsatfractsfha (float a)	L	Function]
accumsatfractsfna (float a)	L	Function]
long accumsatfractsfda (float a)	_	Function]
·	L	Function]
long long accumsatfractsfta (float a)	լուսուսու	r uncomil

saturation.

```
unsigned short fract __satfractsfugg (float a)
                                                              [Runtime Function]
unsigned fract __satfractsfuhg (float a)
                                                              [Runtime Function]
unsigned long fract __satfractsfusq (float a)
                                                              [Runtime Function]
unsigned long long fract __satfractsfudg (float a)
                                                              [Runtime Function]
unsigned short accum __satfractsfuha (float a)
                                                              [Runtime Function]
unsigned accum __satfractsfusa (float a)
                                                              [Runtime Function]
unsigned long accum __satfractsfuda (float a)
                                                              [Runtime Function]
unsigned long long accum __satfractsfuta (float a)
                                                              [Runtime Function]
                                                              [Runtime Function]
short fract __satfractdfqq (double a)
fract __satfractdfhq (double a)
                                                              [Runtime Function]
                                                              [Runtime Function]
long fract __satfractdfsq (double a)
long long fract __satfractdfdq (double a)
                                                              [Runtime Function]
short accum __satfractdfha (double a)
                                                              [Runtime Function]
accum __satfractdfsa (double a)
                                                              [Runtime Function]
long accum __satfractdfda (double a)
                                                              [Runtime Function]
long long accum __satfractdfta (double a)
                                                              [Runtime Function]
unsigned short fract __satfractdfugg (double a)
                                                              [Runtime Function]
unsigned fract __satfractdfuhq (double a)
                                                              [Runtime Function]
unsigned long fract __satfractdfusq (double a)
                                                              [Runtime Function]
unsigned long long fract __satfractdfudq (double a)
                                                              [Runtime Function]
unsigned short accum __satfractdfuha (double a)
                                                              [Runtime Function]
unsigned accum __satfractdfusa (double a)
                                                              [Runtime Function]
unsigned long accum __satfractdfuda (double a)
                                                              [Runtime Function]
unsigned long long accum __satfractdfuta (double a)
                                                              [Runtime Function]
     The functions convert from fractional and signed non-fractionals to fractionals, with
```

unsigned char \_\_fractunsqqqi (short fract a) [Runtime Function] unsigned short \_\_fractunsqqhi (short fract a) [Runtime Function] unsigned int \_\_fractunsqqsi (short fract a) [Runtime Function] unsigned long \_\_fractunsqqdi (short fract a) [Runtime Function] unsigned long long \_\_fractunsqqti (short fract a) [Runtime Function] unsigned char \_\_fractunshqqi (fract a) [Runtime Function] unsigned short \_\_fractunshqhi (fract a) [Runtime Function] unsigned int \_\_fractunshqsi (fract a) [Runtime Function] unsigned long \_\_fractunshqdi (fract a) [Runtime Function] unsigned long long \_\_fractunshqti (fract a) [Runtime Function] unsigned char \_\_fractunssqqi (long fract a) [Runtime Function] unsigned short \_\_fractunssqhi (long fract a) [Runtime Function] unsigned int \_\_fractunssqsi (long fract a) [Runtime Function] unsigned long \_\_fractunssqdi (long fract a) [Runtime Function] unsigned long long \_\_fractunssqti (long fract a) [Runtime Function] unsigned char \_\_fractunsdqqi (long long fract a) [Runtime Function] unsigned short \_\_fractunsdqhi (long long fract a) [Runtime Function] unsigned int \_\_fractunsdqsi (long long fract a) [Runtime Function] unsigned long \_\_fractunsdqdi (long long fract a) [Runtime Function] unsigned long long \_\_fractunsdqti (long long fract a) [Runtime Function]

```
unsigned char __fractunshaqi (short accum a)
                                                              [Runtime Function]
unsigned short __fractunshahi (short accum a)
                                                              [Runtime Function]
unsigned int __fractunshasi (short accum a)
                                                              [Runtime Function]
unsigned long __fractunshadi (short accum a)
                                                              [Runtime Function]
unsigned long long __fractunshati (short accum a)
                                                              [Runtime Function]
unsigned char __fractunssaqi (accum a)
                                                              [Runtime Function]
unsigned short __fractunssahi (accum a)
                                                              [Runtime Function]
unsigned int __fractunssasi (accum a)
                                                              [Runtime Function]
unsigned long __fractunssadi (accum a)
                                                              [Runtime Function]
unsigned long long __fractunssati (accum a)
                                                              [Runtime Function]
unsigned char __fractunsdaqi (long accum a)
                                                              [Runtime Function]
unsigned short __fractunsdahi (long accum a)
                                                              [Runtime Function]
unsigned int __fractunsdasi (long accum a)
                                                              [Runtime Function]
unsigned long __fractunsdadi (long accum a)
                                                              [Runtime Function]
unsigned long long __fractunsdati (long accum a)
                                                              [Runtime Function]
unsigned char __fractunstaqi (long long accum a)
                                                              [Runtime Function]
unsigned short __fractunstahi (long long accum a)
                                                              [Runtime Function]
unsigned int __fractunstasi (long long accum a)
                                                              [Runtime Function]
unsigned long __fractunstadi (long long accum a)
                                                              [Runtime Function]
unsigned long long __fractunstati (long long accum a)
                                                              [Runtime Function]
unsigned char __fractunsuqqqi (unsigned short fract a)
                                                              [Runtime Function]
unsigned short __fractunsuqqhi (unsigned short fract a)
                                                              [Runtime Function]
unsigned int __fractunsuqqsi (unsigned short fract a)
                                                              [Runtime Function]
unsigned long __fractunsuqqdi (unsigned short fract a)
                                                              [Runtime Function]
unsigned long long __fractunsuggti (unsigned short fract
                                                              [Runtime Function]
unsigned char __fractunsuhqqi (unsigned fract a)
                                                              [Runtime Function]
unsigned short __fractunsuhqhi (unsigned fract a)
                                                              [Runtime Function]
unsigned int __fractunsuhqsi (unsigned fract a)
                                                              [Runtime Function]
unsigned long __fractunsuhqdi (unsigned fract a)
                                                              [Runtime Function]
unsigned long long __fractunsuhqti (unsigned fract a)
                                                              [Runtime Function]
unsigned char __fractunsusqqi (unsigned long fract a)
                                                              [Runtime Function]
unsigned short __fractunsusqhi (unsigned long fract a)
                                                              [Runtime Function]
unsigned int __fractunsusqsi (unsigned long fract a)
                                                              [Runtime Function]
unsigned long __fractunsusqdi (unsigned long fract a)
                                                              [Runtime Function]
unsigned long long __fractunsusqti (unsigned long fract
                                                              [Runtime Function]
unsigned char __fractunsudqqi (unsigned long long fract a)
                                                              [Runtime Function]
unsigned short __fractunsudqhi (unsigned long long fract
                                                              [Runtime Function]
         a)
unsigned int __fractunsudqsi (unsigned long long fract a)
                                                              [Runtime Function]
unsigned long __fractunsudqdi (unsigned long long fract a)
                                                              [Runtime Function]
unsigned long long __fractunsudgti (unsigned long long
                                                              [Runtime Function]
         fract a)
unsigned char __fractunsuhaqi (unsigned short accum a)
                                                              [Runtime Function]
unsigned short __fractunsuhahi (unsigned short accum a)
                                                              [Runtime Function]
unsigned int __fractunsuhasi (unsigned short accum a)
                                                              [Runtime Function]
```

unsigned longfractunsuhadi (unsigned short accum a)	[Runtime Function]
unsigned long longfractunsuhati (unsigned short	[Runtime Function]
accum a)	
unsigned charfractunsusaqi (unsigned accum a)	[Runtime Function]
unsigned shortfractunsusahi (unsigned accum a)	[Runtime Function]
unsigned intfractunsusasi (unsigned accum a)	[Runtime Function]
unsigned longfractunsusadi (unsigned accum a)	[Runtime Function]
unsigned long longfractunsusati (unsigned accum a)	[Runtime Function]
unsigned charfractunsudaqi (unsigned long accum a)	[Runtime Function]
unsigned shortfractunsudahi (unsigned long accum a)	[Runtime Function]
unsigned intfractunsudasi (unsigned long accum a)	[Runtime Function]
unsigned longfractunsudadi (unsigned long accum a)	[Runtime Function]
unsigned long longfractunsudati (unsigned long accum a)	[Runtime Function]
unsigned charfractunsutaqi (unsigned long long accum	[Runtime Function]
a)	į j
unsigned shortfractunsutahi (unsigned long long accum	[Runtime Function]
a)	[D .: D .: ]
unsigned intfractursutasi (unsigned long long accum a)	[Runtime Function]
unsigned longfractunsutadi (unsigned long long accum a)	[Runtime Function]
unsigned long longfractunsutati (unsigned long long	[Runtime Function]
accum a)	į j
<pre>short fractfractunsqiqq (unsigned char a)</pre>	[Runtime Function]
<pre>fractfractunsqihq (unsigned char a)</pre>	[Runtime Function]
long fractfractunsqisq (unsigned char a)	[Runtime Function]
long long fractfractunsqidq (unsigned char a)	[Runtime Function]
short accumfractunsqiha (unsigned char a)	[Runtime Function]
accumfractunsqisa (unsigned char a)	[Runtime Function]
long accumfractunsqida (unsigned char a)	[Runtime Function]
long long accumfractunsqita (unsigned char a)	[Runtime Function]
unsigned short fractfractunsqiuqq (unsigned char a)	[Runtime Function]
unsigned fractfractunsqiuhq (unsigned char a)	[Runtime Function]
unsigned long fractfractunsquusq (unsigned char a)	[Runtime Function]
unsigned long long fractfractunsqiudq (unsigned char a)	[Runtime Function]
unsigned short accumfractunsqiuha (unsigned char a)	[Runtime Function]
unsigned accumfractunsqiusa (unsigned char a)	[Runtime Function]
unsigned long accumfractunsqiuda (unsigned char a)	[Runtime Function]
unsigned long long accumfractunsqiuta (unsigned	[Runtime Function]
char a)	(m
short fractfractunshiqq (unsigned short a)	[Runtime Function]
fractfractunshihq (unsigned short a)	[Runtime Function]
long fractfractunshisq (unsigned short a)	[Runtime Function]
long long fractfractunshidq (unsigned short a)	[Runtime Function]
short accumfractunshiha (unsigned short a)	[Runtime Function]
accumfractunshisa (unsigned short a)	[Runtime Function]

long accumfractunshida (unsigned short a)	[Runtime Function]
long long accumfractunshita (unsigned short a)	[Runtime Function]
unsigned short fractfractunshiuqq (unsigned short a)	[Runtime Function]
unsigned fractfractunshiuhq (unsigned short a)	[Runtime Function]
unsigned long fractfractunshiusq (unsigned short a)	[Runtime Function]
unsigned long long fractfractunshiudq (unsigned	[Runtime Function]
short a)	
unsigned short accumfractunshiuha (unsigned short a)	[Runtime Function]
unsigned accumfractunshiusa (unsigned short a)	[Runtime Function]
unsigned long accumfractunshiuda (unsigned short a)	[Runtime Function]
unsigned long long accumfractunshiuta (unsigned	[Runtime Function]
short a)	
<pre>short fractfractunssiqq (unsigned int a)</pre>	[Runtime Function]
<pre>fractfractunssihq (unsigned int a)</pre>	[Runtime Function]
<pre>long fractfractunssisq (unsigned int a)</pre>	[Runtime Function]
<pre>long long fractfractunssidq (unsigned int a)</pre>	[Runtime Function]
short accumfractunssiha (unsigned int a)	[Runtime Function]
accumfractunssisa (unsigned int a)	[Runtime Function]
long accumfractunssida (unsigned int a)	[Runtime Function]
long long accumfractunssita (unsigned int a)	[Runtime Function]
<pre>unsigned short fractfractunssiuqq (unsigned int a)</pre>	[Runtime Function]
unsigned fractfractunssiuhq (unsigned int a)	[Runtime Function]
unsigned long fractfractunssiusq (unsigned int a)	[Runtime Function]
unsigned long long fractfractunssiudq (unsigned int	[Runtime Function]
a)	
unsigned short accumfractunssiuha (unsigned int a)	[Runtime Function]
unsigned accumfractunssiusa (unsigned int a)	[Runtime Function]
unsigned long accumfractunssiuda (unsigned int a)	[Runtime Function]
unsigned long long accumfractunssiuta (unsigned int	[Runtime Function]
a)	
short fractfractunsdiqq (unsigned long a)	[Runtime Function]
fractfractunsdihq (unsigned long a)	[Runtime Function]
<pre>long fractfractunsdisq (unsigned long a)</pre>	[Runtime Function]
long long fractfractunsdidq (unsigned long a)	[Runtime Function]
short accumfractunsdiha (unsigned long a)	[Runtime Function]
accumfractunsdisa (unsigned long a)	[Runtime Function]
long accumfractunsdida (unsigned long a)	[Runtime Function]
long long accumfractunsdita (unsigned long a)	[Runtime Function]
unsigned short fractfractunsdiuqq (unsigned long a)	[Runtime Function]
unsigned fractfractunsdiuhq (unsigned long a)	[Runtime Function]
unsigned long fractfractunsdiusq (unsigned long a)	[Runtime Function]
unsigned long long fractfractunsdiudq (unsigned	[Runtime Function]
long a)	
unsigned short accumfractunsdiuha (unsigned long a)	[Runtime Function]
unsigned accumfractunsdiusa (unsigned long a)	[Runtime Function]
unsigned long accumfractunsdiuda (unsigned long a)	[Runtime Function]

unsigned long long accumfractunsdiuta (unsigned long a)	[Runtime Function]
<pre>short fractfractunstiqq (unsigned long long a)</pre>	[Runtime Function]
fractfractunstihq (unsigned long long a)	[Runtime Function]
<pre>long fractfractunstisq (unsigned long long a)</pre>	[Runtime Function]
long long fractfractunstidq (unsigned long long a)	[Runtime Function]
short accumfractunstiha (unsigned long long a)	[Runtime Function]
accumfractunstisa (unsigned long long a)	[Runtime Function]
long accumfractunstida (unsigned long long a)	[Runtime Function]
long long accumfractunstita (unsigned long long a)	[Runtime Function]
unsigned short fractfractunstiuqq (unsigned long	[Runtime Function]
long a)	
unsigned fractfractunstiuhq (unsigned long long a)	[Runtime Function]
unsigned long fractfractunstiusq (unsigned long long	[Runtime Function]
a)	
unsigned long long fractfractunstiudq (unsigned	[Runtime Function]
long long a)	
unsigned short accumfractunstiuha (unsigned long	[Runtime Function]
long a)	
unsigned accumfractunstiusa (unsigned long long a)	[Runtime Function]
unsigned long accumfractunstiuda (unsigned long long	[Runtime Function]
a)	
unsigned long long accumfractunstiuta (unsigned	[Runtime Function]
long long a)	
These functions convert from fractionals to unsigned non-fraction	onals; and from un-

These functions convert from fractionals to unsigned non-fractionals; and from unsigned non-fractionals to fractionals, without saturation.

```
short fract __satfractunsqiqq (unsigned char a)
                                                             [Runtime Function]
fract __satfractunsqihq (unsigned char a)
                                                             [Runtime Function]
long fract __satfractunsqisq (unsigned char a)
                                                             [Runtime Function]
long long fract __satfractunsqidq (unsigned char a)
                                                             [Runtime Function]
short accum __satfractunsqiha (unsigned char a)
                                                             [Runtime Function]
accum __satfractunsqisa (unsigned char a)
                                                             [Runtime Function]
long accum __satfractunsqida (unsigned char a)
                                                             [Runtime Function]
long long accum __satfractunsqita (unsigned char a)
                                                             [Runtime Function]
unsigned short fract __satfractunsqiuqq (unsigned char
                                                             [Runtime Function]
unsigned fract __satfractunsqiuhq (unsigned char a)
                                                             [Runtime Function]
unsigned long fract __satfractunsqiusq (unsigned char
                                                             [Runtime Function]
unsigned long long fract __satfractunsqiudq
                                                             [Runtime Function]
        (unsigned char a)
unsigned short accum __satfractunsqiuha (unsigned char
                                                             [Runtime Function]
unsigned accum __satfractunsqiusa (unsigned char a)
                                                             [Runtime Function]
unsigned long accum __satfractunsqiuda (unsigned char
                                                             [Runtime Function]
         a)
```

unsigned long long accumsatfractunsqiuta (unsigned char a)	[Runtime Function]
short fractsatfractunshiqq (unsigned short a) fractsatfractunshihq (unsigned short a) long fractsatfractunshisq (unsigned short a) long long fractsatfractunshidq (unsigned short a) short accumsatfractunshiha (unsigned short a) accumsatfractunshisa (unsigned short a) long accumsatfractunshida (unsigned short a) long long accumsatfractunshita (unsigned short a) unsigned short fractsatfractunshiuqq (unsigned short a) short a)	[Runtime Function]
unsigned fractsatfractunshiuhq (unsigned short a) unsigned long fractsatfractunshiusq (unsigned short	[Runtime Function] [Runtime Function]
a) unsigned long long fractsatfractunshiudq (unsigned short a)	[Runtime Function]
unsigned short accumsatfractunshiuha (unsigned short a)	[Runtime Function]
unsigned accumsatfractunshiusa (unsigned short a) unsigned long accumsatfractunshiuda (unsigned short a)	[Runtime Function] [Runtime Function]
unsigned long long accumsatfractunshiuta (unsigned short a)	[Runtime Function]
short fractsatfractunssiqq (unsigned int a) fractsatfractunssihq (unsigned int a) long fractsatfractunssisq (unsigned int a)	[Runtime Function] [Runtime Function]
long long fractsatfractunssidq (unsigned int a) short accumsatfractunssiha (unsigned int a) accumsatfractunssisa (unsigned int a)	[Runtime Function] [Runtime Function] [Runtime Function]
long accumsatfractunssida (unsigned int a) long long accumsatfractunssida (unsigned int a) unsigned short fractsatfractunssida (unsigned int)	[Runtime Function] [Runtime Function] [Runtime Function]
a) unsigned fractsatfractunssiuhq (unsigned int a) unsigned long fractsatfractunssiusq (unsigned int a) unsigned long long fractsatfractunssiudq (unsigned int a)	[Runtime Function] [Runtime Function]
unsigned short accumsatfractunssiuha (unsigned int a)	[Runtime Function]
unsigned accumsatfractunssiusa (unsigned int a) unsigned long accumsatfractunssiuda (unsigned int a) unsigned long long accumsatfractunssiuta (unsigned int a)	[Runtime Function] [Runtime Function] [Runtime Function]
short fractsatfractunsdiqq (unsigned long a) fractsatfractunsdihq (unsigned long a) long fractsatfractunsdisq (unsigned long a)	[Runtime Function] [Runtime Function]

<pre>long long fractsatfractunsdidq (unsigned long a) short accumsatfractunsdiha (unsigned long a) accumsatfractunsdisa (unsigned long a) long accumsatfractunsdida (unsigned long a)</pre>	[Runtime Function] [Runtime Function] [Runtime Function]
· /	
long long accumsatfractunsdita (unsigned long a)	[Runtime Function]
unsigned short fractsatfractunsdiuqq (unsigned long	[Runtime Function]
a)	
<pre>unsigned fractsatfractunsdiuhq (unsigned long a)</pre>	[Runtime Function]
unsigned long fractsatfractunsdiusq (unsigned long	[Runtime Function]
a)	. ,
unsigned long long fractsatfractunsdiudq	[Runtime Function]
(unsigned long a)	
` ,	[Duntima Function]
unsigned short accumsatfractunsdiuha (unsigned long	[Runtime Function]
a)	
unsigned accumsatfractunsdiusa (unsigned long a)	[Runtime Function]
unsigned long accumsatfractunsdiuda (unsigned long	[Runtime Function]
a)	
unsigned long long accumsatfractunsdiuta	[Runtime Function]
(unsigned long a)	
short fractsatfractunstiqq (unsigned long long a)	[Runtime Function]
fractsatfractunstihq (unsigned long long a)	[Runtime Function]
long fractsatfractunstisq (unsigned long long a)	[Runtime Function]
long long fractsatfractunstidq (unsigned long long a)	
· · · · · · · · · · · · · · · ·	[Runtime Function]
short accumsatfractunstiha (unsigned long long a)	[Runtime Function]
accumsatfractunstisa (unsigned long long a)	[Runtime Function]
long accumsatfractunstida (unsigned long long a)	[Runtime Function]
long long accumsatfractunstita (unsigned long long a)	[Runtime Function]
unsigned short fractsatfractunstiuqq (unsigned long	[Runtime Function]
long a)	
<pre>unsigned fractsatfractunstiuhq (unsigned long long a)</pre>	[Runtime Function]
unsigned long fractsatfractunstiusq (unsigned long	[Runtime Function]
long a)	·
unsigned long long fractsatfractunstiudq	[Runtime Function]
(unsigned long long a)	[realisme ranesion]
unsigned short accumsatfractunstiuha (unsigned long	[Duntima Function]
•	[Runtime Function]
long a)	[D D ]
unsigned accumsatfractunstiusa (unsigned long long a)	[Runtime Function]
unsigned long accumsatfractunstiuda (unsigned long	[Runtime Function]
long a)	
unsigned long long accumsatfractunstiuta	[Runtime Function]
(unsigned long long a)	
These functions convent from unsigned non fractional to furtion	ala mith aatumati

These functions convert from unsigned non-fractionals to fractionals, with saturation.

# 4.5 Language-independent routines for exception handling

document me!

\_Unwind\_DeleteException

```
_Unwind_Find_FDE
_Unwind_ForcedUnwind
_Unwind_GetGR
_Unwind_GetIP
_Unwind_GetLanguageSpecificData
_Unwind_GetRegionStart
_Unwind_GetTextRelBase
_Unwind_GetDataRelBase
_Unwind_RaiseException
_Unwind_Resume
_Unwind_SetGR
_Unwind_SetIP
_Unwind_FindEnclosingFunction
_Unwind_SjLj_Register
_Unwind_SjLj_Unregister
_Unwind_SjLj_RaiseException
_Unwind_SjLj_ForcedUnwind
_Unwind_SjLj_Resume
__deregister_frame
__deregister_frame_info
__deregister_frame_info_bases
__register_frame
__register_frame_info
__register_frame_info_bases
__register_frame_info_table
__register_frame_info_table_bases
__register_frame_table
```

# 4.6 Miscellaneous runtime library routines

# 4.6.1 Cache control functions

```
void __clear_cache (char *beg, char *end) [Runtime Function]
This function clears the instruction cache between beg and end.
```

# 4.6.2 Split stack functions and variables

There is no way to iterate over the stack segments of a different thread. However, what is permitted is for one thread to call this with the segment\_arg and sp arguments NULL, to pass next\_segment, next\_sp, and initial\_sp to a different thread, and then to suspend one way or another. A different thread may run the subsequent \_\_splitstack\_find iterations. Of course, this will only work if the first thread is suspended while the second thread is calling \_\_splitstack\_find. If not, the second thread could be looking at the stack while it is changing, and anything could happen.

```
__morestack_segments [Variable]
__morestack_current_segment [Variable]
__morestack_initial_sp [Variable]
Internal variables used by the '-fsplit-stack' implementation.
```

# 5 Language Front Ends in GCC

The interface to front ends for languages in GCC, and in particular the tree structure (see Chapter 11 [GENERIC], page 161), was initially designed for C, and many aspects of it are still somewhat biased towards C and C-like languages. It is, however, reasonably well suited to other procedural languages, and front ends for many such languages have been written for GCC.

Writing a compiler as a front end for GCC, rather than compiling directly to assembler or generating C code which is then compiled by GCC, has several advantages:

- GCC front ends benefit from the support for many different target machines already present in GCC.
- GCC front ends benefit from all the optimizations in GCC. Some of these, such as alias analysis, may work better when GCC is compiling directly from source code then when it is compiling from generated C code.
- Better debugging information is generated when compiling directly from source code than when going via intermediate generated C code.

Because of the advantages of writing a compiler as a GCC front end, GCC front ends have also been created for languages very different from those for which GCC was designed, such as the declarative logic/functional language Mercury. For these reasons, it may also be useful to implement compilers created for specialized purposes (for example, as part of a research project) as GCC front ends.

# 6 Source Tree Structure and Build System

This chapter describes the structure of the GCC source tree, and how GCC is built. The user documentation for building and installing GCC is in a separate manual (http://gcc.gnu.org/install/), with which it is presumed that you are familiar.

# 6.1 Configure Terms and History

The configure and build process has a long and colorful history, and can be confusing to anyone who doesn't know why things are the way they are. While there are other documents which describe the configuration process in detail, here are a few things that everyone working on GCC should know.

There are three system names that the build knows about: the machine you are building on (build), the machine that you are building for (host), and the machine that GCC will produce code for (target). When you configure GCC, you specify these with '--build=', '--host=', and '--target='.

Specifying the host without specifying the build should be avoided, as **configure** may (and once did) assume that the host you specify is also the build, which may not be true.

If build, host, and target are all the same, this is called a *native*. If build and host are the same but target is different, this is called a *cross*. If build, host, and target are all different this is called a *canadian* (for obscure reasons dealing with Canada's political party and the background of the person working on the build at that time). If host and target are the same, but build is different, you are using a cross-compiler to build a native for a different system. Some people call this a *host-x-host*, *crossed native*, or *cross-built native*. If build and target are the same, but host is different, you are using a cross compiler to build a cross compiler that produces code for the machine you're building on. This is rare, so there is no common way of describing it. There is a proposal to call this a *crossback*.

If build and host are the same, the GCC you are building will also be used to build the target libraries (like libstdc++). If build and host are different, you must have already built and installed a cross compiler that will be used to build the target libraries (if you configured with '--target=foo-bar', this compiler will be called foo-bar-gcc).

In the case of target libraries, the machine you're building for is the machine you specified with '--target'. So, build is the machine you're building on (no change there), host is the machine you're building for (the target libraries are built for the target, so host is the target you specified), and target doesn't apply (because you're not building a compiler, you're building libraries). The configure/make process will adjust these variables as needed. It also sets \$with\_cross\_host to the original '--host' value in case you need it.

The libiberty support library is built up to three times: once for the host, once for the target (even if they are the same), and once for the build if build and host are different. This allows it to be used by all programs which are generated in the course of the build process.

# 6.2 Top Level Source Directory

The top level source directory in a GCC distribution contains several files and directories that are shared with other software distributions such as that of GNU Binutils. It also contains several subdirectories that contain parts of GCC and its runtime libraries:

'boehm-gc'

The Boehm conservative garbage collector, optionally used as part of the ObjC runtime library when configured with '--enable-objc-gc'.

'config' Autoconf macros and Makefile fragments used throughout the tree.

'contrib' Contributed scripts that may be found useful in conjunction with GCC. One of these, 'contrib/texi2pod.pl', is used to generate man pages from Texinfo manuals as part of the GCC build process.

'fixincludes'

The support for fixing system headers to work with GCC. See 'fixincludes/README' for more information. The headers fixed by this mechanism are installed in 'libsubdir/include-fixed'. Along with those headers, 'README-fixinc' is also installed, as 'libsubdir/include-fixed/README'.

'gcc' The main sources of GCC itself (except for runtime libraries), including optimizers, support for different target architectures, language front ends, and testsuites. See Section 6.3 [The 'gcc' Subdirectory], page 63, for details.

'gnattools'

Support tools for GNAT.

'include' Headers for the libiberty library.

'intl' GNU libintl, from GNU gettext, for systems which do not include it in libc.

'libada' The Ada runtime library.

'libatomic'

The runtime support library for atomic operations (e.g. for \_\_sync and \_\_ atomic).

'libcpp' The C preprocessor library.

'libdecnumber'

The Decimal Float support library.

'libffi' The libffi library, used as part of the Go runtime library.

'libgcc' The GCC runtime library.

'libgfortran'

The Fortran runtime library.

'libgo' The Go runtime library. The bulk of this library is mirrored from the master Go repository.

'libgomp' The GNU Offloading and Multi Processing Runtime Library.

'libiberty'

The libiberty library, used for portability and for some generally useful data structures and algorithms. See Section "Introduction" in GNU libiberty, for more information about this library.

'libitm' The runtime support library for transactional memory.

'libobjc' The Objective-C and Objective-C++ runtime library.

'libquadmath'

The runtime support library for quad-precision math operations.

'libphobos'

The D standard and runtime library. The bulk of this library is mirrored from the master D repositories.

'libssp' The Stack protector runtime library.

'libstdc++-v3'

The C++ runtime library.

'lto-plugin'

Plugin used by the linker if link-time optimizations are enabled.

'maintainer-scripts'

Scripts used by the gccadmin account on gcc.gnu.org.

'zlib' The zlib compression library, used for compressing and uncompressing GCC's intermediate language in LTO object files.

The build system in the top level directory, including how recursion into subdirectories works and how building runtime libraries for multilibs is handled, is documented in a separate manual, included with GNU Binutils. See Section "GNU configure and build system" in *The GNU configure and build system*, for details.

# 6.3 The 'gcc' Subdirectory

The 'gcc' directory contains many files that are part of the C sources of GCC, other files used as part of the configuration and build process, and subdirectories including documentation and a testsuite. The files that are sources of GCC are documented in a separate chapter. See Chapter 9 [Passes and Files of the Compiler], page 127.

# 6.3.1 Subdirectories of 'gcc'

The 'gcc' directory contains the following subdirectories:

'language'

Subdirectories for various languages. Directories containing a file 'config-lang.in' are language subdirectories. The contents of the subdirectories 'c' (for C), 'cp' (for C++), 'objc' (for Objective-C), 'objcp' (for Objective-C++), and 'lto' (for LTO) are documented in this manual (see Chapter 9 [Passes and Files of the Compiler], page 127); those for other languages are not. See Section 6.3.8 [Anatomy of a Language Front End], page 71, for details of the files in these directories.

'common' Source files shared between the compiler drivers (such as gcc) and the compilers proper (such as 'cc1'). If an architecture defines target hooks shared between those places, it also has a subdirectory in 'common/config'. See Section 18.1 [Target Structure], page 479.

'config' Configuration files for supported architectures and operating systems. See Section 6.3.9 [Anatomy of a Target Back End], page 75, for details of the files in this directory.

'doc' Texinfo documentation for GCC, together with automatically generated man pages and support for converting the installation manual to HTML. See Section 6.3.7 [Documentation], page 69.

## 'ginclude'

System headers installed by GCC, mainly those required by the C standard of freestanding implementations. See Section 6.3.6 [Headers Installed by GCC], page 68, for details of when these and other headers are installed.

'po' Message catalogs with translations of messages produced by GCC into various languages, 'language.po'. This directory also contains 'gcc.pot', the template for these message catalogues, 'exgettext', a wrapper around gettext to extract the messages from the GCC sources and create 'gcc.pot', which is run by 'make gcc.pot', and 'EXCLUDES', a list of files from which messages should not be extracted.

#### 'testsuite'

The GCC testsuites (except for those for runtime libraries). See Chapter 7 [Testsuites], page 79.

# 6.3.2 Configuration in the 'gcc' Directory

The 'gcc' directory is configured with an Autoconf-generated script 'configure'. The 'configure' script is generated from 'configure.ac' and 'aclocal.m4'. From the files 'configure.ac' and 'acconfig.h', Autoheader generates the file 'config.in'. The file 'cstamp-h.in' is used as a timestamp.

# 6.3.2.1 Scripts Used by 'configure'

'configure' uses some other scripts to help in its work:

- The standard GNU 'config.sub' and 'config.guess' files, kept in the top level directory, are used.
- The file 'config.gcc' is used to handle configuration specific to the particular target machine. The file 'config.build' is used to handle configuration specific to the particular build machine. The file 'config.host' is used to handle configuration specific to the particular host machine. (In general, these should only be used for features that cannot reasonably be tested in Autoconf feature tests.) See Section 6.3.2.2 [The 'config.build'; 'config.host'; and 'config.gcc' Files], page 65, for details of the contents of these files.
- Each language subdirectory has a file 'language/config-lang.in' that is used for front-end-specific configuration. See Section 6.3.8.2 [The Front End 'config-lang.in' File], page 73, for details of this file.
- A helper script 'configure.frag' is used as part of creating the output of 'configure'.

# 6.3.2.2 The 'config.build'; 'config.host'; and 'config.gcc' Files

The 'config.build' file contains specific rules for particular systems which GCC is built on. This should be used as rarely as possible, as the behavior of the build system can always be detected by autoconf.

The 'config.host' file contains specific rules for particular systems which GCC will run on. This is rarely needed.

The 'config.gcc' file contains specific rules for particular systems which GCC will generate code for. This is usually needed.

Each file has a list of the shell variables it sets, with descriptions, at the top of the file.

FIXME: document the contents of these files, and what variables should be set to control build, host and target configuration.

# 6.3.2.3 Files Created by configure

Here we spell out what files will be set up by 'configure' in the 'gcc' directory. Some other files are created as temporary files in the configuration process, and are not used in the subsequent build; these are not documented.

- 'Makefile' is constructed from 'Makefile.in', together with the host and target fragments (see Chapter 20 [Makefile Fragments], page 667) 't-target' and 'x-host' from 'config', if any, and language Makefile fragments 'language/Make-lang.in'.
- 'auto-host.h' contains information about the host machine determined by 'configure'. If the host machine is different from the build machine, then 'auto-build.h' is also created, containing such information about the build machine.
- 'config.status' is a script that may be run to recreate the current configuration.
- 'configargs.h' is a header containing details of the arguments passed to 'configure' to configure GCC, and of the thread model used.
- 'cstamp-h' is used as a timestamp.
- If a language 'config-lang.in' file (see Section 6.3.8.2 [The Front End 'config-lang.in' File], page 73) sets outputs, then the files listed in outputs there are also generated.

The following configuration headers are created from the Makefile, using 'mkconfig.sh', rather than directly by 'configure'. 'config.h', 'bconfig.h' and 'tconfig.h' all contain the 'xm-machine.h' header, if any, appropriate to the host, build and target machines respectively, the configuration headers for the target, and some definitions; for the host and build machines, these include the autoconfigured headers generated by 'configure'. The other configuration headers are determined by 'config.gcc'. They also contain the typedefs for rtx, rtvec and tree.

- 'config.h', for use in programs that run on the host machine.
- 'bconfig.h', for use in programs that run on the build machine.
- 'tconfig.h', for use in programs and libraries for the target machine.
- 'tm\_p.h', which includes the header 'machine-protos.h' that contains prototypes for functions in the target 'machine.c' file. The 'machine-protos.h' header is included after the 'rtl.h' and/or 'tree.h' would have been included. The 'tm\_p.h' also includes the header 'tm-preds.h' which is generated by 'genpreds' program during the build to define the declarations and inline functions for the predicate functions.

# 6.3.3 Build System in the 'gcc' Directory

FIXME: describe the build system, including what is built in what stages. Also list the various source files that are used in the build process but aren't source files of GCC itself and so aren't documented below (see Chapter 9 [Passes], page 127).

# 6.3.4 Makefile Targets

These targets are available from the 'gcc' directory:

This is the default target. Depending on what your build/host/target configuration is, it coordinates all the things that need to be built.

doc Produce info-formatted documentation and man pages. Essentially it calls 'make man' and 'make info'.

dvi Produce DVI-formatted documentation.

pdf Produce PDF-formatted documentation.

html Produce HTML-formatted documentation.

man Generate man pages.

info Generate info-formatted pages.

#### mostlyclean

Delete the files made while building the compiler.

clean That, and all the other files built by 'make all'.

#### distclean

That, and all the files created by configure.

#### maintainer-clean

Distclean plus any file that can be generated from other files. Note that additional tools may be required beyond what is normally needed to build GCC.

srcextra Generates files in the source directory that are not version-controlled but should go into a release tarball.

srcinfo

Srcman Copies the info-formatted and manpage documentation into the source directory usually for the purpose of generating a release tarball.

install Installs GCC.

#### uninstall

Deletes installed files, though this is not supported.

Check Run the testsuite. This creates a 'testsuite' subdirectory that has various '.sum' and '.log' files containing the results of the testing. You can run subsets with, for example, 'make check-gcc'. You can specify specific tests by setting RUNTESTFLAGS to be the name of the '.exp' file, optionally followed by (for some tests) an equals and a file wildcard, like:

make check-gcc RUNTESTFLAGS="execute.exp=19980413-\*"

Note that running the testsuite may require additional tools be installed, such as Tcl or DejaGnu.

The toplevel tree from which you start GCC compilation is not the GCC directory, but rather a complex Makefile that coordinates the various steps of the build, including bootstrapping the compiler and using the new compiler to build target libraries.

When GCC is configured for a native configuration, the default action for make is to do a full three-stage bootstrap. This means that GCC is built three times—once with the native compiler, once with the native-built compiler it just built, and once with the compiler it built the second time. In theory, the last two should produce the same results, which 'make compare' can check. Each stage is configured separately and compiled into a separate directory, to minimize problems due to ABI incompatibilities between the native compiler and GCC.

If you do a change, rebuilding will also start from the first stage and "bubble" up the change through the three stages. Each stage is taken from its build directory (if it had been built previously), rebuilt, and copied to its subdirectory. This will allow you to, for example, continue a bootstrap after fixing a bug which causes the stage2 build to crash. It does not provide as good coverage of the compiler as bootstrapping from scratch, but it ensures that the new code is syntactically correct (e.g., that you did not use GCC extensions by mistake), and avoids spurious bootstrap comparison failures<sup>1</sup>.

Other targets available from the top level include:

### bootstrap-lean

Like bootstrap, except that the various stages are removed once they're no longer needed. This saves disk space.

### bootstrap2

#### bootstrap2-lean

Performs only the first two stages of bootstrap. Unlike a three-stage bootstrap, this does not perform a comparison to test that the compiler is running properly. Note that the disk space required by a "lean" bootstrap is approximately independent of the number of stages.

# stageN-bubble (N = 1...4, profile, feedback)

Rebuild all the stages up to N, with the appropriate flags, "bubbling" the changes as described above.

## all-stageN (N = 1...4, profile, feedback)

Assuming that stage N has already been built, rebuild it with the appropriate flags. This is rarely needed.

#### cleanstrap

Remove everything ('make clean') and rebuilds ('make bootstrap').

compare Compares the results of stages 2 and 3. This ensures that the compiler is running properly, since it should produce the same object files regardless of how it itself was compiled.

#### profiledbootstrap

Builds a compiler with profiling feedback information. In this case, the second and third stages are named 'profile' and 'feedback', respectively. For more information, see the installation instructions.

Except if the compiler was buggy and miscompiled some of the files that were not modified. In this case, it's best to use make restrap.

restrap Restart a bootstrap, so that everything that was not built with the system compiler is rebuilt.

stageN-start (N = 1...4, profile, feedback)

For each package that is bootstrapped, rename directories so that, for example, 'gcc' points to the stage NGCC, compiled with the stage N-1 GCC<sup>2</sup>.

You will invoke this target if you need to test or debug the stage N GCC. If you only need to execute GCC (but you need not run 'make' either to rebuild it or to run test suites), you should be able to work directly in the 'stageN-gcc' directory. This makes it easier to debug multiple stages in parallel.

For each package that is bootstrapped, relocate its build directory to indicate its stage. For example, if the 'gcc' directory points to the stage2 GCC, after invoking this target it will be renamed to 'stage2-gcc'.

If you wish to use non-default GCC flags when compiling the stage2 and stage3 compilers, set BOOT\_CFLAGS on the command line when doing 'make'.

Usually, the first stage only builds the languages that the compiler is written in: typically, C and maybe Ada. If you are debugging a miscompilation of a different stage2 front-end (for example, of the Fortran front-end), you may want to have front-ends for other languages in the first stage as well. To do so, set STAGE1\_LANGUAGES on the command line when doing 'make'.

For example, in the aforementioned scenario of debugging a Fortran front-end miscompilation caused by the stage1 compiler, you may need a command like

make stage2-bubble STAGE1\_LANGUAGES=c,fortran

Alternatively, you can use per-language targets to build and test languages that are not enabled by default in stage1. For example, make f951 will build a Fortran compiler even in the stage1 build directory.

# 6.3.5 Library Source Files and Headers under the 'gcc' Directory

FIXME: list here, with explanation, all the C source files and headers under the 'gcc' directory that aren't built into the GCC executable but rather are part of runtime libraries and object files, such as 'crtstuff.c' and 'unwind-dw2.c'. See Section 6.3.6 [Headers Installed by GCC], page 68, for more information about the 'ginclude' directory.

# 6.3.6 Headers Installed by GCC

In general, GCC expects the system C library to provide most of the headers to be used with it. However, GCC will fix those headers if necessary to make them work with GCC, and will install some headers required of freestanding implementations. These headers are installed in 'libsubdir/include'. Headers for non-C runtime libraries are also installed by GCC; these are not documented here. (FIXME: document them somewhere.)

Several of the headers GCC installs are in the 'ginclude' directory. These headers, 'iso646.h', 'stdarg.h', 'stdbool.h', and 'stddef.h', are installed in 'libsubdir/include', unless the target Makefile fragment (see Section 20.1 [Target Fragment], page 667) overrides this by setting USER\_H.

<sup>&</sup>lt;sup>2</sup> Customarily, the system compiler is also termed the 'stage0' GCC.

In addition to these headers and those generated by fixing system headers to work with GCC, some other headers may also be installed in 'libsubdir/include'. 'config.gcc' may set extra\_headers; this specifies additional headers under 'config' to be installed on some systems.

GCC installs its own version of <float.h>, from 'ginclude/float.h'. This is done to cope with command-line options that change the representation of floating point numbers.

GCC also installs its own version of <limits.h>; this is generated from 'glimits.h', together with 'limitx.h' and 'limity.h' if the system also has its own version of <limits.h>. (GCC provides its own header because it is required of ISO C freestanding implementations, but needs to include the system header from its own header as well because other standards such as POSIX specify additional values to be defined in <limits.h>.) The system's limits.h> header is used via 'libsubdir/include/syslimits.h', which is copied from 'gsyslimits.h' if it does not need fixing to work with GCC; if it needs fixing, 'syslimits.h' is the fixed copy.

GCC can also install <tgmath.h>. It will do this when 'config.gcc' sets use\_gcc\_tgmath to yes.

# 6.3.7 Building Documentation

The main GCC documentation is in the form of manuals in Texinfo format. These are installed in Info format; DVI versions may be generated by 'make dvi', PDF versions by 'make pdf', and HTML versions by 'make html'. In addition, some man pages are generated from the Texinfo manuals, there are some other text files with miscellaneous documentation, and runtime libraries have their own documentation outside the 'gcc' directory. FIXME: document the documentation for runtime libraries somewhere.

# 6.3.7.1 Texinfo Manuals

The manuals for GCC as a whole, and the C and C++ front ends, are in files 'doc/\*.texi'. Other front ends have their own manuals in files 'language/\*.texi'. Common files 'doc/include/\*.texi' are provided which may be included in multiple manuals; the following files are in 'doc/include':

```
'fdl.texi'
```

The GNU Free Documentation License.

'funding.texi'

The section "Funding Free Software".

'gcc-common.texi'

Common definitions for manuals.

'gpl\_v3.texi'

The GNU General Public License.

'texinfo.tex'

A copy of 'texinfo.tex' known to work with the GCC manuals.

DVI-formatted manuals are generated by 'make dvi', which uses texi2dvi (via the Makefile macro \$(TEXI2DVI)). PDF-formatted manuals are generated by 'make pdf', which uses texi2pdf (via the Makefile macro \$(TEXI2PDF)). HTML formatted manuals are generated

by 'make html'. Info manuals are generated by 'make info' (which is run as part of a bootstrap); this generates the manuals in the source directory, using makeinfo via the Makefile macro \$(MAKEINFO), and they are included in release distributions.

Manuals are also provided on the GCC web site, in both HTML and PostScript forms. This is done via the script 'maintainer-scripts/update\_web\_docs\_git'. Each manual to be provided online must be listed in the definition of MANUALS in that file; a file 'name.texi' must only appear once in the source tree, and the output manual must have the same name as the source file. (However, other Texinfo files, included in manuals but not themselves the root files of manuals, may have names that appear more than once in the source tree.) The manual file 'name.texi' should only include other files in its own directory or in 'doc/include'. HTML manuals will be generated by 'makeinfo --html', PostScript manuals by texi2dvi and dvips, and PDF manuals by texi2pdf. All Texinfo files that are parts of manuals must be version-controlled, even if they are generated files, for the generation of online manuals to work.

The installation manual, 'doc/install.texi', is also provided on the GCC web site. The HTML version is generated by the script 'doc/install.texi2html'.

# 6.3.7.2 Man Page Generation

Because of user demand, in addition to full Texinfo manuals, man pages are provided which contain extracts from those manuals. These man pages are generated from the Texinfo manuals using 'contrib/texi2pod.pl' and pod2man. (The man page for g++, 'cp/g++.1', just contains a '.so' reference to 'gcc.1', but all the other man pages are generated from Texinfo manuals.)

Because many systems may not have the necessary tools installed to generate the man pages, they are only generated if the 'configure' script detects that recent enough tools are installed, and the Makefiles allow generating man pages to fail without aborting the build. Man pages are also included in release distributions. They are generated in the source directory.

Magic comments in Texinfo files starting '@c man' control what parts of a Texinfo file go into a man page. Only a subset of Texinfo is supported by 'texi2pod.pl', and it may be necessary to add support for more Texinfo features to this script when generating new man pages. To improve the man page output, some special Texinfo macros are provided in 'doc/include/gcc-common.texi' which 'texi2pod.pl' understands:

## @gcctabopt

Use in the form '@table @gcctabopt' for tables of options, where for printed output the effect of '@code' is better than that of '@option' but for man page output a different effect is wanted.

#### @gccoptlist

Use for summary lists of options in manuals.

Use at the end of each line inside '@gccoptlist'. This is necessary to avoid problems with differences in how the '@gccoptlist' macro is handled by different Texinfo formatters.

FIXME: describe the 'texi2pod.pl' input language and magic comments in more detail.

#### 6.3.7.3 Miscellaneous Documentation

In addition to the formal documentation that is installed by GCC, there are several other text files in the 'gcc' subdirectory with miscellaneous documentation:

#### 'ABOUT-GCC-NLS'

Notes on GCC's Native Language Support. FIXME: this should be part of this manual rather than a separate file.

#### 'ABOUT-NLS'

Notes on the Free Translation Project.

'COPYING'

'COPYING3'

The GNU General Public License, Versions 2 and 3.

'COPYING.LIB'

'COPYING3.LIB'

The GNU Lesser General Public License, Versions 2.1 and 3.

'\*ChangeLog\*'

'\*/ChangeLog\*

Change log files for various parts of GCC.

#### 'LANGUAGES'

Details of a few changes to the GCC front-end interface. FIXME: the information in this file should be part of general documentation of the front-end interface in this manual.

'ONEWS' Information about new features in old versions of GCC. (For recent versions, the information is on the GCC web site.)

#### 'README.Portability'

Information about portability issues when writing code in GCC. FIXME: why isn't this part of this manual or of the GCC Coding Conventions?

FIXME: document such files in subdirectories, at least 'config', 'c', 'cp', 'objc', 'testsuite'.

# 6.3.8 Anatomy of a Language Front End

A front end for a language in GCC has the following parts:

- A directory 'language' under 'gcc' containing source files for that front end. See Section 6.3.8.1 [The Front End 'language' Directory], page 72, for details.
- A mention of the language in the list of supported languages in 'gcc/doc/install.texi'.
- A mention of the name under which the language's runtime library is recognized by '--enable-shared=package' in the documentation of that option in 'gcc/doc/install.texi'.
- A mention of any special prerequisites for building the front end in the documentation of prerequisites in 'gcc/doc/install.texi'.
- Details of contributors to that front end in 'gcc/doc/contrib.texi'. If the details are in that front end's own manual then there should be a link to that manual's list in 'contrib.texi'.

- Information about support for that language in 'gcc/doc/frontends.texi'.
- Information about standards for that language, and the front end's support for them, in 'gcc/doc/standards.texi'. This may be a link to such information in the front end's own manual.
- Details of source file suffixes for that language and '-x lang' options supported, in 'gcc/doc/invoke.texi'.
- Entries in default\_compilers in 'gcc.c' for source file suffixes for that language.
- Preferably testsuites, which may be under 'gcc/testsuite' or runtime library directories. FIXME: document somewhere how to write testsuite harnesses.
- Probably a runtime library for the language, outside the 'gcc' directory. FIXME: document this further.
- Details of the directories of any runtime libraries in 'gcc/doc/sourcebuild.texi'.
- Check targets in 'Makefile.def' for the top-level 'Makefile' to check just the compiler or the compiler and runtime library for the language.

If the front end is added to the official GCC source repository, the following are also necessary:

- At least one Bugzilla component for bugs in that front end and runtime libraries. This category needs to be added to the Bugzilla database.
- Normally, one or more maintainers of that front end listed in 'MAINTAINERS'.
- Mentions on the GCC web site in 'index.html' and 'frontends.html', with any relevant links on 'readings.html'. (Front ends that are not an official part of GCC may also be listed on 'frontends.html', with relevant links.)
- A news item on 'index.html', and possibly an announcement on the gcc-announce@gcc.gnu.org mailing list.
- The front end's manuals should be mentioned in 'maintainer-scripts/update\_web\_docs\_git' (see Section 6.3.7.1 [Texinfo Manuals], page 69) and the online manuals should be linked to from 'onlinedocs/index.html'.
- Any old releases or CVS repositories of the front end, before its inclusion in GCC, should be made available on the GCC web site at https://gcc.gnu.org/pub/gcc/old-releases/.
- The release and snapshot script 'maintainer-scripts/gcc\_release' should be updated to generate appropriate tarballs for this front end.
- If this front end includes its own version files that include the current date, 'maintainer-scripts/update\_version' should be updated accordingly.

# 6.3.8.1 The Front End 'language' Directory

A front end 'language' directory contains the source files of that front end (but not of any runtime libraries, which should be outside the 'gcc' directory). This includes documentation, and possibly some subsidiary programs built alongside the front end. Certain files are special and other parts of the compiler depend on their names:

'config-lang.in'

This file is required in all language subdirectories. See Section 6.3.8.2 [The Front End 'config-lang.in' File], page 73, for details of its contents

### 'Make-lang.in'

This file is required in all language subdirectories. See Section 6.3.8.3 [The Front End 'Make-lang.in' File], page 74, for details of its contents.

#### 'lang.opt'

This file registers the set of switches that the front end accepts on the command line, and their '--help' text. See Chapter 8 [Options], page 119.

## 'lang-specs.h'

This file provides entries for default\_compilers in 'gcc.c' which override the default of giving an error that a compiler for that language is not installed.

# 'language-tree.def'

This file, which need not exist, defines any language-specific tree codes.

# 6.3.8.2 The Front End 'config-lang.in' File

Each language subdirectory contains a 'config-lang.in' file. This file is a shell script that may define some variables describing the language:

This definition must be present, and gives the name of the language for some purposes such as arguments to '--enable-languages'.

### lang\_requires

If defined, this variable lists (space-separated) language front ends other than C that this front end requires to be enabled (with the names given being their language settings). For example, the Obj-C++ front end depends on the C++ and ObjC front ends, so sets 'lang\_requires="objc c++"'.

#### subdir\_requires

If defined, this variable lists (space-separated) front end directories other than C that this front end requires to be present. For example, the Objective-C++ front end uses source files from the C++ and Objective-C front ends, so sets 'subdir\_requires="cp objc".

#### target\_libs

If defined, this variable lists (space-separated) targets in the top level 'Makefile' to build the runtime libraries for this language, such as target-libobjc.

# lang\_dirs

If defined, this variable lists (space-separated) top level directories (parallel to 'gcc'), apart from the runtime libraries, that should not be configured if this front end is not built.

# build\_by\_default

If defined to 'no', this language front end is not built unless enabled in a '--enable-languages' argument. Otherwise, front ends are built by default, subject to any special logic in 'configure.ac' (as is present to disable the Ada front end if the Ada compiler is not already installed).

#### boot\_language

If defined to 'yes', this front end is built in stage1 of the bootstrap. This is only relevant to front ends written in their own languages.

compilers

If defined, a space-separated list of compiler executables that will be run by the driver. The names here will each end with '\\$(exeext)'.

outputs If defined, a space-separated list of files that should be generated by 'configure' substituting values in them. This mechanism can be used to create a file 'language/Makefile' from 'language/Makefile.in', but this is deprecated, building everything from the single 'gcc/Makefile' is preferred.

gtfiles If defined, a space-separated list of files that should be scanned by 'gengtype.c' to generate the garbage collection tables and routines for this language. This excludes the files that are common to all front ends. See Chapter 23 [Type Information], page 675.

# 6.3.8.3 The Front End 'Make-lang.in' File

Each language subdirectory contains a 'Make-lang.in' file. It contains targets lang.hook (where lang is the setting of language in 'config-lang.in') for the following values of hook, and any other Makefile rules required to build those targets (which may if necessary use other Makefiles specified in outputs in 'config-lang.in', although this is deprecated). It also adds any testsuite targets that can use the standard rule in 'gcc/Makefile.in' to the variable lang\_checks.

all.cross
start.encap
rest.encap

FIXME: exactly what goes in each of these targets?

tags Build an etags 'TAGS' file in the language subdirectory in the source tree.

info Build info documentation for the front end, in the build directory. This target is only called by 'make bootstrap' if a suitable version of makeinfo is available, so does not need to check for this, and should fail if an error occurs.

dvi Build DVI documentation for the front end, in the build directory. This should be done using \$(TEXI2DVI), with appropriate '-I' arguments pointing to directories of included files.

pdf Build PDF documentation for the front end, in the build directory. This should be done using \$(TEXI2PDF), with appropriate '-I' arguments pointing to directories of included files.

html Build HTML documentation for the front end, in the build directory.

Build generated man pages for the front end from Texinfo manuals (see Section 6.3.7.2 [Man Page Generation], page 70), in the build directory. This target is only called if the necessary tools are available, but should ignore errors so as not to stop the build if errors occur; man pages are optional and the tools involved may be installed in a broken way.

#### install-common

Install everything that is part of the front end, apart from the compiler executables listed in compilers in 'config-lang.in'.

#### install-info

Install info documentation for the front end, if it is present in the source directory. This target should have dependencies on info files that should be installed.

#### install-man

Install man pages for the front end. This target should ignore errors.

#### install-plugin

Install headers needed for plugins.

as a 'configure' option.

STCEXTRA Copies its dependencies into the source directory. This generally should be used for generated files such as Bison output files which are not version-controlled, but should be included in any release tarballs. This target will be executed during a bootstrap if '--enable-generated-files-in-srcdir' was specified

srcinfo

srcman

Copies its dependencies into the source directory. These targets will be executed during a bootstrap if '--enable-generated-files-in-srcdir' was specified as a 'configure' option.

#### uninstall

Uninstall files installed by installing the compiler. This is currently documented not to be supported, so the hook need not do anything.

mostlyclean clean distclean maintainer-clean

The language parts of the standard GNU '\*clean' targets. See Section "Standard Targets for Users" in *GNU Coding Standards*, for details of the standard targets. For GCC, maintainer-clean should delete all generated files in the source directory that are not version-controlled, but should not delete anything that is.

'Make-lang.in' must also define a variable lang\_OBJS to a list of host object files that are used by that language.

# 6.3.9 Anatomy of a Target Back End

A back end for a target architecture in GCC has the following parts:

- A directory 'machine' under 'gcc/config', containing a machine description 'machine.md' file (see Chapter 17 [Machine Descriptions], page 337), header files 'machine.h' and 'machine-protos.h' and a source file 'machine.c' (see Chapter 18 [Target Description Macros and Functions], page 479), possibly a target Makefile fragment 't-machine' (see Section 20.1 [The Target Makefile Fragment], page 667), and maybe some other files. The names of these files may be changed from the defaults given by explicit specifications in 'config.gcc'.
- If necessary, a file 'machine-modes.def' in the 'machine' directory, containing additional machine modes to represent condition codes. See Section 18.15 [Condition Code], page 572, for further details.

- An optional 'machine.opt' file in the 'machine' directory, containing a list of target-specific options. You can also add other option files using the extra\_options variable in 'config.gcc'. See Chapter 8 [Options], page 119.
- Entries in 'config.gcc' (see Section 6.3.2.2 [The 'config.gcc' File], page 65) for the systems with this target architecture.
- Documentation in 'gcc/doc/invoke.texi' for any command-line options supported by this target (see Section 18.3 [Run-time Target Specification], page 486). This means both entries in the summary table of options and details of the individual options.
- Documentation in 'gcc/doc/extend.texi' for any target-specific attributes supported (see Section 18.24 [Defining target-specific uses of \_\_attribute\_\_], page 633), including where the same attribute is already supported on some targets, which are enumerated in the manual.
- Documentation in 'gcc/doc/extend.texi' for any target-specific pragmas supported.
- Documentation in 'gcc/doc/extend.texi' of any target-specific built-in functions supported.
- Documentation in 'gcc/doc/extend.texi' of any target-specific format checking styles supported.
- Documentation in 'gcc/doc/md.texi' of any target-specific constraint letters (see Section 17.8.5 [Constraints for Particular Machines], page 357).
- A note in 'gcc/doc/contrib.texi' under the person or people who contributed the target support.
- Entries in 'gcc/doc/install.texi' for all target triplets supported with this target architecture, giving details of any special notes about installation for this target, or saying that there are no special notes if there are none.
- Possibly other support outside the 'gcc' directory for runtime libraries. FIXME: reference docs for this. The libstdc++ porting manual needs to be installed as info for this to work, or to be a chapter of this manual.

The 'machine.h' header is included very early in GCC's standard sequence of header files, while 'machine-protos.h' is included late in the sequence. Thus 'machine-protos.h' can include declarations referencing types that are not defined when 'machine.h' is included, specifically including those from 'rtl.h' and 'tree.h'. Since both RTL and tree types may not be available in every context where 'machine-protos.h' is included, in this file you should guard declarations using these types inside appropriate #ifdef RTX\_CODE or #ifdef TREE\_CODE conditional code segments.

If the backend uses shared data structures that require GTY markers for garbage collection (see Chapter 23 [Type Information], page 675), you must declare those in 'machine.h' rather than 'machine-protos.h'. Any definitions required for building libgcc must also go in 'machine.h'.

GCC uses the macro IN\_TARGET\_CODE to distinguish between machine-specific '.c' and '.cc' files and machine-independent '.c' and '.cc' files. Machine-specific files should use the directive:

#define IN\_TARGET\_CODE 1
before including config.h.

If the back end is added to the official GCC source repository, the following are also necessary:

- An entry for the target architecture in 'readings.html' on the GCC web site, with any relevant links.
- Details of the properties of the back end and target architecture in 'backends.html' on the GCC web site.
- A news item about the contribution of support for that target architecture, in 'index.html' on the GCC web site.
- Normally, one or more maintainers of that target listed in 'MAINTAINERS'. Some existing architectures may be unmaintained, but it would be unusual to add support for a target that does not have a maintainer when support is added.
- Target triplets covering all 'config.gcc' stanzas for the target, in the list in 'contrib/config-list.mk'.

# 7 Testsuites

GCC contains several testsuites to help maintain compiler quality. Most of the runtime libraries and language front ends in GCC have testsuites. Currently only the C language testsuites are documented here; FIXME: document the others.

# 7.1 Idioms Used in Testsuite Code

In general, C testcases have a trailing '-n.c', starting with '-1.c', in case other testcases with similar names are added later. If the test is a test of some well-defined feature, it should have a name referring to that feature such as 'feature-1.c'. If it does not test a well-defined feature but just happens to exercise a bug somewhere in the compiler, and a bug report has been filed for this bug in the GCC bug database, 'prbug-number-1.c' is the appropriate form of name. Otherwise (for miscellaneous bugs not filed in the GCC bug database), and previously more generally, test cases are named after the date on which they were added. This allows people to tell at a glance whether a test failure is because of a recently found bug that has not yet been fixed, or whether it may be a regression, but does not give any other information about the bug or where discussion of it may be found. Some other language testsuites follow similar conventions.

In the 'gcc.dg' testsuite, it is often necessary to test that an error is indeed a hard error and not just a warning—for example, where it is a constraint violation in the C standard, which must become an error with '-pedantic-errors'. The following idiom, where the first line shown is line line of the file and the line that generates the error, is used for this:

```
/* { dg-bogus "warning" "warning in place of error" } */
/* { dg-error "regexp" "message" { target *-*-* } line } */
```

It may be necessary to check that an expression is an integer constant expression and has a certain value. To check that E has value V, an idiom similar to the following is used:

```
char x[((E) == (V) ? 1 : -1)];
```

In 'gcc.dg' tests, \_\_typeof\_\_ is sometimes used to make assertions about the types of expressions. See, for example, 'gcc.dg/c99-condexpr-1.c'. The more subtle uses depend on the exact rules for the types of conditional expressions in the C standard; see, for example, 'gcc.dg/c99-intconst-1.c'.

It is useful to be able to test that optimizations are being made properly. This cannot be done in all cases, but it can be done where the optimization will lead to code being optimized away (for example, where flow analysis or alias analysis should show that certain code cannot be called) or to functions not being called because they have been expanded as built-in functions. Such tests go in 'gcc.c-torture/execute'. Where code should be optimized away, a call to a nonexistent function such as link\_failure () may be inserted; a definition

```
#ifndef __OPTIMIZE__
void
link_failure (void)
{
   abort ();
}
#endif
```

will also be needed so that linking still succeeds when the test is run without optimization. When all calls to a built-in function should have been optimized and no calls to the non-built-in version of the function should remain, that function may be defined as static to call abort () (although redeclaring a function as static may not work on all targets).

All testcases must be portable. Target-specific testcases must have appropriate code to avoid causing failures on unsupported systems; unfortunately, the mechanisms for this differ by directory.

FIXME: discuss non-C testsuites here.

# 7.2 Directives used within DejaGnu tests

# 7.2.1 Syntax and Descriptions of test directives

Test directives appear within comments in a test source file and begin with dg-. Some of these are defined within DejaGnu and others are local to the GCC testsuite.

The order in which test directives appear in a test can be important: directives local to GCC sometimes override information used by the DejaGnu directives, which know nothing about the GCC directives, so the DejaGnu directives must precede GCC directives.

Several test directives include selectors (see Section 7.2.2 [Selectors], page 84) which are usually preceded by the keyword target or xfail.

# 7.2.1.1 Specify how to build the test

# { dg-do do-what-keyword [{ target/xfail selector }] }

do-what-keyword specifies how the test is compiled and whether it is executed. It is one of:

### preprocess

Compile with '-E' to run only the preprocessor.

compile Compile with '-S' to produce an assembly code file.

assemble Compile with '-c' to produce a relocatable object file.

link Compile, assemble, and link to produce an executable file.

run Produce and run an executable file, which is expected to return an exit code of 0.

The default is compile. That can be overridden for a set of tests by redefining dg-do-what-default within the .exp file for those tests.

If the directive includes the optional '{ target selector}' then the test is skipped unless the target system matches the selector.

If do-what-keyword is run and the directive includes the optional '{ xfail selector}' and the selector is met then the test is expected to fail. The xfail clause is ignored for other values of do-what-keyword; those tests can use directive dg-xfail-if.

# 7.2.1.2 Specify additional compiler options

## { dg-options options [{ target selector }] }

This DejaGnu directive provides a list of compiler options, to be used if the target system matches *selector*, that replace the default options used for this set of tests.

### { dg-add-options feature ... }

Add any compiler options that are needed to access certain features. This directive does nothing on targets that enable the features by default, or that don't provide them at all. It must come after all dg-options directives. For supported values of feature see Section 7.2.4 [Add Options], page 104.

### { dg-additional-options options [{ target selector }] }

This directive provides a list of compiler options, to be used if the target system matches *selector*, that are added to the default options used for this set of tests.

# 7.2.1.3 Modify the test timeout value

The normal timeout limit, in seconds, is found by searching the following in order:

- the value defined by an earlier dg-timeout directive in the test
- variable tool\_timeout defined by the set of tests
- gcc, timeout set in the target board
- 300

# { dg-timeout n [{target selector }] }

Set the time limit for the compilation and for the execution of the test to the specified number of seconds.

#### { dg-timeout-factor x [{ target selector }] }

Multiply the normal time limit for compilation and execution of the test by the specified floating-point factor.

# 7.2.1.4 Skip a test for some targets

### { dg-skip-if comment { selector } [{ include-opts } [{ exclude-opts }]] }

Arguments include-opts and exclude-opts are lists in which each element is a string of zero or more GCC options. Skip the test if all of the following conditions are met:

- the test system is included in selector
- for at least one of the option strings in *include-opts*, every option from that string is in the set of options with which the test would be compiled; use ""\*" for an *include-opts* list that matches any options; that is the default if *include-opts* is not specified
- for each of the option strings in *exclude-opts*, at least one option from that string is not in the set of options with which the test would be compiled; use """ for an empty *exclude-opts* list; that is the default if *exclude-opts* is not specified

For example, to skip a test if option -Os is present:

```
/* { dg-skip-if "" { *-*-* } { "-0s" } { "" } } */

To skip a test if both options -02 and -g are present:
    /* { dg-skip-if "" { *-*-* } { "-02 -g" } { "" } } */

To skip a test if either -02 or -03 is present:
    /* { dg-skip-if "" { *-*-* } { "-02" "-03" } { "" } } */

To skip a test unless option -0s is present:
    /* { dg-skip-if "" { *-*-* } { "*" } { "-0s" } } */
```

To skip a test if either -0.2 or -0.3 is used with -g but not if -fpic is also present:

```
/* { dg-skip-if "" { *-*-* } { "-02 -g" "-03 -g" } { "-fpic" } } */
```

## { dg-require-effective-target keyword [{ selector }] }

Skip the test if the test target, including current multilib flags, is not covered by the effective-target keyword. If the directive includes the optional '{ selector}' then the effective-target test is only performed if the target system matches the selector. This directive must appear after any dg-do directive in the test and before any dg-additional-sources directive. See Section 7.2.3 [Effective-Target Keywords], page 84.

#### { dg-require-support args }

Skip the test if the target does not provide the required support. These directives must appear after any dg-do directive in the test and before any dg-additional-sources directive. They require at least one argument, which can be an empty string if the specific procedure does not examine the argument. See Section 7.2.5 [Require Support], page 106, for a complete list of these directives.

# 7.2.1.5 Expect a test to fail for some targets

- { dg-xfail-if comment { selector } [{ include-opts } [{ exclude-opts }]] }

  Expect the test to fail if the conditions (which are the same as for dg-skip-if)
  are met. This does not affect the execute step.
- { dg-xfail-run-if comment { selector } [{ include-opts } [{ exclude-opts }]] }
  Expect the execute step of a test to fail if the conditions (which are the same as for dg-skip-if) are met.

# 7.2.1.6 Expect the test executable to fail

{ dg-shouldfail comment [{ selector } [{ include-opts } [{ exclude-opts }]]] }

Expect the test executable to return a nonzero exit status if the conditions (which are the same as for dg-skip-if) are met.

# 7.2.1.7 Verify compiler messages

Where *line* is an accepted argument for these commands, a value of '0' can be used if there is no line associated with the message.

```
{ dg-error regexp [comment [{ target/xfail selector } [line] ]] }
```

This DejaGnu directive appears on a source line that is expected to get an error message, or else specifies the source line associated with the message. If there is

no message for that line or if the text of that message is not matched by regexp then the check fails and comment is included in the FAIL message. The check does not look for the string 'error' unless it is part of regexp.

### { dg-warning regexp [comment [{ target/xfail selector } [line] ]] }

This DejaGnu directive appears on a source line that is expected to get a warning message, or else specifies the source line associated with the message. If there is no message for that line or if the text of that message is not matched by regexp then the check fails and comment is included in the FAIL message. The check does not look for the string 'warning' unless it is part of regexp.

# { dg-message regexp [comment [{ target/xfail selector } [line] ]] }

The line is expected to get a message other than an error or warning. If there is no message for that line or if the text of that message is not matched by regexp then the check fails and comment is included in the FAIL message.

# { dg-bogus regexp [comment [{ target/xfail selector } [line] ]] }

This DejaGnu directive appears on a source line that should not get a message matching regexp, or else specifies the source line associated with the bogus message. It is usually used with 'xfail' to indicate that the message is a known problem for a particular set of targets.

### { dg-line linenumvar }

This DejaGnu directive sets the variable *linenumvar* to the line number of the source line. The variable *linenumvar* can then be used in subsequent dg-error, dg-warning, dg-message and dg-bogus directives. For example:

```
int a;  /* { dg-line first_def_a } */
float a; /* { dg-error "conflicting types of" } */
/* { dg-message "previous declaration of" "" { target *-*-* } first_def_a } */
```

### { dg-excess-errors comment [{ target/xfail selector }] }

This DejaGnu directive indicates that the test is expected to fail due to compiler messages that are not handled by 'dg-error', 'dg-warning' or 'dg-bogus'. For this directive 'xfail' has the same effect as 'target'.

### { dg-prune-output regexp }

Prune messages matching regexp from the test output.

# 7.2.1.8 Verify output of the test executable

```
{ dg-output regexp [{ target/xfail selector }] }
```

This DejaGnu directive compares regexp to the combined output that the test executable writes to 'stdout' and 'stderr'.

# 7.2.1.9 Specify environment variables for a test

#### { dg-set-compiler-env-var var\_name "var\_value" }

Specify that the environment variable *var\_name* needs to be set to *var\_value* before invoking the compiler on the test file.

```
{ dg-set-target-env-var var_name "var_value" }
```

Specify that the environment variable *var\_name* needs to be set to *var\_value* before execution of the program created by the test.

# 7.2.1.10 Specify additional files for a test

```
{ dg-additional-files "filelist" }
```

Specify additional files, other than source files, that must be copied to the system where the compiler runs.

```
{ dg-additional-sources "filelist" }
```

Specify additional source files to appear in the compile line following the main test file.

### 7.2.1.11 Add checks at the end of a test

```
{ dg-final { local-directive } }
```

This DejaGnu directive is placed within a comment anywhere in the source file and is processed after the test has been compiled and run. Multiple 'dg-final' commands are processed in the order in which they appear in the source file. See Section 7.2.6 [Final Actions], page 107, for a list of directives that can be used within dg-final.

# 7.2.2 Selecting targets to which a test applies

Several test directives include *selectors* to limit the targets for which a test is run or to declare that a test is expected to fail on particular targets.

A selector is:

- one or more target triplets, possibly including wildcard characters; use '\*-\*-\*' to match any target
- a single effective-target keyword (see Section 7.2.3 [Effective-Target Keywords], page 84)
- a logical expression

Depending on the context, the selector specifies whether a test is skipped and reported as unsupported or is expected to fail. A context that allows either 'target' or 'xfail' also allows '{ target selector1 xfail selector2 }' to skip the test for targets that don't match selector1 and the test to fail for targets that match selector2.

A selector expression appears within curly braces and uses a single logical operator: one of '!', '&&', or '||'. An operand is another selector expression, an effective-target keyword, a single target triplet, or a list of target triplets within quotes or curly braces. For example:

```
{ target { ! "hppa*-*-* ia64*-*-*" } }
{ target { powerpc*-*-* && lp64 } }
{ xfail { lp64 || vect_no_align } }
```

# 7.2.3 Keywords describing target attributes

Effective-target keywords identify sets of targets that support particular functionality. They are used to limit tests to be run only for particular targets, or to specify that particular sets of targets are expected to fail some tests.

Effective-target keywords are defined in 'lib/target-supports.exp' in the GCC testsuite, with the exception of those that are documented as being local to a particular test directory. The 'effective target' takes into account all of the compiler options with which the test will be compiled, including the multilib options. By convention, keywords ending in \_nocache can also include options specified for the particular test in an earlier dg-options or dg-add-options directive.

## 7.2.3.1 Endianness

be Target uses big-endian memory order for multi-byte and multi-word data.

le Target uses little-endian memory order for multi-byte and multi-word data.

# 7.2.3.2 Data type sizes

ilp32 Target has 32-bit int, long, and pointers.

1p64 Target has 32-bit int, 64-bit long and pointers.

Target has 32-bit int and long, 64-bit long long and pointers.

double64 Target has 64-bit double.

double64plus

Target has double that is 64 bits or longer.

longdouble128

Target has 128-bit long double.

int32plus

Target has int that is at 32 bits or longer.

int16 Target has int that is 16 bits or shorter.

longlong64

Target has 64-bit long long.

long\_neq\_int

Target has int and long with different sizes.

int\_eq\_float

Target has int and float with the same size.

ptr\_eq\_long

Target has pointers (void \*) and long with the same size.

large\_double

Target supports double that is longer than float.

large\_long\_double

Target supports long double that is longer than double.

ptr32plus

Target has pointers that are 32 bits or longer.

size20plus

Target has a 20-bit or larger address space, so at least supports 16-bit array and structure sizes.

#### size32plus

Target has a 32-bit or larger address space, so at least supports 24-bit array and structure sizes.

#### 4byte\_wchar\_t

Target has wchar\_t that is at least 4 bytes.

floatn Target has the \_Floatn type.

floatnx Target has the \_Floatnx type.

#### floatn\_runtime

Target has the \_Floatn type, including runtime support for any options added with dg-add-options.

#### floatnx\_runtime

Target has the \_Floatnx type, including runtime support for any options added with dg-add-options.

#### floatn\_nx\_runtime

Target has runtime support for any options added with dg-add-options for any \_Floatn or \_Floatnx type.

inf Target supports floating point infinite (inf) for type double.

# 7.2.3.3 Fortran-specific attributes

# fortran\_integer\_16

Target supports Fortran integer that is 16 bytes or longer.

#### fortran\_real\_10

Target supports Fortran real that is 10 bytes or longer.

#### fortran\_real\_16

Target supports Fortran real that is 16 bytes or longer.

#### fortran\_large\_int

Target supports Fortran integer kinds larger than integer (8).

### fortran\_large\_real

Target supports Fortran real kinds larger than real (8).

# 7.2.3.4 Vector-specific attributes

### vect\_align\_stack\_vars

The target's ABI allows stack variables to be aligned to the preferred vector alignment.

### vect\_avg\_qi

Target supports both signed and unsigned averaging operations on vectors of bytes.

#### vect\_mulhrs\_hi

Target supports both signed and unsigned multiply-high-with-round-and-scale operations on vectors of half-words.

#### vect\_sdiv\_pow2\_si

Target supports signed division by constant power-of-2 operations on vectors of 4-byte integers.

#### vect\_condition

Target supports vector conditional operations.

#### vect\_cond\_mixed

Target supports vector conditional operations where comparison operands have different type from the value operands.

#### vect\_double

Target supports hardware vectors of double.

#### vect\_double\_cond\_arith

Target supports conditional addition, subtraction, multiplication, division, minimum and maximum on vectors of double, via the cond\_ optabs.

### vect\_element\_align\_preferred

The target's preferred vector alignment is the same as the element alignment.

### vect\_float

Target supports hardware vectors of float when '-funsafe-math-optimizations' is in effect.

#### vect\_float\_strict

Target supports hardware vectors of float when '-funsafe-math-optimizations' is not in effect. This implies vect\_float.

vect\_int Target supports hardware vectors of int.

#### vect\_long

Target supports hardware vectors of long.

#### vect\_long\_long

Target supports hardware vectors of long long.

### vect\_check\_ptrs

Target supports the check\_raw\_ptrs and check\_war\_ptrs optabs on vectors.

#### vect\_fully\_masked

Target supports fully-masked (also known as fully-predicated) loops, so that vector loops can handle partial as well as full vectors.

#### vect\_masked\_store

Target supports vector masked stores.

#### vect\_scatter\_store

Target supports vector scatter stores.

## vect\_aligned\_arrays

Target aligns arrays to vector alignment boundary.

#### vect\_hw\_misalign

Target supports a vector misalign access.

#### vect\_no\_align

Target does not support a vector alignment mechanism.

### vect\_peeling\_profitable

Target might require to peel loops for alignment purposes.

#### vect\_no\_int\_min\_max

Target does not support a vector min and max instruction on int.

#### vect\_no\_int\_add

Target does not support a vector add instruction on int.

#### vect\_no\_bitwise

Target does not support vector bitwise instructions.

#### vect\_bool\_cmp

Target supports comparison of bool vectors for at least one vector length.

#### vect\_char\_add

Target supports addition of char vectors for at least one vector length.

### vect\_char\_mult

Target supports vector char multiplication.

#### vect\_short\_mult

Target supports vector short multiplication.

#### vect\_int\_mult

Target supports vector int multiplication.

#### vect\_long\_mult

Target supports 64 bit vector long multiplication.

#### vect\_extract\_even\_odd

Target supports vector even/odd element extraction.

### vect\_extract\_even\_odd\_wide

Target supports vector even/odd element extraction of vectors with elements SImode or larger.

#### vect\_interleave

Target supports vector interleaving.

#### vect\_strided

Target supports vector interleaving and extract even/odd.

#### vect\_strided\_wide

Target supports vector interleaving and extract even/odd for wide element types.

#### vect\_perm

Target supports vector permutation.

#### vect\_perm\_byte

Target supports permutation of vectors with 8-bit elements.

# vect\_perm\_short

Target supports permutation of vectors with 16-bit elements.

#### vect\_perm3\_byte

Target supports permutation of vectors with 8-bit elements, and for the default vector length it is possible to permute:

```
{ a0, a1, a2, b0, b1, b2, ... } to:

{ a0, a0, a0, b0, b0, b0, ... } { a1, a1, a1, b1, b1, b1, ... } { a2, a2, a2, b2, b2, b2, ... }
```

using only two-vector permutes, regardless of how long the sequence is.

#### vect\_perm3\_int

Like vect\_perm3\_byte, but for 32-bit elements.

#### vect\_perm3\_short

Like vect\_perm3\_byte, but for 16-bit elements.

#### vect\_shift

Target supports a hardware vector shift operation.

### vect\_unaligned\_possible

Target prefers vectors to have an alignment greater than element alignment, but also allows unaligned vector accesses in some circumstances.

#### vect\_variable\_length

Target has variable-length vectors.

### vect\_widen\_sum\_hi\_to\_si

Target supports a vector widening summation of short operands into int results, or can promote (unpack) from short to int.

#### vect\_widen\_sum\_qi\_to\_hi

Target supports a vector widening summation of char operands into short results, or can promote (unpack) from char to short.

#### vect\_widen\_sum\_qi\_to\_si

Target supports a vector widening summation of char operands into int results.

#### vect\_widen\_mult\_qi\_to\_hi

Target supports a vector widening multiplication of char operands into short results, or can promote (unpack) from char to short and perform non-widening multiplication of short.

## vect\_widen\_mult\_hi\_to\_si

Target supports a vector widening multiplication of short operands into int results, or can promote (unpack) from short to int and perform non-widening multiplication of int.

#### vect\_widen\_mult\_si\_to\_di\_pattern

Target supports a vector widening multiplication of int operands into long results.

#### vect\_sdot\_qi

Target supports a vector dot-product of signed char.

vect\_udot\_qi

Target supports a vector dot-product of unsigned char.

vect\_sdot\_hi

Target supports a vector dot-product of signed short.

vect\_udot\_hi

Target supports a vector dot-product of unsigned short.

vect\_pack\_trunc

Target supports a vector demotion (packing) of short to char and from int to short using modulo arithmetic.

vect\_unpack

Target supports a vector promotion (unpacking) of char to short and from char to int

vect\_intfloat\_cvt

Target supports conversion from signed int to float.

vect\_uintfloat\_cvt

Target supports conversion from unsigned int to float.

vect\_floatint\_cvt

Target supports conversion from float to signed int.

vect\_floatuint\_cvt

Target supports conversion from float to unsigned int.

vect\_intdouble\_cvt

Target supports conversion from signed int to double.

vect\_doubleint\_cvt

Target supports conversion from double to signed int.

vect\_max\_reduc

Target supports max reduction for vectors.

vect\_sizes\_16B\_8B

Target supports 16- and 8-bytes vectors.

vect\_sizes\_32B\_16B

Target supports 32- and 16-bytes vectors.

vect\_logical\_reduc

Target supports AND, IOR and XOR reduction on vectors.

vect\_fold\_extract\_last

Target supports the fold\_extract\_last optab.

# 7.2.3.5 Thread Local Storage attributes

tls Target supports thread-local storage.

tls\_native

Target supports native (rather than emulated) thread-local storage.

tls\_runtime

Test system supports executing TLS executables.

# 7.2.3.6 Decimal floating point attributes

dfp Targets supports compiling decimal floating point extension to C.

dfp\_nocache

Including the options used to compile this particular test, the target supports compiling decimal floating point extension to C.

dfprt Test system can execute decimal floating point tests.

dfprt\_nocache

Including the options used to compile this particular test, the test system can execute decimal floating point tests.

hard\_dfp Target generates decimal floating point instructions with current options.

# 7.2.3.7 ARM-specific attributes

arm32 ARM target generates 32-bit code.

arm\_little\_endian

ARM target that generates little-endian code.

arm\_eabi ARM target adheres to the ABI for the ARM Architecture.

arm\_fp\_ok

ARM target defines \_\_ARM\_FP using -mfloat-abi=softfp or equivalent options. Some multilibs may be incompatible with these options.

arm\_fp\_dp\_ok

ARM target defines \_\_ARM\_FP with double-precision support using -mfloat-abi=softfp or equivalent options. Some multilibs may be incompatible with these options.

arm\_hf\_eabi

ARM target adheres to the VFP and Advanced SIMD Register Arguments variant of the ABI for the ARM Architecture (as selected with -mfloat-abi=hard).

arm\_softfloat

ARM target uses the soft-float ABI with no floating-point instructions used whatsoever (as selected with -mfloat-abi=soft).

arm\_hard\_vfp\_ok

ARM target supports -mfpu=vfp -mfloat-abi=hard. Some multilibs may be incompatible with these options.

arm\_iwmmxt\_ok

ARM target supports -mcpu=iwmmxt. Some multilibs may be incompatible with this option.

arm\_neon ARM target supports generating NEON instructions.

arm\_tune\_string\_ops\_prefer\_neon

Test CPU tune supports inlining string operations with NEON instructions.

arm\_neon\_hw

Test system supports executing NEON instructions.

#### arm\_neonv2\_hw

Test system supports executing NEON v2 instructions.

#### arm\_neon\_ok

ARM Target supports -mfpu=neon -mfloat-abi=softfp or compatible options. Some multilibs may be incompatible with these options.

#### arm\_neon\_ok\_no\_float\_abi

ARM Target supports NEON with -mfpu=neon, but without any -mfloat-abi=option. Some multilibs may be incompatible with this option.

#### arm\_neonv2\_ok

ARM Target supports -mfpu=neon-vfpv4 -mfloat-abi=softfp or compatible options. Some multilibs may be incompatible with these options.

#### arm\_fp16\_ok

Target supports options to generate VFP half-precision floating-point instructions. Some multilibs may be incompatible with these options. This test is valid for ARM only.

#### arm\_fp16\_hw

Target supports executing VFP half-precision floating-point instructions. This test is valid for ARM only.

#### arm\_neon\_fp16\_ok

ARM Target supports -mfpu=neon-fp16 -mfloat-abi=softfp or compatible options, including -mfp16-format=ieee if necessary to obtain the \_\_fp16 type. Some multilibs may be incompatible with these options.

#### arm\_neon\_fp16\_hw

Test system supports executing Neon half-precision float instructions. (Implies previous.)

#### arm\_fp16\_alternative\_ok

ARM target supports the ARM FP16 alternative format. Some multilibs may be incompatible with the options needed.

### arm\_fp16\_none\_ok

ARM target supports specifying none as the ARM FP16 format.

### arm\_thumb1\_ok

ARM target generates Thumb-1 code for -mthumb.

### arm\_thumb2\_ok

ARM target generates Thumb-2 code for -mthumb.

#### arm\_nothumb

ARM target that is not using Thumb.

## arm\_vfp\_ok

ARM target supports -mfpu=vfp -mfloat-abi=softfp. Some multilibs may be incompatible with these options.

#### arm\_vfp3\_ok

ARM target supports -mfpu=vfp3 -mfloat-abi=softfp. Some multilibs may be incompatible with these options.

#### arm\_arch\_v8a\_hard\_ok

The compiler is targeting arm\*-\*-\* and can compile and assemble code using the options -march=armv8-a -mfpu=neon-fp-armv8 -mfloat-abi=hard. This is not enough to guarantee that linking works.

#### arm\_arch\_v8a\_hard\_multilib

The compiler is targeting arm\*-\*-\* and can build programs using the options -march=armv8-a -mfpu=neon-fp-armv8 -mfloat-abi=hard. The target can also run the resulting binaries.

## arm\_v8\_vfp\_ok

ARM target supports -mfpu=fp-armv8 -mfloat-abi=softfp. Some multilibs may be incompatible with these options.

#### arm\_v8\_neon\_ok

ARM target supports -mfpu=neon-fp-armv8-mfloat-abi=softfp. Some multilibs may be incompatible with these options.

#### arm\_v8\_1a\_neon\_ok

ARM target supports options to generate ARMv8.1-A Adv.SIMD instructions. Some multilibs may be incompatible with these options.

#### arm\_v8\_1a\_neon\_hw

ARM target supports executing ARMv8.1-A Adv.SIMD instructions. Some multilibs may be incompatible with the options needed. Implies arm\_v8\_1a\_neon\_ok.

## arm\_acq\_rel

ARM target supports acquire-release instructions.

# ${\tt arm\_v8\_2a\_fp16\_scalar\_ok}$

ARM target supports options to generate instructions for ARMv8.2-A and scalar instructions from the FP16 extension. Some multilibs may be incompatible with these options.

#### arm\_v8\_2a\_fp16\_scalar\_hw

ARM target supports executing instructions for ARMv8.2-A and scalar instructions from the FP16 extension. Some multilibs may be incompatible with these options. Implies arm\_v8\_2a\_fp16\_neon\_ok.

# arm\_v8\_2a\_fp16\_neon\_ok

ARM target supports options to generate instructions from ARMv8.2-A with the FP16 extension. Some multilibs may be incompatible with these options. Implies arm\_v8\_2a\_fp16\_scalar\_ok.

## arm\_v8\_2a\_fp16\_neon\_hw

ARM target supports executing instructions from ARMv8.2-A with the FP16 extension. Some multilibs may be incompatible with these options. Implies  $arm_v8_2a_fp16_neon_ok$  and  $arm_v8_2a_fp16_scalar_hw$ .

# arm\_v8\_2a\_dotprod\_neon\_ok

ARM target supports options to generate instructions from ARMv8.2-A with the Dot Product extension. Some multilibs may be incompatible with these options.

#### arm\_v8\_2a\_dotprod\_neon\_hw

ARM target supports executing instructions from ARMv8.2-A with the Dot Product extension. Some multilibs may be incompatible with these options. Implies arm\_v8\_2a\_dotprod\_neon\_ok.

## arm\_fp16fml\_neon\_ok

ARM target supports extensions to generate the VFMAL and VFMLS half-precision floating-point instructions available from ARMv8.2-A and onwards. Some multilibs may be incompatible with these options.

# arm\_v8\_2a\_bf16\_neon\_ok

ARM target supports options to generate instructions from ARMv8.2-A with the BFloat16 extension (bf16). Some multilibs may be incompatible with these options.

# arm\_v8\_2a\_i8mm\_ok

ARM target supports options to generate instructions from ARMv8.2-A with the 8-Bit Integer Matrix Multiply extension (i8mm). Some multilibs may be incompatible with these options.

## arm\_v8\_1m\_mve\_ok

ARM target supports options to generate instructions from ARMv8.1-M with the M-Profile Vector Extension (MVE). Some multilibs may be incompatible with these options.

# arm\_v8\_1m\_mve\_fp\_ok

ARM target supports options to generate instructions from ARMv8.1-M with the Half-precision floating-point instructions (HP), Floating-point Extension (FP) along with M-Profile Vector Extension (MVE). Some multilibs may be incompatible with these options.

#### arm\_mve\_hw

Test system supports executing MVE instructions.

# arm\_v8m\_main\_cde

ARM target supports options to generate instructions from ARMv8-M with the Custom Datapath Extension (CDE). Some multilibs may be incompatible with these options.

#### arm\_v8m\_main\_cde\_fp

ARM target supports options to generate instructions from ARMv8-M with the Custom Datapath Extension (CDE) and floating-point (VFP). Some multilibs may be incompatible with these options.

# arm\_v8\_1m\_main\_cde\_mve

ARM target supports options to generate instructions from ARMv8.1-M with the Custom Datapath Extension (CDE) and M-Profile Vector Extension (MVE). Some multilibs may be incompatible with these options.

#### arm\_prefer\_ldrd\_strd

ARM target prefers LDRD and STRD instructions over LDM and STM instructions.

#### arm\_thumb1\_movt\_ok

ARM target generates Thumb-1 code for -mthumb with MOVW and MOVT instructions available.

# arm\_thumb1\_cbz\_ok

ARM target generates Thumb-1 code for -mthumb with CBZ and CBNZ instructions available.

#### arm\_divmod\_simode

ARM target for which divmod transform is disabled, if it supports hardware div instruction.

## arm\_cmse\_ok

ARM target supports ARMv8-M Security Extensions, enabled by the -mcmse option.

## arm\_cmse\_hw

Test system supports executing CMSE instructions.

## arm\_coproc1\_ok

ARM target supports the following coprocessor instructions: CDP, LDC, STC, MCR and MRC.

## arm\_coproc2\_ok

ARM target supports all the coprocessor instructions also listed as supported in [arm\_coproc1\_ok], page 95 in addition to the following: CDP2, LDC2, LDC21, STC2, STC21, MCR2 and MRC2.

#### arm\_coproc3\_ok

ARM target supports all the coprocessor instructions also listed as supported in [arm\_coproc2\_ok], page 95 in addition the following: MCRR and MRRC.

## arm\_coproc4\_ok

ARM target supports all the coprocessor instructions also listed as supported in [arm\_coproc3\_ok], page 95 in addition the following: MCRR2 and MRRC2.

# arm\_simd32\_ok

ARM Target supports options suitable for accessing the SIMD32 intrinsics from arm\_acle.h. Some multilibs may be incompatible with these options.

#### arm\_qbit\_ok

ARM Target supports options suitable for accessing the Q-bit manipulation intrinsics from arm\_acle.h. Some multilibs may be incompatible with these options.

#### arm\_dsp\_ok

ARM Target supports options suitable for accessing the DSP intrinsics from arm\_acle.h. Some multilibs may be incompatible with these options.

#### arm\_softfp\_ok

ARM target supports the -mfloat-abi=softfp option.

#### arm\_hard\_ok

ARM target supports the -mfloat-abi=hard option.

# 7.2.3.8 AArch64-specific attributes

## aarch64\_asm\_<ext>\_ok

AArch64 assembler supports the architecture extension ext via the .arch\_extension pseudo-op.

#### aarch64\_tiny

AArch64 target which generates instruction sequences for tiny memory model.

#### aarch64 small

AArch64 target which generates instruction sequences for small memory model.

#### aarch64\_large

AArch64 target which generates instruction sequences for large memory model.

#### aarch64\_little\_endian

AArch64 target which generates instruction sequences for little endian.

#### aarch64\_big\_endian

AArch64 target which generates instruction sequences for big endian.

# aarch64\_small\_fpic

Binutils installed on test system supports relocation types required by -fpic for AArch64 small memory model.

## aarch64\_sve\_hw

AArch64 target that is able to generate and execute SVE code (regardless of whether it does so by default).

aarch64\_sve128\_hw

aarch64\_sve256\_hw

aarch64\_sve512\_hw

aarch64\_sve1024\_hw

aarch64\_sve2048\_hw

Like aarch64\_sve\_hw, but also test for an exact hardware vector length.

#### aarch64\_fjcvtzs\_hw

AArch64 target that is able to generate and execute armv8.3-a FJCVTZS instruction.

# 7.2.3.9 MIPS-specific attributes

mips64 MIPS target supports 64-bit instructions.

nomips16 MIPS target does not produce MIPS16 code.

# mips16\_attribute

MIPS target can generate MIPS16 code.

#### mips\_loongson

MIPS target is a Loongson-2E or -2F target using an ABI that supports the Loongson vector modes.

mips\_msa MIPS target supports -mmsa, MIPS SIMD Architecture (MSA).

mips\_newabi\_large\_long\_double

MIPS target supports long double larger than double when using the new ABI.

mpaired\_single

MIPS target supports -mpaired-single.

# 7.2.3.10 PowerPC-specific attributes

dfp\_hw PowerPC target supports executing hardware DFP instructions.

p8vector\_hw

PowerPC target supports executing VSX instructions (ISA 2.07).

powerpc64

Test system supports executing 64-bit instructions.

powerpc\_altivec

PowerPC target supports AltiVec.

powerpc\_altivec\_ok

PowerPC target supports -maltivec.

powerpc\_eabi\_ok

PowerPC target supports -meabi.

powerpc\_elfv2

PowerPC target supports -mabi=elfv2.

powerpc\_fprs

PowerPC target supports floating-point registers.

powerpc\_hard\_double

PowerPC target supports hardware double-precision floating-point.

powerpc\_htm\_ok

PowerPC target supports -mhtm

powerpc\_p8vector\_ok

PowerPC target supports -mpower8-vector

powerpc\_popcntb\_ok

PowerPC target supports the popcntb instruction, indicating that this target supports -mcpu=power5.

powerpc\_ppu\_ok

PowerPC target supports -mcpu=cell.

powerpc\_spe

PowerPC target supports PowerPC SPE.

powerpc\_spe\_nocache

Including the options used to compile this particular test, the PowerPC target supports PowerPC SPE.

powerpc\_spu

PowerPC target supports PowerPC SPU.

powerpc\_vsx\_ok

PowerPC target supports -mvsx.

## powerpc\_405\_nocache

Including the options used to compile this particular test, the PowerPC target supports PowerPC 405.

# ppc\_recip\_hw

PowerPC target supports executing reciprocal estimate instructions.

vmx\_hw PowerPC target supports executing AltiVec instructions.

vsx\_hw PowerPC target supports executing VSX instructions (ISA 2.06).

# 7.2.3.11 Other hardware attributes

#### autoincdec

Target supports autoincrement/decrement addressing.

avx Target supports compiling avx instructions.

#### avx\_runtime

Target supports the execution of avx instructions.

avx2 Target supports compiling avx2 instructions.

#### avx2\_runtime

Target supports the execution of avx2 instructions.

avx512f Target supports compiling avx512f instructions.

## avx512f\_runtime

Target supports the execution of avx512f instructions.

## avx512vp2intersect

Target supports the execution of avx512vp2intersect instructions.

#### coldfire\_fpu

Target uses a ColdFire FPU.

divmod Target supporting hardware divmod insn or divmod libcall.

## divmod\_simode

Target supporting hardware divmod inso or divmod libcall for SImode.

## hard\_float

Target supports FPU instructions.

#### non\_strict\_align

Target does not require strict alignment.

## pie\_copyreloc

The x86-64 target linker supports PIE with copy reloc.

rdrand Target supports x86 rdrand instruction.

#### sqrt\_insn

Target has a square root instruction that the compiler can generate.

sse Target supports compiling sse instructions.

#### sse\_runtime

Target supports the execution of sse instructions.

sse2 Target supports compiling sse2 instructions.

#### sse2\_runtime

Target supports the execution of sse2 instructions.

#### sync\_char\_short

Target supports atomic operations on char and short.

# sync\_int\_long

Target supports atomic operations on int and long.

# ultrasparc\_hw

Test environment appears to run executables on a simulator that accepts only EM\_SPARC executables and chokes on EM\_SPARC32PLUS or EM\_SPARCV9 executables.

## vect\_cmdline\_needed

Target requires a command line argument to enable a SIMD instruction set.

xorsign Target supports the xorsign optab expansion.

# 7.2.3.12 Environment attributes

c The language for the compiler under test is C.

c++ The language for the compiler under test is C++.

#### c99\_runtime

Target provides a full C99 runtime.

# correct\_iso\_cpp\_string\_wchar\_protos

Target string.h and wchar.h headers provide C++ required overloads for strchr etc. functions.

#### d\_runtime

Target provides the D runtime.

## d\_runtime\_has\_std\_library

Target provides the D standard library (Phobos).

#### dummy\_wcsftime

Target uses a dummy wcsftime function that always returns zero.

#### fd\_truncate

Target can truncate a file from a file descriptor, as used by 'libgfortran/io/unix.c:fd\_truncate'; i.e. ftruncate or chsize.

fenv Target provides 'fenv.h' include file.

#### fenv\_exceptions

Target supports 'fenv.h' with all the standard IEEE exceptions and floating-point exceptions are raised by arithmetic operations.

Target offers such file I/O library functions as fopen, fclose, tmpnam, and remove. This is a link-time requirement for the presence of the functions in the library; even if they fail at runtime, the requirement is still regarded as satisfied.

#### freestanding

Target is 'freestanding' as defined in section 4 of the C99 standard. Effectively, it is a target which supports no extra headers or libraries other than what is considered essential.

#### gettimeofday

Target supports gettimeofday.

#### init\_priority

Target supports constructors with initialization priority arguments.

## inttypes\_types

Target has the basic signed and unsigned types in inttypes.h. This is for tests that GCC's notions of these types agree with those in the header, as some systems have only inttypes.h.

#### lax\_strtofp

Target might have errors of a few ULP in string to floating-point conversion functions and overflow is not always detected correctly by those functions.

mempcpy Target provides mempcpy function.

mmap Target supports mmap.

newlib Target supports Newlib.

#### newlib\_nano\_io

GCC was configured with --enable-newlib-nano-formatted-io, which reduces the code size of Newlib formatted I/O functions.

pow10 Target provides pow10 function.

pthread Target can compile using pthread.h with no errors or warnings.

# pthread\_h

Target has pthread.h.

#### run\_expensive\_tests

Expensive testcases (usually those that consume excessive amounts of CPU time) should be run on this target. This can be enabled by setting the GCC\_TEST\_RUN\_EXPENSIVE environment variable to a non-empty string.

## simulator

Test system runs executables on a simulator (i.e. slowly) rather than hardware (i.e. fast).

signal Target has signal.h.

stabs Target supports the stabs debugging format.

stdint\_types

Target has the basic signed and unsigned C types in stdint.h. This will be obsolete when GCC ensures a working stdint.h for all targets.

stpcpy Target provides stpcpy function.

trampolines

Target supports trampolines.

uclibc Target supports uClibc.

unwrapped

Target does not use a status wrapper.

vxworks\_kernel

Target is a VxWorks kernel.

vxworks\_rtp

Target is a VxWorks RTP.

wchar Target supports wide characters.

# 7.2.3.13 Other attributes

automatic\_stack\_alignment

Target supports automatic stack alignment.

branch\_cost

Target supports '-branch-cost=N'.

cxa\_atexit

Target uses \_\_cxa\_atexit.

default\_packed

Target has packed layout of structure members by default.

exceptions

Target supports exceptions.

 ${\tt exceptions\_enabled}$ 

Target supports exceptions and they are enabled in the current testing configuration.

fgraphite

Target supports Graphite optimizations.

fixed\_point

Target supports fixed-point extension to C.

fopenacc Target supports OpenACC via '-fopenacc'.

fopenmp Target supports OpenMP via '-fopenmp'.

fpic Target supports '-fpic' and '-fPIC'.

freorder Target supports '-freorder-blocks-and-partition'.

#### fstack\_protector

Target supports '-fstack-protector'.

gas Target uses GNU as.

#### gc\_sections

Target supports '--gc-sections'.

gld Target uses GNU ld.

## keeps\_null\_pointer\_checks

Target keeps null pointer checks, either due to the use of '-fno-delete-null-pointer-checks' or hardwired into the target.

#### llvm\_binutils

Target is using an LLVM assembler and/or linker, instead of GNU Binutils.

1to Compiler has been configured to support link-time optimization (LTO).

#### lto\_incremental

Compiler and linker support link-time optimization relocatable linking with '-r' and '-flto' options.

#### naked\_functions

Target supports the naked function attribute.

#### named\_sections

Target supports named sections.

#### natural\_alignment\_32

Target uses natural alignment (aligned to type size) for types of 32 bits or less.

## target\_natural\_alignment\_64

Target uses natural alignment (aligned to type size) for types of 64 bits or less.

noinit Target supports the noinit variable attribute.

nonpic Target does not generate PIC by default.

#### offload\_gcn

Target has been configured for OpenACC/OpenMP offloading on AMD GCN.

#### pie\_enabled

Target generates PIE by default.

#### pcc\_bitfield\_type\_matters

Target defines PCC\_BITFIELD\_TYPE\_MATTERS.

# pe\_aligned\_commons

Target supports '-mpe-aligned-commons'.

pie Target supports '-pie', '-fpie' and '-fPIE'.

rdynamic Target supports '-rdynamic'.

#### scalar\_all\_fma

Target supports all four fused multiply-add optabs for both float and double. These optabs are: fma\_optab, fms\_optab, fnma\_optab and fnms\_optab.

#### section\_anchors

Target supports section anchors.

#### short\_enums

Target defaults to short enums.

#### stack\_size

Target has limited stack size. The stack size limit can be obtained using the STACK\_SIZE macro defined by [dg-add-options feature stack\_size], page 105.

static Target supports '-static'.

# static\_libgfortran

Target supports statically linking 'libgfortran'.

# string\_merging

Target supports merging string constants at link time.

ucn Target supports compiling and assembling UCN.

#### ucn\_nocache

Including the options used to compile this particular test, the target supports compiling and assembling UCN.

## unaligned\_stack

Target does not guarantee that its STACK\_BOUNDARY is greater than or equal to the required vector alignment.

#### vector\_alignment\_reachable

Vector alignment is reachable for types of 32 bits or less.

## vector\_alignment\_reachable\_for\_64bit

Vector alignment is reachable for types of 64 bits or less.

# wchar\_t\_char16\_t\_compatible

Target supports wchar\_t that is compatible with char16\_t.

## wchar\_t\_char32\_t\_compatible

Target supports wchar\_t that is compatible with char32\_t.

## comdat\_group

Target uses comdat groups.

#### indirect\_calls

Target supports indirect calls, i.e. calls where the target is not constant.

# 7.2.3.14 Local to tests in gcc.target/i386

3dnow Target supports compiling 3dnow instructions.

aes Target supports compiling aes instructions.

fma4 Target supports compiling fma4 instructions.

mfentry Target supports the -mfentry option that alters the position of profiling calls such that they precede the prologue.

## ms\_hook\_prologue

Target supports attribute ms\_hook\_prologue.

pclmul Target supports compiling pclmul instructions.

sse3 Target supports compiling sse3 instructions.

sse4 Target supports compiling sse4 instructions.

sse4a Target supports compiling sse4a instructions.

ssse3 Target supports compiling ssse3 instructions.

vaes Target supports compiling vaes instructions.

vpclmul Target supports compiling vpclmul instructions.

xop Target supports compiling xop instructions.

# 7.2.3.15 Local to tests in gcc.test-framework

no Always returns 0.

yes Always returns 1.

# 7.2.4 Features for dg-add-options

The supported values of feature for directive dg-add-options are:

arm\_fp \_\_ARM\_FP definition. Only ARM targets support this feature, and only then in certain modes; see the [arm\_fp\_ok effective target keyword], page 91.

#### arm\_fp\_dp

\_\_ARM\_FP definition with double-precision support. Only ARM targets support this feature, and only then in certain modes; see the [arm\_fp\_dp\_ok effective target keyword], page 91.

arm\_neon NEON support. Only ARM targets support this feature, and only then in certain modes; see the [arm\_neon\_ok effective target keyword], page 92.

arm\_fp16 VFP half-precision floating point support. This does not select the FP16 format; for that, use [arm\_fp16\_ieee], page 104 or [arm\_fp16\_alternative], page 104 instead. This feature is only supported by ARM targets and then only in certain modes; see the [arm\_fp16\_ok effective target keyword], page 92.

# arm\_fp16\_ieee

ARM IEEE 754-2008 format VFP half-precision floating point support. This feature is only supported by ARM targets and then only in certain modes; see the [arm\_fp16\_ok effective target keyword], page 92.

#### arm\_fp16\_alternative

ARM Alternative format VFP half-precision floating point support. This feature is only supported by ARM targets and then only in certain modes; see the [arm\_fp16\_ok effective target keyword], page 92.

# arm\_neon\_fp16

NEON and half-precision floating point support. Only ARM targets support this feature, and only then in certain modes; see the [arm\_neon\_fp16\_ok effective target keyword], page 92.

arm\_vfp3 arm vfp3 floating point support; see the [arm\_vfp3\_ok effective target keyword], page 92.

#### arm\_arch\_v8a\_hard

Add options for ARMv8-A and the hard-float variant of the AAPCS, if this is supported by the compiler; see the [arm\_arch\_v8a\_hard\_ok], page 93 effective target keyword.

#### arm\_v8\_1a\_neon

Add options for ARMv8.1-A with Adv.SIMD support, if this is supported by the target; see the [arm\_v8\_1a\_neon\_ok], page 93 effective target keyword.

# arm\_v8\_2a\_fp16\_scalar

Add options for ARMv8.2-A with scalar FP16 support, if this is supported by the target; see the [arm\_v8\_2a\_fp16\_scalar\_ok], page 93 effective target keyword.

# arm\_v8\_2a\_fp16\_neon

Add options for ARMv8.2-A with Adv.SIMD FP16 support, if this is supported by the target; see the [arm\_v8\_2a\_fp16\_neon\_ok], page 93 effective target keyword.

#### arm\_v8\_2a\_dotprod\_neon

Add options for ARMv8.2-A with Adv.SIMD Dot Product support, if this is supported by the target; see the [arm\_v8\_2a\_dotprod\_neon\_ok], page 93 effective target keyword.

#### arm\_fp16fml\_neon

Add options to enable generation of the VFMAL and VFMSL instructions, if this is supported by the target; see the [arm\_fp16fml\_neon\_ok], page 94 effective target keyword.

arm\_dsp Add options for ARM DSP intrinsics support, if this is supported by the target; see the [arm\_dsp\_ok effective target keyword], page 95.

#### bind\_pic\_locally

Add the target-specific flags needed to enable functions to bind locally when using pic/PIC passes in the testsuite.

floatn Add the target-specific flags needed to use the \_Floatn type.

floatnx Add the target-specific flags needed to use the \_Floatnx type.

ieee Add the target-specific flags needed to enable full IEEE compliance mode.

# mips16\_attribute

mips16 function attributes. Only MIPS targets support this feature, and only then in certain modes.

#### stack\_size

Add the flags needed to define macro STACK\_SIZE and set it to the stack size limit associated with the [stack\_size effective target], page 103.

## sqrt\_insn

Add the target-specific flags needed to enable hardware square root instructions, if any.

tls Add the target-specific flags needed to use thread-local storage.

# 7.2.5 Variants of dg-require-support

A few of the dg-require directives take arguments.

# dg-require-iconv codeset

Skip the test if the target does not support iconv. *codeset* is the codeset to convert to.

#### dg-require-profiling profopt

Skip the test if the target does not support profiling with option profopt.

#### dg-require-stack-check check

Skip the test if the target does not support the -fstack-check option. If check is "", support for -fstack-check is checked, for -fstack-check=("check") otherwise.

## dg-require-stack-size size

Skip the test if the target does not support a stack size of size.

#### dg-require-visibility vis

Skip the test if the target does not support the visibility attribute. If vis is "", support for visibility("hidden") is checked, for visibility("vis") otherwise.

The original dg-require directives were defined before there was support for effective-target keywords. The directives that do not take arguments could be replaced with effective-target keywords.

# dg-require-alias ""

Skip the test if the target does not support the 'alias' attribute.

# dg-require-ascii-locale ""

Skip the test if the host does not support an ASCII locale.

# dg-require-compat-dfp ""

Skip this test unless both compilers in a 'compat' testsuite support decimal floating point.

# dg-require-cxa-atexit ""

Skip the test if the target does not support \_\_cxa\_atexit. This is equivalent to dg-require-effective-target cxa\_atexit.

## dg-require-dll ""

Skip the test if the target does not support DLL attributes.

## dg-require-dot ""

Skip the test if the host does not have dot.

# dg-require-fork ""

Skip the test if the target does not support fork.

# dg-require-gc-sections ""

Skip the test if the target's linker does not support the --gc-sections flags. This is equivalent to dg-require-effective-target gc-sections.

## dg-require-host-local ""

Skip the test if the host is remote, rather than the same as the build system. Some tests are incompatible with DejaGnu's handling of remote hosts, which involves copying the source file to the host and compiling it with a relative path and "-o a.out".

## dg-require-mkfifo ""

Skip the test if the target does not support mkfifo.

# dg-require-named-sections ""

Skip the test is the target does not support named sections. This is equivalent to dg-require-effective-target named\_sections.

# dg-require-weak ""

Skip the test if the target does not support weak symbols.

# dg-require-weak-override ""

Skip the test if the target does not support overriding weak symbols.

# 7.2.6 Commands for use in dg-final

The GCC testsuite defines the following directives to be used within dg-final.

# 7.2.6.1 Scan a particular file

# scan-file filename regexp [{ target/xfail selector }]

Passes if regexp matches text in filename.

## scan-file-not filename regexp [{ target/xfail selector }]

Passes if regexp does not match text in filename.

## scan-module module regexp [{ target/xfail selector }]

Passes if regexp matches in Fortran module module.

#### dg-check-dot filename

Passes if *filename* is a valid '.dot' file (by running dot -Tpng on it, and verifying the exit code is 0).

# 7.2.6.2 Scan the assembly output

## scan-assembler regex [{ target/xfail selector }]

Passes if regex matches text in the test's assembler output.

## scan-assembler-not regex [{ target/xfail selector }]

Passes if regex does not match text in the test's assembler output.

## scan-assembler-times regex num [{ target/xfail selector }]

Passes if regex is matched exactly num times in the test's assembler output.

## scan-assembler-dem regex [{ target/xfail selector }]

Passes if regex matches text in the test's demangled assembler output.

#### scan-assembler-dem-not regex [{ target/xfail selector }]

Passes if regex does not match text in the test's demangled assembler output.

## scan-hidden symbol [{ target/xfail selector }]

Passes if symbol is defined as a hidden symbol in the test's assembly output.

```
scan-not-hidden symbol [{ target/xfail selector }]
```

Passes if symbol is not defined as a hidden symbol in the test's assembly output.

#### check-function-bodies prefix terminator [options [{ target/xfail selector }]]

Looks through the source file for comments that give the expected assembly output for selected functions. Each line of expected output starts with the prefix string *prefix* and the expected output for a function as a whole is followed by a line that starts with the string *terminator*. Specifying an empty terminator is equivalent to specifying '"\*/".

options, if specified, is a list of regular expressions, each of which matches a full command-line option. A non-empty list prevents the test from running unless all of the given options are present on the command line. This can help if a source file is compiled both with and without optimization, since it is rarely useful to check the full function body for unoptimized code.

The first line of the expected output for a function fn has the form:

```
prefix fn: [{ target/xfail selector }]
```

Subsequent lines of the expected output also start with *prefix*. In both cases, whitespace after *prefix* is not significant.

The test discards assembly directives such as .cfi\_startproc and local label definitions such as .LFBO from the compiler's assembly output. It then matches the result against the expected output for a function as a single regular expression. This means that later lines can use backslashes to refer back to '(...)' captures on earlier lines. For example:

```
/* { dg-final { check-function-bodies "**" "" "-DCHECK_ASM" } } */
...
/*

** add_w0_s8_m:

** mov (z[0-9]+\.b), w0

** add z0\.b, p0/m, z0\.b, \1

** ret

*/

svint8_t add_w0_s8_m (...) { ... }

...
/*

** add_b0_s8_m:

** mov (z[0-9]+\.b), b0

** add z1\.b, p0/m, z1\.b, \1

** ret

*/

svint8_t add_b0_s8_m (...) { ... }
```

checks whether the implementations of add\_w0\_s8\_m and add\_b0\_s8\_m match the regular expressions given. The test only runs when '-DCHECK\_ASM' is passed on the command line.

It is possible to create non-capturing multi-line regular expression groups of the form '(a|b|...)' by putting the '(', '|' and ')' on separate lines (each still using prefix). For example:

/\*

```
** cmple_f16_tied:
** (
** fcmge p0\.h, p0/z, z1\.h, z0\.h
** |
** fcmle p0\.h, p0/z, z0\.h, z1\.h
** )
** ret
*/
svbool_t cmple_f16_tied (...) { ... }
```

checks whether cmple\_f16\_tied is implemented by the fcmge instruction followed by ret or by the fcmle instruction followed by ret. The test is still a single regular rexpression.

A line containing just:

```
prefix ...
```

stands for zero or more unmatched lines; the whitespace after *prefix* is again not significant.

# 7.2.6.3 Scan optimization dump files

These commands are available for kind of tree, ltrans-tree, offload-tree, rtl, offload-rtl, ipa, and wpa-ipa.

```
scan-kind-dump regex suffix [{ target/xfail selector }]
```

Passes if regex matches text in the dump file with suffix suffix.

```
scan-kind-dump-not regex suffix [{ target/xfail selector }]
```

Passes if regex does not match text in the dump file with suffix suffix.

```
scan-kind-dump-times regex num suffix [{ target/xfail selector }]
```

Passes if regex is found exactly num times in the dump file with suffix suffix.

```
scan-kind-dump-dem regex suffix [{ target/xfail selector }]
```

Passes if regex matches demangled text in the dump file with suffix suffix.

```
scan-kind-dump-dem-not regex suffix [{ target/xfail selector }]
```

Passes if regex does not match demangled text in the dump file with suffix suffix.

# 7.2.6.4 Check for output files

```
output-exists [{ target/xfail selector }]
```

Passes if compiler output file exists.

```
output-exists-not [{ target/xfail selector }]
```

Passes if compiler output file does not exist.

```
scan-symbol regexp [{ target/xfail selector }]
```

Passes if the pattern is present in the final executable.

```
scan-symbol-not regexp [{ target/xfail selector }]
```

Passes if the pattern is absent from the final executable.

# 7.2.6.5 Checks for gcov tests

```
run-gcov sourcefile
```

Check line counts in gcov tests.

```
run-gcov [branches] [calls] { opts sourcefile }
```

Check branch and/or call counts, in addition to line counts, in gcov tests.

# 7.2.6.6 Clean up generated test files

Usually the test-framework removes files that were generated during testing. If a testcase, for example, uses any dumping mechanism to inspect a passes dump file, the testsuite recognized the dump option passed to the tool and schedules a final cleanup to remove these files.

There are, however, following additional cleanup directives that can be used to annotate a testcase "manually".

## cleanup-coverage-files

Removes coverage data files generated for this test.

## cleanup-modules "list-of-extra-modules"

Removes Fortran module files generated for this test, excluding the module names listed in keep-modules. Cleaning up module files is usually done automatically by the testsuite by looking at the source files and removing the modules after the test has been executed.

```
module MoD1
end module MoD1
module Mod2
end module Mod2
module moD3
end module moD3
module mod4
end module mod4
! { dg-final { cleanup-modules "mod1 mod2" } } ! redundant
! { dg-final { keep-modules "mod3 mod4" } }
```

#### keep-modules "list-of-modules-not-to-delete"

Whitespace separated list of module names that should not be deleted by cleanup-modules. If the list of modules is empty, all modules defined in this file are kept.

```
module maybe_unneeded
end module maybe_unneeded
module keep1
end module keep1
module keep2
end module keep2
! { dg-final { keep-modules "keep1 keep2" } } ! just keep these two
! { dg-final { keep-modules "" } } ! keep all
```

## dg-keep-saved-temps "list-of-suffixes-not-to-delete"

Whitespace separated list of suffixes that should not be deleted automatically in a testcase that uses '-save-temps'.

```
// { dg-options "-save-temps -fpch-preprocess -I." }
int main() { return 0; }
```

```
// { dg-keep-saved-temps ".s" } ! just keep assembler file
// { dg-keep-saved-temps ".s" ".i" } ! ... and .i
// { dg-keep-saved-temps ".ii" ".o" } ! or just .ii and .o
```

cleanup-profile-file

Removes profiling files generated for this test.

# 7.3 Ada Language Testsuites

The Ada testsuite includes executable tests from the ACATS testsuite, publicly available at http://www.ada-auth.org/acats.html.

These tests are integrated in the GCC testsuite in the 'ada/acats' directory, and enabled automatically when running make check, assuming the Ada language has been enabled when configuring GCC.

You can also run the Ada testsuite independently, using make check-ada, or run a subset of the tests by specifying which chapter to run, e.g.:

```
$ make check-ada CHAPTERS="c3 c9"
```

The tests are organized by directory, each directory corresponding to a chapter of the Ada Reference Manual. So for example, 'c9' corresponds to chapter 9, which deals with tasking features of the language.

The tests are run using two sh scripts: 'run\_acats' and 'run\_all.sh'. To run the tests using a simulator or a cross target, see the small customization section at the top of 'run\_all.sh'.

These tests are run using the build tree: they can be run without doing a make install.

# 7.4 C Language Testsuites

GCC contains the following C language testsuites, in the 'gcc/testsuite' directory:

'gcc.dg' This contains tests of particular features of the C compiler, using the more modern 'dg' harness. Correctness tests for various compiler features should go here if possible.

Magic comments determine whether the file is preprocessed, compiled, linked or run. In these tests, error and warning message texts are compared against expected texts or regular expressions given in comments. These tests are run with the options '-ansi -pedantic' unless other options are given in the test. Except as noted below they are not run with multiple optimization options.

'gcc.dg/compat'

This subdirectory contains tests for binary compatibility using 'lib/compat.exp', which in turn uses the language-independent support (see Section 7.8 [Support for testing binary compatibility], page 115).

```
'gcc.dg/cpp'
```

This subdirectory contains tests of the preprocessor.

#### 'gcc.dg/debug'

This subdirectory contains tests for debug formats. Tests in this subdirectory are run for each debug format that the compiler supports.

#### 'gcc.dg/format'

This subdirectory contains tests of the '-Wformat' format checking. Tests in this directory are run with and without '-DWIDE'.

#### 'gcc.dg/noncompile'

This subdirectory contains tests of code that should not compile and does not need any special compilation options. They are run with multiple optimization options, since sometimes invalid code crashes the compiler with optimization.

## 'gcc.dg/special'

FIXME: describe this.

# 'gcc.c-torture'

This contains particular code fragments which have historically broken easily. These tests are run with multiple optimization options, so tests for features which only break at some optimization levels belong here. This also contains tests to check that certain optimizations occur. It might be worthwhile to separate the correctness tests cleanly from the code quality tests, but it hasn't been done yet.

# 'gcc.c-torture/compat'

FIXME: describe this.

This directory should probably not be used for new tests.

#### 'gcc.c-torture/compile'

This testsuite contains test cases that should compile, but do not need to link or run. These test cases are compiled with several different combinations of optimization options. All warnings are disabled for these test cases, so this directory is not suitable if you wish to test for the presence or absence of compiler warnings. While special options can be set, and tests disabled on specific platforms, by the use of '.x' files, mostly these test cases should not contain platform dependencies. FIXME: discuss how defines such as STACK\_SIZE are used.

#### 'gcc.c-torture/execute'

This testsuite contains test cases that should compile, link and run; otherwise the same comments as for 'gcc.c-torture/compile' apply.

#### 'gcc.c-torture/execute/ieee'

This contains tests which are specific to IEEE floating point.

# 'gcc.c-torture/unsorted'

FIXME: describe this.

This directory should probably not be used for new tests.

## 'gcc.misc-tests'

This directory contains C tests that require special handling. Some of these tests have individual expect files, and others share special-purpose expect files:

#### 'bprob\*.c'

Test '-fbranch-probabilities' using 'gcc.misc-tests/bprob.exp',
which in turn uses the generic, language-independent framework

(see Section 7.7 [Support for testing profile-directed optimizations], page 114).

'gcov\*.c' Test gcov output using 'gcov.exp', which in turn uses the language-independent support (see Section 7.6 [Support for testing gcov], page 114).

'i386-pf-\*.c'

Test i386-specific support for data prefetch using 'i386-prefetch.exp'.

'gcc.test-framework'

'dg-\*.c' Test the testsuite itself using 'gcc.test-framework/test-framework.exp'.

FIXME: merge in 'testsuite/README.gcc' and discuss the format of test cases and magic comments more.

# 7.5 Support for testing link-time optimizations

Tests for link-time optimizations usually require multiple source files that are compiled separately, perhaps with different sets of options. There are several special-purpose test directives used for these tests.

# { dg-lto-do do-what-keyword }

do-what-keyword specifies how the test is compiled and whether it is executed. It is one of:

assemble Compile with '-c' to produce a relocatable object file.

link Compile, assemble, and link to produce an executable file.

run Produce and run an executable file, which is expected to return an exit code of 0.

The default is assemble. That can be overridden for a set of tests by redefining dg-do-what-default within the .exp file for those tests.

Unlike dg-do, dg-lto-do does not support an optional 'target' or 'xfail' list. Use dg-skip-if, dg-xfail-if, or dg-xfail-run-if.

# { dg-lto-options { { options } [{ options }] } [{ target selector }]}

This directive provides a list of one or more sets of compiler options to override *LTO\_OPTIONS*. Each test will be compiled and run with each of these sets of options.

#### { dg-extra-ld-options options [{ target selector }]}

This directive adds options to the linker options used.

#### { dg-suppress-ld-options options [{ target selector }]}

This directive removes options from the set of linker options used.

# 7.6 Support for testing gcov

Language-independent support for testing gcov, and for checking that branch profiling produces expected values, is provided by the expect file 'lib/gcov.exp'. gcov tests also rely on procedures in 'lib/gcc-dg.exp' to compile and run the test program. A typical gcov test contains the following DejaGnu commands within comments:

```
{ dg-options "--coverage" }
{ dg-do run { target native } }
{ dg-final { run-gcov sourcefile } }
```

Checks of gcov output can include line counts, branch percentages, and call return percentages. All of these checks are requested via commands that appear in comments in the test's source file. Commands to check line counts are processed by default. Commands to check branch percentages and call return percentages are processed if the run-gcov command has arguments branches or calls, respectively. For example, the following specifies checking both, as well as passing '-b' to gcov:

```
{ dg-final { run-gcov branches calls { -b sourcefile } } }
```

A line count command appears within a comment on the source line that is expected to get the specified count and has the form count(cnt). A test should only check line counts for lines that will get the same count for any architecture.

Commands to check branch percentages (branch) and call return percentages (returns) are very similar to each other. A beginning command appears on or before the first of a range of lines that will report the percentage, and the ending command follows that range of lines. The beginning command can include a list of percentages, all of which are expected to be found within the range. A range is terminated by the next command of the same kind. A command branch(end) or returns(end) marks the end of a range without starting a new one. For example:

For a call return percentage, the value specified is the percentage of calls reported to return. For a branch percentage, the value is either the expected percentage or 100 minus that value, since the direction of a branch can differ depending on the target or the optimization level.

Not all branches and calls need to be checked. A test should not check for branches that might be optimized away or replaced with predicated instructions. Don't check for calls inserted by the compiler or ones that might be inlined or optimized away.

A single test can check for combinations of line counts, branch percentages, and call return percentages. The command to check a line count must appear on the line that will report that count, but commands to check branch percentages and call return percentages can bracket the lines that report them.

# 7.7 Support for testing profile-directed optimizations

The file 'profopt.exp' provides language-independent support for checking correct execution of a test built with profile-directed optimization. This testing requires that a test program be built and executed twice. The first time it is compiled to generate profile data, and the second time it is compiled to use the data that was generated during the first execution. The second execution is to verify that the test produces the expected results.

To check that the optimization actually generated better code, a test can be built and run a third time with normal optimizations to verify that the performance is better with the profile-directed optimizations. 'profopt.exp' has the beginnings of this kind of support.

'profopt.exp' provides generic support for profile-directed optimizations. Each set of tests that uses it provides information about a specific optimization:

tool tool being tested, e.g., gcc

profile\_option

options used to generate profile data

feedback\_option

options used to optimize using that profile data

prof\_ext suffix of profile data files

PROFOPT\_OPTIONS

list of options with which to run each test, similar to the lists for torture tests

{ dg-final-generate { local-directive } }

This directive is similar to dg-final, but the *local-directive* is run after the generation of profile data.

{ dg-final-use { local-directive } }

The local-directive is run after the profile data have been used.

# 7.8 Support for testing binary compatibility

The file 'compat.exp' provides language-independent support for binary compatibility testing. It supports testing interoperability of two compilers that follow the same ABI, or of multiple sets of compiler options that should not affect binary compatibility. It is intended to be used for testsuites that complement ABI testsuites.

A test supported by this framework has three parts, each in a separate source file: a main program and two pieces that interact with each other to split up the functionality being tested.

'testname\_main.suffix'

Contains the main program, which calls a function in file 'testname\_x.suffix'.

'testname\_x.suffix'

Contains at least one call to a function in 'testname\_y.suffix'.

'testname\_v.suffix'

Shares data with, or gets arguments from, 'testname\_x.suffix'.

Within each test, the main program and one functional piece are compiled by the GCC under test. The other piece can be compiled by an alternate compiler. If no alternate compiler is specified, then all three source files are all compiled by the GCC under test. You can specify pairs of sets of compiler options. The first element of such a pair specifies options used with the GCC under test, and the second element of the pair specifies options used with the alternate compiler. Each test is compiled with each pair of options.

'compat.exp' defines default pairs of compiler options. These can be overridden by defining the environment variable COMPAT\_OPTIONS as:

```
COMPAT_OPTIONS="[list [list {tst1} {alt1}]
...[list {tstn} {altn}]]"
```

where tsti and alti are lists of options, with tsti used by the compiler under test and alti used by the alternate compiler. For example, with [list [list {-g -00} {-03}] [list {-fpic} {-fPIC -02}]], the test is first built with '-g -00' by the compiler under test and with '-03' by the alternate compiler. The test is built a second time using '-fpic' by the compiler under test and '-fPIC -02' by the alternate compiler.

An alternate compiler is specified by defining an environment variable to be the full pathname of an installed compiler; for C define ALT\_CC\_UNDER\_TEST, and for C++ define ALT\_CXX\_UNDER\_TEST. These will be written to the 'site.exp' file used by DejaGnu. The default is to build each test with the compiler under test using the first of each pair of compiler options from COMPAT\_OPTIONS. When ALT\_CC\_UNDER\_TEST or ALT\_CXX\_UNDER\_TEST is same, each test is built using the compiler under test but with combinations of the options from COMPAT\_OPTIONS.

To run only the C++ compatibility suite using the compiler under test and another version of GCC using specific compiler options, do the following from 'objdir/gcc':

```
rm site.exp
make -k \
   ALT_CXX_UNDER_TEST=${alt_prefix}/bin/g++ \
   COMPAT_OPTIONS="lists as shown above" \
   check-c++ \
   RUNTESTFLAGS="compat.exp"
```

A test that fails when the source files are compiled with different compilers, but passes when the files are compiled with the same compiler, demonstrates incompatibility of the generated code or runtime support. A test that fails for the alternate compiler but passes for the compiler under test probably tests for a bug that was fixed in the compiler under test but is present in the alternate compiler.

The binary compatibility tests support a small number of test framework commands that appear within comments in a test file.

#### dg-require-\*

These commands can be used in 'testname\_main.suffix' to skip the test if specific support is not available on the target.

#### dg-options

The specified options are used for compiling this particular source file, appended to the options from COMPAT\_OPTIONS. When this command appears in 'testname\_main.suffix' the options are also used to link the test program.

#### dg-xfail-if

This command can be used in a secondary source file to specify that compilation is expected to fail for particular options on particular targets.

# 7.9 Support for torture testing using multiple options

Throughout the compiler testsuite there are several directories whose tests are run multiple times, each with a different set of options. These are known as torture tests. 'lib/torture-options.exp' defines procedures to set up these lists:

#### torture-init

Initialize use of torture lists.

#### set-torture-options

Set lists of torture options to use for tests with and without loops. Optionally combine a set of torture options with a set of other options, as is done with Objective-C runtime options.

#### torture-finish

Finalize use of torture lists.

The '.exp' file for a set of tests that use torture options must include calls to these three procedures if:

- It calls gcc-dg-runtest and overrides DG\_TORTURE\_OPTIONS.
- It calls \$\{tool\}\-torture or \$\{tool\}\-torture\-execute, where tool is c, fortran, or objc.
- It calls dg-pch.

It is not necessary for a '.exp' file that calls gcc-dg-runtest to call the torture procedures if the tests should use the list in  $DG\_TORTURE\_OPTIONS$  defined in 'gcc-dg.exp'.

uses of torture options lists by defincan override  $_{
m the}$ default ing TORTURE\_OPTIONS or add to the default listby defining ADDI- $TIONAL\_TORTURE\_OPTIONS.$ Define these in a '.dejagnurc' file or add them to the 'site.exp' file; for example

```
set ADDITIONAL_TORTURE_OPTIONS [list \
   { -02 -ftree-loop-linear } \
   { -02 -fpeel-loops } ]
```

# 7.10 Support for testing GIMPLE passes

As of gcc 7, C functions can be tagged with \_\_GIMPLE to indicate that the function body will be GIMPLE, rather than C. The compiler requires the option '-fgimple' to enable this functionality. For example:

```
/* { dg-do compile } */
/* { dg-options "-0 -fgimple" } */

void __GIMPLE (startwith ("dse2")) foo ()
{
  int a;

bb_2:
  if (a > 4)
    goto bb_3;
  else
    goto bb_4;

bb_3:
  a_2 = 10;
  goto bb_5;

bb_4:
  a_3 = 20;
```

```
bb_5:
    a_1 = __PHI (bb_3: a_2, bb_4: a_3);
    a_4 = a_1 + 4;
    return;
}
```

The startwith argument indicates at which pass to begin.

Use the dump modifier -gimple (e.g. '-fdump-tree-all-gimple') to make tree dumps more closely follow the format accepted by the GIMPLE parser.

Example DejaGnu tests of GIMPLE can be seen in the source tree at 'gcc/testsuite/gcc.dg/gimplefe-\*.c'.

The \_\_GIMPLE parser is integrated with the C tokenizer and preprocessor, so it should be possible to use macros to build out test coverage.

# 7.11 Support for testing RTL passes

As of gcc 7, C functions can be tagged with \_\_RTL to indicate that the function body will be RTL, rather than C. For example:

```
double __RTL (startwith ("ira")) test (struct foo *f, const struct bar *b)
{
    (function "test"
        [...snip; various directives go in here...]
    );; function "test"
}
```

The startwith argument indicates at which pass to begin.

The parser expects the RTL body to be in the format emitted by this dumping function:

DEBUG FUNCTION void

```
print_rtx_function (FILE *outfile, function *fn, bool compact);
```

when "compact" is true. So you can capture RTL in the correct format from the debugger using:

```
(gdb) print_rtx_function (stderr, cfun, true);
```

and copy and paste the output into the body of the C function.

Example DejaGnu tests of RTL can be seen in the source tree under 'gcc/testsuite/gcc.dg/rtl'.

The \_\_RTL parser is not integrated with the C tokenizer or preprocessor, and works simply by reading the relevant lines within the braces. In particular, the RTL body must be on separate lines from the enclosing braces, and the preprocessor is not usable within it.

# 8 Option specification files

Most GCC command-line options are described by special option definition files, the names of which conventionally end in .opt. This chapter describes the format of these files.

# 8.1 Option file format

Option files are a simple list of records in which each field occupies its own line and in which the records themselves are separated by blank lines. Comments may appear on their own line anywhere within the file and are preceded by semicolons. Whitespace is allowed before the semicolon.

The files can contain the following types of record:

- A language definition record. These records have two fields: the string 'Language' and the name of the language. Once a language has been declared in this way, it can be used as an option property. See Section 8.2 [Option properties], page 121.
- A target specific save record to save additional information. These records have two fields: the string 'TargetSave', and a declaration type to go in the cl\_target\_option structure.
- A variable record to define a variable used to store option information. These records have two fields: the string 'Variable', and a declaration of the type and name of the variable, optionally with an initializer (but without any trailing ';'). These records may be used for variables used for many options where declaring the initializer in a single option definition record, or duplicating it in many records, would be inappropriate, or for variables set in option handlers rather than referenced by Var properties.
- A variable record to define a variable used to store option information. These records have two fields: the string 'TargetVariable', and a declaration of the type and name of the variable, optionally with an initializer (but without any trailing ';'). 'TargetVariable' is a combination of 'Variable' and 'TargetSave' records in that the variable is defined in the gcc\_options structure, but these variables are also stored in the cl\_target\_option structure. The variables are saved in the target save code and restored in the target restore code.
- A variable record to record any additional files that the 'options.h' file should include. This is useful to provide enumeration or structure definitions needed for target variables. These records have two fields: the string 'HeaderInclude' and the name of the include file.
- A variable record to record any additional files that the 'options.c' or 'options-save.c' file should include. This is useful to provide inline functions needed for target variables and/or #ifdef sequences to properly set up the initialization. These records have two fields: the string 'SourceInclude' and the name of the include file.
- An enumeration record to define a set of strings that may be used as arguments to an option or options. These records have three fields: the string 'Enum', a space-separated list of properties and help text used to describe the set of strings in '--help' output. Properties use the same format as option properties; the following are valid:

#### Name (name)

This property is required; name must be a name (suitable for use in C identifiers) used to identify the set of strings in Enum option properties.

## Type(type)

This property is required; type is the C type for variables set by options using this enumeration together with Var.

## UnknownError(message)

The message message will be used as an error message if the argument is invalid; for enumerations without UnknownError, a generic error message is used. message should contain a single '%qs' format, which will be used to format the invalid argument.

• An enumeration value record to define one of the strings in a set given in an 'Enum' record. These records have two fields: the string 'EnumValue' and a space-separated list of properties. Properties use the same format as option properties; the following are valid:

#### Enum (name)

This property is required; name says which 'Enum' record this 'EnumValue' record corresponds to.

#### String(string)

This property is required; string is the string option argument being described by this record.

#### Value(value)

This property is required; it says what value (representable as int) should be used for the given string.

#### Canonical

This property is optional. If present, it says the present string is the canonical one among all those with the given value. Other strings yielding that value will be mapped to this one so specs do not need to handle them.

#### DriverOnly

This property is optional. If present, the present string will only be accepted by the driver. This is used for cases such as '-march=native' that are processed by the driver so that 'gcc -v' shows how the options chosen depended on the system on which the compiler was run.

- An option definition record. These records have the following fields:
  - 1. the name of the option, with the leading "-" removed
  - 2. a space-separated list of option properties (see Section 8.2 [Option properties], page 121)
  - 3. the help text to use for '--help' (omitted if the second field contains the Undocumented property).

By default, all options beginning with "f", "W" or "m" are implicitly assumed to take a "no-" form. This form should not be listed separately. If an option beginning with one of these letters does not have a "no-" form, you can use the RejectNegative property to reject it.

The help text is automatically line-wrapped before being displayed. Normally the name of the option is printed on the left-hand side of the output and the help text is printed on the right. However, if the help text contains a tab character, the text to the left of the tab is used instead of the option's name and the text to the right of the tab forms the help text. This allows you to elaborate on what type of argument the option takes.

• A target mask record. These records have one field of the form 'Mask(x)'. The options-processing script will automatically allocate a bit in target\_flags (see Section 18.3 [Run-time Target], page 486) for each mask name x and set the macro MASK\_x to the appropriate bitmask. It will also declare a TARGET\_x macro that has the value 1 when bit MASK\_x is set and 0 otherwise.

They are primarily intended to declare target masks that are not associated with user options, either because these masks represent internal switches or because the options are not available on all configurations and yet the masks always need to be defined.

# 8.2 Option properties

The second field of an option record can specify any of the following properties. When an option takes an argument, it is enclosed in parentheses following the option property name. The parser that handles option files is quite simplistic, and will be tricked by any nested parentheses within the argument text itself; in this case, the entire option argument can be wrapped in curly braces within the parentheses to demarcate it, e.g.:

Condition({defined (USE\_CYGWIN\_LIBSTDCXX\_WRAPPERS)})

Common The option is available for all languages and targets.

Target The option is available for all languages but is target-specific.

Driver The option is handled by the compiler driver using code not shared with the compilers proper ('cc1' etc.).

language The option is available when compiling for the given language.

It is possible to specify several different languages for the same option. Each language must have been declared by an earlier Language record. See Section 8.1 [Option file format], page 119.

## RejectDriver

The option is only handled by the compilers proper ('cc1' etc.) and should not be accepted by the driver.

#### RejectNegative

The option does not have a "no-" form. All options beginning with "f", "W" or "m" are assumed to have a "no-" form unless this property is used.

# Negative(othername)

The option will turn off another option othername, which is the option name with the leading "-" removed. This chain action will propagate through the Negative property of the option to be turned off. The driver will prune options, removing those that are turned off by some later option. This pruning is not done for options with Joined or JoinedOrMissing properties, unless the options have either RejectNegative property or the Negative property mentions an option other than itself.

As a consequence, if you have a group of mutually-exclusive options, their Negative properties should form a circular chain. For example, if options '-a', '-b' and '-c' are mutually exclusive, their respective Negative properties should be 'Negative(b)', 'Negative(c)' and 'Negative(a)'.

#### Joined

Separate

The option takes a mandatory argument. Joined indicates that the option and argument can be included in the same argv entry (as with -mflush-func=name, for example). Separate indicates that the option and argument can be separate argv entries (as with -o). An option is allowed to have both of these properties.

# JoinedOrMissing

The option takes an optional argument. If the argument is given, it will be part of the same argv entry as the option itself.

This property cannot be used alongside Joined or Separate.

## MissingArgError(message)

For an option marked Joined or Separate, the message message will be used as an error message if the mandatory argument is missing; for options without MissingArgError, a generic error message is used. message should contain a single '%qs' format, which will be used to format the name of the option passed.

Args(n) For an option marked Separate, indicate that it takes n arguments. The default is 1.

UInteger The option's argument is a non-negative integer consisting of either decimal or hexadecimal digits interpreted as int. Hexadecimal integers may optionally start with the Ox or OX prefix. The option parser validates and converts the argument before passing it to the relevant option handler. UInteger should also be used with options like -falign-loops where both -falign-loops and -falign-loops=n are supported to make sure the saved options are given a full integer. Positive values of the argument in excess of INT\_MAX wrap around zero.

#### Host\_Wide\_Int

The option's argument is a non-negative integer consisting of either decimal or hexadecimal digits interpreted as the widest integer type on the host. As with an UInteger argument, hexadecimal integers may optionally start with the Ox or OX prefix. The option parser validates and converts the argument before passing it to the relevant option handler. Host\_Wide\_Int should be used with options that need to accept very large values. Positive values of the argument in excess of HOST\_WIDE\_INT\_M1U are assigned HOST\_WIDE\_INT\_M1U.

#### IntegerRange(n, m)

The options's arguments are integers of type int. The option's parser validates that the value of an option integer argument is within the closed range [n, m].

ByteSize A property applicable only to UInteger or Host\_Wide\_Int arguments. The option's integer argument is interpreted as if in infinite precision using saturation arithmetic in the corresponding type. The argument may be followed by a 'byte-size' suffix designating a multiple of bytes such as kB and KiB for

kilobyte and kibibyte, respectively, MB and MiB for megabyte and mebibyte, GB and GiB for gigabyte and gigibyte, and so on. ByteSize should be used for with options that take a very large argument representing a size in bytes, such as '-Wlarger-than='.

ToLower The option's argument should be converted to lowercase as part of putting it in canonical form, and before comparing with the strings indicated by any Enum property.

#### NoDriverArg

For an option marked Separate, the option only takes an argument in the compiler proper, not in the driver. This is for compatibility with existing options that are used both directly and via '-Wp,'; new options should not have this property.

Var(var) The state of this option should be stored in variable var (actually a macro for global\_options.x\_var). The way that the state is stored depends on the type of option:

#### WarnRemoved

The option is removed and every usage of such option will result in a warning. We use it option backward compatibility.

#### Var(var, set)

The option controls an integer variable var and is active when var equals set. The option parser will set var to set when the positive form of the option is used and !set when the "no-" form is used.

var is declared in the same way as for the single-argument form described above.

- If the option uses the Mask or InverseMask properties, var is the integer variable that contains the mask.
- If the option is a normal on/off switch, var is an integer variable that is nonzero when the option is enabled. The options parser will set the variable to 1 when the positive form of the option is used and 0 when the "no-" form is used.
- If the option takes an argument and has the UInteger property, var is an integer variable that stores the value of the argument.
- If the option takes an argument and has the Enum property, var is a variable (type given in the Type property of the 'Enum' record whose Name property has the same argument as the Enum property of this option) that stores the value of the argument.
- If the option has the Defer property, var is a pointer to a VEC(cl\_deferred\_option, heap) that stores the option for later processing. (var is declared with type void \* and needs to be cast to VEC(cl\_deferred\_option, heap) before use.)
- Otherwise, if the option takes an argument, var is a pointer to the argument string. The pointer will be null if the argument is optional and wasn't given.

The option-processing script will usually zero-initialize var. You can modify this behavior using Init.

#### Init(value)

The variable specified by the Var property should be statically initialized to value. If more than one option using the same variable specifies Init, all must specify the same initializer.

#### Mask(name)

The option is associated with a bit in the target\_flags variable (see Section 18.3 [Run-time Target], page 486) and is active when that bit is set. You may also specify Var to select a variable other than target\_flags.

The options-processing script will automatically allocate a unique bit for the option. If the option is attached to 'target\_flags', the script will set the macro MASK\_name to the appropriate bitmask. It will also declare a TARGET\_name macro that has the value 1 when the option is active and 0 otherwise. If you use Var to attach the option to a different variable, the bitmask macro with be called OPTION\_MASK\_name.

#### InverseMask(othername)

#### InverseMask(othername, thisname)

The option is the inverse of another option that has the Mask(othername) property. If this name is given, the options-processing script will declare a TARGET\_this name macro that is 1 when the option is active and 0 otherwise.

#### Enum(name)

The option's argument is a string from the set of strings associated with the corresponding 'Enum' record. The string is checked and converted to the integer specified in the corresponding 'EnumValue' record before being passed to option handlers.

Defer The option should be stored in a vector, specified with Var, for later processing.

Alias(opt)

Alias(opt, arg)

Alias(opt, posarg, negarg)

The option is an alias for '-opt' (or the negative form of that option, depending on NegativeAlias). In the first form, any argument passed to the alias is considered to be passed to '-opt', and '-opt' is considered to be negated if the alias is used in negated form. In the second form, the alias may not be negated or have an argument, and posarg is considered to be passed as an argument to '-opt'. In the third form, the alias may not have an argument, if the alias is used in the positive form then posarg is considered to be passed to '-opt', and if the alias is used in the negative form then negarg is considered to be passed to '-opt'.

Aliases should not specify Var or Mask or UInteger. Aliases should normally specify the same languages as the target of the alias; the flags on the target will be used to determine any diagnostic for use of an option for the wrong language, while those on the alias will be used to identify what command-line text is the option and what text is any argument to that option.

When an Alias definition is used for an option, driver specs do not need to handle it and no 'OPT\_' enumeration value is defined for it; only the canonical form of the option will be seen in those places.

#### NegativeAlias

For an option marked with Alias(opt), the option is considered to be an alias for the positive form of '-opt' if negated and for the negative form of '-opt' if not negated. NegativeAlias may not be used with the forms of Alias taking more than one argument.

Ignore This option is ignored apart from printing any warning specified using Warn.

The option will not be seen by specs and no 'OPT\_' enumeration value is defined for it.

# SeparateAlias

For an option marked with Joined, Separate and Alias, the option only acts as an alias when passed a separate argument; with a joined argument it acts as a normal option, with an 'OPT\_' enumeration value. This is for compatibility with the Java '-d' option and should not be used for new options.

## Warn(message)

If this option is used, output the warning message is a format string, either taking a single operand with a '%qs' format which is the option name, or not taking any operands, which is passed to the 'warning' function. If an alias is marked Warn, the target of the alias must not also be marked Warn.

Report The state of the option should be printed by '-fverbose-asm'.

Warning This is a warning option and should be shown as such in '--help' output. This flag does not currently affect anything other than '--help'.

## Optimization

This is an optimization option. It should be shown as such in '--help' output, and any associated variable named using Var should be saved and restored when the optimization level is changed with optimize attributes.

#### PerFunction

This is an option that can be overridden on a per-function basis. Optimization implies PerFunction, but options that do not affect executable code generation may use this flag instead, so that the option is not taken into account in ways that might affect executable code generation.

Param This is an option that is a parameter.

#### Undocumented

The option is deliberately missing documentation and should not be included in the '--help' output.

# Condition(cond)

The option should only be accepted if preprocessor condition *cond* is true. Note that any C declarations associated with the option will be present even if *cond* is false; *cond* simply controls whether the option is accepted and whether it is printed in the '--help' output.

Save Build the cl\_target\_option structure to hold a copy of the option, add the functions cl\_target\_option\_save and cl\_target\_option\_restore to save and restore the options.

#### SetByCombined

The option may also be set by a combined option such as '-ffast-math'. This causes the gcc\_options struct to have a field frontend\_set\_name, where name is the name of the field holding the value of this option (without the leading x\_). This gives the front end a way to indicate that the value has been set explicitly and should not be changed by the combined option. For example, some front ends use this to prevent '-ffast-math' and '-fno-fast-math' from changing the value of '-fmath-errno' for languages that do not use errno.

#### EnabledBy(opt)

EnabledBy(opt || opt2)

EnabledBy(opt && opt2)

If not explicitly set, the option is set to the value of '-opt'; multiple options can be given, separated by ||. The third form using && specifies that the option is only set if both opt and opt2 are set. The options opt and opt2 must have the Common property; otherwise, use LangEnabledBy.

# LangEnabledBy(language, opt)

#### LangEnabledBy(language, opt, posarg, negarg)

When compiling for the given language, the option is set to the value of '-opt', if not explicitly set. opt can be also a list of | | separated options. In the second form, if opt is used in the positive form then posarg is considered to be passed to the option, and if opt is used in the negative form then negarg is considered to be passed to the option. It is possible to specify several different languages. Each language must have been declared by an earlier Language record. See Section 8.1 [Option file format], page 119.

#### NoDWARFRecord

The option is omitted from the producer string written by '-grecord-gcc-switches'.

## PchIgnore

Even if this is a target option, this option will not be recorded / compared to determine if a precompiled header file matches.

CPP(var) The state of this option should be kept in sync with the preprocessor option var. If this property is set, then properties Var and Init must be set as well.

## CppReason(CPP\_W\_Enum)

This warning option corresponds to cpplib.h warning reason code  $CPP_-W_-Enum$ . This should only be used for warning options of the C-family front-ends.

# 9 Passes and Files of the Compiler

This chapter is dedicated to giving an overview of the optimization and code generation passes of the compiler. In the process, it describes some of the language front end interface, though this description is no where near complete.

# 9.1 Parsing pass

The language front end is invoked only once, via lang\_hooks.parse\_file, to parse the entire input. The language front end may use any intermediate language representation deemed appropriate. The C front end uses GENERIC trees (see Chapter 11 [GENERIC], page 161), plus a double handful of language specific tree codes defined in 'c-common.def'. The Fortran front end uses a completely different private representation.

At some point the front end must translate the representation used in the front end to a representation understood by the language-independent portions of the compiler. Current practice takes one of two forms. The C front end manually invokes the gimplifier (see Chapter 12 [GIMPLE], page 209) on each function, and uses the gimplifier callbacks to convert the language-specific tree nodes directly to GIMPLE before passing the function off to be compiled. The Fortran front end converts from a private representation to GENERIC, which is later lowered to GIMPLE when the function is compiled. Which route to choose probably depends on how well GENERIC (plus extensions) can be made to match up with the source language and necessary parsing data structures.

BUG: Gimplification must occur before nested function lowering, and nested function lowering must be done by the front end before passing the data off to cgraph.

TODO: Cgraph should control nested function lowering. It would only be invoked when it is certain that the outer-most function is used.

TODO: Cgraph needs a gimplify\_function callback. It should be invoked when (1) it is certain that the function is used, (2) warning flags specified by the user require some amount of compilation in order to honor, (3) the language indicates that semantic analysis is not complete until gimplification occurs. Hum... this sounds overly complicated. Perhaps we should just have the front end gimplify always; in most cases it's only one function call.

The front end needs to pass all function definitions and top level declarations off to the middle-end so that they can be compiled and emitted to the object file. For a simple procedural language, it is usually most convenient to do this as each top level declaration or definition is seen. There is also a distinction to be made between generating functional code and generating complete debug information. The only thing that is absolutely required for functional code is that function and data definitions be passed to the middle-end. For complete debug information, function, data and type declarations should all be passed as well.

In any case, the front end needs each complete top-level function or data declaration, and each data definition should be passed to rest\_of\_decl\_compilation. Each complete type definition should be passed to rest\_of\_type\_compilation. Each function definition should be passed to cgraph\_finalize\_function.

TODO: I know rest\_of\_compilation currently has all sorts of RTL generation semantics. I plan to move all code generation bits (both Tree and RTL) to compile\_function. Should we hide cgraph from the front ends and move back to rest\_of\_compilation as the official

interface? Possibly we should rename all three interfaces such that the names match in some meaningful way and that is more descriptive than "rest\_of".

The middle-end will, at its option, emit the function and data definitions immediately or queue them for later processing.

# 9.2 Gimplification pass

Gimplification is a whimsical term for the process of converting the intermediate representation of a function into the GIMPLE language (see Chapter 12 [GIMPLE], page 209). The term stuck, and so words like "gimplification", "gimplify", "gimplifier" and the like are sprinkled throughout this section of code.

While a front end may certainly choose to generate GIMPLE directly if it chooses, this can be a moderately complex process unless the intermediate language used by the front end is already fairly simple. Usually it is easier to generate GENERIC trees plus extensions and let the language-independent gimplifier do most of the work.

The main entry point to this pass is gimplify\_function\_tree located in 'gimplify.c'. From here we process the entire function gimplifying each statement in turn. The main workhorse for this pass is gimplify\_expr. Approximately everything passes through here at least once, and it is from here that we invoke the lang\_hooks.gimplify\_expr callback.

The callback should examine the expression in question and return GS\_UNHANDLED if the expression is not a language specific construct that requires attention. Otherwise it should alter the expression in some way to such that forward progress is made toward producing valid GIMPLE. If the callback is certain that the transformation is complete and the expression is valid GIMPLE, it should return GS\_ALL\_DONE. Otherwise it should return GS\_OK, which will cause the expression to be processed again. If the callback encounters an error during the transformation (because the front end is relying on the gimplification process to finish semantic checks), it should return GS\_ERROR.

# 9.3 Pass manager

The pass manager is located in 'passes.c', 'tree-optimize.c' and 'tree-pass.h'. It processes passes as described in 'passes.def'. Its job is to run all of the individual passes in the correct order, and take care of standard bookkeeping that applies to every pass.

The theory of operation is that each pass defines a structure that represents everything we need to know about that pass—when it should be run, how it should be run, what intermediate language form or on-the-side data structures it needs. We register the pass to be run in some particular order, and the pass manager arranges for everything to happen in the correct order.

The actuality doesn't completely live up to the theory at present. Command-line switches and timevar\_id\_t enumerations must still be defined elsewhere. The pass manager validates constraints but does not attempt to (re-)generate data structures or lower intermediate language form based on the requirements of the next pass. Nevertheless, what is present is useful, and a far sight better than nothing at all.

Each pass should have a unique name. Each pass may have its own dump file (for GCC debugging purposes). Passes with a name starting with a star do not dump anything. Sometimes passes are supposed to share a dump file / option name. To still give these

unique names, you can use a prefix that is delimited by a space from the part that is used for the dump file / option name. E.g. When the pass name is "ud dce", the name used for dump file/options is "dce".

TODO: describe the global variables set up by the pass manager, and a brief description of how a new pass should use it. I need to look at what info RTL passes use first...

# 9.4 Inter-procedural optimization passes

The inter-procedural optimization (IPA) passes use call graph information to perform transformations across function boundaries. IPA is a critical part of link-time optimization (LTO) and whole-program (WHOPR) optimization, and these passes are structured with the needs of LTO and WHOPR in mind by dividing their operations into stages. For detailed discussion of the LTO/WHOPR IPA pass stages and interfaces, see Section 25.3 [IPA], page 696.

The following briefly describes the inter-procedural optimization (IPA) passes, which are split into small IPA passes, regular IPA passes, and late IPA passes, according to the LTO/WHOPR processing model.

# 9.4.1 Small IPA passes

A small IPA pass is a pass derived from simple\_ipa\_opt\_pass. As described in Section 25.3 [IPA], page 696, it does everything at once and defines only the *Execute* stage. During this stage it accesses and modifies the function bodies. No generate\_summary, read\_summary, or write\_summary hooks are defined.

- IPA free lang data
  - This pass frees resources that are used by the front end but are not needed once it is done. It is located in 'tree.c' and is described by pass\_ipa\_free\_lang\_data.
- IPA function and variable visibility
  - This is a local function pass handling visibilities of all symbols. This happens before LTO streaming, so '-fwhole-program' should be ignored at this level. It is located in 'ipa-visibility.c' and is described by pass\_ipa\_function\_and\_variable\_visibility.
- IPA remove symbols
  - This pass performs reachability analysis and reclaims all unreachable nodes. It is located in 'passes.c' and is described by pass\_ipa\_remove\_symbols.
- IPA OpenACC
  - This is a pass group for OpenACC processing. It is located in 'tree-ssa-loop.c' and is described by pass\_ipa\_oacc.
- IPA points-to analysis
  - This is a tree-based points-to analysis pass. The idea behind this analyzer is to generate set constraints from the program, then solve the resulting constraints in order to generate the points-to sets. It is located in 'tree-ssa-structalias.c' and is described by pass\_ipa\_pta.
- IPA OpenACC kernels
  - This is a pass group for processing OpenACC kernels regions. It is a subpass of the IPA OpenACC pass group that runs on offloaded functions containing OpenACC kernels loops. It is located in 'tree-ssa-loop.c' and is described by pass\_ipa\_oacc\_kernels.

## • Target clone

This is a pass for parsing functions with multiple target attributes. It is located in 'multiple\_target.c' and is described by pass\_target\_clone.

#### • IPA auto profile

This pass uses AutoFDO profiling data to annotate the control flow graph. It is located in 'auto-profile.c' and is described by pass\_ipa\_auto\_profile.

#### • IPA tree profile

This pass does profiling for all functions in the call graph. It calculates branch probabilities and basic block execution counts. It is located in 'tree-profile.c' and is described by pass\_ipa\_tree\_profile.

# • IPA free function summary

This pass is a small IPA pass when argument small\_p is true. It releases inline function summaries and call summaries. It is located in 'ipa-fnsummary.c' and is described by pass\_ipa\_free\_free\_fn\_summary.

## • IPA increase alignment

This pass increases the alignment of global arrays to improve vectorization. It is located in 'tree-vectorizer.c' and is described by pass\_ipa\_increase\_alignment.

#### • IPA transactional memory

This pass is for transactional memory support. It is located in 'trans-mem.c' and is described by pass\_ipa\_tm.

#### • IPA lower emulated TLS

This pass lowers thread-local storage (TLS) operations to emulation functions provided by libgcc. It is located in 'tree-emutls.c' and is described by pass\_ipa\_lower\_emutls.

## 9.4.2 Regular IPA passes

A regular IPA pass is a pass derived from <code>ipa\_opt\_pass\_d</code> that is executed in WHOPR compilation. Regular IPA passes may have summary hooks implemented in any of the LGEN, WPA or LTRANS stages (see Section 25.3 [IPA], page 696).

## • IPA whole program visibility

This pass performs various optimizations involving symbol visibility with '-fwhole-program', including symbol privatization, discovering local functions, and dismantling comdat groups. It is located in 'ipa-visibility.c' and is described by pass\_ipa\_whole\_program\_visibility.

#### • IPA profile

The IPA profile pass propagates profiling frequencies across the call graph. It is located in 'ipa-profile.c' and is described by pass\_ipa\_profile.

## • IPA identical code folding

This is the inter-procedural identical code folding pass. The goal of this transformation is to discover functions and read-only variables that have exactly the same semantics. It is located in 'ipa-icf.c' and is described by pass\_ipa\_icf.

#### • IPA devirtualization

This pass performs speculative devirtualization based on the type inheritance graph. When a polymorphic call has only one likely target in the unit, it is turned into a speculative call. It is located in 'ipa-devirt.c' and is described by pass\_ipa\_devirt.

## • IPA constant propagation

The goal of this pass is to discover functions that are always invoked with some arguments with the same known constant values and to modify the functions accordingly. It can also do partial specialization and type-based devirtualization. It is located in 'ipa-cp.c' and is described by pass\_ipa\_cp.

## • IPA scalar replacement of aggregates

This pass can replace an aggregate parameter with a set of other parameters representing part of the original, turning those passed by reference into new ones which pass the value directly. It also removes unused function return values and unused function parameters. This pass is located in 'ipa-sra.c' and is described by pass\_ipa\_sra.

# • IPA constructor/destructor merge

This pass merges multiple constructors and destructors for static objects into single functions. It's only run at LTO time unless the target doesn't support constructors and destructors natively. The pass is located in 'ipa.c' and is described by pass\_ipa\_cdtor\_merge.

#### • IPA HSA

This pass is part of the GCC support for HSA (Heterogeneous System Architecture) accelerators. It is responsible for creation of HSA clones and emitting HSAIL instructions for them. It is located in 'ipa-hsa.c' and is described by pass\_ipa\_hsa.

#### • IPA function summary

This pass provides function analysis for inter-procedural passes. It collects estimates of function body size, execution time, and frame size for each function. It also estimates information about function calls: call statement size, time and how often the parameters change for each call. It is located in 'ipa-fnsummary.c' and is described by pass\_ipa\_fn\_summary.

# • IPA inline

The IPA inline pass handles function inlining with whole-program knowledge. Small functions that are candidates for inlining are ordered in increasing badness, bounded by unit growth parameters. Unreachable functions are removed from the call graph. Functions called once and not exported from the unit are inlined. This pass is located in 'ipa-inline.c' and is described by pass\_ipa\_inline.

#### • IPA pure/const analysis

This pass marks functions as being either const (TREE\_READONLY) or pure (DECL\_PURE\_P). The per-function information is produced by pure\_const\_generate\_summary, then the global information is computed by performing a transitive closure over the call graph. It is located in 'ipa-pure-const.c' and is described by pass\_ipa\_pure\_const.

#### • IPA free function summary

This pass is a regular IPA pass when argument small\_p is false. It releases inline function summaries and call summaries. It is located in 'ipa-fnsummary.c' and is described by pass\_ipa\_free\_fn\_summary.

#### • IPA reference

This pass gathers information about how variables whose scope is confined to the compilation unit are used. It is located in 'ipa-reference.c' and is described by pass\_ipa\_reference.

## • IPA single use

This pass checks whether variables are used by a single function. It is located in 'ipa.c' and is described by pass\_ipa\_single\_use.

#### • IPA comdats

This pass looks for static symbols that are used exclusively within one comdat group, and moves them into that comdat group. It is located in 'ipa-comdats.c' and is described by pass\_ipa\_comdats.

# 9.4.3 Late IPA passes

Late IPA passes are simple IPA passes executed after the regular passes. In WHOPR mode the passes are executed after partitioning and thus see just parts of the compiled unit.

#### • Materialize all clones

Once all functions from compilation unit are in memory, produce all clones and update all calls. It is located in 'ipa.c' and is described by pass\_materialize\_all\_clones.

## • IPA points-to analysis

Points-to analysis; this is the same as the points-to-analysis pass run with the small IPA passes (see Section 9.4.1 [Small IPA passes], page 129).

#### • OpenMP simd clone

This is the OpenMP constructs' SIMD clone pass. It creates the appropriate SIMD clones for functions tagged as elemental SIMD functions. It is located in 'omp-simd-clone.c' and is described by pass\_omp\_simd\_clone.

# 9.5 Tree SSA passes

The following briefly describes the Tree optimization passes that are run after gimplification and what source files they are located in.

## • Remove useless statements

This pass is an extremely simple sweep across the gimple code in which we identify obviously dead code and remove it. Here we do things like simplify if statements with constant conditions, remove exception handling constructs surrounding code that obviously cannot throw, remove lexical bindings that contain no variables, and other assorted simplistic cleanups. The idea is to get rid of the obvious stuff quickly rather than wait until later when it's more work to get rid of it. This pass is located in 'tree-cfg.c' and described by pass\_remove\_useless\_stmts.

## • OpenMP lowering

If OpenMP generation ('-fopenmp') is enabled, this pass lowers OpenMP constructs into GIMPLE.

Lowering of OpenMP constructs involves creating replacement expressions for local variables that have been mapped using data sharing clauses, exposing the control flow of most synchronization directives and adding region markers to facilitate the creation

of the control flow graph. The pass is located in 'omp-low.c' and is described by pass\_lower\_omp.

## • OpenMP expansion

If OpenMP generation ('-fopenmp') is enabled, this pass expands parallel regions into their own functions to be invoked by the thread library. The pass is located in 'omp-low.c' and is described by pass\_expand\_omp.

# • Lower control flow

This pass flattens if statements (COND\_EXPR) and moves lexical bindings (BIND\_EXPR) out of line. After this pass, all if statements will have exactly two goto statements in its then and else arms. Lexical binding information for each statement will be found in TREE\_BLOCK rather than being inferred from its position under a BIND\_EXPR. This pass is found in 'gimple-low.c' and is described by pass\_lower\_cf.

# • Lower exception handling control flow

This pass decomposes high-level exception handling constructs (TRY\_FINALLY\_EXPR and TRY\_CATCH\_EXPR) into a form that explicitly represents the control flow involved. After this pass, lookup\_stmt\_eh\_region will return a non-negative number for any statement that may have EH control flow semantics; examine tree\_can\_throw\_internal or tree\_can\_throw\_external for exact semantics. Exact control flow may be extracted from foreach\_reachable\_handler. The EH region nesting tree is defined in 'except.h' and built in 'except.c'. The lowering pass itself is in 'tree-eh.c' and is described by pass\_lower\_eh.

# • Build the control flow graph

This pass decomposes a function into basic blocks and creates all of the edges that connect them. It is located in 'tree-cfg.c' and is described by pass\_build\_cfg.

#### • Find all referenced variables

This pass walks the entire function and collects an array of all variables referenced in the function, referenced\_vars. The index at which a variable is found in the array is used as a UID for the variable within this function. This data is needed by the SSA rewriting routines. The pass is located in 'tree-dfa.c' and is described by pass\_referenced\_vars.

## • Enter static single assignment form

This pass rewrites the function such that it is in SSA form. After this pass, all is\_gimple\_reg variables will be referenced by SSA\_NAME, and all occurrences of other variables will be annotated with VDEFS and VUSES; PHI nodes will have been inserted as necessary for each basic block. This pass is located in 'tree-ssa.c' and is described by pass\_build\_ssa.

## • Warn for uninitialized variables

This pass scans the function for uses of SSA\_NAMEs that are fed by default definition. For non-parameter variables, such uses are uninitialized. The pass is run twice, before and after optimization (if turned on). In the first pass we only warn for uses that are positively uninitialized; in the second pass we warn for uses that are possibly uninitialized. The pass is located in 'tree-ssa.c' and is defined by pass\_early\_warn\_uninitialized and pass\_late\_warn\_uninitialized.

#### • Dead code elimination

This pass scans the function for statements without side effects whose result is unused. It does not do memory life analysis, so any value that is stored in memory is considered used. The pass is run multiple times throughout the optimization process. It is located in 'tree-ssa-dce.c' and is described by pass\_dce.

## • Dominator optimizations

This pass performs trivial dominator-based copy and constant propagation, expression simplification, and jump threading. It is run multiple times throughout the optimization process. It is located in 'tree-ssa-dom.c' and is described by pass\_dominator.

## • Forward propagation of single-use variables

This pass attempts to remove redundant computation by substituting variables that are used once into the expression that uses them and seeing if the result can be simplified. It is located in 'tree-ssa-forwprop.c' and is described by pass\_forwprop.

#### • Copy Renaming

This pass attempts to change the name of compiler temporaries involved in copy operations such that SSA->normal can coalesce the copy away. When compiler temporaries are copies of user variables, it also renames the compiler temporary to the user variable resulting in better use of user symbols. It is located in 'tree-ssa-copyrename.c' and is described by pass\_copyrename.

#### • PHI node optimizations

This pass recognizes forms of PHI inputs that can be represented as conditional expressions and rewrites them into straight line code. It is located in 'tree-ssa-phiopt.c' and is described by pass\_phiopt.

#### • May-alias optimization

This pass performs a flow sensitive SSA-based points-to analysis. The resulting mayalias, must-alias, and escape analysis information is used to promote variables from in-memory addressable objects to non-aliased variables that can be renamed into SSA form. We also update the VDEF/VUSE memory tags for non-renameable aggregates so that we get fewer false kills. The pass is located in 'tree-ssa-alias.c' and is described by pass\_may\_alias.

Interprocedural points-to information is located in 'tree-ssa-structalias.c' and described by pass\_ipa\_pta.

#### Profiling

This pass instruments the function in order to collect runtime block and value profiling data. Such data may be fed back into the compiler on a subsequent run so as to allow optimization based on expected execution frequencies. The pass is located in 'tree-profile.c' and is described by pass\_ipa\_tree\_profile.

#### • Static profile estimation

This pass implements series of heuristics to guess propababilities of branches. The resulting predictions are turned into edge profile by propagating branches across the control flow graphs. The pass is located in 'tree-profile.c' and is described by pass\_profile.

#### • Lower complex arithmetic

This pass rewrites complex arithmetic operations into their component scalar arithmetic operations. The pass is located in 'tree-complex.c' and is described by pass\_lower\_complex.

## • Scalar replacement of aggregates

This pass rewrites suitable non-aliased local aggregate variables into a set of scalar variables. The resulting scalar variables are rewritten into SSA form, which allows subsequent optimization passes to do a significantly better job with them. The pass is located in 'tree-sra.c' and is described by pass\_sra.

#### • Dead store elimination

This pass eliminates stores to memory that are subsequently overwritten by another store, without any intervening loads. The pass is located in 'tree-ssa-dse.c' and is described by pass\_dse.

#### • Tail recursion elimination

This pass transforms tail recursion into a loop. It is located in 'tree-tailcall.c' and is described by pass\_tail\_recursion.

## • Forward store motion

This pass sinks stores and assignments down the flowgraph closer to their use point. The pass is located in 'tree-ssa-sink.c' and is described by pass\_sink\_code.

#### • Partial redundancy elimination

This pass eliminates partially redundant computations, as well as performing load motion. The pass is located in 'tree-ssa-pre.c' and is described by pass\_pre.

Just before partial redundancy elimination, if '-funsafe-math-optimizations' is on, GCC tries to convert divisions to multiplications by the reciprocal. The pass is located in 'tree-ssa-math-opts.c' and is described by pass\_cse\_reciprocal.

#### • Full redundancy elimination

This is a simpler form of PRE that only eliminates redundancies that occur on all paths. It is located in 'tree-ssa-pre.c' and described by pass\_fre.

#### • Loop optimization

The main driver of the pass is placed in 'tree-ssa-loop.c' and described by pass\_loop.

The optimizations performed by this pass are:

Loop invariant motion. This pass moves only invariants that would be hard to handle on RTL level (function calls, operations that expand to nontrivial sequences of insns). With '-funswitch-loops' it also moves operands of conditions that are invariant out of the loop, so that we can use just trivial invariantness analysis in loop unswitching. The pass also includes store motion. The pass is implemented in 'tree-ssa-loop-im.c'.

Canonical induction variable creation. This pass creates a simple counter for number of iterations of the loop and replaces the exit condition of the loop using it, in case when a complicated analysis is necessary to determine the number of iterations. Later optimizations then may determine the number easily. The pass is implemented in 'tree-ssa-loop-ivcanon.c'.

Induction variable optimizations. This pass performs standard induction variable optimizations, including strength reduction, induction variable merging and induction variable elimination. The pass is implemented in 'tree-ssa-loop-ivopts.c'.

Loop unswitching. This pass moves the conditional jumps that are invariant out of the loops. To achieve this, a duplicate of the loop is created for each possible outcome of conditional jump(s). The pass is implemented in 'tree-ssa-loop-unswitch.c'.

Loop splitting. If a loop contains a conditional statement that is always true for one part of the iteration space and false for the other this pass splits the loop into two, one dealing with one side the other only with the other, thereby removing one inner-loop conditional. The pass is implemented in 'tree-ssa-loop-split.c'.

The optimizations also use various utility functions contained in 'tree-ssa-loop-manip.c', 'cfgloop.c', 'cfgloopanal.c' and 'cfgloopmanip.c'.

Vectorization. This pass transforms loops to operate on vector types instead of scalar types. Data parallelism across loop iterations is exploited to group data elements from consecutive iterations into a vector and operate on them in parallel. Depending on available target support the loop is conceptually unrolled by a factor VF (vectorization factor), which is the number of elements operated upon in parallel in each iteration, and the VF copies of each scalar operation are fused to form a vector operation. Additional loop transformations such as peeling and versioning may take place to align the number of iterations, and to align the memory accesses in the loop. The pass is implemented in 'tree-vectorizer.c' (the main driver), 'tree-vect-loop.c' and 'tree-vect-loop-manip.c' (loop specific parts and general loop utilities), 'tree-vect-slp' (loop-aware SLP functionality), 'tree-vect-stmts.c' and 'tree-vect-data-refs.c'. Analysis of data references is in 'tree-data-ref.c'.

SLP Vectorization. This pass performs vectorization of straight-line code. The pass is implemented in 'tree-vectorizer.c' (the main driver), 'tree-vect-slp.c', 'tree-vect-stmts.c' and 'tree-vect-data-refs.c'.

Autoparallelization. This pass splits the loop iteration space to run into several threads. The pass is implemented in 'tree-parloops.c'.

Graphite is a loop transformation framework based on the polyhedral model. Graphite stands for Gimple Represented as Polyhedra. The internals of this infrastructure are documented in <a href="http://gcc.gnu.org/wiki/Graphite">http://gcc.gnu.org/wiki/Graphite</a>. The passes working on this representation are implemented in the various 'graphite-\*' files.

# • Tree level if-conversion for vectorizer

This pass applies if-conversion to simple loops to help vectorizer. We identify if convertible loops, if-convert statements and merge basic blocks in one big block. The idea is to present loop in such form so that vectorizer can have one to one mapping between statements and available vector operations. This pass is located in 'tree-if-conv.c' and is described by pass\_if\_conversion.

## • Conditional constant propagation

This pass relaxes a lattice of values in order to identify those that must be constant even in the presence of conditional branches. The pass is located in 'tree-ssa-ccp.c' and is described by pass\_ccp.

A related pass that works on memory loads and stores, and not just register values, is located in 'tree-ssa-ccp.c' and described by pass\_store\_ccp.

#### • Conditional copy propagation

This is similar to constant propagation but the lattice of values is the "copy-of" relation. It eliminates redundant copies from the code. The pass is located in 'tree-ssa-copy.c' and described by pass\_copy\_prop.

A related pass that works on memory copies, and not just register copies, is located in 'tree-ssa-copy.c' and described by pass\_store\_copy\_prop.

#### • Value range propagation

This transformation is similar to constant propagation but instead of propagating single constant values, it propagates known value ranges. The implementation is based on Patterson's range propagation algorithm (Accurate Static Branch Prediction by Value Range Propagation, J. R. C. Patterson, PLDI '95). In contrast to Patterson's algorithm, this implementation does not propagate branch probabilities nor it uses more than a single range per SSA name. This means that the current implementation cannot be used for branch prediction (though adapting it would not be difficult). The pass is located in 'tree-vrp.c' and is described by pass\_vrp.

## • Folding built-in functions

This pass simplifies built-in functions, as applicable, with constant arguments or with inferable string lengths. It is located in 'tree-ssa-ccp.c' and is described by pass\_fold\_builtins.

## • Split critical edges

This pass identifies critical edges and inserts empty basic blocks such that the edge is no longer critical. The pass is located in 'tree-cfg.c' and is described by pass\_split\_crit\_edges.

## • Control dependence dead code elimination

This pass is a stronger form of dead code elimination that can eliminate unnecessary control flow statements. It is located in 'tree-ssa-dce.c' and is described by pass\_cd\_dce.

## • Tail call elimination

This pass identifies function calls that may be rewritten into jumps. No code transformation is actually applied here, but the data and control flow problem is solved. The code transformation requires target support, and so is delayed until RTL. In the meantime CALL\_EXPR\_TAILCALL is set indicating the possibility. The pass is located in 'tree-tailcall.c' and is described by pass\_tail\_calls. The RTL transformation is handled by fixup\_tail\_calls in 'calls.c'.

#### • Warn for function return without value

For non-void functions, this pass locates return statements that do not specify a value and issues a warning. Such a statement may have been injected by falling off the end of the function. This pass is run last so that we have as much time as possible to prove that the statement is not reachable. It is located in 'tree-cfg.c' and is described by pass\_warn\_function\_return.

## • Leave static single assignment form

This pass rewrites the function such that it is in normal form. At the same time, we eliminate as many single-use temporaries as possible, so the intermediate language is

no longer GIMPLE, but GENERIC. The pass is located in 'tree-outof-ssa.c' and is described by pass\_del\_ssa.

#### • Merge PHI nodes that feed into one another

This is part of the CFG cleanup passes. It attempts to join PHI nodes from a forwarder CFG block into another block with PHI nodes. The pass is located in 'tree-cfgcleanup.c' and is described by pass\_merge\_phi.

#### • Return value optimization

If a function always returns the same local variable, and that local variable is an aggregate type, then the variable is replaced with the return value for the function (i.e., the function's DECL\_RESULT). This is equivalent to the C++ named return value optimization applied to GIMPLE. The pass is located in 'tree-nrv.c' and is described by pass\_nrv.

# • Return slot optimization

If a function returns a memory object and is called as var = foo(), this pass tries to change the call so that the address of var is sent to the caller to avoid an extra memory copy. This pass is located in tree-nrv.c and is described by pass\_return\_slot.

## • Optimize calls to \_\_builtin\_object\_size

This is a propagation pass similar to CCP that tries to remove calls to \_\_builtin\_object\_size when the size of the object can be computed at compile-time. This pass is located in 'tree-object-size.c' and is described by pass\_object\_sizes.

#### • Loop invariant motion

This pass removes expensive loop-invariant computations out of loops. The pass is located in 'tree-ssa-loop.c' and described by pass\_lim.

#### • Loop nest optimizations

This is a family of loop transformations that works on loop nests. It includes loop interchange, scaling, skewing and reversal and they are all geared to the optimization of data locality in array traversals and the removal of dependencies that hamper optimizations such as loop parallelization and vectorization. The pass is located in 'tree-loop-linear.c' and described by pass\_linear\_transform.

## • Removal of empty loops

This pass removes loops with no code in them. The pass is located in 'tree-ssa-loop-ivcanon.c' and described by pass\_empty\_loop.

# • Unrolling of small loops

This pass completely unrolls loops with few iterations. The pass is located in 'tree-ssa-loop-ivcanon.c' and described by pass\_complete\_unroll.

## • Predictive commoning

This pass makes the code reuse the computations from the previous iterations of the loops, especially loads and stores to memory. It does so by storing the values of these computations to a bank of temporary variables that are rotated at the end of loop. To avoid the need for this rotation, the loop is then unrolled and the copies of the loop body are rewritten to use the appropriate version of the temporary variable. This pass is located in 'tree-predcom.c' and described by pass\_predcom.

## • Array prefetching

This pass issues prefetch instructions for array references inside loops. The pass is located in 'tree-ssa-loop-prefetch.c' and described by pass\_loop\_prefetch.

#### Reassociation

This pass rewrites arithmetic expressions to enable optimizations that operate on them, like redundancy elimination and vectorization. The pass is located in 'tree-ssa-reassoc.c' and described by pass\_reassoc.

#### • Optimization of stdarg functions

This pass tries to avoid the saving of register arguments into the stack on entry to stdarg functions. If the function doesn't use any va\_start macros, no registers need to be saved. If va\_start macros are used, the va\_list variables don't escape the function, it is only necessary to save registers that will be used in va\_arg macros. For instance, if va\_arg is only used with integral types in the function, floating point registers don't need to be saved. This pass is located in tree-stdarg.c and described by pass\_stdarg.

# 9.6 RTL passes

The following briefly describes the RTL generation and optimization passes that are run after the Tree optimization passes.

#### • RTL generation

The source files for RTL generation include 'stmt.c', 'calls.c', 'expr.c', 'explow.c', 'expmed.c', 'function.c', 'optabs.c' and 'emit-rtl.c'. Also, the file 'insn-emit.c', generated from the machine description by the program genemit, is used in this pass. The header file 'expr.h' is used for communication within this pass.

The header files 'insn-flags.h' and 'insn-codes.h', generated from the machine description by the programs genflags and gencodes, tell this pass which standard names are available for use and which patterns correspond to them.

## • Generation of exception landing pads

This pass generates the glue that handles communication between the exception handling library routines and the exception handlers within the function. Entry points in the function that are invoked by the exception handling library are called *landing pads*. The code for this pass is located in 'except.c'.

#### • Control flow graph cleanup

This pass removes unreachable code, simplifies jumps to next, jumps to jump, jumps across jumps, etc. The pass is run multiple times. For historical reasons, it is occasionally referred to as the "jump optimization pass". The bulk of the code for this pass is in 'cfgcleanup.c', and there are support routines in 'cfgrtl.c' and 'jump.c'.

## • Forward propagation of single-def values

This pass attempts to remove redundant computation by substituting variables that come from a single definition, and seeing if the result can be simplified. It performs copy propagation and addressing mode selection. The pass is run twice, with values being propagated into loops only on the second run. The code is located in 'fwprop.c'.

## • Common subexpression elimination

This pass removes redundant computation within basic blocks, and optimizes addressing modes based on cost. The pass is run twice. The code for this pass is located in 'cse.c'.

## • Global common subexpression elimination

This pass performs two different types of GCSE depending on whether you are optimizing for size or not (LCM based GCSE tends to increase code size for a gain in speed, while Morel-Renvoise based GCSE does not). When optimizing for size, GCSE is done using Morel-Renvoise Partial Redundancy Elimination, with the exception that it does not try to move invariants out of loops—that is left to the loop optimization pass. If MR PRE GCSE is done, code hoisting (aka unification) is also done, as well as load motion. If you are optimizing for speed, LCM (lazy code motion) based GCSE is done. LCM is based on the work of Knoop, Ruthing, and Steffen. LCM based GCSE also does loop invariant code motion. We also perform load and store motion when optimizing for speed. Regardless of which type of GCSE is used, the GCSE pass also performs global constant and copy propagation. The source file for this pass is 'gcse.c', and the LCM routines are in 'lcm.c'.

## • Loop optimization

This pass performs several loop related optimizations. The source files 'cfgloopanal.c' and 'cfgloopmanip.c' contain generic loop analysis and manipulation code. Initialization and finalization of loop structures is handled by 'loop-init.c'. A loop invariant motion pass is implemented in 'loop-invariant.c'. Basic block level optimizations—unrolling, and peeling loops— are implemented in 'loop-unroll.c'. Replacing of the exit condition of loops by special machine-dependent instructions is handled by 'loop-doloop.c'.

## • Jump bypassing

This pass is an aggressive form of GCSE that transforms the control flow graph of a function by propagating constants into conditional branch instructions. The source file for this pass is 'gcse.c'.

#### • If conversion

This pass attempts to replace conditional branches and surrounding assignments with arithmetic, boolean value producing comparison instructions, and conditional move instructions. In the very last invocation after reload/LRA, it will generate predicated instructions when supported by the target. The code is located in 'ifcvt.c'.

#### • Web construction

This pass splits independent uses of each pseudo-register. This can improve effect of the other transformation, such as CSE or register allocation. The code for this pass is located in 'web.c'.

#### • Instruction combination

This pass attempts to combine groups of two or three instructions that are related by data flow into single instructions. It combines the RTL expressions for the instructions by substitution, simplifies the result using algebra, and then attempts to match the result against the machine description. The code is located in 'combine.c'.

#### • Mode switching optimization

This pass looks for instructions that require the processor to be in a specific "mode" and minimizes the number of mode changes required to satisfy all users. What these modes are, and what they apply to are completely target-specific. The code for this pass is located in 'mode-switching.c'.

#### • Modulo scheduling

This pass looks at innermost loops and reorders their instructions by overlapping different iterations. Modulo scheduling is performed immediately before instruction scheduling. The code for this pass is located in 'modulo-sched.c'.

## • Instruction scheduling

This pass looks for instructions whose output will not be available by the time that it is used in subsequent instructions. Memory loads and floating point instructions often have this behavior on RISC machines. It re-orders instructions within a basic block to try to separate the definition and use of items that otherwise would cause pipeline stalls. This pass is performed twice, before and after register allocation. The code for this pass is located in 'haifa-sched.c', 'sched-deps.c', 'sched-ebb.c', 'sched-rgn.c' and 'sched-vis.c'.

## • Register allocation

These passes make sure that all occurrences of pseudo registers are eliminated, either by allocating them to a hard register, replacing them by an equivalent expression (e.g. a constant) or by placing them on the stack. This is done in several subpasses:

- The integrated register allocator (IRA). It is called integrated because coalescing, register live range splitting, and hard register preferencing are done on-the-fly during coloring. It also has better integration with the reload/LRA pass. Pseudoregisters spilled by the allocator or the reload/LRA have still a chance to get hard-registers if the reload/LRA evicts some pseudo-registers from hard-registers. The allocator helps to choose better pseudos for spilling based on their live ranges and to coalesce stack slots allocated for the spilled pseudo-registers. IRA is a regional register allocator which is transformed into Chaitin-Briggs allocator if there is one region. By default, IRA chooses regions using register pressure but the user can force it to use one region or regions corresponding to all loops.
  - Source files of the allocator are 'ira.c', 'ira-build.c', 'ira-costs.c', 'ira-conflicts.c', 'ira-color.c', 'ira-emit.c', 'ira-lives', plus header files 'ira.h' and 'ira-int.h' used for the communication between the allocator and the rest of the compiler and between the IRA files.
- Reloading. This pass renumbers pseudo registers with the hardware registers numbers they were allocated. Pseudo registers that did not get hard registers are replaced with stack slots. Then it finds instructions that are invalid because a value has failed to end up in a register, or has ended up in a register of the wrong kind. It fixes up these instructions by reloading the problematical values temporarily into registers. Additional instructions are generated to do the copying.

The reload pass also optionally eliminates the frame pointer and inserts instructions to save and restore call-clobbered registers around calls.

Source files are 'reload.c' and 'reload1.c', plus the header 'reload.h' used for communication between them.

• This pass is a modern replacement of the reload pass. Source files are 'lra.c', 'lra-assign.c', 'lra-coalesce.c', 'lra-constraints.c', 'lra-eliminations.c', 'lra-lives.c', 'lra-remat.c', 'lra-spills.c', the header 'lra-int.h' used for communication between them, and the header 'lra.h' used for communication between LRA and the rest of compiler.

Unlike the reload pass, intermediate LRA decisions are reflected in RTL as much as possible. This reduces the number of target-dependent macros and hooks, leaving instruction constraints as the primary source of control.

LRA is run on targets for which TARGET\_LRA\_P returns true.

# • Basic block reordering

This pass implements profile guided code positioning. If profile information is not available, various types of static analysis are performed to make the predictions normally coming from the profile feedback (IE execution frequency, branch probability, etc). It is implemented in the file 'bb-reorder.c', and the various prediction routines are in 'predict.c'.

# • Variable tracking

This pass computes where the variables are stored at each position in code and generates notes describing the variable locations to RTL code. The location lists are then generated according to these notes to debug information if the debugging information format supports location lists. The code is located in 'var-tracking.c'.

## • Delayed branch scheduling

This optional pass attempts to find instructions that can go into the delay slots of other instructions, usually jumps and calls. The code for this pass is located in 'reorg.c'.

#### • Branch shortening

On many RISC machines, branch instructions have a limited range. Thus, longer sequences of instructions must be used for long branches. In this pass, the compiler figures out what how far each instruction will be from each other instruction, and therefore whether the usual instructions, or the longer sequences, must be used for each branch. The code for this pass is located in 'final.c'.

## • Register-to-stack conversion

Conversion from usage of some hard registers to usage of a register stack may be done at this point. Currently, this is supported only for the floating-point registers of the Intel 80387 coprocessor. The code for this pass is located in 'reg-stack.c'.

#### • Final

This pass outputs the assembler code for the function. The source files are 'final.c' plus 'insn-output.c'; the latter is generated automatically from the machine description by the tool 'genoutput'. The header file 'conditions.h' is used for communication between these files.

## • Debugging information output

This is run after final because it must output the stack slot offsets for pseudo registers that did not get hard registers. Source files are 'dbxout.c' for DBX symbol table format, 'dwarfout.c' for DWARF symbol table format, files 'dwarf2out.c' and 'dwarf2asm.c' for DWARF2 symbol table format, and 'vmsdbgout.c' for VMS debug symbol table format.

# 9.7 Optimization info

This section is describes dump infrastructure which is common to both pass dumps as well as optimization dumps. The goal for this infrastructure is to provide both gcc developers and users detailed information about various compiler transformations and optimizations.

# 9.7.1 Dump setup

A dump\_manager class is defined in 'dumpfile.h'. Various passes register dumping pass-specific information via dump\_register in 'passes.c'. During the registration, an optimization pass can select its optimization group (see Section 9.7.2 [Optimization groups], page 143). After that optimization information corresponding to the entire group (presumably from multiple passes) can be output via command-line switches. Note that if a pass does not fit into any of the pre-defined groups, it can select OPTGROUP\_NONE.

Note that in general, a pass need not know its dump output file name, whether certain flags are enabled, etc. However, for legacy reasons, passes could also call dump\_begin which returns a stream in case the particular pass has optimization dumps enabled. A pass could call dump\_end when the dump has ended. These methods should go away once all the passes are converted to use the new dump infrastructure.

The recommended way to setup the dump output is via dump\_start and dump\_end.

# 9.7.2 Optimization groups

The optimization passes are grouped into several categories. Currently defined categories in 'dumpfile.h' are

OPTGROUP\_IPA

IPA optimization passes. Enabled by '-ipa'

OPTGROUP\_LOOP

Loop optimization passes. Enabled by '-loop'.

OPTGROUP\_INLINE

Inlining passes. Enabled by '-inline'.

OPTGROUP\_OMP

OMP (Offloading and Multi Processing) passes. Enabled by '-omp'.

OPTGROUP\_VEC

Vectorization passes. Enabled by '-vec'.

OPTGROUP\_OTHER

All other optimization passes which do not fall into one of the above.

OPTGROUP\_ALL

All optimization passes. Enabled by '-optall'.

By using groups a user could selectively enable optimization information only for a group of passes. By default, the optimization information for all the passes is dumped.

# 9.7.3 Dump files and streams

There are two separate output streams available for outputting optimization information from passes. Note that both these streams accept stderr and stdout as valid streams and

thus it is possible to dump output to standard output or error. This is specially handy for outputting all available information in a single file by redirecting stderr.

pstream

This stream is for pass-specific dump output. For example, '-fdump-tree-vect=foo.v' dumps tree vectorization pass output into the given file name 'foo.v'. If the file name is not provided, the default file name is based on the source file and pass number. Note that one could also use special file names stdout and stderr for dumping to standard output and standard error respectively.

#### alt\_stream

This steam is used for printing optimization specific output in response to the '-fopt-info'. Again a file name can be given. If the file name is not given, it defaults to stderr.

# 9.7.4 Dump output verbosity

The dump verbosity has the following options

'optimized'

Print information when an optimization is successfully applied. It is up to a pass to decide which information is relevant. For example, the vectorizer passes print the source location of loops which got successfully vectorized.

'missed'

Print information about missed optimizations. Individual passes control which information to include in the output. For example,

gcc -02 -ftree-vectorize -fopt-info-vec-missed

will print information about missed optimization opportunities from vectorization passes on stderr.

'note'

Print verbose information about optimizations, such as certain transformations, more detailed messages about decisions etc.

'all' Print detailed optimization information. This includes optimized, missed, and note.

# 9.7.5 Dump types

dump\_printf

This is a generic method for doing formatted output. It takes an additional argument dump\_kind which signifies the type of dump. This method outputs information only when the dumps are enabled for this particular dump\_kind. Note that the caller doesn't need to know if the particular dump is enabled or not, or even the file name. The caller only needs to decide which dump output information is relevant, and under what conditions. This determines the associated flags.

Consider the following example from 'loop-unroll.c' where an informative message about a loop (along with its location) is printed when any of the following flags is enabled

- optimization messages
- RTL dumps

detailed dumps

```
int report_flags = MSG_OPTIMIZED_LOCATIONS | TDF_RTL | TDF_DETAILS;
dump_printf_loc (report_flags, insn,
```

"loop turned into non-loop; it never loops.\n");

dump\_basic\_block

Output basic block.

dump\_generic\_expr

Output generic expression.

dump\_gimple\_stmt

Output gimple statement.

Note that the above methods also have variants prefixed with <code>\_loc</code>, such as <code>dump\_printf\_loc</code>, which are similar except they also output the source location information. The <code>\_loc</code> variants take a <code>const dump\_location\_t &</code>. This class can be constructed from a <code>gimple \*</code> or from a <code>rtx\_insn \*</code>, and so callers can pass a <code>gimple \*</code> or a <code>rtx\_insn \*</code> as the <code>\_loc</code> argument. The <code>dump\_location\_t</code> constructor will extract the source location from the statement or instruction, along with the profile count, and the location in GCC's own source code (or the plugin) from which the dump call was emitted. Only the source location is currently used. There is also a <code>dump\_user\_location\_t</code> class, capturing the source location and profile count, but not the dump emission location, so that locations in the user's code can be passed around. This can also be constructed from a <code>gimple \*</code> and from a <code>rtx\_insn \*</code>, and it too can be passed as the <code>\_loc</code> argument.

# 9.7.6 Dump examples

```
gcc -03 -fopt-info-missed=missed.all
```

outputs missed optimization report from all the passes into 'missed.all'.

As another example,

```
gcc -03 -fopt-info-inline-optimized-missed=inline.txt
```

will output information about missed optimizations as well as optimized locations from all the inlining passes into 'inline.txt'.

If the filename is provided, then the dumps from all the applicable optimizations are concatenated into the 'filename'. Otherwise the dump is output onto 'stderr'. If options is omitted, it defaults to 'optimized-optall', which means dump all information about successful optimizations from all the passes. In the following example, the optimization information is output on to 'stderr'.

```
gcc -03 -fopt-info
```

Note that '-fopt-info-vec-missed' behaves the same as '-fopt-info-missed-vec'. The order of the optimization group names and message types listed after '-fopt-info' does not matter.

As another example, consider

```
gcc -fopt-info-vec-missed=vec.miss -fopt-info-loop-optimized=loop.opt
```

Here the two output file names 'vec.miss' and 'loop.opt' are in conflict since only one output file is allowed. In this case, only the first option takes effect and the subsequent

options are ignored. Thus only the ' $\verb"vec.miss"$ ' is produced which containts dumps from the vectorizer about missed opportunities.

# 10 Sizes and offsets as runtime invariants

GCC allows the size of a hardware register to be a runtime invariant rather than a compiletime constant. This in turn means that various sizes and offsets must also be runtime invariants rather than compile-time constants, such as:

- the size of a general machine\_mode (see Section 14.6 [Machine Modes], page 271);
- the size of a spill slot;
- the offset of something within a stack frame;
- the number of elements in a vector;
- the size and offset of a mem rtx (see Section 14.8 [Regs and Memory], page 282); and
- the byte offset in a subreg rtx (see Section 14.8 [Regs and Memory], page 282).

The motivating example is the Arm SVE ISA, whose vector registers can be any multiple of 128 bits between 128 and 2048 inclusive. The compiler normally produces code that works for all SVE register sizes, with the actual size only being known at runtime.

GCC's main representation of such runtime invariants is the poly\_int class. This chapter describes what poly\_int does, lists the available operations, and gives some general usage guidelines.

# 10.1 Overview of poly\_int

We define indeterminates  $x1, \ldots, xn$  whose values are only known at runtime and use polynomials of the form:

```
c0 + c1 * x1 + ... + cn * xn
```

to represent a size or offset whose value might depend on some of these indeterminates. The coefficients  $c0, \ldots, cn$  are always known at compile time, with the c0 term being the "constant" part that does not depend on any runtime value.

GCC uses the poly\_int class to represent these coefficients. The class has two template parameters: the first specifies the number of coefficients (n + 1) and the second specifies the type of the coefficients. For example, 'poly\_int<2, unsigned short>' represents a polynomial with two coefficients (and thus one indeterminate), with each coefficient having type unsigned short. When n is 0, the class degenerates to a single compile-time constant c0.

The number of coefficients needed for compilation is a fixed property of each target and is specified by the configuration macro NUM\_POLY\_INT\_COEFFS. The default value is 1, since most targets do not have such runtime invariants. Targets that need a different value should #define the macro in their 'cpu-modes.def' file. See Section 6.3.9 [Back End], page 75.

poly\_int makes the simplifying requirement that each indeterminate must be a nonnegative integer. An indeterminate value of 0 should usually represent the minimum possible runtime value, with c0 specifying the value in that case.

For example, when targetting the Arm SVE ISA, the single indeterminate represents the number of 128-bit blocks in a vector beyond the minimum length of 128 bits. Thus the number of 64-bit doublewords in a vector is 2 + 2 \* x1. If an aggregate has a single SVE vector and 16 additional bytes, its total size is 32 + 16 \* x1 bytes.

The header file 'poly-int-types.h' provides typedefs for the most common forms of poly\_int, all having NUM\_POLY\_INT\_COEFFS coefficients:

Since the main purpose of poly\_int is to represent sizes and offsets, the last two typedefs are only rarely used.

# 10.2 Consequences of using poly\_int

The two main consequences of using polynomial sizes and offsets are that:

- there is no total ordering between the values at compile time, and
- some operations might yield results that cannot be expressed as a poly\_int.

For example, if x is a runtime invariant, we cannot tell at compile time whether:

```
3 + 4x <= 1 + 5x
```

since the condition is false when  $x \le 1$  and true when  $x \ge 2$ .

Similarly, poly\_int cannot represent the result of:

```
(3 + 4x) * (1 + 5x)
```

since it cannot (and in practice does not need to) store powers greater than one. It also cannot represent the result of:

```
(3 + 4x) / (1 + 5x)
```

The following sections describe how we deal with these restrictions.

As described earlier, a  $poly_int<1$ , T> has no indeterminates and so degenerates to a compile-time constant of type T. It would be possible in that case to do all normal arithmetic on the T, and to compare the T using the normal C++ operators. We deliberately prevent target-independent code from doing this, since the compiler needs to support other  $poly_int<n$ , T> as well, regardless of the current target's NUM\_POLY\_INT\_COEFFS.

However, it would be very artificial to force target-specific code to follow these restrictions if the target has no runtime indeterminates. There is therefore an implicit conversion from  $poly_int<1$ , T> to T when compiling target-specific translation units.

# 10.3 Comparisons involving poly\_int

In general we need to compare sizes and offsets in two situations: those in which the values need to be ordered, and those in which the values can be unordered. More loosely, the distinction is often between values that have a definite link (usually because they refer to the same underlying register or memory location) and values that have no definite link. An example of the former is the relationship between the inner and outer sizes of a subreg, where we must know at compile time whether the subreg is paradoxical, partial, or complete. An example of the latter is alias analysis: we might want to check whether two arbitrary memory references overlap.

Referring back to the examples in the previous section, it makes sense to ask whether a memory reference of size '3 + 4x' overlaps one of size '1 + 5x', but it does not make sense to have a subreg in which the outer mode has '3 + 4x' bytes and the inner mode has '1 + 5x' bytes (or vice versa). Such subregs are always invalid and should trigger an internal compiler error if formed.

The underlying operators are the same in both cases, but the distinction affects how they are used.

# 10.3.1 Comparison functions for poly\_int

poly\_int provides the following routines for checking whether a particular condition "may be" (might be) true:

For readability, poly\_int also provides "known" inverses of these functions:

```
known_lt (a, b) == !maybe_ge (a, b)
known_le (a, b) == !maybe_gt (a, b)
known_eq (a, b) == !maybe_ne (a, b)
known_ge (a, b) == !maybe_lt (a, b)
known_gt (a, b) == !maybe_le (a, b)
known_ne (a, b) == !maybe_eq (a, b)
```

# 10.3.2 Properties of the poly\_int comparisons

```
All "maybe" relations except maybe_ne are transitive, so for example:
      maybe_lt (a, b) && maybe_lt (b, c) implies maybe_lt (a, c)
  for all a, b and c. maybe_lt, maybe_gt and maybe_ne are irreflexive, so for example:
      !maybe_lt (a, a)
  is true for all a. maybe_le, maybe_eq and maybe_ge are reflexive, so for example:
      maybe_le (a, a)
  is true for all a. maybe_eq and maybe_ne are symmetric, so:
      maybe_eq (a, b) == maybe_eq (b, a)
      maybe_ne(a, b) == maybe_ne(b, a)
  for all a and b. In addition:
      maybe_le(a, b) == maybe_lt(a, b) \mid | maybe_eq(a, b)
     maybe_ge (a, b) == maybe_gt (a, b) \mid \mid maybe_eq (a, b)
     maybe_lt (a, b) == maybe_gt (b, a)
     maybe_le (a, b) == maybe_ge (b, a)
  However:
      maybe_le (a, b) && maybe_le (b, a) does not imply !maybe_ne (a, b) [== known_eq (a, b)]
     maybe_ge (a, b) && maybe_ge (b, a) does not imply !maybe_ne (a, b) [== known_eq (a, b)]
```

One example is again 'a == 3 + 4x' and 'b == 1 + 5x', where 'maybe\_le (a, b)', 'maybe\_ge (a, b)' and 'maybe\_ne (a, b)' all hold. maybe\_le and maybe\_ge are therefore not antisymetric and do not form a partial order.

From the above, it follows that:

- All "known" relations except known\_ne are transitive.
- known\_lt, known\_ne and known\_gt are irreflexive.
- known\_le, known\_eq and known\_ge are reflexive.

Also:

```
known_lt (a, b) == known_gt (b, a)
known_le (a, b) == known_ge (b, a)
known_lt (a, b) implies !known_lt (b, a) [asymmetry]
known_gt (a, b) implies !known_gt (b, a)
known_le (a, b) && known_le (b, a) == known_eq (a, b) [== !maybe_ne (a, b)]
known_ge (a, b) && known_ge (b, a) == known_eq (a, b) [== !maybe_ne (a, b)]
```

known\_le and known\_ge are therefore antisymmetric and are partial orders. However:

```
known_le (a, b) does not imply known_lt (a, b) || known_eq (a, b) known_ge (a, b) does not imply known_gt (a, b) || known_eq (a, b)
```

For example, 'known\_le (4, 4 + 4x)' holds because the runtime indeterminate x is a nonnegative integer, but neither known\_lt (4, 4 + 4x) nor known\_eq (4, 4 + 4x) hold.

## 10.3.3 Comparing potentially-unordered poly\_ints

In cases where there is no definite link between two poly\_ints, we can usually make a conservatively-correct assumption. For example, the conservative assumption for alias analysis is that two references *might* alias.

One way of checking whether [begin1, end1) might overlap [begin2, end2) using the poly\_int comparisons is:

```
maybe_gt (end1, begin2) && maybe_gt (end2, begin1)
and another (equivalent) way is:
  !(known_le (end1, begin2) || known_le (end2, begin1))
```

However, in this particular example, it is better to use the range helper functions instead. See Section 10.3.6 [Range checks on poly\_ints], page 152.

# 10.3.4 Comparing ordered poly\_ints

In cases where there is a definite link between two poly\_ints, such as the outer and inner sizes of subregs, we usually require the sizes to be ordered by the known\_le partial order. poly\_int provides the following utility functions for ordered values:

```
'ordered_p (a, b)'
```

Return true if a and b are ordered by the known\_le partial order.

```
'ordered_min (a, b)'
```

Assert that a and b are ordered by known\_le and return the minimum of the two. When using this function, please add a comment explaining why the values are known to be ordered.

```
'ordered_max (a, b)'
```

Assert that a and b are ordered by known\_le and return the maximum of the two. When using this function, please add a comment explaining why the values are known to be ordered.

For example, if a subreg has an outer mode of size outer and an inner mode of size inner:

- the subreg is complete if known\_eq (inner, outer)
- otherwise, the subreg is paradoxical if known\_le (inner, outer)
- otherwise, the subreg is partial if known\_le (outer, inner)
- otherwise, the subreg is ill-formed

Thus the subreg is only valid if 'ordered\_p (outer, inner)' is true. If this condition is already known to be true then:

- the subreg is complete if known\_eq (inner, outer)
- the subreg is paradoxical if maybe\_lt (inner, outer)
- the subreg is partial if maybe\_lt (outer, inner)

with the three conditions being mutually exclusive.

Code that checks whether a subreg is valid would therefore generally check whether ordered\_p holds (in addition to whatever other checks are required for subreg validity). Code that is dealing with existing subregs can assert that ordered\_p holds and use either of the classifications above.

# 10.3.5 Checking for a poly\_int marker value

It is sometimes useful to have a special "marker value" that is not meant to be taken literally. For example, some code uses a size of -1 to represent an unknown size, rather than having to carry around a separate boolean to say whether the size is known.

The best way of checking whether something is a marker value is known\_eq. Conversely the best way of checking whether something is not a marker value is maybe\_ne.

Thus in the size example just mentioned, 'known\_eq (size, -1)' would check for an unknown size and 'maybe\_ne (size, -1)' would check for a known size.

# 10.3.6 Range checks on poly\_ints

As well as the core comparisons (see Section 10.3.1 [Comparison functions for poly\_int], page 149), poly\_int provides utilities for various kinds of range check. In each case the range is represented by a start position and a size rather than a start position and an end position; this is because the former is used much more often than the latter in GCC. Also, the sizes can be -1 (or all ones for unsigned sizes) to indicate a range with a known start position but an unknown size. All other sizes must be nonnegative. A range of size 0 does not contain anything or overlap anything.

## 'known\_size\_p (size)'

Return true if *size* represents a known range size, false if it is -1 or all ones (for signed and unsigned types respectively).

## 'ranges\_maybe\_overlap\_p (pos1, size1, pos2, size2)'

Return true if the range described by pos1 and size1 might overlap the range described by pos2 and size2 (in other words, return true if we cannot prove that the ranges are disjoint).

# 'ranges\_known\_overlap\_p (pos1, size1, pos2, size2)'

Return true if the range described by *pos1* and *size1* is known to overlap the range described by *pos2* and *size2*.

#### 'known\_subrange\_p (pos1, size1, pos2, size2)'

Return true if the range described by pos1 and size1 is known to be contained in the range described by pos2 and size2.

## 'maybe\_in\_range\_p (value, pos, size)'

Return true if value might be in the range described by pos and size (in other words, return true if we cannot prove that value is outside that range).

#### 'known\_in\_range\_p (value, pos, size)'

Return true if value is known to be in the range described by pos and size.

#### 'endpoint\_representable\_p (pos, size)'

Return true if the range described by *pos* and *size* is open-ended or if the endpoint (*pos* + *size*) is representable in the same type as *pos* and *size*. The function returns false if adding *size* to *pos* makes conceptual sense but could overflow.

There is also a poly\_int version of the IN\_RANGE\_P macro:

# 'coeffs\_in\_range\_p (x, lower, upper)'

Return true if every coefficient of x is in the inclusive range [lower, upper]. This function can be useful when testing whether an operation would cause the values of coefficients to overflow.

Note that the function does not indicate whether x itself is in the given range. x can be either a constant or a poly\_int.

# 10.3.7 Sorting poly\_ints

poly\_int provides the following routine for sorting:

```
'compare_sizes_for_sort (a, b)'
```

Compare a and b in reverse lexicographical order (that is, compare the highest-indexed coefficients first). This can be useful when sorting data structures, since it has the effect of separating constant and non-constant values. If all values are nonnegative, the constant values come first.

Note that the values do not necessarily end up in numerical order. For example,  $^{1} + 1x^{2}$  would come after  $^{1}00^{2}$  in the sort order, but may well be less than  $^{1}00^{2}$  at run time.

# 10.4 Arithmetic on poly\_ints

Addition, subtraction, negation and bit inversion all work normally for poly\_ints. Multiplication by a constant multiplier and left shifting by a constant shift amount also work normally. General multiplication of two poly\_ints is not supported and is not useful in practice.

Other operations are only conditionally supported: the operation might succeed or might fail, depending on the inputs.

This section describes both types of operation.

# 10.4.1 Using poly\_int with C++ arithmetic operators

The following C++ expressions are supported, where p1 and p2 are poly\_ints and where c1 and c2 are scalars:

```
-p1
~p1 + p2
p1 + c2
c1 + p2
p1 - c2
c1 - p2
c1 - p2
p1 - c2
c1 - p2
c1 * p2
p1 * c2
p1 * c2
p1 += c2
p1 -= c2
p1 -= c2
p1 *= c2
p1 *= c2
p1 *= c2
p1 *= c2
```

These arithmetic operations handle integer ranks in a similar way to C++. The main difference is that every coefficient narrower than HOST\_WIDE\_INT promotes to HOST\_WIDE\_INT, whereas in C++ everything narrower than int promotes to int. For example:

In the first two examples, both coefficients are narrower than <code>HOST\_WIDE\_INT</code>, so the result has coefficients of type <code>HOST\_WIDE\_INT</code>. In the other examples, the coefficient with the highest rank "wins".

If one of the operands is wide\_int or poly\_wide\_int, the rules are the same as for wide\_int arithmetic.

# 10.4.2 wi arithmetic on poly\_ints

As well as the C++ operators, poly\_int supports the following wi routines:

```
wi::neg (p1, &overflow)
wi::add (p1, p2)
wi::add (p1, c2)
wi::add (c1, p1)
wi::add (p1, p2, sign, &overflow)
wi::sub (p1, p2)
wi::sub (p1, c2)
wi::sub (c1, p1)
wi::sub (p1, p2, sign, &overflow)
wi::mul (p1, c2)
wi::mul (c1, p1)
wi::mul (c1, p1)
wi::mul (p1, c2, sign, &overflow)
wi::lshift (p1, c2)
```

These routines just check whether overflow occurs on any individual coefficient; it is not possible to know at compile time whether the final runtime value would overflow.

# 10.4.3 Division of poly\_ints

Division of poly\_ints is possible for certain inputs. The functions for division return true if the operation is possible and in most cases return the results by pointer. The routines are:

```
'multiple_p (a, b)'
'multiple_p (a, b, &quotient)'
```

Return true if a is an exact multiple of b, storing the result in *quotient* if so. There are overloads for various combinations of polynomial and constant a, b and *quotient*.

```
'constant_multiple_p (a, b)'
'constant_multiple_p (a, b, &quotient)'
```

Like multiple\_p, but also test whether the multiple is a compile-time constant.

'can\_div\_trunc\_p (a, b, &quotient)'

'can\_div\_trunc\_p (a, b, &quotient, &remainder)'

Return true if we can calculate 'trunc (a / b)' at compile time, storing the result in quotient and remainder if so.

'can\_div\_away\_from\_zero\_p (a, b, &quotient)'

Return true if we can calculate 'a / b' at compile time, rounding away from zero. Store the result in *quotient* if so.

Note that this is true if and only if can\_div\_trunc\_p is true. The only difference is in the rounding of the result.

There is also an asserting form of division:

'exact\_div (a, b)'

Assert that a is a multiple of b and return 'a / b'. The result is a poly\_int if a is a poly\_int.

# 10.4.4 Other poly\_int arithmetic

There are tentative routines for other operations besides division:

'can\_ior\_p (a, b, &result)'

Return true if we can calculate 'a | b' at compile time, storing the result in result if so.

Also, ANDs with a value '(1 << y) - 1' or its inverse can be treated as alignment operations. See Section 10.5 [Alignment of poly\_ints], page 155.

In addition, the following miscellaneous routines are available:

'coeff\_gcd (a)'

Return the greatest common divisor of all nonzero coefficients in a, or zero if a is known to be zero.

'common\_multiple (a, b)'

Return a value that is a multiple of both a and b, where one value is a poly\_int and the other is a scalar. The result will be the least common multiple for some indeterminate values but not necessarily for all.

'force\_common\_multiple (a, b)'

Return a value that is a multiple of both poly\_int a and poly\_int b, asserting that such a value exists. The result will be the least common multiple for some indeterminate values but not necessarily for all.

When using this routine, please add a comment explaining why the assertion is known to hold.

Please add any other operations that you find to be useful.

# 10.5 Alignment of poly\_ints

poly\_int provides various routines for aligning values and for querying misalignments. In each case the alignment must be a power of 2.

## 'can\_align\_p (value, align)'

Return true if we can align value up or down to the nearest multiple of align at compile time. The answer is the same for both directions.

# 'can\_align\_down (value, align, &aligned)'

Return true if can\_align\_p; if so, set aligned to the greatest aligned value that is less than or equal to value.

# 'can\_align\_up (value, align, &aligned)'

Return true if can\_align\_p; if so, set aligned to the lowest aligned value that is greater than or equal to value.

## 'known\_equal\_after\_align\_down (a, b, align)'

Return true if we can align a and b down to the nearest align boundary at compile time and if the two results are equal.

# 'known\_equal\_after\_align\_up (a, b, align)'

Return true if we can align a and b up to the nearest align boundary at compile time and if the two results are equal.

## 'aligned\_lower\_bound (value, align)'

Return a result that is no greater than value and that is aligned to align. The result will the closest aligned value for some indeterminate values but not necessarily for all.

For example, suppose we are allocating an object of *size* bytes in a downward-growing stack whose current limit is given by *limit*. If the object requires *align* bytes of alignment, the new stack limit is given by:

aligned\_lower\_bound (limit - size, align)

#### 'aligned\_upper\_bound (value, align)'

Likewise return a result that is no less than *value* and that is aligned to *align*. This is the routine that would be used for upward-growing stacks in the scenario just described.

# 'known\_misalignment (value, align, &misalign)'

Return true if we can calculate the misalignment of value with respect to align at compile time, storing the result in misalign if so.

#### 'known\_alignment (value)'

Return the minimum alignment that *value* is known to have (in other words, the largest alignment that can be guaranteed whatever the values of the indeterminates turn out to be). Return 0 if *value* is known to be 0.

#### 'force\_align\_down (value, align)'

Assert that value can be aligned down to align at compile time and return the result. When using this routine, please add a comment explaining why the assertion is known to hold.

#### 'force\_align\_up (value, align)'

Likewise, but aligning up.

#### 'force\_align\_down\_and\_div (value, align)'

Divide the result of force\_align\_down by align. Again, please add a comment explaining why the assertion in force\_align\_down is known to hold.

'force\_align\_up\_and\_div (value, align)'
Likewise for force\_align\_up.

'force\_get\_misalignment (value, align)'

Assert that we can calculate the misalignment of value with respect to align at compile time and return the misalignment. When using this function, please add a comment explaining why the assertion is known to hold.

# 10.6 Computing bounds on poly\_ints

poly\_int also provides routines for calculating lower and upper bounds:

'constant\_lower\_bound (a)'

Assert that a is nonnegative and return the smallest value it can have.

'constant\_lower\_bound\_with\_limit (a, b)'

Return the least value a can have, given that the context in which a appears guarantees that the answer is no less than b. In other words, the caller is asserting that a is greater than or equal to b even if 'known\_ge (a, b)' doesn't hold.

'constant\_upper\_bound\_with\_limit (a, b)'

Return the greatest value a can have, given that the context in which a appears guarantees that the answer is no greater than b. In other words, the caller is asserting that a is less than or equal to b even if 'known\_le (a, b)' doesn't hold.

'lower\_bound (a, b)'

Return a value that is always less than or equal to both a and b. It will be the greatest such value for some indeterminate values but necessarily for all.

'upper\_bound (a, b)'

Return a value that is always greater than or equal to both a and b. It will be the least such value for some indeterminate values but necessarily for all.

# 10.7 Converting poly\_ints

A poly\_int<n, T> can be constructed from up to n individual T coefficients, with the remaining coefficients being implicitly zero. In particular, this means that every poly\_int<n, T> can be constructed from a single scalar T, or something compatible with T.

Also, a poly\_int<n, T> can be constructed from a poly\_int<n, U> if T can be constructed from U.

The following functions provide other forms of conversion, or test whether such a conversion would succeed.

'value.is\_constant ()'

Return true if poly\_int value is a compile-time constant.

'value.is\_constant (&c1)'

Return true if  $poly_int$  value is a compile-time constant, storing it in c1 if so. c1 must be able to hold all constant values of value without loss of precision.

## 'value.to\_constant ()'

Assert that value is a compile-time constant and return its value. When using this function, please add a comment explaining why the condition is known to hold (for example, because an earlier phase of analysis rejected non-constants).

#### 'value.to\_shwi (&p2)'

Return true if 'poly\_int<N, T>' value can be represented without loss of precision as a 'poly\_int<N, HOST\_WIDE\_INT>', storing it in that form in p2 if so.

## 'value.to\_uhwi (&p2)'

Return true if 'poly\_int<N, T>' value can be represented without loss of precision as a 'poly\_int<N, unsigned HOST\_WIDE\_INT>', storing it in that form in p2 if so.

## 'value.force\_shwi ()'

Forcibly convert each coefficient of 'poly\_int<N, T>' value to HOST\_WIDE\_INT, truncating any that are out of range. Return the result as a 'poly\_int<N, HOST\_WIDE\_INT>'.

## 'value.force\_uhwi ()'

Forcibly convert each coefficient of 'poly\_int<N, T>' value to unsigned HOST\_WIDE\_INT, truncating any that are out of range. Return the result as a 'poly\_int<N, unsigned HOST\_WIDE\_INT>'.

## 'wi::shwi (value, precision)'

Return a poly\_int with the same value as value, but with the coefficients converted from HOST\_WIDE\_INT to wide\_int. precision specifies the precision of the wide\_int coefficients; if this is wider than a HOST\_WIDE\_INT, the coefficients of value will be sign-extended to fit.

#### 'wi::uhwi (value, precision)'

Like wi::shwi, except that value has coefficients of type unsigned HOST\_WIDE\_INT. If precision is wider than a HOST\_WIDE\_INT, the coefficients of value will be zero-extended to fit.

## 'wi::sext (value, precision)'

Return a poly\_int of the same type as value, sign-extending every coefficient from the low precision bits. This in effect applies wi::sext to each coefficient individually.

## 'wi::zext (value, precision)'

Like wi::sext, but for zero extension.

## 'poly\_wide\_int::from (value, precision, sign)'

Convert value to a poly\_wide\_int in which each coefficient has precision bits. Extend the coefficients according to sign if the coefficients have fewer bits.

## 'poly\_offset\_int::from (value, sign)'

Convert value to a poly\_offset\_int, extending its coefficients according to sign if they have fewer bits than offset\_int.

# 'poly\_widest\_int::from (value, sign)'

Convert value to a poly\_widest\_int, extending its coefficients according to sign if they have fewer bits than widest\_int.

# 10.8 Miscellaneous poly\_int routines

```
'print_dec (value, file, sign)'
'print_dec (value, file)'
```

Print value to file as a decimal value, interpreting the coefficients according to sign. The final argument is optional if value has an inherent sign; for example, poly\_int64 values print as signed by default and poly\_uint64 values print as unsigned by default.

This is a simply a poly\_int version of a wide-int routine.

# 10.9 Guidelines for using poly\_int

One of the main design goals of poly\_int was to make it easy to write target-independent code that handles variable-sized registers even when the current target has fixed-sized registers. There are two aspects to this:

- The set of poly\_int operations should be complete enough that the question in most cases becomes "Can we do this operation on these particular poly\_int values? If not, bail out" rather than "Are these poly\_int values constant? If so, do the operation, otherwise bail out".
- If target-independent code compiles and runs correctly on a target with one value
  of NUM\_POLY\_INT\_COEFFS, and if the code does not use asserting functions like to\_
  constant, it is reasonable to assume that the code also works on targets with other
  values of NUM\_POLY\_INT\_COEFFS. There is no need to check this during everyday development.

So the general principle is: if target-independent code is dealing with a poly\_int value, it is better to operate on it as a poly\_int if at all possible, choosing conservatively-correct behavior if a particular operation fails. For example, the following code handles an index pos into a sequence of vectors that each have nunits elements:

However, there are some contexts in which operating on a poly\_int is not possible or does not make sense. One example is when handling static initializers, since no current target supports the concept of a variable-length static initializer. In these situations, a reasonable fallback is:

```
if (poly_value.is_constant (&const_value))
   {
      ...
      /* Operate on const_value. */
      ...
   }
else
```

```
{
    ...
    /* Conservatively correct fallback. */
    ...
}
```

poly\_int also provides some asserting functions like to\_constant. Please only use these functions if there is a good theoretical reason to believe that the assertion cannot fire. For example, if some work is divided into an analysis phase and an implementation phase, the analysis phase might reject inputs that are not is\_constant, in which case the implementation phase can reasonably use to\_constant on the remaining inputs. The assertions should not be used to discover whether a condition ever occurs "in the field"; in other words, they should not be used to restrict code to constants at first, with the intention of only implementing a poly\_int version if a user hits the assertion.

If a particular asserting function like to\_constant is needed more than once for the same reason, it is probably worth adding a helper function or macro for that situation, so that the justification only needs to be given once. For example:

```
/* Return the size of an element in a vector of size SIZE, given that
    the vector has NELTS elements. The return value is in the same units
    as SIZE (either bits or bytes).

to_constant () is safe in this situation because vector elements are
    always constant-sized scalars. */
#define vector_element_size(SIZE, NELTS) \
    (exact_div (SIZE, NELTS).to_constant ())
```

Target-specific code in 'config/cpu' only needs to handle non-constant poly\_ints if NUM\_POLY\_INT\_COEFFS is greater than one. For other targets, poly\_int degenerates to a compile-time constant and is often interchangable with a normal scalar integer. There are two main exceptions:

- Sometimes an explicit cast to an integer type might be needed, such as to resolve ambiguities in a ?: expression, or when passing values through . . . to things like print functions.
- Target macros are included in target-independent code and so do not have access to the implicit conversion to a scalar integer. If this becomes a problem for a particular target macro, the possible solutions, in order of preference, are:
  - Convert the target macro to a target hook (for all targets).
  - Put the target's implementation of the target macro in its 'cpu.c' file and call it from the target macro in the 'cpu.h' file.
  - Add to\_constant () calls where necessary. The previous option is preferable because it will help with any future conversion of the macro to a hook.

# 11 GENERIC

The purpose of GENERIC is simply to provide a language-independent way of representing an entire function in trees. To this end, it was necessary to add a few new tree codes to the back end, but almost everything was already there. If you can express it with the codes in gcc/tree.def, it's GENERIC.

Early on, there was a great deal of debate about how to think about statements in a tree IL. In GENERIC, a statement is defined as any expression whose value, if any, is ignored. A statement will always have TREE\_SIDE\_EFFECTS set (or it will be discarded), but a non-statement expression may also have side effects. A CALL\_EXPR, for instance.

It would be possible for some local optimizations to work on the GENERIC form of a function; indeed, the adapted tree inliner works fine on GENERIC, but the current compiler performs inlining after lowering to GIMPLE (a restricted form described in the next section). Indeed, currently the frontends perform this lowering before handing off to tree\_rest\_of\_compilation, but this seems inelegant.

# 11.1 Deficiencies

There are many places in which this document is incomplet and incorrekt. It is, as of yet, only *preliminary* documentation.

# 11.2 Overview

The central data structure used by the internal representation is the tree. These nodes, while all of the C type tree, are of many varieties. A tree is a pointer type, but the object to which it points may be of a variety of types. From this point forward, we will refer to trees in ordinary type, rather than in this font, except when talking about the actual C type tree.

You can tell what kind of node a particular tree is by using the TREE\_CODE macro. Many, many macros take trees as input and return trees as output. However, most macros require a certain kind of tree node as input. In other words, there is a type-system for trees, but it is not reflected in the C type-system.

For safety, it is useful to configure GCC with '--enable-checking'. Although this results in a significant performance penalty (since all tree types are checked at run-time), and is therefore inappropriate in a release version, it is extremely helpful during the development process.

Many macros behave as predicates. Many, although not all, of these predicates end in '\_P'. Do not rely on the result type of these macros being of any particular type. You may, however, rely on the fact that the type can be compared to 0, so that statements like

are legal. Macros that return int values now may be changed to return tree values, or other pointers in the future. Even those that continue to return int may return multiple nonzero codes where previously they returned only zero and one. Therefore, you should not write code like

```
if (TEST_P (t) == 1)
```

as this code is not guaranteed to work correctly in the future.

You should not take the address of values returned by the macros or functions described here. In particular, no guarantee is given that the values are lyalues.

In general, the names of macros are all in uppercase, while the names of functions are entirely in lowercase. There are rare exceptions to this rule. You should assume that any macro or function whose name is made up entirely of uppercase letters may evaluate its arguments more than once. You may assume that a macro or function whose name is made up entirely of lowercase letters will evaluate its arguments only once.

The error\_mark\_node is a special tree. Its tree code is ERROR\_MARK, but since there is only ever one node with that code, the usual practice is to compare the tree against error\_mark\_node. (This test is just a test for pointer equality.) If an error has occurred during front-end processing the flag errorcount will be set. If the front end has encountered code it cannot handle, it will issue a message to the user and set sorrycount. When these flags are set, any macro or function which normally returns a tree of a particular kind may instead return the error\_mark\_node. Thus, if you intend to do any processing of erroneous code, you must be prepared to deal with the error\_mark\_node.

Occasionally, a particular tree slot (like an operand to an expression, or a particular field in a declaration) will be referred to as "reserved for the back end". These slots are used to store RTL when the tree is converted to RTL for use by the GCC back end. However, if that process is not taking place (e.g., if the front end is being hooked up to an intelligent editor), then those slots may be used by the back end presently in use.

If you encounter situations that do not match this documentation, such as tree nodes of types not mentioned here, or macros documented to return entities of a particular kind that instead return entities of some different kind, you have found a bug, either in the front end or in the documentation. Please report these bugs as you would any other bug.

## 11.2.1 Trees

All GENERIC trees have two fields in common. First, TREE\_CHAIN is a pointer that can be used as a singly-linked list to other trees. The other is TREE\_TYPE. Many trees store the type of an expression or declaration in this field.

These are some other functions for handling trees:

#### tree\_size

Return the number of bytes a tree takes.

build0
build1
build2
build3
build4
build5
build5

These functions build a tree and supply values to put in each parameter. The basic signature is 'code, type, [operands]'. code is the TREE\_CODE, and type is a tree representing the TREE\_TYPE. These are followed by the operands, each of which is also a tree.

## 11.2.2 Identifiers

An IDENTIFIER\_NODE represents a slightly more general concept than the standard C or C++ concept of identifier. In particular, an IDENTIFIER\_NODE may contain a '\$', or other extraordinary characters.

There are never two distinct IDENTIFIER\_NODEs representing the same identifier. Therefore, you may use pointer equality to compare IDENTIFIER\_NODEs, rather than using a routine like strcmp. Use get\_identifier to obtain the unique IDENTIFIER\_NODE for a supplied string.

You can use the following macros to access identifiers:

#### IDENTIFIER\_POINTER

The string represented by the identifier, represented as a char\*. This string is always NUL-terminated, and contains no embedded NUL characters.

#### IDENTIFIER\_LENGTH

The length of the string returned by IDENTIFIER\_POINTER, not including the trailing NUL. This value of IDENTIFIER\_LENGTH (x) is always the same as strlen (IDENTIFIER\_POINTER (x)).

#### IDENTIFIER\_OPNAME\_P

This predicate holds if the identifier represents the name of an overloaded operator. In this case, you should not depend on the contents of either the IDENTIFIER\_POINTER or the IDENTIFIER\_LENGTH.

#### IDENTIFIER\_TYPENAME\_P

This predicate holds if the identifier represents the name of a user-defined conversion operator. In this case, the TREE\_TYPE of the IDENTIFIER\_NODE holds the type to which the conversion operator converts.

#### 11.2.3 Containers

Two common container data structures can be represented directly with tree nodes. A TREE\_LIST is a singly linked list containing two trees per node. These are the TREE\_PURPOSE and TREE\_VALUE of each node. (Often, the TREE\_PURPOSE contains some kind of tag, or additional information, while the TREE\_VALUE contains the majority of the payload. In other cases, the TREE\_PURPOSE is simply NULL\_TREE, while in still others both the TREE\_PURPOSE and TREE\_VALUE are of equal stature.) Given one TREE\_LIST node, the next node is found by following the TREE\_CHAIN. If the TREE\_CHAIN is NULL\_TREE, then you have reached the end of the list.

A TREE\_VEC is a simple vector. The TREE\_VEC\_LENGTH is an integer (not a tree) giving the number of nodes in the vector. The nodes themselves are accessed using the TREE\_VEC\_ELT macro, which takes two arguments. The first is the TREE\_VEC in question; the second is an integer indicating which element in the vector is desired. The elements are indexed from zero.

# 11.3 Types

All types have corresponding tree nodes. However, you should not assume that there is exactly one tree node corresponding to each type. There are often multiple nodes corresponding to the same type.

For the most part, different kinds of types have different tree codes. (For example, pointer types use a POINTER\_TYPE code while arrays use an ARRAY\_TYPE code.) However, pointers to member functions use the RECORD\_TYPE code. Therefore, when writing a switch statement that depends on the code associated with a particular type, you should take care to handle pointers to member functions under the RECORD\_TYPE case label.

The following functions and macros deal with cv-qualification of types:

#### TYPE\_MAIN\_VARIANT

This macro returns the unqualified version of a type. It may be applied to an unqualified type, but it is not always the identity function in that case.

A few other macros and functions are usable with all types:

#### TYPE\_SIZE

The number of bits required to represent the type, represented as an INTEGER\_CST. For an incomplete type, TYPE\_SIZE will be NULL\_TREE.

#### TYPE\_ALIGN

The alignment of the type, in bits, represented as an int.

#### TYPE\_NAME

This macro returns a declaration (in the form of a TYPE\_DECL) for the type. (Note this macro does *not* return an IDENTIFIER\_NODE, as you might expect, given its name!) You can look at the DECL\_NAME of the TYPE\_DECL to obtain the actual name of the type. The TYPE\_NAME will be NULL\_TREE for a type that is not a built-in type, the result of a typedef, or a named class type.

#### TYPE\_CANONICAL

This macro returns the "canonical" type for the given type node. Canonical types are used to improve performance in the C++ and Objective-C++ front ends by allowing efficient comparison between two type nodes in same\_type\_p: if the TYPE\_CANONICAL values of the types are equal, the types are equivalent; otherwise, the types are not equivalent. The notion of equivalence for canonical types is the same as the notion of type equivalence in the language itself. For instance,

When TYPE\_CANONICAL is NULL\_TREE, there is no canonical type for the given type node. In this case, comparison between this type and any other type requires the compiler to perform a deep, "structural" comparison to see if the two type nodes have the same form and properties.

The canonical type for a node is always the most fundamental type in the equivalence class of types. For instance, int is its own canonical type. A typedef I of int will have int as its canonical type. Similarly, I\* and a typedef IP (defined to I\*) will has int\* as their canonical type. When building a new type node, be sure to set TYPE\_CANONICAL to the appropriate canonical type. If the new type is a compound type (built from other types), and any of those other types require structural equality, use SET\_TYPE\_STRUCTURAL\_EQUALITY to ensure that the new type also requires structural equality. Finally, if for some reason you cannot guarantee that TYPE\_CANONICAL will point to the canonical type, use SET\_TYPE\_STRUCTURAL\_EQUALITY to make sure that the new type—and any type constructed based on it—requires structural equality. If you suspect

that the canonical type system is miscomparing types, pass --param verify-canonical-types=1 to the compiler or configure with --enable-checking to force the compiler to verify its canonical-type comparisons against the structural comparisons; the compiler will then print any warnings if the canonical types miscompare.

## TYPE\_STRUCTURAL\_EQUALITY\_P

This predicate holds when the node requires structural equality checks, e.g., when TYPE\_CANONICAL is NULL\_TREE.

### SET\_TYPE\_STRUCTURAL\_EQUALITY

This macro states that the type node it is given requires structural equality checks, e.g., it sets TYPE\_CANONICAL to NULL\_TREE.

### same\_type\_p

This predicate takes two types as input, and holds if they are the same type. For example, if one type is a typedef for the other, or both are typedefs for the same type. This predicate also holds if the two trees given as input are simply copies of one another; i.e., there is no difference between them at the source level, but, for whatever reason, a duplicate has been made in the representation. You should never use == (pointer equality) to compare types; always use same\_type\_p instead.

Detailed below are the various kinds of types, and the macros that can be used to access them. Although other kinds of types are used elsewhere in G++, the types described here are the only ones that you will encounter while examining the intermediate representation.

## VOID\_TYPE

Used to represent the void type.

### INTEGER\_TYPE

Used to represent the various integral types, including char, short, int, long, and long long. This code is not used for enumeration types, nor for the bool type. The TYPE\_PRECISION is the number of bits used in the representation, represented as an unsigned int. (Note that in the general case this is not the same value as TYPE\_SIZE; suppose that there were a 24-bit integer type, but that alignment requirements for the ABI required 32-bit alignment. Then, TYPE\_SIZE would be an INTEGER\_CST for 32, while TYPE\_PRECISION would be 24.) The integer type is unsigned if TYPE\_UNSIGNED holds; otherwise, it is signed.

The TYPE\_MIN\_VALUE is an INTEGER\_CST for the smallest integer that may be represented by this type. Similarly, the TYPE\_MAX\_VALUE is an INTEGER\_CST for the largest integer that may be represented by this type.

### REAL\_TYPE

Used to represent the float, double, and long double types. The number of bits in the floating-point representation is given by TYPE\_PRECISION, as in the INTEGER\_TYPE case.

## FIXED\_POINT\_TYPE

Used to represent the short \_Fract, \_Fract, long \_Fract, long long \_Fract, short \_Accum, \_Accum, long \_Accum, and long long \_Accum types. The num-

ber of bits in the fixed-point representation is given by TYPE\_PRECISION, as in the INTEGER\_TYPE case. There may be padding bits, fractional bits and integral bits. The number of fractional bits is given by TYPE\_FBIT, and the number of integral bits is given by TYPE\_IBIT. The fixed-point type is unsigned if TYPE\_UNSIGNED holds; otherwise, it is signed. The fixed-point type is saturating if TYPE\_SATURATING holds; otherwise, it is not saturating.

### COMPLEX\_TYPE

Used to represent GCC built-in \_\_complex\_\_ data types. The TREE\_TYPE is the type of the real and imaginary parts.

## ENUMERAL\_TYPE

Used to represent an enumeration type. The TYPE\_PRECISION gives (as an int), the number of bits used to represent the type. If there are no negative enumeration constants, TYPE\_UNSIGNED will hold. The minimum and maximum enumeration constants may be obtained with TYPE\_MIN\_VALUE and TYPE\_MAX\_VALUE, respectively; each of these macros returns an INTEGER\_CST.

The actual enumeration constants themselves may be obtained by looking at the TYPE\_VALUES. This macro will return a TREE\_LIST, containing the constants. The TREE\_PURPOSE of each node will be an IDENTIFIER\_NODE giving the name of the constant; the TREE\_VALUE will be an INTEGER\_CST giving the value assigned to that constant. These constants will appear in the order in which they were declared. The TREE\_TYPE of each of these constants will be the type of enumeration type itself.

### BOOLEAN\_TYPE

Used to represent the bool type.

## POINTER\_TYPE

Used to represent pointer types, and pointer to data member types. The TREE\_TYPE gives the type to which this type points.

## REFERENCE\_TYPE

Used to represent reference types. The TREE\_TYPE gives the type to which this type refers.

### FUNCTION\_TYPE

Used to represent the type of non-member functions and of static member functions. The TREE\_TYPE gives the return type of the function. The TYPE\_ARG\_TYPES are a TREE\_LIST of the argument types. The TREE\_VALUE of each node in this list is the type of the corresponding argument; the TREE\_PURPOSE is an expression for the default argument value, if any. If the last node in the list is void\_list\_node (a TREE\_LIST node whose TREE\_VALUE is the void\_type\_node), then functions of this type do not take variable arguments. Otherwise, they do take a variable number of arguments.

Note that in C (but not in C++) a function declared like void f() is an unprototyped function taking a variable number of arguments; the TYPE\_ARG\_TYPES of such a function will be NULL.

### METHOD\_TYPE

Used to represent the type of a non-static member function. Like a FUNCTION\_TYPE, the return type is given by the TREE\_TYPE. The type of \*this, i.e., the class of which functions of this type are a member, is given by the TYPE\_METHOD\_BASETYPE. The TYPE\_ARG\_TYPES is the parameter list, as for a FUNCTION\_TYPE, and includes the this argument.

### ARRAY\_TYPE

Used to represent array types. The TREE\_TYPE gives the type of the elements in the array. If the array-bound is present in the type, the TYPE\_DOMAIN is an INTEGER\_TYPE whose TYPE\_MIN\_VALUE and TYPE\_MAX\_VALUE will be the lower and upper bounds of the array, respectively. The TYPE\_MIN\_VALUE will always be an INTEGER\_CST for zero, while the TYPE\_MAX\_VALUE will be one less than the number of elements in the array, i.e., the highest value which may be used to index an element in the array.

### RECORD\_TYPE

Used to represent struct and class types, as well as pointers to member functions and similar constructs in other languages. TYPE\_FIELDS contains the items contained in this type, each of which can be a FIELD\_DECL, VAR\_DECL, CONST\_DECL, or TYPE\_DECL. You may not make any assumptions about the ordering of the fields in the type or whether one or more of them overlap.

### UNION\_TYPE

Used to represent union types. Similar to RECORD\_TYPE except that all FIELD\_ DECL nodes in TYPE\_FIELD start at bit position zero.

## QUAL\_UNION\_TYPE

Used to represent part of a variant record in Ada. Similar to UNION\_TYPE except that each FIELD\_DECL has a DECL\_QUALIFIER field, which contains a boolean expression that indicates whether the field is present in the object. The type will only have one field, so each field's DECL\_QUALIFIER is only evaluated if none of the expressions in the previous fields in TYPE\_FIELDS are nonzero. Normally these expressions will reference a field in the outer object using a PLACEHOLDER\_EXPR.

### LANG\_TYPE

This node is used to represent a language-specific type. The front end must handle it.

### OFFSET\_TYPE

This node is used to represent a pointer-to-data member. For a data member X::m the TYPE\_OFFSET\_BASETYPE is X and the TREE\_TYPE is the type of m.

There are variables whose values represent some of the basic types. These include:

### void\_type\_node

A node for void.

### integer\_type\_node

A node for int.

unsigned\_type\_node.

A node for unsigned int.

char\_type\_node.

A node for char.

It may sometimes be useful to compare one of these variables with a type in hand, using same\_type\_p.

## 11.4 Declarations

This section covers the various kinds of declarations that appear in the internal representation, except for declarations of functions (represented by FUNCTION\_DECL nodes), which are described in Section 11.8 [Functions], page 193.

## 11.4.1 Working with declarations

Some macros can be used with any kind of declaration. These include:

DECL\_NAME

This macro returns an IDENTIFIER\_NODE giving the name of the entity.

TREE\_TYPE

This macro returns the type of the entity declared.

### EXPR\_FILENAME

This macro returns the name of the file in which the entity was declared, as a char\*. For an entity declared implicitly by the compiler (like \_\_builtin\_memcpy), this will be the string "<internal>".

### EXPR\_LINENO

This macro returns the line number at which the entity was declared, as an int

### DECL\_ARTIFICIAL

This predicate holds if the declaration was implicitly generated by the compiler. For example, this predicate will hold of an implicitly declared member function, or of the TYPE\_DECL implicitly generated for a class type. Recall that in C++ code like:

```
struct S {};
```

is roughly equivalent to C code like:

```
struct S {};
typedef struct S S;
```

The implicitly generated typedef declaration is represented by a TYPE\_DECL for which DECL\_ARTIFICIAL holds.

The various kinds of declarations include:

### LABEL\_DECL

These nodes are used to represent labels in function bodies. For more information, see Section 11.8 [Functions], page 193. These nodes only appear in block scopes.

### CONST\_DECL

These nodes are used to represent enumeration constants. The value of the constant is given by DECL\_INITIAL which will be an INTEGER\_CST with the same type as the TREE\_TYPE of the CONST\_DECL, i.e., an ENUMERAL\_TYPE.

### RESULT\_DECL

These nodes represent the value returned by a function. When a value is assigned to a RESULT\_DECL, that indicates that the value should be returned, via bitwise copy, by the function. You can use DECL\_SIZE and DECL\_ALIGN on a RESULT\_DECL, just as with a VAR\_DECL.

### TYPE\_DECL

These nodes represent typedef declarations. The TREE\_TYPE is the type declared to have the name given by DECL\_NAME. In some cases, there is no associated name.

VAR\_DECL

These nodes represent variables with namespace or block scope, as well as static data members. The DECL\_SIZE and DECL\_ALIGN are analogous to TYPE\_SIZE and TYPE\_ALIGN. For a declaration, you should always use the DECL\_SIZE and DECL\_ALIGN rather than the TYPE\_SIZE and TYPE\_ALIGN given by the TREE\_TYPE, since special attributes may have been applied to the variable to give it a particular size and alignment. You may use the predicates DECL\_THIS\_STATIC or DECL\_THIS\_EXTERN to test whether the storage class specifiers static or extern were used to declare a variable.

If this variable is initialized (but does not require a constructor), the DECL\_INITIAL will be an expression for the initializer. The initializer should be evaluated, and a bitwise copy into the variable performed. If the DECL\_INITIAL is the error\_mark\_node, there is an initializer, but it is given by an explicit statement later in the code; no bitwise copy is required.

GCC provides an extension that allows either automatic variables, or global variables, to be placed in particular registers. This extension is being used for a particular VAR\_DECL if DECL\_REGISTER holds for the VAR\_DECL, and if DECL\_ASSEMBLER\_NAME is not equal to DECL\_NAME. In that case, DECL\_ASSEMBLER\_NAME is the name of the register into which the variable will be placed.

## PARM\_DECL

Used to represent a parameter to a function. Treat these nodes similarly to VAR\_DECL nodes. These nodes only appear in the DECL\_ARGUMENTS for a FUNCTION\_DECL.

The DECL\_ARG\_TYPE for a PARM\_DECL is the type that will actually be used when a value is passed to this function. It may be a wider type than the TREE\_TYPE of the parameter; for example, the ordinary type might be short while the DECL\_ARG\_TYPE is int.

### DEBUG\_EXPR\_DECL

Used to represent an anonymous debug-information temporary created to hold an expression as it is optimized away, so that its value can be referenced in debug bind statements.

### FIELD\_DECL

These nodes represent non-static data members. The DECL\_SIZE and DECL\_ALIGN behave as for VAR\_DECL nodes. The position of the field within the parent record is specified by a combination of three attributes. DECL\_FIELD\_OFFSET is the position, counting in bytes, of the DECL\_OFFSET\_ALIGN-bit sized word containing the bit of the field closest to the beginning of the structure. DECL\_FIELD\_BIT\_OFFSET is the bit offset of the first bit of the field within this word; this may be nonzero even for fields that are not bit-fields, since DECL\_OFFSET\_ALIGN may be greater than the natural alignment of the field's type.

If DECL\_C\_BIT\_FIELD holds, this field is a bit-field. In a bit-field, DECL\_BIT\_FIELD\_TYPE also contains the type that was originally specified for it, while DECL\_TYPE may be a modified type with lesser precision, according to the size of the bit field.

#### NAMESPACE DECL

Namespaces provide a name hierarchy for other declarations. They appear in the DECL\_CONTEXT of other \_DECL nodes.

## 11.4.2 Internal structure

DECL nodes are represented internally as a hierarchy of structures.

## 11.4.2.1 Current structure hierarchy

### struct tree\_decl\_minimal

This is the minimal structure to inherit from in order for common DECL macros to work. The fields it contains are a unique ID, source location, context, and name.

### struct tree\_decl\_common

This structure inherits from struct tree\_decl\_minimal. It contains fields that most DECL nodes need, such as a field to store alignment, machine mode, size, and attributes.

### struct tree\_field\_decl

This structure inherits from struct tree\_decl\_common. It is used to represent FIELD\_DECL.

### struct tree\_label\_decl

This structure inherits from struct tree\_decl\_common. It is used to represent LABEL\_DECL.

### struct tree\_translation\_unit\_decl

This structure inherits from struct tree\_decl\_common. It is used to represent TRANSLATION\_UNIT\_DECL.

## struct tree\_decl\_with\_rtl

This structure inherits from struct tree\_decl\_common. It contains a field to store the low-level RTL associated with a DECL node.

### struct tree\_result\_decl

This structure inherits from struct tree\_decl\_with\_rtl. It is used to represent RESULT\_DECL.

### struct tree\_const\_decl

This structure inherits from struct tree\_decl\_with\_rtl. It is used to represent CONST\_DECL.

## struct tree\_parm\_decl

This structure inherits from struct tree\_decl\_with\_rtl. It is used to represent PARM\_DECL.

### struct tree\_decl\_with\_vis

This structure inherits from struct tree\_decl\_with\_rtl. It contains fields necessary to store visibility information, as well as a section name and assembler name.

### struct tree\_var\_decl

This structure inherits from struct tree\_decl\_with\_vis. It is used to represent VAR\_DECL.

### struct tree\_function\_decl

This structure inherits from struct tree\_decl\_with\_vis. It is used to represent FUNCTION\_DECL.

## 11.4.2.2 Adding new DECL node types

Adding a new DECL tree consists of the following steps

Add a new tree code for the DECL node

For language specific DECL nodes, there is a '.def' file in each frontend directory where the tree code should be added. For DECL nodes that are part of the middle-end, the code should be added to 'tree.def'.

Create a new structure type for the DECL node

These structures should inherit from one of the existing structures in the language hierarchy by using that structure as the first member.

```
struct tree_foo_decl
{
    struct tree_decl_with_vis common;
}
```

Would create a structure name tree\_foo\_decl that inherits from struct tree\_decl\_with\_vis.

For language specific DECL nodes, this new structure type should go in the appropriate '.h' file. For DECL nodes that are part of the middle-end, the structure type should go in 'tree.h'.

Add a member to the tree structure enumerator for the node

For garbage collection and dynamic checking purposes, each DECL node structure type is required to have a unique enumerator value specified with it. For language specific DECL nodes, this new enumerator value should go in the appropriate '.def' file. For DECL nodes that are part of the middle-end, the enumerator values are specified in 'treestruct.def'.

## Update union tree\_node

In order to make your new structure type usable, it must be added to union tree\_node. For language specific DECL nodes, a new entry should be added to the appropriate '.h' file of the form

```
struct tree_foo_decl GTY ((tag ("TS_VAR_DECL"))) foo_decl;
```

For DECL nodes that are part of the middle-end, the additional member goes directly into union tree\_node in 'tree.h'.

## Update dynamic checking info

In order to be able to check whether accessing a named portion of union tree\_node is legal, and whether a certain DECL node contains one of the enumerated DECL node structures in the hierarchy, a simple lookup table is used. This lookup table needs to be kept up to date with the tree structure hierarchy, or else checking and containment macros will fail inappropriately.

For language specific DECL nodes, their is an <code>init\_ts</code> function in an appropriate '.c' file, which initializes the lookup table. Code setting up the table for new DECL nodes should be added there. For each DECL tree code and enumerator value representing a member of the inheritance hierarchy, the table should contain 1 if that tree code inherits (directly or indirectly) from that member. Thus, a <code>FOO\_DECL</code> node derived from <code>struct decl\_with\_rtl</code>, and enumerator value <code>TS\_FOO\_DECL</code>, would be set up as follows

```
tree_contains_struct[F00_DECL] [TS_F00_DECL] = 1;
tree_contains_struct[F00_DECL] [TS_DECL_WRTL] = 1;
tree_contains_struct[F00_DECL] [TS_DECL_COMMON] = 1;
tree_contains_struct[F00_DECL] [TS_DECL_MINIMAL] = 1;
```

For DECL nodes that are part of the middle-end, the setup code goes into 'tree.c'.

Add macros to access any new fields and flags

Each added field or flag should have a macro that is used to access it, that performs appropriate checking to ensure only the right type of DECL nodes access the field.

These macros generally take the following form

```
#define FOO_DECL_FIELDNAME(NODE) FOO_DECL_CHECK(NODE)->foo_decl.fieldname
```

However, if the structure is simply a base class for further structures, something like the following should be used

```
#define BASE_STRUCT_CHECK(T) CONTAINS_STRUCT_CHECK(T, TS_BASE_STRUCT)
#define BASE_STRUCT_FIELDNAME(NODE) \
    (BASE_STRUCT_CHECK(NODE)->base_struct.fieldname)
```

Reading them from the generated 'all-tree.def' file (which in turn includes all the 'tree.def' files), 'gencheck.c' is used during GCC's build to generate the \*\_CHECK macros for all tree codes.

## 11.5 Attributes in trees

Attributes, as specified using the \_\_attribute\_\_ keyword, are represented internally as a TREE\_LIST. The TREE\_PURPOSE is the name of the attribute, as an IDENTIFIER\_NODE. The TREE\_VALUE is a TREE\_LIST of the arguments of the attribute, if any, or NULL\_TREE if there

are no arguments; the arguments are stored as the TREE\_VALUE of successive entries in the list, and may be identifiers or expressions. The TREE\_CHAIN of the attribute is the next attribute in a list of attributes applying to the same declaration or type, or NULL\_TREE if there are no further attributes in the list.

Attributes may be attached to declarations and to types; these attributes may be accessed with the following macros. All attributes are stored in this way, and many also cause other changes to the declaration or type or to other internal compiler data structures.

## tree DECL\_ATTRIBUTES (tree dec1)

[Tree Macro]

This macro returns the attributes on the declaration decl.

### tree TYPE\_ATTRIBUTES (tree type)

[Tree Macro]

This macro returns the attributes on the type type.

# 11.6 Expressions

The internal representation for expressions is for the most part quite straightforward. However, there are a few facts that one must bear in mind. In particular, the expression "tree" is actually a directed acyclic graph. (For example there may be many references to the integer constant zero throughout the source program; many of these will be represented by the same expression node.) You should not rely on certain kinds of node being shared, nor should you rely on certain kinds of nodes being unshared.

The following macros can be used with all expression nodes:

## TREE\_TYPE

Returns the type of the expression. This value may not be precisely the same type that would be given the expression in the original program.

In what follows, some nodes that one might expect to always have type bool are documented to have either integral or boolean type. At some point in the future, the C front end may also make use of this same intermediate representation, and at this point these nodes will certainly have integral type. The previous sentence is not meant to imply that the C++ front end does not or will not give these nodes integral type.

Below, we list the various kinds of expression nodes. Except where noted otherwise, the operands to an expression are accessed using the TREE\_OPERAND macro. For example, to access the first operand to a binary plus expression expr, use:

```
TREE_OPERAND (expr, 0)
```

As this example indicates, the operands are zero-indexed.

## 11.6.1 Constant expressions

The table below begins with constants, moves on to unary expressions, then proceeds to binary expressions, and concludes with various other kinds of expressions:

### INTEGER\_CST

These nodes represent integer constants. Note that the type of these constants is obtained with TREE\_TYPE; they are not always of type int. In particular, char constants are represented with INTEGER\_CST nodes. The value of the integer constant e is represented in an array of HOST\_WIDE\_INT. There are enough

elements in the array to represent the value without taking extra elements for redundant 0s or -1. The number of elements used to represent e is available via TREE\_INT\_CST\_NUNITS. Element i can be extracted by using TREE\_INT\_CST\_ELT (e, i). TREE\_INT\_CST\_LOW is a shorthand for TREE\_INT\_CST\_ELT (e, 0). The functions tree\_fits\_shwi\_p and tree\_fits\_uhwi\_p can be used to tell if the value is small enough to fit in a signed HOST\_WIDE\_INT or an unsigned

the value is small enough to fit in a signed HOST\_WIDE\_INT or an unsigned HOST\_WIDE\_INT respectively. The value can then be extracted using tree\_to\_shwi and tree\_to\_uhwi.

### REAL\_CST

FIXME: Talk about how to obtain representations of this constant, do comparisons, and so forth.

### FIXED\_CST

These nodes represent fixed-point constants. The type of these constants is obtained with TREE\_TYPE. TREE\_FIXED\_CST\_PTR points to a struct fixed\_value; TREE\_FIXED\_CST returns the structure itself. struct fixed\_value contains data with the size of two HOST\_BITS\_PER\_WIDE\_INT and mode as the associated fixed-point machine mode for data.

### COMPLEX\_CST

These nodes are used to represent complex number constants, that is a \_\_ complex\_\_ whose parts are constant nodes. The TREE\_REALPART and TREE\_IMAGPART return the real and the imaginary parts respectively.

### VECTOR\_CST

These nodes are used to represent vector constants. Each vector constant v is treated as a specific instance of an arbitrary-length sequence that itself contains 'VECTOR\_CST\_NPATTERNS (v)' interleaved patterns. Each pattern has the form:

```
{ base0, base1, base1 + step, base1 + step * 2, ... }
```

The first three elements in each pattern are enough to determine the values of the other elements. However, if all steps are zero, only the first two elements are needed. If in addition each base1 is equal to the corresponding base0, only the first element in each pattern is needed. The number of encoded elements per pattern is given by 'VECTOR\_CST\_NELTS\_PER\_PATTERN (v)'.

For example, the constant:

```
{ 0, 1, 2, 6, 3, 8, 4, 10, 5, 12, 6, 14, 7, 16, 8, 18 }
```

is interpreted as an interleaving of the sequences:

```
{ 0, 2, 3, 4, 5, 6, 7, 8 }
{ 1, 6, 8, 10, 12, 14, 16, 18 }
```

where the sequences are represented by the following patterns:

```
base0 == 0, base1 == 2, step == 1
base0 == 1, base1 == 6, step == 2
```

In this case:

```
VECTOR_CST_NPATTERNS (v) == 2
VECTOR_CST_NELTS_PER_PATTERN (v) == 3
```

The vector is therefore encoded using the first 6 elements ('{ 0, 1, 2, 6, 3, 8 }'), with the remaining 10 elements being implicit extensions of them.

Sometimes this scheme can create two possible encodings of the same vector. For example { 0, 1 } could be seen as two patterns with one element each or one pattern with two elements (base0 and base1). The canonical encoding is always the one with the fewest patterns or (if both encodings have the same number of petterns) the one with the fewest encoded elements.

'vector\_cst\_encoding\_nelts (v)' gives the total number of encoded elements in v, which is 6 in the example above. VECTOR\_CST\_ENCODED\_ELTS (v) gives a pointer to the elements encoded in v and VECTOR\_CST\_ENCODED\_ELT (v, i) accesses the value of encoded element i.

'VECTOR\_CST\_DUPLICATE\_P (v)' is true if v simply contains repeated instances of 'VECTOR\_CST\_NPATTERNS (v)' values. This is a shorthand for testing 'VECTOR\_CST\_NELTS\_PER\_PATTERN (v) == 1'.

'VECTOR\_CST\_STEPPED\_P (v)' is true if at least one pattern in v has a nonzero step. This is a shorthand for testing 'VECTOR\_CST\_NELTS\_PER\_PATTERN (v) == 3'.

The utility function vector\_cst\_elt gives the value of an arbitrary index as a tree. vector\_cst\_int\_elt gives the same value as a wide\_int.

### STRING\_CST

These nodes represent string-constants. The TREE\_STRING\_LENGTH returns the length of the string, as an int. The TREE\_STRING\_POINTER is a char\* containing the string itself. The string may not be NUL-terminated, and it may contain embedded NUL characters. Therefore, the TREE\_STRING\_LENGTH includes the trailing NUL if it is present.

For wide string constants, the TREE\_STRING\_LENGTH is the number of bytes in the string, and the TREE\_STRING\_POINTER points to an array of the bytes of the string, as represented on the target system (that is, as integers in the target endianness). Wide and non-wide string constants are distinguished only by the TREE\_TYPE of the STRING\_CST.

FIXME: The formats of string constants are not well-defined when the target system bytes are not the same width as host system bytes.

### POLY\_INT\_CST

These nodes represent invariants that depend on some target-specific runtime parameters. They consist of NUM\_POLY\_INT\_COEFFS coefficients, with the first coefficient being the constant term and the others being multipliers that are applied to the runtime parameters.

POLY\_INT\_CST\_ELT (x, i) references coefficient number i of POLY\_INT\_CST node x. Each coefficient is an INTEGER\_CST.

## 11.6.2 References to storage

### ARRAY\_REF

These nodes represent array accesses. The first operand is the array; the second is the index. To calculate the address of the memory accessed, you must scale the index by the size of the type of the array elements. The type of these expressions must be the type of a component of the array. The third and

fourth operands are used after gimplification to represent the lower bound and component size but should not be used directly; call array\_ref\_low\_bound and array\_ref\_element\_size instead.

### ARRAY\_RANGE\_REF

These nodes represent access to a range (or "slice") of an array. The operands are the same as that for ARRAY\_REF and have the same meanings. The type of these expressions must be an array whose component type is the same as that of the first operand. The range of that array type determines the amount of data these expressions access.

### TARGET MEM REF

These nodes represent memory accesses whose address directly map to an addressing mode of the target architecture. The first argument is TMR\_SYMBOL and must be a VAR\_DECL of an object with a fixed address. The second argument is TMR\_BASE and the third one is TMR\_INDEX. The fourth argument is TMR\_STEP and must be an INTEGER\_CST. The fifth argument is TMR\_OFFSET and must be an INTEGER\_CST. Any of the arguments may be NULL if the appropriate component does not appear in the address. Address of the TARGET\_MEM\_REF is determined in the following way.

```
&TMR_SYMBOL + TMR_BASE + TMR_INDEX * TMR_STEP + TMR_OFFSET
```

The sixth argument is the reference to the original memory access, which is preserved for the purposes of the RTL alias analysis. The seventh argument is a tag representing the results of tree level alias analysis.

### ADDR\_EXPR

These nodes are used to represent the address of an object. (These expressions will always have pointer or reference type.) The operand may be another expression, or it may be a declaration.

As an extension, GCC allows users to take the address of a label. In this case, the operand of the ADDR\_EXPR will be a LABEL\_DECL. The type of such an expression is void\*.

If the object addressed is not an lvalue, a temporary is created, and the address of the temporary is used.

### INDIRECT\_REF

These nodes are used to represent the object pointed to by a pointer. The operand is the pointer being dereferenced; it will always have pointer or reference type.

MEM\_REF These nodes are used to represent the object pointed to by a pointer offset by a constant. The first operand is the pointer being dereferenced; it will always have pointer or reference type. The second operand is a pointer constant. Its type is specifying the type to be used for type-based alias analysis.

### COMPONENT\_REF

These nodes represent non-static data member accesses. The first operand is the object (rather than a pointer to it); the second operand is the FIELD\_DECL for the data member. The third operand represents the byte offset of the field, but should not be used directly; call component\_ref\_field\_offset instead.

## 11.6.3 Unary and Binary Expressions

### NEGATE\_EXPR

These nodes represent unary negation of the single operand, for both integer and floating-point types. The type of negation can be determined by looking at the type of the expression.

The behavior of this operation on signed arithmetic overflow is controlled by the flag\_wrapv and flag\_trapv variables.

ABS\_EXPR These nodes represent the absolute value of the single operand, for both integer and floating-point types. This is typically used to implement the abs, labs and llabs builtins for integer types, and the fabs, fabsf and fabsl builtins for floating point types. The type of abs operation can be determined by looking at the type of the expression.

This node is not used for complex types. To represent the modulus or complex abs of a complex value, use the BUILT\_IN\_CABS, BUILT\_IN\_CABSF or BUILT\_IN\_CABSL builtins, as used to implement the C99 cabs, cabsf and cabsl built-in functions.

### ABSU\_EXPR

These nodes represent the absolute value of the single operand in equivalent unsigned type such that ABSU\_EXPR of TYPE\_MIN is well defined.

### BIT\_NOT\_EXPR

These nodes represent bitwise complement, and will always have integral type. The only operand is the value to be complemented.

### TRUTH\_NOT\_EXPR

These nodes represent logical negation, and will always have integral (or boolean) type. The operand is the value being negated. The type of the operand and that of the result are always of BOOLEAN\_TYPE or INTEGER\_TYPE.

PREDECREMENT\_EXPR PREINCREMENT\_EXPR POSTDECREMENT\_EXPR POSTINCREMENT\_EXPR

These nodes represent increment and decrement expressions. The value of the single operand is computed, and the operand incremented or decremented. In the case of PREDECREMENT\_EXPR and PREINCREMENT\_EXPR, the value of the expression is the value resulting after the increment or decrement; in the case of POSTDECREMENT\_EXPR and POSTINCREMENT\_EXPR is the value before the increment or decrement occurs. The type of the operand, like that of the result, will be either integral, boolean, or floating-point.

## FIX\_TRUNC\_EXPR

These nodes represent conversion of a floating-point value to an integer. The single operand will have a floating-point type, while the complete expression will have an integral (or boolean) type. The operand is rounded towards zero.

### FLOAT\_EXPR

These nodes represent conversion of an integral (or boolean) value to a floating-point value. The single operand will have integral type, while the complete expression will have a floating-point type.

FIXME: How is the operand supposed to be rounded? Is this dependent on '-mieee'?

### COMPLEX\_EXPR

These nodes are used to represent complex numbers constructed from two expressions of the same (integer or real) type. The first operand is the real part and the second operand is the imaginary part.

### CONJ\_EXPR

These nodes represent the conjugate of their operand.

### REALPART\_EXPR

## IMAGPART\_EXPR

These nodes represent respectively the real and the imaginary parts of complex numbers (their sole argument).

### NON\_LVALUE\_EXPR

These nodes indicate that their one and only operand is not an Ivalue. A back end can treat these identically to the single operand.

NOP\_EXPR These nodes are used to represent conversions that do not require any codegeneration. For example, conversion of a char\* to an int\* does not require any code be generated; such a conversion is represented by a NOP\_EXPR. The single operand is the expression to be converted. The conversion from a pointer to a reference is also represented with a NOP\_EXPR.

### CONVERT\_EXPR

These nodes are similar to NOP\_EXPRs, but are used in those situations where code may need to be generated. For example, if an <code>int\*</code> is converted to an <code>int</code> code may need to be generated on some platforms. These nodes are never used for C++-specific conversions, like conversions between pointers to different classes in an inheritance hierarchy. Any adjustments that need to be made in such cases are always indicated explicitly. Similarly, a user-defined conversion is never represented by a <code>CONVERT\_EXPR</code>; instead, the function calls are made explicit.

### FIXED\_CONVERT\_EXPR

These nodes are used to represent conversions that involve fixed-point values. For example, from a fixed-point value to another fixed-point value, from an integer to a fixed-point value, from a fixed-point value to an integer, from a floating-point value to a fixed-point value, or from a fixed-point value to a floating-point value.

### LSHIFT\_EXPR

### RSHIFT\_EXPR

These nodes represent left and right shifts, respectively. The first operand is the value to shift; it will always be of integral type. The second operand is an expression for the number of bits by which to shift. Right shift should be treated as arithmetic, i.e., the high-order bits should be zero-filled when the expression has unsigned type and filled with the sign bit when the expression has signed type. Note that the result is undefined if the second operand is larger than or equal to the first operand's type size. Unlike most nodes, these can have a vector as first operand and a scalar as second operand.

BIT\_IOR\_EXPR BIT\_XOR\_EXPR BIT\_AND\_EXPR

These nodes represent bitwise inclusive or, bitwise exclusive or, and bitwise and, respectively. Both operands will always have integral type.

TRUTH\_ANDIF\_EXPR
TRUTH\_ORIF\_EXPR

These nodes represent logical "and" and logical "or", respectively. These operators are not strict; i.e., the second operand is evaluated only if the value of the expression is not determined by evaluation of the first operand. The type of the operands and that of the result are always of BOOLEAN\_TYPE or INTEGER\_TYPE.

TRUTH\_AND\_EXPR TRUTH\_OR\_EXPR TRUTH\_XOR\_EXPR

These nodes represent logical and, logical or, and logical exclusive or. They are strict; both arguments are always evaluated. There are no corresponding operators in C or C++, but the front end will sometimes generate these expressions anyhow, if it can tell that strictness does not matter. The type of the operands and that of the result are always of BOOLEAN\_TYPE or INTEGER\_TYPE.

## POINTER\_PLUS\_EXPR

This node represents pointer arithmetic. The first operand is always a pointer/reference type. The second operand is always an unsigned integer type compatible with sizetype. This and POINTER\_DIFF\_EXPR are the only binary arithmetic operators that can operate on pointer types.

### POINTER\_DIFF\_EXPR

This node represents pointer subtraction. The two operands always have pointer/reference type. It returns a signed integer of the same precision as the pointers. The behavior is undefined if the difference of the two pointers, seen as infinite precision non-negative integers, does not fit in the result type. The result does not depend on the pointer type, it is not divided by the size of the pointed-to type.

PLUS\_EXPR MINUS\_EXPR MULT\_EXPR

These nodes represent various binary arithmetic operations. Respectively, these operations are addition, subtraction (of the second operand from the first) and multiplication. Their operands may have either integral or floating type, but there will never be case in which one operand is of floating type and the other is of integral type.

The behavior of these operations on signed arithmetic overflow is controlled by the flag\_wrapv and flag\_trapv variables.

### MULT\_HIGHPART\_EXPR

This node represents the "high-part" of a widening multiplication. For an integral type with b bits of precision, the result is the most significant b bits of the full 2b product.

### RDIV\_EXPR

This node represents a floating point division operation.

TRUNC\_DIV\_EXPR FLOOR\_DIV\_EXPR CEIL\_DIV\_EXPR ROUND\_DIV\_EXPR

These nodes represent integer division operations that return an integer result. TRUNC\_DIV\_EXPR rounds towards zero, FLOOR\_DIV\_EXPR rounds towards negative infinity, CEIL\_DIV\_EXPR rounds towards positive infinity and ROUND\_DIV\_EXPR rounds to the closest integer. Integer division in C and C++ is truncating, i.e. TRUNC\_DIV\_EXPR.

The behavior of these operations on signed arithmetic overflow, when dividing the minimum signed integer by minus one, is controlled by the flag\_wrapv and flag\_trapv variables.

TRUNC\_MOD\_EXPR FLOOR\_MOD\_EXPR CEIL\_MOD\_EXPR ROUND\_MOD\_EXPR

These nodes represent the integer remainder or modulus operation. The integer modulus of two operands a and b is defined as a - (a/b)\*b where the division calculated using the corresponding division operator. Hence for TRUNC\_MOD\_EXPR this definition assumes division using truncation towards zero, i.e. TRUNC\_DIV\_EXPR. Integer remainder in C and C++ uses truncating division, i.e. TRUNC\_MOD\_EXPR.

## EXACT\_DIV\_EXPR

The EXACT\_DIV\_EXPR code is used to represent integer divisions where the numerator is known to be an exact multiple of the denominator. This allows the backend to choose between the faster of TRUNC\_DIV\_EXPR, CEIL\_DIV\_EXPR and FLOOR\_DIV\_EXPR for the current target.

LT\_EXPR LE\_EXPR GT\_EXPR GE\_EXPR LTGT\_EXPR

EQ\_EXPR NE\_EXPR

These nodes represent the less than, less than or equal to, greater than, greater than or equal to, less or greater than, equal, and not equal comparison operators. The first and second operands will either be both of integral type, both of

floating type or both of vector type, except for LTGT\_EXPR where they will only be both of floating type. The result type of these expressions will always be of integral, boolean or signed integral vector type. These operations return the result type's zero value for false, the result type's one value for true, and a vector whose elements are zero (false) or minus one (true) for vectors.

For floating point comparisons, if we honor IEEE NaNs and either operand is NaN, then NE\_EXPR always returns true and the remaining operators always return false. On some targets, comparisons against an IEEE NaN, other than equality and inequality, may generate a floating-point exception.

## ORDERED\_EXPR UNORDERED\_EXPR

These nodes represent non-trapping ordered and unordered comparison operators. These operations take two floating point operands and determine whether they are ordered or unordered relative to each other. If either operand is an IEEE NaN, their comparison is defined to be unordered, otherwise the comparison is defined to be ordered. The result type of these expressions will always be of integral or boolean type. These operations return the result type's zero value for false, and the result type's one value for true.

UNLT\_EXPR UNLE\_EXPR UNGT\_EXPR UNGE\_EXPR UNEQ\_EXPR

These nodes represent the unordered comparison operators. These operations take two floating point operands and determine whether the operands are unordered or are less than, less than or equal to, greater than, greater than or equal to, or equal respectively. For example, UNLT\_EXPR returns true if either operand is an IEEE NaN or the first operand is less than the second. All these operations are guaranteed not to generate a floating point exception. The result type of these expressions will always be of integral or boolean type. These operations return the result type's zero value for false, and the result type's one value for true.

## MODIFY\_EXPR

These nodes represent assignment. The left-hand side is the first operand; the right-hand side is the second operand. The left-hand side will be a VAR\_DECL, INDIRECT\_REF, COMPONENT\_REF, or other lvalue.

These nodes are used to represent not only assignment with '=' but also compound assignments (like '+='), by reduction to '=' assignment. In other words, the representation for 'i += 3' looks just like that for 'i = i + 3'.

### INIT\_EXPR

These nodes are just like MODIFY\_EXPR, but are used only when a variable is initialized, rather than assigned to subsequently. This means that we can assume that the target of the initialization is not used in computing its own value; any reference to the lhs in computing the rhs is undefined.

### COMPOUND\_EXPR

These nodes represent comma-expressions. The first operand is an expression whose value is computed and thrown away prior to the evaluation of the second operand. The value of the entire expression is the value of the second operand.

### COND\_EXPR

These nodes represent ?: expressions. The first operand is of boolean or integral type. If it evaluates to a nonzero value, the second operand should be evaluated, and returned as the value of the expression. Otherwise, the third operand is evaluated, and returned as the value of the expression.

The second operand must have the same type as the entire expression, unless it unconditionally throws an exception or calls a noreturn function, in which case it should have void type. The same constraints apply to the third operand. This allows array bounds checks to be represented conveniently as  $(i \ge 0 \&\& i \le 10)$ ? i : abort().

As a GNU extension, the C language front-ends allow the second operand of the ?: operator may be omitted in the source. For example, x ? : 3 is equivalent to x ? x : 3, assuming that x is an expression without side effects. In the tree representation, however, the second operand is always present, possibly protected by SAVE\_EXPR if the first argument does cause side effects.

### CALL\_EXPR

These nodes are used to represent calls to functions, including non-static member functions. CALL\_EXPRs are implemented as expression nodes with a variable number of operands. Rather than using TREE\_OPERAND to extract them, it is preferable to use the specialized accessor macros and functions that operate specifically on CALL\_EXPR nodes.

CALL\_EXPR\_FN returns a pointer to the function to call; it is always an expression whose type is a POINTER\_TYPE.

The number of arguments to the call is returned by call\_expr\_nargs, while the arguments themselves can be accessed with the CALL\_EXPR\_ARG macro. The arguments are zero-indexed and numbered left-to-right. You can iterate over the arguments using FOR\_EACH\_CALL\_EXPR\_ARG, as in:

```
tree call, arg;
call_expr_arg_iterator iter;
FOR_EACH_CALL_EXPR_ARG (arg, iter, call)
   /* arg is bound to successive arguments of call. */
.
```

For non-static member functions, there will be an operand corresponding to the this pointer. There will always be expressions corresponding to all of the arguments, even if the function is declared with default arguments and some arguments are not explicitly provided at the call sites.

CALL\_EXPRs also have a CALL\_EXPR\_STATIC\_CHAIN operand that is used to implement nested functions. This operand is otherwise null.

### CLEANUP\_POINT\_EXPR

These nodes represent full-expressions. The single operand is an expression to evaluate. Any destructor calls engendered by the creation of temporaries

during the evaluation of that expression should be performed immediately after the expression is evaluated.

### CONSTRUCTOR

These nodes represent the brace-enclosed initializers for a structure or an array. They contain a sequence of component values made out of a vector of constructor\_elt, which is a (INDEX, VALUE) pair.

If the TREE\_TYPE of the CONSTRUCTOR is a RECORD\_TYPE, UNION\_TYPE or QUAL\_UNION\_TYPE then the INDEX of each node in the sequence will be a FIELD\_DECL and the VALUE will be the expression used to initialize that field.

If the TREE\_TYPE of the CONSTRUCTOR is an ARRAY\_TYPE, then the INDEX of each node in the sequence will be an INTEGER\_CST or a RANGE\_EXPR of two INTEGER\_CSTs. A single INTEGER\_CST indicates which element of the array is being assigned to. A RANGE\_EXPR indicates an inclusive range of elements to initialize. In both cases the VALUE is the corresponding initializer. It is reevaluated for each element of a RANGE\_EXPR. If the INDEX is NULL\_TREE, then the initializer is for the next available array element.

In the front end, you should not depend on the fields appearing in any particular order. However, in the middle end, fields must appear in declaration order. You should not assume that all fields will be represented. Unrepresented fields will be cleared (zeroed), unless the CONSTRUCTOR\_NO\_CLEARING flag is set, in which case their value becomes undefined.

### COMPOUND\_LITERAL\_EXPR

These nodes represent ISO C99 compound literals. The COMPOUND\_LITERAL\_EXPR\_DECL\_EXPR is a DECL\_EXPR containing an anonymous VAR\_DECL for the unnamed object represented by the compound literal; the DECL\_INITIAL of that VAR\_DECL is a CONSTRUCTOR representing the brace-enclosed list of initializers in the compound literal. That anonymous VAR\_DECL can also be accessed directly by the COMPOUND\_LITERAL\_EXPR\_DECL macro.

### SAVE EXPR

A SAVE\_EXPR represents an expression (possibly involving side effects) that is used more than once. The side effects should occur only the first time the expression is evaluated. Subsequent uses should just reuse the computed value. The first operand to the SAVE\_EXPR is the expression to evaluate. The side effects should be executed where the SAVE\_EXPR is first encountered in a depth-first preorder traversal of the expression tree.

### TARGET\_EXPR

A TARGET\_EXPR represents a temporary object. The first operand is a VAR\_DECL for the temporary variable. The second operand is the initializer for the temporary. The initializer is evaluated and, if non-void, copied (bitwise) into the temporary. If the initializer is void, that means that it will perform the initialization itself.

Often, a TARGET\_EXPR occurs on the right-hand side of an assignment, or as the second operand to a comma-expression which is itself the right-hand side of an assignment, etc. In this case, we say that the TARGET\_EXPR is "normal"; otherwise, we say it is "orphaned". For a normal TARGET\_EXPR the temporary variable should be treated as an alias for the left-hand side of the assignment, rather than as a new temporary variable.

The third operand to the TARGET\_EXPR, if present, is a cleanup-expression (i.e., destructor call) for the temporary. If this expression is orphaned, then this expression must be executed when the statement containing this expression is complete. These cleanups must always be executed in the order opposite to that in which they were encountered. Note that if a temporary is created on one branch of a conditional operator (i.e., in the second or third operand to a COND\_EXPR), the cleanup must be run only if that branch is actually executed.

### VA\_ARG\_EXPR

This node is used to implement support for the C/C++ variable argument-list mechanism. It represents expressions like va\_arg (ap, type). Its TREE\_TYPE yields the tree representation for type and its sole argument yields the representation for ap.

### ANNOTATE\_EXPR

This node is used to attach markers to an expression. The first operand is the annotated expression, the second is an INTEGER\_CST with a value from enum annot\_expr\_kind, the third is an INTEGER\_CST.

## **11.6.4** Vectors

### VEC\_DUPLICATE\_EXPR

This node has a single operand and represents a vector in which every element is equal to that operand.

### VEC\_SERIES\_EXPR

This node represents a vector formed from a scalar base and step, given as the first and second operands respectively. Element i of the result is equal to 'base + i\*step'.

This node is restricted to integral types, in order to avoid specifying the rounding behavior for floating-point types.

### VEC\_LSHIFT\_EXPR

### VEC\_RSHIFT\_EXPR

These nodes represent whole vector left and right shifts, respectively. The first operand is the vector to shift; it will always be of vector type. The second operand is an expression for the number of bits by which to shift. Note that the result is undefined if the second operand is larger than or equal to the first operand's type size.

## VEC\_WIDEN\_MULT\_HI\_EXPR VEC\_WIDEN\_MULT\_LO\_EXPR

These nodes represent widening vector multiplication of the high and low parts of the two input vectors, respectively. Their operands are vectors that contain the same number of elements (N) of the same integral type. The result is a vector that contains half as many elements, of an integral type whose size is twice as wide. In the case of VEC\_WIDEN\_MULT\_HI\_EXPR the high N/2 elements

of the two vector are multiplied to produce the vector of N/2 products. In the case of VEC\_WIDEN\_MULT\_LO\_EXPR the low N/2 elements of the two vector are multiplied to produce the vector of N/2 products.

## VEC\_UNPACK\_HI\_EXPR VEC\_UNPACK\_LO\_EXPR

These nodes represent unpacking of the high and low parts of the input vector, respectively. The single operand is a vector that contains N elements of the same integral or floating point type. The result is a vector that contains half as many elements, of an integral or floating point type whose size is twice as wide. In the case of VEC\_UNPACK\_HI\_EXPR the high N/2 elements of the vector are extracted and widened (promoted). In the case of VEC\_UNPACK\_LO\_EXPR the low N/2 elements of the vector are extracted and widened (promoted).

## VEC\_UNPACK\_FLOAT\_HI\_EXPR VEC\_UNPACK\_FLOAT\_LO\_EXPR

These nodes represent unpacking of the high and low parts of the input vector, where the values are converted from fixed point to floating point. The single operand is a vector that contains N elements of the same integral type. The result is a vector that contains half as many elements of a floating point type whose size is twice as wide. In the case of VEC\_UNPACK\_FLOAT\_HI\_EXPR the high N/2 elements of the vector are extracted, converted and widened. In the case of VEC\_UNPACK\_FLOAT\_LO\_EXPR the low N/2 elements of the vector are extracted, converted and widened.

## VEC\_UNPACK\_FIX\_TRUNC\_HI\_EXPR VEC\_UNPACK\_FIX\_TRUNC\_LO\_EXPR

These nodes represent unpacking of the high and low parts of the input vector, where the values are truncated from floating point to fixed point. The single operand is a vector that contains N elements of the same floating point type. The result is a vector that contains half as many elements of an integral type whose size is twice as wide. In the case of VEC\_UNPACK\_FIX\_TRUNC\_HI\_EXPR the high N/2 elements of the vector are extracted and converted with truncation. In the case of VEC\_UNPACK\_FIX\_TRUNC\_LO\_EXPR the low N/2 elements of the vector are extracted and converted with truncation.

## VEC\_PACK\_TRUNC\_EXPR

This node represents packing of truncated elements of the two input vectors into the output vector. Input operands are vectors that contain the same number of elements of the same integral or floating point type. The result is a vector that contains twice as many elements of an integral or floating point type whose size is half as wide. The elements of the two vectors are demoted and merged (concatenated) to form the output vector.

### VEC\_PACK\_SAT\_EXPR

This node represents packing of elements of the two input vectors into the output vector using saturation. Input operands are vectors that contain the same number of elements of the same integral type. The result is a vector that contains twice as many elements of an integral type whose size is half as wide.

The elements of the two vectors are demoted and merged (concatenated) to form the output vector.

### VEC\_PACK\_FIX\_TRUNC\_EXPR

This node represents packing of elements of the two input vectors into the output vector, where the values are converted from floating point to fixed point. Input operands are vectors that contain the same number of elements of a floating point type. The result is a vector that contains twice as many elements of an integral type whose size is half as wide. The elements of the two vectors are merged (concatenated) to form the output vector.

### VEC\_PACK\_FLOAT\_EXPR

This node represents packing of elements of the two input vectors into the output vector, where the values are converted from fixed point to floating point. Input operands are vectors that contain the same number of elements of an integral type. The result is a vector that contains twice as many elements of floating point type whose size is half as wide. The elements of the two vectors are merged (concatenated) to form the output vector.

### VEC\_COND\_EXPR

These nodes represent ?: expressions. The three operands must be vectors of the same size and number of elements. The second and third operands must have the same type as the entire expression. The first operand is of signed integral vector type. If an element of the first operand evaluates to a zero value, the corresponding element of the result is taken from the third operand. If it evaluates to a minus one value, it is taken from the second operand. It should never evaluate to any other value currently, but optimizations should not rely on that property. In contrast with a COND\_EXPR, all operands are always evaluated.

SAD\_EXPR This node represents the Sum of Absolute Differences operation. The three operands must be vectors of integral types. The first and second operand must have the same type. The size of the vector element of the third operand must be at lease twice of the size of the vector element of the first and second one. The SAD is calculated between the first and second operands, added to the third operand, and returned.

## 11.7 Statements

Most statements in GIMPLE are assignment statements, represented by GIMPLE\_ASSIGN. No other C expressions can appear at statement level; a reference to a volatile object is converted into a GIMPLE\_ASSIGN.

There are also several varieties of complex statements.

### 11.7.1 Basic Statements

## ASM\_EXPR

Used to represent an inline assembly statement. For an inline assembly statement like:

```
asm ("mov x, y");
```

The ASM\_STRING macro will return a STRING\_CST node for "mov x, y". If the original statement made use of the extended-assembly syntax, then ASM\_OUTPUTS, ASM\_INPUTS, and ASM\_CLOBBERS will be the outputs, inputs, and clobbers for the statement, represented as STRING\_CST nodes. The extended-assembly syntax looks like:

```
asm ("fsinx %1,%0" : "=f" (result) : "f" (angle));
```

The first string is the ASM\_STRING, containing the instruction template. The next two strings are the output and inputs, respectively; this statement has no clobbers. As this example indicates, "plain" assembly statements are merely a special case of extended assembly statements; they have no cv-qualifiers, outputs, inputs, or clobbers. All of the strings will be NUL-terminated, and will contain no embedded NUL-characters.

If the assembly statement is declared volatile, or if the statement was not an extended assembly statement, and is therefore implicitly volatile, then the predicate ASM\_VOLATILE\_P will hold of the ASM\_EXPR.

### DECL\_EXPR

Used to represent a local declaration. The DECL\_EXPR\_DECL macro can be used to obtain the entity declared. This declaration may be a LABEL\_DECL, indicating that the label declared is a local label. (As an extension, GCC allows the declaration of labels with scope.) In C, this declaration may be a FUNCTION\_DECL, indicating the use of the GCC nested function extension. For more information, see Section 11.8 [Functions], page 193.

## LABEL\_EXPR

Used to represent a label. The LABEL\_DECL declared by this statement can be obtained with the LABEL\_EXPR\_LABEL macro. The IDENTIFIER\_NODE giving the name of the label can be obtained from the LABEL\_DECL with DECL\_NAME.

### GOTO\_EXPR

Used to represent a goto statement. The GOTO\_DESTINATION will usually be a LABEL\_DECL. However, if the "computed goto" extension has been used, the GOTO\_DESTINATION will be an arbitrary expression indicating the destination. This expression will always have pointer type.

### RETURN\_EXPR

Used to represent a return statement. Operand 0 represents the value to return. It should either be the RESULT\_DECL for the containing function, or a MODIFY\_EXPR or INIT\_EXPR setting the function's RESULT\_DECL. It will be NULL\_TREE if the statement was just

return;

### LOOP\_EXPR

These nodes represent "infinite" loops. The LOOP\_EXPR\_BODY represents the body of the loop. It should be executed forever, unless an EXIT\_EXPR is encountered.

### EXIT\_EXPR

These nodes represent conditional exits from the nearest enclosing LOOP\_EXPR. The single operand is the condition; if it is nonzero, then the loop should be exited. An EXIT\_EXPR will only appear within a LOOP\_EXPR.

### SWITCH\_STMT

Used to represent a switch statement. The SWITCH\_STMT\_COND is the expression on which the switch is occurring. See the documentation for an IF\_STMT for more information on the representation used for the condition. The SWITCH\_STMT\_BODY is the body of the switch statement. The SWITCH\_STMT\_TYPE is the original type of switch expression as given in the source, before any compiler conversions.

### CASE\_LABEL\_EXPR

Use to represent a case label, range of case labels, or a default label. If CASE\_LOW is NULL\_TREE, then this is a default label. Otherwise, if CASE\_HIGH is NULL\_TREE, then this is an ordinary case label. In this case, CASE\_LOW is an expression giving the value of the label. Both CASE\_LOW and CASE\_HIGH are INTEGER\_CST nodes. These values will have the same type as the condition expression in the switch statement.

Otherwise, if both CASE\_LOW and CASE\_HIGH are defined, the statement is a range of case labels. Such statements originate with the extension that allows users to write things of the form:

```
case 2 ... 5:
```

The first value will be CASE\_LOW, while the second will be CASE\_HIGH.

### DEBUG\_BEGIN\_STMT

Marks the beginning of a source statement, for purposes of debug information generation.

## 11.7.2 Blocks

Block scopes and the variables they declare in GENERIC are expressed using the BIND\_EXPR code, which in previous versions of GCC was primarily used for the C statement-expression extension.

Variables in a block are collected into BIND\_EXPR\_VARS in declaration order through their TREE\_CHAIN field. Any runtime initialization is moved out of DECL\_INITIAL and into a statement in the controlled block. When gimplifying from C or C++, this initialization replaces the DECL\_STMT. These variables will never require cleanups. The scope of these variables is just the body

Variable-length arrays (VLAs) complicate this process, as their size often refers to variables initialized earlier in the block and their initialization involves an explicit stack allocation. To handle this, we add an indirection and replace them with a pointer to stack space allocated by means of alloca. In most cases, we also arrange for this space to be reclaimed when the enclosing BIND\_EXPR is exited, the exception to this being when there is an explicit call to alloca in the source code, in which case the stack is left depressed on exit of the BIND\_EXPR.

A C++ program will usually contain more BIND\_EXPRs than there are syntactic blocks in the source code, since several C++ constructs have implicit scopes associated with them. On the other hand, although the C++ front end uses pseudo-scopes to handle cleanups for objects with destructors, these don't translate into the GIMPLE form; multiple declarations at the same level use the same BIND\_EXPR.

## 11.7.3 Statement Sequences

Multiple statements at the same nesting level are collected into a STATEMENT\_LIST. Statement lists are modified and traversed using the interface in 'tree-iterator.h'.

## 11.7.4 Empty Statements

Whenever possible, statements with no effect are discarded. But if they are nested within another construct which cannot be discarded for some reason, they are instead replaced with an empty statement, generated by build\_empty\_stmt. Initially, all empty statements were shared, after the pattern of the Java front end, but this caused a lot of trouble in practice.

An empty statement is represented as (void)0.

## 11.7.5 Jumps

Other jumps are expressed by either GOTO\_EXPR or RETURN\_EXPR.

The operand of a GOTO\_EXPR must be either a label or a variable containing the address to jump to.

The operand of a RETURN\_EXPR is either NULL\_TREE, RESULT\_DECL, or a MODIFY\_EXPR which sets the return value. It would be nice to move the MODIFY\_EXPR into a separate statement, but the special return semantics in expand\_return make that difficult. It may still happen in the future, perhaps by moving most of that logic into expand\_assignment.

## 11.7.6 Cleanups

Destructors for local C++ objects and similar dynamic cleanups are represented in GIM-PLE by a TRY\_FINALLY\_EXPR. TRY\_FINALLY\_EXPR has two operands, both of which are a sequence of statements to execute. The first sequence is executed. When it completes the second sequence is executed.

The first sequence may complete in the following ways:

- 1. Execute the last statement in the sequence and fall off the end.
- 2. Execute a goto statement (GOTO\_EXPR) to an ordinary label outside the sequence.
- 3. Execute a return statement (RETURN\_EXPR).
- 4. Throw an exception. This is currently not explicitly represented in GIMPLE.

The second sequence is not executed if the first sequence completes by calling setjmp or exit or any other function that does not return. The second sequence is also not executed if the first sequence completes via a non-local goto or a computed goto (in general the compiler does not know whether such a goto statement exits the first sequence or not, so we assume that it doesn't).

After the second sequence is executed, if it completes normally by falling off the end, execution continues wherever the first sequence would have continued, by falling off the end, or doing a goto, etc.

If the second sequence is an EH\_ELSE\_EXPR selector, then the sequence in its first operand is used when the first sequence completes normally, and that in its second operand is used for exceptional cleanups, i.e., when an exception propagates out of the first sequence.

TRY\_FINALLY\_EXPR complicates the flow graph, since the cleanup needs to appear on every edge out of the controlled block; this reduces the freedom to move code across these edges. Therefore, the EH lowering pass which runs before most of the optimization passes eliminates these expressions by explicitly adding the cleanup to each edge. Rethrowing the exception is represented using RESX\_EXPR.

# 11.7.7 OpenMP

All the statements starting with OMP\_ represent directives and clauses used by the OpenMP API https://www.openmp.org.

## OMP\_PARALLEL

Represents #pragma omp parallel [clause1 . . . clauseN]. It has four operands:

Operand OMP\_PARALLEL\_BODY is valid while in GENERIC and High GIMPLE forms. It contains the body of code to be executed by all the threads. During GIMPLE lowering, this operand becomes NULL and the body is emitted linearly after OMP\_PARALLEL.

Operand OMP\_PARALLEL\_CLAUSES is the list of clauses associated with the directive.

Operand OMP\_PARALLEL\_FN is created by pass\_lower\_omp, it contains the FUNCTION\_DECL for the function that will contain the body of the parallel region.

Operand OMP\_PARALLEL\_DATA\_ARG is also created by pass\_lower\_omp. If there are shared variables to be communicated to the children threads, this operand will contain the VAR\_DECL that contains all the shared values and variables.

### OMP\_FOR

Represents #pragma omp for [clause1 . . . clauseN]. It has six operands: Operand OMP\_FOR\_BODY contains the loop body.

Operand OMP\_FOR\_CLAUSES is the list of clauses associated with the directive.

Operand OMP\_FOR\_INIT is the loop initialization code of the form VAR = N1.

Operand OMP\_FOR\_COND is the loop conditional expression of the form VAR {<,>,<=,>=} N2.

Operand OMP\_FOR\_INCR is the loop index increment of the form VAR {+=,-=}

Operand OMP\_FOR\_PRE\_BODY contains side effect code from operands OMP\_FOR\_INIT, OMP\_FOR\_COND and OMP\_FOR\_INC. These side effects are part of the OMP\_FOR block but must be evaluated before the start of loop body.

The loop index variable VAR must be a signed integer variable, which is implicitly private to each thread. Bounds N1 and N2 and the increment expression INCR are required to be loop invariant integer expressions that are evaluated without any synchronization. The evaluation order, frequency of evaluation and side effects are unspecified by the standard.

### OMP\_SECTIONS

Represents #pragma omp sections [clause1 ... clauseN].

Operand OMP\_SECTIONS\_BODY contains the sections body, which in turn contains a set of OMP\_SECTION nodes for each of the concurrent sections delimited by #pragma omp section.

Operand OMP\_SECTIONS\_CLAUSES is the list of clauses associated with the directive.

### OMP\_SECTION

Section delimiter for OMP\_SECTIONS.

## OMP\_SINGLE

Represents #pragma omp single.

Operand OMP\_SINGLE\_BODY contains the body of code to be executed by a single thread.

Operand OMP\_SINGLE\_CLAUSES is the list of clauses associated with the directive.

### OMP\_MASTER

Represents #pragma omp master.

Operand OMP\_MASTER\_BODY contains the body of code to be executed by the master thread.

## OMP\_ORDERED

Represents #pragma omp ordered.

Operand OMP\_ORDERED\_BODY contains the body of code to be executed in the sequential order dictated by the loop index variable.

## OMP\_CRITICAL

Represents #pragma omp critical [name].

Operand OMP\_CRITICAL\_BODY is the critical section.

Operand OMP\_CRITICAL\_NAME is an optional identifier to label the critical section.

### OMP\_RETURN

This does not represent any OpenMP directive, it is an artificial marker to indicate the end of the body of an OpenMP. It is used by the flow graph (tree-cfg.c) and OpenMP region building code (omp-low.c).

## OMP\_CONTINUE

Similarly, this instruction does not represent an OpenMP directive, it is used by OMP\_FOR (and similar codes) as well as OMP\_SECTIONS to mark the place where the code needs to loop to the next iteration, or the next section, respectively.

In some cases, OMP\_CONTINUE is placed right before OMP\_RETURN. But if there are cleanups that need to occur right after the looping body, it will be emitted between OMP\_CONTINUE and OMP\_RETURN.

### OMP\_ATOMIC

Represents #pragma omp atomic.

Operand 0 is the address at which the atomic operation is to be performed.

Operand 1 is the expression to evaluate. The gimplifier tries three alternative code generation strategies. Whenever possible, an atomic update built-in is used. If that fails, a compare-and-swap loop is attempted. If that also fails, a regular critical section around the expression is used.

### OMP\_CLAUSE

Represents clauses associated with one of the OMP\_ directives. Clauses are represented by separate subcodes defined in 'tree.h'. Clauses codes can be one of: OMP\_CLAUSE\_PRIVATE, OMP\_CLAUSE\_SHARED, OMP\_CLAUSE\_FIRSTPRIVATE, OMP\_CLAUSE\_LASTPRIVATE, OMP\_CLAUSE\_COPYIN, OMP\_CLAUSE\_COPYPRIVATE, OMP\_CLAUSE\_IF, OMP\_CLAUSE\_NUM\_THREADS, OMP\_CLAUSE\_SCHEDULE, OMP\_CLAUSE\_NOWAIT, OMP\_CLAUSE\_ORDERED, OMP\_CLAUSE\_DEFAULT, OMP\_CLAUSE\_REDUCTION, OMP\_CLAUSE\_COLLAPSE, OMP\_CLAUSE\_UNTIED, OMP\_CLAUSE\_FINAL, and OMP\_CLAUSE\_MERGEABLE. Each code represents the corresponding OpenMP clause.

Clauses associated with the same directive are chained together via OMP\_CLAUSE\_CHAIN. Those clauses that accept a list of variables are restricted to exactly one, accessed with OMP\_CLAUSE\_VAR. Therefore, multiple variables under the same clause C need to be represented as multiple C clauses chained together. This facilitates adding new clauses during compilation.

## 11.7.8 OpenACC

All the statements starting with OACC\_ represent directives and clauses used by the OpenACC API https://www.openacc.org.

```
OACC_CACHE
```

Represents #pragma acc cache (var ...).

OACC\_DATA

Represents #pragma acc data [clause1 ... clauseN].

OACC\_DECLARE

Represents #pragma acc declare [clause1 ... clauseN].

OACC\_ENTER\_DATA

Represents #pragma acc enter data [clause1 ... clauseN].

OACC\_EXIT\_DATA

Represents #pragma acc exit data [clause1 ... clauseN].

OACC\_HOST\_DATA

Represents #pragma acc host\_data [clause1 ... clauseN].

OACC\_KERNELS

Represents #pragma acc kernels [clause1 ... clauseN].

OACC\_LOOP

Represents #pragma acc loop [clause1 ... clauseN].

See the description of the OMP\_FOR code.

OACC\_PARALLEL

Represents #pragma acc parallel [clause1 ... clauseN].

OACC\_SERIAL

Represents #pragma acc serial [clause1 ... clauseN].

OACC\_UPDATE

Represents #pragma acc update [clause1 ... clauseN].

### 11.8 Functions

A function is represented by a FUNCTION\_DECL node. It stores the basic pieces of the function such as body, parameters, and return type as well as information on the surrounding context, visibility, and linkage.

## 11.8.1 Function Basics

A function has four core parts: the name, the parameters, the result, and the body. The following macros and functions access these parts of a FUNCTION\_DECL as well as other basic features:

DECL\_NAME

This macro returns the unqualified name of the function, as an IDENTIFIER\_NODE. For an instantiation of a function template, the DECL\_NAME is the unqualified name of the template, not something like f<int>. The value of DECL\_NAME is undefined when used on a constructor, destructor, overloaded operator, or type-conversion operator, or any function that is implicitly generated by the compiler. See below for macros that can be used to distinguish these cases.

### DECL\_ASSEMBLER\_NAME

This macro returns the mangled name of the function, also an IDENTIFIER\_NODE. This name does not contain leading underscores on systems that prefix all identifiers with underscores. The mangled name is computed in the same way on all platforms; if special processing is required to deal with the object file format used on a particular platform, it is the responsibility of the back end to perform those modifications. (Of course, the back end should not modify DECL\_ASSEMBLER\_NAME itself.)

Using DECL\_ASSEMBLER\_NAME will cause additional memory to be allocated (for the mangled name of the entity) so it should be used only when emitting assembly code. It should not be used within the optimizers to determine whether or not two declarations are the same, even though some of the existing optimizers do use it in that way. These uses will be removed over time.

### DECL\_ARGUMENTS

This macro returns the PARM\_DECL for the first argument to the function. Subsequent PARM\_DECL nodes can be obtained by following the TREE\_CHAIN links.

DECL\_RESULT

This macro returns the RESULT\_DECL for the function.

### DECL\_SAVED\_TREE

This macro returns the complete body of the function.

TREE\_TYPE

This macro returns the FUNCTION\_TYPE or METHOD\_TYPE for the function.

### DECL\_INITIAL

A function that has a definition in the current translation unit will have a non-NULL DECL\_INITIAL. However, back ends should not make use of the particular value given by DECL\_INITIAL.

It should contain a tree of BLOCK nodes that mirrors the scopes that variables are bound in the function. Each block contains a list of decls declared in a basic block, a pointer to a chain of blocks at the next lower scope level, then a pointer to the next block at the same level and a backpointer to the parent BLOCK or FUNCTION\_DECL. So given a function as follows:

```
void foo()
      {
        int a:
        {
          int b;
        }
        int c;
you would get the following:
      tree foo = FUNCTION_DECL;
      tree decl_a = VAR_DECL;
      tree decl_b = VAR_DECL;
      tree decl_c = VAR_DECL;
      tree block_a = BLOCK;
      tree block_b = BLOCK;
      tree block_c = BLOCK;
      BLOCK_VARS(block_a) = decl_a;
      BLOCK SUBBLOCKS(block a) = block b:
      BLOCK CHAIN(block a) = block c:
      BLOCK SUPERCONTEXT(block a) = foo:
      BLOCK VARS(block b) = decl b:
      BLOCK_SUPERCONTEXT(block_b) = block_a;
      BLOCK_VARS(block_c) = decl_c;
      BLOCK_SUPERCONTEXT(block_c) =
      DECL_INITIAL(foo) = block_a;
```

# 11.8.2 Function Properties

To determine the scope of a function, you can use the DECL\_CONTEXT macro. This macro will return the class (either a RECORD\_TYPE or a UNION\_TYPE) or namespace (a NAMESPACE\_DECL) of which the function is a member. For a virtual function, this macro returns the class in which the function was actually defined, not the base class in which the virtual declaration occurred.

In C, the DECL\_CONTEXT for a function maybe another function. This representation indicates that the GNU nested function extension is in use. For details on the semantics of nested functions, see the GCC Manual. The nested function can refer to local variables in its containing function. Such references are not explicitly marked in the tree structure; back ends must look at the DECL\_CONTEXT for the referenced VAR\_DECL. If the DECL\_CONTEXT for the referenced VAR\_DECL is not the same as the function currently being processed, and

neither DECL\_EXTERNAL nor TREE\_STATIC hold, then the reference is to a local variable in a containing function, and the back end must take appropriate action.

### DECL\_EXTERNAL

This predicate holds if the function is undefined.

### TREE\_PUBLIC

This predicate holds if the function has external linkage.

### TREE\_STATIC

This predicate holds if the function has been defined.

### TREE\_THIS\_VOLATILE

This predicate holds if the function does not return normally.

### TREE\_READONLY

This predicate holds if the function can only read its arguments.

### DECL\_PURE\_P

This predicate holds if the function can only read its arguments, but may also read global memory.

## DECL\_VIRTUAL\_P

This predicate holds if the function is virtual.

### DECL\_ARTIFICIAL

This macro holds if the function was implicitly generated by the compiler, rather than explicitly declared. In addition to implicitly generated class member functions, this macro holds for the special functions created to implement static initialization and destruction, to compute run-time type information, and so forth.

## DECL\_FUNCTION\_SPECIFIC\_TARGET

This macro returns a tree node that holds the target options that are to be used to compile this particular function or NULL\_TREE if the function is to be compiled with the target options specified on the command line.

## DECL\_FUNCTION\_SPECIFIC\_OPTIMIZATION

This macro returns a tree node that holds the optimization options that are to be used to compile this particular function or NULL\_TREE if the function is to be compiled with the optimization options specified on the command line.

# 11.9 Language-dependent trees

Front ends may wish to keep some state associated with various GENERIC trees while parsing. To support this, trees provide a set of flags that may be used by the front end. They are accessed using TREE\_LANG\_FLAG\_n where 'n' is currently 0 through 6.

If necessary, a front end can use some language-dependent tree codes in its GENERIC representation, so long as it provides a hook for converting them to GIMPLE and doesn't expect them to work with any (hypothetical) optimizers that run before the conversion to GIMPLE. The intermediate representation used while parsing C and C++ looks very little like GENERIC, but the C and C++ gimplifier hooks are perfectly happy to take it as input and spit out GIMPLE.

## 11.10 C and C++ Trees

This section documents the internal representation used by GCC to represent C and C++ source programs. When presented with a C or C++ source program, GCC parses the program, performs semantic analysis (including the generation of error messages), and then produces the internal representation described here. This representation contains a complete representation for the entire translation unit provided as input to the front end. This representation is then typically processed by a code-generator in order to produce machine code, but could also be used in the creation of source browsers, intelligent editors, automatic documentation generators, interpreters, and any other programs needing the ability to process C or C++ code.

This section explains the internal representation. In particular, it documents the internal representation for C and C++ source constructs, and the macros, functions, and variables that can be used to access these constructs. The C++ representation is largely a superset of the representation used in the C front end. There is only one construct used in C that does not appear in the C++ front end and that is the GNU "nested function" extension. Many of the macros documented here do not apply in C because the corresponding language constructs do not appear in C.

The C and C++ front ends generate a mix of GENERIC trees and ones specific to C and C++. These language-specific trees are higher-level constructs than the ones in GENERIC to make the parser's job easier. This section describes those trees that aren't part of GENERIC as well as aspects of GENERIC trees that are treated in a language-specific manner.

If you are developing a "back end", be it is a code-generator or some other tool, that uses this representation, you may occasionally find that you need to ask questions not easily answered by the functions and macros available here. If that situation occurs, it is quite likely that GCC already supports the functionality you desire, but that the interface is simply not documented here. In that case, you should ask the GCC maintainers (via mail to gcc@gcc.gnu.org) about documenting the functionality you require. Similarly, if you find yourself writing functions that do not deal directly with your back end, but instead might be useful to other people using the GCC front end, you should submit your patches for inclusion in GCC.

# 11.10.1 Types for C++

In C++, an array type is not qualified; rather the type of the array elements is qualified. This situation is reflected in the intermediate representation. The macros described here will always examine the qualification of the underlying element type when applied to an array type. (If the element type is itself an array, then the recursion continues until a non-array type is found, and the qualification of this type is examined.) So, for example, CP\_TYPE\_CONST\_P will hold of the type const int () [7], denoting an array of seven ints.

The following functions and macros deal with cv-qualification of types:

### cp\_type\_quals

This function returns the set of type qualifiers applied to this type. This value is TYPE\_UNQUALIFIED if no qualifiers have been applied. The TYPE\_QUAL\_CONST bit is set if the type is const-qualified. The TYPE\_QUAL\_VOLATILE bit is set if the type is volatile-qualified. The TYPE\_QUAL\_RESTRICT bit is set if the type is restrict-qualified.

### CP\_TYPE\_CONST\_P

This macro holds if the type is const-qualified.

### CP\_TYPE\_VOLATILE\_P

This macro holds if the type is volatile-qualified.

### CP\_TYPE\_RESTRICT\_P

This macro holds if the type is restrict-qualified.

### CP\_TYPE\_CONST\_NON\_VOLATILE\_P

This predicate holds for a type that is const-qualified, but *not* volatile-qualified; other cv-qualifiers are ignored as well: only the const-ness is tested.

A few other macros and functions are usable with all types:

## TYPE\_SIZE

The number of bits required to represent the type, represented as an INTEGER\_CST. For an incomplete type, TYPE\_SIZE will be NULL\_TREE.

### TYPE\_ALIGN

The alignment of the type, in bits, represented as an int.

### TYPE\_NAME

This macro returns a declaration (in the form of a TYPE\_DECL) for the type. (Note this macro does *not* return an IDENTIFIER\_NODE, as you might expect, given its name!) You can look at the DECL\_NAME of the TYPE\_DECL to obtain the actual name of the type. The TYPE\_NAME will be NULL\_TREE for a type that is not a built-in type, the result of a typedef, or a named class type.

### CP\_INTEGRAL\_TYPE

This predicate holds if the type is an integral type. Notice that in C++, enumerations are *not* integral types.

### ARITHMETIC\_TYPE\_P

This predicate holds if the type is an integral type (in the C++ sense) or a floating point type.

### CLASS\_TYPE\_P

This predicate holds for a class-type.

## TYPE\_BUILT\_IN

This predicate holds for a built-in type.

## TYPE\_PTRDATAMEM\_P

This predicate holds if the type is a pointer to data member.

## TYPE\_PTR\_P

This predicate holds if the type is a pointer type, and the pointee is not a data member.

## TYPE\_PTRFN\_P

This predicate holds for a pointer to function type.

## TYPE\_PTROB\_P

This predicate holds for a pointer to object type. Note however that it does not hold for the generic pointer to object type void \*. You may use TYPE\_PTROBV\_P to test for a pointer to object type as well as void \*.

The table below describes types specific to C and C++ as well as language-dependent info about GENERIC types.

### POINTER\_TYPE

Used to represent pointer types, and pointer to data member types. If TREE\_TYPE is a pointer to data member type, then TYPE\_PTRDATAMEM\_P will hold. For a pointer to data member type of the form 'T X::\*', TYPE\_PTRMEM\_CLASS\_TYPE will be the type X, while TYPE\_PTRMEM\_POINTED\_TO\_TYPE will be the type T.

### RECORD\_TYPE

Used to represent struct and class types in C and C++. If TYPE\_PTRMEMFUNC\_P holds, then this type is a pointer-to-member type. In that case, the TYPE\_PTRMEMFUNC\_FN\_TYPE is a POINTER\_TYPE pointing to a METHOD\_TYPE. The METHOD\_TYPE is the type of a function pointed to by the pointer-to-member function. If TYPE\_PTRMEMFUNC\_P does not hold, this type is a class type. For more information, see Section 11.10.3 [Classes], page 199.

### UNKNOWN\_TYPE

This node is used to represent a type the knowledge of which is insufficient for a sound processing.

### TYPENAME\_TYPE

Used to represent a construct of the form typename T::A. The TYPE\_CONTEXT is T; the TYPE\_NAME is an IDENTIFIER\_NODE for A. If the type is specified via a template-id, then TYPENAME\_TYPE\_FULLNAME yields a TEMPLATE\_ID\_EXPR. The TREE\_TYPE is non-NULL if the node is implicitly generated in support for the implicit typename extension; in which case the TREE\_TYPE is a type node for the base-class.

### TYPEOF\_TYPE

Used to represent the \_\_typeof\_\_ extension. The TYPE\_FIELDS is the expression the type of which is being represented.

## 11.10.2 Namespaces

The root of the entire intermediate representation is the variable global\_namespace. This is the namespace specified with :: in C++ source code. All other namespaces, types, variables, functions, and so forth can be found starting with this namespace.

However, except for the fact that it is distinguished as the root of the representation, the global namespace is no different from any other namespace. Thus, in what follows, we describe namespaces generally, rather than the global namespace in particular.

A namespace is represented by a NAMESPACE\_DECL node.

The following macros and functions can be used on a NAMESPACE\_DECL:

### DECL\_NAME

This macro is used to obtain the IDENTIFIER\_NODE corresponding to the unqualified name of the name of the namespace (see Section 11.2.2 [Identifiers], page 163). The name of the global namespace is '::', even though in C++ the global namespace is unnamed. However, you should use comparison with global\_namespace, rather than DECL\_NAME to determine whether or not a

namespace is the global one. An unnamed namespace will have a DECL\_NAME equal to anonymous\_namespace\_name. Within a single translation unit, all unnamed namespaces will have the same name.

### DECL\_CONTEXT

This macro returns the enclosing namespace. The DECL\_CONTEXT for the global\_namespace is NULL\_TREE.

### DECL\_NAMESPACE\_ALIAS

If this declaration is for a namespace alias, then DECL\_NAMESPACE\_ALIAS is the namespace for which this one is an alias.

Do not attempt to use cp\_namespace\_decls for a namespace which is an alias. Instead, follow DECL\_NAMESPACE\_ALIAS links until you reach an ordinary, non-alias, namespace, and call cp\_namespace\_decls there.

## DECL\_NAMESPACE\_STD\_P

This predicate holds if the namespace is the special ::std namespace.

## cp\_namespace\_decls

This function will return the declarations contained in the namespace, including types, overloaded functions, other namespaces, and so forth. If there are no declarations, this function will return NULL\_TREE. The declarations are connected through their TREE\_CHAIN fields.

Although most entries on this list will be declarations, TREE\_LIST nodes may also appear. In this case, the TREE\_VALUE will be an OVERLOAD. The value of the TREE\_PURPOSE is unspecified; back ends should ignore this value. As with the other kinds of declarations returned by cp\_namespace\_decls, the TREE\_CHAIN will point to the next declaration in this list.

For more information on the kinds of declarations that can occur on this list, See Section 11.4 [Declarations], page 168. Some declarations will not appear on this list. In particular, no FIELD\_DECL, LABEL\_DECL, or PARM\_DECL nodes will appear here.

This function cannot be used with namespaces that have DECL\_NAMESPACE\_ALIAS set.

## 11.10.3 Classes

Besides namespaces, the other high-level scoping construct in C++ is the class. (Throughout this manual the term *class* is used to mean the types referred to in the ANSI/ISO C++ Standard as classes; these include types defined with the class, struct, and union keywords.)

A class type is represented by either a RECORD\_TYPE or a UNION\_TYPE. A class declared with the union tag is represented by a UNION\_TYPE, while classes declared with either the struct or the class tag are represented by RECORD\_TYPEs. You can use the CLASSTYPE\_DECLARED\_CLASS macro to discern whether or not a particular type is a class as opposed to a struct. This macro will be true only for classes declared with the class tag.

Almost all members are available on the TYPE\_FIELDS list. Given one member, the next can be found by following the TREE\_CHAIN. You should not depend in any way on the order in which fields appear on this list. All nodes on this list will be 'DECL' nodes. A

FIELD\_DECL is used to represent a non-static data member, a VAR\_DECL is used to represent a static data member, and a TYPE\_DECL is used to represent a type. Note that the CONST\_DECL for an enumeration constant will appear on this list, if the enumeration type was declared in the class. (Of course, the TYPE\_DECL for the enumeration type will appear here as well.) There are no entries for base classes on this list. In particular, there is no FIELD\_DECL for the "base-class portion" of an object. If a function member is overloaded, each of the overloaded functions appears; no OVERLOAD nodes appear on the TYPE\_FIELDS list. Implicitly declared functions (including default constructors, copy constructors, assignment operators, and destructors) will appear on this list as well.

The TYPE\_VFIELD is a compiler-generated field used to point to virtual function tables. It may or may not appear on the TYPE\_FIELDS list. However, back ends should handle the TYPE\_VFIELD just like all the entries on the TYPE\_FIELDS list.

Every class has an associated binfo, which can be obtained with TYPE\_BINFO. Binfos are used to represent base-classes. The binfo given by TYPE\_BINFO is the degenerate case, whereby every class is considered to be its own base-class. The base binfos for a particular binfo are held in a vector, whose length is obtained with BINFO\_N\_BASE\_BINFOS. The base binfos themselves are obtained with BINFO\_BASE\_BINFO and BINFO\_BASE\_ITERATE. To add a new binfo, use BINFO\_BASE\_APPEND. The vector of base binfos can be obtained with BINFO\_BASE\_BINFOS, but normally you do not need to use that. The class type associated with a binfo is given by BINFO\_TYPE. It is not always the case that BINFO\_TYPE (TYPE\_BINFO (x)), because of typedefs and qualified types. Neither is it the case that TYPE\_BINFO (BINFO\_TYPE (y)) is the same binfo as y. The reason is that if y is a binfo representing a base-class B of a derived class D, then BINFO\_TYPE (y) will be B, and TYPE\_BINFO (BINFO\_TYPE (y)) will be B as its own base-class, rather than as a base-class of D.

The access to a base type can be found with BINFO\_BASE\_ACCESS. This will produce access\_public\_node, access\_private\_node or access\_protected\_node. If bases are always public, BINFO\_BASE\_ACCESSES may be NULL.

BINFO\_VIRTUAL\_P is used to specify whether the binfo is inherited virtually or not. The other flags, BINFO\_FLAG\_0 to BINFO\_FLAG\_6, can be used for language specific use.

The following macros can be used on a tree node representing a class-type.

### LOCAL\_CLASS\_P

This predicate holds if the class is local class i.e. declared inside a function body.

## TYPE\_POLYMORPHIC\_P

This predicate holds if the class has at least one virtual function (declared or inherited).

### TYPE\_HAS\_DEFAULT\_CONSTRUCTOR

This predicate holds whenever its argument represents a class-type with default constructor.

### CLASSTYPE\_HAS\_MUTABLE

### TYPE\_HAS\_MUTABLE\_P

These predicates hold for a class-type having a mutable data member.

### CLASSTYPE\_NON\_POD\_P

This predicate holds only for class-types that are not PODs.

#### TYPE\_HAS\_NEW\_OPERATOR

This predicate holds for a class-type that defines operator new.

#### TYPE\_HAS\_ARRAY\_NEW\_OPERATOR

This predicate holds for a class-type for which operator new[] is defined.

#### TYPE\_OVERLOADS\_CALL\_EXPR

This predicate holds for class-type for which the function call operator() is overloaded.

#### TYPE\_OVERLOADS\_ARRAY\_REF

This predicate holds for a class-type that overloads operator[]

# TYPE\_OVERLOADS\_ARROW

This predicate holds for a class-type for which operator-> is overloaded.

#### 11.10.4 Functions for C++

A function is represented by a FUNCTION\_DECL node. A set of overloaded functions is sometimes represented by an OVERLOAD node.

An OVERLOAD node is not a declaration, so none of the 'DECL\_' macros should be used on an OVERLOAD. An OVERLOAD node is similar to a TREE\_LIST. Use OVL\_CURRENT to get the function associated with an OVERLOAD node; use OVL\_NEXT to get the next OVERLOAD node in the list of overloaded functions. The macros OVL\_CURRENT and OVL\_NEXT are actually polymorphic; you can use them to work with FUNCTION\_DECL nodes as well as with overloads. In the case of a FUNCTION\_DECL, OVL\_CURRENT will always return the function itself, and OVL\_NEXT will always be NULL\_TREE.

To determine the scope of a function, you can use the DECL\_CONTEXT macro. This macro will return the class (either a RECORD\_TYPE or a UNION\_TYPE) or namespace (a NAMESPACE\_DECL) of which the function is a member. For a virtual function, this macro returns the class in which the function was actually defined, not the base class in which the virtual declaration occurred.

If a friend function is defined in a class scope, the DECL\_FRIEND\_CONTEXT macro can be used to determine the class in which it was defined. For example, in

```
class C { friend void f() {} };
```

the DECL\_CONTEXT for f will be the global\_namespace, but the DECL\_FRIEND\_CONTEXT will be the RECORD\_TYPE for C.

The following macros and functions can be used on a FUNCTION\_DECL:

#### DECL\_MAIN\_P

This predicate holds for a function that is the program entry point ::code.

#### DECL\_LOCAL\_FUNCTION\_P

This predicate holds if the function was declared at block scope, even though it has a global scope.

#### DECL\_ANTICIPATED

This predicate holds if the function is a built-in function but its prototype is not yet explicitly declared.

#### DECL\_EXTERN\_C\_FUNCTION\_P

This predicate holds if the function is declared as an 'extern "C"' function.

#### DECL\_LINKONCE\_P

This macro holds if multiple copies of this function may be emitted in various translation units. It is the responsibility of the linker to merge the various copies. Template instantiations are the most common example of functions for which DECL\_LINKONCE\_P holds; G++ instantiates needed templates in all translation units which require them, and then relies on the linker to remove duplicate instantiations.

FIXME: This macro is not yet implemented.

# DECL\_FUNCTION\_MEMBER\_P

This macro holds if the function is a member of a class, rather than a member of a namespace.

#### DECL\_STATIC\_FUNCTION\_P

This predicate holds if the function a static member function.

# DECL\_NONSTATIC\_MEMBER\_FUNCTION\_P

This macro holds for a non-static member function.

#### DECL\_CONST\_MEMFUNC\_P

This predicate holds for a const-member function.

# DECL\_VOLATILE\_MEMFUNC\_P

This predicate holds for a volatile-member function.

#### DECL\_CONSTRUCTOR\_P

This macro holds if the function is a constructor.

# DECL\_NONCONVERTING\_P

This predicate holds if the constructor is a non-converting constructor.

#### DECL\_COMPLETE\_CONSTRUCTOR\_P

This predicate holds for a function which is a constructor for an object of a complete type.

# DECL\_BASE\_CONSTRUCTOR\_P

This predicate holds for a function which is a constructor for a base class subobject.

#### DECL\_COPY\_CONSTRUCTOR\_P

This predicate holds for a function which is a copy-constructor.

#### DECL\_DESTRUCTOR\_P

This macro holds if the function is a destructor.

# DECL\_COMPLETE\_DESTRUCTOR\_P

This predicate holds if the function is the destructor for an object a complete type.

## DECL\_OVERLOADED\_OPERATOR\_P

This macro holds if the function is an overloaded operator.

#### DECL\_CONV\_FN\_P

This macro holds if the function is a type-conversion operator.

#### DECL\_GLOBAL\_CTOR\_P

This predicate holds if the function is a file-scope initialization function.

# DECL\_GLOBAL\_DTOR\_P

This predicate holds if the function is a file-scope finalization function.

#### DECL\_THUNK\_P

This predicate holds if the function is a thunk.

These functions represent stub code that adjusts the this pointer and then jumps to another function. When the jumped-to function returns, control is transferred directly to the caller, without returning to the thunk. The first parameter to the thunk is always the this pointer; the thunk should add THUNK\_DELTA to this value. (The THUNK\_DELTA is an int, not an INTEGER\_CST.)

Then, if THUNK\_VCALL\_OFFSET (an INTEGER\_CST) is nonzero the adjusted this pointer must be adjusted again. The complete calculation is given by the following pseudo-code:

```
this += THUNK_DELTA
if (THUNK_VCALL_OFFSET)
  this += (*((ptrdiff_t **) this))[THUNK_VCALL_OFFSET]
```

Finally, the thunk should jump to the location given by DECL\_INITIAL; this will always be an expression for the address of a function.

#### DECL\_NON\_THUNK\_FUNCTION\_P

This predicate holds if the function is *not* a thunk function.

## GLOBAL\_INIT\_PRIORITY

If either DECL\_GLOBAL\_CTOR\_P or DECL\_GLOBAL\_DTOR\_P holds, then this gives the initialization priority for the function. The linker will arrange that all functions for which DECL\_GLOBAL\_CTOR\_P holds are run in increasing order of priority before main is called. When the program exits, all functions for which DECL\_GLOBAL\_DTOR\_P holds are run in the reverse order.

# TYPE\_RAISES\_EXCEPTIONS

This macro returns the list of exceptions that a (member-)function can raise. The returned list, if non NULL, is comprised of nodes whose TREE\_VALUE represents a type.

#### TYPE\_NOTHROW\_P

This predicate holds when the exception-specification of its arguments is of the form '()'.

#### DECL\_ARRAY\_DELETE\_OPERATOR\_P

This predicate holds if the function an overloaded operator delete[].

#### 11.10.5 Statements for C++

A function that has a definition in the current translation unit will have a non-NULL DECL\_INITIAL. However, back ends should not make use of the particular value given by DECL\_INITIAL.

The DECL\_SAVED\_TREE macro will give the complete body of the function.

#### 11.10.5.1 Statements

There are tree nodes corresponding to all of the source-level statement constructs, used within the C and C++ frontends. These are enumerated here, together with a list of the various macros that can be used to obtain information about them. There are a few macros that can be used with all statements:

#### STMT\_IS\_FULL\_EXPR\_P

In C++, statements normally constitute "full expressions"; temporaries created during a statement are destroyed when the statement is complete. However, G++ sometimes represents expressions by statements; these statements will not have STMT\_IS\_FULL\_EXPR\_P set. Temporaries created during such statements should be destroyed when the innermost enclosing statement with STMT\_IS\_FULL\_EXPR\_P set is exited.

Here is the list of the various statement nodes, and the macros used to access them. This documentation describes the use of these nodes in non-template functions (including instantiations of template functions). In template functions, the same nodes are used, but sometimes in slightly different ways.

Many of the statements have substatements. For example, a while loop will have a body, which is itself a statement. If the substatement is NULL\_TREE, it is considered equivalent to a statement consisting of a single;, i.e., an expression statement in which the expression has been omitted. A substatement may in fact be a list of statements, connected via their TREE\_CHAINS. So, you should always process the statement tree by looping over substatements, like this:

In other words, while the then clause of an if statement in C++ can be only one statement (although that one statement may be a compound statement), the intermediate representation will sometimes use several statements chained together.

#### BREAK\_STMT

Used to represent a break statement. There are no additional fields.

#### CLEANUP\_STMT

Used to represent an action that should take place upon exit from the enclosing scope. Typically, these actions are calls to destructors for local objects,

but back ends cannot rely on this fact. If these nodes are in fact representing such destructors, CLEANUP\_DECL will be the VAR\_DECL destroyed. Otherwise, CLEANUP\_DECL will be NULL\_TREE. In any case, the CLEANUP\_EXPR is the expression to execute. The cleanups executed on exit from a scope should be run in the reverse order of the order in which the associated CLEANUP\_STMTs were encountered.

# CONTINUE\_STMT

Used to represent a continue statement. There are no additional fields.

#### CTOR\_STMT

Used to mark the beginning (if CTOR\_BEGIN\_P holds) or end (if CTOR\_END\_P holds of the main body of a constructor. See also SUBOBJECT for more information on how to use these nodes.

# DO\_STMT

Used to represent a do loop. The body of the loop is given by DO\_BODY while the termination condition for the loop is given by DO\_COND. The condition for a do-statement is always an expression.

# EMPTY\_CLASS\_EXPR

Used to represent a temporary object of a class with no data whose address is never taken. (All such objects are interchangeable.) The TREE\_TYPE represents the type of the object.

#### EXPR\_STMT

Used to represent an expression statement. Use EXPR\_STMT\_EXPR to obtain the expression.

#### FOR\_STMT

Used to represent a for statement. The FOR\_INIT\_STMT is the initialization statement for the loop. The FOR\_COND is the termination condition. The FOR\_EXPR is the expression executed right before the FOR\_COND on each loop iteration; often, this expression increments a counter. The body of the loop is given by FOR\_BODY. Note that FOR\_INIT\_STMT and FOR\_BODY return statements, while FOR\_COND and FOR\_EXPR return expressions.

#### HANDLER

Used to represent a C++ catch block. The HANDLER\_TYPE is the type of exception that will be caught by this handler; it is equal (by pointer equality) to NULL if this handler is for all types. HANDLER\_PARMS is the DECL\_STMT for the catch parameter, and HANDLER\_BODY is the code for the block itself.

# IF\_STMT

Used to represent an if statement. The IF\_COND is the expression.

If the condition is a TREE\_LIST, then the TREE\_PURPOSE is a statement (usually a DECL\_STMT). Each time the condition is evaluated, the statement should be executed. Then, the TREE\_VALUE should be used as the conditional expression itself. This representation is used to handle C++ code like this:

C++ distinguishes between this and COND\_EXPR for handling templates.

if (int 
$$i = 7$$
) ...

where there is a new local variable (or variables) declared within the condition. The THEN\_CLAUSE represents the statement given by the then condition, while the ELSE\_CLAUSE represents the statement given by the else condition.

#### SUBOBJECT

In a constructor, these nodes are used to mark the point at which a subobject of this is fully constructed. If, after this point, an exception is thrown before a CTOR\_STMT with CTOR\_END\_P set is encountered, the SUBOBJECT\_CLEANUP must be executed. The cleanups must be executed in the reverse order in which they appear.

#### SWITCH\_STMT

Used to represent a switch statement. The SWITCH\_STMT\_COND is the expression on which the switch is occurring. See the documentation for an IF\_STMT for more information on the representation used for the condition. The SWITCH\_STMT\_BODY is the body of the switch statement. The SWITCH\_STMT\_TYPE is the original type of switch expression as given in the source, before any compiler conversions.

#### TRY\_BLOCK

Used to represent a try block. The body of the try block is given by TRY\_STMTS. Each of the catch blocks is a HANDLER node. The first handler is given by TRY\_HANDLERS. Subsequent handlers are obtained by following the TREE\_CHAIN link from one handler to the next. The body of the handler is given by HANDLER\_BODY.

If CLEANUP\_P holds of the TRY\_BLOCK, then the TRY\_HANDLERS will not be a HANDLER node. Instead, it will be an expression that should be executed if an exception is thrown in the try block. It must rethrow the exception after executing that code. And, if an exception is thrown while the expression is executing, terminate must be called.

#### USING\_STMT

Used to represent a using directive. The namespace is given by USING\_STMT\_NAMESPACE, which will be a NAMESPACE\_DECL. This node is needed inside template functions, to implement using directives during instantiation.

#### WHILE\_STMT

Used to represent a while loop. The WHILE\_COND is the termination condition for the loop. See the documentation for an IF\_STMT for more information on the representation used for the condition.

The WHILE\_BODY is the body of the loop.

# 11.10.6 C++ Expressions

This section describes expressions specific to the C and C++ front ends.

#### TYPEID\_EXPR

Used to represent a typeid expression.

NEW\_EXPR

VEC\_NEW\_EXPR

Used to represent a call to new and new[] respectively.

DELETE\_EXPR

VEC\_DELETE\_EXPR

Used to represent a call to delete and delete[] respectively.

MEMBER\_REF

Represents a reference to a member of a class.

THROW\_EXPR

Represents an instance of throw in the program. Operand 0, which is the expression to throw, may be NULL\_TREE.

#### AGGR\_INIT\_EXPR

An AGGR\_INIT\_EXPR represents the initialization as the return value of a function call, or as the result of a constructor. An AGGR\_INIT\_EXPR will only appear as a full-expression, or as the second operand of a TARGET\_EXPR. AGGR\_INIT\_EXPRs have a representation similar to that of CALL\_EXPRs. You can use the AGGR\_INIT\_EXPR\_FN and AGGR\_INIT\_EXPR\_ARG macros to access the function to call and the arguments to pass.

If AGGR\_INIT\_VIA\_CTOR\_P holds of the AGGR\_INIT\_EXPR, then the initialization is via a constructor call. The address of the AGGR\_INIT\_EXPR\_SLOT operand, which is always a VAR\_DECL, is taken, and this value replaces the first argument in the argument list.

In either case, the expression is void.

# 12 GIMPLE

GIMPLE is a three-address representation derived from GENERIC by breaking down GENERIC expressions into tuples of no more than 3 operands (with some exceptions like function calls). GIMPLE was heavily influenced by the SIMPLE IL used by the McCAT compiler project at McGill University, though we have made some different choices. For one thing, SIMPLE doesn't support goto.

Temporaries are introduced to hold intermediate values needed to compute complex expressions. Additionally, all the control structures used in GENERIC are lowered into conditional jumps, lexical scopes are removed and exception regions are converted into an on the side exception region tree.

The compiler pass which converts GENERIC into GIMPLE is referred to as the 'gimplifier'. The gimplifier works recursively, generating GIMPLE tuples out of the original GENERIC expressions.

One of the early implementation strategies used for the GIMPLE representation was to use the same internal data structures used by front ends to represent parse trees. This simplified implementation because we could leverage existing functionality and interfaces. However, GIMPLE is a much more restrictive representation than abstract syntax trees (AST), therefore it does not require the full structural complexity provided by the main tree data structure.

The GENERIC representation of a function is stored in the DECL\_SAVED\_TREE field of the associated FUNCTION\_DECL tree node. It is converted to GIMPLE by a call to gimplify\_function\_tree.

If a front end wants to include language-specific tree codes in the tree representation which it provides to the back end, it must provide a definition of LANG\_HOOKS\_GIMPLIFY\_EXPR which knows how to convert the front end trees to GIMPLE. Usually such a hook will involve much of the same code for expanding front end trees to RTL. This function can return fully lowered GIMPLE, or it can return GENERIC trees and let the main gimplifier lower them the rest of the way; this is often simpler. GIMPLE that is not fully lowered is known as "High GIMPLE" and consists of the IL before the pass pass\_lower\_cf. High GIMPLE contains some container statements like lexical scopes (represented by GIMPLE\_BIND) and nested expressions (e.g., GIMPLE\_TRY), while "Low GIMPLE" exposes all of the implicit jumps for control and exception expressions directly in the IL and EH region trees.

The C and C++ front ends currently convert directly from front end trees to GIMPLE, and hand that off to the back end rather than first converting to GENERIC. Their gimplifier hooks know about all the \_STMT nodes and how to convert them to GENERIC forms. There was some work done on a genericization pass which would run first, but the existence of STMT\_EXPR meant that in order to convert all of the C statements into GENERIC equivalents would involve walking the entire tree anyway, so it was simpler to lower all the way. This might change in the future if someone writes an optimization pass which would work better with higher-level trees, but currently the optimizers all expect GIMPLE.

You can request to dump a C-like representation of the GIMPLE form with the flag '-fdump-tree-gimple'.

# 12.1 Tuple representation

GIMPLE instructions are tuples of variable size divided in two groups: a header describing the instruction and its locations, and a variable length body with all the operands. Tuples are organized into a hierarchy with 3 main classes of tuples.

# 12.1.1 gimple (gsbase)

This is the root of the hierarchy, it holds basic information needed by most GIMPLE statements. There are some fields that may not be relevant to every GIMPLE statement, but those were moved into the base structure to take advantage of holes left by other fields (thus making the structure more compact). The structure takes 4 words (32 bytes) on 64 bit hosts:

Field	Size (bits)
code	8
subcode	16
no_warning	1
visited	1
nontemporal_move	1
plf	2
modified	1
has_volatile_ops	1
references_memory_p	1
uid	32
location	32
num_ops	32
bb	64
block	63
Total size	32 bytes

- code Main identifier for a GIMPLE instruction.
- subcode Used to distinguish different variants of the same basic instruction or provide flags applicable to a given code. The subcode flags field has different uses depending on the code of the instruction, but mostly it distinguishes instructions of the same family. The most prominent use of this field is in assignments, where subcode indicates the operation done on the RHS of the assignment. For example, a = b + c is encoded as GIMPLE\_ASSIGN <PLUS\_EXPR, a, b, c>.
- no\_warning Bitflag to indicate whether a warning has already been issued on this statement.
- visited General purpose "visited" marker. Set and cleared by each pass when needed.
- nontemporal\_move Bitflag used in assignments that represent non-temporal moves. Although this bitflag is only used in assignments, it was moved into the base to take advantage of the bit holes left by the previous fields.
- plf Pass Local Flags. This 2-bit mask can be used as general purpose markers by any pass. Passes are responsible for clearing and setting these two flags accordingly.
- modified Bitflag to indicate whether the statement has been modified. Used mainly by the operand scanner to determine when to re-scan a statement for operands.

- has\_volatile\_ops Bitflag to indicate whether this statement contains operands that have been marked volatile.
- references\_memory\_p Bitflag to indicate whether this statement contains memory references (i.e., its operands are either global variables, or pointer dereferences or anything that must reside in memory).
- uid This is an unsigned integer used by passes that want to assign IDs to every statement. These IDs must be assigned and used by each pass.
- location This is a location\_t identifier to specify source code location for this statement. It is inherited from the front end.
- num\_ops Number of operands that this statement has. This specifies the size of the operand vector embedded in the tuple. Only used in some tuples, but it is declared in the base tuple to take advantage of the 32-bit hole left by the previous fields.
- bb Basic block holding the instruction.
- block Lexical block holding this statement. Also used for debug information generation.

# 12.1.2 gimple\_statement\_with\_ops

This tuple is actually split in two: gimple\_statement\_with\_ops\_base and gimple\_statement\_with\_ops. This is needed to accommodate the way the operand vector is allocated. The operand vector is defined to be an array of 1 element. So, to allocate a dynamic number of operands, the memory allocator (gimple\_alloc) simply allocates enough memory to hold the structure itself plus N - 1 operands which run "off the end" of the structure. For example, to allocate space for a tuple with 3 operands, gimple\_alloc reserves sizeof (struct gimple\_statement\_with\_ops) + 2 \* sizeof (tree) bytes.

On the other hand, several fields in this tuple need to be shared with the gimple\_statement\_with\_memory\_ops tuple. So, these common fields are placed in gimple\_statement\_with\_ops\_base which is then inherited from the other two tuples.

```
\begin{array}{lll} {\rm gsbase} & 256 \\ {\rm def\_ops} & 64 \\ {\rm use\_ops} & 64 \\ {\rm op} & {\rm num\_ops} * 64 \\ {\rm Total} & 48 + 8 * {\rm num\_ops} \; {\rm bytes} \\ {\rm size} \end{array}
```

- gsbase Inherited from struct gimple.
- def\_ops Array of pointers into the operand array indicating all the slots that contain a variable written-to by the statement. This array is also used for immediate use chaining. Note that it would be possible to not rely on this array, but the changes required to implement this are pretty invasive.
- use\_ops Similar to def\_ops but for variables read by the statement.
- op Array of trees with num\_ops slots.

# 12.1.3 gimple\_statement\_with\_memory\_ops

This tuple is essentially identical to gimple\_statement\_with\_ops, except that it contains 4 additional fields to hold vectors related memory stores and loads. Similar to the pre-

vious case, the structure is split in two to accommodate for the operand vector (gimple\_statement\_with\_memory\_ops\_base and gimple\_statement\_with\_memory\_ops).

```
Field
            Size (bits)
gsbase
            256
            64
def_ops
            64
use_ops
            64
vdef_ops
            64
vuse_ops
stores
            64
            64
loads
            num_ops*64
oр
            80 + 8 * num_{ops} bytes
Total size
```

- vdef\_ops Similar to def\_ops but for VDEF operators. There is one entry per memory symbol written by this statement. This is used to maintain the memory SSA use-def and def-def chains.
- vuse\_ops Similar to use\_ops but for VUSE operators. There is one entry per memory symbol loaded by this statement. This is used to maintain the memory SSA use-def chains.
- stores Bitset with all the UIDs for the symbols written-to by the statement. This is different than vdef\_ops in that all the affected symbols are mentioned in this set. If memory partitioning is enabled, the vdef\_ops vector will refer to memory partitions. Furthermore, no SSA information is stored in this set.
- loads Similar to stores, but for memory loads. (Note that there is some amount of redundancy here, it should be possible to reduce memory utilization further by removing these sets).

All the other tuples are defined in terms of these three basic ones. Each tuple will add some fields.

# 12.2 Class hierarchy of GIMPLE statements

The following diagram shows the C++ inheritance hierarchy of statement kinds, along with their relationships to GSS\_ values (layouts) and GIMPLE\_ values (codes):

```
code: GIMPLE_DEBUG
     + ggoto
          code: GIMPLE_GOTO
     + glabel
         code: GIMPLE_LABEL
    + gswitch
         code: GIMPLE_SWITCH
 + gimple_statement_with_memory_ops_base
     | layout: GSS_WITH_MEM_OPS_BASE
     + gimple_statement_with_memory_ops
     | | layout: GSS_WITH_MEM_OPS
        + gassign
     | | code GIMPLE_ASSIGN
     + greturn
             code GIMPLE_RETURN
     + gcall
             layout: GSS_CALL, code: GIMPLE_CALL
     + gasm
             layout: GSS_ASM, code: GIMPLE_ASM
     + gtransaction
             layout: GSS_TRANSACTION, code: GIMPLE_TRANSACTION
gimple_statement_omp
layout: GSS_OMP. Used for code GIMPLE_OMP_SECTION
 + gomp_critical
        layout: GSS_OMP_CRITICAL, code: GIMPLE_OMP_CRITICAL
 + gomp_for
     layout: GSS_OMP_FOR, code: GIMPLE_OMP_FOR
 + gomp_parallel_layout
          layout: GSS_OMP_PARALLEL_LAYOUT
 1 1
     + gimple_statement_omp_taskreg
     1 1
        + gomp_parallel
        code: GIMPLE_OMP_PARALLEL
     - 1
     | + gomp_task
          code: GIMPLE_OMP_TASK
    + gimple_statement_omp_target
       code: GIMPLE_OMP_TARGET
  + gomp_sections
          layout: GSS_OMP_SECTIONS, code: GIMPLE_OMP_SECTIONS
```

```
+ gimple_statement_omp_single_layout
          layout: GSS_OMP_SINGLE_LAYOUT
      + gomp_single
              code: GIMPLE_OMP_SINGLE
      + gomp_teams
              code: GIMPLE_OMP_TEAMS
gbind
       layout: GSS_BIND, code: GIMPLE_BIND
gcatch
       layout: GSS_CATCH, code: GIMPLE_CATCH
geh_filter
       layout: GSS_EH_FILTER, code: GIMPLE_EH_FILTER
geh_else
       layout: GSS_EH_ELSE, code: GIMPLE_EH_ELSE
geh_mnt
       layout: GSS_EH_MNT, code: GIMPLE_EH_MUST_NOT_THROW
gphi
       layout: GSS_PHI, code: GIMPLE_PHI
gimple_statement_eh_ctrl
       layout: GSS_EH_CTRL
 + gresx
code: GIMPLE_RESX
 + geh_dispatch
         code: GIMPLE_EH_DISPATCH
gtry
       layout: GSS_TRY, code: GIMPLE_TRY
gimple_statement_wce
      layout: GSS_WCE, code: GIMPLE_WITH_CLEANUP_EXPR
gomp_continue
       layout: GSS_OMP_CONTINUE, code: GIMPLE_OMP_CONTINUE
gomp_atomic_load
       layout: GSS_OMP_ATOMIC_LOAD, code: GIMPLE_OMP_ATOMIC_LOAD
gimple_statement_omp_atomic_store_layout
       layout: GSS_OMP_ATOMIC_STORE_LAYOUT,
       code: GIMPLE_OMP_ATOMIC_STORE
  + gomp_atomic_store
          code: GIMPLE_OMP_ATOMIC_STORE
  + gomp_return
           code: GIMPLE_OMP_RETURN
```

# 12.3 GIMPLE instruction set

The following table briefly describes the GIMPLE instruction set.

O v		
Instruction	High GIMPLE	Low GIMPLE
GIMPLE_ASM	X	X
GIMPLE_ASSIGN	X	X
GIMPLE_BIND	X	
GIMPLE_CALL	X	X
GIMPLE_CATCH	X	
GIMPLE_COND	X	X
GIMPLE_DEBUG	X	X
GIMPLE_EH_FILTER	X	
GIMPLE_GOTO	X	X
GIMPLE_LABEL	X	X
GIMPLE_NOP	X	X
GIMPLE_OMP_ATOMIC_LOAD	X	X
GIMPLE_OMP_ATOMIC_STORE	X	X
GIMPLE_OMP_CONTINUE	X	X
GIMPLE_OMP_CRITICAL	X	X
GIMPLE_OMP_FOR	X	X
GIMPLE_OMP_MASTER	X	X
GIMPLE_OMP_ORDERED	X	X
GIMPLE_OMP_PARALLEL	X	X
GIMPLE_OMP_RETURN	X	X
GIMPLE_OMP_SECTION	X	X
GIMPLE_OMP_SECTIONS	X	X
GIMPLE_OMP_SECTIONS_SWITCH	X	X
GIMPLE_OMP_SINGLE	X	X
GIMPLE_PHI		X
GIMPLE_RESX		X
GIMPLE_RETURN	X	X
GIMPLE_SWITCH	X	X
GIMPLE_TRY	X	

# 12.4 Exception Handling

Other exception handling constructs are represented using GIMPLE\_TRY\_CATCH. GIMPLE\_TRY\_CATCH has two operands. The first operand is a sequence of statements to execute. If executing these statements does not throw an exception, then the second operand is ignored. Otherwise, if an exception is thrown, then the second operand of the GIMPLE\_TRY\_CATCH is checked. The second operand may have the following forms:

- 1. A sequence of statements to execute. When an exception occurs, these statements are executed, and then the exception is rethrown.
- 2. A sequence of GIMPLE\_CATCH statements. Each GIMPLE\_CATCH has a list of applicable exception types and handler code. If the thrown exception matches one of the caught types, the associated handler code is executed. If the handler code falls off the bottom, execution continues after the original GIMPLE\_TRY\_CATCH.

3. A GIMPLE\_EH\_FILTER statement. This has a list of permitted exception types, and code to handle a match failure. If the thrown exception does not match one of the allowed types, the associated match failure code is executed. If the thrown exception does match, it continues unwinding the stack looking for the next handler.

Currently throwing an exception is not directly represented in GIMPLE, since it is implemented by calling a function. At some point in the future we will want to add some way to express that the call will throw an exception of a known type.

Just before running the optimizers, the compiler lowers the high-level EH constructs above into a set of 'goto's, magic labels, and EH regions. Continuing to unwind at the end of a cleanup is represented with a GIMPLE\_RESX.

# 12.5 Temporaries

When gimplification encounters a subexpression that is too complex, it creates a new temporary variable to hold the value of the subexpression, and adds a new statement to initialize it before the current statement. These special temporaries are known as 'expression temporaries', and are allocated using get\_formal\_tmp\_var. The compiler tries to always evaluate identical expressions into the same temporary, to simplify elimination of redundant calculations.

We can only use expression temporaries when we know that it will not be reevaluated before its value is used, and that it will not be otherwise modified<sup>1</sup>. Other temporaries can be allocated using get\_initialized\_tmp\_var or create\_tmp\_var.

Currently, an expression like a=b+5 is not reduced any further. We tried converting it to something like

```
T1 = b + 5;
a = T1;
```

but this bloated the representation for minimal benefit. However, a variable which must live in memory cannot appear in an expression; its value is explicitly loaded into a temporary first. Similarly, storing the value of an expression to a memory variable goes through a temporary.

# 12.6 Operands

In general, expressions in GIMPLE consist of an operation and the appropriate number of simple operands; these operands must either be a GIMPLE rvalue (is\_gimple\_val), i.e. a constant or a register variable. More complex operands are factored out into temporaries, so that

```
a = b + c + d
becomes
T1 = b + c;
a = T1 + d;
```

The same rule holds for arguments to a GIMPLE\_CALL.

The target of an assignment is usually a variable, but can also be a MEM\_REF or a compound lvalue as described below.

<sup>&</sup>lt;sup>1</sup> These restrictions are derived from those in Morgan 4.8.

# 12.6.1 Compound Expressions

The left-hand side of a C comma expression is simply moved into a separate statement.

# 12.6.2 Compound Lvalues

Currently compound l<br/>values involving array and structure field references are not broken down; an expression like a.b[2] = 42 is not reduced any further (though complex array subscripts are). This restriction is a workaround for limitations in later optimizers; if we were to convert this to

```
T1 = &a.b;
T1[2] = 42;
```

alias analysis would not remember that the reference to T1[2] came by way of a.b, so it would think that the assignment could alias another member of a; this broke struct-alias-1.c. Future optimizer improvements may make this limitation unnecessary.

# 12.6.3 Conditional Expressions

A C?: expression is converted into an if statement with each branch assigning to the same temporary. So,

```
a = b ? c : d;
becomes

if (b == 1)
   T1 = c;
else
   T1 = d;
a = T1;
```

The GIMPLE level if-conversion pass re-introduces ?: expression, if appropriate. It is used to vectorize loops with conditions using vector conditional operations.

Note that in GIMPLE, if statements are represented using GIMPLE\_COND, as described below.

# 12.6.4 Logical Operators

Except when they appear in the condition operand of a GIMPLE\_COND, logical 'and' and 'or' operators are simplified as follows: a = b && c becomes

```
T1 = (bool)b;
if (T1 == true)
T1 = (bool)c;
a = T1:
```

Note that T1 in this example cannot be an expression temporary, because it has two different assignments.

# 12.6.5 Manipulating operands

All gimple operands are of type tree. But only certain types of trees are allowed to be used as operand tuples. Basic validation is controlled by the function get\_gimple\_rhs\_class, which given a tree code, returns an enum with the following values of type enum gimple\_rhs\_class

- GIMPLE\_INVALID\_RHS The tree cannot be used as a GIMPLE operand.
- GIMPLE\_TERNARY\_RHS The tree is a valid GIMPLE ternary operation.

- GIMPLE\_BINARY\_RHS The tree is a valid GIMPLE binary operation.
- GIMPLE\_UNARY\_RHS The tree is a valid GIMPLE unary operation.
- GIMPLE\_SINGLE\_RHS The tree is a single object, that cannot be split into simpler operands (for instance, SSA\_NAME, VAR\_DECL, COMPONENT\_REF, etc).

This operand class also acts as an escape hatch for tree nodes that may be flattened out into the operand vector, but would need more than two slots on the RHS. For instance, a COND\_EXPR expression of the form  $(a \ op \ b)$ ? x:y could be flattened out on the operand vector using 4 slots, but it would also require additional processing to distinguish  $c = a \ op \ b$  from  $c = a \ op \ b$ ? x:y. Something similar occurs with ASSERT\_EXPR. In time, these special case tree expressions should be flattened into the operand vector.

For tree nodes in the categories GIMPLE\_TERNARY\_RHS, GIMPLE\_BINARY\_RHS and GIMPLE\_UNARY\_RHS, they cannot be stored inside tuples directly. They first need to be flattened and separated into individual components. For instance, given the GENERIC expression

```
a = b + c
```

its tree representation is:

```
MODIFY_EXPR <VAR_DECL <a>, PLUS_EXPR <VAR_DECL <b>, VAR_DECL <c>>>
```

In this case, the GIMPLE form for this statement is logically identical to its GENERIC form but in GIMPLE, the PLUS\_EXPR on the RHS of the assignment is not represented as a tree, instead the two operands are taken out of the PLUS\_EXPR sub-tree and flattened into the GIMPLE tuple as follows:

```
GIMPLE_ASSIGN <PLUS_EXPR, VAR_DECL <a>, VAR_DECL <b>, VAR_DECL <c>>
```

# 12.6.6 Operand vector allocation

The operand vector is stored at the bottom of the three tuple structures that accept operands. This means, that depending on the code of a given statement, its operand vector will be at different offsets from the base of the structure. To access tuple operands use the following accessors

```
unsigned gimple_num_ops (gimple g)
```

[GIMPLE function]

Returns the number of operands in statement G.

```
tree gimple_op (gimple g, unsigned i)
```

[GIMPLE function]

Returns operand I from statement G.

```
tree * gimple_ops (gimple g)
```

[GIMPLE function]

Returns a pointer into the operand vector for statement G. This is computed using an internal table called gimple\_ops\_offset\_[]. This table is indexed by the gimple code of G.

When the compiler is built, this table is filled-in using the sizes of the structures used by each statement code defined in gimple.def. Since the operand vector is at the bottom of the structure, for a gimple code C the offset is computed as size of (struct-of C) - size of (tree).

This mechanism adds one memory indirection to every access when using gimple\_op(), if this becomes a bottleneck, a pass can choose to memoize the result from gimple\_ops() and use that to access the operands.

# 12.6.7 Operand validation

When adding a new operand to a gimple statement, the operand will be validated according to what each tuple accepts in its operand vector. These predicates are called by the <code>gimple\_name\_set\_...()</code>. Each tuple will use one of the following predicates (Note, this list is not exhaustive):

bool is\_gimple\_val (tree t)

[GIMPLE function]

Returns true if t is a "GIMPLE value", which are all the non-addressable stack variables (variables for which is\_gimple\_reg returns true) and constants (expressions for which is\_gimple\_min\_invariant returns true).

bool is\_gimple\_addressable (tree t)

[GIMPLE function]

Returns true if t is a symbol or memory reference whose address can be taken.

 $\verb|bool is_gimple_asm_val| (tree t)$ 

[GIMPLE function]

Similar to is\_gimple\_val but it also accepts hard registers.

bool is\_gimple\_call\_addr (tree t)

[GIMPLE function]

Return true if t is a valid expression to use as the function called by a GIMPLE\_CALL.

bool is\_gimple\_mem\_ref\_addr (tree t)

[GIMPLE function]

Return true if t is a valid expression to use as first operand of a MEM\_REF expression.

bool is\_gimple\_constant (tree t)

[GIMPLE function]

Return true if t is a valid gimple constant.

bool is\_gimple\_min\_invariant (tree t)

[GIMPLE function]

Return true if t is a valid minimal invariant. This is different from constants, in that the specific value of t may not be known at compile time, but it is known that it doesn't change (e.g., the address of a function local variable).

bool is\_gimple\_ip\_invariant (tree t)

[GIMPLE function]

Return true if t is an interprocedural invariant. This means that t is a valid invariant in all functions (e.g. it can be an address of a global variable but not of a local one).

bool is\_gimple\_ip\_invariant\_address (tree t)

[GIMPLE function]

Return true if t is an ADDR\_EXPR that does not change once the program is running (and which is valid in all functions).

#### 12.6.8 Statement validation

bool is\_gimple\_assign (gimple g)

[GIMPLE function]

Return true if the code of g is GIMPLE\_ASSIGN.

bool is\_gimple\_call (gimple g)

[GIMPLE function]

Return true if the code of g is GIMPLE\_CALL.

bool is\_gimple\_debug (gimple g)

[GIMPLE function]

Return true if the code of g is GIMPLE\_DEBUG.

bool gimple\_assign\_cast\_p (const\_gimple g)

[GIMPLE function]

Return true if g is a GIMPLE\_ASSIGN that performs a type cast operation.

bool gimple\_debug\_bind\_p (gimple g)

[GIMPLE function]

Return true if g is a GIMPLE\_DEBUG that binds the value of an expression to a variable.

bool is\_gimple\_omp (gimple g)

[GIMPLE function]

Return true if g is any of the OpenMP codes.

gimple\_debug\_begin\_stmt\_p (gimple g)

[GIMPLE function]

Return true if g is a GIMPLE\_DEBUG that marks the beginning of a source statement.

gimple\_debug\_inline\_entry\_p (gimple g)

[GIMPLE function]

Return true if g is a GIMPLE\_DEBUG that marks the entry point of an inlined function.

gimple\_debug\_nonbind\_marker\_p (gimple g)

[GIMPLE function]

Return true if g is a GIMPLE\_DEBUG that marks a program location, without any variable binding.

# 12.7 Manipulating GIMPLE statements

This section documents all the functions available to handle each of the GIMPLE instructions.

# 12.7.1 Common accessors

The following are common accessors for gimple statements.

enum gimple\_code gimple\_code (gimple g)

[GIMPLE function]

Return the code for statement G.

basic\_block gimple\_bb (gimple g)

[GIMPLE function]

Return the basic block to which statement G belongs to.

tree gimple\_block (gimple g)

[GIMPLE function]

Return the lexical scope block holding statement G.

tree gimple\_expr\_type (gimple stmt)

[GIMPLE function]

Return the type of the main expression computed by STMT. Return void\_type\_node if STMT computes nothing. This will only return something meaningful for GIMPLE\_ASSIGN, GIMPLE\_COND and GIMPLE\_CALL. For all other tuple codes, it will return void\_type\_node.

enum tree\_code gimple\_expr\_code (gimple stmt)

[GIMPLE function]

Return the tree code for the expression computed by STMT. This is only meaningful for GIMPLE\_CALL, GIMPLE\_ASSIGN and GIMPLE\_COND. If STMT is GIMPLE\_CALL, it will return CALL\_EXPR. For GIMPLE\_COND, it returns the code of the comparison predicate. For GIMPLE\_ASSIGN it returns the code of the operation performed by the RHS of the assignment.

 $\verb"void gimple_set_block" (gimple g, tree block")$ 

[GIMPLE function]

Set the lexical scope block of G to BLOCK.

location\_t gimple\_locus (gimple g)

[GIMPLE function]

Return locus information for statement G.

void gimple\_set\_locus (gimple g, location\_t locus) [GIMPLE function] Set locus information for statement G. [GIMPLE function] bool gimple\_locus\_empty\_p (gimple g) Return true if G does not have locus information. bool gimple\_no\_warning\_p (gimple stmt) [GIMPLE function] Return true if no warnings should be emitted for statement STMT. void gimple\_set\_visited (gimple stmt, bool visited\_p) [GIMPLE function] Set the visited status on statement STMT to VISITED\_P. bool gimple\_visited\_p (gimple stmt) [GIMPLE function] Return the visited status on statement STMT. void gimple\_set\_plf (gimple stmt, enum plf\_mask plf, bool [GIMPLE function] val\_p) Set pass local flag PLF on statement STMT to VAL\_P. unsigned int gimple\_plf (gimple stmt, enum plf\_mask plf) [GIMPLE function] Return the value of pass local flag PLF on statement STMT. bool gimple\_has\_ops (gimple g) [GIMPLE function] Return true if statement G has register or memory operands. bool gimple\_has\_mem\_ops (gimple g) [GIMPLE function] Return true if statement G has memory operands. unsigned gimple\_num\_ops (gimple g) [GIMPLE function] Return the number of operands for statement G. tree \* gimple\_ops (gimple g) [GIMPLE function] Return the array of operands for statement G. tree gimple\_op (gimple g, unsigned i) [GIMPLE function] Return operand I for statement G. [GIMPLE function] tree \* gimple\_op\_ptr (gimple g, unsigned i) Return a pointer to operand I for statement G. void gimple\_set\_op (gimple g, unsigned i, tree op) [GIMPLE function] Set operand I of statement G to OP. bitmap gimple\_addresses\_taken (gimple stmt) [GIMPLE function] Return the set of symbols that have had their address taken by STMT. struct def\_optype\_d \* gimple\_def\_ops (gimple g) [GIMPLE function] Return the set of DEF operands for statement G. void gimple\_set\_def\_ops (gimple g, struct def\_optype\_d [GIMPLE function] \*def)

Set DEF to be the set of DEF operands for statement G.

struct use\_optype\_d \* gimple\_use\_ops (gimple g) [GIMPLE function]
Return the set of USE operands for statement G.

Set  ${\tt USE}$  to be the set of  ${\tt USE}$  operands for statement  ${\tt G}.$ 

Return the set of VUSE operands for statement G. [GIMPLE function]

void gimple\_set\_vuse\_ops (gimple g, struct voptype\_d \*ops) [GIMPLE function] Set OPS to be the set of VUSE operands for statement G.

struct voptype\_d \* gimple\_vdef\_ops (gimple g) [GIMPLE function] Return the set of VDEF operands for statement G.

void gimple\_set\_vdef\_ops (gimple g, struct voptype\_d \*ops) [GIMPLE function] Set OPS to be the set of VDEF operands for statement G.

bitmap gimple\_loaded\_syms (gimple g) [GIMPLE function] Return the set of symbols loaded by statement G. Each element of the set is the DECL\_UID of the corresponding symbol.

bitmap gimple\_stored\_syms (gimple g) [GIMPLE function]
Return the set of symbols stored by statement G. Each element of the set is the DECL\_UID of the corresponding symbol.

bool gimple\_modified\_p (gimple g) [GIMPLE function] Return true if statement G has operands and the modified field has been set.

bool gimple\_has\_volatile\_ops (gimple stmt) [GIMPLE function] Return true if statement STMT contains volatile operands.

void gimple\_set\_has\_volatile\_ops (gimple stmt, bool volatilep)

Return true if statement STMT contains volatile operands.

[GIMPLE function]

void update\_stmt (gimple s) [GIMPLE function]

Mark statement S as modified, and update it.

void update\_stmt\_if\_modified (gimple s) [GIMPLE function]
Update statement S if it has been marked modified.

gimple gimple\_copy (gimple stmt) [GIMPLE function]

Return a deep copy of statement STMT.

# 12.8 Tuple specific accessors

# **12.8.1** GIMPLE\_ASM

- - Build a GIMPLE\_ASM statement. This statement is used for building in-line assembly constructs. STRING is the assembly code. INPUTS, OUTPUTS, CLOBBERS and LABELS are the inputs, outputs, clobbered registers and labels.
- unsigned gimple\_asm\_ninputs ( $const\ gasm\ *g$ ) [GIMPLE function] Return the number of input operands for GIMPLE\_ASM G.
- unsigned gimple\_asm\_noutputs (const gasm \*g) [GIMPLE function]
  Return the number of output operands for GIMPLE\_ASM G.
- unsigned gimple\_asm\_nclobbers (const gasm \*g) [GIMPLE function]
  Return the number of clobber operands for GIMPLE\_ASM G.
- tree gimple\_asm\_input\_op (const gasm \*g, unsigned index) [GIMPLE function] Return input operand INDEX of GIMPLE\_ASM G.
- tree gimple\_asm\_output\_op (const gasm \*g, unsigned index) [GIMPLE function]
  Return output operand INDEX of GIMPLE\_ASM G.
- tree gimple\_asm\_clobber\_op (const gasm \*g, unsigned index) [GIMPLE function] Return clobber operand INDEX of GIMPLE\_ASM G.

- bool gimple\_asm\_volatile\_p (const gasm \*g) [GIMPLE function] Return true if G is an asm statement marked volatile.

# 12.8.2 GIMPLE\_ASSIGN

gassign \*gimple\_build\_assign (tree lhs, tree rhs)
[GIMPLE function]

Build a GIMPLE\_ASSIGN statement. The left-hand side is an Ivalue passed in lhs. The right-hand side can be either a unary or binary tree expression. The expression tree rhs will be flattened and its operands assigned to the corresponding operand slots in the new statement. This function is useful when you already have a tree expression that you want to convert into a tuple. However, try to avoid building expression trees for the sole purpose of calling this function. If you already have the operands in separate trees, it is better to use gimple\_build\_assign with enum tree\_code argument and separate arguments for each operand.

gassign \*gimple\_build\_assign (tree lhs, enum tree\_code [GIMPLE function] subcode, tree op1, tree op2, tree op3)

This function is similar to two operand gimple\_build\_assign, but is used to build a GIMPLE\_ASSIGN statement when the operands of the right-hand side of the assignment are already split into different operands.

The left-hand side is an lvalue passed in lhs. Subcode is the tree\_code for the right-hand side of the assignment. Op1, op2 and op3 are the operands.

Like the above 5 operand gimple\_build\_assign, but with the last argument NULL - this overload should not be used for GIMPLE\_TERNARY\_RHS assignments.

Like the above 4 operand gimple\_build\_assign, but with the last argument NULL - this overload should be used only for GIMPLE\_UNARY\_RHS and GIMPLE\_SINGLE\_RHS assignments.

gimple gimplify\_assign (tree dst, tree src, gimple\_seq [GIMPLE function] \*seq\_p)

Build a new GIMPLE\_ASSIGN tuple and append it to the end of \*SEQ\_P.

DST/SRC are the destination and source respectively. You can pass ungimplified trees in DST or SRC, in which case they will be converted to a gimple operand if necessary.

This function returns the newly created GIMPLE\_ASSIGN tuple.

enum tree\_code gimple\_assign\_rhs\_code ( $gimple\ g$ ) [GIMPLE function] Return the code of the expression computed on the RHS of assignment statement G.

enum gimple\_rhs\_class gimple\_assign\_rhs\_class [GIMPLE function] (gimple g)

Return the gimple rhs class of the code for the expression computed on the rhs of assignment statement G. This will never return GIMPLE\_INVALID\_RHS.

tree gimple\_assign\_lhs (gimple g) [GIMPLE function] Return the LHS of assignment statement G.

- tree \* gimple\_assign\_lhs\_ptr (gimple g) [GIMPLE function]
  Return a pointer to the LHS of assignment statement G.
- tree gimple\_assign\_rhs1 ( $gimple\ g$ ) [GIMPLE function] Return the first operand on the RHS of assignment statement G.
- $\label{eq:continuous} \begin{tabular}{ll} tree * gimple_assign_rhs1\_ptr (gimple g) & [GIMPLE function] \\ Return the address of the first operand on the RHS of assignment statement G. \\ \end{tabular}$
- tree gimple\_assign\_rhs2 ( $gimple\ g$ ) [GIMPLE function] Return the second operand on the RHS of assignment statement G.
- tree \* gimple\_assign\_rhs2\_ptr (gimple g) [GIMPLE function] Return the address of the second operand on the RHS of assignment statement G.
- tree gimple\_assign\_rhs3 ( $gimple\ g$ ) [GIMPLE function] Return the third operand on the RHS of assignment statement G.
- $\label{eq:continuous} \begin{tabular}{ll} tree * gimple_assign_rhs3\_ptr (gimple g) & [GIMPLE function] \\ Return the address of the third operand on the RHS of assignment statement $G$. \\ \end{tabular}$
- void gimple\_assign\_set\_lhs (gimple g, tree lhs) [GIMPLE function] Set LHS to be the LHS operand of assignment statement G.
- void gimple\_assign\_set\_rhs1 (gimple g, tree rhs) [GIMPLE function] Set RHS to be the first operand on the RHS of assignment statement G.
- void gimple\_assign\_set\_rhs2 (gimple g, tree rhs) [GIMPLE function] Set RHS to be the second operand on the RHS of assignment statement G.
- void gimple\_assign\_set\_rhs3 (gimple g, tree rhs) [GIMPLE function] Set RHS to be the third operand on the RHS of assignment statement G.
- bool gimple\_assign\_cast\_p (const\_gimple s) [GIMPLE function] Return true if S is a type-cast assignment.

# 12.8.3 GIMPLE\_BIND

- gbind \*gimple\_build\_bind (tree vars, gimple\_seq body) [GIMPLE function] Build a GIMPLE\_BIND statement with a list of variables in VARS and a body of statements in sequence BODY.
- tree gimple\_bind\_vars (const gbind \*g) [GIMPLE function] Return the variables declared in the GIMPLE\_BIND statement G.
- void gimple\_bind\_set\_vars (gbind \*g, tree vars) [GIMPLE function] Set VARS to be the set of variables declared in the GIMPLE\_BIND statement G.
- void gimple\_bind\_append\_vars (gbind \*g, tree vars) [GIMPLE function] Append VARS to the set of variables declared in the GIMPLE\_BIND statement G.
- gimple\_seq gimple\_bind\_body (gbind \*g) [GIMPLE function] Return the GIMPLE sequence contained in the GIMPLE\_BIND statement G.

- void gimple\_bind\_set\_body (gbind \*g, gimple\_seq seq) [GIMPLE function] Set SEQ to be sequence contained in the GIMPLE\_BIND statement G.
- void gimple\_bind\_add\_stmt (gbind \*gs, gimple stmt) [GIMPLE function]
  Append a statement to the end of a GIMPLE\_BIND's body.
- void gimple\_bind\_add\_seq (gbind \*gs, gimple\_seq seq) [GIMPLE function] Append a sequence of statements to the end of a GIMPLE\_BIND's body.
- tree gimple\_bind\_block (const gbind \*g) [GIMPLE function] Return the TREE\_BLOCK node associated with GIMPLE\_BIND statement G. This is analogous to the BIND\_EXPR\_BLOCK field in trees.
- void gimple\_bind\_set\_block (gbind \*g, tree block) [GIMPLE function]
  Set BLOCK to be the TREE\_BLOCK node associated with GIMPLE\_BIND statement G.

#### 12.8.4 GIMPLE CALL

- gcall \*gimple\_build\_call (tree fn, unsigned nargs, ...) [GIMPLE function] Build a GIMPLE\_CALL statement to function FN. The argument FN must be either a FUNCTION\_DECL or a gimple call address as determined by is\_gimple\_call\_addr. NARGS are the number of arguments. The rest of the arguments follow the argument NARGS, and must be trees that are valid as rvalues in gimple (i.e., each operand is validated with is\_gimple\_operand).
- gcall \*gimple\_build\_call\_from\_tree (tree call\_expr, tree [GIMPLE function] fnptrtype)

Build a GIMPLE\_CALL from a CALL\_EXPR node. The arguments and the function are taken from the expression directly. The type of the GIMPLE\_CALL is set from the second parameter passed by a caller. This routine assumes that call\_expr is already in GIMPLE form. That is, its operands are GIMPLE values and the function call needs no further simplification. All the call flags in call\_expr are copied over to the new GIMPLE\_CALL.

- gcall \*gimple\_build\_call\_vec (tree fn, vec<tree> args) [GIMPLE function] Identical to gimple\_build\_call but the arguments are stored in a vec<tree>.
- tree gimple\_call\_lhs (gimple g) [GIMPLE function]
  Return the LHS of call statement G.
- tree \* gimple\_call\_lhs\_ptr (gimple g) [GIMPLE function]
  Return a pointer to the LHS of call statement G.
- void gimple\_call\_set\_lhs (gimple g, tree lhs) [GIMPLE function] Set LHS to be the LHS operand of call statement G.
- tree gimple\_call\_fn (gimple g) [GIMPLE function] Return the tree node representing the function called by call statement G.
- void gimple\_call\_set\_fn (gcall \*g, tree fn) [GIMPLE function]

  Set FN to be the function called by call statement G. This has to be a gimple value specifying the address of the called function.

tree gimple\_call\_fndecl ( $gimple\ g$ ) [GIMPLE function] If a given GIMPLE\_CALL's callee is a FUNCTION\_DECL, return it. Otherwise return NULL. This function is analogous to get\_callee\_fndecl in GENERIC.

tree gimple\_call\_set\_fndecl (gimple g, tree fndecl) [GIMPLE function] Set the called function to FNDECL.

tree gimple\_call\_return\_type (const gcall \*g) [GIMPLE function]
Return the type returned by call statement G.

tree gimple\_call\_chain (gimple g) [GIMPLE function]

Return the static chain for call statement G.

void gimple\_call\_set\_chain (gcall \*g, tree chain) [GIMPLE function] Set CHAIN to be the static chain for call statement G.

unsigned gimple\_call\_num\_args (gimple g) [GIMPLE function] Return the number of arguments used by call statement G.

tree gimple\_call\_arg (gimple g, unsigned index) [GIMPLE function]
Return the argument at position INDEX for call statement G. The first argument is 0.

tree \* gimple\_call\_arg\_ptr (gimple g, unsigned index) [GIMPLE function]
Return a pointer to the argument at position INDEX for call statement G.

void gimple\_call\_set\_tail (gcall \*s) [GIMPLE function] Mark call statement S as being a tail call (i.e., a call just before the exit of a function). These calls are candidate for tail call optimization.

bool gimple\_call\_tail\_p (gcall \*s) [GIMPLE function]

Return true if GIMPLE\_CALL S is marked as a tail call.

bool gimple\_call\_noreturn\_p (gimple s) [GIMPLE function]

Return true if S is a noreturn call.

gimple gimple\_call\_copy\_skip\_args (gcall \*stmt, bitmap [GIMPLE function] args\_to\_skip)

Build a GIMPLE\_CALL identical to STMT but skipping the arguments in the positions marked by the set ARGS\_TO\_SKIP.

# 12.8.5 GIMPLE\_CATCH

gcatch \*gimple\_build\_catch (tree types, gimple\_seq handler) [GIMPLE function] Build a GIMPLE\_CATCH statement. TYPES are the tree types this catch handles. HANDLER is a sequence of statements with the code for the handler.

tree gimple\_catch\_types (const gcatch \*g) [GIMPLE function]
Return the types handled by GIMPLE\_CATCH statement G.

tree \* gimple\_catch\_types\_ptr (gcatch \*g) [GIMPLE function] Return a pointer to the types handled by GIMPLE\_CATCH statement G.

gimple\_seq gimple\_catch\_handler (gcatch \*g) [GIMPLE function]
Return the GIMPLE sequence representing the body of the handler of GIMPLE\_CATCH statement G.

### 12.8.6 GIMPLE COND

gcond \*gimple\_build\_cond ( enum tree\_code pred\_code, tree [GIMPLE function] lhs, tree rhs, tree t\_label, tree f\_label)

Build a GIMPLE\_COND statement. A GIMPLE\_COND statement compares LHS and RHS and if the condition in PRED\_CODE is true, jump to the label in t\_label, otherwise jump to the label in f\_label. PRED\_CODE are relational operator tree codes like EQ\_EXPR, LT\_EXPR, LE\_EXPR, NE\_EXPR, etc.

gcond \*gimple\_build\_cond\_from\_tree (tree cond, tree [GIMPLE function]

t\_label, tree f\_label)

Build a GIMPLE GOND statement from the conditional supposion tree GOND T LABEL

Build a GIMPLE\_COND statement from the conditional expression tree COND. T\_LABEL and F\_LABEL are as in gimple\_build\_cond.

- enum tree\_code gimple\_cond\_code (gimple g) [GIMPLE function] Return the code of the predicate computed by conditional statement G.
- void gimple\_cond\_set\_code (gcond \*g, enum tree\_code code) [GIMPLE function] Set CODE to be the predicate code for the conditional statement G.
- tree gimple\_cond\_lhs (gimple g) [GIMPLE function] Return the LHS of the predicate computed by conditional statement G.
- void gimple\_cond\_set\_lhs (gcond \*g, tree lhs) [GIMPLE function] Set LHS to be the LHS operand of the predicate computed by conditional statement G.
- tree gimple\_cond\_rhs (gimple g) [GIMPLE function] Return the RHS operand of the predicate computed by conditional G.
- void gimple\_cond\_set\_rhs (gcond \*g, tree rhs) [GIMPLE function]
  Set RHS to be the RHS operand of the predicate computed by conditional statement G.
- tree gimple\_cond\_true\_label (const gcond \*g) [GIMPLE function]

  Return the label used by conditional statement G when its predicate evaluates to true.
- void gimple\_cond\_set\_true\_label (gcond \*g, tree label) [GIMPLE function] Set LABEL to be the label used by conditional statement G when its predicate evaluates to true.

- void gimple\_cond\_set\_false\_label (gcond \*g, tree label) [GIMPLE function] Set LABEL to be the label used by conditional statement G when its predicate evaluates to false.
- tree gimple\_cond\_false\_label (const gcond \*g) [GIMPLE function] Return the label used by conditional statement G when its predicate evaluates to false.
- void gimple\_cond\_make\_false (gcond \*g) [GIMPLE function] Set the conditional COND\_STMT to be of the form 'if (1 == 0)'.
- void gimple\_cond\_make\_true (gcond \*g) [GIMPLE function] Set the conditional COND\_STMT to be of the form 'if (1 == 1)'.

# 12.8.7 GIMPLE\_DEBUG

Build a GIMPLE\_DEBUG statement with GIMPLE\_DEBUG\_BIND subcode. The effect of this statement is to tell debug information generation machinery that the value of user variable var is given by value at that point, and to remain with that value until var runs out of scope, a dynamically-subsequent debug bind statement overrides the binding, or conflicting values reach a control flow merge point. Even if components of the value expression change afterwards, the variable is supposed to retain the same value, though not necessarily the same location.

It is expected that var be most often a tree for automatic user variables (VAR\_DECL or PARM\_DECL) that satisfy the requirements for gimple registers, but it may also be a tree for a scalarized component of a user variable (ARRAY\_REF, COMPONENT\_REF), or a debug temporary (DEBUG\_EXPR\_DECL).

As for value, it can be an arbitrary tree expression, but it is recommended that it be in a suitable form for a gimple assignment RHS. It is not expected that user variables that could appear as var ever appear in value, because in the latter we'd have their SSA\_NAMEs instead, but even if they were not in SSA form, user variables appearing in value are to be regarded as part of the executable code space, whereas those in var are to be regarded as part of the source code space. There is no way to refer to the value bound to a user variable within a value expression.

If value is GIMPLE\_DEBUG\_BIND\_NOVALUE, debug information generation machinery is informed that the variable var is unbound, i.e., that its value is indeterminate, which sometimes means it is really unavailable, and other times that the compiler could not keep track of it.

Block and location information for the newly-created stmt are taken from stmt, if given.

tree gimple\_debug\_bind\_get\_var (gimple stmt) [GIMPLE function]
Return the user variable var that is bound at stmt.

tree gimple\_debug\_bind\_get\_value (gimple stmt) [GIMPLE function]
Return the value expression that is bound to a user variable at stmt.

- tree \* gimple\_debug\_bind\_get\_value\_ptr (gimple stmt) [GIMPLE function] Return a pointer to the value expression that is bound to a user variable at stmt.
- void gimple\_debug\_bind\_set\_var (gimple stmt, tree var) [GIMPLE function]

  Modify the user variable bound at stmt to var.
- void gimple\_debug\_bind\_set\_value (gimple stmt, tree var) [GIMPLE function] Modify the value bound to the user variable bound at stmt to value.
- void gimple\_debug\_bind\_reset\_value (gimple stmt) [GIMPLE function] Modify the value bound to the user variable bound at stmt so that the variable becomes unbound.
- bool gimple\_debug\_bind\_has\_value\_p (gimple stmt) [GIMPLE function] Return TRUE if stmt binds a user variable to a value, and FALSE if it unbinds the variable.

Build a GIMPLE\_DEBUG statement with GIMPLE\_DEBUG\_BEGIN\_STMT subcode. The effect of this statement is to tell debug information generation machinery that the user statement at the given location and block starts at the point at which the statement is inserted. The intent is that side effects (e.g. variable bindings) of all prior user statements are observable, and that none of the side effects of subsequent user statements are.

Build a GIMPLE\_DEBUG statement with GIMPLE\_DEBUG\_INLINE\_ENTRY subcode. The effect of this statement is to tell debug information generation machinery that a function call at location underwent inline substitution, that block is the enclosing lexical block created for the substitution, and that at the point of the program in which the stmt is inserted, all parameters for the inlined function are bound to the respective arguments, and none of the side effects of its stmts are observable.

# 12.8.8 GIMPLE\_EH\_FILTER

Build a GIMPLE\_EH\_FILTER statement. TYPES are the filter's types. FAILURE is a sequence with the filter's failure action.

- tree gimple\_eh\_filter\_types (gimple g) [GIMPLE function]
  Return the types handled by GIMPLE\_EH\_FILTER statement G.
- $\label{eq:continuous} \begin{tabular}{ll} tree * gimple_eh_filter_types_ptr (gimple g) & [GIMPLE function] \\ Return a pointer to the types handled by GIMPLE_EH_FILTER statement G. \\ \end{tabular}$
- gimple\_seq gimple\_eh\_filter\_failure (gimple g) [GIMPLE function] Return the sequence of statement to execute when GIMPLE\_EH\_FILTER statement fails.

void gimple\_eh\_filter\_set\_types (geh\_filter \*g, tree types) [GIMPLE function] Set TYPES to be the set of types handled by GIMPLE\_EH\_FILTER G.

Set FAILURE to be the sequence of statements to execute on failure for GIMPLE\_EH\_ FILTER G.

# 12.8.9 GIMPLE\_LABEL

glabel \*gimple\_build\_label (tree label) [GIMPLE function] Build a GIMPLE\_LABEL statement with corresponding to the tree label, LABEL.

tree gimple\_label\_label ( $const\ glabel\ ^*g$ ) [GIMPLE function] Return the LABEL\_DECL node used by GIMPLE\_LABEL statement G.

void gimple\_label\_set\_label (glabel \*g, tree label) [GIMPLE function] Set LABEL to be the LABEL\_DECL node used by GIMPLE\_LABEL statement G.

# 12.8.10 GIMPLE\_GOTO

ggoto \*gimple\_build\_goto (tree dest) [GIMPLE function]
Build a GIMPLE\_GOTO statement to label DEST.

tree gimple\_goto\_dest (gimple g) [GIMPLE function]
Return the destination of the unconditional jump G.

void gimple\_goto\_set\_dest (ggoto \*g, tree dest) [GIMPLE function] Set DEST to be the destination of the unconditional jump G.

# **12.8.11** GIMPLE\_NOP

gimple gimple\_build\_nop (void) [GIMPLE function]

Build a GIMPLE\_NOP statement.

bool gimple\_nop\_p (gimple g) [GIMPLE function]
Returns TRUE if statement G is a GIMPLE\_NOP.

#### 12.8.12 GIMPLE\_OMP\_ATOMIC\_LOAD

Build a GIMPLE\_OMP\_ATOMIC\_LOAD statement. LHS is the left-hand side of the assignment. RHS is the right-hand side of the assignment.

void gimple\_omp\_atomic\_load\_set\_lhs ( gomp\_atomic\_load [GIMPLE function] \*g, tree lhs) Set the LHS of an atomic load. [GIMPLE function] tree gimple\_omp\_atomic\_load\_lhs ( const gomp\_atomic\_load \*g) Get the LHS of an atomic load. void gimple\_omp\_atomic\_load\_set\_rhs (gomp\_atomic\_load [GIMPLE function] \*g, tree rhs) Set the RHS of an atomic set. tree gimple\_omp\_atomic\_load\_rhs ( const [GIMPLE function] gomp\_atomic\_load \*g) Get the RHS of an atomic set. 12.8.13 GIMPLE\_OMP\_ATOMIC\_STORE gomp\_atomic\_store \*gimple\_build\_omp\_atomic\_store ( [GIMPLE function] tree val) Build a GIMPLE\_OMP\_ATOMIC\_STORE statement. VAL is the value to be stored. void gimple\_omp\_atomic\_store\_set\_val ( [GIMPLE function] gomp\_atomic\_store \*g, tree val) Set the value being stored in an atomic store. [GIMPLE function] tree gimple\_omp\_atomic\_store\_val ( const gomp\_atomic\_store \*g) Return the value being stored in an atomic store. 12.8.14 GIMPLE\_OMP\_CONTINUE gomp\_continue \*gimple\_build\_omp\_continue ( tree [GIMPLE function] control\_def, tree control\_use) Build a GIMPLE\_OMP\_CONTINUE statement. CONTROL\_DEF is the definition of the control variable. CONTROL\_USE is the use of the control variable. tree gimple\_omp\_continue\_control\_def ( const [GIMPLE function] gomp\_continue \*s) Return the definition of the control variable on a GIMPLE\_OMP\_CONTINUE in S. tree gimple\_omp\_continue\_control\_def\_ptr ( [GIMPLE function] gomp\_continue \*s) Same as above, but return the pointer. tree gimple\_omp\_continue\_set\_control\_def ( [GIMPLE function]  $gomp\_continue *s$ ) Set the control variable definition for a GIMPLE\_OMP\_CONTINUE statement in S. [GIMPLE function] tree gimple\_omp\_continue\_control\_use ( const gomp\_continue \*s)

Return the use of the control variable on a GIMPLE\_OMP\_CONTINUE in S.

Same as above, but return the pointer.

Set the control variable use for a GIMPLE\_OMP\_CONTINUE statement in S.

# 12.8.15 GIMPLE\_OMP\_CRITICAL

Build a GIMPLE\_OMP\_CRITICAL statement. BODY is the sequence of statements for which only one thread can execute. NAME is an optional identifier for this critical block.

- tree gimple\_omp\_critical\_name ( const gomp\_critical \*g) [GIMPLE function] Return the name associated with OMP\_CRITICAL statement G.
- tree \* gimple\_omp\_critical\_name\_ptr ( gomp\_critical \*g) [GIMPLE function] Return a pointer to the name associated with OMP critical statement G.

# 12.8.16 GIMPLE\_OMP\_FOR

Build a GIMPLE\_OMP\_FOR statement. BODY is sequence of statements inside the for loop. CLAUSES, are any of the loop construct's clauses. PRE\_BODY is the sequence of statements that are loop invariant. INDEX is the index variable. INITIAL is the initial value of INDEX. FINAL is final value of INDEX. OMP\_FOR\_COND is the predicate used to compare INDEX and FINAL. INCR is the increment expression.

- tree gimple\_omp\_for\_clauses ( $gimple\ g$ ) [GIMPLE function] Return the clauses associated with OMP\_FOR G.
- tree \* gimple\_omp\_for\_clauses\_ptr (gimple g) [GIMPLE function]
  Return a pointer to the OMP\_FOR G.
- void gimple\_omp\_for\_set\_clauses (gimple g, tree clauses) [GIMPLE function] Set CLAUSES to be the list of clauses associated with OMP\_FOR G.
- tree gimple\_omp\_for\_index (gimple g) [GIMPLE function]
  Return the index variable for OMP\_FOR G.
- tree \* gimple\_omp\_for\_index\_ptr (gimple g) [GIMPLE function]
  Return a pointer to the index variable for OMP\_FOR G.

void gimple\_omp\_for\_set\_index (gimple g, tree index) [GIMPLE function] Set INDEX to be the index variable for OMP\_FOR G. tree gimple\_omp\_for\_initial (gimple g) [GIMPLE function] Return the initial value for OMP\_FOR G. [GIMPLE function] tree \* gimple\_omp\_for\_initial\_ptr (gimple g) Return a pointer to the initial value for OMP\_FOR G. void gimple\_omp\_for\_set\_initial (gimple g, tree initial) [GIMPLE function] Set INITIAL to be the initial value for OMP\_FOR G. tree gimple\_omp\_for\_final (gimple g) [GIMPLE function] Return the final value for OMP\_FOR G. [GIMPLE function] tree \* gimple\_omp\_for\_final\_ptr (gimple g) turn a pointer to the final value for OMP\_FOR G. void gimple\_omp\_for\_set\_final (gimple g, tree final) [GIMPLE function] Set FINAL to be the final value for OMP\_FOR G. tree gimple\_omp\_for\_incr (gimple g) [GIMPLE function] Return the increment value for OMP\_FOR G. tree \* gimple\_omp\_for\_incr\_ptr (gimple g) [GIMPLE function] Return a pointer to the increment value for OMP\_FOR G. void gimple\_omp\_for\_set\_incr (gimple g, tree incr) [GIMPLE function] Set INCR to be the increment value for OMP\_FOR G. gimple\_seq gimple\_omp\_for\_pre\_body (gimple g) [GIMPLE function] Return the sequence of statements to execute before the OMP\_FOR statement G starts. void gimple\_omp\_for\_set\_pre\_body (gimple g, gimple\_seq [GIMPLE function] pre\_body) Set PRE\_BODY to be the sequence of statements to execute before the OMP\_FOR statement G starts. void gimple\_omp\_for\_set\_cond (gimple g, enum tree\_code [GIMPLE function] Set COND to be the condition code for OMP\_FOR G. enum tree\_code gimple\_omp\_for\_cond (gimple g) [GIMPLE function]

# 12.8.17 GIMPLE\_OMP\_MASTER

gimple gimple\_build\_omp\_master (gimple\_seq body) [GIMPLE function] Build a GIMPLE\_OMP\_MASTER statement. BODY is the sequence of statements to be executed by just the master.

Return the condition code associated with OMP\_FOR G.

# 12.8.18 GIMPLE\_OMP\_ORDERED

gimple gimple\_build\_omp\_ordered (gimple\_seq body) [GIMPLE function] Build a GIMPLE\_OMP\_ORDERED statement.

BODY is the sequence of statements inside a loop that will executed in sequence.

# 12.8.19 GIMPLE\_OMP\_PARALLEL

BODY is sequence of statements which are executed in parallel. CLAUSES, are the OMP parallel construct's clauses. CHILD\_FN is the function created for the parallel threads to execute. DATA\_ARG are the shared data argument(s).

- bool gimple\_omp\_parallel\_combined\_p (gimple g) [GIMPLE function] Return true if OMP parallel statement G has the GF\_OMP\_PARALLEL\_COMBINED flag set.
- void gimple\_omp\_parallel\_set\_combined\_p (gimple g) [GIMPLE function] Set the GF\_OMP\_PARALLEL\_COMBINED field in OMP parallel statement G.
- $\begin{array}{ll} \texttt{gimple\_seq\ gimple\_omp\_body\ } & [\texttt{GIMPLE\ function}] \\ & \texttt{Return\ the\ body\ for\ the\ OMP\ statement\ G.} \end{array}$
- void gimple\_omp\_set\_body (gimple g, gimple\_seq body)
  Set BODY to be the body for the OMP statement G.
  [GIMPLE function]
- tree gimple\_omp\_parallel\_clauses (gimple g) [GIMPLE function]
  Return the clauses associated with OMP\_PARALLEL G.

- Set CLAUSES to be the list of clauses associated with OMP\_PARALLEL G.

Return a pointer to the child function used to hold the body of OMP\_PARALLEL G.

tree gimple\_omp\_parallel\_data\_arg ( const gomp\_parallel [GIMPLE function] \*g)

Return the artificial argument used to send variables and values from the parent to the children threads in OMP\_PARALLEL G.

Return a pointer to the data argument for  ${\tt OMP\_PARALLEL}$  G.

# 12.8.20 GIMPLE\_OMP\_RETURN

gimple gimple\_build\_omp\_return (bool wait\_p) [GIMPLE function] Build a GIMPLE\_OMP\_RETURN statement. WAIT\_P is true if this is a non-waiting return.

void gimple\_omp\_return\_set\_nowait (gimple s) [GIMPLE function] Set the nowait flag on GIMPLE\_OMP\_RETURN statement S.

bool gimple\_omp\_return\_nowait\_p (gimple g) [GIMPLE function] Return true if OMP return statement G has the GF\_OMP\_RETURN\_NOWAIT flag set.

# 12.8.21 GIMPLE\_OMP\_SECTION

gimple gimple\_build\_omp\_section (gimple\_seq body) [GIMPLE function] Build a GIMPLE\_OMP\_SECTION statement for a sections statement.

BODY is the sequence of statements in the section.

bool gimple\_omp\_section\_last\_p (gimple g) [GIMPLE function] Return true if OMP section statement G has the GF\_OMP\_SECTION\_LAST flag set.

 $\begin{tabular}{ll} \begin{tabular}{ll} \beg$ 

# 12.8.22 GIMPLE\_OMP\_SECTIONS

Build a GIMPLE\_OMP\_SECTIONS statement. BODY is a sequence of section statements. CLAUSES are any of the OMP sections construct's clauses: private, firstprivate, lastprivate, reduction, and nowait.

gimple gimple\_build\_omp\_sections\_switch (void) [GIMPLE function]
Build a GIMPLE\_OMP\_SECTIONS\_SWITCH statement.

tree gimple\_omp\_sections\_control (gimple g) [GIMPLE function] Return the control variable associated with the GIMPLE\_OMP\_SECTIONS in G.

- tree \* gimple\_omp\_sections\_control\_ptr (gimple g) [GIMPLE function] Return a pointer to the clauses associated with the GIMPLE\_OMP\_SECTIONS in G.
- $\begin{tabular}{ll} {\tt void gimple\_omp\_sections\_set\_control} & (gimple\ g,\ tree \\ & control) \end{tabular} \begin{tabular}{ll} {\tt GIMPLE} & (gimple\ g,\ tree \\ & control) \end{tabular}$

Set  ${\tt CONTROL}$  to be the set of clauses associated with the  ${\tt GIMPLE\_OMP\_SECTIONS}$  in G.

- tree gimple\_omp\_sections\_clauses (gimple g) [GIMPLE function] Return the clauses associated with OMP\_SECTIONS G.
- tree \* gimple\_omp\_sections\_clauses\_ptr (gimple g) [GIMPLE function] Return a pointer to the clauses associated with OMP\_SECTIONS G.

# 12.8.23 GIMPLE\_OMP\_SINGLE

- tree gimple\_omp\_single\_clauses ( $gimple\ g$ ) [GIMPLE function] Return the clauses associated with OMP\_SINGLE G.
- $\label{tree * gimple_omp_single_clauses_ptr (gimple g)} \qquad \qquad [GIMPLE \ function] \\ \text{Return a pointer to the clauses associated with OMP_SINGLE G.}$

# 12.8.24 GIMPLE\_PHI

- unsigned gimple\_phi\_capacity (gimple g) [GIMPLE function] Return the maximum number of arguments supported by GIMPLE\_PHI G.
- unsigned gimple\_phi\_num\_args (gimple g) [GIMPLE function] Return the number of arguments in GIMPLE\_PHI G. This must always be exactly the number of incoming edges for the basic block holding G.
- tree gimple\_phi\_result (gimple g) [GIMPLE function] Return the SSA name created by GIMPLE\_PHI G.
- tree \* gimple\_phi\_result\_ptr (gimple g) [GIMPLE function]
  Return a pointer to the SSA name created by GIMPLE\_PHI G.
- void gimple\_phi\_set\_result (gphi \*g, tree result) [GIMPLE function] Set RESULT to be the SSA name created by GIMPLE\_PHI G.

struct phi\_arg\_d \* gimple\_phi\_arg (gimple g, index) [GIMPLE function]
Return the PHI argument corresponding to incoming edge INDEX for GIMPLE\_PHI G.

Set PHIARG to be the argument corresponding to incoming edge INDEX for GIMPLE\_PHI c

# 12.8.25 GIMPLE\_RESX

# gresx \*gimple\_build\_resx (int region)

[GIMPLE function]

Build a GIMPLE\_RESX statement which is a statement. This statement is a placeholder for \_Unwind\_Resume before we know if a function call or a branch is needed. REGION is the exception region from which control is flowing.

int gimple\_resx\_region ( $const\ gresx\ *g$ )

[GIMPLE function]

Return the region number for GIMPLE\_RESX G.

void gimple\_resx\_set\_region (gresx \*g, int region) [GIMPLE function]
Set REGION to be the region number for GIMPLE\_RESX G.

# 12.8.26 GIMPLE RETURN

greturn \*gimple\_build\_return (tree retval)

[GIMPLE function]

Build a GIMPLE\_RETURN statement whose return value is retval.

tree gimple\_return\_retval (const greturn \*g)
Return the return value for GIMPLE\_RETURN G.

[GIMPLE function]

void gimple\_return\_set\_retval (greturn \*g, tree retval)
Set RETVAL to be the return value for GIMPLE\_RETURN G.

[GIMPLE function]

# 12.8.27 GIMPLE\_SWITCH

gswitch \*gimple\_build\_switch (tree index, tree default\_label, vec<tree> \*args)

[GIMPLE function]

Build a GIMPLE\_SWITCH statement. INDEX is the index variable to switch on, and DEFAULT\_LABEL represents the default label. ARGS is a vector of CASE\_LABEL\_EXPR trees that contain the non-default case labels. Each label is a tree of code CASE\_LABEL\_EXPR.

unsigned gimple\_switch\_num\_labels ( const gswitch \*g) [GIMPLE function] Return the number of labels associated with the switch statement G.

Set NLABELS to be the number of labels for the switch statement G.

tree gimple\_switch\_index (const gswitch \*g) [GIMPLE function]
Return the index variable used by the switch statement G.

- void gimple\_switch\_set\_index (gswitch \*g, tree index) [GIMPLE function] Set INDEX to be the index variable for switch statement G.
- tree gimple\_switch\_label (const gswitch \*g, unsigned index) [GIMPLE function] Return the label numbered INDEX. The default label is 0, followed by any labels in a switch statement.
- tree gimple\_switch\_default\_label ( const gswitch \*g) [GIMPLE function] Return the default label for a switch statement.

# 12.8.28 GIMPLE\_TRY

- gtry \*gimple\_build\_try (gimple\_seq eval, gimple\_seq [GIMPLE function] cleanup, unsigned int kind)

  Build a GIMPLE\_TRY statement. EVAL is a sequence with the expression to evaluate. CLEANUP is a sequence of statements to run at clean-up time. KIND is the enumeration value GIMPLE\_TRY\_CATCH if this statement denotes a try/catch construct or GIMPLE\_TRY\_FINALLY if this statement denotes a try/finally construct.
- enum gimple\_try\_flags gimple\_try\_kind (gimple g) [GIMPLE function] Return the kind of try block represented by GIMPLE\_TRY G. This is either GIMPLE\_TRY\_CATCH or GIMPLE\_TRY\_FINALLY.
- bool gimple\_try\_catch\_is\_cleanup (gimple g) [GIMPLE function] Return the GIMPLE\_TRY\_CATCH\_IS\_CLEANUP flag.
- gimple\_seq gimple\_try\_eval (gimple g) [GIMPLE function] Return the sequence of statements used as the body for GIMPLE\_TRY G.
- gimple\_seq gimple\_try\_cleanup (gimple g) [GIMPLE function] Return the sequence of statements used as the cleanup body for GIMPLE\_TRY G.
- void gimple\_try\_set\_eval (gtry \*g, gimple\_seq eval) [GIMPLE function] Set EVAL to be the sequence of statements to use as the body for GIMPLE\_TRY G.
- void gimple\_try\_set\_cleanup (gtry \*g, gimple\_seq cleanup) [GIMPLE function] Set CLEANUP to be the sequence of statements to use as the cleanup body for GIMPLE\_TRY G.

# 12.8.29 GIMPLE\_WITH\_CLEANUP\_EXPR

- gimple gimple\_build\_wce (gimple\_seq cleanup) [GIMPLE function] Build a GIMPLE\_WITH\_CLEANUP\_EXPR statement. CLEANUP is the clean-up expression.
- gimple\_seq gimple\_wce\_cleanup (gimple g) [GIMPLE function]

  Return the cleanup sequence for cleanup statement G.
- void gimple\_wce\_set\_cleanup (gimple g, gimple\_seq cleanup) [GIMPLE function] Set CLEANUP to be the cleanup sequence for G.
- bool gimple\_wce\_cleanup\_eh\_only (gimple g) [GIMPLE function] Return the CLEANUP\_EH\_ONLY flag for a WCE tuple.
- $\begin{tabular}{ll} \begin{tabular}{ll} void gimple_wce_set_cleanup_eh_only (gimple\ g,\ bool \\ eh_only_p) \end{tabular} \begin{tabular}{ll} \begin{tabular}{ll} GIMPLE\ function] \\ \begin{tabular}{ll} \be$

# 12.9 GIMPLE sequences

GIMPLE sequences are the tuple equivalent of STATEMENT\_LIST's used in GENERIC. They are used to chain statements together, and when used in conjunction with sequence iterators, provide a framework for iterating through statements.

GIMPLE sequences are of type struct gimple\_sequence, but are more commonly passed by reference to functions dealing with sequences. The type for a sequence pointer is gimple\_seq which is the same as struct gimple\_sequence \*. When declaring a local sequence, you can define a local variable of type struct gimple\_sequence. When declaring a sequence allocated on the garbage collected heap, use the function gimple\_seq\_alloc documented below

There are convenience functions for iterating through sequences in the section entitled Sequence Iterators.

Below is a list of functions to manipulate and query sequences.

- void gimple\_seq\_add\_stmt (gimple\_seq \*seq, gimple g) [GIMPLE function] Link a gimple statement to the end of the sequence \*SEQ if G is not NULL. If \*SEQ is NULL, allocate a sequence before linking.
- void gimple\_seq\_add\_seq (gimple\_seq \*dest, gimple\_seq src) [GIMPLE function] Append sequence SRC to the end of sequence \*DEST if SRC is not NULL. If \*DEST is NULL, allocate a new sequence before appending.
- gimple\_seq gimple\_seq\_deep\_copy (gimple\_seq src) [GIMPLE function]

  Perform a deep copy of sequence SRC and return the result.
- gimple\_seq gimple\_seq\_reverse (gimple\_seq seq) [GIMPLE function] Reverse the order of the statements in the sequence SEQ. Return SEQ.
- gimple gimple\_seq\_first (gimple\_seq s) [GIMPLE function]
  Return the first statement in sequence S.

```
gimple gimple_seq_last (gimple_seq s)
                                                                  [GIMPLE function]
     Return the last statement in sequence S.
void gimple_seq_set_last (gimple_seq s, gimple last)
                                                                  [GIMPLE function]
     Set the last statement in sequence S to the statement in LAST.
                                                                  [GIMPLE function]
void gimple_seq_set_first (gimple_seq s, gimple first)
     Set the first statement in sequence S to the statement in FIRST.
void gimple_seq_init (gimple_seq s)
                                                                  [GIMPLE function]
     Initialize sequence S to an empty sequence.
gimple_seq_gimple_seq_alloc (void)
                                                                  [GIMPLE function]
     Allocate a new sequence in the garbage collected store and return it.
void gimple_seq_copy (gimple_seq dest, gimple_seq src)
                                                                  [GIMPLE function]
     Copy the sequence SRC into the sequence DEST.
bool gimple_seq_empty_p (gimple_seq s)
                                                                  [GIMPLE function]
     Return true if the sequence S is empty.
gimple_seq bb_seq (basic_block bb)
                                                                  [GIMPLE function]
     Returns the sequence of statements in BB.
void set_bb_seq (basic_block bb, gimple_seq seq)
                                                                  [GIMPLE function]
     Sets the sequence of statements in BB to SEQ.
bool gimple_seq_singleton_p (gimple_seq seq)
                                                                  [GIMPLE function]
     Determine whether SEQ contains exactly one statement.
```

# 12.10 Sequence iterators

Sequence iterators are convenience constructs for iterating through statements in a sequence. Given a sequence SEQ, here is a typical use of gimple sequence iterators:

```
gimple_stmt_iterator gsi;
for (gsi = gsi_start (seq); !gsi_end_p (gsi); gsi_next (&gsi))
   {
     gimple g = gsi_stmt (gsi);
     /* Do something with gimple statement G. */
}
```

Backward iterations are possible:

```
for (gsi = gsi_last (seq); !gsi_end_p (gsi); gsi_prev (&gsi))
```

Forward and backward iterations on basic blocks are possible with gsi\_start\_bb and gsi\_last\_bb.

In the documentation below we sometimes refer to enum gsi\_iterator\_update. The valid options for this enumeration are:

- GSI\_NEW\_STMT Only valid when a single statement is added. Move the iterator to it.
- GSI\_SAME\_STMT Leave the iterator at the same statement.
- GSI\_CONTINUE\_LINKING Move iterator to whatever position is suitable for linking other statements in the same direction.

Below is a list of the functions used to manipulate and use statement iterators.

- gimple\_stmt\_iterator gsi\_start (gimple\_seq seq) [GIMPLE function] Return a new iterator pointing to the sequence SEQ's first statement. If SEQ is empty, the iterator's basic block is NULL. Use gsi\_start\_bb instead when the iterator needs to always have the correct basic block set.
- gimple\_stmt\_iterator gsi\_start\_bb (basic\_block bb) [GIMPLE function] Return a new iterator pointing to the first statement in basic block BB.
- gimple\_stmt\_iterator gsi\_last (gimple\_seq seq) [GIMPLE function] Return a new iterator initially pointing to the last statement of sequence SEQ. If SEQ is empty, the iterator's basic block is NULL. Use gsi\_last\_bb instead when the iterator needs to always have the correct basic block set.
- gimple\_stmt\_iterator gsi\_last\_bb (basic\_block bb) [GIMPLE function] Return a new iterator pointing to the last statement in basic block BB.
- bool gsi\_end\_p (gimple\_stmt\_iterator i) [GIMPLE function] Return TRUE if at the end of I.
- bool gsi\_one\_before\_end\_p (gimple\_stmt\_iterator i) [GIMPLE function]

  Return TRUE if we're one statement before the end of I.
- void gsi\_next (gimple\_stmt\_iterator \*i) [GIMPLE function]
  Advance the iterator to the next gimple statement.
- void gsi\_prev (gimple\_stmt\_iterator \*i) [GIMPLE function]
  Advance the iterator to the previous gimple statement.
- gimple gsi\_stmt (gimple\_stmt\_iterator i) [GIMPLE function]
  Return the current stmt.
- gimple\_stmt\_iterator gsi\_after\_labels (basic\_block bb) [GIMPLE function] Return a block statement iterator that points to the first non-label statement in block BB.
- gimple \* gsi\_stmt\_ptr (gimple\_stmt\_iterator \*i) [GIMPLE function]
  Return a pointer to the current stmt.
- basic\_block gsi\_bb (gimple\_stmt\_iterator i) [GIMPLE function] Return the basic block associated with this iterator.
- gimple\_seq gsi\_seq (gimple\_stmt\_iterator i) [GIMPLE function]

  Return the sequence associated with this iterator.
- void gsi\_remove (gimple\_stmt\_iterator \*i, bool remove\_eh\_info) [GIMPLE function] Remove the current stmt from the sequence. The iterator is updated to point to the next statement. When REMOVE\_EH\_INFO is true we remove the statement pointed to by iterator I from the EH tables. Otherwise we do not modify the EH tables. Generally, REMOVE\_EH\_INFO should be true when the statement is going to be removed from the IL and not reinserted elsewhere.

Links the sequence of statements SEQ before the statement pointed by iterator I. MODE indicates what to do with the iterator after insertion (see enum gsi\_iterator\_update above).

- - Links statement G before the statement pointed-to by iterator I. Updates iterator I according to MODE.

- gimple\_seq gsi\_split\_seq\_after (gimple\_stmt\_iterator i) [GIMPLE function] Move all statements in the sequence after I to a new sequence. Return this new sequence.
- gimple\_seq gsi\_split\_seq\_before (gimple\_stmt\_iterator \*i) [GIMPLE function] Move all statements in the sequence before I to a new sequence. Return this new sequence.
- - Replace the statement pointed-to by I to STMT. If UPDATE\_EH\_INFO is true, the exception handling information of the original statement is moved to the new statement.
- void gsi\_insert\_before (gimple\_stmt\_iterator \*i, gimple [GIMPLE function] stmt, enum gsi\_iterator\_update mode)

  Insert statement STMT before the statement pointed-to by iterator I, update STMT's basic block and scan it for new operands. MODE specifies how to update iterator I after insertion (see enum gsi\_iterator\_update).

after insertion (see enum gsi\_iterator\_update).

gimple\_stmt\_iterator gsi\_for\_stmt (gimple stmt) [GIMPLE function] Finds iterator for STMT.

Move the statement at FROM so it comes right after the statement at TO.

void gsi\_insert\_seq\_on\_edge (edge e, gimple\_seq seq) [GIMPLE function]
Add the sequence of statements in SEQ to the pending list of edge E. No actual insertion is made until a call to gsi\_commit\_edge\_inserts() is made.

Commit insertions pending at edge E. If a new block is created, set NEW\_BB to this block, otherwise set it to NULL.

void gsi\_commit\_edge\_inserts (void) [GIMPLE function]

This routine will commit all pending edge insertions, creating any new basic blocks which are necessary.

# 12.11 Adding a new GIMPLE statement code

be created, it is returned.

The first step in adding a new GIMPLE statement code, is modifying the file gimple.def, which contains all the GIMPLE codes. Then you must add a corresponding gimple subclass located in gimple.h. This in turn, will require you to add a corresponding GTY tag in gsstruct.def, and code to handle this tag in gss\_for\_code which is located in gimple.c.

In order for the garbage collector to know the size of the structure you created in gimple.h, you need to add a case to handle your new GIMPLE statement in gimple\_size which is located in gimple.c.

You will probably want to create a function to build the new gimple statement in gimple.c. The function should be called gimple\_build\_new-tuple-name, and should return the new tuple as a pointer to the appropriate gimple subclass.

If your new statement requires accessors for any members or operands it may have, put simple inline accessors in gimple.h and any non-trivial accessors in gimple.c with a corresponding prototype in gimple.h.

You should add the new statement subclass to the class hierarchy diagram in gimple.texi.

# 12.12 Statement and operand traversals

There are two functions available for walking statements and sequences: walk\_gimple\_stmt and walk\_gimple\_seq, accordingly, and a third function for walking the operands in a statement: walk\_gimple\_op.

This function is used to walk the current statement in GSI, optionally using traversal state stored in WI. If WI is NULL, no state is kept during the traversal.

The callback CALLBACK\_STMT is called. If CALLBACK\_STMT returns true, it means that the callback function has handled all the operands of the statement and it is not necessary to walk its operands.

If CALLBACK\_STMT is NULL or it returns false, CALLBACK\_OP is called on each operand of the statement via walk\_gimple\_op. If walk\_gimple\_op returns non-NULL for any operand, the remaining operands are not scanned.

The return value is that returned by the last call to walk\_gimple\_op, or NULL\_TREE if no CALLBACK\_OP is specified.

Use this function to walk the operands of statement STMT. Every operand is walked via walk\_tree with optional state information in WI.

CALLBACK\_OP is called on each operand of STMT via walk\_tree. Additional parameters to walk\_tree must be stored in WI. For each operand OP, walk\_tree is called as:

walk\_tree (&OP, CALLBACK\_OP, WI, PSET)

If CALLBACK\_OP returns non-NULL for an operand, the remaining operands are not scanned. The return value is that returned by the last call to walk\_tree, or NULL\_TREE if no CALLBACK\_OP is specified.

```
tree walk_gimple_seq (gimple_seq seq, walk_stmt_fn [GIMPLE function] callback_stmt, walk_tree_fn callback_op, struct walk_stmt_info *wi)
```

This function walks all the statements in the sequence SEQ calling walk\_gimple\_stmt on each one. WI is as in walk\_gimple\_stmt. If walk\_gimple\_stmt returns non-NULL, the walk is stopped and the value returned. Otherwise, all the statements are walked and NULL\_TREE returned.

# 13 Analysis and Optimization of GIMPLE tuples

GCC uses three main intermediate languages to represent the program during compilation: GENERIC, GIMPLE and RTL. GENERIC is a language-independent representation generated by each front end. It is used to serve as an interface between the parser and optimizer. GENERIC is a common representation that is able to represent programs written in all the languages supported by GCC.

GIMPLE and RTL are used to optimize the program. GIMPLE is used for target and language independent optimizations (e.g., inlining, constant propagation, tail call elimination, redundancy elimination, etc). Much like GENERIC, GIMPLE is a language independent, tree based representation. However, it differs from GENERIC in that the GIMPLE grammar is more restrictive: expressions contain no more than 3 operands (except function calls), it has no control flow structures and expressions with side effects are only allowed on the right hand side of assignments. See the chapter describing GENERIC and GIMPLE for more details.

This chapter describes the data structures and functions used in the GIMPLE optimizers (also known as "tree optimizers" or "middle end"). In particular, it focuses on all the macros, data structures, functions and programming constructs needed to implement optimization passes for GIMPLE.

# 13.1 Annotations

The optimizers need to associate attributes with variables during the optimization process. For instance, we need to know whether a variable has aliases. All these attributes are stored in data structures called annotations which are then linked to the field ann in struct tree\_common.

# 13.2 SSA Operands

Almost every GIMPLE statement will contain a reference to a variable or memory location. Since statements come in different shapes and sizes, their operands are going to be located at various spots inside the statement's tree. To facilitate access to the statement's operands, they are organized into lists associated inside each statement's annotation. Each element in an operand list is a pointer to a VAR\_DECL, PARM\_DECL or SSA\_NAME tree node. This provides a very convenient way of examining and replacing operands.

Data flow analysis and optimization is done on all tree nodes representing variables. Any node for which SSA\_VAR\_P returns nonzero is considered when scanning statement operands. However, not all SSA\_VAR\_P variables are processed in the same way. For the purposes of optimization, we need to distinguish between references to local scalar variables and references to globals, statics, structures, arrays, aliased variables, etc. The reason is simple, the compiler can gather complete data flow information for a local scalar. On the other hand, a global variable may be modified by a function call, it may not be possible to keep track of all the elements of an array or the fields of a structure, etc.

The operand scanner gathers two kinds of operands: real and virtual. An operand for which is\_gimple\_reg returns true is considered real, otherwise it is a virtual operand. We also distinguish between uses and definitions. An operand is used if its value is loaded by the statement (e.g., the operand at the RHS of an assignment). If the statement assigns a

new value to the operand, the operand is considered a definition (e.g., the operand at the LHS of an assignment).

Virtual and real operands also have very different data flow properties. Real operands are unambiguous references to the full object that they represent. For instance, given

```
{
  int a, b;
  a = b
}
```

Since a and b are non-aliased locals, the statement a = b will have one real definition and one real use because variable a is completely modified with the contents of variable b. Real definition are also known as *killing definitions*. Similarly, the use of b reads all its bits.

In contrast, virtual operands are used with variables that can have a partial or ambiguous reference. This includes structures, arrays, globals, and aliased variables. In these cases, we have two types of definitions. For globals, structures, and arrays, we can determine from a statement whether a variable of these types has a killing definition. If the variable does, then the statement is marked as having a must definition of that variable. However, if a statement is only defining a part of the variable (i.e. a field in a structure), or if we know that a statement might define the variable but we cannot say for sure, then we mark that statement as having a may definition. For instance, given

```
{
  int a, b, *p;
  if (...)
    p = &a;
  else
    p = &b;
  *p = 5;
  return *p;
}
```

The assignment \*p = 5 may be a definition of a or b. If we cannot determine statically where p is pointing to at the time of the store operation, we create virtual definitions to mark that statement as a potential definition site for a and b. Memory loads are similarly marked with virtual use operands. Virtual operands are shown in tree dumps right before the statement that contains them. To request a tree dump with virtual operands, use the '-vops' option to '-fdump-tree':

```
{
  int a, b, *p;
  if (...)
    p = &a;
  else
    p = &b;
  # a = VDEF <a>
  # b = VDEF <b>
  *p = 5;

# VUSE <a>
  # VUSE <b>
  return *p;
}
```

Notice that VDEF operands have two copies of the referenced variable. This indicates that this is not a killing definition of that variable. In this case we refer to it as a may definition or aliased store. The presence of the second copy of the variable in the VDEF operand will become important when the function is converted into SSA form. This will be used to link all the non-killing definitions to prevent optimizations from making incorrect assumptions about them.

Operands are updated as soon as the statement is finished via a call to update\_stmt. If statement elements are changed via SET\_USE or SET\_DEF, then no further action is required (i.e., those macros take care of updating the statement). If changes are made by manipulating the statement's tree directly, then a call must be made to update\_stmt when complete. Calling one of the bsi\_insert routines or bsi\_replace performs an implicit call to update\_stmt.

# 13.2.1 Operand Iterators And Access Routines

Operands are collected by 'tree-ssa-operands.c'. They are stored inside each statement's annotation and can be accessed through either the operand iterators or an access routine.

The following access routines are available for examining operands:

1. SINGLE\_SSA\_{USE,DEF,TREE}\_OPERAND: These accessors will return NULL unless there is exactly one operand matching the specified flags. If there is exactly one operand, the operand is returned as either a tree, def\_operand\_p, or use\_operand\_p.

```
tree t = SINGLE_SSA_TREE_OPERAND (stmt, flags);
use_operand_p u = SINGLE_SSA_USE_OPERAND (stmt, SSA_ALL_VIRTUAL_USES);
def_operand_p d = SINGLE_SSA_DEF_OPERAND (stmt, SSA_OP_ALL_DEFS);
```

2. ZERO\_SSA\_OPERANDS: This macro returns true if there are no operands matching the specified flags.

```
if (ZERO_SSA_OPERANDS (stmt, SSA_OP_ALL_VIRTUALS))
  return:
```

3. NUM\_SSA\_OPERANDS: This macro Returns the number of operands matching 'flags'. This actually executes a loop to perform the count, so only use this if it is really needed.

```
int count = NUM_SSA_OPERANDS (stmt, flags)
```

If you wish to iterate over some or all operands, use the FOR\_EACH\_SSA\_{USE,DEF,TREE}\_OPERAND iterator. For example, to print all the operands for a statement:

```
void
print_ops (tree stmt)
{
   ssa_op_iter;
   tree var;

FOR_EACH_SSA_TREE_OPERAND (var, stmt, iter, SSA_OP_ALL_OPERANDS)
   print_generic_expr (stderr, var, TDF_SLIM);
}
```

How to choose the appropriate iterator:

1. Determine whether you are need to see the operand pointers, or just the trees, and choose the appropriate macro:

```
def_operand_p FOR_EACH_SSA_DEF_OPERAND tree FOR_EACH_SSA_TREE_OPERAND
```

- 2. You need to declare a variable of the type you are interested in, and an ssa\_op\_iter structure which serves as the loop controlling variable.
- 3. Determine which operands you wish to use, and specify the flags of those you are interested in. They are documented in 'tree-ssa-operands.h':

```
#define SSA_OP_USE
                                0x01
                                        /* Real USE operands. */
#define SSA_OP_DEF
                                0x02
                                        /* Real DEF operands. */
#define SSA_OP_VUSE
                                0x04
                                        /* VUSE operands. */
#define SSA_OP_VDEF
                                80x0
                                        /* VDEF operands. */
/* These are commonly grouped operand flags. */
#define SSA_OP_VIRTUAL_USES (SSA_OP_VUSE)
#define SSA_OP_VIRTUAL_DEFS (SSA_OP_VDEF)
                                (SSA_OP_VIRTUAL_USES | SSA_OP_VIRTUAL_DEFS)
#define SSA_OP_ALL_VIRTUALS
#define SSA_OP_ALL_USES (SSA_OP_VIRTUAL_USES | SSA_OP_USE)
#define SSA_OP_ALL_DEFS (SSA_OP_VIRTUAL_DEFS | SSA_OP_DEF)
#define SSA_OP_ALL_OPERANDS (SSA_OP_ALL_USES | SSA_OP_ALL_DEFS)
```

So if you want to look at the use pointers for all the USE and VUSE operands, you would do something like:

```
use_operand_p use_p;
ssa_op_iter iter;

FOR_EACH_SSA_USE_OPERAND (use_p, stmt, iter, (SSA_OP_USE | SSA_OP_VUSE))
    {
        process_use_ptr (use_p);
    }
}
```

The TREE macro is basically the same as the USE and DEF macros, only with the use or def dereferenced via USE\_FROM\_PTR (use\_p) and DEF\_FROM\_PTR (def\_p). Since we aren't using operand pointers, use and defs flags can be mixed.

```
tree var;
ssa_op_iter iter;

FOR_EACH_SSA_TREE_OPERAND (var, stmt, iter, SSA_OP_VUSE)
    {
        print_generic_expr (stderr, var, TDF_SLIM);
    }
```

VDEFs are broken into two flags, one for the DEF portion (SSA\_OP\_VDEF) and one for the USE portion (SSA\_OP\_VUSE).

There are many examples in the code, in addition to the documentation in 'tree-ssa-operands.h' and 'ssa-iterators.h'.

There are also a couple of variants on the stmt iterators regarding PHI nodes.

FOR\_EACH\_PHI\_ARG Works exactly like FOR\_EACH\_SSA\_USE\_OPERAND, except it works over PHI arguments instead of statement operands.

```
/* Look at every virtual PHI use. */
FOR_EACH_PHI_ARG (use_p, phi_stmt, iter, SSA_OP_VIRTUAL_USES)
{
    my_code;
}
/* Look at every real PHI use. */
```

```
FOR_EACH_PHI_ARG (use_p, phi_stmt, iter, SSA_OP_USES)
  my_code;
/* Look at every PHI use. */
FOR_EACH_PHI_ARG (use_p, phi_stmt, iter, SSA_OP_ALL_USES)
  my_code;
```

FOR\_EACH\_PHI\_OR\_STMT\_{USE,DEF} works exactly like FOR\_EACH\_SSA\_{USE,DEF}\_OPERAND, except it will function on either a statement or a PHI node. These should be used when it is appropriate but they are not quite as efficient as the individual FOR\_EACH\_PHI and FOR\_EACH\_SSA routines.

```
FOR_EACH_PHI_OR_STMT_USE (use_operand_p, stmt, iter, flags)
   {
      my_code;
   }
FOR_EACH_PHI_OR_STMT_DEF (def_operand_p, phi, iter, flags)
   {
      my_code;
   }
}
```

# 13.2.2 Immediate Uses

Immediate use information is now always available. Using the immediate use iterators, you may examine every use of any SSA\_NAME. For instance, to change each use of ssa\_var to ssa\_var2 and call fold\_stmt on each stmt after that is done:

```
use_operand_p imm_use_p;
imm_use_iterator iterator;
tree ssa_var, stmt;

FOR_EACH_IMM_USE_STMT (stmt, iterator, ssa_var)
   {
     FOR_EACH_IMM_USE_ON_STMT (imm_use_p, iterator)
        SET_USE (imm_use_p, ssa_var_2);
     fold_stmt (stmt);
}
```

There are 2 iterators which can be used. FOR\_EACH\_IMM\_USE\_FAST is used when the immediate uses are not changed, i.e., you are looking at the uses, but not setting them.

If they do get changed, then care must be taken that things are not changed under the iterators, so use the FOR\_EACH\_IMM\_USE\_STMT and FOR\_EACH\_IMM\_USE\_ON\_STMT iterators. They attempt to preserve the sanity of the use list by moving all the uses for a statement into a controlled position, and then iterating over those uses. Then the optimization can manipulate the stmt when all the uses have been processed. This is a little slower than the FAST version since it adds a placeholder element and must sort through the list a bit for each statement. This placeholder element must be also be removed if the loop is terminated early. The macro BREAK\_FROM\_IMM\_USE\_STMT is provided to do this:

```
FOR_EACH_IMM_USE_STMT (stmt, iterator, ssa_var)
{
   if (stmt == last_stmt)
     BREAK_FROM_IMM_USE_STMT (iterator);

FOR_EACH_IMM_USE_ON_STMT (imm_use_p, iterator)
   SET_USE (imm_use_p, ssa_var_2);
```

```
fold_stmt (stmt);
}
```

There are checks in verify\_ssa which verify that the immediate use list is up to date, as well as checking that an optimization didn't break from the loop without using this macro. It is safe to simply 'break'; from a FOR\_EACH\_IMM\_USE\_FAST traverse.

Some useful functions and macros:

- 1. has\_zero\_uses (ssa\_var): Returns true if there are no uses of ssa\_var.
- 2. has\_single\_use (ssa\_var): Returns true if there is only a single use of ssa\_var.
- 3. single\_imm\_use (ssa\_var, use\_operand\_p \*ptr, tree \*stmt): Returns true if there is only a single use of ssa\_var, and also returns the use pointer and statement it occurs in, in the second and third parameters.
- 4. num\_imm\_uses (ssa\_var): Returns the number of immediate uses of ssa\_var. It is better not to use this if possible since it simply utilizes a loop to count the uses.
- 5. PHI\_ARG\_INDEX\_FROM\_USE (use\_p): Given a use within a PHI node, return the index number for the use. An assert is triggered if the use isn't located in a PHI node.
- 6. USE\_STMT (use\_p): Return the statement a use occurs in.

Note that uses are not put into an immediate use list until their statement is actually inserted into the instruction stream via a bsi\_\* routine.

It is also still possible to utilize lazy updating of statements, but this should be used only when absolutely required. Both alias analysis and the dominator optimizations currently do this.

When lazy updating is being used, the immediate use information is out of date and cannot be used reliably. Lazy updating is achieved by simply marking statements modified via calls to gimple\_set\_modified instead of update\_stmt. When lazy updating is no longer required, all the modified statements must have update\_stmt called in order to bring them up to date. This must be done before the optimization is finished, or verify\_ssa will trigger an abort.

This is done with a simple loop over the instruction stream:

```
block_stmt_iterator bsi;
basic_block bb;
FOR_EACH_BB (bb)
{
   for (bsi = bsi_start (bb); !bsi_end_p (bsi); bsi_next (&bsi))
      update_stmt_if_modified (bsi_stmt (bsi));
}
```

# 13.3 Static Single Assignment

Most of the tree optimizers rely on the data flow information provided by the Static Single Assignment (SSA) form. We implement the SSA form as described in R. Cytron, J. Ferrante, B. Rosen, M. Wegman, and K. Zadeck. Efficiently Computing Static Single Assignment Form and the Control Dependence Graph. ACM Transactions on Programming Languages and Systems, 13(4):451-490, October 1991.

The SSA form is based on the premise that program variables are assigned in exactly one location in the program. Multiple assignments to the same variable create new versions of that variable. Naturally, actual programs are seldom in SSA form initially because variables

tend to be assigned multiple times. The compiler modifies the program representation so that every time a variable is assigned in the code, a new version of the variable is created. Different versions of the same variable are distinguished by subscripting the variable name with its version number. Variables used in the right-hand side of expressions are renamed so that their version number matches that of the most recent assignment.

We represent variable versions using SSA\_NAME nodes. The renaming process in 'tree-ssa.c' wraps every real and virtual operand with an SSA\_NAME node which contains the version number and the statement that created the SSA\_NAME. Only definitions and virtual definitions may create new SSA\_NAME nodes.

Sometimes, flow of control makes it impossible to determine the most recent version of a variable. In these cases, the compiler inserts an artificial definition for that variable called *PHI function* or *PHI node*. This new definition merges all the incoming versions of the variable to create a new name for it. For instance,

```
if (...)
  a_1 = 5;
else if (...)
  a_2 = 2;
else
  a_3 = 13;
# a_4 = PHI <a_1, a_2, a_3>
return a_4;
```

Since it is not possible to determine which of the three branches will be taken at runtime, we don't know which of a\_1, a\_2 or a\_3 to use at the return statement. So, the SSA renamer creates a new version a\_4 which is assigned the result of "merging" a\_1, a\_2 and a\_3. Hence, PHI nodes mean "one of these operands. I don't know which".

The following functions can be used to examine PHI nodes

```
gimple_phi_result (phi)
```

[Function]

Returns the SSA\_NAME created by PHI node phi (i.e., phi's LHS).

```
gimple_phi_num_args (phi)
```

[Function]

Returns the number of arguments in *phi*. This number is exactly the number of incoming edges to the basic block holding *phi*.

```
gimple_phi_arg (phi, i)
```

[Function]

Returns ith argument of phi.

```
gimple_phi_arg_edge (phi, i)
```

[Function]

Returns the incoming edge for the *i*th argument of *phi*.

```
gimple_phi_arg_def (phi, i)
```

[Function]

Returns the SSA\_NAME for the *i*th argument of *phi*.

# 13.3.1 Preserving the SSA form

Some optimization passes make changes to the function that invalidate the SSA property. This can happen when a pass has added new symbols or changed the program so that variables that were previously aliased aren't anymore. Whenever something like this happens,

the affected symbols must be renamed into SSA form again. Transformations that emit new code or replicate existing statements will also need to update the SSA form.

Since GCC implements two different SSA forms for register and virtual variables, keeping the SSA form up to date depends on whether you are updating register or virtual names. In both cases, the general idea behind incremental SSA updates is similar: when new SSA names are created, they typically are meant to replace other existing names in the program.

For instance, given the following code:

```
1 L0:
2 x_1 = PHI (0, x_5)
3 if (x_1 < 10)
4
    if (x_1 > 7)
5
       y_2 = 0
6
7
       y_3 = x_1 + x_7
8
     endif
9
     x_5 = x_1 + 1
10
     goto L0;
11 endif
```

Suppose that we insert new names  $x_10$  and  $x_11$  (lines 4 and 8).

```
1 LO:
2 x_1 = PHI (0, x_5)
  if (x_1 < 10)
3
     x_10 = \dots
5
     if (x_1 > 7)
6
       y_2 = 0
7
     else
8
       x_11 = ...
9
       y_3 = x_1 + x_7
10
     endif
11
     x_5 = x_1 + 1
     goto L0;
13 endif
```

We want to replace all the uses of  $x_1$  with the new definitions of  $x_1$ 0 and  $x_1$ 1. Note that the only uses that should be replaced are those at lines 5, 9 and 11. Also, the use of  $x_7$  at line 9 should *not* be replaced (this is why we cannot just mark symbol x for renaming).

Additionally, we may need to insert a PHI node at line 11 because that is a merge point for  $x\_10$  and  $x\_11$ . So the use of  $x\_1$  at line 11 will be replaced with the new PHI node. The insertion of PHI nodes is optional. They are not strictly necessary to preserve the SSA form, and depending on what the caller inserted, they may not even be useful for the optimizers.

Updating the SSA form is a two step process. First, the pass has to identify which names need to be updated and/or which symbols need to be renamed into SSA form for the first time. When new names are introduced to replace existing names in the program, the mapping between the old and the new names are registered by calling register\_new\_name\_mapping (note that if your pass creates new code by duplicating basic blocks, the call to tree\_duplicate\_bb will set up the necessary mappings automatically).

After the replacement mappings have been registered and new symbols marked for renaming, a call to update\_ssa makes the registered changes. This can be done with an explicit call or by creating TODO flags in the tree\_opt\_pass structure for your pass. There are several TODO flags that control the behavior of update\_ssa:

- TODO\_update\_ssa. Update the SSA form inserting PHI nodes for newly exposed symbols and virtual names marked for updating. When updating real names, only insert PHI nodes for a real name O\_j in blocks reached by all the new and old definitions for O\_j. If the iterated dominance frontier for O\_j is not pruned, we may end up inserting PHI nodes in blocks that have one or more edges with no incoming definition for O\_j. This would lead to uninitialized warnings for O\_j's symbol.
- TODO\_update\_ssa\_no\_phi. Update the SSA form without inserting any new PHI nodes at all. This is used by passes that have either inserted all the PHI nodes themselves or passes that need only to patch use-def and def-def chains for virtuals (e.g., DCE).
- TODO\_update\_ssa\_full\_phi. Insert PHI nodes everywhere they are needed. No pruning of the IDF is done. This is used by passes that need the PHI nodes for O\_j even if it means that some arguments will come from the default definition of O\_j's symbol (e.g., pass\_linear\_transform).
  - WARNING: If you need to use this flag, chances are that your pass may be doing something wrong. Inserting PHI nodes for an old name where not all edges carry a new replacement may lead to silent codegen errors or spurious uninitialized warnings.
- TODO\_update\_ssa\_only\_virtuals. Passes that update the SSA form on their own may want to delegate the updating of virtual names to the generic updater. Since FUD chains are easier to maintain, this simplifies the work they need to do. NOTE: If this flag is used, any OLD->NEW mappings for real names are explicitly destroyed and only the symbols marked for renaming are processed.

# 13.3.2 Examining SSA\_NAME nodes

The following macros can be used to examine SSA\_NAME nodes

# SSA\_NAME\_DEF\_STMT (var)

[Macro]

Returns the statement s that creates the SSA\_NAME var. If s is an empty statement (i.e., IS\_EMPTY\_STMT (s) returns true), it means that the first reference to this variable is a USE or a VUSE.

## SSA\_NAME\_VERSION (var)

[Macro]

Returns the version number of the SSA\_NAME object var.

# 13.3.3 Walking the dominator tree

# void walk\_dominator\_tree (walk\_data, bb)

[Tree SSA function]

This function walks the dominator tree for the current CFG calling a set of callback functions defined in *struct dom\_walk\_data* in 'domwalk.h'. The call back functions you need to define give you hooks to execute custom code at various points during traversal:

- 1. Once to initialize any local data needed while processing bb and its children. This local data is pushed into an internal stack which is automatically pushed and popped as the walker traverses the dominator tree.
- 2. Once before traversing all the statements in the bb.
- 3. Once for every statement inside bb.
- 4. Once after traversing all the statements and before recursing into bb's dominator children.

- 5. It then recurses into all the dominator children of bb.
- 6. After recursing into all the dominator children of bb it can, optionally, traverse every statement in bb again (i.e., repeating steps 2 and 3).
- 7. Once after walking the statements in bb and bb's dominator children. At this stage, the block local data stack is popped.

# 13.4 Alias analysis

Alias analysis in GIMPLE SSA form consists of two pieces. First the virtual SSA web ties conflicting memory accesses and provides a SSA use-def chain and SSA immediate-use chains for walking possibly dependent memory accesses. Second an alias-oracle can be queried to disambiguate explicit and implicit memory references.

# 1. Memory SSA form.

All statements that may use memory have exactly one accompanied use of a virtual SSA name that represents the state of memory at the given point in the IL.

All statements that may define memory have exactly one accompanied definition of a virtual SSA name using the previous state of memory and defining the new state of memory after the given point in the IL.

```
int i;
int foo (void)
{
    # .MEM_3 = VDEF <.MEM_2(D)>
    i = 1;
    # VUSE <.MEM_3>
    return i;
}
```

The virtual SSA names in this case are .MEM\_2(D) and .MEM\_3. The store to the global variable i defines .MEM\_3 invalidating .MEM\_2(D). The load from i uses that new state .MEM\_3.

The virtual SSA web serves as constraints to SSA optimizers preventing illegitimate code-motion and optimization. It also provides a way to walk related memory statements.

## 2. Points-to and escape analysis.

Points-to analysis builds a set of constraints from the GIMPLE SSA IL representing all pointer operations and facts we do or do not know about pointers. Solving this set of constraints yields a conservatively correct solution for each pointer variable in the program (though we are only interested in SSA name pointers) as to what it may possibly point to.

This points-to solution for a given SSA name pointer is stored in the pt\_solution sub-structure of the SSA\_NAME\_PTR\_INFO record. The following accessor functions are available:

- pt\_solution\_includes
- pt\_solutions\_intersect

Points-to analysis also computes the solution for two special set of pointers, ESCAPED and CALLUSED. Those represent all memory that has escaped the scope of analysis or that is used by pure or nested const calls.

## 3. Type-based alias analysis

Type-based alias analysis is frontend dependent though generic support is provided by the middle-end in alias.c. TBAA code is used by both tree optimizers and RTL optimizers.

Every language that wishes to perform language-specific alias analysis should define a function that computes, given a tree node, an alias set for the node. Nodes in different alias sets are not allowed to alias. For an example, see the C front-end function c\_get\_alias\_set.

## 4. Tree alias-oracle

The tree alias-oracle provides means to disambiguate two memory references and memory references against statements. The following queries are available:

- refs\_may\_alias\_p
- ref\_maybe\_used\_by\_stmt\_p
- stmt\_may\_clobber\_ref\_p

In addition to those two kind of statement walkers are available walking statements related to a reference ref. walk\_non\_aliased\_vuses walks over dominating memory defining statements and calls back if the statement does not clobber ref providing the non-aliased VUSE. The walk stops at the first clobbering statement or if asked to. walk\_aliased\_vdefs walks over dominating memory defining statements and calls back on each statement clobbering ref providing its aliasing VDEF. The walk stops if asked to.

# 13.5 Memory model

The memory model used by the middle-end models that of the C/C++ languages. The middle-end has the notion of an effective type of a memory region which is used for type-based alias analysis.

The following is a refinement of ISO C99 6.5/6, clarifying the block copy case to follow common sense and extending the concept of a dynamic effective type to objects with a declared type as required for C++.

The effective type of an object for an access to its stored value is the declared type of the object or the effective type determined by a previous store to it. If a value is stored into an object through an Ivalue having a type that is not a character type, then the type of the Ivalue becomes the effective type of the object for that access and for subsequent accesses that do not modify the stored value. If a value is copied into an object using memcpy or memmove, or is copied as an array of character type, then the effective type of the modified object for that access and for subsequent accesses that do not modify the value is undetermined. For all other accesses to an object, the effective type of the object is simply the type of the Ivalue used for the access.

# 14 RTL Representation

The last part of the compiler work is done on a low-level intermediate representation called Register Transfer Language. In this language, the instructions to be output are described, pretty much one by one, in an algebraic form that describes what the instruction does.

RTL is inspired by Lisp lists. It has both an internal form, made up of structures that point at other structures, and a textual form that is used in the machine description and in printed debugging dumps. The textual form uses nested parentheses to indicate the pointers in the internal form.

# 14.1 RTL Object Types

RTL uses five kinds of objects: expressions, integers, wide integers, strings and vectors. Expressions are the most important ones. An RTL expression ("RTX", for short) is a C structure, but it is usually referred to with a pointer; a type that is given the typedef name rtx.

An integer is simply an int; their written form uses decimal digits. A wide integer is an integral object whose type is HOST\_WIDE\_INT; their written form uses decimal digits.

A string is a sequence of characters. In core it is represented as a char \* in usual C fashion, and it is written in C syntax as well. However, strings in RTL may never be null. If you write an empty string in a machine description, it is represented in core as a null pointer rather than as a pointer to a null character. In certain contexts, these null pointers instead of strings are valid. Within RTL code, strings are most commonly found inside symbol\_ref expressions, but they appear in other contexts in the RTL expressions that make up machine descriptions.

In a machine description, strings are normally written with double quotes, as you would in C. However, strings in machine descriptions may extend over many lines, which is invalid C, and adjacent string constants are not concatenated as they are in C. Any string constant may be surrounded with a single set of parentheses. Sometimes this makes the machine description easier to read.

There is also a special syntax for strings, which can be useful when C code is embedded in a machine description. Wherever a string can appear, it is also valid to write a C-style brace block. The entire brace block, including the outermost pair of braces, is considered to be the string constant. Double quote characters inside the braces are not special. Therefore, if you write string constants in the C code, you need not escape each quote character with a backslash.

A vector contains an arbitrary number of pointers to expressions. The number of elements in the vector is explicitly present in the vector. The written form of a vector consists of square brackets ('[...]') surrounding the elements, in sequence and with whitespace separating them. Vectors of length zero are not created; null pointers are used instead.

Expressions are classified by expression codes (also called RTX codes). The expression code is a name defined in 'rtl.def', which is also (in uppercase) a C enumeration constant. The possible expression codes and their meanings are machine-independent. The code of an RTX can be extracted with the macro GET\_CODE (x) and altered with PUT\_CODE (x, newcode).

The expression code determines how many operands the expression contains, and what kinds of objects they are. In RTL, unlike Lisp, you cannot tell by looking at an operand what kind of object it is. Instead, you must know from its context—from the expression code of the containing expression. For example, in an expression of code subreg, the first operand is to be regarded as an expression and the second operand as a polynomial integer. In an expression of code plus, there are two operands, both of which are to be regarded as expressions. In a symbol\_ref expression, there is one operand, which is to be regarded as a string.

Expressions are written as parentheses containing the name of the expression type, its flags and machine mode if any, and then the operands of the expression (separated by spaces).

Expression code names in the 'md' file are written in lowercase, but when they appear in C code they are written in uppercase. In this manual, they are shown as follows: const\_int.

In a few contexts a null pointer is valid where an expression is normally wanted. The written form of this is (nil).

# 14.2 RTL Classes and Formats

The various expression codes are divided into several *classes*, which are represented by single characters. You can determine the class of an RTX code with the macro GET\_RTX\_CLASS (code). Currently, 'rtl.def' defines these classes:

RTX\_OBJ An RTX code that represents an actual object, such as a register (REG) or a memory location (MEM, SYMBOL\_REF). LO\_SUM) is also included; instead, SUBREG and STRICT\_LOW\_PART are not in this class, but in class RTX\_EXTRA.

## RTX\_CONST\_OBJ

An RTX code that represents a constant object. HIGH is also included in this class.

# RTX\_COMPARE

An RTX code for a non-symmetric comparison, such as GEU or LT.

# RTX\_COMM\_COMPARE

An RTX code for a symmetric (commutative) comparison, such as EQ or ORDERED.

#### RTX\_UNARY

An RTX code for a unary arithmetic operation, such as NEG, NOT, or ABS. This category also includes value extension (sign or zero) and conversions between integer and floating point.

## RTX\_COMM\_ARITH

An RTX code for a commutative binary operation, such as PLUS or AND. NE and EQ are comparisons, so they have class RTX\_COMM\_COMPARE.

# RTX\_BIN\_ARITH

An RTX code for a non-commutative binary operation, such as MINUS, DIV, or ASHIFTRT.

## RTX\_BITFIELD\_OPS

An RTX code for a bit-field operation. Currently only ZERO\_EXTRACT and SIGN\_EXTRACT. These have three inputs and are lvalues (so they can be used for insertion as well). See Section 14.11 [Bit-Fields], page 294.

## RTX\_TERNARY

An RTX code for other three input operations. Currently only IF\_THEN\_ELSE, VEC\_MERGE, SIGN\_EXTRACT, ZERO\_EXTRACT, and FMA.

RTX\_INSN An RTX code for an entire instruction: INSN, JUMP\_INSN, and CALL\_INSN. See Section 14.19 [Insns], page 304.

## RTX\_MATCH

An RTX code for something that matches in insns, such as MATCH\_DUP. These only occur in machine descriptions.

## RTX\_AUTOINC

An RTX code for an auto-increment addressing mode, such as POST\_INC. 'XEXP (x, 0)' gives the auto-modified register.

## RTX\_EXTRA

All other RTX codes. This category includes the remaining codes used only in machine descriptions (DEFINE\_\*, etc.). It also includes all the codes describing side effects (SET, USE, CLOBBER, etc.) and the non-insns that may appear on an insn chain, such as NOTE, BARRIER, and CODE\_LABEL. SUBREG is also part of this class.

For each expression code, 'rtl.def' specifies the number of contained objects and their kinds using a sequence of characters called the *format* of the expression code. For example, the format of subreg is 'ep'.

These are the most commonly used format characters:

- e An expression (actually a pointer to an expression).
- i An integer.
- w A wide integer.
- s A string.
- E A vector of expressions.

A few other format characters are used occasionally:

- u 'u' is equivalent to 'e' except that it is printed differently in debugging dumps. It is used for pointers to insns.
- n 'n' is equivalent to 'i' except that it is printed differently in debugging dumps. It is used for the line number or code number of a **note** insn.
- S 'S' indicates a string which is optional. In the RTL objects in core, 'S' is equivalent to 's', but when the object is read, from an 'md' file, the string value of this operand may be omitted. An omitted string is taken to be the null string.

- V 'V' indicates a vector which is optional. In the RTL objects in core, 'V' is equivalent to 'E', but when the object is read from an 'md' file, the vector value of this operand may be omitted. An omitted vector is effectively the same as a vector of no elements.
- B 'B' indicates a pointer to basic block structure.
- p A polynomial integer. At present this is used only for SUBREG\_BYTE.
- 0 '0' means a slot whose contents do not fit any normal category. '0' slots are not printed at all in dumps, and are often used in special ways by small parts of the compiler.

There are macros to get the number of operands and the format of an expression code:

# GET\_RTX\_LENGTH (code)

Number of operands of an RTX of code code.

# GET\_RTX\_FORMAT (code)

The format of an RTX of code code, as a C string.

Some classes of RTX codes always have the same format. For example, it is safe to assume that all comparison operations have format **ee**.

#### RTX\_UNARY

All codes of this class have format e.

RTX\_BIN\_ARITH

RTX\_COMM\_ARITH

RTX\_COMM\_COMPARE

RTX\_COMPARE

All codes of these classes have format ee.

RTX\_BITFIELD\_OPS

RTX\_TERNARY

All codes of these classes have format eee.

RTX\_INSN All codes of this class have formats that begin with iuueiee. See Section 14.19 [Insns], page 304. Note that not all RTL objects linked onto an insn chain are of class RTX\_INSN.

RTX\_CONST\_OBJ

RTX\_OBJ

RTX\_MATCH

RTX\_EXTRA

You can make no assumptions about the format of these codes.

# 14.3 Access to Operands

Operands of expressions are accessed using the macros XEXP, XINT, XWINT and XSTR. Each of these macros takes two arguments: an expression-pointer (RTX) and an operand number (counting from zero). Thus,

XEXP (x, 2)

accesses operand 2 of expression x, as an expression.

XINT (x, 2)

accesses the same operand as an integer. XSTR, used in the same fashion, would access it as a string.

Any operand can be accessed as an integer, as an expression or as a string. You must choose the correct method of access for the kind of value actually stored in the operand. You would do this based on the expression code of the containing expression. That is also how you would know how many operands there are.

For example, if x is an  $int_list$  expression, you know that it has two operands which can be correctly accessed as XINT (x, 0) and XEXP (x, 1). Incorrect accesses like XEXP (x, 0) and XINT (x, 1) would compile, but would trigger an internal compiler error when rtl checking is enabled. Nothing stops you from writing XEXP (x, 28) either, but this will access memory past the end of the expression with unpredictable results.

Access to operands which are vectors is more complicated. You can use the macro XVEC to get the vector-pointer itself, or the macros XVECEXP and XVECLEN to access the elements and length of a vector.

XVEC (exp, idx)

Access the vector-pointer which is operand number idx in exp.

XVECLEN (exp, idx)

Access the length (number of elements) in the vector which is in operand number idx in exp. This value is an int.

XVECEXP (exp, idx, eltnum)

Access element number *eltnum* in the vector which is in operand number *idx* in *exp*. This value is an RTX.

It is up to you to make sure that *eltnum* is not negative and is less than XVECLEN (exp, idx).

All the macros defined in this section expand into lvalues and therefore can be used to assign the operands, lengths and vector elements as well as to access them.

# 14.4 Access to Special Operands

Some RTL nodes have special annotations associated with them.

MEM

 $MEM_ALIAS_SET(x)$ 

If 0, x is not in any alias set, and may alias anything. Otherwise, x can only alias MEMs in a conflicting alias set. This value is set in a language-dependent manner in the front-end, and should not be altered in the back-end. In some front-ends, these numbers may correspond in some way to types, or other language-level entities, but they need not, and the back-end makes no such assumptions. These set numbers are tested with alias\_sets\_conflict\_p.

#### $MEM_EXPR(x)$

If this register is known to hold the value of some user-level declaration, this is that tree node. It may also be a COMPONENT\_REF,

in which case this is some field reference, and TREE\_OPERAND (x, 0) contains the declaration, or another COMPONENT\_REF, or null if there is no compile-time object associated with the reference.

# MEM\_OFFSET\_KNOWN\_P (x)

True if the offset of the memory reference from MEM\_EXPR is known. 'MEM\_OFFSET (x)' provides the offset if so.

## $MEM_OFFSET(x)$

The offset from the start of MEM\_EXPR. The value is only valid if 'MEM\_OFFSET\_KNOWN\_P (x)' is true.

## MEM\_SIZE\_KNOWN\_P (x)

True if the size of the memory reference is known. 'MEM\_SIZE (x)' provides its size if so.

## $MEM_SIZE(x)$

The size in bytes of the memory reference. This is mostly relevant for BLKmode references as otherwise the size is implied by the mode. The value is only valid if 'MEM\_SIZE\_KNOWN\_P (x)' is true.

#### $MEM\_ALIGN(x)$

The known alignment in bits of the memory reference.

## MEM\_ADDR\_SPACE (x)

The address space of the memory reference. This will commonly be zero for the generic address space.

#### REG

## ORIGINAL\_REGNO (x)

This field holds the number the register "originally" had; for a pseudo register turned into a hard reg this will hold the old pseudo register number.

## $REG_EXPR(x)$

If this register is known to hold the value of some user-level declaration, this is that tree node.

## REG\_OFFSET (x)

If this register is known to hold the value of some user-level declaration, this is the offset into that logical storage.

#### SYMBOL\_REF

#### SYMBOL\_REF\_DECL (x)

If the  $symbol_ref x$  was created for a VAR\_DECL or a FUNCTION\_DECL, that tree is recorded here. If this value is null, then x was created by back end code generation routines, and there is no associated front end symbol table entry.

SYMBOL\_REF\_DECL may also point to a tree of class 'c', that is, some sort of constant. In this case, the <code>symbol\_ref</code> is an entry in the per-file constant pool; again, there is no associated front end symbol table entry.

# SYMBOL\_REF\_CONSTANT (x)

If 'CONSTANT\_POOL\_ADDRESS\_P (x)' is true, this is the constant pool entry for x. It is null otherwise.

# SYMBOL\_REF\_DATA (x)

A field of opaque type used to store SYMBOL\_REF\_DECL or SYMBOL\_REF\_CONSTANT.

## SYMBOL\_REF\_FLAGS (x)

In a symbol\_ref, this is used to communicate various predicates about the symbol. Some of these are common enough to be computed by common code, some are specific to the target. The common bits are:

## SYMBOL\_FLAG\_FUNCTION

Set if the symbol refers to a function.

# SYMBOL\_FLAG\_LOCAL

Set if the symbol is local to this "module". See TARGET\_BINDS\_LOCAL\_P.

# SYMBOL\_FLAG\_EXTERNAL

Set if this symbol is not defined in this translation unit. Note that this is not the inverse of SYMBOL\_FLAG\_LOCAL.

## SYMBOL\_FLAG\_SMALL

Set if the symbol is located in the small data section. See TARGET\_IN\_SMALL\_DATA\_P.

## SYMBOL\_REF\_TLS\_MODEL (x)

This is a multi-bit field accessor that returns the tls\_model to be used for a thread-local storage symbol. It returns zero for non-thread-local symbols.

## SYMBOL\_FLAG\_HAS\_BLOCK\_INFO

Set if the symbol has SYMBOL\_REF\_BLOCK and SYMBOL\_REF\_BLOCK\_OFFSET fields.

# SYMBOL\_FLAG\_ANCHOR

Set if the symbol is used as a section anchor. "Section anchors" are symbols that have a known position within an object\_block and that can be used to access nearby members of that block. They are used to implement '-fsection-anchors'.

If this flag is set, then SYMBOL\_FLAG\_HAS\_BLOCK\_INFO will be too.

Bits beginning with SYMBOL\_FLAG\_MACH\_DEP are available for the target's use.

## SYMBOL\_REF\_BLOCK (x)

If 'SYMBOL\_REF\_HAS\_BLOCK\_INFO\_P (x)', this is the 'object\_block' structure to which the symbol belongs, or NULL if it has not been assigned a block.

# SYMBOL\_REF\_BLOCK\_OFFSET (x)

If 'SYMBOL\_REF\_HAS\_BLOCK\_INFO\_P (x)', this is the offset of x from the first object in 'SYMBOL\_REF\_BLOCK (x)'. The value is negative if x has not yet been assigned to a block, or it has not been given an offset within that block.

# 14.5 Flags in an RTL Expression

RTL expressions contain several flags (one-bit bit-fields) that are used in certain types of expression. Most often they are accessed with the following macros, which expand into lvalues.

# CROSSING\_JUMP\_P (x)

Nonzero in a jump\_insn if it crosses between hot and cold sections, which could potentially be very far apart in the executable. The presence of this flag indicates to other optimizations that this branching instruction should not be "collapsed" into a simpler branching construct. It is used when the optimization to partition basic blocks into hot and cold sections is turned on.

# CONSTANT\_POOL\_ADDRESS\_P (x)

Nonzero in a symbol\_ref if it refers to part of the current function's constant pool. For most targets these addresses are in a .rodata section entirely separate from the function, but for some targets the addresses are close to the beginning of the function. In either case GCC assumes these addresses can be addressed directly, perhaps with the help of base registers. Stored in the unchanging field and printed as '/u'.

# INSN\_ANNULLED\_BRANCH\_P (x)

In a jump\_insn, call\_insn, or insn indicates that the branch is an annulling one. See the discussion under sequence below. Stored in the unchanging field and printed as '/u'.

# INSN\_DELETED\_P (x)

In an insn, call\_insn, jump\_insn, code\_label, jump\_table\_data, barrier, or note, nonzero if the insn has been deleted. Stored in the volatil field and printed as '/v'.

## INSN\_FROM\_TARGET\_P (x)

In an insn or jump\_insn or call\_insn in a delay slot of a branch, indicates that the insn is from the target of the branch. If the branch insn has INSN\_ANNULLED\_BRANCH\_P set, this insn will only be executed if the branch is taken. For annulled branches with INSN\_FROM\_TARGET\_P clear, the insn will be executed only if the branch is not taken. When INSN\_ANNULLED\_BRANCH\_P is not set, this insn will always be executed. Stored in the in\_struct field and printed as '/s'.

## LABEL\_PRESERVE\_P (x)

In a code\_label or note, indicates that the label is referenced by code or data not visible to the RTL of a given function. Labels referenced by a non-local goto will have this bit set. Stored in the in\_struct field and printed as '/s'.

# LABEL\_REF\_NONLOCAL\_P (x)

In label\_ref and reg\_label expressions, nonzero if this is a reference to a non-local label. Stored in the volatil field and printed as '/v'.

## $MEM_KEEP_ALIAS_SET_P(x)$

In mem expressions, 1 if we should keep the alias set for this mem unchanged when we access a component. Set to 1, for example, when we are already in a non-addressable component of an aggregate. Stored in the jump field and printed as '/j'.

## MEM\_VOLATILE\_P (x)

In mem, asm\_operands, and asm\_input expressions, nonzero for volatile memory references. Stored in the volatil field and printed as '/v'.

## MEM\_NOTRAP\_P (x)

In mem, nonzero for memory references that will not trap. Stored in the call field and printed as '/c'.

## $MEM_POINTER(x)$

Nonzero in a mem if the memory reference holds a pointer. Stored in the frame\_related field and printed as '/f'.

# MEM\_READONLY\_P (x)

Nonzero in a mem, if the memory is statically allocated and read-only.

Read-only in this context means never modified during the lifetime of the program, not necessarily in ROM or in write-disabled pages. A common example of the later is a shared library's global offset table. This table is initialized by the runtime loader, so the memory is technically writable, but after control is transferred from the runtime loader to the application, this memory will never be subsequently modified.

Stored in the unchanging field and printed as '/u'.

# PREFETCH\_SCHEDULE\_BARRIER\_P (x)

In a prefetch, indicates that the prefetch is a scheduling barrier. No other INSNs will be moved over it. Stored in the volatil field and printed as '/v'.

# REG\_FUNCTION\_VALUE\_P (x)

Nonzero in a reg if it is the place in which this function's value is going to be returned. (This happens only in a hard register.) Stored in the return\_val field and printed as '/i'.

## REG\_POINTER (x)

Nonzero in a reg if the register holds a pointer. Stored in the frame\_related field and printed as '/f'.

# REG\_USERVAR\_P (x)

In a reg, nonzero if it corresponds to a variable present in the user's source code. Zero for temporaries generated internally by the compiler. Stored in the volatil field and printed as '/v'.

The same hard register may be used also for collecting the values of functions called by this one, but REG\_FUNCTION\_VALUE\_P is zero in this kind of use.

# RTL\_CONST\_CALL\_P (x)

In a call\_insn indicates that the insn represents a call to a const function. Stored in the unchanging field and printed as '/u'.

# RTL\_PURE\_CALL\_P (x)

In a call\_insn indicates that the insn represents a call to a pure function. Stored in the return\_val field and printed as '/i'.

#### RTL\_CONST\_OR\_PURE\_CALL\_P (x)

In a call\_insn, true if RTL\_CONST\_CALL\_P or RTL\_PURE\_CALL\_P is true.

## RTL\_LOOPING\_CONST\_OR\_PURE\_CALL\_P (x)

In a call\_insn indicates that the insn represents a possibly infinite looping call to a const or pure function. Stored in the call field and printed as '/c'. Only true if one of RTL\_CONST\_CALL\_P or RTL\_PURE\_CALL\_P is true.

# RTX\_FRAME\_RELATED\_P (x)

Nonzero in an insn, call\_insn, jump\_insn, barrier, or set which is part of a function prologue and sets the stack pointer, sets the frame pointer, or saves a register. This flag should also be set on an instruction that sets up a temporary register to use in place of the frame pointer. Stored in the frame\_related field and printed as '/f'.

In particular, on RISC targets where there are limits on the sizes of immediate constants, it is sometimes impossible to reach the register save area directly from the stack pointer. In that case, a temporary register is used that is near enough to the register save area, and the Canonical Frame Address, i.e., DWARF2's logical frame pointer, register must (temporarily) be changed to be this temporary register. So, the instruction that sets this temporary register must be marked as RTX\_FRAME\_RELATED\_P.

If the marked instruction is overly complex (defined in terms of what dwarf2out\_frame\_debug\_expr can handle), you will also have to create a REG\_FRAME\_RELATED\_EXPR note and attach it to the instruction. This note should contain a simple expression of the computation performed by this instruction, i.e., one that dwarf2out\_frame\_debug\_expr can handle.

This flag is required for exception handling support on targets with RTL prologues.

# SCHED\_GROUP\_P (x)

During instruction scheduling, in an insn, call\_insn, jump\_insn or jump\_table\_data, indicates that the previous insn must be scheduled together with this insn. This is used to ensure that certain groups of instructions will not be split up by the instruction scheduling pass, for example, use insns before a call\_insn may not be separated from the call\_insn. Stored in the in\_struct field and printed as '/s'.

## $SET_IS_RETURN_P(x)$

For a set, nonzero if it is for a return. Stored in the jump field and printed as '/j'.

# SIBLING\_CALL\_P (x)

For a call\_insn, nonzero if the insn is a sibling call. Stored in the jump field and printed as '/j'.

## STRING\_POOL\_ADDRESS\_P (x)

For a symbol\_ref expression, nonzero if it addresses this function's string constant pool. Stored in the frame\_related field and printed as '/f'.

## SUBREG\_PROMOTED\_UNSIGNED\_P (x)

Returns a value greater then zero for a subreg that has SUBREG\_PROMOTED\_VAR\_P nonzero if the object being referenced is kept zero-extended, zero if it is kept sign-extended, and less then zero if it is extended some other way via the ptr\_extend instruction. Stored in the unchanging field and volatil field, printed as '/u' and '/v'. This macro may only be used to get the value it may not be used to change the value. Use SUBREG\_PROMOTED\_UNSIGNED\_SET to change the value.

## SUBREG\_PROMOTED\_UNSIGNED\_SET (x)

Set the unchanging and volatil fields in a subreg to reflect zero, sign, or other extension. If volatil is zero, then unchanging as nonzero means zero extension and as zero means sign extension. If volatil is nonzero then some other type of extension was done via the ptr\_extend instruction.

#### SUBREG\_PROMOTED\_VAR\_P (x)

Nonzero in a subreg if it was made when accessing an object that was promoted to a wider mode in accord with the PROMOTED\_MODE machine description macro (see Section 18.5 [Storage Layout], page 490). In this case, the mode of the subreg is the declared mode of the object and the mode of SUBREG\_REG is the mode of the register that holds the object. Promoted variables are always either sign- or zero-extended to the wider mode on every assignment. Stored in the in\_struct field and printed as '/s'.

## SYMBOL\_REF\_USED (x)

In a  $symbol_ref$ , indicates that x has been used. This is normally only used to ensure that x is only declared external once. Stored in the used field.

# SYMBOL\_REF\_WEAK (x)

In a symbol\_ref, indicates that x has been declared weak. Stored in the return\_val field and printed as '/i'.

# SYMBOL\_REF\_FLAG (x)

In a symbol\_ref, this is used as a flag for machine-specific purposes. Stored in the volatil field and printed as '/v'.

Most uses of SYMBOL\_REF\_FLAG are historic and may be subsumed by SYMBOL\_REF\_FLAGS. Certainly use of SYMBOL\_REF\_FLAGS is mandatory if the target requires more than one bit of storage.

These are the fields to which the above macros refer:

In a mem, 1 means that the memory reference will not trap.

In a call, 1 means that this pure or const call may possibly infinite loop.

In an RTL dump, this flag is represented as '/c'.

#### frame\_related

In an insn or set expression, 1 means that it is part of a function prologue and sets the stack pointer, sets the frame pointer, saves a register, or sets up a temporary register to use in place of the frame pointer.

In reg expressions, 1 means that the register holds a pointer.

In mem expressions, 1 means that the memory reference holds a pointer.

In symbol\_ref expressions, 1 means that the reference addresses this function's string constant pool.

In an RTL dump, this flag is represented as '/f'.

#### in\_struct

In reg expressions, it is 1 if the register has its entire life contained within the test expression of some loop.

In subreg expressions, 1 means that the subreg is accessing an object that has had its mode promoted from a wider mode.

In label\_ref expressions, 1 means that the referenced label is outside the innermost loop containing the insn in which the label\_ref was found.

In code\_label expressions, it is 1 if the label may never be deleted. This is used for labels which are the target of non-local gotos. Such a label that would have been deleted is replaced with a note of type NOTE\_INSN\_DELETED\_LABEL.

In an insn during dead-code elimination, 1 means that the insn is dead code.

In an insn or jump\_insn during reorg for an insn in the delay slot of a branch, 1 means that this insn is from the target of the branch.

In an insn during instruction scheduling, 1 means that this insn must be scheduled as part of a group together with the previous insn.

In an RTL dump, this flag is represented as '/s'.

#### return\_val

In reg expressions, 1 means the register contains the value to be returned by the current function. On machines that pass parameters in registers, the same register number may be used for parameters as well, but this flag is not set on such uses.

In symbol\_ref expressions, 1 means the referenced symbol is weak.

In call expressions, 1 means the call is pure.

In an RTL dump, this flag is represented as '/i'.

jump In a mem expression, 1 means we should keep the alias set for this mem unchanged when we access a component.

In a set, 1 means it is for a return.

In a call\_insn, 1 means it is a sibling call.

In a jump\_insn, 1 means it is a crossing jump.

In an RTL dump, this flag is represented as '/j'.

#### unchanging

In reg and mem expressions, 1 means that the value of the expression never changes.

In subreg expressions, it is 1 if the subreg references an unsigned object whose mode has been promoted to a wider mode.

In an insn or jump\_insn in the delay slot of a branch instruction, 1 means an annulling branch should be used.

In a symbol\_ref expression, 1 means that this symbol addresses something in the per-function constant pool.

In a call\_insn 1 means that this instruction is a call to a const function.

In an RTL dump, this flag is represented as '/u'.

used

This flag is used directly (without an access macro) at the end of RTL generation for a function, to count the number of times an expression appears in insns. Expressions that appear more than once are copied, according to the rules for shared structure (see Section 14.21 [Sharing], page 314).

For a reg, it is used directly (without an access macro) by the leaf register renumbering code to ensure that each register is only renumbered once.

In a symbol\_ref, it indicates that an external declaration for the symbol has already been written.

volatil

In a mem, asm\_operands, or asm\_input expression, it is 1 if the memory reference is volatile. Volatile memory references may not be deleted, reordered or combined.

In a symbol\_ref expression, it is used for machine-specific purposes.

In a reg expression, it is 1 if the value is a user-level variable. 0 indicates an internal compiler temporary.

In an insn, 1 means the insn has been deleted.

In label\_ref and reg\_label expressions, 1 means a reference to a non-local label.

In prefetch expressions, 1 means that the containing insn is a scheduling barrier.

In an RTL dump, this flag is represented as '/v'.

# 14.6 Machine Modes

A machine mode describes a size of data object and the representation used for it. In the C code, machine modes are represented by an enumeration type, machine\_mode, defined in 'machmode.def'. Each RTL expression has room for a machine mode and so do certain kinds of tree expressions (declarations and types, to be precise).

In debugging dumps and machine descriptions, the machine mode of an RTL expression is written after the expression code with a colon to separate them. The letters 'mode' which appear at the end of each machine mode name are omitted. For example, (reg:SI 38) is a reg expression with machine mode SImode. If the mode is VOIDmode, it is not written at all.

Here is a table of machine modes. The term "byte" below refers to an object of BITS\_PER\_UNIT bits (see Section 18.5 [Storage Layout], page 490).

BImode "Bit" mode represents a single bit, for predicate registers.

QImode "Quarter-Integer" mode represents a single byte treated as an integer.

HImode "Half-Integer" mode represents a two-byte integer.

PSImode "Partial Single Integer" mode represents an integer which occupies four bytes but which doesn't really use all four. On some machines, this is the right mode to use for pointers.

SImode "Single Integer" mode represents a four-byte integer.

PDImode "Partial Double Integer" mode represents an integer which occupies eight bytes but which doesn't really use all eight. On some machines, this is the right mode to use for certain pointers.

DImode "Double Integer" mode represents an eight-byte integer.

TImode "Tetra Integer" (?) mode represents a sixteen-byte integer.

Olmode "Octa Integer" (?) mode represents a thirty-two-byte integer.

XImode "Hexadeca Integer" (?) mode represents a sixty-four-byte integer.

QFmode "Quarter-Floating" mode represents a quarter-precision (single byte) floating point number.

HFmode "Half-Floating" mode represents a half-precision (two byte) floating point number.

TQFmode "Three-Quarter-Floating" (?) mode represents a three-quarter-precision (three byte) floating point number.

"Single Floating" mode represents a four byte floating point number. In the common case, of a processor with IEEE arithmetic and 8-bit bytes, this is a single-precision IEEE floating point number; it can also be used for double-precision (on processors with 16-bit bytes) and single-precision VAX and IBM types.

DFmode "Double Floating" mode represents an eight byte floating point number. In the common case, of a processor with IEEE arithmetic and 8-bit bytes, this is a double-precision IEEE floating point number.

WFmode "Extended Floating" mode represents an IEEE extended floating point number. This mode only has 80 meaningful bits (ten bytes). Some processors require such numbers to be padded to twelve bytes, others to sixteen; this mode is used for either.

SDmode "Single Decimal Floating" mode represents a four byte decimal floating point number (as distinct from conventional binary floating point).

DDmode "Double Decimal Floating" mode represents an eight byte decimal floating point number.

TDmode "Tetra Decimal Floating" mode represents a sixteen byte decimal floating point number all 128 of whose bits are meaningful.

TFmode "Tetra Floating" mode represents a sixteen byte floating point number all 128 of whose bits are meaningful. One common use is the IEEE quad-precision format.

QQmode "Quarter-Fractional" mode represents a single byte treated as a signed fractional number. The default format is "s.7".

HQmode "Half-Fractional" mode represents a two-byte signed fractional number. The default format is "s.15".

"Single Fractional" mode represents a four-byte signed fractional number. The default format is "s.31".

DQmode "Double Fractional" mode represents an eight-byte signed fractional number. The default format is "s.63".

TQmode "Tetra Fractional" mode represents a sixteen-byte signed fractional number. The default format is "s.127".

UQQmode "Unsigned Quarter-Fractional" mode represents a single byte treated as an unsigned fractional number. The default format is ".8".

UHQmode "Unsigned Half-Fractional" mode represents a two-byte unsigned fractional number. The default format is ".16".

"Unsigned Single Fractional" mode represents a four-byte unsigned fractional number. The default format is ".32".

UDQmode "Unsigned Double Fractional" mode represents an eight-byte unsigned fractional number. The default format is ".64".

"Unsigned Tetra Fractional" mode represents a sixteen-byte unsigned fractional number. The default format is ".128".

HAmode "Half-Accumulator" mode represents a two-byte signed accumulator. The default format is "s8.7".

SAmode "Single Accumulator" mode represents a four-byte signed accumulator. The default format is "s16.15".

DAmode "Double Accumulator" mode represents an eight-byte signed accumulator. The default format is "s32.31".

TAmode "Tetra Accumulator" mode represents a sixteen-byte signed accumulator. The default format is "s64.63".

UHAmode "Unsigned Half-Accumulator" mode represents a two-byte unsigned accumulator. The default format is "8.8".

"Unsigned Single Accumulator" mode represents a four-byte unsigned accumulator. The default format is "16.16".

UDAmode "Unsigned Double Accumulator" mode represents an eight-byte unsigned accumulator. The default format is "32.32".

UTAmode "Unsigned Tetra Accumulator" mode represents a sixteen-byte unsigned accumulator. The default format is "64.64".

"Condition Code" mode represents the value of a condition code, which is a machine-specific set of bits used to represent the result of a comparison operation. Other machine-specific modes may also be used for the condition code.

These modes are not used on machines that use cc0 (see Section 18.15 [Condition Code], page 572).

BLKmode "Block" mode represents values that are aggregates to which none of the other modes apply. In RTL, only memory references can have this mode, and only if they appear in string-move or vector instructions. On machines which have no such instructions, BLKmode will not appear in RTL.

VOIDmode Void mode means the absence of a mode or an unspecified mode. For example, RTL expressions of code const\_int have mode VOIDmode because they can be taken to have whatever mode the context requires. In debugging dumps of RTL, VOIDmode is expressed by the absence of any mode.

#### QCmode, HCmode, SCmode, DCmode, XCmode, TCmode

These modes stand for a complex number represented as a pair of floating point values. The floating point values are in QFmode, HFmode, SFmode, DFmode, XFmode, and TFmode, respectively.

CQImode, CHImode, CSImode, CDImode, CTImode, COImode, CPSImode

These modes stand for a complex number represented as a pair of integer values.

The integer values are in QImode, HImode, SImode, DImode, TImode, OImode, and PSImode, respectively.

#### BND32mode BND64mode

These modes stand for bounds for pointer of 32 and 64 bit size respectively. Mode size is double pointer mode size.

The machine description defines Pmode as a C macro which expands into the machine mode used for addresses. Normally this is the mode whose size is BITS\_PER\_WORD, SImode on 32-bit machines.

The only modes which a machine description *must* support are QImode, and the modes corresponding to BITS\_PER\_WORD, FLOAT\_TYPE\_SIZE and DOUBLE\_TYPE\_SIZE. The compiler will attempt to use DImode for 8-byte structures and unions, but this can be prevented by overriding the definition of MAX\_FIXED\_MODE\_SIZE. Alternatively, you can have the compiler use TImode for 16-byte structures and unions. Likewise, you can arrange for the C type short int to avoid using HImode.

Very few explicit references to machine modes remain in the compiler and these few references will soon be removed. Instead, the machine modes are divided into mode classes. These are represented by the enumeration type enum mode\_class defined in 'machmode.h'. The possible mode classes are:

MODE\_INT Integer modes. By default these are BImode, QImode, HImode, SImode, DImode, TImode, and OImode.

# MODE\_PARTIAL\_INT

The "partial integer" modes, PQImode, PHImode, PSImode and PDImode.

#### MODE\_FLOAT

Floating point modes. By default these are QFmode, HFmode, TQFmode, SFmode, DFmode, XFmode and TFmode.

#### MODE\_DECIMAL\_FLOAT

Decimal floating point modes. By default these are SDmode, DDmode and TDmode.

#### MODE\_FRACT

Signed fractional modes. By default these are QQmode, HQmode, SQmode, DQmode and TQmode.

#### MODE\_UFRACT

Unsigned fractional modes. By default these are UQQmode, UHQmode, USQmode, UDQmode and UTQmode.

#### MODE\_ACCUM

Signed accumulator modes. By default these are HAmode, SAmode, DAmode and TAmode.

#### MODE\_UACCUM

Unsigned accumulator modes. By default these are UHAmode, USAmode, UDAmode and UTAmode.

#### MODE\_COMPLEX\_INT

Complex integer modes. (These are not currently implemented).

#### MODE\_COMPLEX\_FLOAT

Complex floating point modes. By default these are QCmode, HCmode, SCmode, DCmode, XCmode, and TCmode.

MODE\_CC Modes representing condition code values. These are CCmode plus any CC\_MODE modes listed in the 'machine-modes.def'. See Section 17.12 [Jump Patterns], page 433, also see Section 18.15 [Condition Code], page 572.

# MODE\_POINTER\_BOUNDS

Pointer bounds modes. Used to represent values of pointer bounds type. Operations in these modes may be executed as NOPs depending on hardware features and environment setup.

#### MODE\_RANDOM

This is a catchall mode class for modes which don't fit into the above classes. Currently VOIDmode and BLKmode are in MODE\_RANDOM.

machmode.h also defines various wrapper classes that combine a machine\_mode with a static assertion that a particular condition holds. The classes are:

#### scalar\_int\_mode

A mode that has class MODE\_INT or MODE\_PARTIAL\_INT.

# scalar\_float\_mode

A mode that has class MODE\_FLOAT or MODE\_DECIMAL\_FLOAT.

## scalar\_mode

A mode that holds a single numerical value. In practice this means that the mode is a scalar\_int\_mode, is a scalar\_float\_mode, or has class MODE\_FRACT, MODE\_UFRACT, MODE\_ACCUM, MODE\_UACCUM or MODE\_POINTER\_BOUNDS.

# complex\_mode

A mode that has class MODE\_COMPLEX\_INT or MODE\_COMPLEX\_FLOAT.

#### fixed\_size\_mode

A mode whose size is known at compile time.

Named modes use the most constrained of the available wrapper classes, if one exists, otherwise they use machine\_mode. For example, QImode is a scalar\_int\_mode, SFmode is a scalar\_float\_mode and BLKmode is a plain machine\_mode. It is possible to refer to any mode as a raw machine\_mode by adding the E\_ prefix, where E stands for "enumeration". For example, the raw machine\_mode names of the modes just mentioned are E\_QImode, E\_SFmode and E\_BLKmode respectively.

The wrapper classes implicitly convert to machine\_mode and to any wrapper class that represents a more general condition; for example scalar\_int\_mode and scalar\_float\_mode both convert to scalar\_mode and all three convert to fixed\_size\_mode. The classes act like machine\_modes that accept only certain named modes.

'machmode.h' also defines a template class opt\_mode<T> that holds a T or nothing, where T can be either machine\_mode or one of the wrapper classes above. The main operations on an opt\_mode<T> x are as follows:

'x.exists()'

Return true if x holds a mode rather than nothing.

'x.exists (&y)'

Return true if x holds a mode rather than nothing, storing the mode in y if so. y must be assignment-compatible with T.

'x.require ()'

Assert that x holds a mode rather than nothing and return that mode.

x' = y' Set x to y, where y is a T or implicitly converts to a T.

The default constructor sets an  $opt_mode < T >$  to nothing. There is also a constructor that takes an initial value of type T.

It is possible to use the 'is-a.h' accessors on a machine\_mode or machine mode wrapper x:

 $'is_a < T > (x)'$ 

Return true if x meets the conditions for wrapper class T.

 $is_a < T > (x, &y)$ 

Return true if x meets the conditions for wrapper class T, storing it in y if so. y must be assignment-compatible with T.

'as\_a  $\langle T \rangle$  (x)'

Assert that x meets the conditions for wrapper class T and return it as a T.

'dyn\_cast  $\langle T \rangle$  (x)'

Return an  $opt_mode < T >$  that holds x if x meets the conditions for wrapper class T and that holds nothing otherwise.

The purpose of these wrapper classes is to give stronger static type checking. For example, if a function takes a scalar\_int\_mode, a caller that has a general machine\_mode must either check or assert that the code is indeed a scalar integer first, using one of the functions above.

The wrapper classes are normal C++ classes, with user-defined constructors. Sometimes it is useful to have a POD version of the same type, particularly if the type appears in a union. The template class  $pod_mode < T > provides$  a POD version of wrapper class T. It is assignment-compatible with T and implicitly converts to both  $machine_mode$  and T.

Here are some C macros that relate to machine modes:

#### $GET_MODE(x)$

Returns the machine mode of the RTX x.

#### PUT MODE (x. newmode)

Alters the machine mode of the RTX x to be newmode.

#### NUM\_MACHINE\_MODES

Stands for the number of machine modes available on the target machine. This is one greater than the largest numeric value of any machine mode.

#### GET\_MODE\_NAME (m)

Returns the name of mode m as a string.

# GET\_MODE\_CLASS (m)

Returns the mode class of mode m.

# GET\_MODE\_WIDER\_MODE (m)

Returns the next wider natural mode. For example, the expression GET\_MODE\_WIDER\_MODE (QImode) returns HImode.

# GET\_MODE\_SIZE (m)

Returns the size in bytes of a datum of mode m.

# GET\_MODE\_BITSIZE (m)

Returns the size in bits of a datum of mode m.

#### GET\_MODE\_IBIT (m)

Returns the number of integral bits of a datum of fixed-point mode m.

# GET\_MODE\_FBIT (m)

Returns the number of fractional bits of a datum of fixed-point mode m.

#### GET\_MODE\_MASK (m)

Returns a bitmask containing 1 for all bits in a word that fit within mode m. This macro can only be used for modes whose bitsize is less than or equal to  ${\tt HOST\_BITS\_PER\_INT}$ .

#### GET\_MODE\_ALIGNMENT (m)

Return the required alignment, in bits, for an object of mode m.

#### GET\_MODE\_UNIT\_SIZE (m)

Returns the size in bytes of the subunits of a datum of mode m. This is the same as  $GET_MODE_SIZE$  except in the case of complex modes. For them, the unit size is the size of the real or imaginary part.

#### GET\_MODE\_NUNITS (m)

Returns the number of units contained in a mode, i.e., GET\_MODE\_SIZE divided by GET\_MODE\_UNIT\_SIZE.

# GET\_CLASS\_NARROWEST\_MODE (c)

Returns the narrowest mode in mode class c.

The following 3 variables are defined on every target. They can be used to allocate buffers that are guaranteed to be large enough to hold any value that can be represented on the

target. The first two can be overridden by defining them in the target's mode.def file, however, the value must be a constant that can determined very early in the compilation process. The third symbol cannot be overridden.

#### BITS\_PER\_UNIT

The number of bits in an addressable storage unit (byte). If you do not define this, the default is 8.

# MAX\_BITSIZE\_MODE\_ANY\_INT

The maximum bitsize of any mode that is used in integer math. This should be overridden by the target if it uses large integers as containers for larger vectors but otherwise never uses the contents to compute integer values.

#### MAX\_BITSIZE\_MODE\_ANY\_MODE

The bitsize of the largest mode on the target. The default value is the largest mode size given in the mode definition file, which is always correct for targets whose modes have a fixed size. Targets that might increase the size of a mode beyond this default should define MAX\_BITSIZE\_MODE\_ANY\_MODE to the actual upper limit in 'machine-modes.def'.

The global variables byte\_mode and word\_mode contain modes whose classes are MODE\_INT and whose bitsizes are either BITS\_PER\_UNIT or BITS\_PER\_WORD, respectively. On 32-bit machines, these are QImode and SImode, respectively.

# 14.7 Constant Expression Types

The simplest RTL expressions are those that represent constant values.

#### (const\_int i)

This type of expression represents the integer value i. i is customarily accessed with the macro INTVAL as in INTVAL (exp), which is equivalent to XWINT (exp, 0)

Constants generated for modes with fewer bits than in HOST\_WIDE\_INT must be sign extended to full width (e.g., with gen\_int\_mode). For constants for modes with more bits than in HOST\_WIDE\_INT the implied high order bits of that constant are copies of the top bit. Note however that values are neither inherently signed nor inherently unsigned; where necessary, signedness is determined by the rtl operation instead.

There is only one expression object for the integer value zero; it is the value of the variable const0\_rtx. Likewise, the only expression for integer value one is found in const1\_rtx, the only expression for integer value two is found in const2\_rtx, and the only expression for integer value negative one is found in constm1\_rtx. Any attempt to create an expression of code const\_int and value zero, one, two or negative one will return const0\_rtx, const1\_rtx, const2\_rtx or constm1\_rtx as appropriate.

Similarly, there is only one object for the integer whose value is STORE\_FLAG\_VALUE. It is found in const\_true\_rtx. If STORE\_FLAG\_VALUE is one, const\_true\_rtx and const1\_rtx will point to the same object. If STORE\_FLAG\_VALUE is -1, const\_true\_rtx and const1\_rtx will point to the same object.

#### (const\_double:m i0 i1 ...)

This represents either a floating-point constant of mode m or (on older ports that do not define TARGET\_SUPPORTS\_WIDE\_INT) an integer constant too large to fit into HOST\_BITS\_PER\_WIDE\_INT bits but small enough to fit within twice that number of bits. In the latter case, m will be VOIDmode. For integral values constants for modes with more bits than twice the number in HOST\_WIDE\_INT the implied high order bits of that constant are copies of the top bit of CONST\_DOUBLE\_HIGH. Note however that integral values are neither inherently signed nor inherently unsigned; where necessary, signedness is determined by the rtl operation instead.

On more modern ports, CONST\_DOUBLE only represents floating point values. New ports define TARGET\_SUPPORTS\_WIDE\_INT to make this designation.

If m is VOIDmode, the bits of the value are stored in i0 and i1. i0 is customarily accessed with the macro CONST\_DOUBLE\_LOW and i1 with CONST\_DOUBLE\_HIGH.

If the constant is floating point (regardless of its precision), then the number of integers used to store the value depends on the size of REAL\_VALUE\_TYPE (see Section 18.22 [Floating Point], page 631). The integers represent a floating point number, but not precisely in the target machine's or host machine's floating point format. To convert them to the precise bit pattern used by the target machine, use the macro REAL\_VALUE\_TO\_TARGET\_DOUBLE and friends (see Section 18.20.2 [Data Output], page 600).

#### (const\_wide\_int:m nunits elt0 ...)

This contains an array of HOST\_WIDE\_INTs that is large enough to hold any constant that can be represented on the target. This form of rtl is only used on targets that define TARGET\_SUPPORTS\_WIDE\_INT to be nonzero and then CONST\_DOUBLES are only used to hold floating-point values. If the target leaves TARGET\_SUPPORTS\_WIDE\_INT defined as 0, CONST\_WIDE\_INTs are not used and CONST\_DOUBLES are as they were before.

The values are stored in a compressed format. The higher-order 0s or -1s are not represented if they are just the logical sign extension of the number that is represented.

# CONST\_WIDE\_INT\_VEC (code)

Returns the entire array of HOST\_WIDE\_INTs that are used to store the value. This macro should be rarely used.

#### CONST\_WIDE\_INT\_NUNITS (code)

The number of <code>HOST\_WIDE\_INTs</code> used to represent the number. Note that this generally is smaller than the number of <code>HOST\_WIDE\_INTs</code> implied by the mode size.

# CONST\_WIDE\_INT\_ELT (code,i)

Returns the ith element of the array. Element 0 is contains the low order bits of the constant.

# (const\_fixed:m ...)

Represents a fixed-point constant of mode m. The operand is a data structure of type struct fixed\_value and is accessed with the macro CONST\_FIXED\_

VALUE. The high part of data is accessed with CONST\_FIXED\_VALUE\_HIGH; the low part is accessed with CONST\_FIXED\_VALUE\_LOW.

```
(const_poly_int:m [c0 c1 ...])
```

Represents a poly\_int-style polynomial integer with coefficients  $c0, c1, \ldots$  The coefficients are wide\_int-based integers rather than rtxes. CONST\_POLY\_INT\_COEFFS gives the values of individual coefficients (which is mostly only useful in low-level routines) and const\_poly\_int\_value gives the full poly\_int value.

```
(const_vector:m [x0 x1 ...])
```

Represents a vector constant. The values in square brackets are elements of the vector, which are always const\_int, const\_wide\_int, const\_double or const\_fixed expressions.

Each vector constant v is treated as a specific instance of an arbitrary-length sequence that itself contains 'CONST\_VECTOR\_NPATTERNS (v)' interleaved patterns. Each pattern has the form:

```
{ base0, base1, base1 + step, base1 + step * 2, ... }
```

The first three elements in each pattern are enough to determine the values of the other elements. However, if all *steps* are zero, only the first two elements are needed. If in addition each *base1* is equal to the corresponding *base0*, only the first element in each pattern is needed. The number of determining elements per pattern is given by 'CONST\_VECTOR\_NELTS\_PER\_PATTERN (v)'.

For example, the constant:

```
{ 0, 1, 2, 6, 3, 8, 4, 10, 5, 12, 6, 14, 7, 16, 8, 18 }
```

is interpreted as an interleaving of the sequences:

```
{ 0, 2, 3, 4, 5, 6, 7, 8 }
{ 1, 6, 8, 10, 12, 14, 16, 18 }
```

where the sequences are represented by the following patterns:

```
base0 == 0, base1 == 2, step == 1
base0 == 1, base1 == 6, step == 2
```

In this case:

```
CONST_VECTOR_NPATTERNS (v) == 2
CONST_VECTOR_NELTS_PER_PATTERN (v) == 3
```

Thus the first 6 elements ('{ 0, 1, 2, 6, 3, 8 }') are enough to determine the whole sequence; we refer to them as the "encoded" elements. They are the only elements present in the square brackets for variable-length const\_vectors (i.e. for const\_vectors whose mode m has a variable number of elements). However, as a convenience to code that needs to handle both const\_vectors and parallels, all elements are present in the square brackets for fixed-length const\_vectors; the encoding scheme simply reduces the amount of work involved in processing constants that follow a regular pattern.

Sometimes this scheme can create two possible encodings of the same vector. For example  $\{0, 1\}$  could be seen as two patterns with one element each or one pattern with two elements (base0 and base1). The canonical encoding is always the one with the fewest patterns or (if both encodings have the same number of petterns) the one with the fewest encoded elements.

'const\_vector\_encoding\_nelts (v)' gives the total number of encoded elements in v, which is 6 in the example above. CONST\_VECTOR\_ENCODED\_ELT (v, i) accesses the value of encoded element i.

'CONST\_VECTOR\_DUPLICATE\_P (v)' is true if v simply contains repeated instances of 'CONST\_VECTOR\_NPATTERNS (v)' values. This is a shorthand for testing 'CONST\_VECTOR\_NELTS\_PER\_PATTERN (v) == 1'.

'CONST\_VECTOR\_STEPPED\_P (v)' is true if at least one pattern in v has a nonzero step. This is a shorthand for testing 'CONST\_VECTOR\_NELTS\_PER\_PATTERN (v) == 3'.

CONST\_VECTOR\_NUNITS (v) gives the total number of elements in v; it is a shorthand for getting the number of units in 'GET\_MODE (v)'.

The utility function const\_vector\_elt gives the value of an arbitrary element as an rtx. const\_vector\_int\_elt gives the same value as a wide\_int.

#### (const\_string str)

Represents a constant string with value *str*. Currently this is used only for insn attributes (see Section 17.19 [Insn Attributes], page 450) since constant strings in C are placed in memory.

# (symbol\_ref:mode symbol)

Represents the value of an assembler label for data. *symbol* is a string that describes the name of the assembler label. If it starts with a '\*', the label is the rest of *symbol* not including the '\*'. Otherwise, the label is *symbol*, usually prefixed with '\_'.

The symbol\_ref contains a mode, which is usually Pmode. Usually that is the only mode for which a symbol is directly valid.

#### (label\_ref:mode label)

Represents the value of an assembler label for code. It contains one operand, an expression, which must be a code\_label or a note of type NOTE\_INSN\_DELETED\_LABEL that appears in the instruction sequence to identify the place where the label should go.

The reason for using a distinct expression type for code label references is so that jump optimization can distinguish them.

The label\_ref contains a mode, which is usually Pmode. Usually that is the only mode for which a label is directly valid.

#### (const:m exp)

Represents a constant that is the result of an assembly-time arithmetic computation. The operand, exp, contains only const\_int, symbol\_ref, label\_ref or unspec expressions, combined with plus and minus. Any such unspecs are target-specific and typically represent some form of relocation operator. m should be a valid address mode.

## (high:m exp)

Represents the high-order bits of exp. The number of bits is machine-dependent and is normally the number of bits specified in an instruction that initializes the high order bits of a register. It is used with lo\_sum to represent the typical two-instruction sequence used in RISC machines to reference large immediate values

and/or link-time constants such as global memory addresses. In the latter case, *m* is Pmode and *exp* is usually a constant expression involving symbol\_ref.

The macro CONSTO\_RTX (mode) refers to an expression with value 0 in mode mode. If mode mode is of mode class MODE\_INT, it returns constO\_rtx. If mode mode is of mode class MODE\_FLOAT, it returns a CONST\_DOUBLE expression in mode mode. Otherwise, it returns a CONST\_VECTOR expression in mode mode. Similarly, the macro CONST1\_RTX (mode) refers to an expression with value 1 in mode mode and similarly for CONST2\_RTX. The CONST1\_RTX and CONST2\_RTX macros are undefined for vector modes.

# 14.8 Registers and Memory

Here are the RTL expression types for describing access to machine registers and to main memory.

(reg:mn) For small values of the integer n (those that are less than FIRST\_PSEUDO\_REGISTER), this stands for a reference to machine register number n: a hard register. For larger values of n, it stands for a temporary value or pseudo register. The compiler's strategy is to generate code assuming an unlimited number of such pseudo registers, and later convert them into hard registers or into memory references.

m is the machine mode of the reference. It is necessary because machines can generally refer to each register in more than one mode. For example, a register may contain a full word but there may be instructions to refer to it as a half word or as a single byte, as well as instructions to refer to it as a floating point number of various precisions.

Even for a register that the machine can access in only one mode, the mode must always be specified.

The symbol FIRST\_PSEUDO\_REGISTER is defined by the machine description, since the number of hard registers on the machine is an invariant characteristic of the machine. Note, however, that not all of the machine registers must be general registers. All the machine registers that can be used for storage of data are given hard register numbers, even those that can be used only in certain instructions or can hold only certain types of data.

A hard register may be accessed in various modes throughout one function, but each pseudo register is given a natural mode and is accessed only in that mode. When it is necessary to describe an access to a pseudo register using a nonnatural mode, a subreg expression is used.

A reg expression with a machine mode that specifies more than one word of data may actually stand for several consecutive registers. If in addition the register number specifies a hardware register, then it actually represents several consecutive hardware registers starting with the specified one.

Each pseudo register number used in a function's RTL code is represented by a unique reg expression.

Some pseudo register numbers, those within the range of FIRST\_VIRTUAL\_REGISTER to LAST\_VIRTUAL\_REGISTER only appear during the RTL generation

phase and are eliminated before the optimization phases. These represent locations in the stack frame that cannot be determined until RTL generation for the function has been completed. The following virtual register numbers are defined:

# VIRTUAL\_INCOMING\_ARGS\_REGNUM

This points to the first word of the incoming arguments passed on the stack. Normally these arguments are placed there by the caller, but the callee may have pushed some arguments that were previously passed in registers.

When RTL generation is complete, this virtual register is replaced by the sum of the register given by ARG\_POINTER\_REGNUM and the value of FIRST\_PARM\_OFFSET.

#### VIRTUAL\_STACK\_VARS\_REGNUM

If FRAME\_GROWS\_DOWNWARD is defined to a nonzero value, this points to immediately above the first variable on the stack. Otherwise, it points to the first variable on the stack.

VIRTUAL\_STACK\_VARS\_REGNUM is replaced with the sum of the register given by FRAME\_POINTER\_REGNUM and the value TARGET\_STARTING\_FRAME\_OFFSET.

#### VIRTUAL\_STACK\_DYNAMIC\_REGNUM

This points to the location of dynamically allocated memory on the stack immediately after the stack pointer has been adjusted by the amount of memory desired.

This virtual register is replaced by the sum of the register given by STACK\_POINTER\_REGNUM and the value STACK\_DYNAMIC\_OFFSET.

# VIRTUAL\_OUTGOING\_ARGS\_REGNUM

This points to the location in the stack at which outgoing arguments should be written when the stack is pre-pushed (arguments pushed using push insns should always use STACK\_POINTER\_REGNUM).

This virtual register is replaced by the sum of the register given by STACK\_POINTER\_REGNUM and the value STACK\_POINTER\_OFFSET.

#### (subreg:m1 reg:m2 bytenum)

subreg expressions are used to refer to a register in a machine mode other than its natural one, or to refer to one register of a multi-part reg that actually refers to several registers.

Each pseudo register has a natural mode. If it is necessary to operate on it in a different mode, the register must be enclosed in a subreg.

There are currently three supported types for the first operand of a subreg:

- pseudo registers This is the most common case. Most subregs have pseudo regs as their first operand.
- mem subregs of mem were common in earlier versions of GCC and are still supported. During the reload pass these are replaced by plain mems. On machines that do not do instruction scheduling, use of subregs of mem are

still used, but this is no longer recommended. Such subregs are considered to be register\_operands rather than memory\_operands before and during reload. Because of this, the scheduling passes cannot properly schedule instructions with subregs of mem, so for machines that do scheduling, subregs of mem should never be used. To support this, the combine and recog passes have explicit code to inhibit the creation of subregs of mem when INSN\_SCHEDULING is defined.

The use of subregs of mem after the reload pass is an area that is not well understood and should be avoided. There is still some code in the compiler to support this, but this code has possibly rotted. This use of subregs is discouraged and will most likely not be supported in the future.

• hard registers It is seldom necessary to wrap hard registers in subregs; such registers would normally reduce to a single reg rtx. This use of subregs is discouraged and may not be supported in the future.

subregs of subregs are not supported. Using simplify\_gen\_subreg is the recommended way to avoid this problem.

subregs come in two distinct flavors, each having its own usage and rules:

# Paradoxical subregs

When m1 is strictly wider than m2, the subreg expression is called paradoxical. The canonical test for this class of subreg is:

```
paradoxical_subreg_p (m1, m2)
```

Paradoxical subregs can be used as both lvalues and rvalues. When used as an lvalue, the low-order bits of the source value are stored in reg and the high-order bits are discarded. When used as an rvalue, the low-order bits of the subreg are taken from reg while the high-order bits may or may not be defined.

The high-order bits of rvalues are defined in the following circumstances:

- subregs of mem When m2 is smaller than a word, the macro LOAD\_EXTEND\_OP, can control how the high-order bits are defined.
- subreg of regs The upper bits are defined when SUBREG\_PROMOTED\_VAR\_P is true. SUBREG\_PROMOTED\_UNSIGNED\_P describes what the upper bits hold. Such subregs usually represent local variables, register variables and parameter pseudo variables that have been promoted to a wider mode.

bytenum is always zero for a paradoxical subreg, even on bigendian targets.

For example, the paradoxical subreg:

```
(set (subreg:SI (reg:HI x) 0) y)
```

stores the lower 2 bytes of y in x and discards the upper 2 bytes. A subsequent:

```
(set z (subreg:SI (reg:HI x) 0))
```

would set the lower two bytes of z to y and set the upper two bytes to an unknown value assuming SUBREG\_PROMOTED\_VAR\_P is false.

# Normal subregs

When m1 is at least as narrow as m2 the subreg expression is called normal.

Normal subregs restrict consideration to certain bits of reg. For this purpose, reg is divided into individually-addressable blocks in which each block has:

```
REGMODE_NATURAL_SIZE (m2)
```

bytes. Usually the value is UNITS\_PER\_WORD; that is, most targets usually treat each word of a register as being independently addressable.

There are two types of normal subreg. If m1 is known to be no bigger than a block, the subreg refers to the least-significant part (or *lowpart*) of one block of reg. If m1 is known to be larger than a block, the subreg refers to two or more complete blocks.

When used as an lvalue, subreg is a block-based accessor. Storing to a subreg modifies all the blocks of reg that overlap the subreg, but it leaves the other blocks of reg alone.

When storing to a normal subreg that is smaller than a block, the other bits of the referenced block are usually left in an undefined state. This laxity makes it easier to generate efficient code for such instructions. To represent an instruction that preserves all the bits outside of those in the subreg, use strict\_low\_part or zero\_extract around the subreg.

bytenum must identify the offset of the first byte of the subreg from the start of reg, assuming that reg is laid out in memory order. The memory order of bytes is defined by two target macros, WORDS\_BIG\_ENDIAN and BYTES\_BIG\_ENDIAN:

- WORDS\_BIG\_ENDIAN, if set to 1, says that byte number zero is part of the most significant word; otherwise, it is part of the least significant word.
- BYTES\_BIG\_ENDIAN, if set to 1, says that byte number zero is the most significant byte within a word; otherwise, it is the least significant byte within a word.

On a few targets, FLOAT\_WORDS\_BIG\_ENDIAN disagrees with WORDS\_BIG\_ENDIAN. However, most parts of the compiler treat floating point values as if they had the same endianness as integer values. This works because they handle them solely as a collection of integer values, with no particular numerical value. Only real.c and the runtime libraries care about FLOAT\_WORDS\_BIG\_ENDIAN.

Thus,

(subreg:HI (reg:SI x) 2)

on a BYTES\_BIG\_ENDIAN, 'UNITS\_PER\_WORD == 4' target is the same as

```
(subreg:HI (reg:SI x) 0)
```

on a little-endian, 'UNITS\_PER\_WORD == 4' target. Both subregs access the lower two bytes of register x.

Note that the byte offset is a polynomial integer; it may not be a compile-time constant on targets with variable-sized modes. However, the restrictions above mean that there are only a certain set of acceptable offsets for a given combination of m1 and m2. The compiler can always tell which blocks a valid subreg occupies, and whether the subreg is a lowpart of a block.

A MODE\_PARTIAL\_INT mode behaves as if it were as wide as the corresponding MODE\_INT mode, except that it has an unknown number of undefined bits. For example:

```
(subreg:PSI (reg:SI 0) 0)
```

accesses the whole of '(reg:SI 0)', but the exact relationship between the PSImode value and the SImode value is not defined. If we assume 'REGMODE\_NATURAL\_SIZE (DImode) <= 4', then the following two subregs:

```
(subreg:PSI (reg:DI 0) 0) (subreg:PSI (reg:DI 0) 4)
```

represent independent 4-byte accesses to the two halves of '(reg:DI 0)'. Both subregs have an unknown number of undefined bits.

If 'REGMODE\_NATURAL\_SIZE (PSImode) <= 2' then these two subregs:

```
(subreg:HI (reg:PSI 0) 0)
(subreg:HI (reg:PSI 0) 2)
```

represent independent 2-byte accesses that together span the whole of '(reg:PSI 0)'. Storing to the first subreg does not affect the value of the second, and vice versa. '(reg:PSI 0)' has an unknown number of undefined bits, so the assignment:

```
(set (subreg:HI (reg:PSI 0) 0) (reg:HI 4))
```

does not guarantee that '(subreg:HI (reg:PSI 0) 0)' has the value '(reg:HI 4)'.

The rules above apply to both pseudo regs and hard regs. If the semantics are not correct for particular combinations of m1, m2 and hard reg, the target-specific code must ensure that those combinations are never used. For example:

```
TARGET_CAN_CHANGE_MODE_CLASS (m2, m1, class)
```

must be false for every class class that includes reg.

GCC must be able to determine at compile time whether a subreg is paradoxical, whether it occupies a whole number of blocks, or whether it is a lowpart of a block. This means that certain combinations of variable-sized mode are not permitted. For example, if m2 holds n SI values, where n is greater than zero, it is not possible to form a DI subreg of it; such a subreg would be paradoxical when n is 1 but not when n is greater than 1.

The first operand of a subreg expression is customarily accessed with the SUBREG\_REG macro and the second operand is customarily accessed with the SUBREG\_BYTE macro.

It has been several years since a platform in which BYTES\_BIG\_ENDIAN not equal to WORDS\_BIG\_ENDIAN has been tested. Anyone wishing to support such a platform in the future may be confronted with code rot.

#### (scratch:m)

This represents a scratch register that will be required for the execution of a single instruction and not used subsequently. It is converted into a reg by either the local register allocator or the reload pass.

scratch is usually present inside a clobber operation (see Section 14.15 [Side Effects], page 297).

- (cc0) This refers to the machine's condition code register. It has no operands and may not have a machine mode. There are two ways to use it:
  - To stand for a complete set of condition code flags. This is best on most machines, where each comparison sets the entire series of flags.
    - With this technique, (cc0) may be validly used in only two contexts: as the destination of an assignment (in test and compare instructions) and in comparison operators comparing against zero (const\_int with value zero; that is to say, const0\_rtx).
  - To stand for a single flag that is the result of a single condition. This is useful on machines that have only a single flag bit, and in which comparison instructions must specify the condition to test.
    - With this technique, (cc0) may be validly used in only two contexts: as the destination of an assignment (in test and compare instructions) where the source is a comparison operator, and as the first operand of if\_then\_else (in a conditional branch).

There is only one expression object of code cc0; it is the value of the variable cc0\_rtx. Any attempt to create an expression of code cc0 will return cc0\_rtx. Instructions can set the condition code implicitly. On many machines, nearly all instructions set the condition code based on the value that they compute or store. It is not necessary to record these actions explicitly in the RTL because the machine description includes a prescription for recognizing the instructions that do so (by means of the macro NOTICE\_UPDATE\_CC). See Section 18.15 [Condition Code], page 572. Only instructions whose sole purpose is to set the condition code, and instructions that use the condition code, need mention (cc0).

On some machines, the condition code register is given a register number and a reg is used instead of (cc0). This is usually the preferable approach if only a small subset of instructions modify the condition code. Other machines store condition codes in general registers; in such cases a pseudo register should be used.

Some machines, such as the SPARC and RS/6000, have two sets of arithmetic instructions, one that sets and one that does not set the condition code. This

is best handled by normally generating the instruction that does not set the condition code, and making a pattern that both performs the arithmetic and sets the condition code register (which would not be (cc0) in this case). For examples, search for 'addcc' and 'andcc' in 'sparc.md'.

(pc) This represents the machine's program counter. It has no operands and may not have a machine mode. (pc) may be validly used only in certain specific contexts in jump instructions.

There is only one expression object of code pc; it is the value of the variable pc\_rtx. Any attempt to create an expression of code pc will return pc\_rtx.

All instructions that do not jump alter the program counter implicitly by incrementing it, but there is no need to mention this in the RTL.

#### (mem:m addr alias)

This RTX represents a reference to main memory at an address represented by the expression addr. m specifies how large a unit of memory is accessed. alias specifies an alias set for the reference. In general two items are in different alias sets if they cannot reference the same memory address.

The construct (mem:BLK (scratch)) is considered to alias all other memories. Thus it may be used as a memory barrier in epilogue stack deallocation patterns.

#### (concatm rtx rtx)

This RTX represents the concatenation of two other RTXs. This is used for complex values. It should only appear in the RTL attached to declarations and during RTL generation. It should not appear in the ordinary inso chain.

#### (concatnm [rtx ...])

This RTX represents the concatenation of all the rtx to make a single value. Like concat, this should only appear in declarations, and not in the insn chain.

# 14.9 RTL Expressions for Arithmetic

Unless otherwise specified, all the operands of arithmetic expressions must be valid for mode m. An operand is valid for mode m if it has mode m, or if it is a const\_int or const\_double and m is a mode of class MODE\_INT.

For commutative binary operations, constants should be placed in the second operand.

```
(plus:m x y)
(ss_plus:m x y)
(us_plus:m x y)
```

These three expressions all represent the sum of the values represented by x and y carried out in machine mode m. They differ in their behavior on overflow of integer modes. plus wraps round modulo the width of m; ss\_plus saturates at the maximum signed value representable in m; us\_plus saturates at the maximum unsigned value.

# $(lo_sum:m \times y)$

This expression represents the sum of x and the low-order bits of y. It is used with high (see Section 14.7 [Constants], page 278) to represent the typical two-instruction sequence used in RISC machines to reference large immediate values

and/or link-time constants such as global memory addresses. In the latter case, m is Pmode and y is usually a constant expression involving symbol\_ref.

The number of low order bits is machine-dependent but is normally the number of bits in mode m minus the number of bits set by high.

```
(minus:m x y)
(ss_minus:m x y)
(us_minus:m x y)
```

These three expressions represent the result of subtracting y from x, carried out in mode M. Behavior on overflow is the same as for the three variants of plus (see above).

```
(compare: m x y)
```

Represents the result of subtracting y from x for purposes of comparison. The result is computed without overflow, as if with infinite precision.

Of course, machines cannot really subtract with infinite precision. However, they can pretend to do so when only the sign of the result will be used, which is the case when the result is stored in the condition code. And that is the *only* way this kind of expression may validly be used: as a value to be stored in the condition codes, either (cc0) or a register. See Section 14.10 [Comparisons], page 292.

The mode m is not related to the modes of x and y, but instead is the mode of the condition code value. If (cc0) is used, it is VOIDmode. Otherwise it is some mode in class MODE\_CC, often CCmode. See Section 18.15 [Condition Code], page 572. If m is VOIDmode or CCmode, the operation returns sufficient information (in an unspecified format) so that any comparison operator can be applied to the result of the COMPARE operation. For other modes in class MODE\_CC, the operation only returns a subset of this information.

Normally, x and y must have the same mode. Otherwise, compare is valid only if the mode of x is in class MODE\_INT and y is a const\_int or const\_double with mode VOIDmode. The mode of x determines what mode the comparison is to be done in; thus it must not be VOIDmode.

If one of the operands is a constant, it should be placed in the second operand and the comparison code adjusted as appropriate.

A compare specifying two VOIDmode constants is not valid since there is no way to know in what mode the comparison is to be performed; the comparison must either be folded during the compilation or the first operand must be loaded into a register while its mode is still known.

```
(neg:m x)
(ss_neg:m x)
(us_neg:m x)
```

These two expressions represent the negation (subtraction from zero) of the value represented by x, carried out in mode m. They differ in the behavior on overflow of integer modes. In the case of neg, the negation of the operand may be a number not representable in mode m, in which case it is truncated to m.  $ss_neg$  and  $us_neg$  ensure that an out-of-bounds result saturates to the maximum or minimum signed or unsigned value.

```
(mult:m x y)
(ss_mult:m x y)
(us_mult:m x y)
```

Represents the signed product of the values represented by x and y carried out in machine mode m.  $ss_mult$  and  $us_mult$  ensure that an out-of-bounds result saturates to the maximum or minimum signed or unsigned value.

Some machines support a multiplication that generates a product wider than the operands. Write the pattern for this as

```
(mult:m (sign_extend:m x) (sign_extend:m y))
```

where m is wider than the modes of x and y, which need not be the same.

For unsigned widening multiplication, use the same idiom, but with zero\_extend instead of sign\_extend.

```
(fma: m x y z)
```

Represents the fma, fmaf, and fmal builtin functions, which compute 'x \* y + z' without doing an intermediate rounding step.

```
(div:m x y)
(ss_div:m x y)
```

Represents the quotient in signed division of x by y, carried out in machine mode m. If m is a floating point mode, it represents the exact quotient; otherwise, the integerized quotient.  $ss\_div$  ensures that an out-of-bounds result saturates to the maximum or minimum signed value.

Some machines have division instructions in which the operands and quotient widths are not all the same; you should represent such instructions using truncate and sign\_extend as in,

```
(truncate:m1 (div:m2 x (sign_extend:m2 y)))
```

```
(udiv:m x y)
(us_div:m x y)
```

Like div but represents unsigned division. us\_div ensures that an out-of-bounds result saturates to the maximum or minimum unsigned value.

```
(mod:m x y)
(umod:m x y)
```

Like div and udiv but represent the remainder instead of the quotient.

```
(smin:m x y)
(smax:m x y)
```

Represents the smaller (for smin) or larger (for smax) of x and y, interpreted as signed values in mode m. When used with floating point, if both operands are zeros, or if either operand is NaN, then it is unspecified which of the two operands is returned as the result.

```
(umin:m x y)
(umax:m x y)
```

Like smin and smax, but the values are interpreted as unsigned integers.

 $(not:m\ x)$  Represents the bitwise complement of the value represented by x, carried out in mode m, which must be a fixed-point machine mode.

 $(and: m \times y)$ 

Represents the bitwise logical-and of the values represented by x and y, carried out in machine mode m, which must be a fixed-point machine mode.

 $(ior: m \times y)$ 

Represents the bitwise inclusive-or of the values represented by x and y, carried out in machine mode m, which must be a fixed-point mode.

 $(xor: m \times y)$ 

Represents the bitwise exclusive-or of the values represented by x and y, carried out in machine mode m, which must be a fixed-point mode.

(ashift:m x c)
(ss\_ashift:m x c)
(us\_ashift:m x c)

These three expressions represent the result of arithmetically shifting x left by c places. They differ in their behavior on overflow of integer modes. An ashift operation is a plain shift with no special behavior in case of a change in the sign bit; ss\_ashift and us\_ashift saturates to the minimum or maximum representable value if any of the bits shifted out differs from the final sign bit. x have mode m, a fixed-point machine mode. c be a fixed-point mode or be a

constant with mode VOIDmode; which mode is determined by the mode called for in the machine description entry for the left-shift instruction. For example, on the VAX, the mode of c is QImode regardless of m.

(lshiftrt:m x c)
(ashiftrt:m x c)

Like ashift but for right shift. Unlike the case for left shift, these two operations are distinct.

 $(rotate:m \times c)$  $(rotatert:m \times c)$ 

Similar but represent left and right rotate. If c is a constant, use rotate.

(abs:m x)

 $(ss_abs:m x)$ 

Represents the absolute value of x, computed in mode m.  $ss_abs$  ensures that an out-of-bounds result saturates to the maximum signed value.

(sqrt:m x)

Represents the square root of x, computed in mode m. Most often m will be a floating point mode.

(ffs:m x) Represents one plus the index of the least significant 1-bit in x, represented as an integer of mode m. (The value is zero if x is zero.) The mode of x must be m or VOIDmode.

(clrsb:mx)

Represents the number of redundant leading sign bits in x, represented as an integer of mode m, starting at the most significant bit position. This is one less than the number of leading sign bits (either 0 or 1), with no special cases. The mode of x must be m or VOIDmode.

- (clz:mx) Represents the number of leading 0-bits in x, represented as an integer of mode m, starting at the most significant bit position. If x is zero, the value is determined by CLZ\_DEFINED\_VALUE\_AT\_ZERO (see Section 18.31 [Misc], page 642). Note that this is one of the few expressions that is not invariant under widening. The mode of x must be m or VOIDmode.
- (ctz:mx) Represents the number of trailing 0-bits in x, represented as an integer of mode m, starting at the least significant bit position. If x is zero, the value is determined by CTZ\_DEFINED\_VALUE\_AT\_ZERO (see Section 18.31 [Misc], page 642). Except for this case, ctz(x) is equivalent to ffs(x) 1. The mode of x must be m or VOIDmode.

# (popcount: m x)

Represents the number of 1-bits in x, represented as an integer of mode m. The mode of x must be m or VOIDmode.

#### (parity:m x)

Represents the number of 1-bits modulo 2 in x, represented as an integer of mode m. The mode of x must be m or VOIDmode.

# (bswap:m x)

Represents the value x with the order of bytes reversed, carried out in mode m, which must be a fixed-point machine mode. The mode of x must be m or VOIDmode.

# 14.10 Comparison Operations

Comparison operators test a relation on two operands and are considered to represent a machine-dependent nonzero value described by, but not necessarily equal to, STORE\_FLAG\_VALUE (see Section 18.31 [Misc], page 642) if the relation holds, or zero if it does not, for comparison operators whose results have a 'MODE\_INT' mode, FLOAT\_STORE\_FLAG\_VALUE (see Section 18.31 [Misc], page 642) if the relation holds, or zero if it does not, for comparison operators that return floating-point values, and a vector of either VECTOR\_STORE\_FLAG\_VALUE (see Section 18.31 [Misc], page 642) if the relation holds, or of zeros if it does not, for comparison operators that return vector results. The mode of the comparison operation is independent of the mode of the data being compared. If the comparison operation is being tested (e.g., the first operand of an if\_then\_else), the mode must be VOIDmode.

There are two ways that comparison operations may be used. The comparison operators may be used to compare the condition codes (cc0) against zero, as in (eq (cc0) (const\_int 0)). Such a construct actually refers to the result of the preceding instruction in which the condition codes were set. The instruction setting the condition code must be adjacent to the instruction using the condition code; only note insns may separate them.

Alternatively, a comparison operation may directly compare two data objects. The mode of the comparison is determined by the operands; they must both be valid for a common machine mode. A comparison with both operands constant would be invalid as the machine mode could not be deduced from it, but such a comparison should never exist in RTL due to constant folding.

In the example above, if (cc0) were last set to (compare x y), the comparison operation is identical to (eq x y). Usually only one style of comparisons is supported on a particular

machine, but the combine pass will try to merge the operations to produce the eq shown in case it exists in the context of the particular insn involved.

Inequality comparisons come in two flavors, signed and unsigned. Thus, there are distinct expression codes gt and gtu for signed and unsigned greater-than. These can produce different results for the same pair of integer values: for example, 1 is signed greater-than -1 but not unsigned greater-than, because -1 when regarded as unsigned is actually 0xffffffffwhich is greater than 1.

The signed comparisons are also used for floating point values. Floating point comparisons are distinguished by the machine modes of the operands.

```
(eq:m \times y)
            STORE_FLAG_VALUE if the values represented by x and y are equal, otherwise 0.
(ne:m \times y)
            STORE_FLAG_VALUE if the values represented by x and y are not equal, otherwise
            0.
(gt:m \times y)
            STORE_FLAG_VALUE if the x is greater than y. If they are fixed-point, the com-
            parison is done in a signed sense.
(gtu: m \times y)
            Like gt but does unsigned comparison, on fixed-point numbers only.
(lt:m x y)
(ltu: m \times y)
            Like gt and gtu but test for "less than".
(ge:m \times y)
(geu: m \times y)
```

Like gt and gtu but test for "greater than or equal".

 $(le:m \times y)$  $(leu: m \times y)$ 

Like gt and gtu but test for "less than or equal".

#### (if\_then\_else cond then else)

This is not a comparison operation but is listed here because it is always used in conjunction with a comparison operation. To be precise, cond is a comparison expression. This expression represents a choice, according to cond, between the value represented by then and the one represented by else.

On most machines, if\_then\_else expressions are valid only to express conditional jumps.

# (cond [test1 value1 test2 value2 ...] default)

Similar to if\_then\_else, but more general. Each of test1, test2, ... is performed in turn. The result of this expression is the value corresponding to the first nonzero test, or default if none of the tests are nonzero expressions.

This is currently not valid for instruction patterns and is supported only for insn attributes. See Section 17.19 [Insn Attributes], page 450.

# 14.11 Bit-Fields

Special expression codes exist to represent bit-field instructions.

# (sign\_extract:m loc size pos)

This represents a reference to a sign-extended bit-field contained or starting in loc (a memory or register reference). The bit-field is size bits wide and starts at bit pos. The compilation option BITS\_BIG\_ENDIAN says which end of the memory unit pos counts from.

If *loc* is in memory, its mode must be a single-byte integer mode. If *loc* is in a register, the mode to use is specified by the operand of the insv or extv pattern (see Section 17.9 [Standard Names], page 392) and is usually a full-word integer mode, which is the default if none is specified.

The mode of pos is machine-specific and is also specified in the insv or extv pattern.

The mode m is the same as the mode that would be used for loc if it were a register.

A sign\_extract cannot appear as an lvalue, or part thereof, in RTL.

# (zero\_extract:m loc size pos)

Like sign\_extract but refers to an unsigned or zero-extended bit-field. The same sequence of bits are extracted, but they are filled to an entire word with zeros instead of by sign-extension.

Unlike sign\_extract, this type of expressions can be lvalues in RTL; they may appear on the left side of an assignment, indicating insertion of a value into the specified bit-field.

# 14.12 Vector Operations

All normal RTL expressions can be used with vector modes; they are interpreted as operating on each part of the vector independently. Additionally, there are a few new expressions to describe specific vector operations.

# (vec\_merge:m vec1 vec2 items)

This describes a merge operation between two vectors. The result is a vector of mode m; its elements are selected from either vec1 or vec2. Which elements are selected is described by items, which is a bit mask represented by a const\_int; a zero bit indicates the corresponding element in the result vector is taken from vec2 while a set bit indicates it is taken from vec1.

#### (vec\_select:m vec1 selection)

This describes an operation that selects parts of a vector. vec1 is the source vector, and selection is a parallel that contains a  $const_int$  (or another expression, if the selection can be made at runtime) for each of the subparts of the result vector, giving the number of the source subpart that should be stored into it. The result mode m is either the submode for a single element of vec1 (if only one subpart is selected), or another vector mode with that element submode (if multiple subparts are selected).

# (vec\_concat:m x1 x2)

Describes a vector concat operation. The result is a concatenation of the vectors or scalars x1 and x2; its length is the sum of the lengths of the two inputs.

#### (vec\_duplicate:m x)

This operation converts a scalar into a vector or a small vector into a larger one by duplicating the input values. The output vector mode must have the same submodes as the input vector mode or the scalar modes, and the number of output parts must be an integer multiple of the number of input parts.

#### (vec\_series:m base step)

This operation creates a vector in which element i is equal to 'base + i\*step'. m must be a vector integer mode.

# 14.13 Conversions

All conversions between machine modes must be represented by explicit conversion operations. For example, an expression which is the sum of a byte and a full word cannot be written as (plus:SI (reg:QI 34) (reg:SI 80)) because the plus operation requires two operands of the same machine mode. Therefore, the byte-sized operand is enclosed in a conversion operation, as in

```
(plus:SI (sign_extend:SI (reg:QI 34)) (reg:SI 80))
```

The conversion operation is not a mere placeholder, because there may be more than one way of converting from a given starting mode to the desired final mode. The conversion operation code says how to do it.

For all conversion operations, x must not be VOIDmode because the mode in which to do the conversion would not be known. The conversion must either be done at compile-time or x must be placed into a register.

# (sign\_extend:m x)

Represents the result of sign-extending the value x to machine mode m. m must be a fixed-point mode and x a fixed-point value of a mode narrower than m.

# (zero\_extend: m x)

Represents the result of zero-extending the value x to machine mode m. m must be a fixed-point mode and x a fixed-point value of a mode narrower than m.

# (float\_extend:m x)

Represents the result of extending the value x to machine mode m. m must be a floating point mode and x a floating point value of a mode narrower than m.

#### (truncate: m x)

Represents the result of truncating the value x to machine mode m. m must be a fixed-point mode and x a fixed-point value of a mode wider than m.

# (ss\_truncate:m x)

Represents the result of truncating the value x to machine mode m, using signed saturation in the case of overflow. Both m and the mode of x must be fixed-point modes.

# (us\_truncate:m x)

Represents the result of truncating the value x to machine mode m, using unsigned saturation in the case of overflow. Both m and the mode of x must be fixed-point modes.

# (float\_truncate:m x)

Represents the result of truncating the value x to machine mode m. m must be a floating point mode and x a floating point value of a mode wider than m.

## (float:m x)

Represents the result of converting fixed point value x, regarded as signed, to floating point mode m.

#### (unsigned\_float:m x)

Represents the result of converting fixed point value x, regarded as unsigned, to floating point mode m.

(fix:mx) When m is a floating-point mode, represents the result of converting floating point value x (valid for mode m) to an integer, still represented in floating point mode m, by rounding towards zero.

When m is a fixed-point mode, represents the result of converting floating point value x to mode m, regarded as signed. How rounding is done is not specified, so this operation may be used validly in compiling C code only for integer-valued operands.

# (unsigned\_fix:m x)

Represents the result of converting floating point value x to fixed point mode m, regarded as unsigned. How rounding is done is not specified.

# (fract\_convert:m x)

Represents the result of converting fixed-point value x to fixed-point mode m, signed integer value x to fixed-point mode m, floating-point value x to fixed-point mode m, fixed-point value x to integer mode m regarded as signed, or fixed-point value x to floating-point mode m. When overflows or underflows happen, the results are undefined.

# (sat\_fract:m x)

Represents the result of converting fixed-point value x to fixed-point mode m, signed integer value x to fixed-point mode m, or floating-point value x to fixed-point mode m. When overflows or underflows happen, the results are saturated to the maximum or the minimum.

# (unsigned\_fract\_convert:m x)

Represents the result of converting fixed-point value x to integer mode m regarded as unsigned, or unsigned integer value x to fixed-point mode m. When overflows or underflows happen, the results are undefined.

# (unsigned\_sat\_fract:m x)

Represents the result of converting unsigned integer value x to fixed-point mode m. When overflows or underflows happen, the results are saturated to the maximum or the minimum.

# 14.14 Declarations

Declaration expression codes do not represent arithmetic operations but rather state assertions about their operands.

# (strict\_low\_part (subreg:m (reg:n r) 0))

This expression code is used in only one context: as the destination operand of a **set** expression. In addition, the operand of this expression must be a non-paradoxical **subreg** expression.

The presence of  $strict_low_part$  says that the part of the register which is meaningful in mode n, but is not part of mode m, is not to be altered. Normally, an assignment to such a subreg is allowed to have undefined effects on the rest of the register when m is smaller than 'REGMODE\_NATURAL\_SIZE (n)'.

# 14.15 Side Effect Expressions

The expression codes described so far represent values, not actions. But machine instructions never produce values; they are meaningful only for their side effects on the state of the machine. Special expression codes are used to represent side effects.

The body of an instruction is always one of these side effect codes; the codes described above, which represent values, appear only as the operands of these.

#### (set lval x)

Represents the action of storing the value of x into the place represented by lval. lval must be an expression representing a place that can be stored in: reg (or subreg, strict\_low\_part or zero\_extract), mem, pc, parallel, or cc0.

If *lval* is a reg, subreg or mem, it has a machine mode; then x must be valid for that mode.

If *lval* is a **reg** whose machine mode is less than the full width of the register, then it means that the part of the register specified by the machine mode is given the specified value and the rest of the register receives an undefined value. Likewise, if *lval* is a **subreg** whose machine mode is narrower than the mode of the register, the rest of the register can be changed in an undefined way.

If *lval* is a **strict\_low\_part** of a subreg, then the part of the register specified by the machine mode of the **subreg** is given the value x and the rest of the register is not changed.

If *lval* is a zero\_extract, then the referenced part of the bit-field (a memory or register reference) specified by the zero\_extract is given the value x and the rest of the bit-field is not changed. Note that sign\_extract cannot appear in *lval*.

If lval is (cc0), it has no machine mode, and x may be either a compare expression or a value that may have any mode. The latter case represents a "test" instruction. The expression (set (cc0) (reg:m n)) is equivalent to (set (cc0) (compare (reg:m n) (const\_int 0))). Use the former expression to save space during the compilation.

If *lval* is a parallel, it is used to represent the case of a function returning a structure in multiple registers. Each element of the parallel is an expr\_list

whose first operand is a **reg** and whose second operand is a **const\_int** representing the offset (in bytes) into the structure at which the data in that register corresponds. The first element may be null to indicate that the structure is also passed partly in memory.

If *lval* is (pc), we have a jump instruction, and the possibilities for x are very limited. It may be a label\_ref expression (unconditional jump). It may be an if\_then\_else (conditional jump), in which case either the second or the third operand must be (pc) (for the case which does not jump) and the other of the two must be a label\_ref (for the case which does jump). x may also be a mem or (plus:SI (pc) y), where y may be a reg or a mem; these unusual patterns are used to represent jumps through branch tables.

If lval is neither (cc0) nor (pc), the mode of lval must not be VOIDmode and the mode of x must be valid for the mode of lval.

lval is customarily accessed with the SET\_DEST macro and x with the SET\_SRC macro.

#### (return)

As the sole expression in a pattern, represents a return from the current function, on machines where this can be done with one instruction, such as VAXen. On machines where a multi-instruction "epilogue" must be executed in order to return from the function, returning is done by jumping to a label which precedes the epilogue, and the return expression code is never used.

Inside an if\_then\_else expression, represents the value to be placed in pc to return to the caller.

Note that an insn pattern of (return) is logically equivalent to (set (pc) (return)), but the latter form is never used.

#### (simple\_return)

Like (return), but truly represents only a function return, while (return) may represent an insn that also performs other functions of the function epilogue. Like (return), this may also occur in conditional jumps.

# (call function nargs)

Represents a function call. *function* is a mem expression whose address is the address of the function to be called. *nargs* is an expression which can be used for two purposes: on some machines it represents the number of bytes of stack argument; on others, it represents the number of argument registers.

Each machine has a standard machine mode which function must have. The machine description defines macro FUNCTION\_MODE to expand into the requisite mode name. The purpose of this mode is to specify what kind of addressing is allowed, on machines where the allowed kinds of addressing depend on the machine mode being addressed.

# (clobber x)

Represents the storing or possible storing of an unpredictable, undescribed value into x, which must be a reg, scratch, parallel or mem expression.

One place this is used is in string instructions that store standard values into particular hard registers. It may not be worth the trouble to describe the values

that are stored, but it is essential to inform the compiler that the registers will be altered, lest it attempt to keep data in them across the string instruction.

If x is (mem:BLK (const\_int 0)) or (mem:BLK (scratch)), it means that all memory locations must be presumed clobbered. If x is a parallel, it has the same meaning as a parallel in a set expression.

Note that the machine description classifies certain hard registers as "call-clobbered". All function call instructions are assumed by default to clobber these registers, so there is no need to use clobber expressions to indicate this fact. Also, each function call is assumed to have the potential to alter any memory location, unless the function is declared const.

If the last group of expressions in a parallel are each a clobber expression whose arguments are reg or match\_scratch (see Section 17.4 [RTL Template], page 339) expressions, the combiner phase can add the appropriate clobber expressions to an insn it has constructed when doing so will cause a pattern to be matched.

This feature can be used, for example, on a machine that whose multiply and add instructions don't use an MQ register but which has an add-accumulate instruction that does clobber the MQ register. Similarly, a combined instruction might require a temporary register while the constituent instructions might not.

When a clobber expression for a register appears inside a parallel with other side effects, the register allocator guarantees that the register is unoccupied both before and after that insn if it is a hard register clobber. For pseudoregister clobber, the register allocator and the reload pass do not assign the same hard register to the clobber and the input operands if there is an insn alternative containing the '&' constraint (see Section 17.8.4 [Modifiers], page 356) for the clobber and the hard register is in register classes of the clobber in the alternative. You can clobber either a specific hard register, a pseudo register, or a scratch expression; in the latter two cases, GCC will allocate a hard register that is available there for use as a temporary.

For instructions that require a temporary register, you should use scratch instead of a pseudo-register because this will allow the combiner phase to add the clobber when required. You do this by coding (clobber (match\_scratch . . . )). If you do clobber a pseudo register, use one which appears nowhere else—generate a new one each time. Otherwise, you may confuse CSE.

There is one other known use for clobbering a pseudo register in a parallel: when one of the input operands of the insn is also clobbered by the insn. In this case, using the same pseudo register in the clobber and elsewhere in the insn produces the expected results.

(use x) Represents the use of the value of x. It indicates that the value in x at this point in the program is needed, even though it may not be apparent why this is so. Therefore, the compiler will not attempt to delete previous instructions whose only effect is to store a value in x. x must be a reg expression.

In some situations, it may be tempting to add a use of a register in a parallel to describe a situation where the value of a special register will modify the behavior of the instruction. A hypothetical example might be a pattern for an addition that can either wrap around or use saturating addition depending on the value of a special control register:

This will not work, several of the optimizers only look at expressions locally; it is very likely that if you have multiple insns with identical inputs to the unspec, they will be optimized away even if register 1 changes in between.

This means that use can *only* be used to describe that the register is live. You should think twice before adding use statements, more often you will want to use unspec instead. The use RTX is most commonly useful to describe that a fixed register is implicitly used in an insn. It is also safe to use in patterns where the compiler knows for other reasons that the result of the whole pattern is variable, such as 'cpymemm' or 'call' patterns.

During the reload phase, an insn that has a use as pattern can carry a reg\_equal note. These use insns will be deleted before the reload phase exits.

During the delayed branch scheduling phase, x may be an insn. This indicates that x previously was located at this place in the code and its data dependencies need to be taken into account. These use insns will be deleted before the delayed branch scheduling phase exits.

```
(parallel [x0 x1 \dots])
```

Represents several side effects performed in parallel. The square brackets stand for a vector; the operand of parallel is a vector of expressions. x0, x1 and so on are individual side effect expressions—expressions of code set, call, return, simple\_return, clobber or use.

"In parallel" means that first all the values used in the individual side-effects are computed, and second all the actual side-effects are performed. For example,

says unambiguously that the values of hard register 1 and the memory location addressed by it are interchanged. In both places where (reg:SI 1) appears as a memory address it refers to the value in register 1 before the execution of the insp

It follows that it is *incorrect* to use parallel and expect the result of one set to be available for the next one. For example, people sometimes attempt to represent a jump-if-zero instruction this way:

But this is incorrect, because it says that the jump condition depends on the condition code value *before* this instruction, not on the new value that is set by this instruction.

Peephole optimization, which takes place together with final assembly code output, can produce insns whose patterns consist of a parallel whose elements

are the operands needed to output the resulting assembler code—often reg, mem or constant expressions. This would not be well-formed RTL at any other stage in compilation, but it is OK then because no further optimization remains to be done. However, the definition of the macro NOTICE\_UPDATE\_CC, if any, must deal with such insns if you define any peephole optimizations.

# (cond\_exec [cond expr])

Represents a conditionally executed expression. The *expr* is executed only if the *cond* is nonzero. The *cond* expression must not have side-effects, but the *expr* may very well have side-effects.

# (sequence [insns ...])

Represents a sequence of insns. If a **sequence** appears in the chain of insns, then each of the *insns* that appears in the sequence must be suitable for appearing in the chain of insns, i.e. must satisfy the INSN\_P predicate.

After delay-slot scheduling is completed, an insn and all the insns that reside in its delay slots are grouped together into a **sequence**. The insn requiring the delay slot is the first insn in the vector; subsequent insns are to be placed in the delay slot.

INSN\_ANNULLED\_BRANCH\_P is set on an insn in a delay slot to indicate that a branch insn should be used that will conditionally annul the effect of the insns in the delay slots. In such a case, INSN\_FROM\_TARGET\_P indicates that the insn is from the target of the branch and should be executed only if the branch is taken; otherwise the insn should be executed only if the branch is not taken. See Section 17.19.8 [Delay Slots], page 459.

Some back ends also use **sequence** objects for purposes other than delay-slot groups. This is not supported in the common parts of the compiler, which treat such sequences as delay-slot groups.

DWARF2 Call Frame Address (CFA) adjustments are sometimes also expressed using sequence objects as the value of a RTX\_FRAME\_RELATED\_P note. This only happens if the CFA adjustments cannot be easily derived from the pattern of the instruction to which the note is attached. In such cases, the value of the note is used instead of best-guesing the semantics of the instruction. The back end can attach notes containing a sequence of set patterns that express the effect of the parent instruction.

These expression codes appear in place of a side effect, as the body of an insn, though strictly speaking they do not always describe side effects as such:

# (asm\_input s)

Represents literal assembler code as described by the string s.

# (unspec [operands ...] index) (unspec\_volatile [operands ...] index)

Represents a machine-specific operation on *operands*. *index* selects between multiple machine-specific operations. unspec\_volatile is used for volatile operations and operations that may trap; unspec is used for other operations.

These codes may appear inside a pattern of an insn, inside a parallel, or inside an expression.

# (addr\_vec:m [1r0 1r1 ...])

Represents a table of jump addresses. The vector elements lr0, etc., are label\_ref expressions. The mode m specifies how much space is given to each address; normally m would be Pmode.

# (addr\_diff\_vec:m base [1r0 1r1 ...] min max flags)

Represents a table of jump addresses expressed as offsets from base. The vector elements lr0, etc., are label\_ref expressions and so is base. The mode m specifies how much space is given to each address-difference. min and max are set up by branch shortening and hold a label with a minimum and a maximum address, respectively. flags indicates the relative position of base, min and max to the containing insn and of min and max to base. See rtl.def for details.

# (prefetch: m addr rw locality)

Represents prefetch of memory at address addr. Operand rw is 1 if the prefetch is for data to be written, 0 otherwise; targets that do not support write prefetches should treat this as a normal prefetch. Operand locality specifies the amount of temporal locality; 0 if there is none or 1, 2, or 3 for increasing levels of temporal locality; targets that do not support locality hints should ignore this.

This insn is used to minimize cache-miss latency by moving data into a cache before it is accessed. It should use only non-faulting data prefetch instructions.

# 14.16 Embedded Side-Effects on Addresses

Six special side-effect expression codes appear as memory addresses.

# (pre\_dec:m x)

Represents the side effect of decrementing x by a standard amount and represents also the value that x has after being decremented. x must be a reg or mem, but most machines allow only a reg. m must be the machine mode for pointers on the machine in use. The amount x is decremented by is the length in bytes of the machine mode of the containing memory reference of which this expression serves as the address. Here is an example of its use:

```
(mem:DF (pre_dec:SI (reg:SI 39)))
```

This says to decrement pseudo register 39 by the length of a DFmode value and use the result to address a DFmode value.

# (pre\_inc:m x)

Similar, but specifies incrementing x instead of decrementing it.

# (post\_dec:m x)

Represents the same side effect as **pre\_dec** but a different value. The value represented here is the value x has *before* being decremented.

# (post\_inc:m x)

Similar, but specifies incrementing x instead of decrementing it.

# (post\_modify:m x y)

Represents the side effect of setting x to y and represents x before x is modified. x must be a reg or mem, but most machines allow only a reg. m must be the machine mode for pointers on the machine in use.

The expression y must be one of three forms: (plus: $m \times z$ ), (minus: $m \times z$ ), or (plus: $m \times i$ ), where z is an index register and i is a constant.

Here is an example of its use:

This says to modify pseudo register 42 by adding the contents of pseudo register 48 to it, after the use of what ever 42 points to.

```
(pre_modify:m x expr)
```

Similar except side effects happen before the use.

These embedded side effect expressions must be used with care. Instruction patterns may not use them. Until the 'flow' pass of the compiler, they may occur only to represent pushes onto the stack. The 'flow' pass finds cases where registers are incremented or decremented in one instruction and used as an address shortly before or after; these cases are then transformed to use pre- or post-increment or -decrement.

If a register used as the operand of these expressions is used in another address in an insn, the original value of the register is used. Uses of the register outside of an address are not permitted within the same insn as a use in an embedded side effect expression because such insns behave differently on different machines and hence must be treated as ambiguous and disallowed.

An instruction that can be represented with an embedded side effect could also be represented using parallel containing an additional set to describe how the address register is altered. This is not done because machines that allow these operations at all typically allow them wherever a memory address is called for. Describing them as additional parallel stores would require doubling the number of entries in the machine description.

# 14.17 Assembler Instructions as Expressions

The RTX code asm\_operands represents a value produced by a user-specified assembler instruction. It is used to represent an asm statement with arguments. An asm statement with a single output operand, like this:

```
asm ("foo %1,%2,%0" : "=a" (outputvar) : "g" (x + y), "di" (*z));
```

is represented using a single asm\_operands RTX which represents the value that is stored in outputvar:

Here the operands of the asm\_operands RTX are the assembler template string, the output-operand's constraint, the index-number of the output operand among the output operands specified, a vector of input operand RTX's, and a vector of input-operand modes and constraints. The mode m1 is the mode of the sum x+y; m2 is that of \*z.

When an asm statement has multiple output values, its insn has several such set RTX's inside of a parallel. Each set contains an asm\_operands; all of these share the same assembler template and vectors, but each contains the constraint for the respective output operand. They are also distinguished by the output-operand index number, which is 0, 1, . . . for successive output operands.

# 14.18 Variable Location Debug Information in RTL

Variable tracking relies on MEM\_EXPR and REG\_EXPR annotations to determine what user variables memory and register references refer to.

Variable tracking at assignments uses these notes only when they refer to variables that live at fixed locations (e.g., addressable variables, global non-automatic variables). For variables whose location may vary, it relies on the following types of notes.

# (var\_location:mode var exp stat)

Binds variable var, a tree, to value exp, an RTL expression. It appears only in NOTE\_INSN\_VAR\_LOCATION and DEBUG\_INSNs, with slightly different meanings. mode, if present, represents the mode of exp, which is useful if it is a modeless expression. stat is only meaningful in notes, indicating whether the variable is known to be initialized or uninitialized.

#### (debug\_expr:mode decl)

Stands for the value bound to the DEBUG\_EXPR\_DECL decl, that points back to it, within value expressions in VAR\_LOCATION nodes.

# (debug\_implicit\_ptr:mode decl)

Stands for the location of a decl that is no longer addressable.

# (entry\_value:mode decl)

Stands for the value a decl had at the entry point of the containing function.

#### (debug\_parameter\_ref:mode decl)

Refers to a parameter that was completely optimized out.

# (debug\_marker: mode)

Marks a program location. With VOIDmode, it stands for the beginning of a statement, a recommended inspection point logically after all prior side effects, and before any subsequent side effects. With BLKmode, it indicates an inline entry point: the lexical block encoded in the INSN\_LOCATION is the enclosing block that encloses the inlined function.

# 14.19 Insns

The RTL representation of the code for a function is a doubly-linked chain of objects called *insns*. Insns are expressions with special codes that are used for no other purpose. Some insns are actual instructions; others represent dispatch tables for switch statements; others represent labels to jump to or various sorts of declarative information.

In addition to its own specific data, each insn must have a unique id-number that distinguishes it from all other insns in the current function (after delayed branch scheduling, copies of an insn with the same id-number may be present in multiple places in a function, but these copies will always be identical and will only appear inside a sequence), and chain pointers to the preceding and following insns. These three fields occupy the same position in every insn, independent of the expression code of the insn. They could be accessed with XEXP and XINT, but instead three special macros are always used:

#### INSN\_UID (i)

Accesses the unique id of insn i.

# PREV\_INSN (i)

Accesses the chain pointer to the insn preceding i. If i is the first insn, this is a null pointer.

#### NEXT\_INSN (i)

Accesses the chain pointer to the insn following i. If i is the last insn, this is a null pointer.

The first insn in the chain is obtained by calling get\_insns; the last insn is the result of calling get\_last\_insn. Within the chain delimited by these insns, the NEXT\_INSN and PREV\_INSN pointers must always correspond: if *insn* is not the first insn,

```
NEXT_INSN (PREV_INSN (insn)) == insn
```

is always true and if *insn* is not the last insn,

```
PREV_INSN (NEXT_INSN (insn)) == insn
```

is always true.

After delay slot scheduling, some of the insns in the chain might be sequence expressions, which contain a vector of insns. The value of NEXT\_INSN in all but the last of these insns is the next insn in the vector; the value of NEXT\_INSN of the last insn in the vector is the same as the value of NEXT\_INSN for the sequence in which it is contained. Similar rules apply for PREV\_INSN.

This means that the above invariants are not necessarily true for insns inside sequence expressions. Specifically, if *insn* is the first insn in a sequence, NEXT\_INSN (PREV\_INSN (*insn*)) is the insn containing the sequence expression, as is the value of PREV\_INSN (NEXT\_INSN (*insn*)) if *insn* is the last insn in the sequence expression. You can use these expressions to find the containing sequence expression.

Every insn has one of the following expression codes:

insn

The expression code insn is used for instructions that do not jump and do not do function calls. sequence expressions are always contained in insns with code insn even if one of those insns should jump or do function calls.

Insps with code insp have four additional fields beyond the three mandatory ones listed above. These four are described in a table below.

#### jump\_insn

The expression code jump\_insn is used for instructions that may jump (or, more generally, may contain label\_ref expressions to which pc can be set in that instruction). If there is an instruction to return from the current function, it is recorded as a jump\_insn.

jump\_insn insns have the same extra fields as insn insns, accessed in the same way and in addition contain a field JUMP\_LABEL which is defined once jump optimization has completed.

For simple conditional and unconditional jumps, this field contains the code\_label to which this insn will (possibly conditionally) branch. In a more complex jump, JUMP\_LABEL records one of the labels that the insn refers to; other jump target labels are recorded as REG\_LABEL\_TARGET notes. The exception is addr\_vec and addr\_diff\_vec, where JUMP\_LABEL is NULL\_RTX and the only way to find the labels is to scan the entire body of the insn.

Return insns count as jumps, but their JUMP\_LABEL is RETURN or SIMPLE\_RETURN.

#### call\_insn

The expression code call\_insn is used for instructions that may do function calls. It is important to distinguish these instructions because they imply that certain registers and memory locations may be altered unpredictably.

call\_insn insns have the same extra fields as insn insns, accessed in the same way and in addition contain a field CALL\_INSN\_FUNCTION\_USAGE, which contains a list (chain of expr\_list expressions) containing use, clobber and sometimes set expressions that denote hard registers and mems used or clobbered by the called function.

A mem generally points to a stack slot in which arguments passed to the libcall by reference (see Section 18.9.7 [Register Arguments], page 536) are stored. If the argument is caller-copied (see Section 18.9.7 [Register Arguments], page 536), the stack slot will be mentioned in clobber and use entries; if it's callee-copied, only a use will appear, and the mem may point to addresses that are not stack slots.

Registers occurring inside a clobber in this list augment registers specified in CALL\_USED\_REGISTERS (see Section 18.7.1 [Register Basics], page 505).

If the list contains a set involving two registers, it indicates that the function returns one of its arguments. Such a set may look like a no-op if the same register holds the argument and the return value.

#### code\_label

A code\_label insn represents a label that a jump insn can jump to. It contains two special fields of data in addition to the three standard ones. CODE\_LABEL\_NUMBER is used to hold the *label number*, a number that identifies this label uniquely among all the labels in the compilation (not just in the current function). Ultimately, the label is represented in the assembler output as an assembler label, usually of the form 'Ln' where n is the label number.

When a code\_label appears in an RTL expression, it normally appears within a label\_ref which represents the address of the label, as a number.

Besides as a code\_label, a label can also be represented as a note of type NOTE\_INSN\_DELETED\_LABEL.

The field LABEL\_NUSES is only defined once the jump optimization phase is completed. It contains the number of times this label is referenced in the current function.

The field LABEL\_KIND differentiates four different types of labels: LABEL\_NORMAL, LABEL\_STATIC\_ENTRY, LABEL\_GLOBAL\_ENTRY, and LABEL\_WEAK\_ENTRY. The only labels that do not have type LABEL\_NORMAL are alternate entry points to the current function. These may be static (visible only in the containing translation unit), global (exposed to all translation units), or weak (global, but can be overridden by another symbol with the same name).

Much of the compiler treats all four kinds of label identically. Some of it needs to know whether or not a label is an alternate entry point; for this purpose,

the macro LABEL\_ALT\_ENTRY\_P is provided. It is equivalent to testing whether 'LABEL\_KIND (label) == LABEL\_NORMAL'. The only place that cares about the distinction between static, global, and weak alternate entry points, besides the front-end code that creates them, is the function output\_alternate\_entry\_point, in 'final.c'.

To set the kind of a label, use the SET\_LABEL\_KIND macro.

# jump\_table\_data

A jump\_table\_data insn is a placeholder for the jump-table data of a casesi or tablejump insn. They are placed after a tablejump\_p insn. A jump\_table\_data insn is not part o a basic blockm but it is associated with the basic block that ends with the tablejump\_p insn. The PATTERN of a jump\_table\_data is always either an addr\_vec or an addr\_diff\_vec, and a jump\_table\_data insn is always preceded by a code\_label. The tablejump\_p insn refers to that code\_label via its JUMP\_LABEL.

barrier

Barriers are placed in the instruction stream when control cannot flow past them. They are placed after unconditional jump instructions to indicate that the jumps are unconditional and after calls to volatile functions, which do not return (e.g., exit). They contain no information beyond the three standard fields.

note

note insns are used to represent additional debugging and declarative information. They contain two nonstandard fields, an integer which is accessed with the macro NOTE\_LINE\_NUMBER and a string accessed with NOTE\_SOURCE\_FILE.

If NOTE\_LINE\_NUMBER is positive, the note represents the position of a source line and NOTE\_SOURCE\_FILE is the source file name that the line came from. These notes control generation of line number data in the assembler output.

Otherwise, NOTE\_LINE\_NUMBER is not really a line number but a code with one of the following values (and NOTE\_SOURCE\_FILE must contain a null pointer):

# NOTE\_INSN\_DELETED

Such a note is completely ignorable. Some passes of the compiler delete insns by altering them into notes of this kind.

#### NOTE\_INSN\_DELETED\_LABEL

This marks what used to be a code\_label, but was not used for other purposes than taking its address and was transformed to mark that no code jumps to it.

NOTE\_INSN\_BLOCK\_BEG

NOTE\_INSN\_BLOCK\_END

These types of notes indicate the position of the beginning and end of a level of scoping of variable names. They control the output of debugging information.

NOTE\_INSN\_EH\_REGION\_BEG

NOTE\_INSN\_EH\_REGION\_END

These types of notes indicate the position of the beginning and end of a level of scoping for exception handling. NOTE\_EH\_HANDLER identifies which region is associated with these notes.

#### NOTE\_INSN\_FUNCTION\_BEG

Appears at the start of the function body, after the function prologue.

# NOTE\_INSN\_VAR\_LOCATION

This note is used to generate variable location debugging information. It indicates that the user variable in its VAR\_LOCATION operand is at the location given in the RTL expression, or holds a value that can be computed by evaluating the RTL expression from that static point in the program up to the next such note for the same user variable.

#### NOTE\_INSN\_BEGIN\_STMT

This note is used to generate is\_stmt markers in line number debuggign information. It indicates the beginning of a user statement.

#### NOTE\_INSN\_INLINE\_ENTRY

This note is used to generate entry\_pc for inlined subroutines in debugging information. It indicates an inspection point at which all arguments for the inlined function have been bound, and before its first statement.

These codes are printed symbolically when they appear in debugging dumps.

# debug\_insn

The expression code debug\_insn is used for pseudo-instructions that hold debugging information for variable tracking at assignments (see '-fvar-tracking-assignments' option). They are the RTL representation of GIMPLE\_DEBUG statements (Section 12.8.7 [GIMPLE\_DEBUG], page 229), with a VAR\_LOCATION operand that binds a user variable tree to an RTL representation of the value in the corresponding statement. A DEBUG\_EXPR in it stands for the value bound to the corresponding DEBUG\_EXPR\_DECL.

GIMPLE\_DEBUG\_BEGIN\_STMT and GIMPLE\_DEBUG\_INLINE\_ENTRY are expanded to RTL as a DEBUG\_INSN with a DEBUG\_MARKER PATTERN; the difference is the RTL mode: the former's DEBUG\_MARKER is VOIDmode, whereas the latter is BLKmode; information about the inlined function can be taken from the lexical block encoded in the INSN\_LOCATION. These DEBUG\_INSNs, that do not carry VAR\_LOCATION information, just DEBUG\_MARKERs, can be detected by testing DEBUG\_MARKER\_INSN\_P, whereas those that do can be recognized as DEBUG\_BIND\_INSN\_P.

Throughout optimization passes, DEBUG\_INSNs are not reordered with respect to each other, particularly during scheduling. Binding information is kept in pseudo-instruction form, so that, unlike notes, it gets the same treatment and adjustments that regular instructions would. It is the variable tracking pass that turns these pseudo-instructions into NOTE\_INSN\_VAR\_LOCATION, NOTE\_INSN\_BEGIN\_STMT and NOTE\_INSN\_INLINE\_ENTRY notes, analyzing control flow, value equivalences and changes to registers and memory referenced in value expressions, propagating the values of debug temporaries and determining expressions that can be used to compute the value of each user variable at as many points (ranges, actually) in the program as possible.

Unlike NOTE\_INSN\_VAR\_LOCATION, the value expression in an INSN\_VAR\_LOCATION denotes a value at that specific point in the program, rather than an expression that can be evaluated at any later point before an overriding VAR\_LOCATION is encountered. E.g., if a user variable is bound to a REG and then a subsequent insn modifies the REG, the note location would keep mapping the user variable to the register across the insn, whereas the insn location would keep the variable bound to the value, so that the variable tracking pass would emit another location note for the variable at the point in which the register is modified.

The machine mode of an insn is normally VOIDmode, but some phases use the mode for various purposes.

The common subexpression elimination pass sets the mode of an insn to QImode when it is the first insn in a block that has already been processed.

The second Haifa scheduling pass, for targets that can multiple issue, sets the mode of an insn to TImode when it is believed that the instruction begins an issue group. That is, when the instruction cannot issue simultaneously with the previous. This may be relied on by later passes, in particular machine-dependent reorg.

Here is a table of the extra fields of insn, jump\_insn and call\_insn insns:

#### PATTERN (i)

An expression for the side effect performed by this insn. This must be one of the following codes: set, call, use, clobber, return, simple\_return, asm\_input, asm\_output, addr\_vec, addr\_diff\_vec, trap\_if, unspec, unspec\_volatile, parallel, cond\_exec, or sequence. If it is a parallel, each element of the parallel must be one these codes, except that parallel expressions cannot be nested and addr\_vec and addr\_diff\_vec are not permitted inside a parallel expression.

## INSN\_CODE (i)

An integer that says which pattern in the machine description matches this insn, or -1 if the matching has not yet been attempted.

Such matching is never attempted and this field remains -1 on an insn whose pattern consists of a single use, clobber, asm\_input, addr\_vec or addr\_diff\_vec expression.

Matching is also never attempted on insns that result from an asm statement. These contain at least one asm\_operands expression. The function asm\_noperands returns a non-negative value for such insns.

In the debugging output, this field is printed as a number followed by a symbolic representation that locates the pattern in the 'md' file as some small positive or negative offset from a named pattern.

### LOG\_LINKS (i)

A list (chain of insn\_list expressions) giving information about dependencies between instructions within a basic block. Neither a jump nor a label may come between the related insns. These are only used by the schedulers and by combine. This is a deprecated data structure. Def-use and use-def chains are now preferred.

## REG\_NOTES (i)

A list (chain of expr\_list, insn\_list and int\_list expressions) giving miscellaneous information about the insn. It is often information pertaining to the registers used in this insn.

The LOG\_LINKS field of an insn is a chain of insn\_list expressions. Each of these has two operands: the first is an insn, and the second is another insn\_list expression (the next one in the chain). The last insn\_list in the chain has a null pointer as second operand. The significant thing about the chain is which insns appear in it (as first operands of insn\_list expressions). Their order is not significant.

This list is originally set up by the flow analysis pass; it is a null pointer until then. Flow only adds links for those data dependencies which can be used for instruction combination. For each insn, the flow analysis pass adds a link to insns which store into registers values that are used for the first time in this insn.

The REG\_NOTES field of an insn is a chain similar to the LOG\_LINKS field but it includes expr\_list and int\_list expressions in addition to insn\_list expressions. There are several kinds of register notes, which are distinguished by the machine mode, which in a register note is really understood as being an enum reg\_note. The first operand op of the note is data whose meaning depends on the kind of note.

The macro REG\_NOTE\_KIND (x) returns the kind of register note. Its counterpart, the macro PUT\_REG\_NOTE\_KIND (x, newkind) sets the register note type of x to be newkind.

Register notes are of three classes: They may say something about an input to an insn, they may say something about an output of an insn, or they may create a linkage between two insns. There are also a set of values that are only used in LOG\_LINKS.

These register notes annotate inputs to an insn:

REG\_DEAD The value in op dies in this insn; that is to say, altering the value immediately after this insn would not affect the future behavior of the program.

It does not follow that the register op has no useful value after this insn since op is not necessarily modified by this insn. Rather, no subsequent instruction uses the contents of op.

### REG\_UNUSED

The register op being set by this inso will not be used in a subsequent inso. This differs from a REG\_DEAD note, which indicates that the value in an input will not be used subsequently. These two notes are independent; both may be present for the same register.

REG\_INC The register op is incremented (or decremented; at this level there is no distinction) by an embedded side effect inside this insn. This means it appears in a post\_inc, pre\_inc, post\_dec or pre\_dec expression.

#### REG\_NONNEG

The register op is known to have a nonnegative value when this insn is reached. This is used by special looping instructions that terminate when the register goes negative.

The REG\_NONNEG note is added only to 'doloop\_end' insns, if its pattern uses a ge condition.

### REG\_LABEL\_OPERAND

This insn uses op, a code\_label or a note of type NOTE\_INSN\_DELETED\_LABEL, but is not a jump\_insn, or it is a jump\_insn that refers to the operand as an ordinary operand. The label may still eventually be a jump target, but if so in an indirect jump in a subsequent insn. The presence of this note allows jump optimization to be aware that op is, in fact, being used, and flow optimization to build an accurate flow graph.

### REG\_LABEL\_TARGET

This insn is a jump\_insn but not an addr\_vec or addr\_diff\_vec. It uses op, a code\_label as a direct or indirect jump target. Its purpose is similar to that of REG\_LABEL\_OPERAND. This note is only present if the insn has multiple targets; the last label in the insn (in the highest numbered insn-field) goes into the JUMP\_LABEL field and does not have a REG\_LABEL\_TARGET note. See Section 14.19 [Insns], page 304.

## REG\_SETJMP

Appears attached to each CALL\_INSN to setjmp or a related function.

The following notes describe attributes of outputs of an insn:

REG\_EQUIV REG\_EQUAL

This note is only valid on an insn that sets only one register and indicates that that register will be equal to op at run time; the scope of this equivalence differs between the two types of notes. The value which the insn explicitly copies into the register may look different from op, but they will be equal at run time. If the output of the single set is a strict\_low\_part or zero\_extract expression, the note refers to the register that is contained in its first operand.

For REG\_EQUIV, the register is equivalent to op throughout the entire function, and could validly be replaced in all its occurrences by op. ("Validly" here refers to the data flow of the program; simple replacement may make some insns invalid.) For example, when a constant is loaded into a register that is never assigned any other value, this kind of note is used.

When a parameter is copied into a pseudo-register at entry to a function, a note of this kind records that the register is equivalent to the stack slot where the parameter was passed. Although in this case the register may be set by other insns, it is still valid to replace the register by the stack slot throughout the function.

A REG\_EQUIV note is also used on an instruction which copies a register parameter into a pseudo-register at entry to a function, if there is a stack slot where that parameter could be stored. Although other insns may set the pseudo-register, it is valid for the compiler to replace the pseudo-register by stack slot throughout the function, provided the compiler ensures that the stack slot is properly initialized by making the replacement in the initial copy instruction as well. This is used on machines for which the calling convention allocates stack space for register parameters. See REG\_PARM\_STACK\_SPACE in Section 18.9.6 [Stack Arguments], page 534.

In the case of REG\_EQUAL, the register that is set by this insn will be equal to op at run time at the end of this insn but not necessarily elsewhere in the function. In this case, op is typically an arithmetic expression. For example, when a sequence of insns such as a library call is used to perform an arithmetic operation, this kind of note is attached to the insn that produces or copies the final value.

These two notes are used in different ways by the compiler passes. REG\_EQUAL is used by passes prior to register allocation (such as common subexpression elimination and loop optimization) to tell them how to think of that value. REG\_EQUIV notes are used by register allocation to indicate that there is an available substitute expression (either a constant or a mem expression for the location of a parameter on the stack) that may be used in place of a register if insufficient registers are available.

Except for stack homes for parameters, which are indicated by a REG\_EQUIV note and are not useful to the early optimization passes and pseudo registers that are equivalent to a memory location throughout their entire life, which is not detected until later in the compilation, all equivalences are initially indicated by an attached REG\_EQUAL note. In the early stages of register allocation, a REG\_EQUAL note is changed into a REG\_EQUIV note if op is a constant and the insn represents the only set of its destination register.

Thus, compiler passes prior to register allocation need only check for REG\_EQUAL notes and passes subsequent to register allocation need only check for REG\_EQUIV notes.

These notes describe linkages between insns. They occur in pairs: one insn has one of a pair of notes that points to a second insn, which has the inverse note pointing back to the first insn.

REG\_CC\_SETTER REG\_CC\_USER

On machines that use cc0, the insns which set and use cc0 set and use cc0 are adjacent. However, when branch delay slot filling is done, this may no longer be true. In this case a REG\_CC\_USER note will be placed on the insn setting cc0 to point to the insn using cc0 and a REG\_CC\_SETTER note will be placed on the insn using cc0 to point to the insn setting cc0.

These values are only used in the LOG\_LINKS field, and indicate the type of dependency that each link represents. Links which indicate a data dependence (a read after write dependence) do not use any code, they simply have mode VOIDmode, and are printed without any descriptive text.

REG\_DEP\_TRUE

This indicates a true dependence (a read after write dependence).

REG\_DEP\_OUTPUT

This indicates an output dependence (a write after write dependence).

REG\_DEP\_ANTI

This indicates an anti dependence (a write after read dependence).

These notes describe information gathered from gcov profile data. They are stored in the REG\_NOTES field of an insn.

### REG\_BR\_PROB

This is used to specify the ratio of branches to non-branches of a branch insn according to the profile data. The note is represented as an int\_list expression whose integer value is an encoding of profile\_probability type. profile\_probability provide member function from\_reg\_br\_prob\_note and to\_reg\_br\_prob\_note to extract and store the probability into the RTL encoding.

#### REG\_BR\_PRED

These notes are found in JUMP insns after delayed branch scheduling has taken place. They indicate both the direction and the likelihood of the JUMP. The format is a bitmask of ATTR\_FLAG\_\* values.

## REG\_FRAME\_RELATED\_EXPR

This is used on an RTX\_FRAME\_RELATED\_P inso wherein the attached expression is used in place of the actual inso pattern. This is done in cases where the pattern is either complex or misleading.

The note REG\_CALL\_NOCF\_CHECK is used in conjunction with the '-fcf-protection=branch' option. The note is set if a nocf\_check attribute is specified for a function type or a pointer to function type. The note is stored in the REG\_NOTES field of an insn.

#### REG\_CALL\_NOCF\_CHECK

Users have control through the nocf\_check attribute to identify which calls to a function should be skipped from control-flow instrumentation when the option '-fcf-protection=branch' is specified. The compiler puts a REG\_CALL\_NOCF\_CHECK note on each CALL\_INSN instruction that has a function type marked with a nocf\_check attribute.

For convenience, the machine mode in an insn\_list or expr\_list is printed using these symbolic codes in debugging dumps.

The only difference between the expression codes insn\_list and expr\_list is that the first operand of an insn\_list is assumed to be an insn and is printed in debugging dumps as the insn's unique id; the first operand of an expr\_list is printed in the ordinary way as an expression.

# 14.20 RTL Representation of Function-Call Insns

Insns that call subroutines have the RTL expression code call\_insn. These insns must satisfy special rules, and their bodies must use a special RTL expression code, call.

A call expression has two operands, as follows:

```
(call (mem:fm addr) nbytes)
```

Here *nbytes* is an operand that represents the number of bytes of argument data being passed to the subroutine, *fm* is a machine mode (which must equal as the definition of the FUNCTION\_MODE macro in the machine description) and *addr* represents the address of the subroutine.

For a subroutine that returns no value, the call expression as shown above is the entire body of the insn, except that the insn might also contain use or clobber expressions.

For a subroutine that returns a value whose mode is not BLKmode, the value is returned in a hard register. If this register's number is r, then the body of the call insu looks like this:

```
(set (reg:m r)
     (call (mem:fm addr) nbytes))
```

This RTL expression makes it clear (to the optimizer passes) that the appropriate register receives a useful value in this insn.

When a subroutine returns a BLKmode value, it is handled by passing to the subroutine the address of a place to store the value. So the call insn itself does not "return" any value, and it has the same RTL form as a call that returns nothing.

On some machines, the call instruction itself clobbers some register, for example to contain the return address. call\_insn insns on these machines should have a body which is a parallel that contains both the call expression and clobber expressions that indicate which registers are destroyed. Similarly, if the call instruction requires some register other than the stack pointer that is not explicitly mentioned in its RTL, a use subexpression should mention that register.

Functions that are called are assumed to modify all registers listed in the configuration macro CALL\_USED\_REGISTERS (see Section 18.7.1 [Register Basics], page 505) and, with the exception of const functions and library calls, to modify all of memory.

Insns containing just use expressions directly precede the call\_insn insn to indicate which registers contain inputs to the function. Similarly, if registers other than those in CALL\_USED\_REGISTERS are clobbered by the called function, insns containing a single clobber follow immediately after the call to indicate which registers.

# 14.21 Structure Sharing Assumptions

The compiler assumes that certain kinds of RTL expressions are unique; there do not exist two distinct objects representing the same value. In other cases, it makes an opposite assumption: that no RTL expression object of a certain kind appears in more than one place in the containing structure.

These assumptions refer to a single function; except for the RTL objects that describe global variables and external functions, and a few standard objects such as small integer constants, no RTL objects are common to two functions.

- Each pseudo-register has only a single reg object to represent it, and therefore only a single machine mode.
- For any symbolic label, there is only one symbol\_ref object referring to it.
- All const\_int expressions with equal values are shared.
- All const\_poly\_int expressions with equal modes and values are shared.
- There is only one pc expression.
- There is only one cc0 expression.
- There is only one const\_double expression with value 0 for each floating point mode. Likewise for values 1 and 2.
- There is only one const\_vector expression with value 0 for each vector mode, be it an integer or a double constant vector.

- No label\_ref or scratch appears in more than one place in the RTL structure; in other words, it is safe to do a tree-walk of all the insns in the function and assume that each time a label\_ref or scratch is seen it is distinct from all others that are seen.
- Only one mem object is normally created for each static variable or stack slot, so these objects are frequently shared in all the places they appear. However, separate but equal objects for these variables are occasionally made.
- When a single asm statement has multiple output operands, a distinct asm\_operands expression is made for each output operand. However, these all share the vector which contains the sequence of input operands. This sharing is used later on to test whether two asm\_operands expressions come from the same statement, so all optimizations must carefully preserve the sharing if they copy the vector at all.
- No RTL object appears in more than one place in the RTL structure except as described above. Many passes of the compiler rely on this by assuming that they can modify RTL objects in place without unwanted side-effects on other insns.
- During initial RTL generation, shared structure is freely introduced. After all the RTL for a function has been generated, all shared structure is copied by unshare\_all\_rtl in 'emit-rtl.c', after which the above rules are guaranteed to be followed.
- During the combiner pass, shared structure within an insn can exist temporarily. However, the shared structure is copied before the combiner is finished with the insn. This is done by calling copy\_rtx\_if\_shared, which is a subroutine of unshare\_all\_rtl.

## 14.22 Reading RTL

To read an RTL object from a file, call read\_rtx. It takes one argument, a stdio stream, and returns a single RTL object. This routine is defined in 'read-rtl.c'. It is not available in the compiler itself, only the various programs that generate the compiler back end from the machine description.

People frequently have the idea of using RTL stored as text in a file as an interface between a language front end and the bulk of GCC. This idea is not feasible.

GCC was designed to use RTL internally only. Correct RTL for a given program is very dependent on the particular target machine. And the RTL does not contain all the information about the program.

The proper way to interface GCC to a new language front end is with the "tree" data structure, described in the files 'tree.h' and 'tree.def'. The documentation for this structure (see Chapter 11 [GENERIC], page 161) is incomplete.

# 15 Control Flow Graph

A control flow graph (CFG) is a data structure built on top of the intermediate code representation (the RTL or GIMPLE instruction stream) abstracting the control flow behavior of a function that is being compiled. The CFG is a directed graph where the vertices represent basic blocks and edges represent possible transfer of control flow from one basic block to another. The data structures used to represent the control flow graph are defined in 'basic-block.h'.

In GCC, the representation of control flow is maintained throughout the compilation process, from constructing the CFG early in pass\_build\_cfg to pass\_free\_cfg (see 'passes.def'). The CFG takes various different modes and may undergo extensive manipulations, but the graph is always valid between its construction and its release. This way, transfer of information such as data flow, a measured profile, or the loop tree, can be propagated through the passes pipeline, and even from GIMPLE to RTL.

Often the CFG may be better viewed as integral part of instruction chain, than structure built on the top of it. Updating the compiler's intermediate representation for instructions cannot be easily done without proper maintenance of the CFG simultaneously.

## 15.1 Basic Blocks

A basic block is a straight-line sequence of code with only one entry point and only one exit. In GCC, basic blocks are represented using the basic\_block data type.

Special basic blocks represent possible entry and exit points of a function. These blocks are called ENTRY\_BLOCK\_PTR and EXIT\_BLOCK\_PTR. These blocks do not contain any code.

The BASIC\_BLOCK array contains all basic blocks in an unspecified order. Each basic\_block structure has a field that holds a unique integer identifier index that is the index of the block in the BASIC\_BLOCK array. The total number of basic blocks in the function is n\_basic\_blocks. Both the basic block indices and the total number of basic blocks may vary during the compilation process, as passes reorder, create, duplicate, and destroy basic blocks. The index for any block should never be greater than last\_basic\_block. The indices 0 and 1 are special codes reserved for ENTRY\_BLOCK and EXIT\_BLOCK, the indices of ENTRY\_BLOCK\_PTR and EXIT\_BLOCK\_PTR.

Two pointer members of the basic\_block structure are the pointers next\_bb and prev\_bb. These are used to keep doubly linked chain of basic blocks in the same order as the underlying instruction stream. The chain of basic blocks is updated transparently by the provided API for manipulating the CFG. The macro FOR\_EACH\_BB can be used to visit all the basic blocks in lexicographical order, except ENTRY\_BLOCK and EXIT\_BLOCK. The macro FOR\_ALL\_BB also visits all basic blocks in lexicographical order, including ENTRY\_BLOCK and EXIT\_BLOCK.

The functions post\_order\_compute and inverted\_post\_order\_compute can be used to compute topological orders of the CFG. The orders are stored as vectors of basic block indices. The BASIC\_BLOCK array can be used to iterate each basic block by index. Dominator traversals are also possible using walk\_dominator\_tree. Given two basic blocks A and B, block A dominates block B if A is always executed before B.

Each basic\_block also contains pointers to the first instruction (the head) and the last instruction (the tail) or end of the instruction stream contained in a basic block. In fact,

since the basic\_block data type is used to represent blocks in both major intermediate representations of GCC (GIMPLE and RTL), there are pointers to the head and end of a basic block for both representations, stored in intermediate representation specific data in the il field of struct basic\_block\_def.

For RTL, these pointers are BB\_HEAD and BB\_END.

In the RTL representation of a function, the instruction stream contains not only the "real" instructions, but also *notes* or *insn notes* (to distinguish them from *reg notes*). Any function that moves or duplicates the basic blocks needs to take care of updating of these notes. Many of these notes expect that the instruction stream consists of linear regions, so updating can sometimes be tedious. All types of insn notes are defined in 'insn-notes.def'.

In the RTL function representation, the instructions contained in a basic block always follow a NOTE\_INSN\_BASIC\_BLOCK, but zero or more CODE\_LABEL nodes can precede the block note. A basic block ends with a control flow instruction or with the last instruction before the next CODE\_LABEL or NOTE\_INSN\_BASIC\_BLOCK. By definition, a CODE\_LABEL cannot appear in the middle of the instruction stream of a basic block.

In addition to notes, the jump table vectors are also represented as "pseudo-instructions" inside the insn stream. These vectors never appear in the basic block and should always be placed just after the table jump instructions referencing them. After removing the table-jump it is often difficult to eliminate the code computing the address and referencing the vector, so cleaning up these vectors is postponed until after liveness analysis. Thus the jump table vectors may appear in the insn stream unreferenced and without any purpose. Before any edge is made fall-thru, the existence of such construct in the way needs to be checked by calling can\_fallthru function.

For the GIMPLE representation, the PHI nodes and statements contained in a basic block are in a gimple\_seq pointed to by the basic block intermediate language specific pointers. Abstract containers and iterators are used to access the PHI nodes and statements in a basic blocks. These iterators are called GIMPLE statement iterators (GSIs). Grep for `gsi in the various 'gimple-\*' and 'tree-\*' files. There is a gimple\_stmt\_iterator type for iterating over all kinds of statement, and a gphi\_iterator subclass for iterating over PHI nodes. The following snippet will pretty-print all PHI nodes the statements of the current function in the GIMPLE representation.

```
basic_block bb;

FOR_EACH_BB (bb)
{
    gphi_iterator pi;
    gimple_stmt_iterator si;

    for (pi = gsi_start_phis (bb); !gsi_end_p (pi); gsi_next (&pi))
        {
            gphi *phi = pi.phi ();
            print_gimple_stmt (dump_file, phi, 0, TDF_SLIM);
        }

    for (si = gsi_start_bb (bb); !gsi_end_p (si); gsi_next (&si))
        {
            gimple stmt = gsi_stmt (si);
            print_gimple_stmt (dump_file, stmt, 0, TDF_SLIM);
        }
    }
}
```

## 15.2 Edges

Edges represent possible control flow transfers from the end of some basic block A to the head of another basic block B. We say that A is a predecessor of B, and B is a successor of A. Edges are represented in GCC with the edge data type. Each edge acts as a link between two basic blocks: The src member of an edge points to the predecessor basic block of the dest basic block. The members preds and succs of the basic\_block data type point to type-safe vectors of edges to the predecessors and successors of the block.

When walking the edges in an edge vector, *edge iterators* should be used. Edge iterators are constructed using the <code>edge\_iterator</code> data structure and several methods are available to operate on them:

- ei\_start This function initializes an edge\_iterator that points to the first edge in a vector of edges.
- ei\_last This function initializes an edge\_iterator that points to the last edge in a vector of edges.
- ei\_end\_p This predicate is true if an edge\_iterator represents the last edge in an edge vector.

### ei\_one\_before\_end\_p

This predicate is true if an edge\_iterator represents the second last edge in an edge vector.

- ei\_next This function takes a pointer to an edge\_iterator and makes it point to the next edge in the sequence.
- ei\_prev This function takes a pointer to an edge\_iterator and makes it point to the previous edge in the sequence.
- ei\_edge This function returns the edge currently pointed to by an edge\_iterator.

### ei\_safe\_safe

This function returns the edge currently pointed to by an edge\_iterator, but returns NULL if the iterator is pointing at the end of the sequence. This function has been provided for existing code makes the assumption that a NULL edge indicates the end of the sequence.

The convenience macro FOR\_EACH\_EDGE can be used to visit all of the edges in a sequence of predecessor or successor edges. It must not be used when an element might be removed during the traversal, otherwise elements will be missed. Here is an example of how to use the macro:

There are various reasons why control flow may transfer from one block to another. One possibility is that some instruction, for example a CODE\_LABEL, in a linearized instruction

stream just always starts a new basic block. In this case a fall-thru edge links the basic block to the first following basic block. But there are several other reasons why edges may be created. The flags field of the edge data type is used to store information about the type of edge we are dealing with. Each edge is of one of the following types:

jump

No type flags are set for edges corresponding to jump instructions. These edges are used for unconditional or conditional jumps and in RTL also for table jumps. They are the easiest to manipulate as they may be freely redirected when the flow graph is not in SSA form.

fall-thru

Fall-thru edges are present in case where the basic block may continue execution to the following one without branching. These edges have the EDGE\_FALLTHRU flag set. Unlike other types of edges, these edges must come into the basic block immediately following in the instruction stream. The function force\_nonfallthru is available to insert an unconditional jump in the case that redirection is needed. Note that this may require creation of a new basic block.

## exception handling

Exception handling edges represent possible control transfers from a trapping instruction to an exception handler. The definition of "trapping" varies. In C++, only function calls can throw, but for Ada exceptions like division by zero or segmentation fault are defined and thus each instruction possibly throwing this kind of exception needs to be handled as control flow instruction. Exception edges have the EDGE\_ABNORMAL and EDGE\_EH flags set.

When updating the instruction stream it is easy to change possibly trapping instruction to non-trapping, by simply removing the exception edge. The opposite conversion is difficult, but should not happen anyway. The edges can be eliminated via purge\_dead\_edges call.

In the RTL representation, the destination of an exception edge is specified by REG\_EH\_REGION note attached to the insn. In case of a trapping call the EDGE\_ABNORMAL\_CALL flag is set too. In the GIMPLE representation, this extra flag is not set.

In the RTL representation, the predicate may\_trap\_p may be used to check whether instruction still may trap or not. For the tree representation, the tree\_could\_trap\_p predicate is available, but this predicate only checks for possible memory traps, as in dereferencing an invalid pointer location.

### sibling calls

Sibling calls or tail calls terminate the function in a non-standard way and thus an edge to the exit must be present. EDGE\_SIBCALL and EDGE\_ABNORMAL are set in such case. These edges only exist in the RTL representation.

### computed jumps

Computed jumps contain edges to all labels in the function referenced from the code. All those edges have EDGE\_ABNORMAL flag set. The edges used to represent computed jumps often cause compile time performance problems, since functions consisting of many taken labels and many computed jumps may have *very* dense flow graphs, so these edges need to be handled with special

care. During the earlier stages of the compilation process, GCC tries to avoid such dense flow graphs by factoring computed jumps. For example, given the following series of jumps,

```
goto *x;
[ ... ]
goto *x;
[ ... ]
goto *x;
[ ... ]
```

factoring the computed jumps results in the following code sequence which has a much simpler flow graph:

```
goto y;
[ ... ]

goto y;
[ ... ]

goto y;
[ ... ]

y:
    goto *x;
```

However, the classic problem with this transformation is that it has a runtime cost in there resulting code: An extra jump. Therefore, the computed jumps are un-factored in the later passes of the compiler (in the pass called pass\_duplicate\_computed\_gotos). Be aware of that when you work on passes in that area. There have been numerous examples already where the compile time for code with unfactored computed jumps caused some serious headaches.

## nonlocal goto handlers

GCC allows nested functions to return into caller using a goto to a label passed to as an argument to the callee. The labels passed to nested functions contain special code to cleanup after function call. Such sections of code are referred to as "nonlocal goto receivers". If a function contains such nonlocal goto receivers, an edge from the call to the label is created with the EDGE\_ABNORMAL and EDGE\_ABNORMAL\_CALL flags set.

### function entry points

By definition, execution of function starts at basic block 0, so there is always an edge from the ENTRY\_BLOCK\_PTR to basic block 0. There is no GIMPLE representation for alternate entry points at this moment. In RTL, alternate entry points are specified by CODE\_LABEL with LABEL\_ALTERNATE\_NAME defined. This feature is currently used for multiple entry point prologues and is limited to post-reload passes only. This can be used by back-ends to emit alternate prologues for functions called from different contexts. In future full support for multiple entry functions defined by Fortran 90 needs to be implemented.

#### function exits

In the pre-reload representation a function terminates after the last instruction in the insn chain and no explicit return instructions are used. This corresponds to the fall-thru edge into exit block. After reload, optimal RTL epilogues are used that use explicit (conditional) return instructions that are represented by edges with no flags set.

## 15.3 Profile information

In many cases a compiler must make a choice whether to trade speed in one part of code for speed in another, or to trade code size for code speed. In such cases it is useful to know information about how often some given block will be executed. That is the purpose for maintaining profile within the flow graph. GCC can handle profile information obtained through profile feedback, but it can also estimate branch probabilities based on statics and heuristics.

The feedback based profile is produced by compiling the program with instrumentation, executing it on a train run and reading the numbers of executions of basic blocks and edges back to the compiler while re-compiling the program to produce the final executable. This method provides very accurate information about where a program spends most of its time on the train run. Whether it matches the average run of course depends on the choice of train data set, but several studies have shown that the behavior of a program usually changes just marginally over different data sets.

When profile feedback is not available, the compiler may be asked to attempt to predict the behavior of each branch in the program using a set of heuristics (see 'predict.def' for details) and compute estimated frequencies of each basic block by propagating the probabilities over the graph.

Each basic\_block contains two integer fields to represent profile information: frequency and count. The frequency is an estimation how often is basic block executed within a function. It is represented as an integer scaled in the range from 0 to BB\_FREQ\_BASE. The most frequently executed basic block in function is initially set to BB\_FREQ\_BASE and the rest of frequencies are scaled accordingly. During optimization, the frequency of the most frequent basic block can both decrease (for instance by loop unrolling) or grow (for instance by cross-jumping optimization), so scaling sometimes has to be performed multiple times.

The count contains hard-counted numbers of execution measured during training runs and is nonzero only when profile feedback is available. This value is represented as the host's widest integer (typically a 64 bit integer) of the special type gcov\_type.

Most optimization passes can use only the frequency information of a basic block, but a few passes may want to know hard execution counts. The frequencies should always match the counts after scaling, however during updating of the profile information numerical error may accumulate into quite large errors.

Each edge also contains a branch probability field: an integer in the range from 0 to REG\_BR\_PROB\_BASE. It represents probability of passing control from the end of the src basic block to the dest basic block, i.e. the probability that control will flow along this edge. The EDGE\_FREQUENCY macro is available to compute how frequently a given edge is taken. There is a count field for each edge as well, representing same information as for a basic block.

The basic block frequencies are not represented in the instruction stream, but in the RTL representation the edge frequencies are represented for conditional jumps (via the REG\_BR\_

PROB macro) since they are used when instructions are output to the assembly file and the flow graph is no longer maintained.

The probability that control flow arrives via a given edge to its destination basic block is called *reverse probability* and is not directly represented, but it may be easily computed from frequencies of basic blocks.

Updating profile information is a delicate task that can unfortunately not be easily integrated with the CFG manipulation API. Many of the functions and hooks to modify the CFG, such as redirect\_edge\_and\_branch, do not have enough information to easily update the profile, so updating it is in the majority of cases left up to the caller. It is difficult to uncover bugs in the profile updating code, because they manifest themselves only by producing worse code, and checking profile consistency is not possible because of numeric error accumulation. Hence special attention needs to be given to this issue in each pass that modifies the CFG.

It is important to point out that REG\_BR\_PROB\_BASE and BB\_FREQ\_BASE are both set low enough to be possible to compute second power of any frequency or probability in the flow graph, it is not possible to even square the count field, as modern CPUs are fast enough to execute \$2^32\$ operations quickly.

## 15.4 Maintaining the CFG

An important task of each compiler pass is to keep both the control flow graph and all profile information up-to-date. Reconstruction of the control flow graph after each pass is not an option, since it may be very expensive and lost profile information cannot be reconstructed at all.

GCC has two major intermediate representations, and both use the basic\_block and edge data types to represent control flow. Both representations share as much of the CFG maintenance code as possible. For each representation, a set of hooks is defined so that each representation can provide its own implementation of CFG manipulation routines when necessary. These hooks are defined in 'cfghooks.h'. There are hooks for almost all common CFG manipulations, including block splitting and merging, edge redirection and creating and deleting basic blocks. These hooks should provide everything you need to maintain and manipulate the CFG in both the RTL and GIMPLE representation.

At the moment, the basic block boundaries are maintained transparently when modifying instructions, so there rarely is a need to move them manually (such as in case someone wants to output instruction outside basic block explicitly).

In the RTL representation, each instruction has a BLOCK\_FOR\_INSN value that represents pointer to the basic block that contains the instruction. In the GIMPLE representation, the function gimple\_bb returns a pointer to the basic block containing the queried statement.

When changes need to be applied to a function in its GIMPLE representation, GIMPLE statement iterators should be used. These iterators provide an integrated abstraction of the flow graph and the instruction stream. Block statement iterators are constructed using the gimple\_stmt\_iterator data structure and several modifiers are available, including the following:

#### gsi\_start

This function initializes a gimple\_stmt\_iterator that points to the first non-empty statement in a basic block.

gsi\_last This function initializes a gimple\_stmt\_iterator that points to the last statement in a basic block.

## gsi\_end\_p

This predicate is true if a gimple\_stmt\_iterator represents the end of a basic block.

gsi\_next This function takes a gimple\_stmt\_iterator and makes it point to its successor.

gsi\_prev This function takes a gimple\_stmt\_iterator and makes it point to its predecessor.

## gsi\_insert\_after

This function inserts a statement after the gimple\_stmt\_iterator passed in. The final parameter determines whether the statement iterator is updated to point to the newly inserted statement, or left pointing to the original statement.

### gsi\_insert\_before

This function inserts a statement before the gimple\_stmt\_iterator passed in. The final parameter determines whether the statement iterator is updated to point to the newly inserted statement, or left pointing to the original statement.

### gsi\_remove

This function removes the gimple\_stmt\_iterator passed in and rechains the remaining statements in a basic block, if any.

In the RTL representation, the macros BB\_HEAD and BB\_END may be used to get the head and end rtx of a basic block. No abstract iterators are defined for traversing the insn chain, but you can just use NEXT\_INSN and PREV\_INSN instead. See Section 14.19 [Insns], page 304.

Usually a code manipulating pass simplifies the instruction stream and the flow of control, possibly eliminating some edges. This may for example happen when a conditional jump is replaced with an unconditional jump. Updating of edges is not transparent and each optimization pass is required to do so manually. However only few cases occur in practice. The pass may call purge\_dead\_edges on a given basic block to remove superfluous edges, if any.

Another common scenario is redirection of branch instructions, but this is best modeled as redirection of edges in the control flow graph and thus use of redirect\_edge\_and\_branch is preferred over more low level functions, such as redirect\_jump that operate on RTL chain only. The CFG hooks defined in 'cfghooks.h' should provide the complete API required for manipulating and maintaining the CFG.

It is also possible that a pass has to insert control flow instruction into the middle of a basic block, thus creating an entry point in the middle of the basic block, which is impossible by definition: The block must be split to make sure it only has one entry point, i.e. the head of the basic block. The CFG hook split\_block may be used when an instruction in the middle of a basic block has to become the target of a jump or branch instruction.

For a global optimizer, a common operation is to split edges in the flow graph and insert instructions on them. In the RTL representation, this can be easily done using the insert\_insn\_on\_edge function that emits an instruction "on the edge", caching it for a later commit\_edge\_insertions call that will take care of moving the inserted instructions

off the edge into the instruction stream contained in a basic block. This includes the creation of new basic blocks where needed. In the GIMPLE representation, the equivalent functions are gsi\_insert\_on\_edge which inserts a block statement iterator on an edge, and gsi\_commit\_edge\_inserts which flushes the instruction to actual instruction stream.

While debugging the optimization pass, the verify\_flow\_info function may be useful to find bugs in the control flow graph updating code.

## 15.5 Liveness information

Liveness information is useful to determine whether some register is "live" at given point of program, i.e. that it contains a value that may be used at a later point in the program. This information is used, for instance, during register allocation, as the pseudo registers only need to be assigned to a unique hard register or to a stack slot if they are live. The hard registers and stack slots may be freely reused for other values when a register is dead.

Liveness information is available in the back end starting with pass\_df\_initialize and ending with pass\_df\_finish. Three flavors of live analysis are available: With LR, it is possible to determine at any point P in the function if the register may be used on some path from P to the end of the function. With UR, it is possible to determine if there is a path from the beginning of the function to P that defines the variable. LIVE is the intersection of the LR and UR and a variable is live at P if there is both an assignment that reaches it from the beginning of the function and a use that can be reached on some path from P to the end of the function.

In general LIVE is the most useful of the three. The macros DF\_[LR,UR,LIVE]\_[IN,OUT] can be used to access this information. The macros take a basic block number and return a bitmap that is indexed by the register number. This information is only guaranteed to be up to date after calls are made to df\_analyze. See the file df-core.c for details on using the dataflow.

The liveness information is stored partly in the RTL instruction stream and partly in the flow graph. Local information is stored in the instruction stream: Each instruction may contain REG\_DEAD notes representing that the value of a given register is no longer needed, or REG\_UNUSED notes representing that the value computed by the instruction is never used. The second is useful for instructions computing multiple values at once.

# 16 Analysis and Representation of Loops

GCC provides extensive infrastructure for work with natural loops, i.e., strongly connected components of CFG with only one entry block. This chapter describes representation of loops in GCC, both on GIMPLE and in RTL, as well as the interfaces to loop-related analyses (induction variable analysis and number of iterations analysis).

## 16.1 Loop representation

This chapter describes the representation of loops in GCC, and functions that can be used to build, modify and analyze this representation. Most of the interfaces and data structures are declared in 'cfgloop.h'. Loop structures are analyzed and this information disposed or updated at the discretion of individual passes. Still most of the generic CFG manipulation routines are aware of loop structures and try to keep them up-to-date. By this means an increasing part of the compilation pipeline is setup to maintain loop structure across passes to allow attaching meta information to individual loops for consumption by later passes.

In general, a natural loop has one entry block (header) and possibly several back edges (latches) leading to the header from the inside of the loop. Loops with several latches may appear if several loops share a single header, or if there is a branching in the middle of the loop. The representation of loops in GCC however allows only loops with a single latch. During loop analysis, headers of such loops are split and forwarder blocks are created in order to disambiguate their structures. Heuristic based on profile information and structure of the induction variables in the loops is used to determine whether the latches correspond to sub-loops or to control flow in a single loop. This means that the analysis sometimes changes the CFG, and if you run it in the middle of an optimization pass, you must be able to deal with the new blocks. You may avoid CFG changes by passing LOOPS\_MAY\_HAVE\_MULTIPLE\_LATCHES flag to the loop discovery, note however that most other loop manipulation functions will not work correctly for loops with multiple latch edges (the functions that only query membership of blocks to loops and subloop relationships, or enumerate and test loop exits, can be expected to work).

Body of the loop is the set of blocks that are dominated by its header, and reachable from its latch against the direction of edges in CFG. The loops are organized in a containment hierarchy (tree) such that all the loops immediately contained inside loop L are the children of L in the tree. This tree is represented by the struct loops structure. The root of this tree is a fake loop that contains all blocks in the function. Each of the loops is represented in a struct loop structure. Each loop is assigned an index (num field of the struct loop structure), and the pointer to the loop is stored in the corresponding field of the larray vector in the loops structure. The indices do not have to be continuous, there may be empty (NULL) entries in the larray created by deleting loops. Also, there is no guarantee on the relative order of a loop and its subloops in the numbering. The index of a loop never changes.

The entries of the larray field should not be accessed directly. The function get\_loop returns the loop description for a loop with the given index. number\_of\_loops function returns number of loops in the function. To traverse all loops, use FOR\_EACH\_LOOP macro. The flags argument of the macro is used to determine the direction of traversal and the set of loops visited. Each loop is guaranteed to be visited exactly once, regardless of the changes to the loop tree, and the loops may be removed during the traversal. The newly

created loops are never traversed, if they need to be visited, this must be done separately after their creation.

Each basic block contains the reference to the innermost loop it belongs to (loop\_father). For this reason, it is only possible to have one struct loops structure initialized at the same time for each CFG. The global variable current\_loops contains the struct loops structure. Many of the loop manipulation functions assume that dominance information is up-to-date.

The loops are analyzed through loop\_optimizer\_init function. The argument of this function is a set of flags represented in an integer bitmask. These flags specify what other properties of the loop structures should be calculated/enforced and preserved later:

- LOOPS\_MAY\_HAVE\_MULTIPLE\_LATCHES: If this flag is set, no changes to CFG will be performed in the loop analysis, in particular, loops with multiple latch edges will not be disambiguated. If a loop has multiple latches, its latch block is set to NULL. Most of the loop manipulation functions will not work for loops in this shape. No other flags that require CFG changes can be passed to loop\_optimizer\_init.
- LOOPS\_HAVE\_PREHEADERS: Forwarder blocks are created in such a way that each loop has only one entry edge, and additionally, the source block of this entry edge has only one successor. This creates a natural place where the code can be moved out of the loop, and ensures that the entry edge of the loop leads from its immediate super-loop.
- LOOPS\_HAVE\_SIMPLE\_LATCHES: Forwarder blocks are created to force the latch block of each loop to have only one successor. This ensures that the latch of the loop does not belong to any of its sub-loops, and makes manipulation with the loops significantly easier. Most of the loop manipulation functions assume that the loops are in this shape. Note that with this flag, the "normal" loop without any control flow inside and with one exit consists of two basic blocks.
- LOOPS\_HAVE\_MARKED\_IRREDUCIBLE\_REGIONS: Basic blocks and edges in the strongly connected components that are not natural loops (have more than one entry block) are marked with BB\_IRREDUCIBLE\_LOOP and EDGE\_IRREDUCIBLE\_LOOP flags. The flag is not set for blocks and edges that belong to natural loops that are in such an irreducible region (but it is set for the entry and exit edges of such a loop, if they lead to/from this region).
- LOOPS\_HAVE\_RECORDED\_EXITS: The lists of exits are recorded and updated for each loop. This makes some functions (e.g., get\_loop\_exit\_edges) more efficient. Some functions (e.g., single\_exit) can be used only if the lists of exits are recorded.

These properties may also be computed/enforced later, using functions create\_preheaders, force\_single\_succ\_latches, mark\_irreducible\_loops and record\_loop\_exits. The properties can be queried using loops\_state\_satisfies\_p.

The memory occupied by the loops structures should be freed with loop\_optimizer\_finalize function. When loop structures are setup to be preserved across passes this function reduces the information to be kept up-to-date to a minimum (only LOOPS\_MAY\_HAVE\_MULTIPLE\_LATCHES set).

The CFG manipulation functions in general do not update loop structures. Specialized versions that additionally do so are provided for the most common tasks. On GIMPLE, cleanup\_tree\_cfg\_loop function can be used to cleanup CFG while updating the loops structures if current\_loops is set.

At the moment loop structure is preserved from the start of GIMPLE loop optimizations until the end of RTL loop optimizations. During this time a loop can be tracked by its struct loop and number.

## 16.2 Loop querying

The functions to query the information about loops are declared in 'cfgloop.h'. Some of the information can be taken directly from the structures. loop\_father field of each basic block contains the innermost loop to that the block belongs. The most useful fields of loop structure (that are kept up-to-date at all times) are:

- header, latch: Header and latch basic blocks of the loop.
- num\_nodes: Number of basic blocks in the loop (including the basic blocks of the sub-loops).
- outer, inner, next: The super-loop, the first sub-loop, and the sibling of the loop in the loops tree.

There are other fields in the loop structures, many of them used only by some of the passes, or not updated during CFG changes; in general, they should not be accessed directly.

The most important functions to query loop structures are:

- loop\_depth: The depth of the loop in the loops tree, i.e., the number of super-loops of the loop.
- flow\_loops\_dump: Dumps the information about loops to a file.
- verify\_loop\_structure: Checks consistency of the loop structures.
- loop\_latch\_edge: Returns the latch edge of a loop.
- loop\_preheader\_edge: If loops have preheaders, returns the preheader edge of a loop.
- flow\_loop\_nested\_p: Tests whether loop is a sub-loop of another loop.
- flow\_bb\_inside\_loop\_p: Tests whether a basic block belongs to a loop (including its sub-loops).
- find\_common\_loop: Finds the common super-loop of two loops.
- superloop\_at\_depth: Returns the super-loop of a loop with the given depth.
- tree\_num\_loop\_insns, num\_loop\_insns: Estimates the number of insns in the loop, on GIMPLE and on RTL.
- loop\_exit\_edge\_p: Tests whether edge is an exit from a loop.
- mark\_loop\_exit\_edges: Marks all exit edges of all loops with EDGE\_LOOP\_EXIT flag.
- get\_loop\_body, get\_loop\_body\_in\_dom\_order, get\_loop\_body\_in\_bfs\_order: Enumerates the basic blocks in the loop in depth-first search order in reversed CFG, ordered by dominance relation, and breath-first search order, respectively.
- single\_exit: Returns the single exit edge of the loop, or NULL if the loop has more than one exit. You can only use this function if LOOPS\_HAVE\_MARKED\_SINGLE\_EXITS property is used.
- get\_loop\_exit\_edges: Enumerates the exit edges of a loop.
- just\_once\_each\_iteration\_p: Returns true if the basic block is executed exactly once during each iteration of a loop (that is, it does not belong to a sub-loop, and it dominates the latch of the loop).

# 16.3 Loop manipulation

The loops tree can be manipulated using the following functions:

- flow\_loop\_tree\_node\_add: Adds a node to the tree.
- flow\_loop\_tree\_node\_remove: Removes a node from the tree.
- add\_bb\_to\_loop: Adds a basic block to a loop.
- remove\_bb\_from\_loops: Removes a basic block from loops.

Most low-level CFG functions update loops automatically. The following functions handle some more complicated cases of CFG manipulations:

- remove\_path: Removes an edge and all blocks it dominates.
- split\_loop\_exit\_edge: Splits exit edge of the loop, ensuring that PHI node arguments remain in the loop (this ensures that loop-closed SSA form is preserved). Only useful on GIMPLE.

Finally, there are some higher-level loop transformations implemented. While some of them are written so that they should work on non-innermost loops, they are mostly untested in that case, and at the moment, they are only reliable for the innermost loops:

- create\_iv: Creates a new induction variable. Only works on GIMPLE. standard\_iv\_increment\_position can be used to find a suitable place for the iv increment.
- duplicate\_loop\_to\_header\_edge, tree\_duplicate\_loop\_to\_header\_edge: These functions (on RTL and on GIMPLE) duplicate the body of the loop prescribed number of times on one of the edges entering loop header, thus performing either loop unrolling or loop peeling. can\_duplicate\_loop\_p (can\_unroll\_loop\_p on GIMPLE) must be true for the duplicated loop.
- loop\_version: This function creates a copy of a loop, and a branch before them that selects one of them depending on the prescribed condition. This is useful for optimizations that need to verify some assumptions in runtime (one of the copies of the loop is usually left unchanged, while the other one is transformed in some way).
- tree\_unroll\_loop: Unrolls the loop, including peeling the extra iterations to make the number of iterations divisible by unroll factor, updating the exit condition, and removing the exits that now cannot be taken. Works only on GIMPLE.

# 16.4 Loop-closed SSA form

Throughout the loop optimizations on tree level, one extra condition is enforced on the SSA form: No SSA name is used outside of the loop in that it is defined. The SSA form satisfying this condition is called "loop-closed SSA form" – LCSSA. To enforce LCSSA, PHI nodes must be created at the exits of the loops for the SSA names that are used outside of them. Only the real operands (not virtual SSA names) are held in LCSSA, in order to save memory.

There are various benefits of LCSSA:

• Many optimizations (value range analysis, final value replacement) are interested in the values that are defined in the loop and used outside of it, i.e., exactly those for that we create new PHI nodes.

- In induction variable analysis, it is not necessary to specify the loop in that the analysis should be performed the scalar evolution analysis always returns the results with respect to the loop in that the SSA name is defined.
- It makes updating of SSA form during loop transformations simpler. Without LCSSA, operations like loop unrolling may force creation of PHI nodes arbitrarily far from the loop, while in LCSSA, the SSA form can be updated locally. However, since we only keep real operands in LCSSA, we cannot use this advantage (we could have local updating of real operands, but it is not much more efficient than to use generic SSA form updating for it as well; the amount of changes to SSA is the same).

However, it also means LCSSA must be updated. This is usually straightforward, unless you create a new value in loop and use it outside, or unless you manipulate loop exit edges (functions are provided to make these manipulations simple). rewrite\_into\_loop\_closed\_ssa is used to rewrite SSA form to LCSSA, and verify\_loop\_closed\_ssa to check that the invariant of LCSSA is preserved.

## 16.5 Scalar evolutions

Scalar evolutions (SCEV) are used to represent results of induction variable analysis on GIMPLE. They enable us to represent variables with complicated behavior in a simple and consistent way (we only use it to express values of polynomial induction variables, but it is possible to extend it). The interfaces to SCEV analysis are declared in 'tree-scalar-evolution.h'. To use scalar evolutions analysis, scev\_initialize must be used. To stop using SCEV, scev\_finalize should be used. SCEV analysis caches results in order to save time and memory. This cache however is made invalid by most of the loop transformations, including removal of code. If such a transformation is performed, scev\_reset must be called to clean the caches.

Given an SSA name, its behavior in loops can be analyzed using the analyze\_scalar\_evolution function. The returned SCEV however does not have to be fully analyzed and it may contain references to other SSA names defined in the loop. To resolve these (potentially recursive) references, instantiate\_parameters or resolve\_mixers functions must be used. instantiate\_parameters is useful when you use the results of SCEV only for some analysis, and when you work with whole nest of loops at once. It will try replacing all SSA names by their SCEV in all loops, including the super-loops of the current loop, thus providing a complete information about the behavior of the variable in the loop nest. resolve\_mixers is useful if you work with only one loop at a time, and if you possibly need to create code based on the value of the induction variable. It will only resolve the SSA names defined in the current loop, leaving the SSA names defined outside unchanged, even if their evolution in the outer loops is known.

The SCEV is a normal tree expression, except for the fact that it may contain several special tree nodes. One of them is SCEV\_NOT\_KNOWN, used for SSA names whose value cannot be expressed. The other one is POLYNOMIAL\_CHREC. Polynomial chrec has three arguments—base, step and loop (both base and step may contain further polynomial chrecs). Type of the expression and of base and step must be the same. A variable has evolution POLYNOMIAL\_CHREC(base, step, loop) if it is (in the specified loop) equivalent to x\_1 in the following example

while (...)

```
{
  x_1 = phi (base, x_2);
  x_2 = x_1 + step;
}
```

Note that this includes the language restrictions on the operations. For example, if we compile C code and x has signed type, then the overflow in addition would cause undefined behavior, and we may assume that this does not happen. Hence, the value with this SCEV cannot overflow (which restricts the number of iterations of such a loop).

In many cases, one wants to restrict the attention just to affine induction variables. In this case, the extra expressive power of SCEV is not useful, and may complicate the optimizations. In this case, simple\_iv function may be used to analyze a value – the result is a loop-invariant base and step.

# 16.6 IV analysis on RTL

The induction variable on RTL is simple and only allows analysis of affine induction variables, and only in one loop at once. The interface is declared in 'cfgloop.h'. Before analyzing induction variables in a loop L, iv\_analysis\_loop\_init function must be called on L. After the analysis (possibly calling iv\_analysis\_loop\_init for several loops) is finished, iv\_analysis\_done should be called. The following functions can be used to access the results of the analysis:

- iv\_analyze: Analyzes a single register used in the given insn. If no use of the register in this insn is found, the following insns are scanned, so that this function can be called on the insn returned by get\_condition.
- iv\_analyze\_result: Analyzes result of the assignment in the given insn.
- iv\_analyze\_expr: Analyzes a more complicated expression. All its operands are analyzed by iv\_analyze, and hence they must be used in the specified insn or one of the following insns.

The description of the induction variable is provided in struct rtx\_iv. In order to handle subregs, the representation is a bit complicated; if the value of the extend field is not UNKNOWN, the value of the induction variable in the i-th iteration is

```
delta + mult * extend_{extend_mode} (subreg_{mode} (base + i * step)),
```

with the following exception: if first\_special is true, then the value in the first iteration (when i is zero) is delta + mult \* base. However, if extend is equal to UNKNOWN, then first\_special must be false, delta 0, mult 1 and the value in the i-th iteration is

```
subreg_{mode} (base + i * step)
```

The function get\_iv\_value can be used to perform these calculations.

# 16.7 Number of iterations analysis

Both on GIMPLE and on RTL, there are functions available to determine the number of iterations of a loop, with a similar interface. The number of iterations of a loop in GCC is defined as the number of executions of the loop latch. In many cases, it is not possible to determine the number of iterations unconditionally – the determined number is correct only if some assumptions are satisfied. The analysis tries to verify these conditions using the information contained in the program; if it fails, the conditions are returned together with the result. The following information and conditions are provided by the analysis:

- assumptions: If this condition is false, the rest of the information is invalid.
- noloop\_assumptions on RTL, may\_be\_zero on GIMPLE: If this condition is true, the loop exits in the first iteration.
- infinite: If this condition is true, the loop is infinite. This condition is only available on RTL. On GIMPLE, conditions for finiteness of the loop are included in assumptions.
- niter\_expr on RTL, niter on GIMPLE: The expression that gives number of iterations. The number of iterations is defined as the number of executions of the loop latch.

Both on GIMPLE and on RTL, it necessary for the induction variable analysis framework to be initialized (SCEV on GIMPLE, loop-iv on RTL). On GIMPLE, the results are stored to struct tree\_niter\_desc structure. Number of iterations before the loop is exited through a given exit can be determined using number\_of\_iterations\_exit function. On RTL, the results are returned in struct niter\_desc structure. The corresponding function is named check\_simple\_exit. There are also functions that pass through all the exits of a loop and try to find one with easy to determine number of iterations - find\_loop\_niter on GIMPLE and find\_simple\_exit on RTL. Finally, there are functions that provide the same information, but additionally cache it, so that repeated calls to number of iterations are not so costly - number\_of\_latch\_executions on GIMPLE and get\_simple\_loop\_desc on RTL.

Note that some of these functions may behave slightly differently than others — some of them return only the expression for the number of iterations, and fail if there are some assumptions. The function number\_of\_latch\_executions works only for single-exit loops. The function number\_of\_cond\_exit\_executions can be used to determine number of executions of the exit condition of a single-exit loop (i.e., the number\_of\_latch\_executions increased by one).

On GIMPLE, below constraint flags affect semantics of some APIs of number of iterations analyzer:

- LOOP\_C\_INFINITE: If this constraint flag is set, the loop is known to be infinite. APIs like number\_of\_iterations\_exit can return false directly without doing any analysis.
- LOOP\_C\_FINITE: If this constraint flag is set, the loop is known to be finite, in other words, loop's number of iterations can be computed with assumptions be true.

Generally, the constraint flags are set/cleared by consumers which are loop optimizers. It's also the consumers' responsibility to set/clear constraints correctly. Failing to do that might result in hard to track down bugs in scev/niter consumers. One typical use case is vectorizer: it drives number of iterations analyzer by setting LOOP\_C\_FINITE and vectorizes possibly infinite loop by versioning loop with analysis result. In return, constraints set by consumers can also help number of iterations analyzer in following optimizers. For example, niter of a loop versioned under assumptions is valid unconditionally.

Other constraints may be added in the future, for example, a constraint indicating that loops' latch must roll thus may\_be\_zero would be false unconditionally.

# 16.8 Data Dependency Analysis

The code for the data dependence analysis can be found in 'tree-data-ref.c' and its interface and data structures are described in 'tree-data-ref.h'. The function that computes the data dependences for all the array and pointer references for a given loop is compute\_data\_dependences\_for\_loop. This function is currently used by the linear loop transform and the vectorization passes. Before calling this function, one has to allocate two vectors: a first vector will contain the set of data references that are contained in the analyzed loop body, and the second vector will contain the dependence relations between the data references. Thus if the vector of data references is of size n, the vector containing the dependence relations will contain n\*n elements. However if the analyzed loop contains side effects, such as calls that potentially can interfere with the data references in the current analyzed loop, the analysis stops while scanning the loop body for data references, and inserts a single chrec\_dont\_know in the dependence relation array.

The data references are discovered in a particular order during the scanning of the loop body: the loop body is analyzed in execution order, and the data references of each statement are pushed at the end of the data reference array. Two data references syntactically occur in the program in the same order as in the array of data references. This syntactic order is important in some classical data dependence tests, and mapping this order to the elements of this array avoids costly queries to the loop body representation.

Three types of data references are currently handled: ARRAY\_REF, INDIRECT\_REF and COMPONENT\_REF. The data structure for the data reference is data\_reference, where data\_reference\_p is a name of a pointer to the data reference structure. The structure contains the following elements:

- base\_object\_info: Provides information about the base object of the data reference and its access functions. These access functions represent the evolution of the data reference in the loop relative to its base, in keeping with the classical meaning of the data reference access function for the support of arrays. For example, for a reference a.b[i][j], the base object is a.b and the access functions, one for each array subscript, are: {i\_init, +i\_step}\_1, {j\_init, +, j\_step}\_2.
- first\_location\_in\_loop: Provides information about the first location accessed by the data reference in the loop and about the access function used to represent evolution relative to this location. This data is used to support pointers, and is not used for arrays (for which we have base objects). Pointer accesses are represented as a one-dimensional access that starts from the first location accessed in the loop. For example:

```
for1 i
for2 j
*((int *)p + i + j) = a[i][j];
```

The access function of the pointer access is {0, +4B}\_for2 relative to p + i. The access functions of the array are {i\_init, + i\_step}\_for1 and {j\_init, +, j\_step}\_for2 relative to a.

Usually, the object the pointer refers to is either unknown, or we cannot prove that the access is confined to the boundaries of a certain object.

Two data references can be compared only if at least one of these two representations has all its fields filled for both data references.

The current strategy for data dependence tests is as follows: If both a and b are represented as arrays, compare a.base\_object and b.base\_object; if they are equal, apply dependence tests (use access functions based on base\_objects). Else if both a and b are represented as pointers, compare a.first\_location and b.first\_location; if they are equal, apply dependence tests (use access functions based on first location). However, if a and b are represented differently, only try to prove that the bases are definitely different.

- Aliasing information.
- Alignment information.

The structure describing the relation between two data references is data\_dependence\_relation and the shorter name for a pointer to such a structure is ddr\_p. This structure contains:

- a pointer to each data reference,
- a tree node are\_dependent that is set to chrec\_known if the analysis has proved that there is no dependence between these two data references, chrec\_dont\_know if the analysis was not able to determine any useful result and potentially there could exist a dependence between these data references, and are\_dependent is set to NULL\_TREE if there exist a dependence relation between the data references, and the description of this dependence relation is given in the subscripts, dir\_vects, and dist\_vects arrays,
- a boolean that determines whether the dependence relation can be represented by a classical distance vector,
- an array subscripts that contains a description of each subscript of the data references. Given two array accesses a subscript is the tuple composed of the access functions for a given dimension. For example, given A[f1][f2][f3] and B[g1][g2][g3], there are three subscripts: (f1, g1), (f2, g2), (f3, g3).
- two arrays dir\_vects and dist\_vects that contain classical representations of the data dependences under the form of direction and distance dependence vectors,
- an array of loops loop\_nest that contains the loops to which the distance and direction vectors refer to.

Several functions for pretty printing the information extracted by the data dependence analysis are available: dump\_ddrs prints with a maximum verbosity the details of a data dependence relations array, dump\_dist\_dir\_vectors prints only the classical distance and direction vectors for a data dependence relations array, and dump\_data\_references prints the details of the data references contained in a data reference array.

# 17 Machine Descriptions

A machine description has two parts: a file of instruction patterns ('.md' file) and a C header file of macro definitions.

The '.md' file for a target machine contains a pattern for each instruction that the target machine supports (or at least each instruction that is worth telling the compiler about). It may also contain comments. A semicolon causes the rest of the line to be a comment, unless the semicolon is inside a quoted string.

See the next chapter for information on the C header file.

## 17.1 Overview of How the Machine Description is Used

There are three main conversions that happen in the compiler:

- 1. The front end reads the source code and builds a parse tree.
- 2. The parse tree is used to generate an RTL insn list based on named instruction patterns.
- 3. The insu list is matched against the RTL templates to produce assembler code.

For the generate pass, only the names of the insns matter, from either a named define\_insn or a define\_expand. The compiler will choose the pattern with the right name and apply the operands according to the documentation later in this chapter, without regard for the RTL template or operand constraints. Note that the names the compiler looks for are hard-coded in the compiler—it will ignore unnamed patterns and patterns with names it doesn't know about, but if you don't provide a named pattern it needs, it will abort.

If a define\_insn is used, the template given is inserted into the insn list. If a define\_expand is used, one of three things happens, based on the condition logic. The condition logic may manually create new insns for the insn list, say via emit\_insn(), and invoke DONE. For certain named patterns, it may invoke FAIL to tell the compiler to use an alternate way of performing that task. If it invokes neither DONE nor FAIL, the template given in the pattern is inserted, as if the define\_expand were a define\_insn.

Once the insn list is generated, various optimization passes convert, replace, and rearrange the insns in the insn list. This is where the define\_split and define\_peephole patterns get used, for example.

Finally, the insn list's RTL is matched up with the RTL templates in the define\_insn patterns, and those patterns are used to emit the final assembly code. For this purpose, each named define\_insn acts like it's unnamed, since the names are ignored.

# 17.2 Everything about Instruction Patterns

A define\_insn expression is used to define instruction patterns to which insns may be matched. A define\_insn expression contains an incomplete RTL expression, with pieces to be filled in later, operand constraints that restrict how the pieces can be filled in, and an output template or C code to generate the assembler output.

A define\_insn is an RTL expression containing four or five operands:

1. An optional name n. When a name is present, the compiler automically generates a C++ function 'gen\_n' that takes the operands of the instruction as arguments and

returns the instruction's rtx pattern. The compiler also assigns the instruction a unique code 'CODE\_FOR\_n', with all such codes belonging to an enum called insn\_code.

These names serve one of two purposes. The first is to indicate that the instruction performs a certain standard job for the RTL-generation pass of the compiler, such as a move, an addition, or a conditional jump. The second is to help the target generate certain target-specific operations, such as when implementing target-specific intrinsic functions.

It is better to prefix target-specific names with the name of the target, to avoid any clash with current or future standard names.

The absence of a name is indicated by writing an empty string where the name should go. Nameless instruction patterns are never used for generating RTL code, but they may permit several simpler insns to be combined later on.

For the purpose of debugging the compiler, you may also specify a name beginning with the '\*' character. Such a name is used only for identifying the instruction in RTL dumps; it is equivalent to having a nameless pattern for all other purposes. Names beginning with the '\*' character are not required to be unique.

The name may also have the form '@n'. This has the same effect as a name 'n', but in addition tells the compiler to generate further helper functions; see Section 17.23.5 [Parameterized Names], page 477 for details.

2. The RTL template: This is a vector of incomplete RTL expressions which describe the semantics of the instruction (see Section 17.4 [RTL Template], page 339). It is incomplete because it may contain match\_operand, match\_operator, and match\_dup expressions that stand for operands of the instruction.

If the vector has multiple elements, the RTL template is treated as a parallel expression.

3. The condition: This is a string which contains a C expression. When the compiler attempts to match RTL against a pattern, the condition is evaluated. If the condition evaluates to true, the match is permitted. The condition may be an empty string, which is treated as always true.

For a named pattern, the condition may not depend on the data in the insn being matched, but only the target-machine-type flags. The compiler needs to test these conditions during initialization in order to learn exactly which named instructions are available in a particular run.

For nameless patterns, the condition is applied only when matching an individual insn, and only after the insn has matched the pattern's recognition template. The insn's operands may be found in the vector operands.

An instruction condition cannot become more restrictive as compilation progresses. If the condition accepts a particular RTL instruction at one stage of compilation, it must continue to accept that instruction until the final pass. For example, '!reload\_completed' and 'can\_create\_pseudo\_p ()' are both invalid instruction conditions, because they are true during the earlier RTL passes and false during the later ones. For the same reason, if a condition accepts an instruction before register allocation, it cannot later try to control register allocation by excluding certain register or value combinations.

Although a condition cannot become more restrictive as compilation progresses, the condition for a nameless pattern *can* become more permissive. For example, a nameless instruction can require 'reload\_completed' to be true, in which case it only matches after register allocation.

- 4. The output template or output statement: This is either a string, or a fragment of C code which returns a string.
  - When simple substitution isn't general enough, you can specify a piece of C code to compute the output. See Section 17.6 [Output Statement], page 344.
- 5. The *insn attributes*: This is an optional vector containing the values of attributes for insns matching this pattern (see Section 17.19 [Insn Attributes], page 450).

## 17.3 Example of define\_insn

Here is an example of an instruction pattern, taken from the machine description for the 68000/68020.

```
(define_insn "tstsi"
        [(set (cc0)
              (match_operand:SI 0 "general_operand" "rm"))]
      {
        if (TARGET_68020 || ! ADDRESS_REG_P (operands[0]))
          return \"tstl %0\";
        return \"cmpl #0,%0\";
This can also be written using braced strings:
      (define_insn "tstsi"
        [(set (cc0)
              (match_operand:SI 0 "general_operand" "rm"))]
      {
        if (TARGET_68020 || ! ADDRESS_REG_P (operands[0]))
          return "tstl %0";
        return "cmpl #0,%0";
      })
```

This describes an instruction which sets the condition codes based on the value of a general operand. It has no condition, so any insn with an RTL description of the form shown may be matched to this pattern. The name 'tstsi' means "test a SImode value" and tells the RTL generation pass that, when it is necessary to test such a value, an insn to do so can be constructed using this pattern.

The output control string is a piece of C code which chooses which output template to return based on the kind of operand and the specific type of CPU for which code is being generated.

""rm" is an operand constraint. Its meaning is explained below.

# 17.4 RTL Template

The RTL template is used to define which insns match the particular pattern and how to find their operands. For named patterns, the RTL template also says how to construct an insn from specified operands.

Construction involves substituting specified operands into a copy of the template. Matching involves determining the values that serve as the operands in the insn being matched. Both of these activities are controlled by special expression types that direct matching and substitution of the operands.

## (match\_operand:m n predicate constraint)

This expression is a placeholder for operand number n of the insn. When constructing an insn, operand number n will be substituted at this point. When matching an insn, whatever appears at this position in the insn will be taken as operand number n; but it must satisfy *predicate* or this instruction pattern will not match at all.

Operand numbers must be chosen consecutively counting from zero in each instruction pattern. There may be only one match\_operand expression in the pattern for each operand number. Usually operands are numbered in the order of appearance in match\_operand expressions. In the case of a define\_expand, any operand numbers used only in match\_dup expressions have higher values than all other operand numbers.

predicate is a string that is the name of a function that accepts two arguments, an expression and a machine mode. See Section 17.7 [Predicates], page 346. During matching, the function will be called with the putative operand as the expression and m as the mode argument (if m is not specified, VOIDmode will be used, which normally causes predicate to accept any mode). If it returns zero, this instruction pattern fails to match. predicate may be an empty string; then it means no test is to be done on the operand, so anything which occurs in this position is valid.

Most of the time, predicate will reject modes other than m—but not always. For example, the predicate address\_operand uses m as the mode of memory ref that the address should be valid for. Many predicates accept const\_int nodes even though their mode is VOIDmode.

constraint controls reloading and the choice of the best register class to use for a value, as explained later (see Section 17.8 [Constraints], page 350). If the constraint would be an empty string, it can be omitted.

People are often unclear on the difference between the constraint and the predicate. The predicate helps decide whether a given insn matches the pattern. The constraint plays no role in this decision; instead, it controls various decisions in the case of an insn which does match.

#### (match\_scratch:m n constraint)

This expression is also a placeholder for operand number n and indicates that operand must be a scratch or reg expression.

When matching patterns, this is equivalent to

(match\_operand:m n "scratch\_operand" constraint)

but, when generating RTL, it produces a (scratch:m) expression.

If the last few expressions in a parallel are clobber expressions whose operands are either a hard register or match\_scratch, the combiner can add or delete them when necessary. See Section 14.15 [Side Effects], page 297.

## (match\_dup n)

This expression is also a placeholder for operand number n. It is used when the operand needs to appear more than once in the insn.

In construction, match\_dup acts just like match\_operand: the operand is substituted into the insn being constructed. But in matching, match\_dup behaves differently. It assumes that operand number n has already been determined by a match\_operand appearing earlier in the recognition template, and it matches only an identical-looking expression.

Note that match\_dup should not be used to tell the compiler that a particular register is being used for two operands (example: add that adds one register to another; the second register is both an input operand and the output operand). Use a matching constraint (see Section 17.8.1 [Simple Constraints], page 350) for those. match\_dup is for the cases where one operand is used in two places in the template, such as an instruction that computes both a quotient and a remainder, where the opcode takes two input operands but the RTL template has to refer to each of those twice; once for the quotient pattern and once for the remainder pattern.

## (match\_operator:m n predicate [operands...])

This pattern is a kind of placeholder for a variable RTL expression code.

When constructing an insn, it stands for an RTL expression whose expression code is taken from that of operand n, and whose operands are constructed from the patterns operands.

When matching an expression, it matches an expression if the function *predicate* returns nonzero on that expression *and* the patterns *operands* match the operands of the expression.

Suppose that the function commutative\_operator is defined as follows, to match any expression whose operator is one of the commutative arithmetic operators of RTL and whose mode is *mode*:

Then the following pattern will match any RTL expression consisting of a commutative operator applied to two general operands:

```
(match_operator:SI 3 "commutative_operator"
  [(match_operand:SI 1 "general_operand" "g")
      (match_operand:SI 2 "general_operand" "g")])
```

Here the vector [operands...] contains two patterns because the expressions to be matched all contain two operands.

When this pattern does match, the two operands of the commutative operator are recorded as operands 1 and 2 of the insn. (This is done by the two instances

of match\_operand.) Operand 3 of the insn will be the entire commutative expression: use GET\_CODE (operands[3]) to see which commutative operator was used.

The machine mode m of match\_operator works like that of match\_operand: it is passed as the second argument to the predicate function, and that function is solely responsible for deciding whether the expression to be matched "has" that mode.

When constructing an insn, argument 3 of the gen-function will specify the operation (i.e. the expression code) for the expression to be made. It should be an RTL expression, whose expression code is copied into a new expression whose operands are arguments 1 and 2 of the gen-function. The subexpressions of argument 3 are not used; only its expression code matters.

When match\_operator is used in a pattern for matching an insn, it usually best if the operand number of the match\_operator is higher than that of the actual operands of the insn. This improves register allocation because the register allocator often looks at operands 1 and 2 of insns to see if it can do register tying.

There is no way to specify constraints in match\_operator. The operand of the insn which corresponds to the match\_operator never has any constraints because it is never reloaded as a whole. However, if parts of its operands are matched by match\_operand patterns, those parts may have constraints of their own

## (match\_op\_dup:m n[operands...])

Like  $\mathtt{match\_dup}$ , except that it applies to operators instead of operands. When constructing an insn, operand number n will be substituted at this point. But in matching,  $\mathtt{match\_op\_dup}$  behaves differently. It assumes that operand number n has already been determined by a  $\mathtt{match\_operator}$  appearing earlier in the recognition template, and it matches only an identical-looking expression.

### (match\_parallel n predicate [subpat...])

This pattern is a placeholder for an insn that consists of a parallel expression with a variable number of elements. This expression should only appear at the top level of an insn pattern.

When constructing an insn, operand number n will be substituted at this point. When matching an insn, it matches if the body of the insn is a parallel expression with at least as many elements as the vector of subpat expressions in the match\_parallel, if each subpat matches the corresponding element of the parallel, and the function predicate returns nonzero on the parallel that is the body of the insn. It is the responsibility of the predicate to validate elements of the parallel beyond those listed in the match\_parallel.

A typical use of match\_parallel is to match load and store multiple expressions, which can contain a variable number of elements in a parallel. For example,

```
(define_insn ""
  [(match_parallel 0 "load_multiple_operation"
       [(set (match_operand:SI 1 "gpc_reg_operand" "=r")
```

```
(match_operand:SI 2 "memory_operand" "m"))
  (use (reg:SI 179))
   (clobber (reg:SI 179))])]
""
"loadm 0,0,%1,%2")
```

This example comes from 'a29k.md'. The function load\_multiple\_operation is defined in 'a29k.c' and checks that subsequent elements in the parallel are the same as the set in the pattern, except that they are referencing subsequent registers and memory locations.

An insn that matches this pattern might look like:

(match\_par\_dup n [subpat...])

Like match\_op\_dup, but for match\_parallel instead of match\_operator.

## 17.5 Output Templates and Operand Substitution

The *output template* is a string which specifies how to output the assembler code for an instruction pattern. Most of the template is a fixed string which is output literally. The character '%' is used to specify where to substitute an operand; it can also be used to identify places where different variants of the assembler require different syntax.

In the simplest case, a % followed by a digit n says to output operand n at that point in the string.

'%' followed by a letter and a digit says to output an operand in an alternate fashion. Four letters have standard, built-in meanings described below. The machine description macro PRINT\_OPERAND can define additional letters with nonstandard meanings.

"%cdigit' can be used to substitute an operand that is a constant value without the syntax that normally indicates an immediate operand.

"%ndigit' is like "%cdigit' except that the value of the constant is negated before printing.

"%adigit' can be used to substitute an operand as if it were a memory reference, with the actual operand treated as the address. This may be useful when outputting a "load address" instruction, because often the assembler syntax for such an instruction requires you to write the operand as if it were a memory reference.

"%ldigit' is used to substitute a label\_ref into a jump instruction.

'%=' outputs a number which is unique to each instruction in the entire compilation. This is useful for making local labels to be referred to more than once in a single template that generates multiple assembler instructions.

"%" followed by a punctuation character specifies a substitution that does not use an operand. Only one case is standard: "%" outputs a "%" into the assembler code. Other

nonstandard cases can be defined in the PRINT\_OPERAND macro. You must also define which punctuation characters are valid with the PRINT\_OPERAND\_PUNCT\_VALID\_P macro.

The template may generate multiple assembler instructions. Write the text for the instructions, with '\;' between them.

When the RTL contains two operands which are required by constraint to match each other, the output template must refer only to the lower-numbered operand. Matching operands are not always identical, and the rest of the compiler arranges to put the proper RTL expression for printing into the lower-numbered operand.

One use of nonstandard letters or punctuation following '%' is to distinguish between different assembler languages for the same machine; for example, Motorola syntax versus MIT syntax for the 68000. Motorola syntax requires periods in most opcode names, while MIT syntax does not. For example, the opcode 'movel' in MIT syntax is 'move.1' in Motorola syntax. The same file of patterns is used for both kinds of output syntax, but the character sequence '%.' is used in each place where Motorola syntax wants a period. The PRINT\_OPERAND macro for Motorola syntax defines the sequence to output a period; the macro for MIT syntax defines it to do nothing.

As a special case, a template consisting of the single character # instructs the compiler to first split the insn, and then output the resulting instructions separately. This helps eliminate redundancy in the output templates. If you have a define\_insn that needs to emit multiple assembler instructions, and there is a matching define\_split already defined, then you can simply use # as the output template instead of writing an output template that emits the multiple assembler instructions.

Note that # only has an effect while generating assembly code; it does not affect whether a split occurs earlier. An associated define\_split must exist and it must be suitable for use after register allocation.

If the macro ASSEMBLER\_DIALECT is defined, you can use construct of the form '{option0|option1|option2}' in the templates. These describe multiple variants of assembler language syntax. See Section 18.20.7 [Instruction Output], page 616.

# 17.6 C Statements for Assembler Output

Often a single fixed template string cannot produce correct and efficient assembler code for all the cases that are recognized by a single instruction pattern. For example, the opcodes may depend on the kinds of operands; or some unfortunate combinations of operands may require extra machine instructions.

If the output control string starts with a '@', then it is actually a series of templates, each on a separate line. (Blank lines and leading spaces and tabs are ignored.) The templates correspond to the pattern's constraint alternatives (see Section 17.8.2 [Multi-Alternative], page 355). For example, if a target machine has a two-address add instruction 'addr' to add into a register and another 'addm' to add a register to memory, you might write this pattern:

```
addr %2,%0 addm %2,%0")
```

If the output control string starts with a '\*', then it is not an output template but rather a piece of C program that should compute a template. It should execute a **return** statement to return the template-string you want. Most such templates use C string literals, which require doublequote characters to delimit them. To include these doublequote characters in the string, prefix each one with '\'.

If the output control string is written as a brace block instead of a double-quoted string, it is automatically assumed to be C code. In that case, it is not necessary to put in a leading asterisk, or to escape the doublequotes surrounding C string literals.

The operands may be found in the array operands, whose C data type is rtx [].

It is very common to select different ways of generating assembler code based on whether an immediate operand is within a certain range. Be careful when doing this, because the result of INTVAL is an integer on the host machine. If the host machine has more bits in an <code>int</code> than the target machine has in the mode in which the constant will be used, then some of the bits you get from INTVAL will be superfluous. For proper results, you must carefully disregard the values of those bits.

It is possible to output an assembler instruction and then go on to output or compute more of them, using the subroutine output\_asm\_insn. This receives two arguments: a template-string and a vector of operands. The vector may be operands, or it may be another array of rtx that you declare locally and initialize yourself.

When an insn pattern has multiple alternatives in its constraints, often the appearance of the assembler code is determined mostly by which alternative was matched. When this is so, the C code can test the variable which\_alternative, which is the ordinal number of the alternative that was actually satisfied (0 for the first, 1 for the second alternative, etc.).

For example, suppose there are two opcodes for storing zero, 'clrreg' for registers and 'clrmem' for memory locations. Here is how a pattern could use which\_alternative to choose between them:

The example above, where the assembler code to generate was *solely* determined by the alternative, could also have been specified as follows, having the output control string start with a '@':

If you just need a little bit of C code in one (or a few) alternatives, you can use '\*' inside of a '@' multi-alternative template:

# 17.7 Predicates

A predicate determines whether a match\_operand or match\_operator expression matches, and therefore whether the surrounding instruction pattern will be used for that combination of operands. GCC has a number of machine-independent predicates, and you can define machine-specific predicates as needed. By convention, predicates used with match\_operand have names that end in '\_operand', and those used with match\_operator have names that end in '\_operator'.

All predicates are boolean functions (in the mathematical sense) of two arguments: the RTL expression that is being considered at that position in the instruction pattern, and the machine mode that the match\_operand or match\_operator specifies. In this section, the first argument is called op and the second argument mode. Predicates can be called from C as ordinary two-argument functions; this can be useful in output templates or other machine-specific code.

Operand predicates can allow operands that are not actually acceptable to the hardware, as long as the constraints give reload the ability to fix them up (see Section 17.8 [Constraints], page 350). However, GCC will usually generate better code if the predicates specify the requirements of the machine instructions as closely as possible. Reload cannot fix up operands that must be constants ("immediate operands"); you must use a predicate that allows only constants, or else enforce the requirement in the extra condition.

Most predicates handle their *mode* argument in a uniform manner. If *mode* is VOIDmode (unspecified), then *op* can have any mode. If *mode* is anything else, then *op* must have the same mode, unless *op* is a CONST\_INT or integer CONST\_DOUBLE. These RTL expressions always have VOIDmode, so it would be counterproductive to check that their mode matches. Instead, predicates that accept CONST\_INT and/or integer CONST\_DOUBLE check that the value stored in the constant will fit in the requested mode.

Predicates with this behavior are called *normal*. **genrecog** can optimize the instruction recognizer based on knowledge of how normal predicates treat modes. It can also diagnose certain kinds of common errors in the use of normal predicates; for instance, it is almost always an error to use a normal predicate without specifying a mode.

Predicates that do something different with their *mode* argument are called *special*. The generic predicates address\_operand and pmode\_register\_operand are special predicates. genrecog does not do any optimizations or diagnosis when special predicates are used.

# 17.7.1 Machine-Independent Predicates

These are the generic predicates available to all back ends. They are defined in 'recog.c'. The first category of predicates allow only constant, or *immediate*, operands.

### immediate\_operand

[Function]

This predicate allows any sort of constant that fits in *mode*. It is an appropriate choice for instructions that take operands that must be constant.

## const\_int\_operand

[Function]

This predicate allows any CONST\_INT expression that fits in *mode*. It is an appropriate choice for an immediate operand that does not allow a symbol or label.

### const\_double\_operand

[Function]

This predicate accepts any CONST\_DOUBLE expression that has exactly mode. If mode is VOIDmode, it will also accept CONST\_INT. It is intended for immediate floating point constants.

The second category of predicates allow only some kind of machine register.

## register\_operand

[Function]

This predicate allows any REG or SUBREG expression that is valid for *mode*. It is often suitable for arithmetic instruction operands on a RISC machine.

### pmode\_register\_operand

[Function]

This is a slight variant on register\_operand which works around a limitation in the machine-description reader.

(match\_operand n "pmode\_register\_operand" constraint)

means exactly what

(match\_operand:P n "register\_operand" constraint)

would mean, if the machine-description reader accepted ':P' mode suffixes. Unfortunately, it cannot, because Pmode is an alias for some other mode, and might vary with machine-specific options. See Section 18.31 [Misc], page 642.

## scratch\_operand

[Function]

This predicate allows hard registers and SCRATCH expressions, but not pseudoregisters. It is used internally by match\_scratch; it should not be used directly.

The third category of predicates allow only some kind of memory reference.

# memory\_operand

[Function]

This predicate allows any valid reference to a quantity of mode *mode* in memory, as determined by the weak form of GO\_IF\_LEGITIMATE\_ADDRESS (see Section 18.13 [Addressing Modes], page 562).

#### address\_operand

[Function]

This predicate is a little unusual; it allows any operand that is a valid expression for the *address* of a quantity of mode *mode*, again determined by the weak form of GO\_IF\_LEGITIMATE\_ADDRESS. To first order, if '(mem:mode (exp))' is acceptable to memory\_operand, then exp is acceptable to address\_operand. Note that exp does not necessarily have the mode *mode*.

### indirect\_operand

[Function]

This is a stricter form of memory\_operand which allows only memory references with a general\_operand as the address expression. New uses of this predicate are discouraged, because general\_operand is very permissive, so it's hard to tell what an indirect\_operand does or does not allow. If a target has different requirements for memory operands for different instructions, it is better to define target-specific predicates which enforce the hardware's requirements explicitly.

push\_operand [Function]

This predicate allows a memory reference suitable for pushing a value onto the stack. This will be a MEM which refers to stack\_pointer\_rtx, with a side effect in its address expression (see Section 14.16 [Incdec], page 302); which one is determined by the STACK\_PUSH\_CODE macro (see Section 18.9.1 [Frame Layout], page 522).

pop\_operand [Function]

This predicate allows a memory reference suitable for popping a value off the stack. Again, this will be a MEM referring to stack\_pointer\_rtx, with a side effect in its address expression. However, this time STACK\_POP\_CODE is expected.

The fourth category of predicates allow some combination of the above operands.

# nonmemory\_operand

[Function]

This predicate allows any immediate or register operand valid for mode.

## nonimmediate\_operand

[Function]

This predicate allows any register or memory operand valid for mode.

general\_operand [Function]

This predicate allows any immediate, register, or memory operand valid for mode.

Finally, there are two generic operator predicates.

#### comparison\_operator

[Function]

This predicate matches any expression which performs an arithmetic comparison in *mode*; that is, COMPARISON\_P is true for the expression code.

#### ordered\_comparison\_operator

[Function]

This predicate matches any expression which performs an arithmetic comparison in *mode* and whose expression code is valid for integer modes; that is, the expression code will be one of eq. ne, lt, ltu, le, leu, gt, gtu, ge, geu.

# 17.7.2 Defining Machine-Specific Predicates

Many machines have requirements for their operands that cannot be expressed precisely using the generic predicates. You can define additional predicates using define\_predicate and define\_special\_predicate expressions. These expressions have three operands:

- The name of the predicate, as it will be referred to in match\_operand or match\_operator expressions.
- An RTL expression which evaluates to true if the predicate allows the operand op, false if it does not. This expression can only use the following RTL codes:

#### MATCH\_OPERAND

When written inside a predicate expression, a MATCH\_OPERAND expression evaluates to true if the predicate it names would allow *op*. The operand number and constraint are ignored. Due to limitations in <code>genrecog</code>, you can only refer to generic predicates and predicates that have already been defined.

#### MATCH\_CODE

This expression evaluates to true if op or a specified subexpression of op has one of a given list of RTX codes.

The first operand of this expression is a string constant containing a comma-separated list of RTX code names (in lower case). These are the codes for which the MATCH\_CODE will be true.

The second operand is a string constant which indicates what subexpression of op to examine. If it is absent or the empty string, op itself is examined. Otherwise, the string constant must be a sequence of digits and/or lowercase letters. Each character indicates a subexpression to extract from the current expression; for the first character this is op, for the second and subsequent characters it is the result of the previous character. A digit n extracts 'XEXP (e, n)'; a letter l extracts 'XVECEXP (e, 0, n)' where n is the alphabetic ordinal of l (0 for 'a', 1 for 'b', and so on). The MATCH\_CODE then examines the RTX code of the subexpression extracted by the complete string. It is not possible to extract components of an rtvec that is not at position 0 within its RTX object.

#### MATCH\_TEST

This expression has one operand, a string constant containing a C expression. The predicate's arguments, op and mode, are available with those names in the C expression. The MATCH\_TEST evaluates to true if the C expression evaluates to a nonzero value. MATCH\_TEST expressions must not have side effects.

AND IOR

NOT

#### IF\_THEN\_ELSE

The basic 'MATCH\_' expressions can be combined using these logical operators, which have the semantics of the C operators '&&', '||', '!', and '?:' respectively. As in Common Lisp, you may give an AND or IOR expression an arbitrary number of arguments; this has exactly the same effect as writing a chain of two-argument AND or IOR expressions.

• An optional block of C code, which should execute 'return true' if the predicate is found to match and 'return false' if it does not. It must not have any side effects. The predicate arguments, op and mode, are available with those names.

If a code block is present in a predicate definition, then the RTL expression must evaluate to true *and* the code block must execute 'return true' for the predicate to allow the operand. The RTL expression is evaluated first; do not re-check anything in the code block that was checked in the RTL expression.

The program genrecog scans define\_predicate and define\_special\_predicate expressions to determine which RTX codes are possibly allowed. You should always make this explicit in the RTL predicate expression, using MATCH\_OPERAND and MATCH\_CODE.

Here is an example of a simple predicate definition, from the IA64 machine description:

Predicates written with define\_predicate automatically include a test that *mode* is VOIDmode, or *op* has the same mode as *mode*, or *op* is a CONST\_INT or CONST\_DOUBLE. They do *not* check specifically for integer CONST\_DOUBLE, nor do they test that the value of either kind of constant fits in the requested mode. This is because target-specific predicates that take constants usually have to do more stringent value checks anyway. If you need the exact same treatment of CONST\_INT or CONST\_DOUBLE that the generic predicates provide, use a MATCH\_OPERAND subexpression to call const\_int\_operand, const\_double\_operand, or immediate\_operand.

Predicates written with define\_special\_predicate do not get any automatic mode checks, and are treated as having special mode handling by genrecog.

The program genpreds is responsible for generating code to test predicates. It also writes a header file containing function declarations for all machine-specific predicates. It is not necessary to declare these predicates in 'cpu-protos.h'.

# 17.8 Operand Constraints

Each match\_operand in an instruction pattern can specify constraints for the operands allowed. The constraints allow you to fine-tune matching within the set of operands allowed by the predicate.

Constraints can say whether an operand may be in a register, and which kinds of register; whether the operand can be a memory reference, and which kinds of address; whether the operand may be an immediate constant, and which possible values it may have. Constraints can also require two operands to match. Side-effects aren't allowed in operands of inline asm, unless '<' or '>' constraints are used, because there is no guarantee that the side effects will happen exactly once in an instruction that can update the addressing register.

# 17.8.1 Simple Constraints

The simplest kind of constraint is a string full of letters, each of which describes one kind of operand that is permitted. Here are the letters that are allowed:

#### whitespace

Whitespace characters are ignored and can be inserted at any position except the first. This enables each alternative for different operands to be visually aligned in the machine description even if they have different number of constraints and modifiers.

'm' A memory operand is allowed, with any kind of address that the machine supports in general. Note that the letter used for the general memory constraint can be re-defined by a back end using the TARGET\_MEM\_CONSTRAINT macro.

'o' A memory operand is allowed, but only if the address is offsettable. This means that adding a small integer (actually, the width in bytes of the operand, as determined by its machine mode) may be added to the address and the result is also a valid memory address.

For example, an address which is constant is offsettable; so is an address that is the sum of a register and a constant (as long as a slightly larger constant is also within the range of address-offsets supported by the machine); but an autoincrement or autodecrement address is not offsettable. More complicated indirect/indexed addresses may or may not be offsettable depending on the other addressing modes that the machine supports.

Note that in an output operand which can be matched by another operand, the constraint letter 'o' is valid only when accompanied by both '<' (if the target machine has predecrement addressing) and '>' (if the target machine has preincrement addressing).

'V' A memory operand that is not offsettable. In other words, anything that would fit the 'm' constraint but not the 'o' constraint.

'<' A memory operand with autodecrement addressing (either predecrement or postdecrement) is allowed. In inline asm this constraint is only allowed if the operand is used exactly once in an instruction that can handle the side effects. Not using an operand with '<' in constraint string in the inline asm pattern at all or using it in multiple instructions isn't valid, because the side effects wouldn't be performed or would be performed more than once. Furthermore, on some targets the operand with '<' in constraint string must be accompanied by special instruction suffixes like %U0 instruction suffix on PowerPC or %P0 on IA-64.

'>' A memory operand with autoincrement addressing (either preincrement or postincrement) is allowed. In inline asm the same restrictions as for '<' apply.

'r' A register operand is allowed provided that it is in a general register.

'i' An immediate integer operand (one with constant value) is allowed. This includes symbolic constants whose values will be known only at assembly time or later.

'n' An immediate integer operand with a known numeric value is allowed. Many systems cannot support assembly-time constants for operands less than a word wide. Constraints for these operands should use 'n' rather than 'i'.

"I', 'J', 'K', ... 'P'

Other letters in the range 'I' through 'P' may be defined in a machine-dependent fashion to permit immediate integer operands with explicit integer values in specified ranges. For example, on the 68000, 'I' is defined to stand for the range of values 1 to 8. This is the range permitted as a shift count in the shift instructions.

- 'E' An immediate floating operand (expression code const\_double) is allowed, but only if the target floating point format is the same as that of the host machine (on which the compiler is running).
- 'F' An immediate floating operand (expression code const\_double or const\_vector) is allowed.
- 'G', 'H' 'G' and 'H' may be defined in a machine-dependent fashion to permit immediate floating operands in particular ranges of values.
- 's' An immediate integer operand whose value is not an explicit integer is allowed. This might appear strange; if an insn allows a constant operand with a value not known at compile time, it certainly must allow any known value. So why use 's' instead of 'i'? Sometimes it allows better code to be generated.

For example, on the 68000 in a fullword instruction it is possible to use an immediate operand; but if the immediate value is between -128 and 127, better code results from loading the value into a register and using the register. This is because the load into the register can be done with a 'moveq' instruction. We arrange for this to happen by defining the letter 'K' to mean "any integer outside the range -128 to 127", and then specifying 'Ks' in the operand constraints.

- 'g' Any register, memory or immediate integer operand is allowed, except for registers that are not general registers.
- 'X' Any operand whatsoever is allowed, even if it does not satisfy general\_operand. This is normally used in the constraint of a match\_scratch when certain alternatives will not actually require a scratch register.

'0', '1', '2', ... '9'

An operand that matches the specified operand number is allowed. If a digit is used together with letters within the same alternative, the digit should come last.

This number is allowed to be more than a single digit. If multiple digits are encountered consecutively, they are interpreted as a single decimal integer. There is scant chance for ambiguity, since to-date it has never been desirable that '10' be interpreted as matching either operand 1 or operand 0. Should this be desired, one can use multiple alternatives instead.

This is called a *matching constraint* and what it really means is that the assembler has only a single operand that fills two roles considered separate in the RTL insn. For example, an add insn has two input operands and one output operand in the RTL, but on most CISC machines an add instruction really has only two operands, one of them an input-output operand:

```
addl #35,r12
```

Matching constraints are used in these circumstances. More precisely, the two operands that match must include one input-only operand and one output-only operand. Moreover, the digit must be a smaller number than the number of the operand that uses it in the constraint.

For operands to match in a particular case usually means that they are identical-looking RTL expressions. But in a few special cases specific kinds of dissimilarity are allowed. For example, \*x as an input operand will match \*x++ as an output operand. For proper results in such cases, the output template should always use the output-operand's number when printing the operand.

'p' An operand that is a valid memory address is allowed. This is for "load address" and "push address" instructions.

'p' in the constraint must be accompanied by address\_operand as the predicate in the match\_operand. This predicate interprets the mode specified in the match\_operand as the mode of the memory reference for which the address would be valid.

other-letters

Other letters can be defined in machine-dependent fashion to stand for particular classes of registers or other arbitrary operand types. 'd', 'a' and 'f' are defined on the 68000/68020 to stand for data, address and floating point registers.

In order to have valid assembler code, each operand must satisfy its constraint. But a failure to do so does not prevent the pattern from applying to an insn. Instead, it directs the compiler to modify the code so that the constraint will be satisfied. Usually this is done by copying an operand into a register.

Contrast, therefore, the two instruction patterns that follow:

which has three operands, two of which are required by a constraint to be identical. If we are considering an insn of the form

the first pattern would not apply at all, because this inso does not contain two identical subexpressions in the right place. The pattern would say, "That does not look like an

add instruction; try other patterns". The second pattern would say, "Yes, that's an add instruction, but there is something wrong with it". It would direct the reload pass of the compiler to generate additional insns to make the constraint true. The results might look like this:

It is up to you to make sure that each operand, in each pattern, has constraints that can handle any RTL expression that could be present for that operand. (When multiple alternatives are in use, each pattern must, for each possible combination of operand expressions, have at least one alternative which can handle that combination of operands.) The constraints don't need to *allow* any possible operand—when this is the case, they do not constrain—but they must at least point the way to reloading any possible operand so that it will fit.

• If the constraint accepts whatever operands the predicate permits, there is no problem: reloading is never necessary for this operand.

For example, an operand whose constraints permit everything except registers is safe provided its predicate rejects registers.

An operand whose predicate accepts only constant values is safe provided its constraints include the letter 'i'. If any possible constant value is accepted, then nothing less than 'i' will do; if the predicate is more selective, then the constraints may also be more selective.

- Any operand expression can be reloaded by copying it into a register. So if an operand's constraints allow some kind of register, it is certain to be safe. It need not permit all classes of registers; the compiler knows how to copy a register into another register of the proper class in order to make an instruction valid.
- A nonoffsettable memory reference can be reloaded by copying the address into a register. So if the constraint uses the letter 'o', all memory references are taken care of
- A constant operand can be reloaded by allocating space in memory to hold it as preinitialized data. Then the memory reference can be used in place of the constant. So if the constraint uses the letters 'o' or 'm', constant operands are not a problem.
- If the constraint permits a constant and a pseudo register used in an insn was not allocated to a hard register and is equivalent to a constant, the register will be replaced with the constant. If the predicate does not permit a constant and the insn is rerecognized for some reason, the compiler will crash. Thus the predicate must always recognize any objects allowed by the constraint.

If the operand's predicate can recognize registers, but the constraint does not permit them, it can make the compiler crash. When this operand happens to be a register, the reload pass will be stymied, because it does not know how to copy a register temporarily into memory. If the predicate accepts a unary operator, the constraint applies to the operand. For example, the MIPS processor at ISA level 3 supports an instruction which adds two registers in SImode to produce a DImode result, but only if the registers are correctly sign extended. This predicate for the input operands accepts a sign\_extend of an SImode register. Write the constraint to indicate the type of register that is required for the operand of the sign\_extend.

# 17.8.2 Multiple Alternative Constraints

Sometimes a single instruction has multiple alternative sets of possible operands. For example, on the 68000, a logical-or instruction can combine register or an immediate value into memory, or it can combine any kind of operand into a register; but it cannot combine one memory location into another.

These constraints are represented as multiple alternatives. An alternative can be described by a series of letters for each operand. The overall constraint for an operand is made from the letters for this operand from the first alternative, a comma, the letters for this operand from the second alternative, a comma, and so on until the last alternative. All operands for a single instruction must have the same number of alternatives. Here is how it is done for fullword logical-or on the 68000:

The first alternative has 'm' (memory) for operand 0, '0' for operand 1 (meaning it must match operand 0), and 'dKs' for operand 2. The second alternative has 'd' (data register) for operand 0, '0' for operand 1, and 'dmKs' for operand 2. The '=' and '%' in the constraints apply to all the alternatives; their meaning is explained in the next section (see Section 17.8.3 [Class Preferences], page 356).

If all the operands fit any one alternative, the instruction is valid. Otherwise, for each alternative, the compiler counts how many instructions must be added to copy the operands so that that alternative applies. The alternative requiring the least copying is chosen. If two alternatives need the same amount of copying, the one that comes first is chosen. These choices can be altered with the '?' and '!' characters:

- ? Disparage slightly the alternative that the '?' appears in, as a choice when no alternative applies exactly. The compiler regards this alternative as one unit more costly for each '?' that appears in it.
- ! Disparage severely the alternative that the '!' appears in. This alternative can still be used if it fits without reloading, but if reloading is needed, some other alternative will be used.
- This constraint is analogous to '?' but it disparages slightly the alternative only if the operand with the '^' needs a reload.
- \$ This constraint is analogous to '!' but it disparages severely the alternative only if the operand with the '\$' needs a reload.

When an insn pattern has multiple alternatives in its constraints, often the appearance of the assembler code is determined mostly by which alternative was matched. When this

is so, the C code for writing the assembler code can use the variable which\_alternative, which is the ordinal number of the alternative that was actually satisfied (0 for the first, 1 for the second alternative, etc.). See Section 17.6 [Output Statement], page 344.

# 17.8.3 Register Class Preferences

The operand constraints have another function: they enable the compiler to decide which kind of hardware register a pseudo register is best allocated to. The compiler examines the constraints that apply to the insns that use the pseudo register, looking for the machine-dependent letters such as 'd' and 'a' that specify classes of registers. The pseudo register is put in whichever class gets the most "votes". The constraint letters 'g' and 'r' also vote: they vote in favor of a general register. The machine description says which registers are considered general.

Of course, on some machines all registers are equivalent, and no register classes are defined. Then none of this complexity is relevant.

#### 17.8.4 Constraint Modifier Characters

Here are constraint modifier characters.

- '=' Means that this operand is written to by this instruction: the previous value is discarded and replaced by new data.
- '+' Means that this operand is both read and written by the instruction.

When the compiler fixes up the operands to satisfy the constraints, it needs to know which operands are read by the instruction and which are written by it. '=' identifies an operand which is only written; '+' identifies an operand that is both read and written; all other operands are assumed to only be read.

If you specify '=' or '+' in a constraint, you put it in the first character of the constraint string.

- '&' Means (in a particular alternative) that this operand is an earlyclobber operand, which is written before the instruction is finished using the input operands. Therefore, this operand may not lie in a register that is read by the instruction or as part of any memory address.
  - '&' applies only to the alternative in which it is written. In constraints with multiple alternatives, sometimes one alternative requires '&' while others do not. See, for example, the 'movdf' insn of the 68000.

A operand which is read by the instruction can be tied to an earlyclobber operand if its only use as an input occurs before the early result is written. Adding alternatives of this form often allows GCC to produce better code when only some of the read operands can be affected by the earlyclobber. See, for example, the 'mulsi3' insn of the ARM.

Furthermore, if the *earlyclobber* operand is also a read/write operand, then that operand is written only after it's used.

'&' does not obviate the need to write '=' or '+'. As earlyclobber operands are always written, a read-only earlyclobber operand is ill-formed and will be rejected by the compiler.

'%' Declares the instruction to be commutative for this operand and the following operand. This means that the compiler may interchange the two operands if that is the cheapest way to make all operands fit the constraints. '%' applies to all alternatives and must appear as the first character in the constraint. Only read-only operands can use '%'.

This is often used in patterns for addition instructions that really have only two operands: the result must go in one of the arguments. Here for example, is how the 68000 halfword-add instruction is defined:

GCC can only handle one commutative pair in an asm; if you use more, the compiler may fail. Note that you need not use the modifier if the two alternatives are strictly identical; this would only waste time in the reload pass. The modifier is not operational after register allocation, so the result of define\_peephole2 and define\_splits performed after reload cannot rely on '%' to make the intended insn match.

- '#' Says that all following characters, up to the next comma, are to be ignored as a constraint. They are significant only for choosing register preferences.
- '\*' Says that the following character should be ignored when choosing register preferences. '\*' has no effect on the meaning of the constraint as a constraint, and no effect on reloading. For LRA '\*' additionally disparages slightly the alternative if the following character matches the operand.

Here is an example: the 68000 has an instruction to sign-extend a halfword in a data register, and can also sign-extend a value by copying it into an address register. While either kind of register is acceptable, the constraints on an address-register destination are less strict, so it is best if register allocation makes an address register its goal. Therefore, '\*' is used so that the 'd' constraint letter (for data register) is ignored when computing register preferences.

#### 17.8.5 Constraints for Particular Machines

Whenever possible, you should use the general-purpose constraint letters in asm arguments, since they will convey meaning more readily to people reading your code. Failing that, use the constraint letters that usually have very similar meanings across architectures. The most commonly used constraints are 'm' and 'r' (for memory and general-purpose registers respectively; see Section 17.8.1 [Simple Constraints], page 350), and 'I', usually the letter indicating the most common immediate-constant format.

Each architecture defines additional constraints. These constraints are used by the compiler itself for instruction generation, as well as for asm statements; therefore, some of the

constraints are not particularly useful for asm. Here is a summary of some of the machinedependent constraints available on some particular machines; it includes both constraints that are useful for asm and constraints that aren't. The compiler source file mentioned in the table heading for each architecture is the definitive reference for the meanings of that architecture's constraints.

# AArch64 family—'config/aarch64/constraints.md'

family- 'conf	ig/aarch64/constraints.md'		
k	The stack pointer register (SP)		
W	Floating point register, Advanced SIMD vector register or SVE vector register		
х	Like w, but restricted to registers 0 to 15 inclusive.		
у	Like w, but restricted to registers 0 to 7 inclusive.		
Upl	One of the low eight SVE predicate registers (P0 to P7)		
Upa	Any of the SVE predicate registers (P0 to P15)		
I	Integer constant that is valid as an immediate operand in an $\mathtt{ADD}$ instruction		
J	Integer constant that is valid as an immediate operand in a ${\tt SUB}$ instruction (once negated)		
K	Integer constant that can be used with a 32-bit logical instruction		
L	Integer constant that can be used with a 64-bit logical instruction		
M	Integer constant that is valid as an immediate operand in a 32-bit $MOV$ pseudo instruction. The $MOV$ may be assembled to one of several different machine instructions depending on the value		
N	Integer constant that is valid as an immediate operand in a 64-bit ${\tt MOV}$ pseudo instruction		
S	An absolute symbolic address or a label reference		
Y	Floating point constant zero		
Z	Integer constant zero		
Ush	The high part (bits 12 and upwards) of the pc-relative address of a symbol within 4GB of the instruction $$		
Q	A memory address which uses a single base register with no offset		
Ump	A memory address suitable for a load/store pair instruction in SI, DI, SF and DF modes		
${ m V-'config/gcn/constraints.md'}$			

## AMD GCN

- Ι Immediate integer in the range -16 to 64
- J Immediate 16-bit signed integer
- Κf Immediate constant -1
- L Immediate 15-bit unsigned integer

A	Immediate constant that can be inlined in an instruction encoding: integer $-1664$ , or float $0.0$ , $+/-0.5$ , $+/-1.0$ , $+/-2.0$ , $+/-4.0$ , $1.0/(2.0*PI)$
В	Immediate 32-bit signed integer that can be attached to an instruction encoding
C	Immediate 32-bit integer in range $-164294967295$ (i.e. 32-bit unsigned integer or 'A' constraint)
DA	Immediate 64-bit constant that can be split into two 'A' constants
DB	Immediate 64-bit constant that can be split into two 'B' constants
U	Any unspec
Y	Any symbol_ref or label_ref
v	VGPR register
Sg	SGPR register
SD	SGPR registers valid for instruction destinations, including VCC, $M0$ and EXEC
SS	SGPR registers valid for instruction sources, including VCC, M0, EXEC and SCC $$
Sm	SGPR registers valid as a source for scalar memory instructions (excludes M0 and EXEC)
Sv	SGPR registers valid as a source or destination for vector instructions (excludes EXEC)
ca	All condition registers: SCC, VCCZ, EXECZ
cs	Scalar condition register: SCC
cV	Vector condition register: VCC, VCC_LO, VCC_HI
е	EXEC register (EXEC_LO and EXEC_HI)
RB	Memory operand with address space suitable for buffer_* instructions
RF	Memory operand with address space suitable for flat_* instructions
RS	Memory operand with address space suitable for $s_*$ instructions
RL	Memory operand with address space suitable for $\mathtt{ds}_*$ LDS instructions
RG	Memory operand with address space suitable for $\mathtt{ds}_{-}*$ GDS instructions
RD	Memory operand with address space suitable for any $ds_*$ instructions

RM Memory operand with address space suitable for global\_\* instructions

#### ARC — 'config/arc/constraints.md'

- q Registers usable in ARCompact 16-bit instructions: r0-r3, r12-r15. This constraint can only match when the '-mq' option is in effect.
- e Registers usable as base-regs of memory addresses in ARCompact 16-bit memory instructions: r0-r3, r12-r15, sp. This constraint can only match when the '-mq' option is in effect.
- D ARC FPX (dpfp) 64-bit registers. D0, D1.
- I A signed 12-bit integer constant.
- constant for arithmetic/logical operations. This might be any constant that can be put into a long immediate by the assmbler or linker without involving a PIC relocation.
- K A 3-bit unsigned integer constant.
- L A 6-bit unsigned integer constant.
- CnL One's complement of a 6-bit unsigned integer constant.
- CmL Two's complement of a 6-bit unsigned integer constant.
- M A 5-bit unsigned integer constant.
- O A 7-bit unsigned integer constant.
- P A 8-bit unsigned integer constant.
- H Any const\_double value.

#### ARM family—'config/arm/constraints.md'

- h In Thumb state, the core registers r8-r15.
- k The stack pointer register.
- In Thumb State the core registers r0-r7. In ARM state this is an alias for the r constraint.
- t VFP floating-point registers s0-s31. Used for 32 bit values.
- VFP floating-point registers d0-d31 and the appropriate subset d0-d15 based on command line options. Used for 64 bit values only.
   Not valid for Thumb1.
- y The iWMMX co-processor registers.
- z The iWMMX GR registers.
- G The floating-point constant 0.0
- I Integer that is valid as an immediate operand in a data processing instruction. That is, an integer in the range 0 to 255 rotated by a multiple of 2

	J	Integer in the range $-4095$ to $4095$
	K	Integer that satisfies constraint 'I' when inverted (ones complement) $$
	L	Integer that satisfies constraint 'I' when negated (two complement) $$
	М	Integer in the range 0 to 32
	Q	A memory reference where the exact address is in a single register ("m" is preferable for ${\tt asm}$ statements)
	R	An item in the constant pool
	S	A symbol in the text segment of the current file
	Uv	A memory reference suitable for VFP load/store insns (reg+constant offset)
	Uy	A memory reference suitable for iWMMXt load/store instructions.
	Uq	A memory reference suitable for the ARMv4 ldrsb instruction.
AVR family	y—'config/	avr/constraints.md'
	1	Registers from r0 to r15
	a	Registers from r16 to r23
	d	Registers from r16 to r31
	W	Registers from r24 to r31. These registers can be used in 'adiw' command
	е	Pointer register (r26–r31)
	Ъ	Base pointer register (r28–r31)
	q	Stack pointer register (SPH:SPL)
	t	Temporary register r0
	x	Register pair X (r27:r26)
	у	Register pair Y (r29:r28)
	z	Register pair Z (r31:r30)
	I	Constant greater than $-1$ , less than $64$
	J	Constant greater than $-64$ , less than 1
	K	Constant integer 2
	L	Constant integer 0
	М	Constant that fits in 8 bits
	N	Constant integer $-1$
	0	Constant integer 8, 16, or 24

	P	Constant integer 1
	G	A floating point constant 0.0
	Q	A memory address based on Y or Z pointer with displacement.
Blackfin far	mily—'confi	ig/bfin/constraints.md'
	a	P register
	d	D register
	z	A call clobbered P register.
	qn	A single register. If $n$ is in the range 0 to 7, the corresponding D register. If it is A, then the register P0.
	D	Even-numbered D register
	W	Odd-numbered D register
	е	Accumulator register.
	A	Even-numbered accumulator register.
	В	Odd-numbered accumulator register.
	b	I register
	v	B register
	f	M register
	С	Registers used for circular buffering, i.e. I, B, or L registers.
	C	The CC register.
	t	LT0 or LT1.
	k	LC0 or LC1.
	u	LB0 or LB1.
	x	Any D, P, B, M, I or L register.
	у	Additional registers typically used only in prologues and epilogues: RETS, RETN, RETI, RETX, RETE, ASTAT, SEQSTAT and USP.
	W	Any register except accumulators or CC.
	Ksh	Signed 16 bit integer (in the range $-32768$ to $32767$ )
	Kuh	Unsigned 16 bit integer (in the range 0 to $65535$ )
	Ks7	Signed 7 bit integer (in the range $-64$ to $63$ )
	Ku7	Unsigned 7 bit integer (in the range 0 to 127)
	Ku5	Unsigned 5 bit integer (in the range 0 to 31)
	Ks4	Signed 4 bit integer (in the range $-8$ to 7)
	Ks3	Signed 3 bit integer (in the range $-3$ to 4)

	Ku3	Unsigned 3 bit integer (in the range 0 to 7)
	Pn	Constant $n$ , where $n$ is a single-digit constant in the range 0 to 4.
	PA	An integer equal to one of the MACFLAG_XXX constants that is suitable for use with either accumulator.
	PB	An integer equal to one of the MACFLAG_XXX constants that is suitable for use only with accumulator A1.
	M1	Constant 255.
	M2	Constant 65535.
	J	An integer constant with exactly a single bit set.
	L	An integer constant with all bits set except exactly one.
	Н	
	Q	Any SYMBOL_REF.
CR16 Arch	itecture—'co	onfig/cr16/cr16.h'
	b	Registers from r0 to r14 (registers without stack pointer)
	t	Register from r0 to r11 (all 16-bit registers)
	p	Register from r12 to r15 (all 32-bit registers)
	I	Signed constant that fits in 4 bits
	J	Signed constant that fits in 5 bits
	K	Signed constant that fits in 6 bits
	L	Unsigned constant that fits in 4 bits
	M	Signed constant that fits in 32 bits
	N	Check for 64 bits wide constants for add/sub instructions
	G	Floating point constant that is legal for store immediate
C-SKY—'c	onfig/csky	v/constraints.md'
	a	The mini registers r0 - r7.
	b	The low registers $r0 - r15$ .
	С	C register.
	У	HI and LO registers.
	1	LO register.
	h	HI register.
	v	Vector registers.
	Z	Stack pointer register (SP).

The C-SKY back end supports a large set of additional constraints that are only useful for instruction selection or splitting rather than inline asm, such as constraints representing constant integer ranges accepted by particular instruction encodings. Refer to the source code for details.

## Epiphany—'config/epiphany/constraints.md'

U16 An unsigned 16-bit constant.

K An unsigned 5-bit constant.

L A signed 11-bit constant.

Cm1 A signed 11-bit constant added to -1. Can only match when the '-m1reg-reg' option is active.

Cl1 Left-shift of -1, i.e., a bit mask with a block of leading ones, the rest being a block of trailing zeroes. Can only match when the '-m1reg-reg' option is active.

Cr1 Right-shift of -1, i.e., a bit mask with a trailing block of ones, the rest being zeroes. Or to put it another way, one less than a power of two. Can only match when the '-m1reg-reg' option is active.

Cal Constant for arithmetic/logical operations. This is like i, except that for position independent code, no symbols / expressions needing relocations are allowed.

Csy Symbolic constant for call/jump instruction.

Rcs The register class usable in short insns. This is a register class constraint, and can thus drive register allocation. This constraint won't match unless '-mprefer-short-insn-regs' is in effect.

Rsc The the register class of registers that can be used to hold a sibcall call address. I.e., a caller-saved register.

Rct Core control register class.

Rgs The register group usable in short insns. This constraint does not use a register class, so that it only passively matches suitable registers, and doesn't drive register allocation.

Car Constant suitable for the addsi3\_r pattern. This is a valid offset For byte, halfword, or word addressing.

Rra Matches the return address if it can be replaced with the link register.

Rcc Matches the integer condition code register.

Sra Matches the return address if it is in a stack slot.

Cfm Matches control register values to switch fp mode, which are encapsulated in UNSPEC\_FP\_MODE.

#### FRV—'config/frv/frv.h'

a Register in the class ACC\_REGS (acc0 to acc7).

b Register in the class EVEN\_ACC\_REGS (acc0 to acc7).

c Register in the class CC\_REGS (fcc0 to fcc3 and icc0 to icc3).

d Register in the class GPR\_REGS (gr0 to gr63).

е	Register in the class EVEN_REGS (gr0 to gr63). Odd registers are excluded not in the class but through the use of a machine mode larger than 4 bytes.
f	Register in the class FPR_REGS (fr0 to fr63).
h	Register in the class FEVEN_REGS (fr0 to fr63). Odd registers are excluded not in the class but through the use of a machine mode larger than 4 bytes.
1	Register in the class LR_REG (the lr register).
q	Register in the class QUAD_REGS (gr2 to gr63). Register numbers not divisible by 4 are excluded not in the class but through the use of a machine mode larger than 8 bytes.
t	Register in the class ICC_REGS (icc0 to icc3).
u	Register in the class FCC_REGS (fcc0 to fcc3).
V	Register in the class ICR_REGS (cc4 to cc7).
W	Register in the class FCR_REGS (cc0 to cc3).
x	Register in the class QUAD_FPR_REGS (fr0 to fr63). Register numbers not divisible by 4 are excluded not in the class but through the use of a machine mode larger than 8 bytes.
z	Register in the class SPR_REGS (lcr and lr).
A	Register in the class QUAD_ACC_REGS (acc0 to acc7).
В	Register in the class ACCG_REGS (accg0 to accg7).
C	Register in the class CR_REGS (cc0 to cc7).
G	Floating point constant zero
I	6-bit signed integer constant
J	10-bit signed integer constant
L	16-bit signed integer constant
M	16-bit unsigned integer constant
N	12-bit signed integer constant that is negative—i.e. in the range of $-2048$ to $-1$

12-bit signed integer constant that is greater than zero—i.e. in the

# FT32—'config/ft32/constraints.md'

0

Р

A An absolute address

Constant zero

range of 1 to 2047.

- B An offset address
- W A register indirect memory operand

е	An offset address.
f	An offset address.
0	The constant zero or one
I	A 16-bit signed constant $(-32768 \dots 32767)$
W	A bitfield mask suitable for bext or bins
x	An inverted bitfield mask suitable for bext or bins
L	A 16-bit unsigned constant, multiple of 4 (0 65532)
S	A 20-bit signed constant $(-524288 \dots 524287)$
b	A constant for a bitfield width (1 16)
KA	A 10-bit signed constant $(-512 \dots 511)$
Hewlett-Packard PA-RA	ISC—'config/pa/pa.h'
a	General register 1
f	Floating point register
q	Shift amount register
x	Floating point register (deprecated)
У	Upper floating point register (32-bit), floating point register (64-bit)
Z	Any register
I	Signed 11-bit integer constant
J	Signed 14-bit integer constant
K	Integer constant that can be deposited with a zdepi instruction
L	Signed 5-bit integer constant
M	Integer constant 0
N	Integer constant that can be loaded with a ldil instruction
0	Integer constant whose value plus one is a power of 2
P	Integer constant that can be used for ${\tt and}$ operations in ${\tt depi}$ and ${\tt extru}$ instructions
S	Integer constant 31
U	Integer constant 63
G	Floating-point constant 0.0
A	A lo_sum data-linkage-table memory operand
Q	A memory operand that can be used as the destination operand of an integer store instruction
R	A scaled or unscaled indexed memory operand

	Т	A memory operand for floating-point loads and stores
	W	A register indirect memory operand
Intel IA-64—'config/ia64/ia64.h'		
	a	General register r0 to r3 for addl instruction
	b	Branch register
	С	Predicate register ('c' as in "conditional")
	d	Application register residing in M-unit
	е	Application register residing in I-unit
	f	Floating-point register
	m	Memory operand. If used together with '<' or '>', the operand can have post increment and postdecrement which require printing with '%Pn' on IA-64.
	G	Floating-point constant 0.0 or 1.0
	I	14-bit signed integer constant
	J	22-bit signed integer constant
	K	8-bit signed integer constant for logical instructions
	L	8-bit adjusted signed integer constant for compare pseudo-ops
	M	6-bit unsigned integer constant for shift counts
	N	9-bit signed integer constant for load and store postincrements
	0	The constant zero
	P	0  or  -1  for dep instruction
	Q	Non-volatile memory for floating-point loads and stores
	R	Integer constant in the range 1 to 4 for shladd instruction
	S	Memory operand except post increment and postdecrement. This is now roughly the same as 'm' when not used together with '<' or '>'.
<i>M32C</i> —'co	nfig/m32c/	m32c.c'
	Rsp	
	Rfb Rsb	'\$sp', '\$fb', '\$sb'.
	Rcr	Any control register, when they're 16 bits wide (nothing if control registers are 24 bits wide)
	Rcl	Any control register, when they're 24 bits wide.
	ROw R1w R2w	¢n0 ¢n1 ¢n2 ¢n2
	R3w	\$r0, \$r1, \$r2, \$r3.

R02 \$r0 or \$r2, or \$r2r0 for 32 bit values.

R13 \$r1 or \$r3, or \$r3r1 for 32 bit values.

Rdi A register that can hold a 64 bit value.

Rhl \$r0 or \$r1 (registers with addressable high/low bytes)

R23 \$r2 or \$r3

Raa Address registers

Raw Address registers when they're 16 bits wide.

Ral Address registers when they're 24 bits wide.

Rqi Registers that can hold QI values.

Rad Registers that can be used with displacements (\$a0, \$a1, \$sb).

Rsi Registers that can hold 32 bit values.

Rhi Registers that can hold 16 bit values.

Rhc Registers chat can hold 16 bit values, including all control registers.

Rra \$r0 through R1, plus \$a0 and \$a1.

Rfl The flags register.

Rmm The memory-based pseudo-registers \$mem0 through \$mem15.

Rpi Registers that can hold pointers (16 bit registers for r8c, m16c; 24 bit registers for m32cm, m32c).

Rpa Matches multiple registers in a PARALLEL to form a larger register. Used to match function return values.

Is3 -8...7

IS1  $-128 \dots 127$ 

IS2  $-32768 \dots 32767$ 

IU2 0 . . . 65535

In4  $-8 \dots -1 \text{ or } 1 \dots 8$ 

In5  $-16 \dots -1 \text{ or } 1 \dots 16$ 

In6  $-32 \dots -1 \text{ or } 1 \dots 32$ 

IM2  $-65536 \dots -1$ 

Ilb An 8 bit value with exactly one bit set.

Ilw A 16 bit value with exactly one bit set.

Sd The common src/dest memory addressing modes.

Sa Memory addressed using \$a0 or \$a1.

Si Memory addressed with immediate addresses.

	Ss	Memory addressed using the stack pointer (\$sp).
	Sf	Memory addressed using the frame base register (\$fb).
	Ss	Memory addressed using the small base register (\$sb).
	S1	\$r1h
MicroBlaze	- 'config/m	nicroblaze/constraints.md'
	d	A general register (r0 to r31).
	z	A status register (rmsr, \$fcc1 to \$fcc7).
MIPS—'cor	nfig/mips/	constraints.md'
	d	A general-purpose register. This is equivalent to <b>r</b> unless generating MIPS16 code, in which case the MIPS16 register set is used.
	f	A floating-point register (if available).
	h	Formerly the hi register. This constraint is no longer supported.
	1	The lo register. Use this register to store values that are no bigger than a word.
	x	The concatenated hi and lo registers. Use this register to store doubleword values.
	С	A register suitable for use in an indirect jump. This will always be \$25 for '-mabicalls'.
	V	Register \$3. Do not use this constraint in new code; it is retained only for compatibility with glibc.
	у	Equivalent to r; retained for backwards compatibility.
	Z	A floating-point condition code register.
	I	A signed 16-bit constant (for arithmetic instructions).
	J	Integer zero.
	K	An unsigned 16-bit constant (for logic instructions).
	L	A signed 32-bit constant in which the lower 16 bits are zero. Such constants can be loaded using lui.
	М	A constant that cannot be loaded using lui, addiu or ori.
	N	A constant in the range $-65535$ to $-1$ (inclusive).
	0	A signed 15-bit constant.
	P	A constant in the range 1 to 65535 (inclusive).
	G	Floating-point zero.
	R	An address that can be used in a non-macro load or store.
	ZC	A memory operand whose address is formed by a base register and offset that is suitable for use in instructions with the same addressing mode as 11 and sc.

ZD An address suitable for a prefetch instruction, or for any other instruction with the same addressing mode as prefetch.

Motorola 680x0—'config/m68k/constraints.md'

$la \ b$	80x0—conf	ig/m68k/constraints.md
	a	Address register
	d	Data register
	f	68881 floating-point register, if available
	I	Integer in the range 1 to 8
	J	16-bit signed number
	K	Signed number whose magnitude is greater than $0x80$
	L	Integer in the range $-8$ to $-1$
	M	Signed number whose magnitude is greater than $0x100$
	N	Range 24 to 31, rotatert:SI 8 to 1 expressed as rotate
	0	16 (for rotate using swap)
	P	Range 8 to 15, rotatert:HI 8 to 1 expressed as rotate
	R	Numbers that mov3q can handle
	G	Floating point constant that is not a 68881 constant
	S	Operands that satisfy 'm' when -mpcrel is in effect
	T	Operands that satisfy 's' when -mpcrel is not in effect
	Q	Address register indirect addressing mode
	U	Register offset addressing
	W	$const\_call\_operand$
	Cs	symbol_ref or const
	Ci	$const\_int$
	CO	const_int 0
	Cj	Range of signed numbers that don't fit in 16 bits
	Cmvq	Integers valid for mvq
	Capsw	Integers valid for a moveq followed by a swap
	Cmvz	Integers valid for mvz
	Cmvs	Integers valid for mvs
	Ap	$push\_operand$

Non-register operands allowed in clr

Moxie-`config/moxie/constraints.md'

A An absolute address

	В	An offset address
	W	A register indirect memory operand
	I	A constant in the range of 0 to 255.
	N	A constant in the range of 0 to $-255$ .
MSP430-	config/msp	o430/constraints.md'
	R12	Register R12.
	R13	Register R13.
	K	Integer constant 1.
	L	Integer constant -1^201^19.
	M	Integer constant 1-4.
	Ya	Memory references which do not require an extended MOVX instruction.
	Yl	Memory reference, labels only.
	Ys	Memory reference, stack only.
NDS32—	config/nds	32/constraints.md'
	W	LOW register class $r0$ to $r7$ constraint for $V3/V3M$ ISA.
	1	LOW register class \$r0 to \$r7.
	d	MIDDLE register class \$r0 to \$r11, \$r16 to \$r19.
	h	HIGH register class \$r12 to \$r14, \$r20 to \$r31.
	t	Temporary assist register \$ta (i.e. \$r15).
	k	Stack register \$sp.
	Iu03	Unsigned immediate 3-bit value.
	In03	Negative immediate 3-bit value in the range of $-7-0$ .
	Iu04	Unsigned immediate 4-bit value.
	Is05	Signed immediate 5-bit value.
	Iu05	Unsigned immediate 5-bit value.
	In05	Negative immediate 5-bit value in the range of $-31$ -0.
	Ip05	Unsigned immediate 5-bit value for movpi45 instruction with range 16–47.
	Iu06	Unsigned immediate 6-bit value constraint for addri36.sp instruction.
	Iu08	Unsigned immediate 8-bit value.
	Iu09	Unsigned immediate 9-bit value.
	Is10	Signed immediate 10-bit value.

	Is11	Signed immediate 11-bit value.
	Is15	Signed immediate 15-bit value.
	Iu15	Unsigned immediate 15-bit value.
	Ic15	A constant which is not in the range of imm15u but ok for bclr instruction.
	Ie15	A constant which is not in the range of imm15u but ok for bset instruction.
	It15	A constant which is not in the range of imm15u but ok for btgl instruction.
	Ii15	A constant whose compliment value is in the range of imm15u and ok for bitci instruction.
	Is16	Signed immediate 16-bit value.
	Is17	Signed immediate 17-bit value.
	Is19	Signed immediate 19-bit value.
	Is20	Signed immediate 20-bit value.
	Ihig	The immediate value that can be simply set high 20-bit.
	Izeb	The immediate value 0xff.
	Izeh	The immediate value 0xffff.
	Ixls	The immediate value $0x01$ .
	Ix11	The immediate value 0x7ff.
	Ibms	The immediate value with power of 2.
	Ifex	The immediate value with power of 2 minus 1.
	U33	Memory constraint for 333 format.
	U45	Memory constraint for 45 format.
	U37	Memory constraint for 37 format.
am	ily—'config	g/nios2/constraints.md'
	I	Integer that is valid as an immediate operand in an instruction taking a signed 16-bit number. Range $-32768$ to $32767$

# Nios II far

- taking a signed 16-bit number. Range -32768 to 32767.
- J Integer that is valid as an immediate operand in an instruction taking an unsigned 16-bit number. Range 0 to 65535.
- Integer that is valid as an immediate operand in an instruction K taking only the upper 16-bits of a 32-bit number. Range 32-bit numbers with the lower 16-bits being 0.
- L Integer that is valid as an immediate operand for a shift instruction. Range 0 to 31.

- M Integer that is valid as an immediate operand for only the value 0. Can be used in conjunction with the format modifier z to use r0 instead of 0 in the assembly output.
- N Integer that is valid as an immediate operand for a custom instruction opcode. Range 0 to 255.
- P An immediate operand for R2 andchi/andci instructions.
- Matches immediates which are addresses in the small data section and therefore can be added to gp as a 16-bit immediate to re-create their 32-bit value.
- U Matches constants suitable as an operand for the rdprs and cache instructions.
- v A memory operand suitable for Nios II R2 load/store exclusive instructions.
- w A memory operand suitable for load/store IO and cache instructions.
- T A const wrapped UNSPEC expression, representing a supported PIC or TLS relocation.

### OpenRISC—'config/or1k/constraints.md'

- I Integer that is valid as an immediate operand in an instruction taking a signed 16-bit number. Range -32768 to 32767.
- K Integer that is valid as an immediate operand in an instruction taking an unsigned 16-bit number. Range 0 to 65535.
- M Signed 16-bit constant shifted left 16 bits. (Used with 1.movhi)
- 0 Zero
- c Register usable for sibcalls.

### PDP-11—'config/pdp11/constraints.md'

- a Floating point registers AC0 through AC3. These can be loaded from/to memory with a single instruction.
- d Odd numbered general registers (R1, R3, R5). These are used for 16-bit multiply operations.
- D A memory reference that is encoded within the opcode, but not auto-increment or auto-decrement.
- f Any of the floating point registers (AC0 through AC5).
- G Floating point constant 0.
- h Floating point registers AC4 and AC5. These cannot be loaded from/to memory with a single instruction.
- I An integer constant that fits in 16 bits.
- J An integer constant whose low order 16 bits are zero.

- K An integer constant that does not meet the constraints for codes 'I' or 'J'.
- L The integer constant 1.
- M The integer constant -1.
- N The integer constant 0.
- O Integer constants 0 through 3; shifts by these amounts are handled as multiple single-bit shifts rather than a single variable-length shift.
- A memory reference which requires an additional word (address or offset) after the opcode.
- R A memory reference that is encoded within the opcode.

#### PowerPC and IBM RS6000—'config/rs6000/constraints.md'

- r A general purpose register (GPR), r0...r31.
- b A base register. Like r, but r0 is not allowed, so r1...r31.
- f A floating point register (FPR), f0...f31.
- d A floating point register. This is the same as f nowadays; historically f was for single-precision and d was for double-precision floating point.
- v An Altivec vector register (VR), v0...v31.
- wa A VSX register (VSR), vs0...vs63. This is either an FPR (vs0...vs31 are f0...f31) or a VR (vs32...vs63 are v0...v31). When using wa, you should use the %x output modifier, so that the correct register number is printed. For example:

```
asm ("xvadddp %x0,%x1,%x2"
: "=wa" (v1)
: "wa" (v2), "wa" (v3));
```

You should not use %x for v operands:

```
asm ("xsaddqp %0,%1,%2"
: "=v" (v1)
: "v" (v2), "v" (v3));
```

- h A special register (vrsave, ctr, or lr).
- c The count register, ctr.
- 1 The link register, 1r.
- x Condition register field 0, cr0.
- y Any condition register field, cr0...cr7.
- z The carry bit, XER[CA].
- we Like wa, if '-mpower9-vector' and '-m64' are used; otherwise, NO\_REGS.

eΙ

G

ported.

one instruction per word.

wn	No register (NO_REGS).
wr	Like r, if '-mpowerpc64' is used; otherwise, NO_REGS.
WX	Like d, if '-mpowerpc-gfxopt' is used; otherwise, NO_REGS.
wA	Like b, if '-mpowerpc64' is used; otherwise, NO_REGS.
wB	Signed 5-bit constant integer that can be loaded into an Altivec register.
wD	Int constant that is the element number of the 64-bit scalar in a vector.
wE	Vector constant that can be loaded with the XXSPLTIB instruction.
wF	Memory operand suitable for power8 GPR load fusion.
wL	Int constant that is the element number mfvsrld accesses in a vector.
wM	Match vector constant with all 1's if the XXLORC instruction is available.
wO	Memory operand suitable for the ISA 3.0 vector d-form instructions.
wQ	Memory operand suitable for the load/store quad instructions.
wS	Vector constant that can be loaded with XXSPLTIB & sign extension.
wY	A memory operand for a DS-form instruction.
wZ	An indexed or indirect memory operand, ignoring the bottom $4$ bits.
I	A signed 16-bit constant.
J	An unsigned 16-bit constant shifted left 16 bits (use L instead for ${\tt SImode}$ constants).
K	An unsigned 16-bit constant.
L	A signed 16-bit constant shifted left 16 bits.
M	An integer constant greater than 31.
N	An exact power of 2.
0	The integer constant zero.
P	A constant whose negation is a signed 16-bit constant.
-	A . 1.041

A signed 34-bit integer constant if prefixed instructions are sup-

A floating point constant that can be loaded into a register with

H A floating point constant that can be loaded into a register using three instructions.

A memory operand. Normally, m does not allow addresses that update the base register. If the < or > constraint is also used, they are allowed and therefore on PowerPC targets in that case it is only safe to use m<> in an asm statement if that asm statement accesses the operand exactly once. The asm statement must also use %U<opno> as a placeholder for the "update" flag in the corresponding load or store instruction. For example:

```
asm ("st%U0 %1,%0" : "=m<>" (mem) : "r" (val));
is correct but:
    asm ("st %1,%0" : "=m<>" (mem) : "r" (val));
is not.
```

A "stable" memory operand; that is, one which does not include any automodification of the base register. This used to be useful when m allowed automodification of the base register, but as those are now only allowed when < or > is used, es is basically the same as m without < and >.

- A memory operand addressed by just a base register.
- Y A memory operand for a DQ-form instruction.
- Z A memory operand accessed with indexed or indirect addressing.
- R An AIX TOC entry.
- a An indexed or indirect address.
- U A V.4 small data reference.
- W A vector constant that does not require memory.
- j The zero vector constant.

## PRU—'config/pru/constraints.md'

- I An unsigned 8-bit integer constant.
- J An unsigned 16-bit integer constant.
- L An unsigned 5-bit integer constant (for shift counts).
- T A text segment (program memory) constant label.
- Z Integer constant zero.

#### RL78—'config/rl78/constraints.md'

- Int3 An integer constant in the range 1 . . . 7.
- Int8 An integer constant in the range 0 . . . 255.
- J An integer constant in the range  $-255 \dots 0$
- K The integer constant 1.

Zint

L	The integer constant -1.
M	The integer constant 0.
N	The integer constant 2.
0	The integer constant -2.
P	An integer constant in the range 1 15.
Qbi	The built-in compare types–eq, ne, gtu, ltu, geu, and leu.
Qsc	The synthetic compare types–gt, lt, ge, and le.
Wab	A memory reference with an absolute address.
Wbc	A memory reference using BC as a base register, with an optional offset.
Wca	A memory reference using ${\tt AX},~{\tt BC},~{\tt DE},~{\tt or}~{\tt HL}$ for the address, for calls.
Wcv	A memory reference using any 16-bit register pair for the address, for calls.
Wd2	A memory reference using DE as a base register, with an optional offset.
Wde	A memory reference using DE as a base register, without any offset.
Wfr	Any memory reference to an address in the far address space.
Wh1	A memory reference using HL as a base register, with an optional one-byte offset.
Whb	A memory reference using $\mathtt{HL}$ as a base register, with $\mathtt{B}$ or $\mathtt{C}$ as the index register.
Whl	A memory reference using HL as a base register, without any offset.
Ws1	A memory reference using SP as a base register, with an optional one-byte offset.
Y	Any memory reference to an address in the near address space.
A	The AX register.
В	The BC register.
D	The DE register.
R	A through L registers.
S	The SP register.
T	The HL register.
Z08W	The 16-bit R8 register.
Z10W	The 16-bit R10 register.
	TI ( ) ( ) ( ) ( ) ( ) ( ) ( ) ( ) ( ) (

The registers reserved for interrupts (R24 to R31).

	a	The A register.		
	Ъ	The B register.		
	С	The C register.		
	d	The D register.		
	е	The E register.		
	h	The H register.		
	1	The L register.		
	v	The virtual registers.		
	W	The PSW register.		
	x	The X register.		
RISC- $V$ -	-'config/ri	scv/constraints.md'		
	f	A floating-point register (if available).		
	I	An I-type 12-bit signed immediate.		
	J	Integer zero.		
	K	A 5-bit unsigned immediate for CSR access instructions.		
	Α	An address that is held in a general-purpose register.		
RX—'config/rx/constraints.md'				
	Q	An address which does not involve register indirect addressing or pre/post increment/decrement addressing.		
	Symbol	A symbol reference.		
	Int08	A constant in the range $-256$ to $255$ , inclusive.		
	Sint08	A constant in the range $-128$ to $127$ , inclusive.		
	a			
	Sint16	A constant in the range $-32768$ to $32767$ , inclusive.		
	Sint16 Sint24	A constant in the range $-32768$ to $32767$ , inclusive. A constant in the range $-8388608$ to $8388607$ , inclusive.		
S/390 and	Sint24 Uint04	A constant in the range $-8388608$ to $8388607$ , inclusive.		
S/390 and	Sint24 Uint04	A constant in the range -8388608 to 8388607, inclusive.  A constant in the range 0 to 15, inclusive.		
S/390 and	Sint24 Uint04 d zSeries—'c	A constant in the range -8388608 to 8388607, inclusive.  A constant in the range 0 to 15, inclusive.  config/s390/s390.h'		
S/390 and	Sint24 Uint04 d zSeries—'c	A constant in the range -8388608 to 8388607, inclusive.  A constant in the range 0 to 15, inclusive.  config/s390/s390.h'  Address register (general purpose register except r0)		
S/390 and	Sint24 Uint04 d zSeries—'c a	A constant in the range -8388608 to 8388607, inclusive.  A constant in the range 0 to 15, inclusive.  config/s390/s390.h'  Address register (general purpose register except r0)  Condition code register		
S/390 and	Sint24 Uint04 d zSeries—'d a c d	A constant in the range -8388608 to 8388607, inclusive.  A constant in the range 0 to 15, inclusive.  config/s390/s390.h'  Address register (general purpose register except r0)  Condition code register  Data register (arbitrary general purpose register)		
S/390 and	Sint24 Uint04 d zSeries—'d a c d	A constant in the range -8388608 to 8388607, inclusive.  A constant in the range 0 to 15, inclusive.  config/s390/s390.h'  Address register (general purpose register except r0)  Condition code register  Data register (arbitrary general purpose register)  Floating-point register		

L Value appropriate as displacement.

(0..4095)

for short displacement

(-524288..524287)

for long displacement

- M Constant integer with a value of 0x7fffffff.
- N Multiple letter constraint followed by 4 parameter letters.
  - 0..9: number of the part counting from most to least significant
  - H,Q: mode of the part
  - D,S,H: mode of the containing operand
  - 0,F: value of the other parts (F—all bits set)

The constraint matches if the specified part of a constant has a value different from its other parts.

- Q Memory reference without index register and with short displacement.
- R Memory reference with index register and short displacement.
- S Memory reference without index register but with long displacement.
- T Memory reference with index register and long displacement.
- U Pointer with short displacement.
- W Pointer with long displacement.
- Y Shift count operand.

# SPARC—'config/sparc/sparc.h'

- f Floating-point register on the SPARC-V8 architecture and lower floating-point register on the SPARC-V9 architecture.
- e Floating-point register. It is equivalent to 'f' on the SPARC-V8 architecture and contains both lower and upper floating-point registers on the SPARC-V9 architecture.
- c Floating-point condition code register.
- d Lower floating-point register. It is only valid on the SPARC-V9 architecture when the Visual Instruction Set is available.
- b Floating-point register. It is only valid on the SPARC-V9 architecture when the Visual Instruction Set is available.
- h 64-bit global or out register for the SPARC-V8+ architecture.
- C The constant all-ones, for floating-point.

А	Signed 5-bit constant		
D	A vector constant		
I	Signed 13-bit constant		
J	Zero		
K	32-bit constant with the low 12 bits clear (a constant that can be loaded with the sethi instruction)		
L	A constant in the range supported by movcc instructions (11-bit signed immediate)		
М	A constant in the range supported by movrcc instructions (10-bit signed immediate)		
N	Same as 'K', except that it verifies that bits that are not in the lower 32-bit range are all zero. Must be used instead of 'K' for modes wider than SImode		
0	The constant 4096		
G	Floating-point zero		
Н	Signed 13-bit constant, sign-extended to 32 or 64 bits		
P	The constant -1		
Q	Floating-point constant whose integral representation can be moved into an integer register using a single sethi instruction		
R	Floating-point constant whose integral representation can be moved into an integer register using a single mov instruction		
S	Floating-point constant whose integral representation can be moved into an integer register using a high/lo_sum instruction sequence		
T	Memory address aligned to an 8-byte boundary		
U	Even register		
W	Memory address for 'e' constraint registers		
w	Memory address with only a base register		
Y	Vector zero		
TI~C6X~family—'config/c6x/constraints.md'			
a	Register file A (A0–A31).		
b	Register file B (B0–B31).		
A	Predicate registers in register file A (A0–A2 on C64X and higher, A1 and A2 otherwise).		
В	Predicate registers in register file B (B0–B2).		
C	A call-used register in register file B (B0–B9, B16–B31).		

Da

```
not C64X or higher).
            Db
                        Register file B, excluding predicate registers (B3–B31).
                        Integer constant in the range 0 \dots 15.
            Iu4
                        Integer constant in the range 0 \dots 31.
            Iu5
            In5
                        Integer constant in the range -31 \dots 0.
            Is5
                        Integer constant in the range -16 \dots 15.
            I5x
                        Integer constant that can be the operand of an ADDA or a SUBA
                        insn.
            IuB
                        Integer constant in the range 0 . . . 65535.
                        Integer constant in the range -32768 \dots 32767.
            IsB
                        Integer constant in the range -2^{20} \dots 2^{20} - 1.
            IsC
                        Integer constant that is a valid mask for the clr instruction.
            Jc
            Js
                        Integer constant that is a valid mask for the set instruction.
            Q
                        Memory location with A base register.
            R
                        Memory location with B base register.
            SO
                        On C64x+ targets, a GP-relative small data reference.
            S1
                        Any kind of SYMBOL_REF, for use in a call address.
                        Any kind of immediate operand, unless it matches the S0 con-
            Si
                        straint.
            Τ
                        Memory location with B base register, but not using a long offset.
            W
                        A memory operand with an address that cannot be used in an
                        unaligned access.
            Z
                        Register B14 (aka DP).
TILE-Gx—'config/tilegx/constraints.md'
            R00
            R01
            R02
            R03
            R04
            R05
            R06
            R07
            R08
            R09
                        Each of these represents a register constraint for an individual reg-
            R10
                        ister, from r0 to r10.
            Ι
                        Signed 8-bit integer constant.
```

Register file A, excluding predicate registers (A3–A31, plus A0 if

```
J
                        Signed 16-bit integer constant.
            K
                        Unsigned 16-bit integer constant.
            L
                        Integer constant that fits in one signed byte when incremented by
                        one (-129 \dots 126).
                        Memory operand. If used together with '<' or '>', the operand can
            m
                        have postincrement which requires printing with '%In' and '%in' on
                        TILE-Gx. For example:
                              asm ("st_add %I0,%1,%i0" : "=m<>" (*mem) : "r" (val));
                        A bit mask suitable for the BFINS instruction.
            М
                        Integer constant that is a byte tiled out eight times.
           N
            0
                        The integer zero constant.
            Ρ
                        Integer constant that is a sign-extended byte tiled out as four shorts.
                        Integer constant that fits in one signed byte when incremented
            Q
                        (-129 \dots 126), but excluding -1.
                        Integer constant that has all 1 bits consecutive and starting at bit
            S
                        0.
            Т
                        A 16-bit fragment of a got, tls, or pc-relative reference.
                        Memory operand except postincrement. This is roughly the same
            U
                        as 'm' when not used together with '<' or '>'.
            W
                        An 8-element vector constant with identical elements.
            Y
                        A 4-element vector constant with identical elements.
            Z0
                        The integer constant 0xfffffff.
            Z1
                        The integer constant 0xfffffff00000000.
TILEPro—'config/tilepro/constraints.md'
            R00
            R01
            R02
            R03
            R04
            R05
            R06
            R07
            R08
            R09
            R10
                        Each of these represents a register constraint for an individual reg-
                        ister, from r0 to r10.
            Ι
                        Signed 8-bit integer constant.
            J
                        Signed 16-bit integer constant.
```

	K	Nonzero integer constant with low 16 bits zero.		
	L	Integer constant that fits in one signed byte when incremented by one $(-129 \dots 126)$ .		
!	m	Memory operand. If used together with '<' or '>', the operand can have postincrement which requires printing with '%In' and '%in' on TILEPro. For example:  asm ("swadd %IO,%1,%iO" : "=m<>" (mem) : "r" (val));		
	M	A bit mask suitable for the MM instruction.		
	N	Integer constant that is a byte tiled out four times.		
	0	The integer zero constant.		
	P	Integer constant that is a sign-extended byte tiled out as two shorts.		
	Q	Integer constant that fits in one signed byte when incremented $(-129\ldots 126)$ , but excluding -1.		
	T	A symbolic operand, or a 16-bit fragment of a got, tls, or pc-relative reference.		
•	U	Memory operand except postincrement. This is roughly the same as 'm' when not used together with '<' or '>'.		
,	W	A 4-element vector constant with identical elements.		
	Y	A 2-element vector constant with identical elements.		
Visium—'config/visium/constraints.md'				
	b	EAM register mdb		
	С	EAM register mdc		
	f	Floating point register		
	k	Register for sibcall optimization		
	1	General register, but not r29, r30 and r31		
	t	Register r1		
	u	Register r2		
	v	Register r3		
	G	Floating-point constant 0.0		
	J	Integer constant in the range 0 $\dots$ 65535 (16-bit immediate)		
	K	Integer constant in the range $1\ldots 31$ (5-bit immediate)		
:	L	Integer constant in the range $-65535\ldots -1$ (16-bit negative immediate)		
	M	Integer constant $-1$		
	0	Integer constant 0		

P Integer constant 32

x86 family—'config/i386/constraints.md'

- R Legacy register—the eight integer registers available on all i386 processors (a, b, c, d, si, di, bp, sp).
- q Any register accessible as r1. In 32-bit mode, a, b, c, and d; in 64-bit mode, any integer register.
- Any register accessible as rh: a, b, c, and d.
- Any register that can be used as the index in a base+index memory access: that is, any general register except the stack pointer.
- a The a register.
- b The b register.
- c The c register.
- d The d register.
- S The si register.
- D The di register.
- A The a and d registers. This class is used for instructions that return double word results in the ax:dx register pair. Single word values will be allocated either in ax or dx. For example on i386 the following implements rdtsc:

```
unsigned long long rdtsc (void)
{
  unsigned long long tick;
  __asm__ _volatile__("rdtsc":"=A"(tick));
  return tick;
}
```

This is not correct on x86-64 as it would allocate tick in either ax or dx. You have to use the following variant instead:

```
unsigned long long rdtsc (void)
{
  unsigned int tickl, tickh;
  __asm__ __volatile__("rdtsc":"=a"(tickl),"=d"(tickh));
  return ((unsigned long long)tickh << 32)|tickl;
}</pre>
```

- U The call-clobbered integer registers.
- f Any 80387 floating-point (stack) register.
- t Top of 80387 floating-point stack (%st(0)).
- u Second from top of 80387 floating-point stack (%st(1)).
- Yk Any mask register that can be used as a predicate, i.e. k1-k7.
- k Any mask register.
- y Any MMX register.

- x Any SSE register.
  v Any EVEX encodable SSE register (%xmm0-%xmm31).
- w Any bound register.

  Yz First SSE register (%xmm0).
- Yi Any SSE register, when SSE2 and inter-unit moves are enabled.
- Yj Any SSE register, when SSE2 and inter-unit moves from vector registers are enabled.
- Ym Any MMX register, when inter-unit moves are enabled.
- Yn Any MMX register, when inter-unit moves from vector registers are enabled.
- Yp Any integer register when TARGET\_PARTIAL\_REG\_STALL is disabled.
- Ya Any integer register when zero extensions with AND are disabled.
- Yb Any register that can be used as the GOT base when calling \_\_\_tls\_get\_addr: that is, any general register except a and sp registers, for '-fno-plt' if linker supports it. Otherwise, b register.
- Yf Any x87 register when 80387 floating-point arithmetic is enabled.
- Yr Lower SSE register when avoiding REX prefix and all SSE registers otherwise.
- Yv For AVX512VL, any EVEX-encodable SSE register (%xmm0-%xmm31), otherwise any SSE register.
- Yh Any EVEX-encodable SSE register, that has number factor of four.
- Bf Flags register operand.
- Bg GOT memory operand.
- Bm Vector memory operand.
- Bc Constant memory operand.
- Bn Memory operand without REX prefix.
- Bs Sibcall memory operand.
- Bw Call memory operand.
- Bz Constant call address operand.
- BC SSE constant -1 operand.
- I Integer constant in the range 0 . . . 31, for 32-bit shifts.
- J Integer constant in the range 0 . . . 63, for 64-bit shifts.
- K Signed 8-bit integer constant.
- L 0xff or 0xffff, for andsi as a zero-extending move.
- M 0, 1, 2, or 3 (shifts for the lea instruction).

N		Unsigned 8-bit integer constant (for in and out instructions).		
	0	Integer constant in the range $0 \dots 127$ , for 128-bit shifts.		
	G	Standard 80387 floating point constant.		
	C	SSE constant zero operand.		
	е	32-bit signed integer constant, or a symbolic reference known to fit that range (for immediate operands in sign-extending x86-64 instructions).		
	We	32-bit signed integer constant, or a symbolic reference known to fit that range (for sign-extending conversion operations that require non-VOIDmode immediate operands).		
	Wz	32-bit unsigned integer constant, or a symbolic reference known to fit that range (for zero-extending conversion operations that require non-VOIDmode immediate operands).		
	Wd	128-bit integer constant where both the high and low 64-bit word satisfy the ${\tt e}$ constraint.		
	Z	32-bit unsigned integer constant, or a symbolic reference known to fit that range (for immediate operands in zero-extending x86-64 instructions).		
	Tv	VSIB address operand.		
	Ts	Address operand without segment register.		
Xstormy16—'config/stormy16/stormy16.h'				
	a	Register r0.		
	b	Register r1.		
	С	Register r2.		
	d	Register r8.		
	е	Registers r0 through r7.		
	t	Registers r0 and r1.		
	у	The carry register.		
	z	Registers r8 and r9.		
	I	A constant between 0 and 3 inclusive.		
	J	A constant that has exactly one bit set.		
	K	A constant that has exactly one bit clear.		
	L	A constant between 0 and 255 inclusive.		

A constant between -255 and 0 inclusive.

A constant between -3 and 0 inclusive.

A constant between 1 and 4 inclusive.

M N

0

P	A constant between $-4$ and $-1$ inclusive.		
Q	A memory reference that is a stack push.		
R	A memory reference that is a stack pop.		
S	A memory reference that refers to a constant address of known value.		
T	The register indicated by Rx (not implemented yet).		
U	A constant that is not between 2 and 15 inclusive.		
Z	The constant 0.		
Xtensa—'config/xtensa/constraints.md'			
a	General-purpose 32-bit register		
b	One-bit boolean register		
A	MAC16 40-bit accumulator register		
I	Signed 12-bit integer constant, for use in MOVI instructions		

# 17.8.6 Disable insn alternatives using the enabled attribute

There are three insn attributes that may be used to selectively disable instruction alternatives:

Integer constant valid for BccI instructions

Unsigned constant valid for BccUI instructions

enabled Says whether an alternative is available on the current subtarget.

#### preferred\_for\_size

J

K

L

Says whether an enabled alternative should be used in code that is optimized for size.

Signed 8-bit integer constant, for use in ADDI instructions

#### preferred\_for\_speed

Says whether an enabled alternative should be used in code that is optimized for speed.

All these attributes should use (const\_int 1) to allow an alternative or (const\_int 0) to disallow it. The attributes must be a static property of the subtarget; they cannot for example depend on the current operands, on the current optimization level, on the location of the insn within the body of a loop, on whether register allocation has finished, or on the current compiler pass.

The enabled attribute is a correctness property. It tells GCC to act as though the disabled alternatives were never defined in the first place. This is useful when adding new instructions to an existing pattern in cases where the new instructions are only available for certain cpu architecture levels (typically mapped to the -march= command-line option).

In contrast, the preferred\_for\_size and preferred\_for\_speed attributes are strong optimization hints rather than correctness properties. preferred\_for\_size tells GCC

which alternatives to consider when adding or modifying an instruction that GCC wants to optimize for size. preferred\_for\_speed does the same thing for speed. Note that things like code motion can lead to cases where code optimized for size uses alternatives that are not preferred for size, and similarly for speed.

Although define\_insns can in principle specify the enabled attribute directly, it is often clearer to have subsiduary attributes for each architectural feature of interest. The define\_insns can then use these subsiduary attributes to say which alternatives require which features. The example below does this for cpu\_facility.

E.g. the following two patterns could easily be merged using the enabled attribute:

```
(define_insn "*movdi_old"
      [(set (match_operand:DI 0 "register_operand" "=d")
            (match_operand:DI 1 "register_operand" " d"))]
      "!TARGET_NEW"
      "lgr %0,%1")
    (define_insn "*movdi_new"
      [(set (match_operand:DI 0 "register_operand" "=d,f,d")
            (match_operand:DI 1 "register_operand" " d,d,f"))]
      "TARGET_NEW"
       lgr %0,%1
      ldgr %0,%1
       lgdr %0,%1")
to:
    (define_insn "*movdi_combined"
      [(set (match_operand:DI 0 "register_operand" "=d,f,d")
            (match_operand:DI 1 "register_operand" " d,d,f"))]
      u @
       lgr %0,%1
       ldgr %0,%1
       lgdr %0,%1"
      [(set_attr "cpu_facility" "*,new,new")])
with the enabled attribute defined like this:
    (define_attr "cpu_facility" "standard,new" (const_string "standard"))
    (define_attr "enabled" ""
      (cond [(eq_attr "cpu_facility" "standard") (const_int 1)
             (and (eq_attr "cpu_facility" "new")
                  (ne (symbol_ref "TARGET_NEW") (const_int 0)))
             (const_int 1)]
            (const_int 0)))
```

# 17.8.7 Defining Machine-Specific Constraints

Machine-specific constraints fall into two categories: register and non-register constraints. Within the latter category, constraints which allow subsets of all possible memory or address operands should be specially marked, to give reload more information.

Machine-specific constraints can be given names of arbitrary length, but they must be entirely composed of letters, digits, underscores ('\_'), and angle brackets ('< >'). Like C identifiers, they must begin with a letter or underscore.

In order to avoid ambiguity in operand constraint strings, no constraint can have a name that begins with any other constraint's name. For example, if x is defined as a constraint name, xy may not be, and vice versa. As a consequence of this rule, no constraint may begin with one of the generic constraint letters: 'E F V X g i m n o p r s'.

Register constraints correspond directly to register classes. See Section 18.8 [Register Classes], page 512. There is thus not much flexibility in their definitions.

# define\_register\_constraint name regclass docstring [MD Expression]

All three arguments are string constants. name is the name of the constraint, as it will appear in match\_operand expressions. If name is a multi-letter constraint its length shall be the same for all constraints starting with the same letter. regclass can be either the name of the corresponding register class (see Section 18.8 [Register Classes], page 512), or a C expression which evaluates to the appropriate register class. If it is an expression, it must have no side effects, and it cannot look at the operand. The usual use of expressions is to map some register constraints to NO\_REGS when the register class is not available on a given subarchitecture.

docstring is a sentence documenting the meaning of the constraint. Docstrings are explained further below.

Non-register constraints are more like predicates: the constraint definition gives a boolean expression which indicates whether the constraint matches.

### define\_constraint name docstring exp

[MD Expression]

The name and docstring arguments are the same as for define\_register\_constraint, but note that the docstring comes immediately after the name for these expressions. exp is an RTL expression, obeying the same rules as the RTL expressions in predicate definitions. See Section 17.7.2 [Defining Predicates], page 348, for details. If it evaluates true, the constraint matches; if it evaluates false, it doesn't. Constraint expressions should indicate which RTL codes they might match, just like predicate expressions.

match\_test C expressions have access to the following variables:

op The RTL object defining the operand.

mode The machine mode of op.

ival 'INTVAL (op)', if op is a const\_int.

hval 'CONST\_DOUBLE\_HIGH (op)', if op is an integer const\_double.

lval 'CONST\_DOUBLE\_LOW (op)', if op is an integer const\_double.

rval 'CONST\_DOUBLE\_REAL\_VALUE (op)', if op is a floating-point

const\_double.

The \*val variables should only be used once another piece of the expression has verified that op is the appropriate kind of RTL object.

Most non-register constraints should be defined with define\_constraint. The remaining two definition expressions are only appropriate for constraints that should be handled specially by reload if they fail to match.

## define\_memory\_constraint name docstring exp

[MD Expression]

Use this expression for constraints that match a subset of all memory operands: that is, reload can make them match by converting the operand to the form '(mem (reg X))', where X is a base register (from the register class specified by BASE\_REG\_CLASS, see Section 18.8 [Register Classes], page 512).

For example, on the S/390, some instructions do not accept arbitrary memory references, but only those that do not make use of an index register. The constraint letter 'Q' is defined to represent a memory address of this type. If 'Q' is defined with define\_memory\_constraint, a 'Q' constraint can handle any memory operand, because reload knows it can simply copy the memory address into a base register if required. This is analogous to the way an 'o' constraint can handle any memory operand.

The syntax and semantics are otherwise identical to define\_constraint.

# define\_special\_memory\_constraint name docstring exp

[MD Expression]

Use this expression for constraints that match a subset of all memory operands: that is, reload cannot make them match by reloading the address as it is described for define\_memory\_constraint or such address reload is undesirable with the performance point of view.

For example, define\_special\_memory\_constraint can be useful if specifically aligned memory is necessary or desirable for some insn operand.

The syntax and semantics are otherwise identical to define\_constraint.

#### define\_address\_constraint name docstring exp

[MD Expression]

Use this expression for constraints that match a subset of all address operands: that is, reload can make the constraint match by converting the operand to the form '(reg X)', again with X a base register.

Constraints defined with define\_address\_constraint can only be used with the address\_operand predicate, or machine-specific predicates that work the same way. They are treated analogously to the generic 'p' constraint.

The syntax and semantics are otherwise identical to define\_constraint.

For historical reasons, names beginning with the letters 'G H' are reserved for constraints that match only const\_doubles, and names beginning with the letters 'I J K L M N O P' are reserved for constraints that match only const\_ints. This may change in the future. For the time being, constraints with these names must be written in a stylized form, so that genpreds can tell you did it correctly:

It is fine to use names beginning with other letters for constraints that match const\_doubles or const\_ints.

Each docstring in a constraint definition should be one or more complete sentences, marked up in Texinfo format. They are currently unused. In the future they will be copied into the GCC manual, in Section 17.8.5 [Machine Constraints], page 357, replacing the hand-maintained tables currently found in that section. Also, in the future the compiler may use this to give more helpful diagnostics when poor choice of asm constraints causes a reload failure.

If you put the pseudo-Texinfo directive '@internal' at the beginning of a docstring, then (in the future) it will appear only in the internals manual's version of the machine-specific constraint tables. Use this for constraints that should not appear in asm statements.

# 17.8.8 Testing constraints from C

It is occasionally useful to test a constraint from C code rather than implicitly via the constraint string in a match\_operand. The generated file 'tm\_p.h' declares a few interfaces for working with constraints. At present these are defined for all constraints except g (which is equivalent to general\_operand).

Some valid constraint names are not valid C identifiers, so there is a mangling scheme for referring to them from C. Constraint names that do not contain angle brackets or underscores are left unchanged. Underscores are doubled, each '<' is replaced with '\_1', and each '>' with '\_g'. Here are some examples:

Original	Mangled
x	x
P42x	P42x
P4_x	P4x
P4>x	P4_gx
P4>>	$P4_gg$
P4_g>	P4g_g

Throughout this section, the variable c is either a constraint in the abstract sense, or a constant from enum constraint\_num; the variable m is a mangled constraint name (usually as part of a larger identifier).

constraint\_num [Enum]

For each constraint except g, there is a corresponding enumeration constant: 'CONSTRAINT\_' plus the mangled name of the constraint. Functions that take an enum constraint\_num as an argument expect one of these constants.

# inline bool satisfies\_constraint\_m (rtx exp)

[Function]

For each non-register constraint m except g, there is one of these functions; it returns true if exp satisfies the constraint. These functions are only visible if 'rtl.h' was included before 'tm\_p.h'.

bool constraint\_satisfied\_p (rtx exp, enum constraint\_num c) [Function] Like the satisfies\_constraint\_m functions, but the constraint to test is given as an argument, c. If c specifies a register constraint, this function will always return false.

```
enum reg_class reg_class_for_constraint (enum constraint_num [Function]
c)
```

Returns the register class associated with c. If c is not a register constraint, or those registers are not available for the currently selected subtarget, returns NO\_REGS.

Here is an example use of satisfies\_constraint\_m. In peephole optimizations (see Section 17.18 [Peephole Definitions], page 446), operand constraint strings are ignored, so if there are relevant constraints, they must be tested in the C condition. In the example, the optimization is applied if operand 2 does *not* satisfy the 'K' constraint. (This is a simplified version of a peephole definition from the i386 machine description.)

# 17.9 Standard Pattern Names For Generation

Here is a table of the instruction names that are meaningful in the RTL generation pass of the compiler. Giving one of these names to an instruction pattern tells the RTL generation pass that it can use the pattern to accomplish a certain task.

'movm'

Here *m* stands for a two-letter machine mode name, in lowercase. This instruction pattern moves data with that machine mode from operand 1 to operand 0. For example, 'movsi' moves full-word data.

If operand 0 is a subreg with mode m of a register whose own mode is wider than m, the effect of this instruction is to store the specified value in the part of the register that corresponds to mode m. Bits outside of m, but which are within the same target word as the subreg are undefined. Bits which are outside the target word are left unchanged.

This class of patterns is special in several ways. First of all, each of these names up to and including full word size *must* be defined, because there is no other way to copy a datum from one place to another. If there are patterns accepting operands in larger modes, 'movm' must be defined for integer modes of those sizes.

Second, these patterns are not used solely in the RTL generation pass. Even the reload pass can generate move insues to copy values from stack slots into temporary registers. When it does so, one of the operands is a hard register and the other is an operand that can need to be reloaded into a register.

Therefore, when given such a pair of operands, the pattern must generate RTL which needs no reloading and needs no temporary registers—no registers other than the operands. For example, if you support the pattern with a define\_

expand, then in such a case the define\_expand mustn't call force\_reg or any other such function which might generate new pseudo registers.

This requirement exists even for subword modes on a RISC machine where fetching those modes from memory normally requires several insns and some temporary registers.

During reload a memory reference with an invalid address may be passed as an operand. Such an address will be replaced with a valid address later in the reload pass. In this case, nothing may be done with the address except to use it as it stands. If it is copied, it will not be replaced with a valid address. No attempt should be made to make such an address into a valid address and no routine (such as change\_address) that will do so may be called. Note that general\_operand will fail when applied to such an address.

The global variable reload\_in\_progress (which must be explicitly declared if required) can be used to determine whether such special handling is required.

The variety of operands that have reloads depends on the rest of the machine description, but typically on a RISC machine these can only be pseudo registers that did not get hard registers, while on other machines explicit memory references will get optional reloads.

If a scratch register is required to move an object to or from memory, it can be allocated using gen\_reg\_rtx prior to life analysis.

If there are cases which need scratch registers during or after reload, you must provide an appropriate secondary\_reload target hook.

The macro can\_create\_pseudo\_p can be used to determine if it is unsafe to create new pseudo registers. If this variable is nonzero, then it is unsafe to call gen\_reg\_rtx to allocate a new pseudo.

The constraints on a 'movm' must permit moving any hard register to any other hard register provided that TARGET\_HARD\_REGNO\_MODE\_OK permits mode m in both registers and TARGET\_REGISTER\_MOVE\_COST applied to their classes returns a value of 2.

It is obligatory to support floating point 'movm' instructions into and out of any registers that can hold fixed point values, because unions and structures (which have modes SImode or DImode) can be in those registers and they may have floating point members.

There may also be a need to support fixed point 'movm' instructions in and out of floating point registers. Unfortunately, I have forgotten why this was so, and I don't know whether it is still true. If TARGET\_HARD\_REGNO\_MODE\_OK rejects fixed point values in floating point registers, then the constraints of the fixed point 'movm' instructions must be designed to avoid ever trying to reload into a floating point register.

'reload\_inm'

'reload\_outm'

These named patterns have been obsoleted by the target hook secondary\_reload.

Like 'movm', but used when a scratch register is required to move between operand 0 and operand 1. Operand 2 describes the scratch register. See the

discussion of the SECONDARY\_RELOAD\_CLASS macro in see Section 18.8 [Register Classes], page 512.

There are special restrictions on the form of the match\_operands used in these patterns. First, only the predicate for the reload operand is examined, i.e., reload\_in examines operand 1, but not the predicates for operand 0 or 2. Second, there may be only one alternative in the constraints. Third, only a single register class letter may be used for the constraint; subsequent constraint letters are ignored. As a special exception, an empty constraint string matches the ALL\_REGS register class. This may relieve ports of the burden of defining an ALL\_REGS constraint letter just for these patterns.

### 'movstrictm'

Like 'movm' except that if operand 0 is a subreg with mode m of a register whose natural mode is wider, the 'movstrictm' instruction is guaranteed not to alter any of the register except the part which belongs to mode m.

#### 'movmisalignm'

This variant of a move pattern is designed to load or store a value from a memory address that is not naturally aligned for its mode. For a store, the memory will be in operand 0; for a load, the memory will be in operand 1. The other operand is guaranteed not to be a memory, so that it's easy to tell whether this is a load or store.

This pattern is used by the autovectorizer, and when expanding a MISALIGNED\_INDIRECT\_REF expression.

## 'load\_multiple'

Load several consecutive memory locations into consecutive registers. Operand 0 is the first of the consecutive registers, operand 1 is the first memory location, and operand 2 is a constant: the number of consecutive registers.

Define this only if the target machine really has such an instruction; do not define this if the most efficient way of loading consecutive registers from memory is to do them one at a time.

On some machines, there are restrictions as to which consecutive registers can be stored into memory, such as particular starting or ending register numbers or only a range of valid counts. For those machines, use a define\_expand (see Section 17.15 [Expander Definitions], page 438) and make the pattern fail if the restrictions are not met.

Write the generated insn as a parallel with elements being a set of one register from the appropriate memory location (you may also need use or clobber elements). Use a match\_parallel (see Section 17.4 [RTL Template], page 339) to recognize the insn. See 'rs6000.md' for examples of the use of this insn pattern.

#### 'store\_multiple'

Similar to 'load\_multiple', but store several consecutive registers into consecutive memory locations. Operand 0 is the first of the consecutive memory locations, operand 1 is the first register, and operand 2 is a constant: the number of consecutive registers.

# 'vec\_load\_lanesmn'

Perform an interleaved load of several vectors from memory operand 1 into register operand 0. Both operands have mode m. The register operand is viewed as holding consecutive vectors of mode n, while the memory operand is a flat array that contains the same number of elements. The operation is equivalent to:

```
int c = GET_MODE_SIZE (m) / GET_MODE_SIZE (n);
for (j = 0; j < GET_MODE_NUNITS (n); j++)
  for (i = 0; i < c; i++)
   operand0[i][j] = operand1[j * c + i];</pre>
```

For example, 'vec\_load\_lanestiv4hi' loads 8 16-bit values from memory into a register of mode 'TI'. The register contains two consecutive vectors of mode 'V4HI'.

This pattern can only be used if:

```
TARGET_ARRAY_MODE_SUPPORTED_P (n, c)
```

is true. GCC assumes that, if a target supports this kind of instruction for some mode n, it also supports unaligned loads for vectors of mode n.

This pattern is not allowed to FAIL.

#### 'vec\_mask\_load\_lanesmn'

Like 'vec\_load\_lanesmn', but takes an additional mask operand (operand 2) that specifies which elements of the destination vectors should be loaded. Other elements of the destination vectors are set to zero. The operation is equivalent to:

```
int c = GET_MODE_SIZE (m) / GET_MODE_SIZE (n);
for (j = 0; j < GET_MODE_NUNITS (n); j++)
  if (operand2[j])
  for (i = 0; i < c; i++)
    operand0[i][j] = operand1[j * c + i];
else
  for (i = 0; i < c; i++)
    operand0[i][j] = 0;</pre>
```

This pattern is not allowed to FAIL.

## 'vec\_store\_lanesmn'

Equivalent to 'vec\_load\_lanesmn', with the memory and register operands reversed. That is, the instruction is equivalent to:

```
int c = GET_MODE_SIZE (m) / GET_MODE_SIZE (n);
for (j = 0; j < GET_MODE_NUNITS (n); j++)
  for (i = 0; i < c; i++)
   operand0[j * c + i] = operand1[i][j];</pre>
```

for a memory operand 0 and register operand 1.

This pattern is not allowed to FAIL.

#### 'vec\_mask\_store\_lanesmn'

Like 'vec\_store\_lanesmn', but takes an additional mask operand (operand 2) that specifies which elements of the source vectors should be stored. The operation is equivalent to:

```
int c = GET_MODE_SIZE (m) / GET_MODE_SIZE (n);
for (j = 0; j < GET_MODE_NUNITS (n); j++)</pre>
```

```
if (operand2[j])
  for (i = 0; i < c; i++)
    operand0[j * c + i] = operand1[i][j];</pre>
```

This pattern is not allowed to FAIL.

## 'gather\_loadmn'

Load several separate memory locations into a vector of mode m. Operand 1 is a scalar base address and operand 2 is a vector of mode n containing offsets from that base. Operand 0 is a destination vector with the same number of elements as n. For each element index i:

- extend the offset element *i* to address width, using zero extension if operand 3 is 1 and sign extension if operand 3 is zero;
- multiply the extended offset by operand 4;
- add the result to the base; and
- load the value at that address into element *i* of operand 0.

The value of operand 3 does not matter if the offsets are already address width.

# 'mask\_gather\_loadmn'

Like 'gather\_loadmn', but takes an extra mask operand as operand 5. Bit i of the mask is set if element i of the result should be loaded from memory and clear if element i of the result should be set to zero.

#### 'scatter\_storemn'

Store a vector of mode m into several distinct memory locations. Operand 0 is a scalar base address and operand 1 is a vector of mode n containing offsets from that base. Operand 4 is the vector of values that should be stored, which has the same number of elements as n. For each element index i:

- extend the offset element i to address width, using zero extension if operand 2 is 1 and sign extension if operand 2 is zero;
- multiply the extended offset by operand 3;
- add the result to the base; and
- store element *i* of operand 4 to that address.

The value of operand 2 does not matter if the offsets are already address width.

# 'mask\_scatter\_storemn'

Like 'scatter\_storemn', but takes an extra mask operand as operand 5. Bit i of the mask is set if element i of the result should be stored to memory.

## 'vec\_setm'

Set given field in the vector value. Operand 0 is the vector to modify, operand 1 is new value of field and operand 2 specify the field index.

## 'vec\_extractmn'

Extract given field from the vector value. Operand 1 is the vector, operand 2 specify field index and operand 0 place to store value into. The n mode is the mode of the field or vector of fields that should be extracted, should be either element mode of the vector mode m, or a vector mode with the same element mode and smaller number of elements. If n is a vector mode, the index is counted in units of that mode.

# 'vec\_initmn'

Initialize the vector to given values. Operand 0 is the vector to initialize and operand 1 is parallel containing values for individual fields. The n mode is the mode of the elements, should be either element mode of the vector mode m, or a vector mode with the same element mode and smaller number of elements.

## 'vec\_duplicatem'

Initialize vector output operand 0 so that each element has the value given by scalar input operand 1. The vector has mode m and the scalar has the mode appropriate for one element of m.

This pattern only handles duplicates of non-constant inputs. Constant vectors go through the movm pattern instead.

This pattern is not allowed to FAIL.

## 'vec\_seriesm'

Initialize vector output operand 0 so that element i is equal to operand 1 plus i times operand 2. In other words, create a linear series whose base value is operand 1 and whose step is operand 2.

The vector output has mode m and the scalar inputs have the mode appropriate for one element of m. This pattern is not used for floating-point vectors, in order to avoid having to specify the rounding behavior for i > 1.

This pattern is not allowed to FAIL.

#### while\_ultmn

Set operand 0 to a mask that is true while incrementing operand 1 gives a value that is less than operand 2. Operand 0 has mode n and operands 1 and 2 are scalar integers of mode m. The operation is equivalent to:

```
operand0[0] = operand1 < operand2;
for (i = 1; i < GET_MODE_NUNITS (n); i++)
  operand0[i] = operand0[i - 1] && (operand1 + i < operand2);</pre>
```

# 'check\_raw\_ptrsm'

Check whether, given two pointers a and b and a length len, a write of len bytes at a followed by a read of len bytes at b can be split into interleaved byte accesses 'a[0], b[0], a[1], b[1], ...' without affecting the dependencies between the bytes. Set operand 0 to true if the split is possible and false otherwise.

Operands 1, 2 and 3 provide the values of a, b and len respectively. Operand 4 is a constant integer that provides the known common alignment of a and b. All inputs have mode m.

This split is possible if:

```
a == b || a + len <= b || b + len <= a
```

You should only define this pattern if the target has a way of accelerating the test without having to do the individual comparisons.

# 'check\_war\_ptrsm'

Like 'check\_raw\_ptrsm', but with the read and write swapped round. The split is possible in this case if:

```
b <= a || a + len <= b
```

# 'vec\_cmpmn'

Output a vector comparison. Operand 0 of mode n is the destination for predicate in operand 1 which is a signed vector comparison with operands of mode m in operands 2 and 3. Predicate is computed by element-wise evaluation of the vector comparison with a truth value of all-ones and a false value of all-zeros.

## 'vec\_cmpumn'

Similar to vec\_cmpmn but perform unsigned vector comparison.

## 'vec\_cmpeqmn'

Similar to vec\_cmpmn but perform equality or non-equality vector comparison only. If vec\_cmpmn or vec\_cmpumn instruction pattern is supported, it will be preferred over vec\_cmpeqmn, so there is no need to define this instruction pattern if the others are supported.

'vcondmn'

#### 'vcondumn'

Similar to vcondmn but performs unsigned vector comparison.

# 'vcondeqmn'

Similar to **vcondmn** but performs equality or non-equality vector comparison only. If **vcondmn** or **vcondumn** instruction pattern is supported, it will be preferred over **vcondeqmn**, so there is no need to define this instruction pattern if the others are supported.

#### 'vcond\_mask\_mn'

Similar to **vcondmn** but operand 3 holds a pre-computed result of vector comparison.

#### 'maskloadmn'

Perform a masked load of vector from memory operand 1 of mode m into register operand 0. Mask is provided in register operand 2 of mode n.

This pattern is not allowed to FAIL.

#### 'maskstoremn'

Perform a masked store of vector from register operand 1 of mode m into memory operand 0. Mask is provided in register operand 2 of mode n.

This pattern is not allowed to FAIL.

#### 'vec\_permm'

Output a (variable) vector permutation. Operand 0 is the destination to receive elements from operand 1 and operand 2, which are of mode m. Operand 3 is the selector. It is an integral mode vector of the same width and number of elements as mode m.

The input elements are numbered from 0 in operand 1 through 2 \* N - 1 in operand 2. The elements of the selector must be computed modulo 2 \* N. Note that if  $\texttt{rtx\_equal\_p(operand1, operand2)}$ , this can be implemented with just operand 1 and selector elements modulo N.

In order to make things easy for a number of targets, if there is no 'vec\_perm' pattern for mode m, but there is for mode q where q is a vector of QImode of the same width as m, the middle-end will lower the mode m VEC\_PERM\_EXPR to mode q.

See also TARGET\_VECTORIZER\_VEC\_PERM\_CONST, which performs the analogous operation for constant selectors.

'pushm1' Output a push instruction. Operand 0 is value to push. Used only when PUSH\_ROUNDING is defined. For historical reason, this pattern may be missing and in such case an mov expander is used instead, with a MEM expression forming the push operation. The mov expander method is deprecated.

'addm3' Add operand 2 and operand 1, storing the result in operand 0. All operands must have mode m. This can be used even on two-address machines, by means of constraints requiring operands 1 and 0 to be the same location.

'ssaddm3', 'usaddm3'
'subm3', 'sssubm3', 'ussubm3'

'mulm3', 'sssubm3', 'ussubm3' 'mulm3', 'ssmulm3', 'usmulm3'

'divm3', 'ssdivm3'

'udivm3', 'usdivm3'

'modm3', 'umodm3'

'uminm3', 'umaxm3'

'andm3', 'iorm3', 'xorm3'

Similar, for other arithmetic operations.

'addvm4' Like addm3 but takes a code\_label as operand 3 and emits code to jump to it if signed overflow occurs during the addition. This pattern is used to implement the built-in functions performing signed integer addition with overflow checking.

'subvm4', 'mulvm4'

Similar, for other signed arithmetic operations.

'uaddvm4' Like addvm4 but for unsigned addition. That is to say, the operation is the same as signed addition but the jump is taken only on unsigned overflow.

'usubvm4', 'umulvm4'

Similar, for other unsigned arithmetic operations.

'addptrm3'

Like addm3 but is guaranteed to only be used for address calculations. The expanded code is not allowed to clobber the condition code. It only needs to be defined if addm3 sets the condition code. If adds used for address calculations and normal adds are not compatible it is required to expand a distinct pattern (e.g. using an unspec). The pattern is used by LRA to emit address calculations. addm3 is used if addptrm3 is not defined.

'fmam4' Multiply operand 2 and operand 1, then add operand 3, storing the result in operand 0 without doing an intermediate rounding step. All operands must have mode m. This pattern is used to implement the fma, fmaf, and fmal builtin functions from the ISO C99 standard.

'fmsm4' Like fmam4, except operand 3 subtracted from the product instead of added to the product. This is represented in the rtl as

```
(fma:m op1 op2 (neg:m op3))
```

'fnmam4' Like fmam4 except that the intermediate product is negated before being added to operand 3. This is represented in the rtl as

```
(fma:m (neg:m op1) op2 op3)
```

'fnmsm4' Like fmsm4 except that the intermediate product is negated before subtracting operand 3. This is represented in the rtl as

```
(fma:m (neg:m op1) op2 (neg:m op3))
```

'sminm3', 'smaxm3'

Signed minimum and maximum operations. When used with floating point, if both operands are zeros, or if either operand is NaN, then it is unspecified which of the two operands is returned as the result.

'fminm3', 'fmaxm3'

IEEE-conformant minimum and maximum operations. If one operand is a quiet NaN, then the other operand is returned. If both operands are quiet NaN, then a quiet NaN is returned. In the case when gcc supports signaling NaN (-fsignaling-nans) an invalid floating point exception is raised and a quiet NaN is returned.

All operands have mode m, which is a scalar or vector floating-point mode. These patterns are not allowed to FAIL.

'reduc\_smin\_scal\_m', 'reduc\_smax\_scal\_m'

Find the signed minimum/maximum of the elements of a vector. The vector is operand 1, and operand 0 is the scalar result, with mode equal to the mode of the elements of the input vector.

'reduc\_umin\_scal\_m', 'reduc\_umax\_scal\_m'

Find the unsigned minimum/maximum of the elements of a vector. The vector is operand 1, and operand 0 is the scalar result, with mode equal to the mode of the elements of the input vector.

'reduc\_plus\_scal\_m'

Compute the sum of the elements of a vector. The vector is operand 1, and operand 0 is the scalar result, with mode equal to the mode of the elements of the input vector.

```
'reduc_and_scal_m'
'reduc_ior_scal_m'
'reduc_xor_scal_m'
```

Compute the bitwise AND/IOR/XOR reduction of the elements of a vector of mode m. Operand 1 is the vector input and operand 0 is the scalar result. The mode of the scalar result is the same as one element of m.

#### extract\_last\_m

Find the last set bit in mask operand 1 and extract the associated element of vector operand 2. Store the result in scalar operand 0. Operand 2 has vector mode m while operand 0 has the mode appropriate for one element of m. Operand 1 has the usual mask mode for vectors of mode m; see TARGET\_VECTORIZE\_GET\_MASK\_MODE.

#### fold\_extract\_last\_m

If any bits of mask operand 2 are set, find the last set bit, extract the associated element from vector operand 3, and store the result in operand 0. Store operand 1 in operand 0 otherwise. Operand 3 has mode m and operands 0 and 1 have the mode appropriate for one element of m. Operand 2 has the usual mask mode for vectors of mode m; see TARGET\_VECTORIZE\_GET\_MASK\_MODE.

## fold\_left\_plus\_m

Take scalar operand 1 and successively add each element from vector operand 2. Store the result in scalar operand 0. The vector has mode m and the scalars have the mode appropriate for one element of m. The operation is strictly in-order: there is no reassociation.

## mask\_fold\_left\_plus\_m

Like 'fold\_left\_plus\_m', but takes an additional mask operand (operand 3) that specifies which elements of the source vector should be added.

'sdot\_prodm'
'udot\_prodm'

Compute the sum of the products of two signed/unsigned elements. Operand 1 and operand 2 are of the same mode. Their product, which is of a wider mode, is computed and added to operand 3. Operand 3 is of a mode equal or wider than the mode of the product. The result is placed in operand 0, which is of the same mode as operand 3.

'ssadm'

'usadm' Compute the sum of absolute differences of two signed/unsigned elements. Operand 1 and operand 2 are of the same mode. Their absolute difference, which is of a wider mode, is computed and added to operand 3. Operand 3 is of a mode equal or wider than the mode of the absolute difference. The result is placed in operand 0, which is of the same mode as operand 3.

'widen\_ssum*m3*'
'widen\_usum*m3*'

Operands 0 and 2 are of the same mode, which is wider than the mode of operand 1. Add operand 1 to operand 2 and place the widened result in operand 0. (This is used express accumulation of elements into an accumulator of a wider mode.)

'smulhsm3'
'umulhsm3'

Signed/unsigned multiply high with scale. This is equivalent to the C code: narrow op0, op1, op2;

```
op0 = (narrow) (((wide) op1 * (wide) op2) >> (N / 2 - 1));
```

where the sign of 'narrow' determines whether this is a signed or unsigned operation, and N is the size of 'wide' in bits.

'smulhrsm3'
'umulhrsm3'

Signed/unsigned multiply high with round and scale. This is equivalent to the C code:

```
narrow op0, op1, op2;
...
op0 = (narrow) (((((wide) op1 * (wide) op2) >> (N / 2 - 2)) + 1) >> 1);
```

where the sign of 'narrow' determines whether this is a signed or unsigned operation, and N is the size of 'wide' in bits.

'sdiv\_pow2m3'
'sdiv\_pow2m3'

Signed division by power-of-2 immediate. Equivalent to:

```
signed op0, op1;
...
op0 = op1 / (1 << imm);</pre>
```

'vec\_shl\_insert\_m'

Shift the elements in vector input operand 1 left one element (i.e. away from element 0) and fill the vacated element 0 with the scalar in operand 2. Store the result in vector output operand 0. Operands 0 and 1 have mode m and operand 2 has the mode appropriate for one element of m.

'vec\_shl\_m'

Whole vector left shift in bits, i.e. away from element 0. Operand 1 is a vector to be shifted. Operand 2 is an integer shift amount in bits. Operand 0 is where the resulting shifted vector is stored. The output and input vectors should have the same modes.

'vec\_shr\_m'

Whole vector right shift in bits, i.e. towards element 0. Operand 1 is a vector to be shifted. Operand 2 is an integer shift amount in bits. Operand 0 is where the resulting shifted vector is stored. The output and input vectors should have the same modes.

'vec\_pack\_trunc\_m'

Narrow (demote) and merge the elements of two vectors. Operands 1 and 2 are vectors of the same mode having N integral or floating point elements of size S. Operand 0 is the resulting vector in which 2\*N elements of size S/2 are concatenated after narrowing them down using truncation.

'vec\_pack\_sbool\_trunc\_m'

Narrow and merge the elements of two vectors. Operands 1 and 2 are vectors of the same type having N boolean elements. Operand 0 is the resulting vector in which 2\*N elements are concatenated. The last operand (operand 3) is the number of elements in the output vector 2\*N as a CONST\_INT. This instruction

pattern is used when all the vector input and output operands have the same scalar mode m and thus using vec\_pack\_trunc\_m would be ambiguous.

# 'vec\_pack\_ssat\_m', 'vec\_pack\_usat\_m'

Narrow (demote) and merge the elements of two vectors. Operands 1 and 2 are vectors of the same mode having N integral elements of size S. Operand 0 is the resulting vector in which the elements of the two input vectors are concatenated after narrowing them down using signed/unsigned saturating arithmetic.

# 'vec\_pack\_sfix\_trunc\_m', 'vec\_pack\_ufix\_trunc\_m'

Narrow, convert to signed/unsigned integral type and merge the elements of two vectors. Operands 1 and 2 are vectors of the same mode having N floating point elements of size S. Operand 0 is the resulting vector in which 2\*N elements of size S/2 are concatenated.

# 'vec\_packs\_float\_m', 'vec\_packu\_float\_m'

Narrow, convert to floating point type and merge the elements of two vectors. Operands 1 and 2 are vectors of the same mode having N signed/unsigned integral elements of size S. Operand 0 is the resulting vector in which 2\*N elements of size S/2 are concatenated.

# 'vec\_unpacks\_hi\_m', 'vec\_unpacks\_lo\_m'

Extract and widen (promote) the high/low part of a vector of signed integral or floating point elements. The input vector (operand 1) has N elements of size S. Widen (promote) the high/low elements of the vector using signed or floating point extension and place the resulting N/2 values of size 2\*S in the output vector (operand 0).

# 'vec\_unpacku\_hi\_m', 'vec\_unpacku\_lo\_m'

Extract and widen (promote) the high/low part of a vector of unsigned integral elements. The input vector (operand 1) has N elements of size S. Widen (promote) the high/low elements of the vector using zero extension and place the resulting N/2 values of size 2\*S in the output vector (operand 0).

#### 'vec\_unpacks\_sbool\_hi\_m', 'vec\_unpacks\_sbool\_lo\_m'

Extract the high/low part of a vector of boolean elements that have scalar mode m. The input vector (operand 1) has N elements, the output vector (operand 0) has N/2 elements. The last operand (operand 2) is the number of elements of the input vector N as a CONST\_INT. These patterns are used if both the input and output vectors have the same scalar mode m and thus using vec\_unpacks\_hi\_m or vec\_unpacks\_lo\_m would be ambiguous.

# 'vec\_unpacks\_float\_hi\_m', 'vec\_unpacks\_float\_lo\_m' 'vec\_unpacku\_float\_hi\_m', 'vec\_unpacku\_float\_lo\_m'

Extract, convert to floating point type and widen the high/low part of a vector of signed/unsigned integral elements. The input vector (operand 1) has N elements of size S. Convert the high/low elements of the vector using floating point conversion and place the resulting N/2 values of size 2\*S in the output vector (operand 0).

```
'vec_unpack_sfix_trunc_hi_m',
'vec_unpack_sfix_trunc_lo_m'
'vec_unpack_ufix_trunc_hi_m'
'vec_unpack_ufix_trunc_lo_m'
```

Extract, convert to signed/unsigned integer type and widen the high/low part of a vector of floating point elements. The input vector (operand 1) has N elements of size S. Convert the high/low elements of the vector to integers and place the resulting N/2 values of size 2\*S in the output vector (operand 0).

```
'vec_widen_umult_hi_m', 'vec_widen_umult_lo_m'
'vec_widen_smult_hi_m', 'vec_widen_smult_lo_m'
'vec_widen_umult_even_m', 'vec_widen_umult_odd_m'
'vec_widen_smult_even_m', 'vec_widen_smult_odd_m'
```

Signed/Unsigned widening multiplication. The two inputs (operands 1 and 2) are vectors with N signed/unsigned elements of size S. Multiply the high/low or even/odd elements of the two vectors, and put the N/2 products of size 2\*S in the output vector (operand 0). A target shouldn't implement even/odd pattern pair if it is less efficient than lo/hi one.

```
'vec_widen_ushiftl_hi_m', 'vec_widen_ushiftl_lo_m' 'vec_widen_sshiftl_hi_m', 'vec_widen_sshiftl_lo_m'
```

Signed/Unsigned widening shift left. The first input (operand 1) is a vector with N signed/unsigned elements of size S. Operand 2 is a constant. Shift the high/low elements of operand 1, and put the N/2 results of size 2\*S in the output vector (operand 0).

#### 'mulhisi3'

Multiply operands 1 and 2, which have mode HImode, and store a SImode product in operand 0.

## 'mulqihi3', 'mulsidi3'

Similar widening-multiplication instructions of other widths.

# 'umulqihi3', 'umulhisi3', 'umulsidi3'

Similar widening-multiplication instructions that do unsigned multiplication.

# 'usmulqihi3', 'usmulhisi3', 'usmulsidi3'

Similar widening-multiplication instructions that interpret the first operand as unsigned and the second operand as signed, then do a signed multiplication.

#### 'smulm3\_highpart'

Perform a signed multiplication of operands 1 and 2, which have mode m, and store the most significant half of the product in operand 0. The least significant half of the product is discarded.

# 'umulm3\_highpart'

Similar, but the multiplication is unsigned.

'maddmn4' Multiply operands 1 and 2, sign-extend them to mode n, add operand 3, and store the result in operand 0. Operands 1 and 2 have mode m and operands 0 and 3 have mode n. Both modes must be integer or fixed-point modes and n must be twice the size of m.

In other words, maddmn4 is like mulmn3 except that it also adds operand 3.

These instructions are not allowed to FAIL.

#### 'umaddmn4'

Like maddmn4, but zero-extend the multiplication operands instead of sign-extending them.

#### 'ssmaddmn4'

Like maddmn4, but all involved operations must be signed-saturating.

#### 'usmaddmn4'

Like umaddmn4, but all involved operations must be unsigned-saturating.

'msubmn4' Multiply operands 1 and 2, sign-extend them to mode n, subtract the result from operand 3, and store the result in operand 0. Operands 1 and 2 have mode m and operands 0 and 3 have mode n. Both modes must be integer or fixed-point modes and n must be twice the size of m.

In other words, msubmn4 is like mulmn3 except that it also subtracts the result from operand 3.

These instructions are not allowed to FAIL.

#### 'umsubmn4'

Like msubmn4, but zero-extend the multiplication operands instead of sign-extending them.

## 'ssmsubmn4'

Like msubmn4, but all involved operations must be signed-saturating.

#### 'usmsubmn4'

Like umsubmn4, but all involved operations must be unsigned-saturating.

## 'divmodm4'

Signed division that produces both a quotient and a remainder. Operand 1 is divided by operand 2 to produce a quotient stored in operand 0 and a remainder stored in operand 3.

For machines with an instruction that produces both a quotient and a remainder, provide a pattern for 'divmodm4' but do not provide patterns for 'divm3' and 'modm3'. This allows optimization in the relatively common case when both the quotient and remainder are computed.

If an instruction that just produces a quotient or just a remainder exists and is more efficient than the instruction that produces both, write the output routine of 'divmodm4' to call find\_reg\_note and look for a REG\_UNUSED note on the quotient or remainder and generate the appropriate instruction.

## 'udivmodm4'

Similar, but does unsigned division.

# 'ashlm3', 'ssashlm3', 'usashlm3'

Arithmetic-shift operand 1 left by a number of bits specified by operand 2, and store the result in operand 0. Here m is the mode of operand 0 and operand 1; operand 2's mode is specified by the instruction pattern, and the compiler will

convert the operand to that mode before generating the instruction. The shift or rotate expander or instruction pattern should explicitly specify the mode of the operand 2, it should never be VOIDmode. The meaning of out-of-range shift counts can optionally be specified by TARGET\_SHIFT\_TRUNCATION\_MASK. See [TARGET\_SHIFT\_TRUNCATION\_MASK], page 645. Operand 2 is always a scalar type.

'ashrm3', 'lshrm3', 'rotlm3', 'rotrm3'

Other shift and rotate instructions, analogous to the ashlm3 instructions. Operand 2 is always a scalar type.

'vashlm3', 'vashrm3', 'vlshrm3', 'vrotlm3', 'vrotrm3'

Vector shift and rotate instructions that take vectors as operand 2 instead of a scalar type.

'avgm3\_floor'
'uavgm3\_floor'

Signed and unsigned average instructions. These instructions add operands 1 and 2 without truncation, divide the result by 2, round towards -Inf, and store the result in operand 0. This is equivalent to the C code:

```
narrow op0, op1, op2;
...
op0 = (narrow) (((wide) op1 + (wide) op2) >> 1);
```

where the sign of 'narrow' determines whether this is a signed or unsigned operation.

'avgm3\_ceil'

'uavgm3\_ceil'

Like 'avgm3\_floor' and 'uavgm3\_floor', but round towards +Inf. This is equivalent to the C code:

```
narrow op0, op1, op2;
...
op0 = (narrow) (((wide) op1 + (wide) op2 + 1) >> 1);
```

'bswapm2' Reverse the order of bytes of operand 1 and store the result in operand 0.

'negm2', 'ssnegm2', 'usnegm2'

Negate operand 1 and store the result in operand 0.

'negvm3' Like negm2 but takes a code\_label as operand 2 and emits code to jump to it if signed overflow occurs during the negation.

'absm2' Store the absolute value of operand 1 into operand 0.

'sqrtm2' Store the square root of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'rsqrtm2' Store the reciprocal of the square root of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

On most architectures this pattern is only approximate, so either its C condition or the TARGET\_OPTAB\_SUPPORTED\_P hook should check for the appropriate

math flags. (Using the C condition is more direct, but using TARGET\_OPTAB\_SUPPORTED\_P can be useful if a target-specific built-in also uses the 'rsqrtm2' pattern.)

This pattern is not allowed to FAIL.

'fmodm3' Store the remainder of dividing operand 1 by operand 2 into operand 0, rounded towards zero to an integer. All operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

#### 'remainderm3'

Store the remainder of dividing operand 1 by operand 2 into operand 0, rounded to the nearest integer. All operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'scalbm3' Raise FLT\_RADIX to the power of operand 2, multiply it by operand 1, and store the result in operand 0. All operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'ldexpm3' Raise 2 to the power of operand 2, multiply it by operand 1, and store the result in operand 0. Operands 0 and 1 have mode m, which is a scalar or vector floating-point mode. Operand 2's mode has the same number of elements as m and each element is wide enough to store an int. The integers are signed.

This pattern is not allowed to FAIL.

'cosm2' Store the cosine of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'sinm2' Store the sine of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

## 'sincosm3'

Store the cosine of operand 2 into operand 0 and the sine of operand 2 into operand 1. All operands have mode m, which is a scalar or vector floating-point mode.

Targets that can calculate the sine and cosine simultaneously can implement this pattern as opposed to implementing individual sinm2 and cosm2 patterns. The sin and cos built-in functions will then be expanded to the sincosm3 pattern, with one of the output values left unused.

'tanm2' Store the tangent of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'asinm2' Store the arc sine of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'acosm2' Store the arc cosine of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'atanm2' Store the arc tangent of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'expm2' Raise e (the base of natural logarithms) to the power of operand 1 and store the result in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'expm1m2' Raise e (the base of natural logarithms) to the power of operand 1, subtract 1, and store the result in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

For inputs close to zero, the pattern is expected to be more accurate than a separate expm2 and subm3 would be.

This pattern is not allowed to FAIL.

'exp10m2' Raise 10 to the power of operand 1 and store the result in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'exp2m2' Raise 2 to the power of operand 1 and store the result in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

' $\log m2$ ' Store the natural logarithm of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

' $\log 1pm2$ ' Add 1 to operand 1, compute the natural logarithm, and store the result in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

For inputs close to zero, the pattern is expected to be more accurate than a separate addm3 and logm2 would be.

This pattern is not allowed to FAIL.

'log10m2' Store the base-10 logarithm of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'log2m2' Store the base-2 logarithm of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'logbm2' Store the base-FLT\_RADIX logarithm of operand 1 into operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

## 'significandm2'

Store the significand of floating-point operand 1 in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'powm3' Store the value of operand 1 raised to the exponent operand 2 into operand 0. All operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'atan2m3' Store the arc tangent (inverse tangent) of operand 1 divided by operand 2 into operand 0, using the signs of both arguments to determine the quadrant of the result. All operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'floorm2' Store the largest integral value not greater than operand 1 in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode. If '-ffp-int-builtin-inexact' is in effect, the "inexact" exception may be raised for noninteger operands; otherwise, it may not.

This pattern is not allowed to FAIL.

#### 'btruncm2'

Round operand 1 to an integer, towards zero, and store the result in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode. If '-ffp-int-builtin-inexact' is in effect, the "inexact" exception may be raised for noninteger operands; otherwise, it may not.

This pattern is not allowed to FAIL.

'roundm2' Round operand 1 to the nearest integer, rounding away from zero in the event of a tie, and store the result in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode. If '-ffp-int-builtin-inexact' is in effect, the "inexact" exception may be raised for noninteger operands; otherwise, it may not.

This pattern is not allowed to FAIL.

'ceilm2' Store the smallest integral value not less than operand 1 in operand 0. Both operands have mode m, which is a scalar or vector floating-point mode. If '-ffp-int-builtin-inexact' is in effect, the "inexact" exception may be raised for noninteger operands; otherwise, it may not.

This pattern is not allowed to FAIL.

### 'nearbyintm2'

Round operand 1 to an integer, using the current rounding mode, and store the result in operand 0. Do not raise an inexact condition when the result is different from the argument. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'rintm2' Round operand 1 to an integer, using the current rounding mode, and store the result in operand 0. Raise an inexact condition when the result is different

from the argument. Both operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

#### 'lrintmn2'

Convert operand 1 (valid for floating point mode m) to fixed point mode n as a signed number according to the current rounding mode and store in operand 0 (which has mode n).

#### 'lroundmn2'

Convert operand 1 (valid for floating point mode m) to fixed point mode n as a signed number rounding to nearest and away from zero and store in operand 0 (which has mode n).

## 'lfloormn2'

Convert operand 1 (valid for floating point mode m) to fixed point mode n as a signed number rounding down and store in operand 0 (which has mode n).

#### 'lceilmn2'

Convert operand 1 (valid for floating point mode m) to fixed point mode n as a signed number rounding up and store in operand 0 (which has mode n).

## 'copysignm3'

Store a value with the magnitude of operand 1 and the sign of operand 2 into operand 0. All operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

### 'xorsignm3'

Equivalent to 'op0 = op1 \* copysign (1.0, op2)': store a value with the magnitude of operand 1 and the sign of operand 2 into operand 0. All operands have mode m, which is a scalar or vector floating-point mode.

This pattern is not allowed to FAIL.

'ffsm2' Store into operand 0 one plus the index of the least significant 1-bit of operand 1. If operand 1 is zero, store zero.

m is either a scalar or vector integer mode. When it is a scalar, operand 1 has mode m but operand 0 can have whatever scalar integer mode is suitable for the target. The compiler will insert conversion instructions as necessary (typically to convert the result to the same width as int). When m is a vector, both operands must have mode m.

This pattern is not allowed to FAIL.

'clrsbm2' Count leading redundant sign bits. Store into operand 0 the number of redundant sign bits in operand 1, starting at the most significant bit position. A redundant sign bit is defined as any sign bit after the first. As such, this count will be one less than the count of leading sign bits.

m is either a scalar or vector integer mode. When it is a scalar, operand 1 has mode m but operand 0 can have whatever scalar integer mode is suitable for the target. The compiler will insert conversion instructions as necessary (typically

to convert the result to the same width as int). When m is a vector, both operands must have mode m.

This pattern is not allowed to FAIL.

'clzm2'

Store into operand 0 the number of leading 0-bits in operand 1, starting at the most significant bit position. If operand 1 is 0, the CLZ\_DEFINED\_VALUE\_AT\_ ZERO (see Section 18.31 [Misc], page 642) macro defines if the result is undefined or has a useful value.

m is either a scalar or vector integer mode. When it is a scalar, operand 1 has mode m but operand 0 can have whatever scalar integer mode is suitable for the target. The compiler will insert conversion instructions as necessary (typically to convert the result to the same width as int). When m is a vector, both operands must have mode m.

This pattern is not allowed to FAIL.

'ctzm2'

Store into operand 0 the number of trailing 0-bits in operand 1, starting at the least significant bit position. If operand 1 is 0, the CTZ\_DEFINED\_VALUE\_AT\_ ZERO (see Section 18.31 [Misc], page 642) macro defines if the result is undefined or has a useful value.

m is either a scalar or vector integer mode. When it is a scalar, operand 1 has mode m but operand 0 can have whatever scalar integer mode is suitable for the target. The compiler will insert conversion instructions as necessary (typically to convert the result to the same width as int). When m is a vector, both operands must have mode m.

This pattern is not allowed to FAIL.

#### 'popcountm2'

Store into operand 0 the number of 1-bits in operand 1.

m is either a scalar or vector integer mode. When it is a scalar, operand 1 has mode m but operand 0 can have whatever scalar integer mode is suitable for the target. The compiler will insert conversion instructions as necessary (typically to convert the result to the same width as int). When m is a vector, both operands must have mode m.

This pattern is not allowed to FAIL.

# 'paritym2'

Store into operand 0 the parity of operand 1, i.e. the number of 1-bits in operand 1 modulo 2.

m is either a scalar or vector integer mode. When it is a scalar, operand 1 has mode m but operand 0 can have whatever scalar integer mode is suitable for the target. The compiler will insert conversion instructions as necessary (typically to convert the result to the same width as int). When m is a vector, both operands must have mode m.

This pattern is not allowed to FAIL.

#### 'one\_cmplm2'

Store the bitwise-complement of operand 1 into operand 0.

'cpymemm'

Block copy instruction. The destination and source blocks of memory are the first two operands, and both are mem:BLKs with an address in mode Pmode.

The number of bytes to copy is the third operand, in mode m. Usually, you specify Pmode for m. However, if you can generate better code knowing the range of valid lengths is smaller than those representable in a full Pmode pointer, you should provide a pattern with a mode corresponding to the range of values you can handle efficiently (e.g., QImode for values in the range 0–127; note we avoid numbers that appear negative) and also a pattern with Pmode.

The fourth operand is the known shared alignment of the source and destination, in the form of a const\_int rtx. Thus, if the compiler knows that both source and destination are word-aligned, it may provide the value 4 for this operand.

Optional operands 5 and 6 specify expected alignment and size of block respectively. The expected alignment differs from alignment in operand 4 in a way that the blocks are not required to be aligned according to it in all cases. This expected alignment is also in bytes, just like operand 4. Expected size, when unknown, is set to (const\_int -1).

Descriptions of multiple  $\operatorname{cpymem} m$  patterns can only be beneficial if the patterns for smaller modes have fewer restrictions on their first, second and fourth operands. Note that the mode m in  $\operatorname{cpymem} m$  does not impose any restriction on the mode of individually copied data units in the block.

The cpymemm patterns need not give special consideration to the possibility that the source and destination strings might overlap. These patterns are used to do inline expansion of \_\_builtin\_memcpy.

'movmemm'

Block move instruction. The destination and source blocks of memory are the first two operands, and both are mem: BLKs with an address in mode Pmode.

The number of bytes to copy is the third operand, in mode m. Usually, you specify Pmode for m. However, if you can generate better code knowing the range of valid lengths is smaller than those representable in a full Pmode pointer, you should provide a pattern with a mode corresponding to the range of values you can handle efficiently (e.g., QImode for values in the range 0–127; note we avoid numbers that appear negative) and also a pattern with Pmode.

The fourth operand is the known shared alignment of the source and destination, in the form of a const\_int rtx. Thus, if the compiler knows that both source and destination are word-aligned, it may provide the value 4 for this operand.

Optional operands 5 and 6 specify expected alignment and size of block respectively. The expected alignment differs from alignment in operand 4 in a way that the blocks are not required to be aligned according to it in all cases. This expected alignment is also in bytes, just like operand 4. Expected size, when unknown, is set to (const\_int -1).

Descriptions of multiple movmemm patterns can only be beneficial if the patterns for smaller modes have fewer restrictions on their first, second and fourth operands. Note that the mode m in movmemm does not impose any restriction on the mode of individually copied data units in the block.

The movmemm patterns must correctly handle the case where the source and destination strings overlap. These patterns are used to do inline expansion of \_\_builtin\_memmove.

'movstr'

String copy instruction, with stpcpy semantics. Operand 0 is an output operand in mode Pmode. The addresses of the destination and source strings are operands 1 and 2, and both are mem:BLKs with addresses in mode Pmode. The execution of the expansion of this pattern should store in operand 0 the address in which the NUL terminator was stored in the destination string.

This pattern has also several optional operands that are same as in setmem.

'setmemm'

Block set instruction. The destination string is the first operand, given as a mem:BLK whose address is in mode Pmode. The number of bytes to set is the second operand, in mode m. The value to initialize the memory with is the third operand. Targets that only support the clearing of memory should reject any value that is not the constant 0. See 'cpymemm' for a discussion of the choice of mode.

The fourth operand is the known alignment of the destination, in the form of a const\_int rtx. Thus, if the compiler knows that the destination is wordaligned, it may provide the value 4 for this operand.

Optional operands 5 and 6 specify expected alignment and size of block respectively. The expected alignment differs from alignment in operand 4 in a way that the blocks are not required to be aligned according to it in all cases. This expected alignment is also in bytes, just like operand 4. Expected size, when unknown, is set to (const\_int -1). Operand 7 is the minimal size of the block and operand 8 is the maximal size of the block (NULL if it cannot be represented as CONST\_INT). Operand 9 is the probable maximal size (i.e. we cannot rely on it for correctness, but it can be used for choosing proper code sequence for a given size).

The use for multiple setmemm is as for cpymemm.

'cmpstrnm'

String compare instruction, with five operands. Operand 0 is the output; it has mode m. The remaining four operands are like the operands of 'cpymemm'. The two memory blocks specified are compared byte by byte in lexicographic order starting at the beginning of each string. The instruction is not allowed to prefetch more than one byte at a time since either string may end in the first byte and reading past that may access an invalid page or segment and cause a fault. The comparison terminates early if the fetched bytes are different or if they are equal to zero. The effect of the instruction is to store a value in operand 0 whose sign indicates the result of the comparison.

'cmpstrm'

String compare instruction, without known maximum length. Operand 0 is the output; it has mode m. The second and third operand are the blocks of memory to be compared; both are mem: BLK with an address in mode Pmode.

The fourth operand is the known shared alignment of the source and destination, in the form of a const\_int rtx. Thus, if the compiler knows that both source and destination are word-aligned, it may provide the value 4 for this operand.

The two memory blocks specified are compared byte by byte in lexicographic order starting at the beginning of each string. The instruction is not allowed to prefetch more than one byte at a time since either string may end in the first byte and reading past that may access an invalid page or segment and cause a fault. The comparison will terminate when the fetched bytes are different or if they are equal to zero. The effect of the instruction is to store a value in operand 0 whose sign indicates the result of the comparison.

'cmpmemm'

Block compare instruction, with five operands like the operands of 'cmpstrm'. The two memory blocks specified are compared byte by byte in lexicographic order starting at the beginning of each block. Unlike 'cmpstrm' the instruction can prefetch any bytes in the two memory blocks. Also unlike 'cmpstrm' the comparison will not stop if both bytes are zero. The effect of the instruction is to store a value in operand 0 whose sign indicates the result of the comparison.

'strlenm'

Compute the length of a string, with three operands. Operand 0 is the result (of mode m), operand 1 is a mem referring to the first character of the string, operand 2 is the character to search for (normally zero), and operand 3 is a constant describing the known alignment of the beginning of the string.

#### 'floatmn2'

Convert signed integer operand 1 (valid for fixed point mode m) to floating point mode n and store in operand 0 (which has mode n).

## 'floatunsmn2'

Convert unsigned integer operand 1 (valid for fixed point mode m) to floating point mode n and store in operand 0 (which has mode n).

'fixmn2' Convert operand 1 (valid for floating point mode m) to fixed point mode n as a signed number and store in operand 0 (which has mode n). This instruction's result is defined only when the value of operand 1 is an integer.

If the machine description defines this pattern, it also needs to define the ftrunc pattern.

#### 'fixunsmn2'

Convert operand 1 (valid for floating point mode m) to fixed point mode n as an unsigned number and store in operand 0 (which has mode n). This instruction's result is defined only when the value of operand 1 is an integer.

#### 'ftruncm2'

Convert operand 1 (valid for floating point mode m) to an integer value, still represented in floating point mode m, and store it in operand 0 (valid for floating point mode m).

## 'fix\_truncmn2'

Like 'fixmn2' but works for any floating point value of mode m by converting the value to an integer.

### 'fixuns\_truncmn2'

Like 'fixunsmn2' but works for any floating point value of mode m by converting the value to an integer.

## 'truncmn2'

Truncate operand 1 (valid for mode m) to mode n and store in operand 0 (which has mode n). Both modes must be fixed point or both floating point.

#### 'extendmn2'

Sign-extend operand 1 (valid for mode m) to mode n and store in operand 0 (which has mode n). Both modes must be fixed point or both floating point.

#### 'zero\_extendmn2'

Zero-extend operand 1 (valid for mode m) to mode n and store in operand 0 (which has mode n). Both modes must be fixed point.

#### 'fractmn2'

Convert operand 1 of mode m to mode n and store in operand 0 (which has mode n). Mode m and mode n could be fixed-point to fixed-point, signed integer to fixed-point, fixed-point to signed integer, floating-point to fixed-point, or fixed-point to floating-point. When overflows or underflows happen, the results are undefined.

#### 'satfractmn2'

Convert operand 1 of mode m to mode n and store in operand 0 (which has mode n). Mode m and mode n could be fixed-point to fixed-point, signed integer to fixed-point, or floating-point to fixed-point. When overflows or underflows happen, the instruction saturates the results to the maximum or the minimum.

## 'fractunsmn2'

Convert operand 1 of mode m to mode n and store in operand 0 (which has mode n). Mode m and mode n could be unsigned integer to fixed-point, or fixed-point to unsigned integer. When overflows or underflows happen, the results are undefined.

#### 'satfractunsmn2'

Convert unsigned integer operand 1 of mode m to fixed-point mode n and store in operand 0 (which has mode n). When overflows or underflows happen, the instruction saturates the results to the maximum or the minimum.

'extvm' Extract a bit-field from register operand 1, sign-extend it, and store it in operand 0. Operand 2 specifies the width of the field in bits and operand 3 the starting bit, which counts from the most significant bit if 'BITS\_BIG\_ENDIAN' is true and from the least significant bit otherwise.

Operands 0 and 1 both have mode m. Operands 2 and 3 have a target-specific mode.

## 'extvmisalignm'

Extract a bit-field from memory operand 1, sign extend it, and store it in operand 0. Operand 2 specifies the width in bits and operand 3 the starting bit. The starting bit is always somewhere in the first byte of operand 1; it counts from the most significant bit if 'BITS\_BIG\_ENDIAN' is true and from the least significant bit otherwise.

Operand 0 has mode m while operand 1 has BLK mode. Operands 2 and 3 have a target-specific mode.

The instruction must not read beyond the last byte of the bit-field.

'extzvm' Like 'extvm' except that the bit-field value is zero-extended.

#### 'extzvmisalignm'

Like 'extymisalignm' except that the bit-field value is zero-extended.

'insvm'

Insert operand 3 into a bit-field of register operand 0. Operand 1 specifies the width of the field in bits and operand 2 the starting bit, which counts from the most significant bit if 'BITS\_BIG\_ENDIAN' is true and from the least significant bit otherwise.

Operands 0 and 3 both have mode m. Operands 1 and 2 have a target-specific mode.

## 'insvmisalignm'

Insert operand 3 into a bit-field of memory operand 0. Operand 1 specifies the width of the field in bits and operand 2 the starting bit. The starting bit is always somewhere in the first byte of operand 0; it counts from the most significant bit if 'BITS\_BIG\_ENDIAN' is true and from the least significant bit otherwise.

Operand 3 has mode m while operand 0 has BLK mode. Operands 1 and 2 have a target-specific mode.

The instruction must not read or write beyond the last byte of the bit-field.

'extv'

Extract a bit-field from operand 1 (a register or memory operand), where operand 2 specifies the width in bits and operand 3 the starting bit, and store it in operand 0. Operand 0 must have mode word\_mode. Operand 1 may have mode byte\_mode or word\_mode; often word\_mode is allowed only for registers. Operands 2 and 3 must be valid for word\_mode.

The RTL generation pass generates this instruction only with constants for operands 2 and 3 and the constant is never zero for operand 2.

The bit-field value is sign-extended to a full word integer before it is stored in operand 0.

This pattern is deprecated; please use 'extvm' and extvmisalignm instead.

'extzv'

Like 'extv' except that the bit-field value is zero-extended.

This pattern is deprecated; please use 'extzvm' and extzvmisalignm instead.

'insv'

Store operand 3 (which must be valid for word\_mode) into a bit-field in operand 0, where operand 1 specifies the width in bits and operand 2 the starting bit. Operand 0 may have mode byte\_mode or word\_mode; often word\_mode is allowed only for registers. Operands 1 and 2 must be valid for word\_mode.

The RTL generation pass generates this instruction only with constants for operands 1 and 2 and the constant is never zero for operand 1.

This pattern is deprecated; please use 'insvm' and insvmisalignm instead.

#### 'movmodecc'

Conditionally move operand 2 or operand 3 into operand 0 according to the comparison in operand 1. If the comparison is true, operand 2 is moved into operand 0, otherwise operand 3 is moved.

The mode of the operands being compared need not be the same as the operands being moved. Some machines, sparc64 for example, have instructions that conditionally move an integer value based on the floating point condition codes and vice versa.

If the machine does not have conditional move instructions, do not define these patterns.

#### 'addmodecc'

Similar to 'movmodecc' but for conditional addition. Conditionally move operand 2 or (operands 2 + operand 3) into operand 0 according to the comparison in operand 1. If the comparison is false, operand 2 is moved into operand 0, otherwise (operand 2 + operand 3) is moved.

```
'cond_addmode'
'cond_submode'
'cond_mulmode'
'cond_divmode'
'cond_udivmode'
'cond_modmode'
'cond_andmode'
'cond_iormode'
'cond_sminmode'
'cond_smaxmode'
'cond_uminmode'
'cond_uminmode'
'cond_uminmode'
```

When operand 1 is true, perform an operation on operands 2 and 3 and store the result in operand 0, otherwise store operand 4 in operand 0. The operation works elementwise if the operands are vectors.

The scalar case is equivalent to:

```
op0 = op1 ? op2 op op3 : op4;
```

while the vector case is equivalent to:

```
for (i = 0; i < GET_MODE_NUNITS (m); i++)
  op0[i] = op1[i] ? op2[i] op op3[i] : op4[i];</pre>
```

where, for example, op is + for 'cond\_addmode'.

When defined for floating-point modes, the contents of 'op3[i]' are not interpreted if 'op1[i]' is false, just like they would not be in a normal C '?:' condition.

Operands 0, 2, 3 and 4 all have mode m. Operand 1 is a scalar integer if m is scalar, otherwise it has the mode returned by TARGET\_VECTORIZE\_GET\_MASK\_MODE.

```
'cond_fmamode'
'cond_fmsmode'
'cond_fnmamode'
'cond_fnmsmode'
```

Like 'cond\_addm', except that the conditional operation takes 3 operands rather than two. For example, the vector form of 'cond\_fmamode' is equivalent to:

```
for (i = 0; i < GET_MODE_NUNITS (m); i++)
  op0[i] = op1[i] ? fma (op2[i], op3[i], op4[i]) : op5[i];</pre>
```

#### 'negmodecc'

Similar to 'movmodecc' but for conditional negation. Conditionally move the negation of operand 2 or the unchanged operand 3 into operand 0 according to the comparison in operand 1. If the comparison is true, the negation of operand 2 is moved into operand 0, otherwise operand 3 is moved.

#### 'notmodecc'

Similar to 'negmodecc' but for conditional complement. Conditionally move the bitwise complement of operand 2 or the unchanged operand 3 into operand 0 according to the comparison in operand 1. If the comparison is true, the complement of operand 2 is moved into operand 0, otherwise operand 3 is moved.

## 'cstoremode4'

Store zero or nonzero in operand 0 according to whether a comparison is true. Operand 1 is a comparison operator. Operand 2 and operand 3 are the first and second operand of the comparison, respectively. You specify the mode that operand 0 must have when you write the match\_operand expression. The compiler automatically sees which mode you have used and supplies an operand of that mode.

The value stored for a true condition must have 1 as its low bit, or else must be negative. Otherwise the instruction is not suitable and you should omit it from the machine description. You describe to the compiler exactly which value is stored by defining the macro STORE\_FLAG\_VALUE (see Section 18.31 [Misc], page 642). If a description cannot be found that can be used for all the possible comparison operators, you should pick one and use a define\_expand to map all results onto the one you chose.

These operations may FAIL, but should do so only in relatively uncommon cases; if they would FAIL for common cases involving integer comparisons, it is best to restrict the predicates to not allow these operands. Likewise if a given comparison operator will always fail, independent of the operands (for floating-point modes, the ordered\_comparison\_operator predicate is often useful in this case).

If this pattern is omitted, the compiler will generate a conditional branch—for example, it may copy a constant one to the target and branching around an assignment of zero to the target—or a libcall. If the predicate for operand 1 only rejects some operators, it will also try reordering the operands and/or inverting the result value (e.g. by an exclusive OR). These possibilities could be cheaper or equivalent to the instructions used for the 'cstoremode4' pattern

followed by those required to convert a positive result from STORE\_FLAG\_VALUE to 1; in this case, you can and should make operand 1's predicate reject some operators in the 'cstoremode4' pattern, or remove the pattern altogether from the machine description.

#### 'cbranchmode4'

Conditional branch instruction combined with a compare instruction. Operand 0 is a comparison operator. Operand 1 and operand 2 are the first and second operands of the comparison, respectively. Operand 3 is the code\_label to jump to.

'jump' A jump inside a function; an unconditional branch. Operand 0 is the code\_label to jump to. This pattern name is mandatory on all machines.

'call' Subroutine call instruction returning no value. Operand 0 is the function to call; operand 1 is the number of bytes of arguments pushed as a const\_int; operand 2 is the number of registers used as operands.

On most machines, operand 2 is not actually stored into the RTL pattern. It is supplied for the sake of some RISC machines which need to put this information into the assembler code; they can put it in the RTL instead of operand 1.

Operand 0 should be a mem RTX whose address is the address of the function. Note, however, that this address can be a symbol\_ref expression even if it would not be a legitimate memory address on the target machine. If it is also not a valid argument for a call instruction, the pattern for this operation should be a define\_expand (see Section 17.15 [Expander Definitions], page 438) that places the address into a register and uses that register in the call instruction.

### 'call\_value'

Subroutine call instruction returning a value. Operand 0 is the hard register in which the value is returned. There are three more operands, the same as the three operands of the 'call' instruction (but with numbers increased by one). Subroutines that return BLKmode objects use the 'call' insn.

## 'call\_pop', 'call\_value\_pop'

Similar to 'call' and 'call\_value', except used if defined and if RETURN\_POPS\_ARGS is nonzero. They should emit a parallel that contains both the function call and a set to indicate the adjustment made to the frame pointer.

For machines where RETURN\_POPS\_ARGS can be nonzero, the use of these patterns increases the number of functions for which the frame pointer can be eliminated, if desired.

## 'untyped\_call'

Subroutine call instruction returning a value of any type. Operand 0 is the function to call; operand 1 is a memory location where the result of calling the function is to be stored; operand 2 is a parallel expression where each element is a set expression that indicates the saving of a function return value into the result block.

This instruction pattern should be defined to support \_\_builtin\_apply on machines where special instructions are needed to call a subroutine with arbitrary arguments or to save the value returned. This instruction pattern is

required on machines that have multiple registers that can hold a return value (i.e. FUNCTION\_VALUE\_REGNO\_P is true for more than one register).

'return'

Subroutine return instruction. This instruction pattern name should be defined only if a single instruction can do all the work of returning from a function.

Like the 'movm' patterns, this pattern is also used after the RTL generation phase. In this case it is to support machines where multiple instructions are usually needed to return from a function, but some class of functions only requires one instruction to implement a return. Normally, the applicable functions are those which do not need to save any registers or allocate stack space.

It is valid for this pattern to expand to an instruction using simple\_return if no epilogue is required.

## 'simple\_return'

Subroutine return instruction. This instruction pattern name should be defined only if a single instruction can do all the work of returning from a function on a path where no epilogue is required. This pattern is very similar to the return instruction pattern, but it is emitted only by the shrink-wrapping optimization on paths where the function prologue has not been executed, and a function return should occur without any of the effects of the epilogue. Additional uses may be introduced on paths where both the prologue and the epilogue have executed.

For such machines, the condition specified in this pattern should only be true when reload\_completed is nonzero and the function's epilogue would only be a single instruction. For machines with register windows, the routine leaf\_function\_p may be used to determine if a register window push is required.

Machines that have conditional return instructions should define patterns such as

where *condition* would normally be the same condition specified on the named 'return' pattern.

#### 'untyped\_return'

Untyped subroutine return instruction. This instruction pattern should be defined to support \_\_builtin\_return on machines where special instructions are needed to return a value of any type.

Operand 0 is a memory location where the result of calling a function with \_\_builtin\_apply is stored; operand 1 is a parallel expression where each element is a set expression that indicates the restoring of a function return value from the result block.

'nop' No-op instruction. This instruction pattern name should always be defined to output a no-op in assembler code. (const\_int 0) will do as an RTL pattern.

## 'indirect\_jump'

An instruction to jump to an address which is operand zero. This pattern name is mandatory on all machines.

'casesi' Instruction to jump through a dispatch table, including bounds checking. This instruction takes five operands:

- 1. The index to dispatch on, which has mode SImode.
- 2. The lower bound for indices in the table, an integer constant.
- 3. The total range of indices in the table—the largest index minus the smallest one (both inclusive).
- 4. A label that precedes the table itself.
- 5. A label to jump to if the index has a value outside the bounds.

The table is an addr\_vec or addr\_diff\_vec inside of a jump\_table\_data. The number of elements in the table is one plus the difference between the upper bound and the lower bound.

## 'tablejump'

Instruction to jump to a variable address. This is a low-level capability which can be used to implement a dispatch table when there is no 'casesi' pattern.

This pattern requires two operands: the address or offset, and a label which should immediately precede the jump table. If the macro CASE\_VECTOR\_PC\_RELATIVE evaluates to a nonzero value then the first operand is an offset which counts from the address of the table; otherwise, it is an absolute address to jump to. In either case, the first operand has mode Pmode.

The 'tablejump' insn is always the last insn before the jump table it uses. Its assembler code normally has no need to use the second operand, but you should incorporate it in the RTL pattern so that the jump optimizer will not delete the table as unreachable code.

### 'doloop\_end'

Conditional branch instruction that decrements a register and jumps if the register is nonzero. Operand 0 is the register to decrement and test; operand 1 is the label to jump to if the register is nonzero. See Section 17.13 [Looping Patterns], page 434.

This optional instruction pattern should be defined for machines with low-overhead looping instructions as the loop optimizer will try to modify suitable loops to utilize it. The target hook TARGET\_CAN\_USE\_DOLOOP\_P controls the conditions under which low-overhead loops can be used.

### 'doloop\_begin'

Companion instruction to doloop\_end required for machines that need to perform some initialization, such as loading a special counter register. Operand 1 is the associated doloop\_end pattern and operand 0 is the register that it decrements.

If initialization insns do not always need to be emitted, use a define\_expand (see Section 17.15 [Expander Definitions], page 438) and make it fail.

## 'canonicalize\_funcptr\_for\_compare'

Canonicalize the function pointer in operand 1 and store the result into operand 0.

Operand 0 is always a reg and has mode Pmode; operand 1 may be a reg, mem, symbol\_ref, const\_int, etc and also has mode Pmode.

Canonicalization of a function pointer usually involves computing the address of the function which would be called if the function pointer were used in an indirect call.

Only define this pattern if function pointers on the target machine can have different values but still call the same function when used in an indirect call.

```
'save_stack_block'
'save_stack_function'
'save_stack_nonlocal'
'restore_stack_block'
'restore_stack_function'
'restore_stack_nonlocal'
```

Most machines save and restore the stack pointer by copying it to or from an object of mode Pmode. Do not define these patterns on such machines.

Some machines require special handling for stack pointer saves and restores. On those machines, define the patterns corresponding to the non-standard cases by using a define\_expand (see Section 17.15 [Expander Definitions], page 438) that produces the required insns. The three types of saves and restores are:

- 1. 'save\_stack\_block' saves the stack pointer at the start of a block that allocates a variable-sized object, and 'restore\_stack\_block' restores the stack pointer when the block is exited.
- 'save\_stack\_function' and 'restore\_stack\_function' do a similar job
  for the outermost block of a function and are used when the function allocates variable-sized objects or calls alloca. Only the epilogue uses the
  restored stack pointer, allowing a simpler save or restore sequence on some
  machines.
- 3. 'save\_stack\_nonlocal' is used in functions that contain labels branched to by nested functions. It saves the stack pointer in such a way that the inner function can use 'restore\_stack\_nonlocal' to restore the stack pointer. The compiler generates code to restore the frame and argument pointer registers, but some machines require saving and restoring additional data such as register window information or stack backchains. Place insns in these patterns to save and restore any such required data.

When saving the stack pointer, operand 0 is the save area and operand 1 is the stack pointer. The mode used to allocate the save area defaults to Pmode but you can override that choice by defining the STACK\_SAVEAREA\_MODE macro (see Section 18.5 [Storage Layout], page 490). You must specify an integral mode, or VOIDmode if no save area is needed for a particular type of save (either because

no save is needed or because a machine-specific save area can be used). Operand 0 is the stack pointer and operand 1 is the save area for restore operations. If 'save\_stack\_block' is defined, operand 0 must not be VOIDmode since these saves can be arbitrarily nested.

A save area is a mem that is at a constant offset from virtual\_stack\_vars\_rtx when the stack pointer is saved for use by nonlocal gotos and a reg in the other two cases.

## 'allocate\_stack'

Subtract (or add if STACK\_GROWS\_DOWNWARD is undefined) operand 1 from the stack pointer to create space for dynamically allocated data.

Store the resultant pointer to this space into operand 0. If you are allocating space from the main stack, do this by emitting a move insn to copy virtual\_stack\_dynamic\_rtx to operand 0. If you are allocating the space elsewhere, generate code to copy the location of the space to operand 0. In the latter case, you must ensure this space gets freed when the corresponding space on the main stack is free.

Do not define this pattern if all that must be done is the subtraction. Some machines require other operations such as stack probes or maintaining the back chain. Define this pattern to emit those operations in addition to updating the stack pointer.

#### 'check\_stack'

If stack checking (see Section 18.9.3 [Stack Checking], page 528) cannot be done on your system by probing the stack, define this pattern to perform the needed check and signal an error if the stack has overflowed. The single operand is the address in the stack farthest from the current stack pointer that you need to validate. Normally, on platforms where this pattern is needed, you would obtain the stack limit from a global or thread-specific variable or register.

### 'probe\_stack\_address'

If stack checking (see Section 18.9.3 [Stack Checking], page 528) can be done on your system by probing the stack but without the need to actually access it, define this pattern and signal an error if the stack has overflowed. The single operand is the memory address in the stack that needs to be probed.

## 'probe\_stack'

If stack checking (see Section 18.9.3 [Stack Checking], page 528) can be done on your system by probing the stack but doing it with a "store zero" instruction is not valid or optimal, define this pattern to do the probing differently and signal an error if the stack has overflowed. The single operand is the memory reference in the stack that needs to be probed.

#### 'nonlocal\_goto'

Emit code to generate a non-local goto, e.g., a jump from one function to a label in an outer function. This pattern has four arguments, each representing a value to be used in the jump. The first argument is to be loaded into the frame pointer, the second is the address to branch to (code to dispatch to the actual label), the third is the address of a location where the stack is saved, and

the last is the address of the label, to be placed in the location for the incoming static chain.

On most machines you need not define this pattern, since GCC will already generate the correct code, which is to load the frame pointer and static chain, restore the stack (using the 'restore\_stack\_nonlocal' pattern, if defined), and jump indirectly to the dispatcher. You need only define this pattern if this code will not work on your machine.

## 'nonlocal\_goto\_receiver'

This pattern, if defined, contains code needed at the target of a nonlocal goto after the code already generated by GCC. You will not normally need to define this pattern. A typical reason why you might need this pattern is if some value, such as a pointer to a global table, must be restored when the frame pointer is restored. Note that a nonlocal goto only occurs within a unit-of-translation, so a global table pointer that is shared by all functions of a given module need not be restored. There are no arguments.

## 'exception\_receiver'

This pattern, if defined, contains code needed at the site of an exception handler that isn't needed at the site of a nonlocal goto. You will not normally need to define this pattern. A typical reason why you might need this pattern is if some value, such as a pointer to a global table, must be restored after control flow is branched to the handler of an exception. There are no arguments.

## 'builtin\_setjmp\_setup'

This pattern, if defined, contains additional code needed to initialize the jmp\_buf. You will not normally need to define this pattern. A typical reason why you might need this pattern is if some value, such as a pointer to a global table, must be restored. Though it is preferred that the pointer value be recalculated if possible (given the address of a label for instance). The single argument is a pointer to the jmp\_buf. Note that the buffer is five words long and that the first three are normally used by the generic mechanism.

## 'builtin\_setjmp\_receiver'

This pattern, if defined, contains code needed at the site of a built-in setjmp that isn't needed at the site of a nonlocal goto. You will not normally need to define this pattern. A typical reason why you might need this pattern is if some value, such as a pointer to a global table, must be restored. It takes one argument, which is the label to which builtin\_longjmp transferred control; this pattern may be emitted at a small offset from that label.

## 'builtin\_longjmp'

This pattern, if defined, performs the entire action of the longjmp. You will not normally need to define this pattern unless you also define builtin\_setjmp\_setup. The single argument is a pointer to the jmp\_buf.

## 'eh\_return'

This pattern, if defined, affects the way \_\_builtin\_eh\_return, and thence the call frame exception handling library routines, are built. It is intended to handle non-trivial actions needed along the abnormal return path.

The address of the exception handler to which the function should return is passed as operand to this pattern. It will normally need to copied by the pattern to some special register or memory location. If the pattern needs to determine the location of the target call frame in order to do so, it may use EH\_RETURN\_STACKADJ\_RTX, if defined; it will have already been assigned.

If this pattern is not defined, the default action will be to simply copy the return address to EH\_RETURN\_HANDLER\_RTX. Either that macro or this pattern needs to be defined if call frame exception handling is to be used.

## 'prologue'

This pattern, if defined, emits RTL for entry to a function. The function entry is responsible for setting up the stack frame, initializing the frame pointer register, saving callee saved registers, etc.

Using a prologue pattern is generally preferred over defining TARGET\_ASM\_FUNCTION\_PROLOGUE to emit assembly code for the prologue.

The prologue pattern is particularly useful for targets which perform instruction scheduling.

## 'window\_save'

This pattern, if defined, emits RTL for a register window save. It should be defined if the target machine has register windows but the window events are decoupled from calls to subroutines. The canonical example is the SPARC architecture.

#### 'epilogue'

This pattern emits RTL for exit from a function. The function exit is responsible for deallocating the stack frame, restoring callee saved registers and emitting the return instruction.

Using an epilogue pattern is generally preferred over defining TARGET\_ASM\_FUNCTION\_EPILOGUE to emit assembly code for the epilogue.

The epilogue pattern is particularly useful for targets which perform instruction scheduling or which have delay slots for their return instruction.

## 'sibcall\_epilogue'

This pattern, if defined, emits RTL for exit from a function without the final branch back to the calling function. This pattern will be emitted before any sibling call (aka tail call) sites.

The sibcall\_epilogue pattern must not clobber any arguments used for parameter passing or any stack slots for arguments passed to the current function.

'trap' This pattern, if defined, signals an error, typically by causing some kind of signal to be raised.

#### 'ctrapMM4'

Conditional trap instruction. Operand 0 is a piece of RTL which performs a comparison, and operands 1 and 2 are the arms of the comparison. Operand 3 is the trap code, an integer.

A typical ctrap pattern looks like

## 'prefetch'

This pattern, if defined, emits code for a non-faulting data prefetch instruction. Operand 0 is the address of the memory to prefetch. Operand 1 is a constant 1 if the prefetch is preparing for a write to the memory address, or a constant 0 otherwise. Operand 2 is the expected degree of temporal locality of the data and is a value between 0 and 3, inclusive; 0 means that the data has no temporal locality, so it need not be left in the cache after the access; 3 means that the data has a high degree of temporal locality and should be left in all levels of cache possible; 1 and 2 mean, respectively, a low or moderate degree of temporal locality.

Targets that do not support write prefetches or locality hints can ignore the values of operands 1 and 2.

## 'blockage'

This pattern defines a pseudo inso that prevents the instruction scheduler and other passes from moving instructions and using register equivalences across the boundary defined by the blockage inso. This needs to be an UNSPEC\_VOLATILE pattern or a volatile ASM.

#### 'memory\_blockage'

This pattern, if defined, represents a compiler memory barrier, and will be placed at points across which RTL passes may not propagate memory accesses. This instruction needs to read and write volatile BLKmode memory. It does not need to generate any machine instruction. If this pattern is not defined, the compiler falls back to emitting an instruction corresponding to asm volatile (""::: "memory").

#### 'memory\_barrier'

If the target memory model is not fully synchronous, then this pattern should be defined to an instruction that orders both loads and stores before the instruction with respect to loads and stores after the instruction. This pattern has no operands.

## 'speculation\_barrier'

If the target can support speculative execution, then this pattern should be defined to an instruction that will block subsequent execution until any prior speculation conditions has been resolved. The pattern must also ensure that the compiler cannot move memory operations past the barrier, so it needs to be an UNSPEC\_VOLATILE pattern. The pattern has no operands.

If this pattern is not defined then the default expansion of \_\_builtin\_speculation\_safe\_value will emit a warning. You can suppress this warning by defining this pattern with a final condition of 0 (zero), which tells the compiler that a speculation barrier is not needed for this target.

## 'sync\_compare\_and\_swapmode'

This pattern, if defined, emits code for an atomic compare-and-swap operation. Operand 1 is the memory on which the atomic operation is performed. Operand 2 is the "old" value to be compared against the current contents of the memory location. Operand 3 is the "new" value to store in the memory if the compare succeeds. Operand 0 is the result of the operation; it should contain the contents of the memory before the operation. If the compare succeeds, this should obviously be a copy of operand 2.

This pattern must show that both operand 0 and operand 1 are modified.

This pattern must issue any memory barrier instructions such that all memory operations before the atomic operation occur before the atomic operation and all memory operations after the atomic operation occur after the atomic operation.

For targets where the success or failure of the compare-and-swap operation is available via the status flags, it is possible to avoid a separate compare operation and issue the subsequent branch or store-flag operation immediately after the compare-and-swap. To this end, GCC will look for a MODE\_CC set in the output of sync\_compare\_and\_swapmode; if the machine description includes such a set, the target should also define special cbranchcc4 and/or cstorecc4 instructions. GCC will then be able to take the destination of the MODE\_CC set and pass it to the cbranchcc4 or cstorecc4 pattern as the first operand of the comparison (the second will be (const\_int 0)).

For targets where the operating system may provide support for this operation via library calls, the sync\_compare\_and\_swap\_optab may be initialized to a function with the same interface as the \_\_sync\_val\_compare\_and\_swap\_n built-in. If the entire set of \_\_sync builtins are supported via library calls, the target can initialize all of the optabs at once with init\_sync\_libfuncs. For the purposes of C++11 std::atomic::is\_lock\_free, it is assumed that these library calls do not use any kind of interruptable locking.

```
'sync_addmode', 'sync_submode'
'sync_iormode', 'sync_andmode'
'sync_xormode', 'sync_nandmode'
```

These patterns emit code for an atomic operation on memory. Operand 0 is the memory on which the atomic operation is performed. Operand 1 is the second operand to the binary operator.

This pattern must issue any memory barrier instructions such that all memory operations before the atomic operation occur before the atomic operation and all memory operations after the atomic operation occur after the atomic operation.

If these patterns are not defined, the operation will be constructed from a compare-and-swap operation, if defined.

```
'sync_old_addmode', 'sync_old_submode'
'sync_old_iormode', 'sync_old_andmode'
'sync_old_xormode', 'sync_old_nandmode'
```

These patterns emit code for an atomic operation on memory, and return the value that the memory contained before the operation. Operand 0 is the result

value, operand 1 is the memory on which the atomic operation is performed, and operand 2 is the second operand to the binary operator.

This pattern must issue any memory barrier instructions such that all memory operations before the atomic operation occur before the atomic operation and all memory operations after the atomic operation occur after the atomic operation.

If these patterns are not defined, the operation will be constructed from a compare-and-swap operation, if defined.

```
'sync_new_addmode', 'sync_new_submode'
```

These patterns are like their sync\_old\_op counterparts, except that they return the value that exists in the memory location after the operation, rather than before the operation.

## 'sync\_lock\_test\_and\_setmode'

This pattern takes two forms, based on the capabilities of the target. In either case, operand 0 is the result of the operand, operand 1 is the memory on which the atomic operation is performed, and operand 2 is the value to set in the lock.

In the ideal case, this operation is an atomic exchange operation, in which the previous value in memory operand is copied into the result operand, and the value operand is stored in the memory operand.

For less capable targets, any value operand that is not the constant 1 should be rejected with FAIL. In this case the target may use an atomic test-and-set bit operation. The result operand should contain 1 if the bit was previously set and 0 if the bit was previously clear. The true contents of the memory operand are implementation defined.

This pattern must issue any memory barrier instructions such that the pattern as a whole acts as an acquire barrier, that is all memory operations after the pattern do not occur until the lock is acquired.

If this pattern is not defined, the operation will be constructed from a compareand-swap operation, if defined.

#### 'sync\_lock\_releasemode'

This pattern, if defined, releases a lock set by sync\_lock\_test\_and\_setmode. Operand 0 is the memory that contains the lock; operand 1 is the value to store in the lock.

If the target doesn't implement full semantics for sync\_lock\_test\_and\_setmode, any value operand which is not the constant 0 should be rejected with FAIL, and the true contents of the memory operand are implementation defined.

This pattern must issue any memory barrier instructions such that the pattern as a whole acts as a release barrier, that is the lock is released only after all previous memory operations have completed.

If this pattern is not defined, then a memory\_barrier pattern will be emitted, followed by a store of the value to the memory operand.

<sup>&#</sup>x27;sync\_new\_iormode', 'sync\_new\_andmode'

<sup>&#</sup>x27;sync\_new\_xormode', 'sync\_new\_nandmode'

## 'atomic\_compare\_and\_swapmode'

This pattern, if defined, emits code for an atomic compare-and-swap operation with memory model semantics. Operand 2 is the memory on which the atomic operation is performed. Operand 0 is an output operand which is set to true or false based on whether the operation succeeded. Operand 1 is an output operand which is set to the contents of the memory before the operation was attempted. Operand 3 is the value that is expected to be in memory. Operand 4 is the value to put in memory if the expected value is found there. Operand 5 is set to 1 if this compare and swap is to be treated as a weak operation. Operand 6 is the memory model to be used if the operation is a success. Operand 7 is the memory model to be used if the operation fails.

If memory referred to in operand 2 contains the value in operand 3, then operand 4 is stored in memory pointed to by operand 2 and fencing based on the memory model in operand 6 is issued.

If memory referred to in operand 2 does not contain the value in operand 3, then fencing based on the memory model in operand 7 is issued.

If a target does not support weak compare-and-swap operations, or the port elects not to implement weak operations, the argument in operand 5 can be ignored. Note a strong implementation must be provided.

If this pattern is not provided, the <code>\_\_atomic\_compare\_exchange</code> built-in functions will utilize the legacy <code>sync\_compare\_and\_swap</code> pattern with an <code>\_\_ATOMIC\_SEQ\_CST</code> memory model.

## 'atomic\_loadmode'

This pattern implements an atomic load operation with memory model semantics. Operand 1 is the memory address being loaded from. Operand 0 is the result of the load. Operand 2 is the memory model to be used for the load operation.

If not present, the <u>\_\_atomic\_load</u> built-in function will either resort to a normal load with memory barriers, or a compare-and-swap operation if a normal load would not be atomic.

### 'atomic\_storemode'

This pattern implements an atomic store operation with memory model semantics. Operand 0 is the memory address being stored to. Operand 1 is the value to be written. Operand 2 is the memory model to be used for the operation.

If not present, the <code>\_\_atomic\_store</code> built-in function will attempt to perform a normal store and surround it with any required memory fences. If the store would not be atomic, then an <code>\_\_atomic\_exchange</code> is attempted with the result being ignored.

#### 'atomic\_exchangemode'

This pattern implements an atomic exchange operation with memory model semantics. Operand 1 is the memory location the operation is performed on. Operand 0 is an output operand which is set to the original value contained in the memory pointed to by operand 1. Operand 2 is the value to be stored. Operand 3 is the memory model to be used.

If this pattern is not present, the built-in function \_\_atomic\_exchange will attempt to preform the operation with a compare and swap loop.

```
'atomic_addmode', 'atomic_submode'
'atomic_ormode', 'atomic_andmode'
'atomic_xormode', 'atomic_nandmode'
```

These patterns emit code for an atomic operation on memory with memory model semantics. Operand 0 is the memory on which the atomic operation is performed. Operand 1 is the second operand to the binary operator. Operand 2 is the memory model to be used by the operation.

If these patterns are not defined, attempts will be made to use legacy sync patterns, or equivalent patterns which return a result. If none of these are available a compare-and-swap loop will be used.

```
'atomic_fetch_addmode', 'atomic_fetch_submode'
'atomic_fetch_ormode', 'atomic_fetch_andmode'
'atomic_fetch_xormode', 'atomic_fetch_nandmode'
```

These patterns emit code for an atomic operation on memory with memory model semantics, and return the original value. Operand 0 is an output operand which contains the value of the memory location before the operation was performed. Operand 1 is the memory on which the atomic operation is performed. Operand 2 is the second operand to the binary operator. Operand 3 is the memory model to be used by the operation.

If these patterns are not defined, attempts will be made to use legacy sync patterns. If none of these are available a compare-and-swap loop will be used.

```
'atomic_add_fetchmode', 'atomic_sub_fetchmode'
'atomic_or_fetchmode', 'atomic_and_fetchmode'
'atomic_xor_fetchmode', 'atomic_nand_fetchmode'
```

These patterns emit code for an atomic operation on memory with memory model semantics and return the result after the operation is performed. Operand 0 is an output operand which contains the value after the operation. Operand 1 is the memory on which the atomic operation is performed. Operand 2 is the second operand to the binary operator. Operand 3 is the memory model to be used by the operation.

If these patterns are not defined, attempts will be made to use legacy sync patterns, or equivalent patterns which return the result before the operation followed by the arithmetic operation required to produce the result. If none of these are available a compare-and-swap loop will be used.

## 'atomic\_test\_and\_set'

This pattern emits code for \_\_builtin\_atomic\_test\_and\_set. Operand 0 is an output operand which is set to true if the previous previous contents of the byte was "set", and false otherwise. Operand 1 is the QImode memory to be modified. Operand 2 is the memory model to be used.

The specific value that defines "set" is implementation defined, and is normally based on what is performed by the native atomic test and set instruction.

```
'atomic_bit_test_and_setmode'
'atomic_bit_test_and_complementmode'
'atomic_bit_test_and_resetmode'
```

These patterns emit code for an atomic bitwise operation on memory with memory model semantics, and return the original value of the specified bit. Operand 0 is an output operand which contains the value of the specified bit from the memory location before the operation was performed. Operand 1 is the memory on which the atomic operation is performed. Operand 2 is the bit within the operand, starting with least significant bit. Operand 3 is the memory model to be used by the operation. Operand 4 is a flag - it is const1\_rtx if operand 0 should contain the original value of the specified bit in the least significant bit of the operand, and const0\_rtx if the bit should be in its original position in the operand. atomic\_bit\_test\_and\_setmode atomically sets the specified bit after remembering its original value, atomic\_bit\_test\_and\_complementmode inverts the specified bit and atomic\_bit\_test\_and\_resetmode clears the specified bit.

If these patterns are not defined, attempts will be made to use atomic\_fetch\_ormode, atomic\_fetch\_xormode or atomic\_fetch\_andmode instruction patterns, or their sync counterparts. If none of these are available a compare-and-swap loop will be used.

## 'mem\_thread\_fence'

This pattern emits code required to implement a thread fence with memory model semantics. Operand 0 is the memory model to be used.

For the \_\_ATOMIC\_RELAXED model no instructions need to be issued and this expansion is not invoked.

The compiler always emits a compiler memory barrier regardless of what expanding this pattern produced.

If this pattern is not defined, the compiler falls back to expanding the memory\_barrier pattern, then to emitting \_\_sync\_synchronize library call, and finally to just placing a compiler memory barrier.

```
'get_thread_pointermode'
'set_thread_pointermode'
```

These patterns emit code that reads/sets the TLS thread pointer. Currently, these are only needed if the target needs to support the \_\_builtin\_thread\_pointer and \_\_builtin\_set\_thread\_pointer builtins.

The get/set patterns have a single output/input operand respectively, with mode intended to be Pmode.

### 'stack\_protect\_combined\_set'

This pattern, if defined, moves a ptr\_mode value from an address whose declaration RTX is given in operand 1 to the memory in operand 0 without leaving the value in a register afterward. If several instructions are needed by the target to perform the operation (eg. to load the address from a GOT entry then load the ptr\_mode value and finally store it), it is the backend's responsibility to ensure no intermediate result gets spilled. This is to avoid leaking the value some

place that an attacker might use to rewrite the stack guard slot after having clobbered it.

If this pattern is not defined, then the address declaration is expanded first in the standard way and a stack\_protect\_set pattern is then generated to move the value from that address to the address in operand 0.

## 'stack\_protect\_set'

This pattern, if defined, moves a ptr\_mode value from the valid memory location in operand 1 to the memory in operand 0 without leaving the value in a register afterward. This is to avoid leaking the value some place that an attacker might use to rewrite the stack guard slot after having clobbered it.

Note: on targets where the addressing modes do not allow to load directly from stack guard address, the address is expanded in a standard way first which could cause some spills.

If this pattern is not defined, then a plain move pattern is generated.

## 'stack\_protect\_combined\_test'

This pattern, if defined, compares a ptr\_mode value from an address whose declaration RTX is given in operand 1 with the memory in operand 0 without leaving the value in a register afterward and branches to operand 2 if the values were equal. If several instructions are needed by the target to perform the operation (eg. to load the address from a GOT entry then load the ptr\_mode value and finally store it), it is the backend's responsibility to ensure no intermediate result gets spilled. This is to avoid leaking the value some place that an attacker might use to rewrite the stack guard slot after having clobbered it.

If this pattern is not defined, then the address declaration is expanded first in the standard way and a stack\_protect\_test pattern is then generated to compare the value from that address to the value at the memory in operand 0.

## 'stack\_protect\_test'

This pattern, if defined, compares a ptr\_mode value from the valid memory location in operand 1 with the memory in operand 0 without leaving the value in a register afterward and branches to operand 2 if the values were equal.

If this pattern is not defined, then a plain compare pattern and conditional branch pattern is used.

#### 'clear\_cache'

This pattern, if defined, flushes the instruction cache for a region of memory. The region is bounded to by the Pmode pointers in operand 0 inclusive and operand 1 exclusive.

If this pattern is not defined, a call to the library function \_\_clear\_cache is used.

## 17.10 When the Order of Patterns Matters

Sometimes an insn can match more than one instruction pattern. Then the pattern that appears first in the machine description is the one used. Therefore, more specific patterns

(patterns that will match fewer things) and faster instructions (those that will produce better code when they do match) should usually go first in the description.

In some cases the effect of ordering the patterns can be used to hide a pattern when it is not valid. For example, the 68000 has an instruction for converting a fullword to floating point and another for converting a byte to floating point. An instruction converting an integer to floating point could match either one. We put the pattern to convert the fullword first to make sure that one will be used rather than the other. (Otherwise a large integer might be generated as a single-byte immediate quantity, which would not work.) Instead of using this pattern ordering it would be possible to make the pattern for convert-a-byte smart enough to deal properly with any constant value.

## 17.11 Interdependence of Patterns

In some cases machines support instructions identical except for the machine mode of one or more operands. For example, there may be "sign-extend halfword" and "sign-extend byte" instructions whose patterns are

```
(set (match_operand:SI 0 ...)
      (extend:SI (match_operand:HI 1 ...)))
(set (match_operand:SI 0 ...)
      (extend:SI (match_operand:QI 1 ...)))
```

Constant integers do not specify a machine mode, so an instruction to extend a constant value could match either pattern. The pattern it actually will match is the one that appears first in the file. For correct results, this must be the one for the widest possible mode (HImode, here). If the pattern matches the QImode instruction, the results will be incorrect if the constant value does not actually fit that mode.

Such instructions to extend constants are rarely generated because they are optimized away, but they do occasionally happen in nonoptimized compilations.

If a constraint in a pattern allows a constant, the reload pass may replace a register with a constant permitted by the constraint in some cases. Similarly for memory references. Because of this substitution, you should not provide separate patterns for increment and decrement instructions. Instead, they should be generated from the same pattern that supports register-register add insns by examining the operands and generating the appropriate machine instruction.

# 17.12 Defining Jump Instruction Patterns

GCC does not assume anything about how the machine realizes jumps. The machine description should define a single pattern, usually a define\_expand, which expands to all the required insns.

Usually, this would be a comparison insn to set the condition code and a separate branch insn testing the condition code and branching or not according to its value. For many machines, however, separating compares and branches is limiting, which is why the more flexible approach with one define\_expand is used in GCC. The machine description becomes clearer for architectures that have compare-and-branch instructions but no condition code. It also works better when different sets of comparison operators are supported by different kinds of conditional branches (e.g. integer vs. floating-point), or by conditional branches with respect to conditional stores.

Two separate insns are always used if the machine description represents a condition code register using the legacy RTL expression (cc0), and on most machines that use a separate condition code register (see Section 18.15 [Condition Code], page 572). For machines that use (cc0), in fact, the set and use of the condition code must be separate and adjacent<sup>1</sup>, thus allowing flags in cc\_status to be used (see Section 18.15 [Condition Code], page 572) and so that the comparison and branch insns could be located from each other by using the functions prev\_cc0\_setter and next\_cc0\_user.

Even in this case having a single entry point for conditional branches is advantageous, because it handles equally well the case where a single comparison instruction records the results of both signed and unsigned comparison of the given operands (with the branch insns coming in distinct signed and unsigned flavors) as in the x86 or SPARC, and the case where there are distinct signed and unsigned compare instructions and only one set of conditional branch instructions as in the PowerPC.

## 17.13 Defining Looping Instruction Patterns

Some machines have special jump instructions that can be utilized to make loops more efficient. A common example is the 68000 'dbra' instruction which performs a decrement of a register and a branch if the result was greater than zero. Other machines, in particular digital signal processors (DSPs), have special block repeat instructions to provide low-overhead loop support. For example, the TI TMS320C3x/C4x DSPs have a block repeat instruction that loads special registers to mark the top and end of a loop and to count the number of loop iterations. This avoids the need for fetching and executing a 'dbra'-like instruction and avoids pipeline stalls associated with the jump.

GCC has two special named patterns to support low overhead looping. They are 'doloop\_begin' and 'doloop\_end'. These are emitted by the loop optimizer for certain well-behaved loops with a finite number of loop iterations using information collected during strength reduction.

The 'doloop\_end' pattern describes the actual looping instruction (or the implicit looping operation) and the 'doloop\_begin' pattern is an optional companion pattern that can be used for initialization needed for some low-overhead looping instructions.

Note that some machines require the actual looping instruction to be emitted at the top of the loop (e.g., the TMS320C3x/C4x DSPs). Emitting the true RTL for a looping instruction at the top of the loop can cause problems with flow analysis. So instead, a dummy doloop insn is emitted at the end of the loop. The machine dependent reorg pass checks for the presence of this doloop insn and then searches back to the top of the loop, where it inserts the true looping insn (provided there are no instructions in the loop which would cause problems). Any additional labels can be emitted at this point. In addition, if the desired special iteration counter register was not allocated, this machine dependent reorg pass could emit a traditional compare and jump instruction pair.

For the 'doloop\_end' pattern, the loop optimizer allocates an additional pseudo register as an iteration counter. This pseudo register cannot be used within the loop (i.e., general induction variables cannot be derived from it), however, in many cases the loop induction variable may become redundant and removed by the flow pass.

<sup>&</sup>lt;sup>1</sup> note insns can separate them, though.

The 'doloop\_end' pattern must have a specific structure to be handled correctly by GCC. The example below is taken (slightly simplified) from the PDP-11 target:

```
(define_expand "doloop_end"
  [(parallel [(set (pc)
                   (if_then_else
                     (ne (match_operand:HI 0 "nonimmediate_operand" "+r,!m")
                         (const_int 1))
                     (label_ref (match_operand 1 "" ""))
                     (pc)))
              (set (match_dup 0)
                   (plus:HI (match_dup 0)
                          (const_int -1)))])]
  "{
   if (GET_MODE (operands[0]) != HImode)
      FAIL;
 }")
(define_insn "doloop_end_insn"
  [(set (pc)
        (if_then_else
         (ne (match_operand:HI 0 "nonimmediate_operand" "+r,!m")
             (const_int 1))
         (label_ref (match_operand 1 "" ""))
         (pc)))
   (set (match_dup 0)
        (plus:HI (match_dup 0)
              (const_int -1)))]
   if (which_alternative == 0)
      return "sob %0,%11";
    /* emulate sob */
   output_asm_insn ("dec %0", operands);
   return "bne %11";
```

The first part of the pattern describes the branch condition. GCC supports three cases for the way the target machine handles the loop counter:

- Loop terminates when the loop register decrements to zero. This is represented by a ne comparison of the register (its old value) with constant 1 (as in the example above).
- Loop terminates when the loop register decrements to -1. This is represented by a ne comparison of the register with constant zero.
- Loop terminates when the loop register decrements to a negative value. This is represented by a ge comparison of the register with constant zero. For this case, GCC will attach a REG\_NONNEG note to the doloop\_end insn if it can determine that the register will be non-negative.

Since the doloop\_end insn is a jump insn that also has an output, the reload pass does not handle the output operand. Therefore, the constraint must allow for that operand to be in memory rather than a register. In the example shown above, that is handled (in the doloop\_end\_insn pattern) by using a loop instruction sequence that can handle memory operands when the memory alternative appears.

GCC does not check the mode of the loop register operand when generating the doloop\_end pattern. If the pattern is only valid for some modes but not others, the pattern should be a define\_expand pattern that checks the operand mode in the preparation code, and issues FAIL if an unsupported mode is found. The example above does this, since the machine instruction to be used only exists for HImode.

If the doloop\_end pattern is a define\_expand, there must also be a define\_insn or define\_insn\_and\_split matching the generated pattern. Otherwise, the compiler will fail during loop optimization.

## 17.14 Canonicalization of Instructions

There are often cases where multiple RTL expressions could represent an operation performed by a single machine instruction. This situation is most commonly encountered with logical, branch, and multiply-accumulate instructions. In such cases, the compiler attempts to convert these multiple RTL expressions into a single canonical form to reduce the number of insn patterns required.

In addition to algebraic simplifications, following canonicalizations are performed:

- For commutative and comparison operators, a constant is always made the second operand. If a machine only supports a constant as the second operand, only patterns that match a constant in the second operand need be supplied.
- For associative operators, a sequence of operators will always chain to the left; for instance, only the left operand of an integer plus can itself be a plus. and, ior, xor, plus, mult, smin, smax, umin, and umax are associative when applied to integers, and sometimes to floating-point.
- For these operators, if only one operand is a neg, not, mult, plus, or minus expression, it will be the first operand.
- In combinations of neg, mult, plus, and minus, the neg operations (if any) will be moved inside the operations as far as possible. For instance, (neg (mult A B)) is canonicalized as (mult (neg A) B), but (plus (mult (neg B) C) A) is canonicalized as (minus A (mult B C)).
- For the compare operator, a constant is always the second operand if the first argument is a condition code register or (cc0).
- For instructions that inherently set a condition code register, the compare operator is always written as the first RTL expression of the parallel instruction pattern. For example,

• An operand of neg, not, mult, plus, or minus is made the first operand under the same conditions as above.

- (ltu (plus a b) b) is converted to (ltu (plus a b) a). Likewise with geu instead of ltu.
- (minus x (const\_int n)) is converted to (plus x (const\_int -n)).
- Within address computations (i.e., inside mem), a left shift is converted into the appropriate multiplication by a power of two.
- De Morgan's Law is used to move bitwise negation inside a bitwise logical-and or logical-or operation. If this results in only one operand being a **not** expression, it will be the first one.

A machine that has an instruction that performs a bitwise logical-and of one operand with the bitwise negation of the other should specify the pattern for that instruction as

Similarly, a pattern for a "NAND" instruction should be written

In both cases, it is not necessary to include patterns for the many logically equivalent RTL expressions.

- The only possible RTL expressions involving both bitwise exclusive-or and bitwise negation are (xor:m x y) and (not:m (xor:m x y)).
- The sum of three items, one of which is a constant, will only appear in the form (plus:m x y) constant)
- Equality comparisons of a group of bits (usually a single bit) with zero will be written using zero\_extract rather than the equivalent and or sign\_extract operations.
- (sign\_extend:m1 (mult:m2 (sign\_extend:m2 x) (sign\_extend:m2 y))) is converted to (mult:m1 (sign\_extend:m1 x) (sign\_extend:m1 y)), and likewise for zero\_extend.
- (sign\_extend:m1 (mult:m2 (ashiftrt:m2 x s) (sign\_extend:m2 y))) is converted to (mult:m1 (sign\_extend:m1 (ashiftrt:m2 x s)) (sign\_extend:m1 y)), and likewise for patterns using zero\_extend and lshiftrt. If the second operand of mult is also a shift, then that is extended also. This transformation is only applied when it can be proven that the original operation had sufficient precision to prevent overflow.

Further canonicalization rules are defined in the function commutative\_operand\_precedence in 'gcc/rtlanal.c'.

## 17.15 Defining RTL Sequences for Code Generation

On some target machines, some standard pattern names for RTL generation cannot be handled with single insn, but a sequence of RTL insns can represent them. For these target machines, you can write a define\_expand to specify how to generate the sequence of RTL.

A define\_expand is an RTL expression that looks almost like a define\_insn; but, unlike the latter, a define\_expand is used only for RTL generation and it can produce more than one RTL insn.

A define\_expand RTX has four operands:

- The name. Each define\_expand must have a name, since the only use for it is to refer to it by name.
- The RTL template. This is a vector of RTL expressions representing a sequence of separate instructions. Unlike define\_insn, there is no implicit surrounding PARALLEL.
- The condition, a string containing a C expression. This expression is used to express how the availability of this pattern depends on subclasses of target machine, selected by command-line options when GCC is run. This is just like the condition of a define\_insn that has a standard name. Therefore, the condition (if present) may not depend on the data in the insn being matched, but only the target-machine-type flags. The compiler needs to test these conditions during initialization in order to learn exactly which named instructions are available in a particular run.
- The preparation statements, a string containing zero or more C statements which are to be executed before RTL code is generated from the RTL template.
  - Usually these statements prepare temporary registers for use as internal operands in the RTL template, but they can also generate RTL insns directly by calling routines such as emit\_insn, etc. Any such insns precede the ones that come from the RTL template.
- Optionally, a vector containing the values of attributes. See Section 17.19 [Insn Attributes], page 450.

Every RTL inso emitted by a define\_expand must match some define\_inso in the machine description. Otherwise, the compiler will crash when trying to generate code for the inso or trying to optimize it.

The RTL template, in addition to controlling generation of RTL insns, also describes the operands that need to be specified when this pattern is used. In particular, it gives a predicate for each operand.

A true operand, which needs to be specified in order to generate RTL from the pattern, should be described with a match\_operand in its first occurrence in the RTL template. This enters information on the operand's predicate into the tables that record such things. GCC uses the information to preload the operand into a register if that is required for valid RTL code. If the operand is referred to more than once, subsequent references should use match\_dup.

The RTL template may also refer to internal "operands" which are temporary registers or labels used only within the sequence made by the define\_expand. Internal operands are substituted into the RTL template with match\_dup, never with match\_operand. The values of the internal operands are not passed in as arguments by the compiler when it requests

use of this pattern. Instead, they are computed within the pattern, in the preparation statements. These statements compute the values and store them into the appropriate elements of operands so that match\_dup can find them.

There are two special macros defined for use in the preparation statements: DONE and FAIL. Use them with a following semicolon, as a statement.

DONE Use the DONE macro to end RTL generation for the pattern. The only RTL insns resulting from the pattern on this occasion will be those already emitted by explicit calls to emit\_insn within the preparation statements; the RTL template will not be generated.

FAIL Make the pattern fail on this occasion. When a pattern fails, it means that the pattern was not truly available. The calling routines in the compiler will try other strategies for code generation using other patterns.

Failure is currently supported only for binary (addition, multiplication, shifting, etc.) and bit-field (extv, extzv, and insv) operations.

If the preparation falls through (invokes neither DONE nor FAIL), then the define\_expand acts like a define\_insn in that the RTL template is used to generate the insn.

The RTL template is not used for matching, only for generating the initial insn list. If the preparation statement always invokes DONE or FAIL, the RTL template may be reduced to a simple list of operands, such as this example:

```
(define_expand "addsi3"
      [(match_operand:SI 0 "register_operand" "")
       (match_operand:SI 1 "register_operand" "")
       (match_operand:SI 2 "register_operand" "")]
    {
     handle_add (operands[0], operands[1], operands[2]);
     DONE;
    }")
Here is an example, the definition of left-shift for the SPUR chip:
    (define_expand "ashlsi3"
      [(set (match_operand:SI 0 "register_operand" "")
            (ashift:SI
              (match_operand:SI 1 "register_operand" "")
              (match_operand:SI 2 "nonmemory_operand" "")))]
      if (GET_CODE (operands[2]) != CONST_INT
          || (unsigned) INTVAL (operands[2]) > 3)
        FAIL:
    ጉ")
```

This example uses define\_expand so that it can generate an RTL insn for shifting when the shift-count is in the supported range of 0 to 3 but fail in other cases where machine insns aren't available. When it fails, the compiler tries another strategy using different patterns (such as, a library call).

If the compiler were able to handle nontrivial condition-strings in patterns with names, then it would be possible to use a define\_insn in that case. Here is another case (zero-extension on the 68000) which makes more use of the power of define\_expand:

Here two RTL insns are generated, one to clear the entire output operand and the other to copy the input operand into its low half. This sequence is incorrect if the input operand refers to [the old value of] the output operand, so the preparation statement makes sure this isn't so. The function make\_safe\_from copies the operands[1] into a temporary register if it refers to operands[0]. It does this by emitting another RTL insn.

Finally, a third example shows the use of an internal operand. Zero-extension on the SPUR chip is done by and-ing the result against a halfword mask. But this mask cannot be represented by a const\_int because the constant value is too large to be legitimate on this machine. So it must be copied into a register with force\_reg and then the register used in the and.

Note: If the define\_expand is used to serve a standard binary or unary arithmetic operation or a bit-field operation, then the last insn it generates must not be a code\_label, barrier or note. It must be an insn, jump\_insn or call\_insn. If you don't need a real insn at the end, emit an insn to copy the result of the operation into itself. Such an insn will generate no code, but it can avoid problems in the compiler.

# 17.16 Defining How to Split Instructions

There are two cases where you should specify how to split a pattern into multiple insns. On machines that have instructions requiring delay slots (see Section 17.19.8 [Delay Slots], page 459) or that have instructions whose output is not available for multiple cycles (see Section 17.19.9 [Processor pipeline description], page 460), the compiler phases that optimize these cases need to be able to move insns into one-instruction delay slots. However, some insns may generate more than one machine instruction. These insns cannot be placed into a delay slot.

Often you can rewrite the single insn as a list of individual insns, each corresponding to one machine instruction. The disadvantage of doing so is that it will cause the compilation to be slower and require more space. If the resulting insns are too complex, it may also suppress some optimizations. The compiler splits the insn if there is a reason to believe that it might improve instruction or delay slot scheduling.

The insn combiner phase also splits putative insns. If three insns are merged into one insn with a complex expression that cannot be matched by some define\_insn pattern, the combiner phase attempts to split the complex pattern into two insns that are recognized. Usually it can break the complex pattern into two patterns by splitting out some subexpression. However, in some other cases, such as performing an addition of a large constant in two insns on a RISC machine, the way to split the addition into two insns is machine-dependent.

The define\_split definition tells the compiler how to split a complex insn into several simpler insns. It looks like this:

```
(define_split
  [insn-pattern]
  "condition"
  [new-insn-pattern-1
   new-insn-pattern-2
   ...]
  "preparation-statements")
```

insn-pattern is a pattern that needs to be split and condition is the final condition to be tested, as in a define\_insn. When an insn matching insn-pattern and satisfying condition is found, it is replaced in the insn list with the insns given by new-insn-pattern-1, new-insn-pattern-2, etc.

The preparation-statements are similar to those statements that are specified for define\_expand (see Section 17.15 [Expander Definitions], page 438) and are executed before the new RTL is generated to prepare for the generated code or emit some insns whose pattern is not fixed. Unlike those in define\_expand, however, these statements must not generate any new pseudo-registers. Once reload has completed, they also must not allocate any space in the stack frame.

There are two special macros defined for use in the preparation statements: DONE and FAIL. Use them with a following semicolon, as a statement.

DONE Use the DONE macro to end RTL generation for the splitter. The only RTL insns generated as replacement for the matched input insn will be those already emitted by explicit calls to <code>emit\_insn</code> within the preparation statements; the replacement pattern is not used.

FAIL Make the define\_split fail on this occasion. When a define\_split fails, it means that the splitter was not truly available for the inputs it was given, and the input insn will not be split.

If the preparation falls through (invokes neither DONE nor FAIL), then the define\_split uses the replacement template.

Patterns are matched against *insn-pattern* in two different circumstances. If an insn needs to be split for delay slot scheduling or insn scheduling, the insn is already known to be valid, which means that it must have been matched by some define\_insn and, if reload\_completed is nonzero, is known to satisfy the constraints of that define\_insn. In that case, the new insn patterns must also be insns that are matched by some define\_insn and, if reload\_completed is nonzero, must also satisfy the constraints of those definitions.

As an example of this usage of define\_split, consider the following example from 'a29k.md', which splits a sign\_extend from HImode to SImode into a pair of shift insns:

When the combiner phase tries to split an insn pattern, it is always the case that the pattern is *not* matched by any define\_insn. The combiner pass first tries to split a single set expression and then the same set expression inside a parallel, but followed by a clobber of a pseudo-reg to use as a scratch register. In these cases, the combiner expects exactly one or two new insn patterns to be generated. It will verify that these patterns match some define\_insn definitions, so you need not do this test in the define\_split (of course, there is no point in writing a define\_split that will never produce insns that match).

Here is an example of this use of define\_split, taken from 'rs6000.md':

Here the predicate non\_add\_cint\_operand matches any const\_int that is *not* a valid operand of a single add insn. The add with the smaller displacement is written so that it can be substituted into the address of a subsequent operation.

An example that uses a scratch register, from the same file, generates an equality comparison of a register and a large constant:

```
(set (match_dup 0) (compare:CC (match_dup 3) (match_dup 5)))]

(
/* Get the constant we are comparing against, C, and see what it looks like sign-extended to 16 bits. Then see what constant could be XOR'ed with C to get the sign-extended value. */

int c = INTVAL (operands[2]);
int sextc = (c << 16) >> 16;
int xorv = c ^ sextc;

operands[4] = GEN_INT (xorv);
operands[5] = GEN_INT (sextc);
}")
```

To avoid confusion, don't write a single define\_split that accepts some insns that match some define\_insn as well as some insns that don't. Instead, write two separate define\_split definitions, one for the insns that are valid and one for the insns that are not valid.

The splitter is allowed to split jump instructions into sequence of jumps or create new jumps in while splitting non-jump instructions. As the control flow graph and branch prediction information needs to be updated, several restriction apply.

Splitting of jump instruction into sequence that over by another jump instruction is always valid, as compiler expect identical behavior of new jump. When new sequence contains multiple jump instructions or new labels, more assistance is needed. Splitter is required to create only unconditional jumps, or simple conditional jump instructions. Additionally it must attach a REG\_BR\_PROB note to each conditional jump. A global variable split\_branch\_probability holds the probability of the original branch in case it was a simple conditional jump, -1 otherwise. To simplify recomputing of edge frequencies, the new sequence is required to have only forward jumps to the newly created labels.

For the common case where the pattern of a define\_split exactly matches the pattern of a define\_insn, use define\_insn\_and\_split. It looks like this:

```
(define_insn_and_split
  [insn-pattern]
  "condition"
  "output-template"
  "split-condition"
  [new-insn-pattern-1
   new-insn-pattern-2
   ...]
  "preparation-statements"
  [insn-attributes])
```

insn-pattern, condition, output-template, and insn-attributes are used as in define\_insn. The new-insn-pattern vector and the preparation-statements are used as in a define\_split. The split-condition is also used as in define\_split, with the additional behavior that if the condition starts with '&&', the condition used for the split will be the constructed as a logical "and" of the split condition with the insn condition. For example, from i386.md:

```
(define_insn_and_split "zero_extendhisi2_and"
  [(set (match_operand:SI 0 "register_operand" "=r")
        (zero_extend:SI (match_operand:HI 1 "register_operand" "0")))
  (clobber (reg:CC 17))]
```

In this case, the actual split condition will be 'TARGET\_ZERO\_EXTEND\_WITH\_AND && !optimize\_size && reload\_completed'.

The define\_insn\_and\_split construction provides exactly the same functionality as two separate define\_insn and define\_split patterns. It exists for compactness, and as a maintenance tool to prevent having to ensure the two patterns' templates match.

It is sometimes useful to have a define\_insn\_and\_split that replaces specific operands of an instruction but leaves the rest of the instruction pattern unchanged. You can do this directly with a define\_insn\_and\_split, but it requires a new-insn-pattern-1 that repeats most of the original insn-pattern. There is also the complication that an implicit parallel in insn-pattern must become an explicit parallel in new-insn-pattern-1, which is easy to overlook. A simpler alternative is to use define\_insn\_and\_rewrite, which is a form of define\_insn\_and\_split that automatically generates new-insn-pattern-1 by replacing each match\_operand in insn-pattern with a corresponding match\_dup, and each match\_operator in the pattern with a corresponding match\_op\_dup. The arguments are otherwise identical to define\_insn\_and\_split:

```
(define_insn_and_rewrite
  [insn-pattern]
  "condition"
  "output-template"
  "split-condition"
  "preparation-statements"
  [insn-attributes])
```

The match\_dups and match\_op\_dups in the new instruction pattern use any new operand values that the *preparation-statements* store in the operands array, as for a normal define\_insn\_and\_split. *preparation-statements* can also emit additional instructions before the new instruction. They can even emit an entirely different sequence of instructions and use DONE to avoid emitting a new form of the original instruction.

The split in a define\_insn\_and\_rewrite is only intended to apply to existing instructions that match insn-pattern. split-condition must therefore start with &&, so that the split condition applies on top of condition.

Here is an example from the AArch64 SVE port, in which operand 1 is known to be equivalent to an all-true constant and isn't used by the output template:

The splitter in this case simply replaces operand 1 with the constant value that it is known to have. The equivalent define\_insn\_and\_split would be:

```
(define_insn_and_split "*while_ult<GPI:mode><PRED_ALL:mode>_cc"
  [(set (reg:CC CC_REGNUM)
        (compare:CC
          (unspec:SI [(match_operand:PRED_ALL 1)
                      (unspec:PRED_ALL
                        [(match_operand:GPI 2 "aarch64_reg_or_zero" "rZ")
                         (match_operand:GPI 3 "aarch64_reg_or_zero" "rZ")]
                        UNSPEC_WHILE_LO)]
                     UNSPEC_PTEST_PTRUE)
          (const_int 0)))
   (set (match_operand:PRED_ALL 0 "register_operand" "=Upa")
        (unspec:PRED_ALL [(match_dup 2)
                          (match_dup 3)]
                         UNSPEC_WHILE_LO))]
  "TARGET_SVE"
  "whilelo\t%0.<PRED_ALL:Vetype>, %<w>2, %<w>3"
  ;; Force the compiler to drop the unused predicate operand, so that we
  ;; don't have an unnecessary PTRUE.
  "&& !CONSTANT_P (operands[1])"
  [(parallel
     [(set (reg:CC CC_REGNUM)
           (compare:CC
             (unspec:SI [(match_dup 1)
                         (unspec:PRED_ALL [(match_dup 2)
                                            (match_dup 3)]
                                           UNSPEC_WHILE_LO)]
                        UNSPEC_PTEST_PTRUE)
             (const_int 0)))
      (set (match_dup 0)
           (unspec:PRED_ALL [(match_dup 2)
                              (match_dup 3)]
                            UNSPEC_WHILE_LO))])]
    operands[1] = CONSTM1_RTX (<MODE>mode);
)
```

## 17.17 Including Patterns in Machine Descriptions.

The include pattern tells the compiler tools where to look for patterns that are in files other than in the file '.md'. This is used only at build time and there is no preprocessing allowed.

```
It looks like:
    (include
        pathname)
For example:
    (include "filestuff")
```

Where pathname is a string that specifies the location of the file, specifies the include file to be in 'gcc/config/target/filestuff'. The directory 'gcc/config/target' is regarded as the default directory.

Machine descriptions may be split up into smaller more manageable subsections and placed into subdirectories.

By specifying:

```
(include "BOGUS/filestuff")
```

the include file is specified to be in 'gcc/config/target/BOGUS/filestuff'.

Specifying an absolute path for the include file such as;

```
(include "/u2/B0GUS/filestuff")
```

is permitted but is not encouraged.

## 17.17.1 RTL Generation Tool Options for Directory Search

The '-Idir' option specifies directories to search for machine descriptions. For example:

```
genrecog -I/p1/abc/proc1 -I/p2/abcd/pro2 target.md
```

Add the directory dir to the head of the list of directories to be searched for header files. This can be used to override a system machine definition file, substituting your own version, since these directories are searched before the default machine description file directories. If you use more than one '-I' option, the directories are scanned in left-to-right order; the standard default directory come after.

# 17.18 Machine-Specific Peephole Optimizers

In addition to instruction patterns the 'md' file may contain definitions of machine-specific peephole optimizations.

The combiner does not notice certain peephole optimizations when the data flow in the program does not suggest that it should try them. For example, sometimes two consecutive insns related in purpose can be combined even though the second one does not appear to

use a register computed in the first one. A machine-specific peephole optimizer can detect such opportunities.

There are two forms of peephole definitions that may be used. The original define\_peephole is run at assembly output time to match insns and substitute assembly text. Use of define\_peephole is deprecated.

A newer define\_peephole2 matches insns and substitutes new insns. The peephole2 pass is run after register allocation but before scheduling, which may result in much better code for targets that do scheduling.

## 17.18.1 RTL to Text Peephole Optimizers

A definition looks like this:

```
(define_peephole
  [insn-pattern-1
   insn-pattern-2
   ...]
  "condition"
  "template"
  "optional-insn-attributes")
```

The last string operand may be omitted if you are not using any machine-specific information in this machine description. If present, it must obey the same rules as in a define\_insn.

In this skeleton, *insn-pattern-1* and so on are patterns to match consecutive insns. The optimization applies to a sequence of insns when *insn-pattern-1* matches the first one, *insn-pattern-2* matches the next, and so on.

Each of the insns matched by a peephole must also match a define\_insn. Peepholes are checked only at the last stage just before code generation, and only optionally. Therefore, any insn which would match a peephole but no define\_insn will cause a crash in code generation in an unoptimized compilation, or at various optimization stages.

The operands of the insns are matched with match\_operands, match\_operator, and match\_dup, as usual. What is not usual is that the operand numbers apply to all the insn patterns in the definition. So, you can check for identical operands in two insns by using match\_operand in one insn and match\_dup in the other.

The operand constraints used in match\_operand patterns do not have any direct effect on the applicability of the peephole, but they will be validated afterward, so make sure your constraints are general enough to apply whenever the peephole matches. If the peephole matches but the constraints are not satisfied, the compiler will crash.

It is safe to omit constraints in all the operands of the peephole; or you can write constraints which serve as a double-check on the criteria previously tested.

Once a sequence of insns matches the patterns, the *condition* is checked. This is a C expression which makes the final decision whether to perform the optimization (we do so if the expression is nonzero). If *condition* is omitted (in other words, the string is empty) then the optimization is applied to every sequence of insns that matches the patterns.

The defined peephole optimizations are applied after register allocation is complete. Therefore, the peephole definition can check which operands have ended up in which kinds of registers, just by looking at the operands.

The way to refer to the operands in *condition* is to write operands [i] for operand number i (as matched by (match\_operand i ...)). Use the variable insn to refer to the last of the insns being matched; use prev\_active\_insn to find the preceding insns.

When optimizing computations with intermediate results, you can use *condition* to match only when the intermediate results are not used elsewhere. Use the C expression dead\_or\_set\_p (insn, op), where insn is the insn in which you expect the value to be used for the last time (from the value of insn, together with use of prev\_nonnote\_insn), and op is the intermediate value (from operands[i]).

Applying the optimization means replacing the sequence of insns with one new insn. The template controls ultimate output of assembler code for this combined insn. It works exactly like the template of a define\_insn. Operand numbers in this template are the same ones used in matching the original sequence of insns.

The result of a defined peephole optimizer does not need to match any of the insn patterns in the machine description; it does not even have an opportunity to match them. The peephole optimizer definition itself serves as the insn pattern to control how the insn is output.

Defined peephole optimizers are run as assembler code is being output, so the insns they produce are never combined or rearranged in any way.

Here is an example, taken from the 68000 machine description:

```
(define_peephole
        [(set (reg:SI 15) (plus:SI (reg:SI 15) (const_int 4)))
         (set (match_operand:DF 0 "register_operand" "=f")
              (match_operand:DF 1 "register_operand" "ad"))]
        "FP_REG_P (operands[0]) && ! FP_REG_P (operands[1])"
      {
        rtx xoperands[2];
        xoperands[1] = gen_rtx_REG (SImode, REGNO (operands[1]) + 1);
      #ifdef MOTOROLA
        output_asm_insn ("move.l %1,(sp)", xoperands);
        output_asm_insn ("move.l %1,-(sp)", operands);
        return "fmove.d (sp)+,%0";
        output_asm_insn ("movel %1,sp@", xoperands);
        output_asm_insn ("movel %1,sp@-", operands);
        return "fmoved sp@+,%0";
  The effect of this optimization is to change
      jbsr _foobar
      addql #4,sp
      movel d1,sp@-
      movel d0,sp@-
      fmoved sp@+,fp0
into
      jbsr _foobar
      movel d1,sp@
      movel d0,sp@-
      fmoved sp@+,fp0
```

insn-pattern-1 and so on look almost like the second operand of define\_insn. There is one important difference: the second operand of define\_insn consists of one or more

RTX's enclosed in square brackets. Usually, there is only one: then the same action can be written as an element of a define\_peephole. But when there are multiple actions in a define\_insn, they are implicitly enclosed in a parallel. Then you must explicitly write the parallel, and the square brackets within it, in the define\_peephole. Thus, if an insn pattern looks like this,

```
(define_insn "divmodsi4"
        [(set (match_operand:SI 0 "general_operand" "=d")
              (div:SI (match_operand:SI 1 "general_operand" "0")
                      (match_operand:SI 2 "general_operand" "dmsK")))
         (set (match_operand:SI 3 "general_operand" "=d")
              (mod:SI (match_dup 1) (match_dup 2)))]
        "TARGET 68020"
        "divsl%.1 %2,%3:%0")
then the way to mention this insn in a peephole is as follows:
      (define_peephole
        [...
         (parallel
          [(set (match_operand:SI 0 "general_operand" "=d")
                (div:SI (match_operand:SI 1 "general_operand" "0")
                        (match_operand:SI 2 "general_operand" "dmsK")))
           (set (match_operand:SI 3 "general_operand" "=d")
                (mod:SI (match_dup 1) (match_dup 2)))])
         ...]
        ...)
```

## 17.18.2 RTL to RTL Peephole Optimizers

The define\_peephole2 definition tells the compiler how to substitute one sequence of instructions for another sequence, what additional scratch registers may be needed and what their lifetimes must be.

```
(define_peephole2
  [insn-pattern-1
   insn-pattern-2
  ...]
  "condition"
  [new-insn-pattern-1
   new-insn-pattern-2
  ...]
  "preparation-statements")
```

The definition is almost identical to define\_split (see Section 17.16 [Insn Splitting], page 440) except that the pattern to match is not a single instruction, but a sequence of instructions.

It is possible to request additional scratch registers for use in the output template. If appropriate registers are not free, the pattern will simply not match.

Scratch registers are requested with a match\_scratch pattern at the top level of the input pattern. The allocated register (initially) will be dead at the point requested within the original sequence. If the scratch is used at more than a single point, a match\_dup pattern at the top level of the input pattern marks the last position in the input sequence at which the register must be available.

Here is an example from the IA-32 machine description:

```
(define_peephole2
  [(match_scratch:SI 2 "r")
```

This pattern tries to split a load from its use in the hopes that we'll be able to schedule around the memory load latency. It allocates a single SImode register of class GENERAL\_REGS ("r") that needs to be live only at the point just before the arithmetic.

A real example requiring extended scratch lifetimes is harder to come by, so here's a silly made-up example:

```
(define_peephole2
  [(match_scratch:SI 4 "r")
  (set (match_operand:SI 0 "" "") (match_operand:SI 1 "" ""))
  (set (match_operand:SI 2 "" "") (match_dup 1))
  (match_dup 4)
  (set (match_operand:SI 3 "" "") (match_dup 1))]
  "/* determine 1 does not overlap 0 and 2 */"
  [(set (match_dup 4) (match_dup 1))
  (set (match_dup 0) (match_dup 4))
  (set (match_dup 2) (match_dup 4))
  (set (match_dup 3) (match_dup 4))]
  "")
```

There are two special macros defined for use in the preparation statements: DONE and FAIL. Use them with a following semicolon, as a statement.

DONE Use the DONE macro to end RTL generation for the peephole. The only RTL insns generated as replacement for the matched input insn will be those already emitted by explicit calls to emit\_insn within the preparation statements; the replacement pattern is not used.

Make the define\_peephole2 fail on this occasion. When a define\_peephole2 fails, it means that the replacement was not truly available for the particular inputs it was given. In that case, GCC may still apply a later define\_peephole2 that also matches the given insn pattern. (Note that this is different from define\_split, where FAIL prevents the input insn from being split at all.)

If the preparation falls through (invokes neither DONE nor FAIL), then the define\_peephole2 uses the replacement template.

If we had not added the (match\_dup 4) in the middle of the input sequence, it might have been the case that the register we chose at the beginning of the sequence is killed by the first or second set.

## 17.19 Instruction Attributes

In addition to describing the instruction supported by the target machine, the 'md' file also defines a group of attributes and a set of values for each. Every generated insn is assigned

a value for each attribute. One possible attribute would be the effect that the insn has on the machine's condition code. This attribute can then be used by NOTICE\_UPDATE\_CC to track the condition codes.

## 17.19.1 Defining Attributes and their Values

The define\_attr expression is used to define each attribute required by the target machine. It looks like:

```
(define_attr name list-of-values default)
```

name is a string specifying the name of the attribute being defined. Some attributes are used in a special way by the rest of the compiler. The enabled attribute can be used to conditionally enable or disable insn alternatives (see Section 17.8.6 [Disable Insn Alternatives], page 387). The predicable attribute, together with a suitable define\_cond\_exec (see Section 17.20 [Conditional Execution], page 466), can be used to automatically generate conditional variants of instruction patterns. The mnemonic attribute can be used to check for the instruction mnemonic (see Section 17.19.7 [Mnemonic Attribute], page 459). The compiler internally uses the names ce\_enabled and nonce\_enabled, so they should not be used elsewhere as alternative names.

list-of-values is either a string that specifies a comma-separated list of values that can be assigned to the attribute, or a null string to indicate that the attribute takes numeric values.

default is an attribute expression that gives the value of this attribute for insns that match patterns whose definition does not include an explicit value for this attribute. See Section 17.19.4 [Attr Example], page 456, for more information on the handling of defaults. See Section 17.19.6 [Constant Attributes], page 458, for information on attributes that do not depend on any particular insn.

For each defined attribute, a number of definitions are written to the 'insn-attr.h' file. For cases where an explicit set of values is specified for an attribute, the following are defined:

- A '#define' is written for the symbol 'HAVE\_ATTR\_name'.
- An enumerated class is defined for 'attr\_name' with elements of the form 'upper-name\_upper-value' where the attribute name and value are first converted to upper-case.
- A function 'get\_attr\_name' is defined that is passed an insn and returns the attribute value for that insn.

If the attribute takes numeric values, no enum type will be defined and the function to obtain the attribute's value will return int.

There are attributes which are tied to a specific meaning. These attributes are not free to use for other purposes:

length The length attribute is used to calculate the length of emitted code chunks. This is especially important when verifying branch distances. See Section 17.19.5 [Insn Lengths], page 457.

enabled The enabled attribute can be defined to prevent certain alternatives of an insn definition from being used during code generation. See Section 17.8.6 [Disable Insn Alternatives], page 387.

mnemonic The mnemonic attribute can be defined to implement instruction specific checks in e.g. the pipeline description. See Section 17.19.7 [Mnemonic Attribute], page 459.

For each of these special attributes, the corresponding 'HAVE\_ATTR\_name' '#define' is also written when the attribute is not defined; in that case, it is defined as '0'.

Another way of defining an attribute is to use:

```
(define_enum_attr "attr" "enum" default)
```

This works in just the same way as define\_attr, except that the list of values is taken from a separate enumeration called *enum* (see [define\_enum], page 472). This form allows you to use the same list of values for several attributes without having to repeat the list each time. For example:

```
(define_enum "processor" [
    model_a
    model_b
    ...
])
(define_enum_attr "arch" "processor"
    (const (symbol_ref "target_arch")))
(define_enum_attr "tune" "processor"
    (const (symbol_ref "target_tune")))

defines the same attributes as:
    (define_attr "arch" "model_a,model_b,..."
        (const (symbol_ref "target_arch")))
    (define_attr "tune" "model_a,model_b,..."
        (const (symbol_ref "target_tune")))
```

but without duplicating the processor list. The second example defines two separate C enums (attr\_arch and attr\_tune) whereas the first defines a single C enum (processor).

## 17.19.2 Attribute Expressions

RTL expressions used to define attributes use the codes described above plus a few specific to attribute definitions, to be discussed below. Attribute value expressions must have one of the following forms:

```
(const_int i)
```

The integer i specifies the value of a numeric attribute. i must be non-negative.

The value of a numeric attribute can be specified either with a const\_int, or as an integer represented as a string in const\_string, eq\_attr (see below), attr, symbol\_ref, simple arithmetic expressions, and set\_attr overrides on specific instructions (see Section 17.19.3 [Tagging Insns], page 455).

#### (const\_string value)

The string value specifies a constant attribute value. If value is specified as ""\*", it means that the default value of the attribute is to be used for the insn containing this expression. ""\*" obviously cannot be used in the default expression of a define\_attr.

If the attribute whose value is being specified is numeric, value must be a string containing a non-negative integer (normally const\_int would be used in this case). Otherwise, it must contain one of the valid values for the attribute.

#### (if\_then\_else test true-value false-value)

test specifies an attribute test, whose format is defined below. The value of this expression is true-value if test is true, otherwise it is false-value.

#### (cond [test1 value1 ...] default)

The first operand of this expression is a vector containing an even number of expressions and consisting of pairs of test and value expressions. The value of the cond expression is that of the value corresponding to the first true test expression. If none of the test expressions are true, the value of the cond expression is that of the default expression.

test expressions can have one of the following forms:

#### (const\_int i)

This test is true if i is nonzero and false otherwise.

(not test)

(ior test1 test2)

(and test1 test2)

These tests are true if the indicated logical function is true.

#### (match\_operand:m n pred constraints)

This test is true if operand n of the insn whose attribute value is being determined has mode m (this part of the test is ignored if m is VOIDmode) and the function specified by the string pred returns a nonzero value when passed operand n and mode m (this part of the test is ignored if pred is the null string).

The constraints operand is ignored and should be the null string.

#### (match\_test c-expr)

The test is true if C expression c-expr is true. In non-constant attributes, c-expr has access to the following variables:

insn The rtl instruction under test.

which\_alternative

The define\_insn alternative that insn matches. See Section 17.6 [Output Statement], page 344.

operands An array of insn's rtl operands.

*c*-expr behaves like the condition in a C if statement, so there is no need to explicitly convert the expression into a boolean 0 or 1 value. For example, the following two tests are equivalent:

These tests are true if the indicated comparison of the two arithmetic expressions is true. Arithmetic expressions are formed with plus, minus, mult, div, mod, abs, neg, and, ior, xor, not, ashift, lshiftrt, and ashiftrt expressions.

const\_int and symbol\_ref are always valid terms (see Section 17.19.5 [Insn Lengths], page 457, for additional forms). symbol\_ref is a string denoting a C expression that yields an int when evaluated by the 'get\_attr\_...' routine. It should normally be a global variable.

#### (eq\_attr name value)

name is a string specifying the name of an attribute.

value is a string that is either a valid value for attribute name, a commaseparated list of values, or '!' followed by a value or list. If value does not begin with a '!', this test is true if the value of the name attribute of the current insn is in the list specified by value. If value begins with a '!', this test is true if the attribute's value is not in the specified list.

For example,

```
(eq_attr "type" "load,store")
is equivalent to
   (ior (eq_attr "type" "load") (eq_attr "type" "store"))
```

If name specifies an attribute of 'alternative', it refers to the value of the compiler variable which\_alternative (see Section 17.6 [Output Statement], page 344) and the values must be small integers. For example,

Note that, for most attributes, an eq\_attr test is simplified in cases where the value of the attribute being tested is known for all insns matching a particular pattern. This is by far the most common case.

#### (attr\_flag name)

The value of an attr\_flag expression is true if the flag specified by name is true for the insn currently being scheduled.

name is a string specifying one of a fixed set of flags to test. Test the flags forward and backward to determine the direction of a conditional branch.

This example describes a conditional branch delay slot which can be nullified for forward branches that are taken (annul-true) or for backward branches which are not taken (annul-false).

```
(define_delay (eq_attr "type" "cbranch")
  [(eq_attr "in_branch_delay" "true")
        (and (eq_attr "in_branch_delay" "true")
              (attr_flag "forward"))
        (and (eq_attr "in_branch_delay" "true")
              (attr_flag "backward"))])
```

The forward and backward flags are false if the current insn being scheduled is not a conditional branch.

attr\_flag is only used during delay slot scheduling and has no meaning to other passes of the compiler.

(attr name)

The value of another attribute is returned. This is most useful for numeric attributes, as eq\_attr and attr\_flag produce more efficient code for non-numeric attributes.

#### 17.19.3 Assigning Attribute Values to Insns

The value assigned to an attribute of an insn is primarily determined by which pattern is matched by that insn (or which define\_peephole generated it). Every define\_insn and define\_peephole can have an optional last argument to specify the values of attributes for matching insns. The value of any attribute not specified in a particular insn is set to the default value for that attribute, as specified in its define\_attr. Extensive use of default values for attributes permits the specification of the values for only one or two attributes in the definition of most insn patterns, as seen in the example in the next section.

The optional last argument of define\_insn and define\_peephole is a vector of expressions, each of which defines the value for a single attribute. The most general way of assigning an attribute's value is to use a set expression whose first operand is an attrexpression giving the name of the attribute being set. The second operand of the set is an attribute expression (see Section 17.19.2 [Expressions], page 452) giving the value of the attribute.

When the attribute value depends on the 'alternative' attribute (i.e., which is the applicable alternative in the constraint of the insn), the set\_attr\_alternative expression can be used. It allows the specification of a vector of attribute expressions, one for each alternative.

When the generality of arbitrary attribute expressions is not required, the simpler set\_attr expression can be used, which allows specifying a string giving either a single attribute value or a list of attribute values, one for each alternative.

The form of each of the above specifications is shown below. In each case, *name* is a string specifying the attribute to be set.

#### (set\_attr name value-string)

value-string is either a string giving the desired attribute value, or a string containing a comma-separated list giving the values for succeeding alternatives.

The number of elements must match the number of alternatives in the constraint of the insn pattern.

Note that it may be useful to specify '\*' for some alternative, in which case the attribute will assume its default value for insns matching that alternative.

```
(set_attr_alternative name [value1 value2 ...])
```

Depending on the alternative of the insn, the value will be one of the specified values. This is a shorthand for using a cond with tests on the 'alternative' attribute.

```
(set (attr name) value)
```

The first operand of this **set** must be the special RTL expression **attr**, whose sole operand is a string giving the name of the attribute being set. *value* is the value of the attribute.

The following shows three different ways of representing the same attribute value specification:

The define\_asm\_attributes expression provides a mechanism to specify the attributes assigned to insns produced from an asm statement. It has the form:

```
(define_asm_attributes [attr-sets])
```

where attr-sets is specified the same as for both the define\_insn and the define\_peephole expressions.

These values will typically be the "worst case" attribute values. For example, they might indicate that the condition code will be clobbered.

A specification for a length attribute is handled specially. The way to compute the length of an asm insn is to multiply the length specified in the expression define\_asm\_attributes by the number of machine instructions specified in the asm statement, determined by counting the number of semicolons and newlines in the string. Therefore, the value of the length attribute specified in a define\_asm\_attributes should be the maximum possible length of a single machine instruction.

#### 17.19.4 Example of Attribute Specifications

The judicious use of defaulting is important in the efficient use of insn attributes. Typically, insns are divided into *types* and an attribute, customarily called **type**, is used to represent this value. This attribute is normally used only to define the default value for other attributes. An example will clarify this usage.

Assume we have a RISC machine with a condition code and in which only full-word operations are performed in registers. Let us assume that we can divide all insns into loads, stores, (integer) arithmetic operations, floating point operations, and branches.

Here we will concern ourselves with determining the effect of an insn on the condition code and will limit ourselves to the following possible effects: The condition code can be set unpredictably (clobbered), not be changed, be set to agree with the results of the operation, or only changed if the item previously set into the condition code has been modified.

Here is part of a sample 'md' file for such a machine:

```
(define_attr "type" "load,store,arith,fp,branch" (const_string "arith"))
(define_attr "cc" "clobber,unchanged,set,change0"
             (cond [(eq_attr "type" "load")
                        (const_string "change0")
                    (eq_attr "type" "store,branch")
                        (const_string "unchanged")
                    (eq_attr "type" "arith")
                        (if_then_else (match_operand:SI 0 "" "")
                                       (const_string "set")
                                       (const_string "clobber"))]
                   (const_string "clobber")))
(define_insn ""
  [(set (match_operand:SI 0 "general_operand" "=r,r,m")
        (match_operand:SI 1 "general_operand" "r,m,r"))]
  "@
  move %0,%1
  load %0.%1
  store %0,%1"
  [(set_attr "type" "arith,load,store")])
```

Note that we assume in the above example that arithmetic operations performed on quantities smaller than a machine word clobber the condition code since they will set the condition code to a value corresponding to the full-word result.

## 17.19.5 Computing the Length of an Insn

For many machines, multiple types of branch instructions are provided, each for different length branch displacements. In most cases, the assembler will choose the correct instruction to use. However, when the assembler cannot do so, GCC can when a special attribute, the length attribute, is defined. This attribute must be defined to have numeric values by specifying a null string in its define\_attr.

In the case of the length attribute, two additional forms of arithmetic terms are allowed in test expressions:

#### (match\_dup n)

This refers to the address of operand n of the current insn, which must be a label\_ref.

(pc) For non-branch instructions and backward branch instructions, this refers to the address of the current insn. But for forward branch instructions, this refers to the address of the next insn, because the length of the current insn is to be computed.

For normal insns, the length will be determined by value of the length attribute. In the case of addr\_vec and addr\_diff\_vec insn patterns, the length is computed as the number of vectors multiplied by the size of each vector.

Lengths are measured in addressable storage units (bytes).

Note that it is possible to call functions via the symbol\_ref mechanism to compute the length of an insn. However, if you use this mechanism you must provide dummy clauses to express the maximum length without using the function call. You can an example of this in the pa machine description for the call\_symref pattern.

The following macros can be used to refine the length computation:

#### ADJUST\_INSN\_LENGTH (insn, length)

If defined, modifies the length assigned to instruction *insn* as a function of the context in which it is used. *length* is an lvalue that contains the initially computed length of the insn and should be updated with the correct length of the insn.

This macro will normally not be required. A case in which it is required is the ROMP. On this machine, the size of an addr\_vec insn must be increased by two to compensate for the fact that alignment may be required.

The routine that returns get\_attr\_length (the value of the length attribute) can be used by the output routine to determine the form of the branch instruction to be written, as the example below illustrates.

As an example of the specification of variable-length branches, consider the IBM 360. If we adopt the convention that a register will be set to the starting address of a function, we can jump to labels within 4k of the start using a four-byte instruction. Otherwise, we need a six-byte sequence to load the address from memory and then branch to it.

On such a machine, a pattern for a branch instruction might be specified as follows:

#### 17.19.6 Constant Attributes

A special form of define\_attr, where the expression for the default value is a const expression, indicates an attribute that is constant for a given run of the compiler. Constant attributes may be used to specify which variety of processor is used. For example,

```
(const_string "slow"))))
```

The routine generated for constant attributes has no parameters as it does not depend on any particular insn. RTL expressions used to define the value of a constant attribute may use the symbol\_ref form, but may not use either the match\_operand form or eq\_attr forms involving insn attributes.

#### 17.19.7 Mnemonic Attribute

The mnemonic attribute is a string type attribute holding the instruction mnemonic for an insn alternative. The attribute values will automatically be generated by the machine description parser if there is an attribute definition in the md file:

```
(define_attr "mnemonic" "unknown" (const_string "unknown"))
```

The default value can be freely chosen as long as it does not collide with any of the instruction mnemonics. This value will be used whenever the machine description parser is not able to determine the mnemonic string. This might be the case for output templates containing more than a single instruction as in "mvcle\t%0,%1,0\; jo\t.-4".

The mnemonic attribute set is not generated automatically if the instruction string is generated via C code.

An existing mnemonic attribute set in an insn definition will not be overriden by the md file parser. That way it is possible to manually set the instruction mnemonics for the cases where the md file parser fails to determine it automatically.

The mnemonic attribute is useful for dealing with instruction specific properties in the pipeline description without defining additional insn attributes.

## 17.19.8 Delay Slot Scheduling

The insn attribute mechanism can be used to specify the requirements for delay slots, if any, on a target machine. An instruction is said to require a *delay slot* if some instructions that are physically after the instruction are executed as if they were located before it. Classic examples are branch and call instructions, which often execute the following instruction before the branch or call is performed.

On some machines, conditional branch instructions can optionally *annul* instructions in the delay slot. This means that the instruction will not be executed for certain branch outcomes. Both instructions that annul if the branch is true and instructions that annul if the branch is false are supported.

Delay slot scheduling differs from instruction scheduling in that determining whether an instruction needs a delay slot is dependent only on the type of instruction being generated, not on data flow between the instructions. See the next section for a discussion of data-dependent instruction scheduling.

The requirement of an insn needing one or more delay slots is indicated via the define\_delay expression. It has the following form:

...])

test is an attribute test that indicates whether this define\_delay applies to a particular insn. If so, the number of required delay slots is determined by the length of the vector specified as the second argument. An insn placed in delay slot n must satisfy attribute test delay-n. annul-true-n is an attribute test that specifies which insns may be annulled if the branch is true. Similarly, annul-false-n specifies which insns in the delay slot may be annulled if the branch is false. If annulling is not supported for that delay slot, (nil) should be coded.

For example, in the common case where branch and call insns require a single delay slot, which may contain any insn other than a branch or call, the following would be placed in the 'md' file:

Multiple define\_delay expressions may be specified. In this case, each such expression specifies different delay slot requirements and there must be no insn for which tests in two define\_delay expressions are both true.

For example, if we have a machine that requires one delay slot for branches but two for calls, no delay slot can contain a branch or call insn, and any valid insn in the delay slot for the branch can be annulled if the branch is true, we might represent this as follows:

#### 17.19.9 Specifying processor pipeline description

To achieve better performance, most modern processors (super-pipelined, superscalar RISC, and VLIW processors) have many functional units on which several instructions can be executed simultaneously. An instruction starts execution if its issue conditions are satisfied. If not, the instruction is stalled until its conditions are satisfied. Such interlock (pipeline) delay causes interruption of the fetching of successor instructions (or demands nop instructions, e.g. for some MIPS processors).

There are two major kinds of interlock delays in modern processors. The first one is a data dependence delay determining instruction latency time. The instruction execution is not started until all source data have been evaluated by prior instructions (there are more complex cases when the instruction execution starts even when the data are not available but will be ready in given time after the instruction execution start). Taking the data dependence delays into account is simple. The data dependence (true, output, and anti-dependence) delay between two instructions is given by a constant. In most cases this approach is adequate. The second kind of interlock delays is a reservation delay. The reservation delay means that two instructions under execution will be in need of shared processors resources, i.e. buses, internal registers, and/or functional units, which are reserved for some time. Taking this kind of delay into account is complex especially for modern RISC processors.

The task of exploiting more processor parallelism is solved by an instruction scheduler. For a better solution to this problem, the instruction scheduler has to have an adequate description of the processor parallelism (or *pipeline description*). GCC machine descriptions describe processor parallelism and functional unit reservations for groups of instructions with the aid of *regular expressions*.

The GCC instruction scheduler uses a *pipeline hazard recognizer* to figure out the possibility of the instruction issue by the processor on a given simulated processor cycle. The pipeline hazard recognizer is automatically generated from the processor pipeline description. The pipeline hazard recognizer generated from the machine description is based on a deterministic finite state automaton (DFA): the instruction issue is possible if there is a transition from one automaton state to another one. This algorithm is very fast, and furthermore, its speed is not dependent on processor complexity<sup>2</sup>.

The rest of this section describes the directives that constitute an automaton-based processor pipeline description. The order of these constructions within the machine description file is not important.

The following optional construction describes names of automata generated and used for the pipeline hazards recognition. Sometimes the generated finite state automaton used by the pipeline hazard recognizer is large. If we use more than one automaton and bind functional units to the automata, the total size of the automata is usually less than the size of the single automaton. If there is no one such construction, only one finite state automaton is generated.

(define\_automaton automata-names)

automata-names is a string giving names of the automata. The names are separated by commas. All the automata should have unique names. The automaton name is used in the constructions define\_cpu\_unit and define\_query\_cpu\_unit.

Each processor functional unit used in the description of instruction reservations should be described by the following construction.

(define\_cpu\_unit unit-names [automaton-name])

*unit-names* is a string giving the names of the functional units separated by commas. Don't use name 'nothing', it is reserved for other goals.

automaton-name is a string giving the name of the automaton with which the unit is bound. The automaton should be described in construction define\_automaton. You should give automaton-name, if there is a defined automaton.

The assignment of units to automata are constrained by the uses of the units in insn reservations. The most important constraint is: if a unit reservation is present on a particular cycle of an alternative for an insn reservation, then some unit from the same automaton must be present on the same cycle for the other alternatives of the insn reservation. The rest of the constraints are mentioned in the description of the subsequent constructions.

The following construction describes CPU functional units analogously to define\_cpu\_unit. The reservation of such units can be queried for an automaton state. The instruction scheduler never queries reservation of functional units for given automaton state. So as

<sup>&</sup>lt;sup>2</sup> However, the size of the automaton depends on processor complexity. To limit this effect, machine descriptions can split orthogonal parts of the machine description among several automata: but then, since each of these must be stepped independently, this does cause a small decrease in the algorithm's performance.

a rule, you don't need this construction. This construction could be used for future code generation goals (e.g. to generate VLIW insn templates).

```
(define_query_cpu_unit unit-names [automaton-name])
```

unit-names is a string giving names of the functional units separated by commas.

automaton-name is a string giving the name of the automaton with which the unit is bound.

The following construction is the major one to describe pipeline characteristics of an instruction.

default\_latency is a number giving latency time of the instruction. There is an important difference between the old description and the automaton based pipeline description. The latency time is used for all dependencies when we use the old description. In the automaton based pipeline description, the given latency time is only used for true dependencies. The cost of anti-dependencies is always zero and the cost of output dependencies is the difference between latency times of the producing and consuming insns (if the difference is negative, the cost is considered to be zero). You can always change the default costs for any description by using the target hook TARGET\_SCHED\_ADJUST\_COST (see Section 18.17 [Scheduling], page 583).

insn-name is a string giving the internal name of the insn. The internal names are used in constructions define\_bypass and in the automaton description file generated for debugging. The internal name has nothing in common with the names in define\_insn. It is a good practice to use insn classes described in the processor manual.

condition defines what RTL insns are described by this construction. You should remember that you will be in trouble if condition for two or more different define\_insn\_reservation constructions is TRUE for an insn. In this case what reservation will be used for the insn is not defined. Such cases are not checked during generation of the pipeline hazards recognizer because in general recognizing that two conditions may have the same value is quite difficult (especially if the conditions contain symbol\_ref). It is also not checked during the pipeline hazard recognizer work because it would slow down the recognizer considerably.

regexp is a string describing the reservation of the cpu's functional units by the instruction. The reservations are described by a regular expression according to the following syntax:

```
| result_name
| "nothing"
| "(" regexp ")"
```

- ',' is used for describing the start of the next cycle in the reservation.
- '|' is used for describing a reservation described by the first regular expression **or** a reservation described by the second regular expression **or** etc.
- '+' is used for describing a reservation described by the first regular expression **and** a reservation described by the second regular expression **and** etc.
- '\*' is used for convenience and simply means a sequence in which the regular expression are repeated *number* times with cycle advancing (see ',').
- 'cpu\_function\_unit\_name' denotes reservation of the named functional unit.
- 'reservation\_name' see description of construction 'define\_reservation'.
- 'nothing' denotes no unit reservations.

Sometimes unit reservations for different insns contain common parts. In such case, you can simplify the pipeline description by describing the common part by the following construction

```
(define_reservation reservation-name regexp)
```

reservation-name is a string giving name of regexp. Functional unit names and reservation names are in the same name space. So the reservation names should be different from the functional unit names and cannot be the reserved name 'nothing'.

The following construction is used to describe exceptions in the latency time for given instruction pair. This is so called bypasses.

number defines when the result generated by the instructions given in string out\_insn\_names will be ready for the instructions given in string in\_insn\_names. Each of these strings is a comma-separated list of filename-style globs and they refer to the names of define\_insn\_reservations. For example:

```
(define_bypass 1 "cpu1_load_*, cpu1_store_*" "cpu1_load_*")
```

defines a bypass between instructions that start with 'cpu1\_load\_' or 'cpu1\_store\_' and those that start with 'cpu1\_load\_'.

guard is an optional string giving the name of a C function which defines an additional guard for the bypass. The function will get the two insns as parameters. If the function returns zero the bypass will be ignored for this case. The additional guard is necessary to recognize complicated bypasses, e.g. when the consumer is only an address of insn 'store' (not a stored value).

If there are more one bypass with the same output and input insns, the chosen bypass is the first bypass with a guard in description whose guard function returns nonzero. If there is no such bypass, then bypass without the guard function is chosen.

The following five constructions are usually used to describe VLIW processors, or more precisely, to describe a placement of small instructions into VLIW instruction slots. They can be used for RISC processors, too.

```
(exclusion_set unit-names unit-names)
(presence_set unit-names patterns)
```

```
(final_presence_set unit-names patterns)
(absence_set unit-names patterns)
(final_absence_set unit-names patterns)
```

unit-names is a string giving names of functional units separated by commas.

patterns is a string giving patterns of functional units separated by comma. Currently pattern is one unit or units separated by white-spaces.

The first construction ('exclusion\_set') means that each functional unit in the first string cannot be reserved simultaneously with a unit whose name is in the second string and vice versa. For example, the construction is useful for describing processors (e.g. some SPARC processors) with a fully pipelined floating point functional unit which can execute simultaneously only single floating point insns or only double floating point insns.

The second construction ('presence\_set') means that each functional unit in the first string cannot be reserved unless at least one of pattern of units whose names are in the second string is reserved. This is an asymmetric relation. For example, it is useful for description that VLIW 'slot1' is reserved after 'slot0' reservation. We could describe it by the following construction

```
(presence_set "slot1" "slot0")
```

Or 'slot1' is reserved only after 'slot0' and unit 'b0' reservation. In this case we could write

```
(presence_set "slot1" "slot0 b0")
```

The third construction ('final\_presence\_set') is analogous to 'presence\_set'. The difference between them is when checking is done. When an instruction is issued in given automaton state reflecting all current and planned unit reservations, the automaton state is changed. The first state is a source state, the second one is a result state. Checking for 'presence\_set' is done on the source state reservation, checking for 'final\_presence\_set' is done on the result reservation. This construction is useful to describe a reservation which is actually two subsequent reservations. For example, if we use

```
(presence_set "slot1" "slot0")
```

the following insn will be never issued (because 'slot1' requires 'slot0' which is absent in the source state).

```
(define_reservation "insn_and_nop" "slot0 + slot1")
```

but it can be issued if we use analogous 'final\_presence\_set'.

The forth construction ('absence\_set') means that each functional unit in the first string can be reserved only if each pattern of units whose names are in the second string is not reserved. This is an asymmetric relation (actually 'exclusion\_set' is analogous to this one but it is symmetric). For example it might be useful in a VLIW description to say that 'slot0' cannot be reserved after either 'slot1' or 'slot2' have been reserved. This can be described as:

```
(absence_set "slot0" "slot1, slot2")
```

Or 'slot2' cannot be reserved if 'slot0' and unit 'b0' are reserved or 'slot1' and unit 'b1' are reserved. In this case we could write

```
(absence_set "slot2" "slot0 b0, slot1 b1")
```

All functional units mentioned in a set should belong to the same automaton.

The last construction ('final\_absence\_set') is analogous to 'absence\_set' but checking is done on the result (state) reservation. See comments for 'final\_presence\_set'.

You can control the generator of the pipeline hazard recognizer with the following construction.

```
(automata_option options)
```

options is a string giving options which affect the generated code. Currently there are the following options:

- no-minimization makes no minimization of the automaton. This is only worth to do when we are debugging the description and need to look more accurately at reservations of states.
- time means printing time statistics about the generation of automata.
- stats means printing statistics about the generated automata such as the number of DFA states, NDFA states and arcs.
- v means a generation of the file describing the result automata. The file has suffix '.dfa' and can be used for the description verification and debugging.
- w means a generation of warning instead of error for non-critical errors.
- no-comb-vect prevents the automaton generator from generating two data structures and comparing them for space efficiency. Using a comb vector to represent transitions may be better, but it can be very expensive to construct. This option is useful if the build process spends an unacceptably long time in genautomata.
- ndfa makes nondeterministic finite state automata. This affects the treatment of operator '|' in the regular expressions. The usual treatment of the operator is to try the first alternative and, if the reservation is not possible, the second alternative. The non-deterministic treatment means trying all alternatives, some of them may be rejected by reservations in the subsequent insns.
- collapse-ndfa modifies the behavior of the generator when producing an automaton. An additional state transition to collapse a nondeterministic NDFA state to a deterministic DFA state is generated. It can be triggered by passing const0\_rtx to state\_transition. In such an automaton, cycle advance transitions are available only for these collapsed states. This option is useful for ports that want to use the ndfa option, but also want to use define\_query\_cpu\_unit to assign units to insns issued in a cycle.
- progress means output of a progress bar showing how many states were generated so far for automaton being processed. This is useful during debugging a DFA description. If you see too many generated states, you could interrupt the generator of the pipeline hazard recognizer and try to figure out a reason for generation of the huge automaton.

As an example, consider a superscalar RISC machine which can issue three insns (two integer insns and one floating point insn) on the cycle but can finish only two insns. To describe this, we define the following functional units.

```
(define_cpu_unit "i0_pipeline, i1_pipeline, f_pipeline")
(define_cpu_unit "port0, port1")
```

All simple integer insns can be executed in any integer pipeline and their result is ready in two cycles. The simple integer insns are issued into the first pipeline unless it is reserved, otherwise they are issued into the second pipeline. Integer division and multiplication insns can be executed only in the second integer pipeline and their results are ready correspondingly in 9 and 4 cycles. The integer division is not pipelined, i.e. the subsequent integer division insn cannot be issued until the current division insn finished. Floating point insns

are fully pipelined and their results are ready in 3 cycles. Where the result of a floating point insn is used by an integer insn, an additional delay of one cycle is incurred. To describe all of this we could specify

#### 17.20 Conditional Execution

A number of architectures provide for some form of conditional execution, or predication. The hallmark of this feature is the ability to nullify most of the instructions in the instruction set. When the instruction set is large and not entirely symmetric, it can be quite tedious to describe these forms directly in the '.md' file. An alternative is the define\_cond\_exec template.

```
(define_cond_exec
  [predicate-pattern]
  "condition"
  "output-template"
  "optional-insn-attribues")
```

predicate-pattern is the condition that must be true for the insn to be executed at runtime and should match a relational operator. One can use match\_operator to match several relational operators at once. Any match\_operand operands must have no more than one alternative.

condition is a C expression that must be true for the generated pattern to match.

output-template is a string similar to the define\_insn output template (see Section 17.5 [Output Template], page 343), except that the '\*' and '@' special cases do not apply. This is only useful if the assembly text for the predicate is a simple prefix to the main insn. In order to handle the general case, there is a global variable current\_insn\_predicate that will contain the entire predicate if the current insn is predicated, and will otherwise be NULL.

optional-insn-attributes is an optional vector of attributes that gets appended to the insn attributes of the produced cond\_exec rtx. It can be used to add some distinguishing attribute to cond\_exec rtxs produced that way. An example usage would be to use this at-

tribute in conjunction with attributes on the main pattern to disable particular alternatives under certain conditions.

When define\_cond\_exec is used, an implicit reference to the predicable instruction attribute is made. See Section 17.19 [Insn Attributes], page 450. This attribute must be a boolean (i.e. have exactly two elements in its list-of-values), with the possible values being no and yes. The default and all uses in the insns must be a simple constant, not a complex expressions. It may, however, depend on the alternative, by using a commaseparated list of values. If that is the case, the port should also define an enabled attribute (see Section 17.8.6 [Disable Insn Alternatives], page 387), which should also allow only no and yes as its values.

For each define\_insn for which the predicable attribute is true, a new define\_insn pattern will be generated that matches a predicated version of the instruction. For example,

```
(define_insn "addsi"
        [(set (match_operand:SI 0 "register_operand" "r")
              (plus:SI (match_operand:SI 1 "register_operand" "r")
                        (match_operand:SI 2 "register_operand" "r")))]
        "+es+1"
        "add %2,%1,%0")
      (define_cond_exec
        [(ne (match_operand:CC 0 "register_operand" "c")
             (const_int 0))]
        "test2"
        "(%0)")
generates a new pattern
      (define_insn ""
        [(cond_exec
           (ne (match_operand:CC 3 "register_operand" "c") (const_int 0))
           (set (match_operand:SI 0 "register_operand" "r")
                 (plus:SI (match_operand:SI 1 "register_operand" "r")
                          (match_operand:SI 2 "register_operand" "r"))))]
        "(test2) && (test1)"
        "(\%3) add \%2,\%1,\%0")
```

# 17.21 RTL Templates Transformations

For some hardware architectures there are common cases when the RTL templates for the instructions can be derived from the other RTL templates using simple transformations. E.g., 'i386.md' contains an RTL template for the ordinary sub instruction—\*subsi\_1, and for the sub instruction with subsequent zero-extension—\*subsi\_1\_zext. Such cases can be easily implemented by a single meta-template capable of generating a modified case based on the initial one:

```
(define_subst "name"
  [input-template]
  "condition"
  [output-template])
```

input-template is a pattern describing the source RTL template, which will be transformed.

condition is a C expression that is conjunct with the condition from the input-template to generate a condition to be used in the output-template.

output-template is a pattern that will be used in the resulting template.

define\_subst mechanism is tightly coupled with the notion of the subst attribute (see Section 17.23.4 [Subst Iterators], page 476). The use of define\_subst is triggered by a reference to a subst attribute in the transforming RTL template. This reference initiates duplication of the source RTL template and substitution of the attributes with their values. The source RTL template is left unchanged, while the copy is transformed by define\_subst. This transformation can fail in the case when the source RTL template is not matched against the input-template of the define\_subst. In such case the copy is deleted.

define\_subst can be used only in define\_insn and define\_expand, it cannot be used in other expressions (e.g. in define\_insn\_and\_split).

#### 17.21.1 define\_subst Example

To illustrate how define\_subst works, let us examine a simple template transformation.

Suppose there are two kinds of instructions: one that touches flags and the other that does not. The instructions of the second type could be generated with the following define\_subst:

This define\_subst can be applied to any RTL pattern containing set of mode SI and generates a copy with clobber when it is applied.

Assume there is an RTL template for a max instruction to be used in define\_subst mentioned above:

To mark the RTL template for define\_subst application, subst-attributes are used. They should be declared in advance:

```
(define_subst_attr "add_clobber_name" "add_clobber_subst" "_noclobber" "_clobber")
```

Here 'add\_clobber\_name' is the attribute name, 'add\_clobber\_subst' is the name of the corresponding define\_subst, the third argument ('\_noclobber') is the attribute value that would be substituted into the unchanged version of the source RTL template, and the last argument ('\_clobber') is the value that would be substituted into the second, transformed, version of the RTL template.

Once the subst-attribute has been defined, it should be used in RTL templates which need to be processed by the define\_subst. So, the original RTL template should be changed:

```
(match_operand:SI 1 "register_operand" "r")
               (match_operand:SI 2 "register_operand" "r")))]
      "max\t{%2, %1, %0|%0, %1, %2}"
The result of the define_subst usage would look like the following:
    (define_insn "maxsi_noclobber"
      [(set (match_operand:SI 0 "register_operand" "=r")
             (max:SI
               (match_operand:SI 1 "register_operand" "r")
               (match_operand:SI 2 "register_operand" "r")))]
      \max\{\frac{2}{1}, \frac{1}{1}, \frac{0}{0}, \frac{1}{1}, \frac{2}{0}\}
     [\ldots]
    (define_insn "maxsi_clobber"
      [(set (match_operand:SI 0 "register_operand" "=r")
             (max:SI
               (match_operand:SI 1 "register_operand" "r")
               (match_operand:SI 2 "register_operand" "r")))
       (clobber (reg:CC FLAGS_REG))]
      "max\t{%2, %1, %0|%0, %1, %2}"
     [\ldots]
```

#### 17.21.2 Pattern Matching in define\_subst

All expressions, allowed in define\_insn or define\_expand, are allowed in the input-template of define\_subst, except match\_par\_dup, match\_scratch, match\_parallel. The meanings of expressions in the input-template were changed:

match\_operand matches any expression (possibly, a subtree in RTL-template), if modes of the match\_operand and this expression are the same, or mode of the match\_operand is VOIDmode, or this expression is match\_dup, match\_op\_dup. If the expression is match\_operand too, and predicate of match\_operand from the input pattern is not empty, then the predicates are compared. That can be used for more accurate filtering of accepted RTL-templates.

match\_operator matches common operators (like plus, minus), unspec, unspec\_volatile operators and match\_operators from the original pattern if the modes match and match\_operator from the input pattern has the same number of operands as the operator from the original pattern.

#### 17.21.3 Generation of output template in define\_subst

If all necessary checks for define\_subst application pass, a new RTL-pattern, based on the output-template, is created to replace the old template. Like in input-patterns, meanings of some RTL expressions are changed when they are used in output-patterns of a define\_subst. Thus, match\_dup is used for copying the whole expression from the original pattern, which matched corresponding match\_operand from the input pattern.

match\_dup N is used in the output template to be replaced with the expression from the original pattern, which matched match\_operand N from the input pattern. As a consequence, match\_dup cannot be used to point to match\_operands from the output pattern, it should always refer to a match\_operand from the input pattern. If a match\_dup N occurs more than once in the output template, its first occurrence is replaced with the expression

from the original pattern, and the subsequent expressions are replaced with match\_dup N, i.e., a reference to the first expression.

In the output template one can refer to the expressions from the original pattern and create new ones. For instance, some operands could be added by means of standard match\_operand.

After replacing match\_dup with some RTL-subtree from the original pattern, it could happen that several match\_operands in the output pattern have the same indexes. It is unknown, how many and what indexes would be used in the expression which would replace match\_dup, so such conflicts in indexes are inevitable. To overcome this issue, match\_operands and match\_operators, which were introduced into the output pattern, are renumerated when all match\_dups are replaced.

Number of alternatives in match\_operands introduced into the output template M could differ from the number of alternatives in the original pattern N, so in the resultant pattern there would be N\*M alternatives. Thus, constraints from the original pattern would be duplicated N times, constraints from the output pattern would be duplicated M times, producing all possible combinations.

#### 17.22 Constant Definitions

Using literal constants inside instruction patterns reduces legibility and can be a maintenance problem.

To overcome this problem, you may use the define\_constants expression. It contains a vector of name-value pairs. From that point on, wherever any of the names appears in the MD file, it is as if the corresponding value had been written instead. You may use define\_constants multiple times; each appearance adds more constants to the table. It is an error to redefine a constant with a different value.

To come back to the a29k load multiple example, instead of

```
(define_insn ""
      [(match_parallel 0 "load_multiple_operation"
         [(set (match_operand:SI 1 "gpc_reg_operand" "=r")
               (match_operand:SI 2 "memory_operand" "m"))
          (use (reg:SI 179))
          (clobber (reg:SI 179))])]
      "loadm 0,0,%1,%2")
You could write:
    (define_constants [
        (R_BP 177)
        (R_FC 178)
        (R_CR 179)
        (R_Q 180)
   1)
   (define_insn ""
      [(match_parallel 0 "load_multiple_operation"
         [(set (match_operand:SI 1 "gpc_reg_operand" "=r")
               (match_operand:SI 2 "memory_operand" "m"))
          (use (reg:SI R_CR))
          (clobber (reg:SI R_CR))])]
```

```
"loadm 0,0,%1,%2")
```

The constants that are defined with a define\_constant are also output in the insn-codes.h header file as #defines.

You can also use the machine description file to define enumerations. Like the constants defined by define\_constant, these enumerations are visible to both the machine description file and the main C code.

The syntax is as follows:

```
(define_c_enum "name" [
  value0
  value1
   ...
  valuen
])
```

This definition causes the equivalent of the following C code to appear in 'insn-constants.h':

```
enum name {
  value0 = 0,
  value1 = 1,
    ...
  valuen = n
};
#define NUM_cname_VALUES (n + 1)
```

where *cname* is the capitalized form of *name*. It also makes each *valuei* available in the machine description file, just as if it had been declared with:

```
(define_constants [(valuei i)])
```

Each value is usually an upper-case identifier and usually begins with cname.

You can split the enumeration definition into as many statements as you like. The above example is directly equivalent to:

```
(define_c_enum "name" [value0])
(define_c_enum "name" [value1])
...
(define_c_enum "name" [valuen])
```

Splitting the enumeration helps to improve the modularity of each individual .md file. For example, if a port defines its synchronization instructions in a separate 'sync.md' file, it is convenient to define all synchronization-specific enumeration values in 'sync.md' rather than in the main '.md' file.

Some enumeration names have special significance to GCC:

unspec\_volatile expressions. For example:

unspecv If an enumeration called unspecv is defined, GCC will use it when printing out

```
(define_c_enum "unspecv" [
        UNSPECV_BLOCKAGE
])
causes GCC to print '(unspec_volatile ... 0)' as:
        (unspec_volatile ... UNSPECV_BLOCKAGE)
```

unspec

If an enumeration called unspec is defined, GCC will use it when printing out unspec expressions. GCC will also use it when printing out unspec\_volatile expressions unless an unspecv enumeration is also defined. You can therefore

decide whether to keep separate enumerations for volatile and non-volatile expressions or whether to use the same enumeration for both.

Another way of defining an enumeration is to use define\_enum:

```
(define_enum "name" [
    value0
    value1
    ...
    valuen
])
This directive implies:
    (define_c_enum "name" [
        cname_cvalue0
        cname_cvalue1
    ...
        cname_cvaluen
])
```

where *cvaluei* is the capitalized form of *valuei*. However, unlike define\_c\_enum, the enumerations defined by define\_enum can be used in attribute specifications (see [define\_enum\_attr], page 452).

#### 17.23 Iterators

Ports often need to define similar patterns for more than one machine mode or for more than one rtx code. GCC provides some simple iterator facilities to make this process easier.

#### 17.23.1 Mode Iterators

Ports often need to define similar patterns for two or more different modes. For example:

- If a processor has hardware support for both single and double floating-point arithmetic, the SFmode patterns tend to be very similar to the DFmode ones.
- If a port uses SImode pointers in one configuration and DImode pointers in another, it will usually have very similar SImode and DImode patterns for manipulating pointers.

Mode iterators allow several patterns to be instantiated from one '.md' file template. They can be used with any type of rtx-based construct, such as a define\_insn, define\_split, or define\_peephole2.

#### 17.23.1.1 Defining Mode Iterators

The syntax for defining a mode iterator is:

```
(define_mode_iterator name [(mode1 "cond1") ... (moden "condn")])
```

This allows subsequent '.md' file constructs to use the mode suffix :name. Every construct that does so will be expanded n times, once with every use of :name replaced by :mode1, once with every use replaced by :mode2, and so on. In the expansion for a particular modei, every C condition will also require that condi be true.

```
For example:
```

```
(define_mode_iterator P [(SI "Pmode == SImode") (DI "Pmode == DImode")])
```

defines a new mode suffix: P. Every construct that uses: P will be expanded twice, once with every: P replaced by: SI and once with every: P replaced by: DI. The: SI version will only apply if Pmode == SImode and the: DI version will only apply if Pmode == DImode.

As with other '.md' conditions, an empty string is treated as "always true". (mode "") can also be abbreviated to mode. For example:

```
(define_mode_iterator GPR [SI (DI "TARGET_64BIT")])
```

means that the :DI expansion only applies if TARGET\_64BIT but that the :SI expansion has no such constraint.

Iterators are applied in the order they are defined. This can be significant if two iterators are used in a construct that requires substitutions. See Section 17.23.1.2 [Substitutions], page 473.

#### 17.23.1.2 Substitution in Mode Iterators

If an '.md' file construct uses mode iterators, each version of the construct will often need slightly different strings or modes. For example:

- When a define\_expand defines several addm3 patterns (see Section 17.9 [Standard Names], page 392), each expander will need to use the appropriate mode name for m.
- When a define\_insn defines several instruction patterns, each instruction will often use a different assembler mnemonic.
- When a define\_insn requires operands with different modes, using an iterator for one of the operand modes usually requires a specific mode for the other operand(s).

GCC supports such variations through a system of "mode attributes". There are two standard attributes: mode, which is the name of the mode in lower case, and MODE, which is the same thing in upper case. You can define other attributes using:

```
(define_mode_attr name [(mode1 "value1") ... (moden "valuen")])
```

where name is the name of the attribute and value is the value associated with modei.

When GCC replaces some :iterator with :mode, it will scan each string and mode in the pattern for sequences of the form <iterator:attr>, where attr is the name of a mode attribute. If the attribute is defined for mode, the whole <...> sequence will be replaced by the appropriate attribute value.

For example, suppose an '.md' file has:

```
(define_mode_iterator P [(SI "Pmode == SImode") (DI "Pmode == DImode")])
(define_mode_attr load [(SI "lw") (DI "ld")])
```

If one of the patterns that uses :P contains the string "<P:load>\t%0,%1", the SI version of that pattern will use "lw\t%0,%1" and the DI version will use "ld\t%0,%1".

Here is an example of using an attribute for a mode:

```
(define_mode_iterator LONG [SI DI])
(define_mode_attr SHORT [(SI "HI") (DI "SI")])
(define_insn ...
  (sign_extend:LONG (match_operand:<LONG:SHORT> ...)) ...)
```

The *iterator*: prefix may be omitted, in which case the substitution will be attempted for every iterator expansion.

#### 17.23.1.3 Mode Iterator Examples

Here is an example from the MIPS port. It defines the following modes and attributes (among others):

```
(define_mode_iterator GPR [SI (DI "TARGET_64BIT")])
    (define_mode_attr d [(SI "") (DI "d")])
and uses the following template to define both subsi3 and subdi3:
   (define insn "sub<mode>3"
      [(set (match_operand:GPR 0 "register_operand" "=d")
            (minus:GPR (match_operand:GPR 1 "register_operand" "d")
                       (match_operand:GPR 2 "register_operand" "d")))]
      "<d>subu\t%0,%1,%2"
      [(set_attr "type" "arith")
       (set_attr "mode" "<MODE>")])
This is exactly equivalent to:
    (define_insn "subsi3"
      [(set (match_operand:SI 0 "register_operand" "=d")
            (minus:SI (match_operand:SI 1 "register_operand" "d")
                      (match_operand:SI 2 "register_operand" "d")))]
      "subu\t%0,%1,%2"
      [(set_attr "type" "arith")
       (set_attr "mode" "SI")])
   (define_insn "subdi3"
      [(set (match_operand:DI 0 "register_operand" "=d")
            (minus:DI (match_operand:DI 1 "register_operand" "d")
                      (match_operand:DI 2 "register_operand" "d")))]
      "dsubu\t%0,%1,%2"
      [(set_attr "type" "arith")
       (set_attr "mode" "DI")])
```

#### 17.23.2 Code Iterators

Code iterators operate in a similar way to mode iterators. See Section 17.23.1 [Mode Iterators], page 472.

The construct:

```
(define_code_iterator name [(code1 "cond1") ... (coden "condn")])
```

defines a pseudo rtx code name that can be instantiated as codei if condition condi is true. Each codei must have the same rtx format. See Section 14.2 [RTL Classes], page 260.

As with mode iterators, each pattern that uses name will be expanded n times, once with all uses of name replaced by code1, once with all uses replaced by code2, and so on. See Section 17.23.1.1 [Defining Mode Iterators], page 472.

It is possible to define attributes for codes as well as for modes. There are two standard code attributes: code, the name of the code in lower case, and CODE, the name of the code in upper case. Other attributes are defined using:

```
(define_code_attr name [(code1 "value1") ... (coden "valuen")])
```

Instruction patterns can use code attributes as rtx codes, which can be useful if two sets of codes act in tandem. For example, the following define\_insn defines two patterns, one calculating a signed absolute difference and another calculating an unsigned absolute difference:

```
(define_code_iterator any_max [smax umax])
(define_code_attr paired_min [(smax "smin") (umax "umin")])
```

The signed version of the instruction uses smax and smin while the unsigned version uses umax and umin. There are no versions that pair smax with umin or umax with smin.

Here's an example of code iterators in action, taken from the MIPS port:

```
(define_code_iterator any_cond [unordered ordered unlt unge uneq ltgt unle ungt
                                     eq ne gt ge lt le gtu geu ltu leu])
    (define_expand "b<code>"
      [(set (pc)
            (if_then_else (any_cond:CC (cc0)
                                        (const_int 0))
                           (label_ref (match_operand 0 ""))
                           (pc)))]
    {
     gen_conditional_branch (operands, <CODE>);
     DONE;
   })
This is equivalent to:
    (define_expand "bunordered"
      [(set (pc)
            (if_then_else (unordered:CC (cc0)
                                         (const_int 0))
                           (label_ref (match_operand 0 ""))
                           (pc)))]
      11 11
    {
     gen_conditional_branch (operands, UNORDERED);
     DONE;
    })
    (define_expand "bordered"
      [(set (pc)
            (if_then_else (ordered:CC (cc0)
                                       (const_int 0))
                           (label_ref (match_operand 0 ""))
                           (pc)))]
      gen_conditional_branch (operands, ORDERED);
     DONE;
   })
```

#### 17.23.3 Int Iterators

Int iterators operate in a similar way to code iterators. See Section 17.23.2 [Code Iterators], page 474.

The construct:

```
(define_int_iterator name [(int1 "cond1") ... (intn "condn")])
```

defines a pseudo integer constant name that can be instantiated as *inti* if condition *condi* is true. Each *int* must have the same rtx format. See Section 14.2 [RTL Classes], page 260. Int iterators can appear in only those rtx fields that have 'i' as the specifier. This means that each *int* has to be a constant defined using define\_constant or define\_c\_enum.

As with mode and code iterators, each pattern that uses name will be expanded n times, once with all uses of name replaced by int1, once with all uses replaced by int2, and so on. See Section 17.23.1.1 [Defining Mode Iterators], page 472.

It is possible to define attributes for ints as well as for codes and modes. Attributes are defined using:

```
(define_int_attr name [(int1 "value1") ... (intn "valuen")])
Here's an example of int iterators in action, taken from the ARM port:
    (define_int_iterator QABSNEG [UNSPEC_VQABS UNSPEC_VQNEG])
    (define_int_attr absneg [(UNSPEC_VQABS "abs") (UNSPEC_VQNEG "neg")])
    (define_insn "neon_vq<absneg><mode>"
      [(set (match_operand:VDQIW 0 "s_register_operand" "=w")
    (unspec: VDQIW [(match_operand: VDQIW 1 "s_register_operand" "w")
           (match_operand:SI 2 "immediate_operand" "i")]
          QABSNEG))]
      "TARGET_NEON"
      "vq<absneg>.<V_s_elem>\t%<V_reg>0, %<V_reg>1"
      [(set_attr "type" "neon_vqneg_vqabs")]
    )
This is equivalent to:
    (define_insn "neon_vqabs<mode>"
      [(set (match_operand:VDQIW 0 "s_register_operand" "=w")
    (unspec:VDQIW [(match_operand:VDQIW 1 "s_register_operand" "w")
           (match_operand:SI 2 "immediate_operand" "i")]
          UNSPEC_VQABS))]
      "TARGET_NEON"
      "vqabs.<V_s_elem>\t%<V_reg>0, %<V_reg>1"
      [(set_attr "type" "neon_vqneg_vqabs")]
    (define_insn "neon_vqneg<mode>"
      [(set (match_operand:VDQIW 0 "s_register_operand" "=w")
    (unspec:VDQIW [(match_operand:VDQIW 1 "s_register_operand" "w")
           (match_operand:SI 2 "immediate_operand" "i")]
          UNSPEC_VQNEG))]
      "TARGET_NEON"
      "vqneg.<V_s_elem>\t%<V_reg>0, %<V_reg>1"
      [(set_attr "type" "neon_vqneg_vqabs")]
    )
```

#### 17.23.4 Subst Iterators

Subst iterators are special type of iterators with the following restrictions: they could not be declared explicitly, they always have only two values, and they do not have explicit dedicated name. Subst-iterators are triggered only when corresponding subst-attribute is used in RTL-pattern.

Subst iterators transform templates in the following way: the templates are duplicated, the subst-attributes in these templates are replaced with the corresponding values, and a new attribute is implicitly added to the given define\_insn/define\_expand. The name of the added attribute matches the name of define\_subst. Such attributes are declared implicitly, and it is not allowed to have a define\_attr named as a define\_subst.

Each subst iterator is linked to a define\_subst. It is declared implicitly by the first appearance of the corresponding define\_subst\_attr, and it is not allowed to define it explicitly.

Declarations of subst-attributes have the following syntax:

```
(define_subst_attr "name"
  "subst-name"
  "no-subst-value"
  "subst-applied-value")
```

name is a string with which the given subst-attribute could be referred to.

subst-name shows which define\_subst should be applied to an RTL-template if the given subst-attribute is present in the RTL-template.

no-subst-value is a value with which subst-attribute would be replaced in the first copy of the original RTL-template.

subst-applied-value is a value with which subst-attribute would be replaced in the second copy of the original RTL-template.

#### 17.23.5 Parameterized Names

Ports sometimes need to apply iterators using C++ code, in order to get the code or RTL pattern for a specific instruction. For example, suppose we have the 'neon\_vq<absney><mode>' pattern given above:

A port might need to generate this pattern for a variable 'QABSNEG' value and a variable 'VDQIW' mode. There are two ways of doing this. The first is to build the rtx for the pattern directly from C++ code; this is a valid technique and avoids any risk of combinatorial explosion. The second is to prefix the instruction name with the special character '@', which tells GCC to generate the four additional functions below. In each case, name is the name of the instruction without the leading '@' character, without the '<...>' placeholders, and with any underscore before a '<...>' placeholder removed if keeping it would lead to a double or trailing underscore.

```
'insn_code maybe_code_for_name (i1, i2, ...)'
```

See whether replacing the first '<...>' placeholder with iterator value i1, the second with iterator value i2, and so on, gives a valid instruction. Return its code if so, otherwise return CODE\_FOR\_nothing.

```
'insn_code code_for_name (i1, i2, ...)'
```

Same, but abort the compiler if the requested instruction does not exist.

```
'rtx maybe_gen_name (i1, i2, ..., op0, op1, ...)'
```

Check for a valid instruction in the same way as maybe\_code\_for\_name. If the instruction exists, generate an instance of it using the operand values given by op0, op1, and so on, otherwise return null.

```
'rtx gen_name (i1, i2, ..., op0, op1, ...)'
```

Same, but abort the compiler if the requested instruction does not exist, or if the instruction generator invoked the FAIL macro.

For example, changing the pattern above to:

would define the same patterns as before, but in addition would generate the four functions below:

```
insn_code maybe_code_for_neon_vq (int, machine_mode);
insn_code code_for_neon_vq (int, machine_mode);
rtx maybe_gen_neon_vq (int, machine_mode, rtx, rtx, rtx);
rtx gen_neon_vq (int, machine_mode, rtx, rtx, rtx);
```

Calling 'code\_for\_neon\_vq (UNSPEC\_VQABS, V8QImode)' would then give CODE\_FOR\_neon\_vqabsv8qi.

It is possible to have multiple '@' patterns with the same name and same types of iterator. For example:

```
(define_insn "@some_arithmetic_op<mode>"
    [(set (match_operand:INTEGER_MODES 0 "register_operand") ...)]
    ...
)
(define_insn "@some_arithmetic_op<mode>"
    [(set (match_operand:FLOAT_MODES 0 "register_operand") ...)]
    ...
)
```

would produce a single set of functions that handles both INTEGER\_MODES and FLOAT\_MODES.

It is also possible for these '@' patterns to have different numbers of operands from each other. For example, patterns with a binary rtl code might take three operands (one output and two inputs) while patterns with a ternary rtl code might take four operands (one output and three inputs). This combination would produce separate 'maybe\_gen\_name' and 'gen\_name' functions for each operand count, but it would still produce a single 'maybe\_code\_for\_name' and a single 'code\_for\_name'.

# 18 Target Description Macros and Functions

In addition to the file 'machine.md', a machine description includes a C header file conventionally given the name 'machine.h' and a C source file named 'machine.c'. The header file defines numerous macros that convey the information about the target machine that does not fit into the scheme of the '.md' file. The file 'tm.h' should be a link to 'machine.h'. The header file 'config.h' includes 'tm.h' and most compiler source files include 'config.h'. The source file defines a variable targetm, which is a structure containing pointers to functions and data relating to the target machine. 'machine.c' should also contain their definitions, if they are not defined elsewhere in GCC, and other functions called through the macros defined in the '.h' file.

## 18.1 The Global targetm Variable

#### struct gcc\_target targetm

[Variable]

The target '.c' file must define the global targetm variable which contains pointers to functions and data relating to the target machine. The variable is declared in 'target.h'; 'target-def.h' defines the macro TARGET\_INITIALIZER which is used to initialize the variable, and macros for the default initializers for elements of the structure. The '.c' file should override those macros for which the default definition is inappropriate. For example:

```
#include "target.h"
#include "target-def.h"

/* Initialize the GCC target structure. */
#undef TARGET_COMP_TYPE_ATTRIBUTES
#define TARGET_COMP_TYPE_ATTRIBUTES machine_comp_type_attributes
struct gcc_target targetm = TARGET_INITIALIZER;
```

Where a macro should be defined in the '.c' file in this manner to form part of the targetm structure, it is documented below as a "Target Hook" with a prototype. Many macros will change in future from being defined in the '.h' file to being part of the targetm structure.

Similarly, there is a targetcm variable for hooks that are specific to front ends for C-family languages, documented as "C Target Hook". This is declared in 'c-family/c-target.h', the initializer TARGETCM\_INITIALIZER in 'c-family/c-target-def.h'. If targets initialize targetcm themselves, they should set target\_has\_targetcm=yes in 'config.gcc'; otherwise a default definition is used.

Similarly, there is a targetm\_common variable for hooks that are shared between the compiler driver and the compilers proper, documented as "Common Target Hook". This is declared in 'common/common-target.h', the initializer TARGETM\_COMMON\_INITIALIZER in 'common/common-target-def.h'. If targets initialize targetm\_common themselves, they should set target\_has\_targetm\_common=yes in 'config.gcc'; otherwise a default definition is used.

Similarly, there is a targetdm variable for hooks that are specific to the D language front end, documented as "D Target Hook". This is declared in 'd/d-target.h', the initializer

TARGETDM\_INITIALIZER in 'd/d-target-def.h'. If targets initialize targetdm themselves, they should set target\_has\_targetdm=yes in 'config.gcc'; otherwise a default definition is used.

# 18.2 Controlling the Compilation Driver, 'gcc'

You can control the compilation driver.

#### DRIVER\_SELF\_SPECS

A list of specs for the driver itself. It should be a suitable initializer for an array of strings, with no surrounding braces.

The driver applies these specs to its own command line between loading default 'specs' files (but not command-line specified ones) and choosing the multilib directory or running any subcommands. It applies them in the order given, so each spec can depend on the options added by earlier ones. It is also possible to remove options using '%<option' in the usual way.

This macro can be useful when a port has several interdependent target options. It provides a way of standardizing the command line so that the other specs are easier to write.

Do not define this macro if it does not need to do anything.

#### OPTION\_DEFAULT\_SPECS

[Macro]

[Macro]

A list of specs used to support configure-time default options (i.e. '--with' options) in the driver. It should be a suitable initializer for an array of structures, each containing two strings, without the outermost pair of surrounding braces.

The first item in the pair is the name of the default. This must match the code in 'config.gcc' for the target. The second item is a spec to apply if a default with this name was specified. The string '%(VALUE)' in the spec will be replaced by the value of the default everywhere it occurs.

The driver will apply these specs to its own command line between loading default 'specs' files and processing DRIVER\_SELF\_SPECS, using the same mechanism as DRIVER\_SELF\_SPECS.

Do not define this macro if it does not need to do anything.

CPP\_SPEC [Macro]

A C string constant that tells the GCC driver program options to pass to CPP. It can also specify how to translate options you give to GCC into options for GCC to pass to the CPP.

Do not define this macro if it does not need to do anything.

#### CPLUSPLUS\_CPP\_SPEC

[Macro]

This macro is just like CPP\_SPEC, but is used for C++, rather than C. If you do not define this macro, then the value of CPP\_SPEC (if any) will be used instead.

CC1\_SPEC [Macro]

A C string constant that tells the GCC driver program options to pass to cc1, cc1plus, f771, and the other language front ends. It can also specify how to translate options you give to GCC into options for GCC to pass to front ends.

Do not define this macro if it does not need to do anything.

CC1PLUS\_SPEC [Macro]

A C string constant that tells the GCC driver program options to pass to cc1plus. It can also specify how to translate options you give to GCC into options for GCC to pass to the cc1plus.

Do not define this macro if it does not need to do anything. Note that everything defined in CC1\_SPEC is already passed to cc1plus so there is no need to duplicate the contents of CC1\_SPEC in CC1PLUS\_SPEC.

ASM\_SPEC [Macro]

A C string constant that tells the GCC driver program options to pass to the assembler. It can also specify how to translate options you give to GCC into options for GCC to pass to the assembler. See the file 'sun3.h' for an example of this.

Do not define this macro if it does not need to do anything.

ASM\_FINAL\_SPEC [Macro]

A C string constant that tells the GCC driver program how to run any programs which cleanup after the normal assembler. Normally, this is not needed. See the file 'mips.h' for an example of this.

Do not define this macro if it does not need to do anything.

#### AS\_NEEDS\_DASH\_FOR\_PIPED\_INPUT

[Macro]

Define this macro, with no value, if the driver should give the assembler an argument consisting of a single dash, '-', to instruct it to read from its standard input (which will be a pipe connected to the output of the compiler proper). This argument is given after any '-o' option specifying the name of the output file.

If you do not define this macro, the assembler is assumed to read its standard input if given no non-option arguments. If your assembler cannot read standard input at all, use a "%{pipe:%e}" construct; see 'mips.h' for instance.

LINK\_SPEC [Macro]

A C string constant that tells the GCC driver program options to pass to the linker. It can also specify how to translate options you give to GCC into options for GCC to pass to the linker.

Do not define this macro if it does not need to do anything.

LIB\_SPEC [Macro]

Another C string constant used much like LINK\_SPEC. The difference between the two is that LIB\_SPEC is used at the end of the command given to the linker.

If this macro is not defined, a default is provided that loads the standard C library from the usual place. See 'gcc.c'.

LIBGCC\_SPEC [Macro]

Another C string constant that tells the GCC driver program how and when to place a reference to 'libgcc.a' into the linker command line. This constant is placed both before and after the value of LIB\_SPEC.

If this macro is not defined, the GCC driver provides a default that passes the string '-lgcc' to the linker.

#### REAL\_LIBGCC\_SPEC

[Macro]

By default, if ENABLE\_SHARED\_LIBGCC is defined, the LIBGCC\_SPEC is not directly used by the driver program but is instead modified to refer to different versions of 'libgcc.a' depending on the values of the command line flags '-static', '-shared', '-static-libgcc', and '-shared-libgcc'. On targets where these modifications are inappropriate, define REAL\_LIBGCC\_SPEC instead. REAL\_LIBGCC\_SPEC tells the driver how to place a reference to 'libgcc' on the link command line, but, unlike LIBGCC\_SPEC, it is used unmodified.

#### USE\_LD\_AS\_NEEDED

[Macro]

A macro that controls the modifications to LIBGCC\_SPEC mentioned in REAL\_LIBGCC\_SPEC. If nonzero, a spec will be generated that uses '--as-needed' or equivalent options and the shared 'libgcc' in place of the static exception handler library, when linking without any of -static, -static-libgcc, or -shared-libgcc.

LINK\_EH\_SPEC [Macro]

If defined, this C string constant is added to LINK\_SPEC. When USE\_LD\_AS\_NEEDED is zero or undefined, it also affects the modifications to LIBGCC\_SPEC mentioned in REAL\_LIBGCC\_SPEC.

#### STARTFILE\_SPEC

[Macro]

Another C string constant used much like LINK\_SPEC. The difference between the two is that STARTFILE\_SPEC is used at the very beginning of the command given to the linker.

If this macro is not defined, a default is provided that loads the standard C startup file from the usual place. See 'gcc.c'.

#### ENDFILE\_SPEC

[Macro]

Another C string constant used much like LINK\_SPEC. The difference between the two is that ENDFILE\_SPEC is used at the very end of the command given to the linker.

Do not define this macro if it does not need to do anything.

#### THREAD\_MODEL\_SPEC

[Macro]

GCC -v will print the thread model GCC was configured to use. However, this doesn't work on platforms that are multilibbed on thread models, such as AIX 4.3. On such platforms, define THREAD\_MODEL\_SPEC such that it evaluates to a string without blanks that names one of the recognized thread models. %\*, the default value of this macro, will expand to the value of thread\_file set in 'config.gcc'.

#### SYSROOT\_SUFFIX\_SPEC

[Macro]

Define this macro to add a suffix to the target sysroot when GCC is configured with a sysroot. This will cause GCC to search for usr/lib, et al, within sysroot+suffix.

#### SYSROOT\_HEADERS\_SUFFIX\_SPEC

|Macro

Define this macro to add a headers\_suffix to the target sysroot when GCC is configured with a sysroot. This will cause GCC to pass the updated sysroot+headers\_suffix to CPP, causing it to search for usr/include, et al, within sysroot+headers\_suffix.

EXTRA\_SPECS [Macro]

Define this macro to provide additional specifications to put in the 'specs' file that can be used in various specifications like CC1\_SPEC.

The definition should be an initializer for an array of structures, containing a string constant, that defines the specification name, and a string constant that provides the specification.

Do not define this macro if it does not need to do anything.

EXTRA\_SPECS is useful when an architecture contains several related targets, which have various ...\_SPECS which are similar to each other, and the maintainer would like one central place to keep these definitions.

For example, the PowerPC System V.4 targets use EXTRA\_SPECS to define either \_ CALL\_SYSV when the System V calling sequence is used or \_CALL\_AIX when the older AIX-based calling sequence is used.

The 'config/rs6000/rs6000.h' target file defines:

#### LINK\_LIBGCC\_SPECIAL\_1

#undef CPP\_SYSV\_DEFAULT

#define CPP\_SYSV\_DEFAULT "-D\_CALL\_AIX"

[Macro]

Define this macro if the driver program should find the library 'libgcc.a'. If you do not define this macro, the driver program will pass the argument '-lgcc' to tell the linker to do the search.

#### LINK\_GCC\_C\_SEQUENCE\_SPEC

[Macro]

The sequence in which libgcc and libc are specified to the linker. By default this is %G %L %G.

POST\_LINK\_SPEC [Macro]

Define this macro to add additional steps to be executed after linker. The default value of this macro is empty string.

#### LINK\_COMMAND\_SPEC

[Macro]

A C string constant giving the complete command line need to execute the linker. When you do this, you will need to update your port each time a change is made to the link command line within 'gcc.c'. Therefore, define this macro only if you need

to completely redefine the command line for invoking the linker and there is no other way to accomplish the effect you need. Overriding this macro may be avoidable by overriding LINK\_GCC\_C\_SEQUENCE\_SPEC instead.

#### bool TARGET\_ALWAYS\_STRIP\_DOTDOT

[Common Target Hook]

True if '..' components should always be removed from directory names computed relative to GCC's internal directories, false (default) if such components should be preserved and directory names containing them passed to other tools such as the linker.

#### MULTILIB\_DEFAULTS

[Macro]

Define this macro as a C expression for the initializer of an array of string to tell the driver program which options are defaults for this target and thus do not need to be handled specially when using MULTILIB\_OPTIONS.

Do not define this macro if MULTILIB\_OPTIONS is not defined in the target makefile fragment or if none of the options listed in MULTILIB\_OPTIONS are set by default. See Section 20.1 [Target Fragment], page 667.

#### RELATIVE\_PREFIX\_NOT\_LINKDIR

[Macro]

Define this macro to tell gcc that it should only translate a '-B' prefix into a '-L' linker option if the prefix indicates an absolute file name.

#### MD\_EXEC\_PREFIX

[Macro]

If defined, this macro is an additional prefix to try after STANDARD\_EXEC\_PREFIX. MD\_EXEC\_PREFIX is not searched when the compiler is built as a cross compiler. If you define MD\_EXEC\_PREFIX, then be sure to add it to the list of directories used to find the assembler in 'configure.ac'.

#### STANDARD\_STARTFILE\_PREFIX

[Macro]

Define this macro as a C string constant if you wish to override the standard choice of libdir as the default prefix to try when searching for startup files such as 'crt0.o'. STANDARD\_STARTFILE\_PREFIX is not searched when the compiler is built as a cross compiler.

#### STANDARD\_STARTFILE\_PREFIX\_1

[Macro]

Define this macro as a C string constant if you wish to override the standard choice of /lib as a prefix to try after the default prefix when searching for startup files such as 'crt0.o'. STANDARD\_STARTFILE\_PREFIX\_1 is not searched when the compiler is built as a cross compiler.

#### STANDARD\_STARTFILE\_PREFIX\_2

[Macro]

Define this macro as a C string constant if you wish to override the standard choice of /lib as yet another prefix to try after the default prefix when searching for startup files such as 'crt0.o'. STANDARD\_STARTFILE\_PREFIX\_2 is not searched when the compiler is built as a cross compiler.

#### MD\_STARTFILE\_PREFIX

[Macro]

If defined, this macro supplies an additional prefix to try after the standard prefixes. MD\_EXEC\_PREFIX is not searched when the compiler is built as a cross compiler.

#### MD\_STARTFILE\_PREFIX\_1

[Macro]

If defined, this macro supplies yet another prefix to try after the standard prefixes. It is not searched when the compiler is built as a cross compiler.

INIT\_ENVIRONMENT [Macro]

Define this macro as a C string constant if you wish to set environment variables for programs called by the driver, such as the assembler and loader. The driver passes the value of this macro to puterv to initialize the necessary environment variables.

LOCAL\_INCLUDE\_DIR

Define this macro as a C string constant if you wish to override the standard choice of '/usr/local/include' as the default prefix to try when searching for local header files. LOCAL\_INCLUDE\_DIR comes before NATIVE\_SYSTEM\_HEADER\_DIR (set in 'config.gcc', normally '/usr/include') in the search order.

Cross compilers do not search either '/usr/local/include' or its replacement.

#### NATIVE\_SYSTEM\_HEADER\_COMPONENT

[Macro]

[Macro]

The "component" corresponding to NATIVE\_SYSTEM\_HEADER\_DIR. See INCLUDE\_DEFAULTS, below, for the description of components. If you do not define this macro, no component is used.

INCLUDE\_DEFAULTS [Macro]

Define this macro if you wish to override the entire default search path for include files. For a native compiler, the default search path usually consists of GCC\_INCLUDE\_DIR, LOCAL\_INCLUDE\_DIR, GPLUSPLUS\_INCLUDE\_DIR, and NATIVE\_SYSTEM\_HEADER\_DIR. In addition, GPLUSPLUS\_INCLUDE\_DIR and GCC\_INCLUDE\_DIR are defined automatically by 'Makefile', and specify private search areas for GCC. The directory GPLUSPLUS\_INCLUDE\_DIR is used only for C++ programs.

The definition should be an initializer for an array of structures. Each array element should have four elements: the directory name (a string constant), the component name (also a string constant), a flag for C++-only directories, and a flag showing that the includes in the directory don't need to be wrapped in extern 'C' when compiling C++. Mark the end of the array with a null element.

The component name denotes what GNU package the include file is part of, if any, in all uppercase letters. For example, it might be 'GCC' or 'BINUTILS'. If the package is part of a vendor-supplied operating system, code the component name as '0'.

For example, here is the definition used for VAX/VMS:

Here is the order of prefixes tried for exec files:

1. Any prefixes specified by the user with '-B'.

- 2. The environment variable GCC\_EXEC\_PREFIX or, if GCC\_EXEC\_PREFIX is not set and the compiler has not been installed in the configure-time *prefix*, the location in which the compiler has actually been installed.
- 3. The directories specified by the environment variable COMPILER\_PATH.
- 4. The macro STANDARD\_EXEC\_PREFIX, if the compiler has been installed in the configured-time prefix.
- 5. The location '/usr/libexec/gcc/', but only if this is a native compiler.
- 6. The location '/usr/lib/gcc/', but only if this is a native compiler.
- 7. The macro MD\_EXEC\_PREFIX, if defined, but only if this is a native compiler.

Here is the order of prefixes tried for startfiles:

- 1. Any prefixes specified by the user with '-B'.
- 2. The environment variable GCC\_EXEC\_PREFIX or its automatically determined value based on the installed toolchain location.
- 3. The directories specified by the environment variable LIBRARY\_PATH (or port-specific name; native only, cross compilers do not use this).
- 4. The macro STANDARD\_EXEC\_PREFIX, but only if the toolchain is installed in the configured prefix or this is a native compiler.
- 5. The location '/usr/lib/gcc/', but only if this is a native compiler.
- 6. The macro MD\_EXEC\_PREFIX, if defined, but only if this is a native compiler.
- 7. The macro MD\_STARTFILE\_PREFIX, if defined, but only if this is a native compiler, or we have a target system root.
- 8. The macro MD\_STARTFILE\_PREFIX\_1, if defined, but only if this is a native compiler, or we have a target system root.
- 9. The macro STANDARD\_STARTFILE\_PREFIX, with any system modifications. If this path is relative it will be prefixed by GCC\_EXEC\_PREFIX and the machine suffix or STANDARD\_EXEC\_PREFIX and the machine suffix.
- 10. The macro STANDARD\_STARTFILE\_PREFIX\_1, but only if this is a native compiler, or we have a target system root. The default for this macro is '/lib/'.
- 11. The macro STANDARD\_STARTFILE\_PREFIX\_2, but only if this is a native compiler, or we have a target system root. The default for this macro is '/usr/lib/'.

# 18.3 Run-time Target Specification

Here are run-time target specifications.

#### TARGET\_CPU\_CPP\_BUILTINS ()

[Macro]

This function-like macro expands to a block of code that defines built-in preprocessor macros and assertions for the target CPU, using the functions builtin\_define, builtin\_define\_std and builtin\_assert. When the front end calls this macro it provides a trailing semicolon, and since it has finished command line option processing your code can use those results freely.

builtin\_assert takes a string in the form you pass to the command-line option '-A', such as cpu=mips, and creates the assertion. builtin\_define takes a string in the form accepted by option '-D' and unconditionally defines the macro.

builtin\_define\_std takes a string representing the name of an object-like macro. If it doesn't lie in the user's namespace, builtin\_define\_std defines it unconditionally. Otherwise, it defines a version with two leading underscores, and another version with two leading and trailing underscores, and defines the original only if an ISO standard was not requested on the command line. For example, passing unix defines \_\_unix, \_\_unix\_\_ and possibly unix; passing \_mips defines \_\_mips, \_\_mips\_\_ and possibly \_mips, and passing \_ABI64 defines only \_ABI64.

You can also test for the C dialect being compiled. The variable c\_language is set to one of clk\_c, clk\_cplusplus or clk\_objective\_c. Note that if we are preprocessing assembler, this variable will be clk\_c but the function-like macro preprocessing\_asm\_p() will return true, so you might want to check for that first. If you need to check for strict ANSI, the variable flag\_iso can be used. The function-like macro preprocessing\_trad\_p() can be used to check for traditional preprocessing.

#### TARGET\_OS\_CPP\_BUILTINS ()

[Macro]

Similarly to TARGET\_CPU\_CPP\_BUILTINS but this macro is optional and is used for the target operating system instead.

#### TARGET\_OBJFMT\_CPP\_BUILTINS ()

[Macro]

Similarly to TARGET\_CPU\_CPP\_BUILTINS but this macro is optional and is used for the target object format. 'elfos.h' uses this macro to define \_\_ELF\_\_, so you probably do not need to define it yourself.

#### extern int target\_flags

[Variable]

This variable is declared in 'options.h', which is included before any target-specific headers.

#### int TARGET\_DEFAULT\_TARGET\_FLAGS

[Common Target Hook]

This variable specifies the initial value of target\_flags. Its default setting is 0.

This hook is called whenever the user specifies one of the target-specific options described by the '.opt' definition files (see Chapter 8 [Options], page 119). It has the opportunity to do some option-specific processing and should return true if the option is valid. The default definition does nothing but return true.

decoded specifies the option and its arguments. opts and opts\_set are the gcc\_options structures to be used for storing option state, and loc is the location at which the option was passed (UNKNOWN\_LOCATION except for options passed via attributes).

# bool TARGET\_HANDLE\_C\_OPTION (size\_t code, const char \*arg, int [C Target Hook] value)

This target hook is called whenever the user specifies one of the target-specific C language family options described by the '.opt' definition files(see Chapter 8 [Options], page 119). It has the opportunity to do some option-specific processing and should return true if the option is valid. The arguments are like for TARGET\_HANDLE\_OPTION. The default definition does nothing but return false.

In general, you should use TARGET\_HANDLE\_OPTION to handle options. However, if processing an option requires routines that are only available in the C (and related language) front ends, then you should use TARGET\_HANDLE\_C\_OPTION instead.

Targets may provide a string object type that can be used within and between C, C++ and their respective Objective-C dialects. A string object might, for example, embed encoding and length information. These objects are considered opaque to the compiler and handled as references. An ideal implementation makes the composition of the string object match that of the Objective-C NSString (NXString for GNUStep), allowing efficient interworking between C-only and Objective-C code. If a target implements string objects then this hook should return a reference to such an object constructed from the normal 'C' string representation provided in string. At present, the hook is used by Objective-C only, to obtain a common-format string object when the target provides one.

# void TARGET\_OBJC\_DECLARE\_UNRESOLVED\_CLASS\_REFERENCE [C Target Hook] (const char \*classname)

Declare that Objective C class classname is referenced by the current TU.

#### 

Declare that Objective C class classname is defined by the current TU.

#### 

If a target implements string objects then this hook should return **true** if *stringref* is a valid reference to such an object.

#### 

If a target implements string objects then this hook should should provide a facility to check the function arguments in  $args\_list$  against the format specifiers in  $format\_arg$  where the type of  $format\_arg$  is one recognized as a valid string reference type.

# void TARGET\_OVERRIDE\_OPTIONS\_AFTER\_CHANGE (void) [Target Hook] This target function is similar to the hook TARGET\_OPTION\_OVERRIDE but is called when the optimize level is changed via an attribute or pragma or when it is reset at the end of the code affected by the attribute or pragma. It is not called at the beginning of compilation when TARGET\_OPTION\_OVERRIDE is called so if you want to perform these actions then, you should have TARGET\_OPTION\_OVERRIDE call TARGET\_OVERRIDE\_OPTIONS\_AFTER\_CHANGE.

#### C\_COMMON\_OVERRIDE\_OPTIONS

[Macro]

This is similar to the TARGET\_OPTION\_OVERRIDE hook but is only used in the C language frontends (C, Objective-C, C++, Objective-C++) and so can be used to alter option flag variables which only exist in those frontends.

# 

[Common Target Hook]

Some machines may desire to change what optimizations are performed for various optimization levels. This variable, if defined, describes options to enable at particular sets of optimization levels. These options are processed once just after the optimization level is determined and before the remainder of the command options have been parsed, so may be overridden by other options passed explicitly.

This processing is run once at program startup and when the optimization options are changed via #pragma GCC optimize or by using the optimize attribute.

# 

[Common Target Hook]

Set target-dependent initial values of fields in opts.

# SWITCHABLE\_TARGET

[Macro]

Some targets need to switch between substantially different subtargets during compilation. For example, the MIPS target has one subtarget for the traditional MIPS architecture and another for MIPS16. Source code can switch between these two subarchitectures using the mips16 and nomips16 attributes.

Such subtargets can differ in things like the set of available registers, the set of available instructions, the costs of various operations, and so on. GCC caches a lot of this type of information in global variables, and recomputing them for each subtarget takes a significant amount of time. The compiler therefore provides a facility for maintaining several versions of the global variables and quickly switching between them; see 'target-globals.h' for details.

Define this macro to 1 if your target needs this facility. The default is 0.

bool TARGET\_FLOAT\_EXCEPTIONS\_ROUNDING\_SUPPORTED\_P (void) [Target Hook] Returns true if the target supports IEEE 754 floating-point exceptions and rounding modes, false otherwise. This is intended to relate to the float and double types, but not necessarily long double. By default, returns true if the adddf3 instruction pattern is available and false otherwise, on the assumption that hardware floating point supports exceptions and rounding modes but software floating point does not.

# 18.4 Defining data structures for per-function information.

If the target needs to store information on a per-function basis, GCC provides a macro and a couple of variables to allow this. Note, just using statics to store the information is a bad idea, since GCC supports nested functions, so you can be halfway through encoding one function when another one comes along.

GCC defines a data structure called struct function which contains all of the data specific to an individual function. This structure contains a field called machine whose type is struct machine\_function \*, which can be used by targets to point to their own specific data.

If a target needs per-function specific data it should define the type struct machine\_function and also the macro INIT\_EXPANDERS. This macro should be used to initialize the function pointer init\_machine\_status. This pointer is explained below.

One typical use of per-function, target specific data is to create an RTX to hold the register containing the function's return address. This RTX can then be used to implement the \_\_builtin\_return\_address function, for level 0.

Note—earlier implementations of GCC used a single data area to hold all of the perfunction information. Thus when processing of a nested function began the old per-function data had to be pushed onto a stack, and when the processing was finished, it had to be popped off the stack. GCC used to provide function pointers called <code>save\_machine\_status</code> and <code>restore\_machine\_status</code> to handle the saving and restoring of the target specific information. Since the single data area approach is no longer used, these pointers are no longer supported.

INIT\_EXPANDERS [Macro]

Macro called to initialize any target specific information. This macro is called once per function, before generation of any RTL has begun. The intention of this macro is to allow the initialization of the function pointer init\_machine\_status.

# void (\*)(struct function \*) init\_machine\_status

[Variable]

If this function pointer is non-NULL it will be called once per function, before function compilation starts, in order to allow the target to perform any target specific initialization of the struct function structure. It is intended that this would be used to initialize the machine of that structure.

struct machine\_function structures are expected to be freed by GC. Generally, any memory that they reference must be allocated by using GC allocation, including the structure itself.

# 18.5 Storage Layout

Note that the definitions of the macros in this table which are sizes or alignments measured in bits do not need to be constant. They can be C expressions that refer to static variables, such as the target\_flags. See Section 18.3 [Run-time Target], page 486.

BITS\_BIG\_ENDIAN [Macro]

Define this macro to have the value 1 if the most significant bit in a byte has the lowest number; otherwise define it to have the value zero. This means that bit-field instructions count from the most significant bit. If the machine has no bit-field instructions, then this must still be defined, but it doesn't matter which value it is defined to. This macro need not be a constant.

This macro does not affect the way structure fields are packed into bytes or words; that is controlled by BYTES\_BIG\_ENDIAN.

# BYTES\_BIG\_ENDIAN [Macro]

Define this macro to have the value 1 if the most significant byte in a word has the lowest number. This macro need not be a constant.

#### WORDS BIG ENDIAN [Macro]

Define this macro to have the value 1 if, in a multiword object, the most significant word has the lowest number. This applies to both memory locations and registers; see REG\_WORDS\_BIG\_ENDIAN if the order of words in memory is not the same as the order in registers. This macro need not be a constant.

# REG\_WORDS\_BIG\_ENDIAN

[Macro]

On some machines, the order of words in a multiword object differs between registers in memory. In such a situation, define this macro to describe the order of words in a register. The macro WORDS\_BIG\_ENDIAN controls the order of words in memory.

# FLOAT\_WORDS\_BIG\_ENDIAN

[Macro]

Define this macro to have the value 1 if DFmode, XFmode or TFmode floating point numbers are stored in memory with the word containing the sign bit at the lowest address; otherwise define it to have the value 0. This macro need not be a constant.

You need not define this macro if the ordering is the same as for multi-word integers.

BITS\_PER\_WORD [Made

Number of bits in a word. If you do not define this macro, the default is BITS\_PER\_UNIT \* UNITS\_PER\_WORD.

# MAX\_BITS\_PER\_WORD

[Macro]

Maximum number of bits in a word. If this is undefined, the default is BITS\_PER\_WORD. Otherwise, it is the constant value that is the largest value that BITS\_PER\_WORD can have at run-time.

UNITS\_PER\_WORD [Macro]

Number of storage units in a word; normally the size of a general-purpose register, a power of two from 1 or 8.

#### MIN\_UNITS\_PER\_WORD

[Macro]

Minimum number of units in a word. If this is undefined, the default is UNITS\_PER\_ WORD. Otherwise, it is the constant value that is the smallest value that UNITS\_PER\_ WORD can have at run-time.

POINTER\_SIZE [Macro]

Width of a pointer, in bits. You must specify a value no wider than the width of Pmode. If it is not equal to the width of Pmode, you must define POINTERS\_EXTEND\_UNSIGNED. If you do not specify a value the default is BITS\_PER\_WORD.

# POINTERS\_EXTEND\_UNSIGNED

[Macro]

A C expression that determines how pointers should be extended from ptr\_mode to either Pmode or word\_mode. It is greater than zero if pointers should be zero-extended, zero if they should be sign-extended, and negative if some other sort of conversion is needed. In the last case, the extension is done by the target's ptr\_extend instruction.

You need not define this macro if the ptr\_mode, Pmode and word\_mode are all the same width.

# PROMOTE\_MODE (m, unsignedp, type)

[Macro]

A macro to update m and unsigned p when an object whose type is type and which has the specified mode and signedness is to be stored in a register. This macro is only called when type is a scalar type.

On most RISC machines, which only have operations that operate on a full register, define this macro to set m to word\_mode if m is an integer mode narrower than

BITS\_PER\_WORD. In most cases, only integer modes should be widened because wider-precision floating-point operations are usually more expensive than their narrower counterparts.

For most machines, the macro definition does not change *unsignedp*. However, some machines, have instructions that preferentially handle either signed or unsigned quantities of certain modes. For example, on the DEC Alpha, 32-bit loads from memory and 32-bit add instructions sign-extend the result to 64 bits. On such machines, set *unsignedp* according to which kind of extension is more efficient.

Do not define this macro if it would never modify m.

#### 

Return a value, with the same meaning as the C99 macro FLT\_EVAL\_METHOD that describes which excess precision should be applied. type is either EXCESS\_PRECISION\_TYPE\_IMPLICIT, EXCESS\_PRECISION\_TYPE\_FAST, or EXCESS\_PRECISION\_TYPE\_STANDARD. For EXCESS\_PRECISION\_TYPE\_IMPLICIT, the target should return which precision and range operations will be implictly evaluated in regardless of the excess precision explicitly added. For EXCESS\_PRECISION\_TYPE\_STANDARD and EXCESS\_PRECISION\_TYPE\_FAST, the target should return the explicit excess precision that should be added depending on the value set for '-fexcess-precision=[standard|fast]'. Note that unpredictable explicit excess precision does not make sense, so a target should never return FLT\_EVAL\_METHOD\_UNPREDICTABLE when type is EXCESS\_PRECISION\_TYPE\_STANDARD or EXCESS\_PRECISION\_TYPE\_FAST.

machine\_mode TARGET\_PROMOTE\_FUNCTION\_MODE (const\_tree type, [Target Hook] machine\_mode mode, int \*punsignedp, const\_tree funtype, int for\_return)
Like PROMOTE\_MODE, but it is applied to outgoing function arguments or function return values. The target hook should return the new mode and possibly change \*punsignedp if the promotion should change signedness. This function is called only for scalar or pointer types.

for\_return allows to distinguish the promotion of arguments and return values. If it is 1, a return value is being promoted and TARGET\_FUNCTION\_VALUE must perform the same promotions done here. If it is 2, the returned mode should be that of the register in which an incoming parameter is copied, or the outgoing result is computed; then the hook should return the same mode as promote\_mode, though the signedness may be different.

type can be NULL when promoting function arguments of libcalls.

The default is to not promote arguments and return values. You can also define the hook to default\_promote\_function\_mode\_always\_promote if you would like to apply the same rules given by PROMOTE\_MODE.

PARM\_BOUNDARY [Macro]

Normal alignment required for function parameters on the stack, in bits. All stack parameters receive at least this much alignment regardless of data type. On most machines, this is the same as the size of an integer.

STACK\_BOUNDARY [Macro]

Define this macro to the minimum alignment enforced by hardware for the stack pointer on this machine. The definition is a C expression for the desired alignment (measured in bits). This value is used as a default if PREFERRED\_STACK\_BOUNDARY is not defined. On most machines, this should be the same as PARM\_BOUNDARY.

# PREFERRED\_STACK\_BOUNDARY

[Macro]

Define this macro if you wish to preserve a certain alignment for the stack pointer, greater than what the hardware enforces. The definition is a C expression for the desired alignment (measured in bits). This macro must evaluate to a value equal to or larger than STACK\_BOUNDARY.

# INCOMING\_STACK\_BOUNDARY

[Macro]

Define this macro if the incoming stack boundary may be different from PREFERRED\_STACK\_BOUNDARY. This macro must evaluate to a value equal to or larger than STACK\_BOUNDARY.

# FUNCTION\_BOUNDARY

[Macro]

Alignment required for a function entry point, in bits.

# BIGGEST\_ALIGNMENT

[Macro]

Biggest alignment that any data type can require on this machine, in bits. Note that this is not the biggest alignment that is supported, just the biggest alignment that, when violated, may cause a fault.

# HOST\_WIDE\_INT TARGET\_ABSOLUTE\_BIGGEST\_ALIGNMENT

[Target Hook]

If defined, this target hook specifies the absolute biggest alignment that a type or variable can have on this machine, otherwise, BIGGEST\_ALIGNMENT is used.

# MALLOC\_ABI\_ALIGNMENT

[Macro]

Alignment, in bits, a C conformant malloc implementation has to provide. If not defined, the default value is BITS\_PER\_WORD.

# ATTRIBUTE\_ALIGNED\_VALUE

[Macro]

Alignment used by the \_\_attribute\_\_ ((aligned)) construct. If not defined, the default value is BIGGEST\_ALIGNMENT.

# MINIMUM\_ATOMIC\_ALIGNMENT

[Macro]

If defined, the smallest alignment, in bits, that can be given to an object that can be referenced in one operation, without disturbing any nearby object. Normally, this is BITS\_PER\_UNIT, but may be larger on machines that don't have byte or half-word store operations.

#### BIGGEST\_FIELD\_ALIGNMENT

[Macro]

Biggest alignment that any structure or union field can require on this machine, in bits. If defined, this overrides BIGGEST\_ALIGNMENT for structure and union fields only, unless the field alignment has been set by the \_\_attribute\_\_ ((aligned (n))) construct.

# ADJUST\_FIELD\_ALIGN (field, type, computed)

[Macro]

An expression for the alignment of a structure field field of type type if the alignment computed in the usual way (including applying of BIGGEST\_ALIGNMENT and BIGGEST\_FIELD\_ALIGNMENT to the alignment) is computed. It overrides alignment only if the field alignment has not been set by the \_\_attribute\_\_ ((aligned (n))) construct. Note that field may be NULL\_TREE in case we just query for the minimum alignment of a field of type type in structure context.

# MAX\_STACK\_ALIGNMENT

[Macro]

Biggest stack alignment guaranteed by the backend. Use this macro to specify the maximum alignment of a variable on stack.

If not defined, the default value is STACK\_BOUNDARY.

# MAX\_OFILE\_ALIGNMENT

[Macro]

Biggest alignment supported by the object file format of this machine. Use this macro to limit the alignment which can be specified using the <code>\_\_attribute\_\_</code> ((aligned (n))) construct for functions and objects with static storage duration. The alignment of automatic objects may exceed the object file format maximum up to the maximum supported by GCC. If not defined, the default value is <code>BIGGEST\_ALIGNMENT</code>.

On systems that use ELF, the default (in 'config/elfos.h') is the largest supported 32-bit ELF section alignment representable on a 32-bit host e.g. '((uint64\_t) 1 << 28) \* 8)'. On 32-bit ELF the largest supported section alignment in bits is '(0x80000000 \* 8)', but this is not representable on 32-bit hosts.

# HOST\_WIDE\_INT TARGET\_STATIC\_RTX\_ALIGNMENT (machine\_mode mode) [Target Hook]

This hook returns the preferred alignment in bits for a statically-allocated rtx, such as a constant pool entry. *mode* is the mode of the rtx. The default implementation returns 'GET\_MODE\_ALIGNMENT (*mode*)'.

# DATA\_ALIGNMENT (type, basic-align)

[Macro]

If defined, a C expression to compute the alignment for a variable in the static store. type is the data type, and basic-align is the alignment that the object would ordinarily have. The value of this macro is used instead of that alignment to align the object.

If this macro is not defined, then basic-align is used.

One use of this macro is to increase alignment of medium-size data to make it all fit in fewer cache lines. Another is to cause character arrays to be word-aligned so that strcpy calls that copy constants to character arrays can be done inline.

# DATA\_ABI\_ALIGNMENT (type, basic-align)

[Macro]

Similar to DATA\_ALIGNMENT, but for the cases where the ABI mandates some alignment increase, instead of optimization only purposes. E.g. AMD x86-64 psABI says that variables with array type larger than 15 bytes must be aligned to 16 byte boundaries.

If this macro is not defined, then basic-align is used.

# HOST\_WIDE\_INT TARGET\_CONSTANT\_ALIGNMENT (const\_tree constant, HOST\_WIDE\_INT basic\_align)

[Target Hook]

This hook returns the alignment in bits of a constant that is being placed in memory. constant is the constant and basic\_align is the alignment that the object would ordinarily have.

The default definition just returns basic\_align.

The typical use of this hook is to increase alignment for string constants to be word aligned so that strcpy calls that copy constants can be done inline. The function constant\_alignment\_word\_strings provides such a definition.

# LOCAL\_ALIGNMENT (type, basic-align)

[Macro]

If defined, a C expression to compute the alignment for a variable in the local store. type is the data type, and basic-align is the alignment that the object would ordinarily have. The value of this macro is used instead of that alignment to align the object.

If this macro is not defined, then basic-align is used.

One use of this macro is to increase alignment of medium-size data to make it all fit in fewer cache lines.

If the value of this macro has a type, it should be an unsigned type.

# HOST\_WIDE\_INT TARGET\_VECTOR\_ALIGNMENT (const\_tree type)

[Target Hook]

This hook can be used to define the alignment for a vector of type type, in order to comply with a platform ABI. The default is to require natural alignment for vector types. The alignment returned by this hook must be a power-of-two multiple of the default alignment of the vector element type.

# STACK\_SLOT\_ALIGNMENT (type, mode, basic-align)

[Macro]

If defined, a C expression to compute the alignment for stack slot. *type* is the data type, *mode* is the widest mode available, and *basic-align* is the alignment that the slot would ordinarily have. The value of this macro is used instead of that alignment to align the slot.

If this macro is not defined, then *basic-align* is used when *type* is NULL. Otherwise, LOCAL\_ALIGNMENT will be used.

This macro is to set alignment of stack slot to the maximum alignment of all possible modes which the slot may have.

If the value of this macro has a type, it should be an unsigned type.

# LOCAL\_DECL\_ALIGNMENT (dec1)

[Macro]

If defined, a C expression to compute the alignment for a local variable decl.

If this macro is not defined, then LOCAL\_ALIGNMENT (TREE\_TYPE (decl), DECL\_ALIGN (decl)) is used.

One use of this macro is to increase alignment of medium-size data to make it all fit in fewer cache lines.

If the value of this macro has a type, it should be an unsigned type.

# MINIMUM\_ALIGNMENT (exp, mode, align)

[Macro]

If defined, a C expression to compute the minimum required alignment for dynamic stack realignment purposes for exp (a type or decl), mode, assuming normal alignment align.

If this macro is not defined, then align will be used.

# EMPTY\_FIELD\_BOUNDARY

[Macro]

Alignment in bits to be given to a structure bit-field that follows an empty field such as int : 0;.

If PCC\_BITFIELD\_TYPE\_MATTERS is true, it overrides this macro.

# STRUCTURE\_SIZE\_BOUNDARY

[Macro]

Number of bits which any structure or union's size must be a multiple of. Each structure or union's size is rounded up to a multiple of this.

If you do not define this macro, the default is the same as BITS\_PER\_UNIT.

# STRICT\_ALIGNMENT

[Macro]

Define this macro to be the value 1 if instructions will fail to work if given data not on the nominal alignment. If instructions will merely go slower in that case, define this macro as 0.

# PCC\_BITFIELD\_TYPE\_MATTERS

[Macro]

Define this if you wish to imitate the way many other C compilers handle alignment of bit-fields and the structures that contain them.

The behavior is that the type written for a named bit-field (int, short, or other integer type) imposes an alignment for the entire structure, as if the structure really did contain an ordinary field of that type. In addition, the bit-field is placed within the structure so that it would fit within such a field, not crossing a boundary for it.

Thus, on most machines, a named bit-field whose type is written as **int** would not cross a four-byte boundary, and would force four-byte alignment for the whole structure. (The alignment used may not be four bytes; it is controlled by the other alignment parameters.)

An unnamed bit-field will not affect the alignment of the containing structure.

If the macro is defined, its definition should be a C expression; a nonzero value for the expression enables this behavior.

Note that if this macro is not defined, or its value is zero, some bit-fields may cross more than one alignment boundary. The compiler can support such references if there are 'insv', 'extv', and 'extzv' insns that can directly reference memory.

The other known way of making bit-fields work is to define STRUCTURE\_SIZE\_BOUNDARY as large as BIGGEST\_ALIGNMENT. Then every structure can be accessed with fullwords.

Unless the machine has bit-field instructions or you define STRUCTURE\_SIZE\_BOUNDARY that way, you must define PCC\_BITFIELD\_TYPE\_MATTERS to have a nonzero value.

If your aim is to make GCC use the same conventions for laying out bit-fields as are used by another compiler, here is how to investigate what the other compiler does. Compile and run this program:

```
struct foo1
  char x;
  char:0;
  char y;
};
struct foo2
  char x;
  int :0;
  char y;
};
main ()
 printf ("Size of foo1 is %d\n",
         sizeof (struct foo1));
 printf ("Size of foo2 is %d\n",
          sizeof (struct foo2));
  exit (0);
```

If this prints 2 and 5, then the compiler's behavior is what you would get from PCC\_BITFIELD\_TYPE\_MATTERS.

# BITFIELD\_NBYTES\_LIMITED

[Macro]

Like PCC\_BITFIELD\_TYPE\_MATTERS except that its effect is limited to aligning a bit-field within the structure.

# bool TARGET\_ALIGN\_ANON\_BITFIELD (void)

[Target Hook]

When PCC\_BITFIELD\_TYPE\_MATTERS is true this hook will determine whether unnamed bitfields affect the alignment of the containing structure. The hook should return true if the structure should inherit the alignment requirements of an unnamed bitfield's type.

# bool TARGET\_NARROW\_VOLATILE\_BITFIELD (void)

[Target Hook]

This target hook should return **true** if accesses to volatile bitfields should use the narrowest mode possible. It should return **false** if these accesses should use the bitfield container type.

The default is false.

# bool TARGET\_MEMBER\_TYPE\_FORCES\_BLK (const\_tree field,

[Target Hook]

machine\_mode mode)

Return true if a structure, union or array containing *field* should be accessed using BLKMODE.

If field is the only field in the structure, mode is its mode, otherwise mode is VOID-mode. mode is provided in the case where structures of one field would require the structure's mode to retain the field's mode.

Normally, this is not needed.

# ROUND\_TYPE\_ALIGN (type, computed, specified)

[Macro]

Define this macro as an expression for the alignment of a type (given by type as a tree node) if the alignment computed in the usual way is computed and the alignment explicitly specified was specified.

The default is to use specified if it is larger; otherwise, use the smaller of computed and  ${\tt BIGGEST\_ALIGNMENT}$ 

# MAX\_FIXED\_MODE\_SIZE

[Macro]

An integer expression for the size in bits of the largest integer machine mode that should actually be used. All integer machine modes of this size or smaller can be used for structures and unions with the appropriate sizes. If this macro is undefined, GET\_MODE\_BITSIZE (DImode) is assumed.

# STACK\_SAVEAREA\_MODE (save\_level)

[Macro]

If defined, an expression of type machine\_mode that specifies the mode of the save area operand of a save\_stack\_level named pattern (see Section 17.9 [Standard Names], page 392). save\_level is one of SAVE\_BLOCK, SAVE\_FUNCTION, or SAVE\_NONLOCAL and selects which of the three named patterns is having its mode specified.

You need not define this macro if it always returns Pmode. You would most commonly define this macro if the save\_stack\_level patterns need to support both a 32- and a 64-bit mode.

STACK\_SIZE\_MODE [Macro]

If defined, an expression of type machine\_mode that specifies the mode of the size increment operand of an allocate\_stack named pattern (see Section 17.9 [Standard Names], page 392).

You need not define this macro if it always returns word\_mode. You would most commonly define this macro if the allocate\_stack pattern needs to support both a 32- and a 64-bit mode.

# scalar\_int\_mode TARGET\_LIBGCC\_CMP\_RETURN\_MODE (void) [Target Hook]

This target hook should return the mode to be used for the return value of compare instructions expanded to libgcc calls. If not defined word\_mode is returned which is the right choice for a majority of targets.

scalar\_int\_mode TARGET\_LIBGCC\_SHIFT\_COUNT\_MODE (void) [Target Hook]

This target hook should return the mode to be used for the shift count operand of shift instructions expanded to libgcc calls. If not defined word\_mode is returned which is the right choice for a majority of targets.

# scalar\_int\_mode TARGET\_UNWIND\_WORD\_MODE (void) [Target Hook] Return machine mode to be used for \_Unwind\_Word type. The default is to use word\_mode.

bool TARGET\_MS\_BITFIELD\_LAYOUT\_P (const\_tree record\_type) [Target Hook]
This target hook returns true if bit-fields in the given record\_type are to be laid out
following the rules of Microsoft Visual C/C++, namely: (i) a bit-field won't share the
same storage unit with the previous bit-field if their underlying types have different

sizes, and the bit-field will be aligned to the highest alignment of the underlying types of itself and of the previous bit-field; (ii) a zero-sized bit-field will affect the alignment of the whole enclosing structure, even if it is unnamed; except that (iii) a zero-sized bit-field will be disregarded unless it follows another bit-field of nonzero size. If this hook returns true, other macros that control bit-field layout are ignored.

When a bit-field is inserted into a packed record, the whole size of the underlying type is used by one or more same-size adjacent bit-fields (that is, if its long:3, 32 bits is used in the record, and any additional adjacent long bit-fields are packed into the same chunk of 32 bits. However, if the size changes, a new field of that size is allocated). In an unpacked record, this is the same as using alignment, but not equivalent when packing.

If both MS bit-fields and '\_\_attribute\_\_((packed))' are used, the latter will take precedence. If '\_\_attribute\_\_((packed))' is used on a single field when MS bit-fields are in use, it will take precedence for that field, but the alignment of the rest of the structure may affect its placement.

# bool TARGET\_DECIMAL\_FLOAT\_SUPPORTED\_P (void)

[Target Hook]

Returns true if the target supports decimal floating point.

# bool TARGET\_FIXED\_POINT\_SUPPORTED\_P (void)

[Target Hook]

Returns true if the target supports fixed-point arithmetic.

# void TARGET\_EXPAND\_TO\_RTL\_HOOK (void)

[Target Hook]

This hook is called just before expansion into rtl, allowing the target to perform additional initializations or analysis before the expansion. For example, the rs6000 port uses it to allocate a scratch stack slot for use in copying SDmode values between memory and floating point registers whenever the function being expanded has any SDmode usage.

# void TARGET\_INSTANTIATE\_DECLS (void)

[Target Hook]

This hook allows the backend to perform additional instantiations on rtl that are not actually in any insns yet, but will be later.

# const char \* TARGET\_MANGLE\_TYPE (const\_tree type)

[Target Hook]

If your target defines any fundamental types, or any types your target uses should be mangled differently from the default, define this hook to return the appropriate encoding for these types as part of a C++ mangled name. The type argument is the tree structure representing the type to be mangled. The hook may be applied to trees which are not target-specific fundamental types; it should return NULL for all such types, as well as arguments it does not recognize. If the return value is not NULL, it must point to a statically-allocated string constant.

Target-specific fundamental types might be new fundamental types or qualified versions of ordinary fundamental types. Encode new fundamental types as 'u n name', where name is the name used for the type in source code, and n is the length of name in decimal. Encode qualified versions of ordinary types as 'U n name code', where name is the name used for the type qualifier in source code, n is the length of name as above, and code is the code used to represent the unqualified version of this type.

(See write\_builtin\_type in 'cp/mangle.c' for the list of codes.) In both cases the spaces are for clarity; do not include any spaces in your string.

This hook is applied to types prior to typedef resolution. If the mangled name for a particular type depends only on that type's main variant, you can perform typedef resolution yourself using TYPE\_MAIN\_VARIANT before mangling.

The default version of this hook always returns NULL, which is appropriate for a target that does not define any new fundamental types.

# 18.6 Layout of Source Language Data Types

These macros define the sizes and other characteristics of the standard basic data types used in programs being compiled. Unlike the macros in the previous section, these apply to specific features of C and related languages, rather than to fundamental aspects of storage layout.

INT\_TYPE\_SIZE [Macro]

A C expression for the size in bits of the type int on the target machine. If you don't define this, the default is one word.

SHORT\_TYPE\_SIZE [Macro]

A C expression for the size in bits of the type **short** on the target machine. If you don't define this, the default is half a word. (If this would be less than one storage unit, it is rounded up to one unit.)

LONG\_TYPE\_SIZE [Macro]

A C expression for the size in bits of the type long on the target machine. If you don't define this, the default is one word.

# ADA\_LONG\_TYPE\_SIZE [Macro]

On some machines, the size used for the Ada equivalent of the type long by a native Ada compiler differs from that used by C. In that situation, define this macro to be a C expression to be used for the size of that type. If you don't define this, the default is the value of LONG\_TYPE\_SIZE.

# LONG\_LONG\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type long long on the target machine. If you don't define this, the default is two words. If you want to support GNU Ada on your machine, the value of this macro must be at least 64.

CHAR\_TYPE\_SIZE [Macro]

A C expression for the size in bits of the type char on the target machine. If you don't define this, the default is BITS\_PER\_UNIT.

BOOL TYPE SIZE [Macro]

A C expression for the size in bits of the C++ type bool and C99 type \_Bool on the target machine. If you don't define this, and you probably shouldn't, the default is CHAR\_TYPE\_SIZE.

FLOAT\_TYPE\_SIZE [Macro]

A C expression for the size in bits of the type float on the target machine. If you don't define this, the default is one word.

# DOUBLE\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type double on the target machine. If you don't define this, the default is two words.

# LONG\_DOUBLE\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type long double on the target machine. If you don't define this, the default is two words.

# SHORT\_FRACT\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type short \_Fract on the target machine. If you don't define this, the default is BITS\_PER\_UNIT.

FRACT\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type \_Fract on the target machine. If you don't define this, the default is BITS\_PER\_UNIT \* 2.

# LONG\_FRACT\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type long \_Fract on the target machine. If you don't define this, the default is BITS\_PER\_UNIT \* 4.

# LONG\_LONG\_FRACT\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type long long \_Fract on the target machine. If you don't define this, the default is BITS\_PER\_UNIT \* 8.

# SHORT\_ACCUM\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type short \_Accum on the target machine. If you don't define this, the default is BITS\_PER\_UNIT \* 2.

# ACCUM\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type \_Accum on the target machine. If you don't define this, the default is BITS\_PER\_UNIT \* 4.

# LONG\_ACCUM\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type long \_Accum on the target machine. If you don't define this, the default is BITS\_PER\_UNIT \* 8.

# LONG\_LONG\_ACCUM\_TYPE\_SIZE

[Macro]

A C expression for the size in bits of the type long long \_Accum on the target machine. If you don't define this, the default is BITS\_PER\_UNIT \* 16.

# LIBGCC2\_GNU\_PREFIX

[Macro]

This macro corresponds to the TARGET\_LIBFUNC\_GNU\_PREFIX target hook and should be defined if that hook is overriden to be true. It causes function names in libgce to be changed to use a \_\_gnu\_ prefix for their name rather than the default \_\_. A port which uses this macro should also arrange to use 't-gnu-prefix' in the libgce 'config.host'.

# WIDEST\_HARDWARE\_FP\_SIZE

[Macro]

A C expression for the size in bits of the widest floating-point format supported by the hardware. If you define this macro, you must specify a value less than or equal to the value of LONG\_DOUBLE\_TYPE\_SIZE. If you do not define this macro, the value of LONG\_DOUBLE\_TYPE\_SIZE is the default.

# DEFAULT\_SIGNED\_CHAR

[Macro]

An expression whose value is 1 or 0, according to whether the type char should be signed or unsigned by default. The user can always override this default with the options '-fsigned-char' and '-funsigned-char'.

# bool TARGET\_DEFAULT\_SHORT\_ENUMS (void)

[Target Hook]

This target hook should return true if the compiler should give an enum type only as many bytes as it takes to represent the range of possible values of that type. It should return false if all enum types should be allocated like int.

The default is to return false.

SIZE\_TYPE [Macro]

A C expression for a string describing the name of the data type to use for size values. The typedef name size\_t is defined using the contents of the string.

The string can contain more than one keyword. If so, separate them with spaces, and write first any length keyword, then unsigned if appropriate, and finally int. The string must exactly match one of the data type names defined in the function c\_common\_nodes\_and\_builtins in the file 'c-family/c-common.c'. You may not omit int or change the order—that would cause the compiler to crash on startup.

If you don't define this macro, the default is "long unsigned int".

SIZETYPE [Macro]

GCC defines internal types (sizetype, ssizetype, bitsizetype and sbitsizetype) for expressions dealing with size. This macro is a C expression for a string describing the name of the data type from which the precision of sizetype is extracted.

The string has the same restrictions as SIZE\_TYPE string.

If you don't define this macro, the default is SIZE\_TYPE.

PTRDIFF\_TYPE [Macro]

A C expression for a string describing the name of the data type to use for the result of subtracting two pointers. The typedef name ptrdiff\_t is defined using the contents of the string. See SIZE\_TYPE above for more information.

If you don't define this macro, the default is "long int".

WCHAR\_TYPE [Macro]

A C expression for a string describing the name of the data type to use for wide characters. The typedef name wchar\_t is defined using the contents of the string. See SIZE\_TYPE above for more information.

If you don't define this macro, the default is "int".

WCHAR\_TYPE\_SIZE [Macro]

A C expression for the size in bits of the data type for wide characters. This is used in cpp, which cannot make use of WCHAR\_TYPE.

WINT\_TYPE [Macro]

A C expression for a string describing the name of the data type to use for wide characters passed to printf and returned from getwc. The typedef name wint\_t is defined using the contents of the string. See SIZE\_TYPE above for more information. If you don't define this macro, the default is "unsigned int".

INTMAX\_TYPE [Macro]

A C expression for a string describing the name of the data type that can represent any value of any standard or extended signed integer type. The typedef name <code>intmax\_t</code> is defined using the contents of the string. See <code>SIZE\_TYPE</code> above for more information.

If you don't define this macro, the default is the first of "int", "long int", or "long long int" that has as much precision as long long int.

UINTMAX\_TYPE [Macro]

A C expression for a string describing the name of the data type that can represent any value of any standard or extended unsigned integer type. The typedef name uintmax\_t is defined using the contents of the string. See SIZE\_TYPE above for more information.

If you don't define this macro, the default is the first of "unsigned int", "long unsigned int", or "long long unsigned int" that has as much precision as long long unsigned int.

SIG_ATOMIC_TYPE	[Macro]
INT8_TYPE	[Macro]
INT16_TYPE	[Macro]
INT32_TYPE	[Macro]
INT64_TYPE	[Macro]
UINT8_TYPE	[Macro]
UINT16_TYPE	[Macro]
UINT32_TYPE	[Macro]
UINT64_TYPE	[Macro]
INT_LEAST8_TYPE	[Macro]
INT_LEAST16_TYPE	[Macro]
INT_LEAST32_TYPE	[Macro]
INT_LEAST64_TYPE	[Macro]
UINT_LEAST8_TYPE	[Macro]
UINT_LEAST16_TYPE	[Macro]
UINT_LEAST32_TYPE	[Macro]
UINT_LEAST64_TYPE	[Macro]
INT_FAST8_TYPE	[Macro]
INT_FAST16_TYPE	[Macro]
INT_FAST32_TYPE	[Macro]
INT_FAST64_TYPE	[Macro]
UINT_FAST8_TYPE	[Macro]
UINT_FAST16_TYPE	[Macro]
UINT_FAST32_TYPE	[Macro]
UINT_FAST64_TYPE	[Macro]
INTPTR_TYPE	[Macro]
UINTPTR_TYPE	Macro

C expressions for the standard types sig\_atomic\_t, int8\_t, int16\_t, int32\_t, int64\_t, uint8\_t, uint16\_t, uint32\_t, uint64\_t, int\_least8\_t, int\_least16\_t, int\_least32\_t, int\_least64\_t, uint\_least8\_t, uint\_least16\_t, uint\_least32\_t, uint\_least64\_t, int\_fast8\_t, int\_fast16\_t, int\_fast32\_t, int\_fast64\_t,

uint\_fast8\_t, uint\_fast16\_t, uint\_fast32\_t, uint\_fast64\_t, intptr\_t, and uintptr\_t. See SIZE\_TYPE above for more information.

If any of these macros evaluates to a null pointer, the corresponding type is not supported; if GCC is configured to provide <stdint.h> in such a case, the header provided may not conform to C99, depending on the type in question. The defaults for all of these macros are null pointers.

# TARGET\_PTRMEMFUNC\_VBIT\_LOCATION

[Macro]

The C++ compiler represents a pointer-to-member-function with a struct that looks like:

```
struct {
  union {
    void (*fn)();
    ptrdiff_t vtable_index;
  };
  ptrdiff_t delta;
};
```

The C++ compiler must use one bit to indicate whether the function that will be called through a pointer-to-member-function is virtual. Normally, we assume that the low-order bit of a function pointer must always be zero. Then, by ensuring that the vtable\_index is odd, we can distinguish which variant of the union is in use. But, on some platforms function pointers can be odd, and so this doesn't work. In that case, we use the low-order bit of the delta field, and shift the remainder of the delta field to the left.

GCC will automatically make the right selection about where to store this bit using the FUNCTION\_BOUNDARY setting for your platform. However, some platforms such as ARM/Thumb have FUNCTION\_BOUNDARY set such that functions always start at even addresses, but the lowest bit of pointers to functions indicate whether the function at that address is in ARM or Thumb mode. If this is the case of your architecture, you should define this macro to ptrmemfunc\_vbit\_in\_delta.

In general, you should not have to define this macro. On architectures in which function addresses are always even, according to FUNCTION\_BOUNDARY, GCC will automatically define this macro to ptrmemfunc\_vbit\_in\_pfn.

# TARGET\_VTABLE\_USES\_DESCRIPTORS

[Macro]

Normally, the C++ compiler uses function pointers in vtables. This macro allows the target to change to use "function descriptors" instead. Function descriptors are found on targets for whom a function pointer is actually a small data structure. Normally the data structure consists of the actual code address plus a data pointer to which the function's data is relative.

If vtables are used, the value of this macro should be the number of words that the function descriptor occupies.

# TARGET\_VTABLE\_ENTRY\_ALIGN

[Macro]

By default, the vtable entries are void pointers, the so the alignment is the same as pointer alignment. The value of this macro specifies the alignment of the vtable entry in bits. It should be defined only when special alignment is necessary. \*/

# TARGET\_VTABLE\_DATA\_ENTRY\_DISTANCE

[Macro]

There are a few non-descriptor entries in the vtable at offsets below zero. If these entries must be padded (say, to preserve the alignment specified by TARGET\_VTABLE\_ENTRY\_ALIGN), set this to the number of words in each data entry.

# 18.7 Register Usage

This section explains how to describe what registers the target machine has, and how (in general) they can be used.

The description of which registers a specific instruction can use is done with register classes; see Section 18.8 [Register Classes], page 512. For information on using registers to access a stack frame, see Section 18.9.4 [Frame Registers], page 530. For passing values in registers, see Section 18.9.7 [Register Arguments], page 536. For returning values in registers, see Section 18.9.8 [Scalar Return], page 544.

# 18.7.1 Basic Characteristics of Registers

Registers have various characteristics.

# FIRST\_PSEUDO\_REGISTER

[Macro]

Number of hardware registers known to the compiler. They receive numbers 0 through FIRST\_PSEUDO\_REGISTER-1; thus, the first pseudo register's number really is assigned the number FIRST\_PSEUDO\_REGISTER.

FIXED\_REGISTERS [Macro]

An initializer that says which registers are used for fixed purposes all throughout the compiled code and are therefore not available for general allocation. These would include the stack pointer, the frame pointer (except on machines where that can be used as a general register when no frame pointer is needed), the program counter on machines where that is considered one of the addressable registers, and any other numbered register with a standard use.

This information is expressed as a sequence of numbers, separated by commas and surrounded by braces. The nth number is 1 if register n is fixed, 0 otherwise.

The table initialized from this macro, and the table initialized by the following one, may be overridden at run time either automatically, by the actions of the macro CONDITIONAL\_REGISTER\_USAGE, or by the user with the command options '-ffixed-reg', '-fcall-used-reg' and '-fcall-saved-reg'.

#### CALL\_USED\_REGISTERS

[Macro]

Like FIXED\_REGISTERS but has 1 for each register that is clobbered (in general) by function calls as well as for fixed registers. This macro therefore identifies the registers that are not available for general allocation of values that must live across function calls.

If a register has 0 in CALL\_USED\_REGISTERS, the compiler automatically saves it on function entry and restores it on function exit, if the register is used within the function.

Exactly one of CALL\_USED\_REGISTERS and CALL\_REALLY\_USED\_REGISTERS must be defined. Modern ports should define CALL\_REALLY\_USED\_REGISTERS.

# CALL\_REALLY\_USED\_REGISTERS

[Macro]

Like CALL\_USED\_REGISTERS except this macro doesn't require that the entire set of FIXED\_REGISTERS be included. (CALL\_USED\_REGISTERS must be a superset of FIXED\_REGISTERS).

Exactly one of CALL\_USED\_REGISTERS and CALL\_REALLY\_USED\_REGISTERS must be defined. Modern ports should define CALL\_REALLY\_USED\_REGISTERS.

# 

[Target Hook]

Return the ABI used by a function with type type; see the definition of predefined\_function\_abi for details of the ABI descriptor. Targets only need to define this hook if they support interoperability between several ABIs in the same translation unit.

#### 

This hook returns a description of the ABI used by the target of call instruction *insn*; see the definition of predefined\_function\_abi for details of the ABI descriptor. Only the global function insn\_callee\_abi should call this hook directly.

Targets only need to define this hook if they support interoperability between several ABIs in the same translation unit.

# bool TARGET\_HARD\_REGNO\_CALL\_PART\_CLOBBERED (unsigned int abi\_id, unsigned int regno, machine\_mode mode) [Target Hook]

ABIs usually specify that calls must preserve the full contents of a particular register, or that calls can alter any part of a particular register. This information is captured by the target macro CALL\_REALLY\_USED\_REGISTERS. However, some ABIs specify that calls must preserve certain bits of a particular register but can alter others. This hook should return true if this applies to at least one of the registers in '(reg:mode regno)', and if as a result the call would alter part of the mode value. For example, if a call preserves the low 32 bits of a 64-bit hard register regno but can clobber the upper 32 bits, this hook should return true for a 64-bit mode but false for a 32-bit mode.

The value of *abi\_id* comes from the **predefined\_function\_abi** structure that describes the ABI of the call; see the definition of the structure for more details. If (as is usual) the target uses the same ABI for all functions in a translation unit, *abi\_id* is always 0.

The default implementation returns false, which is correct for targets that don't have partly call-clobbered registers.

# const char \* TARGET\_GET\_MULTILIB\_ABI\_NAME (void)

[Target Hook]

This hook returns name of multilib ABI name.

# void TARGET\_CONDITIONAL\_REGISTER\_USAGE (void)

[Target Hook]

This hook may conditionally modify five variables fixed\_regs, call\_used\_regs, global\_regs, reg\_names, and reg\_class\_contents, to take into account any dependence of these register sets on target flags. The first three of these are of type char [] (interpreted as boolean vectors). global\_regs is a const char \*[], and reg\_class\_contents is a HARD\_REG\_SET. Before the macro is called,

fixed\_regs, call\_used\_regs, reg\_class\_contents, and reg\_names have been initialized from FIXED\_REGISTERS, CALL\_USED\_REGISTERS, REG\_CLASS\_CONTENTS, and REGISTER\_NAMES, respectively. global\_regs has been cleared, and any '-ffixed-reg', '-fcall-used-reg' and '-fcall-saved-reg' command options have been applied.

If the usage of an entire class of registers depends on the target flags, you may indicate this to GCC by using this macro to modify fixed\_regs and call\_used\_regs to 1 for each of the registers in the classes which should not be used by GCC. Also make define\_register\_constraints return NO\_REGS for constraints that shouldn't be used.

(However, if this class is not included in GENERAL\_REGS and all of the insn patterns whose constraints permit this class are controlled by target switches, then GCC will automatically avoid using these registers when the target switches are opposed to them.)

# INCOMING\_REGNO (out)

[Macro]

Define this macro if the target machine has register windows. This C expression returns the register number as seen by the called function corresponding to the register number *out* as seen by the calling function. Return *out* if register number *out* is not an outbound register.

# OUTGOING\_REGNO (in)

[Macro]

Define this macro if the target machine has register windows. This C expression returns the register number as seen by the calling function corresponding to the register number in as seen by the called function. Return in if register number in is not an inbound register.

# LOCAL\_REGNO (regno)

[Macro]

Define this macro if the target machine has register windows. This C expression returns true if the register is call-saved but is in the register window. Unlike most call-saved registers, such registers need not be explicitly restored on function exit or during non-local gotos.

PC\_REGNUM [Macro]

If the program counter has a register number, define this as that register number. Otherwise, do not define it.

# 18.7.2 Order of Allocation of Registers

Registers are allocated in order.

REG\_ALLOC\_ORDER [Macro]

If defined, an initializer for a vector of integers, containing the numbers of hard registers in the order in which GCC should prefer to use them (from most preferred to least).

If this macro is not defined, registers are used lowest numbered first (all else being equal).

One use of this macro is on machines where the highest numbered registers must always be saved and the save-multiple-registers instruction supports only sequences of consecutive registers. On such machines, define REG\_ALLOC\_ORDER to be an initializer that lists the highest numbered allocable register first.

# ADJUST\_REG\_ALLOC\_ORDER

[Macro]

A C statement (sans semicolon) to choose the order in which to allocate hard registers for pseudo-registers local to a basic block.

Store the desired register order in the array reg\_alloc\_order. Element 0 should be the register to allocate first; element 1, the next register; and so on.

The macro body should not assume anything about the contents of reg\_alloc\_order before execution of the macro.

On most machines, it is not necessary to define this macro.

# HONOR\_REG\_ALLOC\_ORDER

[Macro]

Normally, IRA tries to estimate the costs for saving a register in the prologue and restoring it in the epilogue. This discourages it from using call-saved registers. If a machine wants to ensure that IRA allocates registers in the order given by REG\_ALLOC\_ORDER even if some call-saved registers appear earlier than call-used ones, then define this macro as a C expression to nonzero. Default is 0.

# IRA\_HARD\_REGNO\_ADD\_COST\_MULTIPLIER (regno)

[Macro]

In some case register allocation order is not enough for the Integrated Register Allocator (IRA) to generate a good code. If this macro is defined, it should return a floating point value based on regno. The cost of using regno for a pseudo will be increased by approximately the pseudo's usage frequency times the value returned by this macro. Not defining this macro is equivalent to having it always return 0.0.

On most machines, it is not necessary to define this macro.

# 18.7.3 How Values Fit in Registers

This section discusses the macros that describe which kinds of values (specifically, which machine modes) each register can hold, and how many consecutive registers are needed for a given mode.

# unsigned int TARGET\_HARD\_REGNO\_NREGS (unsigned int regno, machine\_mode mode) [Target Hook]

This hook returns the number of consecutive hard registers, starting at register number regno, required to hold a value of mode mode. This hook must never return zero, even if a register cannot hold the requested mode - indicate that with TARGET\_HARD\_REGNO\_MODE\_OK and/or TARGET\_CAN\_CHANGE\_MODE\_CLASS instead.

The default definition returns the number of words in mode.

# HARD\_REGNO\_NREGS\_HAS\_PADDING (regno, mode)

[Macro]

A C expression that is nonzero if a value of mode *mode*, stored in memory, ends with padding that causes it to take up more space than in registers starting at register number *regno* (as determined by multiplying GCC's notion of the size of the register when containing this mode by the number of registers returned by TARGET\_HARD\_REGNO\_NREGS). By default this is zero.

For example, if a floating-point value is stored in three 32-bit registers but takes up 128 bits in memory, then this would be nonzero.

This macros only needs to be defined if there are cases where subreg\_get\_info would otherwise wrongly determine that a subreg can be represented by an offset to the register number, when in fact such a subreg would contain some of the padding not stored in registers and so not be representable.

# HARD\_REGNO\_NREGS\_WITH\_PADDING (regno, mode)

[Macro]

For values of regno and mode for which HARD\_REGNO\_NREGS\_HAS\_PADDING returns nonzero, a C expression returning the greater number of registers required to hold the value including any padding. In the example above, the value would be four.

# REGMODE\_NATURAL\_SIZE (mode)

Macro

Define this macro if the natural size of registers that hold values of mode mode is not the word size. It is a C expression that should give the natural size in bytes for the specified mode. It is used by the register allocator to try to optimize its results. This happens for example on SPARC 64-bit where the natural size of floating-point registers is still 32-bit.

# bool TARGET\_HARD\_REGNO\_MODE\_OK (unsigned int regno,

[Target Hook]

machine\_mode mode)

This hook returns true if it is permissible to store a value of mode mode in hard register number regno (or in several registers starting with that one). The default definition returns true unconditionally.

You need not include code to check for the numbers of fixed registers, because the allocation mechanism considers them to be always occupied.

On some machines, double-precision values must be kept in even/odd register pairs. You can implement that by defining this hook to reject odd register numbers for such modes.

The minimum requirement for a mode to be OK in a register is that the 'movmode' instruction pattern support moves between the register and other hard register in the same class and that moving a value into the register and back out not alter it.

Since the same instruction used to move word\_mode will work for all narrower integer modes, it is not necessary on any machine for this hook to distinguish between these modes, provided you define patterns 'movhi', etc., to take advantage of this. This is useful because of the interaction between TARGET\_HARD\_REGNO\_MODE\_OK and TARGET\_MODES\_TIEABLE\_P; it is very desirable for all integer modes to be tieable.

Many machines have special registers for floating point arithmetic. Often people assume that floating point machine modes are allowed only in floating point registers. This is not true. Any registers that can hold integers can safely *hold* a floating point machine mode, whether or not floating arithmetic can be done on it in those registers. Integer move instructions can be used to move the values.

On some machines, though, the converse is true: fixed-point machine modes may not go in floating registers. This is true if the floating registers normalize any value stored in them, because storing a non-floating value there would garble it. In this case, TARGET\_HARD\_REGNO\_MODE\_OK should reject fixed-point machine modes in floating registers. But if the floating registers do not automatically normalize, if you can store any bit pattern in one and retrieve it unchanged without a trap, then any machine mode may go in a floating register, so you can define this hook to say so.

The primary significance of special floating registers is rather that they are the registers acceptable in floating point arithmetic instructions. However, this is of no concern to TARGET\_HARD\_REGNO\_MODE\_OK. You handle it by writing the proper constraints for those instructions.

On some machines, the floating registers are especially slow to access, so that it is better to store a value in a stack frame than in such a register if floating point arithmetic is not being done. As long as the floating registers are not in class GENERAL\_REGS, they will not be used unless some pattern's constraint asks for one.

# HARD\_REGNO\_RENAME\_OK (from, to)

[Macro]

A C expression that is nonzero if it is OK to rename a hard register from to another hard register to.

One common use of this macro is to prevent renaming of a register to another register that is not saved by a prologue in an interrupt handler.

The default is always nonzero.

# bool TARGET\_MODES\_TIEABLE\_P (machine\_mode mode1,

[Target Hook]

machine\_mode mode2)

This hook returns true if a value of mode mode1 is accessible in mode mode2 without copying.

If  $TARGET_HARD_REGNO_MODE_OK$  (r, mode1) and  $TARGET_HARD_REGNO_MODE_OK$  (r, mode2) are always the same for any r, then  $TARGET_MODES_TIEABLE_P$  (mode1, mode2) should be true. If they differ for any r, you should define this hook to return false unless some other mechanism ensures the accessibility of the value in a narrower mode.

You should define this hook to return true in as many cases as possible since doing so will allow GCC to perform better register allocation. The default definition returns true unconditionally.

# bool TARGET\_HARD\_REGNO\_SCRATCH\_OK (unsigned int regno)

[Target Hook]

This target hook should return **true** if it is OK to use a hard register *regno* as scratch reg in peephole2.

One common use of this macro is to prevent using of a register that is not saved by a prologue in an interrupt handler.

The default version of this hook always returns true.

# AVOID\_CCMODE\_COPIES

[Macro]

Define this macro if the compiler should avoid copies to/from CCmode registers. You should only define this macro if support for copying to/from CCmode is incomplete.

# 18.7.4 Handling Leaf Functions

On some machines, a leaf function (i.e., one which makes no calls) can run more efficiently if it does not make its own register window. Often this means it is required to receive its arguments in the registers where they are passed by the caller, instead of the registers where they would normally arrive.

The special treatment for leaf functions generally applies only when other conditions are met; for example, often they may use only those registers for its own variables and

temporaries. We use the term "leaf function" to mean a function that is suitable for this special handling, so that functions with no calls are not necessarily "leaf functions".

GCC assigns register numbers before it knows whether the function is suitable for leaf function treatment. So it needs to renumber the registers in order to output a leaf function. The following macros accomplish this.

LEAF\_REGISTERS [Macro]

Name of a char vector, indexed by hard register number, which contains 1 for a register that is allowable in a candidate for leaf function treatment.

If leaf function treatment involves renumbering the registers, then the registers marked here should be the ones before renumbering—those that GCC would ordinarily allocate. The registers which will actually be used in the assembler code, after renumbering, should not be marked with 1 in this vector.

Define this macro only if the target machine offers a way to optimize the treatment of leaf functions.

# LEAF\_REG\_REMAP (regno)

[Macro]

A C expression whose value is the register number to which regno should be renumbered, when a function is treated as a leaf function.

If regno is a register number which should not appear in a leaf function before renumbering, then the expression should yield -1, which will cause the compiler to abort.

Define this macro only if the target machine offers a way to optimize the treatment of leaf functions, and registers need to be renumbered to do this.

TARGET\_ASM\_FUNCTION\_PROLOGUE and TARGET\_ASM\_FUNCTION\_EPILOGUE must usually treat leaf functions specially. They can test the C variable current\_function\_is\_leaf which is nonzero for leaf functions. current\_function\_is\_leaf is set prior to local register allocation and is valid for the remaining compiler passes. They can also test the C variable current\_function\_uses\_only\_leaf\_regs which is nonzero for leaf functions which only use leaf registers. current\_function\_uses\_only\_leaf\_regs is valid after all passes that modify the instructions have been run and is only useful if LEAF\_REGISTERS is defined.

# 18.7.5 Registers That Form a Stack

There are special features to handle computers where some of the "registers" form a stack. Stack registers are normally written by pushing onto the stack, and are numbered relative to the top of the stack.

Currently, GCC can only handle one group of stack-like registers, and they must be consecutively numbered. Furthermore, the existing support for stack-like registers is specific to the 80387 floating point coprocessor. If you have a new architecture that uses stack-like registers, you will need to do substantial work on 'reg-stack.c' and write your machine description to cooperate with it, as well as defining these macros.

STACK\_REGS [Macro]

Define this if the machine has any stack-like registers.

# STACK\_REG\_COVER\_CLASS

[Macro]

This is a cover class containing the stack registers. Define this if the machine has any stack-like registers.

FIRST\_STACK\_REG [Macro]

The number of the first stack-like register. This one is the top of the stack.

LAST\_STACK\_REG [Macro]

The number of the last stack-like register. This one is the bottom of the stack.

# 18.8 Register Classes

On many machines, the numbered registers are not all equivalent. For example, certain registers may not be allowed for indexed addressing; certain registers may not be allowed in some instructions. These machine restrictions are described to the compiler using register classes.

You define a number of register classes, giving each one a name and saying which of the registers belong to it. Then you can specify register classes that are allowed as operands to particular instruction patterns.

In general, each register will belong to several classes. In fact, one class must be named ALL\_REGS and contain all the registers. Another class must be named NO\_REGS and contain no registers. Often the union of two classes will be another class; however, this is not required.

One of the classes must be named GENERAL\_REGS. There is nothing terribly special about the name, but the operand constraint letters 'r' and 'g' specify this class. If GENERAL\_REGS is the same as ALL\_REGS, just define it as a macro which expands to ALL\_REGS.

Order the classes so that if class x is contained in class y then x has a lower class number than y.

The way classes other than GENERAL\_REGS are specified in operand constraints is through machine-dependent operand constraint letters. You can define such letters to correspond to various classes, then use them in operand constraints.

You must define the narrowest register classes for allocatable registers, so that each class either has no subclasses, or that for some mode, the move cost between registers within the class is cheaper than moving a register in the class to or from memory (see Section 18.16 [Costs], page 576).

You should define a class for the union of two classes whenever some instruction allows both classes. For example, if an instruction allows either a floating point (coprocessor) register or a general register for a certain operand, you should define a class FLOAT\_OR\_GENERAL\_REGS which includes both of them. Otherwise you will get suboptimal code, or even internal compiler errors when reload cannot find a register in the class computed via reg\_class\_subunion.

You must also specify certain redundant information about the register classes: for each class, which classes contain it and which ones are contained in it; for each pair of classes, the largest class contained in their union.

When a value occupying several consecutive registers is expected in a certain class, all the registers used must belong to that class. Therefore, register classes cannot be used to enforce a requirement for a register pair to start with an even-numbered register. The way to specify this requirement is with TARGET\_HARD\_REGNO\_MODE\_OK.

Register classes used for input-operands of bitwise-and or shift instructions have a special requirement: each such class must have, for each fixed-point machine mode, a subclass whose registers can transfer that mode to or from memory. For example, on some machines, the operations for single-byte values (QImode) are limited to certain registers. When this is so, each register class that is used in a bitwise-and or shift instruction must have a subclass consisting of registers from which single-byte values can be loaded or stored. This is so that PREFERRED\_RELOAD\_CLASS can always have a possible value to return.

enum reg\_class [Data type]

An enumerated type that must be defined with all the register class names as enumerated values. NO\_REGS must be first. ALL\_REGS must be the last register class, followed by one more enumerated value, LIM\_REG\_CLASSES, which is not a register class but rather tells how many classes there are.

Each register class has a number, which is the value of casting the class name to type int. The number serves as an index in many of the tables described below.

N\_REG\_CLASSES [Macro]

The number of distinct register classes, defined as follows:

#define N\_REG\_CLASSES (int) LIM\_REG\_CLASSES

REG\_CLASS\_NAMES [Macro]

An initializer containing the names of the register classes as C string constants. These names are used in writing some of the debugging dumps.

# REG\_CLASS\_CONTENTS [Macro]

An initializer containing the contents of the register classes, as integers which are bit masks. The *n*th integer specifies the contents of class n. The way the integer mask is interpreted is that register r is in the class if mask & (1 << r) is 1.

When the machine has more than 32 registers, an integer does not suffice. Then the integers are replaced by sub-initializers, braced groupings containing several integers. Each sub-initializer must be suitable as an initializer for the type HARD\_REG\_SET which is defined in 'hard-reg-set.h'. In this situation, the first integer in each sub-initializer corresponds to registers 0 through 31, the second integer to registers 32 through 63, and so on.

# REGNO\_REG\_CLASS (regno)

[Macro]

A C expression whose value is a register class containing hard register regno. In general there is more than one such class; choose a class which is minimal, meaning that no smaller class also contains the register.

BASE\_REG\_CLASS [Macro]

A macro whose definition is the name of the class to which a valid base register must belong. A base register is one used in an address which is the register value plus a displacement.

# MODE\_BASE\_REG\_CLASS (mode)

[Macro]

This is a variation of the BASE\_REG\_CLASS macro which allows the selection of a base register in a mode dependent manner. If *mode* is VOIDmode then it should return the same value as BASE\_REG\_CLASS.

# MODE\_BASE\_REG\_REG\_CLASS (mode)

[Macro]

A C expression whose value is the register class to which a valid base register must belong in order to be used in a base plus index register address. You should define this macro if base plus index addresses have different requirements than other base register uses.

#### 

A C expression whose value is the register class to which a valid base register for a memory reference in mode mode to address space address\_space must belong. outer\_code and index\_code define the context in which the base register occurs. outer\_code is the code of the immediately enclosing expression (MEM for the top level of an address, ADDRESS for something that occurs in an address\_operand). index\_code is the code of the corresponding index expression if outer\_code is PLUS; SCRATCH otherwise.

INDEX\_REG\_CLASS [Macro]

A macro whose definition is the name of the class to which a valid index register must belong. An index register is one used in an address where its value is either multiplied by a scale factor or added to another register (as well as added to a displacement).

# REGNO\_OK\_FOR\_BASE\_P (num)

[Macro]

A C expression which is nonzero if register number *num* is suitable for use as a base register in operand addresses.

# REGNO\_MODE\_OK\_FOR\_BASE\_P (num, mode)

[Macro]

A C expression that is just like REGNO\_OK\_FOR\_BASE\_P, except that that expression may examine the mode of the memory reference in *mode*. You should define this macro if the mode of the memory reference affects whether a register may be used as a base register. If you define this macro, the compiler will use it instead of REGNO\_OK\_FOR\_BASE\_P. The mode may be VOIDmode for addresses that appear outside a MEM, i.e., as an address\_operand.

# REGNO\_MODE\_OK\_FOR\_REG\_BASE\_P (num, mode)

[Macro]

A C expression which is nonzero if register number *num* is suitable for use as a base register in base plus index operand addresses, accessing memory in mode *mode*. It may be either a suitable hard register or a pseudo register that has been allocated such a hard register. You should define this macro if base plus index addresses have different requirements than other base register uses.

Use of this macro is deprecated; please use the more general REGNO\_MODE\_CODE\_OK\_FOR\_BASE\_P.

# REGNO\_MODE\_CODE\_OK\_FOR\_BASE\_P (num, mode, address\_space, outer\_code, index\_code) [Macro]

A C expression which is nonzero if register number num is suitable for use as a base register in operand addresses, accessing memory in mode mode in address space address\_space. This is similar to REGNO\_MODE\_OK\_FOR\_BASE\_P, except that that expression may examine the context in which the register appears in the memory reference. outer\_code is the code of the immediately enclosing expression (MEM if at the top level of the address, ADDRESS for something that occurs in an address\_operand). index\_code is the code of the corresponding index expression if outer\_code is PLUS; SCRATCH otherwise. The mode may be VOIDmode for addresses that appear outside a MEM, i.e., as an address\_operand.

# REGNO\_OK\_FOR\_INDEX\_P (num)

[Macro]

A C expression which is nonzero if register number *num* is suitable for use as an index register in operand addresses. It may be either a suitable hard register or a pseudo register that has been allocated such a hard register.

The difference between an index register and a base register is that the index register may be scaled. If an address involves the sum of two registers, neither one of them scaled, then either one may be labeled the "base" and the other the "index"; but whichever labeling is used must fit the machine's constraints of which registers may serve in each capacity. The compiler will try both labelings, looking for one that is valid, and will reload one or both registers only if neither labeling works.

#### 

A target hook that places additional preference on the register class to use when it is necessary to rename a register in class relass to another class, or perhaps NO\_REGS, if no preferred register class is found or hook preferred\_rename\_class is not implemented. Sometimes returning a more restrictive class makes better code. For example, on ARM, thumb-2 instructions using LO\_REGS may be smaller than instructions using GENERIC\_REGS. By returning LO\_REGS from preferred\_rename\_class, code size can be reduced.

#### 

A target hook that places additional restrictions on the register class to use when it is necessary to copy value x into a register in class rclass. The value is a register class; perhaps rclass, or perhaps another, smaller class.

The default version of this hook always returns value of rclass argument.

Sometimes returning a more restrictive class makes better code. For example, on the 68000, when x is an integer constant that is in range for a 'moveq' instruction, the value of this macro is always DATA\_REGS as long as rclass includes the data registers. Requiring a data register guarantees that a 'moveq' will be used.

One case where TARGET\_PREFERRED\_RELOAD\_CLASS must not return rclass is if x is a legitimate constant which cannot be loaded into some register class. By returning NO\_REGS you can force x into a memory location. For example, rs6000 can load immediate values into general-purpose registers, but does not have an instruction

for loading an immediate value into a floating-point register, so TARGET\_PREFERRED\_RELOAD\_CLASS returns NO\_REGS when x is a floating-point constant. If the constant can't be loaded into any kind of register, code generation will be better if TARGET\_LEGITIMATE\_CONSTANT\_P makes the constant illegitimate instead of using TARGET\_PREFERRED\_RELOAD\_CLASS.

If an insn has pseudos in it after register allocation, reload will go through the alternatives and call repeatedly TARGET\_PREFERRED\_RELOAD\_CLASS to find the best one. Returning NO\_REGS, in this case, makes reload add a ! in front of the constraint: the x86 back-end uses this feature to discourage usage of 387 registers when math is done in the SSE registers (and vice versa).

# PREFERRED\_RELOAD\_CLASS (x, class)

[Macro]

A C expression that places additional restrictions on the register class to use when it is necessary to copy value x into a register in class class. The value is a register class; perhaps class, or perhaps another, smaller class. On many machines, the following definition is safe:

# #define PREFERRED\_RELOAD\_CLASS(X,CLASS) CLASS

Sometimes returning a more restrictive class makes better code. For example, on the 68000, when x is an integer constant that is in range for a 'moveq' instruction, the value of this macro is always DATA\_REGS as long as class includes the data registers. Requiring a data register guarantees that a 'moveq' will be used.

One case where PREFERRED\_RELOAD\_CLASS must not return class is if x is a legitimate constant which cannot be loaded into some register class. By returning NO\_REGS you can force x into a memory location. For example, rs6000 can load immediate values into general-purpose registers, but does not have an instruction for loading an immediate value into a floating-point register, so PREFERRED\_RELOAD\_CLASS returns NO\_REGS when x is a floating-point constant. If the constant cannot be loaded into any kind of register, code generation will be better if TARGET\_LEGITIMATE\_CONSTANT\_P makes the constant illegitimate instead of using TARGET\_PREFERRED\_RELOAD\_CLASS.

If an insn has pseudos in it after register allocation, reload will go through the alternatives and call repeatedly PREFERRED\_RELOAD\_CLASS to find the best one. Returning NO\_REGS, in this case, makes reload add a ! in front of the constraint: the x86 backend uses this feature to discourage usage of 387 registers when math is done in the SSE registers (and vice versa).

#### 

Like TARGET\_PREFERRED\_RELOAD\_CLASS, but for output reloads instead of input reloads.

The default version of this hook always returns value of rclass argument.

You can also use TARGET\_PREFERRED\_OUTPUT\_RELOAD\_CLASS to discourage reload from using some alternatives, like TARGET\_PREFERRED\_RELOAD\_CLASS.

# LIMIT\_RELOAD\_CLASS (mode, class)

Macro

A C expression that places additional restrictions on the register class to use when it is necessary to be able to hold a value of mode mode in a reload register for which class class would ordinarily be used.

Unlike PREFERRED\_RELOAD\_CLASS, this macro should be used when there are certain modes that simply cannot go in certain reload classes.

The value is a register class; perhaps class, or perhaps another, smaller class.

Don't define this macro unless the target machine has limitations which require the macro to do something nontrivial.

Many machines have some registers that cannot be copied directly to or from memory or even from other types of registers. An example is the 'MQ' register, which on most machines, can only be copied to or from general registers, but not memory. Below, we shall be using the term 'intermediate register' when a move operation cannot be performed directly, but has to be done by copying the source into the intermediate register first, and then copying the intermediate register to the destination. An intermediate register always has the same mode as source and destination. Since it holds the actual value being copied, reload might apply optimizations to re-use an intermediate register and eliding the copy from the source when it can determine that the intermediate register still holds the required value.

Another kind of secondary reload is required on some machines which allow copying all registers to and from memory, but require a scratch register for stores to some memory locations (e.g., those with symbolic address on the RT, and those with certain symbolic address on the SPARC when compiling PIC). Scratch registers need not have the same mode as the value being copied, and usually hold a different value than that being copied. Special patterns in the md file are needed to describe how the copy is performed with the help of the scratch register; these patterns also describe the number, register class(es) and mode(s) of the scratch register(s).

In some cases, both an intermediate and a scratch register are required.

For input reloads, this target hook is called with nonzero  $in_p$ , and x is an rtx that needs to be copied to a register of class  $reload_class$  in  $reload_mode$ . For output reloads, this target hook is called with zero  $in_p$ , and a register of class  $reload_class$  needs to be copied to rtx x in  $reload_mode$ .

If copying a register of reload\_class from/to x requires an intermediate register, the hook secondary\_reload should return the register class required for this intermediate register. If no intermediate register is required, it should return NO\_REGS. If more than one intermediate register is required, describe the one that is closest in the copy chain to the reload register.

If scratch registers are needed, you also have to describe how to perform the copy from/to the reload register to/from this closest intermediate register. Or if no intermediate register is required, but still a scratch register is needed, describe the copy from/to the reload register to/from the reload operand x.

You do this by setting sri->icode to the instruction code of a pattern in the md file which performs the move. Operands 0 and 1 are the output and input of this copy, respectively. Operands from operand 2 onward are for scratch operands. These scratch operands must have a mode, and a single-register-class output constraint.

When an intermediate register is used, the **secondary\_reload** hook will be called again to determine how to copy the intermediate register to/from the reload operand x, so your hook must also have code to handle the register class of the intermediate operand.

x might be a pseudo-register or a subreg of a pseudo-register, which could either be in a hard register or in memory. Use true\_regnum to find out; it will return -1 if the pseudo is in memory and the hard register number if it is in a register.

Scratch operands in memory (constraint "=m" / "=&m") are currently not supported. For the time being, you will have to continue to use TARGET\_SECONDARY\_MEMORY\_NEEDED for that purpose.

copy\_cost also uses this target hook to find out how values are copied. If you want it to include some extra cost for the need to allocate (a) scratch register(s), set sri>extra\_cost to the additional cost. Or if two dependent moves are supposed to have a lower cost than the sum of the individual moves due to expected fortuitous scheduling and/or special forwarding logic, you can set sri->extra\_cost to a negative amount.

```
SECONDARY_RELOAD_CLASS (class, mode, x)[Macro]SECONDARY_INPUT_RELOAD_CLASS (class, mode, x)[Macro]SECONDARY_OUTPUT_RELOAD_CLASS (class, mode, x)[Macro]
```

These macros are obsolete, new ports should use the target hook TARGET\_SECONDARY\_RELOAD instead.

These are obsolete macros, replaced by the TARGET\_SECONDARY\_RELOAD target hook. Older ports still define these macros to indicate to the reload phase that it may need to allocate at least one register for a reload in addition to the register to contain the data. Specifically, if copying x to a register class in mode requires an intermediate register, you were supposed to define SECONDARY\_INPUT\_RELOAD\_CLASS to return the largest register class all of whose registers can be used as intermediate registers or scratch registers.

If copying a register class in mode to x requires an intermediate or scratch register, SECONDARY\_OUTPUT\_RELOAD\_CLASS was supposed to be defined be defined to return the largest register class required. If the requirements for input and output reloads were the same, the macro SECONDARY\_RELOAD\_CLASS should have been used instead of defining both macros identically.

The values returned by these macros are often GENERAL\_REGS. Return  $NO_REGS$  if no spare register is needed; i.e., if x can be directly copied to or from a register of class in mode without requiring a scratch register. Do not define this macro if it would always return  $NO_REGS$ .

If a scratch register is required (either with or without an intermediate register), you were supposed to define patterns for 'reload\_inm' or 'reload\_outm', as required (see Section 17.9 [Standard Names], page 392. These patterns, which were normally implemented with a define\_expand, should be similar to the 'movm' patterns, except that operand 2 is the scratch register.

These patterns need constraints for the reload register and scratch register that contain a single register class. If the original reload register (whose class is *class*) can meet the constraint given in the pattern, the value returned by these macros is used

for the class of the scratch register. Otherwise, two additional reload registers are required. Their classes are obtained from the constraints in the insu pattern.

x might be a pseudo-register or a subreg of a pseudo-register, which could either be in a hard register or in memory. Use true\_regnum to find out; it will return -1 if the pseudo is in memory and the hard register number if it is in a register.

These macros should not be used in the case where a particular class of registers can only be copied to memory and not to another class of registers. In that case, secondary reload registers are not needed and would not be helpful. Instead, a stack location must be used to perform the copy and the movm pattern should use memory as an intermediate storage. This case often occurs between floating-point and general registers.

# bool TARGET\_SECONDARY\_MEMORY\_NEEDED (machine\_mode mode, reg\_class\_t class1, reg\_class\_t class2) [Target Hook]

Certain machines have the property that some registers cannot be copied to some other registers without using memory. Define this hook on those machines to return true if objects of mode m in registers of class1 can only be copied to registers of class2 by storing a register of class1 into memory and loading that memory location into a register of class2. The default definition returns false for all inputs.

# SECONDARY\_MEMORY\_NEEDED\_RTX (mode)

[Macro]

Normally when TARGET\_SECONDARY\_MEMORY\_NEEDED is defined, the compiler allocates a stack slot for a memory location needed for register copies. If this macro is defined, the compiler instead uses the memory location defined by this macro.

Do not define this macro if you do not define TARGET\_SECONDARY\_MEMORY\_NEEDED.

# machine\_mode TARGET\_SECONDARY\_MEMORY\_NEEDED\_MODE [Target Hook] (machine\_mode mode)

If TARGET\_SECONDARY\_MEMORY\_NEEDED tells the compiler to use memory when moving between two particular registers of mode *mode*, this hook specifies the mode that the memory should have.

The default depends on TARGET\_LRA\_P. Without LRA, the default is to use a word-sized mode for integral modes that are smaller than a a word. This is right thing to do on most machines because it ensures that all bits of the register are copied and prevents accesses to the registers in a narrower mode, which some machines prohibit for floating-point registers.

However, this default behavior is not correct on some machines, such as the DEC Alpha, that store short integers in floating-point registers differently than in integer registers. On those machines, the default widening will not work correctly and you must define this hook to suppress that widening in some cases. See the file 'alpha.c' for details.

With LRA, the default is to use mode unmodified.

# void TARGET\_SELECT\_EARLY\_REMAT\_MODES (sbitmap modes) [Target Hook]

On some targets, certain modes cannot be held in registers around a standard ABI call and are relatively expensive to spill to the stack. The early rematerialization

pass can help in such cases by aggressively recomputing values after calls, so that they don't need to be spilled.

This hook returns the set of such modes by setting the associated bits in *modes*. The default implementation selects no modes, which has the effect of disabling the early rematerialization pass.

# bool TARGET\_CLASS\_LIKELY\_SPILLED\_P (reg\_class\_t rclass) [Target Hook]

A target hook which returns **true** if pseudos that have been assigned to registers of class rclass would likely be spilled because registers of rclass are needed for spill registers.

The default version of this target hook returns **true** if *rclass* has exactly one register and **false** otherwise. On most machines, this default should be used. For generally register-starved machines, such as i386, or machines with right register constraints, such as SH, this hook can be used to avoid excessive spilling.

This hook is also used by some of the global intra-procedural code transformations to throtle code motion, to avoid increasing register pressure.

# unsigned char TARGET\_CLASS\_MAX\_NREGS (reg\_class\_t rclass, machine\_mode mode) [Target Hook]

A target hook returns the maximum number of consecutive registers of class rclass needed to hold a value of mode mode.

This is closely related to the macro TARGET\_HARD\_REGNO\_NREGS. In fact, the value returned by TARGET\_CLASS\_MAX\_NREGS (rclass, mode) target hook should be the maximum value of TARGET\_HARD\_REGNO\_NREGS (regno, mode) for all regno values in the class rclass.

This target hook helps control the handling of multiple-word values in the reload pass. The default version of this target hook returns the size of *mode* in words.

# CLASS\_MAX\_NREGS (class, mode)

[Macro]

A C expression for the maximum number of consecutive registers of class class needed to hold a value of mode *mode*.

This is closely related to the macro TARGET\_HARD\_REGNO\_NREGS. In fact, the value of the macro CLASS\_MAX\_NREGS (class, mode) should be the maximum value of TARGET\_HARD\_REGNO\_NREGS (regno, mode) for all regno values in the class class.

This macro helps control the handling of multiple-word values in the reload pass.

# bool TARGET\_CAN\_CHANGE\_MODE\_CLASS (machine\_mode from, machine\_mode to, reg\_class\_t rclass) [Target Hook]

This hook returns true if it is possible to bitcast values held in registers of class relass from mode from to mode to and if doing so preserves the low-order bits that are common to both modes. The result is only meaningful if relass has registers that can hold both from and to. The default implementation returns true.

As an example of when such bitcasting is invalid, loading 32-bit integer or floating-point objects into floating-point registers on Alpha extends them to 64 bits. Therefore loading a 64-bit object and then storing it as a 32-bit object does not store the low-order 32 bits, as would be the case for a normal register. Therefore, 'alpha.h' defines TARGET\_CAN\_CHANGE\_MODE\_CLASS to return:

```
(GET_MODE_SIZE (from) == GET_MODE_SIZE (to)
|| !reg_classes_intersect_p (FLOAT_REGS, rclass))
```

Even if storing from a register in mode to would be valid, if both from and raw\_reg\_mode for rclass are wider than word\_mode, then we must prevent to narrowing the mode. This happens when the middle-end assumes that it can load or store pieces of an N-word pseudo, and that the pseudo will eventually be allocated to N word\_mode hard registers. Failure to prevent this kind of mode change will result in the entire raw\_reg\_mode being modified instead of the partial value that the middle-end intended.

#### 

A target hook which can change allocno class for given pseudo from allocno and best class calculated by IRA.

The default version of this target hook always returns given class.

# bool TARGET\_LRA\_P (void)

[Target Hook]

A target hook which returns true if we use LRA instead of reload pass. The default version of this target hook returns true. New ports should use LRA, and existing ports are encouraged to convert.

# int TARGET\_REGISTER\_PRIORITY (int)

[Target Hook]

A target hook which returns the register priority number to which the register hard\_regno belongs to. The bigger the number, the more preferable the hard register usage (when all other conditions are the same). This hook can be used to prefer some hard register over others in LRA. For example, some x86-64 register usage needs additional prefix which makes instructions longer. The hook can return lower priority number for such registers make them less favorable and as result making the generated code smaller. The default version of this target hook returns always zero.

# bool TARGET\_REGISTER\_USAGE\_LEVELING\_P (void)

[Target Hook]

A target hook which returns true if we need register usage leveling. That means if a few hard registers are equally good for the assignment, we choose the least used hard register. The register usage leveling may be profitable for some targets. Don't use the usage leveling for targets with conditional execution or targets with big register files as it hurts if-conversion and cross-jumping optimizations. The default version of this target hook returns always false.

# bool TARGET\_DIFFERENT\_ADDR\_DISPLACEMENT\_P (void)

[Target Hook]

A target hook which returns true if an address with the same structure can have different maximal legitimate displacement. For example, the displacement can depend on memory mode or on operand combinations in the insn. The default version of this target hook returns always false.

# bool TARGET\_CANNOT\_SUBSTITUTE\_MEM\_EQUIV\_P (rtx subst) [Target Hook]

A target hook which returns true if subst can't substitute safely pseudos with equivalent memory values during register allocation. The default version of this target hook returns false. On most machines, this default should be used. For generally machines with non orthogonal register usage for addressing, such as SH, this hook can be used to avoid excessive spilling.

#### 

This hook tries to split address offset orig\_offset into two parts: one that should be added to the base address to create a local anchor point, and an additional offset that can be applied to the anchor to address a value of mode mode. The idea is that the local anchor could be shared by other accesses to nearby locations.

The hook returns true if it succeeds, storing the offset of the anchor from the base in offset1 and the offset of the final address from the anchor in offset2. The default implementation returns false.

- reg\_class\_t TARGET\_SPILL\_CLASS (reg\_class\_t, machine\_mode) [Target Hook]
  This hook defines a class of registers which could be used for spilling pseudos of the given mode and class, or NO\_REGS if only memory should be used. Not defining this hook is equivalent to returning NO\_REGS for all inputs.
- bool TARGET\_ADDITIONAL\_ALLOCNO\_CLASS\_P (reg\_class\_t) [Target Hook] This hook should return true if given class of registers should be an allocno class in any way. Usually RA uses only one register class from all classes containing the same register set. In some complicated cases, you need to have two or more such classes as allocno ones for RA correct work. Not defining this hook is equivalent to returning false for all inputs.
- scalar\_int\_mode TARGET\_CSTORE\_MODE (enum insn\_code icode) [Target Hook] This hook defines the machine mode to use for the boolean result of conditional store patterns. The ICODE argument is the instruction code for the estore being performed. Not definiting this hook is the same as accepting the mode encoded into operand 0 of the estore expander patterns.

#### 

A target hook which lets a backend compute the set of pressure classes to be used by those optimization passes which take register pressure into account, as opposed to letting IRA compute them. It returns the number of register classes stored in the array pressure\_classes.

# 18.9 Stack Layout and Calling Conventions

This describes the stack layout and calling conventions.

# 18.9.1 Basic Stack Layout

Here is the basic stack layout.

# STACK\_GROWS\_DOWNWARD

[Macro]

Define this macro to be true if pushing a word onto the stack moves the stack pointer to a smaller address, and false otherwise.

STACK\_PUSH\_CODE [Macro]

This macro defines the operation used when something is pushed on the stack. In RTL, a push operation will be (set (mem (STACK\_PUSH\_CODE (reg sp))) ...)

The choices are PRE\_DEC, POST\_DEC, PRE\_INC, and POST\_INC. Which of these is correct depends on the stack direction and on whether the stack pointer points to the last item on the stack or whether it points to the space for the next item on the stack.

The default is PRE\_DEC when STACK\_GROWS\_DOWNWARD is true, which is almost always right, and PRE\_INC otherwise, which is often wrong.

# FRAME\_GROWS\_DOWNWARD

[Macro]

Define this macro to nonzero value if the addresses of local variable slots are at negative offsets from the frame pointer.

# ARGS\_GROW\_DOWNWARD

[Macro]

Define this macro if successive arguments to a function occupy decreasing addresses on the stack

# HOST\_WIDE\_INT TARGET\_STARTING\_FRAME\_OFFSET (void)

[Target Hook]

This hook returns the offset from the frame pointer to the first local variable slot to be allocated. If FRAME\_GROWS\_DOWNWARD, it is the offset to end of the first slot allocated, otherwise it is the offset to beginning of the first slot allocated. The default implementation returns 0.

# STACK\_ALIGNMENT\_NEEDED

[Macro]

Define to zero to disable final alignment of the stack during reload. The nonzero default for this macro is suitable for most ports.

On ports where TARGET\_STARTING\_FRAME\_OFFSET is nonzero or where there is a register save block following the local block that doesn't require alignment to STACK\_BOUNDARY, it may be beneficial to disable stack alignment and do it in the backend.

#### STACK\_POINTER\_OFFSET

Macro

Offset from the stack pointer register to the first location at which outgoing arguments are placed. If not specified, the default value of zero is used. This is the proper value for most machines.

If ARGS\_GROW\_DOWNWARD, this is the offset to the location above the first location at which outgoing arguments are placed.

#### FIRST\_PARM\_OFFSET (fundec1)

[Macro]

Offset from the argument pointer register to the first argument's address. On some machines it may depend on the data type of the function.

If ARGS\_GROW\_DOWNWARD, this is the offset to the location above the first argument's address.

# STACK\_DYNAMIC\_OFFSET (fundec1)

[Macro]

Offset from the stack pointer register to an item dynamically allocated on the stack, e.g., by alloca.

The default value for this macro is STACK\_POINTER\_OFFSET plus the length of the outgoing arguments. The default is correct for most machines. See 'function.c' for details.

# INITIAL\_FRAME\_ADDRESS\_RTX

[Macro]

A C expression whose value is RTL representing the address of the initial stack frame. This address is passed to RETURN\_ADDR\_RTX and DYNAMIC\_CHAIN\_ADDRESS. If you don't define this macro, a reasonable default value will be used. Define this macro in order to make frame pointer elimination work in the presence of \_\_builtin\_frame\_address (count) and \_\_builtin\_return\_address (count) for count not equal to zero.

# DYNAMIC\_CHAIN\_ADDRESS (frameaddr)

[Macro]

A C expression whose value is RTL representing the address in a stack frame where the pointer to the caller's frame is stored. Assume that frameaddr is an RTL expression for the address of the stack frame itself.

If you don't define this macro, the default is to return the value of frameaddr—that is, the stack frame address is also the address of the stack word that points to the previous frame.

### SETUP\_FRAME\_ADDRESSES

[Macro]

A C expression that produces the machine-specific code to setup the stack so that arbitrary frames can be accessed. For example, on the SPARC, we must flush all of the register windows to the stack before we can access arbitrary stack frames. You will seldom need to define this macro. The default is to do nothing.

# rtx TARGET\_BUILTIN\_SETJMP\_FRAME\_VALUE (void)

[Target Hook]

This target hook should return an rtx that is used to store the address of the current frame into the built in setjmp buffer. The default value, virtual\_stack\_vars\_rtx, is correct for most machines. One reason you may need to define this target hook is if hard\_frame\_pointer\_rtx is the appropriate value on your machine.

# FRAME\_ADDR\_RTX (frameaddr)

[Macro]

A C expression whose value is RTL representing the value of the frame address for the current frame. frameaddr is the frame pointer of the current frame. This is used for \_builtin\_frame\_address. You need only define this macro if the frame address is not the same as the frame pointer. Most machines do not need to define it.

### RETURN\_ADDR\_RTX (count, frameaddr)

[Macro]

A C expression whose value is RTL representing the value of the return address for the frame count steps up from the current frame, after the prologue. frameaddr is the frame pointer of the count frame, or the frame pointer of the count - 1 frame if RETURN\_ADDR\_IN\_PREVIOUS\_FRAME is nonzero.

The value of the expression must always be the correct address when *count* is zero, but may be NULL\_RTX if there is no way to determine the return address of other frames.

# RETURN\_ADDR\_IN\_PREVIOUS\_FRAME

[Macro]

Define this macro to nonzero value if the return address of a particular stack frame is accessed from the frame pointer of the previous stack frame. The zero default for this macro is suitable for most ports.

### INCOMING\_RETURN\_ADDR\_RTX

[Macro]

A C expression whose value is RTL representing the location of the incoming return address at the beginning of any function, before the prologue. This RTL is either a REG, indicating that the return value is saved in 'REG', or a MEM representing a location in the stack.

You only need to define this macro if you want to support call frame debugging information like that provided by DWARF 2.

If this RTL is a REG, you should also define DWARF\_FRAME\_RETURN\_COLUMN to DWARF\_FRAME\_REGNUM (REGNO).

### DWARF\_ALT\_FRAME\_RETURN\_COLUMN

[Macro]

A C expression whose value is an integer giving a DWARF 2 column number that may be used as an alternative return column. The column must not correspond to any gcc hard register (that is, it must not be in the range of DWARF\_FRAME\_REGNUM).

This macro can be useful if DWARF\_FRAME\_RETURN\_COLUMN is set to a general register, but an alternative column needs to be used for signal frames. Some targets have also used different frame return columns over time.

DWARF\_ZERO\_REG [Macro]

A C expression whose value is an integer giving a DWARF 2 register number that is considered to always have the value zero. This should only be defined if the target has an architected zero register, and someone decided it was a good idea to use that register number to terminate the stack backtrace. New ports should avoid this.

#### 

This target hook allows the backend to emit frame-related insns that contain UN-SPECs or UNSPEC\_VOLATILES. The DWARF 2 call frame debugging info engine will invoke it on insns of the form

```
(set (reg) (unspec [...] UNSPEC_INDEX)) and
```

(set (reg) (unspec\_volatile [...] UNSPECV\_INDEX)).

to let the backend emit the call frame instructions. *label* is the CFI label attached to the insn, *pattern* is the pattern of the insn and *index* is UNSPEC\_INDEX or UNSPECV\_INDEX.

### unsigned int TARGET\_DWARF\_POLY\_INDETERMINATE\_VALUE [Target Hook] (unsigned int i, unsigned int \*factor, int \*offset)

Express the value of poly\_int indeterminate i as a DWARF expression, with i counting from 1. Return the number of a DWARF register R and set '\*factor' and '\*offset' such that the value of the indeterminate is:

```
value_of(R) / factor - offset
```

A target only needs to define this hook if it sets 'NUM\_POLY\_INT\_COEFFS' to a value greater than 1.

### INCOMING\_FRAME\_SP\_OFFSET

[Macro]

A C expression whose value is an integer giving the offset, in bytes, from the value of the stack pointer register to the top of the stack frame at the beginning of any

function, before the prologue. The top of the frame is defined to be the value of the stack pointer in the previous frame, just before the call instruction.

You only need to define this macro if you want to support call frame debugging information like that provided by DWARF 2.

### DEFAULT\_INCOMING\_FRAME\_SP\_OFFSET

[Macro]

Like INCOMING\_FRAME\_SP\_OFFSET, but must be the same for all functions of the same ABI, and when using GAS .cfi\_\* directives must also agree with the default CFI GAS emits. Define this macro only if INCOMING\_FRAME\_SP\_OFFSET can have different values between different functions of the same ABI or when INCOMING\_FRAME\_SP\_OFFSET does not agree with GAS default CFI.

### ARG\_POINTER\_CFA\_OFFSET (fundec1)

[Macro]

A C expression whose value is an integer giving the offset, in bytes, from the argument pointer to the canonical frame address (cfa). The final value should coincide with that calculated by INCOMING\_FRAME\_SP\_OFFSET. Which is unfortunately not usable during virtual register instantiation.

The default value for this macro is FIRST\_PARM\_OFFSET (fundec1) + crt1->args.pretend\_args\_size, which is correct for most machines; in general, the arguments are found immediately before the stack frame. Note that this is not the case on some targets that save registers into the caller's frame, such as SPARC and rs6000, and so such targets need to define this macro.

You only need to define this macro if the default is incorrect, and you want to support call frame debugging information like that provided by DWARF 2.

### FRAME\_POINTER\_CFA\_OFFSET (fundec1)

[Macro]

If defined, a C expression whose value is an integer giving the offset in bytes from the frame pointer to the canonical frame address (cfa). The final value should coincide with that calculated by INCOMING\_FRAME\_SP\_OFFSET.

Normally the CFA is calculated as an offset from the argument pointer, via ARG\_POINTER\_CFA\_OFFSET, but if the argument pointer is variable due to the ABI, this may not be possible. If this macro is defined, it implies that the virtual register instantiation should be based on the frame pointer instead of the argument pointer. Only one of FRAME\_POINTER\_CFA\_OFFSET and ARG\_POINTER\_CFA\_OFFSET should be defined.

### CFA\_FRAME\_BASE\_OFFSET (fundec1)

[Macro]

If defined, a C expression whose value is an integer giving the offset in bytes from the canonical frame address (cfa) to the frame base used in DWARF 2 debug information. The default is zero. A different value may reduce the size of debug information on some ports.

### 18.9.2 Exception Handling Support

### EH\_RETURN\_DATA\_REGNO (N)

[Macro]

A C expression whose value is the Nth register number used for data by exception handlers, or INVALID\_REGNUM if fewer than N registers are usable.

The exception handling library routines communicate with the exception handlers via a set of agreed upon registers. Ideally these registers should be call-clobbered; it is possible to use call-saved registers, but may negatively impact code size. The target must support at least 2 data registers, but should define 4 if there are enough free registers.

You must define this macro if you want to support call frame exception handling like that provided by DWARF 2.

### EH\_RETURN\_STACKADJ\_RTX

[Macro]

A C expression whose value is RTL representing a location in which to store a stack adjustment to be applied before function return. This is used to unwind the stack to an exception handler's call frame. It will be assigned zero on code paths that return normally.

Typically this is a call-clobbered hard register that is otherwise untouched by the epilogue, but could also be a stack slot.

Do not define this macro if the stack pointer is saved and restored by the regular prolog and epilog code in the call frame itself; in this case, the exception handling library routines will update the stack location to be restored in place. Otherwise, you must define this macro if you want to support call frame exception handling like that provided by DWARF 2.

### EH\_RETURN\_HANDLER\_RTX

[Macro]

A C expression whose value is RTL representing a location in which to store the address of an exception handler to which we should return. It will not be assigned on code paths that return normally.

Typically this is the location in the call frame at which the normal return address is stored. For targets that return by popping an address off the stack, this might be a memory address just below the *target* call frame rather than inside the current call frame. If defined, EH\_RETURN\_STACKADJ\_RTX will have already been assigned, so it may be used to calculate the location of the target call frame.

Some targets have more complex requirements than storing to an address calculable during initial code generation. In that case the eh\_return instruction pattern should be used instead.

If you want to support call frame exception handling, you must define either this macro or the eh\_return instruction pattern.

### RETURN\_ADDR\_OFFSET

[Macro]

If defined, an integer-valued C expression for which rtl will be generated to add it to the exception handler address before it is searched in the exception handling tables, and to subtract it again from the address before using it to return to the exception handler.

### ASM\_PREFERRED\_EH\_DATA\_FORMAT (code, global)

[Macro]

This macro chooses the encoding of pointers embedded in the exception handling sections. If at all possible, this should be defined such that the exception handling section will not require dynamic relocations, and so may be read-only.

code is 0 for data, 1 for code labels, 2 for function pointers. global is true if the symbol may be affected by dynamic relocations. The macro should return a combination of the DW\_EH\_PE\_\* defines as found in 'dwarf2.h'.

If this macro is not defined, pointers will not be encoded but represented directly.

### ASM\_MAYBE\_OUTPUT\_ENCODED\_ADDR\_RTX (file, encoding, size, addr, done) [Macro]

This macro allows the target to emit whatever special magic is required to represent the encoding chosen by ASM\_PREFERRED\_EH\_DATA\_FORMAT. Generic code takes care of pc-relative and indirect encodings; this must be defined if the target uses text-relative or data-relative encodings.

This is a C statement that branches to *done* if the format was handled. *encoding* is the format chosen, *size* is the number of bytes that the format occupies, *addr* is the SYMBOL\_REF to be emitted.

### MD\_FALLBACK\_FRAME\_STATE\_FOR (context, fs)

[Macro]

This macro allows the target to add CPU and operating system specific code to the call-frame unwinder for use when there is no unwind data available. The most common reason to implement this macro is to unwind through signal frames.

This macro is called from uw\_frame\_state\_for in 'unwind-dw2.c', 'unwind-dw2-xtensa.c' and 'unwind-ia64.c'. context is an \_Unwind\_Context; fs is an \_Unwind\_FrameState. Examine context->ra for the address of the code being executed and context->cfa for the stack pointer value. If the frame can be decoded, the register save addresses should be updated in fs and the macro should evaluate to \_URC\_NO\_REASON. If the frame cannot be decoded, the macro should evaluate to \_URC\_END\_OF\_STACK.

For proper signal handling in Java this macro is accompanied by MAKE\_THROW\_FRAME, defined in 'libjava/include/\*-signal.h' headers.

#### MD\_HANDLE\_UNWABI (context, fs)

[Macro]

This macro allows the target to add operating system specific code to the call-frame unwinder to handle the IA-64 .unwabi unwinding directive, usually used for signal or interrupt frames.

This macro is called from uw\_update\_context in libgcc's 'unwind-ia64.c'. context is an \_Unwind\_Context; fs is an \_Unwind\_FrameState. Examine fs->unwabi for the abi and context in the .unwabi directive. If the .unwabi directive can be handled, the register save addresses should be updated in fs.

### TARGET\_USES\_WEAK\_UNWIND\_INFO

Macro

A C expression that evaluates to true if the target requires unwind info to be given comdat linkage. Define it to be 1 if comdat linkage is necessary. The default is 0.

### 18.9.3 Specifying How Stack Checking is Done

GCC will check that stack references are within the boundaries of the stack, if the option '-fstack-check' is specified, in one of three ways:

1. If the value of the STACK\_CHECK\_BUILTIN macro is nonzero, GCC will assume that you have arranged for full stack checking to be done at appropriate places in the configuration files. GCC will not do other special processing.

- 2. If STACK\_CHECK\_BUILTIN is zero and the value of the STACK\_CHECK\_STATIC\_BUILTIN macro is nonzero, GCC will assume that you have arranged for static stack checking (checking of the static stack frame of functions) to be done at appropriate places in the configuration files. GCC will only emit code to do dynamic stack checking (checking on dynamic stack allocations) using the third approach below.
- 3. If neither of the above are true, GCC will generate code to periodically "probe" the stack pointer using the values of the macros defined below.

If neither STACK\_CHECK\_BUILTIN nor STACK\_CHECK\_STATIC\_BUILTIN is defined, GCC will change its allocation strategy for large objects if the option '-fstack-check' is specified: they will always be allocated dynamically if their size exceeds STACK\_CHECK\_MAX\_VAR\_SIZE bytes.

#### STACK CHECK BUILTIN

[Macro]

A nonzero value if stack checking is done by the configuration files in a machine-dependent manner. You should define this macro if stack checking is required by the ABI of your machine or if you would like to do stack checking in some more efficient way than the generic approach. The default value of this macro is zero.

### STACK\_CHECK\_STATIC\_BUILTIN

[Macro]

A nonzero value if static stack checking is done by the configuration files in a machine-dependent manner. You should define this macro if you would like to do static stack checking in some more efficient way than the generic approach. The default value of this macro is zero.

### STACK\_CHECK\_PROBE\_INTERVAL\_EXP

[Macro]

An integer specifying the interval at which GCC must generate stack probe instructions, defined as 2 raised to this integer. You will normally define this macro so that the interval be no larger than the size of the "guard pages" at the end of a stack area. The default value of 12 (4096-byte interval) is suitable for most systems.

### STACK\_CHECK\_MOVING\_SP

[Macro]

An integer which is nonzero if GCC should move the stack pointer page by page when doing probes. This can be necessary on systems where the stack pointer contains the bottom address of the memory area accessible to the executing thread at any point in time. In this situation an alternate signal stack is required in order to be able to recover from a stack overflow. The default value of this macro is zero.

### STACK\_CHECK\_PROTECT

[Macro]

The number of bytes of stack needed to recover from a stack overflow, for languages where such a recovery is supported. The default value of 4KB/8KB with the setjmp/longjmp-based exception handling mechanism and 8KB/12KB with other exception handling mechanisms should be adequate for most architectures and operating systems.

The following macros are relevant only if neither STACK\_CHECK\_BUILTIN nor STACK\_CHECK\_STATIC\_BUILTIN is defined; you can omit them altogether in the opposite case.

### STACK\_CHECK\_MAX\_FRAME\_SIZE

[Macro]

The maximum size of a stack frame, in bytes. GCC will generate probe instructions in non-leaf functions to ensure at least this many bytes of stack are available. If a stack frame is larger than this size, stack checking will not be reliable and GCC will issue a warning. The default is chosen so that GCC only generates one instruction on most systems. You should normally not change the default value of this macro.

### STACK\_CHECK\_FIXED\_FRAME\_SIZE

[Macro]

GCC uses this value to generate the above warning message. It represents the amount of fixed frame used by a function, not including space for any callee-saved registers, temporaries and user variables. You need only specify an upper bound for this amount and will normally use the default of four words.

### STACK\_CHECK\_MAX\_VAR\_SIZE

[Macro]

The maximum size, in bytes, of an object that GCC will place in the fixed area of the stack frame when the user specifies '-fstack-check'. GCC computed the default from the values of the above macros and you will normally not need to override that default.

### HOST\_WIDE\_INT

[Target Hook]

### TARGET\_STACK\_CLASH\_PROTECTION\_ALLOCA\_PROBE\_RANGE (void)

Some targets have an ABI defined interval for which no probing needs to be done. When a probe does need to be done this same interval is used as the probe distance up when doing stack clash protection for alloca. On such targets this value can be set to override the default probing up interval. Define this variable to return nonzero if such a probe range is required or zero otherwise. Defining this hook also requires your functions which make use of alloca to have at least 8 byesof outgoing arguments. If this is not the case the stack will be corrupted. You need not define this macro if it would always have the value zero.

### 18.9.4 Registers That Address the Stack Frame

This discusses registers that address the stack frame.

### STACK\_POINTER\_REGNUM

[Macro]

The register number of the stack pointer register, which must also be a fixed register according to FIXED\_REGISTERS. On most machines, the hardware determines which register this is.

### FRAME\_POINTER\_REGNUM

[Macro]

The register number of the frame pointer register, which is used to access automatic variables in the stack frame. On some machines, the hardware determines which register this is. On other machines, you can choose any register you wish for this purpose.

### HARD\_FRAME\_POINTER\_REGNUM

[Macro]

On some machines the offset between the frame pointer and starting offset of the automatic variables is not known until after register allocation has been done (for example, because the saved registers are between these two locations). On those

machines, define FRAME\_POINTER\_REGNUM the number of a special, fixed register to be used internally until the offset is known, and define HARD\_FRAME\_POINTER\_REGNUM to be the actual hard register number used for the frame pointer.

You should define this macro only in the very rare circumstances when it is not possible to calculate the offset between the frame pointer and the automatic variables until after register allocation has been completed. When this macro is defined, you must also indicate in your definition of ELIMINABLE\_REGS how to eliminate FRAME\_POINTER\_REGNUM into either HARD\_FRAME\_POINTER\_REGNUM or STACK\_POINTER\_REGNUM.

Do not define this macro if it would be the same as FRAME\_POINTER\_REGNUM.

### ARG\_POINTER\_REGNUM

[Macro]

The register number of the arg pointer register, which is used to access the function's argument list. On some machines, this is the same as the frame pointer register. On some machines, the hardware determines which register this is. On other machines, you can choose any register you wish for this purpose. If this is not the same register as the frame pointer register, then you must mark it as a fixed register according to FIXED\_REGISTERS, or arrange to be able to eliminate it (see Section 18.9.5 [Elimination], page 533).

#### HARD\_FRAME\_POINTER\_IS\_FRAME\_POINTER

[Macro]

Define this to a preprocessor constant that is nonzero if hard\_frame\_pointer\_rtx and frame\_pointer\_rtx should be the same. The default definition is '(HARD\_FRAME\_POINTER\_REGNUM == FRAME\_POINTER\_REGNUM)'; you only need to define this macro if that definition is not suitable for use in preprocessor conditionals.

### HARD\_FRAME\_POINTER\_IS\_ARG\_POINTER

[Macro]

Define this to a preprocessor constant that is nonzero if hard\_frame\_pointer\_rtx and arg\_pointer\_rtx should be the same. The default definition is '(HARD\_FRAME\_POINTER\_REGNUM == ARG\_POINTER\_REGNUM)'; you only need to define this macro if that definition is not suitable for use in preprocessor conditionals.

### RETURN\_ADDRESS\_POINTER\_REGNUM

Macro

The register number of the return address pointer register, which is used to access the current function's return address from the stack. On some machines, the return address is not at a fixed offset from the frame pointer or stack pointer or argument pointer. This register can be defined to point to the return address on the stack, and then be converted by ELIMINABLE\_REGS into either the frame pointer or stack pointer.

Do not define this macro unless there is no other way to get the return address from the stack.

### STATIC\_CHAIN\_REGNUM

[Macro]

### STATIC\_CHAIN\_INCOMING\_REGNUM

[Macro]

Register numbers used for passing a function's static chain pointer. If register windows are used, the register number as seen by the called function is STATIC\_CHAIN\_INCOMING\_REGNUM, while the register number as seen by the calling function is STATIC\_CHAIN\_REGNUM. If these registers are the same, STATIC\_CHAIN\_INCOMING\_REGNUM need not be defined.

The static chain register need not be a fixed register.

If the static chain is passed in memory, these macros should not be defined; instead, the TARGET\_STATIC\_CHAIN hook should be used.

### rtx TARGET\_STATIC\_CHAIN (const\_tree fndecl\_or\_type, bool incoming\_p) [Target Hook]

This hook replaces the use of STATIC\_CHAIN\_REGNUM et al for targets that may use different static chain locations for different nested functions. This may be required if the target has function attributes that affect the calling conventions of the function and those calling conventions use different static chain locations.

The default version of this hook uses STATIC\_CHAIN\_REGNUM et al.

If the static chain is passed in memory, this hook should be used to provide rtx giving mem expressions that denote where they are stored. Often the mem expression as seen by the caller will be at an offset from the stack pointer and the mem expression as seen by the callee will be at an offset from the frame pointer. The variables stack\_pointer\_rtx, frame\_pointer\_rtx, and arg\_pointer\_rtx will have been initialized and should be used to refer to those items.

### DWARF\_FRAME\_REGISTERS

[Macro]

This macro specifies the maximum number of hard registers that can be saved in a call frame. This is used to size data structures used in DWARF2 exception handling. Prior to GCC 3.0, this macro was needed in order to establish a stable exception handling ABI in the face of adding new hard registers for ISA extensions. In GCC 3.0 and later, the EH ABI is insulated from changes in the number of hard registers. Nevertheless, this macro can still be used to reduce the runtime memory requirements of the exception handling routines, which can be substantial if the ISA contains a lot of registers that are not call-saved.

If this macro is not defined, it defaults to FIRST\_PSEUDO\_REGISTER.

### PRE\_GCC3\_DWARF\_FRAME\_REGISTERS

[Macro]

This macro is similar to DWARF\_FRAME\_REGISTERS, but is provided for backward compatibility in pre GCC 3.0 compiled code.

If this macro is not defined, it defaults to DWARF\_FRAME\_REGISTERS.

### DWARF\_REG\_TO\_UNWIND\_COLUMN (regno)

[Macro]

Define this macro if the target's representation for dwarf registers is different than the internal representation for unwind column. Given a dwarf register, this macro should return the internal unwind column number to use instead.

### DWARF\_FRAME\_REGNUM (regno)

[Macro]

Define this macro if the target's representation for dwarf registers used in .eh\_frame or .debug\_frame is different from that used in other debug info sections. Given a GCC hard register number, this macro should return the .eh\_frame register number. The default is DBX\_REGISTER\_NUMBER (regno).

### DWARF2\_FRAME\_REG\_OUT (regno, for\_eh)

[Macro]

Define this macro to map register numbers held in the call frame info that GCC has collected using DWARF\_FRAME\_REGNUM to those that should be output in .debug\_frame (for\_eh is zero) and .eh\_frame (for\_eh is nonzero). The default is to return regno.

### REG\_VALUE\_IN\_UNWIND\_CONTEXT

[Macro]

Define this macro if the target stores register values as \_Unwind\_Word type in unwind context. It should be defined if target register size is larger than the size of void \*. The default is to store register values as void \* type.

### ASSUME\_EXTENDED\_UNWIND\_CONTEXT

[Macro]

Define this macro to be 1 if the target always uses extended unwind context with version, args\_size and by\_value fields. If it is undefined, it will be defined to 1 when REG\_VALUE\_IN\_UNWIND\_CONTEXT is defined and 0 otherwise.

### DWARF\_LAZY\_REGISTER\_VALUE (regno, value)

[Macro]

Define this macro if the target has pseudo DWARF registers whose values need to be computed lazily on demand by the unwinder (such as when referenced in a CFA expression). The macro returns true if regno is such a register and stores its value in '\*value' if so.

### 18.9.5 Eliminating Frame Pointer and Arg Pointer

This is about eliminating the frame pointer and arg pointer.

### bool TARGET\_FRAME\_POINTER\_REQUIRED (void)

[Target Hook]

This target hook should return **true** if a function must have and use a frame pointer. This target hook is called in the reload pass. If its return value is **true** the function will have a frame pointer.

This target hook can in principle examine the current function and decide according to the facts, but on most machines the constant false or the constant true suffices. Use false when the machine allows code to be generated with no frame pointer, and doing so saves some time or space. Use true when there is no possible advantage to avoiding a frame pointer.

In certain cases, the compiler does not know how to produce valid code without a frame pointer. The compiler recognizes those cases and automatically gives the function a frame pointer regardless of what targetm.frame\_pointer\_required returns. You don't need to worry about them.

In a function that does not require a frame pointer, the frame pointer register can be allocated for ordinary usage, unless you mark it as a fixed register. See FIXED\_REGISTERS for more information.

Default return value is false.

#### ELIMINABLE\_REGS

[Macro]

This macro specifies a table of register pairs used to eliminate unneeded registers that point into the stack frame.

The definition of this macro is a list of structure initializations, each of which specifies an original and replacement register.

On some machines, the position of the argument pointer is not known until the compilation is completed. In such a case, a separate hard register must be used for the argument pointer. This register can be eliminated by replacing it with either the frame pointer or the argument pointer, depending on whether or not the frame pointer has been eliminated.

In this case, you might specify:

```
#define ELIMINABLE_REGS \
{{ARG_POINTER_REGNUM, STACK_POINTER_REGNUM}, \
{ARG_POINTER_REGNUM, FRAME_POINTER_REGNUM}, \
{FRAME_POINTER_REGNUM, STACK_POINTER_REGNUM}}
```

Note that the elimination of the argument pointer with the stack pointer is specified first since that is the preferred elimination.

### bool TARGET\_CAN\_ELIMINATE (const int from\_reg, const int to\_reg) [Target Hook]

This target hook should return **true** if the compiler is allowed to try to replace register number from\_reg with register number to\_reg. This target hook will usually be **true**, since most of the cases preventing register elimination are things that the compiler already knows about.

Default return value is true.

### INITIAL\_ELIMINATION\_OFFSET (from-reg, to-reg, offset-var) [Macro]

This macro returns the initial difference between the specified pair of registers. The value would be computed from information such as the result of get\_frame\_size () and the tables of registers df\_regs\_ever\_live\_p and call\_used\_regs.

### void TARGET\_COMPUTE\_FRAME\_LAYOUT (void)

[Target Hook]

This target hook is called once each time the frame layout needs to be recalculated. The calculations can be cached by the target and can then be used by INITIAL\_ELIMINATION\_OFFSET instead of re-computing the layout on every invocation of that hook. This is particularly useful for targets that have an expensive frame layout function. Implementing this callback is optional.

### 18.9.6 Passing Function Arguments on the Stack

The macros in this section control how arguments are passed on the stack. See the following section for other macros that control passing certain arguments in registers.

### bool TARGET\_PROMOTE\_PROTOTYPES (const\_tree fntype)

[Target Hook]

This target hook returns true if an argument declared in a prototype as an integral type smaller than int should actually be passed as an int. In addition to avoiding errors in certain cases of mismatch, it also makes for better code on certain machines. The default is to not promote prototypes.

PUSH\_ARGS [Macro]

A C expression. If nonzero, push insns will be used to pass outgoing arguments. If the target machine does not have a push instruction, set it to zero. That directs GCC to use an alternate strategy: to allocate the entire argument block and then store the arguments into it. When PUSH\_ARGS is nonzero, PUSH\_ROUNDING must be defined too.

### PUSH\_ARGS\_REVERSED

[Macro]

A C expression. If nonzero, function arguments will be evaluated from last to first, rather than from first to last. If this macro is not defined, it defaults to PUSH\_ARGS on targets where the stack and args grow in opposite directions, and 0 otherwise.

### PUSH\_ROUNDING (npushed)

[Macro]

A C expression that is the number of bytes actually pushed onto the stack when an instruction attempts to push *npushed* bytes.

On some machines, the definition

#define PUSH\_ROUNDING(BYTES) (BYTES)

will suffice. But on other machines, instructions that appear to push one byte actually push two bytes in an attempt to maintain alignment. Then the definition should be

#define PUSH\_ROUNDING(BYTES) (((BYTES) + 1) & ~1)

If the value of this macro has a type, it should be an unsigned type.

### ACCUMULATE\_OUTGOING\_ARGS

[Macro]

A C expression. If nonzero, the maximum amount of space required for outgoing arguments will be computed and placed into crtl->outgoing\_args\_size. No space will be pushed onto the stack for each call; instead, the function prologue should increase the stack frame size by this amount.

Setting both PUSH\_ARGS and ACCUMULATE\_OUTGOING\_ARGS is not proper.

### REG\_PARM\_STACK\_SPACE (fndec1)

[Macro]

Define this macro if functions should assume that stack space has been allocated for arguments even when their values are passed in registers.

The value of this macro is the size, in bytes, of the area reserved for arguments passed in registers for the function represented by *fndecl*, which can be zero if GCC is calling a library function. The argument *fndecl* can be the FUNCTION\_DECL, or the type itself of the function.

This space can be allocated by the caller, or be a part of the machine-dependent stack frame: OUTGOING\_REG\_PARM\_STACK\_SPACE says which.

### INCOMING\_REG\_PARM\_STACK\_SPACE (fndec1)

[Macro]

Like REG\_PARM\_STACK\_SPACE, but for incoming register arguments. Define this macro if space guaranteed when compiling a function body is different to space required when making a call, a situation that can arise with K&R style function definitions.

### OUTGOING\_REG\_PARM\_STACK\_SPACE (fntype)

[Macro]

Define this to a nonzero value if it is the responsibility of the caller to allocate the area reserved for arguments passed in registers when calling a function of *fntype*. *fntype* may be NULL if the function called is a library function.

If ACCUMULATE\_OUTGOING\_ARGS is defined, this macro controls whether the space for these arguments counts in the value of crtl->outgoing\_args\_size.

### STACK\_PARMS\_IN\_REG\_PARM\_AREA

[Macro]

Define this macro if REG\_PARM\_STACK\_SPACE is defined, but the stack parameters don't skip the area specified by it.

Normally, when a parameter is not passed in registers, it is placed on the stack beyond the REG\_PARM\_STACK\_SPACE area. Defining this macro suppresses this behavior and causes the parameter to be passed on the stack in its natural location.

### poly\_int64 TARGET\_RETURN\_POPS\_ARGS (tree fundec1, tree [Target Hook] funtype, poly\_int64 size)

This target hook returns the number of bytes of its own arguments that a function pops on returning, or 0 if the function pops no arguments and the caller must therefore pop them all after the function returns.

fundecl is a C variable whose value is a tree node that describes the function in question. Normally it is a node of type FUNCTION\_DECL that describes the declaration of the function. From this you can obtain the DECL\_ATTRIBUTES of the function.

funtype is a C variable whose value is a tree node that describes the function in question. Normally it is a node of type FUNCTION\_TYPE that describes the data type of the function. From this it is possible to obtain the data types of the value and arguments (if known).

When a call to a library function is being considered, fundecl will contain an identifier node for the library function. Thus, if you need to distinguish among various library functions, you can do so by their names. Note that "library function" in this context means a function used to perform arithmetic, whose name is known specially in the compiler and was not mentioned in the C code being compiled.

size is the number of bytes of arguments passed on the stack. If a variable number of bytes is passed, it is zero, and argument popping will always be the responsibility of the calling function.

On the VAX, all functions always pop their arguments, so the definition of this macro is size. On the 68000, using the standard calling convention, no functions pop their arguments, so the value of the macro is always 0 in this case. But an alternative calling convention is available in which functions that take a fixed number of arguments pop them but other functions (such as printf) pop nothing (the caller pops all). When this convention is in use, funtype is examined to determine whether a function takes a fixed number of arguments.

#### CALL\_POPS\_ARGS (cum)

Macro

A C expression that should indicate the number of bytes a call sequence pops off the stack. It is added to the value of RETURN\_POPS\_ARGS when compiling a function call. *cum* is the variable in which all arguments to the called function have been accumulated.

On certain architectures, such as the SH5, a call trampoline is used that pops certain registers off the stack, depending on the arguments that have been passed to the function. Since this is a property of the call site, not of the called function, RETURN\_POPS\_ARGS is not appropriate.

### 18.9.7 Passing Arguments in Registers

This section describes the macros which let you control how various types of arguments are passed in registers or how they are arranged in the stack.

### rtx TARGET\_FUNCTION\_ARG (cumulative\_args\_t ca, const function\_arg\_info &arg) [Target Hook]

Return an RTX indicating whether function argument arg is passed in a register and if so, which register. Argument ca summarizes all the previous arguments.

The return value is usually either a reg RTX for the hard register in which to pass the argument, or zero to pass the argument on the stack.

The return value can be a <code>const\_int</code> which means argument is passed in a target specific slot with specified number. Target hooks should be used to store or load argument in such case. See <code>TARGET\_STORE\_BOUNDS\_FOR\_ARG</code> and <code>TARGET\_LOAD\_BOUNDS\_FOR\_ARG</code> for more information.

The value of the expression can also be a parallel RTX. This is used when an argument is passed in multiple locations. The mode of the parallel should be the mode of the entire argument. The parallel holds any number of expr\_list pairs; each one describes where part of the argument is passed. In each expr\_list the first operand must be a reg RTX for the hard register in which to pass this part of the argument, and the mode of the register RTX indicates how large this part of the argument is. The second operand of the expr\_list is a const\_int which gives the offset in bytes into the entire argument of where this part starts. As a special exception the first expr\_list in the parallel RTX may have a first operand of zero. This indicates that the entire argument is also stored on the stack.

The last time this hook is called, it is called with MODE == VOIDmode, and its result is passed to the call or call\_value pattern as operands 2 and 3 respectively.

The usual way to make the ISO library 'stdarg.h' work on a machine where some arguments are usually passed in registers, is to cause nameless arguments to be passed on the stack instead. This is done by making TARGET\_FUNCTION\_ARG return 0 whenever named is false.

You may use the hook targetm.calls.must\_pass\_in\_stack in the definition of this macro to determine if this argument is of a type that must be passed in the stack. If REG\_PARM\_STACK\_SPACE is not defined and TARGET\_FUNCTION\_ARG returns nonzero for such an argument, the compiler will abort. If REG\_PARM\_STACK\_SPACE is defined, the argument will be computed in the stack and then loaded into a register.

bool TARGET\_MUST\_PASS\_IN\_STACK (const function\_arg\_info &arg) [Target Hook]
This target hook should return true if we should not pass arg solely in registers.
The file 'expr.h' defines a definition that is usually appropriate, refer to 'expr.h' for additional documentation.

### rtx TARGET\_FUNCTION\_INCOMING\_ARG (cumulative\_args\_t ca, const [Target Hook] function\_arg\_info & arg)

Define this hook if the caller and callee on the target have different views of where arguments are passed. Also define this hook if there are functions that are never directly called, but are invoked by the hardware and which have nonstandard calling conventions.

In this case TARGET\_FUNCTION\_ARG computes the register in which the caller passes the value, and TARGET\_FUNCTION\_INCOMING\_ARG should be defined in a similar fashion to tell the function being called where the arguments will arrive.

TARGET\_FUNCTION\_INCOMING\_ARG can also return arbitrary address computation using hard register, which can be forced into a register, so that it can be used to pass special arguments.

If TARGET\_FUNCTION\_INCOMING\_ARG is not defined, TARGET\_FUNCTION\_ARG serves both purposes.

### bool TARGET\_USE\_PSEUDO\_PIC\_REG (void)

[Target Hook]

This hook should return 1 in case pseudo register should be created for pic\_offset\_table\_rtx during function expand.

### void TARGET\_INIT\_PIC\_REG (void)

[Target Hook]

Perform a target dependent initialization of pic\_offset\_table\_rtx. This hook is called at the start of register allocation.

### int TARGET\_ARG\_PARTIAL\_BYTES (cumulative\_args\_t cum, const function\_arg\_info &arg)

[Target Hook]

This target hook returns the number of bytes at the beginning of an argument that must be put in registers. The value must be zero for arguments that are passed entirely in registers or that are entirely pushed on the stack.

On some machines, certain arguments must be passed partially in registers and partially in memory. On these machines, typically the first few words of arguments are passed in registers, and the rest on the stack. If a multi-word argument (a double or a structure) crosses that boundary, its first few words must be passed in registers and the rest must be pushed. This macro tells the compiler when this occurs, and how many bytes should go in registers.

TARGET\_FUNCTION\_ARG for these arguments should return the first register to be used by the caller for this argument; likewise TARGET\_FUNCTION\_INCOMING\_ARG, for the called function.

### bool TARGET\_PASS\_BY\_REFERENCE (cumulative\_args\_t cum, const function\_arg\_info &arg)

[Target Hook]

This target hook should return true if argument arg at the position indicated by cum should be passed by reference. This predicate is queried after target independent reasons for being passed by reference, such as TREE\_ADDRESSABLE (arg.type).

If the hook returns true, a copy of that argument is made in memory and a pointer to the argument is passed instead of the argument itself. The pointer is passed in whatever way is appropriate for passing a pointer to that type.

### bool TARGET\_CALLEE\_COPIES (cumulative\_args\_t cum, const

[Target Hook]

function\_arg\_info &arg)

The function argument described by the parameters to this hook is known to be passed by reference. The hook should return true if the function argument should be copied by the callee instead of copied by the caller.

For any argument for which the hook returns true, if it can be determined that the argument is not modified, then a copy need not be generated.

The default version of this hook always returns false.

### CUMULATIVE\_ARGS

[Macro]

A C type for declaring a variable that is used as the first argument of TARGET\_FUNCTION\_ARG and other related values. For some target machines, the type int suffices and can hold the number of bytes of argument so far.

There is no need to record in CUMULATIVE\_ARGS anything about the arguments that have been passed on the stack. The compiler has other variables to keep track of that. For target machines on which all arguments are passed on the stack, there is no need to store anything in CUMULATIVE\_ARGS; however, the data structure must exist and should not be empty, so use int.

### OVERRIDE\_ABI\_FORMAT (fndec1)

[Macro]

If defined, this macro is called before generating any code for a function, but after the cfun descriptor for the function has been created. The back end may use this macro to update cfun to reflect an ABI other than that which would normally be used by default. If the compiler is generating code for a compiler-generated function, findecl may be NULL.

#### 

A C statement (sans semicolon) for initializing the variable cum for the state at the beginning of the argument list. The variable has type CUMULATIVE\_ARGS. The value of fntype is the tree node for the data type of the function which will receive the args, or 0 if the args are to a compiler support library function. For direct calls that are not libcalls, fndecl contain the declaration node of the function. fndecl is also set when INIT\_CUMULATIVE\_ARGS is used to find arguments for the function being compiled.  $n\_named\_args$  is set to the number of named arguments, including a structure return address if it is passed as a parameter, when making a call. When processing incoming arguments,  $n\_named\_args$  is set to -1.

When processing a call to a compiler support library function, *libname* identifies which one. It is a symbol\_ref rtx which contains the name of the function, as a string. *libname* is 0 when an ordinary C function call is being processed. Thus, each time this macro is called, either *libname* or *fntype* is nonzero, but never both of them at once.

### INIT\_CUMULATIVE\_LIBCALL\_ARGS (cum, mode, libname)

Macro

Like INIT\_CUMULATIVE\_ARGS but only used for outgoing libcalls, it gets a MODE argument instead of *fntype*, that would be NULL. *indirect* would always be zero, too. If this macro is not defined, INIT\_CUMULATIVE\_ARGS (cum, NULL\_RTX, libname, 0) is used instead.

### INIT\_CUMULATIVE\_INCOMING\_ARGS (cum, fntype, libname)

[Macro]

Like INIT\_CUMULATIVE\_ARGS but overrides it for the purposes of finding the arguments for the function being compiled. If this macro is undefined, INIT\_CUMULATIVE\_ARGS is used instead.

The value passed for *libname* is always 0, since library routines with special calling conventions are never compiled with GCC. The argument *libname* exists for symmetry with INIT\_CUMULATIVE\_ARGS.

#### 

This hook updates the summarizer variable pointed to by ca to advance past argument arg in the argument list. Once this is done, the variable cum is suitable for analyzing the following argument with TARGET\_FUNCTION\_ARG, etc.

This hook need not do anything if the argument in question was passed on the stack. The compiler knows how to track the amount of stack space used for arguments without any special help.

#### 

This hook returns the number of bytes to add to the offset of an argument of type type and mode mode when passed in memory. This is needed for the SPU, which passes char and short arguments in the preferred slot that is in the middle of the quad word instead of starting at the top. The default implementation returns 0.

#### 

This hook determines whether, and in which direction, to pad out an argument of mode mode and type type. It returns PAD\_UPWARD to insert padding above the argument, PAD\_DOWNWARD to insert padding below the argument, or PAD\_NONE to inhibit padding.

The amount of padding is not controlled by this hook, but by TARGET\_FUNCTION\_ARG\_ROUND\_BOUNDARY. It is always just enough to reach the next multiple of that boundary.

This hook has a default definition that is right for most systems. For little-endian machines, the default is to pad upward. For big-endian machines, the default is to pad downward for an argument of constant size shorter than an int, and upward otherwise.

PAD\_VARARGS\_DOWN [Macro]

If defined, a C expression which determines whether the default implementation of va\_arg will attempt to pad down before reading the next argument, if that argument is smaller than its aligned space as controlled by PARM\_BOUNDARY. If this macro is not defined, all such arguments are padded down if BYTES\_BIG\_ENDIAN is true.

### BLOCK\_REG\_PADDING (mode, type, first)

[Macro]

Specify padding for the last element of a block move between registers and memory. first is nonzero if this is the only element. Defining this macro allows better control of register function parameters on big-endian machines, without using PARALLEL rtl. In particular, MUST\_PASS\_IN\_STACK need not test padding and mode of types in registers, as there is no longer a "wrong" part of a register; For example, a three byte aggregate may be passed in the high part of a register if so required.

### unsigned int TARGET\_FUNCTION\_ARG\_BOUNDARY (machine\_mode [Target Hook] mode, const\_tree type)

This hook returns the alignment boundary, in bits, of an argument with the specified mode and type. The default hook returns PARM\_BOUNDARY for all arguments.

### unsigned int TARGET\_FUNCTION\_ARG\_ROUND\_BOUNDARY [Target Hook]

(machine\_mode mode, const\_tree type)

Normally, the size of an argument is rounded up to PARM\_BOUNDARY, which is the default value for this hook. You can define this hook to return a different value if an argument size must be rounded to a larger value.

### FUNCTION\_ARG\_REGNO\_P (regno)

[Macro]

A C expression that is nonzero if regno is the number of a hard register in which function arguments are sometimes passed. This does not include implicit arguments such as the static chain and the structure-value address. On many machines, no registers can be used for this purpose since all function arguments are pushed on the stack.

### bool TARGET\_SPLIT\_COMPLEX\_ARG (const\_tree type)

[Target Hook]

This hook should return true if parameter of type *type* are passed as two scalar parameters. By default, GCC will attempt to pack complex arguments into the target's word size. Some ABIs require complex arguments to be split and treated as their individual components. For example, on AIX64, complex floats should be passed in a pair of floating point registers, even though a complex float would fit in one 64-bit floating point register.

The default value of this hook is NULL, which is treated as always false.

### tree TARGET\_BUILD\_BUILTIN\_VA\_LIST (void)

[Target Hook]

This hook returns a type node for va\_list for the target. The default version of the hook returns void\*.

#### 

This target hook is used in function  $c\_common\_nodes\_and\_builtins$  to iterate through the target specific builtin types for va\_list. The variable idx is used as iterator. pname has to be a pointer to a const char \* and ptree a pointer to a tree typed variable. The arguments pname and ptree are used to store the result of this macro and are set to the name of the va\_list builtin type and its internal type. If the return value of this macro is zero, then there is no more element. Otherwise the IDX should be increased for the next call of this macro to iterate through all types.

### tree TARGET\_FN\_ABI\_VA\_LIST (tree fndec1)

[Target Hook]

This hook returns the va\_list type of the calling convention specified by *fndecl*. The default version of this hook returns va\_list\_type\_node.

### tree TARGET\_CANONICAL\_VA\_LIST\_TYPE (tree type)

[Target Hook]

This hook returns the va\_list type of the calling convention specified by the type of type. If type is not a valid va\_list type, it returns NULL\_TREE.

### tree TARGET\_GIMPLIFY\_VA\_ARG\_EXPR (tree valist, tree type, gimple\_seq \*pre\_p, gimple\_seq \*post\_p) [Target Hook]

This hook performs target-specific gimplification of VA\_ARG\_EXPR. The first two parameters correspond to the arguments to va\_arg; the latter two are as in gimplify.c:gimplify\_expr.

### bool TARGET\_VALID\_POINTER\_MODE (scalar\_int\_mode mode) [T

[Target Hook]

Define this to return nonzero if the port can handle pointers with machine mode mode. The default version of this hook returns true for both ptr\_mode and Pmode.

### bool TARGET\_REF\_MAY\_ALIAS\_ERRNO (ao\_ref \*ref)

[Target Hook]

Define this to return nonzero if the memory reference ref may alias with the system C library errno location. The default version of this hook assumes the system C library errno location is either a declaration of type int or accessed by dereferencing a pointer to int.

### machine\_mode TARGET\_TRANSLATE\_MODE\_ATTRIBUTE

[Target Hook]

(machine\_mode mode)

Define this hook if during mode attribute processing, the port should translate machine\_mode mode to another mode. For example, rs6000's KFmode, when it is the same as TFmode.

The default version of the hook returns that mode that was passed in.

### bool TARGET\_SCALAR\_MODE\_SUPPORTED\_P (scalar\_mode mode)

[Target Hook]

Define this to return nonzero if the port is prepared to handle insns involving scalar mode *mode*. For a scalar mode to be considered supported, all the basic arithmetic and comparisons must work.

The default version of this hook returns true for any mode required to handle the basic C types (as defined by the port). Included here are the double-word arithmetic supported by the code in 'optabs.c'.

bool TARGET\_VECTOR\_MODE\_SUPPORTED\_P (machine\_mode mode) [Target Hook]

Define this to return nonzero if the port is prepared to handle insns involving vector mode mode. At the very least, it must have move patterns for this mode.

### bool TARGET\_COMPATIBLE\_VECTOR\_TYPES\_P (const\_tree type1, const\_tree type2) [Target Hook]

Return true if there is no target-specific reason for treating vector types type1 and type2 as distinct types. The caller has already checked for target-independent reasons, meaning that the types are known to have the same mode, to have the same number of elements, and to have what the caller considers to be compatible element types.

The main reason for defining this hook is to reject pairs of types that are handled differently by the target's calling convention. For example, when a new N-bit vector architecture is added to a target, the target may want to handle normal N-bit VECTOR\_TYPE arguments and return values in the same way as before, to maintain backwards compatibility. However, it may also provide new, architecture-specific VECTOR\_TYPEs that are passed and returned in a more efficient way. It is then important to maintain a distinction between the "normal" VECTOR\_TYPEs and the new architecture-specific ones.

The default implementation returns true, which is correct for most targets.

### opt\_machine\_mode TARGET\_ARRAY\_MODE (machine\_mode mode, unsigned HOST\_WIDE\_INT nelems) [Target Hook]

Return the mode that GCC should use for an array that has nelems elements, with each element having mode mode. Return no mode if the target has no special requirements. In the latter case, GCC looks for an integer mode of the appropriate size if available and uses BLKmode otherwise. Usually the search for the integer mode is

limited to MAX\_FIXED\_MODE\_SIZE, but the TARGET\_ARRAY\_MODE\_SUPPORTED\_P hook allows a larger mode to be used in specific cases.

The main use of this hook is to specify that an array of vectors should also have a vector mode. The default implementation returns no mode.

### bool TARGET\_ARRAY\_MODE\_SUPPORTED\_P (machine\_mode mode, unsigned HOST\_WIDE\_INT nelems) [Target Hook]

Return true if GCC should try to use a scalar mode to store an array of *nelems* elements, given that each element has mode *mode*. Returning true here overrides the usual MAX\_FIXED\_MODE limit and allows GCC to use any defined integer mode.

One use of this hook is to support vector load and store operations that operate on several homogeneous vectors. For example, ARM NEON has operations like:

```
int8x8x3_t vld3_s8 (const int8_t *)
where the return type is defined as:
    typedef struct int8x8x3_t
    {
        int8x8_t val[3];
    } int8x8x3_t;
```

If this hook allows val to have a scalar mode, then int8x8x3\_t can have the same mode. GCC can then store int8x8x3\_ts in registers rather than forcing them onto the stack.

### bool TARGET\_LIBGCC\_FLOATING\_MODE\_SUPPORTED\_P [Target Hook] (scalar\_float\_mode mode)

Define this to return nonzero if libgcc provides support for the floating-point mode mode, which is known to pass TARGET\_SCALAR\_MODE\_SUPPORTED\_P. The default version of this hook returns true for all of SFmode, DFmode, XFmode and TFmode, if such modes exist.

### opt\_scalar\_float\_mode TARGET\_FLOATN\_MODE (int n, bool extended) [Target Hook]

Define this to return the machine mode to use for the type \_Floatn, if extended is false, or the type \_Floatnx, if extended is true. If such a type is not supported, return opt\_scalar\_float\_mode (). The default version of this hook returns SFmode for \_Float32, DFmode for \_Float64 and \_Float32x and TFmode for \_Float128, if those modes exist and satisfy the requirements for those types and pass TARGET\_SCALAR\_MODE\_SUPPORTED\_P and TARGET\_LIBGCC\_FLOATING\_MODE\_SUPPORTED\_P; for \_Float64x, it returns the first of XFmode and TFmode that exists and satisfies the same requirements; for other types, it returns opt\_scalar\_float\_mode (). The hook is only called for values of n and extended that are valid according to ISO/IEC TS 18661-3:2015; that is, n is one of 32, 64, 128, or, if extended is false, 16 or greater than 128 and a multiple of 32.

### bool TARGET\_FLOATN\_BUILTIN\_P (int func)

[Target Hook]

Define this to return true if the \_Floatn and \_Floatnx built-in functions should implicitly enable the built-in function without the \_\_builtin\_ prefix in addition to the normal built-in function with the \_\_builtin\_ prefix. The default is to only enable built-in functions without the \_\_builtin\_ prefix for the GNU C language. In strict

ANSI/ISO mode, the built-in function without the \_\_builtin\_ prefix is not enabled. The argument FUNC is the enum built\_in\_function id of the function to be enabled.

### bool TARGET\_SMALL\_REGISTER\_CLASSES\_FOR\_MODE\_P [Target Hook] (machine\_mode mode)

Define this to return nonzero for machine modes for which the port has small register classes. If this target hook returns nonzero for a given mode, the compiler will try to minimize the lifetime of registers in mode. The hook may be called with VOIDmode as argument. In this case, the hook is expected to return nonzero if it returns nonzero for any mode.

On some machines, it is risky to let hard registers live across arbitrary insns. Typically, these machines have instructions that require values to be in specific registers (like an accumulator), and reload will fail if the required hard register is used for another purpose across such an insn.

Passes before reload do not know which hard registers will be used in an instruction, but the machine modes of the registers set or used in the instruction are already known. And for some machines, register classes are small for, say, integer registers but not for floating point registers. For example, the AMD x86-64 architecture requires specific registers for the legacy x86 integer instructions, but there are many SSE registers for floating point operations. On such targets, a good strategy may be to return nonzero from this hook for INTEGRAL\_MODE\_P machine modes but zero for the SSE register classes.

The default version of this hook returns false for any mode. It is always safe to redefine this hook to return with a nonzero value. But if you unnecessarily define it, you will reduce the amount of optimizations that can be performed in some cases. If you do not define this hook to return a nonzero value when it is required, the compiler will run out of spill registers and print a fatal error message.

### 18.9.8 How Scalar Function Values Are Returned

This section discusses the macros that control returning scalars as values—values that can fit in registers.

### rtx TARGET\_FUNCTION\_VALUE (const\_tree ret\_type, const\_tree [Target Hook] fn\_decl\_or\_type, bool outgoing)

Define this to return an RTX representing the place where a function returns or receives a value of data type  $ret\_type$ , a tree node representing a data type.  $fn\_decl\_or\_type$  is a tree node representing FUNCTION\_DECL or FUNCTION\_TYPE of a function being called. If outgoing is false, the hook should compute the register in which the caller will see the return value. Otherwise, the hook should return an RTX representing the place where a function returns a value.

On many machines, only TYPE\_MODE (ret\_type) is relevant. (Actually, on most machines, scalar values are returned in the same place regardless of mode.) The value of the expression is usually a reg RTX for the hard register where the return value is stored. The value can also be a parallel RTX, if the return value is in multiple places. See TARGET\_FUNCTION\_ARG for an explanation of the parallel form. Note that the callee will populate every location specified in the parallel, but if the

first element of the parallel contains the whole return value, callers will use that element as the canonical location and ignore the others. The m68k port uses this type of parallel to return pointers in both '%a0' (the canonical location) and '%d0'.

If TARGET\_PROMOTE\_FUNCTION\_RETURN returns true, you must apply the same promotion rules specified in PROMOTE\_MODE if valtype is a scalar type.

If the precise function being called is known, *func* is a tree node (FUNCTION\_DECL) for it; otherwise, *func* is a null pointer. This makes it possible to use a different value-returning convention for specific functions when all their calls are known.

Some target machines have "register windows" so that the register in which a function returns its value is not the same as the one in which the caller sees the value. For such machines, you should return different RTX depending on *outgoing*.

TARGET\_FUNCTION\_VALUE is not used for return values with aggregate data types, because these are returned in another way. See TARGET\_STRUCT\_VALUE\_RTX and related macros, below.

### FUNCTION\_VALUE (valtype, func)

[Macro]

This macro has been deprecated. Use TARGET\_FUNCTION\_VALUE for a new target instead.

### LIBCALL\_VALUE (mode)

[Macro]

A C expression to create an RTX representing the place where a library function returns a value of mode *mode*.

Note that "library function" in this context means a compiler support routine, used to perform arithmetic, whose name is known specially by the compiler and was not mentioned in the C code being compiled.

rtx TARGET\_LIBCALL\_VALUE (machine\_mode mode, const\_rtx fun) [Target Hook]

Define this hook if the back-end needs to know the name of the libcall function in order to determine where the result should be returned.

The mode of the result is given by *mode* and the name of the called library function is given by *fun*. The hook should return an RTX representing the place where the library function result will be returned.

If this hook is not defined, then LIBCALL\_VALUE will be used.

### FUNCTION\_VALUE\_REGNO\_P (regno)

[Macro]

A C expression that is nonzero if *regno* is the number of a hard register in which the values of called function may come back.

A register whose use for returning values is limited to serving as the second of a pair (for a value of type double, say) need not be recognized by this macro. So for most machines, this definition suffices:

```
#define FUNCTION_VALUE_REGNO_P(N) ((N) == 0)
```

If the machine has register windows, so that the caller and the called function use different registers for the return value, this macro should recognize only the caller's register numbers.

This macro has been deprecated. Use TARGET\_FUNCTION\_VALUE\_REGNO\_P for a new target instead.

### bool TARGET\_FUNCTION\_VALUE\_REGNO\_P (const unsigned int regno)

[Target Hook]

A target hook that return **true** if *regno* is the number of a hard register in which the values of called function may come back.

A register whose use for returning values is limited to serving as the second of a pair (for a value of type double, say) need not be recognized by this target hook.

If the machine has register windows, so that the caller and the called function use different registers for the return value, this target hook should recognize only the caller's register numbers.

If this hook is not defined, then FUNCTION\_VALUE\_REGNO\_P will be used.

### APPLY\_RESULT\_SIZE

[Macro]

Define this macro if 'untyped\_call' and 'untyped\_return' need more space than is implied by FUNCTION\_VALUE\_REGNO\_P for saving and restoring an arbitrary return value.

### bool TARGET\_OMIT\_STRUCT\_RETURN\_REG

[Target Hook]

Normally, when a function returns a structure by memory, the address is passed as an invisible pointer argument, but the compiler also arranges to return the address from the function like it would a normal pointer return value. Define this to true if that behavior is undesirable on your target.

### bool TARGET\_RETURN\_IN\_MSB (const\_tree type)

[Target Hook]

This hook should return true if values of type type are returned at the most significant end of a register (in other words, if they are padded at the least significant end). You can assume that type is returned in a register; the caller is required to check this.

Note that the register provided by TARGET\_FUNCTION\_VALUE must be able to hold the complete return value. For example, if a 1-, 2- or 3-byte structure is returned at the most significant end of a 4-byte register, TARGET\_FUNCTION\_VALUE should provide an SImode rtx.

### 18.9.9 How Large Values Are Returned

When a function value's mode is BLKmode (and in some other cases), the value is not returned according to TARGET\_FUNCTION\_VALUE (see Section 18.9.8 [Scalar Return], page 544). Instead, the caller passes the address of a block of memory in which the value should be stored. This address is called the *structure value address*.

This section describes how to control returning structure values in memory.

### bool TARGET\_RETURN\_IN\_MEMORY (const\_tree type, const\_tree fntype)

[Target Hook]

This target hook should return a nonzero value to say to return the function value in memory, just as large structures are always returned. Here *type* will be the data type of the value, and *fntype* will be the type of the function doing the returning, or NULL for libcalls.

Note that values of mode BLKmode must be explicitly handled by this function. Also, the option '-fpcc-struct-return' takes effect regardless of this macro. On most

systems, it is possible to leave the hook undefined; this causes a default definition to be used, whose value is the constant 1 for BLKmode values, and 0 otherwise.

Do not use this hook to indicate that structures and unions should always be returned in memory. You should instead use <code>DEFAULT\_PCC\_STRUCT\_RETURN</code> to indicate this.

### DEFAULT\_PCC\_STRUCT\_RETURN

[Macro]

Define this macro to be 1 if all structure and union return values must be in memory. Since this results in slower code, this should be defined only if needed for compatibility with other compilers or with an ABI. If you define this macro to be 0, then the conventions used for structure and union return values are decided by the TARGET\_RETURN\_IN\_MEMORY target hook.

If not defined, this defaults to the value 1.

### rtx TARGET\_STRUCT\_VALUE\_RTX (tree fndecl, int incoming)

[Target Hook]

This target hook should return the location of the structure value address (normally a mem or reg), or 0 if the address is passed as an "invisible" first argument. Note that *fndecl* may be NULL, for libcalls. You do not need to define this target hook if the address is always passed as an "invisible" first argument.

On some architectures the place where the structure value address is found by the called function is not the same place that the caller put it. This can be due to register windows, or it could be because the function prologue moves it to a different place. *incoming* is 1 or 2 when the location is needed in the context of the called function, and 0 in the context of the caller.

If *incoming* is nonzero and the address is to be found on the stack, return a mem which refers to the frame pointer. If *incoming* is 2, the result is being used to fetch the structure value address at the beginning of a function. If you need to emit adjusting code, you should do it at this point.

### PCC\_STATIC\_STRUCT\_RETURN

Macro

Define this macro if the usual system convention on the target machine for returning structures and unions is for the called function to return the address of a static variable containing the value.

Do not define this if the usual system convention is for the caller to pass an address to the subroutine.

This macro has effect in '-fpcc-struct-return' mode, but it does nothing when you use '-freg-struct-return' mode.

### fixed\_size\_mode TARGET\_GET\_RAW\_RESULT\_MODE (int regno)

[Target Hook]

This target hook returns the mode to be used when accessing raw return registers in \_\_builtin\_return. Define this macro if the value in reg\_raw\_mode is not correct.

### fixed\_size\_mode TARGET\_GET\_RAW\_ARG\_MODE (int regno)

[Target Hook]

This target hook returns the mode to be used when accessing raw argument registers in \_\_builtin\_apply\_args. Define this macro if the value in reg\_raw\_mode is not correct.

### bool TARGET\_EMPTY\_RECORD\_P (const\_tree type)

[Target Hook]

This target hook returns true if the type is an empty record. The default is to return false.

#### 

This target hook warns about the change in empty class parameter passing ABI.

### 18.9.10 Caller-Saves Register Allocation

If you enable it, GCC can save registers around function calls. This makes it possible to use call-clobbered registers to hold variables that must live across calls.

### HARD\_REGNO\_CALLER\_SAVE\_MODE (regno, nregs)

[Macro]

A C expression specifying which mode is required for saving *nregs* of a pseudo-register in call-clobbered hard register *regno*. If *regno* is unsuitable for caller save, VOIDmode should be returned. For most machines this macro need not be defined since GCC will select the smallest suitable mode.

### 18.9.11 Function Entry and Exit

This section describes the macros that output function entry (prologue) and exit (epilogue) code.

### void TARGET\_ASM\_PRINT\_PATCHABLE\_FUNCTION\_ENTRY (FILE [Target Hook] \*file, unsigned HOST\_WIDE\_INT patch\_area\_size, bool record\_p)

Generate a patchable area at the function start, consisting of patch\_area\_size NOP instructions. If the target supports named sections and if record\_p is true, insert a pointer to the current location in the table of patchable functions. The default implementation of the hook places the table of pointers in the special section named \_\_patchable\_function\_entries.

### void TARGET\_ASM\_FUNCTION\_PROLOGUE (FILE \*file) [Target Hook]

If defined, a function that outputs the assembler code for entry to a function. The prologue is responsible for setting up the stack frame, initializing the frame pointer register, saving registers that must be saved, and allocating *size* additional bytes of storage for the local variables. *file* is a stdio stream to which the assembler code should be output.

The label for the beginning of the function need not be output by this macro. That has already been done when the macro is run.

To determine which registers to save, the macro can refer to the array  $regs\_ever\_$  live: element r is nonzero if hard register r is used anywhere within the function. This implies the function prologue should save register r, provided it is not one of the call-used registers. (TARGET\_ASM\_FUNCTION\_EPILOGUE must likewise use  $regs\_ever\_$  live.)

On machines that have "register windows", the function entry code does not save on the stack the registers that are in the windows, even if they are supposed to be preserved by function calls; instead it takes appropriate steps to "push" the register stack, if any non-call-used registers are used in the function.

On machines where functions may or may not have frame-pointers, the function entry code must vary accordingly; it must set up the frame pointer if one is wanted, and not otherwise. To determine whether a frame pointer is in wanted, the macro can refer

to the variable frame\_pointer\_needed. The variable's value will be 1 at run time in a function that needs a frame pointer. See Section 18.9.5 [Elimination], page 533.

The function entry code is responsible for allocating any stack space required for the function. This stack space consists of the regions listed below. In most cases, these regions are allocated in the order listed, with the last listed region closest to the top of the stack (the lowest address if STACK\_GROWS\_DOWNWARD is defined, and the highest address if it is not defined). You can use a different order for a machine if doing so is more convenient or required for compatibility reasons. Except in cases where required by standard or by a debugger, there is no reason why the stack layout used by GCC need agree with that used by other compilers for a machine.

# void TARGET\_ASM\_FUNCTION\_END\_PROLOGUE (FILE \*file) [Target Hook] If defined, a function that outputs assembler code at the end of a prologue. This should be used when the function prologue is being emitted as RTL, and you have some extra assembler that needs to be emitted. See [prologue instruction pattern], page 425.

void TARGET\_ASM\_FUNCTION\_BEGIN\_EPILOGUE (FILE \*file) [Target Hook] If defined, a function that outputs assembler code at the start of an epilogue. This should be used when the function epilogue is being emitted as RTL, and you have some extra assembler that needs to be emitted. See [epilogue instruction pattern], page 425.

### void TARGET\_ASM\_FUNCTION\_EPILOGUE (FILE \*file) [Target Hook]

If defined, a function that outputs the assembler code for exit from a function. The epilogue is responsible for restoring the saved registers and stack pointer to their values when the function was called, and returning control to the caller. This macro takes the same argument as the macro <code>TARGET\_ASM\_FUNCTION\_PROLOGUE</code>, and the registers to restore are determined from <code>regs\_ever\_live</code> and <code>CALL\_USED\_REGISTERS</code> in the same way.

On some machines, there is a single instruction that does all the work of returning from the function. On these machines, give that instruction the name 'return' and do not define the macro TARGET\_ASM\_FUNCTION\_EPILOGUE at all.

Do not define a pattern named 'return' if you want the TARGET\_ASM\_FUNCTION\_ EPILOGUE to be used. If you want the target switches to control whether return instructions or epilogues are used, define a 'return' pattern with a validity condition that tests the target switches appropriately. If the 'return' pattern's validity condition is false, epilogues will be used.

On machines where functions may or may not have frame-pointers, the function exit code must vary accordingly. Sometimes the code for these two cases is completely different. To determine whether a frame pointer is wanted, the macro can refer to the variable frame\_pointer\_needed. The variable's value will be 1 when compiling a function that needs a frame pointer.

Normally, TARGET\_ASM\_FUNCTION\_PROLOGUE and TARGET\_ASM\_FUNCTION\_EPILOGUE must treat leaf functions specially. The C variable current\_function\_is\_leaf is nonzero for such a function. See Section 18.7.4 [Leaf Functions], page 510.

On some machines, some functions pop their arguments on exit while others leave that for the caller to do. For example, the 68020 when given '-mrtd' pops arguments in functions that take a fixed number of arguments.

Your definition of the macro RETURN\_POPS\_ARGS decides which functions pop their own arguments. TARGET\_ASM\_FUNCTION\_EPILOGUE needs to know what was decided. The number of bytes of the current function's arguments that this function should pop is available in crtl->args.pops\_args. See Section 18.9.8 [Scalar Return], page 544.

- A region of crtl->args.pretend\_args\_size bytes of uninitialized space just underneath the first argument arriving on the stack. (This may not be at the very start of the allocated stack region if the calling sequence has pushed anything else since pushing the stack arguments. But usually, on such machines, nothing else has been pushed yet, because the function prologue itself does all the pushing.) This region is used on machines where an argument may be passed partly in registers and partly in memory, and, in some cases to support the features in <stdarg.h>.
- An area of memory used to save certain registers used by the function. The size of this
  area, which may also include space for such things as the return address and pointers
  to previous stack frames, is machine-specific and usually depends on which registers
  have been used in the function. Machines with register windows often do not require a
  save area.
- A region of at least *size* bytes, possibly rounded up to an allocation boundary, to contain the local variables of the function. On some machines, this region and the save area may occur in the opposite order, with the save area closer to the top of the stack.
- Optionally, when ACCUMULATE\_OUTGOING\_ARGS is defined, a region of crt1->outgoing\_args\_size bytes to be used for outgoing argument lists of the function. See Section 18.9.6 [Stack Arguments], page 534.

### EXIT\_IGNORE\_STACK

[Macro] etion or the

Define this macro as a C expression that is nonzero if the return instruction or the function epilogue ignores the value of the stack pointer; in other words, if it is safe to delete an instruction to adjust the stack pointer before a return from the function. The default is 0.

Note that this macro's value is relevant only for functions for which frame pointers are maintained. It is never safe to delete a final stack adjustment in a function that has no frame pointer, and the compiler knows this regardless of EXIT\_IGNORE\_STACK.

### EPILOGUE\_USES (regno)

[Macro]

[Macro]

Define this macro as a C expression that is nonzero for registers that are used by the epilogue or the 'return' pattern. The stack and frame pointer registers are already assumed to be used as needed.

### EH\_USES (regno)

Define this macro as a C expression that is nonzero for registers that are used by the exception handling mechanism, and so should be considered live on entry to an exception edge.

#### 

A function that outputs the assembler code for a thunk function, used to implement C++ virtual function calls with multiple inheritance. The thunk acts as a wrapper around a virtual function, adjusting the implicit object parameter before handing control off to the real function.

First, emit code to add the integer *delta* to the location that contains the incoming first argument. Assume that this argument contains a pointer, and is the one used to pass the **this** pointer in C++. This is the incoming argument *before* the function prologue, e.g. '%o0' on a sparc. The addition must preserve the values of all other incoming arguments.

Then, if  $vcall\_offset$  is nonzero, an additional adjustment should be made after adding delta. In particular, if p is the adjusted pointer, the following adjustment should be made:

```
p += (*((ptrdiff_t **)p))[vcall_offset/sizeof(ptrdiff_t)]
```

After the additions, emit code to jump to function, which is a FUNCTION\_DECL. This is a direct pure jump, not a call, and does not touch the return address. Hence returning from FUNCTION will return to whoever called the current 'thunk'.

The effect must be as if function had been called directly with the adjusted first argument. This macro is responsible for emitting all of the code for a thunk function; TARGET\_ASM\_FUNCTION\_PROLOGUE and TARGET\_ASM\_FUNCTION\_EPILOGUE are not invoked.

The thunk\_fndecl is redundant. (delta and function have already been extracted from it.) It might possibly be useful on some targets, but probably not.

If you do not define this macro, the target-independent code in the C++ front end will generate a less efficient heavyweight thunk that calls *function* instead of jumping to it. The generic approach does not support varargs.

```
bool TARGET_ASM_CAN_OUTPUT_MI_THUNK (const_tree [Target Hook] thunk_fndec1, HOST_WIDE_INT delta, HOST_WIDE_INT vcall_offset, const_tree function) [Target Hook]
```

A function that returns true if TARGET\_ASM\_OUTPUT\_MI\_THUNK would be able to output the assembler code for the thunk function specified by the arguments it is passed, and false otherwise. In the latter case, the generic approach will be used by the C++ front end, with the limitations previously exposed.

### 18.9.12 Generating Code for Profiling

These macros will help you generate code for profiling.

### FUNCTION\_PROFILER (file, labelno)

[Macro]

A C statement or compound statement to output to file some assembler code to call the profiling subroutine mcount.

The details of how mcount expects to be called are determined by your operating system environment, not by GCC. To figure them out, compile a small program for

profiling using the system's installed C compiler and look at the assembler code that results.

Older implementations of mcount expect the address of a counter variable to be loaded into some register. The name of this variable is 'LP' followed by the number *labelno*, so you would generate the name using 'LP%d' in a fprintf.

PROFILE\_HOOK [Macro]

A C statement or compound statement to output to file some assembly code to call the profiling subroutine mcount even the target does not support profiling.

### NO\_PROFILE\_COUNTERS

[Macro]

Define this macro to be an expression with a nonzero value if the mcount subroutine on your system does not need a counter variable allocated for each function. This is true for almost all modern implementations. If you define this macro, you must not use the *labelno* argument to FUNCTION\_PROFILER.

### PROFILE\_BEFORE\_PROLOGUE

[Macro]

Define this macro if the code for function profiling should come before the function prologue. Normally, the profiling code comes after.

### bool TARGET\_KEEP\_LEAF\_WHEN\_PROFILED (void)

[Target Hook]

This target hook returns true if the target wants the leaf flag for the current function to stay true even if it calls mount. This might make sense for targets using the leaf flag only to determine whether a stack frame needs to be generated or not and for which the call to mount is generated before the function prologue.

### 18.9.13 Permitting tail calls

bool TARGET\_FUNCTION\_OK\_FOR\_SIBCALL (tree decl, tree exp) [Target Hook] True if it is OK to do sibling call optimization for the specified call expression exp. decl will be the called function, or NULL if this is an indirect call.

It is not uncommon for limitations of calling conventions to prevent tail calls to functions outside the current unit of translation, or during PIC compilation. The hook is used to enforce these restrictions, as the sibcall md pattern cannot fail, or fall over to a "normal" call. The criteria for successful sibling call optimization may vary greatly between different architectures.

### void TARGET\_EXTRA\_LIVE\_ON\_ENTRY (bitmap regs)

[Target Hook]

Add any hard registers to regs that are live on entry to the function. This hook only needs to be defined to provide registers that cannot be found by examination of FUNCTION\_ARG\_REGNO\_P, the callee saved registers, STATIC\_CHAIN\_INCOMING\_REGNUM, STATIC\_CHAIN\_REGNUM, TARGET\_STRUCT\_VALUE\_RTX, FRAME\_POINTER\_REGNUM, EH\_USES, FRAME\_POINTER\_REGNUM, ARG\_POINTER\_REGNUM, and the PIC\_OFFSET\_TABLE\_REGNUM.

### 

This hook should add additional registers that are computed by the prologue to the hard regset for shrink-wrapping optimization purposes.

### bool TARGET\_WARN\_FUNC\_RETURN (tree)

[Target Hook]

True if a function's return statements should be checked for matching the function's return type. This includes checking for falling off the end of a non-void function. Return false if no such check should be made.

### 18.9.14 Shrink-wrapping separate components

The prologue may perform a variety of target dependent tasks such as saving callee-saved registers, saving the return address, aligning the stack, creating a stack frame, initializing the PIC register, setting up the static chain, etc.

On some targets some of these tasks may be independent of others and thus may be shrink-wrapped separately. These independent tasks are referred to as components and are handled generically by the target independent parts of GCC.

Using the following hooks those prologue or epilogue components can be shrink-wrapped separately, so that the initialization (and possibly teardown) those components do is not done as frequently on execution paths where this would unnecessary.

What exactly those components are is up to the target code; the generic code treats them abstractly, as a bit in an sbitmap. These sbitmaps are allocated by the shrink\_wrap.get\_separate\_components and shrink\_wrap.components\_for\_bb hooks, and deallocated by the generic code.

- sbitmap TARGET\_SHRINK\_WRAP\_GET\_SEPARATE\_COMPONENTS (void) [Target Hook] This hook should return an sbitmap with the bits set for those components that can be separately shrink-wrapped in the current function. Return NULL if the current function should not get any separate shrink-wrapping. Don't define this hook if it would always return NULL. If it is defined, the other hooks in this group have to be defined as well.
- sbitmap TARGET\_SHRINK\_WRAP\_COMPONENTS\_FOR\_BB (basic\_block) [Target Hook] This hook should return an sbitmap with the bits set for those components where either the prologue component has to be executed before the basic\_block, or the epilogue component after it, or both.

### void TARGET\_SHRINK\_WRAP\_DISQUALIFY\_COMPONENTS (sbitmap components, edge e, sbitmap edge\_components, bool is\_prologue)

This hook should clear the bits in the *components* bitmap for those components in *edge\_components* that the target cannot handle on edge *e*, where *is\_prologue* says if this is for a prologue or an epilogue instead.

#### 

Emit prologue insps for the components indicated by the parameter.

### void TARGET\_SHRINK\_WRAP\_EMIT\_EPILOGUE\_COMPONENTS [Target Hook] (sbitmap)

Emit epilogue insns for the components indicated by the parameter.

void TARGET\_SHRINK\_WRAP\_SET\_HANDLED\_COMPONENTS (sbitmap) [Target Hook] Mark the components in the parameter as handled, so that the prologue and epilogue named patterns know to ignore those components. The target code should not hang on to the sbitmap, it will be deleted after this call.

### 18.9.15 Stack smashing protection

### tree TARGET\_STACK\_PROTECT\_GUARD (void)

[Target Hook]

This hook returns a DECL node for the external variable to use for the stack protection guard. This variable is initialized by the runtime to some random value and is used to initialize the guard value that is placed at the top of the local stack frame. The type of this variable must be ptr\_type\_node.

The default version of this hook creates a variable called '\_\_stack\_chk\_guard', which is normally defined in 'libgcc2.c'.

### tree TARGET\_STACK\_PROTECT\_FAIL (void)

[Target Hook]

This hook returns a CALL\_EXPR that alerts the runtime that the stack protect guard variable has been modified. This expression should involve a call to a noreturn function.

The default version of this hook invokes a function called '\_\_stack\_chk\_fail', taking no arguments. This function is normally defined in 'libgcc2.c'.

### bool TARGET\_STACK\_PROTECT\_RUNTIME\_ENABLED\_P (void)

[Target Hook]

Returns true if the target wants GCC's default stack protect runtime support, otherwise return false. The default implementation always returns true.

### bool TARGET\_SUPPORTS\_SPLIT\_STACK (bool report,

[Common Target Hook]

struct gcc\_options \*opts)

Whether this target supports splitting the stack when the options described in *opts* have been passed. This is called after options have been parsed, so the target may reject splitting the stack in some configurations. The default version of this hook returns false. If *report* is true, this function may issue a warning or error; if *report* is false, it must simply return a value

#### 

The hook is used for options that have a non-trivial list of possible option values. OPTION\_CODE is option code of opt\_code enum type. PREFIX is used for bash completion and allows an implementation to return more specific completion based on the prefix. All string values should be allocated from heap memory and consumers should release them. The result will be pruned to cases with PREFIX if not NULL.

### 18.9.16 Miscellaneous register hooks

### bool TARGET\_CALL\_FUSAGE\_CONTAINS\_NON\_CALLEE\_CLOBBERS [Target Hook] Set to true if each call that binds to a local definition explicitly clobbers or sets

all non-fixed registers modified by performing the call. That is, by the call pattern itself, or by code that might be inserted by the linker (e.g. stubs, veneers,

branch islands), but not including those modifiable by the callee. The affected registers may be mentioned explicitly in the call pattern, or included as clobbers in CALL\_INSN\_FUNCTION\_USAGE. The default version of this hook is set to false. The purpose of this hook is to enable the fipa-ra optimization.

### 18.10 Implementing the Varargs Macros

GCC comes with an implementation of <varargs.h> and <stdarg.h> that work without change on machines that pass arguments on the stack. Other machines require their own implementations of varargs, and the two machine independent header files must have conditionals to include it.

ISO <stdarg.h> differs from traditional <varargs.h> mainly in the calling convention for va\_start. The traditional implementation takes just one argument, which is the variable in which to store the argument pointer. The ISO implementation of va\_start takes an additional second argument. The user is supposed to write the last named argument of the function here.

However, va\_start should not use this argument. The way to find the end of the named arguments is with the built-in functions described below.

### \_\_builtin\_saveregs ()

[Macro]

Use this built-in function to save the argument registers in memory so that the varargs mechanism can access them. Both ISO and traditional versions of va\_start must use \_\_builtin\_saveregs, unless you use TARGET\_SETUP\_INCOMING\_VARARGS (see below) instead.

On some machines, \_\_builtin\_saveregs is open-coded under the control of the target hook TARGET\_EXPAND\_BUILTIN\_SAVEREGS. On other machines, it calls a routine written in assembler language, found in 'libgcc2.c'.

Code generated for the call to \_\_builtin\_saveregs appears at the beginning of the function, as opposed to where the call to \_\_builtin\_saveregs is written, regardless of what the code is. This is because the registers must be saved before the function starts to use them for its own purposes.

### \_\_builtin\_next\_arg (lastarg)

[Macro]

This builtin returns the address of the first anonymous stack argument, as type void \*. If ARGS\_GROW\_DOWNWARD, it returns the address of the location above the first anonymous stack argument. Use it in va\_start to initialize the pointer for fetching arguments from the stack. Also use it in va\_start to verify that the second parameter lastarg is the last named argument of the current function.

### \_\_builtin\_classify\_type (object)

[Macro]

Since each machine has its own conventions for which data types are passed in which kind of register, your implementation of va\_arg has to embody these conventions. The easiest way to categorize the specified data type is to use \_\_builtin\_classify\_type together with sizeof and \_\_alignof\_\_.

\_\_builtin\_classify\_type ignores the value of *object*, considering only its data type. It returns an integer describing what kind of type that is—integer, floating, pointer, structure, and so on.

The file 'typeclass.h' defines an enumeration that you can use to interpret the values of \_\_builtin\_classify\_type.

These machine description macros help implement varargs:

### rtx TARGET\_EXPAND\_BUILTIN\_SAVEREGS (void)

[Target Hook]

If defined, this hook produces the machine-specific code for a call to \_\_builtin\_saveregs. This code will be moved to the very beginning of the function, before any parameter access are made. The return value of this function should be an RTX that contains the value to use as the return of \_\_builtin\_saveregs.

### void TARGET\_SETUP\_INCOMING\_VARARGS (cumulative\_args\_t

[Target Hook]

args\_so\_far, const function\_arg\_info &arg, int \*pretend\_args\_size, int
second\_time)

This target hook offers an alternative to using \_\_builtin\_saveregs and defining the hook TARGET\_EXPAND\_BUILTIN\_SAVEREGS. Use it to store the anonymous register arguments into the stack so that all the arguments appear to have been passed consecutively on the stack. Once this is done, you can use the standard implementation of varargs that works for machines that pass all their arguments on the stack.

The argument  $args\_so\_far$  points to the CUMULATIVE\_ARGS data structure, containing the values that are obtained after processing the named arguments. The argument arg describes the last of these named arguments.

The target hook should do two things: first, push onto the stack all the argument registers *not* used for the named arguments, and second, store the size of the data thus pushed into the int-valued variable pointed to by *pretend\_args\_size*. The value that you store here will serve as additional offset for setting up the stack frame.

Because you must generate code to push the anonymous arguments at compile time without knowing their data types, TARGET\_SETUP\_INCOMING\_VARARGS is only useful on machines that have just a single category of argument register and use it uniformly for all data types.

If the argument <code>second\_time</code> is nonzero, it means that the arguments of the function are being analyzed for the second time. This happens for an inline function, which is not actually compiled until the end of the source file. The hook <code>TARGET\_SETUP\_INCOMING\_VARARGS</code> should not generate any instructions in this case.

bool TARGET\_STRICT\_ARGUMENT\_NAMING (cumulative\_args\_t ca) [Target Hook]

Define this hook to return true if the location where a function argument is passed depends on whether or not it is a named argument.

This hook controls how the *named* argument to TARGET\_FUNCTION\_ARG is set for varargs and stdarg functions. If this hook returns true, the *named* argument is always true for named arguments, and false for unnamed arguments. If it returns false, but TARGET\_PRETEND\_OUTGOING\_VARARGS\_NAMED returns true, then all arguments are treated as named. Otherwise, all named arguments except the last are treated as named.

You need not define this hook if it always returns false.

### void TARGET\_CALL\_ARGS (rtx, tree)

[Target Hook]

While generating RTL for a function call, this target hook is invoked once for each argument passed to the function, either a register returned by TARGET\_FUNCTION\_ARG or a memory location. It is called just before the point where argument registers are stored. The type of the function to be called is also passed as the second argument; it is NULL\_TREE for libcalls. The TARGET\_END\_CALL\_ARGS hook is invoked just after the code to copy the return reg has been emitted. This functionality can be used to perform special setup of call argument registers if a target needs it. For functions without arguments, the hook is called once with pc\_rtx passed instead of an argument register. Most ports do not need to implement anything for this hook.

### void TARGET\_END\_CALL\_ARGS (void)

[Target Hook]

This target hook is invoked while generating RTL for a function call, just after the point where the return reg is copied into a pseudo. It signals that all the call argument and return registers for the just emitted call are now no longer in use. Most ports do not need to implement anything for this hook.

### bool TARGET\_PRETEND\_OUTGOING\_VARARGS\_NAMED

[Target Hook]

(cumulative\_args\_t ca)

If you need to conditionally change ABIs so that one works with TARGET\_SETUP\_INCOMING\_VARARGS, but the other works like neither TARGET\_SETUP\_INCOMING\_VARARGS nor TARGET\_STRICT\_ARGUMENT\_NAMING was defined, then define this hook to return true if TARGET\_SETUP\_INCOMING\_VARARGS is used, false otherwise. Otherwise, you should not define this hook.

### rtx TARGET\_LOAD\_BOUNDS\_FOR\_ARG (rtx slot, rtx arg, rtx slot\_no)

[Target Hook]

This hook is used by expand pass to emit inso to load bounds of arg passed in slot. Expand pass uses this hook in case bounds of arg are not passed in register. If slot is a memory, then bounds are loaded as for regular pointer loaded from memory. If slot is not a memory then slot\_no is an integer constant holding number of the target dependent special slot which should be used to obtain bounds. Hook returns RTX holding loaded bounds.

### void TARGET\_STORE\_BOUNDS\_FOR\_ARG (rtx arg, rtx slot, rtx bounds, rtx slot\_no) [Target Hook]

This hook is used by expand pass to emit insns to store bounds of arg passed in slot. Expand pass uses this hook in case bounds of arg are not passed in register. If slot is a memory, then bounds are stored as for regular pointer stored in memory. If slot is not a memory then slot\_no is an integer constant holding number of the target dependent special slot which should be used to store bounds.

### rtx TARGET\_LOAD\_RETURNED\_BOUNDS (rtx slot)

[Target Hook]

This hook is used by expand pass to emit inso to load bounds returned by function call in *slot*. Hook returns RTX holding loaded bounds.

## void TARGET\_STORE\_RETURNED\_BOUNDS (rtx slot, rtx bounds) [Target Hook] This hook is used by expand pass to emit insn to store bounds returned by function call into slot.

### 18.11 Support for Nested Functions

Taking the address of a nested function requires special compiler handling to ensure that the static chain register is loaded when the function is invoked via an indirect call.

GCC has traditionally supported nested functions by creating an executable *trampoline* at run time when the address of a nested function is taken. This is a small piece of code which normally resides on the stack, in the stack frame of the containing function. The trampoline loads the static chain register and then jumps to the real address of the nested function.

The use of trampolines requires an executable stack, which is a security risk. To avoid this problem, GCC also supports another strategy: using descriptors for nested functions. Under this model, taking the address of a nested function results in a pointer to a non-executable function descriptor object. Initializing the static chain from the descriptor is handled at indirect call sites.

On some targets, including HPPA and IA-64, function descriptors may be mandated by the ABI or be otherwise handled in a target-specific way by the back end in its code generation strategy for indirect calls. GCC also provides its own generic descriptor implementation to support the '-fno-trampolines' option. In this case runtime detection of function descriptors at indirect call sites relies on descriptor pointers being tagged with a bit that is never set in bare function addresses. Since GCC's generic function descriptors are not ABI-compliant, this option is typically used only on a per-language basis (notably by Ada) or when it can otherwise be applied to the whole program.

Define the following hook if your backend either implements ABI-specified descriptor support, or can use GCC's generic descriptor implementation for nested functions.

### int TARGET\_CUSTOM\_FUNCTION\_DESCRIPTORS

|Target Hook

If the target can use GCC's generic descriptor mechanism for nested functions, define this hook to a power of 2 representing an unused bit in function pointers which can be used to differentiate descriptors at run time. This value gives the number of bytes by which descriptor pointers are misaligned compared to function pointers. For example, on targets that require functions to be aligned to a 4-byte boundary, a value of either 1 or 2 is appropriate unless the architecture already reserves the bit for another purpose, such as on ARM.

Define this hook to 0 if the target implements ABI support for function descriptors in its standard calling sequence, like for example HPPA or IA-64.

Using descriptors for nested functions eliminates the need for trampolines that reside on the stack and require it to be made executable.

The following macros tell GCC how to generate code to allocate and initialize an executable trampoline. You can also use this interface if your back end needs to create ABI-specified non-executable descriptors; in this case the "trampoline" created is the descriptor containing data only.

The instructions in an executable trampoline must do two things: load a constant address into the static chain register, and jump to the real address of the nested function. On CISC machines such as the m68k, this requires two instructions, a move immediate and a jump. Then the two addresses exist in the trampoline as word-long immediate operands. On RISC

machines, it is often necessary to load each address into a register in two parts. Then pieces of each address form separate immediate operands.

The code generated to initialize the trampoline must store the variable parts—the static chain value and the function address—into the immediate operands of the instructions. On a CISC machine, this is simply a matter of copying each address to a memory reference at the proper offset from the start of the trampoline. On a RISC machine, it may be necessary to take out pieces of the address and store them separately.

#### void TARGET\_ASM\_TRAMPOLINE\_TEMPLATE (FILE \*f)

[Target Hook]

This hook is called by  $assemble_trampoline_template$  to output, on the stream f, assembler code for a block of data that contains the constant parts of a trampoline. This code should not include a label—the label is taken care of automatically.

If you do not define this hook, it means no template is needed for the target. Do not define this hook on systems where the block move code to copy the trampoline into place would be larger than the code to generate it on the spot.

#### TRAMPOLINE\_SECTION

[Macro]

Return the section into which the trampoline template is to be placed (see Section 18.18 [Sections], page 590). The default value is readonly\_data\_section.

TRAMPOLINE\_SIZE

[Macro]

A C expression for the size in bytes of the trampoline, as an integer.

### TRAMPOLINE\_ALIGNMENT

[Macro]

Alignment required for trampolines, in bits.

If you don't define this macro, the value of FUNCTION\_ALIGNMENT is used for aligning trampolines.

#### 

This hook is called to initialize a trampoline.  $m_t tramp$  is an RTX for the memory block for the trampoline; fndecl is the FUNCTION\_DECL for the nested function;  $static\_chain$  is an RTX for the static chain value that should be passed to the function when it is called.

If the target defines TARGET\_ASM\_TRAMPOLINE\_TEMPLATE, then the first thing this hook should do is emit a block move into  $m\_tramp$  from the memory block returned by assemble\_trampoline\_template. Note that the block move need only cover the constant parts of the trampoline. If the target isolates the variable parts of the trampoline to the end, not all TRAMPOLINE\_SIZE bytes need be copied.

If the target requires any other actions, such as flushing caches or enabling stack execution, these actions should be performed after initializing the trampoline proper.

### rtx TARGET\_TRAMPOLINE\_ADJUST\_ADDRESS (rtx addr)

[Target Hook]

This hook should perform any machine-specific adjustment in the address of the trampoline. Its argument contains the address of the memory block that was passed to TARGET\_TRAMPOLINE\_INIT. In case the address to be used for a function call should be different from the address at which the template was stored, the different address should be returned; otherwise addr should be returned unchanged. If this hook is not defined, addr will be used for function calls.

Implementing trampolines is difficult on many machines because they have separate instruction and data caches. Writing into a stack location fails to clear the memory in the instruction cache, so when the program jumps to that location, it executes the old contents.

Here are two possible solutions. One is to clear the relevant parts of the instruction cache whenever a trampoline is set up. The other is to make all trampolines identical, by having them jump to a standard subroutine. The former technique makes trampoline execution faster; the latter makes initialization faster.

To clear the instruction cache when a trampoline is initialized, define the following macro.

### CLEAR\_INSN\_CACHE (beg, end)

[Macro]

If defined, expands to a C expression clearing the *instruction cache* in the specified interval. The definition of this macro would typically be a series of asm statements. Both beg and end are both pointer expressions.

To use a standard subroutine, define the following macro. In addition, you must make sure that the instructions in a trampoline fill an entire cache line with identical instructions, or else ensure that the beginning of the trampoline code is always aligned at the same point in its cache line. Look in 'm68k.h' as a guide.

### TRANSFER\_FROM\_TRAMPOLINE

[Macro]

Define this macro if trampolines need a special subroutine to do their work. The macro should expand to a series of asm statements which will be compiled with GCC. They go in a library function named \_\_transfer\_from\_trampoline.

If you need to avoid executing the ordinary prologue code of a compiled C function when you jump to the subroutine, you can do so by placing a special label of your own in the assembler code. Use one asm statement to generate an assembler label, and another to make the label global. Then trampolines can use that label to jump directly to your special assembler code.

### 18.12 Implicit Calls to Library Routines

Here is an explanation of implicit calls to library routines.

### DECLARE\_LIBRARY\_RENAMES

[Macro]

This macro, if defined, should expand to a piece of C code that will get expanded when compiling functions for libgcc.a. It can be used to provide alternate names for GCC's internal library functions if there are ABI-mandated names that the compiler should provide.

### void TARGET\_INIT\_LIBFUNCS (void)

[Target Hook]

This hook should declare additional library routines or rename existing ones, using the functions set\_optab\_libfunc and init\_one\_libfunc defined in 'optabs.c'. init\_optabs calls this macro after initializing all the normal library routines.

The default is to do nothing. Most ports don't need to define this hook.

#### bool TARGET\_LIBFUNC\_GNU\_PREFIX

[Target Hook]

If false (the default), internal library routines start with two underscores. If set to true, these routines start with <code>\_\_gnu\_</code> instead. E.g., <code>\_\_muldi3</code> changes to <code>\_\_gnu\_</code> muldi3. This currently only affects functions defined in 'libgcc2.c'. If this is set to true, the 'tm.h' file must also #define LIBGCC2\_GNU\_PREFIX.

### FLOAT\_LIB\_COMPARE\_RETURNS\_BOOL (mode, comparison)

[Macro]

This macro should return **true** if the library routine that implements the floating point comparison operator *comparison* in mode *mode* will return a boolean, and *false* if it will return a tristate.

GCC's own floating point libraries return tristates from the comparison operators, so the default returns false always. Most ports don't need to define this macro.

### TARGET\_LIB\_INT\_CMP\_BIASED

[Macro]

This macro should evaluate to true if the integer comparison functions (like  $\_$ cmpdi2) return 0 to indicate that the first operand is smaller than the second, 1 to indicate that they are equal, and 2 to indicate that the first operand is greater than the second. If this macro evaluates to false the comparison functions return -1, 0, and 1 instead of 0, 1, and 2. If the target uses the routines in 'libgcc.a', you do not need to define this macro.

### TARGET\_HAS\_NO\_HW\_DIVIDE

[Macro]

This macro should be defined if the target has no hardware divide instructions. If this macro is defined, GCC will use an algorithm which make use of simple logical and arithmetic operations for 64-bit division. If the macro is not defined, GCC will use an algorithm which make use of a 64-bit by 32-bit divide primitive.

TARGET\_EDOM [Macro]

The value of EDOM on the target machine, as a C integer constant expression. If you don't define this macro, GCC does not attempt to deposit the value of EDOM into errno directly. Look in '/usr/include/errno.h' to find the value of EDOM on your system.

If you do not define TARGET\_EDOM, then compiled code reports domain errors by calling the library function and letting it report the error. If mathematical functions on your system use matherr when there is an error, then you should leave TARGET\_EDOM undefined so that matherr is used normally.

GEN\_ERRNO\_RTX [Macro]

Define this macro as a C expression to create an rtl expression that refers to the global "variable" errno. (On certain systems, errno may not actually be a variable.) If you don't define this macro, a reasonable default is used.

bool TARGET\_LIBC\_HAS\_FUNCTION (enum function\_class fn\_class) [Target Hook] This hook determines whether a function from a class of functions fn\_class is present in the target C library.

### bool TARGET\_LIBC\_HAS\_FAST\_FUNCTION (int fcode)

[Target Hook]

This hook determines whether a function from a class of functions (enum function\_class) fcode has a fast implementation.

### NEXT\_OBJC\_RUNTIME

|Macro|

Set this macro to 1 to use the "NeXT" Objective-C message sending conventions by default. This calling convention involves passing the object, the selector and the method arguments all at once to the method-lookup library function. This is the usual

setting when targeting Darwin/Mac OS X systems, which have the NeXT runtime installed.

If the macro is set to 0, the "GNU" Objective-C message sending convention will be used by default. This convention passes just the object and the selector to the method-lookup function, which returns a pointer to the method.

In either case, it remains possible to select code-generation for the alternate scheme, by means of compiler command line switches.

### 18.13 Addressing Modes

This is about addressing modes.

HAVE\_PRE\_INCREMENT [Macro]
HAVE\_PRE\_DECREMENT [Macro]
HAVE\_POST\_INCREMENT [Macro]
HAVE\_POST\_DECREMENT [Macro]

A C expression that is nonzero if the machine supports pre-increment, pre-decrement, post-increment, or post-decrement addressing respectively.

### HAVE\_PRE\_MODIFY\_DISP

[Macro]

HAVE\_POST\_MODIFY\_DISP

[Macro]

A C expression that is nonzero if the machine supports pre- or post-address side-effect generation involving constants other than the size of the memory operand.

### HAVE\_PRE\_MODIFY\_REG

[Macro]

HAVE\_POST\_MODIFY\_REG

[Macro]

A C expression that is nonzero if the machine supports pre- or post-address side-effect generation involving a register displacement.

### CONSTANT\_ADDRESS\_P (x)

[Macro]

A C expression that is 1 if the RTX x is a constant which is a valid address. On most machines the default definition of (CONSTANT\_P (x) && GET\_CODE (x) != CONST\_DOUBLE) is acceptable, but a few machines are more restrictive as to which constant addresses are supported.

 $CONSTANT_P(x)$  [Macro]

CONSTANT\_P, which is defined by target-independent code, accepts integer-values expressions whose values are not explicitly known, such as symbol\_ref, label\_ref, and high expressions and const arithmetic expressions, in addition to const\_int and const\_double expressions.

### MAX\_REGS\_PER\_ADDRESS

[Macro]

A number, the maximum number of registers that can appear in a valid memory address. Note that it is up to you to specify a value equal to the maximum number that TARGET\_LEGITIMATE\_ADDRESS\_P would ever accept.

### bool TARGET\_LEGITIMATE\_ADDRESS\_P (machine\_mode mode, rtx x, bool strict) [Target Hook]

A function that returns whether x (an RTX) is a legitimate memory address on the target machine for a memory operand of mode mode.

Legitimate addresses are defined in two variants: a strict variant and a non-strict one. The *strict* parameter chooses which variant is desired by the caller.

The strict variant is used in the reload pass. It must be defined so that any pseudo-register that has not been allocated a hard register is considered a memory reference. This is because in contexts where some kind of register is required, a pseudo-register with no hard register must be rejected. For non-hard registers, the strict variant should look up the reg\_renumber array; it should then proceed using the hard register number in the array, or treat the pseudo as a memory reference if the array holds -1.

The non-strict variant is used in other passes. It must be defined to accept all pseudoregisters in every context where some kind of register is required.

Normally, constant addresses which are the sum of a symbol\_ref and an integer are stored inside a const RTX to mark them as constant. Therefore, there is no need to recognize such sums specifically as legitimate addresses. Normally you would simply recognize any const as legitimate.

Usually PRINT\_OPERAND\_ADDRESS is not prepared to handle constant sums that are not marked with const. It assumes that a naked plus indicates indexing. If so, then you *must* reject such naked constant sums as illegitimate addresses, so that none of them will be given to PRINT\_OPERAND\_ADDRESS.

On some machines, whether a symbolic address is legitimate depends on the section that the address refers to. On these machines, define the target hook TARGET\_ENCODE\_SECTION\_INFO to store the information into the symbol\_ref, and then check for it here. When you see a const, you will have to look inside it to find the symbol\_ref in order to determine the section. See Section 18.20 [Assembler Format], page 596.

Some ports are still using a deprecated legacy substitute for this hook, the GO\_IF\_LEGITIMATE\_ADDRESS macro. This macro has this syntax:

```
#define GO_IF_LEGITIMATE_ADDRESS (mode, x, label)
```

and should goto label if the address x is a valid address on the target machine for a memory operand of mode mode.

Compiler source files that want to use the strict variant of this macro define the macro REG\_OK\_STRICT. You should use an #ifdef REG\_OK\_STRICT conditional to define the strict variant in that case and the non-strict variant otherwise.

Using the hook is usually simpler because it limits the number of files that are recompiled when changes are made.

### TARGET\_MEM\_CONSTRAINT

[Macro]

A single character to be used instead of the default 'm' character for general memory addresses. This defines the constraint letter which matches the memory addresses accepted by TARGET\_LEGITIMATE\_ADDRESS\_P. Define this macro if you want to support new address formats in your back end without changing the semantics of the 'm' constraint. This is necessary in order to preserve functionality of inline assembly constructs using the 'm' constraint.

### $FIND_BASE_TERM(x)$

|Macro|

A C expression to determine the base term of address x, or to provide a simplified version of x from which 'alias.c' can easily find the base term. This macro is used in only two places: find\_base\_value and find\_base\_term in 'alias.c'.

It is always safe for this macro to not be defined. It exists so that alias analysis can understand machine-dependent addresses.

The typical use of this macro is to handle addresses containing a label\_ref or symbol\_ref within an UNSPEC.

### rtx TARGET\_LEGITIMIZE\_ADDRESS (rtx x, rtx oldx, machine\_mode [Target Hook] mode)

This hook is given an invalid memory address x for an operand of mode mode and should try to return a valid memory address.

x will always be the result of a call to break\_out\_memory\_refs, and oldx will be the operand that was given to that function to produce x.

The code of the hook should not alter the substructure of x. If it transforms x into a more legitimate form, it should return the new x.

It is not necessary for this hook to come up with a legitimate address, with the exception of native TLS addresses (see Section 18.25 [Emulated TLS], page 636). The compiler has standard ways of doing so in all cases. In fact, if the target supports only emulated TLS, it is safe to omit this hook or make it return x if it cannot find a valid way to legitimize the address. But often a machine-dependent strategy can generate better code.

### LEGITIMIZE\_RELOAD\_ADDRESS (x, mode, opnum, type, ind\_levels, win) [Macro]

A C compound statement that attempts to replace x, which is an address that needs reloading, with a valid memory address for an operand of mode *mode*. win will be a C statement label elsewhere in the code. It is not necessary to define this macro, but it might be useful for performance reasons.

For example, on the i386, it is sometimes possible to use a single reload register instead of two by reloading a sum of two pseudo registers into a register. On the other hand, for number of RISC processors offsets are limited so that often an intermediate address needs to be generated in order to address a stack slot. By defining LEGITIMIZE\_RELOAD\_ADDRESS appropriately, the intermediate addresses generated for adjacent some stack slots can be made identical, and thus be shared.

*Note*: This macro should be used with caution. It is necessary to know something of how reload works in order to effectively use this, and it is quite easy to produce macros that build in too much knowledge of reload internals.

*Note*: This macro must be able to reload an address created by a previous invocation of this macro. If it fails to handle such addresses then the compiler may generate incorrect code or abort.

The macro definition should use push\_reload to indicate parts that need reloading; opnum, type and ind\_levels are usually suitable to be passed unaltered to push\_reload.

The code generated by this macro must not alter the substructure of x. If it transforms x into a more legitimate form, it should assign x (which will always be a C variable) a new value. This also applies to parts that you change indirectly by calling push\_reload.

The macro definition may use strict\_memory\_address\_p to test if the address has become legitimate.

If you want to change only a part of x, one standard way of doing this is to use copy\_rtx. Note, however, that it unshares only a single level of rtl. Thus, if the part to be changed is not at the top level, you'll need to replace first the top level. It is not necessary for this macro to come up with a legitimate address; but often a machine-dependent strategy can generate better code.

#### 

This hook returns **true** if memory address addr in address space addrspace can have different meanings depending on the machine mode of the memory reference it is used for or if the address is valid for some modes but not others.

Autoincrement and autodecrement addresses typically have mode-dependent effects because the amount of the increment or decrement is the size of the operand being addressed. Some machines have other mode-dependent addresses. Many RISC machines have no mode-dependent addresses.

You may assume that addr is a valid address for the machine.

The default version of this hook returns false.

### 

This hook returns true if x is a legitimate constant for a mode-mode immediate operand on the target machine. You can assume that x satisfies CONSTANT\_P, so you need not check this.

The default definition returns true.

### rtx TARGET\_DELEGITIMIZE\_ADDRESS (rtx x)

[Target Hook]

This hook is used to undo the possibly obfuscating effects of the LEGITIMIZE\_ADDRESS and LEGITIMIZE\_RELOAD\_ADDRESS target macros. Some backend implementations of these macros wrap symbol references inside an UNSPEC rtx to represent PIC or similar addressing modes. This target hook allows GCC's optimizers to understand the semantics of these opaque UNSPECs by converting them back into their original form.

### bool TARGET\_CONST\_NOT\_OK\_FOR\_DEBUG\_P (rtx x) [Target Hook] This hook should return true if x should not be emitted into debug sections.

### bool TARGET\_CANNOT\_FORCE\_CONST\_MEM (machine\_mode mode, rtx [Target Hook] x)

This hook should return true if x is of a form that cannot (or should not) be spilled to the constant pool. mode is the mode of x.

The default version of this hook returns false.

The primary reason to define this hook is to prevent reload from deciding that a non-legitimate constant would be better reloaded from the constant pool instead of spilling and reloading a register holding the constant. This restriction is often true of addresses of TLS symbols for various targets.

### bool TARGET\_USE\_BLOCKS\_FOR\_CONSTANT\_P (machine\_mode mode, const\_rtx x) [Target Hook]

This hook should return true if pool entries for constant x can be placed in an object\_block structure. *mode* is the mode of x.

The default version returns false for all constants.

### bool TARGET\_USE\_BLOCKS\_FOR\_DECL\_P (const\_tree decl)

[Target Hook]

This hook should return true if pool entries for *decl* should be placed in an object\_block structure.

The default version returns true for all decls.

### tree TARGET\_BUILTIN\_RECIPROCAL (tree fndec1)

[Target Hook]

This hook should return the DECL of a function that implements the reciprocal of the machine-specific builtin function *fndecl*, or NULL\_TREE if such a function is not available.

### tree TARGET\_VECTORIZE\_BUILTIN\_MASK\_FOR\_LOAD (void)

[Target Hook]

This hook should return the DECL of a function f that given an address addr as an argument returns a mask m that can be used to extract from two vectors the relevant data that resides in addr in case addr is not properly aligned.

The autovectorizer, when vectorizing a load operation from an address addr that may be unaligned, will generate two vector loads from the two aligned addresses around addr. It then generates a REALIGN\_LOAD operation to extract the relevant data from the two loaded vectors. The first two arguments to REALIGN\_LOAD, v1 and v2, are the two vectors, each of size VS, and the third argument, OFF, defines how the data will be extracted from these two vectors: if OFF is 0, then the returned vector is v2; otherwise, the returned vector is composed from the last VS-OFF elements of v1 concatenated to the first OFF elements of v2.

If this hook is defined, the autovectorizer will generate a call to f (using the DECL tree that this hook returns) and will use the return value of f as the argument OFF to REALIGN\_LOAD. Therefore, the mask m returned by f should comply with the semantics expected by REALIGN\_LOAD described above. If this hook is not defined, then addr will be used as the argument OFF to REALIGN\_LOAD, in which case the low  $\log 2(VS) - 1$  bits of addr will be considered.

#### 

Returns cost of different scalar or vector statements for vectorization cost model. For vector memory operations the cost may depend on type (vectype) and misalignment value (misalign).

### poly\_uint64 TARGET\_VECTORIZE\_PREFERRED\_VECTOR\_ALIGNMENT [Target Hook] (const\_tree type)

This hook returns the preferred alignment in bits for accesses to vectors of type type in vectorized code. This might be less than or greater than the ABI-defined value returned by TARGET\_VECTOR\_ALIGNMENT. It can be equal to the alignment of a single element, in which case the vectorizer will not try to optimize for alignment.

The default hook returns TYPE\_ALIGN (type), which is correct for most targets.

### bool TARGET\_VECTORIZE\_VECTOR\_ALIGNMENT\_REACHABLE

[Target Hook]

(const\_tree type, bool is\_packed)

Return true if vector alignment is reachable (by peeling N iterations) for the given scalar type type. is\_packed is false if the scalar access using type is known to be naturally aligned.

### bool TARGET\_VECTORIZE\_VEC\_PERM\_CONST (machine\_mode mode,

[Target Hook]

rtx output, rtx in0, rtx in1, const vec\_perm\_indices &sel)

This hook is used to test whether the target can permute up to two vectors of mode mode using the permutation vector sel, and also to emit such a permutation. In the former case in0, in1 and out are all null. In the latter case in0 and in1 are the source vectors and out is the destination vector; all three are registers of mode mode. in1 is the same as in0 if sel describes a permutation on one vector instead of two.

Return true if the operation is possible, emitting instructions for it if rtxes are provided.

If the hook returns false for a mode with multibyte elements, GCC will try the equivalent byte operation. If that also fails, it will try forcing the selector into a register and using the *vec\_permmode* instruction pattern. There is no need for the hook to handle these two implementation approaches itself.

### tree TARGET\_VECTORIZE\_BUILTIN\_VECTORIZED\_FUNCTION

[Target Hook]

(unsigned code, tree vec\_type\_out, tree vec\_type\_in)

This hook should return the decl of a function that implements the vectorized variant of the function with the combined\_fn code code or NULL\_TREE if such a function is not available. The return type of the vectorized function shall be of vector type  $vec\_type\_out$  and the argument types should be  $vec\_type\_in$ .

### tree TARGET\_VECTORIZE\_BUILTIN\_MD\_VECTORIZED\_FUNCTION

[Target Hook]

(tree fndec1, tree vec\_type\_out, tree vec\_type\_in)

This hook should return the decl of a function that implements the vectorized variant of target built-in function fndecl. The return type of the vectorized function shall be of vector type  $vec_type_out$  and the argument types should be  $vec_type_in$ .

### bool TARGET\_VECTORIZE\_SUPPORT\_VECTOR\_MISALIGNMENT

[Target Hook]

(machine\_mode mode, const\_tree type, int misalignment, bool is\_packed)

This hook should return true if the target supports misaligned vector store/load of a specific factor denoted in the *misalignment* parameter. The vector store/load should be of machine mode mode and the elements in the vectors should be of type type. is\_packed parameter is true if the memory access is defined in a packed struct.

### machine\_mode TARGET\_VECTORIZE\_PREFERRED\_SIMD\_MODE

[Target Hook]

(scalar\_mode mode)

This hook should return the preferred mode for vectorizing scalar mode *mode*. The default is equal to word\_mode, because the vectorizer can do some transformations even in absence of specialized SIMD hardware.

### machine\_mode TARGET\_VECTORIZE\_SPLIT\_REDUCTION

[Target Hook]

(machine\_mode)

This hook should return the preferred mode to split the final reduction step on *mode* to. The reduction is then carried out reducing upper against lower halves of vectors recursively until the specified mode is reached. The default is *mode* which means no splitting.

unsigned int [Target Hook]

TARGET\_VECTORIZE\_AUTOVECTORIZE\_VECTOR\_MODES (vector\_modes \*modes, bool all)

If using the mode returned by TARGET\_VECTORIZE\_PREFERRED\_SIMD\_MODE is not the only approach worth considering, this hook should add one mode to *modes* for each useful alternative approach. These modes are then passed to TARGET\_VECTORIZE\_RELATED\_MODE to obtain the vector mode for a given element mode.

The modes returned in *modes* should use the smallest element mode possible for the vectorization approach that they represent, preferring integer modes over floating-poing modes in the event of a tie. The first mode should be the TARGET\_VECTORIZE\_PREFERRED\_SIMD\_MODE for its element mode.

If all is true, add suitable vector modes even when they are generally not expected to be worthwhile.

The hook returns a bitmask of flags that control how the modes in *modes* are used. The flags are:

### VECT\_COMPARE\_COSTS

Tells the loop vectorizer to try all the provided modes and pick the one with the lowest cost. By default the vectorizer will choose the first mode that works.

The hook does not need to do anything if the vector returned by TARGET\_VECTORIZE\_PREFERRED\_SIMD\_MODE is the only one relevant for autovectorization. The default implementation adds no modes and returns 0.

### opt\_machine\_mode TARGET\_VECTORIZE\_RELATED\_MODE

[Target Hook]

(machine\_mode vector\_mode, scalar\_mode element\_mode, poly\_uint64 nunits)

If a piece of code is using vector mode vector\_mode and also wants to operate on elements of mode element\_mode, return the vector mode it should use for those elements. If nunits is nonzero, ensure that the mode has exactly nunits elements, otherwise pick whichever vector size pairs the most naturally with vector\_mode. Return an empty opt\_machine\_mode if there is no supported vector mode with the required properties.

There is no prescribed way of handling the case in which nunits is zero. One common choice is to pick a vector mode with the same size as vector\_mode; this is the natural choice if the target has a fixed vector size. Another option is to choose a vector mode with the same number of elements as vector\_mode; this is the natural choice if the target has a fixed number of elements. Alternatively, the hook might choose a middle ground, such as trying to keep the number of elements as similar as possible while applying maximum and minimum vector sizes.

The default implementation uses mode\_for\_vector to find the requested mode, returning a mode with the same size as vector\_mode when nunits is zero. This is the correct behavior for most targets.

### opt\_machine\_mode TARGET\_VECTORIZE\_GET\_MASK\_MODE [Target Hook] (machine\_mode mode)

Return the mode to use for a vector mask that holds one boolean result for each element of vector mode *mode*. The returned mask mode can be a vector of integers (class MODE\_VECTOR\_INT), a vector of booleans (class MODE\_VECTOR\_BOOL) or a scalar integer (class MODE\_INT). Return an empty opt\_machine\_mode if no such mask mode exists.

The default implementation returns a MODE\_VECTOR\_INT with the same size and number of elements as *mode*, if such a mode exists.

#### 

This hook returns true if masked internal function *ifn* (really of type internal\_fn) should be considered expensive when the mask is all zeros. GCC can then try to branch around the instruction instead.

void \* TARGET\_VECTORIZE\_INIT\_COST (class loop \*loop\_info) [Target Hook]
This hook should initialize target-specific data structures in preparation for modeling
the costs of vectorizing a loop or basic block. The default allocates three unsigned
integers for accumulating costs for the prologue, body, and epilogue of the loop or
basic block. If loop\_info is non-NULL, it identifies the loop being vectorized; otherwise
a single block is being vectorized.

# unsigned TARGET\_VECTORIZE\_ADD\_STMT\_COST (void \*data, int [Target Hook] count, enum vect\_cost\_for\_stmt kind, class \_stmt\_vec\_info \*stmt\_info, int misalign, enum vect\_cost\_model\_location where)

This hook should update the target-specific data in response to adding count copies of the given kind of statement to a loop or basic block. The default adds the builtin vectorizer cost for the copies of the statement to the accumulator specified by where, (the prologue, body, or epilogue) and returns the amount added. The return value should be viewed as a tentative cost that may later be revised.

#### 

This hook should complete calculations of the cost of vectorizing a loop or basic block based on *data*, and return the prologue, body, and epilogue costs as unsigned integers. The default returns the value of the three accumulators.

# void TARGET\_VECTORIZE\_DESTROY\_COST\_DATA (void \*data) [Target Hook] This hook should release data and any related data structures allocated by TARGET\_VECTORIZE\_INIT\_COST. The default releases the accumulator.

### tree TARGET\_VECTORIZE\_BUILTIN\_GATHER (const\_tree

[Target Hook]

mem\_vectype, const\_tree index\_type, int scale)

Target builtin that implements vector gather operation.  $mem\_vectype$  is the vector type of the load and  $index\_type$  is scalar type of the index, scaled by scale. The default is NULL\_TREE which means to not vectorize gather loads.

### tree TARGET\_VECTORIZE\_BUILTIN\_SCATTER (const\_tree vectype, const\_tree index\_type, int scale) [Target Hook]

Target builtin that implements vector scatter operation. *vectype* is the vector type of the store and *index\_type* is scalar type of the index, scaled by *scale*. The default is NULL\_TREE which means to not vectorize scatter stores.

#### 

This hook should set  $vecsize\_mangle$ ,  $vecsize\_int$ ,  $vecsize\_float$  fields in  $simd\_clone$  structure pointed by  $clone\_info$  argument and also simdlen field if it was previously 0. The hook should return 0 if SIMD clones shouldn't be emitted, or number of  $vecsize\_mangle$  variants that should be emitted.

## void TARGET\_SIMD\_CLONE\_ADJUST (struct cgraph\_node \*) [Target Hook] This hook should add implicit attribute(target("...")) attribute to SIMD clone node if needed.

# int TARGET\_SIMD\_CLONE\_USABLE (struct cgraph\_node \*) [Target Hook] This hook should return -1 if SIMD clone node shouldn't be used in vectorized loops in current function, or non-negative number if it is usable. In that case, the smaller the number is, the more desirable it is to use it.

### int TARGET\_SIMT\_VF (void)

[Target Hook]

Return number of threads in SIMT thread group on the target.

### int TARGET\_OMP\_DEVICE\_KIND\_ARCH\_ISA (enum

[Target Hook]

omp\_device\_kind\_arch\_isa trait, const char \*name)

Return 1 if *trait name* is present in the OpenMP context's device trait set, return 0 if not present in any OpenMP context in the whole translation unit, or -1 if not present in the current OpenMP context but might be present in another OpenMP context in the same TU.

### bool TARGET\_GOACC\_VALIDATE\_DIMS (tree decl, int \*dims, int fn\_level, unsigned used) [Target Hook]

This hook should check the launch dimensions provided for an OpenACC compute region, or routine. Defaulted values are represented as -1 and non-constant values as 0. The *fn\_level* is negative for the function corresponding to the compute region. For a routine it is the outermost level at which partitioned execution may be spawned. The hook should verify non-default values. If DECL is NULL, global defaults are being validated and unspecified defaults should be filled in. Diagnostics should be issued as appropriate. Return true, if changes have been made. You must override this hook to provide dimensions larger than 1.

### int TARGET\_GOACC\_DIM\_LIMIT (int axis)

[Target Hook]

This hook should return the maximum size of a particular dimension, or zero if unbounded.

bool TARGET\_GOACC\_FORK\_JOIN (gcall \*call, const int \*dims, bool [Target Hook] is\_fork)

This hook can be used to convert IFN\_GOACC\_FORK and IFN\_GOACC\_JOIN function calls to target-specific gimple, or indicate whether they should be retained. It is executed during the oacc\_device\_lower pass. It should return true, if the call should be retained. It should return false, if it is to be deleted (either because target-specific gimple has been inserted before it, or there is no need for it). The default hook returns false, if there are no RTL expanders for them.

### void TARGET\_GOACC\_REDUCTION (gcall \*call)

[Target Hook]

This hook is used by the oacc\_transform pass to expand calls to the GOACC\_REDUCTION internal function, into a sequence of gimple instructions. call is gimple statement containing the call to the function. This hook removes statement call after the expanded sequence has been inserted. This hook is also responsible for allocating any storage for reductions when necessary.

tree TARGET\_PREFERRED\_ELSE\_VALUE (unsigned ifn, tree type, unsigned nops, tree \*ops) [Target Hook]

This hook returns the target's preferred final argument for a call to conditional internal function *ifn* (really of type <code>internal\_fn</code>). *type* specifies the return type of the function and *ops* are the operands to the conditional operation, of which there are *nops*.

For example, if *ifn* is IFN\_COND\_ADD, the hook returns a value of type *type* that should be used when 'ops[0]' and 'ops[1]' are conditionally added together.

This hook is only relevant if the target supports conditional patterns like cond\_addm. The default implementation returns a zero constant of type type.

### 18.14 Anchored Addresses

GCC usually addresses every static object as a separate entity. For example, if we have:

```
static int a, b, c;
int foo (void) { return a + b + c; }
```

the code for foo will usually calculate three separate symbolic addresses: those of a, b and c. On some targets, it would be better to calculate just one symbolic address and access the three variables relative to it. The equivalent pseudocode would be something like:

```
int foo (void)
{
  register int *xr = &x;
  return xr[&a - &x] + xr[&b - &x] + xr[&c - &x];
}
```

(which isn't valid C). We refer to shared addresses like x as "section anchors". Their use is controlled by '-fsection-anchors'.

The hooks below describe the target properties that GCC needs to know in order to make effective use of section anchors. It won't use section anchors at all unless either TARGET\_MIN\_ANCHOR\_OFFSET or TARGET\_MAX\_ANCHOR\_OFFSET is set to a nonzero value.

### HOST\_WIDE\_INT TARGET\_MIN\_ANCHOR\_OFFSET

[Target Hook]

The minimum offset that should be applied to a section anchor. On most targets, it should be the smallest offset that can be applied to a base register while still giving a legitimate address for every mode. The default value is 0.

### HOST\_WIDE\_INT TARGET\_MAX\_ANCHOR\_OFFSET

[Target Hook]

Like TARGET\_MIN\_ANCHOR\_OFFSET, but the maximum (inclusive) offset that should be applied to section anchors. The default value is 0.

### void TARGET\_ASM\_OUTPUT\_ANCHOR (rtx x)

[Target Hook]

Write the assembly code to define section anchor x, which is a SYMBOL\_REF for which 'SYMBOL\_REF\_ANCHOR\_P (x)' is true. The hook is called with the assembly output position set to the beginning of SYMBOL\_REF\_BLOCK (x).

If  $ASM_OUTPUT_DEF$  is available, the hook's default definition uses it to define the symbol as '. + SYMBOL\_REF\_BLOCK\_OFFSET (x)'. If  $ASM_OUTPUT_DEF$  is not available, the hook's default definition is NULL, which disables the use of section anchors altogether.

### bool TARGET\_USE\_ANCHORS\_FOR\_SYMBOL\_P (const\_rtx x)

[Target Hook]

Return true if GCC should attempt to use anchors to access SYMBOL\_REF x. You can assume 'SYMBOL\_REF\_HAS\_BLOCK\_INFO\_P (x)' and '!SYMBOL\_REF\_ANCHOR\_P (x)'.

The default version is correct for most targets, but you might need to intercept this hook to handle things like target-specific attributes or target-specific sections.

### 18.15 Condition Code Status

The macros in this section can be split in two families, according to the two ways of representing condition codes in GCC.

The first representation is the so called (cc0) representation (see Section 17.12 [Jump Patterns], page 433), where all instructions can have an implicit clobber of the condition codes. The second is the condition code register representation, which provides better schedulability for architectures that do have a condition code register, but on which most instructions do not affect it. The latter category includes most RISC machines.

The implicit clobbering poses a strong restriction on the placement of the definition and use of the condition code. In the past the definition and use were always adjacent. However, recent changes to support trapping arithmatic may result in the definition and user being in different blocks. Thus, there may be a NOTE\_INSN\_BASIC\_BLOCK between them. Additionally, the definition may be the source of exception handling edges.

These restrictions can prevent important optimizations on some machines. For example, on the IBM RS/6000, there is a delay for taken branches unless the condition code register is set three instructions earlier than the conditional branch. The instruction scheduler cannot perform this optimization if it is not permitted to separate the definition and use of the condition code register.

For this reason, it is possible and suggested to use a register to represent the condition code for new ports. If there is a specific condition code register in the machine, use a hard register. If the condition code or comparison result can be placed in any general register, or if there are multiple condition registers, use a pseudo register. Registers used to store the condition code value will usually have a mode that is in class MODE\_CC.

Alternatively, you can use BImode if the comparison operator is specified already in the compare instruction. In this case, you are not interested in most macros in this section.

### 18.15.1 Representation of condition codes using (cc0)

The file 'conditions.h' defines a variable cc\_status to describe how the condition code was computed (in case the interpretation of the condition code depends on the instruction that it was set by). This variable contains the RTL expressions on which the condition code is currently based, and several standard flags.

Sometimes additional machine-specific flags must be defined in the machine description header file. It can also add additional machine-specific information by defining CC\_STATUS\_MDEP.

CC\_STATUS\_MDEP [Macro]

C code for a data type which is used for declaring the mdep component of cc\_status. It defaults to int.

This macro is not used on machines that do not use cc0.

### CC\_STATUS\_MDEP\_INIT

[Macro]

A C expression to initialize the mdep field to "empty". The default definition does nothing, since most machines don't use the field anyway. If you want to use the field, you should probably define this macro to initialize it.

This macro is not used on machines that do not use cc0.

### NOTICE\_UPDATE\_CC (exp, insn)

[Macro]

A C compound statement to set the components of cc\_status appropriately for an insn *insn* whose body is *exp*. It is this macro's responsibility to recognize insns that set the condition code as a byproduct of other activity as well as those that explicitly set (cc0).

This macro is not used on machines that do not use cc0.

If there are insns that do not set the condition code but do alter other machine registers, this macro must check to see whether they invalidate the expressions that the condition code is recorded as reflecting. For example, on the 68000, insns that store in address registers do not set the condition code, which means that usually NOTICE\_UPDATE\_CC can leave cc\_status unaltered for such insns. But suppose that the previous insn set the condition code based on location 'a40(102)' and the current insn stores a new value in 'a4'. Although the condition code is not changed by this, it will no longer be true that it reflects the contents of 'a40(102)'. Therefore, NOTICE\_UPDATE\_CC must alter cc\_status in this case to say that nothing is known about the condition code value.

The definition of NOTICE\_UPDATE\_CC must be prepared to deal with the results of peephole optimization: insns whose patterns are parallel RTXs containing various reg, mem or constants which are just the operands. The RTL structure of these insns is not sufficient to indicate what the insns actually do. What NOTICE\_UPDATE\_CC should do when it sees one is just to run CC\_STATUS\_INIT.

A possible definition of NOTICE\_UPDATE\_CC is to call a function that looks at an attribute (see Section 17.19 [Insn Attributes], page 450) named, for example, 'cc'.

This avoids having detailed information about patterns in two places, the 'md' file and in NOTICE\_UPDATE\_CC.

### 18.15.2 Representation of condition codes using registers

### SELECT\_CC\_MODE (op, x, y)

[Macro]

On many machines, the condition code may be produced by other instructions than compares, for example the branch can use directly the condition code set by a subtract instruction. However, on some machines when the condition code is set this way some bits (such as the overflow bit) are not set in the same way as a test instruction, so that a different branch instruction must be used for some conditional branches. When this happens, use the machine mode of the condition code register to record different formats of the condition code register. Modes can also be used to record which compare instruction (e.g. a signed or an unsigned comparison) produced the condition codes.

If other modes than CCmode are required, add them to 'machine-modes.def' and define SELECT\_CC\_MODE to choose a mode given an operand of a compare. This is needed because the modes have to be chosen not only during RTL generation but also, for example, by instruction combination. The result of SELECT\_CC\_MODE should be consistent with the mode used in the patterns; for example to support the case of the add on the SPARC discussed above, we have the pattern

together with a SELECT\_CC\_MODE that returns CCNZmode for comparisons whose argument is a plus:

Another reason to use modes is to retain information on which operands were used by the comparison; see REVERSIBLE\_CC\_MODE later in this section.

You should define this macro if and only if you define extra CC modes in 'machine-modes.def'.

```
void TARGET_CANONICALIZE_COMPARISON (int *code, rtx *op0, rtx [Target Hook] *op1, bool op0_preserve_value)
```

On some machines not all possible comparisons are defined, but you can convert an invalid comparison into a valid one. For example, the Alpha does not have a GT comparison, but you can use an LT comparison instead and swap the order of the operands.

On such machines, implement this hook to do any required conversions. code is the initial comparison code and op0 and op1 are the left and right operands of the comparison, respectively. If  $op0\_preserve\_value$  is true the implementation is not allowed to change the value of op0 since the value might be used in RTXs which aren't comparisons. E.g. the implementation is not allowed to swap operands in that case.

GCC will not assume that the comparison resulting from this macro is valid but will see if the resulting insn matches a pattern in the 'md' file.

You need not to implement this hook if it would never change the comparison code or operands.

### REVERSIBLE\_CC\_MODE (mode)

[Macro]

A C expression whose value is one if it is always safe to reverse a comparison whose mode is *mode*. If SELECT\_CC\_MODE can ever return *mode* for a floating-point inequality comparison, then REVERSIBLE\_CC\_MODE (*mode*) must be zero.

You need not define this macro if it would always returns zero or if the floating-point format is anything other than IEEE\_FLOAT\_FORMAT. For example, here is the definition used on the SPARC, where floating-point inequality comparisons are given either CCFPEmode or CCFPmode:

```
#define REVERSIBLE_CC_MODE(MODE) \
  ((MODE) != CCFPEmode && (MODE) != CCFPmode)
```

### REVERSE\_CONDITION (code, mode)

[Macro]

A C expression whose value is reversed condition code of the *code* for comparison done in CC\_MODE *mode*. The macro is used only in case REVERSIBLE\_CC\_MODE (*mode*) is nonzero. Define this macro in case machine has some non-standard way how to reverse certain conditionals. For instance in case all floating point conditions are non-trapping, compiler may freely convert unordered compares to ordered ones. Then definition may look like:

```
#define REVERSE_CONDITION(CODE, MODE) \
  ((MODE) != CCFPmode ? reverse_condition (CODE) \
    : reverse_condition_maybe_unordered (CODE))
```

### bool TARGET\_FIXED\_CONDITION\_CODE\_REGS (unsigned int \*p1, unsigned int \*p2) [Target Hook]

On targets which do not use (cc0), and which use a hard register rather than a pseudo-register to hold condition codes, the regular CSE passes are often not able to identify cases in which the hard register is set to a common value. Use this hook to enable a small pass which optimizes such cases. This hook should return true to enable this pass, and it should set the integers to which its arguments point to the hard register numbers used for condition codes. When there is only one such register, as is true on most systems, the integer pointed to by p2 should be set to INVALID\_REGNUM.

The default version of this hook returns false.

#### 

On targets which use multiple condition code modes in class MODE\_CC, it is sometimes the case that a comparison can be validly done in more than one mode. On such a system, define this target hook to take two mode arguments and to return a mode in which both comparisons may be validly done. If there is no such mode, return VOIDmode.

The default version of this hook checks whether the modes are the same. If they are, it returns that mode. If they are different, it returns VOIDmode.

### unsigned int TARGET\_FLAGS\_REGNUM

[Target Hook]

If the target has a dedicated flags register, and it needs to use the post-reload comparison elimination pass, or the delay slot filler pass, then this value should be set appropriately.

### 18.16 Describing Relative Costs of Operations

These macros let you describe the relative speed of various operations on the target machine.

### REGISTER\_MOVE\_COST (mode, from, to)

[Macro]

A C expression for the cost of moving data of mode mode from a register in class from to one in class to. The classes are expressed using the enumeration values such as GENERAL\_REGS. A value of 2 is the default; other values are interpreted relative to that.

It is not required that the cost always equal 2 when from is the same as to; on some machines it is expensive to move between registers if they are not general registers.

If reload sees an insn consisting of a single set between two hard registers, and if REGISTER\_MOVE\_COST applied to their classes returns a value of 2, reload does not check to ensure that the constraints of the insn are met. Setting a cost of other than 2 will allow reload to verify that the constraints are met. You should do this if the 'movm' pattern's constraints do not allow such copying.

These macros are obsolete, new ports should use the target hook TARGET\_REGISTER\_MOVE\_COST instead.

#### 

This target hook should return the cost of moving data of mode *mode* from a register in class *from* to one in class *to*. The classes are expressed using the enumeration values such as GENERAL\_REGS. A value of 2 is the default; other values are interpreted relative to that.

It is not required that the cost always equal 2 when from is the same as to; on some machines it is expensive to move between registers if they are not general registers.

If reload sees an insn consisting of a single set between two hard registers, and if TARGET\_REGISTER\_MOVE\_COST applied to their classes returns a value of 2, reload does not check to ensure that the constraints of the insn are met. Setting a cost of other than 2 will allow reload to verify that the constraints are met. You should do this if the 'movm' pattern's constraints do not allow such copying.

The default version of this function returns 2.

### MEMORY\_MOVE\_COST (mode, class, in)

[Macro]

A C expression for the cost of moving data of mode mode between a register of class class and memory; in is zero if the value is to be written to memory, nonzero if it is to

be read in. This cost is relative to those in REGISTER\_MOVE\_COST. If moving between registers and memory is more expensive than between two registers, you should define this macro to express the relative cost.

If you do not define this macro, GCC uses a default cost of 4 plus the cost of copying via a secondary reload register, if one is needed. If your machine requires a secondary reload register to copy between memory and a register of *class* but the reload mechanism is more complex than copying via an intermediate, define this macro to reflect the actual cost of the move.

GCC defines the function memory\_move\_secondary\_cost if secondary reloads are needed. It computes the costs due to copying via a secondary register. If your machine copies from memory using a secondary register in the conventional way but the default base value of 4 is not correct for your machine, define this macro to add some other value to the result of that function. The arguments to that function are the same as to this macro.

These macros are obsolete, new ports should use the target hook TARGET\_MEMORY\_MOVE\_COST instead.

#### 

This target hook should return the cost of moving data of mode mode between a register of class rclass and memory; in is false if the value is to be written to memory, true if it is to be read in. This cost is relative to those in TARGET\_REGISTER\_MOVE\_COST. If moving between registers and memory is more expensive than between two registers, you should add this target hook to express the relative cost.

If you do not add this target hook, GCC uses a default cost of 4 plus the cost of copying via a secondary reload register, if one is needed. If your machine requires a secondary reload register to copy between memory and a register of relass but the reload mechanism is more complex than copying via an intermediate, use this target hook to reflect the actual cost of the move.

GCC defines the function memory\_move\_secondary\_cost if secondary reloads are needed. It computes the costs due to copying via a secondary register. If your machine copies from memory using a secondary register in the conventional way but the default base value of 4 is not correct for your machine, use this target hook to add some other value to the result of that function. The arguments to that function are the same as to this target hook.

### BRANCH\_COST (speed\_p, predictable\_p)

[Macro]

A C expression for the cost of a branch instruction. A value of 1 is the default; other values are interpreted relative to that. Parameter  $speed_p$  is true when the branch in question should be optimized for speed. When it is false, BRANCH\_COST should return a value optimal for code size rather than performance.  $predictable_p$  is true for well-predicted branches. On many architectures the BRANCH\_COST can be reduced then.

Here are additional macros which do not specify precise relative costs, but only that certain actions are more expensive than GCC would ordinarily expect.

### SLOW\_BYTE\_ACCESS [Macro]

Define this macro as a C expression which is nonzero if accessing less than a word of memory (i.e. a char or a short) is no faster than accessing a word of memory, i.e., if such access require more than one instruction or if there is no difference in cost between byte and (aligned) word loads.

When this macro is not defined, the compiler will access a field by finding the smallest containing object; when it is defined, a fullword load will be used if alignment permits. Unless bytes accesses are faster than word accesses, using word accesses is preferable since it may eliminate subsequent memory access if subsequent accesses occur to other fields in the same word of the structure, but to different bytes.

### bool TARGET\_SLOW\_UNALIGNED\_ACCESS (machine\_mode mode, unsigned int align) [Target Hook]

This hook returns true if memory accesses described by the *mode* and *alignment* parameters have a cost many times greater than aligned accesses, for example if they are emulated in a trap handler. This hook is invoked only for unaligned accesses, i.e. when *alignment* < GET\_MODE\_ALIGNMENT (*mode*).

When this hook returns true, the compiler will act as if STRICT\_ALIGNMENT were true when generating code for block moves. This can cause significantly more instructions to be produced. Therefore, do not make this hook return true if unaligned accesses only add a cycle or two to the time for a memory access.

The hook must return true whenever STRICT\_ALIGNMENT is true. The default implementation returns STRICT\_ALIGNMENT.

### MOVE\_RATIO (speed) [Macro]

The threshold of number of scalar memory-to-memory move insns, below which a sequence of insns should be generated instead of a string move insn or a library call. Increasing the value will always make code faster, but eventually incurs high cost in increased code size.

Note that on machines where the corresponding move insn is a define\_expand that emits a sequence of insns, this macro counts the number of such sequences.

The parameter *speed* is true if the code is currently being optimized for speed rather than size.

If you don't define this, a reasonable default is used.

#### 

GCC will attempt several strategies when asked to copy between two areas of memory, or to set, clear or store to memory, for example when copying a struct. The by\_pieces infrastructure implements such memory operations as a sequence of load, store or move insns. Alternate strategies are to expand the cpymem or setmem optabs, to emit a library call, or to emit unit-by-unit, loop-based operations.

This target hook should return true if, for a memory operation with a given size and alignment, using the by\_pieces infrastructure is expected to result in better code generation. Both size and alignment are measured in terms of storage units.

The parameter op is one of: CLEAR\_BY\_PIECES, MOVE\_BY\_PIECES, SET\_BY\_PIECES, STORE\_BY\_PIECES or COMPARE\_BY\_PIECES. These describe the type of memory operation under consideration.

The parameter  $speed_p$  is true if the code is currently being optimized for speed rather than size.

Returning true for higher values of size can improve code generation for speed if the target does not provide an implementation of the cpymem or setmem standard names, if the cpymem or setmem implementation would be more expensive than a sequence of insns, or if the overhead of a library call would dominate that of the body of the memory operation.

Returning true for higher values of size may also cause an increase in code size, for example where the number of insns emitted to perform a move would be greater than that of a library call.

### 

When expanding a block comparison in MODE, gcc can try to reduce the number of branches at the expense of more memory operations. This hook allows the target to override the default choice. It should return the factor by which branches should be reduced over the plain expansion with one comparison per *mode*-sized piece. A port can also prevent a particular mode from being used for block comparisons by returning a negative number from this hook.

MOVE\_MAX\_PIECES [Macro]

A C expression used by move\_by\_pieces to determine the largest unit a load or store used to copy memory is. Defaults to MOVE\_MAX.

STORE\_MAX\_PIECES [Macro]

A C expression used by store\_by\_pieces to determine the largest unit a store used to memory is. Defaults to MOVE\_MAX\_PIECES, or two times the size of HOST\_WIDE\_INT, whichever is smaller.

### COMPARE\_MAX\_PIECES [Macro]

A C expression used by compare\_by\_pieces to determine the largest unit a load or store used to compare memory is. Defaults to MOVE\_MAX\_PIECES.

### CLEAR\_RATIO (speed)

[Macro]

The threshold of number of scalar move insns, below which a sequence of insns should be generated to clear memory instead of a string clear insn or a library call. Increasing the value will always make code faster, but eventually incurs high cost in increased code size.

The parameter *speed* is true if the code is currently being optimized for speed rather than size.

If you don't define this, a reasonable default is used.

### SET\_RATIO (speed) [Macro]

The threshold of number of scalar move insns, below which a sequence of insns should be generated to set memory to a constant value, instead of a block set insn or a library

call. Increasing the value will always make code faster, but eventually incurs high cost in increased code size.

The parameter *speed* is true if the code is currently being optimized for speed rather than size.

If you don't define this, it defaults to the value of MOVE\_RATIO.

### USE\_LOAD\_POST\_INCREMENT (mode)

[Macro]

A C expression used to determine whether a load postincrement is a good thing to use for a given mode. Defaults to the value of HAVE\_POST\_INCREMENT.

### USE\_LOAD\_POST\_DECREMENT (mode)

[Macro]

A C expression used to determine whether a load postdecrement is a good thing to use for a given mode. Defaults to the value of HAVE\_POST\_DECREMENT.

### USE\_LOAD\_PRE\_INCREMENT (mode)

[Macro]

A C expression used to determine whether a load preincrement is a good thing to use for a given mode. Defaults to the value of HAVE\_PRE\_INCREMENT.

### USE\_LOAD\_PRE\_DECREMENT (mode)

[Macro]

A C expression used to determine whether a load predecrement is a good thing to use for a given mode. Defaults to the value of HAVE\_PRE\_DECREMENT.

### USE\_STORE\_POST\_INCREMENT (mode)

[Macro]

A C expression used to determine whether a store postincrement is a good thing to use for a given mode. Defaults to the value of HAVE\_POST\_INCREMENT.

### USE\_STORE\_POST\_DECREMENT (mode)

[Macro]

A C expression used to determine whether a store postdecrement is a good thing to use for a given mode. Defaults to the value of HAVE\_POST\_DECREMENT.

### USE\_STORE\_PRE\_INCREMENT (mode)

[Macro]

This macro is used to determine whether a store preincrement is a good thing to use for a given mode. Defaults to the value of HAVE\_PRE\_INCREMENT.

### USE\_STORE\_PRE\_DECREMENT (mode)

[Macro]

This macro is used to determine whether a store predecrement is a good thing to use for a given mode. Defaults to the value of HAVE\_PRE\_DECREMENT.

### NO\_FUNCTION\_CSE

[Macro]

Define this macro to be true if it is as good or better to call a constant function address than to call an address kept in a register.

### LOGICAL\_OP\_NON\_SHORT\_CIRCUIT

Macro

Define this macro if a non-short-circuit operation produced by 'fold\_range\_test ()' is optimal. This macro defaults to true if BRANCH\_COST is greater than or equal to the value 2.

### bool TARGET\_OPTAB\_SUPPORTED\_P (int op, machine\_mode mode1, machine\_mode mode2, optimization\_type opt\_type) [Target Hook]

Return true if the optimizers should use optab op with modes mode1 and mode2 for optimization type  $opt\_type$ . The optab is known to have an associated '.md'

instruction whose C condition is true. mode2 is only meaningful for conversion optabs; for direct optabs it is a copy of mode1.

For example, when called with *op* equal to rint\_optab and *mode1* equal to DFmode, the hook should say whether the optimizers should use optab rintdf2.

The default hook returns true for all inputs.

### bool TARGET\_RTX\_COSTS (rtx x, machine\_mode mode, int outer\_code, int opno, int \*total, bool speed)

[Target Hook]

This target hook describes the relative costs of RTL expressions.

The cost may depend on the precise form of the expression, which is available for examination in x, and the fact that x appears as operand opno of an expression with rtx code  $outer\_code$ . That is, the hook can assume that there is some rtx y such that 'GET\_CODE  $(y) = outer\_code$ ' and such that either (a) 'XEXP (y, opno) = x' or (b) 'XVEC (y, opno)' contains x.

mode is x's machine mode, or for cases like const\_int that do not have a mode, the mode in which x is used.

In implementing this hook, you can use the construct  $COSTS_N_INSNS$  (n) to specify a cost equal to n fast instructions.

On entry to the hook, \*total contains a default estimate for the cost of the expression. The hook should modify this value as necessary. Traditionally, the default costs are COSTS\_N\_INSNS (5) for multiplications, COSTS\_N\_INSNS (7) for division and modulus operations, and COSTS\_N\_INSNS (1) for all other operations.

When optimizing for code size, i.e. when **speed** is false, this target hook should be used to estimate the relative size cost of an expression, again relative to COSTS\_N\_INSNS.

The hook returns true when all subexpressions of x have been processed, and false when  $rtx\_cost$  should recurse.

#### 

This hook computes the cost of an addressing mode that contains address. If not defined, the cost is computed from the address expression and the TARGET\_RTX\_COST hook.

For most CISC machines, the default cost is a good approximation of the true cost of the addressing mode. However, on RISC machines, all instructions normally have the same length and execution time. Hence all addresses will have equal costs.

In cases where more than one form of an address is known, the form with the lowest cost will be used. If multiple forms have the same, lowest, cost, the one that is the most complex will be used.

For example, suppose an address that is equal to the sum of a register and a constant is used twice in the same basic block. When this macro is not defined, the address will be computed in a register and memory references will be indirect through that register. On machines where the cost of the addressing mode containing the sum is no higher than that of a simple indirect reference, this will produce an additional instruction and possibly require an additional register. Proper specification of this macro eliminates this overhead for such machines.

This hook is never called with an invalid address.

On machines where an address involving more than one register is as cheap as an address computation involving only one register, defining TARGET\_ADDRESS\_COST to reflect this can cause two registers to be live over a region of code where only one would have been if TARGET\_ADDRESS\_COST were not defined in that manner. This effect should be considered in the definition of this macro. Equivalent costs should probably only be given to addresses with different numbers of registers on machines with lots of registers.

### int TARGET\_INSN\_COST (rtx\_insn \*insn, bool speed)

[Target Hook]

This target hook describes the relative costs of RTL instructions.

In implementing this hook, you can use the construct  $COSTS_N_INSNS$  (n) to specify a cost equal to n fast instructions.

When optimizing for code size, i.e. when **speed** is false, this target hook should be used to estimate the relative size cost of an expression, again relative to **COSTS\_N\_INSNS**.

### unsigned int TARGET\_MAX\_NOCE\_IFCVT\_SEQ\_COST (edge e) [Target Hook]

This hook returns a value in the same units as TARGET\_RTX\_COSTS, giving the maximum acceptable cost for a sequence generated by the RTL if-conversion pass when conditional execution is not available. The RTL if-conversion pass attempts to convert conditional operations that would require a branch to a series of unconditional operations and movmodecc insns. This hook returns the maximum cost of the unconditional instructions and the movmodecc insns. RTL if-conversion is cancelled if the cost of the converted sequence is greater than the value returned by this hook.

e is the edge between the basic block containing the conditional branch to the basic block which would be executed if the condition were true.

The default implementation of this hook uses the max-rtl-if-conversion-[un]predictable parameters if they are set, and uses a multiple of BRANCH\_COST otherwise.

### bool TARGET\_NOCE\_CONVERSION\_PROFITABLE\_P (rtx\_insn \*seq, struct noce\_if\_info \*if\_info) [Target Hook]

This hook returns true if the instruction sequence **seq** is a good candidate as a replacement for the if-convertible sequence described in **if\_info**.

# This predicate controls the use of the eager delay slot filler to disallow speculatively executed instructions being placed in delay slots. Targets such as certain MIPS architectures possess both branches with and without delay slots. As the eager delay slot filler can decrease performance, disabling it is beneficial when ordinary branches are available. Use of delay slot branches filled using the basic filler is often still desirable

as the delay slot can hide a pipeline bubble.

HOST\_WIDE\_INT TARGET\_ESTIMATED\_POLY\_VALUE (poly\_int64 val)

 $\texttt{E} (poly\_int64 \text{ val}) \qquad [\texttt{Target Hook}]$ 

Return an estimate of the runtime value of val, for use in things like cost calculations or profiling frequencies. The default implementation returns the lowest possible value of val.

### 18.17 Adjusting the Instruction Scheduler

The instruction scheduler may need a fair amount of machine-specific adjustment in order to produce good code. GCC provides several target hooks for this purpose. It is usually enough to define just a few of them: try the first ones in this list first.

### int TARGET\_SCHED\_ISSUE\_RATE (void)

[Target Hook]

This hook returns the maximum number of instructions that can ever issue at the same time on the target machine. The default is one. Although the insn scheduler can define itself the possibility of issue an insn on the same cycle, the value can serve as an additional constraint to issue insns on the same simulated processor cycle (see hooks 'TARGET\_SCHED\_REORDER' and 'TARGET\_SCHED\_REORDER2'). This value must be constant over the entire compilation. If you need it to vary depending on what the instructions are, you must use 'TARGET\_SCHED\_VARIABLE\_ISSUE'.

#### 

This hook is executed by the scheduler after it has scheduled an insn from the ready list. It should return the number of insns which can still be issued in the current cycle. The default is 'more - 1' for insns other than CLOBBER and USE, which normally are not counted against the issue rate. You should define this hook if some insns take more machine resources than others, so that fewer insns can follow them in the same cycle. file is either a null pointer, or a stdio stream to write any debug output to. verbose is the verbose level provided by '-fsched-verbose-n'. insn is the instruction that was scheduled.

### int TARGET\_SCHED\_ADJUST\_COST (rtx\_insn \*insn, int dep\_type1, rtx\_insn \*dep\_insn, int cost, unsigned int dw) [Target Hook]

This function corrects the value of cost based on the relationship between insn and dep\_insn through a dependence of type dep\_type, and strength dw. It should return the new value. The default is to make no adjustment to cost. This can be used for example to specify to the scheduler using the traditional pipeline description that an output- or anti-dependence does not incur the same cost as a data-dependence. If the scheduler using the automaton based pipeline description, the cost of anti-dependence is zero and the cost of output-dependence is maximum of one and the difference of latency times of the first and the second insns. If these values are not acceptable, you could use the hook to modify them too. See also see Section 17.19.9 [Processor pipeline description], page 460.

#### 

This hook adjusts the integer scheduling priority priority of insn. It should return the new priority. Increase the priority to execute insn earlier, reduce the priority to execute insn later. Do not define this hook if you do not need to adjust the scheduling priorities of insns.

#### 

This hook is executed by the scheduler after it has scheduled the ready list, to allow the machine description to reorder it (for example to combine two small instructions together on 'VLIW' machines). file is either a null pointer, or a stdio stream to write any debug output to. verbose is the verbose level provided by '-fsched-verbose-n'. ready is a pointer to the ready list of instructions that are ready to be scheduled.  $n_{-}$ readyp is a pointer to the number of elements in the ready list. The scheduler reads the ready list in reverse order, starting with  $ready[*n_{-}$ readyp - 1] and going to ready[0]. clock is the timer tick of the scheduler. You may modify the ready list and the number of ready insns. The return value is the number of insns that can issue this cycle; normally this is just issue\_rate. See also 'TARGET\_SCHED\_REORDER2'.

#### 

Like 'TARGET\_SCHED\_REORDER', but called at a different time. That function is called whenever the scheduler starts a new cycle. This one is called once per iteration over a cycle, immediately after 'TARGET\_SCHED\_VARIABLE\_ISSUE'; it can reorder the ready list and return the number of insns to be scheduled in the same cycle. Defining this hook can be useful if there are frequent situations where scheduling one insn causes other insns to become ready in the same cycle. These other insns can then be taken into account properly.

### bool TARGET\_SCHED\_MACRO\_FUSION\_P (void) [Target Hook] This hook is used to check whether target platform supports macro fusion.

### bool TARGET\_SCHED\_MACRO\_FUSION\_PAIR\_P (rtx\_insn \*prev, rtx\_insn \*curr) [Target Hook]

This hook is used to check whether two insns should be macro fused for a target microarchitecture. If this hook returns true for the given insn pair (prev and curr), the scheduler will put them into a sched group, and they will not be scheduled apart. The two insns will be either two SET insns or a compare and a conditional jump and this hook should validate any dependencies needed to fuse the two insns together.

#### 

This hook is called after evaluation forward dependencies of insns in chain given by two parameter values (*head* and *tail* correspondingly) but before insns scheduling of the insn chain. For example, it can be used for better insn classification if it requires analysis of dependencies. This hook can use backward and forward dependencies of the insn scheduler because they are already calculated.

#### 

This hook is executed by the scheduler at the beginning of each block of instructions that are to be scheduled. *file* is either a null pointer, or a stdio stream to write any debug output to. *verbose* is the verbose level provided by '-fsched-verbose-n'. *max\_ready* is the maximum number of insns in the current scheduling region that can be live at the same time. This can be used to allocate scratch space if it is needed, e.g. by 'TARGET\_SCHED\_REORDER'.

# void TARGET\_SCHED\_FINISH (FILE \*file, int verbose) [Target Hook] This hook is executed by the scheduler at the end of each block of instructions that are to be scheduled. It can be used to perform cleanup of any actions done by the

other scheduling hooks. *file* is either a null pointer, or a stdio stream to write any debug output to. *verbose* is the verbose level provided by '-fsched-verbose-n'.

#### 

This hook is executed by the scheduler after function level initializations. *file* is either a null pointer, or a stdio stream to write any debug output to. *verbose* is the verbose level provided by '-fsched-verbose-n'. *old\_max\_uid* is the maximum insn uid when scheduling begins.

- void TARGET\_SCHED\_FINISH\_GLOBAL (FILE \*file, int verbose) [Target Hook]
  This is the cleanup hook corresponding to TARGET\_SCHED\_INIT\_GLOBAL. file is either a null pointer, or a stdio stream to write any debug output to. verbose is the verbose level provided by '-fsched-verbose-n'.
- TARGET\_SCHED\_DFA\_PRE\_CYCLE\_INSN (void) [Target Hook]

  The hook returns an RTL insn. The automaton state used in the pipeline hazard recognizer is changed as if the insn were scheduled when the new simulated processor cycle starts. Usage of the hook may simplify the automaton pipeline description for some VLIW processors. If the hook is defined, it is used only for the automaton based pipeline description. The default is not to change the state when the new simulated processor cycle starts.
- void TARGET\_SCHED\_INIT\_DFA\_PRE\_CYCLE\_INSN (void) [Target Hook]

  The hook can be used to initialize data used by the previous hook.
- rtx\_insn \* TARGET\_SCHED\_DFA\_POST\_CYCLE\_INSN (void) [Target Hook]
  The hook is analogous to 'TARGET\_SCHED\_DFA\_PRE\_CYCLE\_INSN' but used to changed
  the state as if the insn were scheduled when the new simulated processor cycle finishes.
- void TARGET\_SCHED\_INIT\_DFA\_POST\_CYCLE\_INSN (void) [Target Hook]
  The hook is analogous to 'TARGET\_SCHED\_INIT\_DFA\_PRE\_CYCLE\_INSN' but used to initialize data used by the previous hook.
- void TARGET\_SCHED\_DFA\_PRE\_ADVANCE\_CYCLE (void) [Target Hook]
  The hook to notify target that the current simulated cycle is about to finish. The hook is analogous to 'TARGET\_SCHED\_DFA\_PRE\_CYCLE\_INSN' but used to change the state in more complicated situations e.g., when advancing state on a single insn is not enough.
- void TARGET\_SCHED\_DFA\_POST\_ADVANCE\_CYCLE (void) [Target Hook]

  The hook to notify target that new simulated cycle has just started. The hook is analogous to 'TARGET\_SCHED\_DFA\_POST\_CYCLE\_INSN' but used to change the state in more complicated situations e.g., when advancing state on a single insn is not enough.
- int TARGET\_SCHED\_FIRST\_CYCLE\_MULTIPASS\_DFA\_LOOKAHEAD [Target Hook] (void)

This hook controls better choosing an insn from the ready insn queue for the DFAbased insn scheduler. Usually the scheduler chooses the first insn from the queue. If the hook returns a positive value, an additional scheduler code tries all permutations of 'TARGET\_SCHED\_FIRST\_CYCLE\_MULTIPASS\_DFA\_LOOKAHEAD ()' subsequent ready insns to choose an insn whose issue will result in maximal number of issued insns on the same cycle. For the VLIW processor, the code could actually solve the problem of packing simple insns into the VLIW insn. Of course, if the rules of VLIW packing are described in the automaton.

This code also could be used for superscalar RISC processors. Let us consider a superscalar RISC processor with 3 pipelines. Some insns can be executed in pipelines A or B, some insns can be executed only in pipelines B or C, and one insn can be executed in pipeline B. The processor may issue the 1st insn into A and the 2nd one into B. In this case, the 3rd insn will wait for freeing B until the next cycle. If the scheduler issues the 3rd insn the first, the processor could issue all 3 insns per cycle.

Actually this code demonstrates advantages of the automaton based pipeline hazard recognizer. We try quickly and easy many insn schedules to choose the best one.

The default is no multipass scheduling.

int [Target Hook]

### TARGET\_SCHED\_FIRST\_CYCLE\_MULTIPASS\_DFA\_LOOKAHEAD\_GUARD (rtx\_insn \*insn, int ready\_index)

This hook controls what insns from the ready insn queue will be considered for the multipass insn scheduling. If the hook returns zero for *insn*, the insn will be considered in multipass scheduling. Positive return values will remove *insn* from consideration on the current round of multipass scheduling. Negative return values will remove *insn* from consideration for given number of cycles. Backends should be careful about returning non-zero for highest priority instruction at position 0 in the ready list. ready\_index is passed to allow backends make correct judgements.

The default is that any ready insns can be chosen to be issued.

void TARGET\_SCHED\_FIRST\_CYCLE\_MULTIPASS\_BEGIN (void [Target Hook]

\*data, signed char \*ready\_try, int n\_ready, bool first\_cycle\_insn\_p)

This hook prepares the target backend for a new round of multipass scheduling.

This hook is called when multipass scheduling evaluates instruction INSN.

This is called when multipass scheduling backtracks from evaluation of an instruction.

This hook notifies the target about the result of the concluded current round of multipass scheduling.

void TARGET\_SCHED\_FIRST\_CYCLE\_MULTIPASS\_INIT (void \*data) [Target Hook] This hook initializes target-specific data used in multipass scheduling.

void TARGET\_SCHED\_FIRST\_CYCLE\_MULTIPASS\_FINI (void \*data) [Target Hook] This hook finalizes target-specific data used in multipass scheduling.

### int TARGET\_SCHED\_DFA\_NEW\_CYCLE (FILE \*dump, int verbose, rtx\_insn \*insn, int last\_clock, int clock, int \*sort\_p) [Target Hook]

This hook is called by the insn scheduler before issuing *insn* on cycle *clock*. If the hook returns nonzero, *insn* is not issued on this processor cycle. Instead, the processor cycle is advanced. If \*sort\_p is zero, the insn ready queue is not sorted on the new cycle start as usually. *dump* and *verbose* specify the file and verbosity level to use for debugging output. *last\_clock* and *clock* are, respectively, the processor cycle on which the previous insn has been issued, and the current processor cycle.

### bool TARGET\_SCHED\_IS\_COSTLY\_DEPENDENCE (struct \_dep \*\_dep, int cost, int distance) [Target Hook]

This hook is used to define which dependences are considered costly by the target, so costly that it is not advisable to schedule the insns that are involved in the dependence too close to one another. The parameters to this hook are as follows: The first parameter \_dep is the dependence being evaluated. The second parameter cost is the cost of the dependence as estimated by the scheduler, and the third parameter distance is the distance in cycles between the two insns. The hook returns true if considering the distance between the two insns the dependence between them is considered costly by the target, and false otherwise.

Defining this hook can be useful in multiple-issue out-of-order machines, where (a) it's practically hopeless to predict the actual data/resource delays, however: (b) there's a better chance to predict the actual grouping that will be formed, and (c) correctly emulating the grouping can be very important. In such targets one may want to allow issuing dependent insns closer to one another—i.e., closer than the dependence distance; however, not in cases of "costly dependences", which this hooks allows to define.

### void TARGET\_SCHED\_H\_I\_D\_EXTENDED (void)

[Target Hook]

This hook is called by the insn scheduler after emitting a new instruction to the instruction stream. The hook notifies a target backend to extend its per instruction data structures.

void \* TARGET\_SCHED\_ALLOC\_SCHED\_CONTEXT (void) [Target Hook] Return a pointer to a store large enough to hold target scheduling context.

#### 

Initialize store pointed to by tc to hold target scheduling context. It  $clean_p$  is true then initialize tc as if scheduler is at the beginning of the block. Otherwise, copy the current context into tc.

### void TARGET\_SCHED\_SET\_SCHED\_CONTEXT (void \*tc) [Target Hook] Copy target scheduling context pointed to by tc to the current context.

void TARGET\_SCHED\_CLEAR\_SCHED\_CONTEXT (void \*tc) [Target Hook]

Deallocate internal data in target scheduling context pointed to by tc.

### void TARGET\_SCHED\_FREE\_SCHED\_CONTEXT (void \*tc)

[Target Hook]

Deallocate a store for target scheduling context pointed to by tc.

### int TARGET\_SCHED\_SPECULATE\_INSN (rtx\_insn \*insn, unsigned int [Target Hook] dep\_status, rtx \*new\_pat)

This hook is called by the insn scheduler when *insn* has only speculative dependencies and therefore can be scheduled speculatively. The hook is used to check if the pattern of *insn* has a speculative version and, in case of successful check, to generate that speculative pattern. The hook should return 1, if the instruction has a speculative form, or -1, if it doesn't. request describes the type of requested speculation. If the return value equals 1 then  $new_pat$  is assigned the generated speculative pattern.

bool TARGET\_SCHED\_NEEDS\_BLOCK\_P (unsigned int dep\_status) [Target Hook]
This hook is called by the insn scheduler during generation of recovery code for insn. It should return true, if the corresponding check instruction should branch to recovery code, or false otherwise.

### rtx TARGET\_SCHED\_GEN\_SPEC\_CHECK (rtx\_insn \*insn, rtx\_insn [Target Hook] \*label, unsigned int ds)

This hook is called by the insn scheduler to generate a pattern for recovery check instruction. If  $mutate_p$  is zero, then insn is a speculative instruction for which the check should be generated. label is either a label of a basic block, where recovery code should be emitted, or a null pointer, when requested check doesn't branch to recovery code (a simple check). If  $mutate_p$  is nonzero, then a pattern for a branchy check corresponding to a simple check denoted by insn should be generated. In this case label can't be null.

#### 

This hook is used by the insn scheduler to find out what features should be enabled/used. The structure \*spec\_info should be filled in by the target. The structure describes speculation types that can be used in the scheduler.

# bool TARGET\_SCHED\_CAN\_SPECULATE\_INSN (rtx\_insn \*insn) [Target Hook] Some instructions should never be speculated by the schedulers, usually because the instruction is too expensive to get this wrong. Often such instructions have long latency, and often they are not fully modeled in the pipeline descriptions. This hook should return false if insn should not be speculated.

### int TARGET\_SCHED\_SMS\_RES\_MII (struct ddg \*g)

[Target Hook]

This hook is called by the swing modulo scheduler to calculate a resource-based lower bound which is based on the resources available in the machine and the resources required by each instruction. The target backend can use g to calculate such bound. A very simple lower bound will be used in case this hook is not implemented: the total number of instructions divided by the issue rate.

## bool TARGET\_SCHED\_DISPATCH (rtx\_insn \*insn, int x) [Target Hook] This hook is called by Haifa Scheduler. It returns true if dispatch scheduling is supported in hardware and the condition specified in the parameter is true.

### void TARGET\_SCHED\_DISPATCH\_DO (rtx\_insn \*insn, int x)

[Target Hook]

This hook is called by Haifa Scheduler. It performs the operation specified in its second parameter.

### bool TARGET\_SCHED\_EXPOSED\_PIPELINE

[Target Hook]

True if the processor has an exposed pipeline, which means that not just the order of instructions is important for correctness when scheduling, but also the latencies of operations.

### int TARGET\_SCHED\_REASSOCIATION\_WIDTH (unsigned int opc,

[Target Hook]

machine\_mode mode)

This hook is called by tree reassociator to determine a level of parallelism required in output calculations chain.

### void TARGET\_SCHED\_FUSION\_PRIORITY (rtx\_insn \*insn, int max\_pri, int \*fusion\_pri, int \*pri)

[Target Hook]

This hook is called by scheduling fusion pass. It calculates fusion priorities for each instruction passed in by parameter. The priorities are returned via pointer parame-

insn is the instruction whose priorities need to be calculated. max\_pri is the maximum priority can be returned in any cases. fusion\_pri is the pointer parameter through which insn's fusion priority should be calculated and returned. pri is the pointer parameter through which insn's priority should be calculated and returned.

Same fusion\_pri should be returned for instructions which should be scheduled together. Different pri should be returned for instructions with same fusion\_pri. fusion\_pri is the major sort key, pri is the minor sort key. All instructions will be scheduled according to the two priorities. All priorities calculated should be between 0 (exclusive) and max\_pri (inclusive). To avoid false dependencies, fusion\_pri of instructions which need to be scheduled together should be smaller than fusion\_pri of irrelevant instructions.

Given below example:

```
ldr r10, [r1, 4]
add r4, r4, r10
ldr r15, [r2, 8]
sub r5, r5, r15
ldr r11, [r1, 0]
add r4, r4, r11
ldr r16, [r2, 12]
sub r5, r5, r16
```

On targets like ARM/AArch64, the two pairs of consecutive loads should be merged. Since peephole pass can't help in this case unless consecutive loads are actually next to each other in instruction flow. That's where this scheduling fusion pass works. This hook calculates priority for each instruction based on its fustion type, like:

```
ldr r10, [r1, 4] ; fusion_pri=99, pri=96
add r4, r4, r10
                ; fusion_pri=100, pri=100
ldr r15, [r2, 8] ; fusion_pri=98, pri=92
sub r5, r5, r15 ; fusion_pri=100, pri=100
ldr r11, [r1, 0] ; fusion_pri=99, pri=100
add r4, r4, r11 ; fusion_pri=100, pri=100
```

```
ldr r16, [r2, 12] ; fusion_pri=98, pri=88
sub r5, r5, r16 ; fusion_pri=100, pri=100
```

Scheduling fusion pass then sorts all ready to issue instructions according to the priorities. As a result, instructions of same fusion type will be pushed together in instruction flow, like:

```
ldr r11, [r1, 0]
ldr r10, [r1, 4]
ldr r15, [r2, 8]
ldr r16, [r2, 12]
add r4, r4, r10
sub r5, r5, r15
add r4, r4, r11
sub r5, r5, r16
```

Now peephole2 pass can simply merge the two pairs of loads.

Since scheduling fusion pass relies on peephole2 to do real fusion work, it is only enabled by default when peephole2 is in effect.

This is firstly introduced on ARM/AArch64 targets, please refer to the hook implementation for how different fusion types are supported.

```
void TARGET_EXPAND_DIVMOD_LIBFUNC (rtx libfunc, [Target Hook]

machine_mode mode, rtx op0, rtx op1, rtx *quot, rtx *rem)

Define this hook for enabling divmed transform if the port does not have hardware
```

Define this hook for enabling divmod transform if the port does not have hardware divmod insn but defines target-specific divmod libfuncs.

### 18.18 Dividing the Output into Sections (Texts, Data, ...)

An object file is divided into sections containing different types of data. In the most common case, there are three sections: the *text section*, which holds instructions and read-only data; the *data section*, which holds initialized writable data; and the *bss section*, which holds uninitialized data. Some systems have other kinds of sections.

'varasm.c' provides several well-known sections, such as text\_section, data\_section and bss\_section. The normal way of controlling a foo\_section variable is to define the associated FOO\_SECTION\_ASM\_OP macro, as described below. The macros are only read once, when 'varasm.c' initializes itself, so their values must be run-time constants. They may however depend on command-line flags.

*Note:* Some run-time files, such 'crtstuff.c', also make use of the FOO\_SECTION\_ASM\_OP macros, and expect them to be string literals.

Some assemblers require a different string to be written every time a section is selected. If your assembler falls into this category, you should define the TARGET\_ASM\_INIT\_SECTIONS hook and use get\_unnamed\_section to set up the sections.

You must always create a text\_section, either by defining TEXT\_SECTION\_ASM\_OP or by initializing text\_section in TARGET\_ASM\_INIT\_SECTIONS. The same is true of data\_section and DATA\_SECTION\_ASM\_OP. If you do not create a distinct readonly\_data\_section, the default is to reuse text\_section.

All the other 'varasm.c' sections are optional, and are null if the target does not provide them.

### TEXT\_SECTION\_ASM\_OP

[Macro]

A C expression whose value is a string, including spacing, containing the assembler operation that should precede instructions and read-only data. Normally "\t.text" is right.

### HOT\_TEXT\_SECTION\_NAME

[Macro]

If defined, a C string constant for the name of the section containing most frequently executed functions of the program. If not defined, GCC will provide a default definition if the target supports named sections.

### UNLIKELY\_EXECUTED\_TEXT\_SECTION\_NAME

[Macro]

If defined, a C string constant for the name of the section containing unlikely executed functions in the program.

### DATA\_SECTION\_ASM\_OP

[Macro]

A C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as writable initialized data. Normally "\t.data" is right.

### SDATA\_SECTION\_ASM\_OP

[Macro]

If defined, a C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as initialized, writable small data.

### READONLY\_DATA\_SECTION\_ASM\_OP

[Macro]

A C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as read-only initialized data.

### BSS\_SECTION\_ASM\_OP

[Macro]

If defined, a C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as uninitialized global data. If not defined, and ASM\_OUTPUT\_ALIGNED\_BSS not defined, uninitialized global data will be output in the data section if '-fno-common' is passed, otherwise ASM\_OUTPUT\_COMMON will be used.

### SBSS\_SECTION\_ASM\_OP

[Macro]

If defined, a C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as uninitialized, writable small data.

### TLS\_COMMON\_ASM\_OP

[Macro]

If defined, a C expression whose value is a string containing the assembler operation to identify the following data as thread-local common data. The default is ".tls\_common".

### TLS\_SECTION\_ASM\_FLAG

[Macro]

If defined, a C expression whose value is a character constant containing the flag used to mark a section as a TLS section. The default is 'T'.

### INIT\_SECTION\_ASM\_OP

[Macro]

If defined, a C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as initialization code. If not defined,

GCC will assume such a section does not exist. This section has no corresponding init\_section variable; it is used entirely in runtime code.

#### FINI\_SECTION\_ASM\_OP

[Macro]

If defined, a C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as finalization code. If not defined, GCC will assume such a section does not exist. This section has no corresponding fini\_section variable; it is used entirely in runtime code.

### INIT\_ARRAY\_SECTION\_ASM\_OP

[Macro]

If defined, a C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as part of the .init\_array (or equivalent) section. If not defined, GCC will assume such a section does not exist. Do not define both this macro and INIT\_SECTION\_ASM\_OP.

### FINI\_ARRAY\_SECTION\_ASM\_OP

[Macro]

If defined, a C expression whose value is a string, including spacing, containing the assembler operation to identify the following data as part of the .fini\_array (or equivalent) section. If not defined, GCC will assume such a section does not exist. Do not define both this macro and FINI\_SECTION\_ASM\_OP.

### MACH\_DEP\_SECTION\_ASM\_FLAG

[Macro]

If defined, a C expression whose value is a character constant containing the flag used to mark a machine-dependent section. This corresponds to the SECTION\_MACH\_DEP section flag.

### CRT\_CALL\_STATIC\_FUNCTION (section\_op, function)

[Macro]

If defined, an ASM statement that switches to a different section via  $section\_op$ , calls function, and switches back to the text section. This is used in 'crtstuff.c' if INIT\_SECTION\_ASM\_OP or FINI\_SECTION\_ASM\_OP to calls to initialization and finalization functions from the init and fini sections. By default, this macro uses a simple function call. Some ports need hand-crafted assembly code to avoid dependencies on registers initialized in the function prologue or to ensure that constant pools don't end up too far way in the text section.

### TARGET\_LIBGCC\_SDATA\_SECTION

[Macro]

If defined, a string which names the section into which small variables defined in crtstuff and libgcc should go. This is useful when the target has options for optimizing access to small data, and you want the crtstuff and libgcc routines to be conservative in what they expect of your application yet liberal in what your application expects. For example, for targets with a .sdata section (like MIPS), you could compile crtstuff with -G 0 so that it doesn't require small data support from your application, but use this macro to put small data into .sdata so that your application can access these variables whether it uses small data or not.

### FORCE\_CODE\_SECTION\_ALIGN

Macro

If defined, an ASM statement that aligns a code section to some arbitrary boundary. This is used to force all fragments of the .init and .fini sections to have to same alignment and thus prevent the linker from having to add any padding.

### JUMP\_TABLES\_IN\_TEXT\_SECTION

[Macro]

Define this macro to be an expression with a nonzero value if jump tables (for tablejump insns) should be output in the text section, along with the assembler instructions. Otherwise, the readonly data section is used.

This macro is irrelevant if there is no separate readonly data section.

### void TARGET\_ASM\_INIT\_SECTIONS (void)

[Target Hook]

Define this hook if you need to do something special to set up the 'varasm.c' sections, or if your target has some special sections of its own that you need to create.

GCC calls this hook after processing the command line, but before writing any assembly code, and before calling any of the section-returning hooks described below.

### int TARGET\_ASM\_RELOC\_RW\_MASK (void)

[Target Hook]

Return a mask describing how relocations should be treated when selecting sections. Bit 1 should be set if global relocations should be placed in a read-write section; bit 0 should be set if local relocations should be placed in a read-write section.

The default version of this function returns 3 when '-fpic' is in effect, and 0 otherwise. The hook is typically redefined when the target cannot support (some kinds of) dynamic relocations in read-only sections even in executables.

### bool TARGET\_ASM\_GENERATE\_PIC\_ADDR\_DIFF\_VEC (void)

[Target Hook]

Return true to generate ADDR\_DIF\_VEC table or false to generate ADDR\_VEC table for jumps in case of -fPIC.

The default version of this function returns true if flag\_pic equals true and false otherwise

### section \* TARGET\_ASM\_SELECT\_SECTION (tree exp, int reloc, unsigned HOST\_WIDE\_INT align) [Target Hook]

Return the section into which exp should be placed. You can assume that exp is either a VAR\_DECL node or a constant of some sort. reloc indicates whether the initial value of exp requires link-time relocations. Bit 0 is set when variable contains local relocations only, while bit 1 is set for global relocations. align is the constant alignment in bits.

The default version of this function takes care of putting read-only variables in readonly\_data\_section.

See also USE\_SELECT\_SECTION\_FOR\_FUNCTIONS.

### USE\_SELECT\_SECTION\_FOR\_FUNCTIONS

[Macro]

Define this macro if you wish TARGET\_ASM\_SELECT\_SECTION to be called for FUNCTION\_DECLs as well as for variables and constants.

In the case of a FUNCTION\_DECL, reloc will be zero if the function has been determined to be likely to be called, and nonzero if it is unlikely to be called.

### void TARGET\_ASM\_UNIQUE\_SECTION (tree decl, int reloc)

[Target Hook]

Build up a unique section name, expressed as a STRING\_CST node, and assign it to 'DECL\_SECTION\_NAME (decl)'. As with TARGET\_ASM\_SELECT\_SECTION, reloc indicates whether the initial value of exp requires link-time relocations.

The default version of this function appends the symbol name to the ELF section name that would normally be used for the symbol. For example, the function foo would be placed in .text.foo. Whatever the actual target object format, this is often good enough.

- Return the readonly data section associated with 'DECL\_SECTION\_NAME (decl)'. The default version of this function selects .gnu.linkonce.r.name if the function's section is .gnu.linkonce.t.name, .rodata.name if function is in .text.name, and the normal readonly-data section otherwise.
- const char \* TARGET\_ASM\_MERGEABLE\_RODATA\_PREFIX [Target Hook]
  Usually, the compiler uses the prefix ".rodata" to construct section names for mergeable constant data. Define this macro to override the string if a different section name should be used.
- section \* TARGET\_ASM\_TM\_CLONE\_TABLE\_SECTION (void) [Target Hook] Return the section that should be used for transactional memory clone tables.

Return the section into which a constant x, of mode mode, should be placed. You can assume that x is some kind of constant in RTL. The argument mode is redundant except in the case of a const\_int rtx. align is the constant alignment in bits.

The default version of this function takes care of putting symbolic constants in flag\_pic mode in data\_section and everything else in readonly\_data\_section.

- tree TARGET\_MANGLE\_DECL\_ASSEMBLER\_NAME (tree dec1, tree id) [Target Hook] Define this hook if you need to postprocess the assembler name generated by target-independent code. The id provided to this hook will be the computed name (e.g., the macro DECL\_NAME of the decl in C, or the mangled name of the decl in C++). The return value of the hook is an IDENTIFIER\_NODE for the appropriate mangled name on your target system. The default implementation of this hook just returns the id provided.

Define this hook if references to a symbol or a constant must be treated differently depending on something about the variable or function named by the symbol (such as what section it is in).

The hook is executed immediately after rtl has been created for decl, which may be a variable or function declaration or an entry in the constant pool. In either case, rtl is the rtl in question. Do not use DECL\_RTL (decl) in this hook; that field may not have been initialized yet.

In the case of a constant, it is safe to assume that the rtl is a mem whose address is a symbol\_ref. Most decls will also have this form, but that is not guaranteed. Global register variables, for instance, will have a reg for their rtl. (Normally the right thing to do with such unusual rtl is leave it alone.)

The new\_decl\_p argument will be true if this is the first time that TARGET\_ENCODE\_ SECTION\_INFO has been invoked on this decl. It will be false for subsequent invocations, which will happen for duplicate declarations. Whether or not anything must be done for the duplicate declaration depends on whether the hook examines DECL\_ ATTRIBUTES. new\_decl\_p is always true when the hook is called for a constant.

The usual thing for this hook to do is to record flags in the symbol\_ref, using SYMBOL\_REF\_FLAG or SYMBOL\_REF\_FLAGS. Historically, the name string was modified if it was necessary to encode more than one bit of information, but this practice is now discouraged; use SYMBOL\_REF\_FLAGS.

The default definition of this hook, default\_encode\_section\_info in 'varasm.c', sets a number of commonly-useful bits in SYMBOL\_REF\_FLAGS. Check whether the default does what you need before overriding it.

const char \* TARGET\_STRIP\_NAME\_ENCODING (const char \*name) [Target Hook]

Decode name and return the real name part, sans the characters that TARGET\_ENCODE\_

SECTION\_INFO may have added.

### bool TARGET\_IN\_SMALL\_DATA\_P (const\_tree exp)

[Target Hook]

Returns true if exp should be placed into a "small data" section. The default version of this hook always returns false.

### bool TARGET\_HAVE\_SRODATA\_SECTION

[Target Hook]

Contains the value true if the target places read-only "small data" into a separate section. The default value is false.

### bool TARGET\_PROFILE\_BEFORE\_PROLOGUE (void)

[Target Hook]

It returns true if target wants profile code emitted before prologue.

The default version of this hook use the target macro PROFILE\_BEFORE\_PROLOGUE.

### bool TARGET\_BINDS\_LOCAL\_P (const\_tree exp)

[Target Hook]

Returns true if exp names an object for which name resolution rules must resolve to the current "module" (dynamic shared library or executable image).

The default version of this hook implements the name resolution rules for ELF, which has a looser model of global name binding than other currently supported object file formats.

### bool TARGET\_HAVE\_TLS

[Target Hook]

Contains the value true if the target supports thread-local storage. The default value is false.

### 18.19 Position Independent Code

This section describes macros that help implement generation of position independent code. Simply defining these macros is not enough to generate valid PIC; you must also add support to the hook TARGET\_LEGITIMATE\_ADDRESS\_P and to the macro PRINT\_OPERAND\_ADDRESS, as well as LEGITIMIZE\_ADDRESS. You must modify the definition of 'movsi' to do something appropriate when the source operand contains a symbolic address. You may also need to alter the handling of switch statements so that they use relative addresses.

### PIC\_OFFSET\_TABLE\_REGNUM

[Macro]

The register number of the register used to address a table of static data addresses in memory. In some cases this register is defined by a processor's "application binary interface" (ABI). When this macro is defined, RTL is generated for this register once, as with the stack pointer and frame pointer registers. If this macro is not defined, it is up to the machine-dependent files to allocate such a register (if necessary). Note that this register must be fixed when in use (e.g. when flag\_pic is true).

### PIC\_OFFSET\_TABLE\_REG\_CALL\_CLOBBERED

[Macro]

A C expression that is nonzero if the register defined by PIC\_OFFSET\_TABLE\_REGNUM is clobbered by calls. If not defined, the default is zero. Do not define this macro if PIC\_OFFSET\_TABLE\_REGNUM is not defined.

### LEGITIMATE\_PIC\_OPERAND\_P (x)

[Macro]

A C expression that is nonzero if x is a legitimate immediate operand on the target machine when generating position independent code. You can assume that x satisfies CONSTANT\_P, so you need not check this. You can also assume flag\_pic is true, so you need not check it either. You need not define this macro if all constants (including SYMBOL\_REF) can be immediate operands when generating position independent code.

### 18.20 Defining the Output Assembler Language

This section describes macros whose principal purpose is to describe how to write instructions in assembler language—rather than what the instructions do.

### 18.20.1 The Overall Framework of an Assembler File

This describes the overall framework of an assembly file.

### void TARGET\_ASM\_FILE\_START (void)

[Target Hook]

Output to asm\_out\_file any text which the assembler expects to find at the beginning of a file. The default behavior is controlled by two flags, documented below. Unless your target's assembler is quite unusual, if you override the default, you should call default\_file\_start at some point in your target hook. This lets other target files rely on these variables.

#### bool TARGET\_ASM\_FILE\_START\_APP\_OFF

[Target Hook]

If this flag is true, the text of the macro ASM\_APP\_OFF will be printed as the very first line in the assembly file, unless '-fverbose-asm' is in effect. (If that macro has been defined to the empty string, this variable has no effect.) With the normal definition of ASM\_APP\_OFF, the effect is to notify the GNU assembler that it need not bother stripping comments or extra whitespace from its input. This allows it to work a bit faster.

The default is false. You should not set it to true unless you have verified that your port does not generate any extra whitespace or comments that will cause GAS to issue errors in NO\_APP mode.

### bool TARGET\_ASM\_FILE\_START\_FILE\_DIRECTIVE

[Target Hook]

If this flag is true, output\_file\_directive will be called for the primary source file, immediately after printing ASM\_APP\_OFF (if that is enabled). Most ELF assemblers expect this to be done. The default is false.

## void TARGET\_ASM\_FILE\_END (void)

[Target Hook]

Output to asm\_out\_file any text which the assembler expects to find at the end of a file. The default is to output nothing.

#### void file\_end\_indicate\_exec\_stack ()

[Function]

Some systems use a common convention, the '.note.GNU-stack' special section, to indicate whether or not an object file relies on the stack being executable. If your system uses this convention, you should define TARGET\_ASM\_FILE\_END to this function. If you need to do other things in that hook, have your hook function call this function.

#### void TARGET\_ASM\_LTO\_START (void)

[Target Hook]

Output to asm\_out\_file any text which the assembler expects to find at the start of an LTO section. The default is to output nothing.

#### void TARGET\_ASM\_LTO\_END (void)

[Target Hook]

Output to asm\_out\_file any text which the assembler expects to find at the end of an LTO section. The default is to output nothing.

## void TARGET\_ASM\_CODE\_END (void)

[Target Hook]

Output to asm\_out\_file any text which is needed before emitting unwind info and debug info at the end of a file. Some targets emit here PIC setup thunks that cannot be emitted at the end of file, because they couldn't have unwind info then. The default is to output nothing.

#### ASM\_COMMENT\_START

[Macro]

A C string constant describing how to begin a comment in the target assembler language. The compiler assumes that the comment will end at the end of the line.

### ASM\_APP\_ON

Macro

A C string constant for text to be output before each asm statement or group of consecutive ones. Normally this is "#APP", which is a comment that has no effect on most assemblers but tells the GNU assembler that it must check the lines that follow for all valid assembler constructs.

## ASM\_APP\_OFF

[Macro]

A C string constant for text to be output after each asm statement or group of consecutive ones. Normally this is "#NO\_APP", which tells the GNU assembler to resume making the time-saving assumptions that are valid for ordinary compiler output.

### ASM\_OUTPUT\_SOURCE\_FILENAME (stream, name)

Macro

A C statement to output COFF information or DWARF debugging information which indicates that filename *name* is the current source file to the stdio stream *stream*.

This macro need not be defined if the standard form of output for the file format in use is appropriate.

#### 

Output DWARF debugging information which indicates that filename name is the current source file to the stdio stream file.

This target hook need not be defined if the standard form of output for the file format in use is appropriate.

## void TARGET\_ASM\_OUTPUT\_IDENT (const char \*name)

[Target Hook]

Output a string based on *name*, suitable for the '#ident' directive, or the equivalent directive or pragma in non-C-family languages. If this hook is not defined, nothing is output for the '#ident' directive.

## OUTPUT\_QUOTED\_STRING (stream, string)

[Macro]

A C statement to output the string string to the stdio stream stream. If you do not call the function output\_quoted\_string in your config files, GCC will only call it to output filenames to the assembler source. So you can use it to canonicalize the format of the filename using this macro.

#### 

Output assembly directives to switch to section name. The section should have attributes as specified by flags, which is a bit mask of the SECTION\_\* flags defined in 'output.h'. If decl is non-NULL, it is the VAR\_DECL or FUNCTION\_DECL with which this section is associated.

# bool TARGET\_ASM\_ELF\_FLAGS\_NUMERIC (unsigned int flags, unsigned int \*num) [Target Hook]

This hook can be used to encode ELF section flags for which no letter code has been defined in the assembler. It is called by default\_asm\_named\_section whenever the section flags need to be emitted in the assembler output. If the hook returns true, then the numerical value for ELF section flags should be calculated from flags and saved in \*num; the value is printed out instead of the normal sequence of letter codes. If the hook is not defined, or if it returns false, then num is ignored and the traditional letter sequence is emitted.

#### 

Return preferred text (sub)section for function decl. Main purpose of this function is to separate cold, normal and hot functions. startup is true when function is known to be used only at startup (from static constructors or it is main()). exit is true when function is known to be used only at exit (from static destructors). Return NULL if function should go to default text section.

#### 

Used by the target to emit any assembler directives or additional labels needed when a function is partitioned between different sections. Output should be written to file. The function decl is available as decl and the new section is 'cold' if new\_is\_cold is true.

#### bool TARGET\_HAVE\_NAMED\_SECTIONS

[Common Target Hook]

This flag is true if the target supports TARGET\_ASM\_NAMED\_SECTION. It must not be modified by command-line option processing.

#### bool TARGET\_HAVE\_SWITCHABLE\_BSS\_SECTIONS

[Target Hook]

This flag is true if we can create zeroed data by switching to a BSS section and then using ASM\_OUTPUT\_SKIP to allocate the space. This is true on most ELF targets.

# unsigned int TARGET\_SECTION\_TYPE\_FLAGS (tree decl, const char \*name, int reloc) [Target Hook]

Choose a set of section attributes for use by TARGET\_ASM\_NAMED\_SECTION based on a variable or function decl, a section name, and whether or not the declaration's initializer may contain runtime relocations. *decl* may be null, in which case readwrite data should be assumed.

The default version of this function handles choosing code vs data, read-only vs read-write data, and flag\_pic. You should only need to override this if your target has special flags that might be set via \_\_attribute\_\_.

#### 

Provides the target with the ability to record the gcc command line switches that have been passed to the compiler, and options that are enabled. The *type* argument specifies what is being recorded. It can take the following values:

#### SWITCH\_TYPE\_PASSED

text is a command line switch that has been set by the user.

#### SWITCH\_TYPE\_ENABLED

text is an option which has been enabled. This might be as a direct result of a command line switch, or because it is enabled by default or because it has been enabled as a side effect of a different command line switch. For example, the '-02' switch enables various different individual optimization passes.

#### SWITCH\_TYPE\_DESCRIPTIVE

text is either NULL or some descriptive text which should be ignored. If text is NULL then it is being used to warn the target hook that either recording is starting or ending. The first time type is SWITCH\_TYPE\_DESCRIPTIVE and text is NULL, the warning is for start up and the second time the warning is for wind down. This feature is to allow the target hook to make any necessary preparations before it starts to record switches and to perform any necessary tidying up after it has finished recording switches.

#### SWITCH\_TYPE\_LINE\_START

This option can be ignored by this target hook.

#### SWITCH\_TYPE\_LINE\_END

This option can be ignored by this target hook.

The hook's return value must be zero. Other return values may be supported in the future.

By default this hook is set to NULL, but an example implementation is provided for ELF based targets. Called *elf\_record\_gcc\_switches*, it records the switches as ASCII text inside a new, string mergeable section in the assembler output file. The name of the new section is provided by the TARGET\_ASM\_RECORD\_GCC\_SWITCHES\_SECTION target hook.

const char \* TARGET\_ASM\_RECORD\_GCC\_SWITCHES\_SECTION [Target Hook] This is the name of the section that will be created by the example ELF implementation of the TARGET\_ASM\_RECORD\_GCC\_SWITCHES target hook.

# 18.20.2 Output of Data

const	char	*	TARGET_ASM_BYTE_OP	[Target	Hook]
const	char	*	TARGET_ASM_ALIGNED_HI_OP	[Target	Hook]
const	char	*	TARGET_ASM_ALIGNED_PSI_OP	[Target	Hook]
const	char	*	TARGET_ASM_ALIGNED_SI_OP	[Target	Hook]
const	char	*	TARGET_ASM_ALIGNED_PDI_OP	[Target	Hook]
const	char	*	TARGET_ASM_ALIGNED_DI_OP	[Target	Hook]
const	char	*	TARGET_ASM_ALIGNED_PTI_OP	[Target	Hook]
const	char	*	TARGET_ASM_ALIGNED_TI_OP	[Target	Hook]
const	char	*	TARGET_ASM_UNALIGNED_HI_OP	[Target	Hook]
const	char	*	TARGET_ASM_UNALIGNED_PSI_OP	[Target	Hook]
const	char	*	TARGET_ASM_UNALIGNED_SI_OP	[Target	Hook]
const	char	*	TARGET_ASM_UNALIGNED_PDI_OP	[Target	Hook]
const	char	*	TARGET_ASM_UNALIGNED_DI_OP	[Target	Hook]
const	char	*	TARGET_ASM_UNALIGNED_PTI_OP	[Target	Hook]
const	char	*	TARGET_ASM_UNALIGNED_TI_OP	[Target	Hook]
_					

These hooks specify assembly directives for creating certain kinds of integer object. The TARGET\_ASM\_BYTE\_OP directive creates a byte-sized object, the TARGET\_ASM\_ALIGNED\_HI\_OP one creates an aligned two-byte object, and so on. Any of the hooks may be NULL, indicating that no suitable directive is available.

The compiler will print these strings at the start of a new line, followed immediately by the object's initial value. In most cases, the string should contain a tab, a pseudo-op, and then another tab.

# bool TARGET\_ASM\_INTEGER (rtx x, unsigned int size, int aligned\_p) [Target Hook]

The assemble\_integer function uses this hook to output an integer object. x is the object's value, size is its size in bytes and aligned\_p indicates whether it is aligned. The function should return true if it was able to output the object. If it returns false, assemble\_integer will try to split the object into smaller parts.

The default implementation of this hook will use the TARGET\_ASM\_BYTE\_OP family of strings, returning false when the relevant string is NULL.

# void TARGET\_ASM\_DECL\_END (void) [Target Hook]

Define this hook if the target assembler requires a special marker to terminate an initialized variable declaration.

bool TARGET\_ASM\_OUTPUT\_ADDR\_CONST\_EXTRA (FILE \*file, rtx x) [Target Hook] A target hook to recognize rtx patterns that output\_addr\_const can't deal with, and output assembly code to file corresponding to the pattern x. This may be used to allow machine-dependent UNSPECs to appear within constants.

If target hook fails to recognize a pattern, it must return false, so that a standard error message is printed. If it prints an error message itself, by calling, for example, output\_operand\_lossage, it may just return true.

### ASM\_OUTPUT\_ASCII (stream, ptr, len)

[Macro]

A C statement to output to the stdio stream stream an assembler instruction to assemble a string constant containing the len bytes at ptr. ptr will be a C expression of type char \* and len a C expression of type int.

If the assembler has a .ascii pseudo-op as found in the Berkeley Unix assembler, do not define the macro ASM\_OUTPUT\_ASCII.

#### ASM\_OUTPUT\_FDESC (stream, decl, n)

[Macro]

A C statement to output word n of a function descriptor for decl. This must be defined if TARGET\_VTABLE\_USES\_DESCRIPTORS is defined, and is otherwise unused.

#### CONSTANT\_POOL\_BEFORE\_FUNCTION

[Macro]

You may define this macro as a C expression. You should define the expression to have a nonzero value if GCC should output the constant pool for a function before the code for the function, or a zero value if GCC should output the constant pool after the function. If you do not define this macro, the usual case, GCC will output the constant pool before the function.

#### ASM\_OUTPUT\_POOL\_PROLOGUE (file, funname, fundecl, size)

[Macro]

A C statement to output assembler commands to define the start of the constant pool for a function. funname is a string giving the name of the function. Should the return type of the function be required, it can be obtained via fundecl. size is the size, in bytes, of the constant pool that will be written immediately after this call.

If no constant-pool prefix is required, the usual case, this macro need not be defined.

# ASM\_OUTPUT\_SPECIAL\_POOL\_ENTRY (file, x, mode, align, labelno, jumpto) [Macro]

A C statement (with or without semicolon) to output a constant in the constant pool, if it needs special treatment. (This macro need not do anything for RTL expressions that can be output normally.)

The argument file is the standard I/O stream to output the assembler code on. x is the RTL expression for the constant to output, and mode is the machine mode (in case x is a 'const\_int'). align is the required alignment for the value x; you should output an assembler directive to force this much alignment.

The argument *labelno* is a number to use in an internal label for the address of this pool entry. The definition of this macro is responsible for outputting the label definition at the proper place. Here is how to do this:

```
(*targetm.asm_out.internal_label) (file, "LC", labelno);
```

When you output a pool entry specially, you should end with a **goto** to the label *jumpto*. This will prevent the same pool entry from being output a second time in the usual manner.

You need not define this macro if it would do nothing.

## ASM\_OUTPUT\_POOL\_EPILOGUE (file funname fundecl size)

[Macro]

A C statement to output assembler commands to at the end of the constant pool for a function. *funname* is a string giving the name of the function. Should the return type of the function be required, you can obtain it via *fundecl. size* is the size, in bytes, of the constant pool that GCC wrote immediately before this call.

If no constant-pool epilogue is required, the usual case, you need not define this macro.

#### IS\_ASM\_LOGICAL\_LINE\_SEPARATOR (C, STR)

[Macro]

Define this macro as a C expression which is nonzero if C is used as a logical line separator by the assembler. STR points to the position in the string where C was found; this can be used if a line separator uses multiple characters.

If you do not define this macro, the default is that only the character ';' is treated as a logical line separator.

```
const char * TARGET_ASM_OPEN_PAREN
const char * TARGET_ASM_CLOSE_PAREN
```

[Target Hook]

[Target Hook]

These target hooks are C string constants, describing the syntax in the assembler for grouping arithmetic expressions. If not overridden, they default to normal parentheses, which is correct for most assemblers.

These macros are provided by 'real.h' for writing the definitions of ASM\_OUTPUT\_DOUBLE and the like:

```
REAL_VALUE_TO_TARGET_SINGLE (x, 1) [Macro] REAL_VALUE_TO_TARGET_DOUBLE (x, 1) [Macro] REAL_VALUE_TO_TARGET_LONG_DOUBLE (x, 1) [Macro] REAL_VALUE_TO_TARGET_DECIMAL32 (x, 1) [Macro] REAL_VALUE_TO_TARGET_DECIMAL64 (x, 1) [Macro] REAL_VALUE_TO_TARGET_DECIMAL128 (x, 1) [Macro]
```

These translate x, of type REAL\_VALUE\_TYPE, to the target's floating point representation, and store its bit pattern in the variable l. For REAL\_VALUE\_TO\_TARGET\_SINGLE and REAL\_VALUE\_TO\_TARGET\_DECIMAL32, this variable should be a simple long int. For the others, it should be an array of long int. The number of elements in this array is determined by the size of the desired target floating point data type: 32 bits of it go in each long int array element. Each array element holds 32 bits of the result, even if long int is wider than 32 bits on the host machine.

The array element values are designed so that you can print them out using fprintf in the order they should appear in the target machine's memory.

# 18.20.3 Output of Uninitialized Variables

Each of the macros in this section is used to do the whole job of outputting a single uninitialized variable.

## ASM\_OUTPUT\_COMMON (stream, name, size, rounded)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream the assembler definition of a common-label named name whose size is size bytes. The variable rounded is the size rounded up to whatever alignment the caller wants. It is possible

that size may be zero, for instance if a struct with no other member than a zerolength array is defined. In this case, the backend must output a symbol definition that allocates at least one byte, both so that the address of the resulting object does not compare equal to any other, and because some object formats cannot even express the concept of a zero-sized common symbol, as that is how they represent an ordinary undefined external.

Use the expression assemble\_name (stream, name) to output the name itself; before and after that, output the additional assembler syntax for defining the name, and a newline.

This macro controls how the assembler definitions of uninitialized common global variables are output.

## ASM\_OUTPUT\_ALIGNED\_COMMON (stream, name, size, alignment) [Macro]

Like ASM\_OUTPUT\_COMMON except takes the required alignment as a separate, explicit argument. If you define this macro, it is used in place of ASM\_OUTPUT\_COMMON, and gives you more flexibility in handling the required alignment of the variable. The alignment is specified as the number of bits.

# ASM\_OUTPUT\_ALIGNED\_DECL\_COMMON (stream, decl, name, size, alignment) [Macro]

Like ASM\_OUTPUT\_ALIGNED\_COMMON except that *decl* of the variable to be output, if there is one, or NULL\_TREE if there is no corresponding variable. If you define this macro, GCC will use it in place of both ASM\_OUTPUT\_COMMON and ASM\_OUTPUT\_ALIGNED\_COMMON. Define this macro when you need to see the variable's decl in order to chose what to output.

## ASM\_OUTPUT\_ALIGNED\_BSS (stream, decl, name, size, alignment) [Macro]

A C statement (sans semicolon) to output to the stdio stream stream the assembler definition of uninitialized global decl named name whose size is size bytes. The variable alignment is the alignment specified as the number of bits.

Try to use function asm\_output\_aligned\_bss defined in file 'varasm.c' when defining this macro. If unable, use the expression assemble\_name (stream, name) to output the name itself; before and after that, output the additional assembler syntax for defining the name, and a newline.

There are two ways of handling global BSS. One is to define this macro. The other is to have TARGET\_ASM\_SELECT\_SECTION return a switchable BSS section (see [TARGET\_HAVE\_SWITCHABLE\_BSS\_SECTIONS], page 598). You do not need to do both.

Some languages do not have **common** data, and require a non-common form of global BSS in order to handle uninitialized globals efficiently. C++ is one example of this. However, if the target does not support global BSS, the front end may choose to make globals common in order to save space in the object file.

## ASM\_OUTPUT\_LOCAL (stream, name, size, rounded)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream the assembler definition of a local-common-label named name whose size is size bytes. The variable rounded is the size rounded up to whatever alignment the caller wants.

Use the expression assemble\_name (stream, name) to output the name itself; before and after that, output the additional assembler syntax for defining the name, and a newline.

This macro controls how the assembler definitions of uninitialized static variables are output.

## ASM\_OUTPUT\_ALIGNED\_LOCAL (stream, name, size, alignment)

[Macro]

Like ASM\_OUTPUT\_LOCAL except takes the required alignment as a separate, explicit argument. If you define this macro, it is used in place of ASM\_OUTPUT\_LOCAL, and gives you more flexibility in handling the required alignment of the variable. The alignment is specified as the number of bits.

# ASM\_OUTPUT\_ALIGNED\_DECL\_LOCAL (stream, decl, name, size, alignment) [Macro]

Like ASM\_OUTPUT\_ALIGNED\_LOCAL except that decl of the variable to be output, if there is one, or NULL\_TREE if there is no corresponding variable. If you define this macro, GCC will use it in place of both ASM\_OUTPUT\_LOCAL and ASM\_OUTPUT\_ALIGNED\_LOCAL. Define this macro when you need to see the variable's decl in order to chose what to output.

# 18.20.4 Output and Generation of Labels

This is about outputting labels.

## ASM\_OUTPUT\_LABEL (stream, name)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream the assembler definition of a label named name. Use the expression assemble\_name (stream, name) to output the name itself; before and after that, output the additional assembler syntax for defining the name, and a newline. A default definition of this macro is provided which is correct for most systems.

## ASM\_OUTPUT\_FUNCTION\_LABEL (stream, name, decl)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream the assembler definition of a label named name of a function. Use the expression assemble\_name (stream, name) to output the name itself; before and after that, output the additional assembler syntax for defining the name, and a newline. A default definition of this macro is provided which is correct for most systems.

If this macro is not defined, then the function name is defined in the usual manner as a label (by means of ASM\_OUTPUT\_LABEL).

#### ASM\_OUTPUT\_INTERNAL\_LABEL (stream, name)

[Macro]

Identical to ASM\_OUTPUT\_LABEL, except that name is known to refer to a compiler-generated label. The default definition uses assemble\_name\_raw, which is like assemble\_name except that it is more efficient.

SIZE\_ASM\_OP [Macro]

A C string containing the appropriate assembler directive to specify the size of a symbol, without any arguments. On systems that use ELF, the default (in 'config/elfos.h') is '"\t.size\t"'; on other systems, the default is not to define this macro.

Define this macro only if it is correct to use the default definitions of ASM\_OUTPUT\_SIZE\_DIRECTIVE and ASM\_OUTPUT\_MEASURED\_SIZE for your system. If you need your own custom definitions of those macros, or if you do not need explicit symbol sizes at all, do not define this macro.

## ASM\_OUTPUT\_SIZE\_DIRECTIVE (stream, name, size)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream a directive telling the assembler that the size of the symbol name is size. size is a HOST\_WIDE\_INT. If you define SIZE\_ASM\_OP, a default definition of this macro is provided.

### ASM\_OUTPUT\_MEASURED\_SIZE (stream, name)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream a directive telling the assembler to calculate the size of the symbol name by subtracting its address from the current address.

If you define SIZE\_ASM\_OP, a default definition of this macro is provided. The default assumes that the assembler recognizes a special '.' symbol as referring to the current address, and can calculate the difference between this and another symbol. If your assembler does not recognize '.' or cannot do calculations with it, you will need to redefine ASM\_OUTPUT\_MEASURED\_SIZE to use some other technique.

#### NO\_DOLLAR\_IN\_LABEL

[Macro]

Define this macro if the assembler does not accept the character '\$' in label names. By default constructors and destructors in G++ have '\$' in the identifiers. If this macro is defined, '.' is used instead.

## NO\_DOT\_IN\_LABEL

[Macro]

Define this macro if the assembler does not accept the character '.' in label names. By default constructors and destructors in G++ have names that use '.'. If this macro is defined, these names are rewritten to avoid '.'.

TYPE\_ASM\_OP [Macro]

A C string containing the appropriate assembler directive to specify the type of a symbol, without any arguments. On systems that use ELF, the default (in 'config/elfos.h') is '"\t.type\t"'; on other systems, the default is not to define this macro.

Define this macro only if it is correct to use the default definition of ASM\_OUTPUT\_TYPE\_DIRECTIVE for your system. If you need your own custom definition of this macro, or if you do not need explicit symbol types at all, do not define this macro.

### TYPE\_OPERAND\_FMT

Macro

A C string which specifies (using printf syntax) the format of the second operand to TYPE\_ASM\_OP. On systems that use ELF, the default (in 'config/elfos.h') is '"@%s"'; on other systems, the default is not to define this macro.

Define this macro only if it is correct to use the default definition of ASM\_OUTPUT\_ TYPE\_DIRECTIVE for your system. If you need your own custom definition of this macro, or if you do not need explicit symbol types at all, do not define this macro.

### ASM\_OUTPUT\_TYPE\_DIRECTIVE (stream, type)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream a directive telling the assembler that the type of the symbol name is type. type is a C string; currently, that string is always either '"function" or '"object", but you should not count on this.

If you define TYPE\_ASM\_OP and TYPE\_OPERAND\_FMT, a default definition of this macro is provided.

#### ASM\_DECLARE\_FUNCTION\_NAME (stream, name, decl)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream any text necessary for declaring the name name of a function which is being defined. This macro is responsible for outputting the label definition (perhaps using ASM\_OUTPUT\_FUNCTION\_LABEL). The argument decl is the FUNCTION\_DECL tree node representing the function.

If this macro is not defined, then the function name is defined in the usual manner as a label (by means of ASM\_OUTPUT\_FUNCTION\_LABEL).

You may wish to use ASM\_OUTPUT\_TYPE\_DIRECTIVE in the definition of this macro.

## ASM\_DECLARE\_FUNCTION\_SIZE (stream, name, decl)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream any text necessary for declaring the size of a function which is being defined. The argument name is the name of the function. The argument decl is the FUNCTION\_DECL tree node representing the function.

If this macro is not defined, then the function size is not defined.

You may wish to use ASM\_OUTPUT\_MEASURED\_SIZE in the definition of this macro.

## ASM\_DECLARE\_COLD\_FUNCTION\_NAME (stream, name, decl)

Macro

A C statement (sans semicolon) to output to the stdio stream stream any text necessary for declaring the name name of a cold function partition which is being defined. This macro is responsible for outputting the label definition (perhaps using ASM\_OUTPUT\_FUNCTION\_LABEL). The argument decl is the FUNCTION\_DECL tree node representing the function.

If this macro is not defined, then the cold partition name is defined in the usual manner as a label (by means of ASM\_OUTPUT\_LABEL).

You may wish to use ASM\_OUTPUT\_TYPE\_DIRECTIVE in the definition of this macro.

## ASM\_DECLARE\_COLD\_FUNCTION\_SIZE (stream, name, decl)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream any text necessary for declaring the size of a cold function partition which is being defined. The argument name is the name of the cold partition of the function. The argument decl is the FUNCTION\_DECL tree node representing the function.

If this macro is not defined, then the partition size is not defined.

You may wish to use ASM\_OUTPUT\_MEASURED\_SIZE in the definition of this macro.

#### ASM\_DECLARE\_OBJECT\_NAME (stream, name, decl)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream any text necessary for declaring the name name of an initialized variable which is being defined.

This macro must output the label definition (perhaps using ASM\_OUTPUT\_LABEL). The argument decl is the VAR\_DECL tree node representing the variable.

If this macro is not defined, then the variable name is defined in the usual manner as a label (by means of ASM\_OUTPUT\_LABEL).

You may wish to use ASM\_OUTPUT\_TYPE\_DIRECTIVE and/or ASM\_OUTPUT\_SIZE\_DIRECTIVE in the definition of this macro.

# void TARGET\_ASM\_DECLARE\_CONSTANT\_NAME (FILE \*file, const char \*name, const\_tree expr, HOST\_WIDE\_INT size) [Target Hook]

A target hook to output to the stdio stream file any text necessary for declaring the name name of a constant which is being defined. This target hook is responsible for outputting the label definition (perhaps using assemble\_label). The argument exp is the value of the constant, and size is the size of the constant in bytes. The name will be an internal label.

The default version of this target hook, define the *name* in the usual manner as a label (by means of assemble\_label).

You may wish to use ASM\_OUTPUT\_TYPE\_DIRECTIVE in this target hook.

## ASM\_DECLARE\_REGISTER\_GLOBAL (stream, decl, regno, name)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream any text necessary for claiming a register regno for a global variable decl with name name.

If you don't define this macro, that is equivalent to defining it to do nothing.

# ASM\_FINISH\_DECLARE\_OBJECT (stream, decl, toplevel, atend) [Macro]

A C statement (sans semicolon) to finish up declaring a variable name once the compiler has processed its initializer fully and thus has had a chance to determine the size of an array when controlled by an initializer. This is used on systems where it's necessary to declare something about the size of the object.

If you don't define this macro, that is equivalent to defining it to do nothing.

You may wish to use ASM\_OUTPUT\_SIZE\_DIRECTIVE and/or ASM\_OUTPUT\_MEASURED\_SIZE in the definition of this macro.

## 

This target hook is a function to output to the stdio stream stream some commands that will make the label name global; that is, available for reference from other files.

The default implementation relies on a proper definition of GLOBAL\_ASM\_OP.

#### 

This target hook is a function to output to the stdio stream stream some commands that will make the name associated with decl global; that is, available for reference from other files.

The default implementation uses the TARGET\_ASM\_GLOBALIZE\_LABEL target hook.

#### 

This target hook is a function to output to the stdio stream stream some commands that will declare the name associated with decl which is not defined in the current translation unit. Most assemblers do not require anything to be output in this case.

#### ASM\_WEAKEN\_LABEL (stream, name)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream some commands that will make the label name weak; that is, available for reference from other files but only used if no other definition is available. Use the expression assemble\_name (stream, name) to output the name itself; before and after that, output the additional assembler syntax for making that name weak, and a newline.

If you don't define this macro or ASM\_WEAKEN\_DECL, GCC will not support weak symbols and you should not define the SUPPORTS\_WEAK macro.

## ASM\_WEAKEN\_DECL (stream, decl, name, value)

[Macro]

Combines (and replaces) the function of ASM\_WEAKEN\_LABEL and ASM\_OUTPUT\_WEAK\_ALIAS, allowing access to the associated function or variable decl. If value is not NULL, this C statement should output to the stdio stream stream assembler code which defines (equates) the weak symbol name to have the value value. If value is NULL, it should output commands to make name weak.

## ASM\_OUTPUT\_WEAKREF (stream, decl, name, value)

[Macro]

Outputs a directive that enables name to be used to refer to symbol value with weak-symbol semantics. decl is the declaration of name.

SUPPORTS\_WEAK [Macro]

A preprocessor constant expression which evaluates to true if the target supports weak symbols.

If you don't define this macro, 'defaults.h' provides a default definition. If either ASM\_WEAKEN\_LABEL or ASM\_WEAKEN\_DECL is defined, the default definition is '1'; otherwise, it is '0'.

## TARGET\_SUPPORTS\_WEAK

[Macro]

A C expression which evaluates to true if the target supports weak symbols.

If you don't define this macro, 'defaults.h' provides a default definition. The default definition is '(SUPPORTS\_WEAK)'. Define this macro if you want to control weak symbol support with a compiler flag such as '-melf'.

## MAKE\_DECL\_ONE\_ONLY (dec1)

[Macro]

A C statement (sans semicolon) to mark decl to be emitted as a public symbol such that extra copies in multiple translation units will be discarded by the linker. Define this macro if your object file format provides support for this concept, such as the 'COMDAT' section flags in the Microsoft Windows PE/COFF format, and this support requires changes to decl, such as putting it in a separate section.

#### SUPPORTS\_ONE\_ONLY

[Macro]

A C expression which evaluates to true if the target supports one-only semantics.

If you don't define this macro, 'varasm.c' provides a default definition. If MAKE\_DECL\_ONE\_ONLY is defined, the default definition is '1'; otherwise, it is '0'. Define this macro if you want to control one-only symbol support with a compiler flag, or if setting the DECL\_ONE\_ONLY flag is enough to mark a declaration to be emitted as one-only.

# 

[Target Hook]

This target hook is a function to output to  $asm_out_file$  some commands that will make the symbol(s) associated with decl have hidden, protected or internal visibility as specified by visibility.

#### TARGET\_WEAK\_NOT\_IN\_ARCHIVE\_TOC

[Macro]

A C expression that evaluates to true if the target's linker expects that weak symbols do not appear in a static archive's table of contents. The default is 0.

Leaving weak symbols out of an archive's table of contents means that, if a symbol will only have a definition in one translation unit and will have undefined references from other translation units, that symbol should not be weak. Defining this macro to be nonzero will thus have the effect that certain symbols that would normally be weak (explicit template instantiations, and vtables for polymorphic classes with noninline key methods) will instead be nonweak.

The C++ ABI requires this macro to be zero. Define this macro for targets where full C++ ABI compliance is impossible and where linker restrictions require weak symbols to be left out of a static archive's table of contents.

#### ASM\_OUTPUT\_EXTERNAL (stream, decl, name)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream any text necessary for declaring the name of an external symbol named name which is referenced in this compilation but not defined. The value of decl is the tree node for the declaration.

This macro need not be defined if it does not need to output anything. The GNU assembler and most Unix assemblers don't require anything.

### void TARGET\_ASM\_EXTERNAL\_LIBCALL (rtx symref)

[Target Hook]

This target hook is a function to output to  $asm_out_file$  an assembler pseudo-op to declare a library function name external. The name of the library function is given by symref, which is a  $symbol_ref$ .

void TARGET\_ASM\_MARK\_DECL\_PRESERVED (const char \*symbol) [Target Hook]
This target hook is a function to output to asm\_out\_file an assembler directive to annotate symbol as used. The Darwin target uses the .no\_dead\_code\_strip directive.

#### ASM\_OUTPUT\_LABELREF (stream, name)

[Macro]

A C statement (sans semicolon) to output to the stdio stream stream a reference in assembler syntax to a label named name. This should add '\_' to the front of the name, if that is customary on your operating system, as it is in most Berkeley Unix systems. This macro is used in assemble\_name.

## tree TARGET\_MANGLE\_ASSEMBLER\_NAME (const char \*name)

[Target Hook]

Given a symbol *name*, perform same mangling as varasm.c's assemble\_name, but in memory rather than to a file stream, returning result as an IDENTIFIER\_NODE. Required for correct LTO symtabs. The default implementation calls the TARGET\_STRIP\_NAME\_ENCODING hook and then prepends the USER\_LABEL\_PREFIX, if any.

## ASM\_OUTPUT\_SYMBOL\_REF (stream, sym)

[Macro]

A C statement (sans semicolon) to output a reference to SYMBOL\_REF sym. If not defined, assemble\_name will be used to output the name of the symbol. This macro may be used to modify the way a symbol is referenced depending on information encoded by TARGET\_ENCODE\_SECTION\_INFO.

#### ASM\_OUTPUT\_LABEL\_REF (stream, buf)

[Macro]

A C statement (sans semicolon) to output a reference to buf, the result of ASM\_GENERATE\_INTERNAL\_LABEL. If not defined, assemble\_name will be used to output the name of the symbol. This macro is not used by output\_asm\_label, or the %1 specifier that calls it; the intention is that this macro should be set when it is necessary to output a label differently when its address is being taken.

#### 

A function to output to the stdio stream stream a label whose name is made from the string prefix and the number labelno.

It is absolutely essential that these labels be distinct from the labels used for user-level functions and variables. Otherwise, certain programs will have name conflicts with internal labels.

It is desirable to exclude internal labels from the symbol table of the object file. Most assemblers have a naming convention for labels that should be excluded; on many systems, the letter 'L' at the beginning of a label has this effect. You should find out what convention your system uses, and follow it.

The default version of this function utilizes ASM\_GENERATE\_INTERNAL\_LABEL.

### ASM\_OUTPUT\_DEBUG\_LABEL (stream, prefix, num)

[Macro]

A C statement to output to the stdio stream stream a debug info label whose name is made from the string prefix and the number num. This is useful for VLIW targets, where debug info labels may need to be treated differently than branch target labels. On some systems, branch target labels must be at the beginning of instruction bundles, but debug info labels can occur in the middle of instruction bundles.

If this macro is not defined, then (\*targetm.asm\_out.internal\_label) will be used.

#### ASM\_GENERATE\_INTERNAL\_LABEL (string, prefix, num)

[Macro]

A C statement to store into the string *string* a label whose name is made from the string *prefix* and the number *num*.

This string, when output subsequently by assemble\_name, should produce the output that (\*targetm.asm\_out.internal\_label) would produce with the same *prefix* and *num*.

If the string begins with '\*', then assemble\_name will output the rest of the string unchanged. It is often convenient for ASM\_GENERATE\_INTERNAL\_LABEL to use '\*' in

this way. If the string doesn't start with '\*', then ASM\_OUTPUT\_LABELREF gets to output the string, and may change it. (Of course, ASM\_OUTPUT\_LABELREF is also part of your machine description, so you should know what it does on your machine.)

## ASM\_FORMAT\_PRIVATE\_NAME (outvar, name, number)

[Macro]

A C expression to assign to *outvar* (which is a variable of type char \*) a newly allocated string made from the string name and the number number, with some suitable punctuation added. Use alloca to get space for the string.

The string will be used as an argument to ASM\_OUTPUT\_LABELREF to produce an assembler label for an internal static variable whose name is *name*. Therefore, the string must be such as to result in valid assembler code. The argument *number* is different each time this macro is executed; it prevents conflicts between similarly-named internal static variables in different scopes.

Ideally this string should not be a valid C identifier, to prevent any conflict with the user's own symbols. Most assemblers allow periods or percent signs in assembler symbols; putting at least one of these between the name and the number will suffice.

If this macro is not defined, a default definition will be provided which is correct for most systems.

#### ASM\_OUTPUT\_DEF (stream, name, value)

[Macro]

A C statement to output to the stdio stream stream assembler code which defines (equates) the symbol name to have the value value.

If SET\_ASM\_OP is defined, a default definition is provided which is correct for most systems.

# ASM\_OUTPUT\_DEF\_FROM\_DECLS (stream, decl\_of\_name, decl\_of\_value)

[Macro]

A C statement to output to the stdio stream stream assembler code which defines (equates) the symbol whose tree node is  $decl\_of\_name$  to have the value of the tree node  $decl\_of\_value$ . This macro will be used in preference to 'ASM\_OUTPUT\_DEF' if it is defined and if the tree nodes are available.

If SET\_ASM\_OP is defined, a default definition is provided which is correct for most systems.

# TARGET\_DEFERRED\_OUTPUT\_DEFS (decl\_of\_name, decl\_of\_value) [Macro]

A C statement that evaluates to true if the assembler code which defines (equates) the symbol whose tree node is  $decl\_of\_name$  to have the value of the tree node  $decl\_of\_value$  should be emitted near the end of the current compilation unit. The default is to not defer output of defines. This macro affects defines output by 'ASM\_OUTPUT\_DEF' and 'ASM\_OUTPUT\_DEF\_FROM\_DECLS'.

#### ASM\_OUTPUT\_WEAK\_ALIAS (stream, name, value)

[Macro]

A C statement to output to the stdio stream stream assembler code which defines (equates) the weak symbol name to have the value value. If value is NULL, it defines name as an undefined weak symbol.

Define this macro if the target only supports weak aliases; define ASM\_OUTPUT\_DEF instead if possible.

# OBJC\_GEN\_METHOD\_LABEL (buf, is\_inst, class\_name, cat\_name, sel\_name) [Macro]

Define this macro to override the default assembler names used for Objective-C methods

The default name is a unique method number followed by the name of the class (e.g. '\_1\_Foo'). For methods in categories, the name of the category is also included in the assembler name (e.g. '\_1\_Foo\_Bar').

These names are safe on most systems, but make debugging difficult since the method's selector is not present in the name. Therefore, particular systems define other ways of computing names.

buf is an expression of type char \* which gives you a buffer in which to store the name; its length is as long as class\_name, cat\_name and sel\_name put together, plus 50 characters extra.

The argument *is\_inst* specifies whether the method is an instance method or a class method; *class\_name* is the name of the class; *cat\_name* is the name of the category (or NULL if the method is not in a category); and *sel\_name* is the name of the selector.

On systems where the assembler can handle quoted names, you can use this macro to provide more human-readable names.

### 18.20.5 How Initialization Functions Are Handled

The compiled code for certain languages includes *constructors* (also called *initialization routines*)—functions to initialize data in the program when the program is started. These functions need to be called before the program is "started"—that is to say, before main is called.

Compiling some languages generates destructors (also called termination routines) that should be called when the program terminates.

To make the initialization and termination functions work, the compiler must output something in the assembler code to cause those functions to be called at the appropriate time. When you port the compiler to a new system, you need to specify how to do this.

There are two major ways that GCC currently supports the execution of initialization and termination functions. Each way has two variants. Much of the structure is common to all four variations.

The linker must build two lists of these functions—a list of initialization functions, called \_\_CTOR\_LIST\_\_, and a list of termination functions, called \_\_DTOR\_LIST\_\_.

Each list always begins with an ignored function pointer (which may hold 0, -1, or a count of the function pointers after it, depending on the environment). This is followed by a series of zero or more function pointers to constructors (or destructors), followed by a function pointer containing zero.

Depending on the operating system and its executable file format, either 'crtstuff.c' or 'libgcc2.c' traverses these lists at startup time and exit time. Constructors are called in reverse order of the list; destructors in forward order.

The best way to handle static constructors works only for object file formats which provide arbitrarily-named sections. A section is set aside for a list of constructors, and another for a list of destructors. Traditionally these are called '.ctors' and '.dtors'. Each object file

that defines an initialization function also puts a word in the constructor section to point to that function. The linker accumulates all these words into one contiguous '.ctors' section. Termination functions are handled similarly.

This method will be chosen as the default by 'target-def.h' if TARGET\_ASM\_NAMED\_SECTION is defined. A target that does not support arbitrary sections, but does support special designated constructor and destructor sections may define CTORS\_SECTION\_ASM\_OP and DTORS\_SECTION\_ASM\_OP to achieve the same effect.

When arbitrary sections are available, there are two variants, depending upon how the code in 'crtstuff.c' is called. On systems that support a .init section which is executed at program startup, parts of 'crtstuff.c' are compiled into that section. The program is linked by the gcc driver like this:

```
ld -o output_file crti.o crtbegin.o ... -lgcc crtend.o crtn.o
```

The prologue of a function (\_\_init) appears in the .init section of 'crti.o'; the epilogue appears in 'crtn.o'. Likewise for the function \_\_fini in the .fini section. Normally these files are provided by the operating system or by the GNU C library, but are provided by GCC for a few targets.

The objects 'crtbegin.o' and 'crtend.o' are (for most targets) compiled from 'crtstuff.c'. They contain, among other things, code fragments within the .init and .fini sections that branch to routines in the .text section. The linker will pull all parts of a section together, which results in a complete \_\_init function that invokes the routines we need at startup.

To use this variant, you must define the INIT\_SECTION\_ASM\_OP macro properly.

If no init section is available, when GCC compiles any function called main (or more accurately, any function designated as a program entry point by the language front end calling expand\_main\_function), it inserts a procedure call to \_\_main as the first executable code after the function prologue. The \_\_main function is defined in 'libgcc2.c' and runs the global constructors.

In file formats that don't support arbitrary sections, there are again two variants. In the simplest variant, the GNU linker (GNU 1d) and an 'a.out' format must be used. In this case, TARGET\_ASM\_CONSTRUCTOR is defined to produce a .stabs entry of type 'N\_SETT', referencing the name \_\_CTOR\_LIST\_\_, and with the address of the void function containing the initialization code as its value. The GNU linker recognizes this as a request to add the value to a set; the values are accumulated, and are eventually placed in the executable as a vector in the format described above, with a leading (ignored) count and a trailing zero element. TARGET\_ASM\_DESTRUCTOR is handled similarly. Since no init section is available, the absence of INIT\_SECTION\_ASM\_OP causes the compilation of main to call \_\_main as above, starting the initialization process.

The last variant uses neither arbitrary sections nor the GNU linker. This is preferable when you want to do dynamic linking and when using file formats which the GNU linker does not support, such as 'ECOFF'. In this case, TARGET\_HAVE\_CTORS\_DTORS is false, initialization and termination functions are recognized simply by their names. This requires an extra program in the linkage step, called collect2. This program pretends to be the linker, for use with GCC; it does its job by running the ordinary linker, but also arranges to include the vectors of initialization and termination functions. These functions are called

via \_\_main as described above. In order to use this method, use\_collect2 must be defined in the target in 'config.gcc'.

## 18.20.6 Macros Controlling Initialization Routines

Here are the macros that control how the compiler handles initialization and termination functions:

## INIT\_SECTION\_ASM\_OP

[Macro]

If defined, a C string constant, including spacing, for the assembler operation to identify the following data as initialization code. If not defined, GCC will assume such a section does not exist. When you are using special sections for initialization and termination functions, this macro also controls how 'crtstuff.c' and 'libgcc2.c' arrange to run the initialization functions.

## HAS\_INIT\_SECTION

[Macro]

If defined, main will not call \_\_main as described above. This macro should be defined for systems that control start-up code on a symbol-by-symbol basis, such as OSF/1, and should not be defined explicitly for systems that support INIT\_SECTION\_ASM\_OP.

#### LD\_INIT\_SWITCH

Macro

If defined, a C string constant for a switch that tells the linker that the following symbol is an initialization routine.

#### LD\_FINI\_SWITCH

[Macro]

If defined, a C string constant for a switch that tells the linker that the following symbol is a finalization routine.

## COLLECT\_SHARED\_INIT\_FUNC (stream, func)

[Macro]

If defined, a C statement that will write a function that can be automatically called when a shared library is loaded. The function should call *func*, which takes no arguments. If not defined, and the object format requires an explicit initialization function, then a function called <code>\_GLOBAL\_\_DI</code> will be generated.

This function and the following one are used by collect when linking a shared library that needs constructors or destructors, or has DWARF2 exception tables embedded in the code.

### COLLECT\_SHARED\_FINI\_FUNC (stream, func)

[Macro]

If defined, a C statement that will write a function that can be automatically called when a shared library is unloaded. The function should call *func*, which takes no arguments. If not defined, and the object format requires an explicit finalization function, then a function called <code>\_GLOBAL\_\_DD</code> will be generated.

## INVOKE\_\_main

[Macro]

If defined, main will call \_\_main despite the presence of INIT\_SECTION\_ASM\_OP. This macro should be defined for systems where the init section is not actually run automatically, but is still useful for collecting the lists of constructors and destructors.

## SUPPORTS\_INIT\_PRIORITY

|Macro

If nonzero, the C++ init\_priority attribute is supported and the compiler should emit instructions to control the order of initialization of objects. If zero, the compiler will issue an error message upon encountering an init\_priority attribute.

### bool TARGET\_HAVE\_CTORS\_DTORS

[Target Hook]

This value is true if the target supports some "native" method of collecting constructors and destructors to be run at startup and exit. It is false if we must use collect2.

## void TARGET\_ASM\_CONSTRUCTOR (rtx symbol, int priority)

[Target Hook]

If defined, a function that outputs assembler code to arrange to call the function referenced by *symbol* at initialization time.

Assume that *symbol* is a SYMBOL\_REF for a function taking no arguments and with no return value. If the target supports initialization priorities, *priority* is a value between 0 and MAX\_INIT\_PRIORITY; otherwise it must be DEFAULT\_INIT\_PRIORITY.

If this macro is not defined by the target, a suitable default will be chosen if (1) the target supports arbitrary section names, (2) the target defines CTORS\_SECTION\_ASM\_OP, or (3) USE\_COLLECT2 is not defined.

# void TARGET\_ASM\_DESTRUCTOR (rtx symbol, int priority)

[Target Hook]

This is like TARGET\_ASM\_CONSTRUCTOR but used for termination functions rather than initialization functions.

If TARGET\_HAVE\_CTORS\_DTORS is true, the initialization routine generated for the generated object file will have static linkage.

If your system uses collect2 as the means of processing constructors, then that program normally uses nm to scan an object file for constructor functions to be called.

On certain kinds of systems, you can define this macro to make collect2 work faster (and, in some cases, make it work at all):

#### OBJECT\_FORMAT\_COFF

[Macro]

Define this macro if the system uses COFF (Common Object File Format) object files, so that collect2 can assume this format and scan object files directly for dynamic constructor/destructor functions.

This macro is effective only in a native compiler; collect2 as part of a cross compiler always uses nm for the target machine.

#### REAL\_NM\_FILE\_NAME

[Macro]

Define this macro as a C string constant containing the file name to use to execute nm. The default is to search the path normally for nm.

### NM\_FLAGS

[Macro]

collect2 calls nm to scan object files for static constructors and destructors and LTO info. By default, '-n' is passed. Define NM\_FLAGS to a C string constant if other options are needed to get the same output format as GNU nm -n produces.

If your system supports shared libraries and has a program to list the dynamic dependencies of a given library or executable, you can define these macros to enable support for running initialization and termination functions in shared libraries:

#### LDD\_SUFFIX

[Macro]

Define this macro to a C string constant containing the name of the program which lists dynamic dependencies, like 1dd under SunOS 4.

## PARSE\_LDD\_OUTPUT (ptr)

[Macro]

Define this macro to be C code that extracts filenames from the output of the program denoted by LDD\_SUFFIX. ptr is a variable of type char \* that points to the beginning of a line of output from LDD\_SUFFIX. If the line lists a dynamic dependency, the code must advance ptr to the beginning of the filename on that line. Otherwise, it must set ptr to NULL.

SHLIB\_SUFFIX [Macro]

Define this macro to a C string constant containing the default shared library extension of the target (e.g., '".so"'). collect2 strips version information after this suffix when generating global constructor and destructor names. This define is only needed on targets that use collect2 to process constructors and destructors.

## 18.20.7 Output of Assembler Instructions

This describes assembler instruction output.

REGISTER\_NAMES [Macro]

A C initializer containing the assembler's names for the machine registers, each one as a C string constant. This is what translates register numbers in the compiler into assembler language.

### ADDITIONAL\_REGISTER\_NAMES

Macro

If defined, a C initializer for an array of structures containing a name and a register number. This macro defines additional names for hard registers, thus allowing the asm option in declarations to refer to registers using alternate names.

#### OVERLAPPING\_REGISTER\_NAMES

Macro

If defined, a C initializer for an array of structures containing a name, a register number and a count of the number of consecutive machine registers the name overlaps. This macro defines additional names for hard registers, thus allowing the asm option in declarations to refer to registers using alternate names. Unlike ADDITIONAL\_REGISTER\_NAMES, this macro should be used when the register name implies multiple underlying registers.

This macro should be used when it is important that a clobber in an asm statement clobbers all the underlying values implied by the register name. For example, on ARM, clobbering the double-precision VFP register "d0" implies clobbering both single-precision registers "s0" and "s1".

#### ASM\_OUTPUT\_OPCODE (stream, ptr)

[Macro]

Define this macro if you are using an unusual assembler that requires different names for the machine instructions.

The definition is a C statement or statements which output an assembler instruction opcode to the stdio stream stream. The macro-operand ptr is a variable of type char \* which points to the opcode name in its "internal" form—the form that is written in the machine description. The definition should output the opcode name to stream, performing any translation you desire, and increment the variable ptr to point at the end of the opcode so that it will not be output twice.

In fact, your macro definition may process less than the entire opcode name, or more than the opcode name; but if you want to process text that includes '%'-sequences to substitute operands, you must take care of the substitution yourself. Just be sure to increment *ptr* over whatever text should not be output normally.

If you need to look at the operand values, they can be found as the elements of recog\_data.operand.

If the macro definition does nothing, the instruction is output in the usual way.

# FINAL\_PRESCAN\_INSN (insn, opvec, noperands)

[Macro]

If defined, a C statement to be executed just prior to the output of assembler code for *insn*, to modify the extracted operands so they will be output differently.

Here the argument opvec is the vector containing the operands extracted from insn, and noperands is the number of elements of the vector which contain meaningful data for this insn. The contents of this vector are what will be used to convert the insn template into assembler code, so you can change the assembler output by changing the contents of the vector.

This macro is useful when various assembler syntaxes share a single file of instruction patterns; by defining this macro differently, you can cause a large class of instructions to be output differently (such as with rearranged operands). Naturally, variations in assembler syntax affecting individual insn patterns ought to be handled by writing conditional output routines in those patterns.

If this macro is not defined, it is equivalent to a null statement.

#### 

If defined, this target hook is a function which is executed just after the output of assembler code for *insn*, to change the mode of the assembler if necessary.

Here the argument *opvec* is the vector containing the operands extracted from *insn*, and *noperands* is the number of elements of the vector which contain meaningful data for this insn. The contents of this vector are what was used to convert the insn template into assembler code, so you can change the assembler mode by checking the contents of the vector.

#### PRINT\_OPERAND (stream, x, code)

[Macro]

A C compound statement to output to stdio stream stream the assembler syntax for an instruction operand x. x is an RTL expression.

code is a value that can be used to specify one of several ways of printing the operand. It is used when identical operands must be printed differently depending on the context. code comes from the '%' specification that was used to request printing of the operand. If the specification was just '%digit' then code is 0; if the specification was '%ltr digit' then code is the ASCII code for ltr.

If x is a register, this macro should print the register's name. The names can be found in an array reg\_names whose type is char \*[]. reg\_names is initialized from REGISTER NAMES.

When the machine description has a specification '%punct' (a '%' followed by a punctuation character), this macro is called with a null pointer for x and the punctuation character for code.

## PRINT\_OPERAND\_PUNCT\_VALID\_P (code)

[Macro]

A C expression which evaluates to true if *code* is a valid punctuation character for use in the PRINT\_OPERAND macro. If PRINT\_OPERAND\_PUNCT\_VALID\_P is not defined, it means that no punctuation characters (except for the standard one, '%') are used in this way.

## PRINT\_OPERAND\_ADDRESS (stream, x)

[Macro]

A C compound statement to output to stdio stream stream the assembler syntax for an instruction operand that is a memory reference whose address is x. x is an RTL expression.

On some machines, the syntax for a symbolic address depends on the section that the address refers to. On these machines, define the hook TARGET\_ENCODE\_SECTION\_INFO to store the information into the symbol\_ref, and then check for it here. See Section 18.20 [Assembler Format], page 596.

## DBR\_OUTPUT\_SEQEND (file)

[Macro]

A C statement, to be executed after all slot-filler instructions have been output. If necessary, call dbr\_sequence\_length to determine the number of slots filled in a sequence (zero if not currently outputting a sequence), to decide how many no-ops to output, or whatever.

Don't define this macro if it has nothing to do, but it is helpful in reading assembly output if the extent of the delay sequence is made explicit (e.g. with white space).

Note that output routines for instructions with delay slots must be prepared to deal with not being output as part of a sequence (i.e. when the scheduling pass is not run, or when no slot fillers could be found.) The variable final\_sequence is null when not processing a sequence, otherwise it contains the sequence rtx being output.

REGISTER\_PREFIX [Macro]
LOCAL\_LABEL\_PREFIX [Macro]
USER\_LABEL\_PREFIX [Macro]
IMMEDIATE\_PREFIX [Macro]

If defined, C string expressions to be used for the '%R', '%L', '%U', and '%I' options of asm\_fprintf (see 'final.c'). These are useful when a single 'md' file must support multiple assembler formats. In that case, the various 'tm.h' files can define these macros differently.

## ASM\_FPRINTF\_EXTENSIONS (file, argptr, format)

[Macro]

If defined this macro should expand to a series of case statements which will be parsed inside the switch statement of the asm\_fprintf function. This allows targets to define extra printf formats which may useful when generating their assembler statements. Note that uppercase letters are reserved for future generic extensions to asm\_fprintf, and so are not available to target specific code. The output file is given by the parameter file. The varargs input pointer is argptr and the rest of the format string, starting the character after the one that is being switched upon, is pointed to by format.

## ASSEMBLER\_DIALECT

[Macro]

If your target supports multiple dialects of assembler language (such as different opcodes), define this macro as a C expression that gives the numeric index of the assembler language dialect to use, with zero as the first variant.

If this macro is defined, you may use constructs of the form

```
'{option0|option1|option2...}'
```

in the output templates of patterns (see Section 17.5 [Output Template], page 343) or in the first argument of asm\_fprintf. This construct outputs 'option0', 'option1', 'option2', etc., if the value of ASSEMBLER\_DIALECT is zero, one, two, etc. Any special characters within these strings retain their usual meaning. If there are fewer alternatives within the braces than the value of ASSEMBLER\_DIALECT, the construct outputs nothing. If it's needed to print curly braces or '|' character in assembler output directly, '%{', '%}' and '%|' can be used.

If you do not define this macro, the characters '{', '|' and '}' do not have any special meaning when used in templates or operands to asm\_fprintf.

Define the macros REGISTER\_PREFIX, LOCAL\_LABEL\_PREFIX, USER\_LABEL\_PREFIX and IMMEDIATE\_PREFIX if you can express the variations in assembler language syntax with that mechanism. Define ASSEMBLER\_DIALECT and use the '{option0|option1}' syntax if the syntax variant are larger and involve such things as different opcodes or operand order.

## ASM\_OUTPUT\_REG\_PUSH (stream, regno)

[Macro]

A C expression to output to *stream* some assembler code which will push hard register number *regno* onto the stack. The code need not be optimal, since this macro is used only when profiling.

## ASM\_OUTPUT\_REG\_POP (stream, regno)

[Macro]

A C expression to output to *stream* some assembler code which will pop hard register number *regno* off of the stack. The code need not be optimal, since this macro is used only when profiling.

# 18.20.8 Output of Dispatch Tables

This concerns dispatch tables.

## ASM\_OUTPUT\_ADDR\_DIFF\_ELT (stream, body, value, rel)

[Macro]

A C statement to output to the stdio stream stream an assembler pseudo-instruction to generate a difference between two labels. value and rel are the numbers of two internal labels. The definitions of these labels are output using (\*targetm.asm\_out.internal\_label), and they must be printed in the same way here. For example,

You must provide this macro on machines where the addresses in a dispatch table are relative to the table's own address. If defined, GCC will also use this macro on all machines when producing PIC. body is the body of the ADDR\_DIFF\_VEC; it is provided so that the mode and flags can be read.

## ASM\_OUTPUT\_ADDR\_VEC\_ELT (stream, value)

[Macro]

This macro should be provided on machines where the addresses in a dispatch table are absolute.

The definition should be a C statement to output to the stdio stream stream an assembler pseudo-instruction to generate a reference to a label. value is the number of an internal label whose definition is output using (\*targetm.asm\_out.internal\_label). For example,

fprintf (stream, "\t.word L%d\n", value)

## ASM\_OUTPUT\_CASE\_LABEL (stream, prefix, num, table)

[Macro]

Define this if the label before a jump-table needs to be output specially. The first three arguments are the same as for (\*targetm.asm\_out.internal\_label); the fourth argument is the jump-table which follows (a jump\_table\_data containing an addr\_vec or addr\_diff\_vec).

This feature is used on system V to output a swbeg statement for the table.

If this macro is not defined, these labels are output with (\*targetm.asm\_out.internal\_label).

## ASM\_OUTPUT\_CASE\_END (stream, num, table)

[Macro]

Define this if something special must be output at the end of a jump-table. The definition should be a C statement to be executed after the assembler code for the table is written. It should write the appropriate code to stdio stream *stream*. The argument *table* is the jump-table insn, and *num* is the label-number of the preceding label.

If this macro is not defined, nothing special is output at the end of the jump-table.

# void TARGET\_ASM\_POST\_CFI\_STARTPROC (FILE \*, tree)

|Target Hook

This target hook is used to emit assembly strings required by the target after the .cfi\_startproc directive. The first argument is the file stream to write the strings to and the second argument is the function's declaration. The expected use is to add more .cfi\_\* directives.

The default is to not output any assembly strings.

#### 

This target hook emits a label at the beginning of each FDE. It should be defined on targets where FDEs need special labels, and it should write the appropriate label, for the FDE associated with the function declaration decl, to the stdio stream stream. The third argument, for\_eh, is a boolean: true if this is for an exception table. The fourth argument, empty, is a boolean: true if this is a placeholder label for an omitted FDE.

The default is that FDEs are not given nonlocal labels.

## void TARGET\_ASM\_EMIT\_EXCEPT\_TABLE\_LABEL (FILE \*stream) [Target Hook]

This target hook emits a label at the beginning of the exception table. It should be defined on targets where it is desirable for the table to be broken up according to function.

The default is that no label is emitted.

- void TARGET\_ASM\_EMIT\_EXCEPT\_PERSONALITY (rtx personality) [Target Hook] If the target implements TARGET\_ASM\_UNWIND\_EMIT, this hook may be used to emit a directive to install a personality hook into the unwind info. This hook should not be used if dwarf2 unwind info is used.
- void TARGET\_ASM\_UNWIND\_EMIT (FILE \*stream, rtx\_insn \*insn) [Target Hook]
  This target hook emits assembly directives required to unwind the given instruction.
  This is only used when TARGET\_EXCEPT\_UNWIND\_INFO returns UI\_TARGET.

## bool TARGET\_ASM\_UNWIND\_EMIT\_BEFORE\_INSN

[Target Hook]

True if the TARGET\_ASM\_UNWIND\_EMIT hook should be called before the assembly for insn has been emitted, false if the hook should be called afterward.

## bool TARGET\_ASM\_SHOULD\_RESTORE\_CFA\_STATE (void)

[Target Hook]

For DWARF-based unwind frames, two CFI instructions provide for save and restore of register state. GCC maintains the current frame address (CFA) separately from the register bank but the unwinder in libgcc preserves this state along with the registers (and this is expected by the code that writes the unwind frames). This hook allows the target to specify that the CFA data is not saved/restored along with the registers by the target unwinder so that suitable additional instructions should be emitted to restore it.

## 18.20.9 Assembler Commands for Exception Regions

This describes commands marking the start and the end of an exception region.

#### EH\_FRAME\_SECTION\_NAME

[Macro]

If defined, a C string constant for the name of the section containing exception handling frame unwind information. If not defined, GCC will provide a default definition if the target supports named sections. 'crtstuff.c' uses this macro to switch to the appropriate section.

You should define this symbol if your target supports DWARF 2 frame unwind information and the default definition does not work.

## EH\_FRAME\_THROUGH\_COLLECT2

[Macro]

If defined, DWARF 2 frame unwind information will identified by specially named labels. The collect2 process will locate these labels and generate code to register the frames.

This might be necessary, for instance, if the system linker will not place the eh\_frames in-between the sentinals from 'crtstuff.c', or if the system linker does garbage collection and sections cannot be marked as not to be collected.

## EH\_TABLES\_CAN\_BE\_READ\_ONLY

[Macro]

Define this macro to 1 if your target is such that no frame unwind information encoding used with non-PIC code will ever require a runtime relocation, but the linker may not support merging read-only and read-write sections into a single read-write section.

#### MASK\_RETURN\_ADDR [Macro]

An rtx used to mask the return address found via RETURN\_ADDR\_RTX, so that it does not contain any extraneous set bits in it.

#### DWARF2\_UNWIND\_INFO

[Macro]

Define this macro to 0 if your target supports DWARF 2 frame unwind information, but it does not yet work with exception handling. Otherwise, if your target supports this information (if it defines INCOMING\_RETURN\_ADDR\_RTX and OBJECT\_FORMAT\_ELF), GCC will provide a default definition of 1.

#### 

This hook defines the mechanism that will be used for exception handling by the target. If the target has ABI specified unwind tables, the hook should return UI\_TARGET. If the target is to use the setjmp/longjmp-based exception handling scheme, the hook should return UI\_SJLJ. If the target supports DWARF 2 frame unwind information, the hook should return UI\_DWARF2.

A target may, if exceptions are disabled, choose to return UI\_NONE. This may end up simplifying other parts of target-specific code. The default implementation of this hook never returns UI\_NONE.

Note that the value returned by this hook should be constant. It should not depend on anything except the command-line switches described by *opts*. In particular, the setting UI\_SJLJ must be fixed at compiler start-up as C pre-processor macros and builtin functions related to exception handling are set up depending on this setting.

The default implementation of the hook first honors the '--enable-sjlj-exceptions' configure option, then DWARF2\_UNWIND\_INFO, and finally defaults to UI\_SJLJ. If DWARF2\_UNWIND\_INFO depends on command-line options, the target must define this hook so that *opts* is used correctly.

### bool TARGET\_UNWIND\_TABLES\_DEFAULT

[Common Target Hook]

This variable should be set to **true** if the target ABI requires unwinding tables even when exceptions are not used. It must not be modified by command-line option processing.

## DONT\_USE\_BUILTIN\_SETJMP

|Macro

Define this macro to 1 if the setjmp/longjmp-based scheme should use the setjmp/longjmp functions from the C library instead of the \_\_builtin\_setjmp/\_\_builtin\_longjmp machinery.

JMP\_BUF\_SIZE [Macro]

This macro has no effect unless DONT\_USE\_BUILTIN\_SETJMP is also defined. Define this macro if the default size of jmp\_buf buffer for the setjmp/longjmp-based exception handling mechanism is not large enough, or if it is much too large. The default size is FIRST\_PSEUDO\_REGISTER \* sizeof(void \*).

#### DWARF\_CIE\_DATA\_ALIGNMENT

[Macro]

This macro need only be defined if the target might save registers in the function prologue at an offset to the stack pointer that is not aligned to UNITS\_PER\_WORD. The definition should be the negative minimum alignment if STACK\_GROWS\_DOWNWARD is true, and the positive minimum alignment otherwise. See Section 18.21.5 [DWARF], page 629. Only applicable if the target supports DWARF 2 frame unwind information.

## bool TARGET\_TERMINATE\_DW2\_EH\_FRAME\_INFO

[Target Hook]

Contains the value true if the target should add a zero word onto the end of a Dwarf-2 frame info section when used for exception handling. Default value is false if EH\_FRAME\_SECTION\_NAME is defined, and true otherwise.

## rtx TARGET\_DWARF\_REGISTER\_SPAN (rtx reg)

[Target Hook]

Given a register, this hook should return a parallel of registers to represent where to find the register pieces. Define this hook if the register and its mode are represented in Dwarf in non-contiguous locations, or if the register should be represented in more than one register in Dwarf. Otherwise, this hook should return NULL\_RTX. If not defined, the default is to return NULL\_RTX.

#### machine\_mode TARGET\_DWARF\_FRAME\_REG\_MODE (int regno)

[Target Hook]

Given a register, this hook should return the mode which the corresponding Dwarf frame register should have. This is normally used to return a smaller mode than the raw mode to prevent call clobbered parts of a register altering the frame register size

## void TARGET\_INIT\_DWARF\_REG\_SIZES\_EXTRA (tree address)

Target Hook

If some registers are represented in Dwarf-2 unwind information in multiple pieces, define this hook to fill in information about the sizes of those pieces in the table used by the unwinder at runtime. It will be called by expand\_builtin\_init\_dwarf\_reg\_sizes after filling in a single size corresponding to each hard register; address is the address of the table.

#### bool TARGET\_ASM\_TTYPE (rtx sym)

[Target Hook]

This hook is used to output a reference from a frame unwinding table to the type\_info object identified by sym. It should return true if the reference was output. Returning false will cause the reference to be output using the normal Dwarf2 routines.

#### bool TARGET\_ARM\_EABI\_UNWINDER

[Target Hook]

This flag should be set to **true** on targets that use an ARM EABI based unwinding library, and **false** on other targets. This effects the format of unwinding tables, and how the unwinder in entered after running a cleanup. The default is **false**.

## 18.20.10 Assembler Commands for Alignment

This describes commands for alignment.

#### JUMP\_ALIGN (label)

[Macro]

The alignment (log base 2) to put in front of *label*, which is a common destination of jumps and has no fallthru incoming edge.

This macro need not be defined if you don't want any special alignment to be done at such a time. Most machine descriptions do not currently define the macro.

Unless it's necessary to inspect the *label* parameter, it is better to set the variable *align\_jumps* in the target's TARGET\_OPTION\_OVERRIDE. Otherwise, you should try to honor the user's selection in *align\_jumps* in a JUMP\_ALIGN implementation.

#### LABEL\_ALIGN\_AFTER\_BARRIER (label)

[Macro]

The alignment (log base 2) to put in front of label, which follows a BARRIER.

This macro need not be defined if you don't want any special alignment to be done at such a time. Most machine descriptions do not currently define the macro.

## LOOP\_ALIGN (label)

[Macro]

The alignment (log base 2) to put in front of *label* that heads a frequently executed basic block (usually the header of a loop).

This macro need not be defined if you don't want any special alignment to be done at such a time. Most machine descriptions do not currently define the macro.

Unless it's necessary to inspect the *label* parameter, it is better to set the variable align\_loops in the target's TARGET\_OPTION\_OVERRIDE. Otherwise, you should try to honor the user's selection in align\_loops in a LOOP\_ALIGN implementation.

# LABEL\_ALIGN (label)

[Macro]

The alignment (log base 2) to put in front of *label*. If LABEL\_ALIGN\_AFTER\_BARRIER / LOOP\_ALIGN specify a different alignment, the maximum of the specified values is used.

Unless it's necessary to inspect the *label* parameter, it is better to set the variable align\_labels in the target's TARGET\_OPTION\_OVERRIDE. Otherwise, you should try to honor the user's selection in align\_labels in a LABEL\_ALIGN implementation.

## ASM\_OUTPUT\_SKIP (stream, nbytes)

[Macro]

A C statement to output to the stdio stream stream an assembler instruction to advance the location counter by *nbytes* bytes. Those bytes should be zero when loaded. *nbytes* will be a C expression of type unsigned HOST\_WIDE\_INT.

### ASM\_NO\_SKIP\_IN\_TEXT

[Macro]

Define this macro if ASM\_OUTPUT\_SKIP should not be used in the text section because it fails to put zeros in the bytes that are skipped. This is true on many Unix systems, where the pseudo—op to skip bytes produces no-op instructions rather than zeros when used in the text section.

### ASM\_OUTPUT\_ALIGN (stream, power)

[Macro]

A C statement to output to the stdio stream an assembler command to advance the location counter to a multiple of 2 to the *power* bytes. *power* will be a C expression of type int.

### ASM\_OUTPUT\_ALIGN\_WITH\_NOP (stream, power)

[Macro]

Like ASM\_OUTPUT\_ALIGN, except that the "nop" instruction is used for padding, if necessary.

### ASM\_OUTPUT\_MAX\_SKIP\_ALIGN (stream, power, max\_skip)

[Macro]

A C statement to output to the stdio stream an assembler command to advance the location counter to a multiple of 2 to the *power* bytes, but only if *max\_skip* or fewer bytes are needed to satisfy the alignment request. *power* and *max\_skip* will be a C expression of type int.

# 18.21 Controlling Debugging Information Format

This describes how to specify debugging information.

# 18.21.1 Macros Affecting All Debugging Formats

These macros affect all debugging formats.

## DBX\_REGISTER\_NUMBER (regno)

[Macro]

A C expression that returns the DBX register number for the compiler register number regno. In the default macro provided, the value of this expression will be regno itself. But sometimes there are some registers that the compiler knows about and DBX does not, or vice versa. In such cases, some register may need to have one number in the compiler and another for DBX.

If two registers have consecutive numbers inside GCC, and they can be used as a pair to hold a multiword value, then they *must* have consecutive numbers after renumbering with DBX\_REGISTER\_NUMBER. Otherwise, debuggers will be unable to access such a pair, because they expect register pairs to be consecutive in their own numbering scheme.

If you find yourself defining DBX\_REGISTER\_NUMBER in way that does not preserve register pairs, then what you must do instead is redefine the actual register numbering scheme.

#### DEBUGGER\_AUTO\_OFFSET (x)

[Macro]

A C expression that returns the integer offset value for an automatic variable having address x (an RTL expression). The default computation assumes that x is based on the frame-pointer and gives the offset from the frame-pointer. This is required for targets that produce debugging output for DBX and allow the frame-pointer to be eliminated when the '-g' option is used.

## DEBUGGER\_ARG\_OFFSET (offset, x)

[Macro]

A C expression that returns the integer offset value for an argument having address x (an RTL expression). The nominal offset is offset.

#### PREFERRED\_DEBUGGING\_TYPE

[Macro]

A C expression that returns the type of debugging output GCC should produce when the user specifies just '-g'. Define this if you have arranged for GCC to support more than one format of debugging output. Currently, the allowable values are DBX\_DEBUG, DWARF2\_DEBUG, XCOFF\_DEBUG, VMS\_DEBUG, and VMS\_AND\_DWARF2\_DEBUG.

When the user specifies '-ggdb', GCC normally also uses the value of this macro to select the debugging output format, but with two exceptions. If DWARF2\_DEBUGGING\_INFO is defined, GCC uses the value DWARF2\_DEBUG. Otherwise, if DBX\_DEBUGGING\_INFO is defined, GCC uses DBX\_DEBUG.

The value of this macro only affects the default debugging output; the user can always get a specific type of output by using '-gstabs', '-gdwarf-2', '-gxcoff', or '-gvms'.

# 18.21.2 Specific Options for DBX Output

These are specific options for DBX output.

### DBX\_DEBUGGING\_INFO

[Macro]

Define this macro if GCC should produce debugging output for DBX in response to the '-g' option.

#### XCOFF\_DEBUGGING\_INFO

[Macro]

Define this macro if GCC should produce XCOFF format debugging output in response to the '-g' option. This is a variant of DBX format.

#### DEFAULT\_GDB\_EXTENSIONS

[Macro]

Define this macro to control whether GCC should by default generate GDB's extended version of DBX debugging information (assuming DBX-format debugging information is enabled at all). If you don't define the macro, the default is 1: always generate the extended information if there is any occasion to.

#### DEBUG\_SYMS\_TEXT

[Macro]

Define this macro if all .stabs commands should be output while in the text section.

#### ASM\_STABS\_OP

[Macro]

A C string constant, including spacing, naming the assembler pseudo op to use instead of "\t.stabs\t" to define an ordinary debugging symbol. If you don't define this macro, "\t.stabs\t" is used. This macro applies only to DBX debugging information format.

ASM\_STABD\_OP

[Macro]

A C string constant, including spacing, naming the assembler pseudo op to use instead of "\t.stabd\t" to define a debugging symbol whose value is the current location. If you don't define this macro, "\t.stabd\t" is used. This macro applies only to DBX debugging information format.

### ASM\_STABN\_OP

[Macro]

A C string constant, including spacing, naming the assembler pseudo op to use instead of "\t.stabn\t" to define a debugging symbol with no name. If you don't define this macro, "\t.stabn\t" is used. This macro applies only to DBX debugging information format.

DBX\_NO\_XREFS

[Macro]

Define this macro if DBX on your system does not support the construct 'xstagname'. On some systems, this construct is used to describe a forward reference to a structure named tagname. On other systems, this construct is not supported at all.

#### DBX\_CONTIN\_LENGTH

|Macro|

A symbol name in DBX-format debugging information is normally continued (split into two separate .stabs directives) when it exceeds a certain length (by default, 80 characters). On some operating systems, DBX requires this splitting; on others, splitting must not be done. You can inhibit splitting by defining this macro with the value zero. You can override the default splitting-length by defining this macro as an expression for the length you desire.

## DBX\_CONTIN\_CHAR [Macro]

Normally continuation is indicated by adding a '\' character to the end of a .stabs string when a continuation follows. To use a different character instead, define this macro as a character constant for the character you want to use. Do not define this macro if backslash is correct for your system.

#### DBX\_STATIC\_STAB\_DATA\_SECTION

[Macro]

Define this macro if it is necessary to go to the data section before outputting the '.stabs' pseudo-op for a non-global static variable.

### DBX\_TYPE\_DECL\_STABS\_CODE

[Macro]

The value to use in the "code" field of the .stabs directive for a typedef. The default is N\_LSYM.

### DBX\_STATIC\_CONST\_VAR\_CODE

[Macro]

The value to use in the "code" field of the .stabs directive for a static variable located in the text section. DBX format does not provide any "right" way to do this. The default is N\_FUN.

#### DBX\_REGPARM\_STABS\_CODE

[Macro]

The value to use in the "code" field of the .stabs directive for a parameter passed in registers. DBX format does not provide any "right" way to do this. The default is N\_RSYM.

#### DBX\_REGPARM\_STABS\_LETTER

Macro

The letter to use in DBX symbol data to identify a symbol as a parameter passed in registers. DBX format does not customarily provide any way to do this. The default is 'P'.

#### DBX\_FUNCTION\_FIRST

[Macro]

Define this macro if the DBX information for a function and its arguments should precede the assembler code for the function. Normally, in DBX format, the debugging information entirely follows the assembler code.

#### DBX\_BLOCKS\_FUNCTION\_RELATIVE

[Macro]

Define this macro, with value 1, if the value of a symbol describing the scope of a block (N\_LBRAC or N\_RBRAC) should be relative to the start of the enclosing function. Normally, GCC uses an absolute address.

#### DBX\_LINES\_FUNCTION\_RELATIVE

[Macro]

Define this macro, with value 1, if the value of a symbol indicating the current line number (N\_SLINE) should be relative to the start of the enclosing function. Normally, GCC uses an absolute address.

## DBX\_USE\_BINCL

[Macro]

Define this macro if GCC should generate N\_BINCL and N\_EINCL stabs for included header files, as on Sun systems. This macro also directs GCC to output a type number as a pair of a file number and a type number within the file. Normally, GCC does not generate N\_BINCL or N\_EINCL stabs, and it outputs a single number for a type number.

# 18.21.3 Open-Ended Hooks for DBX Format

These are hooks for DBX format.

## DBX\_OUTPUT\_SOURCE\_LINE (stream, line, counter)

[Macro]

A C statement to output DBX debugging information before code for line number *line* of the current source file to the stdio stream *stream*. *counter* is the number of time the macro was invoked, including the current invocation; it is intended to generate unique labels in the assembly output.

This macro should not be defined if the default output is correct, or if it can be made correct by defining DBX\_LINES\_FUNCTION\_RELATIVE.

## NO\_DBX\_FUNCTION\_END

[Macro]

Some stabs encapsulation formats (in particular ECOFF), cannot handle the .stabs "",N\_FUN,,0,0,Lscope-function-1 gdb dbx extension construct. On those machines, define this macro to turn this feature off without disturbing the rest of the gdb extensions.

#### NO\_DBX\_BNSYM\_ENSYM

[Macro]

Some assemblers cannot handle the .stabd BNSYM/ENSYM,0,0 gdb dbx extension construct. On those machines, define this macro to turn this feature off without disturbing the rest of the gdb extensions.

## 18.21.4 File Names in DBX Format

This describes file names in DBX format.

## DBX\_OUTPUT\_MAIN\_SOURCE\_FILENAME (stream, name)

[Macro]

A C statement to output DBX debugging information to the stdio stream *stream*, which indicates that file *name* is the main source file—the file specified as the input file for compilation. This macro is called only once, at the beginning of compilation.

This macro need not be defined if the standard form of output for DBX debugging information is appropriate.

It may be necessary to refer to a label equal to the beginning of the text section. You can use 'assemble\_name (stream, ltext\_label\_name)' to do so. If you do this, you must also set the variable used\_ltext\_label\_name to true.

#### NO\_DBX\_MAIN\_SOURCE\_DIRECTORY

[Macro]

Define this macro, with value 1, if GCC should not emit an indication of the current directory for compilation and current source language at the beginning of the file.

### NO\_DBX\_GCC\_MARKER

[Macro]

Define this macro, with value 1, if GCC should not emit an indication that this object file was compiled by GCC. The default is to emit an N\_OPT stab at the beginning of every source file, with 'gcc2\_compiled.' for the string and value 0.

# DBX\_OUTPUT\_MAIN\_SOURCE\_FILE\_END (stream, name)

[Macro]

A C statement to output DBX debugging information at the end of compilation of the main source file name. Output should be written to the stdio stream stream.

If you don't define this macro, nothing special is output at the end of compilation, which is correct for most machines.

## DBX\_OUTPUT\_NULL\_N\_SO\_AT\_MAIN\_SOURCE\_FILE\_END

[Macro]

Define this macro *instead of* defining DBX\_OUTPUT\_MAIN\_SOURCE\_FILE\_END, if what needs to be output at the end of compilation is an N\_SO stab with an empty string, whose value is the highest absolute text address in the file.

## 18.21.5 Macros for DWARF Output

Here are macros for DWARF output.

#### DWARF2\_DEBUGGING\_INFO

[Macro]

Define this macro if GCC should produce dwarf version 2 format debugging output in response to the '-g' option.

## 

[Target Hook]

Define this to enable the dwarf attribute DW\_AT\_calling\_convention to be emitted for each function. Instead of an integer return the enum value for the DW\_CC\_ tag.

To support optional call frame debugging information, you must also define INCOMING\_RETURN\_ADDR\_RTX and either set RTX\_FRAME\_RELATED\_P on the prologue insns if you use RTL for the prologue, or call dwarf2out\_def\_cfa and dwarf2out\_reg\_save as appropriate from TARGET\_ASM\_FUNCTION\_PROLOGUE if you don't.

#### DWARF2\_FRAME\_INFO

[Macro]

Define this macro to a nonzero value if GCC should always output Dwarf 2 frame information. If TARGET\_EXCEPT\_UNWIND\_INFO (see Section 18.20.9 [Exception Region Output], page 621) returns UI\_DWARF2, and exceptions are enabled, GCC will output this information not matter how you define DWARF2\_FRAME\_INFO.

# enum unwind\_info\_type TARGET\_DEBUG\_UNWIND\_INFO (void) [Target Hook] This hook defines the mechanism that will be used for describing frame unwind in-

formation to the debugger. Normally the hook will return UI\_DWARF2 if DWARF 2 debug information is enabled, and return UI\_NONE otherwise.

A target may return UI\_DWARF2 even when DWARF 2 debug information is disabled in order to always output DWARF 2 frame information.

A target may return UI\_TARGET if it has ABI specified unwind tables. This will suppress generation of the normal debug frame unwind information.

## DWARF2\_ASM\_LINE\_DEBUG\_INFO

[Macro]

Define this macro to be a nonzero value if the assembler can generate Dwarf 2 line debug info sections. This will result in much more compact line number tables, and hence is desirable if it works.

## DWARF2\_ASM\_VIEW\_DEBUG\_INFO

[Macro]

Define this macro to be a nonzero value if the assembler supports view assignment and verification in .loc. If it does not, but the user enables location views, the compiler may have to fallback to internal line number tables.

## int TARGET\_RESET\_LOCATION\_VIEW (rtx\_insn \*)

[Target Hook]

This hook, if defined, enables -ginternal-reset-location-views, and uses its result to override cases in which the estimated min insn length might be nonzero even when a PC advance (i.e., a view reset) cannot be taken for granted.

If the hook is defined, it must return a positive value to indicate the insn definitely advances the PC, and so the view number can be safely assumed to be reset; a negative value to mean the insn definitely does not advance the PC, and os the view number must not be reset; or zero to decide based on the estimated insn length.

If insn length is to be regarded as reliable, set the hook to hook\_int\_rtx\_insn\_0.

### bool TARGET\_WANT\_DEBUG\_PUB\_SECTIONS

[Target Hook]

True if the .debug\_pubtypes and .debug\_pubnames sections should be emitted. These sections are not used on most platforms, and in particular GDB does not use them.

#### bool TARGET\_DELAY\_SCHED2

[Target Hook]

True if sched2 is not to be run at its normal place. This usually means it will be run as part of machine-specific reorg.

#### bool TARGET\_DELAY\_VARTRACK

[Target Hook]

True if vartrack is not to be run at its normal place. This usually means it will be run as part of machine-specific reorg.

#### bool TARGET\_NO\_REGISTER\_ALLOCATION

[Target Hook]

True if register allocation and the passes following it should not be run. Usually true only for virtual assembler targets.

### ASM\_OUTPUT\_DWARF\_DELTA (stream, size, label1, label2)

[Macro]

A C statement to issue assembly directives that create a difference lab1 minus lab2, using an integer of the given size.

## ASM\_OUTPUT\_DWARF\_VMS\_DELTA (stream, size, label1, label2)

[Macro]

A C statement to issue assembly directives that create a difference between the two given labels in system defined units, e.g. instruction slots on IA64 VMS, using an integer of the given size.

#### ASM\_OUTPUT\_DWARF\_OFFSET (stream, size, label, offset, section) [Macro]

A C statement to issue assembly directives that create a section-relative reference to the given *label* plus *offset*, using an integer of the given *size*. The label is known to be defined in the given *section*.

## ASM\_OUTPUT\_DWARF\_PCREL (stream, size, label)

[Macro]

A C statement to issue assembly directives that create a self-relative reference to the given *label*, using an integer of the given *size*.

#### ASM\_OUTPUT\_DWARF\_DATAREL (stream, size, label)

[Macro]

A C statement to issue assembly directives that create a reference to the given *label* relative to the dbase, using an integer of the given *size*.

### ASM\_OUTPUT\_DWARF\_TABLE\_REF (label)

[Macro]

A C statement to issue assembly directives that create a reference to the DWARF table identifier *label* from the current section. This is used on some systems to avoid garbage collecting a DWARF table which is referenced by a function.

void TARGET\_ASM\_OUTPUT\_DWARF\_DTPREL (FILE \*file, int size, [Target Hook] rtx x)

If defined, this target hook is a function which outputs a DTP-relative reference to the given TLS symbol of the specified size.

# 18.21.6 Macros for VMS Debug Format

Here are macros for VMS debug format.

#### VMS\_DEBUGGING\_INFO

[Macro]

Define this macro if GCC should produce debugging output for VMS in response to the '-g' option. The default behavior for VMS is to generate minimal debug info for a traceback in the absence of '-g' unless explicitly overridden with '-g0'. This behavior is controlled by TARGET\_OPTION\_OPTIMIZATION and TARGET\_OPTION\_OVERRIDE.

# 18.22 Cross Compilation and Floating Point

While all modern machines use twos-complement representation for integers, there are a variety of representations for floating point numbers. This means that in a cross-compiler the representation of floating point numbers in the compiled program may be different from that used in the machine doing the compilation.

Because different representation systems may offer different amounts of range and precision, all floating point constants must be represented in the target machine's format. Therefore, the cross compiler cannot safely use the host machine's floating point arithmetic; it must emulate the target's arithmetic. To ensure consistency, GCC always uses emulation to work with floating point values, even when the host and target floating point formats are identical.

The following macros are provided by 'real.h' for the compiler to use. All parts of the compiler which generate or optimize floating-point calculations must use these macros. They may evaluate their operands more than once, so operands must not have side effects.

REAL\_VALUE\_TYPE [Macro]

The C data type to be used to hold a floating point value in the target machine's format. Typically this is a struct containing an array of HOST\_WIDE\_INT, but all code should treat it as an opaque quantity.

## HOST\_WIDE\_INT REAL\_VALUE\_FIX (REAL\_VALUE\_TYPE x)

[Macro]

Truncates x to a signed integer, rounding toward zero.

# unsigned HOST\_WIDE\_INT REAL\_VALUE\_UNSIGNED\_FIX (REAL\_VALUE\_TYPE x)

[Macro]

Truncates x to an unsigned integer, rounding toward zero. If x is negative, returns zero.

# REAL\_VALUE\_TYPE REAL\_VALUE\_ATOF (const char \*string, machine\_mode [Macro] mode)

Converts *string* into a floating point number in the target machine's representation for mode *mode*. This routine can handle both decimal and hexadecimal floating point constants, using the syntax defined by the C language for both.

## int REAL\_VALUE\_NEGATIVE (REAL\_VALUE\_TYPE x)

[Macro]

Returns 1 if x is negative (including negative zero), 0 otherwise.

## int REAL\_VALUE\_ISINF (REAL\_VALUE\_TYPE x)

[Macro]

Determines whether x represents infinity (positive or negative).

## int REAL\_VALUE\_ISNAN (REAL\_VALUE\_TYPE x)

[Macro]

Determines whether x represents a "NaN" (not-a-number).

# ${\tt REAL\_VALUE\_TYPE\ REAL\_VALUE\_NEGATE\ } (REAL\_VALUE\_TYPE\ {\tt x})$

[Macro]

Returns the negative of the floating point value x.

# REAL\_VALUE\_TYPE REAL\_VALUE\_ABS (REAL\_VALUE\_TYPE x)

[Macro]

Returns the absolute value of x.

# 18.23 Mode Switching Instructions

The following macros control mode switching optimizations:

## OPTIMIZE\_MODE\_SWITCHING (entity)

[Macro]

Define this macro if the port needs extra instructions inserted for mode switching in an optimizing compilation.

For an example, the SH4 can perform both single and double precision floating point operations, but to perform a single precision operation, the FPSCR PR bit has to be cleared, while for a double precision operation, this bit has to be set. Changing the PR bit requires a general purpose register as a scratch register, hence these FPSCR sets have to be inserted before reload, i.e. you cannot put this into instruction emitting or TARGET\_MACHINE\_DEPENDENT\_REORG.

You can have multiple entities that are mode-switched, and select at run time which entities actually need it. OPTIMIZE\_MODE\_SWITCHING should return nonzero for any entity that needs mode-switching. If you define this macro, you also have to define NUM\_MODES\_FOR\_MODE\_SWITCHING, TARGET\_MODE\_NEEDED, TARGET\_MODE\_PRIORITY and TARGET\_MODE\_EMIT. TARGET\_MODE\_AFTER, TARGET\_MODE\_ENTRY, and TARGET\_MODE\_EXIT are optional.

#### NUM\_MODES\_FOR\_MODE\_SWITCHING

[Macro]

If you define OPTIMIZE\_MODE\_SWITCHING, you have to define this as initializer for an array of integers. Each initializer element N refers to an entity that needs mode switching, and specifies the number of different modes that might need to be set for this entity. The position of the initializer in the initializer—starting counting at zero—determines the integer that is used to refer to the mode-switched entity in question. In macros that take mode arguments / yield a mode result, modes are represented as numbers  $0 \dots N-1$ . N is used to specify that no mode switch is needed / supplied.

# void TARGET\_MODE\_EMIT (int entity, int mode, int prev\_mode, HARD\_REG\_SET regs\_live) [Target Hook]

Generate one or more insns to set *entity* to *mode*. *hard\_reg\_live* is the set of hard registers live at the point where the insn(s) are to be inserted. *prev\_moxde* indicates the mode to switch from. Sets of a lower numbered entity will be emitted before sets of a higher numbered entity to a mode of the same or lower priority.

- int TARGET\_MODE\_NEEDED (int entity, rtx\_insn \*insn) [Target Hook] entity is an integer specifying a mode-switched entity. If OPTIMIZE\_MODE\_SWITCHING is defined, you must define this macro to return an integer value not larger than the corresponding element in NUM\_MODES\_FOR\_MODE\_SWITCHING, to denote the mode that entity must be switched into prior to the execution of insn.
- int TARGET\_MODE\_AFTER (int entity, int mode, rtx\_insn \*insn) [Target Hook] entity is an integer specifying a mode-switched entity. If this macro is defined, it is evaluated for every insn during mode switching. It determines the mode that an insn results in (if different from the incoming mode).
- int TARGET\_MODE\_ENTRY (int entity) [Target Hook] If this macro is defined, it is evaluated for every entity that needs mode switching. It should evaluate to an integer, which is a mode that entity is assumed to be switched to at function entry. If TARGET\_MODE\_ENTRY is defined then TARGET\_MODE\_EXIT must be defined.
- int TARGET\_MODE\_EXIT (int entity) [Target Hook] If this macro is defined, it is evaluated for every entity that needs mode switching. It should evaluate to an integer, which is a mode that entity is assumed to be switched to at function exit. If TARGET\_MODE\_EXIT is defined then TARGET\_MODE\_ENTRY must be defined.

### 18.24 Defining target-specific uses of \_\_attribute\_\_

Target-specific attributes may be defined for functions, data and types. These are described using the following target hooks; they also need to be documented in 'extend.texi'.

- const struct attribute\_spec \* TARGET\_ATTRIBUTE\_TABLE [Target Hook]

  If defined, this target hook points to an array of 'struct attribute\_spec' (defined in 'tree-core.h') specifying the machine specific attributes for this target and some of the restrictions on the entities to which these attributes are applied and the arguments they take.
- bool TARGET\_ATTRIBUTE\_TAKES\_IDENTIFIER\_P (const\_tree name) [Target Hook] If defined, this target hook is a function which returns true if the machine-specific attribute named name expects an identifier given as its first argument to be passed on

as a plain identifier, not subjected to name lookup. If this is not defined, the default is false for all machine-specific attributes.

# int TARGET\_COMP\_TYPE\_ATTRIBUTES (const\_tree type1, const\_tree type2) [Target Hook]

If defined, this target hook is a function which returns zero if the attributes on *type1* and *type2* are incompatible, one if they are compatible, and two if they are nearly compatible (which causes a warning to be generated). If this is not defined, machine-specific attributes are supposed always to be compatible.

- void TARGET\_SET\_DEFAULT\_TYPE\_ATTRIBUTES (tree type) [Target Hook] If defined, this target hook is a function which assigns default attributes to the newly defined type.
- tree TARGET\_MERGE\_TYPE\_ATTRIBUTES (tree type1, tree type2) [Target Hook] Define this target hook if the merging of type attributes needs special handling. If defined, the result is a list of the combined TYPE\_ATTRIBUTES of type1 and type2. It is assumed that comptypes has already been called and returned 1. This function may call merge\_attributes to handle machine-independent merging.

#### 

Define this target hook if the merging of decl attributes needs special handling. If defined, the result is a list of the combined DECL\_ATTRIBUTES of olddecl and newdecl. newdecl is a duplicate declaration of olddecl. Examples of when this is needed are when one attribute overrides another, or when an attribute is nullified by a subsequent definition. This function may call merge\_attributes to handle machine-independent merging.

If the only target-specific handling you require is 'dllimport' for Microsoft Windows targets, you should define the macro TARGET\_DLLIMPORT\_DECL\_ATTRIBUTES to 1. The compiler will then define a function called merge\_dllimport\_decl\_attributes which can then be defined as the expansion of TARGET\_MERGE\_DECL\_ATTRIBUTES. You can also add handle\_dll\_attribute in the attribute table for your port to perform initial processing of the 'dllimport' and 'dllexport' attributes. This is done in 'i386/cygwin.h' and 'i386/i386.c', for example.

bool TARGET\_VALID\_DLLIMPORT\_ATTRIBUTE\_P (const\_tree decl) [Target Hook] decl is a variable or function with \_\_attribute\_\_((dllimport)) specified. Use this hook if the target needs to add extra validation checks to handle\_dll\_attribute.

TARGET\_DECLSPEC [Macro]

Define this macro to a nonzero value if you want to treat \_\_declspec(X) as equivalent to \_\_attribute((X)). By default, this behavior is enabled only for targets that define TARGET\_DLLIMPORT\_DECL\_ATTRIBUTES. The current implementation of \_\_declspec is via a built-in macro, but you should not rely on this implementation detail.

void TARGET\_INSERT\_ATTRIBUTES (tree node, tree \*attr\_ptr) [Target Hook]

Define this target hook if you want to be able to add attributes to a decl when it is being created. This is normally useful for back ends which wish to implement a

pragma by using the attributes which correspond to the pragma's effect. The node argument is the decl which is being created. The attr\_ptr argument is a pointer to the attribute list for this decl. The list itself should not be modified, since it may be shared with other decls, but attributes may be chained on the head of the list and \*attr\_ptr modified to point to the new attributes, or a copy of the list may be made if further changes are needed.

# tree TARGET\_HANDLE\_GENERIC\_ATTRIBUTE (tree \*node, tree name, tree args, int flags, bool \*no\_add\_attrs) [Target Hook]

Define this target hook if you want to be able to perform additional target-specific processing of an attribute which is handled generically by a front end. The arguments are the same as those which are passed to attribute handlers. So far this only affects the *noinit* and *section* attribute.

### 

This target hook returns **true** if it is OK to inline *fndecl* into the current function, despite its having target-specific attributes, **false** otherwise. By default, if a function has a target specific attribute attached to it, it will not be inlined.

#### 

This hook is called to parse attribute(target("...")), which allows setting target-specific options on individual functions. These function-specific options may differ from the options specified on the command line. The hook should return true if the options are valid.

The hook should set the DECL\_FUNCTION\_SPECIFIC\_TARGET field in the function declaration to hold a pointer to a target-specific struct cl\_target\_option structure.

#### 

This hook is called to save any additional target-specific information in the struct cl\_target\_option structure for function-specific options from the struct gcc\_options structure. See Section 8.1 [Option file format], page 119.

#### 

This hook is called to restore any additional target-specific information in the struct cl\_target\_option structure for function-specific options to the struct gcc\_options structure.

### 

This hook is called to update target-specific information in the struct cl\_target\_option structure after it is streamed in from LTO bytecode.

#### 

This hook is called to print any additional target-specific information in the struct cl\_target\_option structure for function-specific options.

bool TARGET\_OPTION\_PRAGMA\_PARSE (tree args, tree pop\_target) [Target Hook] This target hook parses the options for #pragma GCC target, which sets the target-specific options for functions that occur later in the input stream. The options accepted should be the same as those handled by the TARGET\_OPTION\_VALID\_ATTRIBUTE\_P hook.

### void TARGET\_OPTION\_OVERRIDE (void)

[Target Hook]

Sometimes certain combinations of command options do not make sense on a particular target machine. You can override the hook TARGET\_OPTION\_OVERRIDE to take account of this. This hooks is called once just after all the command options have been parsed.

Don't use this hook to turn on various extra optimizations for '-0'. That is what TARGET\_OPTION\_OPTIMIZATION is for.

If you need to do something whenever the optimization level is changed via the optimize attribute or pragma, see TARGET\_OVERRIDE\_OPTIONS\_AFTER\_CHANGE

# bool TARGET\_OPTION\_FUNCTION\_VERSIONS (tree decl1, tree decl2) [Target Hook]

This target hook returns **true** if *DECL1* and *DECL2* are versions of the same function. *DECL1* and *DECL2* are function versions if and only if they have the same function signature and different target specific attributes, that is, they are compiled for different target machines.

### bool TARGET\_CAN\_INLINE\_P (tree caller, tree callee) [Target Hook]

This target hook returns false if the *caller* function cannot inline *callee*, based on target specific information. By default, inlining is not allowed if the callee function has function specific target options and the caller does not use the same options.

### void TARGET\_RELAYOUT\_FUNCTION (tree fndec1)

[Target Hook]

This target hook fixes function *findecl* after attributes are processed. Default does nothing. On ARM, the default function's alignment is updated with the attribute target.

### 18.25 Emulating TLS

For targets whose psABI does not provide Thread Local Storage via specific relocations and instruction sequences, an emulation layer is used. A set of target hooks allows this emulation layer to be configured for the requirements of a particular target. For instance the psABI may in fact specify TLS support in terms of an emulation layer.

The emulation layer works by creating a control object for every TLS object. To access the TLS object, a lookup function is provided which, when given the address of the control object, will return the address of the current thread's instance of the TLS object.

### const char \* TARGET\_EMUTLS\_GET\_ADDRESS

[Target Hook]

Contains the name of the helper function that uses a TLS control object to locate a TLS instance. The default causes libgcc's emulated TLS helper function to be used.

### const char \* TARGET\_EMUTLS\_REGISTER\_COMMON

[Target Hook]

Contains the name of the helper function that should be used at program startup to register TLS objects that are implicitly initialized to zero. If this is NULL, all TLS objects will have explicit initializers. The default causes libgcc's emulated TLS registration function to be used.

### const char \* TARGET\_EMUTLS\_VAR\_SECTION

[Target Hook]

Contains the name of the section in which TLS control variables should be placed. The default of NULL allows these to be placed in any section.

### const char \* TARGET\_EMUTLS\_TMPL\_SECTION

[Target Hook]

Contains the name of the section in which TLS initializers should be placed. The default of NULL allows these to be placed in any section.

### const char \* TARGET\_EMUTLS\_VAR\_PREFIX

[Target Hook]

Contains the prefix to be prepended to TLS control variable names. The default of NULL uses a target-specific prefix.

### const char \* TARGET\_EMUTLS\_TMPL\_PREFIX

[Target Hook]

Contains the prefix to be prepended to TLS initializer objects. The default of NULL uses a target-specific prefix.

### tree TARGET\_EMUTLS\_VAR\_FIELDS (tree type, tree \*name)

[Target Hook]

Specifies a function that generates the FIELD\_DECLs for a TLS control object type. type is the RECORD\_TYPE the fields are for and name should be filled with the structure tag, if the default of \_\_emutls\_object is unsuitable. The default creates a type suitable for libgcc's emulated TLS function.

### 

[Target Hook]

Cmpr\_auur) Specifies a functio

Specifies a function that generates the CONSTRUCTOR to initialize a TLS control object. var is the TLS control object, decl is the TLS object and tmpl\_addr is the address of the initializer. The default initializes libgcc's emulated TLS control object.

### bool TARGET\_EMUTLS\_VAR\_ALIGN\_FIXED

[Target Hook]

Specifies whether the alignment of TLS control variable objects is fixed and should not be increased as some backends may do to optimize single objects. The default is false.

### bool TARGET\_EMUTLS\_DEBUG\_FORM\_TLS\_ADDRESS

[Target Hook]

Specifies whether a DWARF DW\_OP\_form\_tls\_address location descriptor may be used to describe emulated TLS control objects.

### 18.26 Defining coprocessor specifics for MIPS targets.

The MIPS specification allows MIPS implementations to have as many as 4 coprocessors, each with as many as 32 private registers. GCC supports accessing these registers and transferring values between the registers and memory using asm-ized variables. For example:

```
register unsigned int cp0count asm ("c0r1");
unsigned int d;
d = cp0count + 3;
```

("c0r1" is the default name of register 1 in coprocessor 0; alternate names may be added as described below, or the default names may be overridden entirely in SUBTARGET\_CONDITIONAL\_REGISTER\_USAGE.)

Coprocessor registers are assumed to be epilogue-used; sets to them will be preserved even if it does not appear that the register is used again later in the function.

Another note: according to the MIPS spec, coprocessor 1 (if present) is the FPU. One accesses COP1 registers through standard mips floating-point support; they are not included in this mechanism.

### 18.27 Parameters for Precompiled Header Validity Checking

### void \* TARGET\_GET\_PCH\_VALIDITY (size\_t \*sz)

[Target Hook]

This hook returns a pointer to the data needed by TARGET\_PCH\_VALID\_P and sets '\*sz' to the size of the data in bytes.

const char \* TARGET\_PCH\_VALID\_P (const void \*data, size\_t sz) [Target Hook]
This hook checks whether the options used to create a PCH file are compatible with
the current settings. It returns NULL if so and a suitable error message if not. Error
messages will be presented to the user and must be localized using '\_(msg)'.

data is the data that was returned by TARGET\_GET\_PCH\_VALIDITY when the PCH file was created and sz is the size of that data in bytes. It's safe to assume that the data was created by the same version of the compiler, so no format checking is needed.

The default definition of default\_pch\_valid\_p should be suitable for most targets.

#### 

If this hook is nonnull, the default implementation of TARGET\_PCH\_VALID\_P will use it to check for compatible values of target\_flags. pch\_flags specifies the value that target\_flags had when the PCH file was created. The return value is the same as for TARGET\_PCH\_VALID\_P.

### void TARGET\_PREPARE\_PCH\_SAVE (void)

[Target Hook]

Called before writing out a PCH file. If the target has some garbage-collected data that needs to be in a particular state on PCH loads, it can use this hook to enforce that state. Very few targets need to do anything here.

### 18.28 C++ ABI parameters

### tree TARGET\_CXX\_GUARD\_TYPE (void)

[Target Hook]

Define this hook to override the integer type used for guard variables. These are used to implement one-time construction of static objects. The default is long\_long\_integer\_type\_node.

### bool TARGET\_CXX\_GUARD\_MASK\_BIT (void)

[Target Hook]

This hook determines how guard variables are used. It should return false (the default) if the first byte should be used. A return value of true indicates that only the least significant bit should be used.

### tree TARGET\_CXX\_GET\_COOKIE\_SIZE (tree type)

[Target Hook]

This hook returns the size of the cookie to use when allocating an array whose elements have the indicated *type*. Assumes that it is already known that a cookie is needed. The default is max(sizeof (size\_t), alignof(type)), as defined in section 2.7 of the IA64/Generic C++ ABI.

### bool TARGET\_CXX\_COOKIE\_HAS\_SIZE (void)

[Target Hook]

This hook should return **true** if the element size should be stored in array cookies. The default is to return **false**.

### 

[Target Hook]

If defined by a backend this hook allows the decision made to export class type to be overruled. Upon entry  $import\_export$  will contain 1 if the class is going to be exported, -1 if it is going to be imported and 0 otherwise. This function should return the modified value and perform any other actions necessary to support the backend's targeted operating system.

### bool TARGET\_CXX\_CDTOR\_RETURNS\_THIS (void)

[Target Hook]

This hook should return **true** if constructors and destructors return the address of the object created/destroyed. The default is to return **false**.

### bool TARGET\_CXX\_KEY\_METHOD\_MAY\_BE\_INLINE (void)

[Target Hook]

This hook returns true if the key method for a class (i.e., the method which, if defined in the current translation unit, causes the virtual table to be emitted) may be an inline function. Under the standard Itanium C++ ABI the key method may be an inline function so long as the function is not declared inline in the class definition. Under some variants of the ABI, an inline function can never be the key method. The default is to return true.

#### 

decl is a virtual table, virtual table table, typeinfo object, or other similar implicit class data object that will be emitted with external linkage in this translation unit. No ELF visibility has been explicitly specified. If the target needs to specify a visibility other than that of the containing class, use this hook to set DECL\_VISIBILITY and DECL\_VISIBILITY\_SPECIFIED.

### bool TARGET\_CXX\_CLASS\_DATA\_ALWAYS\_COMDAT (void)

[Target Hook]

This hook returns true (the default) if virtual tables and other similar implicit class data objects are always COMDAT if they have external linkage. If this hook returns false, then class data for classes whose virtual table will be emitted in only one translation unit will not be COMDAT.

### bool TARGET\_CXX\_LIBRARY\_RTTI\_COMDAT (void)

[Target Hook]

This hook returns true (the default) if the RTTI information for the basic types which is defined in the C++ runtime should always be COMDAT, false if it should not be COMDAT.

### bool TARGET\_CXX\_USE\_AEABI\_ATEXIT (void)

[Target Hook]

This hook returns true if \_\_aeabi\_atexit (as defined by the ARM EABI) should be used to register static destructors when '-fuse-cxa-atexit' is in effect. The default is to return false to use \_\_cxa\_atexit.

### bool TARGET\_CXX\_USE\_ATEXIT\_FOR\_CXA\_ATEXIT (void)

[Target Hook]

This hook returns true if the target atexit function can be used in the same manner as \_\_cxa\_atexit to register C++ static destructors. This requires that atexit-registered functions in shared libraries are run in the correct order when the libraries are unloaded. The default is to return false.

void TARGET\_CXX\_ADJUST\_CLASS\_AT\_DEFINITION (tree type) [Target Hook] type is a C++ class (i.e., RECORD\_TYPE or UNION\_TYPE) that has just been defined. Use this hook to make adjustments to the class (eg, tweak visibility or perform any other required target modifications).

tree TARGET\_CXX\_DECL\_MANGLING\_CONTEXT (const\_tree decl)
Return target-specific mangling context of decl or NULL\_TREE.

[Target Hook]

18.29 D ABI parameters

# void TARGET\_D\_CPU\_VERSIONS (void)

[D Target Hook]

Declare all environmental version identifiers relating to the target CPU using the function builtin\_version, which takes a string representing the name of the version. Version identifiers predefined by this hook apply to all modules that are being compiled and imported.

### void TARGET\_D\_OS\_VERSIONS (void)

[D Target Hook]

Similarly to TARGET\_D\_CPU\_VERSIONS, but is used for versions relating to the target operating system.

### unsigned TARGET\_D\_CRITSEC\_SIZE (void)

[D Target Hook]

Returns the size of the data structure used by the target operating system for critical sections and monitors. For example, on Microsoft Windows this would return the sizeof(CRITICAL\_SECTION), while other platforms that implement pthreads would return sizeof(pthread\_mutex\_t).

## 18.30 Adding support for named address spaces

The draft technical report of the ISO/IEC JTC1 S22 WG14 N1275 standards committee, Programming Languages - C - Extensions to support embedded processors, specifies a syntax for embedded processors to specify alternate address spaces. You can configure a GCC port to support section 5.1 of the draft report to add support for address spaces other than the default address space. These address spaces are new keywords that are similar to the volatile and const type attributes.

Pointers to named address spaces can have a different size than pointers to the generic address space.

For example, the SPU port uses the \_\_ea address space to refer to memory in the host processor, rather than memory local to the SPU processor. Access to memory in the \_\_ea address space involves issuing DMA operations to move data between the host processor and the local processor memory address space. Pointers in the \_\_ea address space are either 32 bits or 64 bits based on the '-mea32' or '-mea64' switches (native SPU pointers are always 32 bits).

Internally, address spaces are represented as a small integer in the range 0 to 15 with address space 0 being reserved for the generic address space.

To register a named address space qualifier keyword with the C front end, the target may call the c\_register\_addr\_space routine. For example, the SPU port uses the following to declare \_\_ea as the keyword for named address space #1:

```
#define ADDR_SPACE_EA 1
c_register_addr_space ("__ea", ADDR_SPACE_EA);
```

### scalar\_int\_mode TARGET\_ADDR\_SPACE\_POINTER\_MODE

[Target Hook]

(addr\_space\_t address\_space)

Define this to return the machine mode to use for pointers to address\_space if the target supports named address spaces. The default version of this hook returns ptr\_mode

# scalar\_int\_mode TARGET\_ADDR\_SPACE\_ADDRESS\_MODE

[Target Hook]

(addr\_space\_t address\_space)

Define this to return the machine mode to use for addresses in address\_space if the target supports named address spaces. The default version of this hook returns Pmode.

# bool TARGET\_ADDR\_SPACE\_VALID\_POINTER\_MODE (scalar\_int\_mode [Target Hook] mode, addr\_space\_t as)

Define this to return nonzero if the port can handle pointers with machine mode *mode* to address space as. This target hook is the same as the TARGET\_VALID\_POINTER\_MODE target hook, except that it includes explicit named address space support. The default version of this hook returns true for the modes returned by either the TARGET\_ADDR\_SPACE\_POINTER\_MODE or TARGET\_ADDR\_SPACE\_ADDRESS\_MODE target hooks for the given address space.

### bool TARGET\_ADDR\_SPACE\_LEGITIMATE\_ADDRESS\_P

[Target Hook]

(machine\_mode mode, rtx exp, bool strict, addr\_space\_t as)

Define this to return true if  $\exp$  is a valid address for mode mode in the named address space as. The strict parameter says whether strict addressing is in effect after reload has finished. This target hook is the same as the TARGET\_LEGITIMATE\_ADDRESS\_P target hook, except that it includes explicit named address space support.

# rtx TARGET\_ADDR\_SPACE\_LEGITIMIZE\_ADDRESS (rtx x, rtx oldx, machine\_mode mode, addr\_space\_t as) [Target Hook]

Define this to modify an invalid address x to be a valid address with mode *mode* in the named address space as. This target hook is the same as the TARGET\_LEGITIMIZE\_ADDRESS target hook, except that it includes explicit named address space support.

# bool TARGET\_ADDR\_SPACE\_SUBSET\_P (addr\_space\_t subset, addr\_space\_t superset) [Target Hook]

Define this to return whether the *subset* named address space is contained within the *superset* named address space. Pointers to a named address space that is a subset of another named address space will be converted automatically without a cast if used together in arithmetic operations. Pointers to a superset address space can be converted to pointers to a subset address space via explicit casts.

### 

Define this to modify the default handling of address 0 for the address space. Return true if 0 should be considered a valid address.

# rtx TARGET\_ADDR\_SPACE\_CONVERT (rtx op, tree from\_type, tree [Target Hook] to\_type)

Define this to convert the pointer expression represented by the RTL op with type from\_type that points to a named address space to a new pointer expression with type to\_type that points to a different named address space. When this hook it called, it is guaranteed that one of the two address spaces is a subset of the other, as determined by the TARGET\_ADDR\_SPACE\_SUBSET\_P target hook.

# int TARGET\_ADDR\_SPACE\_DEBUG (addr\_space\_t as) [Target Hook] Define this to define how the address space is encoded in dwarf. The result is the

Define this to define how the address space is encoded in dwarf. The result is the value to be used with DW\_AT\_address\_class.

#### 

Define this hook if the availability of an address space depends on command line options and some diagnostics should be printed when the address space is used. This hook is called during parsing and allows to emit a better diagnostic compared to the case where the address space was not registered with c\_register\_addr\_space. as is the address space as registered with c\_register\_addr\_space. loc is the location of the address space qualifier token. The default implementation does nothing.

### 18.31 Miscellaneous Parameters

Here are several miscellaneous parameters.

### HAS\_LONG\_COND\_BRANCH

[Macro]

Define this boolean macro to indicate whether or not your architecture has conditional branches that can span all of memory. It is used in conjunction with an optimization that partitions hot and cold basic blocks into separate sections of the executable. If this macro is set to false, gcc will convert any conditional branches that attempt to cross between sections into unconditional branches or indirect jumps.

### HAS\_LONG\_UNCOND\_BRANCH

|Macro|

Define this boolean macro to indicate whether or not your architecture has unconditional branches that can span all of memory. It is used in conjunction with an optimization that partitions hot and cold basic blocks into separate sections of the

executable. If this macro is set to false, gcc will convert any unconditional branches that attempt to cross between sections into indirect jumps.

### CASE\_VECTOR\_MODE [Macro]

An alias for a machine mode name. This is the machine mode that elements of a jump-table should have.

### CASE\_VECTOR\_SHORTEN\_MODE (min\_offset, max\_offset, body) [Macro]

Optional: return the preferred mode for an addr\_diff\_vec when the minimum and maximum offset are known. If you define this, it enables extra code in branch shortening to deal with addr\_diff\_vec. To make this work, you also have to define INSN\_ALIGN and make the alignment for addr\_diff\_vec explicit. The body argument is provided so that the offset\_unsigned and scale flags can be updated.

### CASE\_VECTOR\_PC\_RELATIVE

[Macro]

Define this macro to be a C expression to indicate when jump-tables should contain relative addresses. You need not define this macro if jump-tables never contain relative addresses, or jump-tables should contain relative addresses only when '-fPIC' or '-fPIC' is in effect.

### unsigned int TARGET\_CASE\_VALUES\_THRESHOLD (void)

[Target Hook]

This function return the smallest number of different values for which it is best to use a jump-table instead of a tree of conditional branches. The default is four for machines with a casesi instruction and five otherwise. This is best for most machines.

### WORD\_REGISTER\_OPERATIONS

Macro

Define this macro to 1 if operations between registers with integral mode smaller than a word are always performed on the entire register. To be more explicit, if you start with a pair of word\_mode registers with known values and you do a subword, for example QImode, addition on the low part of the registers, then the compiler may consider that the result has a known value in word\_mode too if the macro is defined to 1. Most RISC machines have this property and most CISC machines do not.

### unsigned int TARGET\_MIN\_ARITHMETIC\_PRECISION (void)

[Target Hook]

On some RISC architectures with 64-bit registers, the processor also maintains 32-bit condition codes that make it possible to do real 32-bit arithmetic, although the operations are performed on the full registers.

On such architectures, defining this hook to 32 tells the compiler to try using 32-bit arithmetical operations setting the condition codes instead of doing full 64-bit arithmetic.

More generally, define this hook on RISC architectures if you want the compiler to try using arithmetical operations setting the condition codes with a precision lower than the word precision.

You need not define this hook if WORD\_REGISTER\_OPERATIONS is not defined to 1.

### LOAD\_EXTEND\_OP (mem\_mode)

[Macro]

Define this macro to be a C expression indicating when insns that read memory in  $mem\_mode$ , an integral mode narrower than a word, set the bits outside of  $mem\_mode$ 

to be either the sign-extension or the zero-extension of the data read. Return SIGN\_ EXTEND for values of mem\_mode for which the insn sign-extends, ZERO\_EXTEND for which it zero-extends, and UNKNOWN for other modes.

This macro is not called with mem\_mode non-integral or with a width greater than or equal to BITS\_PER\_WORD, so you may return any value in this case. Do not define this macro if it would always return UNKNOWN. On machines where this macro is defined, you will normally define it as the constant SIGN\_EXTEND or ZERO\_EXTEND.

You may return a non-UNKNOWN value even if for some hard registers the sign extension is not performed, if for the REGNO\_REG\_CLASS of these hard registers TARGET\_CAN\_ CHANGE\_MODE\_CLASS returns false when the from mode is mem\_mode and the to mode is any integral mode larger than this but not larger than word\_mode.

You must return UNKNOWN if for some hard registers that allow this mode, TARGET\_ CAN\_CHANGE\_MODE\_CLASS says that they cannot change to word\_mode, but that they can change to another integral mode that is larger than mem\_mode but still smaller than word\_mode.

### SHORT\_IMMEDIATES\_SIGN\_EXTEND

[Macro]

Define this macro to 1 if loading short immediate values into registers sign extends.

### unsigned int TARGET\_MIN\_DIVISIONS\_FOR\_RECIP\_MUL (machine\_mode mode)

[Target Hook]

When '-ffast-math' is in effect, GCC tries to optimize divisions by the same divisor,

by turning them into multiplications by the reciprocal. This target hook specifies the minimum number of divisions that should be there for GCC to perform the optimization for a variable of mode mode. The default implementation returns 3 if the machine has an instruction for the division, and 2 if it does not.

MOVE\_MAX [Macro]

The maximum number of bytes that a single instruction can move quickly between memory and registers or between two memory locations.

MAX\_MOVE\_MAX [Macro]

The maximum number of bytes that a single instruction can move quickly between memory and registers or between two memory locations. If this is undefined, the default is MOVE\_MAX. Otherwise, it is the constant value that is the largest value that MOVE\_MAX can have at run-time.

### SHIFT COUNT TRUNCATED

[Macro]

A C expression that is nonzero if on this machine the number of bits actually used for the count of a shift operation is equal to the number of bits needed to represent the size of the object being shifted. When this macro is nonzero, the compiler will assume that it is safe to omit a sign-extend, zero-extend, and certain bitwise 'and' instructions that truncates the count of a shift operation. On machines that have instructions that act on bit-fields at variable positions, which may include 'bit test' instructions, a nonzero SHIFT\_COUNT\_TRUNCATED also enables deletion of truncations of the values that serve as arguments to bit-field instructions.

If both types of instructions truncate the count (for shifts) and position (for bit-field operations), or if no variable-position bit-field instructions exist, you should define this macro.

However, on some machines, such as the 80386 and the 680x0, truncation only applies to shift operations and not the (real or pretended) bit-field operations. Define SHIFT\_COUNT\_TRUNCATED to be zero on such machines. Instead, add patterns to the 'md' file that include the implied truncation of the shift instructions.

You need not define this macro if it would always have the value of zero.

# unsigned HOST\_WIDE\_INT TARGET\_SHIFT\_TRUNCATION\_MASK [Target Hook] (machine\_mode mode)

This function describes how the standard shift patterns for *mode* deal with shifts by negative amounts or by more than the width of the mode. See [shift patterns], page 405.

On many machines, the shift patterns will apply a mask m to the shift count, meaning that a fixed-width shift of x by y is equivalent to an arbitrary-width shift of x by y & m. If this is true for mode mode, the function should return m, otherwise it should return 0. A return value of 0 indicates that no particular behavior is guaranteed.

Note that, unlike SHIFT\_COUNT\_TRUNCATED, this function does *not* apply to general shift rtxes; it applies only to instructions that are generated by the named shift patterns.

The default implementation of this function returns GET\_MODE\_BITSIZE (mode) - 1 if SHIFT\_COUNT\_TRUNCATED and 0 otherwise. This definition is always safe, but if SHIFT\_COUNT\_TRUNCATED is false, and some shift patterns nevertheless truncate the shift count, you may get better code by overriding it.

# bool TARGET\_TRULY\_NOOP\_TRUNCATION (poly\_uint64 outprec, poly\_uint64 inprec) [Target Hook]

This hook returns true if it is safe to "convert" a value of *inprec* bits to one of *outprec* bits (where *outprec* is smaller than *inprec*) by merely operating on it as if it had only *outprec* bits. The default returns true unconditionally, which is correct for most machines.

If TARGET\_MODES\_TIEABLE\_P returns false for a pair of modes, suboptimal code can result if this hook returns true for the corresponding mode sizes. Making this hook return false in such cases may improve things.

#### 

The representation of an integral mode can be such that the values are always extended to a wider integral mode. Return SIGN\_EXTEND if values of mode are represented in sign-extended form to rep\_mode. Return UNKNOWN otherwise. (Currently, none of the targets use zero-extended representation this way so unlike LOAD\_EXTEND\_OP, TARGET\_MODE\_REP\_EXTENDED is expected to return either SIGN\_EXTEND or UNKNOWN. Also no target extends mode to rep\_mode so that rep\_mode is not the next widest integral mode and currently we take advantage of this fact.)

Similarly to LOAD\_EXTEND\_OP you may return a non-UNKNOWN value even if the extension is not performed on certain hard registers as long as for the REGNO\_REG\_CLASS of these hard registers TARGET\_CAN\_CHANGE\_MODE\_CLASS returns false.

Note that TARGET\_MODE\_REP\_EXTENDED and LOAD\_EXTEND\_OP describe two related properties. If you define TARGET\_MODE\_REP\_EXTENDED (mode, word\_mode) you probably also want to define LOAD\_EXTEND\_OP (mode) to return the same type of extension. In order to enforce the representation of mode, TARGET\_TRULY\_NOOP\_TRUNCATION should return false when truncating to mode.

bool TARGET\_SETJMP\_PRESERVES\_NONVOLATILE\_REGS\_P (void) [Target Hook] On some targets, it is assumed that the compiler will spill all pseudos that are live across a call to setjmp, while other targets treat setjmp calls as normal function calls.

This hook returns false if setjmp calls do not preserve all non-volatile registers so that gcc that must spill all pseudos that are live across setjmp calls. Define this to return true if the target does not need to spill all pseudos live across setjmp calls. The default implementation conservatively assumes all pseudos must be spilled across setjmp calls.

STORE\_FLAG\_VALUE [Macro]

A C expression describing the value returned by a comparison operator with an integral mode and stored by a store-flag instruction ('cstoremode4') when the condition is true. This description must apply to *all* the 'cstoremode4' patterns and all the comparison operators whose results have a MODE\_INT mode.

A value of 1 or -1 means that the instruction implementing the comparison operator returns exactly 1 or -1 when the comparison is true and 0 when the comparison is false. Otherwise, the value indicates which bits of the result are guaranteed to be 1 when the comparison is true. This value is interpreted in the mode of the comparison operation, which is given by the mode of the first operand in the 'cstoremode4' pattern. Either the low bit or the sign bit of STORE\_FLAG\_VALUE be on. Presently, only those bits are used by the compiler.

If STORE\_FLAG\_VALUE is neither 1 or -1, the compiler will generate code that depends only on the specified bits. It can also replace comparison operators with equivalent operations if they cause the required bits to be set, even if the remaining bits are undefined. For example, on a machine whose comparison operators return an SImode value and where STORE\_FLAG\_VALUE is defined as '0x80000000', saying that just the sign bit is relevant, the expression

```
(ne:SI (and:SI x (const_int power-of-2)) (const_int 0))
can be converted to
    (ashift:SI x (const_int n))
```

where n is the appropriate shift count to move the bit being tested into the sign bit. There is no way to describe a machine that always sets the low-order bit for a true value, but does not guarantee the value of any other bits, but we do not know of any machine that has such an instruction. If you are trying to port GCC to such a machine, include an instruction to perform a logical-and of the result with 1 in the

Often, a machine will have multiple instructions that obtain a value from a comparison (or the condition codes). Here are rules to guide the choice of value for STORE\_FLAG\_VALUE, and hence the instructions to be used:

pattern for the comparison operators and let us know at gcc@gcc.gnu.org.

- Use the shortest sequence that yields a valid definition for STORE\_FLAG\_VALUE. It is more efficient for the compiler to "normalize" the value (convert it to, e.g., 1 or 0) than for the comparison operators to do so because there may be opportunities to combine the normalization with other operations.
- For equal-length sequences, use a value of 1 or -1, with -1 being slightly preferred on machines with expensive jumps and 1 preferred on other machines.
- As a second choice, choose a value of '0x80000001' if instructions exist that set both the sign and low-order bits but do not define the others.
- Otherwise, use a value of '0x80000000'.

Many machines can produce both the value chosen for STORE\_FLAG\_VALUE and its negation in the same number of instructions. On those machines, you should also define a pattern for those cases, e.g., one matching

```
(set A (neg:m (ne:m B C)))
```

Some machines can also perform and or plus operations on condition code values with less instructions than the corresponding 'cstoremode4' insn followed by and or plus. On those machines, define the appropriate patterns. Use the names incscc and decscc, respectively, for the patterns which perform plus or minus operations on condition code values. See 'rs6000.md' for some examples. The GNU Superoptimizer can be used to find such instruction sequences on other machines.

If this macro is not defined, the default value, 1, is used. You need not define STORE\_FLAG\_VALUE if the machine has no store-flag instructions, or if the value generated by these instructions is 1.

### FLOAT\_STORE\_FLAG\_VALUE (mode)

[Macro]

A C expression that gives a nonzero REAL\_VALUE\_TYPE value that is returned when comparison operators with floating-point results are true. Define this macro on machines that have comparison operations that return floating-point values. If there are no such operations, do not define this macro.

### VECTOR\_STORE\_FLAG\_VALUE (mode)

[Macro]

A C expression that gives a rtx representing the nonzero true element for vector comparisons. The returned rtx should be valid for the inner mode of *mode* which is guaranteed to be a vector mode. Define this macro on machines that have vector comparison operations that return a vector result. If there are no such operations, do not define this macro. Typically, this macro is defined as const1\_rtx or constm1\_rtx. This macro may return NULL\_RTX to prevent the compiler optimizing such vector comparison operations for the given mode.

# CLZ\_DEFINED\_VALUE\_AT\_ZERO (mode, value) CTZ\_DEFINED\_VALUE\_AT\_ZERO (mode, value)

[Macro] [Macro]

A C expression that indicates whether the architecture defines a value for clz or ctz with a zero operand. A result of 0 indicates the value is undefined. If the value is defined for only the RTL expression, the macro should evaluate to 1; if the value applies also to the corresponding optab entry (which is normally the case if it expands directly into the corresponding RTL), then the macro should evaluate to 2. In the cases where the value is defined, value should be set to this value.

If this macro is not defined, the value of clz or ctz at zero is assumed to be undefined.

This macro must be defined if the target's expansion for ffs relies on a particular value to get correct results. Otherwise it is not necessary, though it may be used to optimize some corner cases, and to provide a default expansion for the ffs optab.

Note that regardless of this macro the "definedness" of clz and ctz at zero do not extend to the builtin functions visible to the user. Thus one may be free to adjust the value at will to match the target expansion of these operations without fear of breaking the API.

Pmode [Macro]

An alias for the machine mode for pointers. On most machines, define this to be the integer mode corresponding to the width of a hardware pointer; SImode on 32-bit machine or DImode on 64-bit machines. On some machines you must define this to be one of the partial integer modes, such as PSImode.

The width of Pmode must be at least as large as the value of POINTER\_SIZE. If it is not equal, you must define the macro POINTERS\_EXTEND\_UNSIGNED to specify how pointers are extended to Pmode.

FUNCTION\_MODE [Macro]

An alias for the machine mode used for memory references to functions being called, in call RTL expressions. On most CISC machines, where an instruction can begin at any byte address, this should be QImode. On most RISC machines, where all instructions have fixed size and alignment, this should be a mode with the same size and alignment as the machine instruction words - typically SImode or HImode.

### STDC\_O\_IN\_SYSTEM\_HEADERS

[Macro]

In normal operation, the preprocessor expands \_\_STDC\_\_ to the constant 1, to signify that GCC conforms to ISO Standard C. On some hosts, like Solaris, the system compiler uses a different convention, where \_\_STDC\_\_ is normally 0, but is 1 if the user specifies strict conformance to the C Standard.

Defining STDC\_0\_IN\_SYSTEM\_HEADERS makes GNU CPP follows the host convention when processing system header files, but when processing user files \_\_STDC\_\_ will always expand to 1.

### const char \* TARGET\_C\_PREINCLUDE (void)

[C Target Hook]

Define this hook to return the name of a header file to be included at the start of all compilations, as if it had been included with #include <file>. If this hook returns NULL, or is not defined, or the header is not found, or if the user specifies '-ffreestanding' or '-nostdinc', no header is included.

This hook can be used together with a header provided by the system C library to implement ISO C requirements for certain macros to be predefined that describe properties of the whole implementation rather than just the compiler.

### bool TARGET\_CXX\_IMPLICIT\_EXTERN\_C (const char\*)

[C Target Hook]

Define this hook to add target-specific C++ implicit extern C functions. If this function returns true for the name of a file-scope function, that function implicitly gets extern "C" linkage rather than whatever language linkage the declaration would normally have. An example of such function is WinMain on Win32 targets.

### SYSTEM\_IMPLICIT\_EXTERN\_C

[Macro]

Define this macro if the system header files do not support C++. This macro handles system header files by pretending that system header files are enclosed in 'extern "C" {...}'.

### REGISTER\_TARGET\_PRAGMAS ()

[Macro]

Define this macro if you want to implement any target-specific pragmas. If defined, it is a C expression which makes a series of calls to c\_register\_pragma or c\_register\_pragma\_with\_expansion for each pragma. The macro may also do any setup required for the pragmas.

The primary reason to define this macro is to provide compatibility with other compilers for the same target. In general, we discourage definition of target-specific pragmas for GCC.

If the pragma can be implemented by attributes then you should consider defining the target hook 'TARGET\_INSERT\_ATTRIBUTES' as well.

Preprocessor macros that appear on pragma lines are not expanded. All '#pragma' directives that do not match any registered pragma are silently ignored, unless the user specifies '-Wunknown-pragmas'.

Each call to c\_register\_pragma or c\_register\_pragma\_with\_expansion establishes one pragma. The *callback* routine will be called when the preprocessor encounters a pragma of the form

#pragma [space] name ...

space is the case-sensitive namespace of the pragma, or NULL to put the pragma in the global namespace. The callback routine receives *pfile* as its first argument, which can be passed on to cpplib's functions if necessary. You can lex tokens after the *name* by calling pragma\_lex. Tokens that are not read by the callback will be silently ignored. The end of the line is indicated by a token of type CPP\_EOF. Macro expansion occurs on the arguments of pragmas registered with c\_register\_pragma\_with\_expansion but not on the arguments of pragmas registered with c\_register\_pragma.

Note that the use of pragma\_lex is specific to the C and C++ compilers. It will not work in the Java or Fortran compilers, or any other language compilers for that matter. Thus if pragma\_lex is going to be called from target-specific code, it must only be done so when building the C and C++ compilers. This can be done by defining the variables c\_target\_objs and cxx\_target\_objs in the target entry in the 'config.gcc' file. These variables should name the target-specific, language-specific object file which contains the code that uses pragma\_lex. Note it will also be necessary to add a rule to the makefile fragment pointed to by tmake\_file that shows how to build this object file.

### HANDLE\_PRAGMA\_PACK\_WITH\_EXPANSION

[Macro]

Define this macro if macros should be expanded in the arguments of '#pragma pack'.

### TARGET\_DEFAULT\_PACK\_STRUCT

[Macro]

If your target requires a structure packing default other than 0 (meaning the machine default), define this macro to the necessary value (in bytes). This must be a value that would also be valid to use with '#pragma pack()' (that is, a small power of two).

### DOLLARS\_IN\_IDENTIFIERS

[Macro]

Define this macro to control use of the character '\$' in identifier names for the C family of languages. 0 means '\$' is not allowed by default; 1 means it is allowed. 1 is the default; there is no need to define this macro in that case.

### INSN\_SETS\_ARE\_DELAYED (insn)

[Macro]

Define this macro as a C expression that is nonzero if it is safe for the delay slot scheduler to place instructions in the delay slot of *insn*, even if they appear to use a resource set or clobbered in *insn*. *insn* is always a jump\_insn or an insn; GCC knows that every call\_insn has this behavior. On machines where some insn or jump\_insn is really a function call and hence has this behavior, you should define this macro.

You need not define this macro if it would always return zero.

### INSN\_REFERENCES\_ARE\_DELAYED (insn)

[Macro]

Define this macro as a C expression that is nonzero if it is safe for the delay slot scheduler to place instructions in the delay slot of *insn*, even if they appear to set or clobber a resource referenced in *insn*. *insn* is always a <code>jump\_insn</code> or an <code>insn</code>. On machines where some <code>insn</code> or <code>jump\_insn</code> is really a function call and its operands are registers whose use is actually in the subroutine it calls, you should define this macro. Doing so allows the delay slot scheduler to move instructions which copy arguments into the argument registers into the delay slot of *insn*.

You need not define this macro if it would always return zero.

### MULTIPLE\_SYMBOL\_SPACES

[Macro]

Define this macro as a C expression that is nonzero if, in some cases, global symbols from one translation unit may not be bound to undefined symbols in another translation unit without user intervention. For instance, under Microsoft Windows symbols must be explicitly imported from shared libraries (DLLs).

You need not define this macro if it would always evaluate to zero.

#### 

This target hook may add *clobbers* to *clobbers* and *clobbered\_regs* for any hard regs the port wishes to automatically clobber for an asm. The *outputs* and *inputs* may be inspected to avoid clobbering a register that is already used by the asm.

It may modify the *outputs*, *inputs*, and *constraints* as necessary for other pre-processing. In this case the return value is a sequence of insus to emit after the asm.

MATH\_LIBRARY [Macro]

Define this macro as a C string constant for the linker argument to link in the system math library, minus the initial '"-1"', or '""' if the target does not have a separate math library.

You need only define this macro if the default of "m" is wrong.

### LIBRARY\_PATH\_ENV [Macro]

Define this macro as a C string constant for the environment variable that specifies where the linker should look for libraries.

You need only define this macro if the default of ""LIBRARY\_PATH"' is wrong.

TARGET\_POSIX\_IO [Macro]

Define this macro if the target supports the following POSIX file functions, access, mkdir and file locking with fcntl / F\_SETLKW. Defining TARGET\_POSIX\_IO will enable the test coverage code to use file locking when exiting a program, which avoids race conditions if the program has forked. It will also create directories at run-time for cross-profiling.

### MAX\_CONDITIONAL\_EXECUTE

[Macro]

A C expression for the maximum number of instructions to execute via conditional execution instructions instead of a branch. A value of BRANCH\_COST+1 is the default if the machine does not use cc0, and 1 if it does use cc0.

### IFCVT\_MODIFY\_TESTS (ce\_info, true\_expr, false\_expr)

[Macro]

Used if the target needs to perform machine-dependent modifications on the conditionals used for turning basic blocks into conditionally executed code. ce\_info points to a data structure, struct ce\_if\_block, which contains information about the currently processed blocks. true\_expr and false\_expr are the tests that are used for converting the then-block and the else-block, respectively. Set either true\_expr or false\_expr to a null pointer if the tests cannot be converted.

#### 

Like IFCVT\_MODIFY\_TESTS, but used when converting more complicated if-statements into conditions combined by and and or operations. *bb* contains the basic block that contains the test that is currently being processed and about to be turned into a condition.

### IFCVT\_MODIFY\_INSN (ce\_info, pattern, insn)

[Macro]

A C expression to modify the *PATTERN* of an *INSN* that is to be converted to conditional execution format. *ce\_info* points to a data structure, **struct ce\_if\_ block**, which contains information about the currently processed blocks.

### IFCVT\_MODIFY\_FINAL (ce\_info)

|Macro|

A C expression to perform any final machine dependent modifications in converting code to conditional execution. The involved basic blocks can be found in the struct ce\_if\_block structure that is pointed to by ce\_info.

### IFCVT\_MODIFY\_CANCEL (ce\_info)

[Macro]

A C expression to cancel any machine dependent modifications in converting code to conditional execution. The involved basic blocks can be found in the struct ce\_if\_block structure that is pointed to by ce\_info.

### IFCVT\_MACHDEP\_INIT (ce\_info)

[Macro]

A C expression to initialize any machine specific data for if-conversion of the if-block in the struct ce\_if\_block structure that is pointed to by ce\_info.

### void TARGET\_MACHINE\_DEPENDENT\_REORG (void)

[Target Hook]

If non-null, this hook performs a target-specific pass over the instruction stream. The compiler will run it at all optimization levels, just before the point at which it normally does delayed-branch scheduling.

The exact purpose of the hook varies from target to target. Some use it to do transformations that are necessary for correctness, such as laying out in-function constant pools or avoiding hardware hazards. Others use it as an opportunity to do some machine-dependent optimizations.

You need not implement the hook if it has nothing to do. The default definition is null.

### void TARGET\_INIT\_BUILTINS (void)

[Target Hook]

Define this hook if you have any machine-specific built-in functions that need to be defined. It should be a function that performs the necessary setup.

Machine specific built-in functions can be useful to expand special machine instructions that would otherwise not normally be generated because they have no equivalent in the source language (for example, SIMD vector instructions or prefetch instructions).

To create a built-in function, call the function lang\_hooks.builtin\_function which is defined by the language front end. You can use any type nodes set up by build\_common\_tree\_nodes; only language front ends that use those two functions will call 'TARGET\_INIT\_BUILTINS'.

tree TARGET\_BUILTIN\_DECL (unsigned code, bool initialize\_p) [Target Hook] Define this hook if you have any machine-specific built-in functions that need to be defined. It should be a function that returns the builtin function declaration for the builtin function code code. If there is no such builtin and it cannot be initialized at this time if initialize\_p is true the function should return NULL\_TREE. If code is out of range the function should return error\_mark\_node.

# rtx TARGET\_EXPAND\_BUILTIN (tree exp, rtx target, rtx subtarget, machine\_mode mode, int ignore)

[Target Hook]

Expand a call to a machine specific built-in function that was set up by 'TARGET\_INIT\_BUILTINS'. exp is the expression for the function call; the result should go to target if that is convenient, and have mode mode if that is convenient. subtarget may be used as the target for computing one of exp's operands. ignore is nonzero if the value is to be ignored. This function should return the result of the call to the built-in function.

tree TARGET\_RESOLVE\_OVERLOADED\_BUILTIN (unsigned int loc, tree fndecl, void \*arglist) [Target Hook]

Select a replacement for a machine specific built-in function that was set up by 'TARGET\_INIT\_BUILTINS'. This is done *before* regular type checking, and so allows the target to implement a crude form of function overloading. *findecl* is the declaration of the built-in function. *arglist* is the list of arguments passed to the built-in function. The result is a complete expression that implements the operation, usually another CALL\_EXPR. *arglist* really has type 'VEC(tree,gc)\*'

bool TARGET\_CHECK\_BUILTIN\_CALL (location\_t loc, vec<location\_t> [Target Hook] arg\_loc, tree fndecl, tree orig\_fndecl, unsigned int nargs, tree \*args)

Perform semantic checking on a call to a machine-specific built-in function after its arguments have been constrained to the function signature. Return true if the call is valid, otherwise report an error and return false.

This hook is called after TARGET\_RESOLVE\_OVERLOADED\_BUILTIN. The call was originally to built-in function *orig\_fndecl*, but after the optional TARGET\_RESOLVE\_OVERLOADED\_BUILTIN step is now to built-in function *fndecl*. *loc* is the location of the call and *args* is an array of function arguments, of which there are *nargs*. *arg\_loc* specifies the location of each argument.

tree TARGET\_FOLD\_BUILTIN (tree fndec1, int n\_args, tree \*argp, bool ignore) [Target Hook]

Fold a call to a machine specific built-in function that was set up by 'TARGET\_INIT\_BUILTINS'. fndecl is the declaration of the built-in function.  $n\_args$  is the number of arguments passed to the function; the arguments themselves are pointed to by argp. The result is another tree, valid for both GIMPLE and GENERIC, containing a simplified expression for the call's result. If ignore is true the value will be ignored.

- bool TARGET\_GIMPLE\_FOLD\_BUILTIN (gimple\_stmt\_iterator \*gsi) [Target Hook] Fold a call to a machine specific built-in function that was set up by 'TARGET\_INIT\_BUILTINS'. gsi points to the gimple statement holding the function call. Returns true if any change was made to the GIMPLE stream.
- int TARGET\_COMPARE\_VERSION\_PRIORITY (tree decl1, tree decl2) [Target Hook] This hook is used to compare the target attributes in two functions to determine which function's features get higher priority. This is used during function multi-versioning to figure out the order in which two versions must be dispatched. A function version with a higher priority is checked for dispatching earlier. decl1 and decl2 are the two function decls that will be compared.
- tree TARGET\_GET\_FUNCTION\_VERSIONS\_DISPATCHER (void \*dec1) [Target Hook]
  This hook is used to get the dispatcher function for a set of function versions. The dispatcher function is called to invoke the right function version at run-time. decl is one version from a set of semantically identical versions.
- tree TARGET\_GENERATE\_VERSION\_DISPATCHER\_BODY (void \*arg) [Target Hook]
  This hook is used to generate the dispatcher logic to invoke the right function version at run-time for a given set of function versions. arg points to the callgraph node of the dispatcher function whose body must be generated.

### bool TARGET\_PREDICT\_DOLOOP\_P (class loop \*loop)

[Target Hook]

Return true if we can predict it is possible to use a low-overhead loop for a particular loop. The parameter *loop* is a pointer to the loop. This target hook is required only when the target supports low-overhead loops, and will help ivopts to make some decisions. The default version of this hook returns false.

### bool TARGET\_HAVE\_COUNT\_REG\_DECR\_P

[Target Hook]

Return true if the target supports hardware count register for decrement and branch. The default value is false.

### int64\_t TARGET\_DOLOOP\_COST\_FOR\_GENERIC

[Target Hook]

One IV candidate dedicated for doloop is introduced in IVOPTs, we can calculate the computation cost of adopting it to any generic IV use by function get\_computation\_cost as before. But for targets which have hardware count register support for decrement and branch, it may have to move IV value from hardware count register to general purpose register while doloop IV candidate is used for generic IV uses. It probably takes expensive penalty. This hook allows target owners to define the cost for this especially for generic IV uses. The default value is zero.

### int64\_t TARGET\_DOLOOP\_COST\_FOR\_ADDRESS

[Target Hook]

One IV candidate dedicated for doloop is introduced in IVOPTs, we can calculate the computation cost of adopting it to any address IV use by function get\_computation\_cost as before. But for targets which have hardware count register support for decrement and branch, it may have to move IV value from hardware count register to general purpose register while doloop IV candidate is used for address IV uses. It probably takes expensive penalty. This hook allows target owners to define the cost for this escpecially for address IV uses. The default value is zero.

# bool TARGET\_CAN\_USE\_DOLOOP\_P (const widest\_int &iterations, const widest\_int &iterations\_max, unsigned int loop\_depth, bool entered\_at\_top) [Target Hook]

Return true if it is possible to use low-overhead loops (doloop\_end and doloop\_begin) for a particular loop. *iterations* gives the exact number of iterations, or 0 if not known. *iterations\_max* gives the maximum number of iterations, or 0 if not known. *loop\_depth* is the nesting depth of the loop, with 1 for innermost loops, 2 for loops that contain innermost loops, and so on. *entered\_at\_top* is true if the loop is only entered from the top.

This hook is only used if doloop\_end is available. The default implementation returns true. You can use can\_use\_doloop\_if\_innermost if the loop must be the innermost, and if there are no other restrictions.

#### 

Take an instruction in *insn* and return NULL if it is valid within a low-overhead loop, otherwise return a string explaining why doloop could not be applied.

Many targets use special registers for low-overhead looping. For any instruction that clobbers these this function should return a string indicating the reason why the doloop could not be applied. By default, the RTL loop optimizer does not use a

present doloop pattern for loops containing function calls or branch on table instructions.

- bool TARGET\_LEGITIMATE\_COMBINED\_INSN (rtx\_insn \*insn) [Target Hook]

  Take an instruction in insn and return false if the instruction is not appropriate as a combination of two or more instructions. The default is to accept all instructions.
- bool TARGET\_CAN\_FOLLOW\_JUMP (const rtx\_insn \*follower, const rtx\_insn \*followee) [Target Hook]

FOLLOWER and FOLLOWEE are JUMP\_INSN instructions; return true if FOLLOWER may be modified to follow FOLLOWEE; false, if it can't. For example, on some targets, certain kinds of branches can't be made to follow through a hot/cold partitioning.

- bool TARGET\_COMMUTATIVE\_P (const\_rtx x, int outer\_code) [Target Hook]
  This target hook returns true if x is considered to be commutative. Usually, this is just COMMUTATIVE\_P (x), but the HP PA doesn't consider PLUS to be commutative inside a MEM. outer\_code is the rtx code of the enclosing rtl, if known, otherwise it is UNKNOWN.
- TARGET\_ALLOCATE\_INITIAL\_VALUE (rtx hard\_reg) [Target Hook] When the initial value of a hard register has been copied in a pseudo register, it is often not necessary to actually allocate another register to this pseudo register, because the original hard register or a stack slot it has been saved into can be used. TARGET\_ALLOCATE\_INITIAL\_VALUE is called at the start of register allocation once for each hard register that had its initial value copied by using get\_func\_hard\_reg\_initial\_val or get\_hard\_reg\_initial\_val. Possible values are NULL\_RTX, if you don't want to do any special allocation, a REG rtx—that would typically be the hard register itself, if it is known not to be clobbered—or a MEM. If you are returning a MEM, this is only a hint for the allocator; it might decide to use another register anyways. You may use current\_function\_is\_leaf or REG\_N\_SETS in the hook to determine if the hard register in question will not be clobbered. The default value of this hook is NULL, which disables any special allocation.
- int TARGET\_UNSPEC\_MAY\_TRAP\_P (const\_rtx x, unsigned flags) [Target Hook] This target hook returns nonzero if x, an unspec or unspec\_volatile operation, might cause a trap. Targets can use this hook to enhance precision of analysis for unspec and unspec\_volatile operations. You may call may\_trap\_p\_1 to analyze inner elements of x in which case flags should be passed along.
- void TARGET\_SET\_CURRENT\_FUNCTION (tree dec1) [Target Hook]

  The compiler invokes this hook whenever it changes its current function context (cfun). You can define this function if the back end needs to perform any initialization or reset actions on a per-function basis. For example, it may be used to implement function attributes that affect register usage or code generation patterns. The argument dec1 is the declaration for the new function context, and may be null to indicate that the compiler has left a function context and is returning to processing at the top level. The default hook function does nothing.

GCC sets cfun to a dummy function context during initialization of some parts of the back end. The hook function is not invoked in this situation; you need not worry about the hook being invoked recursively, or when the back end is in a partiallyinitialized state. cfun might be NULL to indicate processing at top level, outside of any function scope.

### TARGET\_OBJECT\_SUFFIX

[Macro]

Define this macro to be a C string representing the suffix for object files on your target machine. If you do not define this macro, GCC will use '.o' as the suffix for object files.

### TARGET\_EXECUTABLE\_SUFFIX

[Macro]

Define this macro to be a C string representing the suffix to be automatically added to executable files on your target machine. If you do not define this macro, GCC will use the null string as the suffix for executable files.

### COLLECT\_EXPORT\_LIST

[Macro]

If defined, collect2 will scan the individual object files specified on its command line and create an export list for the linker. Define this macro for systems like AIX, where the linker discards object files that are not referenced from main and uses export lists.

### bool TARGET\_CANNOT\_MODIFY\_JUMPS\_P (void)

[Target Hook]

This target hook returns **true** past the point in which new jump instructions could be created. On machines that require a register for every jump such as the SHmedia ISA of SH5, this point would typically be reload, so this target hook should be defined to a function such as:

```
static bool
cannot_modify_jumps_past_reload_p ()
{
   return (reload_completed || reload_in_progress);
}
```

### bool TARGET\_HAVE\_CONDITIONAL\_EXECUTION (void)

[Target Hook]

This target hook returns true if the target supports conditional execution. This target hook is required only when the target has several different modes and they have different conditional execution capability, such as ARM.

# rtx TARGET\_GEN\_CCMP\_FIRST (rtx\_insn \*\*prep\_seq, rtx\_insn [Target Hook] \*\*gen\_seq, int code, tree op0, tree op1)

This function prepares to emit a comparison insn for the first compare in a sequence of conditional comparisions. It returns an appropriate comparison with CC for passing to gen\_ccmp\_next or cbranch\_optab. The insns to prepare the compare are saved in  $prep\_seq$  and the compare insns are saved in  $gen\_seq$ . They will be emitted when all the compares in the conditional comparision are generated without error. code is the  $rtx\_code$  of the compare for op0 and op1.

```
rtx TARGET_GEN_CCMP_NEXT (rtx_insn **prep_seq, rtx_insn [Target Hook] 
**gen_seq, rtx prev, int cmp_code, tree op0, tree op1, int bit_code)
```

This function prepares to emit a conditional comparison within a sequence of conditional comparisons. It returns an appropriate comparison with CC for passing to

gen\_ccmp\_next or cbranch\_optab. The insns to prepare the compare are saved in prep\_seq and the compare insns are saved in gen\_seq. They will be emitted when all the compares in the conditional comparision are generated without error. The prev expression is the result of a prior call to gen\_ccmp\_first or gen\_ccmp\_next. It may return NULL if the combination of prev and this comparison is not supported, otherwise the result must be appropriate for passing to gen\_ccmp\_next or cbranch\_optab. code is the rtx\_code of the compare for op0 and op1. bit\_code is AND or IOR, which is the op on the compares.

# unsigned TARGET\_LOOP\_UNROLL\_ADJUST (unsigned nunrol1, class [Target Hook] loop \*loop)

This target hook returns a new value for the number of times *loop* should be unrolled. The parameter *nunroll* is the number of times the loop is to be unrolled. The parameter *loop* is a pointer to the loop, which is going to be checked for unrolling. This target hook is required only when the target has special constraints like maximum number of memory accesses.

POWI\_MAX\_MULTS [Macro]

If defined, this macro is interpreted as a signed integer C expression that specifies the maximum number of floating point multiplications that should be emitted when expanding exponentiation by an integer constant inline. When this value is defined, exponentiation requiring more than this number of multiplications is implemented by calling the system library's pow, powf or powl routines. The default value places no upper bound on the multiplication count.

#### 

This target hook should register any extra include files for the target. The parameter *stdinc* indicates if normal include files are present. The parameter *sysroot* is the system root directory. The parameter *iprefix* is the prefix for the gcc directory.

#### 

This target hook should register any extra include files for the target before any standard headers. The parameter *stdinc* indicates if normal include files are present. The parameter *sysroot* is the system root directory. The parameter *iprefix* is the prefix for the gcc directory.

### void TARGET\_OPTF (char \*path)

[Macro]

This target hook should register special include paths for the target. The parameter path is the include to register. On Darwin systems, this is used for Framework includes, which have semantics that are different from '-I'.

### bool TARGET\_USE\_LOCAL\_THUNK\_ALIAS\_P (tree fndec1) [Macro]

This target macro returns true if it is safe to use a local alias for a virtual function findecl when constructing thunks, false otherwise. By default, the macro returns true for all functions, if a target supports aliases (i.e. defines ASM\_OUTPUT\_DEF), false otherwise,

### TARGET\_FORMAT\_TYPES

[Macro]

If defined, this macro is the name of a global variable containing target-specific format checking information for the '-Wformat' option. The default is to have no target-specific format checks.

### TARGET\_N\_FORMAT\_TYPES

[Macro]

If defined, this macro is the number of entries in TARGET\_FORMAT\_TYPES.

### TARGET\_OVERRIDES\_FORMAT\_ATTRIBUTES

[Macro]

If defined, this macro is the name of a global variable containing target-specific format overrides for the '-Wformat' option. The default is to have no target-specific format overrides. If defined, TARGET\_FORMAT\_TYPES must be defined, too.

### TARGET\_OVERRIDES\_FORMAT\_ATTRIBUTES\_COUNT

[Macro]

If defined, this macro specifies the number of entries in TARGET\_OVERRIDES\_FORMAT\_ATTRIBUTES.

### TARGET\_OVERRIDES\_FORMAT\_INIT

[Macro]

If defined, this macro specifies the optional initialization routine for target specific customizations of the system printf and scanf formatter settings.

### const char \* TARGET\_INVALID\_ARG\_FOR\_UNPROTOTYPED\_FN

[Target Hook]

(const\_tree typelist, const\_tree funcdec1, const\_tree val)

If defined, this macro returns the diagnostic message when it is illegal to pass argument val to function functed with prototype typelist.

# const char \* TARGET\_INVALID\_CONVERSION (const\_tree

[Target Hook]

fromtype, const\_tree totype)

If defined, this macro returns the diagnostic message when it is invalid to convert from from type to totype, or NULL if validity should be determined by the front end.

### 

[Target Hook]

If defined, this macro returns the diagnostic message when it is invalid to apply operation op (where unary plus is denoted by CONVERT\_EXPR) to an operand of type type, or NULL if validity should be determined by the front end.

# 

[Target Hook]

type1, const\_tree type2)

If defined, this macro returns the diagnostic message when it is invalid to apply operation op to operands of types type1 and type2, or NULL if validity should be determined by the front end.

### tree TARGET\_PROMOTED\_TYPE (const\_tree type)

[Target Hook]

If defined, this target hook returns the type to which values of *type* should be promoted when they appear in expressions, analogous to the integer promotions, or NULL\_TREE to use the front end's normal promotion rules. This hook is useful when there are target-specific types with special promotion rules. This is currently used only by the C and C++ front ends.

### tree TARGET\_CONVERT\_TO\_TYPE (tree type, tree expr)

[Target Hook]

If defined, this hook returns the result of converting *expr* to *type*. It should return the converted expression, or NULL\_TREE to apply the front end's normal conversion rules. This hook is useful when there are target-specific types with special conversion rules. This is currently used only by the C and C++ front ends.

### bool TARGET\_VERIFY\_TYPE\_CONTEXT (location\_t loc,

[Target Hook]

type\_context\_kind context, const\_tree type, bool silent\_p)

If defined, this hook returns false if there is a target-specific reason why type type cannot be used in the source language context described by context. When silent\_p is false, the hook also reports an error against loc for invalid uses of type.

Calls to this hook should be made through the global function verify\_type\_context, which makes the *silent\_p* parameter default to false and also handles error\_mark\_node.

The default implementation always returns true.

OBJC\_JBLEN

[Macro]

This macro determines the size of the objective C jump buffer for the NeXT runtime. By default, OBJC\_JBLEN is defined to an innocuous value.

### LIBGCC2\_UNWIND\_ATTRIBUTE

[Macro]

Define this macro if any target-specific attributes need to be attached to the functions in 'libgcc' that provide low-level support for call stack unwinding. It is used in declarations in 'unwind-generic.h' and the associated definitions of those functions.

### void TARGET\_UPDATE\_STACK\_BOUNDARY (void)

[Target Hook]

Define this macro to update the current function stack boundary if necessary.

### rtx TARGET\_GET\_DRAP\_RTX (void)

[Target Hook]

This hook should return an rtx for Dynamic Realign Argument Pointer (DRAP) if a different argument pointer register is needed to access the function's argument list due to stack realignment. Return NULL if no DRAP is needed.

### bool TARGET\_ALLOCATE\_STACK\_SLOTS\_FOR\_ARGS (void)

[Target Hook]

When optimization is disabled, this hook indicates whether or not arguments should be allocated to stack slots. Normally, GCC allocates stacks slots for arguments when not optimizing in order to make debugging easier. However, when a function is declared with <code>\_\_attribute\_\_((naked))</code>, there is no stack frame, and the compiler cannot safely move arguments from the registers in which they are passed to the stack. Therefore, this hook should return true in general, but false for naked functions. The default implementation always returns true.

### unsigned HOST\_WIDE\_INT TARGET\_CONST\_ANCHOR

[Target Hook]

On some architectures it can take multiple instructions to synthesize a constant. If there is another constant already in a register that is close enough in value then it is preferable that the new constant is computed from this register using immediate addition or subtraction. We accomplish this through CSE. Besides the value of the constant we also add a lower and an upper constant anchor to the available expressions. These are then queried when encountering new constants. The anchors are computed by rounding the constant up and down to a multiple of the value of TARGET\_CONST\_ANCHOR. TARGET\_CONST\_ANCHOR should be the maximum positive value accepted by immediate-add plus one. We currently assume that the value of TARGET\_CONST\_ANCHOR is a power of 2. For example, on MIPS, where add-immediate takes a 16-bit signed value, TARGET\_CONST\_ANCHOR is set to '0x8000'. The default value is zero, which disables this optimization.

unsigned HOST\_WIDE\_INT TARGET\_ASAN\_SHADOW\_OFFSET (void) [Target Hook] Return the offset bitwise ored into shifted address to get corresponding Address Sanitizer shadow memory address. NULL if Address Sanitizer is not supported by the target.

# unsigned HOST\_WIDE\_INT TARGET\_MEMMODEL\_CHECK (unsigned HOST\_WIDE\_INT val) [Target Hook]

Validate target specific memory model mask bits. When NULL no target specific memory model bits are allowed.

# unsigned char TARGET\_ATOMIC\_TEST\_AND\_SET\_TRUEVAL [Target Hook] This value should be set if the result written by atomic\_test\_and\_set is not exactly 1, i.e. the bool true.

### bool TARGET\_HAS\_IFUNC\_P (void)

[Target Hook]

It returns true if the target supports GNU indirect functions. The support includes the assembler, linker and dynamic linker. The default value of this hook is based on target's libc.

#### 

If defined, this function returns an appropriate alignment in bits for an atomic object of machine\_mode *mode*. If 0 is returned then the default alignment for the specified mode is used.

#### 

ISO C11 requires atomic compound assignments that may raise floating-point exceptions to raise exceptions corresponding to the arithmetic operation whose result was successfully stored in a compare-and-exchange sequence. This requires code equivalent to calls to feholdexcept, feclearexcept and feupdateenv to be generated at appropriate points in the compare-and-exchange sequence. This hook should set \*hold to an expression equivalent to the call to feholdexcept, \*clear to an expression equivalent to the call to feclearexcept and \*update to an expression equivalent to the call to feupdateenv. The three expressions are NULL\_TREE on entry to the hook and may be left as NULL\_TREE if no code is required in a particular place. The default implementation leaves all three expressions as NULL\_TREE. The \_\_atomic\_feraiseexcept function from libatomic may be of use as part of the code generated in \*update.

### void TARGET\_RECORD\_OFFLOAD\_SYMBOL (tree)

[Target Hook]

Used when offloaded functions are seen in the compilation unit and no named sections are available. It is called once for each symbol that must be recorded in the offload function and variable table.

### char \* TARGET\_OFFLOAD\_OPTIONS (void)

[Target Hook]

Used when writing out the list of options into an LTO file. It should translate any relevant target-specific options (such as the ABI in use) into one of the '-foffload' options that exist as a common interface to express such options. It should return a string containing these options, separated by spaces, which the caller will free.

### TARGET\_SUPPORTS\_WIDE\_INT

[Macro]

On older ports, large integers are stored in CONST\_DOUBLE rtl objects. Newer ports define TARGET\_SUPPORTS\_WIDE\_INT to be nonzero to indicate that large integers are stored in CONST\_WIDE\_INT rtl objects. The CONST\_WIDE\_INT allows very large integer constants to be represented. CONST\_DOUBLE is limited to twice the size of the host's HOST\_WIDE\_INT representation.

Converting a port mostly requires looking for the places where CONST\_DOUBLES are used with VOIDmode and replacing that code with code that accesses CONST\_WIDE\_INTs. '"grep -i const\_double"' at the port level gets you to 95% of the changes that need to be made. There are a few places that require a deeper look.

- There is no equivalent to hval and lval for CONST\_WIDE\_INTs. This would be difficult to express in the md language since there are a variable number of elements.
  - Most ports only check that hval is either 0 or -1 to see if the value is small. As mentioned above, this will no longer be necessary since small constants are always CONST\_INT. Of course there are still a few exceptions, the alpha's constraint used by the zap instruction certainly requires careful examination by C code. However, all the current code does is pass the hval and lval to C code, so evolving the c code to look at the CONST\_WIDE\_INT is not really a large change.
- Because there is no standard template that ports use to materialize constants, there is likely to be some futzing that is unique to each port in this code.
- The rtx costs may have to be adjusted to properly account for larger constants that are represented as CONST\_WIDE\_INT.

All and all it does not take long to convert ports that the maintainer is familiar with.

### bool TARGET\_HAVE\_SPECULATION\_SAFE\_VALUE (bool active)

[Target Hook]

This hook is used to determine the level of target support for \_\_builtin\_speculation\_safe\_value. If called with an argument of false, it returns true if the target has been modified to support this builtin. If called with an argument of true, it returns true if the target requires active mitigation execution might be speculative. The default implementation returns false if the target does not define a pattern named speculation\_barrier. Else it returns true for the first case and whether the pattern is enabled for the current compilation for the second case.

For targets that have no processors that can execute instructions speculatively an alternative implemenation of this hook is available: simply redefine this hook to speculation\_safe\_value\_not\_needed along with your other target hooks.

# rtx TARGET\_SPECULATION\_SAFE\_VALUE (machine\_mode mode, rtx result, rtx val, rtx failval) [Target Hook]

This target hook can be used to generate a target-specific code sequence that implements the \_\_builtin\_speculation\_safe\_value built-in function. The function must always return val in result in mode mode when the cpu is not executing speculatively, but must never return that when speculating until it is known that the speculation will not be unwound. The hook supports two primary mechanisms for implementing the requirements. The first is to emit a speculation barrier which forces the processor to wait until all prior speculative operations have been resolved; the second is to use a target-specific mechanism that can track the speculation state and to return failval if it can determine that speculation must be unwound at a later time.

The default implementation simply copies val to result and emits a speculation\_barrier instruction if that is defined.

### void TARGET\_RUN\_TARGET\_SELFTESTS (void)

[Target Hook]

If selftests are enabled, run any selftests for this target.

## 19 Host Configuration

Most details about the machine and system on which the compiler is actually running are detected by the configure script. Some things are impossible for configure to detect; these are described in two ways, either by macros defined in a file named 'xm-machine.h' or by hook functions in the file specified by the out\_host\_hook\_obj variable in 'config.gcc'. (The intention is that very few hosts will need a header file but nearly every fully supported host will need to override some hooks.)

If you need to define only a few macros, and they have simple definitions, consider using the xm\_defines variable in your 'config.gcc' entry instead of creating a host configuration header. See Section 6.3.2.2 [System Config], page 65.

### 19.1 Host Common

Some things are just not portable, even between similar operating systems, and are too difficult for autoconf to detect. They get implemented using hook functions in the file specified by the *host\_hook\_obj* variable in 'config.gcc'.

### void HOST\_HOOKS\_EXTRA\_SIGNALS (void)

[Host Hook]

This host hook is used to set up handling for extra signals. The most common thing to do in this hook is to detect stack overflow.

void \* HOST\_HOOKS\_GT\_PCH\_GET\_ADDRESS (size\_t size, int fd) [Host Hook]
This host hook returns the address of some space that is likely to be free in some subsequent invocation of the compiler. We intend to load the PCH data at this address such that the data need not be relocated. The area should be able to hold size bytes. If the host uses mmap, fd is an open file descriptor that can be used for probing.

#### 

This host hook is called when a PCH file is about to be loaded. We want to load size bytes from fd at offset into memory at address. The given address will be the result of a previous invocation of HOST\_HOOKS\_GT\_PCH\_GET\_ADDRESS. Return -1 if we couldn't allocate size bytes at address. Return 0 if the memory is allocated but the data is not loaded. Return 1 if the hook has performed everything.

If the implementation uses reserved address space, free any reserved space beyond size, regardless of the return value. If no PCH will be loaded, this hook may be called with size zero, in which case all reserved address space should be freed.

Do not try to handle values of address that could not have been returned by this executable; just return -1. Such values usually indicate an out-of-date PCH file (built by some other GCC executable), and such a PCH file won't work.

### size\_t HOST\_HOOKS\_GT\_PCH\_ALLOC\_GRANULARITY (void); [Host Hool

This host hook returns the alignment required for allocating virtual memory. Usually this is the same as getpagesize, but on some hosts the alignment for reserving memory differs from the pagesize for committing memory.

### 19.2 Host Filesystem

GCC needs to know a number of things about the semantics of the host machine's filesystem. Filesystems with Unix and MS-DOS semantics are automatically detected. For other systems, you can define the following macros in 'xm-machine.h'.

### HAVE\_DOS\_BASED\_FILE\_SYSTEM

This macro is automatically defined by 'system.h' if the host file system obeys the semantics defined by MS-DOS instead of Unix. DOS file systems are case insensitive, file specifications may begin with a drive letter, and both forward slash and backslash ('/' and '\') are directory separators.

### DIR\_SEPARATOR

### DIR\_SEPARATOR\_2

If defined, these macros expand to character constants specifying separators for directory names within a file specification. 'system.h' will automatically give them appropriate values on Unix and MS-DOS file systems. If your file system is neither of these, define one or both appropriately in 'xm-machine.h'.

However, operating systems like VMS, where constructing a pathname is more complicated than just stringing together directory names separated by a special character, should not define either of these macros.

### PATH\_SEPARATOR

If defined, this macro should expand to a character constant specifying the separator for elements of search paths. The default value is a colon (':'). DOS-based systems usually, but not always, use semicolon (';').

VMS Define this macro if the host system is VMS.

### HOST\_OBJECT\_SUFFIX

Define this macro to be a C string representing the suffix for object files on your host machine. If you do not define this macro, GCC will use '.o' as the suffix for object files.

### HOST\_EXECUTABLE\_SUFFIX

Define this macro to be a C string representing the suffix for executable files on your host machine. If you do not define this macro, GCC will use the null string as the suffix for executable files.

### HOST\_BIT\_BUCKET

A pathname defined by the host operating system, which can be opened as a file and written to, but all the information written is discarded. This is commonly known as a bit bucket or null device. If you do not define this macro, GCC will use '/dev/null' as the bit bucket. If the host does not support a bit bucket, define this macro to an invalid filename.

### UPDATE\_PATH\_HOST\_CANONICALIZE (path)

If defined, a C statement (sans semicolon) that performs host-dependent canonicalization when a path used in a compilation driver or preprocessor is canonicalized. path is a malloc-ed path to be canonicalized. If the C statement does canonicalize path into a different buffer, the old path should be freed and the new buffer should have been allocated with malloc.

### DUMPFILE\_FORMAT

Define this macro to be a C string representing the format to use for constructing the index part of debugging dump file names. The resultant string must fit in fifteen bytes. The full filename will be the concatenation of: the prefix of the assembler file name, the string resulting from applying this format to an index number, and a string unique to each dump file kind, e.g. 'rtl'.

If you do not define this macro, GCC will use '.%02d.'. You should define this macro if using the default will create an invalid file name.

#### DELETE\_IF\_ORDINARY

Define this macro to be a C statement (sans semicolon) that performs host-dependent removal of ordinary temp files in the compilation driver.

If you do not define this macro, GCC will use the default version. You should define this macro if the default version does not reliably remove the temp file as, for example, on VMS which allows multiple versions of a file.

### HOST\_LACKS\_INODE\_NUMBERS

Define this macro if the host filesystem does not report meaningful inode numbers in struct stat.

### 19.3 Host Misc

### FATAL\_EXIT\_CODE

A C expression for the status code to be returned when the compiler exits after serious errors. The default is the system-provided macro 'EXIT\_FAILURE', or '1' if the system doesn't define that macro. Define this macro only if these defaults are incorrect.

### SUCCESS\_EXIT\_CODE

A C expression for the status code to be returned when the compiler exits without serious errors. (Warnings are not serious errors.) The default is the system-provided macro 'EXIT\_SUCCESS', or '0' if the system doesn't define that macro. Define this macro only if these defaults are incorrect.

### USE\_C\_ALLOCA

Define this macro if GCC should use the C implementation of alloca provided by 'libiberty.a'. This only affects how some parts of the compiler itself allocate memory. It does not change code generation.

When GCC is built with a compiler other than itself, the C alloca is always used. This is because most other implementations have serious bugs. You should define this macro only on a system where no stack-based alloca can possibly work. For instance, if a system has a small limit on the size of the stack, GCC's builtin alloca will not work reliably.

### COLLECT2\_HOST\_INITIALIZATION

If defined, a C statement (sans semicolon) that performs host-dependent initialization when collect2 is being initialized.

### GCC\_DRIVER\_HOST\_INITIALIZATION

If defined, a C statement (sans semicolon) that performs host-dependent initialization when a compilation driver is being initialized.

### HOST\_LONG\_LONG\_FORMAT

If defined, the string used to indicate an argument of type long long to functions like printf. The default value is "ll".

### HOST\_LONG\_FORMAT

If defined, the string used to indicate an argument of type long to functions like printf. The default value is "1".

### HOST\_PTR\_PRINTF

If defined, the string used to indicate an argument of type void \* to functions like printf. The default value is "%p".

In addition, if configure generates an incorrect definition of any of the macros in 'auto-host.h', you can override that definition in a host configuration header. If you need to do this, first see if it is possible to fix configure.

# 20 Makefile Fragments

When you configure GCC using the 'configure' script, it will construct the file 'Makefile' from the template file 'Makefile.in'. When it does this, it can incorporate makefile fragments from the 'config' directory. These are used to set Makefile parameters that are not amenable to being calculated by autoconf. The list of fragments to incorporate is set by 'config.gcc' (and occasionally 'config.build' and 'config.host'); See Section 6.3.2.2 [System Config], page 65.

Fragments are named either 't-target' or 'x-host', depending on whether they are relevant to configuring GCC to produce code for a particular target, or to configuring GCC to run on a particular host. Here target and host are mnemonics which usually have some relationship to the canonical system name, but no formal connection.

If these files do not exist, it means nothing needs to be added for a given target or host. Most targets need a few 't-target' fragments, but needing 'x-host' fragments is rare.

### 20.1 Target Makefile Fragments

Target makefile fragments can set these Makefile variables.

### LIBGCC2\_CFLAGS

Compiler flags to use when compiling 'libgcc2.c'.

#### LIB2FUNCS\_EXTRA

A list of source file names to be compiled or assembled and inserted into 'libgcc.a'.

### CRTSTUFF\_T\_CFLAGS

Special flags used when compiling 'crtstuff.c'. See Section 18.20.5 [Initialization], page 612.

### CRTSTUFF\_T\_CFLAGS\_S

Special flags used when compiling 'crtstuff.c' for shared linking. Used if you use 'crtbeginS.o' and 'crtendS.o' in EXTRA-PARTS. See Section 18.20.5 [Initialization], page 612.

### MULTILIB\_OPTIONS

For some targets, invoking GCC in different ways produces objects that cannot be linked together. For example, for some targets GCC produces both big and little endian code. For these targets, you must arrange for multiple versions of 'libgcc.a' to be compiled, one for each set of incompatible options. When GCC invokes the linker, it arranges to link in the right version of 'libgcc.a', based on the command line options used.

The MULTILIB\_OPTIONS macro lists the set of options for which special versions of 'libgcc.a' must be built. Write options that are mutually incompatible side by side, separated by a slash. Write options that may be used together separated by a space. The build procedure will build all combinations of compatible options.

For example, if you set MULTILIB\_OPTIONS to 'm68000/m68020 msoft-float', 'Makefile' will build special versions of 'libgcc.a' using the following sets of

options: '-m68000', '-m68020', '-msoft-float', '-m68000 -msoft-float', and '-m68020 -msoft-float'.

### MULTILIB\_DIRNAMES

If MULTILIB\_OPTIONS is used, this variable specifies the directory names that should be used to hold the various libraries. Write one element in MULTILIB\_DIRNAMES for each element in MULTILIB\_OPTIONS. If MULTILIB\_DIRNAMES is not used, the default value will be MULTILIB\_OPTIONS, with all slashes treated as spaces.

MULTILIB\_DIRNAMES describes the multilib directories using GCC conventions and is applied to directories that are part of the GCC installation. When multilib-enabled, the compiler will add a subdirectory of the form prefix/multilib before each directory in the search path for libraries and crt files.

For example, if MULTILIB\_OPTIONS is set to 'm68000/m68020 msoft-float', then the default value of MULTILIB\_DIRNAMES is 'm68000 m68020 msoft-float'. You may specify a different value if you desire a different set of directory names.

### MULTILIB\_MATCHES

Sometimes the same option may be written in two different ways. If an option is listed in MULTILIB\_OPTIONS, GCC needs to know about any synonyms. In that case, set MULTILIB\_MATCHES to a list of items of the form 'option=option' to describe all relevant synonyms. For example, 'm68000=mc68000 m68020=mc68020'.

#### MULTILIB\_EXCEPTIONS

Sometimes when there are multiple sets of MULTILIB\_OPTIONS being specified, there are combinations that should not be built. In that case, set MULTILIB\_EXCEPTIONS to be all of the switch exceptions in shell case syntax that should not be built.

For example the ARM processor cannot execute both hardware floating point instructions and the reduced size THUMB instructions at the same time, so there is no need to build libraries with both of these options enabled. Therefore MULTILIB\_EXCEPTIONS is set to:

\*mthumb/\*mhard-float\*

### MULTILIB\_REQUIRED

Sometimes when there are only a few combinations are required, it would be a big effort to come up with a MULTILIB\_EXCEPTIONS list to cover all undesired ones. In such a case, just listing all the required combinations in MULTILIB\_REQUIRED would be more straightforward.

The way to specify the entries in MULTILIB\_REQUIRED is same with the way used for MULTILIB\_EXCEPTIONS, only this time what are required will be specified. Suppose there are multiple sets of MULTILIB\_OPTIONS and only two combinations are required, one for ARMv7-M and one for ARMv7-R with hard floating-point ABI and FPU, the MULTILIB\_REQUIRED can be set to:

```
MULTILIB_REQUIRED = mthumb/march=armv7-m

MULTILIB_REQUIRED += march=armv7-r/mfloat-abi=hard/mfpu=vfpv3-d16
```

The MULTILIB\_REQUIRED can be used together with MULTILIB\_EXCEPTIONS. The option combinations generated from MULTILIB\_OPTIONS will be filtered by MULTILIB\_EXCEPTIONS and then by MULTILIB\_REQUIRED.

#### MULTILIB\_REUSE

Sometimes it is desirable to reuse one existing multilib for different sets of options. Such kind of reuse can minimize the number of multilib variants. And for some targets it is better to reuse an existing multilib than to fall back to default multilib when there is no corresponding multilib. This can be done by adding reuse rules to MULTILIB\_REUSE.

A reuse rule is comprised of two parts connected by equality sign. The left part is the option set used to build multilib and the right part is the option set that will reuse this multilib. Both parts should only use options specified in MULTILIB\_OPTIONS and the equality signs found in options name should be replaced with periods. An explicit period in the rule can be escaped by preceding it with a backslash. The order of options in the left part matters and should be same with those specified in MULTILIB\_REQUIRED or aligned with the order in MULTILIB\_OPTIONS. There is no such limitation for options in the right part as we don't build multilib from them.

MULTILIB\_REUSE is different from MULTILIB\_MATCHES in that it sets up relations between two option sets rather than two options. Here is an example to demo how we reuse libraries built in Thumb mode for applications built in ARM mode:

MULTILIB\_REUSE = mthumb/march.armv7-r=marm/march.armv7-r

Before the advent of MULTILIB\_REUSE, GCC select multilib by comparing command line options with options used to build multilib. The MULTILIB\_REUSE is complementary to that way. Only when the original comparison matches nothing it will work to see if it is OK to reuse some existing multilib.

#### MULTILIB\_EXTRA\_OPTS

Sometimes it is desirable that when building multiple versions of 'libgcc.a' certain options should always be passed on to the compiler. In that case, set MULTILIB\_EXTRA\_OPTS to be the list of options to be used for all builds. If you set this, you should probably set CRTSTUFF\_T\_CFLAGS to a dash followed by it.

#### MULTILIB OSDIRNAMES

If MULTILIB\_OPTIONS is used, this variable specifies a list of subdirectory names, that are used to modify the search path depending on the chosen multilib. Unlike MULTILIB\_DIRNAMES, MULTILIB\_OSDIRNAMES describes the multilib directories using operating systems conventions, and is applied to the directories such as lib or those in the LIBRARY\_PATH environment variable. The format is either the same as of MULTILIB\_DIRNAMES, or a set of mappings. When it is the same as MULTILIB\_DIRNAMES, it describes the multilib directories using operating system conventions, rather than GCC conventions. When it is a set of mappings of the form gccdir=osdir, the left side gives the GCC convention and the right gives the equivalent OS defined location. If the osdir part begins with a '!', GCC will not search in the non-multilib directory and use exclusively the multilib directory. Otherwise, the compiler will examine the search path for

libraries and crt files twice; the first time it will add multilib to each directory in the search path, the second it will not.

For configurations that support both multilib and multiarch, MULTILIB\_OSDIRNAMES also encodes the multiarch name, thus subsuming MULTIARCH\_DIRNAME. The multiarch name is appended to each directory name, separated by a colon (e.g. '../lib32:i386-linux-gnu').

Each multiarch subdirectory will be searched before the corresponding OS multilib directory, for example '/lib/i386-linux-gnu' before '/lib/../lib32'. The multiarch name will also be used to modify the system header search path, as explained for MULTIARCH\_DIRNAME.

#### MULTIARCH\_DIRNAME

This variable specifies the multiarch name for configurations that are multiarchenabled but not multilibbed configurations.

The multiarch name is used to augment the search path for libraries, crt files and system header files with additional locations. The compiler will add a multiarch subdirectory of the form prefix/multiarch before each directory in the library and crt search path. It will also add two directories LOCAL\_INCLUDE\_DIR/multiarch and NATIVE\_SYSTEM\_HEADER\_DIR/multiarch) to the system header search path, respectively before LOCAL\_INCLUDE\_DIR and NATIVE\_SYSTEM\_HEADER\_DIR.

MULTIARCH\_DIRNAME is not used for configurations that support both multilib and multiarch. In that case, multiarch names are encoded in MULTILIB\_OSDIRNAMES instead.

More documentation about multiarch can be found at https://wiki.debian.org/Multiarch.

Unfortunately, setting MULTILIB\_EXTRA\_OPTS is not enough, since it does not affect the build of target libraries, at least not the build of the default multilib. One possible work-around is to use DRIVER\_SELF\_SPECS to bring options from the 'specs' file as if they had been passed in the compiler driver command line. However, you don't want to be adding these options after the toolchain is installed, so you can instead tweak the 'specs' file that will be used during the toolchain build, while you still install the original, built-in 'specs'. The trick is to set SPECS to some other filename (say 'specs.install'), that will then be created out of the built-in specs, and introduce a 'Makefile' rule to generate the 'specs' file that's going to be used at build time out of your 'specs.install'.

T\_CFLAGS These are extra flags to pass to the C compiler. They are used both when building GCC, and when compiling things with the just-built GCC. This variable is deprecated and should not be used.

### 20.2 Host Makefile Fragments

The use of 'x-host' fragments is discouraged. You should only use it for makefile dependencies.

### 21 collect2

GCC uses a utility called collect2 on nearly all systems to arrange to call various initialization functions at start time.

The program collect2 works by linking the program once and looking through the linker output file for symbols with particular names indicating they are constructor functions. If it finds any, it creates a new temporary '.c' file containing a table of them, compiles it, and links the program a second time including that file.

The actual calls to the constructors are carried out by a subroutine called \_\_main, which is called (automatically) at the beginning of the body of main (provided main was compiled with GNU CC). Calling \_\_main is necessary, even when compiling C code, to allow linking C and C++ object code together. (If you use '-nostdlib', you get an unresolved reference to \_\_main, since it's defined in the standard GCC library. Include '-lgcc' at the end of your compiler command line to resolve this reference.)

The program collect2 is installed as ld in the directory where the passes of the compiler are installed. When collect2 needs to find the real ld, it tries the following file names:

- a hard coded linker file name, if GCC was configured with the '--with-ld' option.
- 'real-ld' in the directories listed in the compiler's search directories.
- 'real-ld' in the directories listed in the environment variable PATH.
- The file specified in the REAL\_LD\_FILE\_NAME configuration macro, if specified.
- 'ld' in the compiler's search directories, except that collect2 will not execute itself recursively.
- 'ld' in PATH.

"The compiler's search directories" means all the directories where gcc searches for passes of the compiler. This includes directories that you specify with '-B'.

Cross-compilers search a little differently:

- 'real-ld' in the compiler's search directories.
- 'target-real-ld' in PATH.
- The file specified in the REAL\_LD\_FILE\_NAME configuration macro, if specified.
- 'ld' in the compiler's search directories.
- 'target-ld' in PATH.

collect2 explicitly avoids running 1d using the file name under which collect2 itself was invoked. In fact, it remembers up a list of such names—in case one copy of collect2 finds another copy (or version) of collect2 installed as 1d in a second place in the search path.

collect2 searches for the utilities nm and strip using the same algorithm as above for ld.

### 22 Standard Header File Directories

GCC\_INCLUDE\_DIR means the same thing for native and cross. It is where GCC stores its private include files, and also where GCC stores the fixed include files. A cross compiled GCC runs fixincludes on the header files in '\$(tooldir)/include'. (If the cross compilation header files need to be fixed, they must be installed before GCC is built. If the cross compilation header files are already suitable for GCC, nothing special need be done).

GPLUSPLUS\_INCLUDE\_DIR means the same thing for native and cross. It is where g++ looks first for header files. The C++ library installs only target independent header files in that directory.

LOCAL\_INCLUDE\_DIR is used only by native compilers. GCC doesn't install anything there. It is normally '/usr/local/include'. This is where local additions to a packaged system should place header files.

CROSS\_INCLUDE\_DIR is used only by cross compilers. GCC doesn't install anything there.

TOOL\_INCLUDE\_DIR is used for both native and cross compilers. It is the place for other packages to install header files that GCC will use. For a cross-compiler, this is the equivalent of '/usr/include'. When you build a cross-compiler, fixincludes processes any header files in this directory.

# 23 Memory Management and Type Information

GCC uses some fairly sophisticated memory management techniques, which involve determining information about GCC's data structures from GCC's source code and using this information to perform garbage collection and implement precompiled headers.

A full C++ parser would be too complicated for this task, so a limited subset of C++ is interpreted and special markers are used to determine what parts of the source to look at. All struct, union and template structure declarations that define data structures that are allocated under control of the garbage collector must be marked. All global variables that hold pointers to garbage-collected memory must also be marked. Finally, all global variables that need to be saved and restored by a precompiled header must be marked. (The precompiled header mechanism can only save static variables if they're scalar. Complex data structures must be allocated in garbage-collected memory to be saved in a precompiled header.)

```
The full format of a marker is

GTY (([option] [(param)], [option] [(param)] ...))
```

but in most cases no options are needed. The outer double parentheses are still necessary, though: GTY(()). Markers can appear:

- In a structure definition, before the open brace;
- In a global variable declaration, after the keyword static or extern; and
- In a structure field definition, before the name of the field.

Here are some examples of marking simple data structures and globals.

```
struct GTY(()) tag
{
   fields...
};

typedef struct GTY(()) tag
{
   fields...
} *typename;

static GTY(()) struct tag *list; /* points to GC memory */
static GTY(()) int counter; /* save counter in a PCH */
```

The parser understands simple typedefs such as typedef struct tag \*name; and typedef int name;. These don't need to be marked.

Since gengtype's understanding of C++ is limited, there are several constructs and declarations that are not supported inside classes/structures marked for automatic GC code generation. The following C++ constructs produce a gengtype error on structures/classes marked for automatic GC code generation:

- Type definitions inside classes/structures are not supported.
- Enumerations inside classes/structures are not supported.

If you have a class or structure using any of the above constructs, you need to mark that class as GTY ((user)) and provide your own marking routines (see section Section 23.3 [User GC], page 680 for details).

It is always valid to include function definitions inside classes. Those are always ignored by gengtype, as it only cares about data members.

#### 23.1 The Inside of a GTY(())

Sometimes the C code is not enough to fully describe the type structure. Extra information can be provided with GTY options and additional markers. Some options take a parameter, which may be either a string or a type name, depending on the parameter. If an option takes no parameter, it is acceptable either to omit the parameter entirely, or to provide an empty string as a parameter. For example, GTY ((skip)) and GTY ((skip (""))) are equivalent.

When the parameter is a string, often it is a fragment of C code. Four special escapes may be used in these strings, to refer to pieces of the data structure being marked:

- %h The current structure.
- The structure that immediately contains the current structure.
- %0 The outermost structure that contains the current structure.
- %a A partial expression of the form [i1][i2]... that indexes the array item currently being marked.

For instance, suppose that you have a structure of the form

```
struct A {
    ...
};
struct B {
    struct A foo[12];
}:
```

and b is a variable of type struct B. When marking 'b.foo[11]', %h would expand to 'b.foo[11]', %0 and %1 would both expand to 'b', and %a would expand to '[11]'.

As in ordinary C, adjacent strings will be concatenated; this is helpful when you have a complicated expression.

The available options are:

#### length ("expression")

There are two places the type machinery will need to be explicitly told the length of an array of non-atomic objects. The first case is when a structure ends in a variable-length array, like this:

```
struct GTY(()) rtvec_def {
  int num_elem;    /* number of elements */
  rtx GTY ((length ("%h.num_elem"))) elem[1];
};
```

In this case, the length option is used to override the specified array length (which should usually be 1). The parameter of the option is a fragment of C code that calculates the length.

The second case is when a structure or a global variable contains a pointer to an array, like this:

```
struct gimple_omp_for_iter * GTY((length ("%h.collapse"))) iter;
```

In this case, iter has been allocated by writing something like

```
x->iter = ggc_alloc_cleared_vec_gimple_omp_for_iter (collapse);
```

and the collapse provides the length of the field.

This second use of length also works on global variables, like:

```
static GTY((length("reg_known_value_size"))) rtx *reg_known_value;
```

Note that the length option is only meant for use with arrays of non-atomic objects, that is, objects that contain pointers pointing to other GTY-managed objects. For other GC-allocated arrays and strings you should use atomic.

skip

If skip is applied to a field, the type machinery will ignore it. This is somewhat dangerous; the only safe use is in a union when one field really isn't ever used.

for\_user

Use this to mark types that need to be marked by user gc routines, but are not refered to in a template argument. So if you have some user gc type T1 and a non user gc type T2 you can give T2 the for\_user option so that the marking functions for T1 can call non mangled functions to mark T2.

```
desc ("expression")
tag ("constant")
default
```

The type machinery needs to be told which field of a union is currently active. This is done by giving each field a constant tag value, and then specifying a discriminator using desc. The value of the expression given by desc is compared against each tag value, each of which should be different. If no tag is matched, the field marked with default is used if there is one, otherwise no field in the union will be marked.

In the desc option, the "current structure" is the union that it discriminates. Use %1 to mean the structure containing it. There are no escapes available to the tag option, since it is a constant.

For example,

```
struct GTY(()) tree_binding
{
  struct tree_common common;
  union tree_binding_u {
    tree GTY ((tag ("0"))) scope;
    struct cp_binding_level * GTY ((tag ("1"))) level;
  } GTY ((desc ("BINDING_HAS_LEVEL_P ((tree)&%0)"))) xscope;
  tree value;
};
```

In this example, the value of BINDING\_HAS\_LEVEL\_P when applied to a struct tree\_binding \* is presumed to be 0 or 1. If 1, the type mechanism will treat the field level as being present and if 0, will treat the field scope as being present.

The desc and tag options can also be used for inheritance to denote which subclass an instance is. See Section 23.2 [Inheritance and GTY], page 680 for more information.

#### cache

When the cache option is applied to a global variable gt\_clear\_cache is called on that variable between the mark and sweep phases of garbage collection. The gt\_clear\_cache function is free to mark blocks as used, or to clear pointers in the variable.

#### deletable

deletable, when applied to a global variable, indicates that when garbage collection runs, there's no need to mark anything pointed to by this variable, it can just be set to NULL instead. This is used to keep a list of free structures around for re-use.

#### maybe\_undef

When applied to a field, maybe\_undef indicates that it's OK if the structure that this fields points to is never defined, so long as this field is always NULL. This is used to avoid requiring backends to define certain optional structures. It doesn't work with language frontends.

#### nested\_ptr (type, "to expression", "from expression")

The type machinery expects all pointers to point to the start of an object. Sometimes for abstraction purposes it's convenient to have a pointer which points inside an object. So long as it's possible to convert the original object to and from the pointer, such pointers can still be used. type is the type of the original object, the to expression returns the pointer given the original object, and the from expression returns the original object given the pointer. The pointer will be available using the %h escape.

```
chain_next ("expression")
chain_prev ("expression")
chain_circular ("expression")
```

It's helpful for the type machinery to know if objects are often chained together in long lists; this lets it generate code that uses less stack space by iterating along the list instead of recursing down it. chain\_next is an expression for the next item in the list, chain\_prev is an expression for the previous item. For singly linked lists, use only chain\_next; for doubly linked lists, use both. The machinery requires that taking the next item of the previous item gives the original item. chain\_circular is similar to chain\_next, but can be used for circular single linked lists.

#### reorder ("function name")

Some data structures depend on the relative ordering of pointers. If the precompiled header machinery needs to change that ordering, it will call the function referenced by the reorder option, before changing the pointers in the object that's pointed to by the field the option applies to. The function must take four arguments, with the signature 'void\*, void\*, gt\_pointer\_operator, void\*. The first parameter is a pointer to the structure that contains the object being updated, or the object itself if there is no containing structure. The second parameter is a cookie that should be ignored. The third parameter is a routine that, given a pointer, will

update it to its correct new value. The fourth parameter is a cookie that must be passed to the second parameter.

PCH cannot handle data structures that depend on the absolute values of pointers. reorder functions can be expensive. When possible, it is better to depend on properties of the data, like an ID number or the hash of a string instead.

atomic

The atomic option can only be used with pointers. It informs the GC machinery that the memory that the pointer points to does not contain any pointers, and hence it should be treated by the GC and PCH machinery as an "atomic" block of memory that does not need to be examined when scanning memory for pointers. In particular, the machinery will not scan that memory for pointers to mark them as reachable (when marking pointers for GC) or to relocate them (when writing a PCH file).

The atomic option differs from the skip option. atomic keeps the memory under Garbage Collection, but makes the GC ignore the contents of the memory. skip is more drastic in that it causes the pointer and the memory to be completely ignored by the Garbage Collector. So, memory marked as atomic is automatically freed when no longer reachable, while memory marked as skip is not.

The atomic option must be used with great care, because all sorts of problem can occur if used incorrectly, that is, if the memory the pointer points to does actually contain a pointer.

Here is an example of how to use it:

```
struct GTY(()) my_struct {
  int number_of_elements;
  unsigned int * GTY ((atomic)) elements;
}:
```

In this case, elements is a pointer under GC, and the memory it points to needs to be allocated using the Garbage Collector, and will be freed automatically by the Garbage Collector when it is no longer referenced. But the memory that the pointer points to is an array of unsigned int elements, and the GC must not try to scan it to find pointers to mark or relocate, which is why it is marked with the atomic option.

Note that, currently, global variables cannot be marked with atomic; only fields of a struct can. This is a known limitation. It would be useful to be able to mark global pointers with atomic to make the PCH machinery aware of them so that they are saved and restored correctly to PCH files.

#### special ("name")

The special option is used to mark types that have to be dealt with by special case machinery. The parameter is the name of the special case. See 'gengtype.c' for further details. Avoid adding new special cases unless there is no other alternative.

The user option indicates that the code to mark structure fields is completely handled by user-provided routines. See section Section 23.3 [User GC], page 680 for details on what functions need to be provided.

### 23.2 Support for inheritance

gengtype has some support for simple class hierarchies. You can use this to have gengtype autogenerate marking routines, provided:

- There must be a concrete base class, with a discriminator expression that can be used to identify which subclass an instance is.
- Only single inheritance is used.
- None of the classes within the hierarchy are templates.

If your class hierarchy does not fit in this pattern, you must use Section 23.3 [User GC], page 680 instead.

The base class and its discriminator must be identified using the "desc" option. Each concrete subclass must use the "tag" option to identify which value of the discriminator it corresponds to.

Every class in the hierarchy must have a GTY(()) marker, as gengtype will only attempt to parse classes that have such a marker<sup>1</sup>.

```
class GTY((desc("%h.kind"), tag("0"))) example_base
{
public:
    int kind;
    tree a;
};

class GTY((tag("1"))) some_subclass : public example_base
{
public:
    tree b;
};

class GTY((tag("2"))) some_other_subclass : public example_base
{
public:
    tree c;
}.
```

The generated marking routines for the above will contain a "switch" on "kind", visiting all appropriate fields. For example, if kind is 2, it will cast to "some\_other\_subclass" and visit fields a, b, and c.

# 23.3 Support for user-provided GC marking routines

The garbage collector supports types for which no automatic marking code is generated. For these types, the user is required to provide three functions: one to act as a marker for

<sup>&</sup>lt;sup>1</sup> Classes lacking such a marker will not be identified as being part of the hierarchy, and so the marking routines will not handle them, leading to a assertion failure within the marking routines due to an unknown tag value (assuming that assertions are enabled).

garbage collection, and two functions to act as marker and pointer walker for pre-compiled headers.

Given a structure struct GTY((user)) my\_struct, the following functions should be defined to mark my\_struct:

```
void gt_ggc_mx (my_struct *p)
{
    /* This marks field 'fld'. */
    gt_ggc_mx (p->fld);
}

void gt_pch_nx (my_struct *p)
{
    /* This marks field 'fld'. */
    gt_pch_nx (tp->fld);
}

void gt_pch_nx (my_struct *p, gt_pointer_operator op, void *cookie)
{
    /* For every field 'fld', call the given pointer operator. */
    op (&(tp->fld), cookie);
}
```

In general, each marker M should call M for every pointer field in the structure. Fields that are not allocated in GC or are not pointers must be ignored.

For embedded lists (e.g., structures with a next or prev pointer), the marker must follow the chain and mark every element in it.

Note that the rules for the pointer walker gt\_pch\_nx (my\_struct \*, gt\_pointer\_operator, void \*) are slightly different. In this case, the operation op must be applied to the address of every pointer field.

#### 23.3.1 User-provided marking routines for template types

When a template type TP is marked with GTY, all instances of that type are considered user-provided types. This means that the individual instances of TP do not need to be marked with GTY. The user needs to provide template functions to mark all the fields of the type.

The following code snippets represent all the functions that need to be provided. Note that type TP may reference to more than one type. In these snippets, there is only one type T, but there could be more.

```
template<typename T>
void gt_ggc_mx (TP<T> *tp)
{
   extern void gt_ggc_mx (T&);

   /* This marks field 'fld' of type 'T'. */
   gt_ggc_mx (tp->fld);
}

template<typename T>
void gt_pch_nx (TP<T> *tp)
{
   extern void gt_pch_nx (T&);

   /* This marks field 'fld' of type 'T'. */
   gt_pch_nx (tp->fld);
```

```
template<typename T>
void gt_pch_nx (TP<T *> *tp, gt_pointer_operator op, void *cookie)
{
   /* For every field 'fld' of 'tp' with type 'T *', call the given pointer operator. */
   op (&(tp->fld), cookie);
}

template<typename T>
void gt_pch_nx (TP<T> *tp, gt_pointer_operator, void *cookie)
{
   extern void gt_pch_nx (T *, gt_pointer_operator, void *);

   /* For every field 'fld' of 'tp' with type 'T', call the pointer walker for all the fields of T. */
   gt_pch_nx (&(tp->fld), op, cookie);
}
```

Support for user-defined types is currently limited. The following restrictions apply:

- 1. Type TP and all the argument types T must be marked with GTY.
- 2. Type TP can only have type names in its argument list.
- 3. The pointer walker functions are different for TP<T> and TP<T \*>. In the case of TP<T>, references to T must be handled by calling gt\_pch\_nx (which will, in turn, walk all the pointers inside fields of T). In the case of TP<T \*>, references to T \* must be handled by calling the op function on the address of the pointer (see the code snippets above).

### 23.4 Marking Roots for the Garbage Collector

In addition to keeping track of types, the type machinery also locates the global variables (roots) that the garbage collector starts at. Roots must be declared using one of the following syntaxes:

- extern GTY(([options])) type name;
- static GTY(([options])) type name;

The syntax

• GTY(([options])) type name;

is *not* accepted. There should be an **extern** declaration of such a variable in a header somewhere—mark that, not the definition. Or, if the variable is only used in one file, make it **static**.

# 23.5 Source Files Containing Type Information

Whenever you add GTY markers to a source file that previously had none, or create a new source file containing GTY markers, there are three things you need to do:

- 1. You need to add the file to the list of source files the type machinery scans. There are four cases:
  - a. For a back-end file, this is usually done automatically; if not, you should add it to target\_gtfiles in the appropriate port's entries in 'config.gcc'.

- b. For files shared by all front ends, add the filename to the GTFILES variable in 'Makefile.in'.
- c. For files that are part of one front end, add the filename to the gtfiles variable defined in the appropriate 'config-lang.in'. Headers should appear before nonheaders in this list.
- d. For files that are part of some but not all front ends, add the filename to the gtfiles variable of *all* the front ends that use it.
- 2. If the file was a header file, you'll need to check that it's included in the right place to be visible to the generated files. For a back-end header file, this should be done automatically. For a front-end header file, it needs to be included by the same file that includes 'gtype-lang.h'. For other header files, it needs to be included in 'gtype-desc.c', which is a generated file, so add it to ifiles in open\_base\_file in 'gengtype.c'.

For source files that aren't header files, the machinery will generate a header file that should be included in the source file you just changed. The file will be called 'gt-path.h' where path is the pathname relative to the 'gcc' directory with slashes replaced by -, so for example the header file to be included in 'cp/parser.c' is called 'gt-cp-parser.c'. The generated header file should be included after everything else in the source file. Don't forget to mention this file as a dependency in the 'Makefile'!

For language frontends, there is another file that needs to be included somewhere. It will be called 'gtype-lang.h', where lang is the name of the subdirectory the language is contained in.

Plugins can add additional root tables. Run the gengtype utility in plugin mode as gengtype -P pluginout.h source-dir file-list plugin\*.c with your plugin files plugin\*.c using GTY to generate the pluginout.h file. The GCC build tree is needed to be present in that mode.

### 23.6 How to invoke the garbage collector

The GCC garbage collector GGC is only invoked explicitly. In contrast with many other garbage collectors, it is not implicitly invoked by allocation routines when a lot of memory has been consumed. So the only way to have GGC reclaim storage is to call the ggc\_collect function explicitly. This call is an expensive operation, as it may have to scan the entire heap. Beware that local variables (on the GCC call stack) are not followed by such an invocation (as many other garbage collectors do): you should reference all your data from static or external GTY-ed variables, and it is advised to call ggc\_collect with a shallow call stack. The GGC is an exact mark and sweep garbage collector (so it does not scan the call stack for pointers). In practice GCC passes don't often call ggc\_collect themselves, because it is called by the pass manager between passes.

At the time of the ggc\_collect call all pointers in the GC-marked structures must be valid or NULL. In practice this means that there should not be uninitialized pointer fields in the structures even if your code never reads or writes those fields at a particular instance. One way to ensure this is to use cleared versions of allocators unless all the fields are initialized manually immediately after allocation.

### 23.7 Troubleshooting the garbage collector

With the current garbage collector implementation, most issues should show up as GCC compilation errors. Some of the most commonly encountered issues are described below.

- Gengtype does not produce allocators for a GTY-marked type. Gengtype checks if there is at least one possible path from GC roots to at least one instance of each type before outputting allocators. If there is no such path, the GTY markers will be ignored and no allocators will be output. Solve this by making sure that there exists at least one such path. If creating it is unfeasible or raises a "code smell", consider if you really must use GC for allocating such type.
- Link-time errors about undefined gt\_ggc\_r\_foo\_bar and similarly-named symbols. Check if your 'foo\_bar' source file has #include "gt-foo\_bar.h" as its very last line.

# 24 Plugins

GCC plugins are loadable modules that provide extra features to the compiler. Like GCC itself they can be distributed in source and binary forms.

GCC plugins provide developers with a rich subset of the GCC API to allow them to extend GCC as they see fit. Whether it is writing an additional optimization pass, transforming code, or analyzing information, plugins can be quite useful.

### 24.1 Loading Plugins

Plugins are supported on platforms that support '-ldl -rdynamic' as well as Windows/MinGW. They are loaded by the compiler using dlopen or equivalent and invoked at pre-determined locations in the compilation process.

Plugins are loaded with

```
'-fplugin=/path/to/name.ext' '-fplugin-arg-name-key1[=value1]'
```

Where name is the plugin name and ext is the platform-specific dynamic library extension. It should be dll on Windows/MinGW, dylib on Darwin/Mac OS X, and so on all other platforms. The plugin arguments are parsed by GCC and passed to respective plugins as key-value pairs. Multiple plugins can be invoked by specifying multiple '-fplugin' arguments.

A plugin can be simply given by its short name (no dots or slashes). When simply passing '-fplugin=name', the plugin is loaded from the 'plugin' directory, so '-fplugin=name' is the same as '-fplugin='gcc -print-file-name=plugin'/name.ext', using backquote shell syntax to query the 'plugin' directory.

# 24.2 Plugin API

Plugins are activated by the compiler at specific events as defined in 'gcc-plugin.h'. For each event of interest, the plugin should call register\_callback specifying the name of the event and address of the callback function that will handle that event.

The header 'gcc-plugin.h' must be the first gcc header to be included.

### 24.2.1 Plugin license check

Every plugin should define the global symbol plugin\_is\_GPL\_compatible to assert that it has been licensed under a GPL-compatible license. If this symbol does not exist, the compiler will emit a fatal error and exit with the error message:

```
fatal error: plugin name is not licensed under a GPL-compatible license name: undefined symbol: plugin_is_GPL_compatible compilation terminated
```

The declared type of the symbol should be int, to match a forward declaration in 'gcc-plugin.h' that suppresses C++ mangling. It does not need to be in any allocated section, though. The compiler merely asserts that the symbol exists in the global scope. Something like this is enough:

```
int plugin_is_GPL_compatible;
```

#### 24.2.2 Plugin initialization

Every plugin should export a function called plugin\_init that is called right after the plugin is loaded. This function is responsible for registering all the callbacks required by the plugin and do any other required initialization.

This function is called from compile\_file right before invoking the parser. The arguments to plugin\_init are:

- plugin\_info: Plugin invocation information.
- version: GCC version.

The plugin\_info struct is defined as follows:

If initialization fails, plugin\_init must return a non-zero value. Otherwise, it should return 0.

The version of the GCC compiler loading the plugin is described by the following structure:

```
struct plugin_gcc_version
{
  const char *basever;
  const char *datestamp;
  const char *devphase;
  const char *revision;
  const char *configuration_arguments;
};
```

The function plugin\_default\_version\_check takes two pointers to such structure and compare them field by field. It can be used by the plugin's plugin\_init function.

The version of GCC used to compile the plugin can be found in the symbol gcc\_version defined in the header 'plugin-version.h'. The recommended version check to perform looks like

but you can also check the individual fields if you want a less strict check.

#### 24.2.3 Plugin callbacks

```
Callback functions have the following prototype:
      /* The prototype for a plugin callback function.
            gcc_data - event-specific data provided by GCC
           user_data - plugin-specific data provided by the plug-in. */
      typedef void (*plugin_callback_func)(void *gcc_data, void *user_data);
  Callbacks can be invoked at the following pre-determined events:
      enum plugin_event
        PLUGIN_START_PARSE_FUNCTION, /* Called before parsing the body of a function. */
        PLUGIN_FINISH_PARSE_FUNCTION, /* After finishing parsing a function. */
        PLUGIN_PASS_MANAGER_SETUP, /* To hook into pass manager. */
        PLUGIN_FINISH_TYPE,
                                        /* After finishing parsing a type. */
        PLUGIN_FINISH_DECL, /* After finishing parsing a declaration. */
PLUGIN_FINISH_UNIT, /* Useful for summary processing. */
PLUGIN_PRE_GENERICIZE, /* Allows to see low level AST in C and C++ frontends. */
PLUGIN_FINISH, /* Called before GCC exits. */
                                       /* Information about the plugin. */
        PLUGIN_INFO,
                                /* Called at start of GCC Garbage Collection. */
/* Extend the GGC marking. */
/* Called at end of GGC. */
        PLUGIN_GGC_START,
        PLUGIN_GGC_MARKING,
        PLUGIN_GGC_END,
        PLUGIN_REGISTER_GGC_ROOTS, /* Register an extra GGC root table. */
        PLUGIN_ATTRIBUTES,
                                       /* Called during attribute registration */
        PLUGIN_START_UNIT,
                                        /* Called before processing a translation unit. */
                                        /* Called during pragma registration. */
        PLUGIN_PRAGMAS,
        /* Called before first pass from all_passes. */
        PLUGIN_ALL_PASSES_START,
        /* Called after last pass from all_passes. */
        PLUGIN_ALL_PASSES_END,
        /* Called before first ipa pass. */
        PLUGIN_ALL_IPA_PASSES_START,
        /* Called after last ipa pass. */
        PLUGIN_ALL_IPA_PASSES_END,
        /* Allows to override pass gate decision for current_pass. */
        PLUGIN_OVERRIDE_GATE,
        /* Called before executing a pass. */
        PLUGIN_PASS_EXECUTION,
        /* Called before executing subpasses of a GIMPLE_PASS in
            execute_ipa_pass_list. */
        PLUGIN_EARLY_GIMPLE_PASSES_START,
        /* Called after executing subpasses of a GIMPLE_PASS in
           execute_ipa_pass_list. */
        PLUGIN_EARLY_GIMPLE_PASSES_END,
        /* Called when a pass is first instantiated. */
        PLUGIN_NEW_PASS,
      /* Called when a file is #include-d or given via the #line directive.
         This could happen many times. The event data is the included file path,
         as a const char* pointer. */
        PLUGIN_INCLUDE_FILE,
        PLUGIN_EVENT_FIRST_DYNAMIC
                                         /* Dummy event used for indexing callback
                                            array. */
      };
```

In addition, plugins can also look up the enumerator of a named event, and / or generate new events dynamically, by calling the function get\_named\_event\_id.

To register a callback, the plugin calls register\_callback with the arguments:

- char \*name: Plugin name.
- int event: The event code.
- plugin\_callback\_func callback: The function that handles event.
- void \*user\_data: Pointer to plugin-specific data.

For the  $PLUGIN\_PASS\_MANAGER\_SETUP$ ,  $PLUGIN\_INFO$ , and  $PLUGIN\_REGISTER\_GGC\_ROOTS$  pseudo-events the callback should be null, and the user\_data is specific.

When the *PLUGIN\_PRAGMAS* event is triggered (with a null pointer as data from GCC), plugins may register their own pragmas. Notice that pragmas are not available from 'lto1', so plugins used with -flto option to GCC during link-time optimization cannot use pragmas and do not even see functions like c\_register\_pragma or pragma\_lex.

The *PLUGIN\_INCLUDE\_FILE* event, with a const char\* file path as GCC data, is triggered for processing of #include or #line directives.

The *PLUGIN\_FINISH* event is the last time that plugins can call GCC functions, notably emit diagnostics with warning, error etc.

### 24.3 Interacting with the pass manager

There needs to be a way to add/reorder/remove passes dynamically. This is useful for both analysis plugins (plugging in after a certain pass such as CFG or an IPA pass) and optimization plugins.

Basic support for inserting new passes or replacing existing passes is provided. A plugin registers a new pass with GCC by calling register\_callback with the PLUGIN\_PASS\_MANAGER\_SETUP event and a pointer to a struct register\_pass\_info object defined as follows

```
enum pass_positioning_ops
 PASS_POS_INSERT_AFTER, // Insert after the reference pass.
 PASS_POS_INSERT_BEFORE, // Insert before the reference pass.
 PASS_POS_REPLACE
                     // Replace the reference pass.
struct register_pass_info
                                   /* New pass provided by the plugin. */
 struct opt_pass *pass;
 const char *reference_pass_name; /* Name of the reference pass for hooking
                                      up the new pass. */
 int ref_pass_instance_number;
                                   /* Insert the pass at the specified
                                      instance number of the reference pass. */
                                   /* Do it for every instance if it is 0. */
 enum pass_positioning_ops pos_op; /* how to insert the new pass. */  
};
/* Sample plugin code that registers a new pass. */
plugin_init (struct plugin_name_args *plugin_info,
            struct plugin_gcc_version *version)
 struct register_pass_info pass_info;
```

```
/* Code to fill in the pass_info object with new pass information. */
...

/* Register the new pass. */
register_callback (plugin_info->base_name, PLUGIN_PASS_MANAGER_SETUP, NULL, &pass_info);
...
}
```

### 24.4 Interacting with the GCC Garbage Collector

Some plugins may want to be informed when GGC (the GCC Garbage Collector) is running. They can register callbacks for the PLUGIN\_GGC\_START and PLUGIN\_GGC\_END events (for which the callback is called with a null gcc\_data) to be notified of the start or end of the GCC garbage collection.

Some plugins may need to have GGC mark additional data. This can be done by registering a callback (called with a null gcc\_data) for the PLUGIN\_GGC\_MARKING event. Such callbacks can call the ggc\_set\_mark routine, preferably through the ggc\_mark macro (and conversely, these routines should usually not be used in plugins outside of the PLUGIN\_GGC\_MARKING event). Plugins that wish to hold weak references to gc data may also use this event to drop weak references when the object is about to be collected. The ggc\_marked\_p function can be used to tell if an object is marked, or is about to be collected. The gt\_clear\_cache overloads which some types define may also be of use in managing weak references.

Some plugins may need to add extra GGC root tables, e.g. to handle their own GTY-ed data. This can be done with the PLUGIN\_REGISTER\_GGC\_ROOTS pseudo-event with a null callback and the extra root table (of type struct ggc\_root\_tab\*) as user\_data. Running the gengtype -p source-dir file-list plugin\*.c . . . utility generates these extra root tables.

You should understand the details of memory management inside GCC before using PLUGIN\_GGC\_MARKING or PLUGIN\_REGISTER\_GGC\_ROOTS.

### 24.5 Giving information about a plugin

A plugin should give some information to the user about itself. This uses the following structure:

```
struct plugin_info
{
  const char *version;
  const char *help;
};
```

Such a structure is passed as the user\_data by the plugin's init routine using register\_callback with the PLUGIN\_INFO pseudo-event and a null callback.

# 24.6 Registering custom attributes or pragmas

For analysis (or other) purposes it is useful to be able to add custom attributes or pragmas.

The PLUGIN\_ATTRIBUTES callback is called during attribute registration. Use the register\_attribute function to register custom attributes.

```
/* Attribute handler callback */
static tree
handle_user_attribute (tree *node, tree name, tree args,
                       int flags, bool *no_add_attrs)
 return NULL_TREE;
/* Attribute definition */
static struct attribute_spec user_attr =
  { "user", 1, 1, false, false, false, false, handle_user_attribute, NULL };
/* Plugin callback called during attribute registration.
Registered with register_callback (plugin_name, PLUGIN_ATTRIBUTES, register_attributes, NULL)
*/
static void
register_attributes (void *event_data, void *data)
 warning (0, G_("Callback to register attributes"));
 register_attribute (&user_attr);
}
```

The *PLUGIN\_PRAGMAS* callback is called once during pragmas registration. Use the c\_register\_pragma, c\_register\_pragma\_with\_data, c\_register\_pragma\_with\_expansion, c\_register\_pragma\_with\_expansion\_and\_data functions to register custom pragmas and their handlers (which often want to call pragma\_lex) from 'c-family/c-pragma.h'.

It is suggested to pass "GCCPLUGIN" (or a short name identifying your plugin) as the "space" argument of your pragma.

Pragmas registered with c\_register\_pragma\_with\_expansion or c\_register\_pragma\_with\_expansion\_and\_data support preprocessor expansions. For example:

```
#define NUMBER 10
#pragma GCCPLUGIN foothreshold (NUMBER)
```

### 24.7 Recording information about pass execution

The event PLUGIN\_PASS\_EXECUTION passes the pointer to the executed pass (the same as current\_pass) as gcc\_data to the callback. You can also inspect cfun to find out about which function this pass is executed for. Note that this event will only be invoked if the gate check (if applicable, modified by PLUGIN\_OVERRIDE\_GATE) succeeds. You can use other hooks, like PLUGIN\_ALL\_PASSES\_START, PLUGIN\_ALL\_PASSES\_END,

PLUGIN\_ALL\_IPA\_PASSES\_START, PLUGIN\_ALL\_IPA\_PASSES\_END, PLUGIN\_EARLY\_GIMPLE\_PASSES\_START, and/or PLUGIN\_EARLY\_GIMPLE\_PASSES\_END to manipulate global state in your plugin(s) in order to get context for the pass execution.

### 24.8 Controlling which passes are being run

After the original gate function for a pass is called, its result - the gate status - is stored as an integer. Then the event PLUGIN\_OVERRIDE\_GATE is invoked, with a pointer to the gate status in the gcc\_data parameter to the callback function. A nonzero value of the gate status means that the pass is to be executed. You can both read and write the gate status via the passed pointer.

### 24.9 Keeping track of available passes

When your plugin is loaded, you can inspect the various pass lists to determine what passes are available. However, other plugins might add new passes. Also, future changes to GCC might cause generic passes to be added after plugin loading. When a pass is first added to one of the pass lists, the event PLUGIN\_NEW\_PASS is invoked, with the callback parameter gcc\_data pointing to the new pass.

### 24.10 Building GCC plugins

If plugins are enabled, GCC installs the headers needed to build a plugin (somewhere in the installation tree, e.g. under '/usr/local'). In particular a 'plugin/include' directory is installed, containing all the header files needed to build plugins.

On most systems, you can query this plugin directory by invoking gcc -print-file-name=plugin (replace if needed gcc with the appropriate program path).

Inside plugins, this plugin directory name can be queried by calling default\_plugin\_dir\_name ().

Plugins may know, when they are compiled, the GCC version for which 'plugin-version.h' is provided. The constant macros GCCPLUGIN\_VERSION\_MAJOR, GCCPLUGIN\_VERSION\_PATCHLEVEL, GCCPLUGIN\_VERSION are integer numbers, so a plugin could ensure it is built for GCC 4.7 with

```
#if GCCPLUGIN_VERSION != 4007
#error this GCC plugin is for GCC 4.7
#endif
```

The following GNU Makefile excerpt shows how to build a simple plugin:

```
HOST_GCC=g++
TARGET_GCC=gcc
PLUGIN_SOURCE_FILES= plugin1.c plugin2.cc
GCCPLUGINS_DIR:= $(shell $(TARGET_GCC) -print-file-name=plugin)
CXXFLAGS+= -I$(GCCPLUGINS_DIR)/include -fPIC -fno-rtti -02
plugin.so: $(PLUGIN_SOURCE_FILES)
    $(HOST_GCC) -shared $(CXXFLAGS) $^ -o $@
```

A single source file plugin may be built with g++-I'gcc-print-file-name=plugin'/include-fPIC-shared-fno-rtti-02 plugin.c-o plugin.so, using backquote shell syntax to query the 'plugin' directory.

Plugin support on Windows/MinGW has a number of limitations and additional requirements. When building a plugin on Windows we have to link an import library for the corresponding backend executable, for example, 'cc1.exe', 'cc1plus.exe', etc., in order to gain access to the symbols provided by GCC. This means that on Windows a plugin is language-specific, for example, for C, C++, etc. If you wish to use your plugin with multiple languages, then you will need to build multiple plugin libraries and either instruct your users on how to load the correct version or provide a compiler wrapper that does this automatically.

Additionally, on Windows the plugin library has to export the plugin\_is\_GPL\_compatible and plugin\_init symbols. If you do not wish to modify the source code of your plugin, then you can use the '-Wl,--export-all-symbols' option or provide a suitable DEF file. Alternatively, you can export just these two symbols by decorating them with \_\_declspec(dllexport), for example:

```
#ifdef _WIN32
__declspec(dllexport)
#endif
int plugin_is_GPL_compatible;

#ifdef _WIN32
__declspec(dllexport)
#endif
int plugin_init (plugin_name_args *, plugin_gcc_version *)
```

The import libraries are installed into the plugin directory and their names are derived by appending the .a extension to the backend executable names, for example, 'cc1.exe.a', 'cc1plus.exe.a', etc. The following command line shows how to build the single source file plugin on Windows to be used with the C++ compiler:

```
g++ -I'gcc -print-file-name=plugin'/include -shared -Wl,--export-all-symbols \
-o plugin.dll plugin.c 'gcc -print-file-name=plugin'/cc1plus.exe.a
```

When a plugin needs to use gengtype, be sure that both 'gengtype' and 'gtype.state' have the same version as the GCC for which the plugin is built.

# 25 Link Time Optimization

Link Time Optimization (LTO) gives GCC the capability of dumping its internal representation (GIMPLE) to disk, so that all the different compilation units that make up a single executable can be optimized as a single module. This expands the scope of inter-procedural optimizations to encompass the whole program (or, rather, everything that is visible at link time).

### 25.1 Design Overview

Link time optimization is implemented as a GCC front end for a bytecode representation of GIMPLE that is emitted in special sections of .o files. Currently, LTO support is enabled in most ELF-based systems, as well as darwin, cygwin and mingw systems.

Since GIMPLE bytecode is saved alongside final object code, object files generated with LTO support are larger than regular object files. This "fat" object format makes it easy to integrate LTO into existing build systems, as one can, for instance, produce archives of the files. Additionally, one might be able to ship one set of fat objects which could be used both for development and the production of optimized builds. A, perhaps surprising, side effect of this feature is that any mistake in the toolchain leads to LTO information not being used (e.g. an older libtool calling ld directly). This is both an advantage, as the system is more robust, and a disadvantage, as the user is not informed that the optimization has been disabled.

The current implementation only produces "fat" objects, effectively doubling compilation time and increasing file sizes up to 5x the original size. This hides the problem that some tools, such as ar and nm, need to understand symbol tables of LTO sections. These tools were extended to use the plugin infrastructure, and with these problems solved, GCC will also support "slim" objects consisting of the intermediate code alone.

At the highest level, LTO splits the compiler in two. The first half (the "writer") produces a streaming representation of all the internal data structures needed to optimize and generate code. This includes declarations, types, the callgraph and the GIMPLE representation of function bodies.

When '-flto' is given during compilation of a source file, the pass manager executes all the passes in all\_lto\_gen\_passes. Currently, this phase is composed of two IPA passes:

- pass\_ipa\_lto\_gimple\_out This pass executes the function lto\_output in 'lto-streamer-out.c', which traverses the call graph encoding every reachable declaration, type and function. This generates a memory representation of all the file sections described below.
- pass\_ipa\_lto\_finish\_out This pass executes the function produce\_asm\_for\_decls in 'lto-streamer-out.c', which takes the memory image built in the previous pass and encodes it in the corresponding ELF file sections.

The second half of LTO support is the "reader". This is implemented as the GCC front end 'lto1' in 'lto/lto.c'. When 'collect2' detects a link set of .o/.a files with LTO information and the '-flto' is enabled, it invokes 'lto1' which reads the set of files and aggregates them into a single translation unit for optimization. The main entry point for the reader is 'lto/lto.c':lto\_main.

#### 25.1.1 LTO modes of operation

One of the main goals of the GCC link-time infrastructure was to allow effective compilation of large programs. For this reason GCC implements two link-time compilation modes.

- 1. LTO mode, in which the whole program is read into the compiler at link-time and optimized in a similar way as if it were a single source-level compilation unit.
- 2. WHOPR or partitioned mode, designed to utilize multiple CPUs and/or a distributed compilation environment to quickly link large applications. WHOPR stands for WHOle Program optimizeR (not to be confused with the semantics of '-fwhole-program'). It partitions the aggregated callgraph from many different .o files and distributes the compilation of the sub-graphs to different CPUs.

Note that distributed compilation is not implemented yet, but since the parallelism is facilitated via generating a Makefile, it would be easy to implement.

#### WHOPR splits LTO into three main stages:

- 1. Local generation (LGEN) This stage executes in parallel. Every file in the program is compiled into the intermediate language and packaged together with the local call-graph and summary information. This stage is the same for both the LTO and WHOPR compilation mode.
- 2. Whole Program Analysis (WPA) WPA is performed sequentially. The global call-graph is generated, and a global analysis procedure makes transformation decisions. The global call-graph is partitioned to facilitate parallel optimization during phase 3. The results of the WPA stage are stored into new object files which contain the partitions of program expressed in the intermediate language and the optimization decisions.
- 3. Local transformations (LTRANS) This stage executes in parallel. All the decisions made during phase 2 are implemented locally in each partitioned object file, and the final object code is generated. Optimizations which cannot be decided efficiently during the phase 2 may be performed on the local call-graph partitions.

WHOPR can be seen as an extension of the usual LTO mode of compilation. In LTO, WPA and LTRANS are executed within a single execution of the compiler, after the whole program has been read into memory.

When compiling in WHOPR mode, the callgraph is partitioned during the WPA stage. The whole program is split into a given number of partitions of roughly the same size. The compiler tries to minimize the number of references which cross partition boundaries. The main advantage of WHOPR is to allow the parallel execution of LTRANS stages, which are the most time-consuming part of the compilation process. Additionally, it avoids the need to load the whole program into memory.

#### 25.2 LTO file sections

LTO information is stored in several ELF sections inside object files. Data structures and enum codes for sections are defined in 'lto-streamer.h'.

These sections are emitted from 'lto-streamer-out.c' and mapped in all at once from 'lto/lto.c':lto\_file\_read. The individual functions dealing with the reading/writing of each section are described below.

• Command line options (.gnu.lto\_.opts)

This section contains the command line options used to generate the object files. This is used at link time to determine the optimization level and other settings when they are not explicitly specified at the linker command line.

Currently, GCC does not support combining LTO object files compiled with different set of the command line options into a single binary. At link time, the options given on the command line and the options saved on all the files in a link-time set are applied globally. No attempt is made at validating the combination of flags (other than the usual validation done by option processing). This is implemented in 'lto/lto.c':lto\_read\_all\_file\_options.

• Symbol table (.gnu.lto\_.symtab)

This table replaces the ELF symbol table for functions and variables represented in the LTO IL. Symbols used and exported by the optimized assembly code of "fat" objects might not match the ones used and exported by the intermediate code. This table is necessary because the intermediate code is less optimized and thus requires a separate symbol table.

Additionally, the binary code in the "fat" object will lack a call to a function, since the call was optimized out at compilation time after the intermediate language was streamed out. In some special cases, the same optimization may not happen during link-time optimization. This would lead to an undefined symbol if only one symbol table was used.

The symbol table is emitted in 'lto-streamer-out.c':produce\_symtab.

• Global declarations and types (.gnu.lto\_.decls)

This section contains an intermediate language dump of all declarations and types required to represent the callgraph, static variables and top-level debug info.

The contents of this section are emitted in 'lto-streamer-out.c':produce\_asm\_for\_decls. Types and symbols are emitted in a topological order that preserves the sharing of pointers when the file is read back in ('lto.c':read\_cgraph\_and\_symbols).

• The callgraph (.gnu.lto\_.cgraph)

This section contains the basic data structure used by the GCC inter-procedural optimization infrastructure. This section stores an annotated multi-graph which represents the functions and call sites as well as the variables, aliases and top-level asm statements.

This section is emitted in 'lto-streamer-out.c':output\_cgraph and read in 'lto-cgraph.c':input\_cgraph.

• IPA references (.gnu.lto\_.refs)

This section contains references between function and static variables. It is emitted by 'lto-cgraph.c':output\_refs and read by 'lto-cgraph.c':input\_refs.

• Function bodies (.gnu.lto\_.function\_body.<name>)

This section contains function bodies in the intermediate language representation. Every function body is in a separate section to allow copying of the section independently to different object files or reading the function on demand.

Functions are emitted in 'lto-streamer-out.c':output\_function and read in 'lto-streamer-in.c':input\_function.

- Static variable initializers (.gnu.lto\_.vars)

  This section contains all the symbols in the global variable pool. It is emitted by 'lto-cgraph.c':output\_varpool and read in 'lto-cgraph.c':input\_cgraph.
- Summaries and optimization summaries used by IPA passes (.gnu.lto\_.<xxx>, where <xxx> is one of jmpfuncs, pureconst or reference)

These sections are used by IPA passes that need to emit summary information during LTO generation to be read and aggregated at link time. Each pass is responsible for implementing two pass manager hooks: one for writing the summary and another for reading it in. The format of these sections is entirely up to each individual pass. The only requirement is that the writer and reader hooks agree on the format.

### 25.3 Using summary information in IPA passes

Programs are represented internally as a *callgraph* (a multi-graph where nodes are functions and edges are call sites) and a *varpool* (a list of static and external variables in the program).

The inter-procedural optimization is organized as a sequence of individual passes, which operate on the callgraph and the varpool. To make the implementation of WHOPR possible, every inter-procedural optimization pass is split into several stages that are executed at different times during WHOPR compilation:

#### • LGEN time

- 1. Generate summary (generate\_summary in struct ipa\_opt\_pass\_d). This stage analyzes every function body and variable initializer is examined and stores relevant information into a pass-specific data structure.
- 2. Write summary (write\_summary in struct ipa\_opt\_pass\_d). This stage writes all the pass-specific information generated by generate\_summary. Summaries go into their own LTO\_section\_\* sections that have to be declared in 'lto-streamer.h':enum lto\_section\_type. A new section is created by calling create\_output\_block and data can be written using the lto\_output\_\* routines.

#### • WPA time

- 1. Read summary (read\_summary in struct ipa\_opt\_pass\_d). This stage reads all the pass-specific information in exactly the same order that it was written by write\_summary.
- 2. Execute (execute in struct opt\_pass). This performs inter-procedural propagation. This must be done without actual access to the individual function bodies or variable initializers. Typically, this results in a transitive closure operation over the summary information of all the nodes in the callgraph.
- 3. Write optimization summary (write\_optimization\_summary in struct ipa\_opt\_pass\_d). This writes the result of the inter-procedural propagation into the object file. This can use the same data structures and helper routines used in write\_summary.

#### • LTRANS time

1. Read optimization summary (read\_optimization\_summary in struct ipa\_opt\_ pass\_d). The counterpart to write\_optimization\_summary. This reads the interprocedural optimization decisions in exactly the same format emitted by write\_optimization\_summary.

2. Transform (function\_transform and variable\_transform in struct ipa\_opt\_ pass\_d). The actual function bodies and variable initializers are updated based on the information passed down from the Execute stage.

The implementation of the inter-procedural passes are shared between LTO, WHOPR and classic non-LTO compilation.

- During the traditional file-by-file mode every pass executes its own *Generate summary*, *Execute*, and *Transform* stages within the single execution context of the compiler.
- In LTO compilation mode, every pass uses *Generate summary* and *Write summary* stages at compilation time, while the *Read summary*, *Execute*, and *Transform* stages are executed at link time.
- In WHOPR mode all stages are used.

To simplify development, the GCC pass manager differentiates between normal interprocedural passes (see Section 9.4.2 [Regular IPA passes], page 130), small inter-procedural passes (see Section 9.4.1 [Small IPA passes], page 129) and late inter-procedural passes (see Section 9.4.3 [Late IPA passes], page 132). A small or late IPA pass (SIMPLE\_IPA\_PASS) does everything at once and thus cannot be executed during WPA in WHOPR mode. It defines only the *Execute* stage and during this stage it accesses and modifies the function bodies. Such passes are useful for optimization at LGEN or LTRANS time and are used, for example, to implement early optimization before writing object files. The simple inter-procedural passes can also be used for easier prototyping and development of a new inter-procedural pass.

#### 25.3.1 Virtual clones

One of the main challenges of introducing the WHOPR compilation mode was addressing the interactions between optimization passes. In LTO compilation mode, the passes are executed in a sequence, each of which consists of analysis (or *Generate summary*), propagation (or *Execute*) and *Transform* stages. Once the work of one pass is finished, the next pass sees the updated program representation and can execute. This makes the individual passes dependent on each other.

In WHOPR mode all passes first execute their *Generate summary* stage. Then summary writing marks the end of the LGEN stage. At WPA time, the summaries are read back into memory and all passes run the *Execute* stage. Optimization summaries are streamed and sent to LTRANS, where all the passes execute the *Transform* stage.

Most optimization passes split naturally into analysis, propagation and transformation stages. But some do not. The main problem arises when one pass performs changes and the following pass gets confused by seeing different callgraphs between the *Transform* stage and the *Generate summary* or *Execute* stage. This means that the passes are required to communicate their decisions with each other.

To facilitate this communication, the GCC callgraph infrastructure implements *virtual clones*, a method of representing the changes performed by the optimization passes in the callgraph without needing to update function bodies.

A *virtual clone* in the callgraph is a function that has no associated body, just a description of how to create its body based on a different function (which itself may be a virtual clone).

The description of function modifications includes adjustments to the function's signature (which allows, for example, removing or adding function arguments), substitutions to

perform on the function body, and, for inlined functions, a pointer to the function that it will be inlined into.

It is also possible to redirect any edge of the callgraph from a function to its virtual clone. This implies updating of the call site to adjust for the new function signature.

Most of the transformations performed by inter-procedural optimizations can be represented via virtual clones. For instance, a constant propagation pass can produce a virtual clone of the function which replaces one of its arguments by a constant. The inliner can represent its decisions by producing a clone of a function whose body will be later integrated into a given function.

Using *virtual clones*, the program can be easily updated during the *Execute* stage, solving most of pass interactions problems that would otherwise occur during *Transform*.

Virtual clones are later materialized in the LTRANS stage and turned into real functions. Passes executed after the virtual clone were introduced also perform their *Transform* stage on new functions, so for a pass there is no significant difference between operating on a real function or a virtual clone introduced before its *Execute* stage.

Optimization passes then work on virtual clones introduced before their *Execute* stage as if they were real functions. The only difference is that clones are not visible during the *Generate Summary* stage.

To keep function summaries updated, the callgraph interface allows an optimizer to register a callback that is called every time a new clone is introduced as well as when the actual function or variable is generated or when a function or variable is removed. These hooks are registered in the *Generate summary* stage and allow the pass to keep its information intact until the *Execute* stage. The same hooks can also be registered during the *Execute* stage to keep the optimization summaries updated for the *Transform* stage.

#### 25.3.2 IPA references

GCC represents IPA references in the callgraph. For a function or variable A, the *IPA reference* is a list of all locations where the address of A is taken and, when A is a variable, a list of all direct stores and reads to/from A. References represent an oriented multi-graph on the union of nodes of the callgraph and the varpool. See 'ipa-reference.c':ipa\_reference\_write\_optimization\_summary and 'ipa-reference.c':ipa\_reference\_read\_optimization\_summary for details.

#### 25.3.3 Jump functions

Suppose that an optimization pass sees a function A and it knows the values of (some of) its arguments. The *jump function* describes the value of a parameter of a given function call in function A based on this knowledge.

Jump functions are used by several optimizations, such as the inter-procedural constant propagation pass and the devirtualization pass. The inliner also uses jump functions to perform inlining of callbacks.

# 25.4 Whole program assumptions, linker plugin and symbol visibilities

Link-time optimization gives relatively minor benefits when used alone. The problem is that propagation of inter-procedural information does not work well across functions and variables that are called or referenced by other compilation units (such as from a dynamically linked library). We say that such functions and variables are *externally visible*.

To make the situation even more difficult, many applications organize themselves as a set of shared libraries, and the default ELF visibility rules allow one to overwrite any externally visible symbol with a different symbol at runtime. This basically disables any optimizations across such functions and variables, because the compiler cannot be sure that the function body it is seeing is the same function body that will be used at runtime. Any function or variable not declared static in the sources degrades the quality of inter-procedural optimization.

To avoid this problem the compiler must assume that it sees the whole program when doing link-time optimization. Strictly speaking, the whole program is rarely visible even at link-time. Standard system libraries are usually linked dynamically or not provided with the link-time information. In GCC, the whole program option ('-fwhole-program') asserts that every function and variable defined in the current compilation unit is static, except for function main (note: at link time, the current unit is the union of all objects compiled with LTO). Since some functions and variables need to be referenced externally, for example by another DSO or from an assembler file, GCC also provides the function and variable attribute externally\_visible which can be used to disable the effect of '-fwhole-program' on a specific symbol.

The whole program mode assumptions are slightly more complex in C++, where inline functions in headers are put into *COMDAT* sections. COMDAT function and variables can be defined by multiple object files and their bodies are unified at link-time and dynamic link-time. COMDAT functions are changed to local only when their address is not taken and thus un-sharing them with a library is not harmful. COMDAT variables always remain externally visible, however for readonly variables it is assumed that their initializers cannot be overwritten by a different value.

GCC provides the function and variable attribute visibility that can be used to specify the visibility of externally visible symbols (or alternatively an '-fdefault-visibility' command line option). ELF defines the default, protected, hidden and internal visibilities.

The most commonly used is visibility is hidden. It specifies that the symbol cannot be referenced from outside of the current shared library. Unfortunately, this information cannot be used directly by the link-time optimization in the compiler since the whole shared library also might contain non-LTO objects and those are not visible to the compiler.

GCC solves this problem using linker plugins. A *linker plugin* is an interface to the linker that allows an external program to claim the ownership of a given object file. The linker then performs the linking procedure by querying the plugin about the symbol table of the claimed objects and once the linking decisions are complete, the plugin is allowed to provide the final object file before the actual linking is made. The linker plugin obtains the symbol resolution information which specifies which symbols provided by the claimed objects are bound from the rest of a binary being linked.

GCC is designed to be independent of the rest of the toolchain and aims to support linkers without plugin support. For this reason it does not use the linker plugin by default. Instead, the object files are examined by collect2 before being passed to the linker and objects found to have LTO sections are passed to lto1 first. This mode does not work for library archives.

The decision on what object files from the archive are needed depends on the actual linking and thus GCC would have to implement the linker itself. The resolution information is missing too and thus GCC needs to make an educated guess based on '-fwhole-program'. Without the linker plugin GCC also assumes that symbols are declared hidden and not referred by non-LTO code by default.

### 25.5 Internal flags controlling 1to1

The following flags are passed into lto1 and are not meant to be used directly from the command line.

- -fwpa This option runs the serial part of the link-time optimizer performing the interprocedural propagation (WPA mode). The compiler reads in summary information from all inputs and performs an analysis based on summary information only. It generates object files for subsequent runs of the link-time optimizer where individual object files are optimized using both summary information from the WPA mode and the actual function bodies. It then drives the LTRANS phase.
- -fltrans This option runs the link-time optimizer in the local-transformation (LTRANS) mode, which reads in output from a previous run of the LTO in WPA mode. In the LTRANS mode, LTO optimizes an object and produces the final assembly.
- -fltrans-output-list=file This option specifies a file to which the names of LTRANS output files are written. This option is only meaningful in conjunction with '-fwpa'.
- -fresolution=file This option specifies the linker resolution file. This option is only meaningful in conjunction with '-fwpa' and as option to pass through to the LTO linker plugin.

# 26 Match and Simplify

The GIMPLE and GENERIC pattern matching project match-and-simplify tries to address several issues.

- 1. unify expression simplifications currently spread and duplicated over separate files like fold-const.c, gimple-fold.c and builtins.c
- 2. allow for a cheap way to implement building and simplifying non-trivial GIMPLE expressions, avoiding the need to go through building and simplifying GENERIC via fold\_buildN and then gimplifying via force\_gimple\_operand

To address these the project introduces a simple domain specific language to write expression simplifications from which code targeting GIMPLE and GENERIC is auto-generated. The GENERIC variant follows the fold\_buildN API while for the GIMPLE variant and to address 2) new APIs are introduced.

#### 26.1 GIMPLE API

```
tree gimple_simplify (enum tree_code, tree, tree, gimple_seq
                                                                   [GIMPLE function]
          *, tree (*)(tree))
tree gimple_simplify (enum tree_code, tree, tree, tree,
                                                                    [GIMPLE function]
         gimple_seq *, tree (*)(tree))
tree gimple_simplify (enum tree_code, tree, tree, tree, tree,
                                                                   [GIMPLE function]
         gimple_seq *, tree (*)(tree))
tree gimple_simplify (enum built_in_function, tree, tree,
                                                                   [GIMPLE function]
         gimple_seq *, tree (*)(tree))
tree gimple_simplify (enum built_in_function, tree, tree,
                                                                   [GIMPLE function]
         tree, gimple_seq *, tree (*)(tree))
tree gimple_simplify (enum built_in_function, tree, tree,
                                                                   [GIMPLE function]
         tree, tree, gimple_seq *, tree (*)(tree))
     The main GIMPLE API entry to the expression simplifications mimicing that of the
     GENERIC fold_{unary,binary,ternary} functions.
```

thus providing n-ary overloads for operation or function. The additional arguments are a gimple\_seq where built statements are inserted on (if NULL then simplifications requiring new statements are not performed) and a valueization hook that can be used to tie simplifications to a SSA lattice.

In addition to those APIs fold\_stmt is overloaded with a valueization hook:

```
tree gimple_build (gimple_seq *, location_t, enum built_in_function, tree, tree, tree (*valueize) (tree) = NULL);

tree gimple_build (gimple_seq *, location_t, enum [GIMPLE function] built_in_function, tree, tree, tree, tree (*valueize) (tree) = NULL);

tree gimple_build (gimple_seq *, location_t, enum [GIMPLE function] built_in_function, tree, tree, tree, tree, tree (*valueize) (tree) = NULL);

tree gimple_convert (gimple_seq *, location_t, tree, tree); [GIMPLE function] which is supposed to replace force_gimple_operand (fold_buildN (...), ...) and calls to fold_convert. Overloads without the location_t argument exist. Built statements are inserted on the provided sequence and simplification is performed using the optional valueization hook.
```

### 26.2 The Language

The language to write expression simplifications in resembles other domain-specific languages GCC uses. Thus it is lispy. Lets start with an example from the match.pd file:

```
(simplify
  (bit_and @0 integer_all_onesp)
  @0)
```

This example contains all required parts of an expression simplification. A simplification is wrapped inside a (simplify...) expression. That contains at least two operands - an expression that is matched with the GIMPLE or GENERIC IL and a replacement expression that is returned if the match was successful.

Expressions have an operator ID, bit\_and in this case. Expressions can be lower-case tree codes with \_expr stripped off or builtin function code names in all-caps, like BUILT\_IN\_SQRT.

On denotes a so-called capture. It captures the operand and lets you refer to it in other places of the match-and-simplify. In the above example it is referred to in the replacement expression. Captures are O followed by a number or an identifier.

```
(simplify
  (bit_xor @0 @0)
  { build_zero_cst (type); })
```

In this example **@0** is mentioned twice which constrains the matched expression to have two equal operands. Usually matches are constraint to equal types. If operands may be constants and conversions are involved matching by value might be preferred in which case use **@00** to denote a by value match and the specific operand you want to refer to in the result part. This example also introduces operands written in C code. These can be used in the expression replacements and are supposed to evaluate to a tree node which has to be a valid GIMPLE operand (so you cannot generate expressions in C code).

```
(simplify
  (trunc_mod integer_zerop@0 @1)
  (if (!integer_zerop (@1))
    @0))
```

Here @O captures the first operand of the trunc\_mod expression which is also predicated with integer\_zerop. Expression operands may be either expressions, predicates or captures. Captures can be unconstrained or capture expressions or predicates.

This example introduces an optional operand of simplify, the if-expression. This condition is evaluated after the expression matched in the IL and is required to evaluate to true to

enable the replacement expression in the second operand position. The expression operand of the if is a standard C expression which may contain references to captures. The if has an optional third operand which may contain the replacement expression that is enabled when the condition evaluates to false.

A if expression can be used to specify a common condition for multiple simplify patterns, avoiding the need to repeat that multiple times:

```
(if (!TYPE_SATURATING (type)
    && !FLOAT_TYPE_P (type) && !FIXED_POINT_TYPE_P (type))
(simplify
    (minus (plus @0 @1) @0)
    @1)
(simplify
    (minus (minus @0 @1) @0)
    (negate @1)))
```

Note that ifs in outer position do not have the optional else clause but instead have multiple then clauses.

Ifs can be nested.

There exists a switch expression which can be used to chain conditions avoiding nesting ifs too much:

```
(simplify
     (simple_comparison @O REAL_CST@1)
     (switch
      /* a CMP (-0) -> a CMP 0 */
      (if (REAL_VALUE_MINUS_ZERO (TREE_REAL_CST (@1)))
       (cmp @0 { build_real (TREE_TYPE (@1), dconst0); }))
      /* x != NaN is always true, other ops are always false. */
      (if (REAL_VALUE_ISNAN (TREE_REAL_CST (01))
           && ! HONOR_SNANS (@1))
       { constant_boolean_node (cmp == NE_EXPR, type); })))
Is equal to
    (simplify
     (simple_comparison @O REAL_CST@1)
     (switch
      /* a CMP (-0) -> a CMP 0 */
      (if (REAL_VALUE_MINUS_ZERO (TREE_REAL_CST (@1)))
       (cmp @0 { build_real (TREE_TYPE (@1), dconst0); })
       /* x != NaN is always true, other ops are always false. */
       (if (REAL_VALUE_ISNAN (TREE_REAL_CST (@1))
            && ! HONOR_SNANS (@1))
        { constant_boolean_node (cmp == NE_EXPR, type); }))))
```

which has the second if in the else operand of the first. The switch expression takes if expressions as operands (which may not have else clauses) and as a last operand a replacement expression which should be enabled by default if no other condition evaluated to true

Captures can also be used for capturing results of sub-expressions.

```
#if GIMPLE
(simplify
  (pointer_plus (addr@2 @0) INTEGER_CST_P@1)
  (if (is_gimple_min_invariant (@2)))
  {
    poly_int64 off;
```

In the above example, @2 captures the result of the expression (addr @0). For outermost expression only its type can be captured, and the keyword type is reserved for this purpose. The above example also gives a way to conditionalize patterns to only apply to GIMPLE or GENERIC by means of using the pre-defined preprocessor macros GIMPLE and GENERIC and using preprocessor directives.

```
(simplify
  (bit_and:c integral_op_p@0 (bit_ior:c (bit_not @0) @1))
  (bit_and @1 @0))
```

Here we introduce flags on match expressions. The flag used above, c, denotes that the expression should be also matched commutated. Thus the above match expression is really the following four match expressions:

```
(bit_and integral_op_p@0 (bit_ior (bit_not @0) @1)) (bit_and (bit_ior (bit_not @0) @1) integral_op_p@0) (bit_and integral_op_p@0 (bit_ior @1 (bit_not @0))) (bit_and (bit_ior @1 (bit_not @0)) integral_op_p@0)
```

Usual canonicalizations you know from GENERIC expressions are applied before matching, so for example constant operands always come second in commutative expressions.

The second supported flag is s which tells the code generator to fail the pattern if the expression marked with s does have more than one use and the simplification results in an expression with more than one operator. For example in

```
(simplify
  (pointer_plus (pointer_plus:s @0 @1) @3)
  (pointer_plus @0 (plus @1 @3)))
```

this avoids the association if (pointer\_plus 00 01) is used outside of the matched expression and thus it would stay live and not trivially removed by dead code elimination. Now consider ((x + 3) + -3) with the temporary holding (x + 3) used elsewhere. This simplifies down to x which is desirable and thus flagging with s does not prevent the transform. Now consider ((x + 3) + 1) which simplifies to (x + 4). Despite being flagged with s the simplification will be performed. The simplification of ((x + a) + 1) to (x + (a + 1)) will not performed in this case though.

More features exist to avoid too much repetition.

```
(for op (plus pointer_plus minus bit_ior bit_xor)
  (simplify
    (op @O integer_zerop)
    @O))
```

A for expression can be used to repeat a pattern for each operator specified, substituting op. for can be nested and a for can have multiple operators to iterate.

```
(for opa (plus minus) opb (minus plus)
```

```
(for opc (plus minus)
  (simplify...
```

In this example the pattern will be repeated four times with opa, opb, opc being plus, minus, plus; plus, minus, minus; minus, plus; minus, plus, minus.

To avoid repeating operator lists in for you can name them via

```
(define_operator_list pmm plus minus mult)
```

and use them in for operator lists where they get expanded.

```
(for opa (pmm trunc_div)
  (simplify...
```

So this example iterates over plus, minus, mult and trunc\_div.

Using operator lists can also remove the need to explicitly write a for. All operator list uses that appear in a simplify or match pattern in operator positions will implicitly be added to a new for. For example

fors and operator lists can include the special identifier null that matches nothing and can never be generated. This can be used to pad an operator list so that it has a standard form, even if there isn't a suitable operator for every form.

Another building block are with expressions in the result expression which nest the generated code in a new C block followed by its argument:

```
(simplify
  (convert (mult @0 @1))
  (with { tree utype = unsigned_type_for (type); }
    (convert (mult (convert:utype @0) (convert:utype @1)))))
```

This allows code nested in the with to refer to the declared variables. In the above case we use the feature to specify the type of a generated expression with the :type syntax where type needs to be an identifier that refers to the desired type. Usually the types of the generated result expressions are determined from the context, but sometimes like in the above case it is required that you specify them explicitely.

As intermediate conversions are often optional there is a way to avoid the need to repeat patterns both with and without such conversions. Namely you can mark a conversion as being optional with a ?:

```
(simplify
  (eq (convert@0 @1) (convert? @2))
  (eq @1 (convert @2)))
```

which will match both (eq (convert @1) (convert @2)) and (eq (convert @1) @2). The optional converts are supposed to be all either present or not, thus (eq (convert?

@1) (convert? @2)) will result in two patterns only. If you want to match all four combinations you have access to two additional conditional converts as in (eq (convert1? @1) (convert2? @2)).

The support for ? marking extends to all unary operations including predicates you declare yourself with match.

Predicates available from the GCC middle-end need to be made available explicitely via define\_predicates:

```
(define_predicates
  integer_onep integer_zerop integer_all_onesp)
```

You can also define predicates using the pattern matching language and the match form:

This shows that for match expressions there is t available which captures the outermost expression (something not possible in the simplify context). As you can see match has an identifier as first operand which is how you refer to the predicate in patterns. Multiple match for the same identifier add additional cases where the predicate matches.

Predicates can also match an expression in which case you need to provide a template specifying the identifier and where to get its operands from:

```
(match (logical_inverted_value @0)
   (eq @0 integer_zerop))
   (match (logical_inverted_value @0)
      (bit_not truth_valued_p@0))

You can use the above predicate like
   (simplify
      (bit_and @0 (logical_inverted_value @0))
      { build_zero_cst (type); })
```

Which will match a bitwise and of an operand with its logical inverted value.

# 27 Static Analyzer

# 27.1 Analyzer Internals

### 27.1.1 Overview

The analyzer implementation works on the gimple-SSA representation. (I chose this in the hopes of making it easy to work with LTO to do whole-program analysis).

The implementation is read-only: it doesn't attempt to change anything, just emit warnings.

The gimple representation can be seen using '-fdump-ipa-analyzer'.

First, we build a supergraph which combines the callgraph and all of the CFGs into a single directed graph, with both interprocedural and intraprocedural edges. The nodes and edges in the supergraph are called "supernodes" and "superedges", and often referred to in code as snodes and sedges. Basic blocks in the CFGs are split at interprocedural calls, so there can be more than one supernode per basic block. Most statements will be in just one supernode, but a call statement can appear in two supernodes: at the end of one for the call, and again at the start of another for the return.

The supergraph can be seen using '-fdump-analyzer-supergraph'.

We then build an analysis\_plan which walks the callgraph to determine which calls might be suitable for being summarized (rather than fully explored) and thus in what order to explore the functions.

Next is the heart of the analyzer: we use a worklist to explore state within the supergraph, building an "exploded graph". Nodes in the exploded graph correspond to <point, state> pairs, as in "Precise Interprocedural Dataflow Analysis via Graph Reachability" (Thomas Reps, Susan Horwitz and Mooly Sagiv).

We reuse nodes for <point, state> pairs we've already seen, and avoid tracking state too closely, so that (hopefully) we rapidly converge on a final exploded graph, and terminate the analysis. We also bail out if the number of exploded <end-of-basic-block, state> nodes gets larger than a particular multiple of the total number of basic blocks (to ensure termination in the face of pathological state-explosion cases, or bugs). We also stop exploring a point once we hit a limit of states for that point.

We can identify problems directly when processing a <point, state> instance. For example, if we're finding the successors of

```
<point: before-stmt: "free (ptr);",
  state: {"ptr": freed}>
```

then we can detect a double-free of "ptr". We can then emit a path to reach the problem by finding the simplest route through the graph.

Program points in the analysis are much more fine-grained than in the CFG and super-graph, with points (and thus potentially exploded nodes) for various events, including before individual statements. By default the exploded graph merges multiple consecutive statements in a supernode into one exploded edge to minimize the size of the exploded graph. This can be suppressed via '-fanalyzer-fine-grained'. The fine-grained approach seems to make things simpler and more debuggable that other approaches I tried, in that each point is responsible for one thing.

Program points in the analysis also have a "call string" identifying the stack of callsites below them, so that paths in the exploded graph correspond to interprocedurally valid paths: we always return to the correct call site, propagating state information accordingly. We avoid infinite recursion by stopping the analysis if a callsite appears more than analyzer-max-recursion-depth in a callstring (defaulting to 2).

# 27.1.2 Graphs

Nodes and edges in the exploded graph are called "exploded nodes" and "exploded edges" and often referred to in the code as enodes and eedges (especially when distinguishing them from the snodes and sedges in the supergraph).

Each graph numbers its nodes, giving unique identifiers - supernodes are referred to throughout dumps in the form 'SN': index' and exploded nodes in the form 'EN: index' (e.g. 'SN: 2' and 'EN:29').

The supergraph can be seen using '-fdump-analyzer-supergraph-graph'.

The exploded graph can be seen using '-fdump-analyzer-exploded-graph' and other dump options. Exploded nodes are color-coded in the .dot output based on state-machine states to make it easier to see state changes at a glance.

# 27.1.3 State Tracking

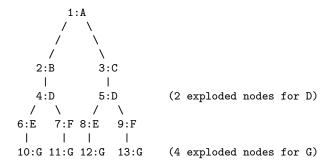
There's a tension between:

- precision of analysis in the straight-line case, vs
- exponential blow-up in the face of control flow.

For example, in general, given this CFG:



we want to avoid differences in state-tracking in B and C from leading to blow-up. If we don't prevent state blowup, we end up with exponential growth of the exploded graph like this:



Similar issues arise with loops.

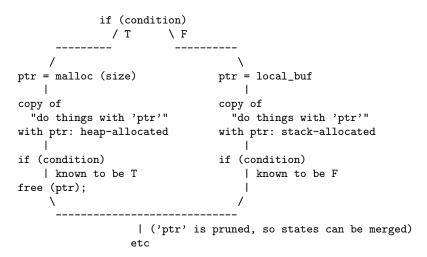
To prevent this, we follow various approaches:

- a. state pruning: which tries to discard state that won't be relevant later on withing the function. This can be disabled via '-fno-analyzer-state-purge'.
- b. state merging. We can try to find the commonality between two program\_state instances to make a third, simpler program\_state. We have two strategies here:
  - 1. the worklist keeps new nodes for the same program\_point together, and tries to merge them before processing, and thus before they have successors. Hence, in the above, the two nodes for D (4 and 5) reach the front of the worklist together, and we create a node for D with the merger of the incoming states.
  - 2. try merging with the state of existing enodes for the program\_point (which may have already been explored). There will be duplication, but only one set of duplication; subsequent duplicates are more likely to hit the cache. In particular, (hopefully) all merger chains are finite, and so we guarantee termination. This is intended to help with loops: we ought to explore the first iteration, and then have a "subsequent iterations" exploration, which uses a state merged from that of the first, to be more abstract.

We avoid merging pairs of states that have state-machine differences, as these are the kinds of differences that are likely to be most interesting. So, for example, given:

```
if (condition)
  ptr = malloc (size);
else
  ptr = local_buf;
.... do things with 'ptr'
if (condition)
  free (ptr);
....etc
```

then we end up with an exploded graph that looks like this:



where some duplication has occurred, but only for the places where the the different paths are worth exploringly separately. Merging can be disabled via '-fno-analyzer-state-merge'.

# 27.1.4 Region Model

Part of the state stored at a exploded\_node is a region\_model. This is an implementation of the region-based ternary model described in "A Memory Model for Static Analysis of C Programs" (Zhongxing Xu, Ted Kremenek, and Jian Zhang).

A region\_model encapsulates a representation of the state of memory, with a tree of region instances, along with their associated values. The representation is graph-like because values can be pointers to regions. It also stores a constraint\_manager, capturing relationships between the values.

Because each node in the exploded\_graph has a region\_model, and each of the latter is graph-like, the exploded\_graph is in some ways a graph of graphs.

Here's an example of printing a region\_model, showing the ASCII-art used to visualize the region hierarchy (colorized when printing to stderr):

```
(gdb) call debug (*this)
r0: {kind: 'root', parent: null, sval: null}
|-stack: r1: {kind: 'stack', parent: r0, sval: sv1}
| |: sval: sv1: {poisoned: uninit}
  -frame for 'test': r2: {kind: 'frame', parent: r1, sval: null, map: {'ptr_3': r3}, func-
tion: 'test', depth: 0}
| | '-'ptr_3': r3: {kind: 'map', parent: r2, sval: sv3, type: 'void *', map: {}}
        |: sval: sv3: {type: 'void *', unknown}
        |: type: 'void *'
  '-frame for 'calls_malloc': r4: {kind: 'frame', parent: r1, sval: null, map: {'result_3': r7, '_4': r8
mous>': r5}, function: 'calls_malloc', depth: 1}
     |-'<anonymous>': r5: {kind: 'map', parent: r4, sval: sv4, type: 'void *', map: {}}
     | |: sval: sv4: {type: 'void *', &r6}
     | |: type: 'void *'
    |-'result_3': r7: {kind: 'map', parent: r4, sval: sv4, type: 'void *', map: {}}
    | |: sval: sv4: {type: 'void *', &r6}
     | |: type: 'void *
     '-'_4': r8: {kind: 'map', parent: r4, sval: sv4, type: 'void *', map: {}}
      |: sval: sv4: {type: 'void *', &r6}
      |: type: 'void *'
'-heap: r9: {kind: 'heap', parent: r0, sval: sv2}
  |: sval: sv2: {poisoned: uninit}
  '-r6: {kind: 'symbolic', parent: r9, sval: null, map: {}}
svalues:
 sv0: {type: 'size_t', '1024'}
 sv1: {poisoned: uninit}
 sv2: {poisoned: uninit}
 sv3: {type: 'void *', unknown}
  sv4: {type: 'void *', &r6}
constraint manager:
  equiv classes:
    ec0: {sv0 == '1024'}
    ec1: {sv4}
```

This is the state at the point of returning from calls\_malloc back to test in the following:

```
void *
calls_malloc (void)
{
```

```
void *result = malloc (1024);
return result;
}

void test (void)
{
  void *ptr = calls_malloc ();
  /* etc. */
}
```

The "root" region ("r0") has a "stack" child ("r1"), with two children: a frame for test ("r2"), and a frame for calls\_malloc ("r4"). These frame regions have child regions for storing their local variables. For example, the return region and that of various other regions within the "calls\_malloc" frame all have value "sv4", a pointer to a heap-allocated region "r6". Within the parent frame, ptr\_3 has value "sv3", an unknown void \*.

### 27.1.5 Analyzer Paths

We need to explain to the user what the problem is, and to persuade them that there really is a problem. Hence having a diagnostic\_path isn't just an incidental detail of the analyzer; it's required.

Paths ought to be:

- interprocedurally-valid
- feasible

Without state-merging, all paths in the exploded graph are feasible (in terms of constraints being satisified). With state-merging, paths in the exploded graph can be infeasible.

We collate warnings and only emit them for the simplest path e.g. for a bug in a utility function, with lots of routes to calling it, we only emit the simplest path (which could be intraprocedural, if it can be reproduced without a caller). We apply a check that each duplicate warning's shortest path is feasible, rejecting any warnings for which the shortest path is infeasible (which could lead to false negatives).

We use the shortest feasible exploded\_path through the exploded\_graph (a list of exploded\_edge \*) to build a diagnostic\_path (a list of events for the diagnostic subsystem) - specifically a checker\_path.

Having built the checker\_path, we prune it to try to eliminate events that aren't relevant, to minimize how much the user has to read.

After pruning, we notify each event in the path of its ID and record the IDs of interesting events, allowing for events to refer to other events in their descriptions. The pending\_diagnostic class has various vfuncs to support emitting more precise descriptions, so that e.g.

• a deref-of-unchecked-malloc diagnostic might use:

```
returning possibly-NULL pointer to 'make_obj' from 'allocator'
```

for a **return\_event** to make it clearer how the unchecked value moves from callee back to caller

• a double-free diagnostic might use:

```
second 'free' here; first 'free' was at (3) and a use-after-free might use
```

```
use after 'free' here; memory was freed at (2)
```

At this point we can emit the diagnostic.

### 27.1.6 Limitations

- Only for C so far
- The implementation of call summaries is currently very simplistic.
- Lack of function pointer analysis
- The constraint-handling code assumes reflexivity in some places (that values are equal to themselves), which is not the case for NaN. As a simple workaround, constraints on floating-point values are currently ignored.
- The region model code creates lots of little mutable objects at each region\_model (and thus per exploded\_node) rather than sharing immutable objects and having the mutable state in the program\_state or region\_model. The latter approach might be more efficient, and might avoid dealing with IDs rather than pointers (which requires us to impose an ordering to get meaningful equality).
- The region model code doesn't yet support memcpy. At the gimple-ssa level these have been optimized to statements like this:

```
_10 = MEM <long unsigned int> [(char * {ref-all})&c]
MEM <long unsigned int> [(char * {ref-all})&d] = _10;
```

Perhaps they could be supported via a new compound\_svalue type.

- There are various other limitations in the region model (grep for TODO/xfail in the testsuite).
- The constraint\_manager's implementation of transitivity is currently too expensive to enable by default and so must be manually enabled via '-fanalyzer-transitivity').
- The checkers are currently hardcoded and don't allow for user extensibility (e.g. adding allocate/release pairs).
- Although the analyzer's test suite has a proof-of-concept test case for LTO, LTO support hasn't had extensive testing. There are various lang-specific things in the analyzer that assume C rather than LTO. For example, SSA names are printed to the user in "raw" form, rather than printing the underlying variable name.

Some ideas for other checkers

- File-descriptor-based APIs
- Linux kernel internal APIs
- Signal handling

# 27.2 Debugging the Analyzer

# 27.2.1 Special Functions for Debugging the Analyzer

The analyzer recognizes various special functions by name, for use in debugging the analyzer. Declarations can be seen in the testsuite in 'analyzer-decls.h'. None of these functions are actually implemented.

Add:

```
__analyzer_break ();
```

to the source being analyzed to trigger a breakpoint in the analyzer when that source is reached. By putting a series of these in the source, it's much easier to effectively step through the program state as it's analyzed.

```
__analyzer_dump ();
```

will dump the copious information about the analyzer's state each time it reaches the call in its traversal of the source.

```
__analyzer_dump_path ();
```

will emit a placeholder "note" diagnostic with a path to that call site, if the analyzer finds a feasible path to it.

The builtin \_\_analyzer\_dump\_exploded\_nodes will emit a warning after analysis containing information on all of the exploded nodes at that program point:

```
__analyzer_dump_exploded_nodes (0);
```

will output the number of "processed" nodes, and the IDs of both "processed" and "merger" nodes, such as:

```
warning: 2 processed enodes: [EN: 56, EN: 58] merger(s): [EN: 54-55, EN: 57, EN: 59]
With a non-zero argument
    __analyzer_dump_exploded_nodes (1);
it will also dump all of the states within the "processed" nodes.
    __analyzer_dump_region_model ();
will dump the region_model's state to stderr.
    __analyzer_eval (expr);
```

will emit a warning with text "TRUE", FALSE" or "UNKNOWN" based on the truth-fulness of the argument. This is useful for writing DejaGnu tests.

# 27.2.2 Other Debugging Techniques

One approach when tracking down where a particular bogus state is introduced into the exploded\_graph is to add custom code to region\_model::validate.

For example, this custom code (added to region\_model::validate) breaks with an assertion failure when a variable called ptr acquires a value that's unknown, using region\_model::get\_value\_by\_name to locate the variable

```
/* Find a variable matching "ptr". */
    svalue_id sid = get_value_by_name ("ptr");
    if (!sid.null_p ())
     {
    svalue *sval = get_svalue (sid);
    gcc_assert (sval->get_kind () != SK_UNKNOWN);
    }
```

making it easier to investigate further in a debugger when this occurs.

# 28 User Experience Guidelines

To borrow a slogan from Elm,

Compilers should be assistants, not adversaries. A compiler should not just detect bugs, it should then help you understand why there is a bug. It should not berate you in a robot voice, it should give you specific hints that help you write better code. Ultimately, a compiler should make programming faster and more fun!

—Evan Czaplicki

This chapter provides guidelines on how to implement diagnostics and command-line options in ways that we hope achieve the above ideal.

# 28.1 Guidelines for Diagnostics

### 28.1.1 Talk in terms of the user's code

Diagnostics should be worded in terms of the user's source code, and the source language, rather than GCC's own implementation details.

# 28.1.2 Diagnostics are actionable

A good diagnostic is actionable: it should assist the user in taking action.

Consider what an end user will want to do when encountering a diagnostic.

Given an error, an end user will think: "How do I fix this?"

Given a warning, an end user will think:

- "Is this a real problem?"
- "Do I care?"
- if they decide it's genuine: "How do I fix this?"

A good diagnostic provides pertinent information to allow the user to easily answer the above questions.

# 28.1.3 The user's attention is important

A perfect compiler would issue a warning on every aspect of the user's source code that ought to be fixed, and issue no other warnings. Naturally, this ideal is impossible to achieve.

Warnings should have a good signal-to-noise ratio: we should have few false positives (falsely issuing a warning when no warning is warranted) and few false negatives (failing to issue a warning when one is justified).

Note that a false positive can mean, in practice, a warning that the user doesn't agree with. Ideally a diagnostic should contain enough information to allow the user to make an informed choice about whether they should care (and how to fix it), but a balance must be drawn against overloading the user with irrelevant data.

# 28.1.4 Precision of Wording

Provide the user with details that allow them to identify what the problem is. For example, the vaguely-worded message:

doesn't tell the user why the attribute was ignored, or what kind of entity the compiler thought the attribute was being applied to (the source location for the diagnostic is also poor; see [discussion of input\_location], page 718). A better message would be:

which spells out the missing information (and fixes the location information, as discussed below).

The above example uses a note to avoid a combinatorial explosion of possible messages.

# 28.1.5 Try the diagnostic on real-world code

It's worth testing a new warning on many instances of real-world code, written by different people, and seeing what it complains about, and what it doesn't complain about.

This may suggest heuristics that silence common false positives.

It may also suggest ways to improve the precision of the message.

### 28.1.6 Make mismatches clear

Many diagnostics relate to a mismatch between two different places in the user's source code. Examples include:

- a type mismatch, where the type at a usage site does not match the type at a declaration
- the argument count at a call site does not match the parameter count at the declaration
- something is erroneously duplicated (e.g. an error, due to breaking a uniqueness requirement, or a warning, if it's suggestive of a bug)
- an "opened" syntactic construct (such as an open-parenthesis) is not closed

In each case, the diagnostic should indicate **both** pertinent locations (so that the user can easily see the problem and how to fix it).

The standard way to do this is with a note (via inform). For example:

The inform call should be guarded by the return value from the warning\_at call so that the note isn't emitted when the warning is suppressed.

For cases involving punctuation where the locations might be near each other, they can be conditionally consolidated via gcc\_rich\_location::add\_location\_if\_nearby:

### 28.1.7 Location Information

GCC's location\_t type can support both ordinary locations, and locations relating to a macro expansion.

As of GCC 6, ordinary locations changed from supporting just a point in the user's source code to supporting three points: the *caret* location, plus a start and a finish:

Tokens coming out of libcpp have locations of the form caret == start, such as for foo here:

Compound expressions should be reported using the location of the expression as a whole, rather than just of one token within it.

For example, in -Wformat, rather than underlining just the first token of a bad argument:

the whole of the expression should be underlined, so that the user can easily identify what is being referred to:

Avoid using the input\_location global, and the diagnostic functions that implicitly use it—use error\_at and warning\_at rather than error and warning, and provide the most appropriate location\_t value available at that phase of the compilation. It's possible to supply secondary location\_t values via rich\_location.

For example, in the example of imprecise wording above, generating the diagnostic using warning:

which thus happened to use the location of the int token, rather than that of the attribute. Using warning\_at with the location of the attribute, providing the location of the declaration in question as a secondary location, and adding a note:

## 28.1.8 Coding Conventions

See the diagnostics section of the GCC coding conventions.

In the C++ front end, when comparing two types in a message, use '%H' and '%I' rather than '%T', as this allows the diagnostics subsystem to highlight differences between template-based types. For example, rather than using '%qT':

### 28.1.9 Group logically-related diagnostics

Use auto\_diagnostic\_group when issuing multiple related diagnostics (seen in various examples on this page). This informs the diagnostic subsystem that all diagnostics issued within the lifetime of the auto\_diagnostic\_group are related. For example, '-fdiagnostics-format=json' will treat the first diagnostic emitted within the group as a top-level diagnostic, and all subsequent diagnostics within the group as its children.

# 28.1.10 Quoting

Text should be quoted by either using the 'q' modifier in a directive such as '%qE', or by enclosing the quoted text in a pair of '%<' and '%>' directives, and never by using explicit quote characters. The directives handle the appropriate quote characters for each language and apply the correct color or highlighting.

The following elements should be quoted in GCC diagnostics:

- Language keywords.
- Tokens.
- Boolean, numerical, character, and string constants that appear in the source code.
- Identifiers, including function, macro, type, and variable names.

Other elements such as numbers that do not refer to numeric constants that appear in the source code should not be quoted. For example, in the message:

```
argument %d of %qE must be a pointer type
```

since the argument number does not refer to a numerical constant in the source code it should not be quoted.

# 28.1.11 Spelling and Terminology

See the terminology and markup section of the GCC coding conventions.

### 28.1.12 Fix-it hints

GCC's diagnostic subsystem can emit fix-it hints: small suggested edits to the user's source code.

They are printed by default underneath the code in question. They can also be viewed via '-fdiagnostics-generate-patch' and '-fdiagnostics-parseable-fixits'. With the latter, an IDE ought to be able to offer to automatically apply the suggested fix.

Fix-it hints contain code fragments, and thus they should not be marked for translation.

Fix-it hints can be added to a diagnostic by using a rich\_location rather than a location\_t - the fix-it hints are added to the rich\_location using one of the various add\_fixit member functions of rich\_location. They are documented with rich\_location in 'libcpp/line-map.h'. It's easiest to use the gcc\_rich\_location subclass of rich\_location found in 'gcc-rich-location.h', as this implicitly supplies the line\_table variable.

For example:

```
if (const char *suggestion = hint.suggestion ())
  {
    gcc_rich_location richloc (location);
    richloc.add_fixit_replace (suggestion);
```

Non-trivial edits can be built up by adding multiple fix-it hints to one rich\_location. It's best to express the edits in terms of the locations of individual tokens. Various handy functions for adding fix-it hints for idiomatic C and C++ can be seen in 'gcc-rich-location.h'.

### 28.1.12.1 Fix-it hints should work

When implementing a fix-it hint, please verify that the suggested edit leads to fixed, compilable code. (Unfortunately, this currently must be done by hand using '-fdiagnostics-generate-patch'. It would be good to have an automated way of verifying that fix-it hints actually fix the code).

For example, a "gotcha" here is to forget to add a space when adding a missing reserved word. Consider a C++ fix-it hint that adds typename in front of a template declaration. A naive way to implement this might be:

```
gcc_rich_location richloc (loc);
      // BAD: insertion is missing a trailing space
      richloc.add_fixit_insert_before ("typename");
      error_at (&richloc, "need %<typename%> before %<%T::%E%> because "
                           "%qT is a dependent scope",
                           parser->scope, id, parser->scope);
When applied to the code, this might lead to:
      T::type x;
being "corrected" to:
      typenameT::type x;
In this case, the correct thing to do is to add a trailing space after typename:
      gcc_rich_location richloc (loc);
      // OK: note that here we have a trailing space
      richloc.add_fixit_insert_before ("typename");
      error_at (&richloc, "need %<typename%> before %<%T::%E%> because "
                           "%qT is a dependent scope",
                           parser->scope, id, parser->scope);
leading to this corrected code:
      typename T::type x;
```

# 28.1.12.2 Express deletion in terms of deletion, not replacement

It's best to express deletion suggestions in terms of deletion fix-it hints, rather than replacement fix-it hints. For example, consider this:

```
auto_diagnostic_group d;
gcc_rich_location richloc (location_of (retval));
tree name = DECL_NAME (arg);
richloc.add_fixit_replace (IDENTIFIER_POINTER (name));
```

where the change has been expressed as replacement, replacing with the name of the declaration. This works for simple cases, but consider this case:

```
#ifdef SOME_CONFIG_FLAG
# define CONFIGURY_GLOBAL global_a
#else
# define CONFIGURY_GLOBAL global_b
#endif
int fn ()
{
   return std::move (CONFIGURY_GLOBAL /* some comment */);
}
```

The above implementation erroneously strips out the macro and the comment in the fix-it hint:

```
return std::move (CONFIGURY_GLOBAL /* some comment */);
global_a
this regulting code:
```

and thus this resulting code:

```
return global_a;
```

It's better to do deletions in terms of deletions; deleting the std::move ( and the trailing close-paren, leading to this:

```
return std::move (CONFIGURY_GLOBAL /* some comment */);

CONFIGURY_GLOBAL /* some comment */

and thus this result:

return CONFIGURY_GLOBAL /* some comment */;
```

Unfortunately, the pertinent location\_t values are not always available.

# 28.1.12.3 Multiple suggestions

In the rare cases where you need to suggest more than one mutually exclusive solution to a problem, this can be done by emitting multiple notes and calling rich\_location::fixits\_cannot\_be\_auto\_applied on each note's rich\_location. If this is called, then the fix-it hints in the rich\_location will be printed, but will not be added to generated patches.

# 28.2 Guidelines for Options

# Funding Free Software

If you want to have more free software a few years from now, it makes sense for you to help encourage people to contribute funds for its development. The most effective approach known is to encourage commercial redistributors to donate.

Users of free software systems can boost the pace of development by encouraging forafee distributors to donate part of their selling price to free software developers—the Free Software Foundation, and others.

The way to convince distributors to do this is to demand it and expect it from them. So when you compare distributors, judge them partly by how much they give to free software development. Show distributors they must compete to be the one who gives the most.

To make this approach work, you must insist on numbers that you can compare, such as, "We will donate ten dollars to the Frobnitz project for each disk sold." Don't be satisfied with a vague promise, such as "A portion of the profits are donated," since it doesn't give a basis for comparison.

Even a precise fraction "of the profits from this disk" is not very meaningful, since creative accounting and unrelated business decisions can greatly alter what fraction of the sales price counts as profit. If the price you pay is \$50, ten percent of the profit is probably less than a dollar; it might be a few cents, or nothing at all.

Some redistributors do development work themselves. This is useful too; but to keep everyone honest, you need to inquire how much they do, and what kind. Some kinds of development make much more long-term difference than others. For example, maintaining a separate version of a program contributes very little; maintaining the standard version of a program for the whole community contributes much. Easy new ports contribute little, since someone else would surely do them; difficult ports such as adding a new CPU to the GNU Compiler Collection contribute more; major new features or packages contribute the most.

By establishing the idea that supporting further development is "the proper thing to do" when distributing free software for a fee, we can assure a steady flow of resources into making more free software.

Copyright © 1994 Free Software Foundation, Inc. Verbatim copying and redistribution of this section is permitted without royalty; alteration is not permitted.

# The GNU Project and GNU/Linux

The GNU Project was launched in 1984 to develop a complete Unix-like operating system which is free software: the GNU system. (GNU is a recursive acronym for "GNU's Not Unix"; it is pronounced "guh-NEW".) Variants of the GNU operating system, which use the kernel Linux, are now widely used; though these systems are often referred to as "Linux", they are more accurately called GNU/Linux systems.

For more information, see:

```
http://www.gnu.org/
http://www.gnu.org/gnu/linux-and-gnu.html
```

# GNU General Public License

Version 3, 29 June 2007

Copyright © 2007 Free Software Foundation, Inc. http://fsf.org/

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

### Preamble

The GNU General Public License is a free, copyleft license for software and other kinds of works.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change all versions of a program—to make sure it remains free software for all its users. We, the Free Software Foundation, use the GNU General Public License for most of our software; it applies also to any other work released this way by its authors. You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things.

To protect your rights, we need to prevent others from denying you these rights or asking you to surrender the rights. Therefore, you have certain responsibilities if you distribute copies of the software, or if you modify it: responsibilities to respect the freedom of others.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must pass on to the recipients the same freedoms that you received. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

Developers that use the GNU GPL protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License giving you legal permission to copy, distribute and/or modify it.

For the developers' and authors' protection, the GPL clearly explains that there is no warranty for this free software. For both users' and authors' sake, the GPL requires that modified versions be marked as changed, so that their problems will not be attributed erroneously to authors of previous versions.

Some devices are designed to deny users access to install or run modified versions of the software inside them, although the manufacturer can do so. This is fundamentally incompatible with the aim of protecting users' freedom to change the software. The systematic pattern of such abuse occurs in the area of products for individuals to use, which is precisely where it is most unacceptable. Therefore, we have designed this version of the GPL to prohibit the practice for those products. If such problems arise substantially in other domains, we stand ready to extend this provision to those domains in future versions of the GPL, as needed to protect the freedom of users.

Finally, every program is threatened constantly by software patents. States should not allow patents to restrict development and use of software on general-purpose computers, but in those that do, we wish to avoid the special danger that patents applied to a free program could make it effectively proprietary. To prevent this, the GPL assures that patents cannot be used to render the program non-free.

The precise terms and conditions for copying, distribution and modification follow.

### TERMS AND CONDITIONS

#### 0. Definitions.

"This License" refers to version 3 of the GNU General Public License.

"Copyright" also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

"The Program" refers to any copyrightable work licensed under this License. Each licensee is addressed as "you". "Licensees" and "recipients" may be individuals or organizations.

To "modify" a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting work is called a "modified version" of the earlier work or a work "based on" the earlier work.

A "covered work" means either the unmodified Program or a work based on the Program.

To "propagate" a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To "convey" a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays "Appropriate Legal Notices" to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

#### 1. Source Code.

The "source code" for a work means the preferred form of the work for making modifications to it. "Object code" means any non-source form of a work.

A "Standard Interface" means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language. The "System Libraries" of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A "Major Component", in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The "Corresponding Source" for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work's System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding Source includes interface definition files associated with source files for the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

### 2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

### 3. Protecting Users' Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

### 4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-permissive terms added in accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

### 5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

- a. The work must carry prominent notices stating that you modified it, and giving a relevant date.
- b. The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to "keep intact all notices".
- c. You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.
- d. If the work has interactive user interfaces, each must display Appropriate Legal Notices; however, if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or on a volume of a storage or distribution medium, is called an "aggregate" if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation's users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

### 6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

- a. Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.
- b. Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.
- c. Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.
- d. Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.
- e. Convey the object code using peer-to-peer transmission, provided you inform other peers where the object code and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A "User Product" is either (1) a "consumer product", which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, "normally used" refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

"Installation Information" for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source.

The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key for unpacking, reading or copying.

### 7. Additional Terms.

"Additional permissions" are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permissions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

- a. Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or
- b. Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it: or
- c. Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or

- d. Limiting the use for publicity purposes of names of licensors or authors of the material; or
- e. Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or
- f. Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered "further restrictions" within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

### 8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

### 9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance.

However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

### 10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An "entity transaction" is a transaction transferring control of an organization, or substantially all assets of one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party's predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate litigation (including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

### 11. Patents.

A "contributor" is a copyright holder who authorizes use under this License of the Program or a work on which the Program is based. The work thus licensed is called the contributor's "contributor version".

A contributor's "essential patent claims" are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, "control" includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor's essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a "patent license" is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To "grant" such a patent license to a party means to make such an agreement or commitment not to enforce a patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so

available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. "Knowingly relying" means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient's use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is "discriminatory" if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted under this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

### 12. No Surrender of Others' Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

### 13. Use with the GNU Affero General Public License.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU Affero General Public License into a single combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the special requirements of the GNU Affero General Public License, section 13, concerning interaction through a network will apply to the combination as such.

### 14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU General Public License "or any later version" applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU General Public License can be used, that proxy's public statement of acceptance of a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

### 15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

#### 16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

### 17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.

### END OF TERMS AND CONDITIONS

# How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively state the exclusion of warranty; and each file should have at least the "copyright" line and a pointer to where the full notice is found.

one line to give the program's name and a brief idea of what it does. Copyright (C) year name of author

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see http://www.gnu.org/licenses/.

Also add information on how to contact you by electronic and paper mail.

If the program does terminal interaction, make it output a short notice like this when it starts in an interactive mode:

```
program Copyright (C) year name of author
This program comes with ABSOLUTELY NO WARRANTY; for details type 'show w'.
This is free software, and you are welcome to redistribute it
under certain conditions; type 'show c' for details.
```

The hypothetical commands 'show w' and 'show c' should show the appropriate parts of the General Public License. Of course, your program's commands might be different; for a GUI interface, you would use an "about box".

You should also get your employer (if you work as a programmer) or school, if any, to sign a "copyright disclaimer" for the program, if necessary. For more information on this, and how to apply and follow the GNU GPL, see http://www.gnu.org/licenses/.

The GNU General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Lesser General Public License instead of this License. But first, please read https://www.gnu.org/licenses/why-not-lgpl.html.

# **GNU Free Documentation License**

Version 1.3, 3 November 2008

Copyright © 2000, 2001, 2002, 2007, 2008 Free Software Foundation, Inc. http://fsf.org/

Everyone is permitted to copy and distribute verbatim copies of this license document, but changing it is not allowed.

### 0. PREAMBLE

The purpose of this License is to make a manual, textbook, or other functional and useful document free in the sense of freedom: to assure everyone the effective freedom to copy and redistribute it, with or without modifying it, either commercially or non-commercially. Secondarily, this License preserves for the author and publisher a way to get credit for their work, while not being considered responsible for modifications made by others.

This License is a kind of "copyleft", which means that derivative works of the document must themselves be free in the same sense. It complements the GNU General Public License, which is a copyleft license designed for free software.

We have designed this License in order to use it for manuals for free software, because free software needs free documentation: a free program should come with manuals providing the same freedoms that the software does. But this License is not limited to software manuals; it can be used for any textual work, regardless of subject matter or whether it is published as a printed book. We recommend this License principally for works whose purpose is instruction or reference.

### 1. APPLICABILITY AND DEFINITIONS

This License applies to any manual or other work, in any medium, that contains a notice placed by the copyright holder saying it can be distributed under the terms of this License. Such a notice grants a world-wide, royalty-free license, unlimited in duration, to use that work under the conditions stated herein. The "Document", below, refers to any such manual or work. Any member of the public is a licensee, and is addressed as "you". You accept the license if you copy, modify or distribute the work in a way requiring permission under copyright law.

A "Modified Version" of the Document means any work containing the Document or a portion of it, either copied verbatim, or with modifications and/or translated into another language.

A "Secondary Section" is a named appendix or a front-matter section of the Document that deals exclusively with the relationship of the publishers or authors of the Document to the Document's overall subject (or to related matters) and contains nothing that could fall directly within that overall subject. (Thus, if the Document is in part a textbook of mathematics, a Secondary Section may not explain any mathematics.) The relationship could be a matter of historical connection with the subject or with related matters, or of legal, commercial, philosophical, ethical or political position regarding them.

The "Invariant Sections" are certain Secondary Sections whose titles are designated, as being those of Invariant Sections, in the notice that says that the Document is released

under this License. If a section does not fit the above definition of Secondary then it is not allowed to be designated as Invariant. The Document may contain zero Invariant Sections. If the Document does not identify any Invariant Sections then there are none.

The "Cover Texts" are certain short passages of text that are listed, as Front-Cover Texts or Back-Cover Texts, in the notice that says that the Document is released under this License. A Front-Cover Text may be at most 5 words, and a Back-Cover Text may be at most 25 words.

A "Transparent" copy of the Document means a machine-readable copy, represented in a format whose specification is available to the general public, that is suitable for revising the document straightforwardly with generic text editors or (for images composed of pixels) generic paint programs or (for drawings) some widely available drawing editor, and that is suitable for input to text formatters or for automatic translation to a variety of formats suitable for input to text formatters. A copy made in an otherwise Transparent file format whose markup, or absence of markup, has been arranged to thwart or discourage subsequent modification by readers is not Transparent. An image format is not Transparent if used for any substantial amount of text. A copy that is not "Transparent" is called "Opaque".

Examples of suitable formats for Transparent copies include plain ASCII without markup, Texinfo input format, LaT<sub>E</sub>X input format, SGML or XML using a publicly available DTD, and standard-conforming simple HTML, PostScript or PDF designed for human modification. Examples of transparent image formats include PNG, XCF and JPG. Opaque formats include proprietary formats that can be read and edited only by proprietary word processors, SGML or XML for which the DTD and/or processing tools are not generally available, and the machine-generated HTML, PostScript or PDF produced by some word processors for output purposes only.

The "Title Page" means, for a printed book, the title page itself, plus such following pages as are needed to hold, legibly, the material this License requires to appear in the title page. For works in formats which do not have any title page as such, "Title Page" means the text near the most prominent appearance of the work's title, preceding the beginning of the body of the text.

The "publisher" means any person or entity that distributes copies of the Document to the public.

A section "Entitled XYZ" means a named subunit of the Document whose title either is precisely XYZ or contains XYZ in parentheses following text that translates XYZ in another language. (Here XYZ stands for a specific section name mentioned below, such as "Acknowledgements", "Dedications", "Endorsements", or "History".) To "Preserve the Title" of such a section when you modify the Document means that it remains a section "Entitled XYZ" according to this definition.

The Document may include Warranty Disclaimers next to the notice which states that this License applies to the Document. These Warranty Disclaimers are considered to be included by reference in this License, but only as regards disclaiming warranties: any other implication that these Warranty Disclaimers may have is void and has no effect on the meaning of this License.

### 2. VERBATIM COPYING

You may copy and distribute the Document in any medium, either commercially or noncommercially, provided that this License, the copyright notices, and the license notice saying this License applies to the Document are reproduced in all copies, and that you add no other conditions whatsoever to those of this License. You may not use technical measures to obstruct or control the reading or further copying of the copies you make or distribute. However, you may accept compensation in exchange for copies. If you distribute a large enough number of copies you must also follow the conditions in section 3.

You may also lend copies, under the same conditions stated above, and you may publicly display copies.

#### 3. COPYING IN QUANTITY

If you publish printed copies (or copies in media that commonly have printed covers) of the Document, numbering more than 100, and the Document's license notice requires Cover Texts, you must enclose the copies in covers that carry, clearly and legibly, all these Cover Texts: Front-Cover Texts on the front cover, and Back-Cover Texts on the back cover. Both covers must also clearly and legibly identify you as the publisher of these copies. The front cover must present the full title with all words of the title equally prominent and visible. You may add other material on the covers in addition. Copying with changes limited to the covers, as long as they preserve the title of the Document and satisfy these conditions, can be treated as verbatim copying in other respects.

If the required texts for either cover are too voluminous to fit legibly, you should put the first ones listed (as many as fit reasonably) on the actual cover, and continue the rest onto adjacent pages.

If you publish or distribute Opaque copies of the Document numbering more than 100, you must either include a machine-readable Transparent copy along with each Opaque copy, or state in or with each Opaque copy a computer-network location from which the general network-using public has access to download using public-standard network protocols a complete Transparent copy of the Document, free of added material. If you use the latter option, you must take reasonably prudent steps, when you begin distribution of Opaque copies in quantity, to ensure that this Transparent copy will remain thus accessible at the stated location until at least one year after the last time you distribute an Opaque copy (directly or through your agents or retailers) of that edition to the public.

It is requested, but not required, that you contact the authors of the Document well before redistributing any large number of copies, to give them a chance to provide you with an updated version of the Document.

#### 4. MODIFICATIONS

You may copy and distribute a Modified Version of the Document under the conditions of sections 2 and 3 above, provided that you release the Modified Version under precisely this License, with the Modified Version filling the role of the Document, thus licensing distribution and modification of the Modified Version to whoever possesses a copy of it. In addition, you must do these things in the Modified Version:

A. Use in the Title Page (and on the covers, if any) a title distinct from that of the Document, and from those of previous versions (which should, if there were any,

- be listed in the History section of the Document). You may use the same title as a previous version if the original publisher of that version gives permission.
- B. List on the Title Page, as authors, one or more persons or entities responsible for authorship of the modifications in the Modified Version, together with at least five of the principal authors of the Document (all of its principal authors, if it has fewer than five), unless they release you from this requirement.
- C. State on the Title page the name of the publisher of the Modified Version, as the publisher.
- D. Preserve all the copyright notices of the Document.
- E. Add an appropriate copyright notice for your modifications adjacent to the other copyright notices.
- F. Include, immediately after the copyright notices, a license notice giving the public permission to use the Modified Version under the terms of this License, in the form shown in the Addendum below.
- G. Preserve in that license notice the full lists of Invariant Sections and required Cover Texts given in the Document's license notice.
- H. Include an unaltered copy of this License.
- I. Preserve the section Entitled "History", Preserve its Title, and add to it an item stating at least the title, year, new authors, and publisher of the Modified Version as given on the Title Page. If there is no section Entitled "History" in the Document, create one stating the title, year, authors, and publisher of the Document as given on its Title Page, then add an item describing the Modified Version as stated in the previous sentence.
- J. Preserve the network location, if any, given in the Document for public access to a Transparent copy of the Document, and likewise the network locations given in the Document for previous versions it was based on. These may be placed in the "History" section. You may omit a network location for a work that was published at least four years before the Document itself, or if the original publisher of the version it refers to gives permission.
- K. For any section Entitled "Acknowledgements" or "Dedications", Preserve the Title of the section, and preserve in the section all the substance and tone of each of the contributor acknowledgements and/or dedications given therein.
- L. Preserve all the Invariant Sections of the Document, unaltered in their text and in their titles. Section numbers or the equivalent are not considered part of the section titles.
- M. Delete any section Entitled "Endorsements". Such a section may not be included in the Modified Version.
- N. Do not retitle any existing section to be Entitled "Endorsements" or to conflict in title with any Invariant Section.
- O. Preserve any Warranty Disclaimers.

If the Modified Version includes new front-matter sections or appendices that qualify as Secondary Sections and contain no material copied from the Document, you may at your option designate some or all of these sections as invariant. To do this, add their titles to the list of Invariant Sections in the Modified Version's license notice. These titles must be distinct from any other section titles.

You may add a section Entitled "Endorsements", provided it contains nothing but endorsements of your Modified Version by various parties—for example, statements of peer review or that the text has been approved by an organization as the authoritative definition of a standard.

You may add a passage of up to five words as a Front-Cover Text, and a passage of up to 25 words as a Back-Cover Text, to the end of the list of Cover Texts in the Modified Version. Only one passage of Front-Cover Text and one of Back-Cover Text may be added by (or through arrangements made by) any one entity. If the Document already includes a cover text for the same cover, previously added by you or by arrangement made by the same entity you are acting on behalf of, you may not add another; but you may replace the old one, on explicit permission from the previous publisher that added the old one.

The author(s) and publisher(s) of the Document do not by this License give permission to use their names for publicity for or to assert or imply endorsement of any Modified Version.

#### 5. COMBINING DOCUMENTS

You may combine the Document with other documents released under this License, under the terms defined in section 4 above for modified versions, provided that you include in the combination all of the Invariant Sections of all of the original documents, unmodified, and list them all as Invariant Sections of your combined work in its license notice, and that you preserve all their Warranty Disclaimers.

The combined work need only contain one copy of this License, and multiple identical Invariant Sections may be replaced with a single copy. If there are multiple Invariant Sections with the same name but different contents, make the title of each such section unique by adding at the end of it, in parentheses, the name of the original author or publisher of that section if known, or else a unique number. Make the same adjustment to the section titles in the list of Invariant Sections in the license notice of the combined work.

In the combination, you must combine any sections Entitled "History" in the various original documents, forming one section Entitled "History"; likewise combine any sections Entitled "Acknowledgements", and any sections Entitled "Dedications". You must delete all sections Entitled "Endorsements."

#### 6. COLLECTIONS OF DOCUMENTS

You may make a collection consisting of the Document and other documents released under this License, and replace the individual copies of this License in the various documents with a single copy that is included in the collection, provided that you follow the rules of this License for verbatim copying of each of the documents in all other respects.

You may extract a single document from such a collection, and distribute it individually under this License, provided you insert a copy of this License into the extracted document, and follow this License in all other respects regarding verbatim copying of that document.

#### 7. AGGREGATION WITH INDEPENDENT WORKS

A compilation of the Document or its derivatives with other separate and independent documents or works, in or on a volume of a storage or distribution medium, is called an "aggregate" if the copyright resulting from the compilation is not used to limit the legal rights of the compilation's users beyond what the individual works permit. When the Document is included in an aggregate, this License does not apply to the other works in the aggregate which are not themselves derivative works of the Document.

If the Cover Text requirement of section 3 is applicable to these copies of the Document, then if the Document is less than one half of the entire aggregate, the Document's Cover Texts may be placed on covers that bracket the Document within the aggregate, or the electronic equivalent of covers if the Document is in electronic form. Otherwise they must appear on printed covers that bracket the whole aggregate.

#### 8. TRANSLATION

Translation is considered a kind of modification, so you may distribute translations of the Document under the terms of section 4. Replacing Invariant Sections with translations requires special permission from their copyright holders, but you may include translations of some or all Invariant Sections in addition to the original versions of these Invariant Sections. You may include a translation of this License, and all the license notices in the Document, and any Warranty Disclaimers, provided that you also include the original English version of this License and the original versions of those notices and disclaimers. In case of a disagreement between the translation and the original version of this License or a notice or disclaimer, the original version will prevail.

If a section in the Document is Entitled "Acknowledgements", "Dedications", or "History", the requirement (section 4) to Preserve its Title (section 1) will typically require changing the actual title.

#### 9. TERMINATION

You may not copy, modify, sublicense, or distribute the Document except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense, or distribute it is void, and will automatically terminate your rights under this License.

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, receipt of a copy of some or all of the same material does not give you any rights to use it.

#### 10. FUTURE REVISIONS OF THIS LICENSE

The Free Software Foundation may publish new, revised versions of the GNU Free Documentation License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns. See http://www.gnu.org/copyleft/.

Each version of the License is given a distinguishing version number. If the Document specifies that a particular numbered version of this License "or any later version" applies to it, you have the option of following the terms and conditions either of that specified version or of any later version that has been published (not as a draft) by the Free Software Foundation. If the Document does not specify a version number of this License, you may choose any version ever published (not as a draft) by the Free Software Foundation. If the Document specifies that a proxy can decide which future versions of this License can be used, that proxy's public statement of acceptance of a version permanently authorizes you to choose that version for the Document.

#### 11. RELICENSING

"Massive Multiauthor Collaboration Site" (or "MMC Site") means any World Wide Web server that publishes copyrightable works and also provides prominent facilities for anybody to edit those works. A public wiki that anybody can edit is an example of such a server. A "Massive Multiauthor Collaboration" (or "MMC") contained in the site means any set of copyrightable works thus published on the MMC site.

"CC-BY-SA" means the Creative Commons Attribution-Share Alike 3.0 license published by Creative Commons Corporation, a not-for-profit corporation with a principal place of business in San Francisco, California, as well as future copyleft versions of that license published by that same organization.

"Incorporate" means to publish or republish a Document, in whole or in part, as part of another Document.

An MMC is "eligible for relicensing" if it is licensed under this License, and if all works that were first published under this License somewhere other than this MMC, and subsequently incorporated in whole or in part into the MMC, (1) had no cover texts or invariant sections, and (2) were thus incorporated prior to November 1, 2008.

The operator of an MMC Site may republish an MMC contained in the site under CC-BY-SA on the same site at any time before August 1, 2009, provided the MMC is eligible for relicensing.

### ADDENDUM: How to use this License for your documents

To use this License in a document you have written, include a copy of the License in the document and put the following copyright and license notices just after the title page:

Copyright (C) year your name.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts. A copy of the license is included in the section entitled ''GNU Free Documentation License''.

If you have Invariant Sections, Front-Cover Texts and Back-Cover Texts, replace the "with...Texts." line with this:

with the Invariant Sections being list their titles, with the Front-Cover Texts being list, and with the Back-Cover Texts being list.

If you have Invariant Sections without Cover Texts, or some other combination of the three, merge those two alternatives to suit the situation.

If your document contains nontrivial examples of program code, we recommend releasing these examples in parallel under your choice of free software license, such as the GNU General Public License, to permit their use in free software.

## Contributors to GCC

The GCC project would like to thank its many contributors. Without them the project would not have been nearly as successful as it has been. Any omissions in this list are accidental. Feel free to contact law@redhat.com or gerald@pfeifer.com if you have been left out or some of your contributions are not listed. Please keep this list in alphabetical order.

- Analog Devices helped implement the support for complex data types and iterators.
- John David Anglin for threading-related fixes and improvements to libstdc++-v3, and the HP-UX port.
- James van Artsdalen wrote the code that makes efficient use of the Intel 80387 register stack.
- Abramo and Roberto Bagnara for the SysV68 Motorola 3300 Delta Series port.
- Alasdair Baird for various bug fixes.
- Giovanni Bajo for analyzing lots of complicated C++ problem reports.
- Peter Barada for his work to improve code generation for new ColdFire cores.
- Gerald Baumgartner added the signature extension to the C++ front end.
- Godmar Back for his Java improvements and encouragement.
- Scott Bambrough for help porting the Java compiler.
- Wolfgang Bangerth for processing tons of bug reports.
- Jon Beniston for his Microsoft Windows port of Java and port to Lattice Mico32.
- Daniel Berlin for better DWARF 2 support, faster/better optimizations, improved alias analysis, plus migrating GCC to Bugzilla.
- Geoff Berry for his Java object serialization work and various patches.
- David Binderman tests weekly snapshots of GCC trunk against Fedora Rawhide for several architectures.
- Laurynas Biveinis for memory management work and DJGPP port fixes.
- Uros Bizjak for the implementation of x87 math built-in functions and for various middle end and i386 back end improvements and bug fixes.
- Eric Blake for helping to make GCJ and libgcj conform to the specifications.
- Janne Blomqvist for contributions to GNU Fortran.
- Hans-J. Boehm for his garbage collector, IA-64 libffi port, and other Java work.
- Segher Boessenkool for helping maintain the PowerPC port and the instruction combiner plus various contributions to the middle end.
- Neil Booth for work on cpplib, lang hooks, debug hooks and other miscellaneous cleanups.
- Steven Bosscher for integrating the GNU Fortran front end into GCC and for contributing to the tree-ssa branch.
- Eric Botcazou for fixing middle- and backend bugs left and right.
- Per Bothner for his direction via the steering committee and various improvements to the infrastructure for supporting new languages. Chill front end implementation.

Initial implementations of cpplib, fix-header, config.guess, libio, and past C++ library (libg++) maintainer. Dreaming up, designing and implementing much of GCJ.

- Devon Bowen helped port GCC to the Tahoe.
- Don Bowman for mips-vxworks contributions.
- James Bowman for the FT32 port.
- Dave Brolley for work on cpplib and Chill.
- Paul Brook for work on the ARM architecture and maintaining GNU Fortran.
- Robert Brown implemented the support for Encore 32000 systems.
- Christian Bruel for improvements to local store elimination.
- Herman A.J. ten Brugge for various fixes.
- Joerg Brunsmann for Java compiler hacking and help with the GCJ FAQ.
- Joe Buck for his direction via the steering committee from its creation to 2013.
- Iain Buclaw for the D frontend.
- Craig Burley for leadership of the G77 Fortran effort.
- Tobias Burnus for contributions to GNU Fortran.
- Stephan Buys for contributing Doxygen notes for libstdc++.
- Paolo Carlini for libstdc++ work: lots of efficiency improvements to the C++ strings, streambufs and formatted I/O, hard detective work on the frustrating localization issues, and keeping up with the problem reports.
- John Carr for his alias work, SPARC hacking, infrastructure improvements, previous contributions to the steering committee, loop optimizations, etc.
- Stephane Carrez for 68HC11 and 68HC12 ports.
- Steve Chamberlain for support for the Renesas SH and H8 processors and the PicoJava processor, and for GCJ config fixes.
- Glenn Chambers for help with the GCJ FAQ.
- John-Marc Chandonia for various libgcj patches.
- Denis Chertykov for contributing and maintaining the AVR port, the first GCC port for an 8-bit architecture.
- Kito Cheng for his work on the RISC-V port, including bringing up the test suite and maintenance.
- Scott Christley for his Objective-C contributions.
- Eric Christopher for his Java porting help and clean-ups.
- Branko Cibej for more warning contributions.
- The GNU Classpath project for all of their merged runtime code.
- Nick Clifton for arm, mcore, fr30, v850, m32r, msp430 rx work, '--help', and other random hacking.
- Michael Cook for libstdc++ cleanup patches to reduce warnings.
- R. Kelley Cook for making GCC buildable from a read-only directory as well as other miscellaneous build process and documentation clean-ups.
- Ralf Corsepius for SH testing and minor bug fixing.

- François-Xavier Coudert for contributions to GNU Fortran.
- Stan Cox for care and feeding of the x86 port and lots of behind the scenes hacking.
- Alex Crain provided changes for the 3b1.
- Ian Dall for major improvements to the NS32k port.
- Paul Dale for his work to add uClinux platform support to the m68k backend.
- Palmer Dabbelt for his work maintaining the RISC-V port.
- Dario Dariol contributed the four varieties of sample programs that print a copy of their source.
- Russell Davidson for fstream and stringstream fixes in libstdc++.
- Bud Davis for work on the G77 and GNU Fortran compilers.
- Mo DeJong for GCJ and libgcj bug fixes.
- Jerry DeLisle for contributions to GNU Fortran.
- DJ Delorie for the DJGPP port, build and libiberty maintenance, various bug fixes, and the M32C, MeP, MSP430, and RL78 ports.
- Arnaud Desitter for helping to debug GNU Fortran.
- Gabriel Dos Reis for contributions to G++, contributions and maintenance of GCC diagnostics infrastructure, libstdc++-v3, including valarray<>, complex<>, maintaining the numerics library (including that pesky imits> :-) and keeping up-to-date anything to do with numbers.
- Ulrich Drepper for his work on glibc, testing of GCC using glibc, ISO C99 support, CFG dumping support, etc., plus support of the C++ runtime libraries including for all kinds of C interface issues, contributing and maintaining complex<>, sanity checking and disbursement, configuration architecture, libio maintenance, and early math work.
- François Dumont for his work on libstdc++-v3, especially maintaining and improving debug-mode and associative and unordered containers.
- Zdenek Dvorak for a new loop unroller and various fixes.
- Michael Eager for his work on the Xilinx MicroBlaze port.
- Richard Earnshaw for his ongoing work with the ARM.
- David Edelsohn for his direction via the steering committee, ongoing work with the RS6000/PowerPC port, help cleaning up Haifa loop changes, doing the entire AIX port of libstdc++ with his bare hands, and for ensuring GCC properly keeps working on AIX.
- Kevin Ediger for the floating point formatting of num\_put::do\_put in libstdc++.
- Phil Edwards for libstdc++ work including configuration hackery, documentation maintainer, chief breaker of the web pages, the occasional iostream bug fix, and work on shared library symbol versioning.
- Paul Eggert for random hacking all over GCC.
- Mark Elbrecht for various DJGPP improvements, and for libstdc++ configuration support for locales and fstream-related fixes.
- Vadim Egorov for libstdc++ fixes in strings, streambufs, and iostreams.
- Christian Ehrhardt for dealing with bug reports.

- Ben Elliston for his work to move the Objective-C runtime into its own subdirectory and for his work on autoconf.
- Revital Eres for work on the PowerPC 750CL port.
- Marc Espie for OpenBSD support.
- Doug Evans for much of the global optimization framework, arc, m32r, and SPARC work.
- Christopher Faylor for his work on the Cygwin port and for caring and feeding the gcc.gnu.org box and saving its users tons of spam.
- Fred Fish for BeOS support and Ada fixes.
- Ivan Fontes Garcia for the Portuguese translation of the GCJ FAQ.
- Peter Gerwinski for various bug fixes and the Pascal front end.
- Kaveh R. Ghazi for his direction via the steering committee, amazing work to make '-W-Wall-W\*-Werror' useful, and testing GCC on a plethora of platforms. Kaveh extends his gratitude to the CAIP Center at Rutgers University for providing him with computing resources to work on Free Software from the late 1980s to 2010.
- John Gilmore for a donation to the FSF earmarked improving GNU Java.
- Judy Goldberg for c++ contributions.
- Torbjorn Granlund for various fixes and the c-torture testsuite, multiply- and divideby-constant optimization, improved long long support, improved leaf function register allocation, and his direction via the steering committee.
- Jonny Grant for improvements to collect2's '--help' documentation.
- Anthony Green for his '-Os' contributions, the moxie port, and Java front end work.
- Stu Grossman for gdb hacking, allowing GCJ developers to debug Java code.
- Michael K. Gschwind contributed the port to the PDP-11.
- Richard Biener for his ongoing middle-end contributions and bug fixes and for release management.
- Ron Guilmette implemented the protoize and unprotoize tools, the support for DWARF 1 symbolic debugging information, and much of the support for System V Release 4. He has also worked heavily on the Intel 386 and 860 support.
- Sumanth Gundapaneni for contributing the CR16 port.
- Mostafa Hagog for Swing Modulo Scheduling (SMS) and post reload GCSE.
- Bruno Haible for improvements in the runtime overhead for EH, new warnings and assorted bug fixes.
- Andrew Haley for his amazing Java compiler and library efforts.
- Chris Hanson assisted in making GCC work on HP-UX for the 9000 series 300.
- Michael Hayes for various thankless work he's done trying to get the c30/c40 ports functional. Lots of loop and unroll improvements and fixes.
- Dara Hazeghi for wading through myriads of target-specific bug reports.
- Kate Hedstrom for staking the G77 folks with an initial testsuite.
- Richard Henderson for his ongoing SPARC, alpha, ia32, and ia64 work, loop opts, and generally fixing lots of old problems we've ignored for years, flow rewrite and lots of further stuff, including reviewing tons of patches.

- Aldy Hernandez for working on the PowerPC port, SIMD support, and various fixes.
- Nobuyuki Hikichi of Software Research Associates, Tokyo, contributed the support for the Sony NEWS machine.
- Kazu Hirata for caring and feeding the Renesas H8/300 port and various fixes.
- Katherine Holcomb for work on GNU Fortran.
- Manfred Hollstein for his ongoing work to keep the m88k alive, lots of testing and bug fixing, particularly of GCC configury code.
- Steve Holmgren for MachTen patches.
- Mat Hostetter for work on the TILE-Gx and TILEPro ports.
- Jan Hubicka for his x86 port improvements.
- Falk Hueffner for working on C and optimization bug reports.
- Bernardo Innocenti for his m68k work, including merging of ColdFire improvements and uClinux support.
- Christian Iseli for various bug fixes.
- Kamil Iskra for general m68k hacking.
- Lee Iverson for random fixes and MIPS testing.
- Balaji V. Iyer for Cilk+ development and merging.
- Andreas Jaeger for testing and benchmarking of GCC and various bug fixes.
- Martin Jambor for his work on inter-procedural optimizations, the switch conversion pass, and scalar replacement of aggregates.
- Jakub Jelinek for his SPARC work and sibling call optimizations as well as lots of bug fixes and test cases, and for improving the Java build system.
- Janis Johnson for ia64 testing and fixes, her quality improvement sidetracks, and web page maintenance.
- Kean Johnston for SCO OpenServer support and various fixes.
- Tim Josling for the sample language treelang based originally on Richard Kenner's "toy" language.
- Nicolai Josuttis for additional libstdc++ documentation.
- Klaus Kaempf for his ongoing work to make alpha-vms a viable target.
- Steven G. Kargl for work on GNU Fortran.
- David Kashtan of SRI adapted GCC to VMS.
- Ryszard Kabatek for many, many libstdc++ bug fixes and optimizations of strings, especially member functions, and for auto\_ptr fixes.
- Geoffrey Keating for his ongoing work to make the PPC work for GNU/Linux and his automatic regression tester.
- Brendan Kehoe for his ongoing work with G++ and for a lot of early work in just about every part of libstdc++.
- Oliver M. Kellogg of Deutsche Aerospace contributed the port to the MIL-STD-1750A.
- Richard Kenner of the New York University Ultracomputer Research Laboratory wrote the machine descriptions for the AMD 29000, the DEC Alpha, the IBM RT PC, and the IBM RS/6000 as well as the support for instruction attributes. He also made

changes to better support RISC processors including changes to common subexpression elimination, strength reduction, function calling sequence handling, and condition code support, in addition to generalizing the code for frame pointer elimination and delay slot scheduling. Richard Kenner was also the head maintainer of GCC for several years.

- Mumit Khan for various contributions to the Cygwin and Mingw32 ports and maintaining binary releases for Microsoft Windows hosts, and for massive libstdc++ porting work to Cygwin/Mingw32.
- Robin Kirkham for cpu32 support.
- Mark Klein for PA improvements.
- Thomas Koenig for various bug fixes.
- Bruce Korb for the new and improved fixincludes code.
- Benjamin Kosnik for his G++ work and for leading the libstdc++-v3 effort.
- Maxim Kuvyrkov for contributions to the instruction scheduler, the Android and m68k/Coldfire ports, and optimizations.
- Charles LaBrec contributed the support for the Integrated Solutions 68020 system.
- Asher Langton and Mike Kumbera for contributing Cray pointer support to GNU Fortran, and for other GNU Fortran improvements.
- Jeff Law for his direction via the steering committee, coordinating the entire egcs project and GCC 2.95, rolling out snapshots and releases, handling merges from GCC2, reviewing tons of patches that might have fallen through the cracks else, and random but extensive hacking.
- Walter Lee for work on the TILE-Gx and TILEPro ports.
- Marc Lehmann for his direction via the steering committee and helping with analysis and improvements of x86 performance.
- Victor Leikehman for work on GNU Fortran.
- Ted Lemon wrote parts of the RTL reader and printer.
- Kriang Lerdsuwanakij for C++ improvements including template as template parameter support, and many C++ fixes.
- Warren Levy for tremendous work on libgcj (Java Runtime Library) and random work on the Java front end.
- Alain Lichnewsky ported GCC to the MIPS CPU.
- Oskar Liljeblad for hacking on AWT and his many Java bug reports and patches.
- Robert Lipe for OpenServer support, new testsuites, testing, etc.
- Chen Liqin for various S+core related fixes/improvement, and for maintaining the S+core port.
- Martin Liska for his work on identical code folding, the sanitizers, HSA, general bug fixing and for running automated regression testing of GCC and reporting numerous bugs.
- Weiwen Liu for testing and various bug fixes.
- Manuel López-Ibáñez for improving '-Wconversion' and many other diagnostics fixes and improvements.
- Dave Love for his ongoing work with the Fortran front end and runtime libraries.

Martin von Löwis for internal consistency checking infrastructure, various C++ improvements including namespace support, and tons of assistance with libstdc++/compiler merges.

- H.J. Lu for his previous contributions to the steering committee, many x86 bug reports, prototype patches, and keeping the GNU/Linux ports working.
- Greg McGary for random fixes and (someday) bounded pointers.
- Andrew MacLeod for his ongoing work in building a real EH system, various code generation improvements, work on the global optimizer, etc.
- Vladimir Makarov for hacking some ugly i960 problems, PowerPC hacking improvements to compile-time performance, overall knowledge and direction in the area of instruction scheduling, design and implementation of the automaton based instruction scheduler and design and implementation of the integrated and local register allocators.
- David Malcolm for his work on improving GCC diagnostics, JIT, self-tests and unit testing.
- Bob Manson for his behind the scenes work on dejagnu.
- John Marino for contributing the DragonFly BSD port.
- Philip Martin for lots of libstdc++ string and vector iterator fixes and improvements, and string clean up and testsuites.
- Michael Matz for his work on dominance tree discovery, the x86-64 port, link-time optimization framework and general optimization improvements.
- All of the Mauve project contributors for Java test code.
- Bryce McKinlay for numerous GCJ and libgcj fixes and improvements.
- Adam Megacz for his work on the Microsoft Windows port of GCJ.
- Michael Meissner for LRS framework, ia32, m32r, v850, m88k, MIPS, powerpc, haifa, ECOFF debug support, and other assorted hacking.
- Jason Merrill for his direction via the steering committee and leading the G++ effort.
- Martin Michlmayr for testing GCC on several architectures using the entire Debian archive.
- David Miller for his direction via the steering committee, lots of SPARC work, improvements in jump.c and interfacing with the Linux kernel developers.
- Gary Miller ported GCC to Charles River Data Systems machines.
- Alfred Minarik for libstdc++ string and ios bug fixes, and turning the entire libstdc++ testsuite namespace-compatible.
- Mark Mitchell for his direction via the steering committee, mountains of C++ work, load/store hoisting out of loops, alias analysis improvements, ISO C restrict support, and serving as release manager from 2000 to 2011.
- Alan Modra for various GNU/Linux bits and testing.
- Toon Moene for his direction via the steering committee, Fortran maintenance, and his ongoing work to make us make Fortran run fast.
- Jason Molenda for major help in the care and feeding of all the services on the gcc.gnu.org (formerly egcs.cygnus.com) machine—mail, web services, ftp services, etc etc. Doing all this work on scrap paper and the backs of envelopes would have been... difficult.

- Catherine Moore for fixing various ugly problems we have sent her way, including the haifa bug which was killing the Alpha & PowerPC Linux kernels.
- Mike Moreton for his various Java patches.
- David Mosberger-Tang for various Alpha improvements, and for the initial IA-64 port.
- Stephen Moshier contributed the floating point emulator that assists in cross-compilation and permits support for floating point numbers wider than 64 bits and for ISO C99 support.
- Bill Moyer for his behind the scenes work on various issues.
- Philippe De Muyter for his work on the m68k port.
- Joseph S. Myers for his work on the PDP-11 port, format checking and ISO C99 support, and continuous emphasis on (and contributions to) documentation.
- Nathan Myers for his work on libstdc++-v3: architecture and authorship through the first three snapshots, including implementation of locale infrastructure, string, shadow C headers, and the initial project documentation (DESIGN, CHECKLIST, and so forth). Later, more work on MT-safe string and shadow headers.
- Felix Natter for documentation on porting libstdc++.
- Nathanael Nerode for cleaning up the configuration/build process.
- NeXT, Inc. donated the front end that supports the Objective-C language.
- Hans-Peter Nilsson for the CRIS and MMIX ports, improvements to the search engine setup, various documentation fixes and other small fixes.
- Geoff Noer for his work on getting cygwin native builds working.
- Vegard Nossum for running automated regression testing of GCC and reporting numerous bugs.
- Diego Novillo for his work on Tree SSA, OpenMP, SPEC performance tracking web pages, GIMPLE tuples, and assorted fixes.
- David O'Brien for the FreeBSD/alpha, FreeBSD/AMD x86-64, FreeBSD/ARM, FreeBSD/PowerPC, and FreeBSD/SPARC64 ports and related infrastructure improvements.
- Alexandre Oliva for various build infrastructure improvements, scripts and amazing testing work, including keeping libtool issues sane and happy.
- Stefan Olsson for work on mt\_alloc.
- Melissa O'Neill for various NeXT fixes.
- Rainer Orth for random MIPS work, including improvements to GCC's o32 ABI support, improvements to dejagnu's MIPS support, Java configuration clean-ups and porting work, and maintaining the IRIX, Solaris 2, and Tru64 UNIX ports.
- Steven Pemberton for his contribution of 'enquire' which allowed GCC to determine various properties of the floating point unit and generate 'float.h' in older versions of GCC.
- Hartmut Penner for work on the s390 port.
- Paul Petersen wrote the machine description for the Alliant FX/8.
- Alexandre Petit-Bianco for implementing much of the Java compiler and continued Java maintainership.

- Matthias Pfaller for major improvements to the NS32k port.
- Gerald Pfeifer for his direction via the steering committee, pointing out lots of problems we need to solve, maintenance of the web pages, and taking care of documentation maintenance in general.
- Marek Polacek for his work on the C front end, the sanitizers and general bug fixing.
- Andrew Pinski for processing bug reports by the dozen.
- Ovidiu Predescu for his work on the Objective-C front end and runtime libraries.
- Jerry Quinn for major performance improvements in C++ formatted I/O.
- Ken Raeburn for various improvements to checker, MIPS ports and various cleanups in the compiler.
- Rolf W. Rasmussen for hacking on AWT.
- David Reese of Sun Microsystems contributed to the Solaris on PowerPC port.
- John Regehr for running automated regression testing of GCC and reporting numerous bugs.
- Volker Reichelt for running automated regression testing of GCC and reporting numerous bugs and for keeping up with the problem reports.
- Joern Rennecke for maintaining the sh port, loop, regmove & reload hacking and developing and maintaining the Epiphany port.
- Loren J. Rittle for improvements to libstdc++-v3 including the FreeBSD port, threading fixes, thread-related configury changes, critical threading documentation, and solutions to really tricky I/O problems, as well as keeping GCC properly working on FreeBSD and continuous testing.
- Craig Rodrigues for processing tons of bug reports.
- Ola Rönnerup for work on mt\_alloc.
- Gavin Romig-Koch for lots of behind the scenes MIPS work.
- David Ronis inspired and encouraged Craig to rewrite the G77 documentation in texinfo format by contributing a first pass at a translation of the old 'g77-0.5.16/f/D0C' file.
- Ken Rose for fixes to GCC's delay slot filling code.
- Ira Rosen for her contributions to the auto-vectorizer.
- Paul Rubin wrote most of the preprocessor.
- Pétur Runólfsson for major performance improvements in C++ formatted I/O and large file support in C++ filebuf.
- Chip Salzenberg for libstdc++ patches and improvements to locales, traits, Makefiles, libio, libtool hackery, and "long long" support.
- Juha Sarlin for improvements to the H8 code generator.
- Greg Satz assisted in making GCC work on HP-UX for the 9000 series 300.
- Roger Sayle for improvements to constant folding and GCC's RTL optimizers as well as for fixing numerous bugs.
- Bradley Schatz for his work on the GCJ FAQ.
- Peter Schauer wrote the code to allow debugging to work on the Alpha.
- William Schelter did most of the work on the Intel 80386 support.

- Tobias Schlüter for work on GNU Fortran.
- Bernd Schmidt for various code generation improvements and major work in the reload pass, serving as release manager for GCC 2.95.3, and work on the Blackfin and C6X ports.
- Peter Schmid for constant testing of libstdc++—especially application testing, going above and beyond what was requested for the release criteria—and libstdc++ header file tweaks.
- Jason Schroeder for jcf-dump patches.
- Andreas Schwab for his work on the m68k port.
- Lars Segerlund for work on GNU Fortran.
- Dodji Seketeli for numerous C++ bug fixes and debug info improvements.
- Tim Shen for major work on <regex>.
- Joel Sherrill for his direction via the steering committee, RTEMS contributions and RTEMS testing.
- Nathan Sidwell for many C++ fixes/improvements.
- Jeffrey Siegal for helping RMS with the original design of GCC, some code which handles the parse tree and RTL data structures, constant folding and help with the original VAX & m68k ports.
- Kenny Simpson for prompting libstdc++ fixes due to defect reports from the LWG (thereby keeping GCC in line with updates from the ISO).
- Franz Sirl for his ongoing work with making the PPC port stable for GNU/Linux.
- Andrey Slepuhin for assorted AIX hacking.
- Trevor Smigiel for contributing the SPU port.
- Christopher Smith did the port for Convex machines.
- Danny Smith for his major efforts on the Mingw (and Cygwin) ports. Retired from GCC maintainership August 2010, having mentored two new maintainers into the role.
- Randy Smith finished the Sun FPA support.
- Ed Smith-Rowland for his continuous work on libstdc++-v3, special functions, <random>, and various improvements to C++11 features.
- Scott Snyder for queue, iterator, istream, and string fixes and libstdc++ testsuite entries. Also for providing the patch to G77 to add rudimentary support for INTEGER\*1, INTEGER\*2, and LOGICAL\*1.
- Zdenek Sojka for running automated regression testing of GCC and reporting numerous bugs.
- Arseny Solokha for running automated regression testing of GCC and reporting numerous bugs.
- Jayant Sonar for contributing the CR16 port.
- Brad Spencer for contributions to the GLIBCPP\_FORCE\_NEW technique.
- Richard Stallman, for writing the original GCC and launching the GNU project.
- Jan Stein of the Chalmers Computer Society provided support for Genix, as well as part of the 32000 machine description.

• Gerhard Steinmetz for running automated regression testing of GCC and reporting numerous bugs.

- Nigel Stephens for various mips16 related fixes/improvements.
- Jonathan Stone wrote the machine description for the Pyramid computer.
- Graham Stott for various infrastructure improvements.
- John Stracke for his Java HTTP protocol fixes.
- Mike Stump for his Elxsi port, G++ contributions over the years and more recently his vxworks contributions
- Jeff Sturm for Java porting help, bug fixes, and encouragement.
- Zhendong Su for running automated regression testing of GCC and reporting numerous bugs.
- Chengnian Sun for running automated regression testing of GCC and reporting numerous bugs.
- Shigeya Suzuki for this fixes for the bsdi platforms.
- Ian Lance Taylor for the Go frontend, the initial mips16 and mips64 support, general configury hacking, fixincludes, etc.
- Holger Teutsch provided the support for the Clipper CPU.
- Gary Thomas for his ongoing work to make the PPC work for GNU/Linux.
- Paul Thomas for contributions to GNU Fortran.
- Philipp Thomas for random bug fixes throughout the compiler
- Jason Thorpe for thread support in libstdc++ on NetBSD.
- Kresten Krab Thorup wrote the run time support for the Objective-C language and the fantastic Java bytecode interpreter.
- Michael Tiemann for random bug fixes, the first instruction scheduler, initial C++ support, function integration, NS32k, SPARC and M88k machine description work, delay slot scheduling.
- Andreas Tobler for his work porting libgcj to Darwin.
- Teemu Torma for thread safe exception handling support.
- Leonard Tower wrote parts of the parser, RTL generator, and RTL definitions, and of the VAX machine description.
- Daniel Towner and Hariharan Sandanagobalane contributed and maintain the picoChip port.
- Tom Tromey for internationalization support and for his many Java contributions and libgcj maintainership.
- Lassi Tuura for improvements to config.guess to determine HP processor types.
- Petter Urkedal for libstdc++ CXXFLAGS, math, and algorithms fixes.
- Andy Vaught for the design and initial implementation of the GNU Fortran front end.
- Brent Verner for work with the libstdc++ cshadow files and their associated configure steps.
- Todd Vierling for contributions for NetBSD ports.
- Andrew Waterman for contributing the RISC-V port, as well as maintaining it.

- Jonathan Wakely for contributing libstdc++ Doxygen notes and XHTML guidance and maintaining libstdc++.
- Dean Wakerley for converting the install documentation from HTML to texinfo in time for GCC 3.0.
- Krister Walfridsson for random bug fixes.
- Feng Wang for contributions to GNU Fortran.
- Stephen M. Webb for time and effort on making libstdc++ shadow files work with the tricky Solaris 8+ headers, and for pushing the build-time header tree. Also, for starting and driving the <regex> effort.
- John Wehle for various improvements for the x86 code generator, related infrastructure improvements to help x86 code generation, value range propagation and other work, WE32k port.
- Ulrich Weigand for work on the s390 port.
- Janus Weil for contributions to GNU Fortran.
- Zack Weinberg for major work on cpplib and various other bug fixes.
- Matt Welsh for help with Linux Threads support in GCJ.
- Urban Widmark for help fixing java.io.
- Mark Wielaard for new Java library code and his work integrating with Classpath.
- Dale Wiles helped port GCC to the Tahoe.
- Bob Wilson from Tensilica, Inc. for the Xtensa port.
- Jim Wilson for his direction via the steering committee, tackling hard problems in various places that nobody else wanted to work on, strength reduction and other loop optimizations.
- Paul Woegerer and Tal Agmon for the CRX port.
- Carlo Wood for various fixes.
- Tom Wood for work on the m88k port.
- Chung-Ju Wu for his work on the Andes NDS32 port.
- Cangun Yang for work on GNU Fortran.
- Masanobu Yuhara of Fujitsu Laboratories implemented the machine description for the Tron architecture (specifically, the Gmicro).
- Kevin Zachmann helped port GCC to the Tahoe.
- Ayal Zaks for Swing Modulo Scheduling (SMS).
- Qirun Zhang for running automated regression testing of GCC and reporting numerous bugs.
- Xiaoqiang Zhang for work on GNU Fortran.
- Gilles Zunino for help porting Java to Irix.

The following people are recognized for their contributions to GNAT, the Ada front end of GCC:

- Bernard Banner
- Romain Berrendonner

- Geert Bosch
- Emmanuel Briot
- Joel Brobecker
- Ben Brosgol
- Vincent Celier
- Arnaud Charlet
- Chien Chieng
- Cyrille Comar
- Cyrille Crozes
- Robert Dewar
- Gary Dismukes
- Robert Duff
- Ed Falis
- Ramon Fernandez
- Sam Figueroa
- Vasiliy Fofanov
- Michael Friess
- Franco Gasperoni
- Ted Giering
- Matthew Gingell
- Laurent Guerby
- Jerome Guitton
- Olivier Hainque
- Jerome Hugues
- Hristian Kirtchev
- Jerome Lambourg
- Bruno Leclerc
- Albert Lee
- Sean McNeil
- Javier Miranda
- Laurent Nana
- Pascal Obry
- Dong-Ik Oh
- Laurent Pautet
- Brett Porter
- Thomas Quinot
- Nicolas Roche
- Pat Rogers
- Jose Ruiz

- Douglas Rupp
- Sergey Rybin
- Gail Schenker
- Ed Schonberg
- Nicolas Setton
- Samuel Tardieu

The following people are recognized for their contributions of new features, bug reports, testing and integration of classpath/libgcj for GCC version 4.1:

- Lillian Angel for JTree implementation and lots Free Swing additions and bug fixes.
- Wolfgang Baer for GapContent bug fixes.
- Anthony Balkissoon for JList, Free Swing 1.5 updates and mouse event fixes, lots of Free Swing work including JTable editing.
- Stuart Ballard for RMI constant fixes.
- Goffredo Baroncelli for HTTPURLConnection fixes.
- Gary Benson for MessageFormat fixes.
- Daniel Bonniot for Serialization fixes.
- Chris Burdess for lots of gnu.xml and http protocol fixes, StAX and DOM xml:id support.
- Ka-Hing Cheung for TreePath and TreeSelection fixes.
- Archie Cobbs for build fixes, VM interface updates, URLClassLoader updates.
- Kelley Cook for build fixes.
- Martin Cordova for Suggestions for better SocketTimeoutException.
- David Daney for BitSet bug fixes, HttpURLConnection rewrite and improvements.
- Thomas Fitzsimmons for lots of upgrades to the gtk+ AWT and Cairo 2D support. Lots of imageio framework additions, lots of AWT and Free Swing bug fixes.
- Jeroen Frijters for ClassLoader and nio cleanups, serialization fixes, better Proxy support, bug fixes and IKVM integration.
- Santiago Gala for AccessControlContext fixes.
- Nicolas Geoffray for VMClassLoader and AccessController improvements.
- David Gilbert for basic and metal icon and plaf support and lots of documenting, Lots of Free Swing and metal theme additions. MetallconFactory implementation.
- Anthony Green for MIDI framework, ALSA and DSSI providers.
- Andrew Haley for Serialization and URLClassLoader fixes, gcj build speedups.
- Kim Ho for JFileChooser implementation.
- Andrew John Hughes for Locale and net fixes, URI RFC2986 updates, Serialization fixes, Properties XML support and generic branch work, VMIntegration guide update.
- Bastiaan Huisman for TimeZone bug fixing.
- Andreas Jaeger for mprec updates.
- Paul Jenner for better '-Werror' support.
- Ito Kazumitsu for NetworkInterface implementation and updates.

• Roman Kennke for BoxLayout, GrayFilter and SplitPane, plus bug fixes all over. Lots of Free Swing work including styled text.

- Simon Kitching for String cleanups and optimization suggestions.
- Michael Koch for configuration fixes, Locale updates, bug and build fixes.
- Guilhem Lavaux for configuration, thread and channel fixes and Kaffe integration. JCL native Pointer updates. Logger bug fixes.
- David Lichteblau for JCL support library global/local reference cleanups.
- Aaron Luchko for JDWP updates and documentation fixes.
- Ziga Mahkovec for Graphics2D upgraded to Cairo 0.5 and new regex features.
- Sven de Marothy for BMP imageio support, CSS and TextLayout fixes. GtkImage rewrite, 2D, awt, free swing and date/time fixes and implementing the Qt4 peers.
- Casey Marshall for crypto algorithm fixes, FileChannel lock, SystemLogger and FileHandler rotate implementations, NIO FileChannel.map support, security and policy updates.
- Bryce McKinlay for RMI work.
- Audrius Meskauskas for lots of Free Corba, RMI and HTML work plus testing and documenting.
- Kalle Olavi Niemitalo for build fixes.
- Rainer Orth for build fixes.
- Andrew Overholt for File locking fixes.
- Ingo Proetel for Image, Logger and URLClassLoader updates.
- Olga Rodimina for MenuSelectionManager implementation.
- Jan Roehrich for BasicTreeUI and JTree fixes.
- Julian Scheid for documentation updates and gidoc support.
- Christian Schlichtherle for zip fixes and cleanups.
- Robert Schuster for documentation updates and beans fixes, TreeNode enumerations and ActionCommand and various fixes, XML and URL, AWT and Free Swing bug fixes.
- Keith Seitz for lots of JDWP work.
- Christian Thalinger for 64-bit cleanups, Configuration and VM interface fixes and CACAO integration, fdlibm updates.
- Gael Thomas for VMClassLoader boot packages support suggestions.
- Andreas Tobler for Darwin and Solaris testing and fixing, Qt4 support for Darwin/OS X, Graphics2D support, gtk+ updates.
- Dalibor Topic for better DEBUG support, build cleanups and Kaffe integration. Qt4 build infrastructure, SHA1PRNG and GdkPixbugDecoder updates.
- Tom Tromey for Eclipse integration, generics work, lots of bug fixes and gcj integration including coordinating The Big Merge.
- Mark Wielaard for bug fixes, packaging and release management, Clipboard implementation, system call interrupts and network timeouts and GdkPixpufDecoder fixes.

In addition to the above, all of which also contributed time and energy in testing GCC, we would like to thank the following for their contributions to testing:

- Michael Abd-El-Malek
- Thomas Arend
- Bonzo Armstrong
- Steven Ashe
- Chris Baldwin
- David Billinghurst
- Jim Blandy
- Stephane Bortzmeyer
- Horst von Brand
- Frank Braun
- Rodney Brown
- Sidney Cadot
- Bradford Castalia
- Robert Clark
- Jonathan Corbet
- Ralph Doncaster
- Richard Emberson
- Levente Farkas
- Graham Fawcett
- Mark Fernyhough
- Robert A. French
- Jörgen Freyh
- Mark K. Gardner
- Charles-Antoine Gauthier
- Yung Shing Gene
- David Gilbert
- Simon Gornall
- Fred Gray
- John Griffin
- Patrik Hagglund
- Phil Hargett
- Amancio Hasty
- Takafumi Hayashi
- Bryan W. Headley
- Kevin B. Hendricks
- Joep Jansen
- Christian Joensson
- Michel Kern
- David Kidd

- Tobias Kuipers
- Anand Krishnaswamy
- A. O. V. Le Blanc
- llewelly
- Damon Love
- Brad Lucier
- Matthias Klose
- Martin Knoblauch
- Rick Lutowski
- Jesse Macnish
- Stefan Morrell
- Anon A. Mous
- Matthias Mueller
- Pekka Nikander
- Rick Niles
- Jon Olson
- Magnus Persson
- Chris Pollard
- Richard Polton
- Derk Reefman
- David Rees
- Paul Reilly
- Tom Reilly
- Torsten Rueger
- Danny Sadinoff
- Marc Schifer
- Erik Schnetter
- Wayne K. Schroll
- David Schuler
- Vin Shelton
- Tim Souder
- Adam Sulmicki
- Bill Thorson
- George Talbot
- Pedro A. M. Vazquez
- Gregory Warnes
- Ian Watson
- David E. Young
- And many others

And finally we'd like to thank everyone who uses the compiler, provides feedback and generally reminds us why we're doing this work in the first place.

Option Index 765

## **Option Index**

GCC's command line options are indexed here without any initial '-' or '--'. Where an option has both positive and negative forms (such as '-foption' and '-fno-option'), relevant entries in the manual are indexed under the most appropriate form; it may sometimes be useful to look up both forms.

${f F}$		fwpa	700
fltrans		3.6	
fltrans-output-list70	00	$\mathbf{M}$	
fresolution	00	msoft-float	12

!	*gimple_build_omp_single 237
'!' in constraint	*gimple_build_resx
	*gimple_build_return
11	*gimple_build_switch
#	*gimple_build_try
'#' in constraint	
# in template	+
#pragma	
	'+' in constraint
\$	
Φ	-
'\$' in constraint	'-fsection-anchors'
%	,
'%' in constraint	'/c' in RTL dump
% in GTY option	'/f' in RTL dump
'%' in template 343	'/i' in RTL dump
	'/j' in RTL dump
&	'/s' in RTL dump
<b>&amp;</b>	'/u' in RTL dump
'&' in constraint	'/v' in RTL dump
(	<
(	
(gimple	'<' in constraint
(gimple_stmt_iterator	
(nil)	=
	'=' in constraint
*	
'*' in constraint	>
* in template	·
*gimple_build_asm_vec	'>' in constraint
*gimple_build_assign	
*gimple_build_bind	?
*gimple_build_call 226	•
*gimple_build_call_from_tree 226	'?' in constraint
*gimple_build_call_vec	
*gimple_build_catch	@
*gimple_build_cond	
*gimple_build_cond_from_tree	'@' in instruction pattern names 477
*gimple_build_debug_bind	
*gimple_build_eh_filter 230	^
*gimple_build_goto	
*gimple_build_label	'^' in constraint
*gimple_build_omp_atomic_load	
*gimple_build_omp_atomic_store	
*gimple_build_omp_continue	-
*gimple_build_omp_critical	absvdi211
*gimple_build_omp_for	absvsi211
*gimple_build_omp_parallel	addda322
*gimple_build_omp_sections	adddf3 12

04440	22	bid_eqtd2	21
adddq3		bid_extendddtd2	
addha3		bid_extendddtf	
		bid_extendddxf	
addqq3		bid_extenddfdd	
addsa3			
addsf3		bid_extenddftd	
addsq3		bid_extendsddd2	
addta3		bid_extendsddf	
addtf3		bid_extendsdtd2	
adduda3		bid_extendsdtf	
addudq3		bid_extendsdxf	
adduha3		bid_extendsfdd	
adduhq3		bid_extendsfsd	
adduqq3		bid_extendsftd	
addusa3		bid_extendtftd	
addusq3		bid_extendxftd	
adduta3	23	bid_fixdddi	
addvdi3		bid_fixddsi	
addvsi3	11	bid_fixsddi	19
addxf3	12	bid_fixsdsi	19
ashlda3	28	bid_fixtddi	19
ashldi3	. 9	bid_fixtdsi	19
ashldq3	28	bid_fixunsdddi	19
ashlha3	28	bid_fixunsddsi	19
ashlhq3	28	bid_fixunssddi	19
ashlqq3		bid_fixunssdsi	19
ashlsa3		bid_fixunstddi	
ashlsi3		bid_fixunstdsi	
ashlsq3		bid_floatdidd	
ashlta3		bid_floatdisd	
ashlti3		bid_floatditd	
ashluda3		bid_floatsidd	
ashludq3		bid_floatsisd	
ashluha3		bid_floatsitd	
ashluhq3		bid_floatunsdidd	
ashluqq3		bid_floatunsdisd	
ashlusa3		bid_floatunsditd	
ashlusq3		bid_floatunssidd	
ashluta3		bid_floatunssisd	
ashrda3		bid_floatunssitd	
ashrdi3		bid_gedd2	
ashrdq3		bid_gesd2	
ashrha3		bid_getd2	
ashrhq3		bid_gtdd2	
ashrqq3		bid_gtsd2	
ashrsa3		bid_gttd2	
ashrsi3		bid_ledd2	
ashrsq3		bid_lesd2	
ashrta3		bid_letd2	
ashrti3		bid_ltdd2	
bid_adddd3		bid_1tsd2	
bid_addsd3		bid_lttd2	
bid_addtd3		bid_muldd3	
bid_divdd3		bid_mulsd3	
bid_divsd3		bid_multd3	
bid_divtd3	11	bid_nedd2	
bid_eqdd2		bid_negdd2	
hid easd?	21	hid negsd2	17

		_
bid_negtd2	divda3	
bid_nesd2 21	divdc3	
bid_netd2 21	divdf3	
bid_subdd3	divdi3	
bid_subsd3	divdq3	
bid_subtd3	divha3	
bid_truncdddf	divhq3	
bid_truncddsd2	divqq3	
bid_truncddsf	divsa3	
bid_truncdfsd	divsc3	
bid_truncsdsf	divsf3	
bid_trunctddd2	divsi3	
bid_trunctddf	divsq3	
bid_trunctdsd2	divta3	
bid_trunctdsf	divtc3	
bid_trunctdtf	divtf3	
bid_trunctdxf	divti3	
bid_trunctfdd	divxc3	
bid_trunctfsd	divxf3	
bid_truncxfdd	dpd_adddd3	
bid_truncxfsd		
bid_unorddd2	dpd_addtd3	
bid_unordtd2	dpd_divsd3	
bswapdi2	dpd_divtd3	
bswapsi2	dpd_eqdd2	
builtin_classify_type	dpd_eqsd2	
builtin_crassify_type	dpd_eqtd2	
builtin_next_arg	dpd_extendddtd2	
clear_cache	dpd_extendddtf	
clzdi2	dpd_extendddxf	
clzsi2	dpd_extenddax1	
clzti2	dpd_extenddftd	
cmpda230	dpd_extendsddd2	
cmpdf2	dpd_extendsddf	
cmpdi210	dpd_extendsdtd2	
cmpdq230	dpd_extendsdtf	
cmpha230	dpd_extendsdxf	
cmphq230	dpd_extendsfdd	
cmpqq230	dpd_extendsfsd	
cmpsa230	dpd_extendsftd	
cmpsf2	dpd_extendtftd	
cmpsq230	dpd_extendxftd	
cmpta230	dpd_fixdddi	
cmptf2	dpd_fixddsi	
cmpti210	dpd_fixsddi	
cmpuda230	dpd_fixsdsi	
cmpudq230	dpd_fixtddi	
cmpuha230	dpd_fixtdsi	
cmpuhq230	dpd_fixunsdddi	19
cmpuqq230	dpd_fixunsddsi	
cmpusa2 30	dpd_fixunssddi	
cmpusq230	dpd_fixunssdsi	
cmputa2 30	dpd_fixunstddi	
CTOR_LIST	dpd_fixunstdsi	
ctzdi211	dpd_floatdidd	
ctzsi211	dpd_floatdisd	20
ctzti?	dnd floatditd	20

dpd_floatsidd		extendsftf2	
$\verb dpd_floatsisd$		extendsfxf2	13
dpd_floatsitd	20	ffsdi2	11
dpd_floatunsdidd	20	ffsti2	11
dpd_floatunsdisd	20	fixdfdi	13
dpd_floatunsditd	20	fixdfsi	13
dpd_floatunssidd		fixdfti	13
dpd_floatunssisd		fixsfdi	13
dpd_floatunssitd		fixsfsi	
dpd_gedd2		fixsfti	
dpd_gesd2		fixtfdi	
dpd_getd2		fixtfsi	
dpd_gtdd2		fixtfti	
dpd_gtsd2		fixunsdfdi	
dpd_gttd2		fixunsdfsi	
dpd_ledd2		fixunsdfti	
dpd_lesd2		fixunssfdi	
dpd_letd2		fixunssfsi	
dpd_ltdd2		fixunssfti	
dpd_ltsd2		fixunstfdi	
dpd_lttd2	21	fixunstfsi	13
dpd_muldd3	17	fixunstfti	14
dpd_mulsd3	17	fixunsxfdi	13
dpd_multd3	17	fixunsxfsi	13
dpd_nedd2	21	fixunsxfti	14
dpd_negdd2		fixxfdi	13
dpd_negsd2		fixxfsi	13
dpd_negtd2		fixxfti	
dpd_nesd2		floatdidf	
dpd_netd2		floatdisf	
dpd_subdd3		floatditf	
dpd_subsd3		floatdixf	
dpd_subtd3		floatsidf	
dpd_truncdddf		floatsisf	
dpd_truncddsd2		floatsitf	
dpd_truncddsf		floatsixf	
dpd_truncdfsd		floattidf	
dpd_truncsdsf		floattisf	
dpd_trunctddd2		floattitf	
dpd_trunctddf		floattixf	
dpd_trunctdsd2		floatundidf	
dpd_trunctdsf		floatundisf	
dpd_trunctdtf		floatunditf	
dpd_trunctdxf		floatundixf	
dpd_trunctfdd		floatunsidf	
dpd_trunctfsd		floatunsisf	
dpd_truncxfdd		floatunsitf	
dpd_truncxfsd		floatunsixf	
dpd_unorddd2		floatuntidf	
dpd_unordsd2		floatuntisf	
$\verb dpd_unordtd2$	20	floatuntitf	
DTOR_LIST		floatuntixf	
eqdf2		fractdadf	-
eqsf2		fractdadi	
eqtf2		fractdadq	
extenddftf2		fractdaha2	
extenddfxf2		fractdahi	
extendsfdf?	13	fractdahg	33

	9.9		20
fractdaqi		fractdqsi	
fractdaqq		fractdqsq2	
fractdasa2		fractdqta	
fractdasf		fractdqti	
fractdasi	34	fractdquda	
fractdasq	33	fractdqudq	
fractdata2	33	fractdquha	
fractdati	34	fractdquhq	32
fractdauda	33	fractdquqq	
fractdaudq	33	fractdqusa	
fractdauha	33	fractdqusq	
fractdauhq		fractdquta	32
fractdauqq		fracthada2	
fractdausa		fracthadf	
fractdausq		fracthadi	
fractdauta		fracthadq	
fractdfda		fracthahi	
fractdfdq		fracthahq	
fractdfha		fracthaqi	
fractdfhq		fracthaqq	
fractdfqq		fracthasa2	
fractdfsa		fracthasf	
fractdfsq		fracthasi	
fractdfta		fracthasq	
fractdfuda		fracthata2	
fractdfudq		fracthati	
fractdfuha		fracthauda	
fractdfuhq		fracthaudq	
fractdfuqq	41	fracthauha	
fractdfusa	41	fracthauhq	
fractdfusq	41	fracthauqq	
fractdfuta	41	fracthausa	
fractdida	40	fracthausq	32
fractdidq	40	fracthauta	33
fractdiha	40	fracthida	39
fractdihq	40	fracthidq	39
fractdiqq	40	fracthiha	39
fractdisa	40	fracthihq	39
fractdisq	40	fracthiqq	39
fractdita	40	fracthisa	39
fractdiuda	40	fracthisq	39
fractdiudq	40	fracthita	
fractdiuha	4.0	fracthiuda	
fractdiuhq		fracthiudq	
fractdiuqq		fracthiuha	
fractdiusa		fracthiuhq	
fractdiusq		fracthiuqq	
fractdiuta		fracthiusa	
fractdqda		fracthiusq	
fractdqdf		fracthiuta	
fractdqdi		fracthqda	
fractdqha		fracthqdf	
fractdqhi		fracthqdi	
fractdqhq2		fracthqdq2	
fractdqqi		fracthqha	
fractdqqq2		fracthqhi	
fractdqsa		fracthqqi	
fractdasf	32	fracthaga?	31

	0.4		~ ~
fracthqsa		fractsahq	
fracthqsf		fractsaqi	
fracthqsi		fractsaqq	
fracthqsq2		fractsasf	
fracthqta		fractsasi	
fracthqti		fractsasq	
fracthquda	31	fractsata2	33
fracthqudq	31	fractsati	
fracthquha	31	fractsauda	33
fracthquhq	31	fractsaudq	33
fracthquqq	31	fractsauha	33
fracthqusa	31	fractsauhq	33
fracthqusq	31	fractsauqq	33
fracthquta	31	fractsausa	33
fractqida		fractsausq	33
fractqidq		fractsauta	
fractqiha		fractsfda	
fractqihq		fractsfdq	
fractqiqq		fractsfha	
fractqisa		fractsfhq	
fractqisq		fractsfqq	
fractqita		fractsfsa	
fractqiuda		fractsfsq	
fractqiudq		fractsfta	
fractqiuha		fractsfuda	
fractqiuhq		fractsfudq	
fractqiuqq		fractsfuha	
		fractsfuhq	
fractqiusa		fractsfuqq	
fractqiusq			
fractqiuta		fractsfusa	
fractqqda		fractsfusq	
fractqqdf		fractsfuta	
fractqqdi		fractsida	
fractqqdq2		fractsidq	
fractqqha		fractsiha	
fractqqhi		fractsihq	
fractqqhq2		fractsiqq	
fractqqqi		fractsisa	
fractqqsa		fractsisq	
fractqqsf		fractsita	
fractqqsi		fractsiuda	
fractqqsq2	~ ~	fractsiudq	
fractqqta		fractsiuha	
fractqqti		fractsiuhq	
fractqquda		fractsiuqq	
fractqqudq		fractsiusa	40
fractqquha	31	fractsiusq	40
fractqquhq	30	fractsiuta	40
fractqquqq	30	fractsqda	
fractqqusa	31	fractsqdf	32
fractqqusq	30	fractsqdi	
fractqquta		fractsqdq2	31
fractsada2		fractsqha	
fractsadf		fractsqhi	
fractsadi	33	fractsqhq2	
fractsadq		fractsqqi	32
fractsaha2		fractsqqq2	
fractsahi	33		31

fractsqsf	32	fractudaqq	37
fractsqsi	32	fractudasa	
fractsqta	31	fractudasf	
fractsqti	32	fractudasi	38
fractsquda	32	fractudasq	38
fractsqudq	31	fractudata	38
fractsquha	31	fractudati	38
fractsquhq	31	fractudaudq	38
fractsquqq		fractudauha2	
fractsqusa		fractudauhq	38
fractsqusq		fractudauqq	
fractsquta		fractudausa2	
fracttada2		fractudausq	
fracttadf		fractudauta2	
fracttadi		fractudqda	
fracttadq		fractudqdf	
fracttaha2		fractudqdi	
fracttahi		fractudqdq	
fracttahq		fractudqha	
fracttaqi		fractudqhi	
fracttaqq		fractudqhq	
fracttasa2		fractudqqi	
fracttasf		fractudqqq	
fracttasi		fractudqsa	
fracttasq		fractudqsf	
fracttati		fractudqsi	
fracttauda		fractudqsq	
fracttaudq		fractudqta	
fracttauha		fractudqti	
fracttauhq		fractudquda	
fracttauqq		fractudquha	
fracttausa		fractudquhq2	
fracttausq		fractudquqq2	
fracttauta		fractudqusa	
fracttida		fractudqusq2	
		fractudquta	
fracttidq		fractuhada	
fracttiha			
fracttihq		fractuhadf fractuhadi	
		<del></del>	
fracttisa		fractuhadq	
fracttisq		fractuhaha	
fracttita		fractuhahi	
		fractuhahq	
fracttiudq		fractuhaqi	
fracttiuha		fractuhaqq	
fracttiuhq		fractuhasa	
fracttiuqq		fractuhasf	
fracttiusa		fractuhasi	
fracttiusq		fractuhasq	
fracttiuta		fractuhata	
fractudada		fractuhati	
fractudadf		fractuhauda2	
fractudadi		fractuhaudq	
fractudadq		fractuhauhq	
fractudaha		fractuhauqq	
fractudahi		fractuhausa2	
fractudahq		fractuhausq	
fractudaqi	38	fractuhauta2	37

fractuhqda	35	fractunshisa	52
fractuhqdf		fractunshisq	
fractuhqdi	35	fractunshita	53
fractuhqdq	35	fractunshiuda	53
fractuhqha	35	fractunshiudq	53
fractuhqhi	35	fractunshiuha	53
fractuhqhq	35	fractunshiuhq	53
fractuhqqi	35	fractunshiuqq	53
fractuhqqq	35	fractunshiusa	53
fractuhqsa	35	fractunshiusq	53
fractuhqsf	35	fractunshiuta	53
fractuhqsi	35	fractunshqdi	50
fractuhqsq	35	fractunshqhi	50
fractuhqta	35	fractunshqqi	50
fractuhqti		fractunshqsi	
fractuhquda	35	fractunshqti	50
fractuhqudq2		fractunsqida	52
fractuhquha		fractunsqidq	
fractuhquqq2		fractunsqiha	
fractuhqusa		fractunsqihq	
fractuhqusq2	35	fractunsqiqq	
fractuhquta		fractunsqisa	
fractunsdadi		fractunsqisq	
fractunsdahi		fractunsqita	
fractunsdaqi	51	fractunsqiuda	
fractunsdasi		fractunsqiudq	
fractunsdati		fractunsqiuha	
fractunsdida		fractunsqiuhq	
fractunsdidq	53	fractunsqiuqq	
fractunsdiha		fractunsqiusa	
fractunsdihq		fractunsqiusq	
fractunsdiqq		fractunsqiuta	
fractunsdisa		fractunsqqdi	
fractunsdisq	53	fractunsqqhi	
fractunsdita		fractunsqqqi	
fractunsdiuda		fractunsqqsi	
fractunsdiudq		fractunsqqti	
fractunsdiuha		fractunssadi	
fractunsdiuhq		fractunssahi	
fractunsdiuqq		fractunssaqi	
fractunsdiusa		fractunssasi	
fractunsdiusq	53	fractunssati	
fractunsdiuta		fractunssida	
fractunsdqdi		fractunssidq	
fractunsdqhi		fractunssiha	
fractunsdqqi		fractunssihq	53
fractunsdqsi		fractunssiqq	53
fractunsdqti		fractunssisa	
fractunshadi		fractunssisq	53
fractunshahi		fractunssita	
fractunshaqi		fractunssiuda	
fractunshasi		fractunssiudq	
fractunshati		fractunssiuha	
fractunshida		fractunssiuhq	
fractunshidq		fractunssiuqq	
fractunshiha		fractunssiusa	
fractunshihq		fractunssiusq	
fractunshigg		fractunssiuta	

fractunssqdi	50	fractunsusqqi	51
fractunssqhi	50	fractunsusqsi	
fractunssqqi	50	fractunsusqti	51
fractunssqsi	50	fractunsutadi	
fractunssqti		fractunsutahi	
fractunstadi	51	fractunsutaqi	52
fractunstahi	51	fractunsutasi	52
fractunstaqi	51	fractunsutati	52
fractunstasi	51	fractuqqda	34
fractunstati	51	fractuqqdf	
fractunstida	54	fractuqqdi	35
fractunstidq	54	fractuqqdq	34
fractunstiha	54	fractuqqha	
fractunstihq	54	fractuqqhi	
fractunstiqq		fractuqqhq	
fractunstisa		fractuqqqi	
fractunstisq		fractuqqqq	
fractunstita		fractuqqsa	
fractunstiuda		fractuqqsf	
fractunstiudq		fractuqqsi	
fractunstiuha		fractuqqsq	
fractunstiuhq		fractuqqta	
fractunstiuqq		fractuqqti	
fractunstiusa		fractuqquda	
fractunstiusq		fractuqqudq2	
fractunstiuta		fractuqquha	
fractunsudadi		fractuqquhq2	
fractunsudahi		fractuqqusa	
fractunsudaqi		fractuqqusq2	
fractunsudasi		fractuqquta	
fractunsudati		fractusada	
fractunsudqdi		fractusadf	
fractunsudqhi		fractusadi	
fractunsudqqi		fractusadq	
fractunsudqsi		fractusaha	
fractunsudqti		fractusahi	
fractunsuhadi		fractusahq	
fractunsuhahi		fractusaqi	
fractunsuhaqi		fractusaqq	
fractunsuhasi		fractusasa	
fractunsuhati		fractusasf	
fractunsuhqdi		fractusasi	
fractunsuhqhi		fractusasq	
fractunsuhqqi		fractusata	
fractunsuhqsi		fractusati	
fractunsuhqti		fractusauda2	
fractunsuqqdi		fractusaudq	
fractunsuqqhi		fractusauha2	
fractunsuqqqi		fractusauhq	
fractunsuqqsi		fractusauqq	
fractunsuqqti		fractusausq	
fractunsusadi		fractusauta2	
fractunsusahi		fractusqda	
fractunsusaqi		fractusqdf	
fractunsusasi		fractusqdi	
fractunsusati		fractusqdq	
	<u> </u>		50
fractunsusqdi	51	fractusqha	35

fractusqhq	35	ltdf2	
fractusqqi		ltsf2	15
fractusqqq	35	lttf2	
fractusqsa		main	
fractusqsf		moddi3	
fractusqsi		modsi3	
fractusqsq	35	modti3	
fractusqta	35	morestack_current_segment	58
fractusqti		morestack_initial_sp	
fractusquda		morestack_segments	
fractusqudq2		mulda3	
fractusquha	36	muldc3	
fractusquhq2		muldf3	
fractusquqq2	35	muldi3	10
fractusqusa	36	muldq3	25
fractusquta	36	mulha3	25
fractutada	38	mulhq3	25
fractutadf	39	mulqq3	25
fractutadi	39	mulsa3	25
fractutadq	38	mulsc3	16
fractutaha	38	mulsf3	12
fractutahi	39	mulsi3	10
fractutahq	38	mulsq3	25
fractutaqi	38	multa3	25
fractutaqq	38	multc3	16
fractutasa	38	multf3	12
fractutasf	39	multi3	10
fractutasi	39	muluda3	25
fractutasq	38	muludq3	25
fractutata	38	muluha3	25
fractutati	39	muluhq3	25
fractutauda2	38	muluqq3	25
fractutaudq	38	mulusa3	25
fractutauha2	38	mulusq3	25
fractutauhq	38	muluta3	25
fractutauqq	38	mulvdi3	11
fractutausa2	38	mulvsi3	
fractutausq	38	mulxc3	16
gedf2		mulxf3	12
gesf2		nedf2	15
getf2	15	negda2	27
gtdf2		negdf2	
gtsf2	16	negdi2	10
gttf2	16	negdq2	27
ledf2	16	negha2	27
lesf2	16	neghq2	27
letf2		negqq2	27
lshrdi3		negsa2	27
lshrsi3	. 9	negsf2	12
lshrti3	. 9	negsq2	
lshruda3		negta2	27
lshrudq3	29	negtf2	12
lshruha3	29	negti2	10
lshruhq3	29	neguda2	27
lshruqq3		negudq2	27
lshrusa3		neguha2	
lshrusq3		neguhq2	
lshruta3		_neguqq2	27

•	0.7		40
negusa2		satfractdiudq	
negusq2		satfractdiuha	
neguta2		satfractdiuhq	
negvdi2		satfractdiuqq	
negvsi2		satfractdiusa	
negxf2		satfractdiusq	
nesf2		satfractdiuta	
netf2		satfractdqda	
paritydi2		satfractdqha	
paritysi2		satfractdqhq2	
parityti2		satfractdqqq2	
popcountdi2		satfractdqsa	
popcountsi2		satfractdqsq2	
popcountti2		satfractdqta	
powidf2		satfractdquda	
powisf2		satfractdqudq	
powitf2		satfractdquha	
powixf2		satfractdquhq	
$\verb satfractdadq$		satfractdquqq	
satfractdaha2		satfractdqusa	
satfractdahq	43	satfractdqusq	
satfractdaqq	43	satfractdquta	
satfractdasa2		satfracthada2	43
satfractdasq		satfracthadq	43
satfractdata2	43	satfracthahq	42
satfractdauda	43	satfracthaqq	
satfractdaudq	43	satfracthasa2	43
satfractdauha	43	satfracthasq	42
satfractdauhq	43	satfracthata2	43
satfractdauqq	43	satfracthauda	43
satfractdausa	43	satfracthaudq	43
satfractdausq		satfracthauha	43
satfractdauta	43	satfracthauhq	43
satfractdfda	50	satfracthauqq	43
satfractdfdq		satfracthausa	43
satfractdfha	50	satfracthausq	43
satfractdfhq	50	satfracthauta	43
satfractdfqq	50	satfracthida	48
satfractdfsa		satfracthidq	48
satfractdfsq		satfracthiha	48
satfractdfta	50	satfracthihq	48
satfractdfuda	50	satfracthiqq	
satfractdfudq		satfracthisa	
satfractdfuha	50	satfracthisq	
satfractdfuhq	50	satfracthita	48
satfractdfuqq		satfracthiuda	48
satfractdfusa		satfracthiudq	48
satfractdfusq		satfracthiuha	
satfractdfuta		satfracthiuhq	
satfractdida		satfracthiuqq	
satfractdidq		satfracthiusa	
satfractdiha		satfracthiusq	
satfractdihq		satfracthiuta	
satfractdiqq		satfracthqda	
satfractdisa		satfracthqdq2	
satfractdisq		satfracthqha	
satfractdita	49	satfracthqqq2	
satfractdiuda	49		42

$\verb satfracthqsq2$		satfractsfha	
satfracthqta		satfractsfhq	
satfracthquda		satfractsfqq	49
$\verb satfracthqudq$		satfractsfsa	
satfracthquha		satfractsfsq	
satfracthquhq	42	satfractsfta	49
satfracthquqq	42	satfractsfuda	50
satfracthqusa	42	satfractsfudq	50
satfracthqusq	42	satfractsfuha	50
satfracthquta	42	satfractsfuhq	50
satfractqida		satfractsfuqq	49
satfractqidq		satfractsfusa	
satfractqiha		satfractsfusq	
satfractqihq		satfractsfuta	
satfractqiqq		satfractsida	
satfractqisa		satfractsidq	
satfractqisq		satfractsiha	
satfractqita		satfractsihq	
satfractqiuda		satfractsiqq	
satfractqiudq		satfractsisa	
satfractqiuha		satfractsisq	
satfractqiuhq		satfractsita	
satfractqiuqq		satfractsiuda	
satfractqiusa		satfractsiudq	
satfractqiusq		satfractsiuha	
satfractqiuta		satfractsiuhq	
satfractqqda		satfractsiuqq	
satfractqqdq2		satfractsiusa	
satfractqqha		satfractsiusq	
satfractqqhq2		satfractsiuta	
satfractqqsa		satfractsqda	
satfractqqsq2		satfractsqdq2	
satfractqqta		satfractsqha	
satfractqquda		satfractsqhq2	
satfractqqudq		satfractsqqq2	
satfractqquha		satfractsqsa	
satfractqquhq		satfractsqta	
satfractqquqq		satfractsquda	
satfractqqusa		satfractsqudq	
satfractqqusq		satfractsquha	
satfractqquta		satfractsquhq	
satfractsada2		satfractsquqq	
satfractsadq		satfractsqusa	
satfractsaha2		satfractsqusq	
satfractsahq		satfractsquta	
satfractsaqq		satfracttada2	
satfractsasq		satfracttadq	
satfractsata2		satfracttaha2	
satfractsauda	43	satfracttahq	44
$\verb satfractsaudq$		satfracttaqq	
satfractsauha		satfracttasa2	
$\verb satfractsauhq$		satfracttasq	
satfractsauqq		satfracttauda	44
satfractsausa		satfracttaudq	44
satfractsausq		satfracttauha	44
satfractsauta		satfracttauhq	44
satfractsfda	49	satfracttauqq	44
satfractsfdo	49	satfracttausa	44

satfracttausq		satfractuhauhq	46
satfracttauta	44	satfractuhauqq	
satfracttida	49	satfractuhausa2	46
satfracttidq	49	satfractuhausq	46
satfracttiha	49	satfractuhauta2	46
satfracttihq	49	satfractuhqda	44
satfracttiqq	49	satfractuhqdq	
satfracttisa		satfractuhqha	
satfracttisq		satfractuhqhq	
satfracttita		satfractuhqqq	
satfracttiuda		satfractuhqsa	
satfracttiudq		satfractuhqsq	
satfracttiuha		satfractuhqta	
satfracttiuhq		satfractuhquda	
satfracttiuqq		satfractuhqudq2	
satfracttiusa		satfractuhquha	
satfracttiusq		satfractuhquqq2	
satfracttiuta		satfractuhqusa	
satfractudada		satfractuhqusq2	
satfractudadq		satfractuhquta	
satfractudaha		satfractunsdida	
satfractudahq		satfractunsdidq	
satfractudaqq		satfractunsdiha	
satfractudasa		satfractunsdihq	
satfractudasq		satfractunsdiqq	
satfractudata		satfractunsdisa	
satfractudaudq		satfractunsdisq	
satfractudauha2		satfractunsdita	
		satfractunsdita	
satfractudauhq			
satfractudauqq		satfractunsdiudq	
satfractudausa2		satfractunsdiuha	
satfractudausq		satfractunsdiuhq	
satfractudauta2		satfractunsdiuqq	
satfractudqda		satfractunsdiusa	
satfractudqdq		satfractunsdiusq	
satfractudqha		satfractunsdiuta	
satfractudqhq		satfractunshida	
satfractudqqq		satfractunshidq	
satfractudqsa		satfractunshiha	
satfractudqsq		satfractunshihq	
satfractudqta		satfractunshiqq	
satfractudquda		satfractunshisa	
satfractudquha		satfractunshisq	
satfractudquhq2		satfractunshita	
satfractudquqq2		satfractunshiuda	
satfractudqusa		satfractunshiudq	
satfractudqusq2		satfractunshiuha	
satfractudquta		satfractunshiuhq	
satfractuhada		satfractunshiuqq	
satfractuhadq		satfractunshiusa	
satfractuhaha		satfractunshiusq	
$\verb satfractuhahq$		satfractunshiuta	
$\verb satfractuhaqq$		satfractunsqida	
satfractuhasa		$\verb satfractunsqidq$	
$\verb satfractuhasq$		$\verb satfractunsqiha$	
satfractuhata		$\verb satfractunsqihq$	
satfractuhauda2		$\verb satfractunsqiqq$	
satfractuhaudg	46	satfractunsgisa	54

$\verb satfractunsqisq$		satfractusadq	
satfractunsqita		satfractusaha	
satfractunsqiuda		satfractusahq	
satfractunsqiudq		satfractusaqq	
satfractunsqiuha		satfractusasa	
satfractunsqiuhq	54	satfractusasq	46
satfractunsqiuqq	54	satfractusata	46
satfractunsqiusa	54	satfractusauda2	
satfractunsqiusq	54	satfractusaudq	
satfractunsqiuta	54	satfractusauha2	46
satfractunssida	55	satfractusauhq	46
satfractunssidq	55	satfractusauqq	46
satfractunssiha	55	satfractusausq	46
satfractunssihq	55	satfractusauta2	47
satfractunssiqq	55	satfractusqda	45
satfractunssisa	55	satfractusqdq	45
satfractunssisq	55	satfractusqha	45
satfractunssita		satfractusqhq	
satfractunssiuda		satfractusqqq	
satfractunssiudq		satfractusqsa	
satfractunssiuha		satfractusqsq	
satfractunssiuhq		satfractusqta	
satfractunssiuqq		satfractusquda	
satfractunssiusa		satfractusqudq2	
satfractunssiusq		satfractusquha	
satfractunssiuta		satfractusquhq2	
satfractunstida		satfractusquqq2	
satfractunstidq		satfractusqusa	
satfractunstiha		satfractusquta	
satfractunstihq		satfractutada	
satfractunstiqq		satfractutadq	
satfractunstisa		satfractutaha	
satfractunstisq		satfractutahq	
satfractunstita		satfractutaqq	
satfractunstiuda		satfractutasa	
satfractunstiudq		satfractutasq	
satfractunstiuha		satfractutata	
satfractunstiuhq		satfractutauda2	
satfractunstiuqq		satfractutaudq	
satfractunstiusa		satfractutauha2	
satfractunstiusq		satfractutauhq	
satfractunstiuta		satfractutauqq	
satfractuqqda		satfractutausa2	
satfractuqqdq		satfractutausq	
satfractuqqha		splitstack_find	
satfractuqqhq		ssaddda3	
satfractuqqqq		ssadddq3	
satfractuqqsa		ssaddha3	
satfractuqqsq		ssaddhq3	
satfractuqqta		ssaddqq3ssaddsa3	
satfractuqquda		ssaddsq3	
satfractuqquaq2		ssaddta3	
		ssashlda3	
satfractuqquhq2			
satfractuqqusa		ssashldq3ssashlha3	
satfractuqqusq2		ssashlhq3	
satfractusada	44 46	-	29 29
		220211200	7.54

ssashlsq3	20	trunctfsf2	12
ssashlta3		truncxfdf2	
ssdivda3		truncxfsf2	
ssdivdq3		ucmpdi2	
ssdivha3		ucmpti2udivdi3	
ssdivhq3			
ssdivqq3		udivmoddi4	
ssdivsa3		udivmodti4	
ssdivsq3		udivsi3	
ssdivta3		udivti3	
ssmulda3		udivuda3	
ssmuldq3		udivudq3	
ssmulha3		udivuha3	
ssmulhq3		udivuhq3	
ssmulqq3		udivuqq3	26
ssmulsa3	25	udivusa3	26
ssmulsq3	25	udivusq3	26
ssmulta3	25	udivuta3	26
ssnegda2	28	umoddi3	10
ssnegdq2		umodsi3	10
ssnegha2		umodti3	10
ssneghq2		unorddf2	
ssnegqq2		unordsf2	
ssnegsa2		unordtf2	
ssnegsq2		usadduda3	
ssnegta2		usaddudq3	
sssubda3		usadduha3	
sssubdq3		usadduhq3	
sssubha3		usadduqq3	
		usaddusa3	
sssubhq3			
sssubqq3		usaddusq3	
sssubsa3			
sssubsq3		usashluda3	
sssubta3		usashludq3	
subda3		usashluha3	
subdf3		usashluhq3	
subdq3		usashluqq3	
subha3		usashlusa3	
subhq3		usashlusq3	
subqq3	23	usashluta3	
subsa3	24	usdivuda3	
subsf3		usdivudq3	
subsq3	23	usdivuha3	27
subta3		usdivuhq3	27
subtf3	12	usdivuqq3	27
subuda3	24	usdivusa3	27
subudq3	24	usdivusq3	27
subuha3		usdivuta3	27
subuhq3		usmuluda3	
subuqq3		usmuludq3	
subusa3		usmuluha3	
subusq3		usmuluhq3	
subuta3		usmuluqq3	
subvdi3		usmulusa3	
subvsi3		usmulusq3	
subxf3		usmuluta3	
truncdfsf2		usneguda2	
trunctfdf?	13		28

usneguha2	28	ADJUST_REG_ALLOC_ORDER	508
usneguhq2		aggregates as return values	
usneguqq2		alias	
usnegusa2		ALL_REGS	
usnegusq2		allocate_stack instruction pattern	
usneguta2		alternate entry points	
ussubuda3		analyzer	
ussubudq3		analyzer, debugging	
ussubuha3		analyzer, internals	
ussubuhq3		anchored addresses	
ussubuqq3		and	
ussubusa3		and and attributes	
ussubusq3		and, canonicalization of	
ussubuta3		andm3 instruction pattern	
ussubutao		ANNOTATE_EXPR	
		annotations	
\		APPLY_RESULT_SIZE	
\3		arg_pointer_rtx	
		ARG_POINTER_CFA_OFFSET	
<b>0</b>		ARG_POINTER_REGNUM	
		ARG_POINTER_REGNUM and virtual registers	
'0' in constraint		ARGS_GROW_DOWNWARD	
		argument passing	
$\mathbf{A}$		arguments in registers	
		arguments on stack	
abort		arithmetic library	
abs 2		arithmetic shift	
abs and attributes4		arithmetic shift with signed saturation	
ABS_EXPR1		arithmetic shift with unsigned saturation	
absence_set 4		arithmetic, in RTL	
absm2 instruction pattern 4	106	ARITHMETIC_TYPE_P	197
absolute value	291	array	163
ABSU_EXPR1	.77	ARRAY_RANGE_REF	175
access to operands		ARRAY_REF	
access to special operands	263	ARRAY_TYPE	163
accessors		AS_NEEDS_DASH_FOR_PIPED_INPUT	
ACCUM_TYPE_SIZE5	501	ashift	291
ACCUMULATE_OUTGOING_ARGS 5	35	ashift and attributes	454
ACCUMULATE_OUTGOING_ARGS and stack frames		ashiftrt	291
5	550	ashiftrt and attributes	454
acosm2 instruction pattern	107	ashlm3 instruction pattern	405
ADA_LONG_TYPE_SIZE 5		ashrm3 instruction pattern	
Adding a new GIMPLE statement code 2		asinm2 instruction pattern	407
ADDITIONAL_REGISTER_NAMES 6		asm_fprintf	
addm3 instruction pattern		asm_input	
addmodecc instruction pattern 4		asm_input and '/v'	
addptrm3 instruction pattern 3		asm_noperands	
addr_diff_vec		asm_operands and '/v'	
addr_diff_vec, length of		asm_operands, RTL sharing	
addr_vec		asm_operands, usage	
addr_vec, length of		ASM_APP_OFF	
ADDR_EXPR		ASM_APP_ON	
address constraints		ASM_COMMENT_START	
address_operand		ASM_DECLARE_COLD_FUNCTION_NAME	
addressing modes		ASM_DECLARE_COLD_FUNCTION_NAME  ASM_DECLARE_COLD_FUNCTION_SIZE	
addvm4 instruction pattern		ASM_DECLARE_FUNCTION_NAME	
ADJUST_FIELD_ALIGN 4		ASM_DECLARE_FUNCTION_SIZE	
	194 158		606
MIZ 111.3 1 T.N.3 N. T. P.N. + I.D.		A SECULIE A DE LE LE LE LE LE MANTE.	4 11 11.7

ASM_DECLARE_REGISTER_GLOBAL	607	ASM_WEAKEN_DECL	
ASM_FINAL_SPEC		ASM_WEAKEN_LABEL	
ASM_FINISH_DECLARE_OBJECT	607	assemble_name	604
ASM_FORMAT_PRIVATE_NAME		assemble_name_raw	
ASM_FPRINTF_EXTENSIONS		assembler format	
ASM_GENERATE_INTERNAL_LABEL	610	assembler instructions in RTL	303
ASM_MAYBE_OUTPUT_ENCODED_ADDR_RTX	528	ASSEMBLER_DIALECT	
ASM_NO_SKIP_IN_TEXT	624	assigning attribute values to insns	
ASM_OUTPUT_ADDR_DIFF_ELT	619	ASSUME_EXTENDED_UNWIND_CONTEXT	533
ASM_OUTPUT_ADDR_VEC_ELT		asterisk in template	
ASM_OUTPUT_ALIGN		atan2m3 instruction pattern	
ASM_OUTPUT_ALIGN_WITH_NOP		atanm2 instruction pattern	408
ASM_OUTPUT_ALIGNED_BSS	603	atomic	679
ASM_OUTPUT_ALIGNED_COMMON	603	<pre>atomic_add_fetchmode instruction pattern</pre>	430
ASM_OUTPUT_ALIGNED_DECL_COMMON	603	atomic_addmode instruction pattern	430
ASM_OUTPUT_ALIGNED_DECL_LOCAL	604	<pre>atomic_and_fetchmode instruction pattern</pre>	430
ASM_OUTPUT_ALIGNED_LOCAL	604	atomic_andmode instruction pattern	430
ASM_OUTPUT_ASCII	601	atomic_bit_test_and_complementmode instruct	ion
ASM_OUTPUT_CASE_END	620	pattern	430
ASM_OUTPUT_CASE_LABEL	620	<pre>atomic_bit_test_and_resetmode instruction</pre>	
ASM_OUTPUT_COMMON	602	pattern	430
ASM_OUTPUT_DEBUG_LABEL	610	atomic_bit_test_and_setmode instruction patt	ern
ASM_OUTPUT_DEF	611		430
ASM_OUTPUT_DEF_FROM_DECLS	611	$\verb"atomic_compare_and_swap" mode instruction patt$	ern
ASM_OUTPUT_DWARF_DATAREL	630		428
ASM_OUTPUT_DWARF_DELTA	630	atomic_exchangemode instruction pattern	429
ASM_OUTPUT_DWARF_OFFSET	630	atomic_fetch_addmode instruction pattern	430
ASM_OUTPUT_DWARF_PCREL	630	atomic_fetch_andmode instruction pattern	430
ASM_OUTPUT_DWARF_TABLE_REF	631	<pre>atomic_fetch_nandmode instruction pattern</pre>	430
ASM_OUTPUT_DWARF_VMS_DELTA	630	atomic_fetch_ormode instruction pattern	430
ASM_OUTPUT_EXTERNAL	609	atomic_fetch_submode instruction pattern	430
ASM_OUTPUT_FDESC		atomic_fetch_xormode instruction pattern	430
ASM_OUTPUT_FUNCTION_LABEL		atomic_loadmode instruction pattern	429
ASM_OUTPUT_INTERNAL_LABEL		atomic_nand_fetchmode instruction pattern	430
ASM_OUTPUT_LABEL		atomic_nandmode instruction pattern	430
ASM_OUTPUT_LABEL_REF		atomic_or_fetchmode instruction pattern	
ASM_OUTPUT_LABELREF		atomic_ormode instruction pattern	
ASM_OUTPUT_LOCAL		atomic_storemode instruction pattern	429
ASM_OUTPUT_MAX_SKIP_ALIGN		atomic_sub_fetchmode instruction pattern	
ASM_OUTPUT_MEASURED_SIZE		atomic_submode instruction pattern	430
ASM_OUTPUT_OPCODE		atomic_test_and_set instruction pattern	
ASM_OUTPUT_POOL_EPILOGUE		atomic_xor_fetchmode instruction pattern	
ASM_OUTPUT_POOL_PROLOGUE	601	atomic_xormode instruction pattern	
ASM_OUTPUT_REG_POP		attr	
ASM_OUTPUT_REG_PUSH		attr_flag	454
ASM_OUTPUT_SIZE_DIRECTIVE		attribute expressions	
ASM_OUTPUT_SKIP		attribute specifications	
ASM_OUTPUT_SOURCE_FILENAME		attribute specifications example	
ASM_OUTPUT_SPECIAL_POOL_ENTRY		ATTRIBUTE_ALIGNED_VALUE	
ASM_OUTPUT_SYMBOL_REF		attributes	
ASM_OUTPUT_TYPE_DIRECTIVE		attributes, defining	
ASM_OUTPUT_WEAK_ALIAS		attributes, target-specific	
ASM_OUTPUT_WEAKREF		autoincrement addressing, availability	
ASM_PREFERRED_EH_DATA_FORMAT		autoincrement/decrement addressing	
ASM_SPEC		automata_option	
ASM_STABD_OP		automaton based pipeline description 460,	
ASM_STABN_OP		automaton based scheduler	
ASM STABS OP		avgm3 ceil instruction pattern	

2 fleer instruction nottons	40 <i>c</i>	build3	160
avgm3_floor instruction pattern			
AVOID_CCMODE_COPIES	510	build4	
		build5	
В		build6	
D		builtin_longjmp instruction pattern	424
backslash	344	builtin_setjmp_receiver instruction pattern	
barrier	307		
barrier and '/f'	268	<pre>builtin_setjmp_setup instruction pattern</pre>	424
barrier and '/v'	266	byte_mode	278
BASE_REG_CLASS		BYTES_BIG_ENDIAN	490
basic block		BYTES_BIG_ENDIAN, effect on subreg	285
Basic Statements		,	
basic-block.h			
basic_block		$\mathbf{C}$	
BASIC_BLOCK			640
		c_register_pragma	
bb_seq		c_register_pragma_with_expansion	
BB_HEAD, BB_END		C statements for assembler output	
BIGGEST_ALIGNMENT		C_COMMON_OVERRIDE_OPTIONS	
BIGGEST_FIELD_ALIGNMENT		cache	
BImode		call	
BIND_EXPR		call instruction pattern	
BINFO_TYPE	199	call usage	313
bit-fields	294	call, in call_insn	268
BIT_AND_EXPR	177	call, in mem	267
BIT_IOR_EXPR	177	call-clobbered register 505,	506
BIT_NOT_EXPR	177	call-saved register	
BIT_XOR_EXPR	177	call-used register	
BITFIELD_NBYTES_LIMITED		call_insn	
BITS_BIG_ENDIAN		call_insn and '/c'	
BITS_BIG_ENDIAN, effect on sign_extract		call_insn and '/f'	
BITS_PER_UNIT		call_insn and '/i'	
BITS_PER_WORD		call_insn and '/j'	
bitwise complement		call_insn and '/s'	
bitwise exclusive-or			
		call_insn and '/u'	
bitwise inclusive-or		call_insn and '/u' or '/i'	
bitwise logical-and		call_insn and '/v'	
BLKmode		call_pop instruction pattern	
BLKmode, and function return values		call_used_regs	
BLOCK_FOR_INSN, gimple_bb	323	call_value instruction pattern	
BLOCK_REG_PADDING		call_value_pop instruction pattern	
blockage instruction pattern		CALL_EXPR	
Blocks	188	CALL_INSN_FUNCTION_USAGE	306
BND32mode	274	CALL_POPS_ARGS	
BND64mode	274	CALL_REALLY_USED_REGISTERS	506
bool	657	CALL_USED_REGISTERS	
BOOL_TYPE_SIZE	500	calling conventions	522
BOOLEAN_TYPE	163	calling functions in RTL	313
branch prediction	322	can_create_pseudo_p	
BRANCH_COST		can_fallthru	
break_out_memory_refs		canadian	
BREAK_STMT		canonicalization of instructions	
BSS_SECTION_ASM_OP		canonicalize_funcptr_for_compare instruction	
bswap		pattern	422
bswapm2 instruction pattern		caret	
		CASE_VECTOR_MODE	
btruncm2 instruction patternbuild0			
build1		CASE_VECTOR_SHORTEN_MODE	
build2	162	casesi instruction pattern	421

cbranchmode4 instruction pattern	419	code_label and '/v'	266
cc_status	573	CODE_LABEL	318
cc0		CODE_LABEL_NUMBER	306
cc0, RTL sharing	314	codes, RTL expression	259
cc0_rtx	287	COImode	274
CC_STATUS_MDEP	573	COLLECT_EXPORT_LIST	656
CC_STATUS_MDEP_INIT		COLLECT_SHARED_FINI_FUNC	614
CC1_SPEC	480	COLLECT_SHARED_INIT_FUNC	614
CC1PLUS_SPEC		COLLECT2_HOST_INITIALIZATION	665
CCmode	574	command-line options, guidelines for	
CDImode		commit_edge_insertions	
CEIL_DIV_EXPR		compare	
CEIL_MOD_EXPR		compare, canonicalization of	
ceilm2 instruction pattern		COMPARE_MAX_PIECES	
CFA_FRAME_BASE_OFFSET		comparison_operator	
CFG verification		compiler passes and files	
CFG, Control Flow Graph		complement, bitwise	
cfghooks.h		complex_mode	
cgraph_finalize_function		COMPLEX_CST	
chain_circular		COMPLEX EXPR	
chain_next		COMPLEX_TYPE	
chain_prev		COMPONENT_REF	
change_address		Compound Expressions	
CHAR_TYPE_SIZE		Compound Lyalues	
check_raw_ptrsm instruction pattern		COMPOUND_EXPR	
check_stack instruction pattern		COMPOUND_LITERAL_EXPR	
check_war_ptrsm instruction pattern		COMPOUND_LITERAL_EXPR_DECL	
CHImode			
		COMPOUND_LITERAL_EXPR_DECL_EXPR computed jump	
class definitions, register		computing the length of an insn	
class, scope		concat	
CLASS_MAX_NREGS		concatn	
CLASS_TYPE_P		cond	
classes of RTX codes		cond and attributes	
CLASSTYPE_DECLARED_CLASS		cond_addmode instruction pattern	
CLASSTYPE_HAS_MUTABLE		cond_andmode instruction pattern	
CLASSTYPE_NON_POD_P		cond_divmode instruction pattern	
CLEANUP_DECL		cond_exec	
CLEANUP_EXPR		cond_fmamode instruction pattern	
CLEANUP_POINT_EXPR		cond_fmsmode instruction pattern	
CLEANUP_STMT		<pre>cond_fnmamode instruction pattern</pre>	
Cleanups		cond_fnmsmode instruction pattern	
clear_cache instruction pattern		<pre>cond_iormode instruction pattern</pre>	
CLEAR_INSN_CACHE		cond_modmode instruction pattern	
CLEAR_RATIO		${\tt cond\_mulmode} \ instruction \ pattern \dots \dots$	
clobber	298	${\tt cond\_smax} {\tt mode} \ {\rm instruction} \ {\rm pattern}$	
clrsb		cond_sminmode instruction pattern	
clrsbm2 instruction pattern		cond_submode instruction pattern	
clz		$\verb"cond_udiv" mode instruction pattern$	417
CLZ_DEFINED_VALUE_AT_ZERO		${\tt cond\_umax\it mode} \ instruction \ pattern$	
clzm2 instruction pattern		${\tt cond\_umin} {\tt mode} \ {\rm instruction} \ {\rm pattern}$	417
cmpmemm instruction pattern		$\verb"cond_umod mode" instruction pattern$	
cmpstrm instruction pattern	413	$\verb"cond_xormode" instruction pattern$	417
cmpstrnm instruction pattern	413	COND_EXPR	177
code generation RTL sequences		condition code register	287
code iterators in '.md' files		condition code status	
code_label		condition codes	
code label and '/i'		conditional execution	

G 1111 1 I I	015		_
Conditional Expressions		conventions, run-time	
conditions, in patterns		conversions	
configuration file		CONVERT_EXPR	
configure terms		copy_rtx	
CONJ_EXPR		copy_rtx_if_shared	
const		copysignm3 instruction pattern	
const_double		cosm2 instruction pattern	
const_double, RTL sharing		costs of instructions	
const_double_operand		cp_namespace_decls	
const_fixed		cp_type_quals	
${\tt const\_int}$		CP_INTEGRAL_TYPE	
const_int and attribute tests		CP_TYPE_CONST_NON_VOLATILE_P	
const_int and attributes	452	CP_TYPE_CONST_P	
const_int, RTL sharing	314	CP_TYPE_RESTRICT_P	197
const_int_operand	347	CP_TYPE_VOLATILE_P	197
const_poly_int		CPLUSPLUS_CPP_SPEC	480
const_poly_int, RTL sharing	314	CPP_SPEC	480
const_string		CPSImode	274
const_string and attributes	452	cpymemm instruction pattern	411
const_true_rtx		CQImode	274
const_vector	280	cross compilation and floating point	631
const_vector, RTL sharing		CROSSING_JUMP_P	
const0_rtx		CRT_CALL_STATIC_FUNCTION	
const1_rtx		crtl->args.pops_args	
const2_rtx		crtl->args.pretend_args_size	
CONST_DECL		crtl->outgoing_args_size	
CONST_DOUBLE_LOW		CRTSTUFF_T_CFLAGS	
CONST_WIDE_INT		CRTSTUFF_T_CFLAGS_S	
CONST_WIDE_INT_ELT		CSImode	
CONST_WIDE_INT_EET		cstoremode4 instruction pattern	
CONST_WIDE_INT_VEC		CTImode	
		ctrapMM4 instruction pattern	
CONSTO_RTX		ctz	
CONST1_RTX		CTZ_DEFINED_VALUE_AT_ZERO	
CONST2_RTX			
constant attributes		ctzm2 instruction pattern	
constant definitions		CUMULATIVE_ARGS	
CONSTANT_ADDRESS_P		current_function_is_leaf	
CONSTANT_P		current_function_uses_only_leaf_regs	
CONSTANT_POOL_ADDRESS_P		current_insn_predicate	400
CONSTANT_POOL_BEFORE_FUNCTION			
constants in constraints		D	
constm1_rtx		D	
constraint modifier characters		DAmode	273
constraint, matching	352	data bypass 462,	463
constraint_num	391	data dependence delays	460
${\tt constraint\_satisfied\_p}$	391	Data Dependency Analysis	334
constraints	350	data structures	489
constraints, defining	388	DATA_ABI_ALIGNMENT	494
constraints, machine specific		DATA_ALIGNMENT	494
constraints, testing		DATA_SECTION_ASM_OP	
constructors, automatic calls		dbr_sequence_length	
constructors, output of			618
CONSTRUCTOR		DBX_BLOCKS_FUNCTION_RELATIVE	
container		DBX_CONTIN_CHAR	
CONTINUE_STMT		DBX_CONTIN_LENGTH	
contributors			626
controlling register usage		DBX_FUNCTION_FIRST	
controlling the compilation driver		DBX_LINES_FUNCTION_RELATIVE	
controlling one compilation driver	100	221-1110-1 01101 1011-1WHA11 VH	041

DBX_NO_XREFS	626	DECL_NON_THUNK_FUNCTION_P	203
DBX_OUTPUT_MAIN_SOURCE_FILE_END	628	DECL_NONCONVERTING_P	202
DBX_OUTPUT_MAIN_SOURCE_FILENAME	628	DECL_NONSTATIC_MEMBER_FUNCTION_P	202
DBX_OUTPUT_NULL_N_SO_AT_MAIN_SOURCE_FILE	_	DECL_OVERLOADED_OPERATOR_P	202
END	629	DECL_PURE_P	195
DBX_OUTPUT_SOURCE_LINE	628	DECL_RESULT	193
DBX_REGISTER_NUMBER	625	DECL_SAVED_TREE	193
DBX_REGPARM_STABS_CODE	627	DECL_SIZE	168
DBX_REGPARM_STABS_LETTER	627	DECL_STATIC_FUNCTION_P	202
DBX_STATIC_CONST_VAR_CODE	627	DECL_STMT	203
DBX_STATIC_STAB_DATA_SECTION	627	DECL_STMT_DECL	203
DBX_TYPE_DECL_STABS_CODE	627	DECL_THUNK_P	203
DBX_USE_BINCL	627	DECL_VIRTUAL_P	195
DCmode	274	DECL_VOLATILE_MEMFUNC_P	202
DDmode	272	declaration	
De Morgan's law	437	declarations, RTL	297
dead_or_set_p	448	DECLARE_LIBRARY_RENAMES	
debug_expr	304	default	677
debug_implicit_ptr		default_file_start	
debug_insn		DEFAULT_GDB_EXTENSIONS	626
debug_marker		DEFAULT_INCOMING_FRAME_SP_OFFSET	
debug_parameter_ref		DEFAULT_PCC_STRUCT_RETURN	
DEBUG_EXPR_DECL		DEFAULT_SIGNED_CHAR	
DEBUG_SYMS_TEXT		define_address_constraint	
DEBUGGER_ARG_OFFSET		define_asm_attributes	
DEBUGGER_AUTO_OFFSET		define_attr	
decimal float library		define_automaton	
DECL_ALIGN		define_bypass	
DECL_ANTICIPATED		define_c_enum	
DECL_ARGUMENTS		define_code_attr	
DECL_ARRAY_DELETE_OPERATOR_P		define_code_iterator	
DECL_ARTIFICIAL 168, 193,		define_cond_exec	
DECL_ASSEMBLER_NAME		define_constants	
DECL_ATTRIBUTES		define_constraint	
DECL_BASE_CONSTRUCTOR_P		define_cpu_unit	
DECL_COMPLETE_CONSTRUCTOR_P		define_delay	
DECL_COMPLETE_DESTRUCTOR_P		define_enum	
DECL_CONST_MEMFUNC_P		define_enum_attr	
DECL_CONSTRUCTOR_P		define_expand	
DECL_CONTEXT		define_insn	
DECL_CONV_FN_P		define_insn example	
DECL_COPY_CONSTRUCTOR_P		define_insn_and_rewrite	
DECL_DESTRUCTOR_P		define_insn_and_split	
DECL_EXTERN_C_FUNCTION_P		define_insn_reservation	
DECL_EXTERNAL		define_int_attr	
DECL_FUNCTION_MEMBER_P		define_int_iterator	475
DECL_FUNCTION_SPECIFIC_OPTIMIZATION		define_memory_constraint	
195	,	define_mode_attr	
DECL_FUNCTION_SPECIFIC_TARGET 193,	195	define_mode_iterator	
DECL_GLOBAL_CTOR_P		define_peephole	
DECL_GLOBAL_DTOR_P		define_peephole2	
DECL_INITIAL 168,		define_predicate	
DECL_LINKONCE_P		define_query_cpu_unit	
DECL_LOCAL_FUNCTION_P		define_register_constraint	
DECL_MAIN_P		define_reservation	
DECL_NAME		define_special_memory_constraint	
DECL_NAMESPACE_ALIAS		define_special_predicate	
DECL NAMESPACE STD P		define split	

$define\_subst$		DWARF_ALT_FRAME_RETURN_COLUMN	
define_subst_attr		DWARF_CIE_DATA_ALIGNMENT	622
defining attributes and their values 4		DWARF_FRAME_REGISTERS	532
defining constraints	388	DWARF_FRAME_REGNUM	532
defining jump instruction patterns	433	DWARF_LAZY_REGISTER_VALUE	533
defining looping instruction patterns 4	434	DWARF_REG_TO_UNWIND_COLUMN	532
defining peephole optimizers	446	DWARF_ZERO_REG	525
defining predicates		DWARF2_ASM_LINE_DEBUG_INFO	629
defining RTL sequences for code generation 4		DWARF2_ASM_VIEW_DEBUG_INFO	
delay slots, defining		DWARF2_DEBUGGING_INFO	
deletable(		DWARF2_FRAME_INFO	
DELETE_IF_ORDINARY		DWARF2_FRAME_REG_OUT	
Dependent Patterns		DWARF2_UNWIND_INFO	
desc		DYNAMIC_CHAIN_ADDRESS	
descriptors for nested functions		DINAIIIO_OHAIN_ADDINEDO	024
destructors, output of			
deterministic finite state automaton 460, 4		$\mathbf{E}$	
DFmode	104		050
		'E' in constraint	
diagnostics guidelines, fix-it hints		earlyclobber operand	
diagnostics, actionable		edge	
diagnostics, false positive		edge in the flow graph	
diagnostics, guidelines for		edge iterators	
diagnostics, locations		edge splitting	
diagnostics, true positive		EDGE_ABNORMAL	
digits in constraint		EDGE_ABNORMAL, EDGE_ABNORMAL_CALL	321
DImode		EDGE_ABNORMAL, EDGE_EH	
DIR_SEPARATOR6		EDGE_ABNORMAL, EDGE_SIBCALL	
DIR_SEPARATOR_2		EDGE_FALLTHRU, force_nonfallthru	320
directory options .md		EDOM, implicit usage	561
disabling certain registers		eh_return instruction pattern	424
dispatch table	319	EH_FRAME_SECTION_NAME	621
div 2	290	EH_FRAME_THROUGH_COLLECT2	621
div and attributes	454	EH_RETURN_DATA_REGNO	526
division	290	EH_RETURN_HANDLER_RTX	527
divm3 instruction pattern	399	EH_RETURN_STACKADJ_RTX	527
divmodm4 instruction pattern 4	405	EH_TABLES_CAN_BE_READ_ONLY	621
DO_BODY	203	EH_USES	550
DO_COND		ei_edge	
DO_STMT	203	ei_end_p	319
dollar sign 3	355	ei_last	319
DOLLARS_IN_IDENTIFIERS 6	650	ei_next	319
doloop_begin instruction pattern 4		ei_one_before_end_p	319
doloop_end instruction pattern 4	421	ei_prev	319
DONE 439, 441, 4		ei_safe_safe	319
DONT_USE_BUILTIN_SETJMP 6	522	ei_start	319
DOUBLE_TYPE_SIZE	501	ELIMINABLE_REGS	533
DQmode	273	ELSE_CLAUSE	203
driver	480	Embedded C	22
DRIVER_SELF_SPECS 4	480	Empty Statements	189
dump examples	145	EMPTY_CLASS_EXPR	203
dump setup 1		EMPTY_FIELD_BOUNDARY	
dump types		Emulated TLS	
dump verbosity		enabled	387
dump_basic_block1		ENDFILE_SPEC	482
dump_generic_expr 1		endianness	
dump_gimple_stmt1		entry_value	304
dump_printf 1	144	ENTRY_BLOCK_PTR, EXIT_BLOCK_PTR	317
DUMPFILE FORMAT		enum reg class	513

ENUMERAL_TYPE		final_absence_set	
enumerations		final_presence_set	463
epilogue		final_sequence	
epilogue instruction pattern		FINAL_PRESCAN_INSN	
EPILOGUE_USES		FIND_BASE_TERM	
eq		FINI_ARRAY_SECTION_ASM_OP	
eq and attributes		FINI_SECTION_ASM_OP	
eq_attr		finite state automaton minimization	
EQ_EXPR		FIRST_PARM_OFFSET	
equal		FIRST_PARM_OFFSET and virtual registers	
errno, implicit usage		FIRST_PSEUDO_REGISTER	
EXACT_DIV_EXPRexamining SSA_NAMEs		FIRST_STACK_REG	
9		FIRST_VIRTUAL_REGISTER	
exception handling		fixfix-it hints	
exclamation point		fix_truncmn2 instruction pattern	
exclusion_set		FIX_TRUNC_EXPR	
exclusive-or, bitwise			
EXIT_EXPR		fixed register fixed-point fractional library	
EXIT_IGNORE_STACK		fixed_regs	
exp10m2 instruction pattern		fixed_regsfixed_size_mode	
exp2m2 instruction pattern		FIXED_CONVERT_EXPR	
expander definitions		FIXED_CST	
expm1m2 instruction pattern			
expm2 instruction pattern		FIXED_POINT_TYPE	
expr_list		fixmn2 instruction pattern	
EXPR_FILENAME		fixums_truncmn2 instruction pattern	
EXPR_LINENO			
EXPR_STMT		fixunsmn2 instruction pattern	
EXPR_STMT_EXPR		float	
expression		float_extend	
expression codes			
extendmn2 instruction pattern		float_truncate	
extensible constraints		FLOAT_LIB_COMPARE_RETURNS_BOOL	
EXTRA_SPECS		FLOAT_STORE_FLAG_VALUE	
extract_last_m instruction pattern		FLOAT_TYPE_SIZE	
extv instruction pattern		FLOAT_WORDS_BIG_ENDIAN	
extvm instruction pattern		FLOAT_WORDS_BIG_ENDIAN, (lack of) effect on	491
extvmisalignm instruction pattern		subreg	285
extzv instruction pattern		floating point and cross compilation	
extzvm instruction pattern		floatmn2 instruction pattern	
extzvmisalignm instruction pattern		floatunsmn2 instruction pattern	
S as a f		FLOOR_DIV_EXPR	
-		FLOOR_MOD_EXPR.	
$\mathbf{F}$		floorm2 instruction pattern	
'F' in constraint	352	flow-insensitive alias analysis	
FAIL		flow-sensitive alias analysis	
fall-thru		fma	
false positive		fmam4 instruction pattern	
FATAL_EXIT_CODE		fmaxm3 instruction pattern	
FDL, GNU Free Documentation License		fminm3 instruction pattern	
features, optional, in system conventions		fmodm3 instruction pattern	
ffs		fmsm4 instruction pattern	
ffsm2 instruction pattern		fnmam4 instruction pattern	
FIELD_DECL		fnmsm4 instruction pattern	
file_end_indicate_exec_stack		fold_extract_last_m instruction pattern	
files and passes of the compiler		fold_left_plus_m instruction pattern	
files, generated		for_user	

FOR_BODY	ge and attributes	454
FOR_COND	GE_EXPR	177
FOR_EXPR	GEN_ERRNO_RTX	561
FOR_INIT_STMT	gencodes	139
FOR_STMT	general_operand	348
force_reg	GENERAL_REGS	512
FORCE_CODE_SECTION_ALIGN 592	generated files	682
fract_convert	generating assembler output	344
FRACT_TYPE_SIZE	generating insns	339
fractional types	generic predicates	346
fractmn2 instruction pattern 415	GENERIC	161
fractunsmn2 instruction pattern	genflags	139
frame layout 522	get_attr	454
frame_pointer_needed	get_attr_length	458
frame_pointer_rtx 532	get_insns	
frame_related	get_last_insn	305
frame_related, in insn, call_insn, jump_insn,	<pre>get_thread_pointermode instruction pattern</pre>	
barrier, and set		431
frame_related, in mem	GET_CLASS_NARROWEST_MODE	277
frame_related, in reg	GET_CODE	259
frame_related, in symbol_ref 268	GET_MODE	277
FRAME_ADDR_RTX524	GET_MODE_ALIGNMENT	277
FRAME_GROWS_DOWNWARD523	GET_MODE_BITSIZE	277
FRAME_GROWS_DOWNWARD and virtual registers 283	GET_MODE_CLASS	277
FRAME_POINTER_CFA_OFFSET	GET_MODE_FBIT	277
FRAME_POINTER_REGNUM530	GET_MODE_IBIT	277
FRAME_POINTER_REGNUM and virtual registers 283	GET_MODE_MASK	277
frequency, count, BB_FREQ_BASE	GET_MODE_NAME	277
ftruncm2 instruction pattern 414	GET_MODE_NUNITS	277
function	GET_MODE_SIZE	277
function call conventions 7	GET_MODE_UNIT_SIZE	277
function entry and exit 548	GET_MODE_WIDER_MODE	277
function entry point, alternate function entry point	GET_RTX_CLASS	260
321	GET_RTX_FORMAT	262
function properties	GET_RTX_LENGTH	262
function-call insns	geu	293
FUNCTION_ARG_REGNO_P541	geu and attributes	
FUNCTION_BOUNDARY	ggc_collect	
FUNCTION_DECL 193, 201	ĞĞC	
FUNCTION_MODE648	gimple	210
FUNCTION_PROFILER	gimple_addresses_taken	221
FUNCTION_TYPE	gimple_asm_clobber_op	
FUNCTION_VALUE	gimple_asm_input_op	223
FUNCTION_VALUE_REGNO_P545	gimple_asm_nclobbers	
functions, leaf	gimple_asm_ninputs	
fundamental type	gimple_asm_noutputs	
	gimple_asm_output_op	
	gimple_asm_set_clobber_op	
$\mathbf{G}$	gimple_asm_set_input_op	
'g' in constraint	gimple_asm_set_output_op	
'G' in constraint	gimple_asm_set_volatile	
garbage collector, invocation	gimple_asm_string	
garbage collector, troubleshooting	gimple_asm_volatile_p	
gather_loadmn instruction pattern 396	gimple_assign_cast_p	
GCC and portability	gimple_assign_lhs	
GCC_DRIVER_HOST_INITIALIZATION 665	gimple_assign_lhs_ptr	
gcov_type322	gimple_assign_rhs_class	
ge 293		224

gimple_assign_rhs1	225	gimple_cond_make_false	229
gimple_assign_rhs1_ptr		gimple_cond_make_true	
gimple_assign_rhs2	225	gimple_cond_rhs	228
${\tt gimple\_assign\_rhs2\_ptr}$	225	gimple_cond_set_code	228
gimple_assign_rhs3		<pre>gimple_cond_set_false_label</pre>	229
${\tt gimple\_assign\_rhs3\_ptr}$		gimple_cond_set_lhs	228
${\tt gimple\_assign\_set\_lhs}$	225	gimple_cond_set_rhs	
gimple_assign_set_rhs1	225	<pre>gimple_cond_set_true_label</pre>	228
${\tt gimple\_assign\_set\_rhs2}$		gimple_cond_true_label	
${\tt gimple\_assign\_set\_rhs3}$		gimple_convert	
${\tt gimple\_bb}$	220	gimple_copy	
gimple_bind_add_seq	226	<pre>gimple_debug_bind_get_value</pre>	
gimple_bind_add_stmt	226	<pre>gimple_debug_bind_get_value_ptr</pre>	230
gimple_bind_append_vars		<pre>gimple_debug_bind_get_var</pre>	229
gimple_bind_block	226	<pre>gimple_debug_bind_has_value_p</pre>	230
gimple_bind_body	225	gimple_debug_bind_p	220
gimple_bind_set_block	226	<pre>gimple_debug_bind_reset_value</pre>	230
gimple_bind_set_body	226	<pre>gimple_debug_bind_set_value</pre>	230
gimple_bind_set_vars	225	<pre>gimple_debug_bind_set_var</pre>	230
gimple_bind_vars	225	gimple_def_ops	221
gimple_block	220	gimple_eh_filter_failure	230
gimple_build 701,		gimple_eh_filter_set_failure	231
<pre>gimple_build_debug_begin_stmt</pre>	230	<pre>gimple_eh_filter_set_types</pre>	231
<pre>gimple_build_debug_inline_entry</pre>	230	gimple_eh_filter_types	230
gimple_build_nop	231	gimple_eh_filter_types_ptr	230
gimple_build_omp_master	234	gimple_eh_must_not_throw_fndecl	231
gimple_build_omp_ordered	235	<pre>gimple_eh_must_not_throw_set_fndecl</pre>	231
gimple_build_omp_return	236	gimple_expr_code	220
gimple_build_omp_section	236	gimple_expr_type	220
gimple_build_omp_sections_switch	236	gimple_goto_dest	231
gimple_build_wce	240	gimple_goto_set_dest	
gimple_call_arg		gimple_has_mem_ops	
gimple_call_arg_ptr		gimple_has_ops	
gimple_call_chain	227	gimple_has_volatile_ops	222
gimple_call_copy_skip_args		gimple_label_label	
gimple_call_fn		gimple_label_set_label	
gimple_call_fndecl		gimple_loaded_syms	
gimple_call_lhs	226	gimple_locus	
gimple_call_lhs_ptr		gimple_locus_empty_p	
gimple_call_noreturn_p		gimple_modified_p	
gimple_call_num_args	227	gimple_no_warning_p	
gimple_call_return_type		gimple_nop_p	
gimple_call_set_arg		gimple_num_ops	
gimple_call_set_chain		gimple_omp_atomic_load_lhs	
gimple_call_set_fn	226	gimple_omp_atomic_load_rhs	
gimple_call_set_fndecl		gimple_omp_atomic_load_set_lhs	
gimple_call_set_lhs	226	gimple_omp_atomic_load_set_rhs	
gimple_call_set_tail		gimple_omp_atomic_store_set_val	
gimple_call_tail_p	227	gimple_omp_atomic_store_val	232
gimple_catch_handler		gimple_omp_body	
gimple_catch_set_handler			
gimple_catch_set_types		gimple_omp_continue_control_def_ptr	
gimple_catch_types		gimple_omp_continue_control_use	
gimple_catch_types_ptr		gimple_omp_continue_control_use_ptr	
gimple_code		gimple_omp_continue_set_control_def	
gimple_cond_code		gimple_omp_continue_set_control_use	
gimple_cond_false_label		gimple_omp_critical_name	
gimple cond lhs		gimple_omp_critical_name_ptr	

gimple_omp_critical_set_name	233	gimple_resx_region	238
gimple_omp_for_clauses	233	gimple_resx_set_region	
gimple_omp_for_clauses_ptr		gimple_return_retval	
gimple_omp_for_cond		gimple_return_set_retval	238
gimple_omp_for_final		gimple_seq_add_seq	
gimple_omp_for_final_ptr		gimple_seq_add_stmt	
gimple_omp_for_incr		gimple_seq_alloc	
gimple_omp_for_incr_ptr		gimple_seq_copy	
gimple_omp_for_index		gimple_seq_deep_copy	
gimple_omp_for_index_ptr		gimple_seq_empty_p	
gimple_omp_for_initial		gimple_seq_first	
gimple_omp_for_initial_ptr		gimple_seq_init	
gimple_omp_for_pre_body		gimple_seq_last	
gimple_omp_for_set_clauses		gimple_seq_reverse	
gimple_omp_for_set_cond		gimple_seq_set_first	
gimple_omp_for_set_final		gimple_seq_set_last	
gimple_omp_for_set_incr		gimple_seq_singleton_p	
gimple_omp_for_set_index		gimple_set_block	
gimple_omp_for_set_initial		gimple_set_def_ops	
gimple_omp_for_set_pre_body		gimple_set_has_volatile_ops	
gimple_omp_parallel_child_fn		gimple_set_locus	
gimple_omp_parallel_child_fn_ptr		gimple_set_op	
gimple_omp_parallel_clauses		gimple_set_plf	
gimple_omp_parallel_clauses_ptr		gimple_set_use_ops	
gimple_omp_parallel_combined_p		gimple_set_vdef_ops	
gimple_omp_parallel_data_arg		gimple_set_visited	
gimple_omp_parallel_data_arg_ptr		gimple_set_vuse_ops	
gimple_omp_parallel_set_child_fn		gimple_simplify	
gimple_omp_parallel_set_clauses		gimple_statement_with_ops	
gimple_omp_parallel_set_combined_p		gimple_stored_syms	
gimple_omp_parallel_set_data_arg		gimple_switch_default_label	
gimple_omp_return_nowait_p		gimple_switch_index	
gimple_omp_return_set_nowait		gimple_switch_label	
gimple_omp_section_last_p		gimple_switch_num_labels	
gimple_omp_section_set_last		gimple_switch_set_default_label	
gimple_omp_sections_clauses		gimple_switch_set_index	
gimple_omp_sections_clauses_ptr		gimple_switch_set_label	
gimple_omp_sections_control		gimple_switch_set_num_labels	
gimple_omp_sections_control_ptr		gimple_try_catch_is_cleanup	
gimple_omp_sections_set_clauses		gimple_try_cleanup	
gimple_omp_sections_set_control		gimple_try_eval	
gimple_omp_set_body		gimple_try_kind	
gimple_omp_single_clauses		gimple_try_set_catch_is_cleanup	
gimple_omp_single_clauses_ptr		gimple_try_set_cleanup	
gimple_omp_single_set_clauses		gimple_try_set_eval	
gimple_op		gimple_use_ops	
gimple_op_ptr		gimple_vdef_ops	
gimple_ops		gimple_visited_p	
gimple_phi_arg		gimple_vuse_ops	
gimple_phi_arg_def		gimple_wce_cleanup	
gimple_phi_arg_edge		gimple_wce_cleanup_eh_only	
gimple_phi_capacity		gimple_wce_set_cleanup	
gimple_phi_num_args 237,		gimple_wce_set_cleanup_eh_only	
gimple_phi_result		GIMPLE	
gimple_phi_result_ptr		GIMPLE API	
gimple_phi_set_arg		GIMPLE class hierarchy	
gimple_phi_set_result		GIMPLE Exception Handling	
gimple plf		GIMPLE instruction set	

GIMPLE sequences		gsi_link_after	
GIMPLE statement iterators 318,		gsi_link_before	
GIMPLE_ASM		gsi_link_seq_after	
GIMPLE_ASSIGN		gsi_link_seq_before	
GIMPLE_BIND		gsi_move_after	
GIMPLE_CALL		gsi_move_before	
GIMPLE_CATCH		gsi_move_to_bb_end	
GIMPLE_COND		gsi_next	
GIMPLE_DEBUG		gsi_one_before_end_p	
GIMPLE_DEBUG_BEGIN_STMT		gsi_prev	
GIMPLE_DEBUG_BIND		gsi_remove	
GIMPLE_DEBUG_INLINE_ENTRY		gsi_replace	
GIMPLE_EH_FILTER		gsi_seq	
GIMPLE_GOTO		gsi_split_seq_after	
GIMPLE_LABEL		<pre>gsi_split_seq_before</pre>	
GIMPLE_NOP		gsi_start	
GIMPLE_OMP_ATOMIC_LOAD	231	gsi_start_bb	
GIMPLE_OMP_ATOMIC_STORE		gsi_stmt	
GIMPLE_OMP_CONTINUE	232	gsi_stmt_ptr	
GIMPLE_OMP_CRITICAL	233	gt	293
GIMPLE_OMP_FOR	233	gt and attributes	454
GIMPLE_OMP_MASTER	234	GT_EXPR	
GIMPLE_OMP_ORDERED	235	gtu	
GIMPLE_OMP_PARALLEL	235	gtu and attributes	454
GIMPLE_OMP_RETURN	236	GTY	
GIMPLE_OMP_SECTION	236	guidelines for diagnostics	
GIMPLE_OMP_SECTIONS	236	guidelines for options	
GIMPLE_OMP_SINGLE	237	guidelines, user experience	715
GIMPLE_PHI	237		
GIMPLE_RESX	238	тт	
GIMPLE_RESX		Н	
	238	H 'H' in constraint	352
GIMPLE_RETURN	238 238		
GIMPLE_RETURN	238 238 239	'H' in constraint	273
GIMPLE_RETURN	238 238 239 240	'H' in constraint	273 649
GIMPLE_RETURN	238 238 239 240 128	'H' in constraint	273 649 203
GIMPLE_RETURN	238 238 239 240 128 127	'H' in constraint	273 649 203 203
GIMPLE_RETURN	238 238 239 240 128 127 224	'H' in constraint	273 649 203 203 203
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification 127, gimplifier  gimplify_assign	238 238 239 240 128 127 224 128	'H' in constraint	273 649 203 203 203 282
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification 127, gimplifier  gimplify_assign  gimplify_expr	238 238 239 240 128 127 224 128 128	'H' in constraint	273 649 203 203 203 282 531
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification. 127,  gimplify_assign.  gimplify_expr  gimplify_function_tree  global_regs	238 238 239 240 128 127 224 128 128 506	'H' in constraint	273 649 203 203 203 282 531 531
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification 127, gimplify_assign gimplify_expr gimplify_function_tree	238 238 239 240 128 127 224 128 128 506 203	'H' in constraint	273 649 203 203 282 531 531 530
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563	'H' in constraint	273 649 203 203 282 531 531 530 548
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293	'H' in constraint	273 649 203 203 282 531 531 530 548 508
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242	'H' in constraint	273 649 203 203 282 531 530 548 508 509
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 242	'H' in constraint  HAmode  HANDLE_PRAGMA_PACK_WITH_EXPANSION  HANDLER  HANDLER_BODY  HANDLER_PARMS  hard registers  HARD_FRAME_POINTER_IS_ARG_POINTER  HARD_FRAME_POINTER_IS_FRAME_POINTER  HARD_FRAME_POINTER_REGNUM  HARD_REGNO_CALLER_SAVE_MODE  HARD_REGNO_NREGS_HAS_PADDING  HARD_REGNO_NREGS_WITH_PADDING  HARD_REGNO_RENAME_OK	273 649 203 203 282 531 530 548 508 509 510
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 242 324	'H' in constraint  HAmode  HANDLE_PRAGMA_PACK_WITH_EXPANSION  HANDLER  HANDLER_BODY  HANDLER_PARMS  hard registers  HARD_FRAME_POINTER_IS_ARG_POINTER  HARD_FRAME_POINTER_IS_FRAME_POINTER  HARD_FRAME_POINTER_REGNUM  HARD_REGNO_CALLER_SAVE_MODE  HARD_REGNO_NREGS_HAS_PADDING  HARD_REGNO_NREGS_WITH_PADDING  HARD_REGNO_RENAME_OK  HAS_INIT_SECTION	273 649 203 203 282 531 530 548 508 509 510 614
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 242 324 244	'H' in constraint	273 649 203 203 203 282 531 530 548 508 510 614 642
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 1128 127 224 128 128 506 203 563 293 242 242 324 244 324	'H' in constraint.  HAmode.  HANDLE_PRAGMA_PACK_WITH_EXPANSION.  HANDLER.  HANDLER_BODY.  HANDLER_PARMS.  hard registers.  HARD_FRAME_POINTER_IS_ARG_POINTER.  HARD_FRAME_POINTER_IS_FRAME_POINTER.  HARD_FRAME_POINTER_REGNUM.  HARD_REGNO_CALLER_SAVE_MODE.  HARD_REGNO_NREGS_HAS_PADDING.  HARD_REGNO_NREGS_WITH_PADDING.  HARD_REGNO_RENAME_OK.  HAS_INIT_SECTION.  HAS_LONG_COND_BRANCH.  HAS_LONG_UNCOND_BRANCH.	273 649 203 203 203 282 531 530 548 509 510 614 642 642
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 242 324 244 324 244	'H' in constraint	273 649 203 203 282 531 530 548 508 510 614 642 642 664
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 242 244 324 244 324	'H' in constraint.  HAmode.  HANDLE_PRAGMA_PACK_WITH_EXPANSION.  HANDLER.  HANDLER_BODY.  HANDLER_PARMS.  hard registers.  HARD_FRAME_POINTER_IS_ARG_POINTER.  HARD_FRAME_POINTER_IS_FRAME_POINTER.  HARD_FRAME_POINTER_REGNUM.  HARD_REGNO_CALLER_SAVE_MODE.  HARD_REGNO_NREGS_HAS_PADDING.  HARD_REGNO_NREGS_WITH_PADDING.  HARD_REGNO_RENAME_OK.  HAS_INIT_SECTION.  HAS_LONG_COND_BRANCH.  HAS_LONG_UNCOND_BRANCH.  HAVE_DOS_BASED_FILE_SYSTEM.	273 649 203 203 282 531 530 548 508 510 614 642 664 562
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 242 244 324 244 324 324 324	'H' in constraint  HAmode  HANDLE_PRAGMA_PACK_WITH_EXPANSION  HANDLER  HANDLER_BODY  HANDLER_PARMS  hard registers  HARD_FRAME_POINTER_IS_ARG_POINTER  HARD_FRAME_POINTER_IS_FRAME_POINTER  HARD_FRAME_POINTER_REGNUM  HARD_REGNO_CALLER_SAVE_MODE  HARD_REGNO_NREGS_HAS_PADDING  HARD_REGNO_NREGS_WITH_PADDING  HARD_REGNO_RENAME_OK  HAS_INIT_SECTION  HAS_LONG_COND_BRANCH  HAS_LONG_UNCOND_BRANCH  HAVE_POST_DECREMENT  HAVE_POST_DECREMENT	273 649 203 203 282 531 530 548 509 510 614 642 664 562 562
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 242 244 324 244 324 324 324 324 32	'H' in constraint  HAmode  HANDLE_PRAGMA_PACK_WITH_EXPANSION  HANDLER  HANDLER_BODY  HANDLER_PARMS  hard registers.  HARD_FRAME_POINTER_IS_ARG_POINTER  HARD_FRAME_POINTER_IS_FRAME_POINTER  HARD_FRAME_POINTER_REGNUM  HARD_REGNO_CALLER_SAVE_MODE  HARD_REGNO_NREGS_HAS_PADDING  HARD_REGNO_NREGS_WITH_PADDING  HARD_REGNO_RENAME_OK  HAS_INIT_SECTION  HAS_LONG_COND_BRANCH  HAS_LONG_UNCOND_BRANCH  HAVE_DOS_BASED_FILE_SYSTEM  HAVE_POST_DECREMENT	273 649 203 203 282 531 530 548 509 510 614 642 664 562 562
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 242 324 244 324 244 324 324 324 32	'H' in constraint  HAmode  HANDLE_PRAGMA_PACK_WITH_EXPANSION  HANDLER  HANDLER_BODY  HANDLER_PARMS  hard registers  HARD_FRAME_POINTER_IS_ARG_POINTER  HARD_FRAME_POINTER_IS_FRAME_POINTER  HARD_FRAME_POINTER_REGNUM  HARD_REGNO_CALLER_SAVE_MODE  HARD_REGNO_NREGS_HAS_PADDING  HARD_REGNO_NREGS_WITH_PADDING  HARD_REGNO_RENAME_OK  HAS_INIT_SECTION  HAS_LONG_COND_BRANCH  HAS_LONG_UNCOND_BRANCH  HAVE_POST_DECREMENT  HAVE_POST_DECREMENT  HAVE_POST_INCREMENT  HAVE_POST_MODIFY_DISP	273 649 203 203 203 282 531 530 548 509 614 642 664 562 562 562 562
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 244 244 324 244 324 324 324 244 24	'H' in constraint  HAmode  HANDLE_PRAGMA_PACK_WITH_EXPANSION  HANDLER  HANDLER_BODY  HANDLER_PARMS  hard registers.  HARD_FRAME_POINTER_IS_ARG_POINTER  HARD_FRAME_POINTER_IS_FRAME_POINTER  HARD_FRAME_POINTER_REGNUM  HARD_REGNO_CALLER_SAVE_MODE  HARD_REGNO_NREGS_HAS_PADDING  HARD_REGNO_NREGS_WITH_PADDING  HARD_REGNO_RENAME_OK  HAS_INIT_SECTION  HAS_LONG_COND_BRANCH  HAS_LONG_UNCOND_BRANCH  HAVE_POST_DECREMENT  HAVE_POST_DECREMENT  HAVE_POST_INCREMENT  HAVE_POST_MODIFY_DISP  HAVE_POST_MODIFY_REG  HAVE_PRE_DECREMENT	273 649 203 203 203 282 531 530 548 509 510 614 642 562 562 562 562 562 562
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 244 244 324 244 324 324 324 244 24	'H' in constraint  HAmode  HANDLE_PRAGMA_PACK_WITH_EXPANSION  HANDLER  HANDLER_BODY  HANDLER_PARMS  hard registers  HARD_FRAME_POINTER_IS_ARG_POINTER  HARD_FRAME_POINTER_IS_FRAME_POINTER  HARD_FRAME_POINTER_REGNUM  HARD_REGNO_CALLER_SAVE_MODE  HARD_REGNO_NREGS_HAS_PADDING  HARD_REGNO_NREGS_WITH_PADDING  HARD_REGNO_RENAME_OK  HAS_INIT_SECTION  HAS_LONG_COND_BRANCH  HAS_LONG_UNCOND_BRANCH  HAVE_DOS_BASED_FILE_SYSTEM  HAVE_POST_DECREMENT  HAVE_POST_INCREMENT  HAVE_POST_MODIFY_DISP  HAVE_POST_MODIFY_REG  HAVE_PRE_DECREMENT  HAVE_PRE_DECREMENT  HAVE_PRE_INCREMENT  HAVE_PRE_DECREMENT  HAVE_PRE_INCREMENT  HAVE_PRE_INCREMENT	273 649 203 203 203 282 531 530 548 509 614 642 562 562 562 562 562 562
GIMPLE_RETURN  GIMPLE_SWITCH  GIMPLE_TRY  GIMPLE_WITH_CLEANUP_EXPR  gimplification	238 238 239 240 128 127 224 128 128 506 203 563 293 242 244 244 324 324 324 324 324 244 24	'H' in constraint  HAmode  HANDLE_PRAGMA_PACK_WITH_EXPANSION  HANDLER  HANDLER_BODY  HANDLER_PARMS  hard registers.  HARD_FRAME_POINTER_IS_ARG_POINTER  HARD_FRAME_POINTER_IS_FRAME_POINTER  HARD_FRAME_POINTER_REGNUM  HARD_REGNO_CALLER_SAVE_MODE  HARD_REGNO_NREGS_HAS_PADDING  HARD_REGNO_NREGS_WITH_PADDING  HARD_REGNO_RENAME_OK  HAS_INIT_SECTION  HAS_LONG_COND_BRANCH  HAS_LONG_UNCOND_BRANCH  HAVE_POST_DECREMENT  HAVE_POST_DECREMENT  HAVE_POST_INCREMENT  HAVE_POST_MODIFY_DISP  HAVE_POST_MODIFY_REG  HAVE_PRE_DECREMENT	273 649 203 203 203 282 531 530 548 509 510 614 562 562 562 562 562 562 562 562

HFmode	272	inclusive on hituriae	201
		inclusive-or, bitwise	
high		INCOMING_FRAME_SP_OFFSET	
HImode		INCOMING_REG_PARM_STACK_SPACE	
HImode, in insn		INCOMING_REGNO	
HONOR_REG_ALLOC_ORDER		INCOMING_RETURN_ADDR_RTX	
host configuration		INCOMING_STACK_BOUNDARY	
host functions		INDEX_REG_CLASS	
host hooks		${\tt indirect\_jump} \ instruction \ pattern$	
host makefile fragment		indirect_operand	
HOST_BIT_BUCKET		INDIRECT_REF	
HOST_EXECUTABLE_SUFFIX		init_machine_status	490
HOST_HOOKS_EXTRA_SIGNALS	663	$\verb init_one_libfunc $	560
HOST_HOOKS_GT_PCH_ALLOC_GRANULARITY	663	INIT_ARRAY_SECTION_ASM_OP	592
HOST_HOOKS_GT_PCH_GET_ADDRESS	663	INIT_CUMULATIVE_ARGS	539
HOST_HOOKS_GT_PCH_USE_ADDRESS	663	<pre>INIT_CUMULATIVE_INCOMING_ARGS</pre>	539
HOST_LACKS_INODE_NUMBERS	665	INIT_CUMULATIVE_LIBCALL_ARGS	539
HOST_LONG_FORMAT	666	INIT_ENVIRONMENT	
HOST_LONG_FORMAT	666	INIT_EXPANDERS	490
HOST_OBJECT_SUFFIX		INIT_EXPR	177
HOST_PTR_PRINTF		INIT_SECTION_ASM_OP 591,	
HOT_TEXT_SECTION_NAME		INITIAL_ELIMINATION_OFFSET	
HQmode		INITIAL_FRAME_ADDRESS_RTX	
		initialization routines	
		inlining	
I		insert_insn_on_edge	
(2) inint	951	insn	
'i' in constraint		insn and '/f'	
'I' in constraint			
identifier		insn and '/j'	
IDENTIFIER_LENGTH		insn and '/s'	
IDENTIFIER_NODE		insn and '/u'	
IDENTIFIER_OPNAME_P		insn and '/v'	
IDENTIFIER_POINTER		insn attributes	
IDENTIFIER_TYPENAME_P		insn canonicalization	
IEEE 754-2008		insn includes	
if_then_else		insn lengths, computing	
if_then_else and attributes		insn notes, notes	
if_then_else usage		insn splitting	
IF_COND		insn-attr.h	
IF_STMT		insn_list	
IFCVT_MACHDEP_INIT		INSN_ANNULLED_BRANCH_P	
IFCVT_MODIFY_CANCEL		INSN_CODE	
IFCVT_MODIFY_FINAL		INSN_DELETED_P	
IFCVT_MODIFY_INSN		INSN_FROM_TARGET_P	
IFCVT_MODIFY_MULTIPLE_TESTS		INSN_REFERENCES_ARE_DELAYED	
IFCVT_MODIFY_TESTS		INSN_SETS_ARE_DELAYED	
IMAGPART_EXPR		INSN_UID	
Immediate Uses	_	INSN_VAR_LOCATION	
<pre>immediate_operand</pre>		insns	
IMMEDIATE_PREFIX		, 0	
in_struct		insns, recognizing	
<pre>in_struct, in code_label and note</pre>		instruction attributes	
in_struct, in insn and jump_insn and call_in		instruction latency time	
		1	
in_struct, in insn, call_insn, jump_insn and		instruction splitting	
jump_table_data		insv instruction pattern	
in_struct, in subreg		insvm instruction pattern	
include		<pre>insvmisalignm instruction pattern</pre>	
INCLUDE_DEFAULTS	485	int iterators in '.md' files	475

INT_FAST16_TYPE	jump, in mem	266
INT_FAST32_TYPE	jump_insn	305
INT_FAST64_TYPE	jump_insn and '/f'	
INT_FAST8_TYPE	jump_insn and '/j'	266
INT_LEAST16_TYPE	jump_insn and '/s'	268
INT_LEAST32_TYPE	jump_insn and '/u'	266
INT_LEAST64_TYPE	jump_insn and '/v'	266
INT_LEAST8_TYPE	jump_table_data	307
INT_TYPE_SIZE	jump_table_data and '/s'	268
INT16_TYPE 503	jump_table_data and '/v'	266
INT32_TYPE 503	JUMP_ALIGN	623
INT64_TYPE 503	JUMP_LABEL	
INT8_TYPE	JUMP_TABLES_IN_TEXT_SECTION	
INTEGER_CST	Jumps	
INTEGER_TYPE		
inter-procedural optimization passes 129		
Interdependence of Patterns 433	$\mathbf{L}$	
interfacing to GCC output	1-1-1	001
interlock delays	label_ref	
intermediate representation lowering	label_ref and '/v'	
INTMAX_TYPE	label_ref, RTL sharing	
	LABEL_ALIGN	
INTPTR_TYPE	LABEL_ALIGN_AFTER_BARRIER	
introduction 1	LABEL_ALT_ENTRY_P	
INVOKEmain	LABEL_ALTERNATE_NAME	321
ior	LABEL_DECL	168
ior and attributes	LABEL_KIND	306
ior, canonicalization of	LABEL_NUSES	306
iorm3 instruction pattern	LABEL_PRESERVE_P	266
IPA passes	LABEL_REF_NONLOCAL_P	
IRA_HARD_REGNO_ADD_COST_MULTIPLIER 508	lang_hooks.gimplify_expr	
is_a	lang_hooks.parse_file	
is_gimple_addressable	language-dependent trees	
is_gimple_asm_val	language-independent intermediate representati	
is_gimple_assign	imigaage maependent intermediate representati	
is_gimple_call	large return values	
is_gimple_call_addr	LAST_STACK_REG	
is_gimple_constant	LAST_VIRTUAL_REGISTER	
is_gimple_debug		
is_gimple_ip_invariant	late IPA passes	
is_gimple_ip_invariant_address		
is_gimple_mem_ref_addr	LCSSA	
is_gimple_min_invariant	LD_FINI_SWITCH	
	LD_INIT_SWITCH	
is_gimple_omp	LDD_SUFFIX	
is_gimple_val	ldexpm3 instruction pattern	
IS_ASM_LOGICAL_LINE_SEPARATOR	le	
iterators in '.md' files	le and attributes	454
IV analysis on GIMPLE	LE_EXPR	177
IV analysis on RTL	leaf functions	510
	leaf_function_p	420
т	LEAF_REG_REMAP	
J	LEAF_REGISTERS	
JMP_BUF_SIZE	left rotate	
jump	left shift	
jump instruction pattern	LEGITIMATE_PIC_OPERAND_P	
jump instruction patterns	LEGITIMIZE_RELOAD_ADDRESS	
jump instructions and set	length	
jump, in call_insn	less than	
jump, in carr_insh	less than or equal	
jump, 111 IIISII	ress than of equal	∠93

leu	202	lowering, language-dependent intermediate	
leu and attributes			107
	_	representation	
lfloormn2		lrintmn2	
LIB_SPEC		1roundmn2	
LIB2FUNCS_EXTRA		LSHIFT_EXPR	
LIBCALL_VALUE		lshiftrt	
'libgcc.a'		lshiftrt and attributes	
LIBGCC_SPEC		lshrm3 instruction pattern	
LIBGCC2_CFLAGS		lt	
LIBGCC2_GNU_PREFIX	501	1t and attributes	454
LIBGCC2_UNWIND_ATTRIBUTE	659	LT_EXPR	177
library subroutine names	560	LTGT_EXPR	177
LIBRARY_PATH_ENV	651	lto	693
LIMIT_RELOAD_CLASS	516	ltrans	693
LINK_COMMAND_SPEC		ltu	293
LINK_EH_SPEC			
LINK_GCC_C_SEQUENCE_SPEC		T. 4	
LINK_LIBGCC_SPECIAL_1		$\mathbf{M}$	
LINK_SPEC		'm' in constraint	351
list		MACH_DEP_SECTION_ASM_FLAG	
Liveness representation		machine attributes	
lo_sum		machine description macros	
load address instruction		machine description macros	
load_multiple instruction pattern		machine mode conversions	
LOAD_EXTEND_OP		machine mode wrapper classes	
Local Register Allocator (LRA)		machine modes	
LOCAL_ALIGNMENT		machine specific constraints	
LOCAL_CLASS_P		machine-independent predicates	
LOCAL_DECL_ALIGNMENT		machine_mode	
LOCAL_INCLUDE_DIR		macros, target description	
LOCAL_LABEL_PREFIX		maddmn4 instruction pattern	
LOCAL_REGNO		make_safe_from	
location information		MAKE_DECL_ONE_ONLY	
log10m2 instruction pattern		makefile fragment	
log1pm2 instruction pattern		makefile targets	
log2m2 instruction pattern		MALLOC_ABI_ALIGNMENT	
LOG_LINKS		Manipulating GIMPLE statements	
logbm2 instruction pattern		marking roots	
Logical Operators		<pre>mask_fold_left_plus_m instruction pattern</pre>	
${\it logical-and, bitwise} \ldots \ldots \ldots$	290	mask_gather_loadmn instruction pattern	
LOGICAL_OP_NON_SHORT_CIRCUIT		<pre>mask_scatter_storemn instruction pattern</pre>	
logm2 instruction pattern		MASK_RETURN_ADDR	
LONG_ACCUM_TYPE_SIZE		maskloadmn instruction pattern	
LONG_DOUBLE_TYPE_SIZE		maskstoremn instruction pattern	
LONG_FRACT_TYPE_SIZE	501	Match and Simplify	
LONG_LONG_ACCUM_TYPE_SIZE	501	${\tt match\_dup$	
LONG_LONG_FRACT_TYPE_SIZE	501	match_dup and attributes	457
LONG_LONG_TYPE_SIZE	500	match_op_dup	
LONG_TYPE_SIZE	500	match_operand	340
longjmp and automatic variables	7	match_operand and attributes	453
Loop analysis	327	match_operator	341
Loop manipulation	330	match_par_dup	343
Loop querying	329	match_parallel	342
Loop representation	327	match_scratch 340,	
Loop-closed SSA form	330	match_test and attributes	453
LOOP_ALIGN	624	matching constraint	352
LOOP_EXPR	177	matching operands	344
looping instruction patterns	434	math library	

math, in RTL		MINUS_EXPR	
MATH_LIBRARY	651	MIPS coprocessor-definition macros	
matherr		miscellaneous register hooks	
MAX_BITS_PER_WORD	491	mnemonic attribute	459
MAX_BITSIZE_MODE_ANY_INT		$\verb"mod"\ldots$	
MAX_BITSIZE_MODE_ANY_MODE	278	mod and attributes	454
MAX_CONDITIONAL_EXECUTE	651	mode classes	
MAX_FIXED_MODE_SIZE	498	mode iterators in '.md' files	472
MAX_MOVE_MAX	644	mode switching	632
MAX_OFILE_ALIGNMENT	494	MODE_ACCUM	275
MAX_REGS_PER_ADDRESS	562	MODE_BASE_REG_CLASS	514
MAX_STACK_ALIGNMENT	494	MODE_BASE_REG_REG_CLASS	514
maxm3 instruction pattern	400	MODE_CC	574
may_trap_p, tree_could_trap_p	320	MODE_CODE_BASE_REG_CLASS	514
maybe_undef		MODE_COMPLEX_FLOAT	
mcount		MODE_COMPLEX_INT	275
MD_EXEC_PREFIX	484	MODE_DECIMAL_FLOAT	
MD_FALLBACK_FRAME_STATE_FOR		MODE_FLOAT	
MD_HANDLE_UNWABI		MODE_FRACT	
MD_STARTFILE_PREFIX		MODE_INT	
MD_STARTFILE_PREFIX_1		MODE_PARTIAL_INT	
mem		MODE_POINTER_BOUNDS	
mem and '/c'		MODE_RANDOM	
mem and '/f'		MODE_UACCUM	
mem and '/j'		MODE_UFRACT	
mem and '/u'		modifiers in constraints	
mem and '/v'		MODIFY_EXPR	
mem, RTL sharing		modm3 instruction pattern	
mem_thread_fence instruction pattern		modulo scheduling	
MEM_ADDR_SPACE		MOVE_MAX	
MEM_ALIAS_SET		MOVE_MAX_PIECES	
MEM_ALIGN		MOVE_RATIO	
MEM_EXPR		movm instruction pattern	
MEM_KEEP_ALIAS_SET_P		movmemm instruction pattern	
MEM_NOTRAP_P		movmisalignm instruction pattern	
MEM_OFFSET			
		movmodecc instruction pattern	
MEM_OFFSET_KNOWN_P		movstr instruction pattern	
MEM_POINTER		movstrictm instruction pattern	
MEM_READONLY_P		msubmn4 instruction pattern	
MEM_REF		mulhisi3 instruction pattern	
MEM_SIZE		mulm3 instruction pattern	
MEM_SIZE_KNOWN_P		mulqihi3 instruction pattern	
MEM_VOLATILE_P		mulsidi3 instruction pattern	
memory model		mult	
memory reference, nonoffsettable		mult and attributes	
memory references in constraints		mult, canonicalization of	
memory_barrier instruction pattern		MULT_EXPR	
memory_blockage instruction pattern		MULT_HIGHPART_EXPR	
memory_operand		MULTIARCH_DIRNAME	
MEMORY_MOVE_COST		MULTILIB_DEFAULTS	
METHOD_TYPE		MULTILIB_DIRNAMES	
MIN_UNITS_PER_WORD		MULTILIB_EXCEPTIONS	
MINIMUM_ALIGNMENT		MULTILIB_EXTRA_OPTS	
MINIMUM_ATOMIC_ALIGNMENT		MULTILIB_MATCHES	
minm3 instruction pattern		MULTILIB_OPTIONS	
minus		MULTILIB_OSDIRNAMES	
minus and attributes		MULTILIB_REQUIRED	
minus, canonicalization of	436	MULTILIB_REUSE	669

multiple alternative constraints		nop instruction pattern	
MULTIPLE_SYMBOL_SPACES		NOP_EXPR	
multiplication		normal predicates	
multiplication with signed saturation		not	
$\   \text{multiplication with unsigned saturation} \ldots \ldots$		not and attributes	
mulvm4 instruction pattern	399	not equal	
		not, canonicalization of	436
<b>N</b> T		note	
N		note and '/i'	
'n' in constraint	351	note and '/v'	266
N_REG_CLASSES	513	NOTE_INSN_BASIC_BLOCK	318
name	163	NOTE_INSN_BEGIN_STMT	308
named address spaces		NOTE_INSN_BLOCK_BEG	307
named patterns and conditions		NOTE_INSN_BLOCK_END	307
names, pattern		NOTE_INSN_DELETED	307
namespace, scope		NOTE_INSN_DELETED_LABEL	307
NAMESPACE_DECL		NOTE_INSN_EH_REGION_BEG	307
NATIVE_SYSTEM_HEADER_COMPONENT		NOTE_INSN_EH_REGION_END	307
ne		NOTE_INSN_FUNCTION_BEG	307
ne and attributes		NOTE_INSN_INLINE_ENTRY	308
NE_EXPR		NOTE_INSN_VAR_LOCATION	308
nearbyintm2 instruction pattern		NOTE_LINE_NUMBER	
neg		NOTE_SOURCE_FILE	
		NOTE_VAR_LOCATION	
neg and attributes		NOTICE_UPDATE_CC	
neg, canonicalization of		notmodecc instruction pattern	
NEGATE_EXPR		NUM_MACHINE_MODES	
negation		NOTI_TIAOTITNE_TIODED	211
			632
$\begin{array}{cccccccccccccccccccccccccccccccccccc$	289	NUM_MODES_FOR_MODE_SWITCHING	
negation with signed saturationnegation with unsigned saturation	289 289	NUM_MODES_FOR_MODE_SWITCHING	147
negation with signed saturationnegation with unsigned saturationnegm2 instruction pattern	289 289 406	NUM_MODES_FOR_MODE_SWITCHING	147
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern	289 289 406 418	NUM_MODES_FOR_MODE_SWITCHING	147
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern	289 289 406 418 406	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis	147
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for	289 289 406 418 406 558	NUM_MODES_FOR_MODE_SWITCHING	147 332
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr	289 289 406 418 406 558	NUM_MODES_FOR_MODE_SWITCHING	147 332 351
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB	289 289 406 418 406 558 678	NUM_MODES_FOR_MODE_SWITCHING	147 332 351 192
negation with signed saturation	289 289 406 418 406 558 678	NUM_MODES_FOR_MODE_SWITCHING	147 332 351 192 192
negation with signed saturation	289 289 406 418 406 558 678 317 305	NUM_MODES_FOR_MODE_SWITCHING	351 192 192 192
negation with signed saturation	289 289 406 418 406 558 678 317 305 561	NUM_MODES_FOR_MODE_SWITCHING.  NUM_POLY_INT_COEFFS.  Number of iterations analysis.  O  'o' in constraint.  OACC_CACHE.  OACC_DATA.  OACC_DECLARE.  OACC_ENTER_DATA.	351 192 192 192
negation with signed saturation	289 289 406 418 406 558 678 317 305 561 260	NUM_MODES_FOR_MODE_SWITCHING	351 192 192 192
negation with signed saturation	289 289 406 418 406 558 678 317 305 561 260 615	NUM_MODES_FOR_MODE_SWITCHING.  NUM_POLY_INT_COEFFS.  Number of iterations analysis.  O  'o' in constraint.  OACC_CACHE.  OACC_DATA.  OACC_DECLARE.  OACC_ENTER_DATA.	351 192 192 192 192 192
negation with signed saturation	289 289 406 418 406 558 678 317 305 561 260 615	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_HOST_DATA OACC_KERNELS	351 192 192 192 192 192 192 192
negation with signed saturation	289 289 406 418 406 558 678 317 305 561 260 615 628	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_HOST_DATA	351 192 192 192 192 192 192 192
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM	289 289 406 418 406 558 678 317 305 561 260 615 628 628	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_HOST_DATA OACC_KERNELS	351 192 192 192 192 192 192 192 192
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_HOST_DATA OACC_KERNELS OACC_LOOP	351 192 192 192 192 192 192 192 192
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 628	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL	351 192 192 192 192 192 192 192 192 192 19
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 628 605	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL	351 192 192 192 192 192 192 192 192 192 19
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_DOT_IN_LABEL	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 628 605 605	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL	351 192 192 192 192 192 192 192 192 192 19
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 628 605 605 580	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN	351 192 192 192 192 192 192 192 192 192 612 659
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_DOT_IN_LABEL NO_FUNCTION_CSE NO_PROFILE_COUNTERS	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 605 605 580 552	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN OBJCC_FORMAT_COFF	351 192 192 192 192 192 192 192 662 6659 615
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_FUNCTION_CSE	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 628 605 605 580 552 512	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN	351 192 192 192 192 192 192 192 612 659 615 163
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_DOT_IN_LABEL NO_FUNCTION_CSE NO_PROFILE_COUNTERS NO_REGS	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 605 605 580 552 512 177	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN OBJCCT_FORMAT_COFF OFFSET_TYPE	351 192 192 192 192 192 192 192 192 192 19
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_FUNCTION_CSE NO_PROFILE_COUNTERS NO_REGS NON_LVALUE_EXPR nondeterministic finite state automaton	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 628 605 605 580 552 512 177 464	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN OBJCCT_FORMAT_COFF OFFSET_TYPE offsettable address	351 192 192 192 192 192 192 192 192 192 19
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_FUNCTION_CSE NO_PROFILE_COUNTERS NO_REGS NON_LVALUE_EXPR nondeterministic finite state automaton nonimmediate_operand	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 605 605 580 552 512 177 464 348	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN OBJCCT_FORMAT_COFF OFFSET_TYPE offsettable address OImode	351 192 192 192 192 192 192 192 612 659 615 163 351 272 190
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_FUNCTION_CSE NO_PROFILE_COUNTERS NO_REGS NON_LVALUE_EXPR nondeterministic finite state automaton nonimmediate_operand nonlocal goto handler	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 605 605 580 552 512 177 464 348 321	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENITE_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN OBJCCT_FORMAT_COFF OFFSET_TYPE offsettable address OImode OMP_ATOMIC	351 192 192 192 192 192 192 192 192 192 19
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_FUNCTION_CSE NO_PROFILE_COUNTERS NO_REGS NON_LVALUE_EXPR nondeterministic finite state automaton nonimmediate_operand nonlocal_goto instruction pattern	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 605 605 580 552 512 177 464 348 321	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENTER_DATA OACC_ENIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN OBJC_JBLEN OBJCCT_FORMAT_COFF OFFSET_TYPE offsettable address OImode OMP_ATOMIC OMP_CLAUSE OMP_CONTINUE	351 192 192 192 192 192 192 192 192 192 19
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_FUNCTION_CSE NO_PROFILE_COUNTERS NO_REGS NON_LVALUE_EXPR nondeterministic finite state automaton nonimmediate_operand nonlocal goto handler	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 628 605 580 552 512 177 464 348 321 423	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENITE_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN OBJCC_JBLEN OBJECT_FORMAT_COFF OFFSET_TYPE offsettable address OImode OMP_ATOMIC OMP_CLAUSE OMP_CONTINUE OMP_CRITICAL	351 192 192 192 192 192 192 192 192 192 19
negation with signed saturation negation with unsigned saturation negm2 instruction pattern negmodecc instruction pattern negvm3 instruction pattern nested functions, support for nested_ptr next_bb, prev_bb, FOR_EACH_BB, FOR_ALL_BB  NEXT_INSN NEXT_OBJC_RUNTIME nil NM_FLAGS NO_DBX_BNSYM_ENSYM NO_DBX_FUNCTION_END NO_DBX_GCC_MARKER NO_DBX_MAIN_SOURCE_DIRECTORY NO_DOLLAR_IN_LABEL NO_FUNCTION_CSE NO_PROFILE_COUNTERS NO_REGS NON_LVALUE_EXPR nondeterministic finite state automaton nonimmediate_operand nonlocal_goto instruction pattern nonlocal_goto_receiver instruction pattern	289 289 406 418 406 558 678 317 305 561 260 615 628 628 628 628 605 580 552 512 177 464 348 321 423 424	NUM_MODES_FOR_MODE_SWITCHING NUM_POLY_INT_COEFFS Number of iterations analysis  O  'o' in constraint OACC_CACHE OACC_DATA OACC_DECLARE OACC_ENITE_DATA OACC_EXIT_DATA OACC_EXIT_DATA OACC_KERNELS OACC_LOOP OACC_PARALLEL OACC_SERIAL OACC_UPDATE OBJC_GEN_METHOD_LABEL OBJC_JBLEN OBJCC_JBLEN OBJCCT_FORMAT_COFF OFFSET_TYPE offsettable address OImode OMP_ATOMIC OMP_CLAUSE OMP_CONTINUE	351 192 192 192 192 192 192 192 192 615 665 615 163 351 272 190 190 190

OMP_PARALLEL		parallel	
OMP_RETURN		parameters, c++ abi	
OMP_SECTION		parameters, d abi	
OMP_SECTIONS		parameters, miscellaneous	
OMP_SINGLE		parameters, precompiled headers	
one_cmplm2 instruction pattern		parity	
operand access		paritym2 instruction pattern	
Operand Access Routines		PARM_BOUNDARY	
operand constraints		PARM_DECL	
Operand Iterators		PARSE_LDD_OUTPUT	
operand predicates		pass dumps	
operand substitution		pass_duplicate_computed_gotos	
Operands		passes and files of the compiler	
operands		passing arguments	
operator predicates		PATH_SEPARATOR	
opt_mode		pattern conditions	
'optc-gen.awk'		pattern names	
OPTGROUP_ALL		Pattern Ordering	
OPTGROUP_INLINE		patterns	
OPTGROUP_IPA		PATTERN	
OPTGROUP_LOOP		pc	
OPTGROUP_OMP		pc and attributes	
OPTGROUP_OTHER		pc, RTL sharing	
OPTGROUP_VEC		pc_rtx	288
optimization dumps		PC_REGNUM	
optimization groups		PCC_BITFIELD_TYPE_MATTERS	
optimization info file names		PCC_STATIC_STRUCT_RETURN	
Optimization infrastructure for GIMPLE	247	PDImode	
OPTIMIZE_MODE_SWITCHING		peephole optimization, RTL representation	300
option specification files		peephole optimizer definitions	446
OPTION_DEFAULT_SPECS		per-function data	489
optional hardware or system features		percent sign	
options, directory search		PHI nodes	
options, guidelines for		PIC	595
order of register allocation		PIC_OFFSET_TABLE_REG_CALL_CLOBBERED	596
ordered_comparison_operator		PIC_OFFSET_TABLE_REGNUM	
ORDERED_EXPR		pipeline hazard recognizer 460,	
Ordering of Patterns		Plugins	
ORIGINAL_REGNO		plus	
other register constraints		plus and attributes	454
outgoing_args_size		plus, canonicalization of	436
OUTGOING_REG_PARM_STACK_SPACE		PLUS_EXPR	
OUTGOING_REGNO		Pmode	
output of assembler code		pmode_register_operand	347
output statements	344	pointer	163
output templates		POINTER_DIFF_EXPR	177
output_asm_insn		POINTER_PLUS_EXPR	177
OUTPUT_QUOTED_STRING	598	POINTER_SIZE	491
OVERLAPPING_REGISTER_NAMES	616	POINTER_TYPE	163
OVERLOAD	201	POINTERS_EXTEND_UNSIGNED	491
OVERRIDE_ABI_FORMAT	539	poly_int	
OVL_CURRENT	201	poly_int, invariant range	
OVL_NEXT	201	poly_int, main typedefs	147
		poly_int, runtime value	
D		poly_int, template parameters	
P		poly_int, use in target-independent code	
'p' in constraint	353	poly_int, use in target-specific code	
PAD_VARARGS_DOWN		POLY_INT_CST	

polynomial integers	47	PROMOTE_MODE	491
pop_operand 34	48	pseudo registers	282
popcount	92	PSImode	272
popcountm2 instruction pattern 41	11	PTRDIFF_TYPE	502
pops_args 55	50	purge_dead_edges320,	324
portability	5	push address instruction	353
position independent code 59	95	push_operand	348
post_dec		push_reload	
post_inc		PUSH_ARGS	534
post_modify30		PUSH_ARGS_REVERSED	534
post_order_compute,		PUSH_ROUNDING	535
inverted_post_order_compute,		pushm1 instruction pattern	399
walk_dominator_tree31		PUT_CODE	
POST_LINK_SPEC48	83	PUT_MODE	277
POSTDECREMENT_EXPR		PUT_REG_NOTE_KIND	310
POSTINCREMENT_EXPR 17			
POWI_MAX_MULTS6	57		
powm3 instruction pattern		$\mathbf{Q}$	
pragma		QCmode	274
pre_dec		QFmode	
pre_inc		QImode	
pre_modify 30		QImode, in insn	
PRE_GCC3_DWARF_FRAME_REGISTERS		QQmode	
PREDECREMENT_EXPR		qualified type	
predefined macros		querying function unit reservations	
predicates		question mark	
predicates and machine modes		quotient	
predication		questone	
predict.def	22		
PREFERRED_DEBUGGING_TYPE		${ m R}$	
PREFERRED_RELOAD_CLASS		'r' in constraint	351
PREFERRED_STACK_BOUNDARY		RDIV_EXPR	
prefetch		READONLY_DATA_SECTION_ASM_OP	
prefetch and '/v'	~-	real operands	
prefetch instruction pattern 42		REAL_CST	
PREFETCH_SCHEDULE_BARRIER_P		REAL_LIBGCC_SPEC	
PREINCREMENT_EXPR		REAL_NM_FILE_NAME	
presence_set 40		REAL TYPE	
preserving SSA form		REAL_VALUE_ABS	
pretend_args_size		REAL_VALUE_ATOF	
prev_active_insn		REAL_VALUE_FIX	
PREV_INSN		REAL_VALUE_ISINF	
PRINT_OPERAND		REAL_VALUE_ISNAN	
PRINT_OPERAND_ADDRESS		REAL_VALUE_ISNAN	
PRINT_OPERAND_PUNCT_VALID_P61		REAL_VALUE_NEGATIVE	
probe_stack instruction pattern		REAL_VALUE_TO_TARGET_DECIMAL128	
probe_stack_address instruction pattern 42		REAL_VALUE_TO_TARGET_DECIMAL32	
processor functional units		REAL_VALUE_TO_TARGET_DECIMAL64	
processor pipeline description		REAL_VALUE_TO_TARGET_DOUBLE	
product		REAL_VALUE_TO_TARGET_LONG_DOUBLE	
profile feedback		REAL_VALUE_TO_TARGET_LONG_DOOBLE	
profile representation		REAL_VALUE_TYPE	
PROFILE_BEFORE_PROLOGUE		REAL_VALUE_UNSIGNED_FIX	
PROFILE_HOOK		REALPART_EXPR	
profiling, code generation		recog_data.operand	
program counter		recognizing insns	
prologue		RECORD_TYPE	
prologue instruction pattern		redirect edge and branch	
PI OIOGUO IIISUI UCUIOII PAUUCIII	40	LOGILOGO CARECAMA DI AMOM	040

<pre>redirect_edge_and_branch, redirect_jump</pre>		REG_WORDS_BIG_ENDIAN	491
	324	register allocation order	507
<pre>reduc_and_scal_m instruction pattern</pre>	400	register class definitions	512
reduc_ior_scal_m instruction pattern	400	register class preference constraints	356
<pre>reduc_plus_scal_m instruction pattern</pre>	400	register pairs	509
<pre>reduc_smax_scal_m instruction pattern</pre>	400	Register Transfer Language (RTL)	259
reduc_smin_scal_m instruction pattern	400	register usage	
reduc_umax_scal_m instruction pattern	400	register_operand	347
reduc_umin_scal_m instruction pattern	400	REGISTER_MOVE_COST	
reduc_xor_scal_m instruction pattern		REGISTER_NAMES	
reference		REGISTER_PREFIX	
REFERENCE_TYPE		REGISTER_TARGET_PRAGMAS	
reg		registers arguments	
reg and '/f'		registers in constraints	
reg and '/i'		REGMODE_NATURAL_SIZE 285, 286,	
reg and '/v'		REGNO_MODE_CODE_OK_FOR_BASE_P	
reg, RTL sharing		REGNO_MODE_OK_FOR_BASE_P	
reg_class_contents		REGNO_MODE_OK_FOR_REG_BASE_P	
reg_class_for_constraint		REGNO_OK_FOR_BASE_P	
reg_label and '/v'		REGNO_OK_FOR_INDEX_P	
reg_names		REGNO_REG_CLASS	
REG_ALLOC_ORDER		regs_ever_live	
REG_BR_PRED		regular expressions	
REG_BR_PROB		regular IPA passes	
REG_BR_PROB_BASE, BB_FREQ_BASE, count		relative costs	
REG_BR_PROB_BASE, EDGE_FREQUENCY		RELATIVE_PREFIX_NOT_LINKDIR	
REG_CALL_NOCF_CHECK		reload_completed	
REG_CC_SETTER		reload_in instruction pattern	
REG_CC_USER		reload_in_progress	
REG_CLASS_CONTENTS		reload_out instruction pattern	
REG_CLASS_NAMES		reloading	
REG_DEAD.		remainder	
REG_DEAD, REG_UNUSED		remainderm3 instruction pattern	
REG_DEP_ANTI		reorder	
REG_DEP_OUTPUT		representation of RTL	
REG_DEP_TRUE		reservation delays	
		rest_of_decl_compilation	
REG_EH_REGION, EDGE_ABNORMAL_CALL		rest_of_type_compilation	
REG_EQUIV.		restore_stack_block instruction pattern	
REG_EXPR.		restore_stack_function instruction pattern restore_stack_function instruction pattern	422
REG_FRAME_RELATED_EXPR.		restore_stack_runction instruction pattern	499
REG_FUNCTION_VALUE_P		restore_stack_nonlocal instruction pattern	422
REG_INC		restore_stack_nonrocal instruction pattern	499
REG_LABEL_OPERAND		RESULT_DECL	
REG_LABEL_TARGET		return	
REG_NONNEG		return instruction pattern	
REG_NOTE_KIND		return values in registers	
REG_NOTES		return_val	
_	264	return_val, in call_insn	
REG_OK_STRICT		return_val, in reg	
REG_PARM_STACK_SPACE		return_val, in symbol_ref	
REG_PARM_STACK_SPACE, and TARGET_FUNCTION_A		RETURN_ADDR_IN_PREVIOUS_FRAME	
		RETURN_ADDR_OFFSET	
REG_POINTER		RETURN_ADDREGG_DOINTED_REGNIM	
REG_SETJMP		RETURN_ADDRESS_POINTER_REGNUM	
REG_UNUSED		RETURN_EXPR	
REG_USERVAR_P		RETURN_STMT	
REC VALUE IN UNWIND CONTEXT	533	returning aggregate values	546

returning structures and unions		RTX_FRAME_RELATED_P	
reverse probability		run-time conventions	
REVERSE_CONDITION		run-time target specification	486
REVERSIBLE_CC_MODE			
right rotate		S	
right shift		S	
<pre>rintm2 instruction pattern</pre>		's' in constraint	352
RISC		SAD_EXPR	184
roots, marking	682	same_type_p	
rotate	291	SAmode	
rotatert	291	sat_fract	296
rotlm3 instruction pattern	406	satfractmn2 instruction pattern	
rotrm3 instruction pattern	406	satfractunsmn2 instruction pattern	
ROUND_DIV_EXPR	177	satisfies_constraint_m	
ROUND_MOD_EXPR	177	save_stack_block instruction pattern	
ROUND_TYPE_ALIGN		save_stack_function instruction pattern	
roundm2 instruction pattern		save_stack_nonlocal instruction pattern	
RSHIFT_EXPR		SAVE_EXPR	
rsqrtm2 instruction pattern		SBSS_SECTION_ASM_OP	
RTL addition		Scalar evolutions	
RTL addition with signed saturation		scalar_float_mode	
RTL addition with unsigned saturation		scalar_int_mode	
RTL classes			
RTL comparison		scalar_mode	
RTL comparison operations		scalars, returned as values	
RTL constant expression types		scalbm3 instruction pattern	
RTL constants		scatter_storemn instruction pattern	
		SCHED_GROUP_P	
RTL declarations		SCmode	
RTL difference		scratch	
RTL expression		scratch operands	
RTL expressions for arithmetic		scratch, RTL sharing	
RTL format		${\tt scratch\_operand}$	
RTL format characters		SDATA_SECTION_ASM_OP	
RTL function-call insns		<pre>sdiv_pow2m3 instruction pattern</pre>	
RTL insn template		SDmode	
RTL integers	259	<pre>sdot_prodm instruction pattern</pre>	
RTL memory expressions		search options	
RTL object types		SECONDARY_INPUT_RELOAD_CLASS	518
RTL postdecrement	302	SECONDARY_MEMORY_NEEDED_RTX	519
RTL postincrement	302	SECONDARY_OUTPUT_RELOAD_CLASS	518
RTL predecrement	302	SECONDARY_RELOAD_CLASS	
RTL preincrement	302	SELECT_CC_MODE	
RTL register expressions	282	sequence	301
RTL representation		Sequence iterators	
RTL side effect expressions		set	
RTL strings		set and '/f'	
RTL structure sharing assumptions		set_attr	
RTL subtraction		set_attr_alternative	
RTL subtraction with signed saturation		set_bb_seq	
RTL subtraction with unsigned saturation		set_optab_libfunc	
RTL sum		set_thread_pointermode instruction pattern	500
RTL vectors.		set_thread_pointermode instruction pattern	431
RTL_CONST_CALL_P.		SET_ASM_OP	
RTL_CONST_CALL_F		SET_DEST	
		SET_IS_RETURN_P	
RTL_LOOPING_CONST_OR_PURE_CALL_PRTL_PURE_CALL_P			
		SET_LABEL_KIND	
RTX (See RTL)		SET_RATIO	
RTX codes, classes of	∠0U	SET_SRC	∠98

SET_TYPE_STRUCTURAL_EQUALITY 163, 165	ss_abs 2	91
setmemm instruction pattern 413	ss_ashift 2	91
SETUP_FRAME_ADDRESSES	ss_div 2	90
SFmode	ss_minus 2	89
sharing of RTL components	ss_mult 2	89
shift	ss_neg 2	89
SHIFT_COUNT_TRUNCATED644	ss_plus 2	88
SHLIB_SUFFIX616	ss_truncate 2	
SHORT_ACCUM_TYPE_SIZE501	SSA_NAME_DEF_STMT 2	
SHORT_FRACT_TYPE_SIZE501	SSA_NAME_VERSION2	
SHORT_IMMEDIATES_SIGN_EXTEND644	ssaddm3 instruction pattern 3	
SHORT_TYPE_SIZE500	ssadm instruction pattern 4	
shrink-wrapping separate components 553	ssashlm3 instruction pattern 4	
sibcall_epilogue instruction pattern 425	SSA	
sibling call	ssdivm3 instruction pattern 3	
SIBLING_CALL_P	ssmaddmn4 instruction pattern 4	
SIG_ATOMIC_TYPE503	ssmsubmn4 instruction pattern 4	
sign_extend	ssmulm3 instruction pattern	
sign_extract         294	ssnegm2 instruction pattern 4	
sign_extract, canonicalization of	sssubm3 instruction pattern	
signal-to-noise ratio (metaphorical usage for	stack arguments	
diagnostics)	stack frame layout	
signed division	stack mashing protection	
signed division with signed saturation 290	stack_pointer_rtx5	
signed maximum	stack_protect_combined_set instruction pattern	
signed maximum 290 signed minimum 290	stack_protect_combined_set instruction pattern	
significandm2 instruction pattern	stack_protect_combined_test instruction patte	
-	stack_protect_combined_test instruction parte	
SImode		
simple constraints	stack_protect_set instruction pattern 4	
simple_return	stack_protect_test instruction pattern 4	
simple_return instruction pattern	STACK_ALIGNMENT_NEEDED	
sincosm3 instruction pattern 407	STACK_BOUNDARY4	
sinm2 instruction pattern	STACK_CHECK_BUILTIN	
SIZE_ASM_OP	STACK_CHECK_FIXED_FRAME_SIZE	
SIZE_TYPE	STACK_CHECK_MAX_FRAME_SIZE	
SIZETYPE	STACK_CHECK_MAX_VAR_SIZE	
skip	STACK_CHECK_MOVING_SP5	
SLOW_BYTE_ACCESS	STACK_CHECK_PROBE_INTERVAL_EXP 5	
small IPA passes	STACK_CHECK_PROTECT 5	$^{29}$
smax		
smin	STACK_CHECK_STATIC_BUILTIN 5	29
	STACK_DYNAMIC_OFFSET	$\frac{29}{23}$
sms, swing, software pipelining	STACK_DYNAMIC_OFFSET	29 23 283
smulhrsm3 instruction pattern 402	STACK_DYNAMIC_OFFSET	529 523 283 522
smulhrsm3 instruction pattern    402      smulhsm3 instruction pattern    401	STACK_DYNAMIC_OFFSET	29 23 83 622 635
smulhrsm3 instruction pattern 402	STACK_DYNAMIC_OFFSET	29 23 83 22 35 23
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12	STACK_DYNAMIC_OFFSET	29 23 83 22 35 23
smulhrsm3 instruction pattern       402         smulhsm3 instruction pattern       401         smulm3_highpart instruction pattern       404	STACK_DYNAMIC_OFFSET	29 623 622 635 623 630
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679	STACK_DYNAMIC_OFFSET5STACK_DYNAMIC_OFFSET and virtual registers2STACK_GROWS_DOWNWARD5STACK_PARMS_IN_REG_PARM_AREA5STACK_POINTER_OFFSET5STACK_POINTER_OFFSET and virtual registers2	29 623 622 635 623 630
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679special predicates346	STACK_DYNAMIC_OFFSET	29 23 83 22 35 23 83 30 83
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679	STACK_DYNAMIC_OFFSET	29 23 83 622 635 623 83 630 83
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679special predicates346	STACK_DYNAMIC_OFFSET	629 623 622 635 623 630 83 622 612
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679special predicates346SPECS670speculation_barrier instruction pattern426	STACK_DYNAMIC_OFFSET	529 523 522 535 523 530 530 530 522 512
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679special predicates346SPECS670speculation_barrier instruction pattern426	STACK_DYNAMIC_OFFSET	29 23 83 22 35 23 83 62 61 61 61 61 61 61
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679special predicates346SPECS670speculation_barrier instruction pattern426speed of instructions576	STACK_DYNAMIC_OFFSET	29 23 83 22 35 23 30 83 62 61 61 61 19 8
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679special predicates346SPECS670speculation_barrier instruction pattern426speed of instructions576split_block324	STACK_DYNAMIC_OFFSET	629 623 83 622 635 623 83 622 612 611 98
smulhrsm3 instruction pattern402smulhsm3 instruction pattern401smulm3_highpart instruction pattern404soft float library12source code, location information717special679special predicates346SPECS670speculation_barrier instruction pattern426speed of instructions576split_block324splitting instructions440	STACK_DYNAMIC_OFFSET	129 123 123 122 135 123 130 183 198 198 198
smulhrsm3 instruction pattern       402         smulhsm3 instruction pattern       401         smulm3_highpart instruction pattern       404         soft float library       12         source code, location information       717         special       679         special predicates       346         SPECS       670         speculation_barrier instruction pattern       426         speed of instructions       576         split_block       324         splitting instructions       440         SQmode       273	STACK_DYNAMIC_OFFSET	629 623 622 635 623 623 623 623 612 612 612 612 612 612 612 612 612 612

STARTFILE_SPEC	symbol_ref and '/u'
Statement and operand traversals	symbol_ref and '/v'
Statement Sequences	symbol_ref, RTL sharing
statements	SYMBOL_FLAG_ANCHOR
Statements	SYMBOL_FLAG_EXTERNAL 265
static analysis	SYMBOL_FLAG_FUNCTION
static analyzer	SYMBOL_FLAG_HAS_BLOCK_INFO
static analyzer, debugging	SYMBOL_FLAG_LOCAL
static analyzer, internals	SYMBOL_FLAG_SMALL 265
Static profile estimation	SYMBOL_FLAG_TLS_SHIFT
$ \begin{array}{llllllllllllllllllllllllllllllllllll$	SYMBOL_REF_ANCHOR_P
	SYMBOL_REF_BLOCK
STATIC_CHAIN_REGNUM	SYMBOL_REF_BLOCK_OFFSET
STDC_0_IN_SYSTEM_HEADERS	SYMBOL_REF_CONSTANT
STMT_EXPR	SYMBOL_REF_DATA
STMT_IS_FULL_EXPR_P	SYMBOL_REF_DECL
storage layout	SYMBOL_REF_EXTERNAL_P
'store_multiple' instruction pattern	SYMBOL_REF_FLAG
STORE_FLAG_VALUE	SYMBOL_REF_FLAG, in
	TARGET_ENCODE_SECTION_INFO595
STORE_MAX_PIECES	SYMBOL_REF_FLAGS
strcpy	SYMBOL_REF_FUNCTION_P
strict_low_part	SYMBOL_REF_HAS_BLOCK_INFO_P
strict_memory_address_p 564	SYMBOL_REF_LOCAL_P
STRICT_ALIGNMENT         496           STRING_CST         173	SYMBOL_REF_SMALL_P
STRING_POOL_ADDRESS_P	SYMBOL_REF_TLS_MODEL
strlenm instruction pattern	SYMBOL_REF_USED
structure value address	SYMBOL_REF_WEAK
STRUCTURE_SIZE_BOUNDARY	symbolic label
structures, returning	sync_addmode instruction pattern 427
subm3 instruction pattern	sync_andmode instruction pattern 427
SUBOBJECT	sync_compare_and_swapmode instruction pattern
SUBOBJECT_CLEANUP	426
subreg	sync_iormode instruction pattern 427
subreg and '/s'	sync_lock_releasemode instruction pattern 428
subreg and '/u'	sync_lock_test_and_setmode instruction pattern
subreg and '/u' and '/v'	
subreg, in strict_low_part 297	sync_nandmode instruction pattern 427
SUBREG_BYTE	sync_new_addmode instruction pattern 428
SUBREG_PROMOTED_UNSIGNED_P	sync_new_andmode instruction pattern 428
SUBREG_PROMOTED_UNSIGNED_SET	sync_new_iormode instruction pattern 428
SUBREG_PROMOTED_VAR_P	sync_new_nandmode instruction pattern 428
SUBREG_REG	sync_new_submode instruction pattern 428
subst iterators in '.md' files	sync_new_xormode instruction pattern 428
subvm4 instruction pattern	sync_old_addmode instruction pattern 427
SUCCESS_EXIT_CODE	sync_old_andmode instruction pattern 427
support for nested functions	sync_old_iormode instruction pattern 427
SUPPORTS_INIT_PRIORITY614	sync_old_nandmode instruction pattern 427
SUPPORTS_ONE_ONLY	sync_old_submode instruction pattern 427
SUPPORTS_WEAK	sync_old_xormode instruction pattern 427 sync_old_xormode instruction pattern 427
SWITCH_BODY	sync_submode instruction pattern 427 sync_submode instruction pattern
SWITCH_COND	sync_submode instruction pattern
SWITCH_STMT 203	SYSROOT_HEADERS_SUFFIX_SPEC
SWITCHABLE_TARGET 489	SYSROOT_SUFFIX_SPEC
symbol_ref	SYSTEM_IMPLICIT_EXTERN_C
symbol_ref and '/f'	DIDIER_IRFLICII_EAIEMN_C 049
symbol_ref and '/i'	

$\mathbf{T}$	TARGET_ASM_EMIT_UNWIND_LABEL	620
	TARGET_ASM_ENTI_ONWIND_LABEL	
't-target' 667		
table jump	TARGET_ASM_FILE_END	
tablejump instruction pattern 421	TARGET_ASM_FILE_START	
tag 677	TARGET_ASM_FILE_START_APP_OFF	
tagging insns	TARGET_ASM_FILE_START_FILE_DIRECTIVE	
tail calls	TARGET_ASM_FINAL_POSTSCAN_INSN	617
TAmode	TARGET_ASM_FUNCTION_BEGIN_EPILOGUE	549
tanm2 instruction pattern	TARGET_ASM_FUNCTION_END_PROLOGUE	549
target attributes	TARGET_ASM_FUNCTION_EPILOGUE	549
target description macros	TARGET_ASM_FUNCTION_PROLOGUE	548
target functions	TARGET_ASM_FUNCTION_RODATA_SECTION	594
target hooks	TARGET_ASM_FUNCTION_SECTION	
	TARGET_ASM_FUNCTION_SWITCHED_TEXT_SECTION	
target makefile fragment		
target specifications	TARGET_ASM_GENERATE_PIC_ADDR_DIFF_VEC	
target_flags	TARGET_ASM_GLOBALIZE_DECL_NAME	
TARGET_ABSOLUTE_BIGGEST_ALIGNMENT 493	TARGET_ASM_GLOBALIZE_DEGL_NAME	
TARGET_ADDITIONAL_ALLOCNO_CLASS_P 522		
TARGET_ADDR_SPACE_ADDRESS_MODE 641	TARGET_ASM_INIT_SECTIONS	
TARGET_ADDR_SPACE_CONVERT 642	TARGET_ASM_INTEGER	
TARGET_ADDR_SPACE_DEBUG 642	TARGET_ASM_INTERNAL_LABEL	
TARGET_ADDR_SPACE_DIAGNOSE_USAGE 642	TARGET_ASM_LTO_END	
TARGET_ADDR_SPACE_LEGITIMATE_ADDRESS_P 641	TARGET_ASM_LTO_START	
TARGET_ADDR_SPACE_LEGITIMIZE_ADDRESS 641	TARGET_ASM_MARK_DECL_PRESERVED	
TARGET_ADDR_SPACE_POINTER_MODE 641	TARGET_ASM_MERGEABLE_RODATA_PREFIX	594
TARGET_ADDR_SPACE_SUBSET_P642	TARGET_ASM_NAMED_SECTION	598
TARGET_ADDR_SPACE_VALID_POINTER_MODE 641	TARGET_ASM_OPEN_PAREN	602
TARGET_ADDR_SPACE_ZERO_ADDRESS_VALID 642	TARGET_ASM_OUTPUT_ADDR_CONST_EXTRA	600
TARGET_ADDRESS_COST	TARGET_ASM_OUTPUT_ANCHOR	572
TARGET_ALIGN_ANON_BITFIELD	TARGET_ASM_OUTPUT_DWARF_DTPREL	
TARGET_ALLOCATE_INITIAL_VALUE	TARGET_ASM_OUTPUT_IDENT	
TARGET_ALLOCATE_STACK_SLOTS_FOR_ARGS 659	TARGET_ASM_OUTPUT_MI_THUNK	
	TARGET_ASM_OUTPUT_SOURCE_FILENAME	
TARGET_ALWAYS_STRIP_DOTDOT	TARGET_ASM_POST_CFI_STARTPROC	
TARGET_ARG_PARTIAL_BYTES	TARGET_ASM_PRINT_PATCHABLE_FUNCTION_ENTRY	
TARGET_ARM_EABI_UNWINDER	TARGET_ASM_FRINT_FATGRADEE_FUNCTION_ENTRI	
TARGET_ARRAY_MODE		
TARGET_ARRAY_MODE_SUPPORTED_P	TARGET_ASM_RECORD_GCC_SWITCHES	
TARGET_ASAN_SHADOW_OFFSET	TARGET_ASM_RECORD_GCC_SWITCHES_SECTION	
TARGET_ASM_ALIGNED_DI_OP	TARGET_ASM_RELOC_RW_MASK	
TARGET_ASM_ALIGNED_HI_OP	TARGET_ASM_SELECT_RTX_SECTION	
TARGET_ASM_ALIGNED_PDI_OP 600	TARGET_ASM_SELECT_SECTION	
TARGET_ASM_ALIGNED_PSI_OP 600	TARGET_ASM_SHOULD_RESTORE_CFA_STATE	
TARGET_ASM_ALIGNED_PTI_OP 600	TARGET_ASM_TM_CLONE_TABLE_SECTION	
TARGET_ASM_ALIGNED_SI_OP 600	TARGET_ASM_TRAMPOLINE_TEMPLATE	
TARGET_ASM_ALIGNED_TI_OP 600	TARGET_ASM_TTYPE	
TARGET_ASM_ASSEMBLE_UNDEFINED_DECL 608	TARGET_ASM_UNALIGNED_DI_OP	600
TARGET_ASM_ASSEMBLE_VISIBILITY 609	TARGET_ASM_UNALIGNED_HI_OP	600
TARGET_ASM_BYTE_OP	TARGET_ASM_UNALIGNED_PDI_OP	600
TARGET_ASM_CAN_OUTPUT_MI_THUNK 551	TARGET_ASM_UNALIGNED_PSI_OP	600
TARGET_ASM_CLOSE_PAREN	TARGET_ASM_UNALIGNED_PTI_OP	
TARGET_ASM_CODE_END	TARGET_ASM_UNALIGNED_SI_OP	
TARGET_ASM_CONSTRUCTOR	TARGET_ASM_UNALIGNED_TI_OP	
TARGET_ASM_DECL_END	TARGET_ASM_UNIQUE_SECTION	
TARGET_ASM_DECLARE_CONSTANT_NAME	TARGET_ASM_UNWIND_EMIT	
TARGET_ASM_DESTRUCTOR	TARGET_ASM_UNWIND_EMIT_BEFORE_INSN	
TARGET_ASM_ELF_FLAGS_NUMERIC	TARGET_ATOMIC_ALIGN_FOR_MODE	
TARGET_ASM_EMIT_EXCEPT_PERSONALITY 621	TARGET_ATOMIC_ASSIGN_EXPAND_FENV	
TARGET ASM EMIT EXCEPT TABLE LABEL 620	TARGET ATOMIC TEST AND SET TRUEVAL	บบบ

TARGET_ATTRIBUTE_TABLE633	TARGET_CXX_IMPORT_EXPORT_CLASS	
TARGET_ATTRIBUTE_TAKES_IDENTIFIER_P 633	TARGET_CXX_KEY_METHOD_MAY_BE_INLINE	
TARGET_BINDS_LOCAL_P	TARGET_CXX_LIBRARY_RTTI_COMDAT	
TARGET_BUILD_BUILTIN_VA_LIST541	TARGET_CXX_USE_AEABI_ATEXIT	
TARGET_BUILTIN_DECL	TARGET_CXX_USE_ATEXIT_FOR_CXA_ATEXIT	
TARGET_BUILTIN_RECIPROCAL 566	TARGET_D_CPU_VERSIONS	
TARGET_BUILTIN_SETJMP_FRAME_VALUE 524	TARGET_D_CRITSEC_SIZE	
TARGET_C_EXCESS_PRECISION 492	TARGET_D_OS_VERSIONS	
TARGET_C_PREINCLUDE 648	TARGET_DEBUG_UNWIND_INFO	629
TARGET_CALL_ARGS557	TARGET_DECIMAL_FLOAT_SUPPORTED_P	499
TARGET_CALL_FUSAGE_CONTAINS_NON_CALLEE_	TARGET_DECLSPEC	634
CLOBBERS	TARGET_DEFAULT_PACK_STRUCT	650
TARGET_CALLEE_COPIES538	TARGET_DEFAULT_SHORT_ENUMS	502
TARGET_CAN_CHANGE_MODE_CLASS	TARGET_DEFAULT_TARGET_FLAGS	
TARGET_CAN_CHANGE_MODE_CLASS and subreg	TARGET_DEFERRED_OUTPUT_DEFS	
semantics	TARGET_DELAY_SCHED2	
TARGET_CAN_ELIMINATE	TARGET_DELAY_VARTRACK	
TARGET_CAN_FOLLOW_JUMP	TARGET_DELEGITIMIZE_ADDRESS	
TARGET_CAN_INLINE_P	TARGET_DIFFERENT_ADDR_DISPLACEMENT_P	
TARGET_CAN_USE_DOLOOP_P	TARGET_DLLIMPORT_DECL_ATTRIBUTES	
TARGET_CANNOT_FORCE_CONST_MEM	TARGET_DOLOOP_COST_FOR_ADDRESS	
TARGET_CANNOT_MODIFY_JUMPS_P	TARGET_DOLOOP_COST_FOR_GENERIC	
TARGET_CANNOT_SUBSTITUTE_MEM_EQUIV_P 521	TARGET_DWARF_CALLING_CONVENTION	
TARGET_CANONICAL_VA_LIST_TYPE	TARGET_DWARF_FRAME_REG_MODE	
TARGET_CANONICALIZE_COMPARISON	TARGET_DWARF_HANDLE_FRAME_UNSPEC	
	TARGET_DWARF_POLY_INDETERMINATE_VALUE	
TARGET_CASE_VALUES_THRESHOLD		
TARGET_CC_MODES_COMPATIBLE	TARGET_DWARF_REGISTER_SPAN	
TARGET_CHECK_BUILTIN_CALL	TARGET_EDOM	
TARGET_CHECK_PCH_TARGET_FLAGS	TARGET_EMPTY_RECORD_P	
TARGET_CHECK_STRING_OBJECT_FORMAT_ARG 488	TARGET_EMUTLS_DEBUG_FORM_TLS_ADDRESS	
TARGET_CLASS_LIKELY_SPILLED_P	TARGET_EMUTLS_GET_ADDRESS	
TARGET_CLASS_MAX_NREGS	TARGET_EMUTLS_REGISTER_COMMON	
TARGET_COMMUTATIVE_P655	TARGET_EMUTLS_TMPL_PREFIX	
TARGET_COMP_TYPE_ATTRIBUTES634	TARGET_EMUTLS_TMPL_SECTION	
TARGET_COMPARE_BY_PIECES_BRANCH_RATIO 579	TARGET_EMUTLS_VAR_ALIGN_FIXED	
TARGET_COMPARE_VERSION_PRIORITY 653	TARGET_EMUTLS_VAR_FIELDS	
TARGET_COMPATIBLE_VECTOR_TYPES_P 542	TARGET_EMUTLS_VAR_INIT	
TARGET_COMPUTE_FRAME_LAYOUT	TARGET_EMUTLS_VAR_PREFIX	
TARGET_COMPUTE_PRESSURE_CLASSES 522	TARGET_EMUTLS_VAR_SECTION	
TARGET_CONDITIONAL_REGISTER_USAGE 506	TARGET_ENCODE_SECTION_INFO	594
TARGET_CONST_ANCHOR659	TARGET_ENCODE_SECTION_INFO and address	
TARGET_CONST_NOT_OK_FOR_DEBUG_P 565	validation	
TARGET_CONSTANT_ALIGNMENT 495	TARGET_ENCODE_SECTION_INFO usage	618
TARGET_CONVERT_TO_TYPE659	TARGET_END_CALL_ARGS	557
TARGET_CPU_CPP_BUILTINS 486	TARGET_ENUM_VA_LIST_P	541
TARGET_CSTORE_MODE	TARGET_ESTIMATED_POLY_VALUE	582
TARGET_CUSTOM_FUNCTION_DESCRIPTORS 558	TARGET_EXCEPT_UNWIND_INFO	
TARGET_CXX_ADJUST_CLASS_AT_DEFINITION 640	TARGET_EXECUTABLE_SUFFIX	656
TARGET_CXX_CDTOR_RETURNS_THIS	TARGET_EXPAND_BUILTIN	
TARGET_CXX_CLASS_DATA_ALWAYS_COMDAT 639	TARGET_EXPAND_BUILTIN_SAVEREGS	
TARGET_CXX_COOKIE_HAS_SIZE		
TARGET_CXX_DECL_MANGLING_CONTEXT		590
	TARGET_EXPAND_DIVMOD_LIBFUNC	
	TARGET_EXPAND_DIVMOD_LIBFUNC  TARGET_EXPAND_TO_RTL_HOOK	499
TARGET_CXX_DETERMINE_CLASS_DATA_VISIBILITY	TARGET_EXPAND_DIVMOD_LIBFUNCTARGET_EXPAND_TO_RTL_HOOKTARGET_EXPR	499 177
TARGET_CXX_DETERMINE_CLASS_DATA_VISIBILITY 639	TARGET_EXPAND_DIVMOD_LIBFUNC  TARGET_EXPAND_TO_RTL_HOOK  TARGET_EXPR  TARGET_EXTRA_INCLUDES	$499 \\ 177 \\ 657$
TARGET_CXX_DETERMINE_CLASS_DATA_VISIBILITY	TARGET_EXPAND_DIVMOD_LIBFUNC TARGET_EXPAND_TO_RTL_HOOK TARGET_EXPR TARGET_EXTRA_INCLUDES TARGET_EXTRA_LIVE_ON_ENTRY	499 177 657 552
TARGET_CXX_DETERMINE_CLASS_DATA_VISIBILITY 639	TARGET_EXPAND_DIVMOD_LIBFUNC  TARGET_EXPAND_TO_RTL_HOOK  TARGET_EXPR  TARGET_EXTRA_INCLUDES	499 177 657 552 657

TARGET_FLAGS_REGNUM	TARGET_INIT_DWARF_REG_SIZES_EXTRA	
TARGET_FLOAT_EXCEPTIONS_ROUNDING_SUPPORTED_	TARGET_INIT_LIBFUNCS	
P	TARGET_INIT_PIC_REG	
TARGET_FLOATN_BUILTIN_P 543	TARGET_INSERT_ATTRIBUTES	
TARGET_FLOATN_MODE	TARGET_INSN_CALLEE_ABI	
TARGET_FN_ABI_VA_LIST541	TARGET_INSN_COST	
TARGET_FNTYPE_ABI	TARGET_INSTANTIATE_DECLS	
TARGET_FOLD_BUILTIN	TARGET_INVALID_ARG_FOR_UNPROTOTYPED_FN	
TARGET_FORMAT_TYPES 658	TARGET_INVALID_BINARY_OP	
TARGET_FRAME_POINTER_REQUIRED533	TARGET_INVALID_CONVERSION	
TARGET_FUNCTION_ARG 536	TARGET_INVALID_UNARY_OP	
TARGET_FUNCTION_ARG_ADVANCE539	TARGET_INVALID_WITHIN_DOLOOP	
TARGET_FUNCTION_ARG_BOUNDARY540	TARGET_IRA_CHANGE_PSEUDO_ALLOCNO_CLASS	521
TARGET_FUNCTION_ARG_OFFSET 540	TARGET_KEEP_LEAF_WHEN_PROFILED	552
TARGET_FUNCTION_ARG_PADDING540	TARGET_LEGITIMATE_ADDRESS_P	562
TARGET_FUNCTION_ARG_ROUND_BOUNDARY 540	TARGET_LEGITIMATE_COMBINED_INSN	655
TARGET_FUNCTION_ATTRIBUTE_INLINABLE_P 635	TARGET_LEGITIMATE_CONSTANT_P	565
TARGET_FUNCTION_INCOMING_ARG537	TARGET_LEGITIMIZE_ADDRESS	564
TARGET_FUNCTION_OK_FOR_SIBCALL 552	TARGET_LEGITIMIZE_ADDRESS_DISPLACEMENT	522
TARGET_FUNCTION_VALUE544	TARGET_LIB_INT_CMP_BIASED	561
TARGET_FUNCTION_VALUE_REGNO_P546	TARGET_LIBC_HAS_FAST_FUNCTION	561
TARGET_GEN_CCMP_FIRST656	TARGET_LIBC_HAS_FUNCTION	561
TARGET_GEN_CCMP_NEXT656	TARGET_LIBCALL_VALUE	545
TARGET_GENERATE_VERSION_DISPATCHER_BODY	TARGET_LIBFUNC_GNU_PREFIX	
653	TARGET_LIBGCC_CMP_RETURN_MODE	498
TARGET_GET_DRAP_RTX	TARGET_LIBGCC_FLOATING_MODE_SUPPORTED_P	
TARGET_GET_FUNCTION_VERSIONS_DISPATCHER		543
653	TARGET_LIBGCC_SDATA_SECTION	592
TARGET_GET_MULTILIB_ABI_NAME	TARGET_LIBGCC_SHIFT_COUNT_MODE	
TARGET_GET_PCH_VALIDITY	TARGET_LOAD_BOUNDS_FOR_ARG	
TARGET_GET_RAW_ARG_MODE 547	TARGET_LOAD_RETURNED_BOUNDS	557
TARGET_GET_RAW_RESULT_MODE 547	TARGET_LOOP_UNROLL_ADJUST	
TARGET_GET_VALID_OPTION_VALUES 554	TARGET_LRA_P	
TARGET_GIMPLE_FOLD_BUILTIN 653	TARGET_MACHINE_DEPENDENT_REORG	
TARGET_GIMPLIFY_VA_ARG_EXPR541	TARGET_MANGLE_ASSEMBLER_NAME	
TARGET_GOACC_DIM_LIMIT	TARGET_MANGLE_DECL_ASSEMBLER_NAME	
TARGET_GOACC_FORK_JOIN	TARGET_MANGLE_TYPE	
TARGET_GOACC_REDUCTION	TARGET_MAX_ANCHOR_OFFSET	
TARGET_GOACC_VALIDATE_DIMS	TARGET_MAX_NOCE_IFCVT_SEQ_COST	
TARGET_HANDLE_C_OPTION	TARGET_MD_ASM_ADJUST	
TARGET_HANDLE_GENERIC_ATTRIBUTE 635	TARGET_MEM_CONSTRAINT	
TARGET_HANDLE_OPTION	TARGET_MEM_REF	
TARGET_HARD_REGNO_CALL_PART_CLOBBERED 506	TARGET_MEMBER_TYPE_FORCES_BLK	
TARGET_HARD_REGNO_MODE_OK 509	TARGET_MEMMODEL_CHECK	
TARGET_HARD_REGNO_NREGS	TARGET_MEMORY_MOVE_COST	
TARGET_HARD_REGNO_SCRATCH_OK	TARGET_MERGE_DECL_ATTRIBUTES	
TARGET_HAS_IFUNC_P	TARGET_MERGE_TYPE_ATTRIBUTES	
TARGET_HAS_NO_HW_DIVIDE	TARGET_MIN_ANCHOR_OFFSET	
TARGET_HAVE_CONDITIONAL_EXECUTION	TARGET_MIN_ARITHMETIC_PRECISION	
TARGET_HAVE_COUNT_REG_DECR_P	TARGET_MIN_DIVISIONS_FOR_RECIP_MUL	
TARGET_HAVE_CTORS_DTORS	TARGET_MODE_AFTER	
TARGET_HAVE_NAMED_SECTIONS	TARGET_MODE_DEPENDENT_ADDRESS_P	
TARGET_HAVE_SPECULATION_SAFE_VALUE 661	TARGET_MODE_EMIT	
TARGET_HAVE_SRODATA_SECTION	TARGET_MODE_ENTRY	
TARGET_HAVE_SWITCHABLE_BSS_SECTIONS 598	TARGET_MODE_EXIT	
TARGET_HAVE_TLS	TARGET_MODE_NEEDED	
TARGET_IN_SMALL_DATA_P	TARGET_MODE_PRIORITY	
TARGET INIT BUILTINS	TARGET_MODE_FRIORITI	
INIVIDE INTERPOLITION		0.40

TARGET_MODES_TIEABLE_P510	TARGET_RESET_LOCATION_VIEW	630
TARGET_MS_BITFIELD_LAYOUT_P	TARGET_RESOLVE_OVERLOADED_BUILTIN	653
TARGET_MUST_PASS_IN_STACK 537	TARGET_RETURN_IN_MEMORY	546
TARGET_MUST_PASS_IN_STACK, and	TARGET_RETURN_IN_MSB	
TARGET_FUNCTION_ARG	TARGET_RETURN_POPS_ARGS	
TARGET_N_FORMAT_TYPES658	TARGET_RTX_COSTS	
TARGET_NARROW_VOLATILE_BITFIELD 497	TARGET_RUN_TARGET_SELFTESTS	
TARGET_NO_REGISTER_ALLOCATION	TARGET_SCALAR_MODE_SUPPORTED_P	
TARGET_NO_SPECULATION_IN_DELAY_SLOTS_P 582	TARGET_SCHED_ADJUST_COST	
TARGET_NOCE_CONVERSION_PROFITABLE_P 582	TARGET_SCHED_ADJUST_PRIORITY	
TARGET_OBJC_CONSTRUCT_STRING_OBJECT 488	TARGET_SCHED_ALLOC_SCHED_CONTEXT	
TARGET_OBJC_DECLARE_CLASS_DEFINITION 488	TARGET_SCHED_CAN_SPECULATE_INSN	
TARGET_OBJC_DECLARE_UNRESOLVED_CLASS_	TARGET_SCHED_CLEAR_SCHED_CONTEXT	
REFERENCE	TARGET_SCHED_DEPENDENCIES_EVALUATION_HOOK	
TARGET_OBJECT_SUFFIX	TARGET_SORES_SER ENDEROTES_EVALUATION_HOUS	
TARGET_OBJFMT_CPP_BUILTINS	TARGET_SCHED_DFA_NEW_CYCLE	
TARGET_OFFLOAD_OPTIONS	TARGET_SCHED_DFA_NEW_GTGLE  TARGET_SCHED_DFA_POST_ADVANCE_CYCLE	
TARGET_OMIT_STRUCT_RETURN_REG	TARGET_SCHED_DFA_FOST_ADVANCE_CTCLE  TARGET_SCHED_DFA_POST_CYCLE_INSN	
TARGET_OMP_DEVICE_KIND_ARCH_ISA		
	TARGET_SCHED_DFA_PRE_ADVANCE_CYCLE	
TARGET_OPTAB_SUPPORTED_P	TARGET_SCHED_DFA_PRE_CYCLE_INSN	
TARGET_OPTF	TARGET_SCHED_DISPATCH	
TARGET_OPTION_FUNCTION_VERSIONS	TARGET_SCHED_DISPATCH_DO	
TARGET_OPTION_INIT_STRUCT	TARGET_SCHED_EXPOSED_PIPELINE	
TARGET_OPTION_OPTIMIZATION_TABLE	TARGET_SCHED_FINISH	
TARGET_OPTION_OVERRIDE636	TARGET_SCHED_FINISH_GLOBAL	585
TARGET_OPTION_POST_STREAM_IN	TARGET_SCHED_FIRST_CYCLE_MULTIPASS_	
TARGET_OPTION_PRAGMA_PARSE	BACKTRACK	586
TARGET_OPTION_PRINT 635	TARGET_SCHED_FIRST_CYCLE_MULTIPASS_BEGIN	
TARGET_OPTION_RESTORE635		586
TARGET_OPTION_SAVE 635	TARGET_SCHED_FIRST_CYCLE_MULTIPASS_DFA_	
TARGET_OPTION_VALID_ATTRIBUTE_P 635	LOOKAHEAD	585
TARGET_OS_CPP_BUILTINS	TARGET_SCHED_FIRST_CYCLE_MULTIPASS_DFA_	
TARGET_OVERRIDE_OPTIONS_AFTER_CHANGE 488	LOOKAHEAD_GUARD	
TARGET_OVERRIDES_FORMAT_ATTRIBUTES 658	TARGET_SCHED_FIRST_CYCLE_MULTIPASS_END	586
TARGET_OVERRIDES_FORMAT_ATTRIBUTES_COUNT	TARGET_SCHED_FIRST_CYCLE_MULTIPASS_FINI	
658		587
TARGET_OVERRIDES_FORMAT_INIT658	TARGET_SCHED_FIRST_CYCLE_MULTIPASS_INIT	
TARGET_PASS_BY_REFERENCE 538		586
TARGET_PCH_VALID_P	TARGET_SCHED_FIRST_CYCLE_MULTIPASS_ISSUE	
TARGET_POSIX_IO		586
TARGET_PREDICT_DOLOOP_P 654	TARGET_SCHED_FREE_SCHED_CONTEXT	588
TARGET_PREFERRED_ELSE_VALUE	TARGET_SCHED_FUSION_PRIORITY	589
TARGET_PREFERRED_OUTPUT_RELOAD_CLASS 516	TARGET_SCHED_GEN_SPEC_CHECK	588
TARGET_PREFERRED_RELOAD_CLASS	TARGET_SCHED_H_I_D_EXTENDED	587
TARGET_PREFERRED_RENAME_CLASS	TARGET_SCHED_INIT	
TARGET_PREPARE_PCH_SAVE	TARGET_SCHED_INIT_DFA_POST_CYCLE_INSN	
TARGET_PRETEND_OUTGOING_VARARGS_NAMED 557	TARGET_SCHED_INIT_DFA_PRE_CYCLE_INSN	
TARGET_PROFILE_BEFORE_PROLOGUE	TARGET_SCHED_INIT_GLOBAL	
TARGET_PROMOTE_FUNCTION_MODE	TARGET_SCHED_INIT_SCHED_CONTEXT	
TARGET_PROMOTE_PROTOTYPES	TARGET_SCHED_IS_COSTLY_DEPENDENCE	
TARGET_PROMOTED_TYPE	TARGET_SCHED_ISSUE_RATE	
TARGET_PTRMEMFUNC_VBIT_LOCATION 504	TARGET_SCHED_MACRO_FUSION_P	
TARGET_RECORD_OFFLOAD_SYMBOL	TARGET_SCHED_MACRO_FUSION_PAIR_P	
TARGET_REF_MAY_ALIAS_ERRNO	TARGET_SCHED_NEEDS_BLOCK_P	
TARGET_REGISTER_MOVE_COST	TARGET_SCHED_REASSOCIATION_WIDTH	
TARGET_REGISTER_PRIORITY	TARGET_SCHED_REORDER	
TARGET_REGISTER_PRIORITY	TARGET_SCHED_REORDER2	
TARGET RELAYOUT FUNCTION	TARGET_SCHED_REURDERZTARGET_SCHED_SET_SCHED_CONTEXT	
IARGEL RELATION FUNCTION	IMAGEI_OUREN_OEI_OUREN_CUNIEXI	001

TARGET_SCHED_SET_SCHED_FLAGS	TARGET_TRAMPOLINE_ADJUST_ADDRESS	559
TARGET_SCHED_SMS_RES_MII 588	TARGET_TRAMPOLINE_INIT	559
TARGET_SCHED_SPECULATE_INSN	TARGET_TRANSLATE_MODE_ATTRIBUTE	542
TARGET_SCHED_VARIABLE_ISSUE	TARGET_TRULY_NOOP_TRUNCATION	645
TARGET_SECONDARY_MEMORY_NEEDED 519	TARGET_UNSPEC_MAY_TRAP_P	655
TARGET_SECONDARY_MEMORY_NEEDED_MODE 519	TARGET_UNWIND_TABLES_DEFAULT	
TARGET_SECONDARY_RELOAD	TARGET_UNWIND_WORD_MODE	
TARGET_SECTION_TYPE_FLAGS	TARGET_UPDATE_STACK_BOUNDARY	
TARGET_SELECT_EARLY_REMAT_MODES 519	TARGET_USE_ANCHORS_FOR_SYMBOL_P	
TARGET_SET_CURRENT_FUNCTION	TARGET_USE_BLOCKS_FOR_CONSTANT_P	
TARGET_SET_DEFAULT_TYPE_ATTRIBUTES 634	TARGET_USE_BLOCKS_FOR_DECL_P	
TARGET_SET_UP_BY_PROLOGUE	TARGET_USE_BY_PIECES_INFRASTRUCTURE_P	
TARGET_SETJMP_PRESERVES_NONVOLATILE_REGS_P	TARGET_USE_PSEUDO_PIC_REG	
	TARGET_USES_WEAK_UNWIND_INFO	
TARGET_SETUP_INCOMING_VARARGS	TARGET_VALID_DLLIMPORT_ATTRIBUTE_P	
TARGET_SHIFT_TRUNCATION_MASK	TARGET_VALID_POINTER_MODE	
TARGET_SHRINK_WRAP_COMPONENTS_FOR_BB 553	TARGET_VECTOR_ALIGNMENT	
TARGET_SHRINK_WRAP_DISQUALIFY_COMPONENTS	TARGET_VECTOR_MODE_SUPPORTED_P	
553	TARGET_VECTORIZE_ADD_STMT_COST	
TARGET_SHRINK_WRAP_EMIT_EPILOGUE_COMPONENTS	TARGET_VECTORIZE_AUTOVECTORIZE_VECTOR_MOD	
553		
TARGET_SHRINK_WRAP_EMIT_PROLOGUE_COMPONENTS	TARGET_VECTORIZE_BUILTIN_GATHER	
553	TARGET_VECTORIZE_BUILTIN_MASK_FOR_LOAD	
TARGET_SHRINK_WRAP_GET_SEPARATE_COMPONENTS	TARGET_VECTORIZE_BUILTIN_MD_VECTORIZED_	500
553	FUNCTION	567
TARGET_SHRINK_WRAP_SET_HANDLED_COMPONENTS	TARGET_VECTORIZE_BUILTIN_SCATTER	
	TARGET_VECTORIZE_BUILTIN_SCATTER	
TARGET_SIMD_CLONE_ADJUST	TARGET_VECTORIZE_BOTETIN_VECTORIZATION_CC	
TARGET_SIMD_CLONE_COMPUTE_VECSIZE_AND_	TARGET_VECTORIZE_BUILTIN_VECTORIZED_	500
SIMDLEN	FUNCTION	567
TARGET_SIMD_CLONE_USABLE         570           TARGET_SIMT_VF         570	TARGET_VECTORIZE_DESTROY_COST_DATA	509
	TARGET_VECTORIZE_EMPTY_MASK_IS_EXPENSIVE	E60
TARGET_SLOW_UNALIGNED_ACCESS	TARGET REGERRATE EINIGH COCT	
TARGET_SMALL_REGISTER_CLASSES_FOR_MODE_P	TARGET_VECTORIZE_FINISH_COST	
TARGET_SPECULATION_SAFE_VALUE	TARGET_VECTORIZE_GET_MASK_MODE	
TARGET_SPILL_CLASS	TARGET_VECTORIZE_INIT_COST	
	TARGET_VECTORIZE_PREFERRED_SIMD_MODE	
TARGET_SPLIT_COMPLEX_ARG	TARGET_VECTORIZE_PREFERRED_VECTOR_ALIGNME	
TARGET_STACK_CLASH_PROTECTION_ALLOCA_PROBE_	TARGET VEGTORIZE DELATER MORE	
RANGE	TARGET_VECTORIZE_RELATED_MODE	
TARGET_STACK_PROTECT_FAIL	TARGET_VECTORIZE_SPLIT_REDUCTION	508
TARGET_STACK_PROTECT_GUARD	TARGET_VECTORIZE_SUPPORT_VECTOR_	-0-
TARGET_STACK_PROTECT_RUNTIME_ENABLED_P 554	MISALIGNMENT	
TARGET_STARTING_FRAME_OFFSET	TARGET_VECTORIZE_VEC_PERM_CONST	
TARGET_STARTING_FRAME_OFFSET and virtual	TARGET_VECTORIZE_VECTOR_ALIGNMENT_REACHAE	
registers		
TARGET_STATIC_CHAIN	TARGET_VERIFY_TYPE_CONTEXT	
TARGET_STATIC_RTX_ALIGNMENT	TARGET_VTABLE_DATA_ENTRY_DISTANCE	
TARGET_STORE_BOUNDS_FOR_ARG	TARGET_VTABLE_ENTRY_ALIGN	
TARGET_STORE_RETURNED_BOUNDS	TARGET_VTABLE_USES_DESCRIPTORS	
TARGET_STRICT_ARGUMENT_NAMING	TARGET_WANT_DEBUG_PUB_SECTIONS	
TARGET_STRING_OBJECT_REF_TYPE_P 488	TARGET_WARN_FUNC_RETURN	
TARGET_STRIP_NAME_ENCODING	TARGET_WARN_PARAMETER_PASSING_ABI	
TARGET_STRUCT_VALUE_RTX	TARGET_WEAK_NOT_IN_ARCHIVE_TOC	
TARGET_SUPPORTS_SPLIT_STACK	targetm	
TARGET_SUPPORTS_WEAK	targets, makefile	
TARGET_SUPPORTS_WIDE_INT 661	TCmode	
TARGET TERMINATE DW2 EH FRAME INFO 623	TDmode	272

TEMPLATE_DECL	168	TRUTH_AND_EXPR	177
Temporaries	216	TRUTH_ANDIF_EXPR	177
termination routines		TRUTH_NOT_EXPR	
testing constraints	391	TRUTH_OR_EXPR	
TEXT_SECTION_ASM_OP		TRUTH_ORIF_EXPR	
TFmode		TRUTH_XOR_EXPR	
The Language		TRY_BLOCK.	
THEN_CLAUSE	203	TRY_HANDLERS	
THREAD_MODEL_SPEC		TRY_STMTS	
THROW_EXPR		Tuple specific accessors	
THUNK_DECL		tuples	
THUNK_DELTA		type	
TImode	272	type declaration	
TImode, in insn	309	TYPE_ALIGN	
TLS_COMMON_ASM_OP		TYPE_ARG_TYPES	
TLS_SECTION_ASM_FLAG		TYPE_ASM_OP	
'tm.h' macros		TYPE_ATTRIBUTES	
TQFmode		TYPE_BINFO	
TQmode		TYPE_BUILT_IN	
TRAMPOLINE_ALIGNMENT		TYPE_BUILI_IN	
TRAMPOLINE_SECTION	559		
TRAMPOLINE_SIZE		TYPE_CONTEXT	
trampolines for nested functions		TYPE_DECL	
TRANSFER_FROM_TRAMPOLINE		TYPE_FIELDS	
trap instruction pattern		TYPE_HAS_ARRAY_NEW_OPERATOR	
tree		TYPE_HAS_DEFAULT_CONSTRUCTOR	
Tree SSA		TYPE_HAS_MUTABLE_P	
tree_fits_shwi_p		TYPE_HAS_NEW_OPERATOR	
tree_fits_uhwi_p		TYPE_MAIN_VARIANT	
tree_int_cst_equal		TYPE_MAX_VALUE	
tree_int_cst_lt		$\texttt{TYPE\_METHOD\_BASETYPE} \dots 163,$	
tree_size		TYPE_MIN_VALUE	
tree_to_shwi		TYPE_NAME 163, 164, 196,	
tree_to_uhwi		TYPE_NOTHROW_P	
TREE_CHAIN		$\texttt{TYPE\_OFFSET\_BASETYPE} \dots 163,$	
TREE_CODE		TYPE_OPERAND_FMT	
TREE_INT_CST_ELT		TYPE_OVERLOADS_ARRAY_REF	
TREE_INT_CST_LOW		TYPE_OVERLOADS_ARROW	
TREE_INT_CST_NUNITS		TYPE_OVERLOADS_CALL_EXPR	
TREE_LIST		TYPE_POLYMORPHIC_P	
TREE_OPERAND		TYPE_PRECISION	
TREE_PUBLIC		TYPE_PTR_P	
TREE_PURPOSE		TYPE_PTRDATAMEM_P	
TREE_READONLY		TYPE_PTRFN_P	197
TREE_STATIC		TYPE_PTROB_P	197
TREE_STRING_LENGTH		TYPE_PTROBV_P	196
TREE_STRING_POINTER		TYPE_QUAL_CONST	196
TREE_THIS_VOLATILE		TYPE_QUAL_RESTRICT 163,	196
TREE_TYPE 162, 163, 168, 173, 194,		TYPE_QUAL_VOLATILE 163,	196
TREE_VALUE		TYPE_RAISES_EXCEPTIONS	203
TREE_VEC		TYPE_SIZE 163, 164, 196,	197
TREE_VEC_ELT		TYPE_STRUCTURAL_EQUALITY_P 163,	165
TREE_VEC_LENGTH		TYPE_UNQUALIFIED	
true positive		TYPE_VFIELD	199
TRUNC_DIV_EXPR		TYPENAME_TYPE	196
TRUNC_MOD_EXPR		TYPENAME_TYPE_FULLNAME	
truncate		TYPEOF_TYPE	196
truncmn2 instruction pattern			

$\mathbf{U}$	UNORDERED_EXPR	177
uaddvm4 instruction pattern	unshare_all_rtl	315
uavgm3_ceil instruction pattern	unsigned division	290
uavgm3_floor instruction pattern	unsigned division with unsigned saturation	290
UDAmode	unsigned greater than	293
udiv	unsigned less than	293
udivm3 instruction pattern	unsigned minimum and maximum	290
udivmodm4 instruction pattern	unsigned_fix	296
udot_prodm instruction pattern 401	unsigned_float	296
UDQmode	unsigned_fract_convert	296
UHAmode	unsigned_sat_fract	296
UHQmode	unspec 301,	
UINT_FAST16_TYPE	unspec_volatile	
UINT_FAST32_TYPE	untyped_call instruction pattern	419
UINT_FAST64_TYPE	untyped_return instruction pattern	
UINT_FAST8_TYPE	update_ssa	
UINT_LEAST16_TYPE	update_stmt	
UINT_LEAST32_TYPE	update_stmt_if_modified	222
UINT_LEAST64_TYPE	UPDATE_PATH_HOST_CANONICALIZE (path)	
UINT_LEAST8_TYPE	UQQmode	
UINT16_TYPE	us_ashift	
UINT32_TYPE	us_minus	
UINT64_TYPE	us_mult	
UINT8_TYPE	us_neg	
UINTMAX_TYPE	us_plus	
UINTPTR_TYPE	us_truncate	
umaddmn4 instruction pattern 405	usaddm3 instruction pattern	
umax	usadm instruction pattern	
umaxm3 instruction pattern	USAmode	
umin	usashlm3 instruction pattern	
uminm3 instruction pattern	usdivm3 instruction pattern	
umod	use	
umodm3 instruction pattern	USE_C_ALLOCA	665
umsubmn4 instruction pattern 405	USE_LD_AS_NEEDED	482
umulhisi3 instruction pattern	USE_LOAD_POST_DECREMENT	
umulhrsm3 instruction pattern 402	USE_LOAD_POST_INCREMENT	
umulhsm3 instruction pattern 401	USE_LOAD_PRE_DECREMENT	
umulm3_highpart instruction pattern 404	USE_LOAD_PRE_INCREMENT	
umulqihi3 instruction pattern 404	USE_SELECT_SECTION_FOR_FUNCTIONS	
umulsidi3 instruction pattern 404	USE_STORE_POST_DECREMENT	
umulvm4 instruction pattern 399	USE_STORE_POST_INCREMENT	580
unchanging	USE_STORE_PRE_DECREMENT	
unchanging, in call_insn	USE_STORE_PRE_INCREMENT	
unchanging, in jump_insn, call_insn and insn	used	
	used, in symbol_ref	269
unchanging, in mem	user	
unchanging, in subreg	user experience guidelines	715
unchanging, in symbol_ref	user gc	
UNEQ_EXPR	USER_LABEL_PREFIX	
UNGE_EXPR	USING_STMT	
UNGT_EXPR	usmaddmn4 instruction pattern	
UNION_TYPE	usmsubmn4 instruction pattern	
unions, returning	usmulhisi3 instruction pattern	
UNITS_PER_WORD	usmulm3 instruction pattern	
UNKNOWN_TYPE	usmulqihi3 instruction pattern	
UNLE_EXPR	usmulsidi3 instruction pattern	
UNLIKELY_EXECUTED_TEXT_SECTION_NAME 591	usnegm2 instruction pattern	
UNLT EXPR	USQmode	

ussubm3 instruction pattern 399	
usubvm4 instruction pattern 399	
UTAmode	<pre>3 vec_unpack_ufix_trunc_lo_m instruction pattern</pre>
UTQmode	3 vec_unpacks_float_hi_m instruction pattern
$\mathbf{V}$	vec_unpacks_float_lo_m instruction pattern
'V' in constraint	1 403
VA_ARG_EXPR	vec_unpacks_hi_m instruction pattern 403
values, returned by functions	vec_unpacks_lo_m instruction pattern 403
var_location	vec_unpacks_sbool_hi_m instruction pattern
VAR_DECL	8
varargs implementation	vec_unpacks_sbool_lo_m instruction pattern
variable	0
Variable Location Debug Information in RTL	vec_unpacku_iloat_ni_m instruction pattern
304	403
vashlm3 instruction pattern 400	vec_unpacku_float_lo_m instruction pattern
vashrm3 instruction pattern 400	e
vcond_mask_mn instruction pattern	vec_unpacku_nr_m instruction pattern 403
vcondeqmn instruction pattern 398	vec_unpacku_10_m instruction pattern 403
vcondmn instruction pattern 398	vec_widen_smuit_even_m instruction pattern
vcondumn instruction pattern 398	
vec_cmpeqmn instruction pattern	vec_widen_smuit_ni_m instruction pattern 404
vec_cmpmn instruction pattern	vec_widen_smallc_io_m instruction pattern 404
vec_cmpumn instruction pattern	vec_widen_smalle_odd_m institution pattern 404
vec_concat	vec_widen_ssniitt_ni_m instruction pattern
vec_duplicate	101
vec_duplicatem instruction pattern 39'	_ vec_widen_ssmilti_io_m institution pattern
vec_extractmn instruction pattern	
	voo_widon_dmdio_ovon_m institution pattern
vec_initmn instruction pattern	
vec_mask_load_lanesmn instruction pattern 39	
vec_mask_store_lanesmn instruction pattern598	voo_wadau_amaato_ao_m moordeenen peeteeni
vec_mask_store_ranesmn instruction pattern	vec_widen_umult_odd_m instruction pattern 404
	vec_widen_dshiiici_mi_minstruction pattern
vec_merge	
vec_pack_sbool_trunc_m instruction pattern	vec_widen_ushiftl_lo_m instruction pattern
vec_pack_sfix_trunc_m instruction pattern 40:	
vec_pack_ssat_m instruction pattern 400	<del>-</del> -
vec_pack_ssat_m instruction pattern 400 vec_pack_trunc_m instruction pattern 400	
vec_pack_ufix_trunc_m instruction pattern 40.	
vec_pack_urix_trunc_m instruction pattern 400 vec_pack_usat_m instruction pattern 400	
vec_packs_float_m instruction pattern 400	_
vec_packu_float_m instruction pattern 4000 vec_packu_float_m instruction pattern	
vec_packu_iioat_m instruction pattern	
vec_select	
_	
vec_series	
vec_seriesm instruction pattern	
vec_setm instruction pattern	
vec_shl_insert_m instruction pattern 400	
vec_shl_m instruction pattern	
vec_shr_m instruction pattern	
vec_store_lanesmn instruction pattern 399	VEC_WIDEN_MULT_LO_EXPR
<pre>vec_unpack_sfix_trunc_hi_m instruction pattern</pre>	
400	
vec_unpack_sfix_trunc_lo_m instruction pattern 40:	vector operations         294           3 VECTOR CST         173
40,	v=0.100_πU1160_πU10Ξv = 1/3

VECTOR_STORE_FLAG_VALUE	WHILE_STMT 20	05
verify_flow_info	whopr	93
virtual operands	widen_ssumm3 instruction pattern 40	01
VIRTUAL_INCOMING_ARGS_REGNUM	widen_usumm3 instruction pattern 40	01
VIRTUAL_OUTGOING_ARGS_REGNUM	WIDEST_HARDWARE_FP_SIZE 50	01
VIRTUAL_STACK_DYNAMIC_REGNUM	window_save instruction pattern 4	
VIRTUAL_STACK_VARS_REGNUM	WINT_TYPE50	
VLIW 460, 463	word_mode2	
vlshrm3 instruction pattern 406	WORD_REGISTER_OPERATIONS	
VMS 664	WORDS_BIG_ENDIAN	
VMS_DEBUGGING_INFO	WORDS_BIG_ENDIAN, effect on subreg	
VOID_TYPE	wpa	
VOIDmode	wpa	Je
volatil		
volatil, in insn, call_insn, jump_insn,	X	
code_label, jump_table_data, barrier, and		
note	'x-host'	
volatil, in label_ref and reg_label 266	'X' in constraint	-
volatil, in mem, asm_operands, and asm_input	XCmode	
	XCOFF_DEBUGGING_INFO65	
volatil, in reg	XEXP 20	
volatil, in subreg	XFmode	$7^{2}$
volatil, in symbol_ref 269	XImode	$7^{2}$
volatile memory references	XINT 20	62
volatile, in prefetch	'xm-machine.h'	65
voting between constraint alternatives	xor 29	91
vrotlm3 instruction pattern 406	xor, canonicalization of 43	37
vrotrm3 instruction pattern 406	xorm3 instruction pattern3	96
1	xorsignm3 instruction pattern 4	10
	XSTR 20	62
$\mathbf{W}$	XVEC 20	
walk_dominator_tree	XVECEXP	63
walk_gimple_op	XVECLEN	63
walk_gimple_seq	XWINT	62
walk_gimple_stmt		
WCHAR_TYPE	_	
WCHAR_TYPE_SIZE	${f Z}$	
which_alternative	zero_extend	g.
while_ultmn instruction pattern	zero_extendmn2 instruction pattern 4	
WHILE_BODY	zero_extract	
WHILE COND	zero extract. canonicalization of	
WILLE COND 205	zero extract. Canonicanzation of 4.	.) [